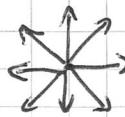


# EXAM JUNE 5TH 2018

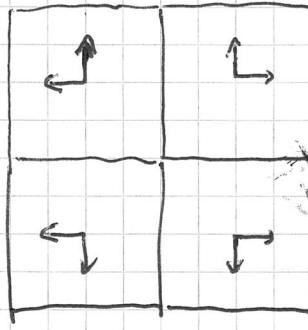
1)

a)

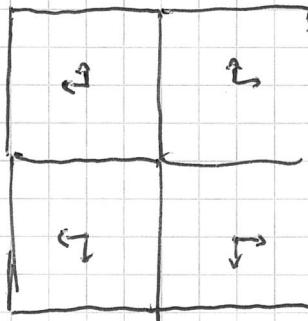
Vector bouquet: (8 bins):



Vector bouquets for P (same order of cells):



Vector bouquets for Q:



There are fewer pixels with gradients, so the entries in the gradient histogram are smaller for Q than for P

There are 4 cells, each with 8 histogram bins

$$\rightarrow 4 \cdot 8 = 32 \text{ dimensions}$$

b)

distances to		ratio
1st neighbor	2nd neighbor	
$d_1$	0.1 ( $q_1$ )	<del>0.2</del> $0.5 = \frac{1}{2}$
$d_2$	0.1 ( $q_2$ )	$0.25 = \frac{1}{4}$
$d_3$	2 ( $q_3$ )	$0.25 = \frac{1}{4}$
$d_4$	2 ( $q_3$ )	$1/6$

According to the Lowe ratio, the match  $(d_4, q_3)$  is most certain.

c) The density of the SIFT descriptor space is non-uniform. There are regions, e.g., corresponding to repeating structures, that are denser. Thus, absolute distances are unreliable.  $\rightarrow$  ~~standard~~ Rather, we need to normalize the distances.

2 |

a) Loss for training example  $I_i$  with true label 5:

$$L(I_i) = -\ln(p_5(I_i))$$

b) For simplicity:  $p_5(I_i) = p_5$ ,  $p_j(I_i) = p_j$

Softmax function:

$$p_5 = \frac{e^{y_5}}{\sum_i e^{y_i}}$$

$y_i$  are outputs of fully connected layer

We want to compute the gradient

$$\frac{\partial L(I_i)}{\partial y_5}$$

Using the chain rule, we get

$$\frac{\partial L(I_i)}{\partial y_5} = \frac{\partial L(I_i)}{\partial p_5} \frac{\partial p_5}{\partial y_5}$$

$$\text{Where } \frac{\partial L(I_i)}{\partial p_5} = -\frac{1}{p_5}$$

We use the quotient rule  $\frac{f'(x)g(x) - f(x)g'(x)}{(g(x))^2}$

to compute  $\frac{\partial p_5}{\partial y_5}$ :

$$\frac{\partial p_5}{\partial y_5} \quad \text{Here: } \cancel{f(x)g(x)}$$

$$\text{Here: } f(y_5) = e^{y_5}, f'(y_5) = e^{y_5}$$

$$g(y_5) = \sum_i e^{y_i}, g'(y_5) = e^{y_5}$$

this gives us:

$$\frac{\partial p_5}{\partial y_5} = \frac{e^{y_5}g(y_5) - e^{y_5}g'(y_5)}{(g(y_5))^2}$$

$$= \frac{e^{y_5}}{g(y_5)} \cdot \frac{g(y_5) - e^{y_5}}{g(y_5)}$$

## 2b Continued

$$= \frac{e^{y_5}}{\sum_i e^{y_i}} \cdot \frac{\sum_i e^{y_i} - e^{y_5}}{\sum_i e^{y_i}} = p_5 (1 - p_5)$$

$$\Rightarrow \frac{\partial L(\pi_i)}{\partial y_5} = -\frac{1}{p_5} p_5 (1 - p_5) = p_5 - 1$$

c) Input dimensions:

$$\underbrace{46 \times 46 \times 3}_{\substack{\text{original} \\ \text{input image}}} \rightarrow 44 \times 44 \times 10 \xrightarrow{\substack{\uparrow \\ \text{no} \\ \text{due to padding}}} 44 \times 44 \times 10 \xrightarrow{\substack{\text{after} \\ \text{ReLU}}}$$

$$\rightarrow 22 \times 22 \times 10 \xrightarrow{\substack{\text{after max.} \\ \text{pooling}}} \underbrace{20 \times 20 \times 20}_{\substack{\downarrow \\ \text{due to no} \\ \text{padding}}}$$

$$\rightarrow \underbrace{20 \times 20 \times 20}_{\substack{\text{after ReLU}}} \xrightarrow{\substack{\uparrow \\ \text{after max.} \\ \text{pooling}}} 10 \times 10 \times 20$$

$$\rightarrow 10 \rightarrow 10$$

Number of parameters:

for the bias term

$$\times 10 \text{ } 3 \times 3 \text{ convolutions: } 10 \cdot (3 \cdot 3 \cdot 3 + 1) = 280$$

$$\times 20 \text{ } 3 \times 3 \text{ convolutions: } 20 \cdot (10 \cdot 3 \cdot 3 + 1) = 1820$$

$\times$  Fully connected layer:

$$\text{input size: } 10 \times 10 \times 20 = 2000$$

$$\text{output size: } 10$$

$$\Rightarrow 2000 \times 10 = 20000 \text{ parameters}$$

d) Replace the FC layer by  $\times 10 \text{ } 10 \times 10$  convolutions.

This will result in exactly the same output  
for an input image of size  $46 \times 46$ .

3)

a) Re-write the system as a function of the parameters of the affine transform:

$$\begin{pmatrix} \tilde{x} & \tilde{y} & \tilde{z} & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & \tilde{x} & \tilde{y} & \tilde{z} & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & \tilde{x} & \tilde{y} & \tilde{z} & 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} a \\ b \\ c \\ d \\ e \\ f \\ g \\ h \\ i \\ j \\ tx \\ ty \\ tz \end{pmatrix} = \begin{pmatrix} x \\ y \\ z \end{pmatrix}$$

One match gives 3 equations in 12

unknowns  $\rightarrow$  we need 4 matches.

Stack the systems into a system  $M\theta = b$ ,

where  $M \in \mathbb{R}^{12 \times 12}$ ,  $b \in \mathbb{R}^{12}$  and solve

in Matlab as  $\theta = M^{-1}b$ .

b) The minimal solver will create a solution  
such that the  $K$  points used to  
compute it are inliers. Thus, we will  
always have <sup>at least</sup>  $K$  inliers. A solution with  $K$   
inliers is thus not ~~not~~ trustworthy

c) In 2D, only 3 matches are needed.

Probability of selecting 3 inliers:  $(0.5)^3$

$\hookrightarrow$  On average  $\frac{1}{(0.5)^3}$  samples needed

$\hookrightarrow$  In 3D:  $\frac{1}{(0.5)^4}$  samples needed

$\hookrightarrow$  Finding a good transformation takes

$$\frac{(0.5)^3}{(0.5)^4} = 2 \text{ times longer.}$$

4|

a) ransac-triangulation ( $\{P_i\}, \{u_i\}, t$ )

$$I_{\max} = \lceil \log(0.01) / \log(1 - (\frac{t}{n})^2) \rceil$$

$I_{\max} = 0$   
for  $I = 1 : I_{\max}$

$P_s, u_s$  = random sample of size 2 from  
 $\{P_i\}, \{u_i\}$

$h_{\text{hat}} = \text{minimal\_solver}(P_s, u_s)$

reprojection-errors = reprojection-errors ( $\{P\}, \{u\}, h_{\text{hat}}$ )

if ~~sum(reprojection-errors)~~  
 $I = \text{sum}(\text{reprojection-errors}) \leq t$

if  $I > I_{\max}$

$$I_{\max} = I$$

$h_{\text{best}} = h_{\text{hat}}$

$$\varepsilon = \frac{I}{n}$$

$$I_{\max} = \lceil \log(0.01) / \log(1 - \varepsilon^2) \rceil$$

end  
 end  
 return  $h_{\text{best}}, I_{\max}$ ;

b) The projected point in homogeneous coordinates is given as

$$\begin{pmatrix} x' \\ y' \\ z' \end{pmatrix} = P \begin{pmatrix} u \\ v \\ 1 \end{pmatrix}$$

This yields the following image coordinates:

$$\begin{pmatrix} u' \\ v' \end{pmatrix} = \frac{1}{z'} \begin{pmatrix} x' \\ y' \end{pmatrix}$$

This yields the following residuals:

$$r_x = x - u' = x - \frac{1}{z'} x'$$

$$r_y = y - v' = y - \frac{1}{z'} y'$$

The equations are valid for  $z' > 0$ , i.e.,  
 if  $u$  is in front of the camera

51

a) Using the chain rule, we can

express  $\frac{\partial L_i}{\partial w_k}$  as

$$\frac{\partial L_i}{\partial w_n} = \frac{\partial L_i}{\partial y_n} \frac{\partial y_n}{\partial w_n}$$

Since  $y_1 = x_1 w_1$ , we have  $\frac{\partial y_1}{\partial w_1} = x_1$

$L_i$  depends on the  $\tilde{y}_j$ , applying the chain rule, we obtain

$$\frac{\partial L_i}{\partial w_n} = x_n \sum_j \frac{\partial L_i}{\partial y_j}$$

$\tilde{y}_j$  depends on  $y_i$  through  $t$  and  $\tilde{x}$ .

We thus obtain

$$\frac{\partial \tilde{y}_j}{\partial y_1} = \frac{\partial \tilde{y}_j}{\partial \tilde{x}} \frac{\partial \tilde{x}}{\partial z} \frac{\partial z}{\partial y_1}$$

$$W_2 \text{ have } Z = \sum_m y_m \Rightarrow \frac{\partial Z}{\partial y_2} = 1$$

$$\text{We have } \tilde{x} = \xi(z) = \frac{e^z}{1+e^z}$$

$$\text{Thus } \frac{dF}{dz} = \frac{e^z (1+e^z) - (e^z)^2}{(1+e^z)^2} = \frac{e^z}{(1+e^z)^2}$$

We have have  $\hat{y}_j = \tilde{w}_j \tilde{x}$

$$\Rightarrow \frac{\partial y_i}{\partial x} = \tilde{w}_j$$

$$\Rightarrow \frac{\partial L_i}{\partial w_1} = x_1 \sum_j \tilde{w}_j \frac{e^z}{(1+e^z)^2} \cdot \frac{\partial L_i}{\partial \tilde{y}_j}$$

b) Plugging the values given into the formula from (a) yields

~~30/07/2023~~ ~~23/07/2023~~  
~~28/07/2023~~

$$z = 2 - 4 + 2 = 0 \Rightarrow \frac{e^z}{(1+e^z)^2} = \frac{1}{4}$$

$$\Rightarrow \frac{\partial C_i}{\partial w_1} = \cancel{2w_1} \quad \frac{2}{4} \left( 2 \cdot 1 + 4 \cdot 1 + 1 \cdot 2 \right)$$

$$= \frac{1}{2} (2+4+2) = 4$$

c) The next value for  $w_1$  will be

$$1 - 0.01 \cdot 4 = 0.96$$

Q1 Equation for perpendicular distance:

$$r(x, y) = \frac{ax + by + c}{\sqrt{a^2 + b^2}}$$

$$\text{or } r(x, y) = x \cos \alpha + y \sin \alpha + c$$

Sum of least squares:

$$\sum (r(x_i, y_i))^2 = \sum (x_i \cos \alpha + y_i \sin \alpha + c)^2 = f(\alpha, c)$$

The line goes through the mean point:

$$\frac{\partial f(\alpha, c)}{\partial c} = \sum 2(x_i \cos \alpha + y_i \sin \alpha + c) = 0$$

$$\Leftrightarrow (\cos \alpha \cdot \sum x_i + \sin \alpha \cdot \sum y_i) + c \sum 1 = 0$$

$$\Leftrightarrow \cos \alpha \cdot \frac{\sum x_i}{n} + \sin \alpha \cdot \frac{\sum y_i}{n} + c = 0$$

$\Rightarrow$  work on centered data, i.e., shift ~~data points~~

center into origin

$\Rightarrow$  new line  $x \cos \alpha + y \sin \alpha + c' = 0$  with  $c' = 0$

2 ways to solve for the normal  $(\cos \alpha, \sin \alpha)$  of the line:

1) Closed form solution:

$$\text{Let } x'_i = x_i - \mu_x, \quad y'_i = y_i - \mu_y$$

$$\mu_x = \frac{\sum x'_i}{n}, \quad \mu_y = \frac{\sum y'_i}{n}$$

$\Rightarrow$  Consider centered line:

$$f(\alpha) = \sum (x'_i \cos \alpha + y'_i \sin \alpha)^2$$

$$\begin{aligned} \frac{\partial f(\alpha)}{\partial \alpha} &= \sum 2(x'_i \cos \alpha + y'_i \sin \alpha)(y'_i \cos \alpha - x'_i \sin \alpha) \\ &\stackrel{!}{=} 0 \end{aligned}$$

$$\Rightarrow \sum \cancel{\cos^2 \alpha} x'_i y'_i + y'^2 \sin \alpha \cos \alpha - x'^2 \cos \alpha \sin \alpha - \sin^2 \alpha x'_i y'_i = 0$$
$$\stackrel{!}{=} 0$$

$$\Leftrightarrow \cos^2 \alpha \sum x'_i y'_i + \cos \alpha \sin \alpha \sum (y'^2 - x'^2) - \sin^2 \alpha \sum x'_i y'_i = 0$$

$$\Leftrightarrow \underbrace{\sum x'_i y'_i}_{a} (\cos^2 \alpha - \sin^2 \alpha) + \cos \alpha \sin \alpha \underbrace{\sum (y'^2 - x'^2)}_{b} = 0$$

$$\Leftrightarrow (\cos^2 \alpha - \sin^2 \alpha) a + b \cos \alpha \sin \alpha = 0$$

## 6 continued

$$\Rightarrow 0 = \frac{a \cos^2 \alpha}{\cos^2 \alpha} - \frac{a \sin^2 \alpha}{\cos^2 \alpha} + b \frac{\cos \alpha \sin \alpha}{\cos^2 \alpha}$$

$\Rightarrow$   $0 = a - a \tan^2 \alpha + b \tan \alpha$   
 $\tan \alpha = \frac{\sin \alpha}{\cos \alpha}$

$$\Rightarrow \tan \alpha_1 = \frac{-b + \sqrt{b^2 + 4a^2}}{-2a} = \frac{b - \sqrt{b^2 + 4a^2}}{2a}$$

$$\text{or } \tan \alpha_2 = \frac{-b - \sqrt{b^2 + 4a^2}}{-2a} = \frac{b + \sqrt{b^2 + 4a^2}}{2a}$$

$\Rightarrow$  Select  $\alpha_i$  that minimizes errors, then

Compute  $c$  using  $\alpha_i$

2) Alternative solution:

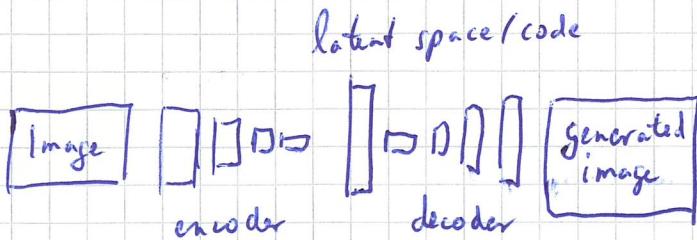
Compute the covariance matrix

$$\Sigma = \sum_{\text{vector } k} \begin{pmatrix} x_i^2 & x_i y_i \\ x_i y_i & y_i^2 \end{pmatrix}$$

The eigenvalue corresponding to the smaller eigenvalue of  $\Sigma$  corresponds to  $(\cos \alpha, \sin \alpha)$

## Example Question: Variational Autoencoders

a)



Given an image  $x$  as input, the encoder (typically a convolutional neural network) maps it to a point  $z$  in the latent space.

The encoder thus defines a probability distribution  $q(z|x)$  in the latent space.

Given a point  $z$  in the latent space as input, the decoder (typically a CNN) maps  $z$  to an image  $x$ . The decoder thus defines a probability distribution  $p(x|z)$  in the space of images.

b) Given a probability distribution  $p(z)$  simple on the latent space (typically a normal or distribution), we sample a point  $z$  in latent space according to  $p(z)$ . We then use the decoder to map it to an image.

c) The term  $D_{KL}(q(z|x) \parallel p(z))$  minimizes the Kullback - Leibler (KL) divergence between the distributions  $q(z|x)$  and  $p(z)$  defined above. The KL divergence term tries to align  $q(z|x)$  (which is learned through the encoder) with the chosen distribution  $p(z)$ . Intuitively, the term tries to ensure that points  $z$  that are likely to be

Sampled ~~according~~ according to  $p(z)$  correspond to points that likely result in valid generated images.