

SSY098 - Image Analysis

Lecture 1 - Introduction and Linear Classifiers

*Torsten Sattler
(slides adapted from Olof Enqvist)*

Team



Rasmus Kjær Høier
hier@chalmers.se



Carl Toft
carl.toft@chalmers.se



José Pedro Lopes Iglesias
jose.iglesias@chalmers.se



Torsten Sattler torsat@chalmers.se

Course Information

- No exam!
- Exercises (optional, non-graded)
- Four mandatory lab assignments
 - Hand in code and short report
 - Pass or fail
 - Recommended to do labs in pairs or groups of 3
 - Deadlines (23:59): February 2nd, February 12th, February 23rd, March 4th
- Mandatory mini project (four choices):
 - Hand in code, report
 - Graded
 - Has to be done individually
- Deadlines are strict, but you have 8 late days in total
- Revisions of labs possible

Grading

- Grade 3:
 - Pass all 4 labs
 - Pass basic part of the project
- Grade 4:
 - Pass all 4 labs, pass basic part of the project
 - Theoretical questions from project **or** advanced part
- Grade 5:
 - Pass all 4 labs, pass basic part of the project
 - Theoretical questions from project **and** advanced part
- Labs 2-4 have theoretical questions
 - Get 12 out of 18 points to get a better grade (only if pass)
 - Have to be answered individually
- Late days also apply to project (8 in total)
- Theoretical questions from lab cannot be revised
- More details on projects on Canvas page and later in course

Check Canvas Course Page Regularly

≡ SSY098 > Syllabus

Ip3 VT20

Image analysis

[Jump to today](#)

[Home](#)

Syllabus

[Files](#)

[Assignments](#)

[Discussions](#)

[Grades](#)

[People](#)

Course-PM

SSY098 Image analysis Ip3 VT20 (7.5 hp)

Course is offered by the department of Electrical Engineering

Contact details

- examiner, lecturer: Torsten Sattler, torsat@chalmers.se, room 7316, EDIT building
- teaching assistants:
 - Rasmus Kjær Høier, hier@chalmers.se, room 7318, EDIT building
 - José Pedro Lopes Iglesias, jose.iglesias@chalmers.se, room 7412, EDIT building
 - Carl Toft, carl.toft@chalmers.se, room 7314, EDIT building

Course purpose

The main aim of the course is to give a basic introduction to the algorithms and mathematical methods used in image analysis, to an extent that will allow the student to handle industrial image analysis problems. In addition the aim is to help the student

Learning Materials

- Lecture Notes from SSY097
 - Do not cover all material covered in the lecture
 - Cover material not covered in the lecture
- Book: Szeliski, Computer Vision: Algorithms and Applications
 - More detailed than lecture notes, provides pointers to literature
 - Use to read up on technical details
- Book: Goodfellow et al., Deep Learning
 - Very detailed introduction to Deep Learning
 - Covers generative neural networks (not covered in lecture notes)
- Books are freely available, see Canvas for links

Ask Questions

- Feel free to ask questions any time
- Please give us feedback (positive and negative) about the lecture, the labs, ...
 - Talk to me before or after the lecture
 - Write me an email
- Happy to adapt based on your suggestions

Today

- What is Image Analysis
- Modeling Images
- Filtering and Linear Classifiers


What is Image Analysis?

Image Classification



Is there a robin in this image?

Image Classification



IN CS, IT CAN BE HARD TO EXPLAIN
THE DIFFERENCE BETWEEN THE EASY
AND THE VIRTUALLY IMPOSSIBLE.

© Randall Munroe (xkcd.com)
<https://xkcd.com/1425/>

Computer Vision

MASSACHUSETTS INSTITUTE OF TECHNOLOGY
PROJECT MAC

Artificial Intelligence Group
Vision Memo. No. 100.


July 7, 1966

THE SUMMER VISION PROJECT

Seymour Papert

The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

Object Detection / Recognition




ICRA 2018 plenary talk by Rodney Brooks

Object Detection



Where is the robin?

Object Detection



[YOLOv3](#)

Image Segmentation

Full-Resolution Residual Networks for Semantic Segmentation in Street Scenes

Tobias Pohlen, Alexander Hermans,
Markus Mathias, Bastian Leibe

Visual Computing Institute, Computer Vision Group
RWTH Aachen University




Visual Computing Institute
Computer Vision
Prof. Dr. Bastian Leibe

RWTHAACHEN
UNIVERSITY



[Full-Resolution Residual Networks \(FRRNs\) for
Semantic Image Segmentation in Street Scenes](#)

Instance-Level Segmentation




[Masked R-CNN](#)

3D Segmentation



3D Reconstruction

Dubrovnik



[Structure-from-Motion Revisited](#)


3D Reconstruction

South Building


128 images
51 million points

[Pixelwise View Selection for
Unstructured Multi-View Stereo](#)

Photogrammetry



Photogrammetry



Prague in the 19th century

Mixed Reality




[Microsoft HoloLens](#)

21:28 ● ↗



Your location
Home

13 min 42 min 44 min 13 min 14 mi



44 min (3.5 km)

via Fjällgatan

Start

Start AR

Steps & more

Mixed Reality


Google Maps

Part 1. Basic image handling and sparse features





Lab 1


Part 2. Machine learning for classification, detection, and segmentation




Lab 2
Project A




Part 3. Robust model estimation and image registration



Part 4. Camera geometry and basic 3D reconstruction



Part 5. Generative models for machine learning and machine learning in 3D vision



Project D

PROGRESSIVE GROWING OF GANs FOR IMPROVED QUALITY, STABILITY, AND VARIATION

Tero Karras
NVIDIA

Timo Aila
NVIDIA

Samuli Laine
NVIDIA

Jaakko Lehtinen
NVIDIA
Aalto University




Course Schedule


Jan. 20	Introduction, Linear classifiers and filtering	
Jan. 23	Filtering, gradients, scale	Lab 1
Jan. 27	Local features	
Jan. 30	Learning a classifier	
Feb. 3	Convolutional neural networks	Lab 2
Feb. 6	More convolutional neural networks	
Feb. 10	Robust model fitting and RANSAC	
Feb. 13	Image registration	Lab 3
Feb. 17	Camera Geometry	
Feb. 20	More camera geometry	
Feb. 24	Generative neural networks	
Feb. 27	Generative neural networks	
Mar. 2	TBA	
Mar. 9	TBA	

Modeling Images

What is an Image?



What is an Image?



255	255	255	255	255	0	0	0	0	0	255	255	255	255	255	255
255	255	255	0	0	150	150	150	150	255	255	0	0	255	255	255
255	255	0	255	255	150	150	150	150	255	255	255	255	0	255	255
255	0	255	255	150	150	150	150	150	255	255	255	255	255	0	255
255	0	255	150	150	255	255	255	255	150	150	255	255	255	0	255
0	150	150	150	255	255	255	255	255	150	150	150	150	150	150	0
0	150	150	150	255	255	255	255	255	150	150	150	255	255	150	0
0	255	150	150	255	255	255	255	255	150	255	255	255	255	255	0
0	255	255	150	150	255	255	255	255	150	150	255	255	255	255	0
0	255	255	150	150	150	150	150	150	150	150	255	255	150	150	0
0	255	150	150	0	0	0	0	0	0	0	150	150	150	150	0
255	0	0	0	255	255	0	255	255	0	255	255	0	0	0	255
255	255	0	255	255	255	0	255	255	0	255	255	255	0	255	255
255	255	0	255	255	255	255	255	255	255	255	255	255	0	255	255
255	255	255	0	255	255	255	255	255	255	255	255	255	0	255	255
255	255	255	255	0	0	0	0	0	0	0	0	0	255	255	255

Images are Vectors

Scalar multiplication, dot product

$2 *$

0	0	0	0	0	0	0	0
0	0	150	150	150	150	0	0
0	150	150	150	150	150	150	0
0	150	150	150	150	150	150	0
0	150	150	150	150	150	150	0
0	150	150	150	150	150	150	0
0	0	150	150	150	150	0	0
0	0	0	0	0	0	0	0

=

0	0	0	0	0	0	0	0
0	0	300	300	300	300	0	0
0	300	300	300	300	300	300	0
0	300	300	300	300	300	300	0
0	300	300	300	300	300	300	0
0	300	300	300	300	300	300	0
0	0	300	300	300	300	0	0
0	0	0	0	0	0	0	0

$$\begin{matrix} 0 & 255 & 0 \\ 255 & 255 & 255 \end{matrix}$$

*

$$\begin{matrix} 0 & 255 & 0 \\ 0 & 255 & 0 \\ 0 & 0 & 0 \end{matrix}$$

$$= 2 * 255^2$$

$$\begin{matrix} 0 & 255 & 0 \end{matrix}$$

Images are Vectors



Average
speed limit sign

Images are Vectors



Average
speed limit sign



image source: [pinterest](#)

Average
president

Color



Color



Color





RGB



54	122	124	145	153	100	123	109
54	233	232	211	198	188	177	155
255	121	121	121	78	13	87	76
255	56	78	89	23	23	12	12
255	78	18	24	25	53	23	54
255	255	23	34	24	35	54	23
255	244	34	53	35	35	52	43
255	34	56	34	87	76	56	87
244	98	9	76	58	48	59	28
67	28	38	28	28	38	38	38
28	8	0	29	9	82	0	

Channel-wise operations

Digital Cameras





Bayer filter

Digital Cameras



Digital Cameras



Digital Cameras



Demosaicing

Lighting and White Balance



RGB-D Images







image credit: Thomas Schöps


3D Images



Computed tomography



3D Images




Computed tomography






Filtering and Linear Classifiers

Count the Nuclei





Classification

nuclei



other




21 x 21 patches

A Linear Classifier

Target /
template

$w \cdot$



$> \tau$


Decision
threshold

$w \cdot$



$> \tau$

$w \cdot$



$< \tau$


$w \cdot$



$< \tau$


Designing a Classifier

high weight



low
weight

.



$> \tau$




$> \tau$

Not a good classifier!

Designing a Better Classifier

$$w = \frac{1}{N} (w' - \mu_{w'} \mathbf{1})$$




w'

Motivation: uniform patches should have response ~ 0


Designing a Better Classifier

Example:


$$w = 1/9 * ($$




$$- 224 *$$



$$w^* = 4604.44$$




$$w^* = 0$$




Choosing A Template

$w =$




some nucleus

$w =$



mean nucleus


Choosing a Threshold




1.72




1.88



1.26



1.45



0.56



-0.07




-0.06




0.75

Count the Nuclei




Sliding Window Classification




$$\cdot w < \tau$$

Sliding Window Classification




- $w < \tau$

Sliding Window Classification




- $w < \tau$

Sliding Window Classification




- $w < \tau$

Sliding Window Classification




$$\cdot w < \tau$$

Sliding Window Classification




- $w > \tau$

Sliding Window Classification



- $w < \tau$

Sliding Window Classification



- $w < \tau$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

			0			

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0						

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0	1					

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0	1					

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1	2			

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1	2	1	0	
	1	5	6	4	1	
	2	6				

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1	2	1	0	
	1	5	6	4	1	
	2	6	8			

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1	2	1	0	
	1	5	6	4	1	
	2	6	8	4	1	
	1	4	4	2	0	
	0	1	1	0	0	

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1	2	1	0	
1	5	6	4	1		
2	6	8	4	1		
1	4	4	2	0		
0	1	1	0	0		

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0							
0	0	0	0	0	0	0			
0	0	0	0	0	0	0			
0	0	5	10	5	0	0			
0	0	10	10	5	0	0			
0	0	5	5	0	0	0			
0	0	0	0	0	0	0			
0	0	0	0	0	0	0			

0									
	0	1	2	1	0				
	1	5	6	4	1				
	2	6	8	4	1				
	1	4	4	2	0				
	0	1	1	0	0				

I

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0							
0	0	0	0	0	0	0			
0	0	0	0	0	0	0			
0	0	5	10	5	0	0			
0	0	10	10	5	0	0			
0	0	5	5	0	0	0			
0	0	0	0	0	0	0			
0	0	0	0	0	0	0			

0	0								
	0	1	2	1	0				
1	5	6	4	1					
2	6	8	4	1					
1	4	4	2	0					
0	1	1	0	0					

I

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5} \begin{bmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0	0	0	0	0	0	0
0	0	1	2	1	0	0
0	1	5	6	4	1	0
0	2	6	8	4	1	0
0	1	4	4	2	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0


0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	1	1	0	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0


I

$I \otimes w > 4$




Boundary Padding



#OTDSTOTD#OTD
AOLLOADAOI
#OTDSTOTD#OTD




Sliding Window Classification

 \otimes  $=$ 

Per-Pixel Classification Results



Per-Pixel Classification



Maximum Filter / Non-Maximum Suppression

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

0						

result

Nonlinear Filters

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

	0	0				

result

Nonlinear Filters

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

	0	0	0	0	0	
	0	1				

result

Nonlinear Filters

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

	0	0	0	0	0	
	0	1	0			

result

Nonlinear Filters

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

	0	0	0	0	0	
	0	1	0	0	0	
	0	0	0	0	0	
	0	0	0	1		

result

Nonlinear Filters

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

	0	0	0	0	0	
	0	1	0	0	0	
	0	0	0	0	0	
	0	0	0	1	0	
	0	0	0	0	0	

result

Nonlinear Filters

Is this pixel larger than its neighbors?


1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input


	0	0	0	0	0
0	1	0	0	0	
	0	0	0	0	0
0	0	0	1	0	
	0	0	0	0	0

result

A Sliding-Window Detector




Result



Count: 52

Manual: 48

Instance Level Segmentation



Result from applying Masked R-CNN

Lessons Learned

- Main lessons from this lecture
 - Images are vectors: Addition, subtraction, dot product
 - Linear filters: Mean, dot product-based
 - Nonlinear filters: Non-maximum suppression
 - Sliding window classification
- Next lecture: More filtering, gradients, multi-scale processing

Next Lecture

Jan. 20	Introduction, Linear classifiers and filtering	
Jan. 23	Filtering, gradients, scale	Lab 1
Jan. 27	Local features	
Jan. 30	Learning a classifier	
Feb. 3	Convolutional neural networks	Lab 2
Feb. 6	More convolutional neural networks	
Feb. 10	Robust model fitting and RANSAC	
Feb. 13	Image registration	Lab 3
Feb. 17	Camera Geometry	
Feb. 20	More camera geometry	
Feb. 24	Generative neural networks	
Feb. 27	Generative neural networks	
Mar. 2	TBA	
Mar. 9	TBA	