

SSY097 - Image Analysis

Lecture 2 - Filtering, gradients and scale

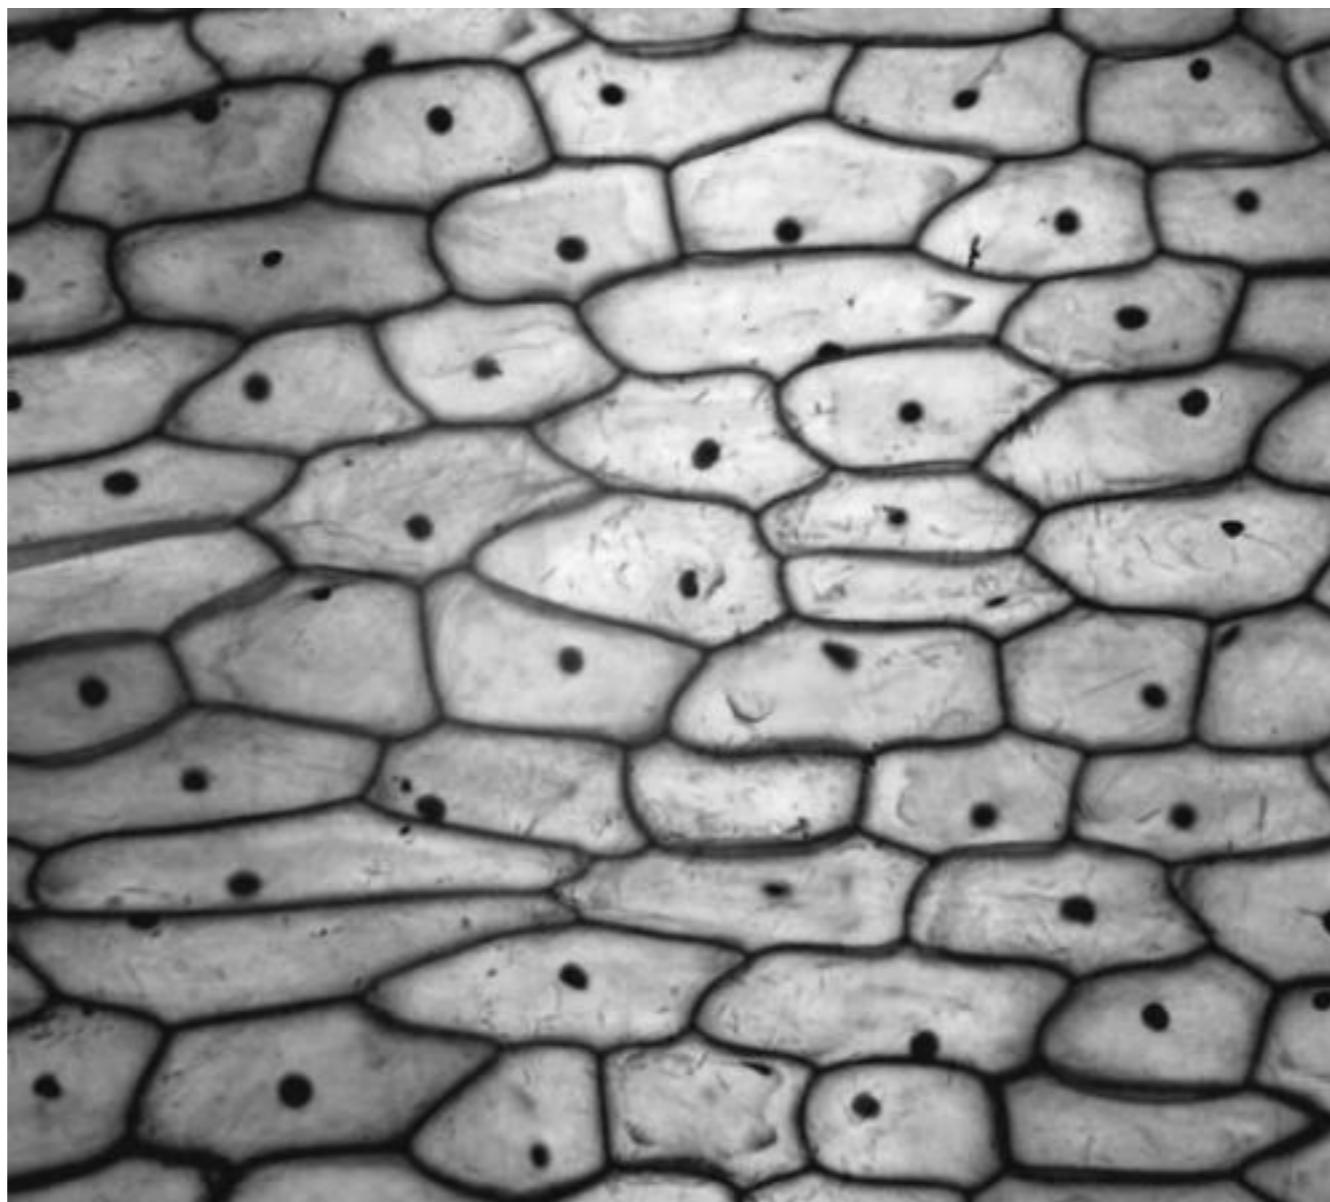
*Torsten Sattler
(slides adapted from Olof Enqvist)*

Last Lecture

Jan. 20	Introduction, Linear classifiers and filtering	
Jan. 23	Filtering, gradients, scale	Lab 1
Jan. 27	Local features	
Jan. 30	Learning a classifier	
Feb. 3	Convolutional neural networks	Lab 2
Feb. 6	More convolutional neural networks	
Feb. 10	Robust model fitting and RANSAC	
Feb. 13	Image registration	Lab 3
Feb. 17	Camera Geometry	
Feb. 20	More camera geometry	
Feb. 24	Generative neural networks	
Feb. 27	Generative neural networks	
Mar. 2	TBA	
Mar. 9	TBA	

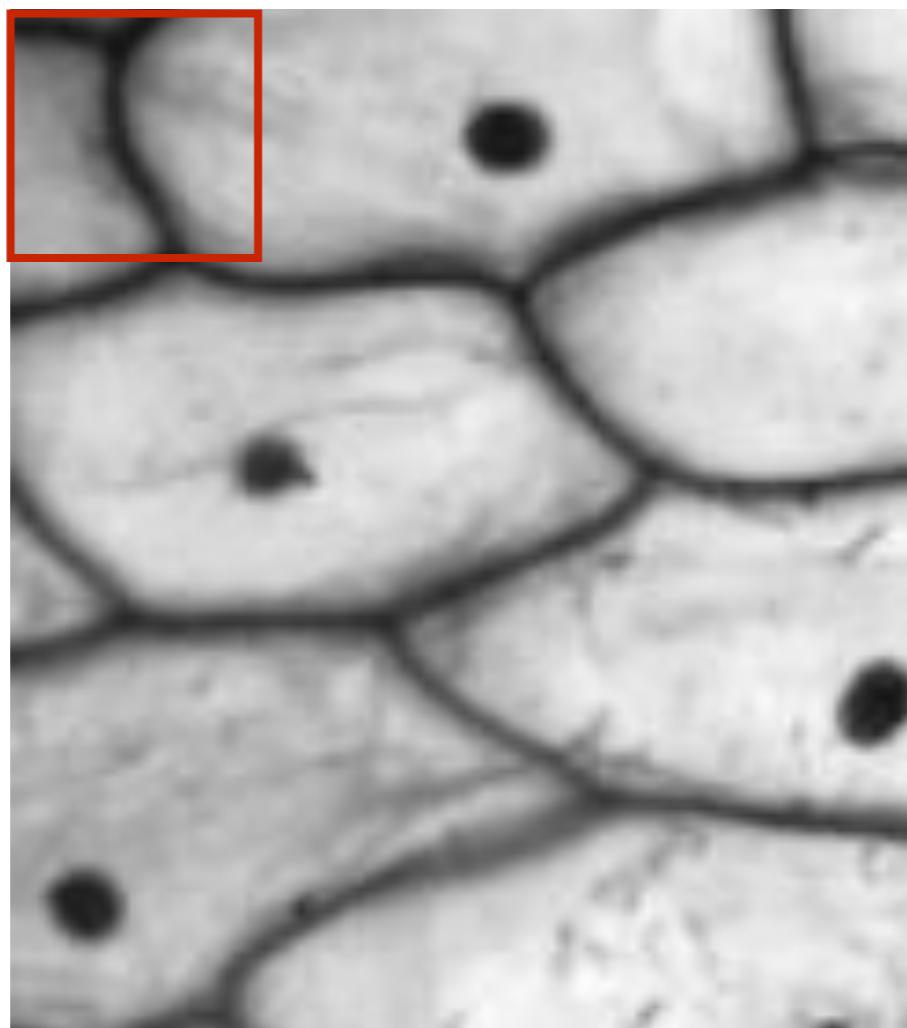
Last Lecture

Application: Counting of objects (nuclei)



Last Lecture

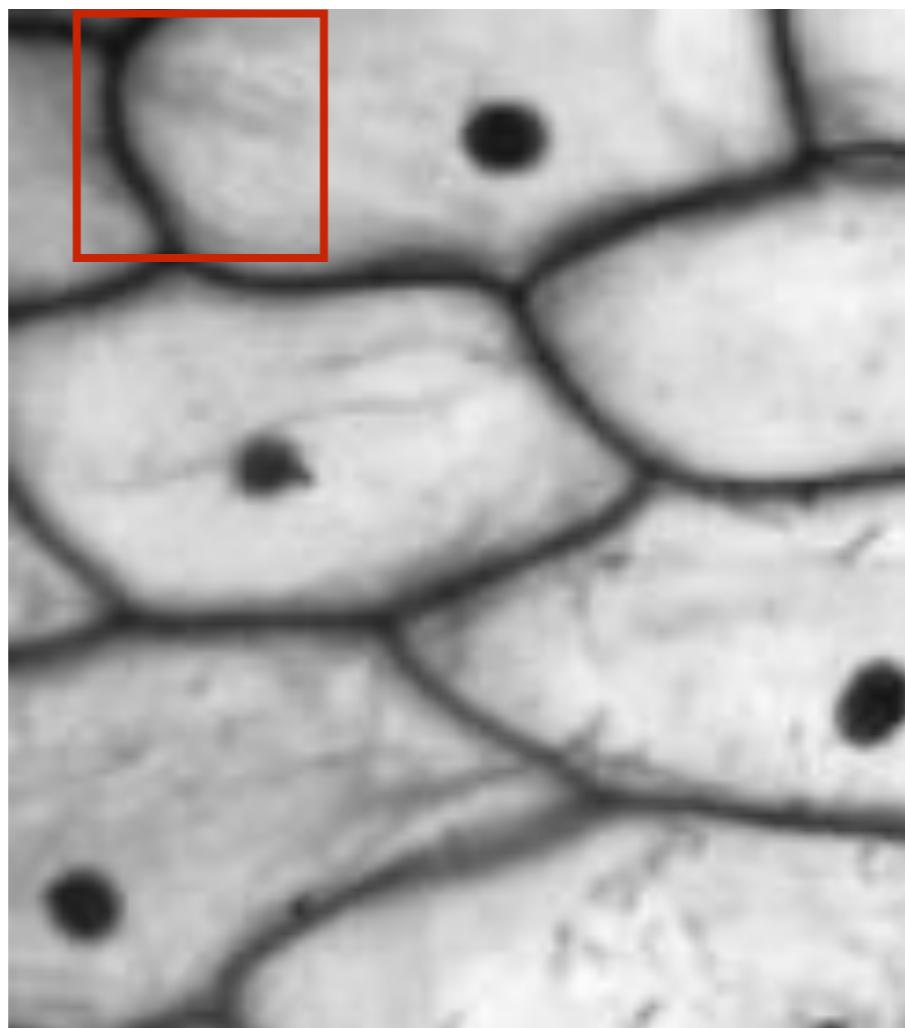
Sliding window classification



- $w < \tau$

Last Lecture

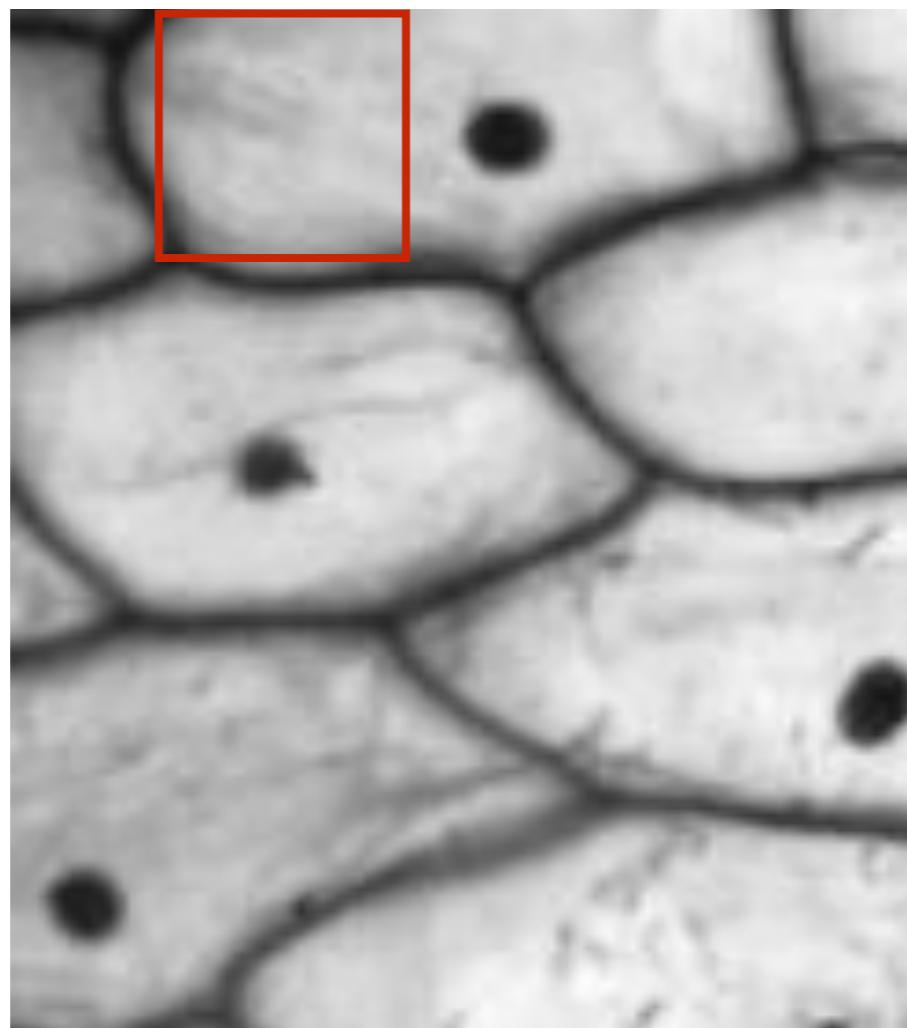
Sliding window classification



- $w < \tau$

Last Lecture

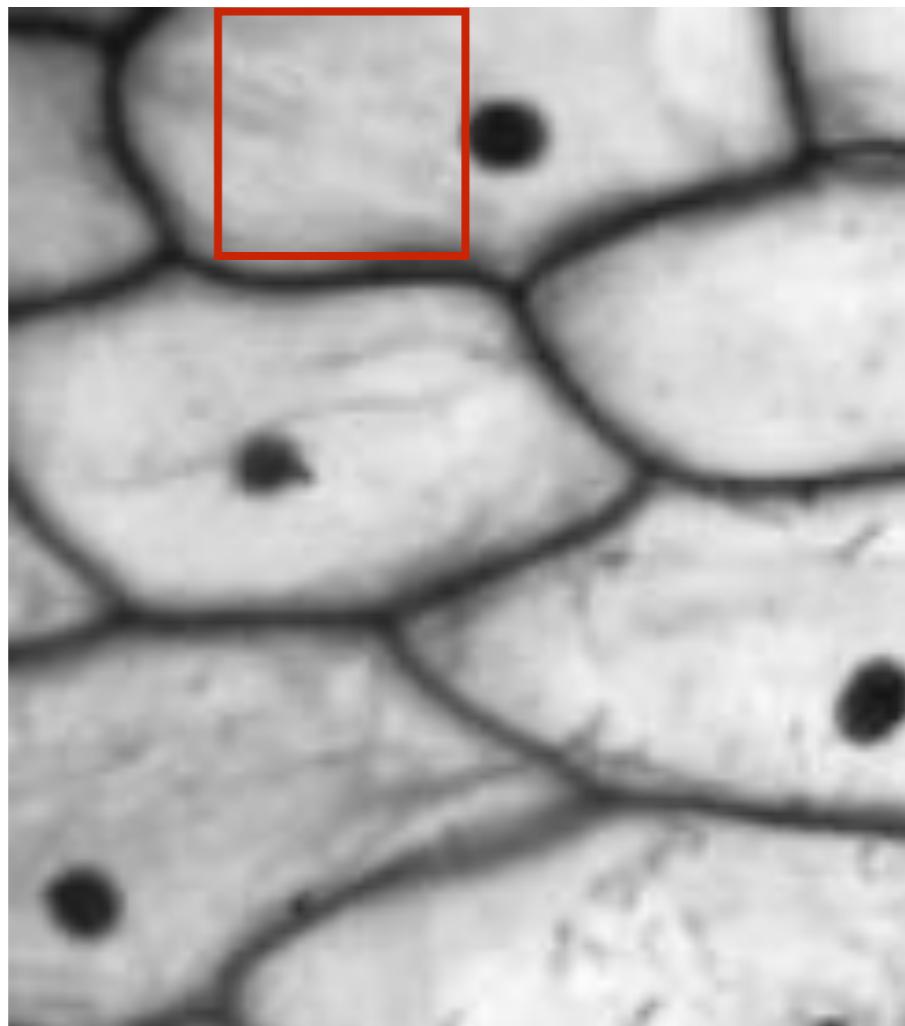
Sliding window classification



- $w < \tau$

Last Lecture

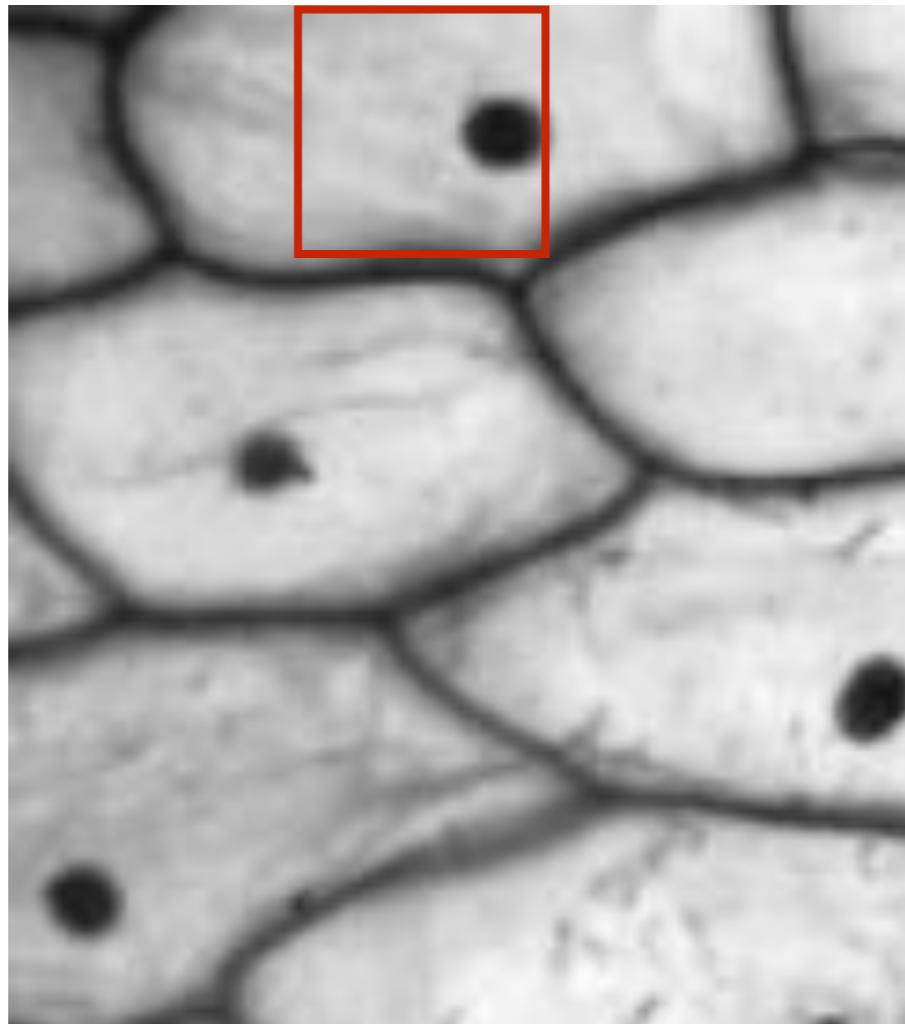
Sliding window classification



- $w < \tau$

Last Lecture

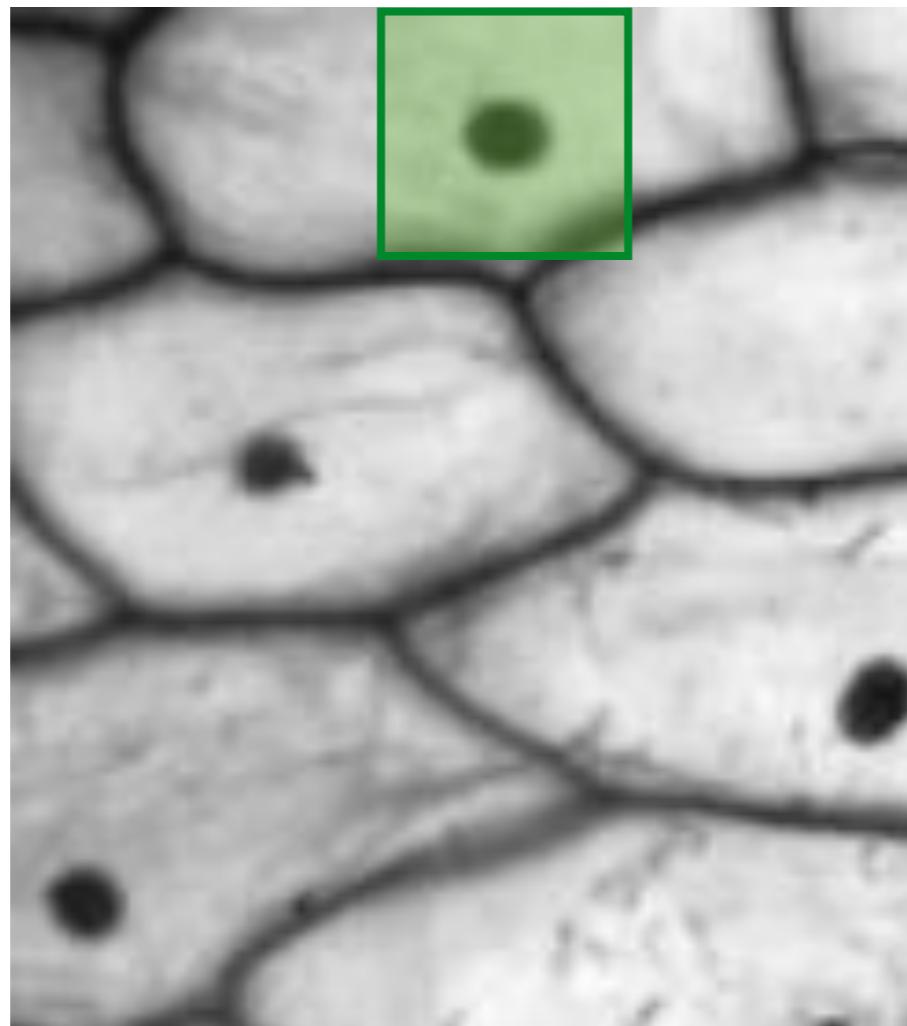
Sliding window classification



- $w < \tau$

Last Lecture

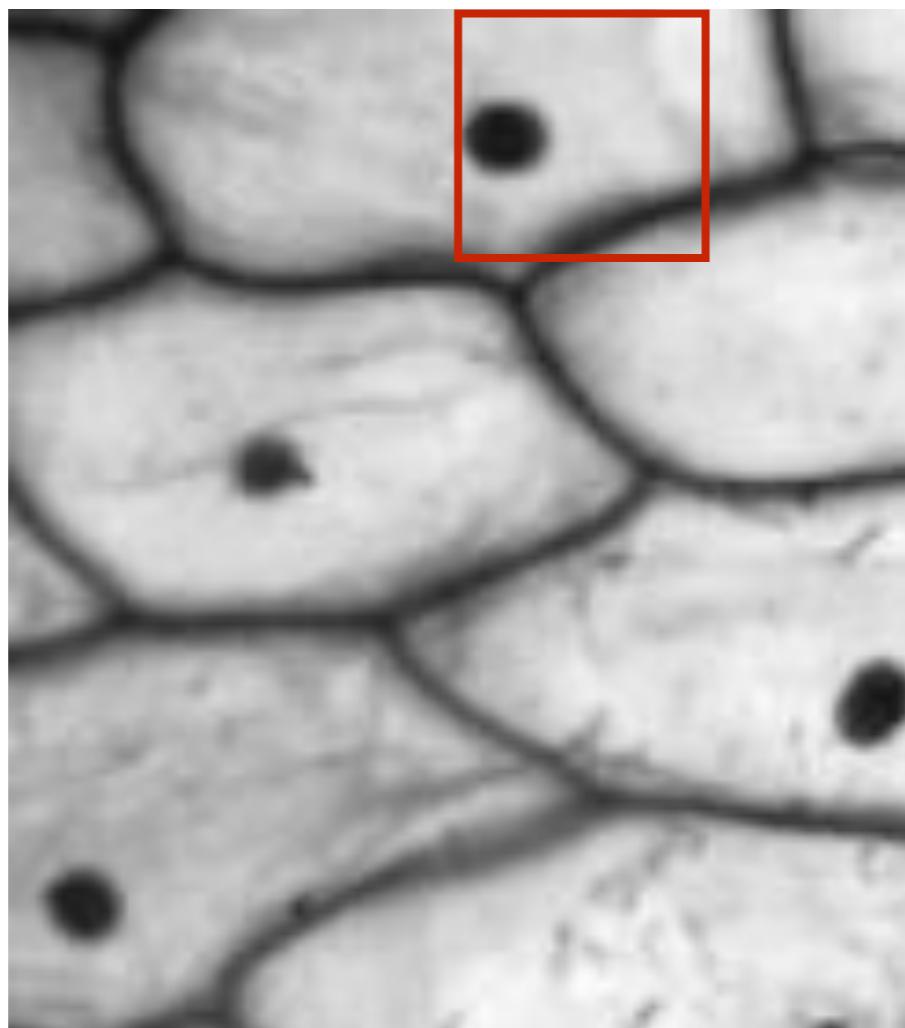
Sliding window classification



- $w > \tau$

Last Lecture

Sliding window classification

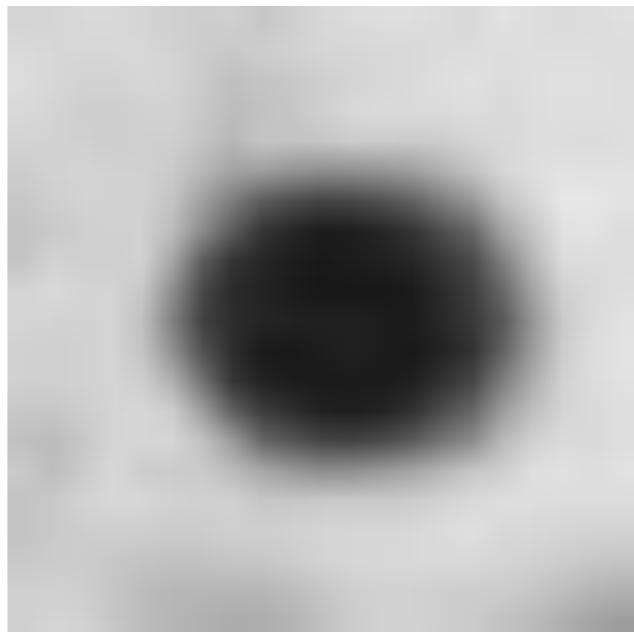


- $w < \tau$

Last Lecture

Linear Filters

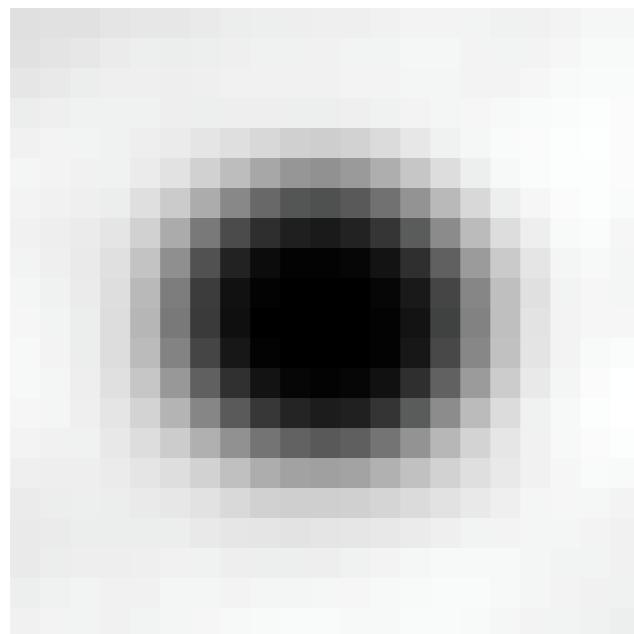
$w \cdot$



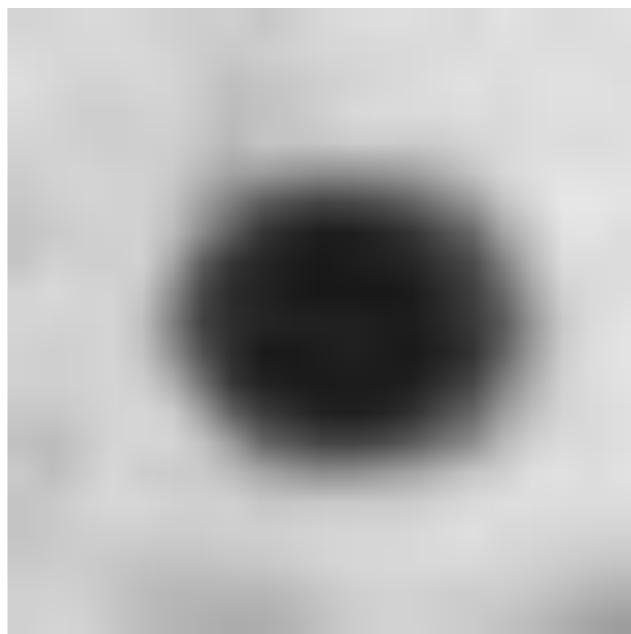
$> \tau$

Last Lecture

Linear Filters



•

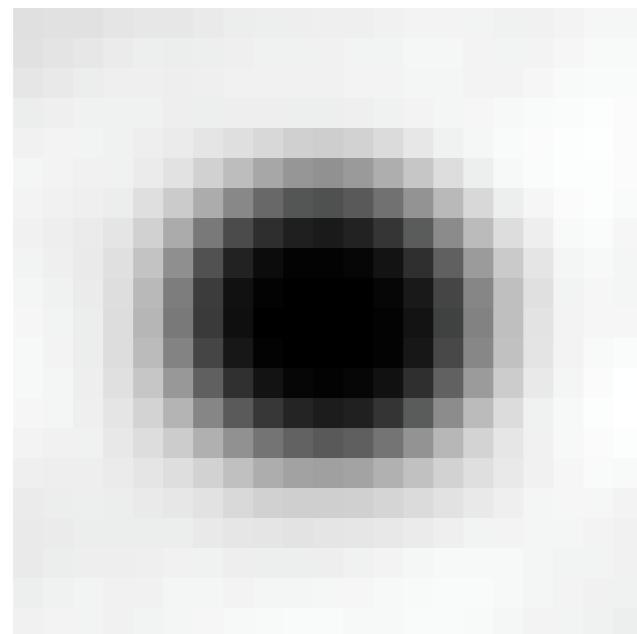


$> \tau$

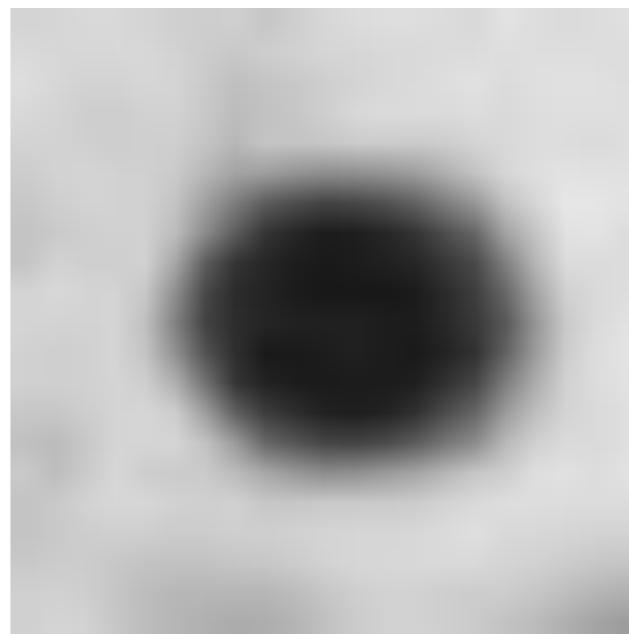
Mean positive / foreground
patch

Last Lecture

Linear Filters



•

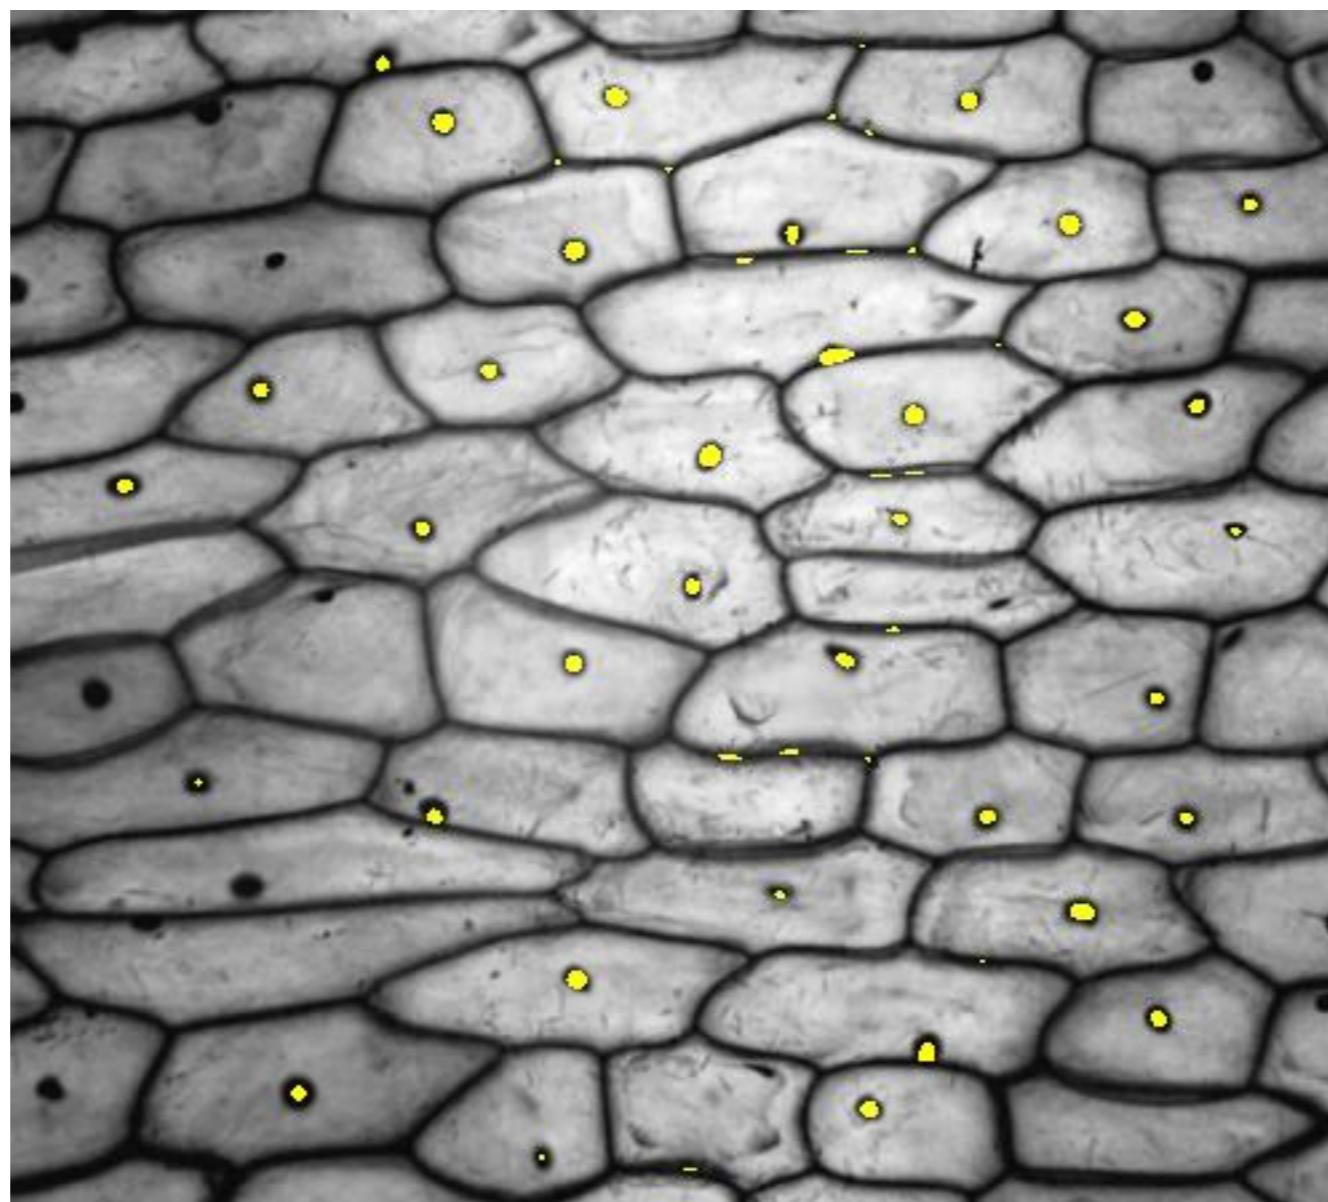


$> \tau$

$$\frac{1}{N_{\text{elements}}} (T - \mu_T \mathbf{1})$$

Last Lecture

Nonlinear filters: Non-maximum suppression



Last Lecture

Nonlinear filters: Non-maximum suppression

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

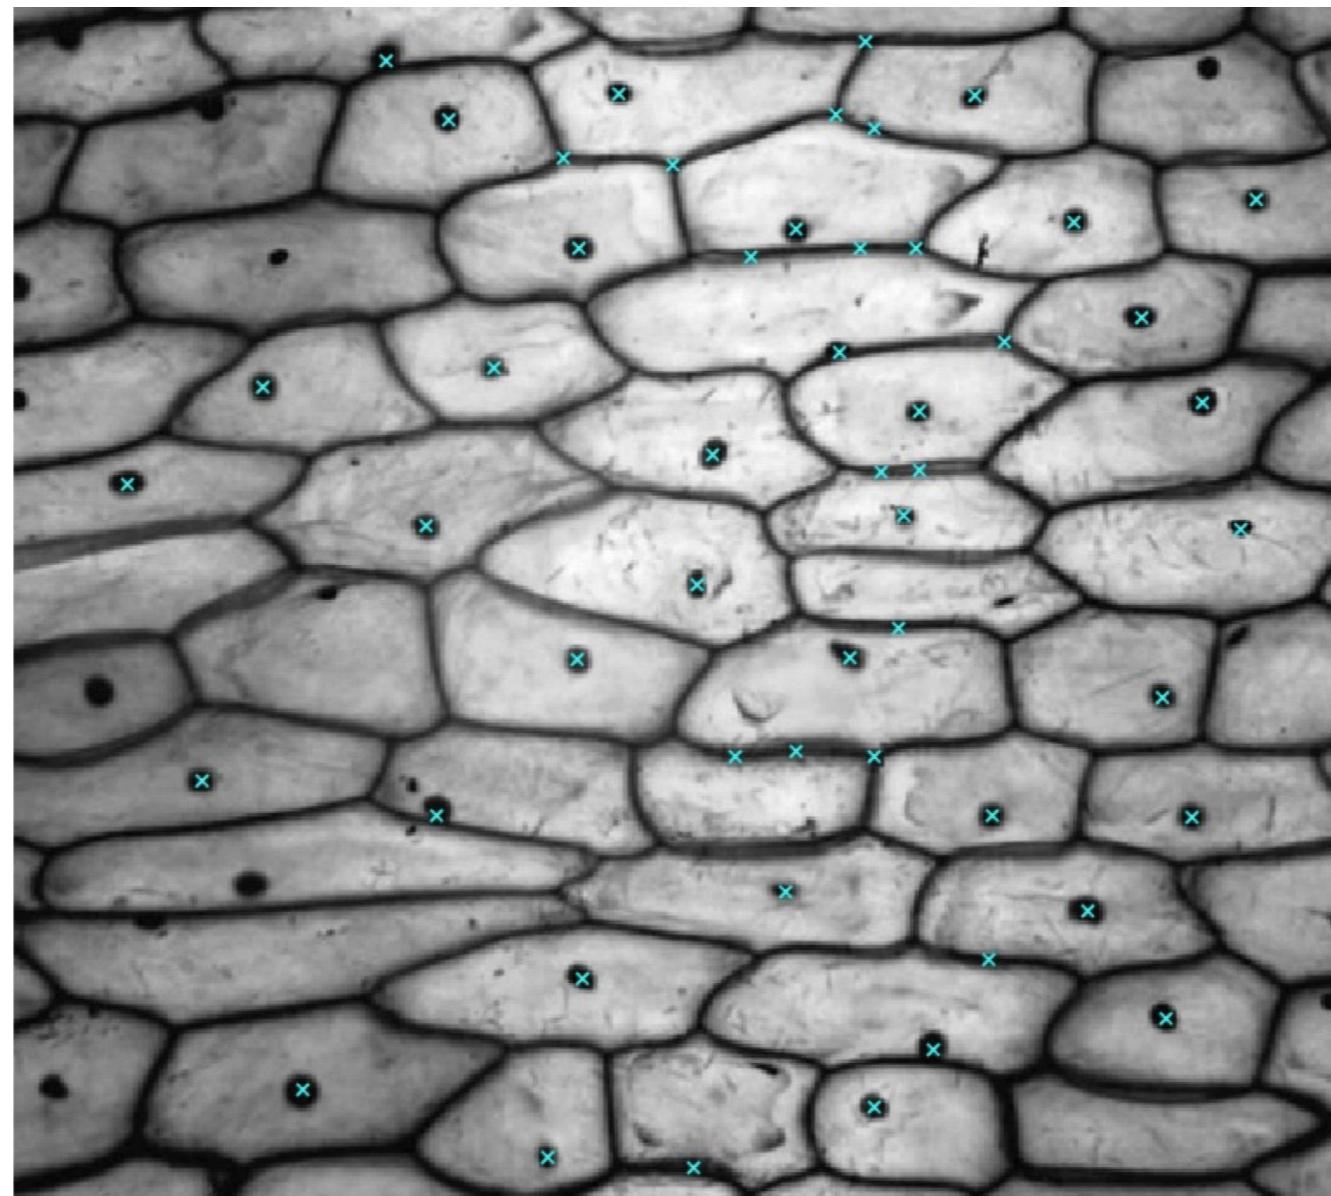
input

	0	0	0	0	0	
0	1					

result

Last Lecture

Nonlinear filters: Non-maximum suppression



Today

- More Filters
- Similarity Measures
- Multi-Scale Processing

More Filters

Applying a Filter

image I

filter f

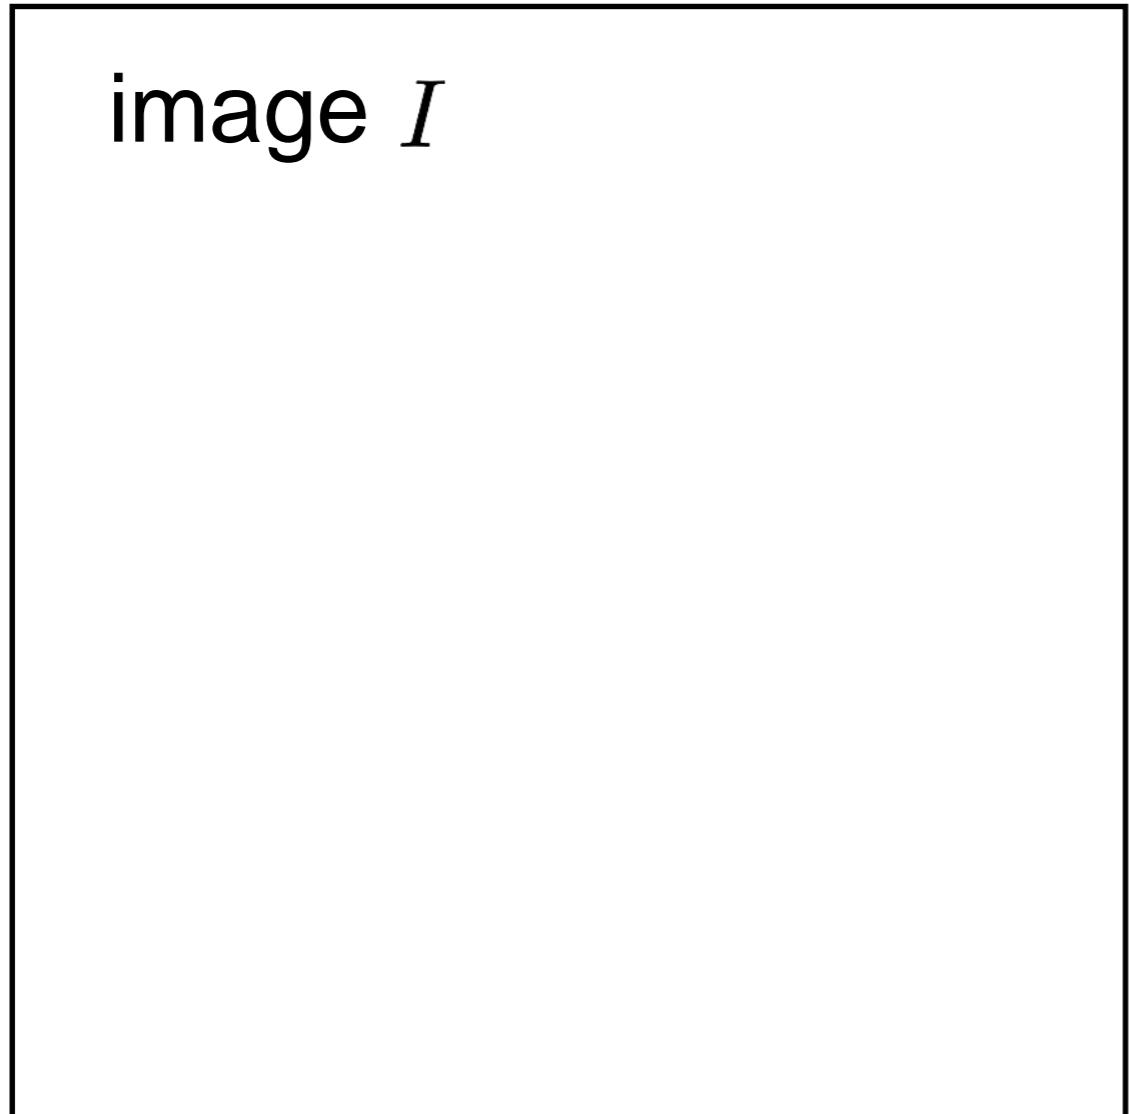
*

$= J$

Applying a Filter

$(0, 0)$

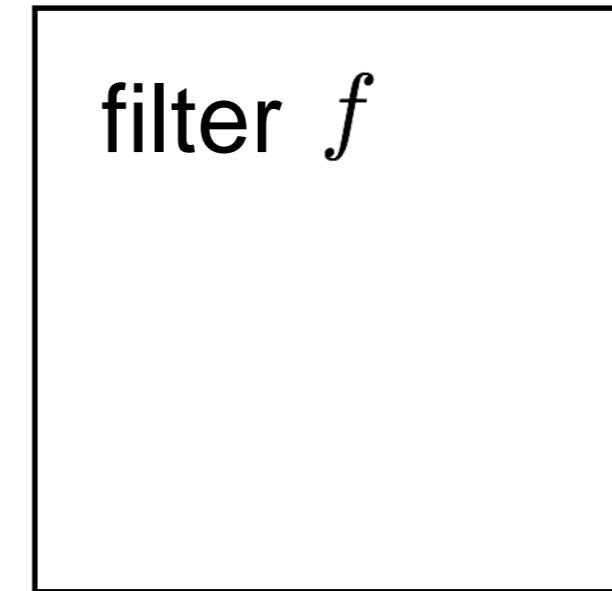
image I



$(-m, -n)$

*

filter f



$= J$

(m, n)

(w, h)

Applying a Filter

(0, 0)

image I

(-m, -n)

filter f

*

= J

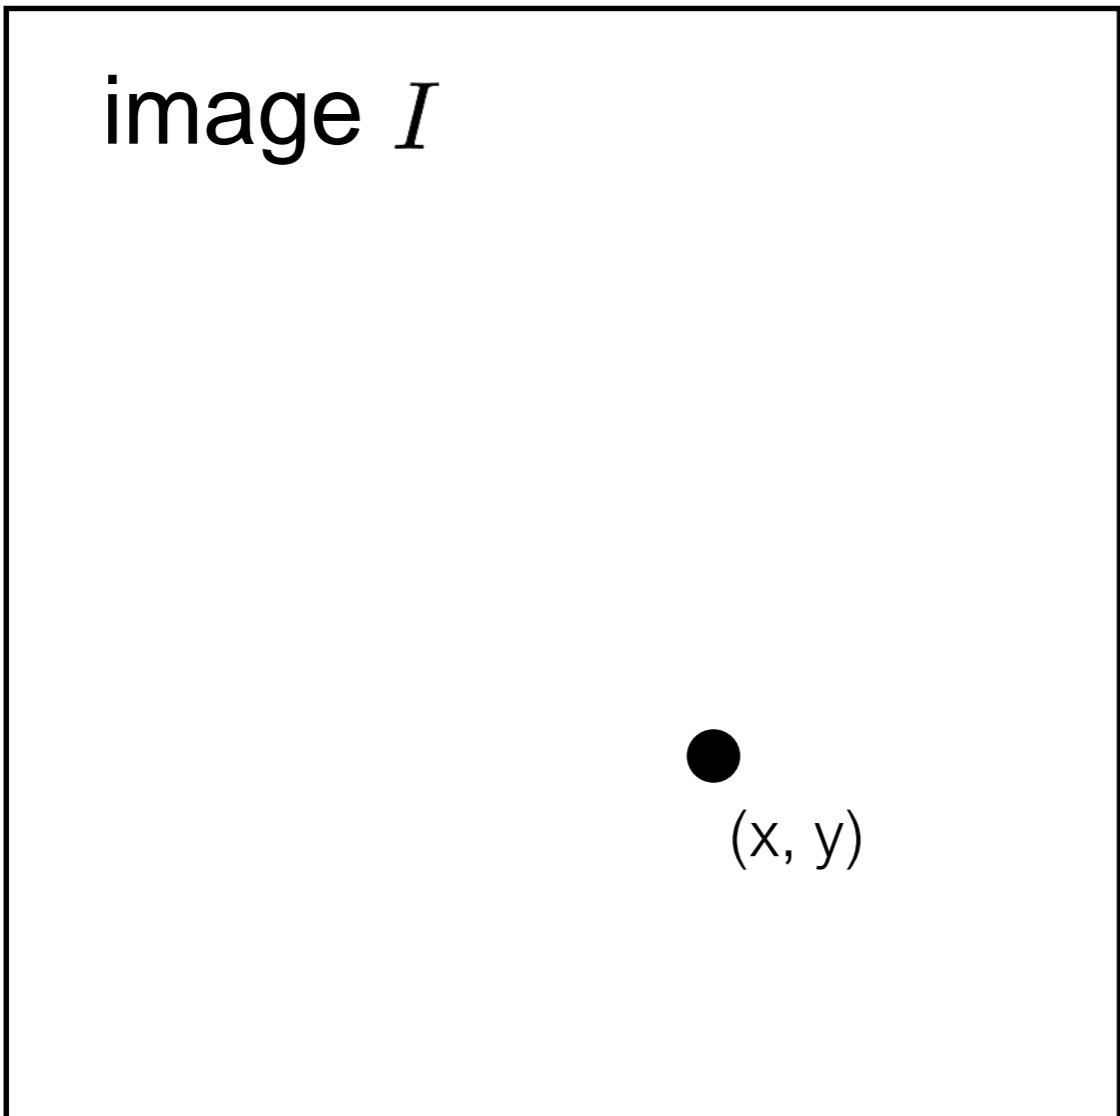
(m, n)

(w, h)

$$J(x, y) = \sum_{u=-m}^m \sum_{v=-n}^n I(x+u, y+v) \cdot f(u, v)$$

Applying a Filter

(0, 0)



(-m, -n)

filter f

*

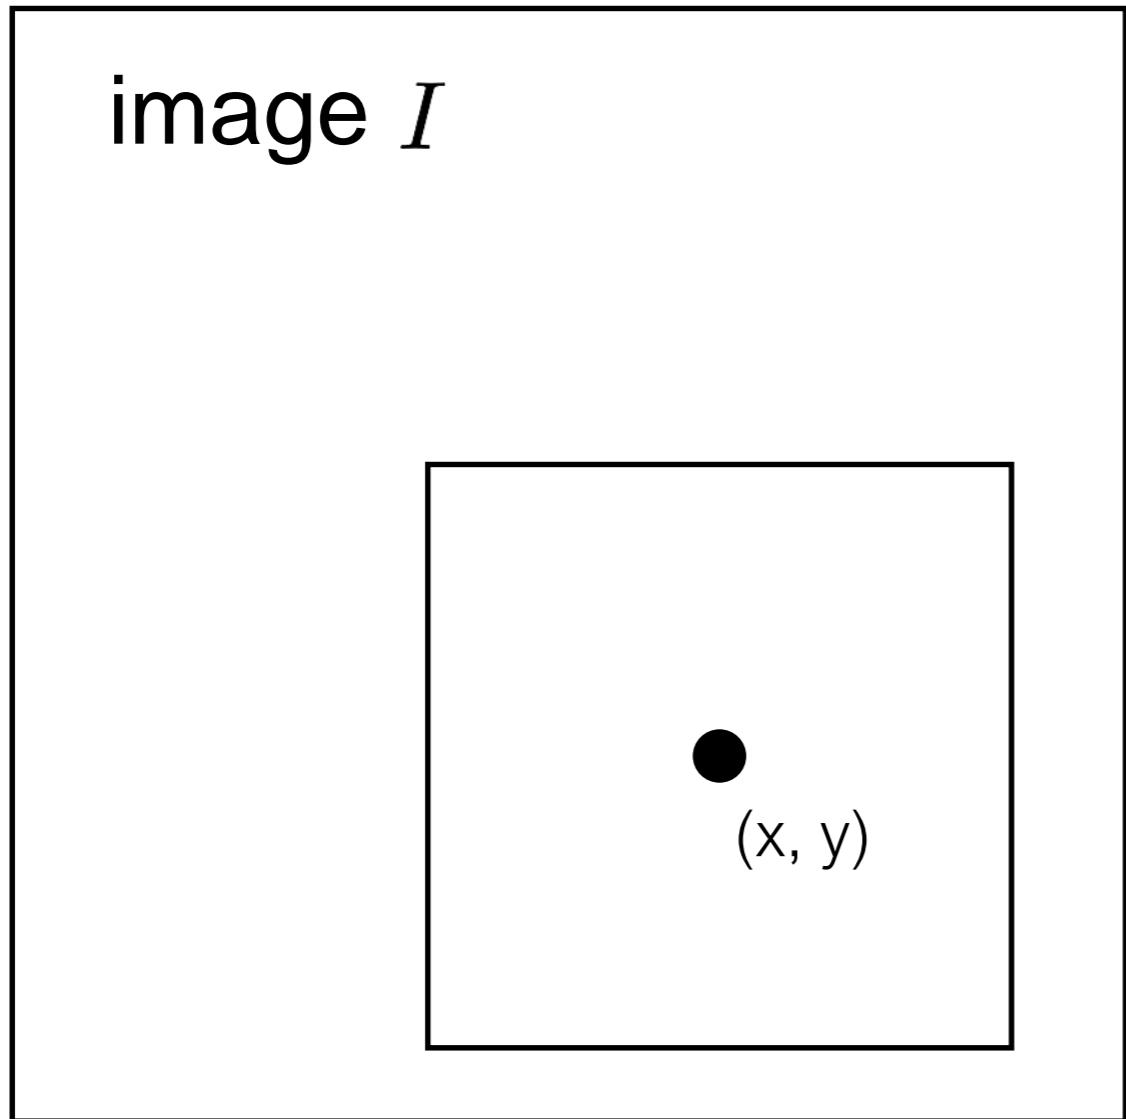
$= J$

(m, n)

$$J(x, y) = \sum_{u=-m}^m \sum_{v=-n}^n I(x+u, y+v) \cdot f(u, v)$$

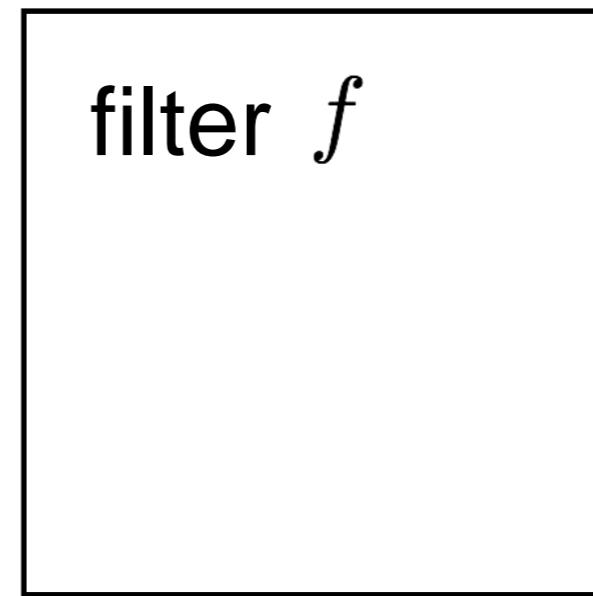
Applying a Filter

(0, 0)



(w, h)

(-m, -n)



*

= J

$$J(x, y) = \sum_{u=-m}^m \sum_{v=-n}^n I(x+u, y+v) \cdot f(u, v)$$

Applying a Filter: A Warning

We are applying filters via cross-correlation:

$$J(x, y) = \sum_{u=-m}^{m} \sum_{v=-n}^{n} I(x + u, y + v) \cdot f(u, v)$$

Often denotes as $J = I \otimes f$

Applying a Filter: A Warning

We are applying filters via cross-correlation:

$$J(x, y) = \sum_{u=-m}^m \sum_{v=-n}^n I(x + u, y + v) \cdot f(u, v)$$

Often denotes as $J = I \otimes f$

Similar to mathematical definition of convolution:

$$J(x, y) = \sum_{u=-m}^m \sum_{v=-n}^n I(x - u, y - v) \cdot f(u, v)$$

Applying a Filter: A Warning

We are applying filters via cross-correlation:

$$J(x, y) = \sum_{u=-m}^m \sum_{v=-n}^n I(x + u, y + v) \cdot f(u, v)$$

Often denotes as $J = I \otimes f$

Similar to mathematical definition of convolution:

$$J(x, y) = \sum_{u=-m}^m \sum_{v=-n}^n I(x - u, y - v) \cdot f(u, v)$$

Convolution typically denoted as $J = I \star f$

Applying a Filter: A Warning

We are applying filters via cross-correlation:

$$J(x, y) = \sum_{u=-m}^m \sum_{v=-n}^n I(x + u, y + v) \cdot f(u, v)$$

Often denotes as $J = I \otimes f$

Similar to mathematical definition of convolution:

$$J(x, y) = \sum_{u=-m}^m \sum_{v=-n}^n I(x - u, y - v) \cdot f(u, v)$$

Convolution typically denoted as $J = I \star f$

But, we will use $J = I \star f$ to denote correlation here!

Average Filter

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Average Filter



$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

image source: <http://graphics.stanford.edu/data/3Dscanrep/>

Average Filter



$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

image source: <http://graphics.stanford.edu/data/3Dscanrep/>

Average Filter



image source: <http://graphics.stanford.edu/data/3Dscanrep/>

Average Filter



image source: <http://graphics.stanford.edu/data/3Dscanrep/>

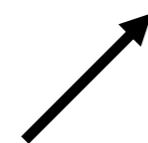
Gaussian Filter

$$G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right)$$

Gaussian Filter

$$G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right)$$

standard deviation



Gaussian Filter

$$G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right)$$

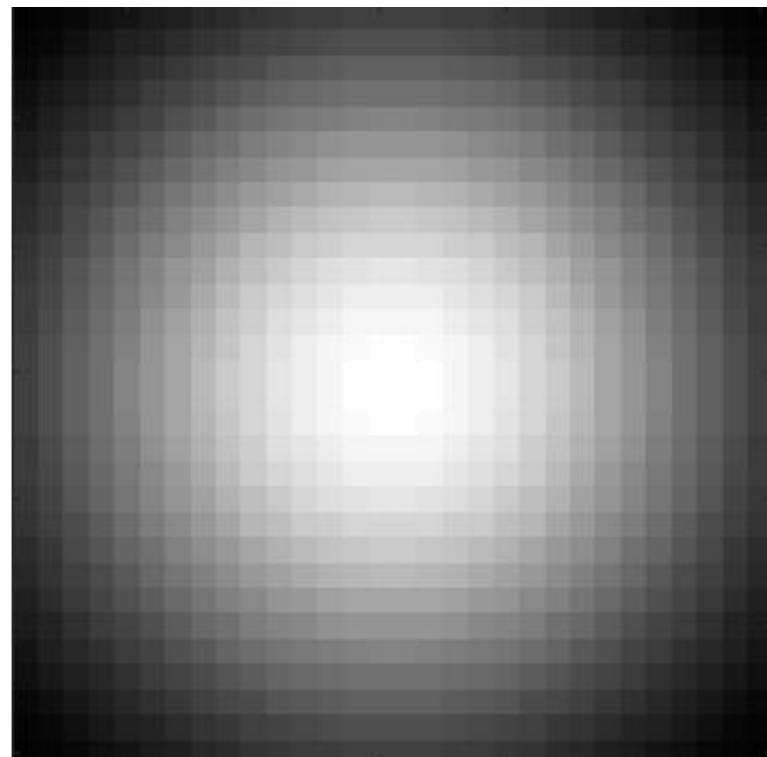
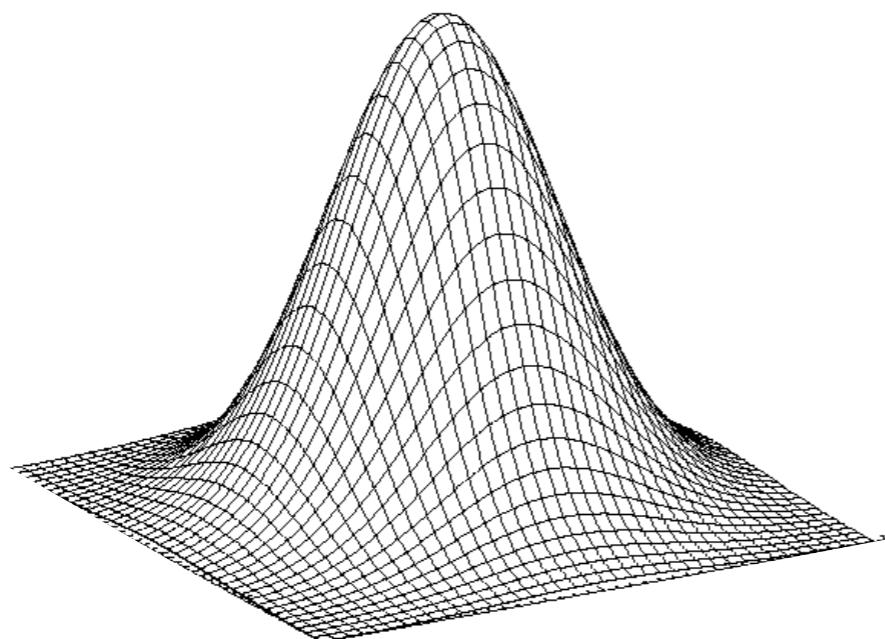
standard deviation 

normalization factor 

Gaussian Filter

$$G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right)$$

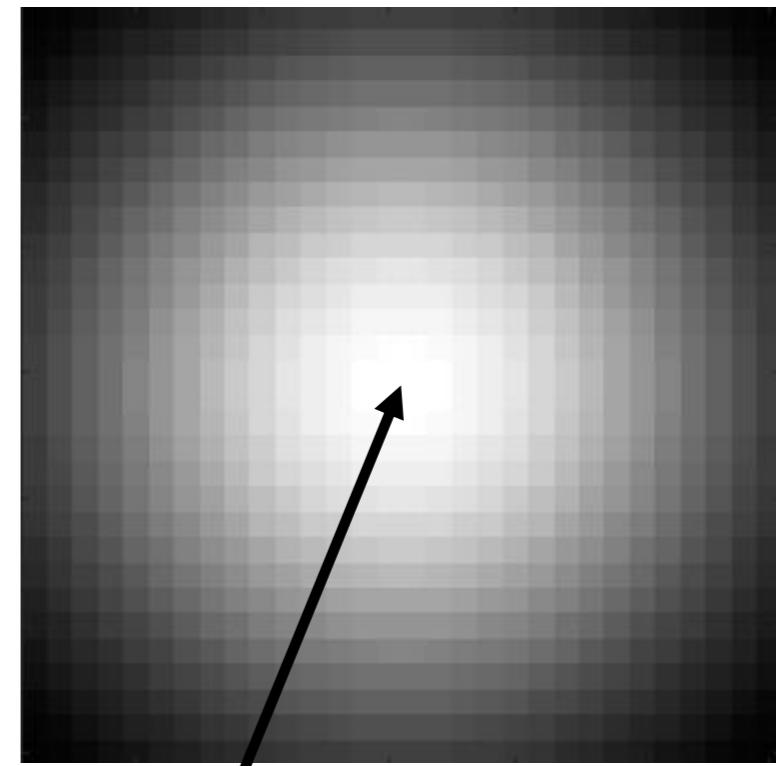
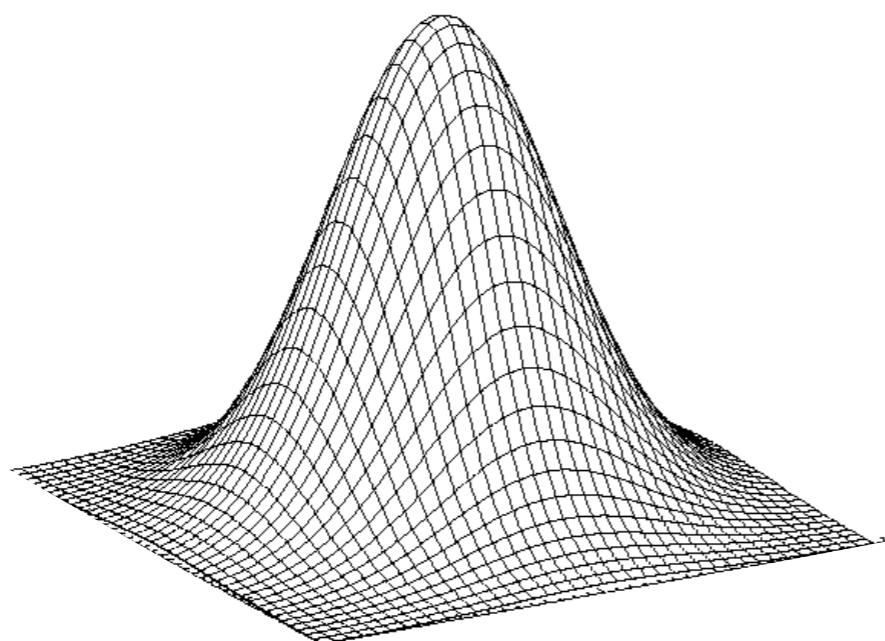
standard deviation normalization factor



Gaussian Filter

$$G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right)$$

standard deviation normalization factor



(0, 0)

Gaussian Filter



Average filter



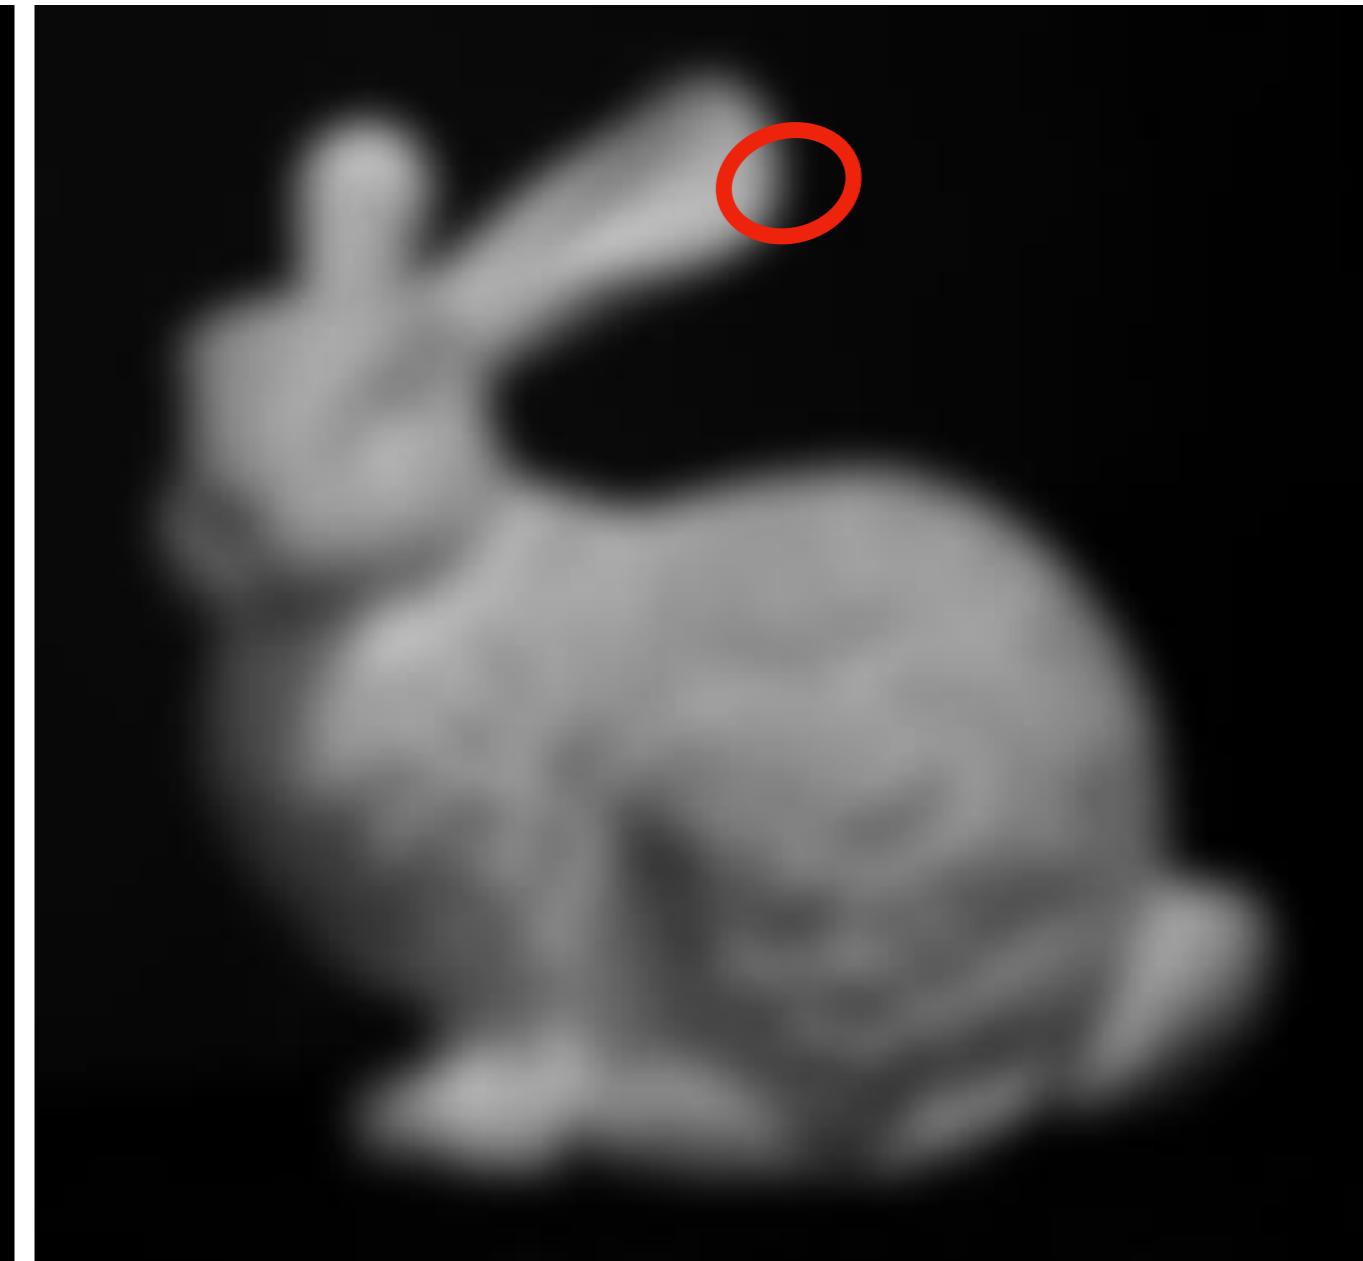
Gaussian filter

image source: <http://graphics.stanford.edu/data/3Dscanrep/>

Gaussian Filter

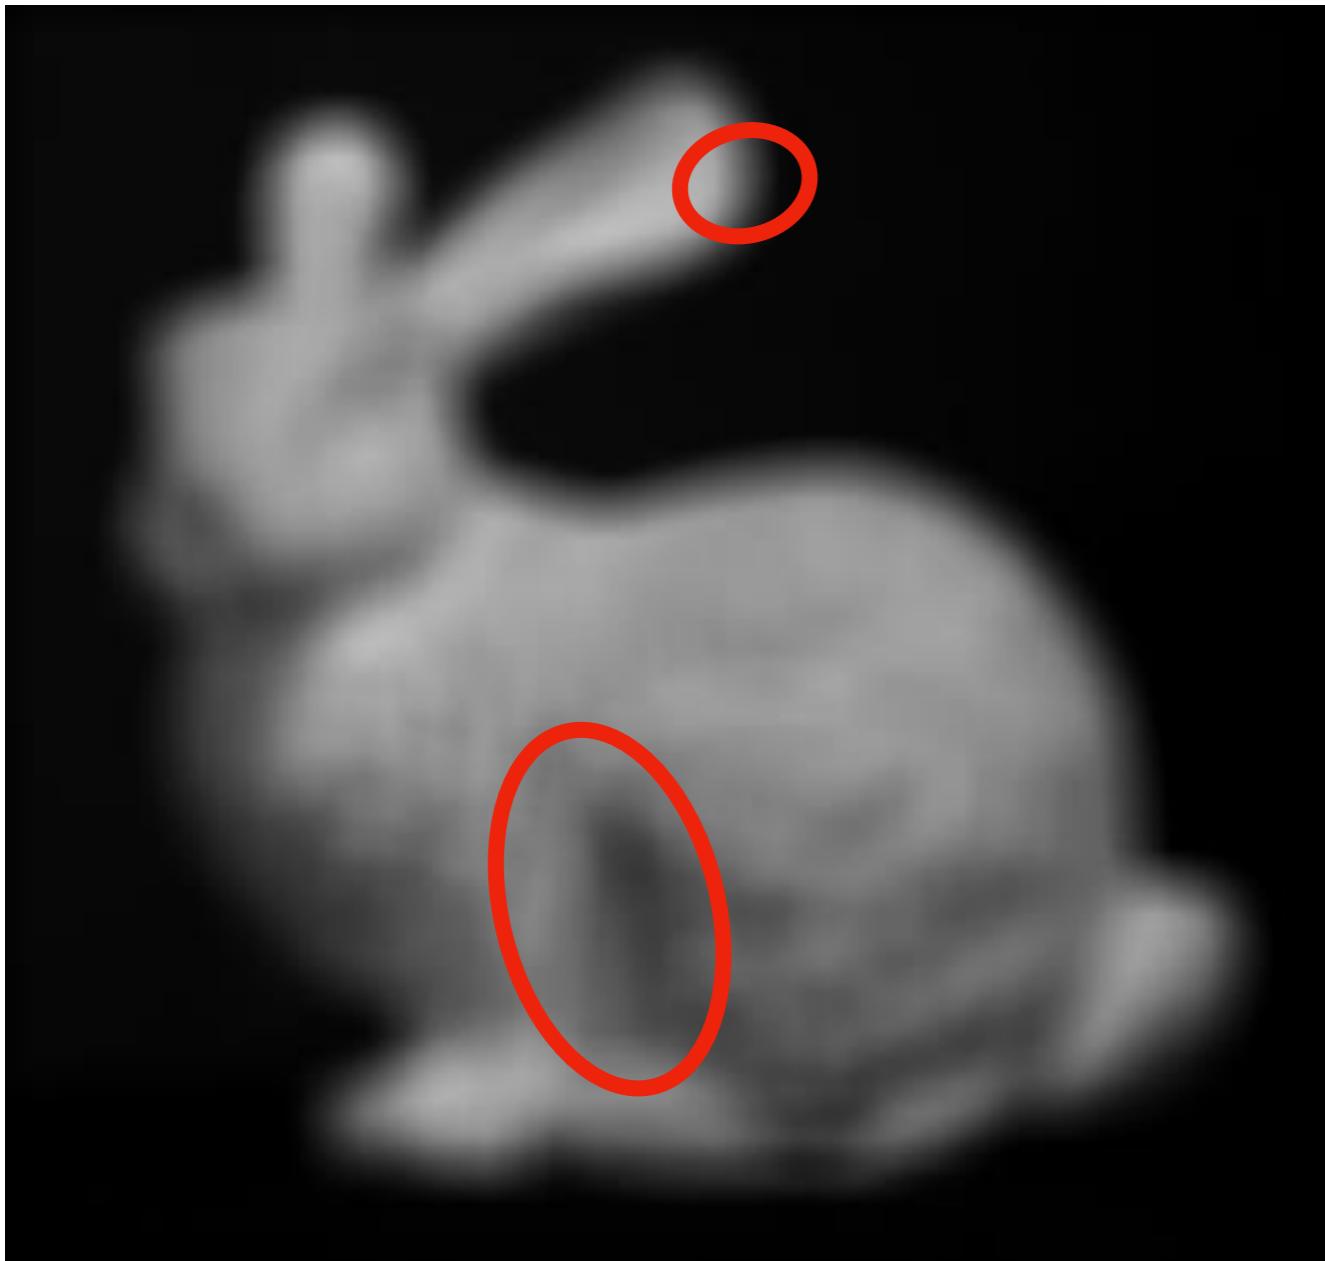


Average filter

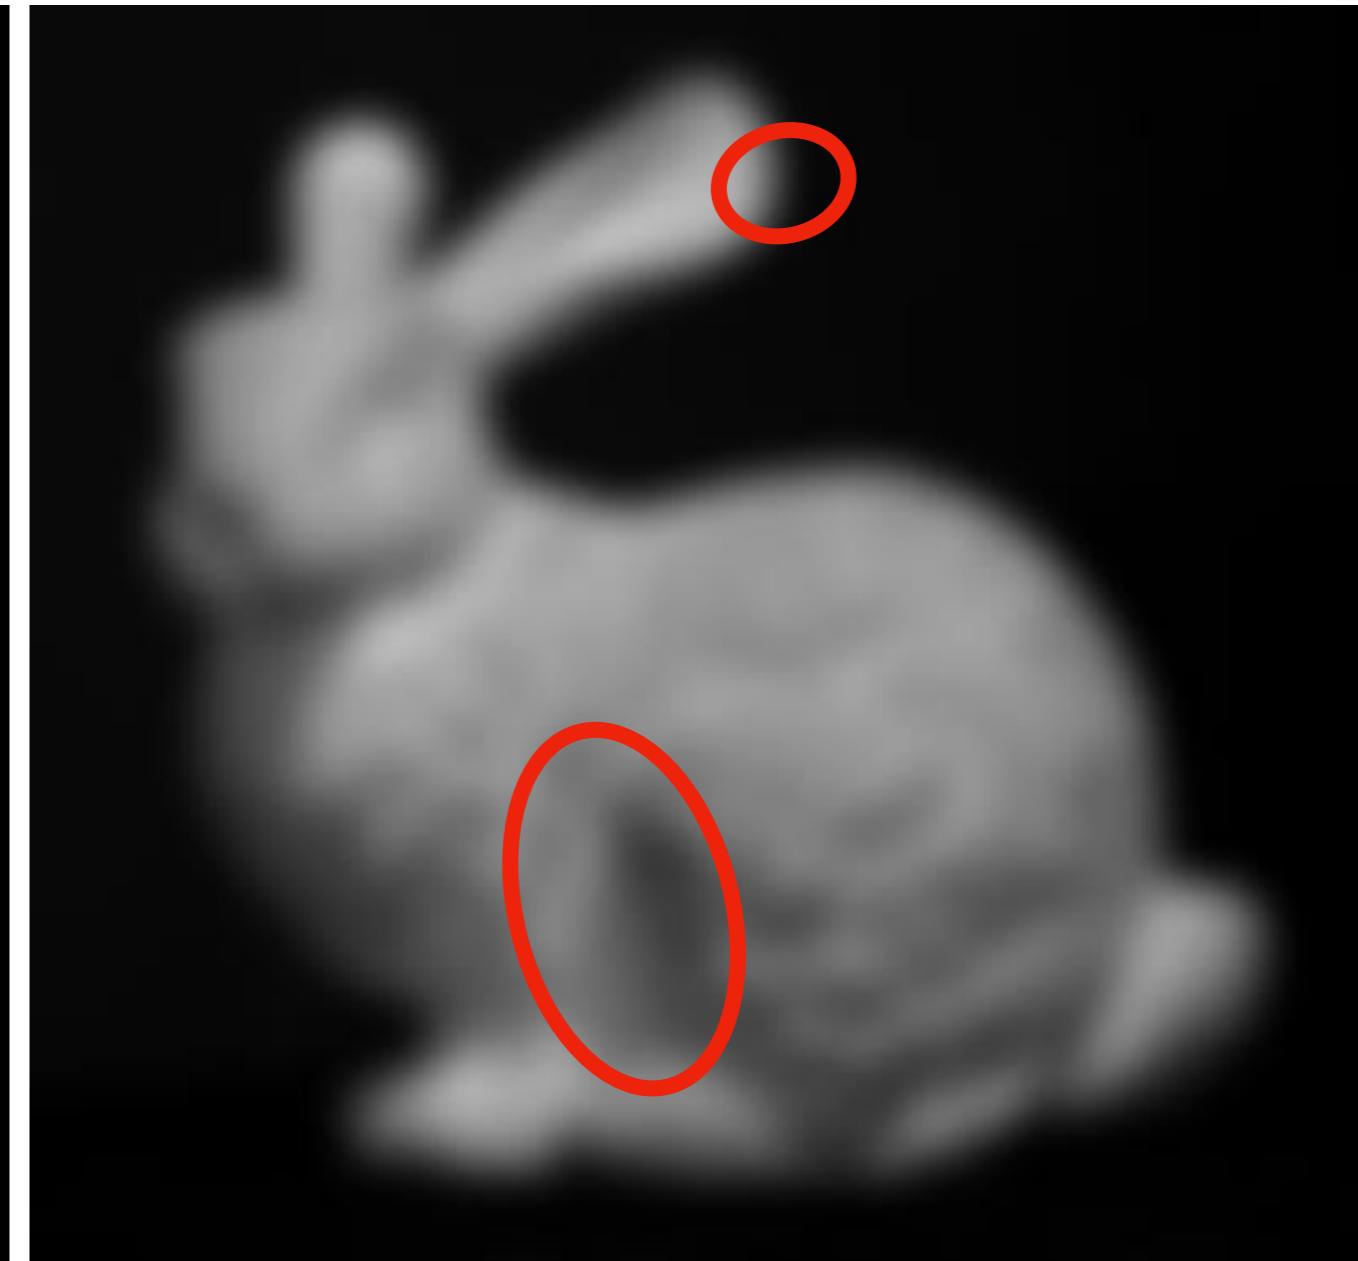


Gaussian filter

Gaussian Filter



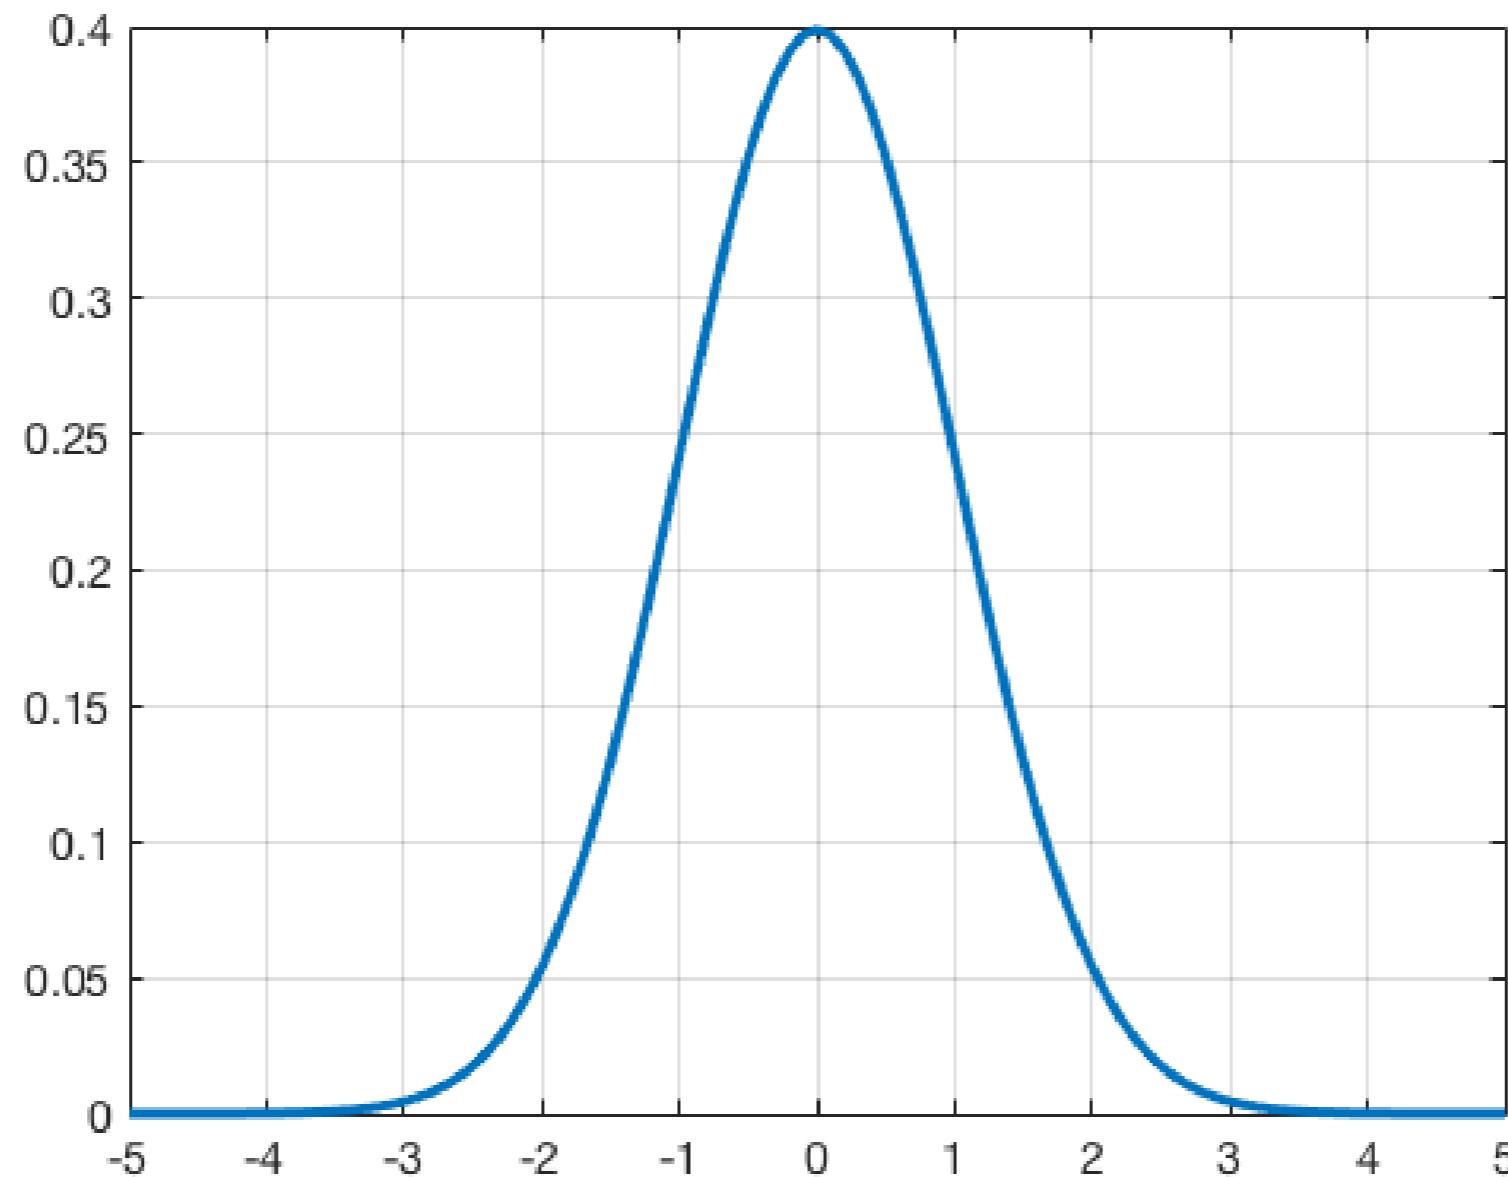
Average filter



Gaussian filter

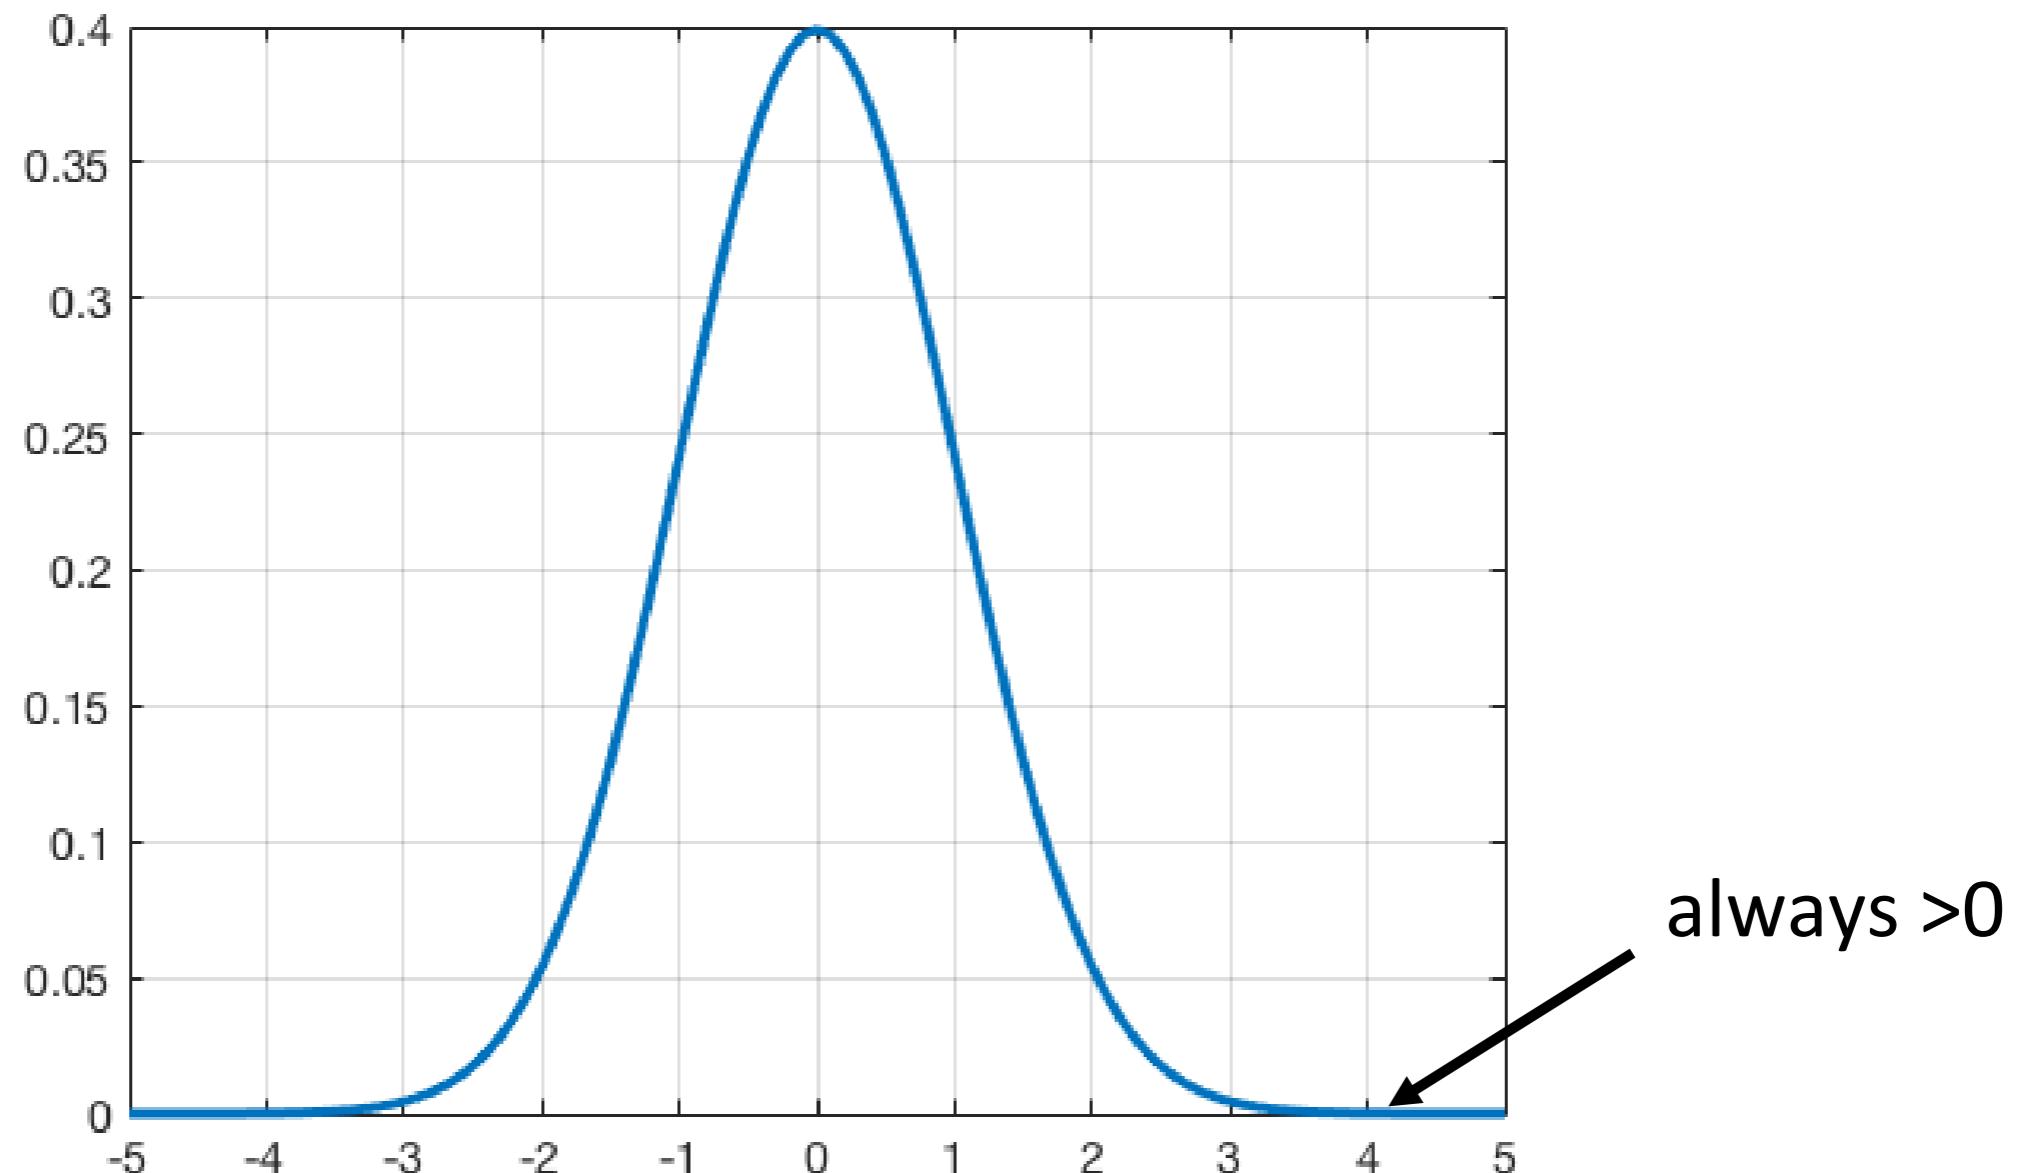
Gaussian Filter

Gaussian function has infinite support



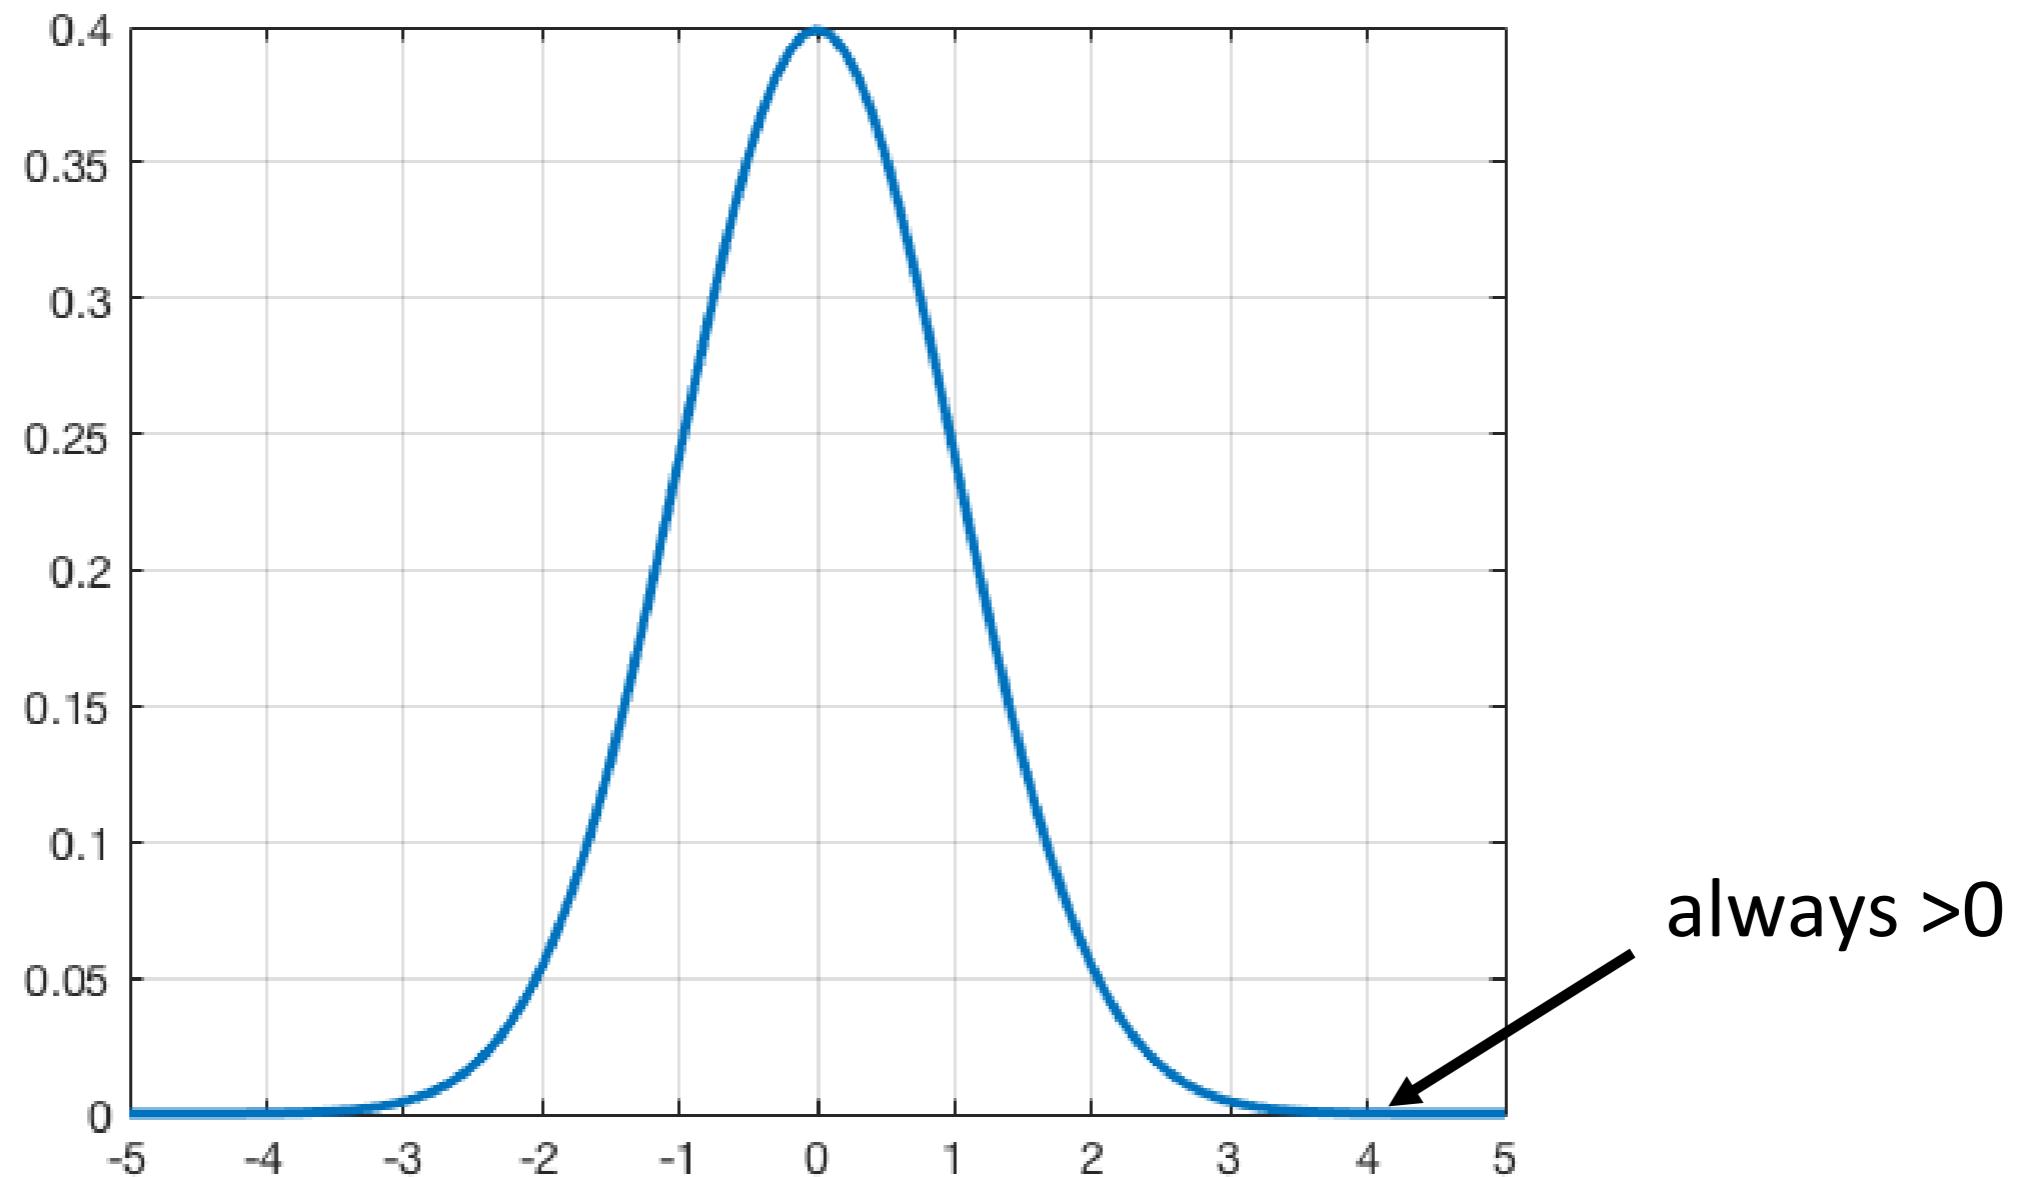
Gaussian Filter

Gaussian function has infinite support



Gaussian Filter

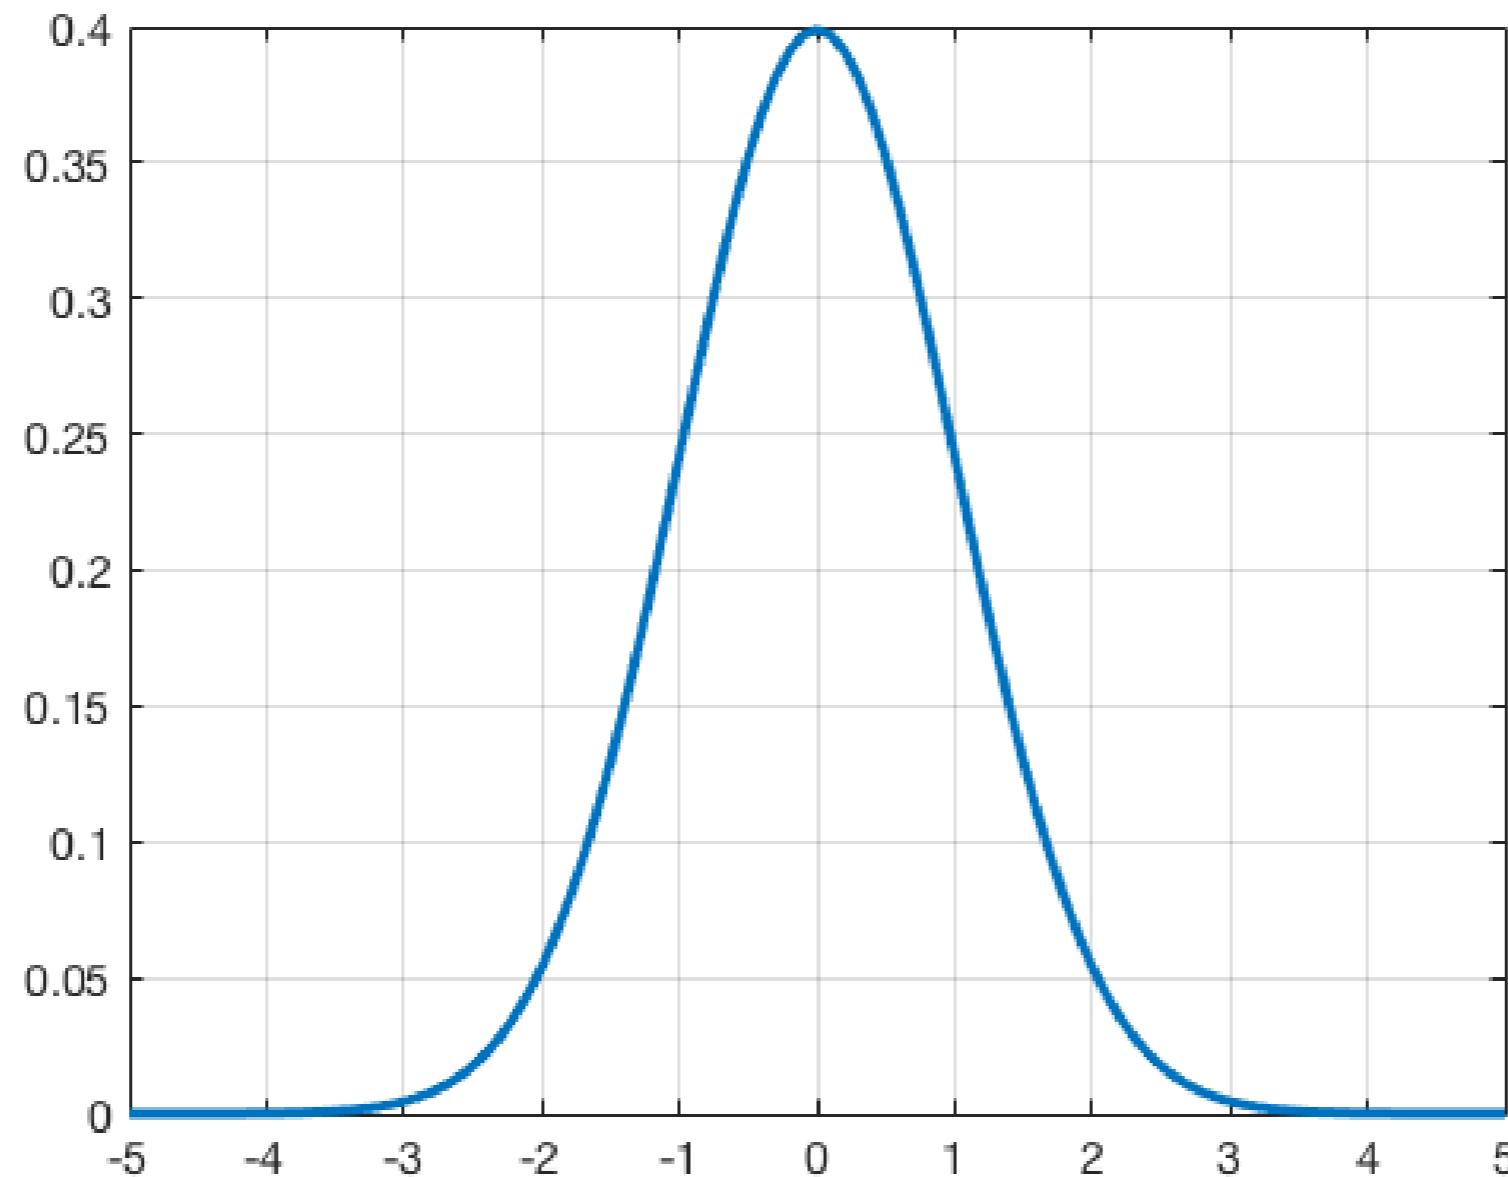
Gaussian function has infinite support



How to obtain a finite filter?

Gaussian Filter

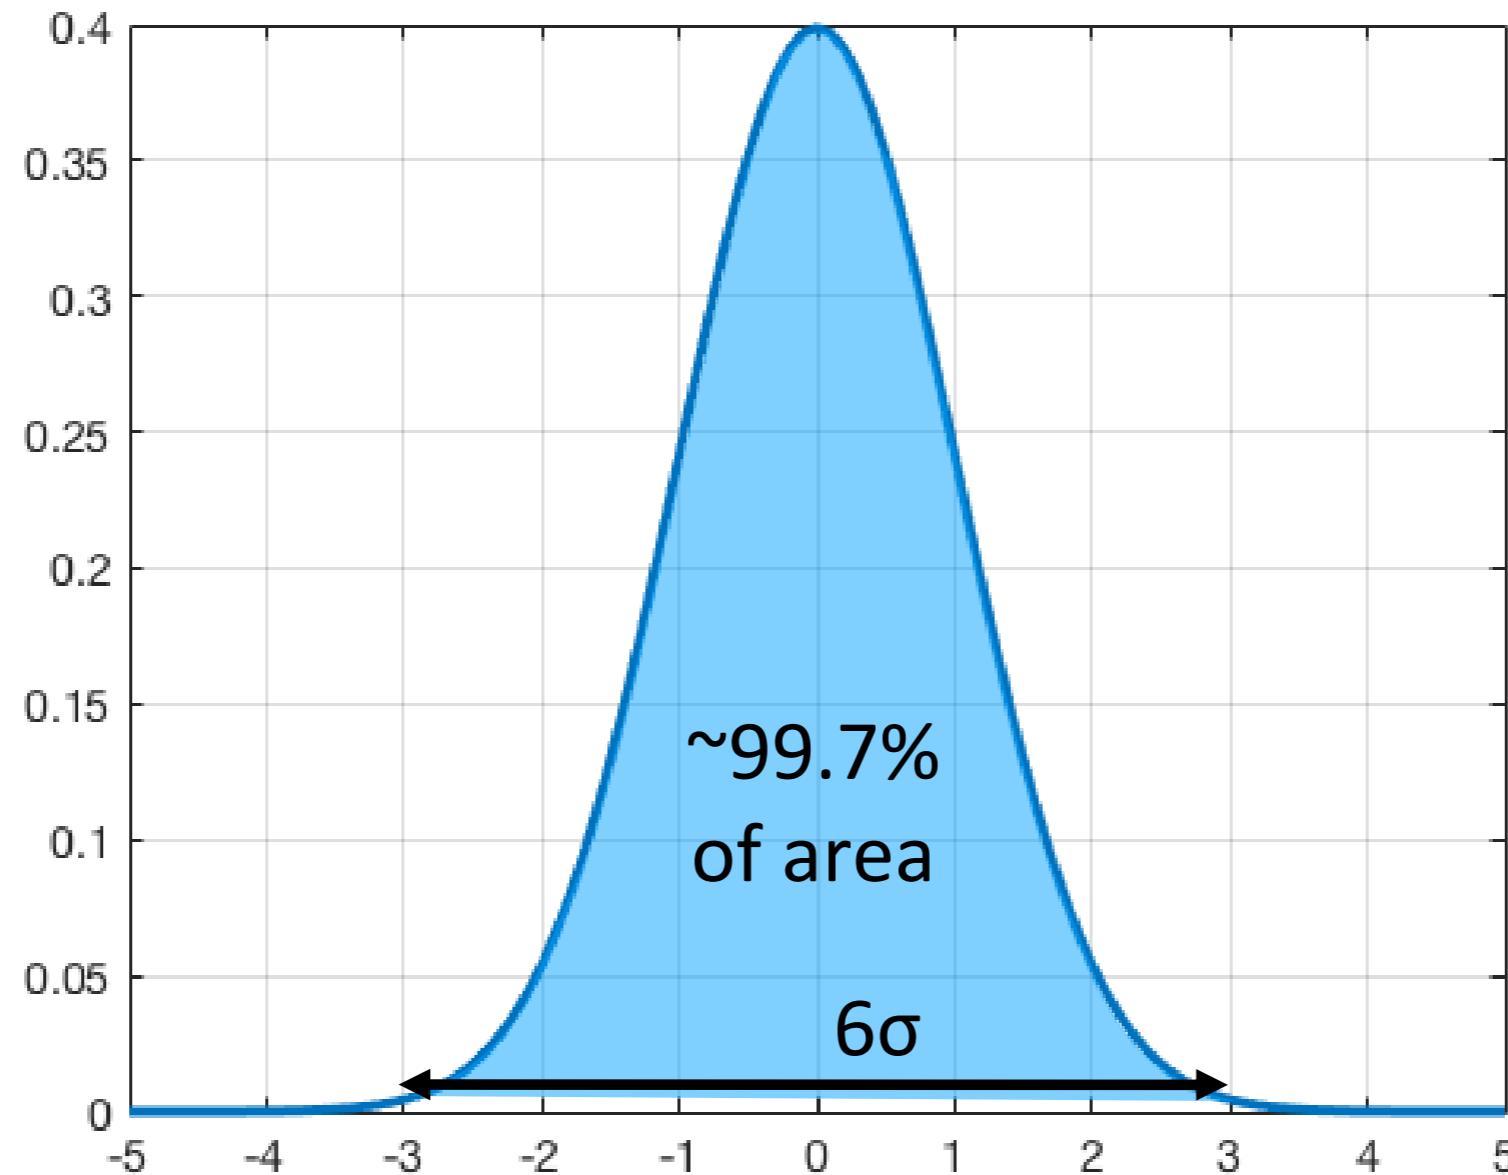
Gaussian function has infinite support



How to obtain a finite filter?

Gaussian Filter

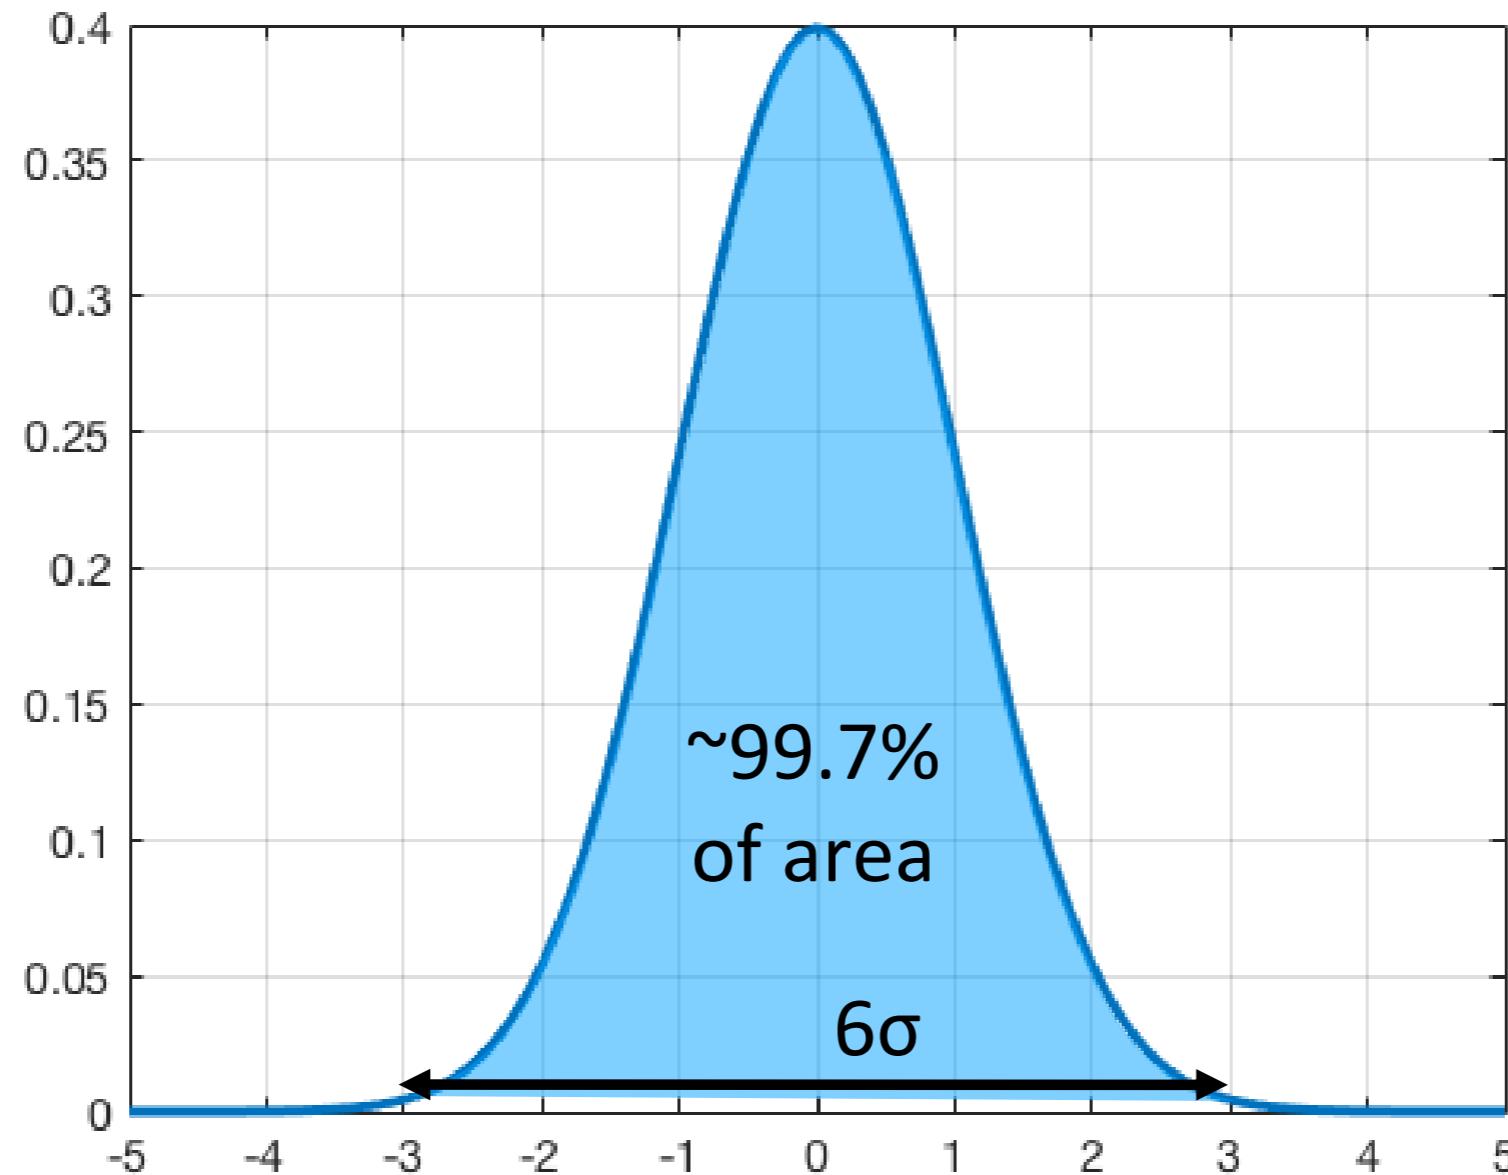
Gaussian function has infinite support



How to obtain a finite filter?

Gaussian Filter

Gaussian function has infinite support



Use filter size $\sim 6\sigma$

Gaussian Filter

Discrete Example for $\sigma=1$

1	4	6	4	1
4	16	24	16	4
6	24	36	24	6
4	16	24	16	4
1	4	6	4	1

$\frac{1}{256}$

Scaled such that smallest value of $G_\sigma(x,y)$ corresponds to 1

Gaussian Filter

Gaussian filters are **separable**

1	4	6	4	1
4	16	24	16	4
6	24	36	24	6
4	16	24	16	4
1	4	6	4	1

$$\begin{matrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{matrix} = \begin{matrix} 1 \\ 4 \\ 6 \\ 4 \\ 1 \end{matrix} * \begin{matrix} 1 & 4 & 6 & 4 & 1 \end{matrix}$$

2D Gaussian

1D Gaussian
y-direction

1D Gaussian
x-direction

Gaussian Filter

Gaussian filters are **separable**

$$\boxed{\text{image}} \star G_\sigma(x, y) = \boxed{\text{image}} \star G_\sigma(y) \star G_\sigma(x)$$

Gaussian Filter

Gaussian filters are **separable**

$$h \begin{array}{|c|} \hline \text{image} \\ \hline w \end{array} * G_\sigma(x, y) = \begin{array}{|c|} \hline \text{image} \\ \hline k \times k \end{array} * G_\sigma(y) * G_\sigma(x)$$

Gaussian Filter

Gaussian filters are **separable**

$$h \begin{array}{|c|} \hline \text{image} \\ \hline w \end{array} * G_\sigma(x, y) = \begin{array}{|c|} \hline \text{image} \\ \hline k \times k \end{array} * G_\sigma(y) * G_\sigma(x)$$

w * h * k * k operations

Gaussian Filter

Gaussian filters are **separable**

$$h \begin{matrix} \text{image} \\ \boxed{} \\ w \end{matrix} * G_\sigma(x, y) = \begin{matrix} \text{image} \\ \boxed{} \\ k \times k \end{matrix} * G_\sigma(y) * G_\sigma(x)$$
$$k \times 1 \quad 1 \times k$$

w * h * k * k operations

Gaussian Filter

Gaussian filters are **separable**

$$h \begin{matrix} \text{image} \\ \boxed{} \\ w \end{matrix} * G_\sigma(x, y) = \begin{matrix} \text{image} \\ \boxed{} \\ k \times k \end{matrix} * G_\sigma(y) * G_\sigma(x)$$
$$\qquad\qquad\qquad \begin{matrix} k \times 1 \\ \boxed{} \end{matrix} \qquad \begin{matrix} 1 \times k \\ \boxed{} \end{matrix}$$

w * h * k * k operations

2 * w * h * k operations

Gaussian Filter

Gaussian filters are **separable**

$$h \begin{matrix} \text{image} \\ \boxed{} \\ w \end{matrix} * G_\sigma(x, y) = \begin{matrix} \text{image} \\ \boxed{} \\ k \times k \end{matrix} * G_\sigma(y) * G_\sigma(x) \quad k \times 1 \quad 1 \times k$$

$w * h * k * k$ operations > $2 * w * h * k$ operations

Gaussian Filter

Gaussian filters are **separable**

$$h \begin{matrix} \text{image} \\ \boxed{} \\ w \end{matrix} * G_\sigma(x, y) = \begin{matrix} \text{image} \\ \boxed{} \\ k \times k \end{matrix} * G_\sigma(y) * G_\sigma(x)$$
$$k \times 1 \quad 1 \times k$$

$$w * h * k * k \text{ operations} > 2 * w * h * k \text{ operations}$$

Separable filter can be applied efficiently

Gaussian Filter

Iterative application

$$I \star G_{\sigma_1}(x, y) \star G_{\sigma_2}(x, y) = I \star G_{\sqrt{\sigma_1^2 + \sigma_2^2}}(x, y)$$

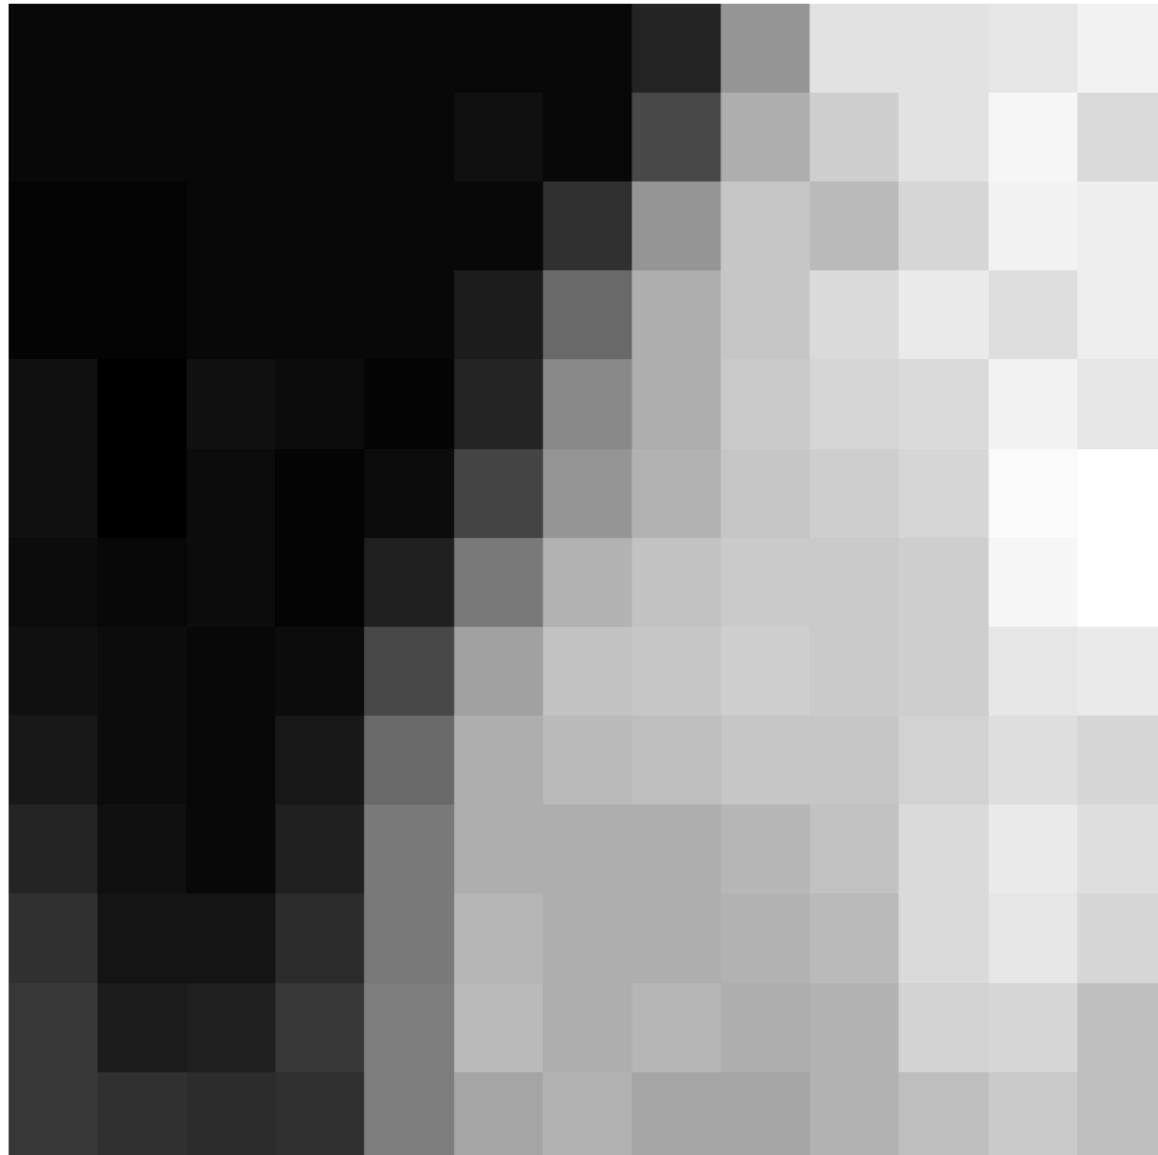
Gaussian Filter

Iterative application

$$I \star G_{\sigma_1}(x, y) \star G_{\sigma_2}(x, y) = I \star G_{\sqrt{\sigma_1^2 + \sigma_2^2}}(x, y)$$

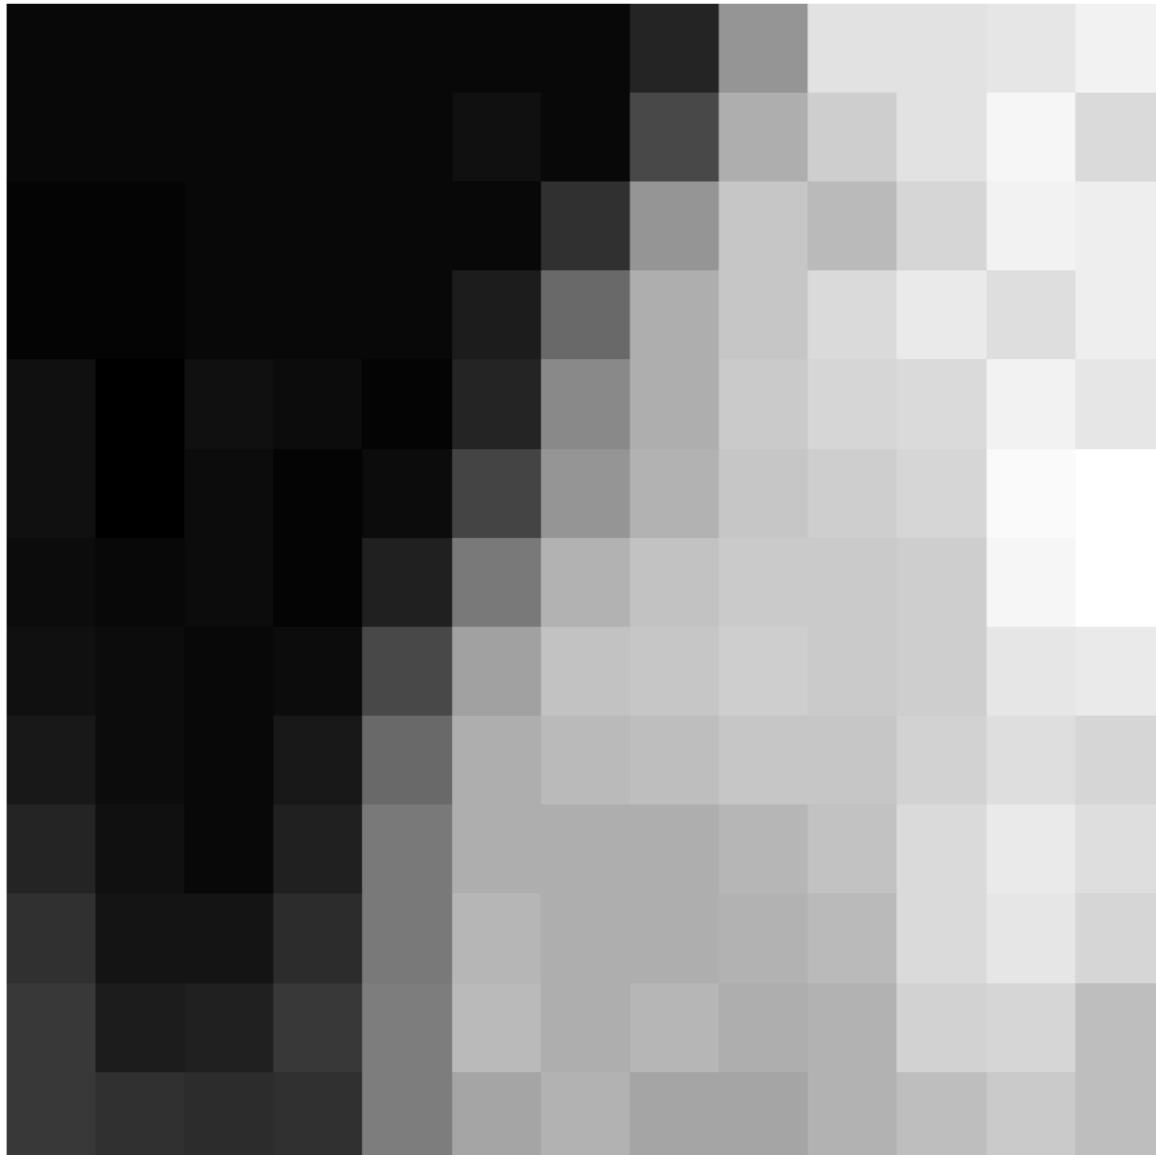
Useful if we need multiple smoothed version of an image
(more on this later)

Image Gradients



$$\nabla I(x, y) = \begin{pmatrix} I'_x \\ I'_y \end{pmatrix}$$

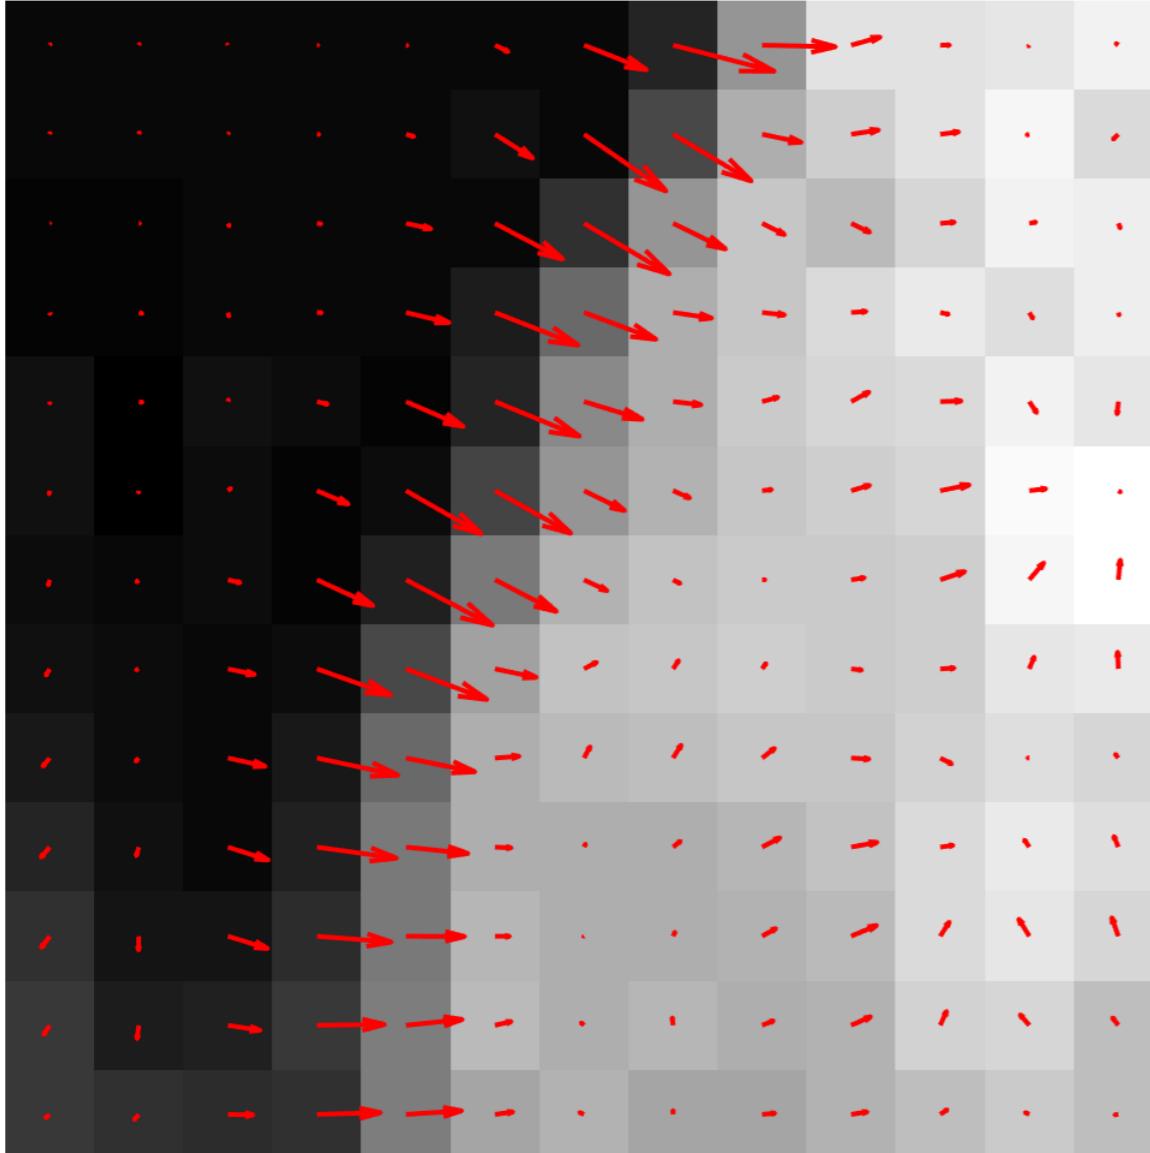
Image Gradients



$$\nabla I(x, y) = \begin{pmatrix} I'_x \\ I'_y \end{pmatrix}$$

$$I'_x(x, y) \approx \frac{I(x + 1, y) - I(x - 1, y)}{2}$$

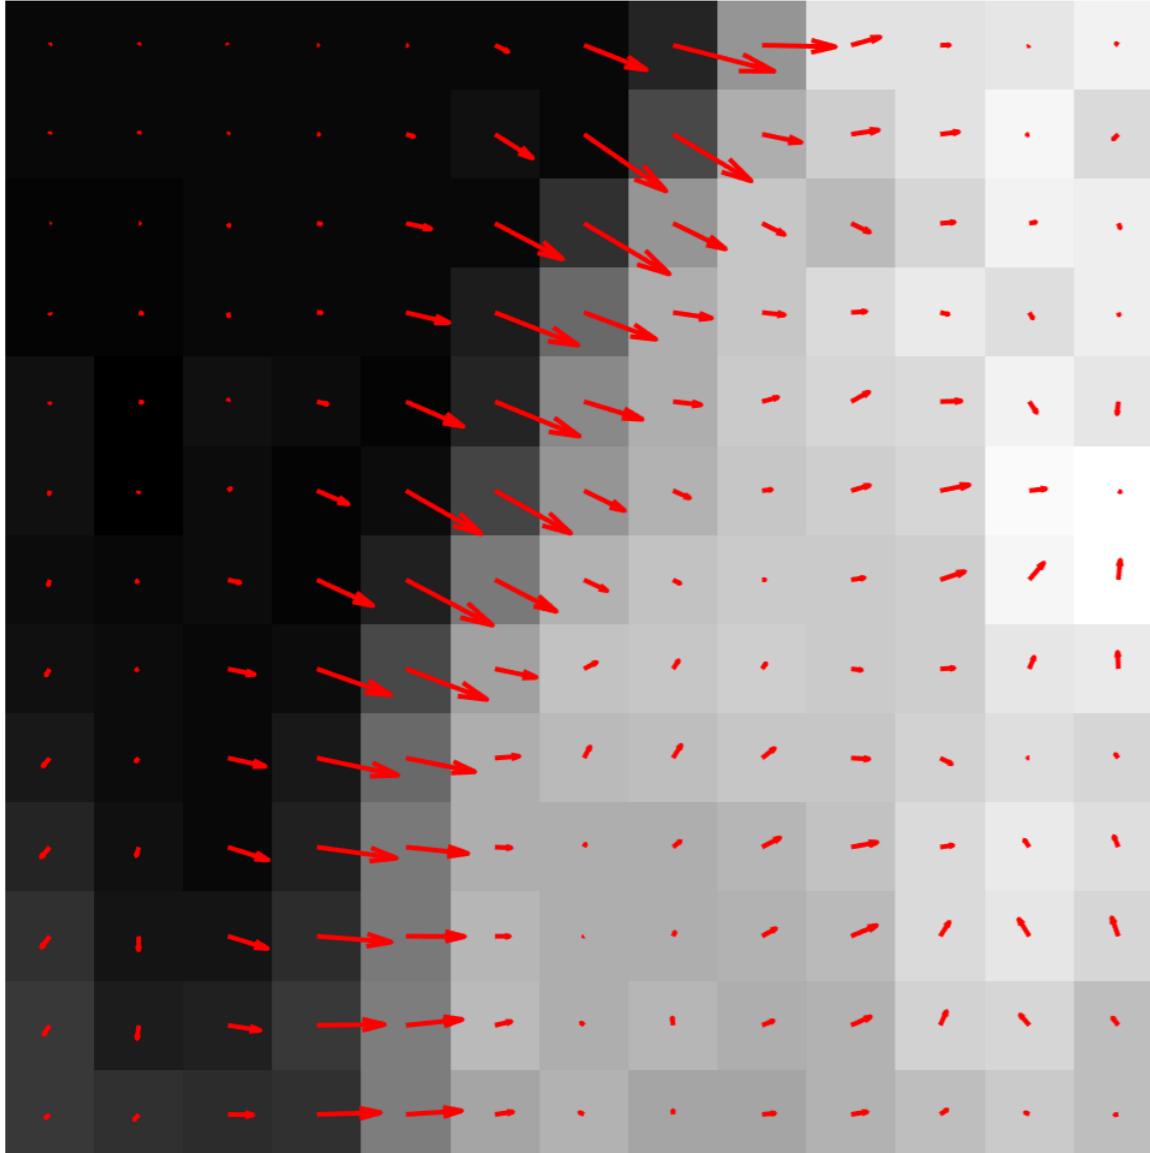
Image Gradients



$$\nabla I(x, y) = \begin{pmatrix} I'_x \\ I'_y \end{pmatrix}$$

$$I'_x(x, y) \approx \frac{I(x+1, y) - I(x-1, y)}{2}$$

Image Gradients

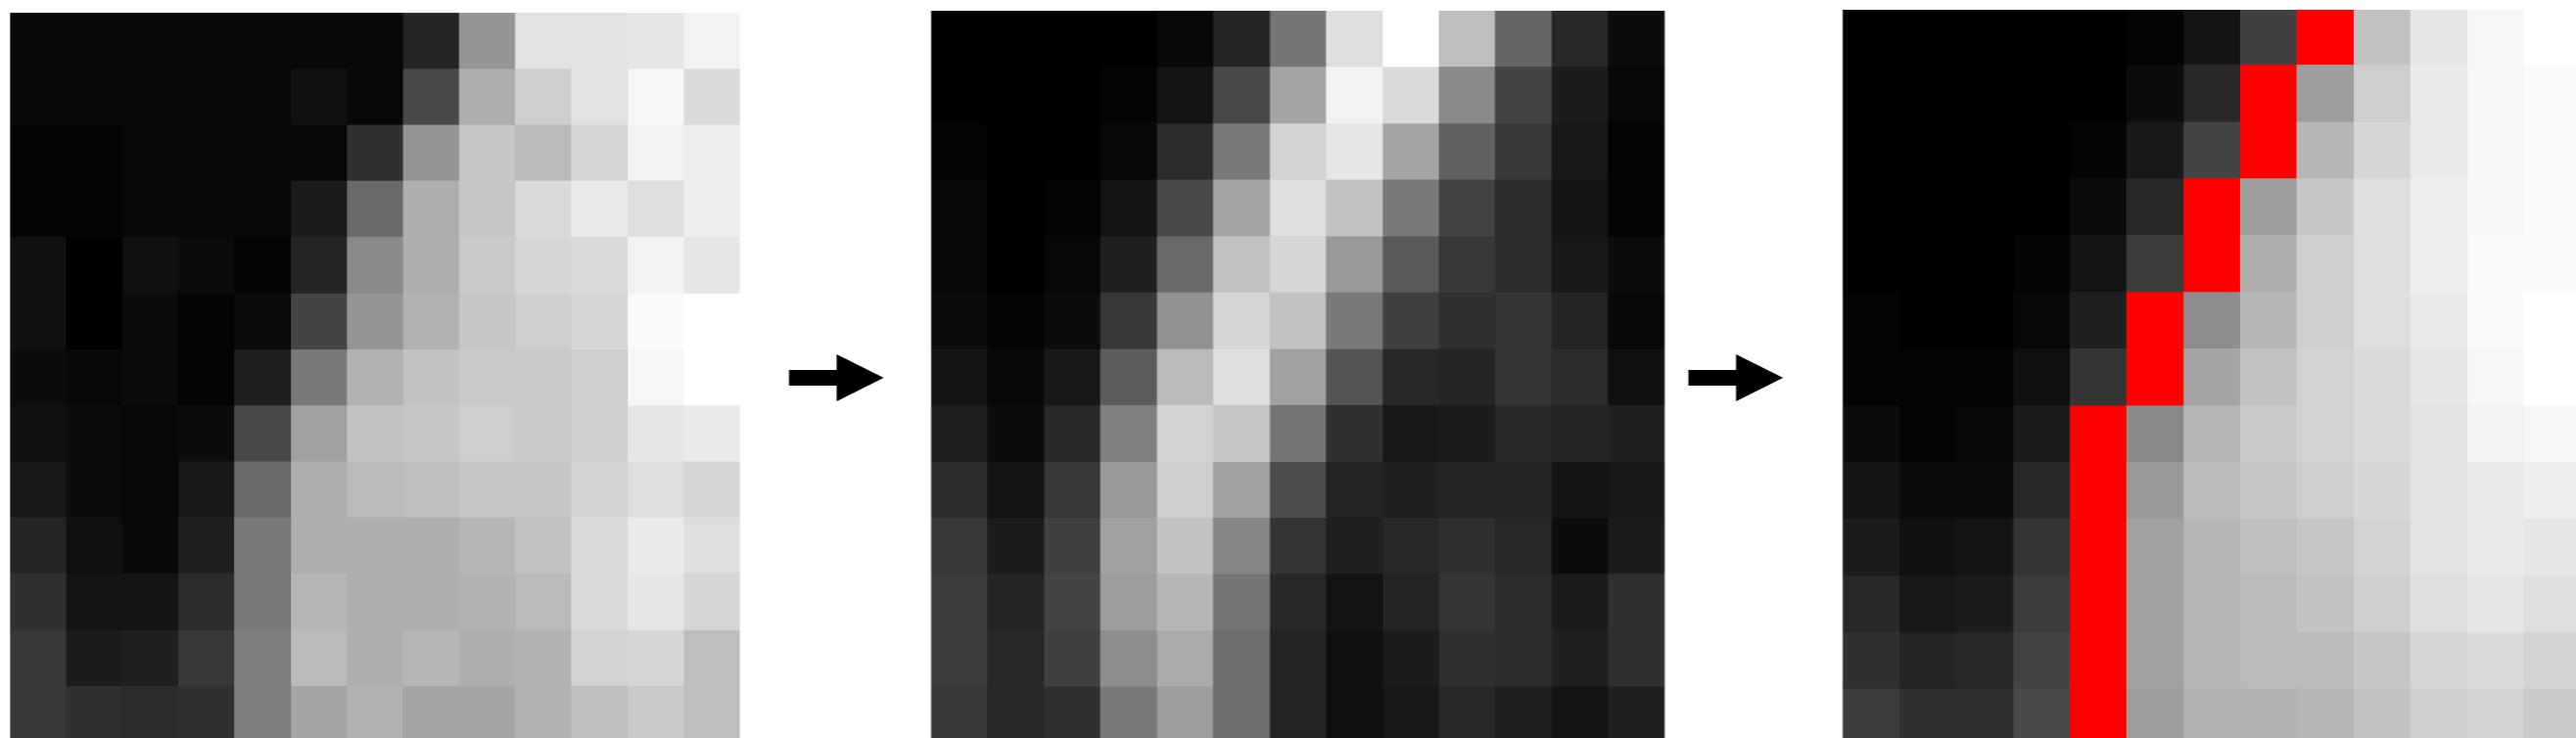


$$\nabla I(x, y) = \begin{pmatrix} I'_x \\ I'_y \end{pmatrix}$$

$$I'_x(x, y) \approx \frac{I(x+1, y) - I(x-1, y)}{2}$$

$$I'_x = I \star (-0.5 \quad 0 \quad 0.5)$$

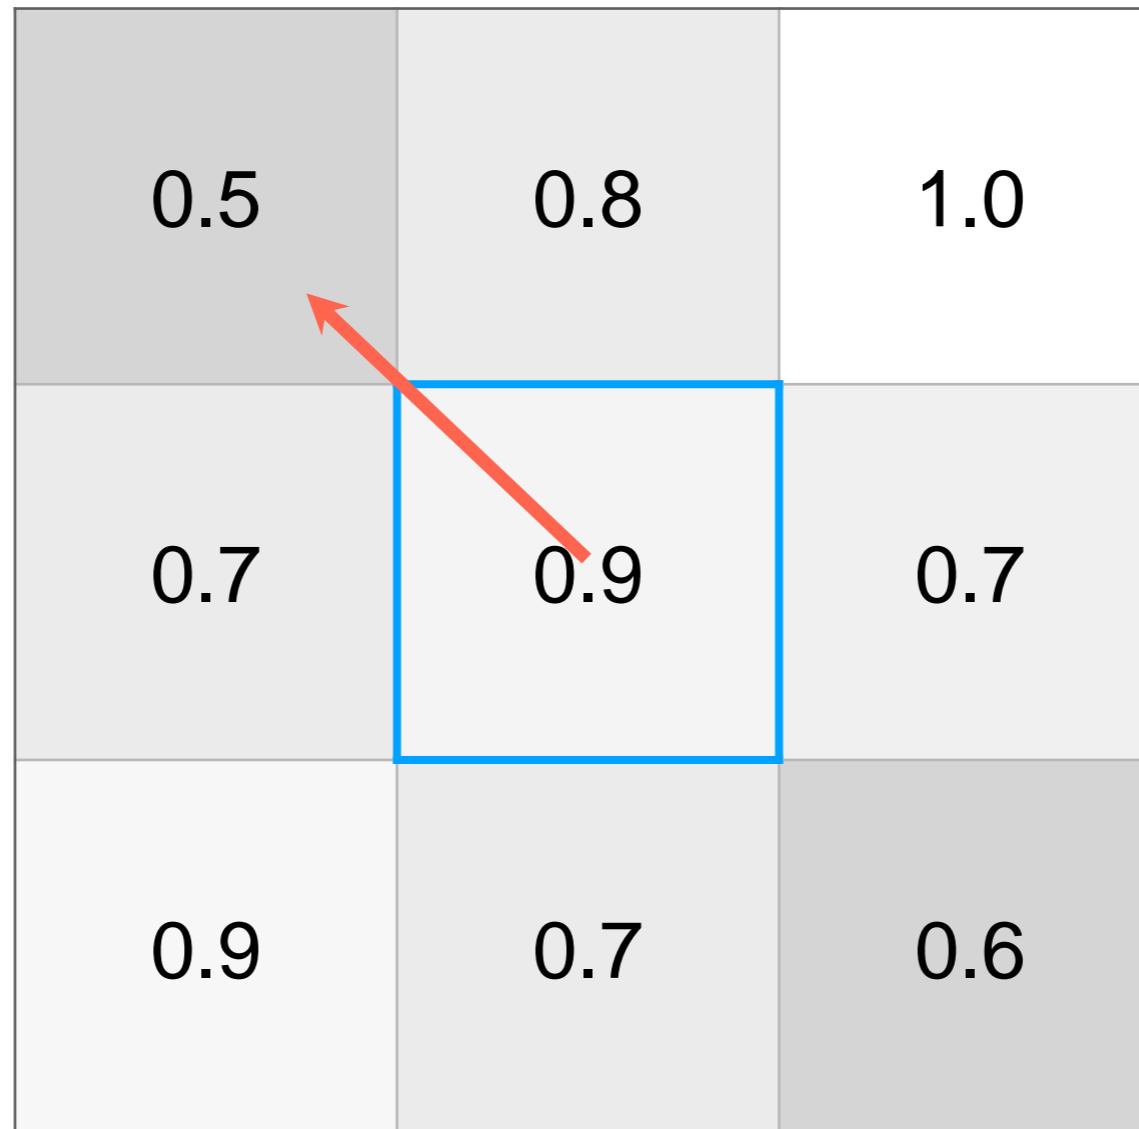
Application: Edge Detection



Magnitude of
gradients

Thresholding and
non-maximum
suppression

Non-Maximum Suppression



Non-Maximum Suppression



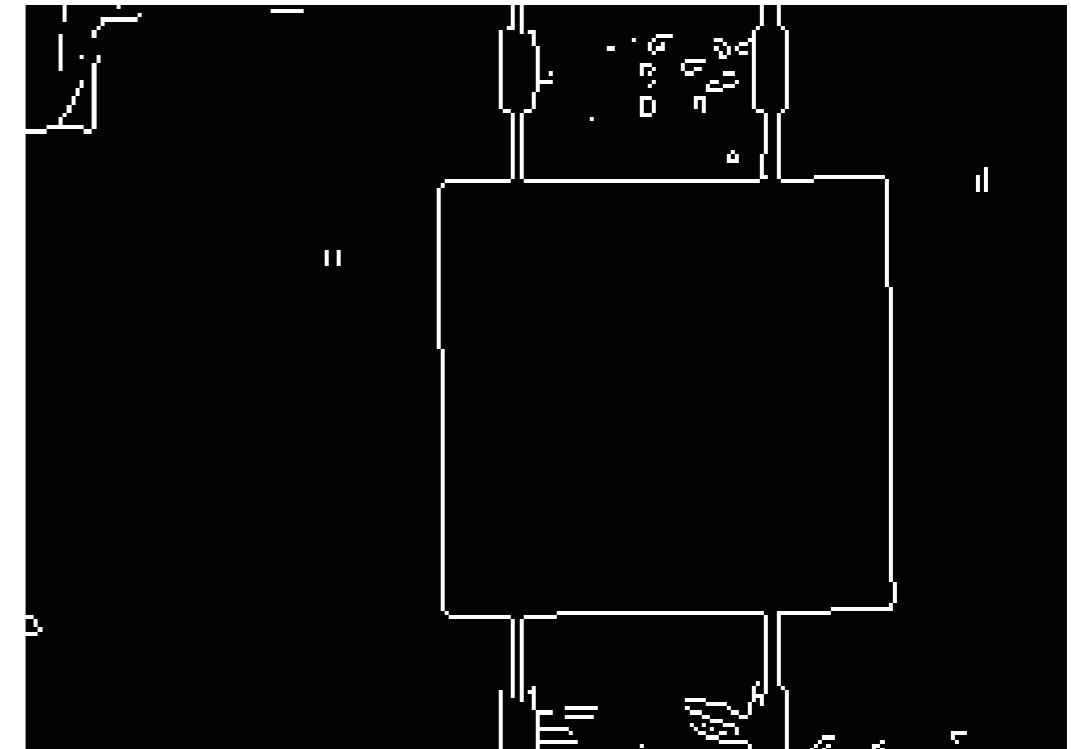
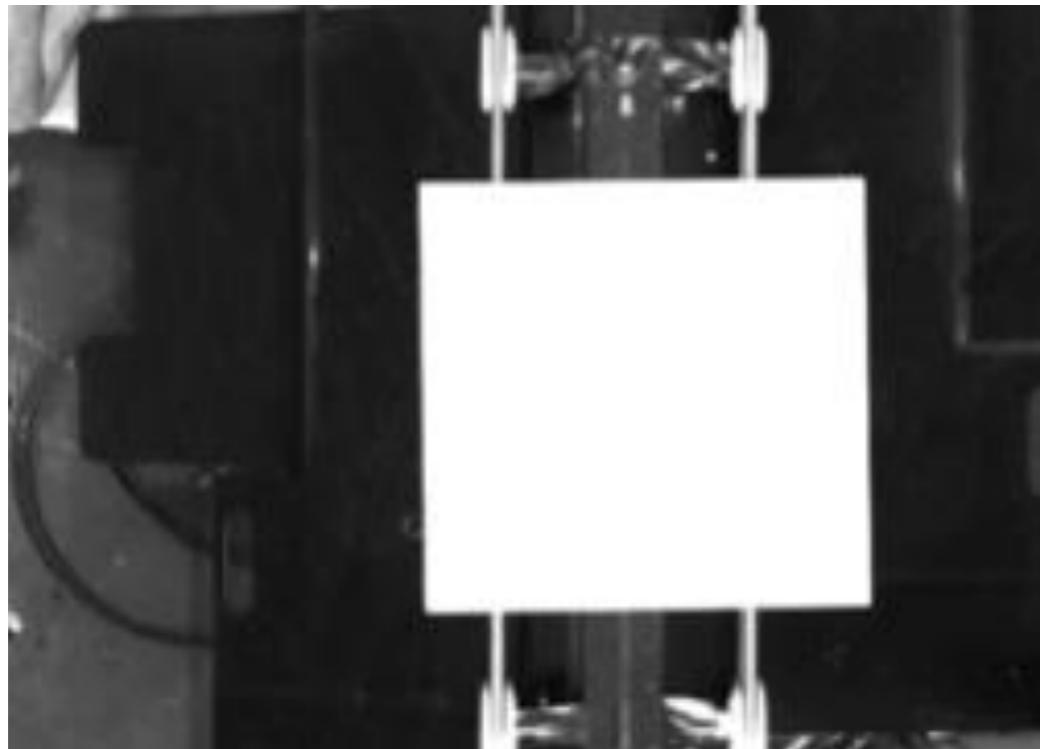
Only compare to neighbors along the gradient.

Non-Maximum Suppression



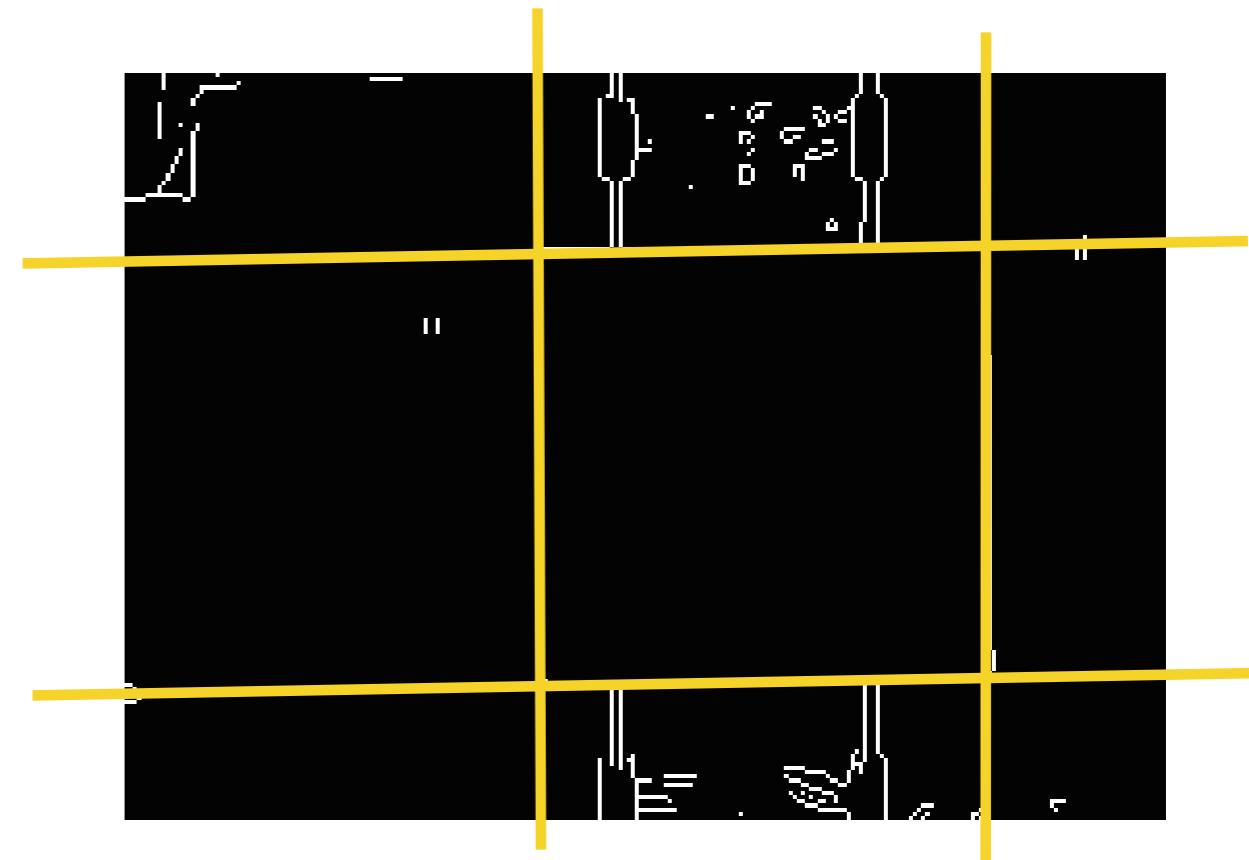
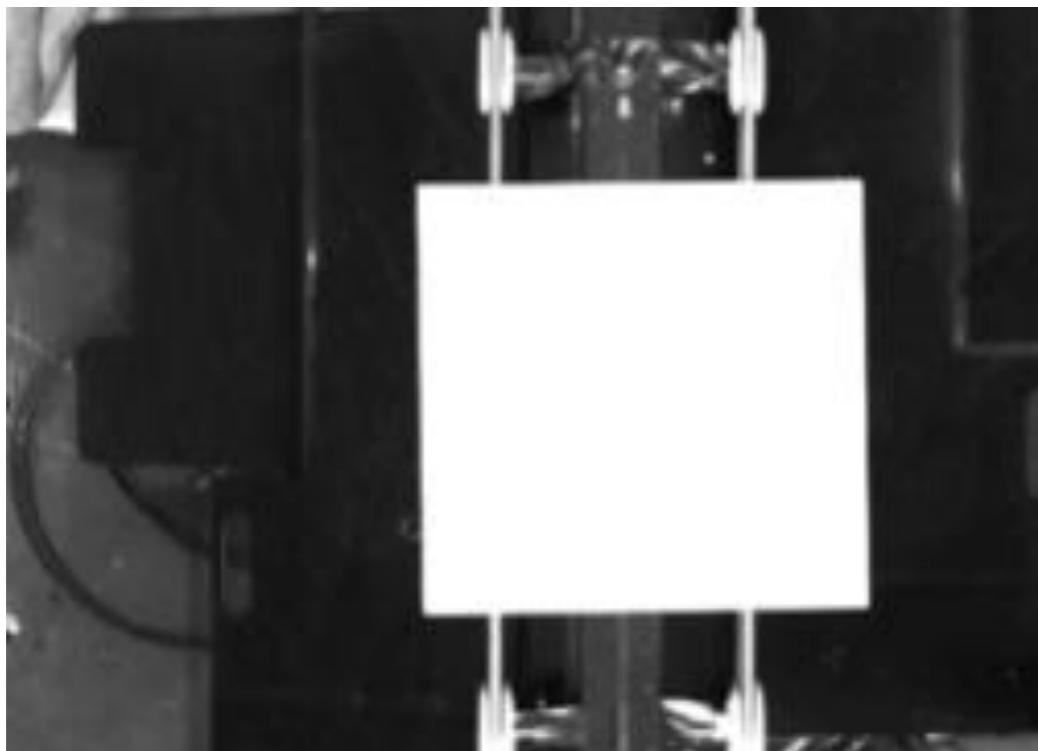
Only compare to neighbors along the gradient.
Keep if larger response than these.

Application: Edge Detection

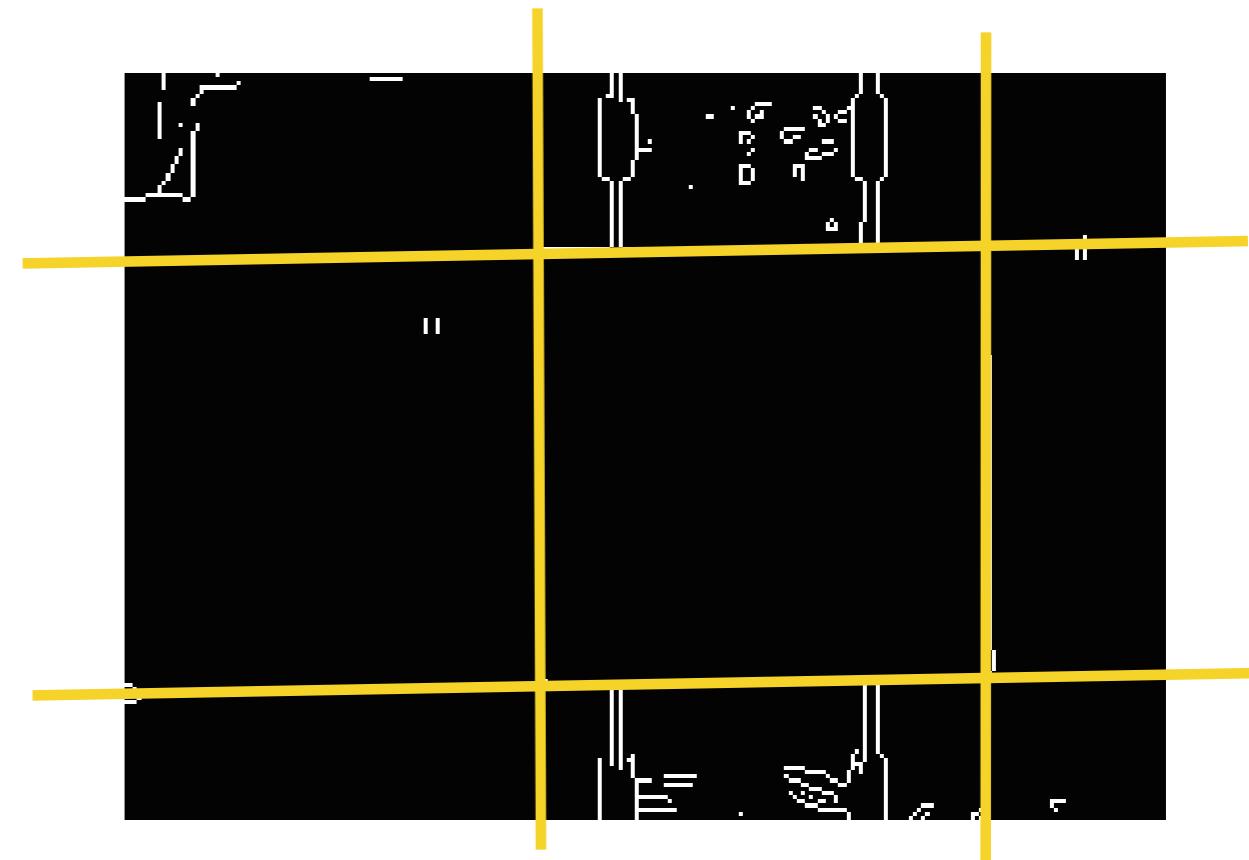
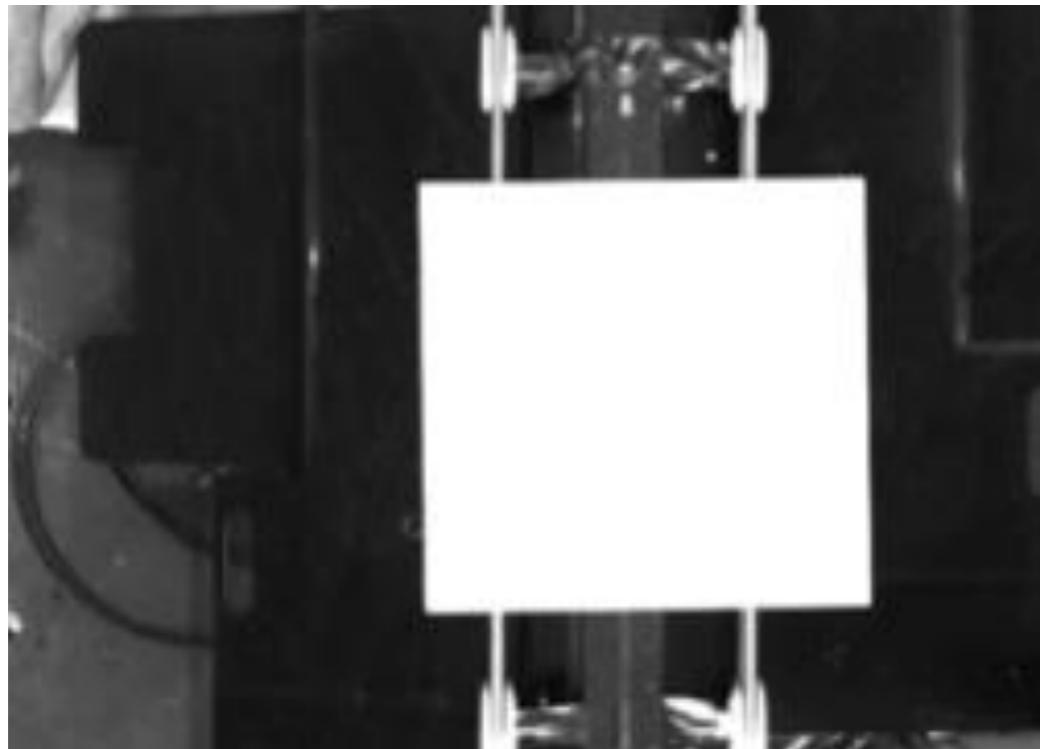


Using Matlab with default thresholds

Application: Edge Detection

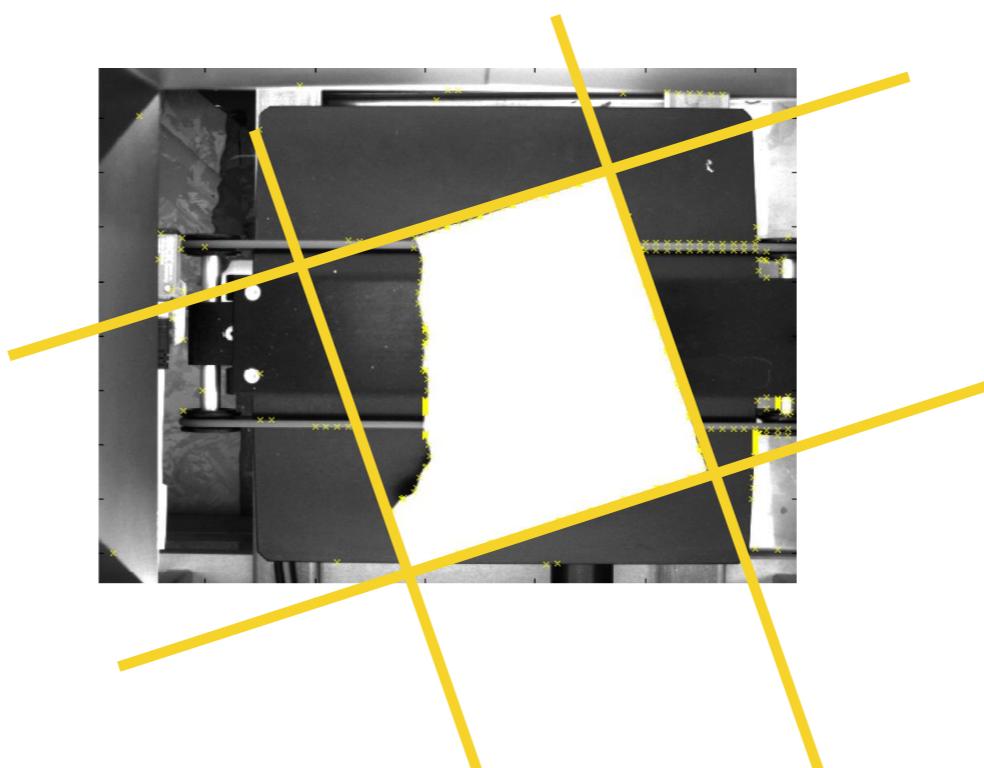
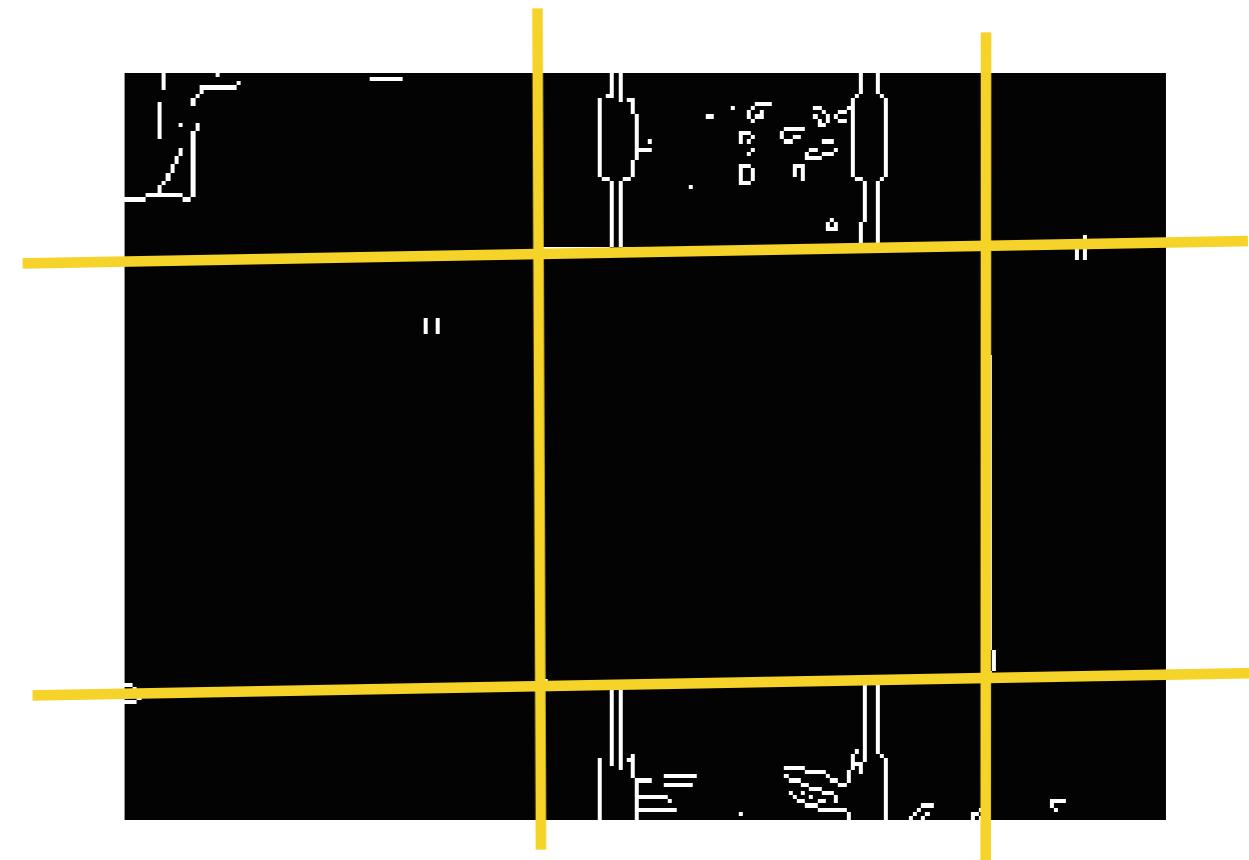
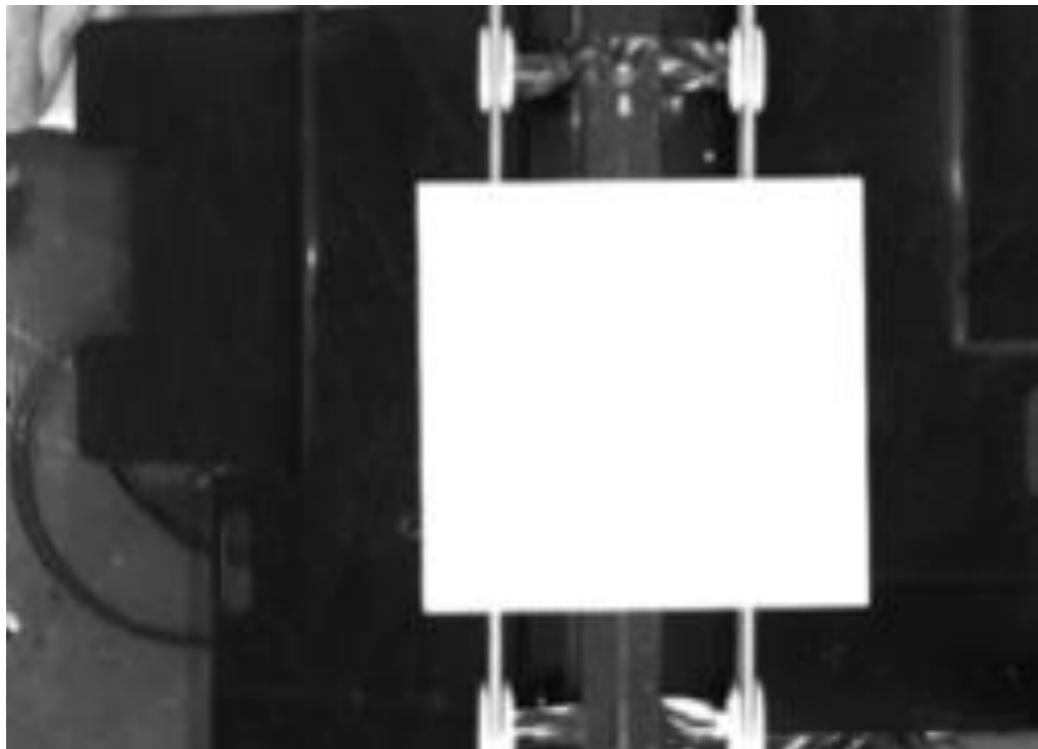


Application: Edge Detection



Fitting lines into the remaining pixels (see Lecture 7)

Application: Edge Detection



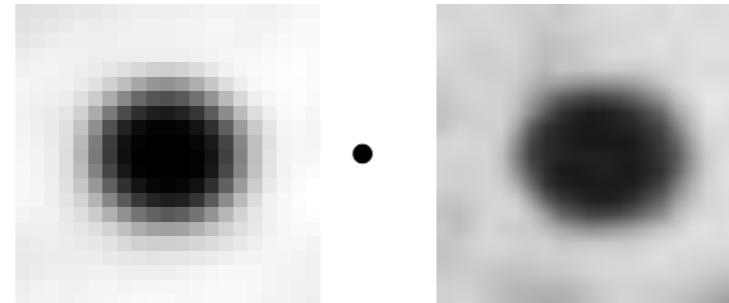
Similarity Measures

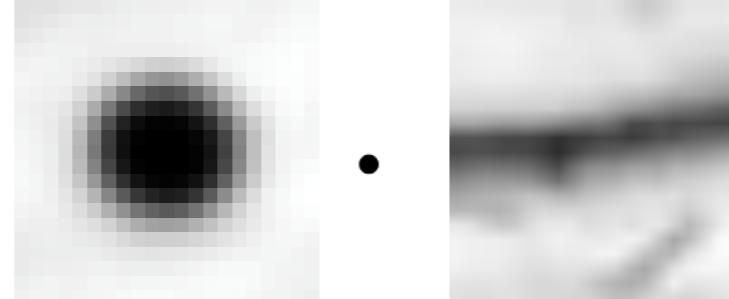
Measuring Similarity

Measuring similarity between images / patches central problem

Measuring Similarity

Measuring similarity between images / patches central problem


$$\cdot > \tau$$


$$\cdot < \tau$$

Measuring Similarity

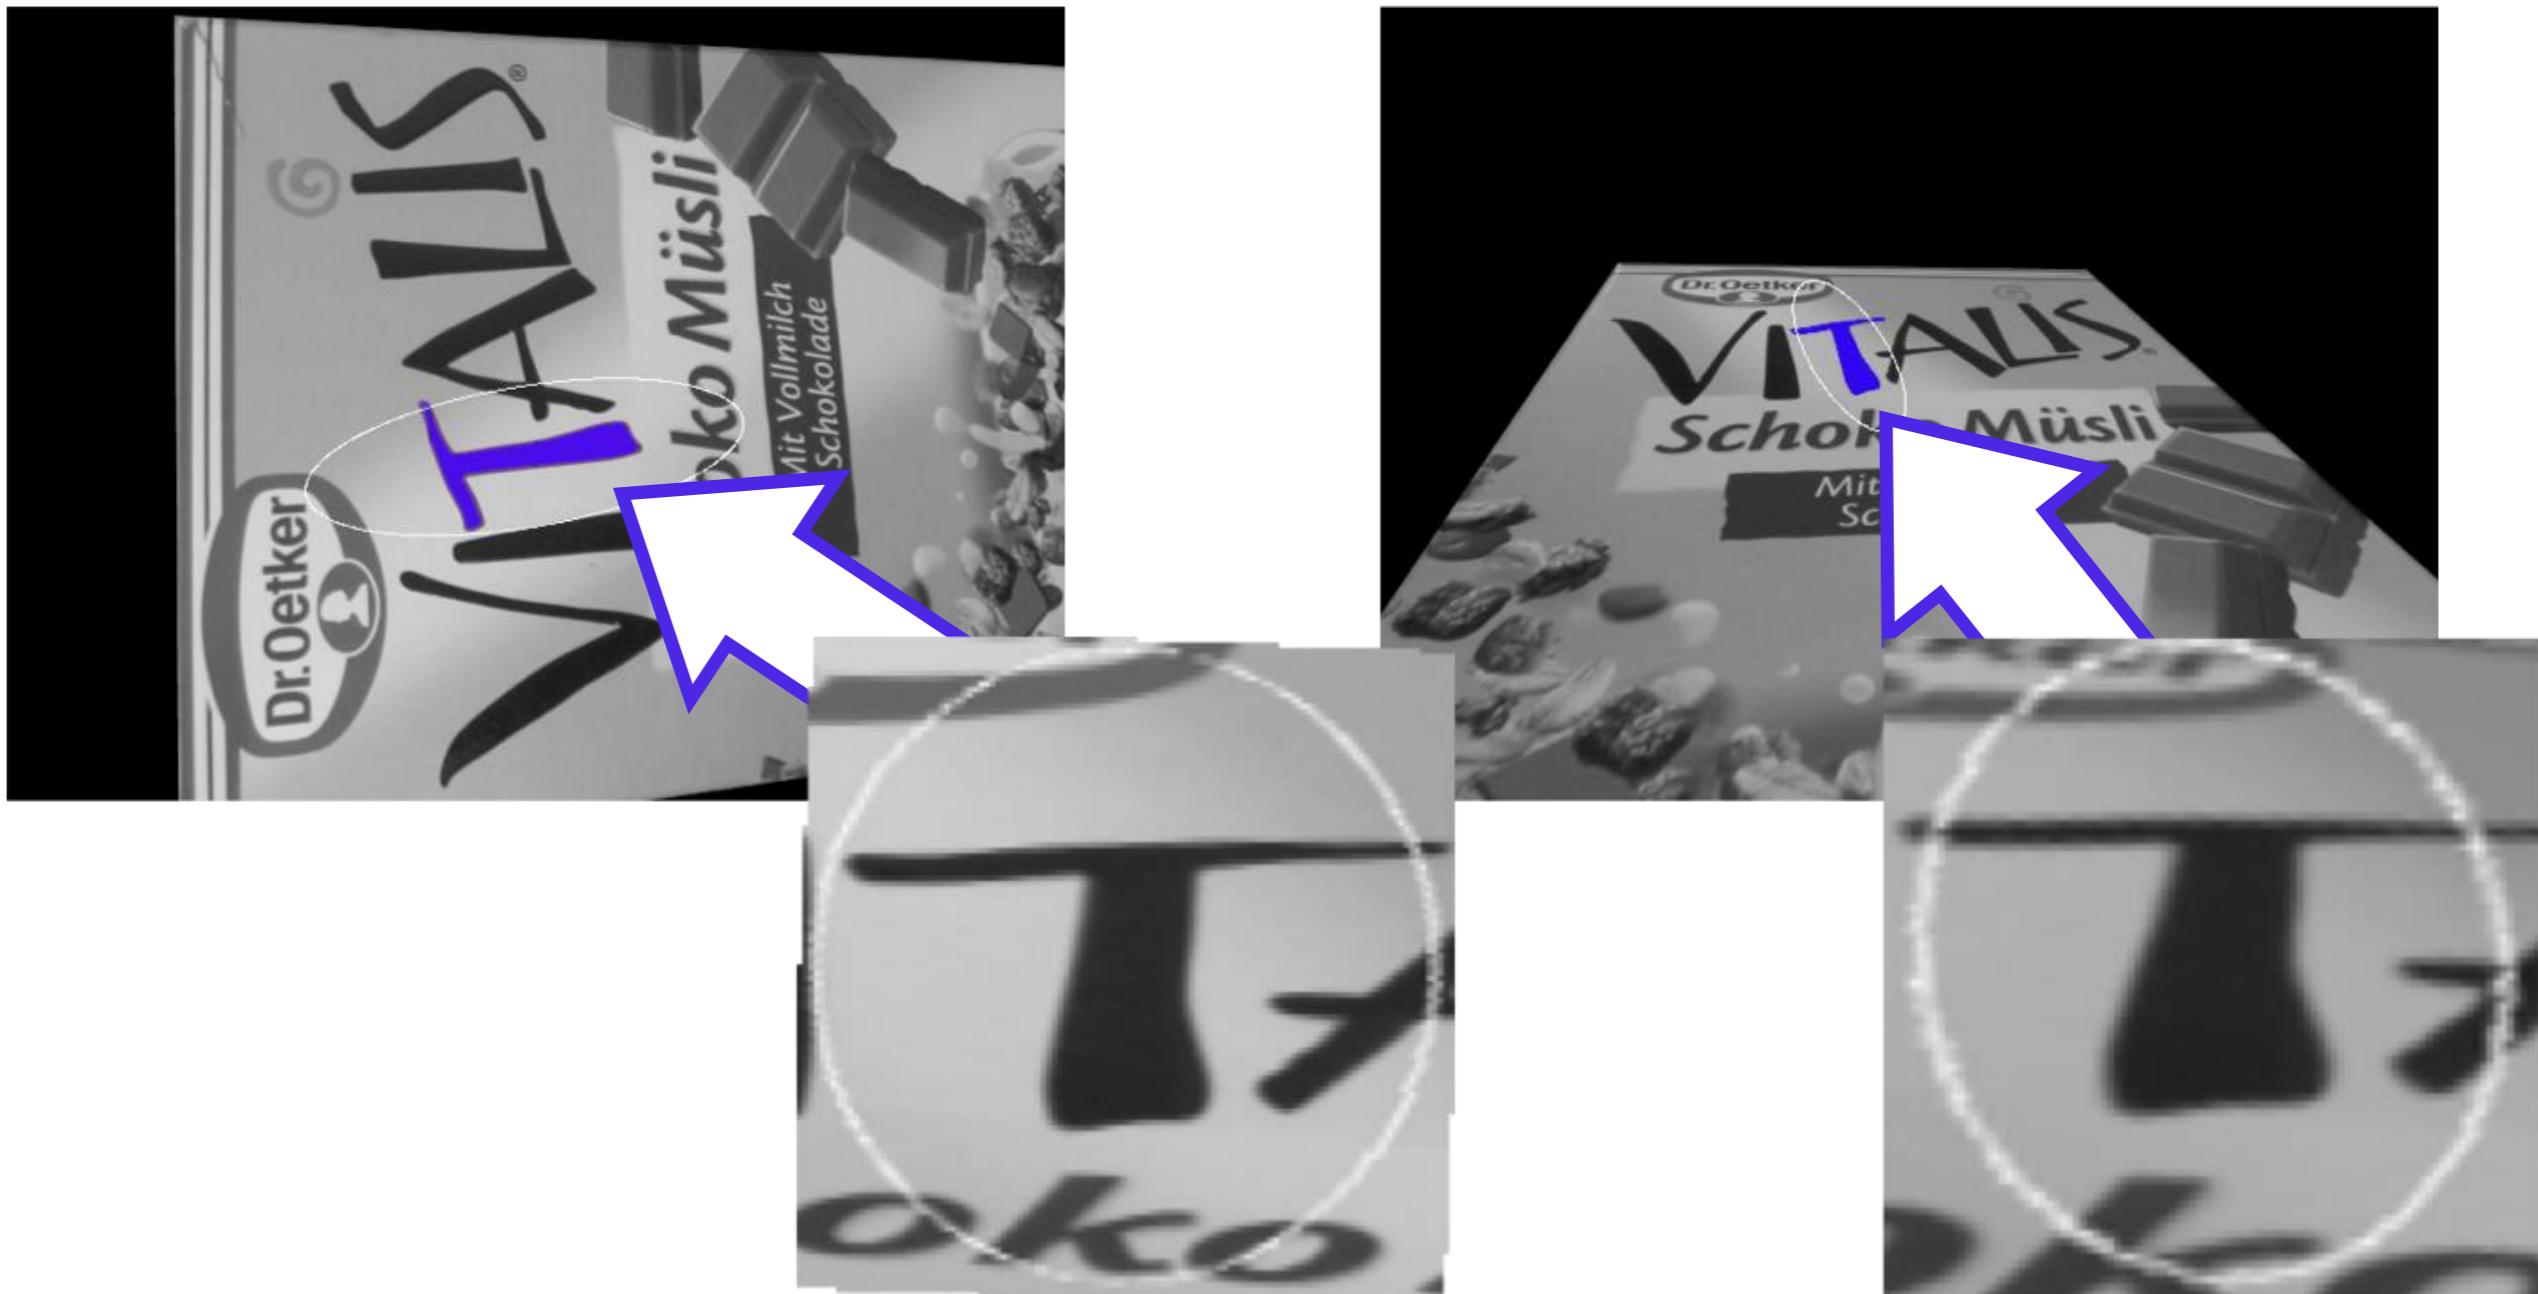
Measuring similarity between images / patches central problem



slide credit: Marc Pollefeys, Kevin Köser

Measuring Similarity

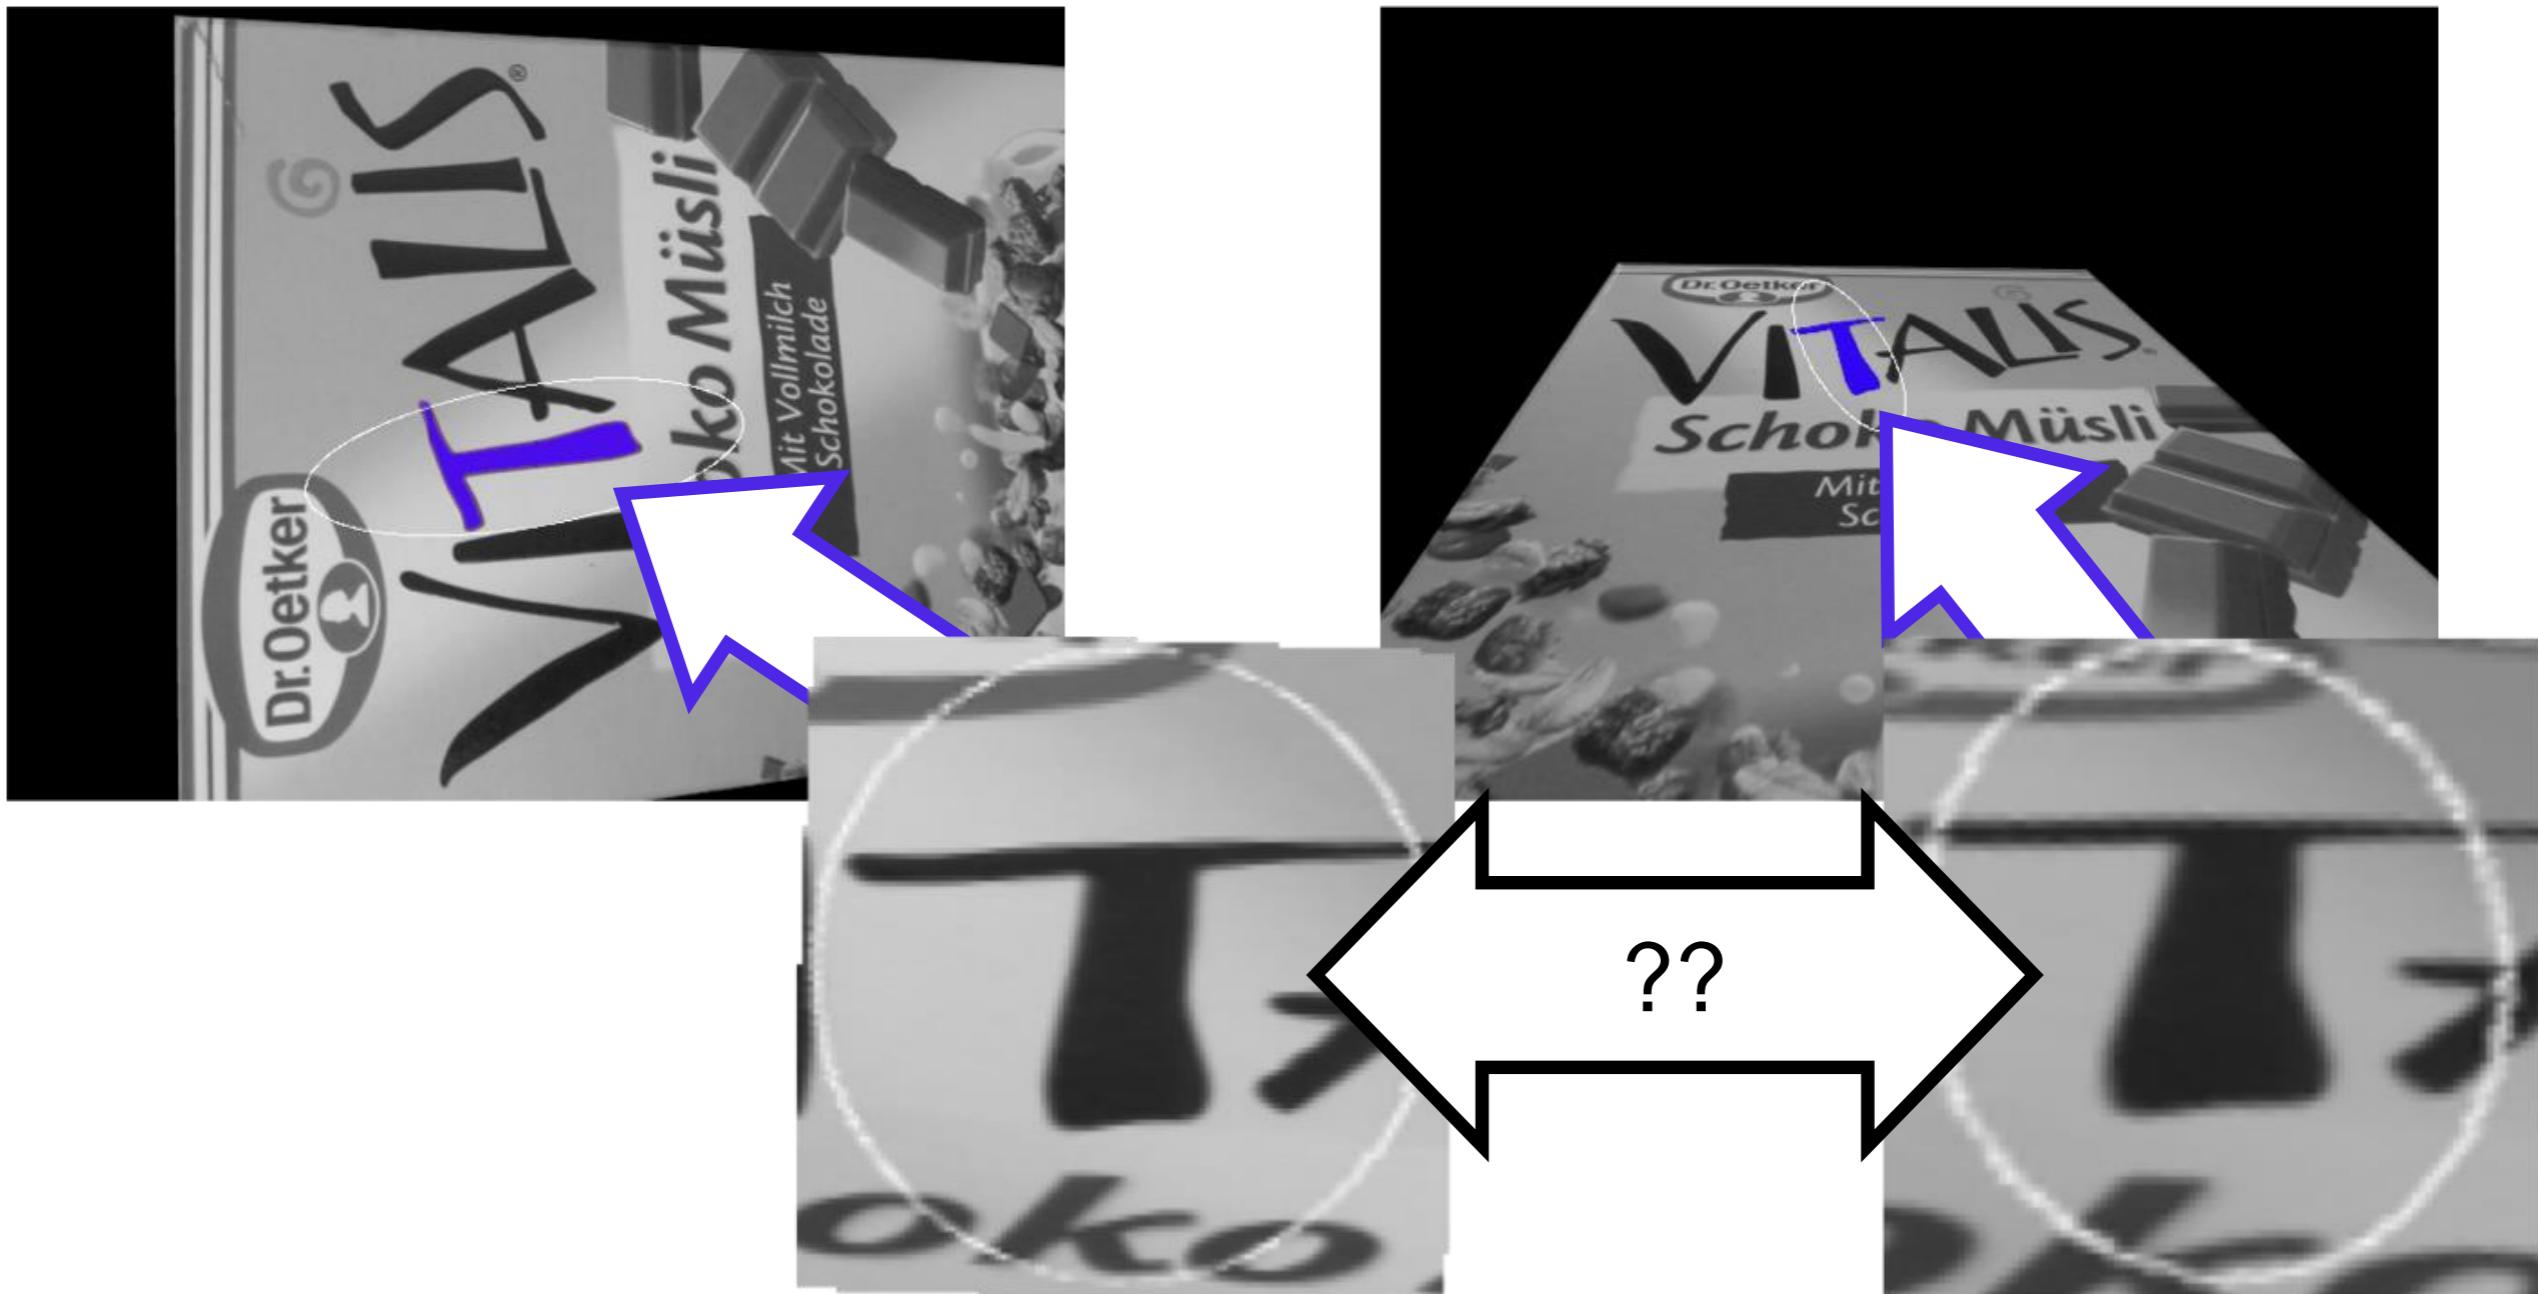
Measuring similarity between images / patches central problem



slide credit: Marc Pollefeys, Kevin Köser

Measuring Similarity

Measuring similarity between images / patches central problem



slide credit: Marc Pollefeys, Kevin Köser

Measuring Similarity

Covariance:

$$P_1 =$$



$$P_2 =$$



Measuring Similarity

Covariance:

$$P_1 =$$



$$P_2 =$$



$$\text{Cov}(P_1, P_2) = \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1}) (P_2 - \mu_2 \mathbb{1})$$

Measuring Similarity

Covariance:

$$P_1 =$$



$$P_2 =$$



$$\begin{aligned}\text{Cov}(P_1, P_2) &= \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1}) (P_2 - \mu_2 \mathbb{1}) \\ &= \frac{1}{N_{\text{elements}}} ((P_1 - \mu_1 \mathbb{1}) P_2 - \mu_2 (P_1 \mathbb{1} - \mu_1 \mathbb{1} \mathbb{1}))\end{aligned}$$

Measuring Similarity

Covariance:

$$P_1 =$$



$$P_2 =$$



$$\begin{aligned}\text{Cov}(P_1, P_2) &= \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1}) (P_2 - \mu_2 \mathbb{1}) \\ &= \frac{1}{N_{\text{elements}}} ((P_1 - \mu_1 \mathbb{1}) P_2 - \mu_2 (P_1 \mathbb{1} - \mu_1 \mathbb{1} \mathbb{1})) \\ &= \frac{1}{N_{\text{elements}}} ((P_1 - \mu_1 \mathbb{1}) P_2 - \mu_2 (N_{\text{elements}} \cdot \mu_1 - \mu_1 \cdot N_{\text{elements}}))\end{aligned}$$

Measuring Similarity

Covariance:

$$P_1 =$$



$$P_2 =$$



$$\begin{aligned}\text{Cov}(P_1, P_2) &= \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1}) (P_2 - \mu_2 \mathbb{1}) \\ &= \frac{1}{N_{\text{elements}}} ((P_1 - \mu_1 \mathbb{1}) P_2 - \mu_2 (P_1 \mathbb{1} - \mu_1 \mathbb{1} \mathbb{1})) \\ &= \frac{1}{N_{\text{elements}}} ((P_1 - \mu_1 \mathbb{1}) P_2 - \mu_2 (N_{\text{elements}} \cdot \mu_1 - \mu_1 \cdot N_{\text{elements}})) \\ &= \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1}) P_2\end{aligned}$$

Measuring Similarity

Covariance:

$$P_1 =$$



$$P_2 =$$



$$\begin{aligned}\text{Cov}(P_1, P_2) &= \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1}) (P_2 - \mu_2 \mathbb{1}) \\ &= \frac{1}{N_{\text{elements}}} ((P_1 - \mu_1 \mathbb{1}) P_2 - \mu_2 (P_1 \mathbb{1} - \mu_1 \mathbb{1} \mathbb{1})) \\ &= \frac{1}{N_{\text{elements}}} ((P_1 - \mu_1 \mathbb{1}) P_2 - \mu_2 (N_{\text{elements}} \cdot \mu_1 - \mu_1 \cdot N_{\text{elements}})) \\ &= \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1}) P_2 \quad \text{We have used this before!}\end{aligned}$$

Measuring Similarity

Is Covariance a good similarity measure?

Measuring Similarity

Is Covariance a good similarity measure?

$$P_1 =$$



$$P_2 =$$



$$= k \cdot P_1$$

Measuring Similarity

Is Covariance a good similarity measure?

$$P_1 =$$



$$P_2 =$$



$$= k \cdot P_1$$

$$\text{Cov}(P_1, P_2) = \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1})(k \cdot P_1 - k \cdot \mu_1 \mathbb{1})$$

Measuring Similarity

Is Covariance a good similarity measure?

$$P_1 =$$



$$P_2 =$$



$$= k \cdot P_1$$

$$\begin{aligned}\text{Cov}(P_1, P_2) &= \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1})(k \cdot P_1 - k \cdot \mu_1 \mathbb{1}) \\ &= k \cdot \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1})(P_1 - \mu_1 \mathbb{1})\end{aligned}$$

Measuring Similarity

Is Covariance a good similarity measure?

$$P_1 =$$



$$P_2 =$$



$$= k \cdot P_1$$

$$\begin{aligned}\text{Cov}(P_1, P_2) &= \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1})(k \cdot P_1 - k \cdot \mu_1 \mathbb{1}) \\ &= k \cdot \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1})(P_1 - \mu_1 \mathbb{1}) \\ &= k \cdot \text{Cov}(P_1, P_1)\end{aligned}$$

Measuring Similarity

Is Covariance a good similarity measure?

$$P_1 =$$



$$P_2 =$$



$$= k \cdot P_1$$

$$\begin{aligned}\text{Cov}(P_1, P_2) &= \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1})(k \cdot P_1 - k \cdot \mu_1 \mathbb{1}) \\ &= k \cdot \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1})(P_1 - \mu_1 \mathbb{1}) \\ &= k \cdot \text{Cov}(P_1, P_1) > \text{Cov}(P_1, P_1)\end{aligned}$$

Measuring Similarity

Is Covariance a good similarity measure?

$$P_1 =$$



$$P_2 =$$



$$= k \cdot P_1$$

$$\begin{aligned}\text{Cov}(P_1, P_2) &= \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1})(k \cdot P_1 - k \cdot \mu_1 \mathbb{1}) \\ &= k \cdot \frac{1}{N_{\text{elements}}} (P_1 - \mu_1 \mathbb{1})(P_1 - \mu_1 \mathbb{1}) \\ &= k \cdot \text{Cov}(P_1, P_1) > \text{Cov}(P_1, P_1)\end{aligned}$$

This is not what we want!

Measuring Similarity

Correlation:



$$\text{Corr}(P_1, P_2) = \frac{\text{Cov}(P_1, P_2)}{\sqrt{\text{Cov}(P_1, P_1)} \cdot \sqrt{\text{Cov}(P_2, P_2)}}$$

Also known as Zero-Mean Normalized Cross-Correlation (ZNCC)

Measuring Similarity

When does ZNCC work?

Measuring Similarity

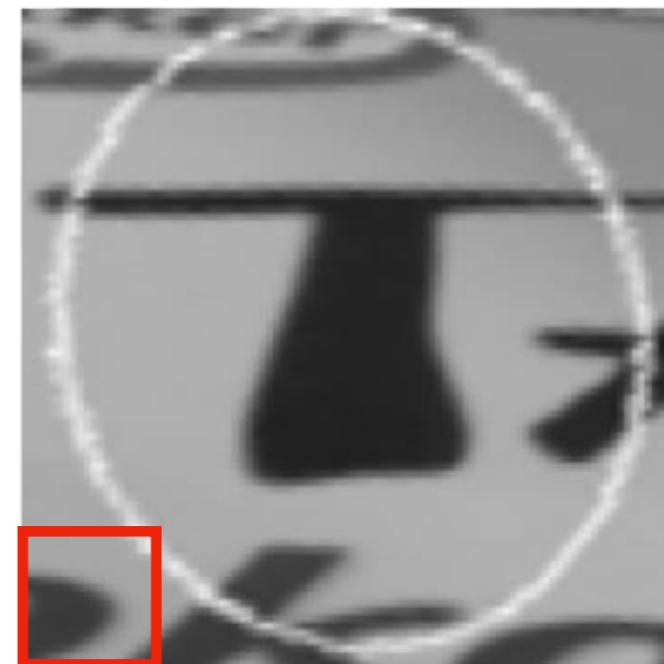
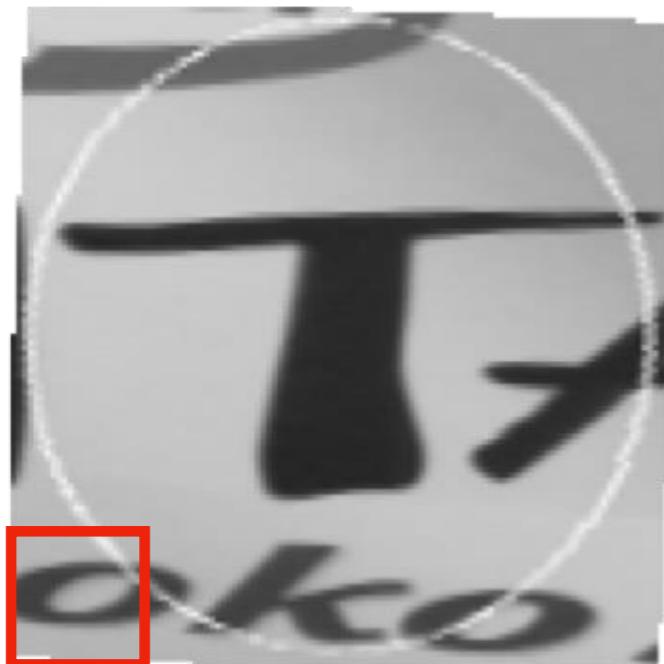
When does ZNCC work?



slide credit: Marc Pollefeys, Kevin Köser

Measuring Similarity

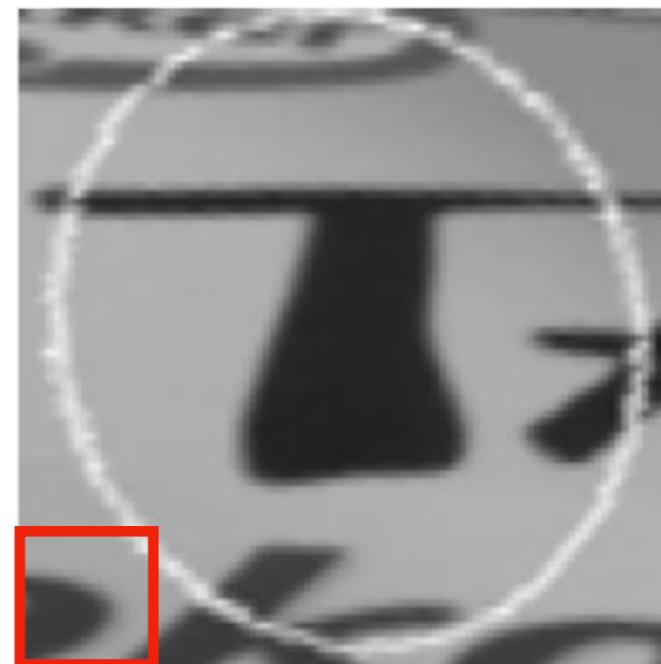
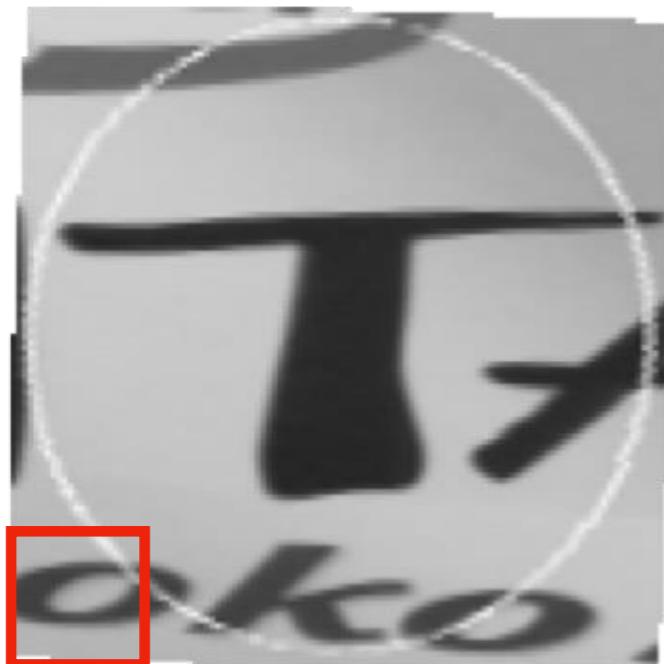
When does ZNCC work?



slide credit: Marc Pollefeys, Kevin Köser

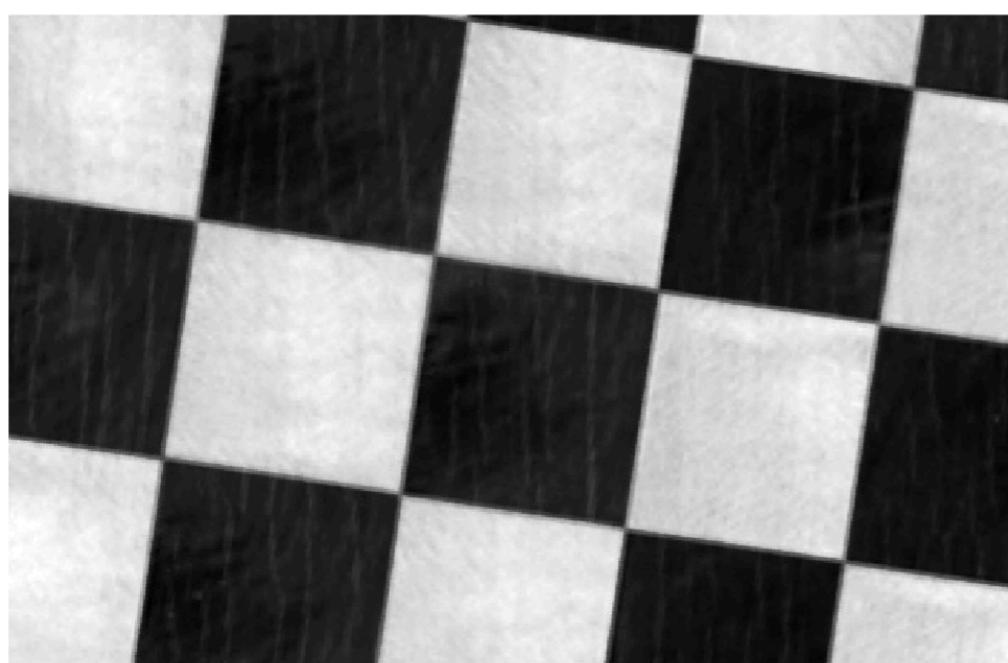
Measuring Similarity

When does ZNCC work?

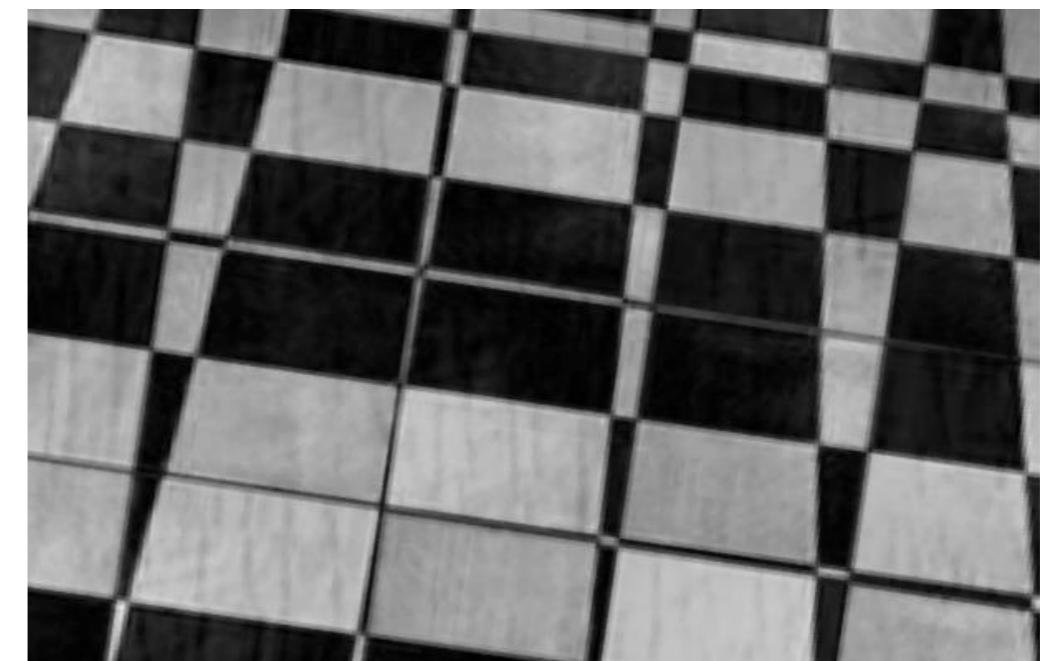
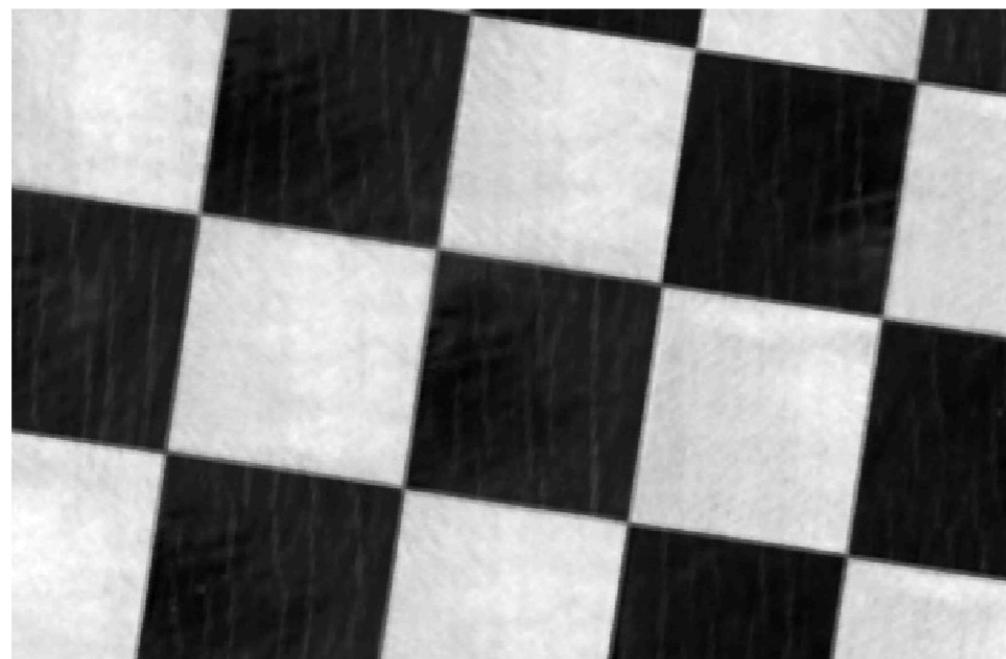


Pixels need to be (roughly) aligned!

Measuring Similarity

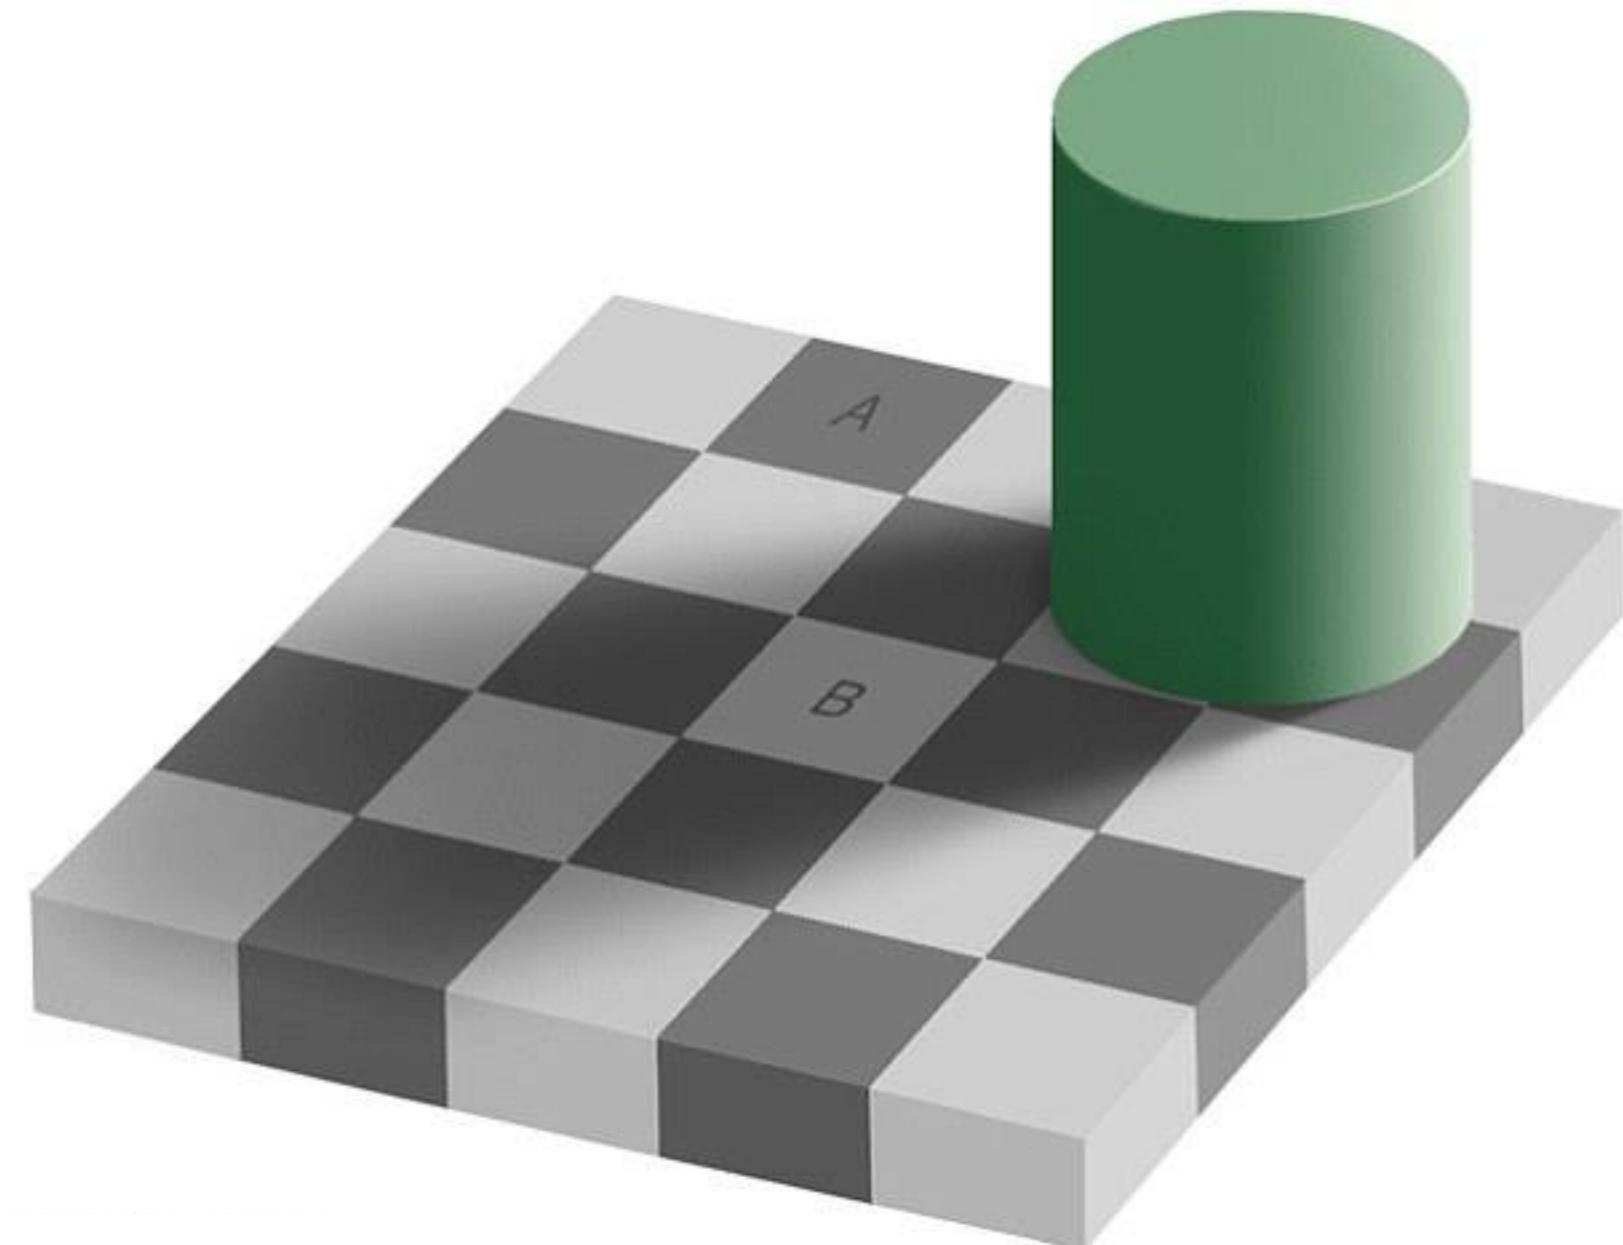


Measuring Similarity

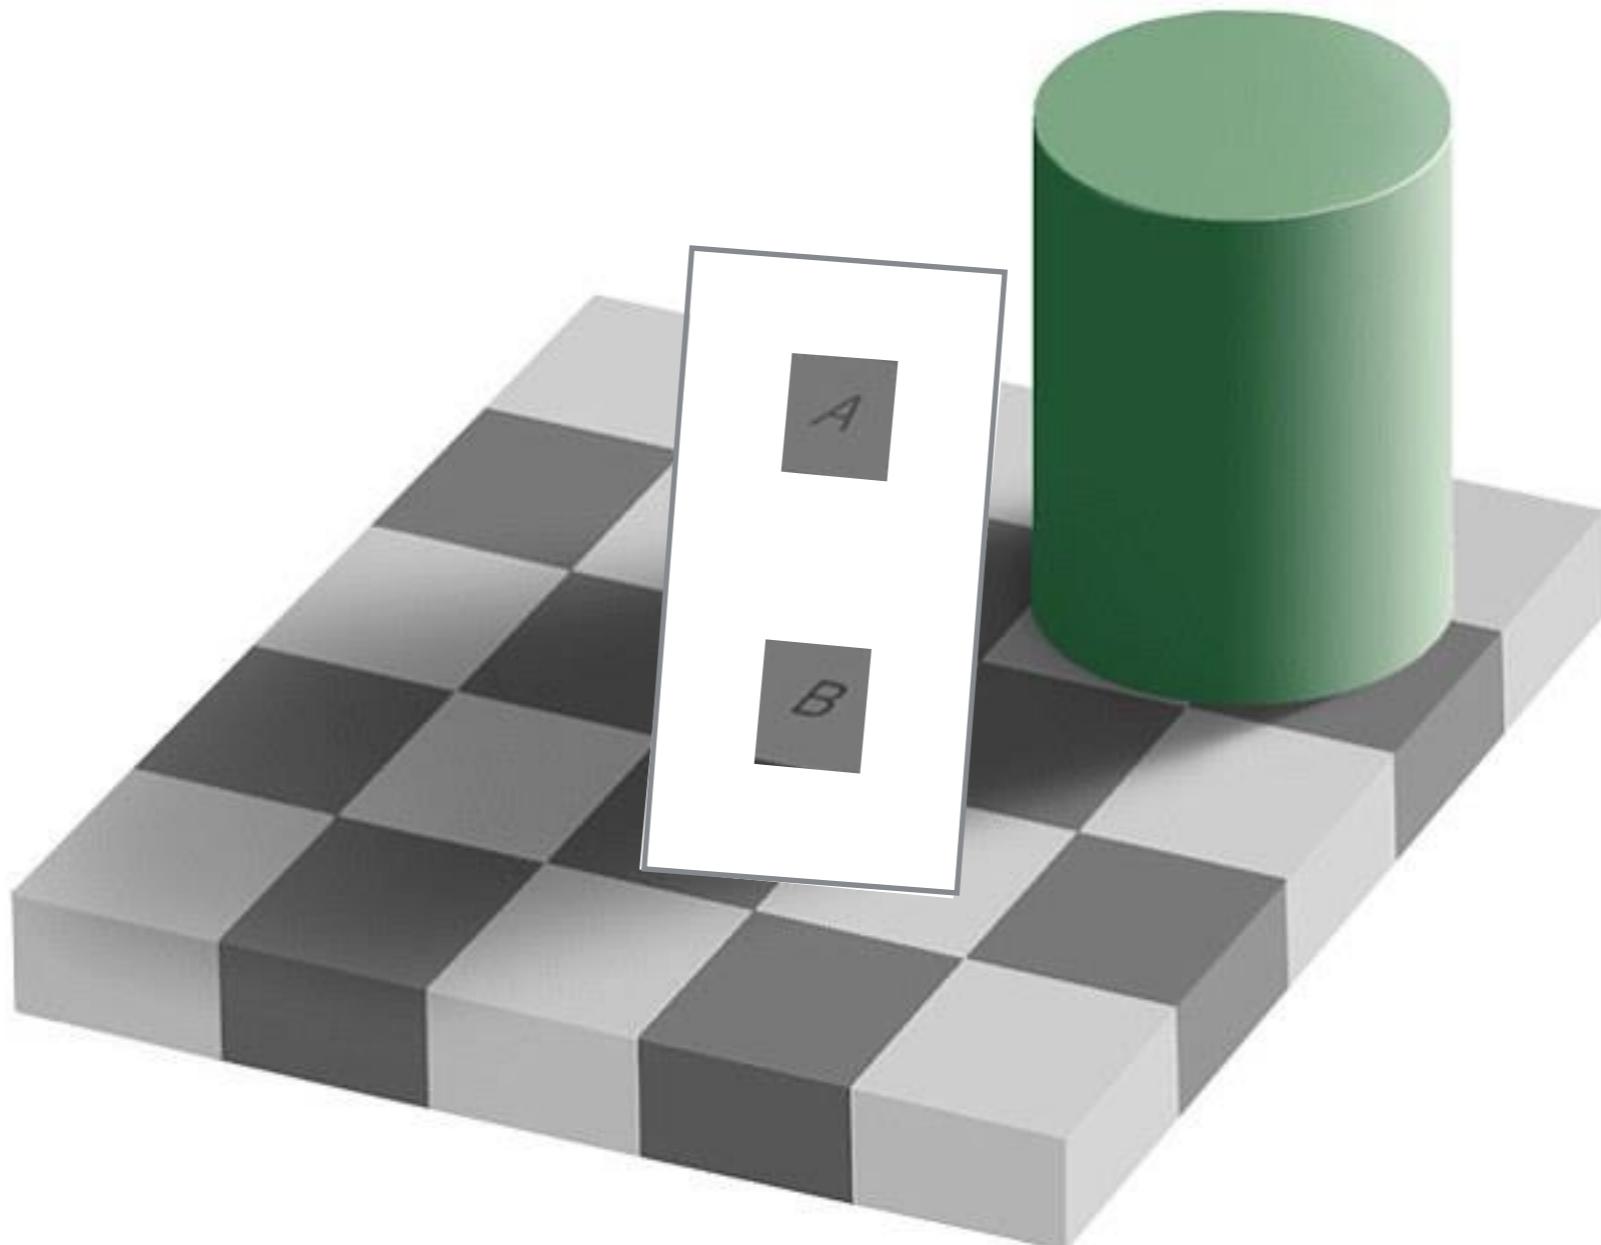


Difference image
black = no difference
white = difference

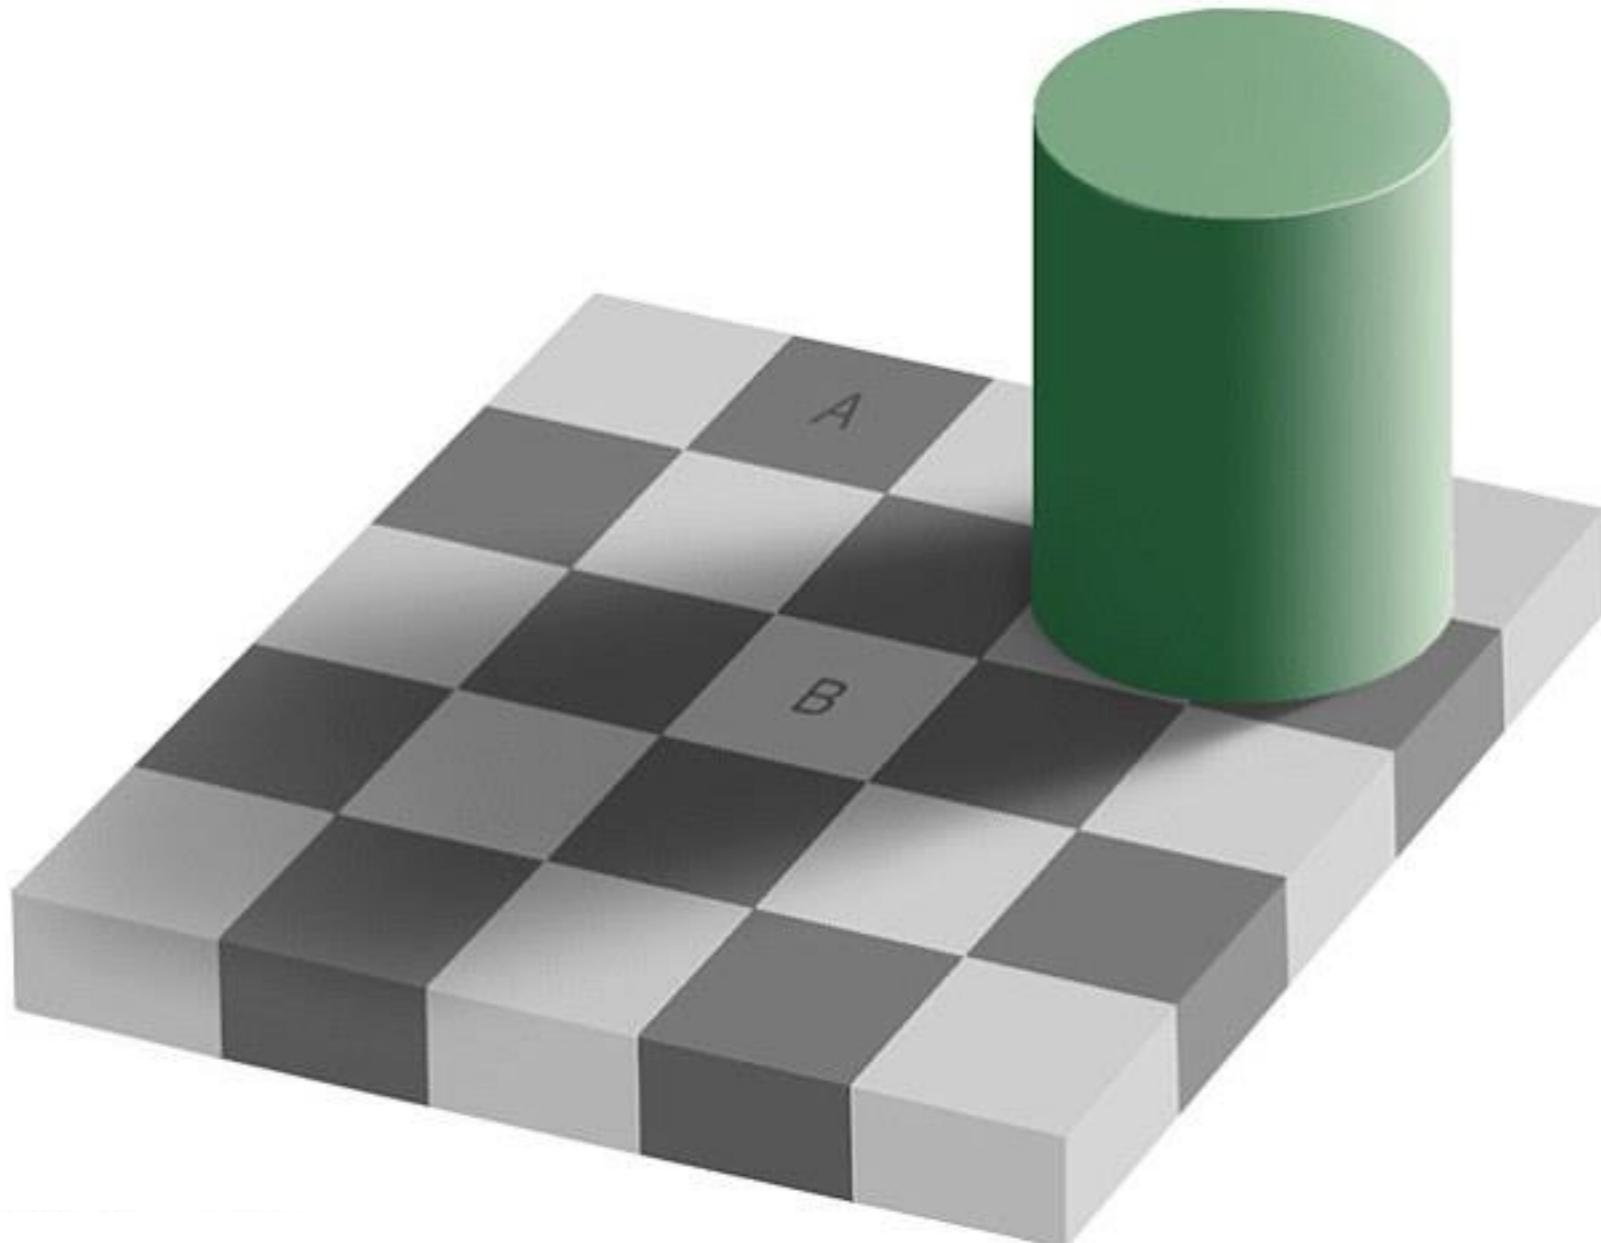
A Hint?



A Hint?

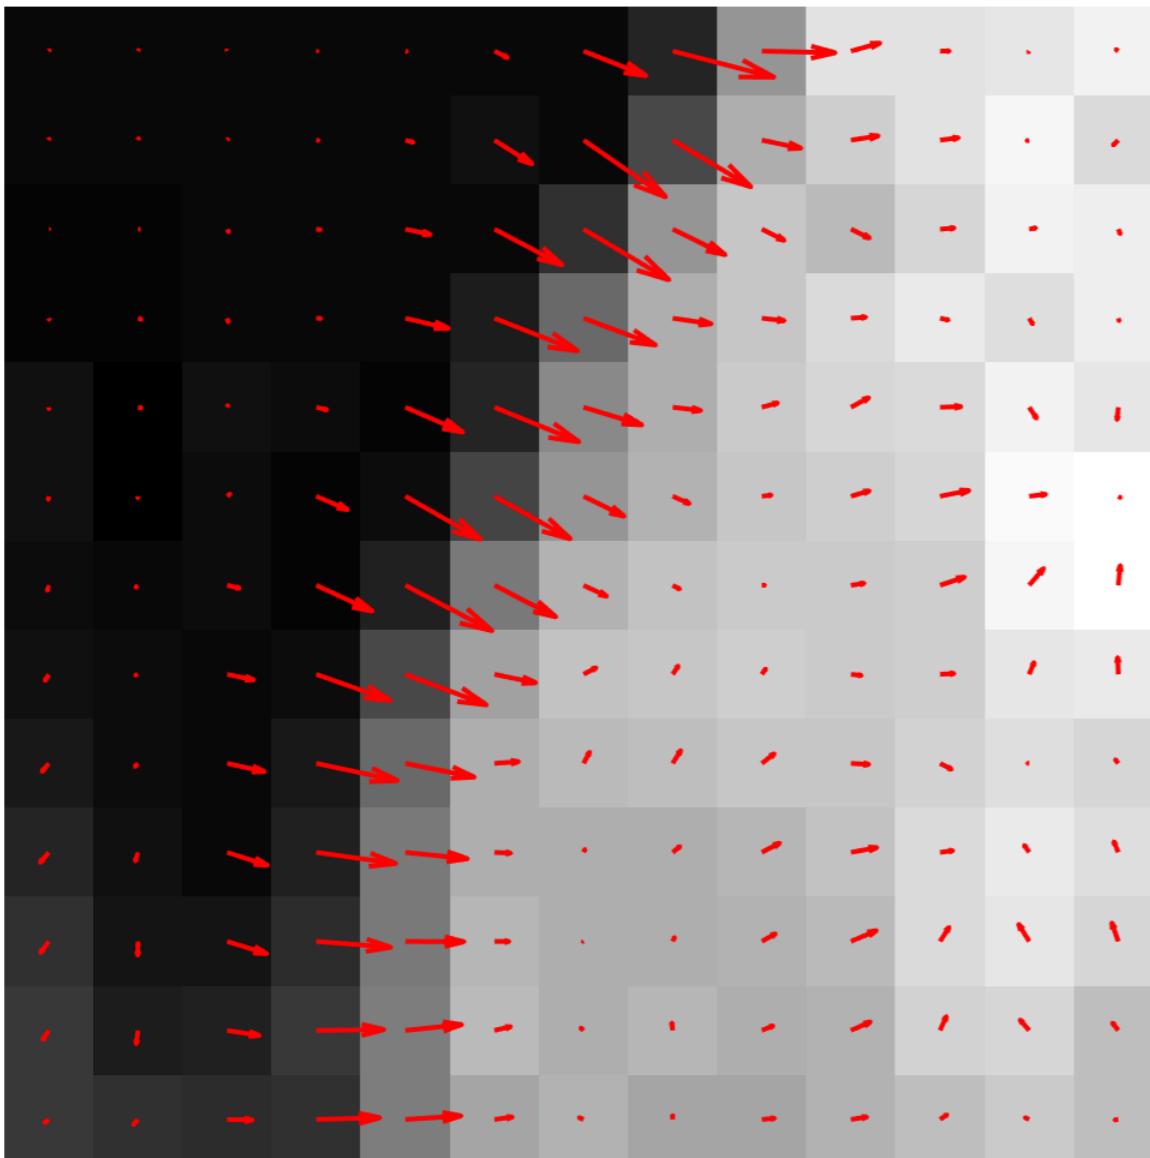


A Hint?

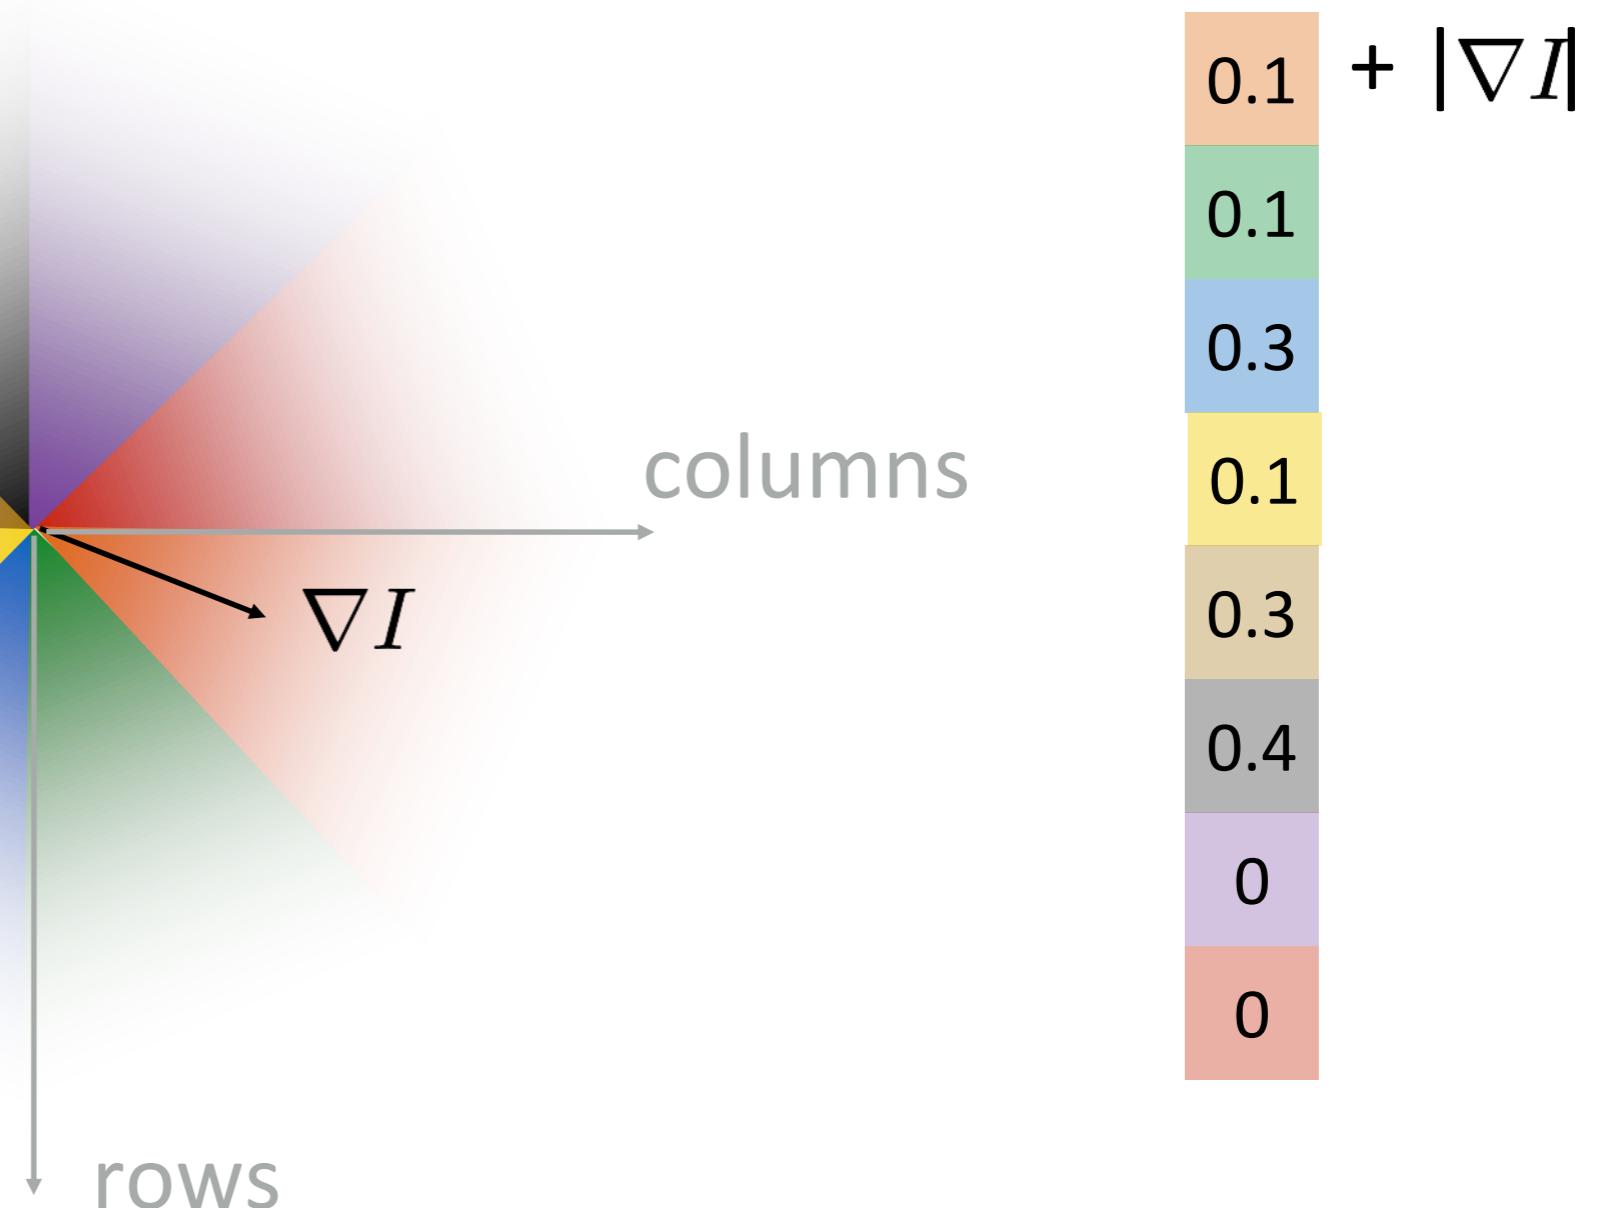


Humans seem to focus on gradients rather than absolute values.

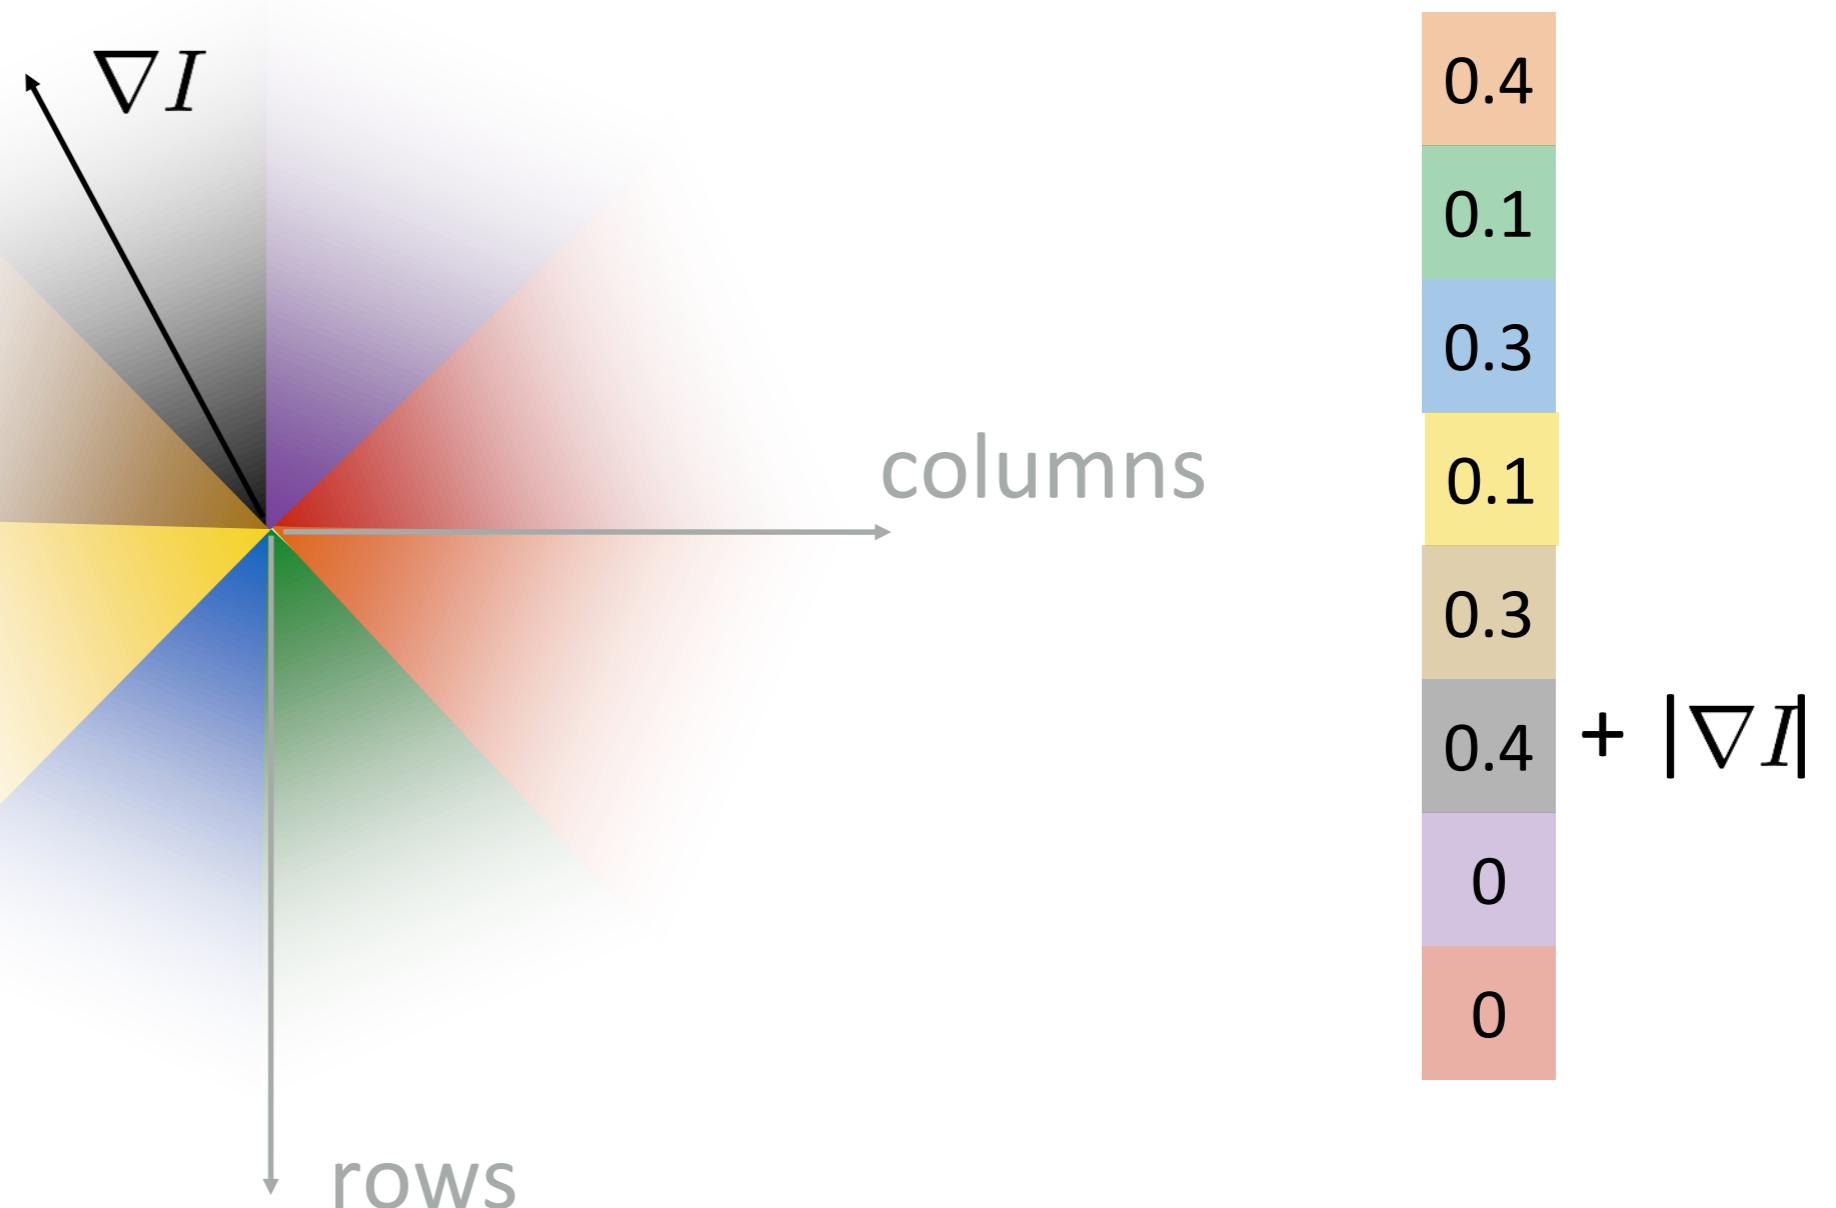
Gradient Histogram



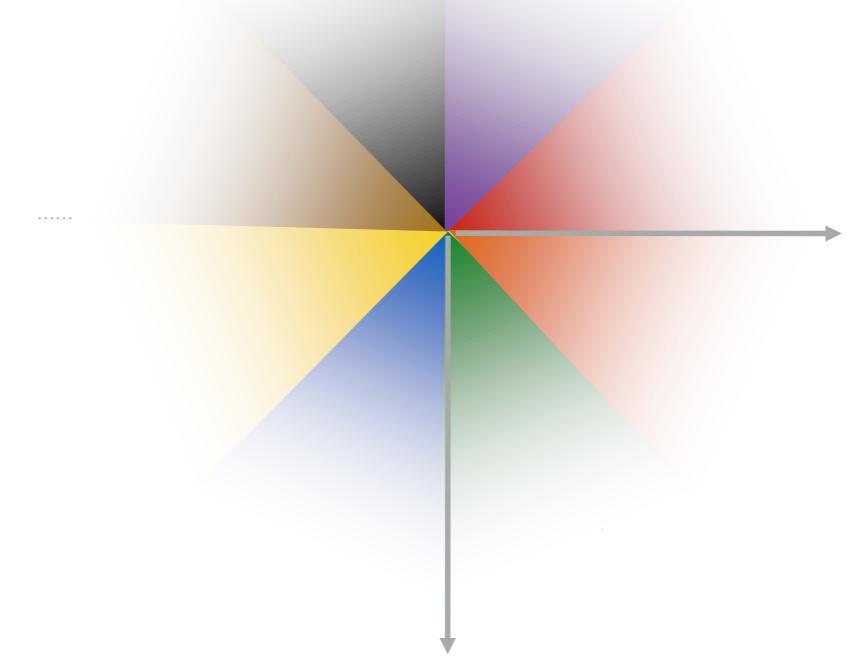
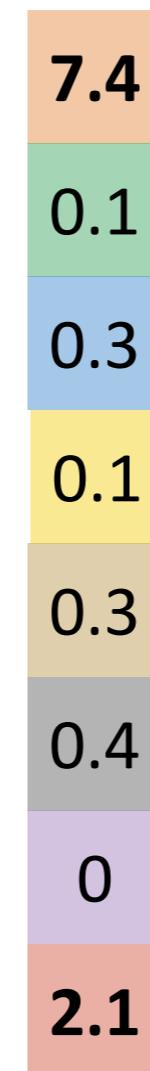
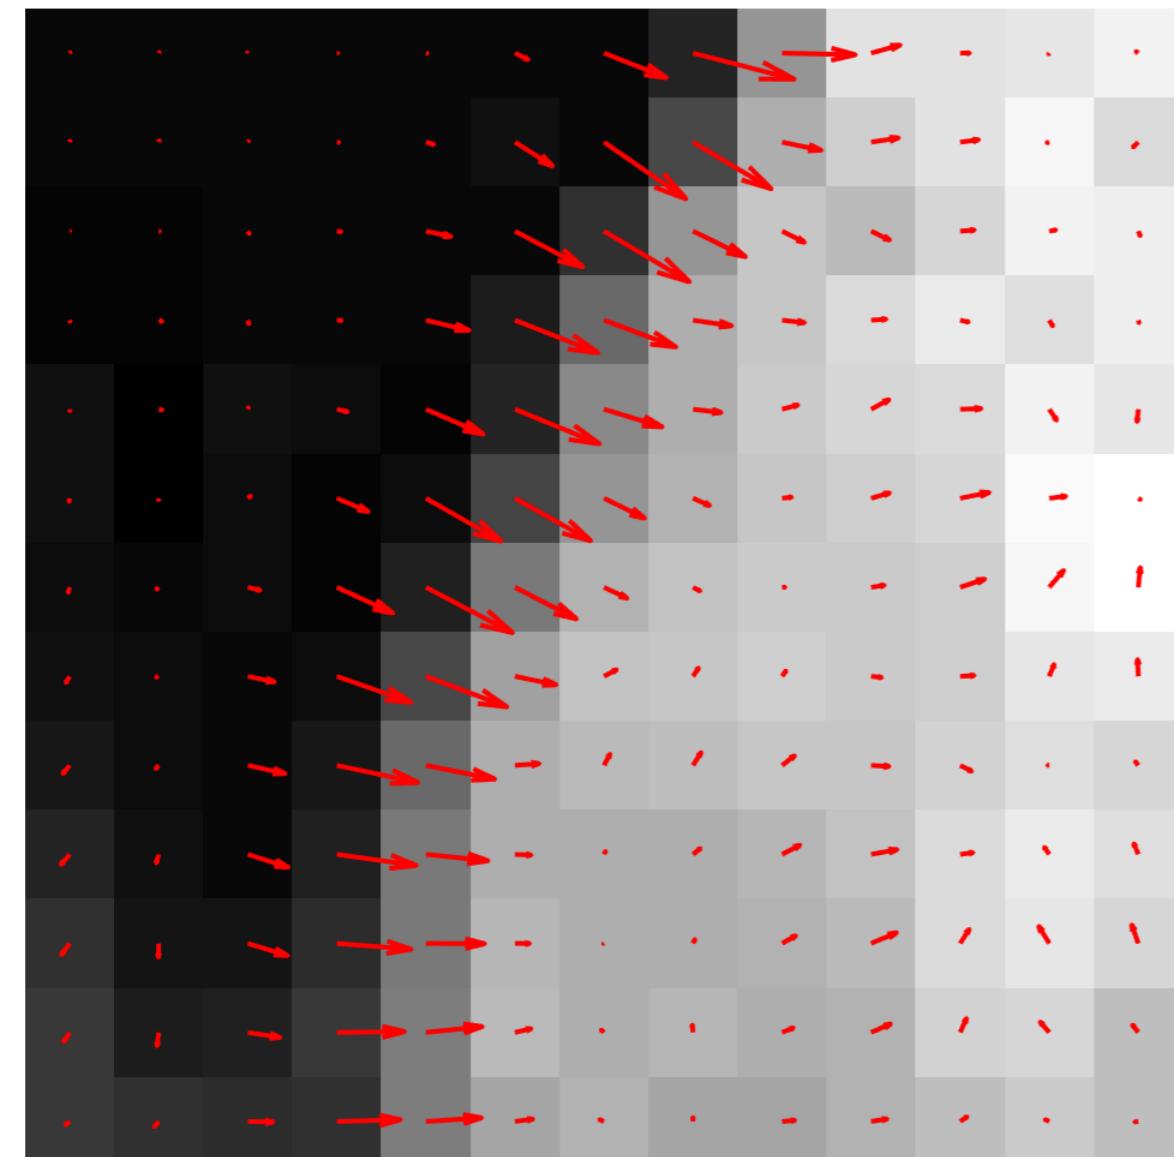
Gradient Histogram



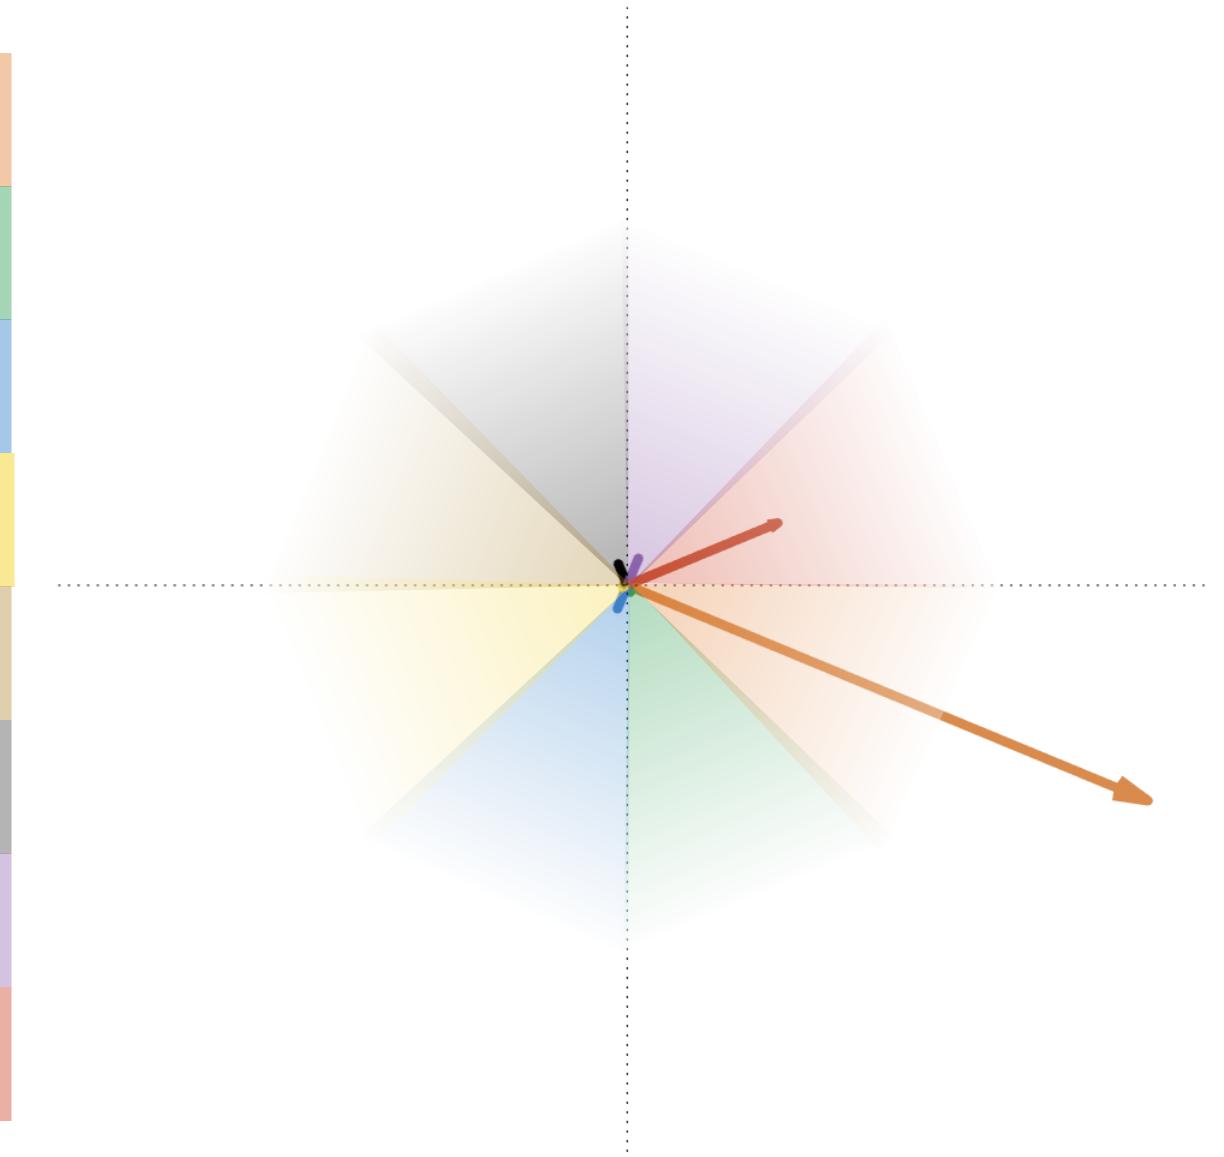
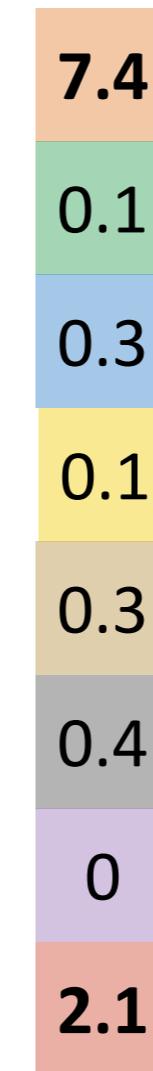
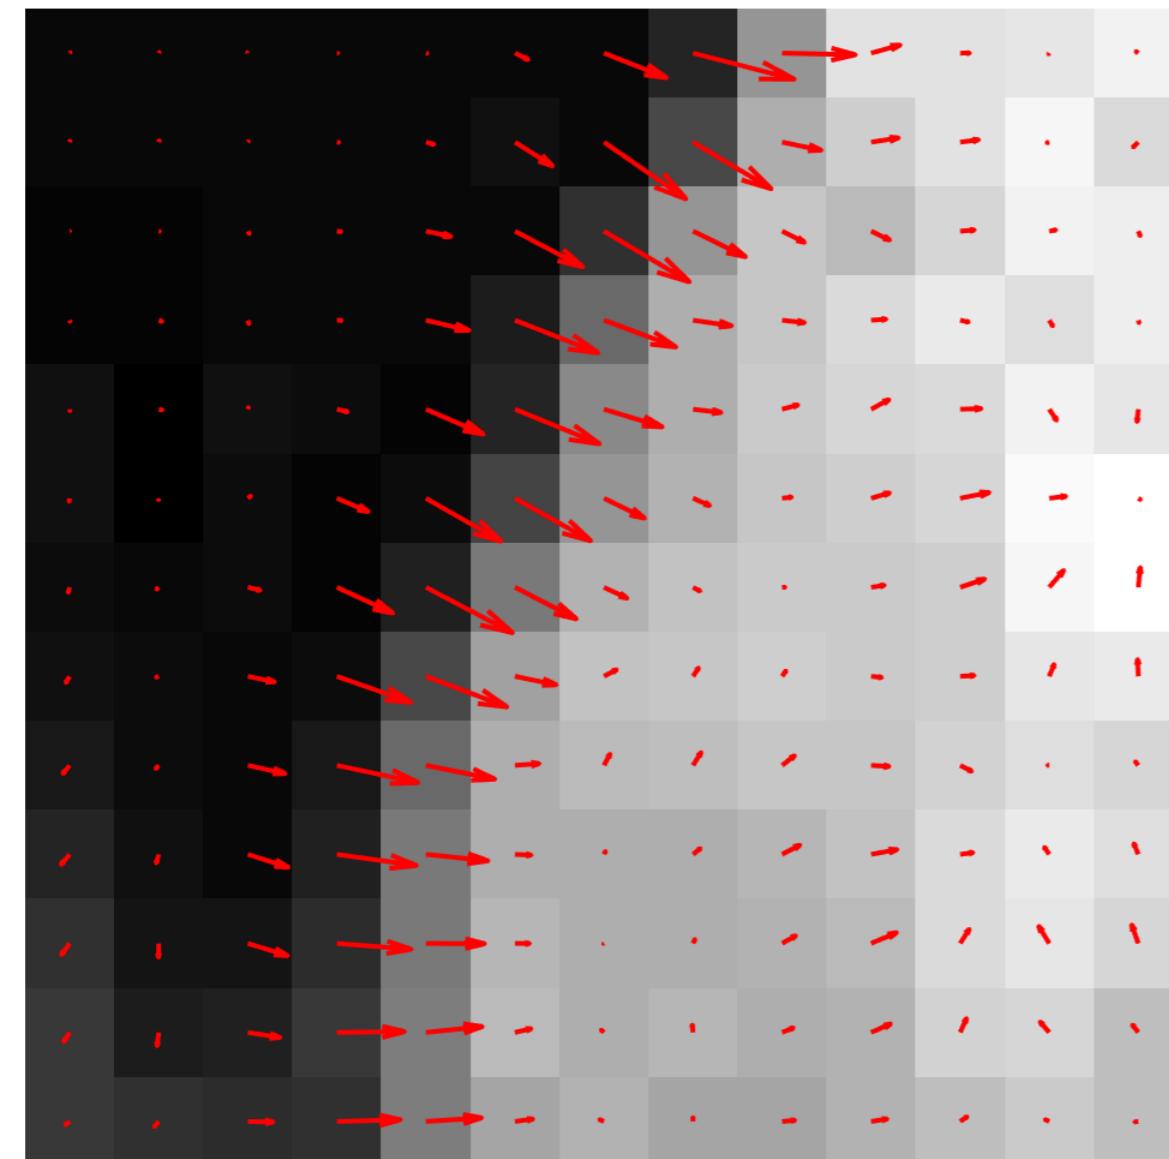
Gradient Histogram



Gradient Histogram

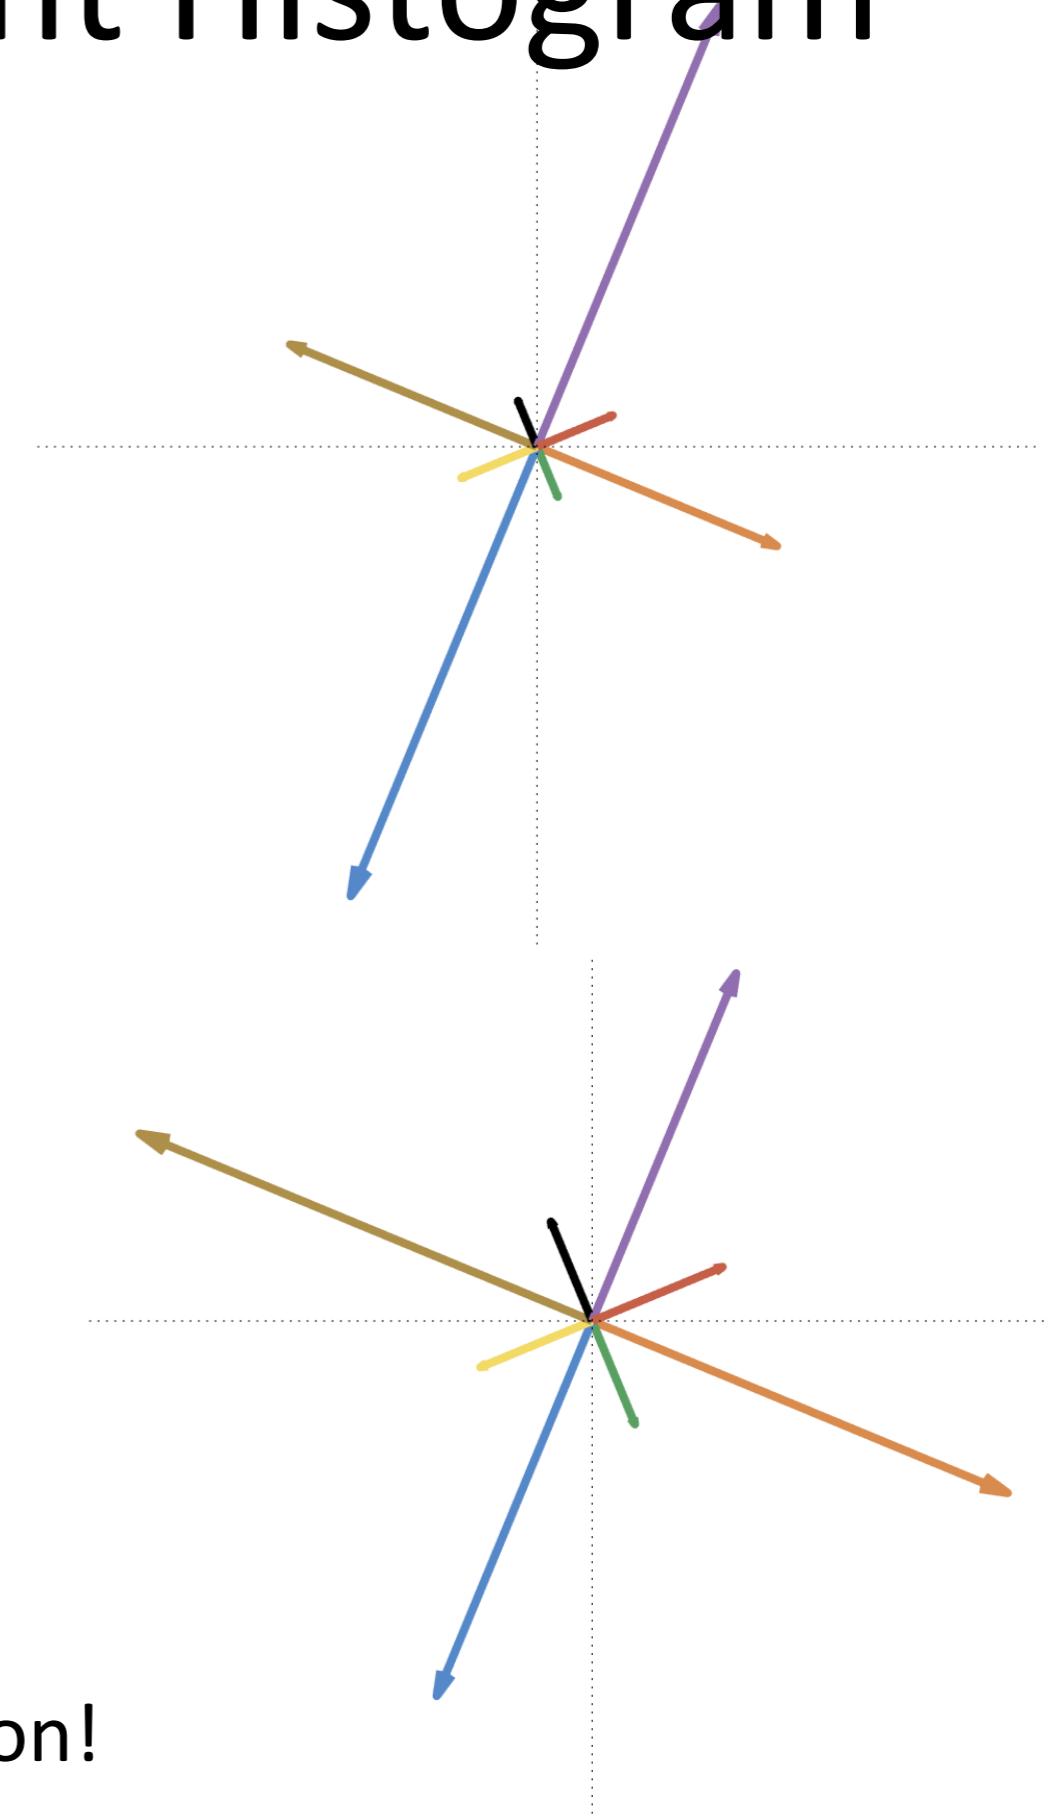
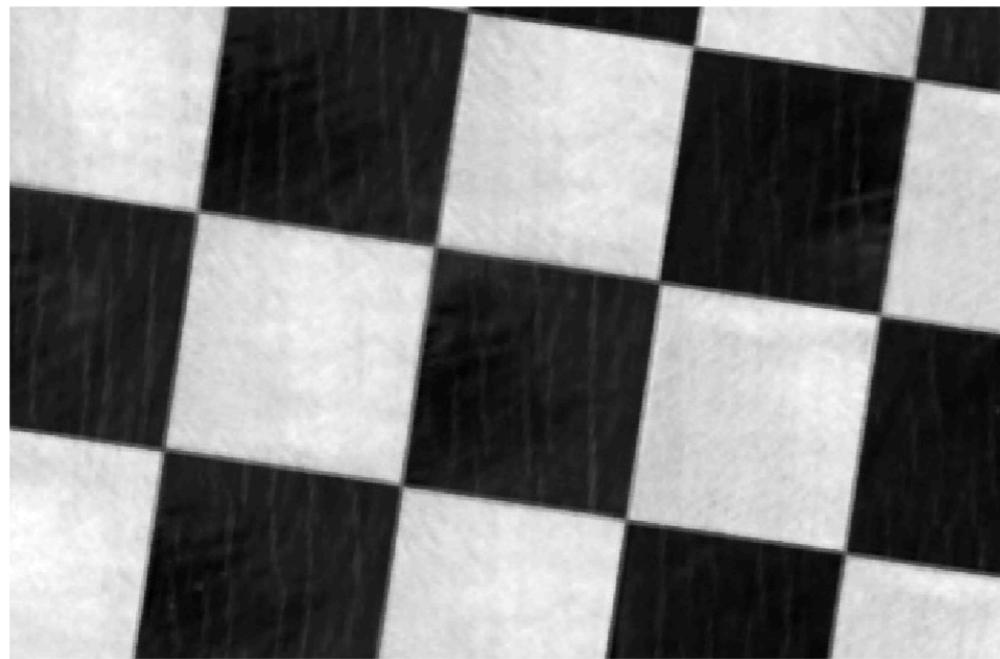


Gradient Histogram



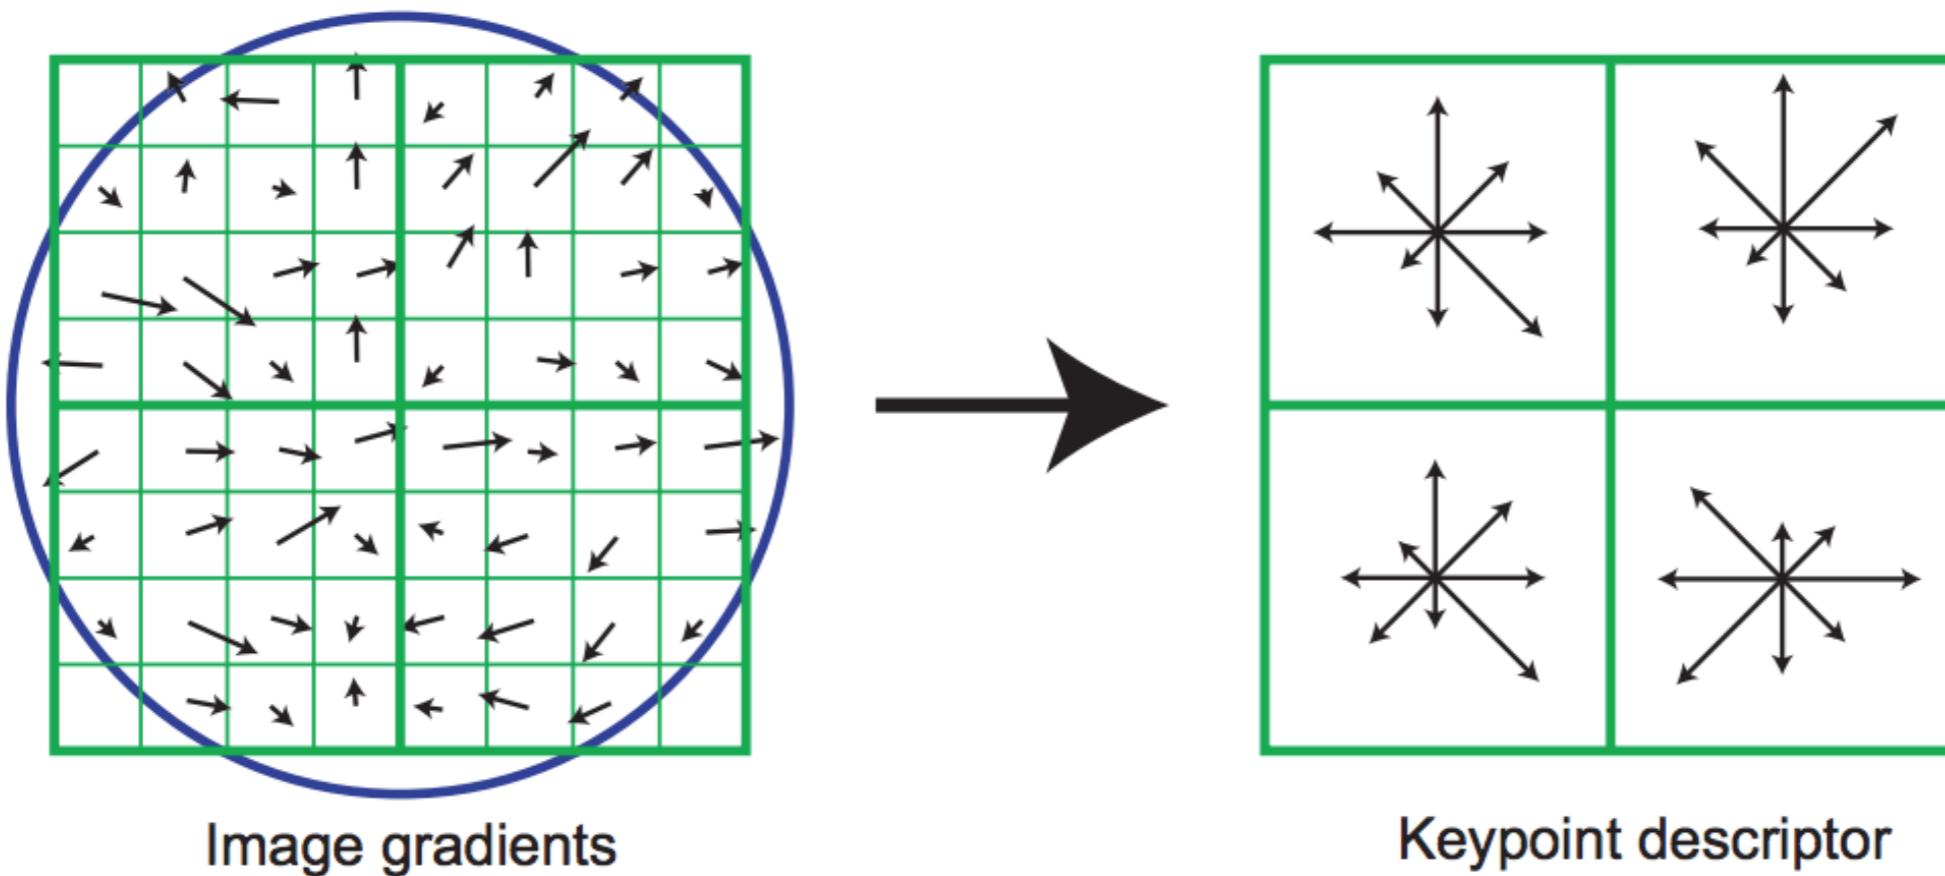
“bouquet of vectors”

Gradient Histogram



Much more similar than pixel-wise comparison!

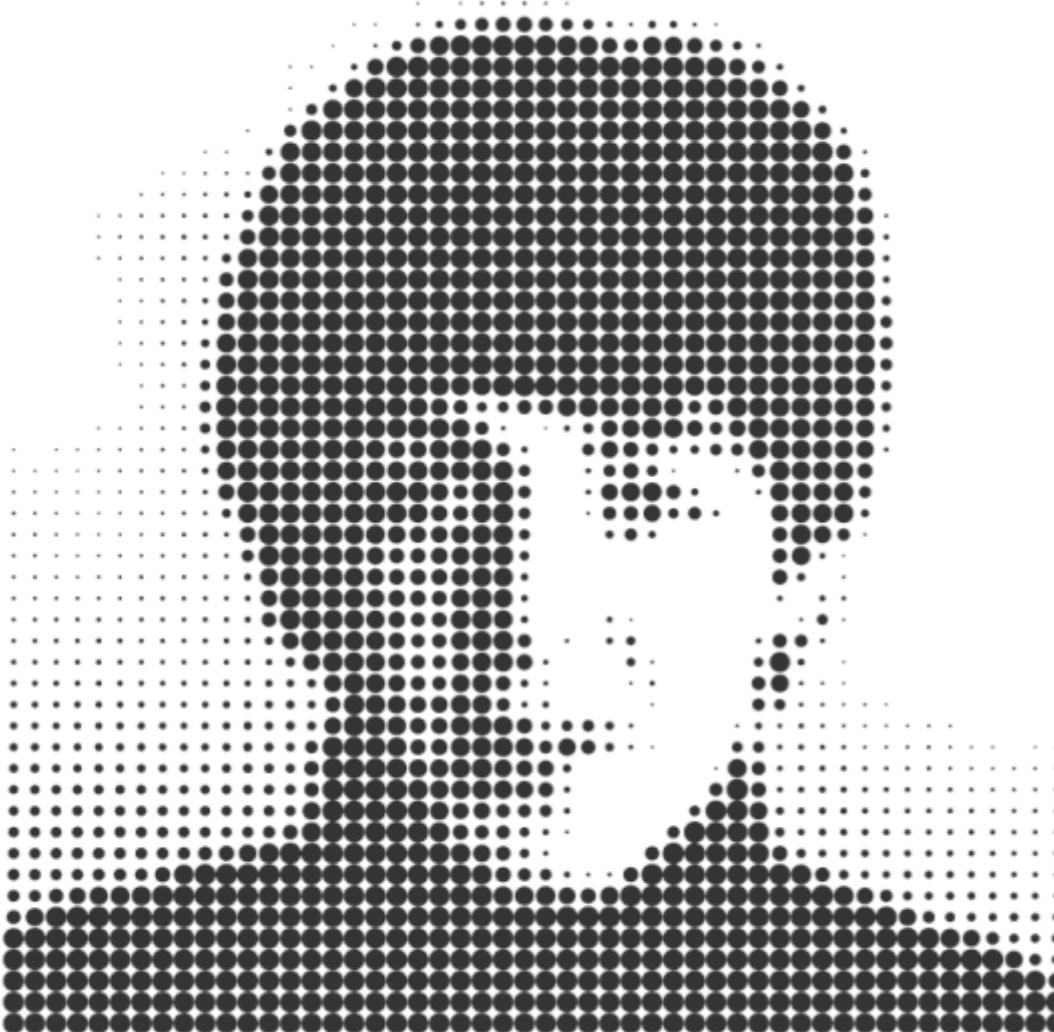
SIFT / HOG



Known as **SIFT** or Histogram of Oriented Gradients (**HOG**)
More details in next lecture

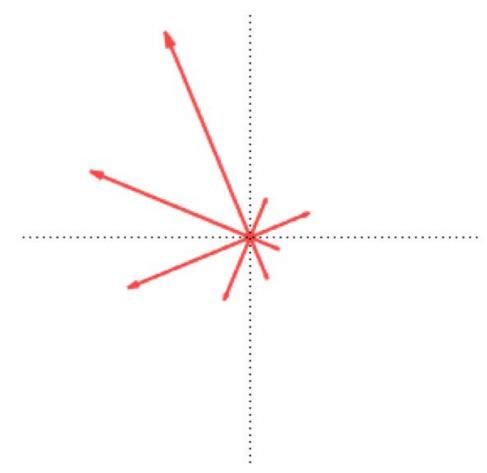
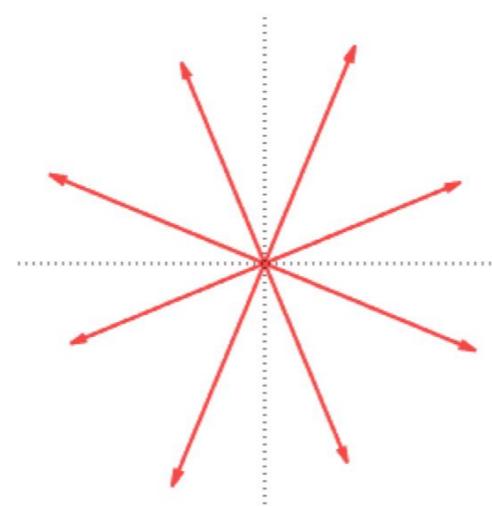
Multi-Scale Processing

Scale

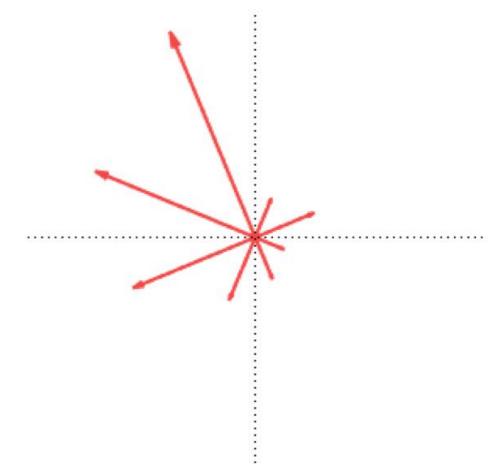
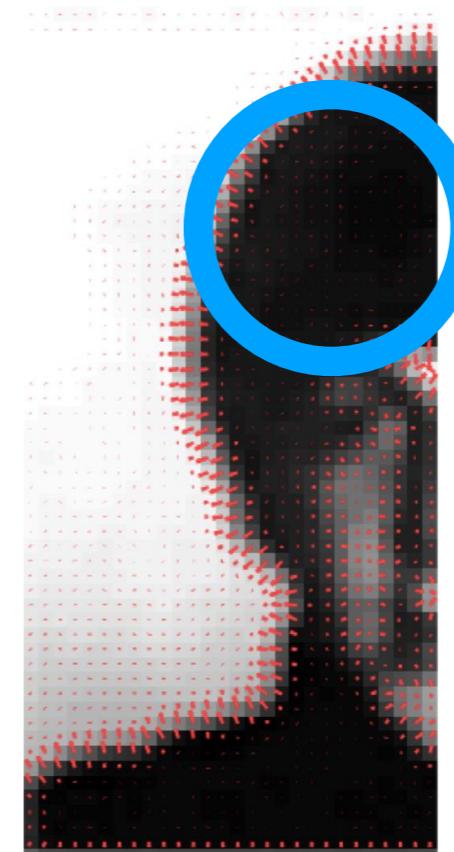
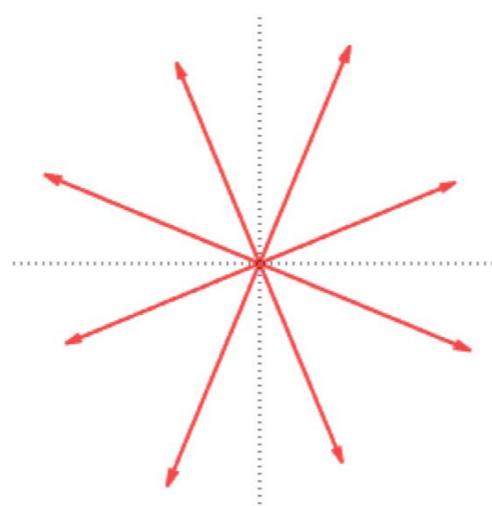
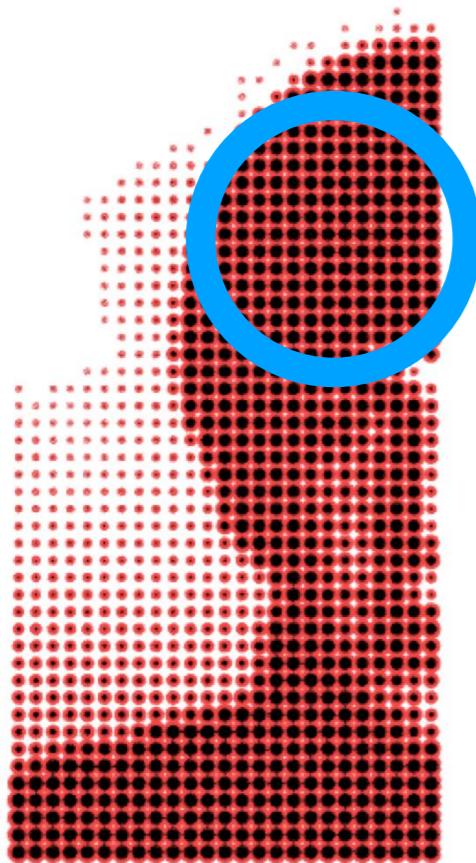


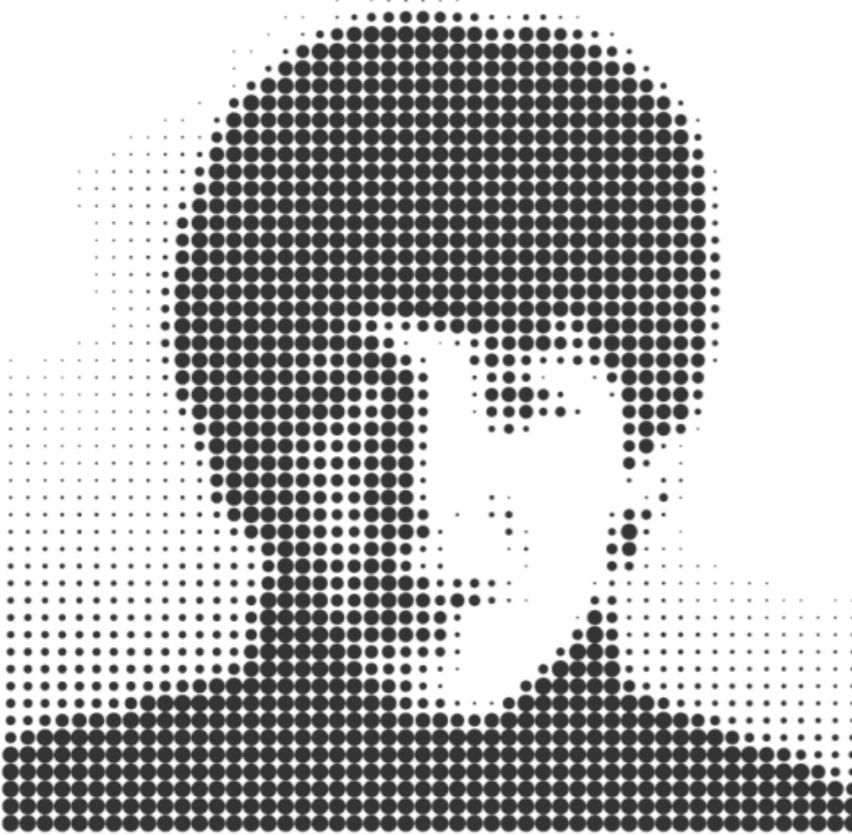
How to recognize objects at different distances?

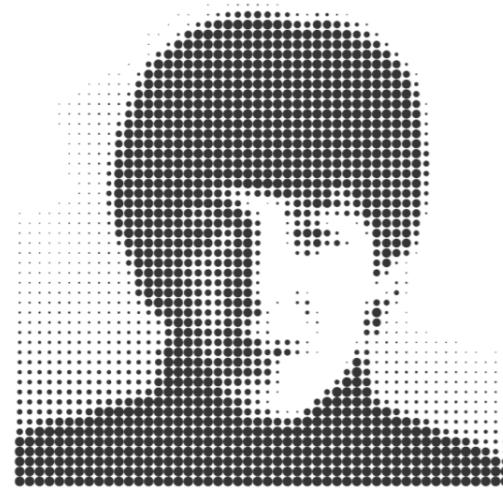
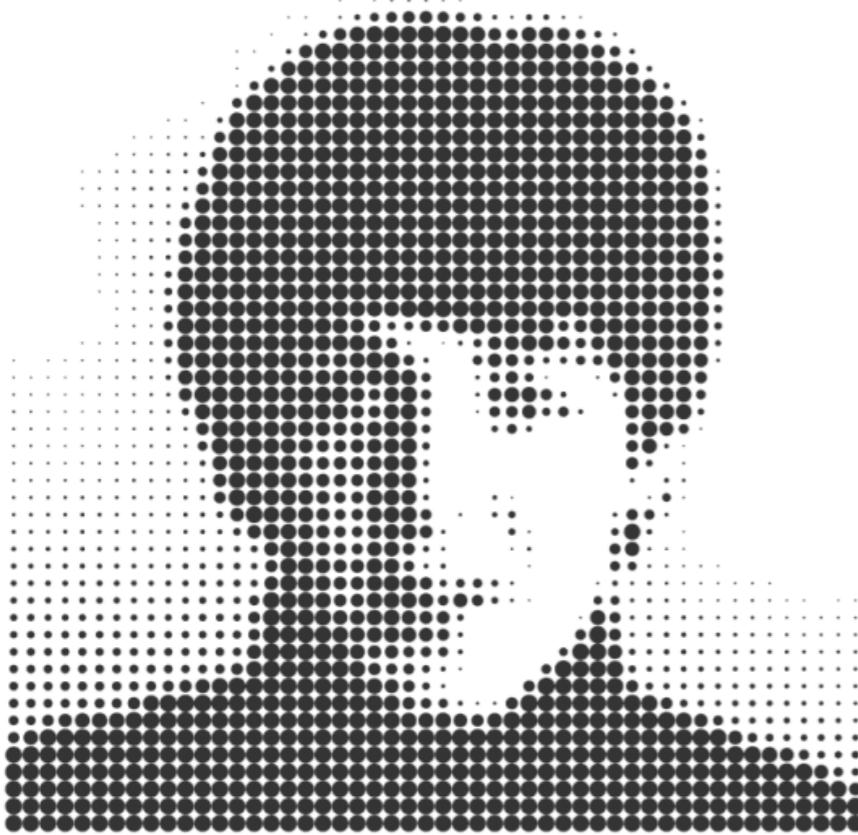
Gradient histogram

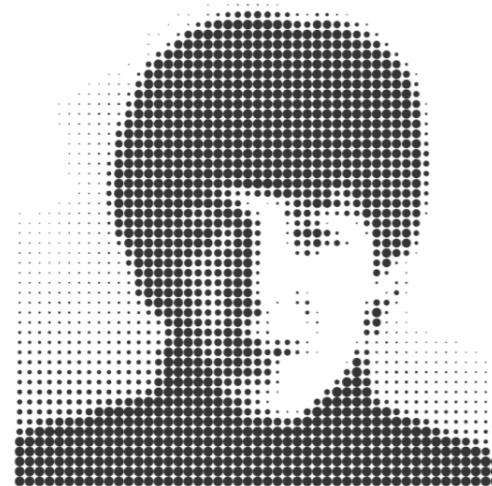
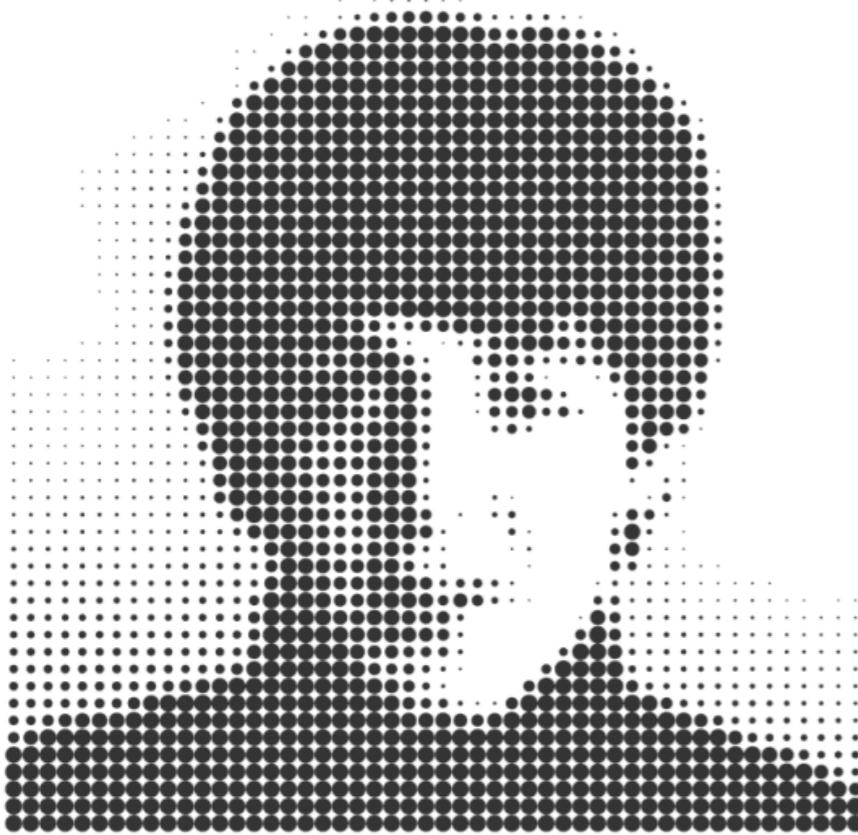


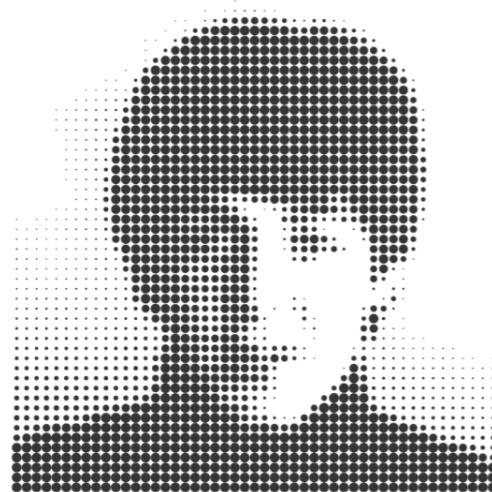
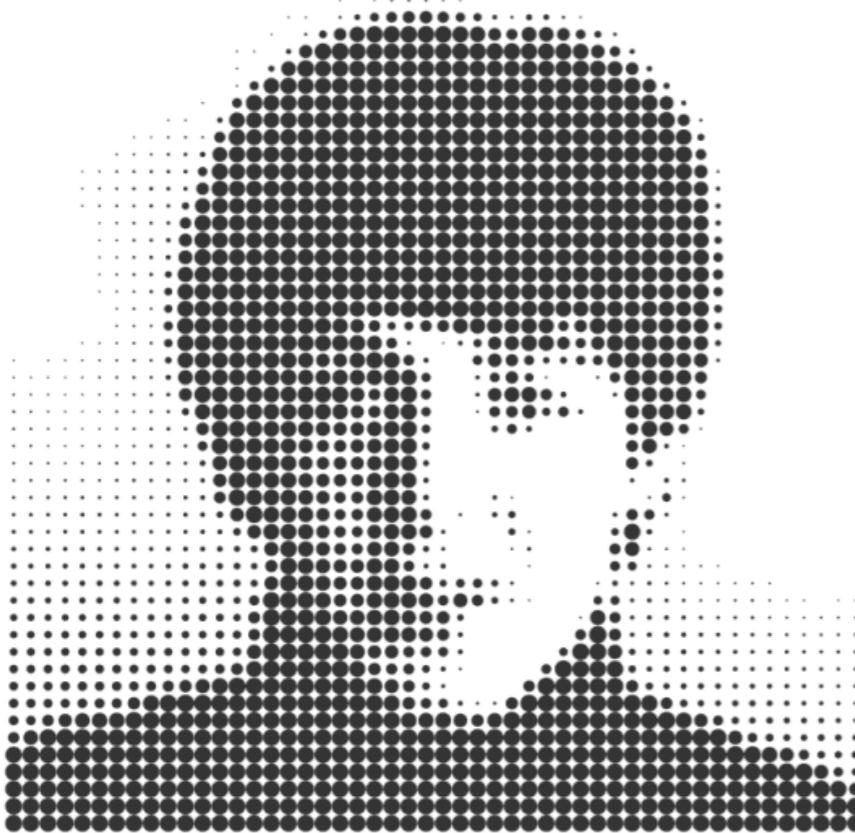
Gradient histogram

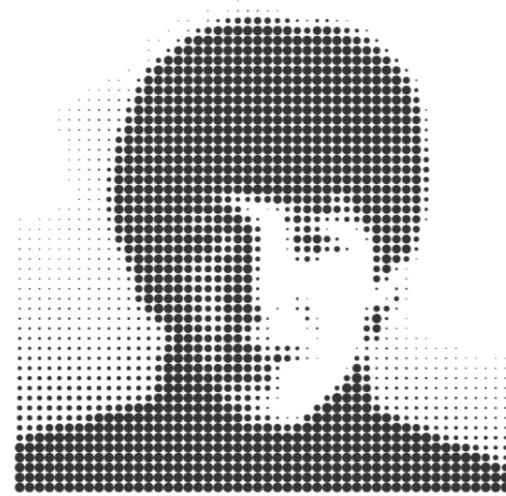
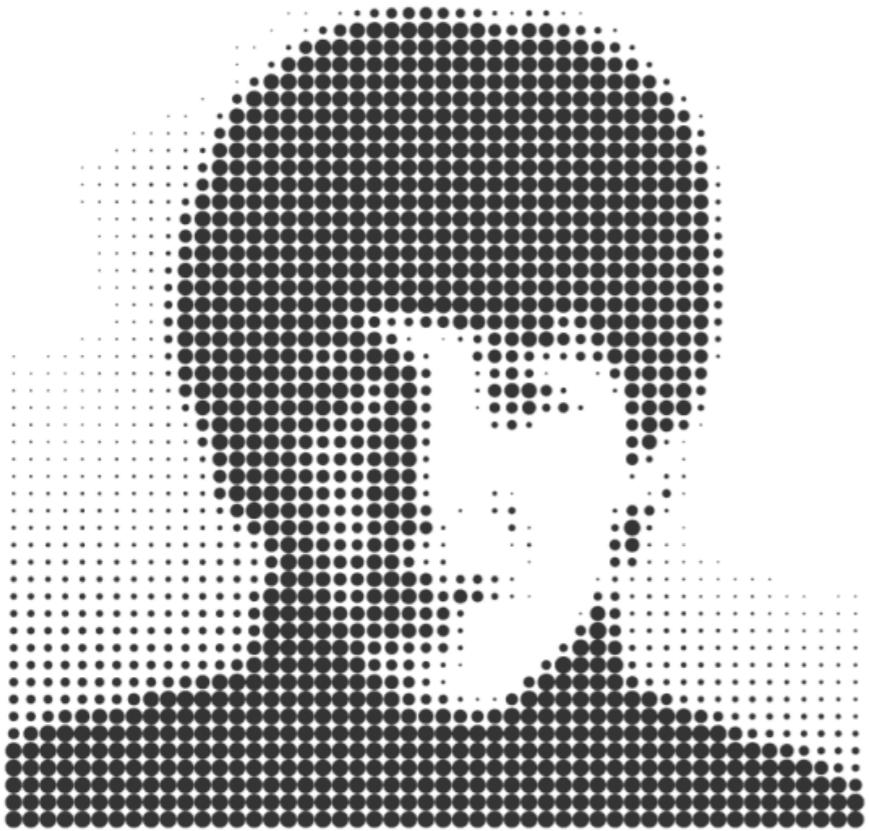


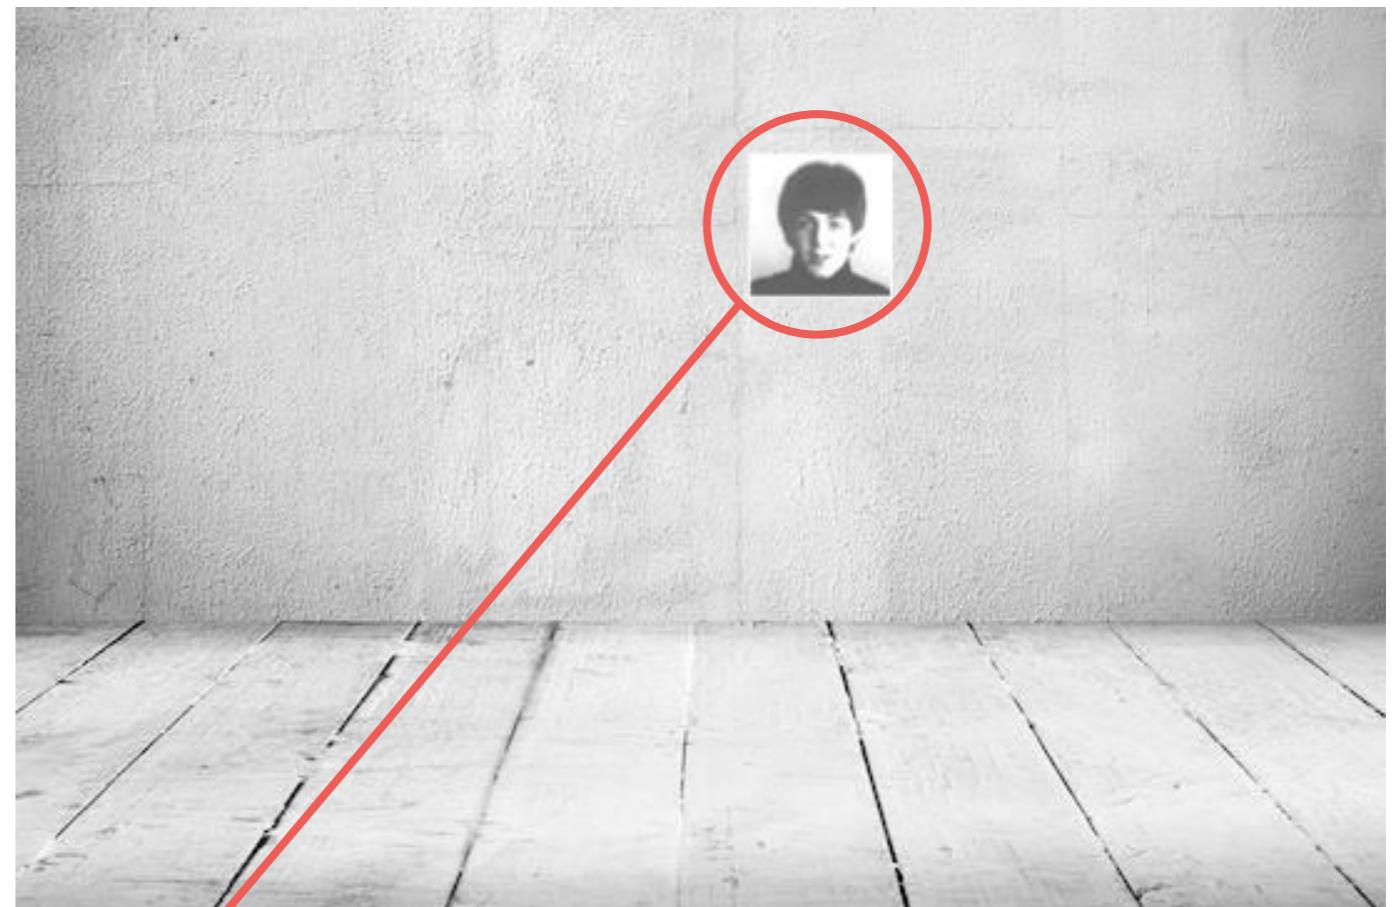
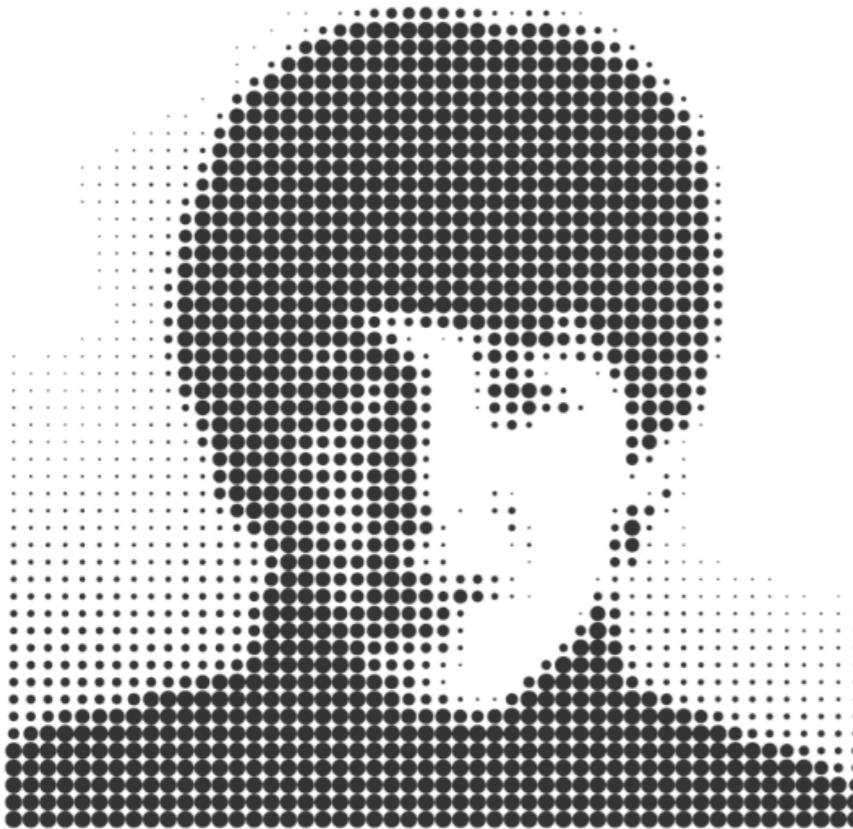




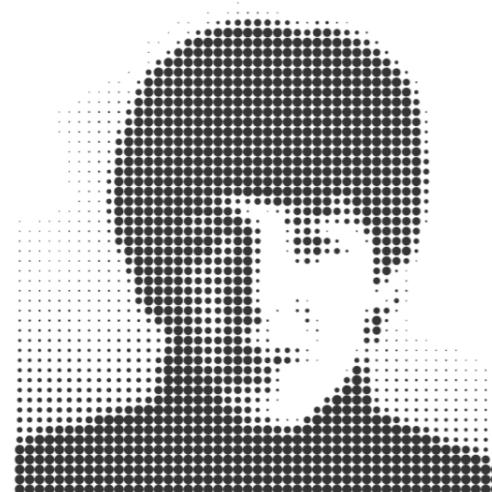






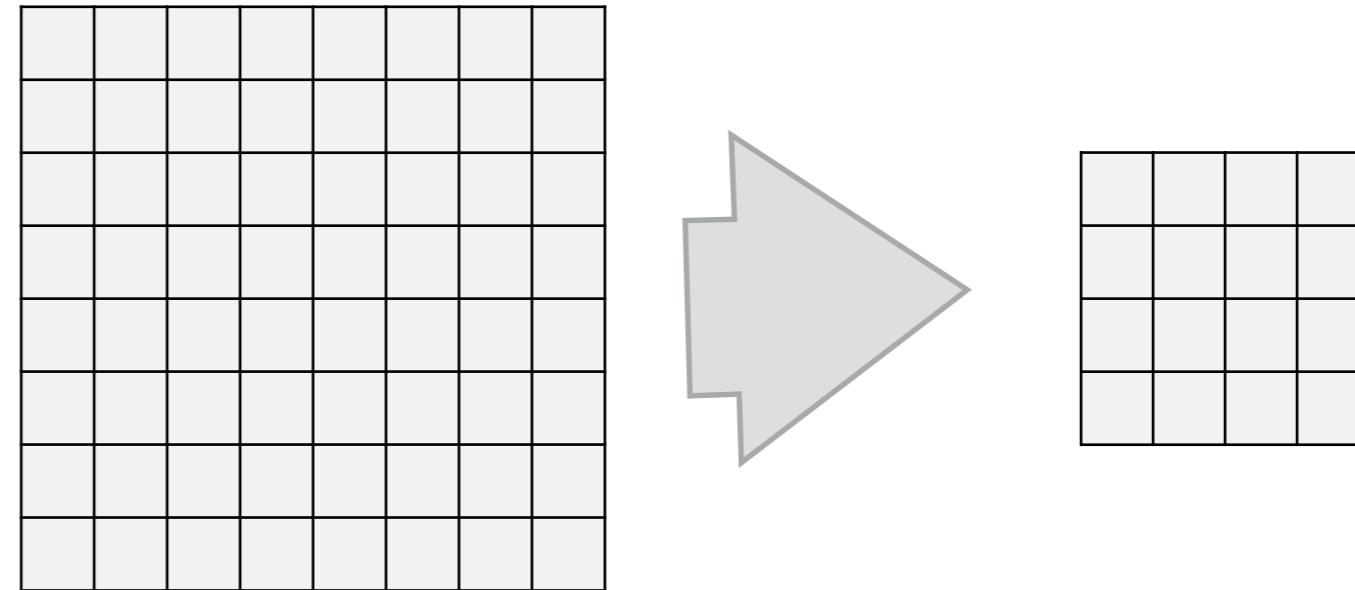


Scale Space
Representation



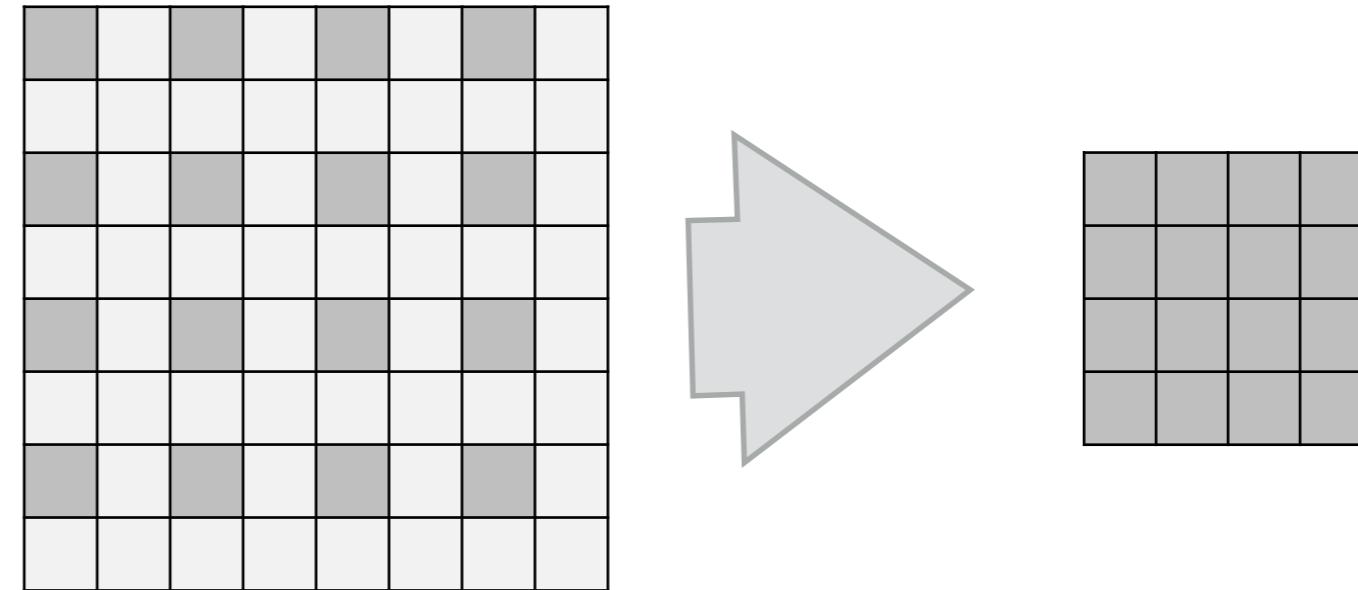
Building a Scale Space Representation

Downsampling:



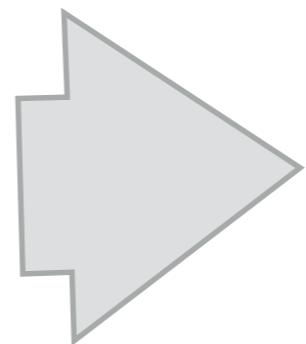
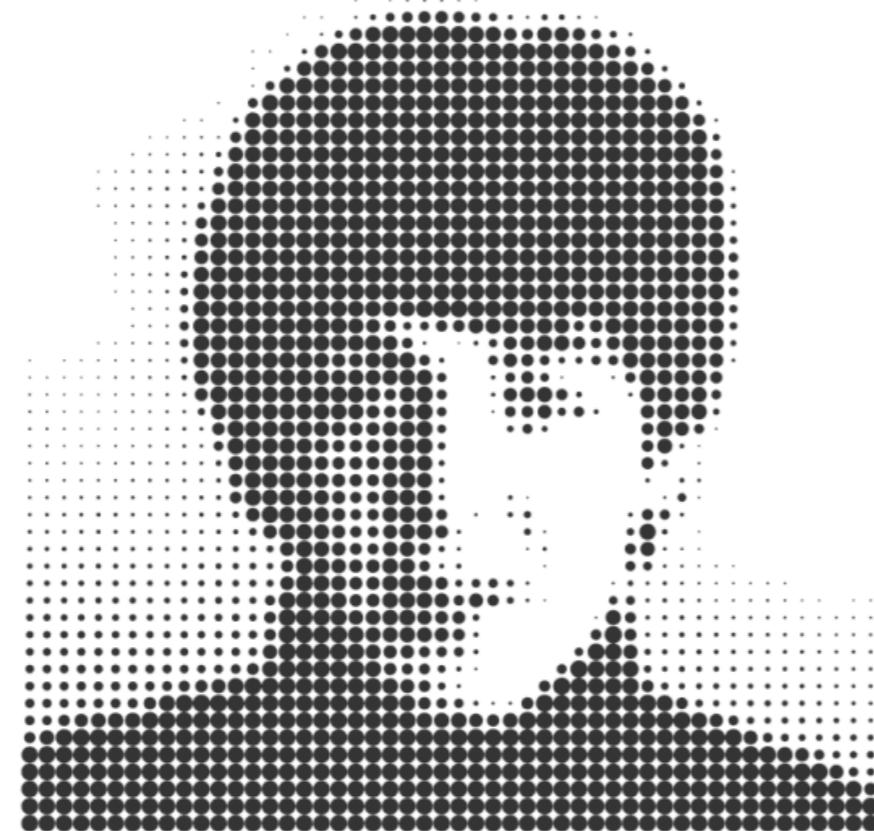
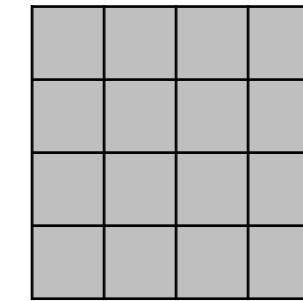
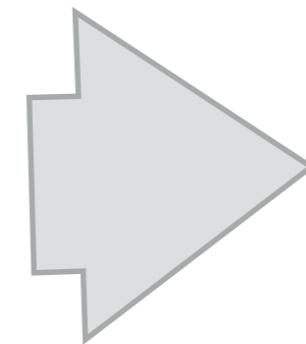
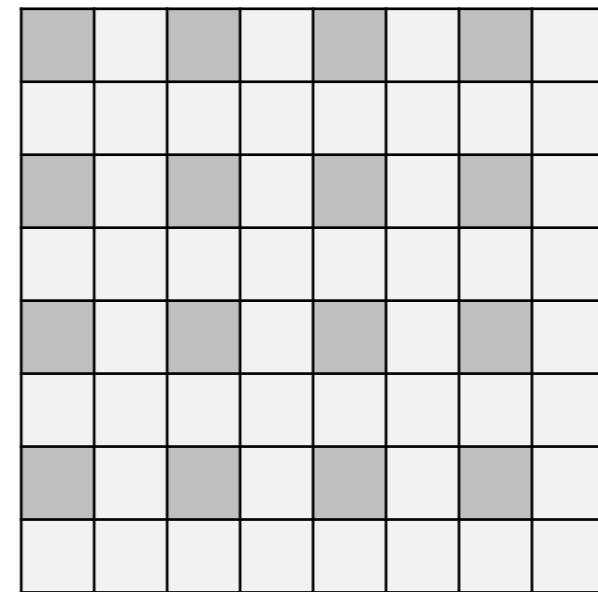
Building a Scale Space Representation

Downsampling:



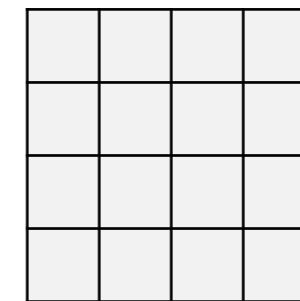
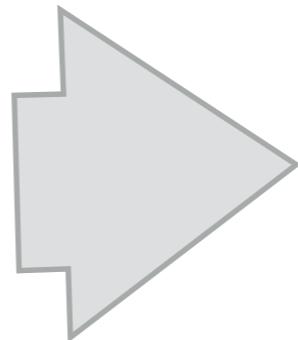
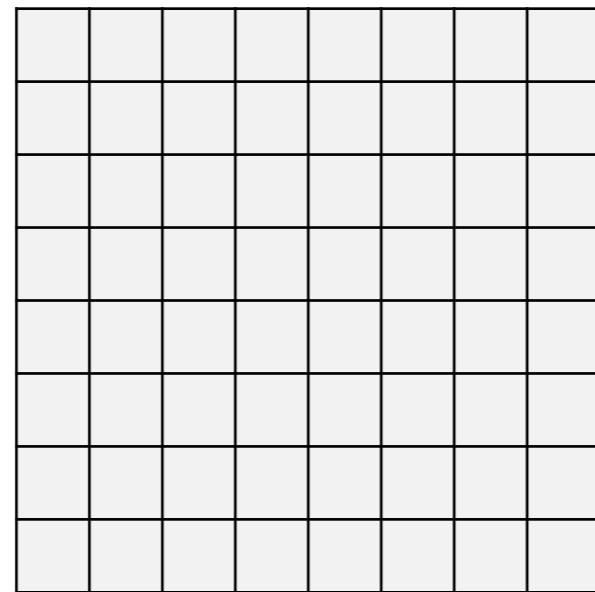
Building a Scale Space Representation

Downsampling:



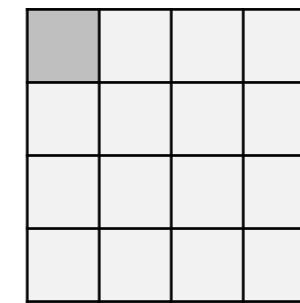
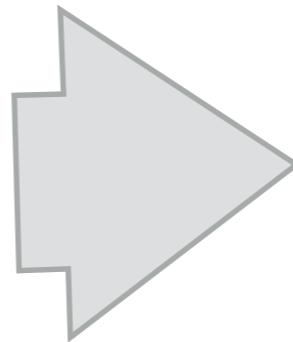
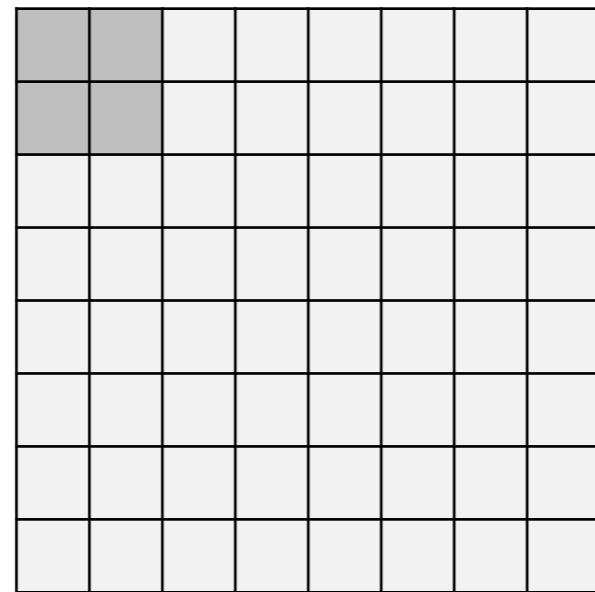
Building a Scale Space Representation

Better: Multiple pixels in larger image contribute to pixel in smaller image

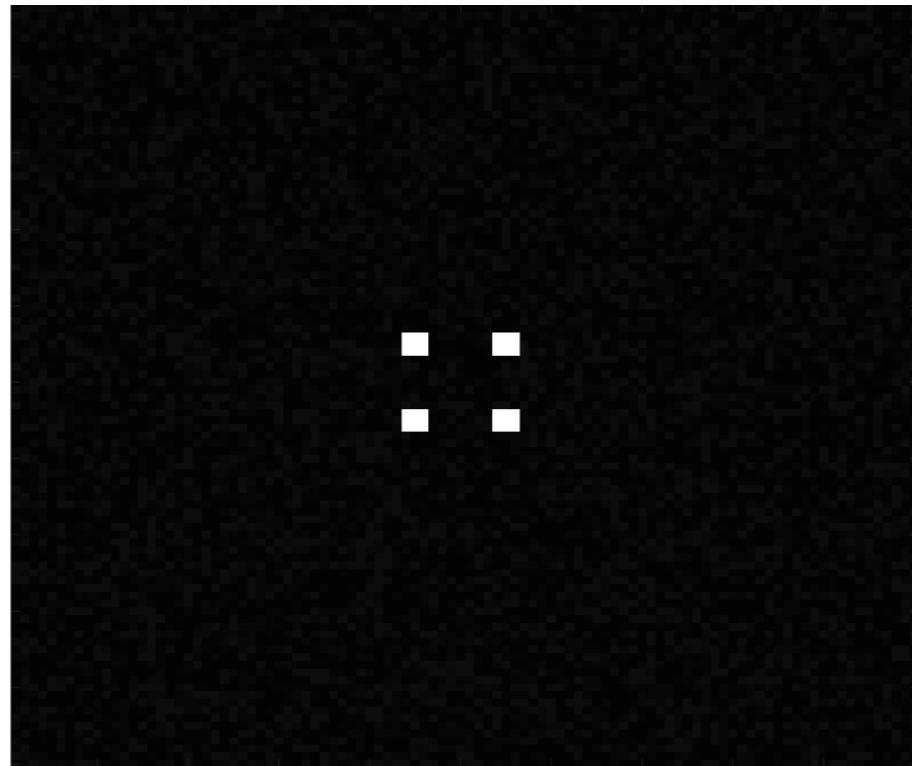


Building a Scale Space Representation

Better: Multiple pixels in larger image contribute to pixel in smaller image



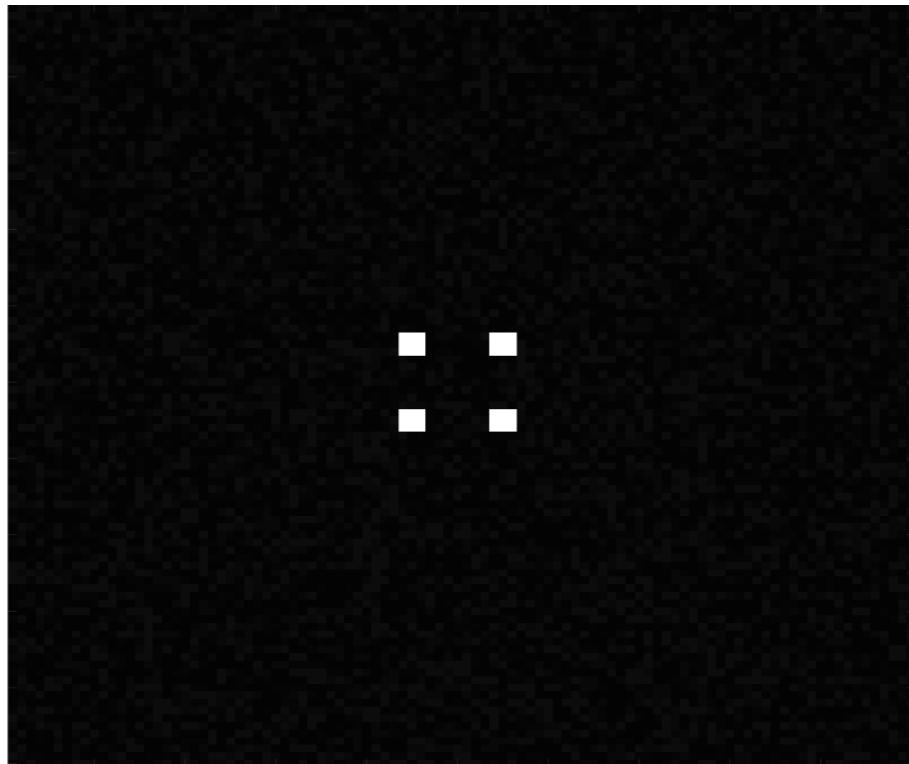
Averaging



$$\ast \quad \frac{1}{25}$$

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

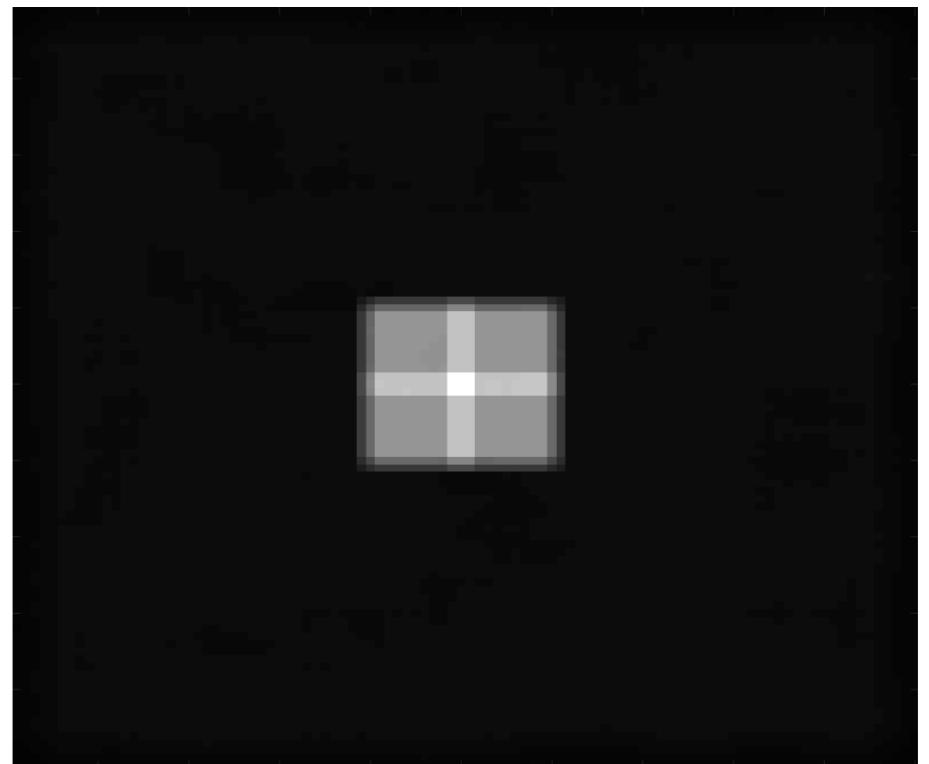
Averaging



$$* \frac{1}{25}$$

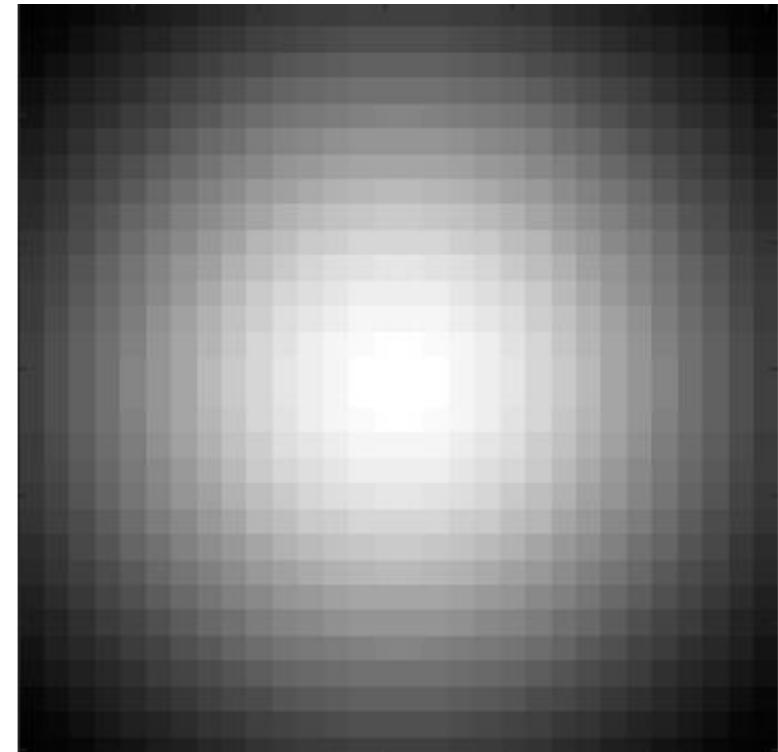
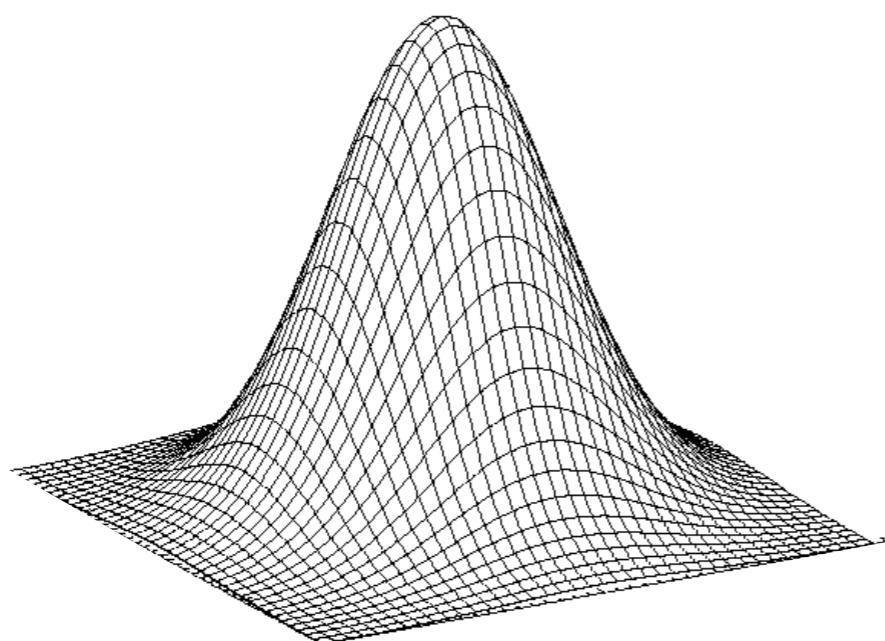
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

=

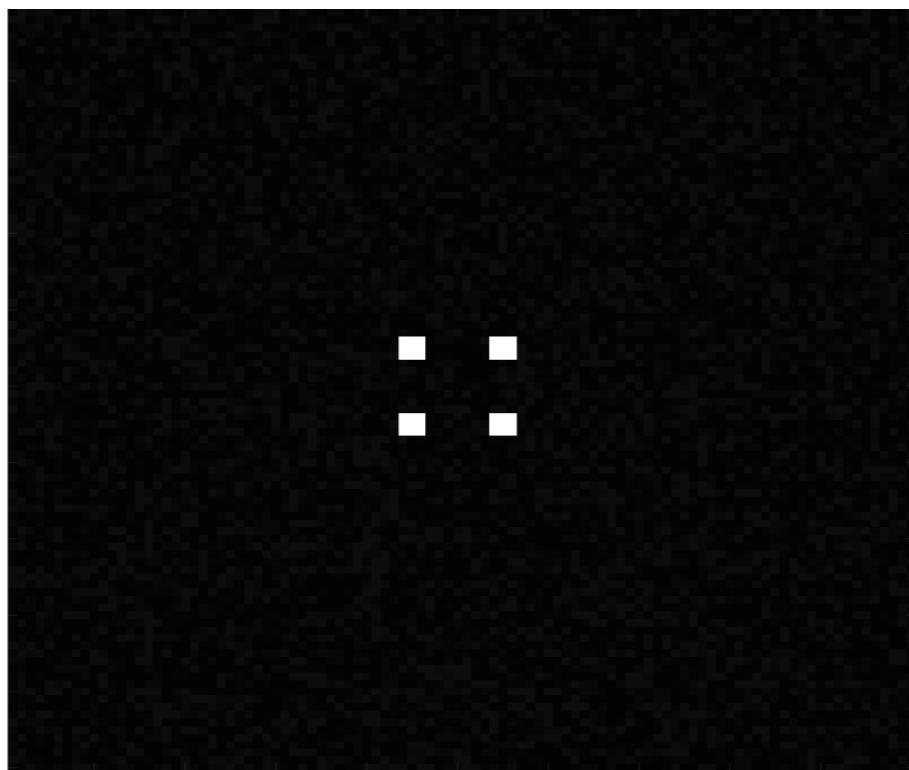


Gaussian Filter

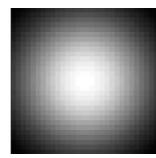
$$G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right)$$



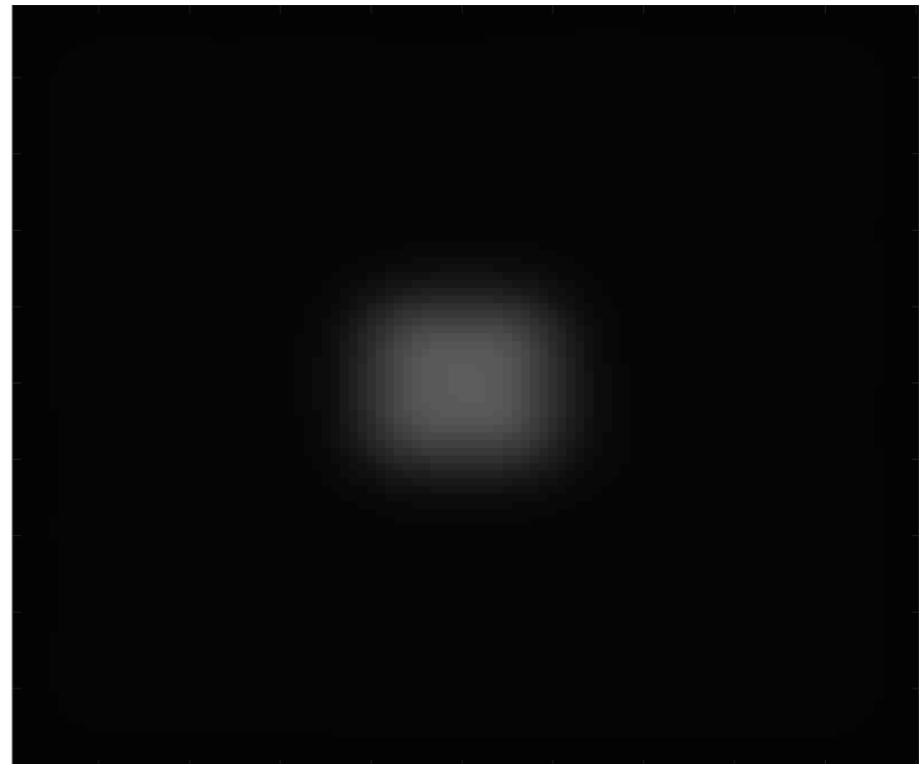
Gaussian Filter



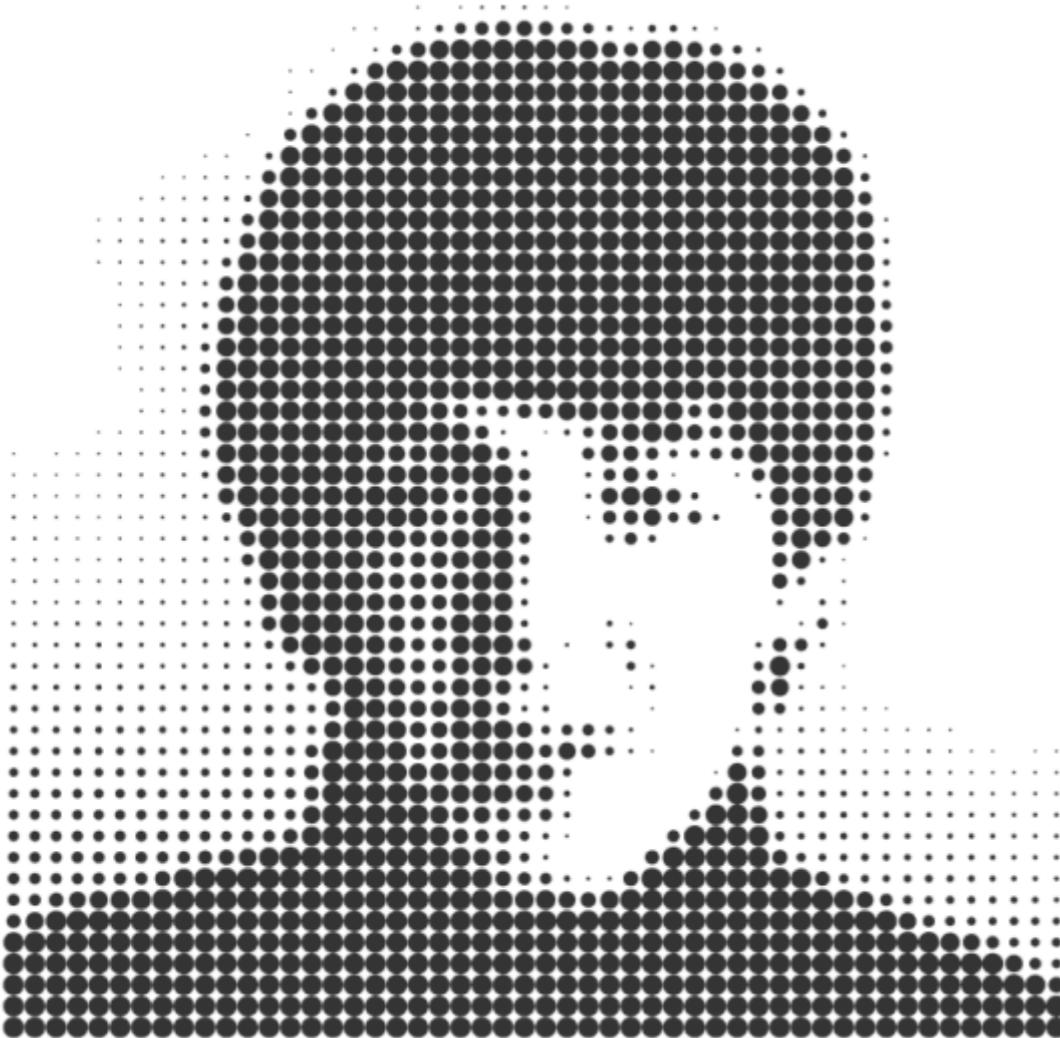
*



=



Scale



Scale



after Gaussian filtering



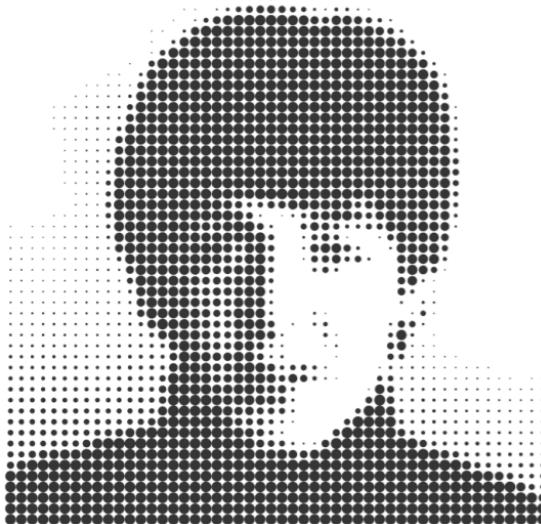
Scale Space Representation

3D space defined by pixel position and blur-level: $L(x, y, \sigma^2)$

Scale Space Representation

3D space defined by pixel position and blur-level: $L(x, y, \sigma^2)$

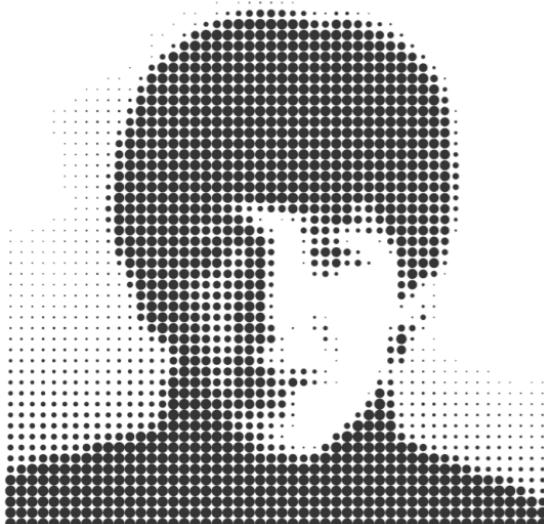
$$L(x, y, 1^2)$$



Scale Space Representation

3D space defined by pixel position and blur-level: $L(x, y, \sigma^2)$

$$L(x, y, 1^2)$$



σ

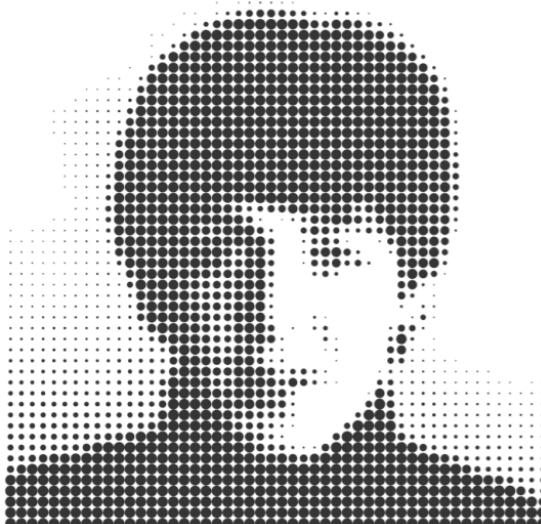


$$L(x, y, 2^2)$$

Scale Space Representation

3D space defined by pixel position and blur-level: $L(x, y, \sigma^2)$

$$L(x, y, 1^2)$$



$$L(x, y, 4^2)$$



σ

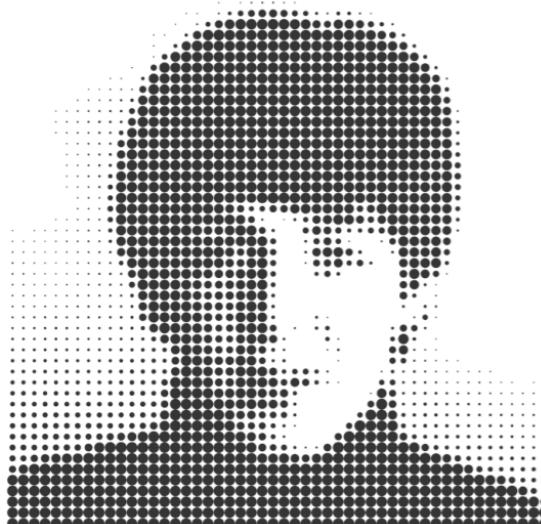


$$L(x, y, 2^2)$$

Scale Space Representation

3D space defined by pixel position and blur-level: $L(x, y, \sigma^2)$

$$L(x, y, 1^2)$$



$$L(x, y, 4^2)$$



σ



$$L(x, y, 2^2)$$

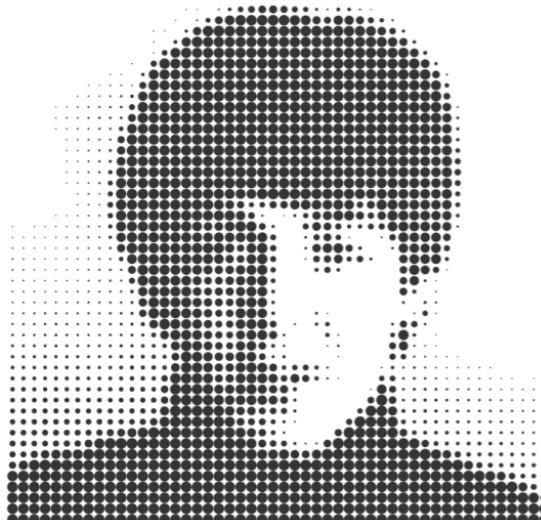


$$L(x, y, 8^2)$$

Scale Space Representation

3D space defined by pixel position and blur-level: $L(x, y, \sigma^2)$

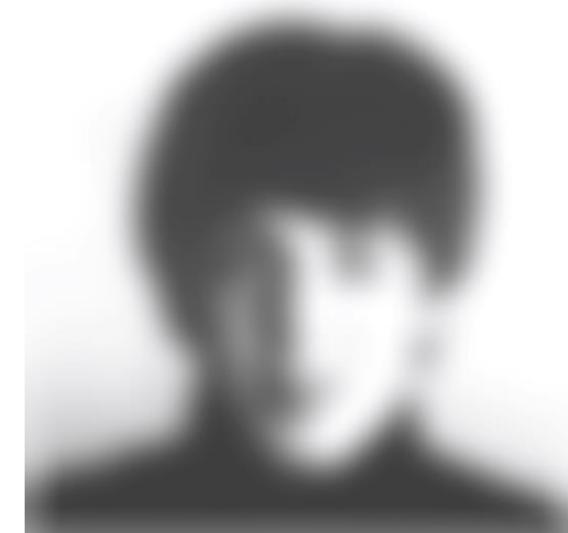
$$L(x, y, 1^2)$$



$$L(x, y, 4^2)$$



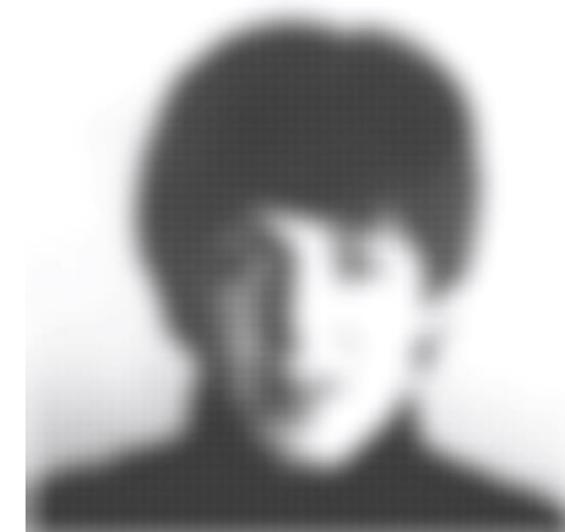
$$L(x, y, 16^2)$$



σ



$$L(x, y, 2^2)$$



$$L(x, y, 8^2)$$

Scale Space Pyramid

σ



$L(x, y, 8^2)$



$L(x, y, 4^2)$

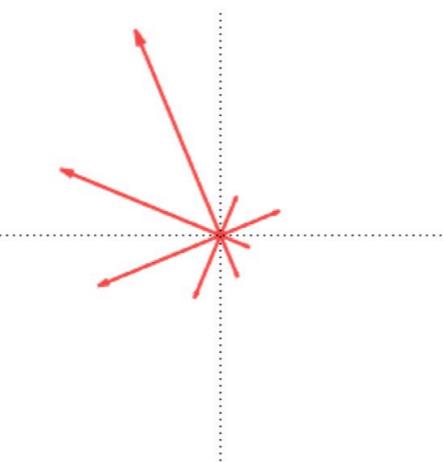
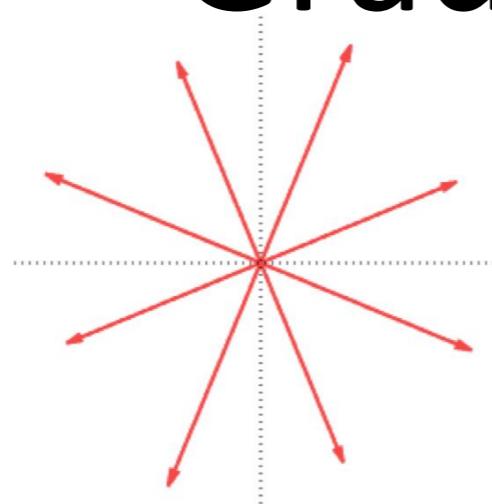


$L(x, y, 2^2)$

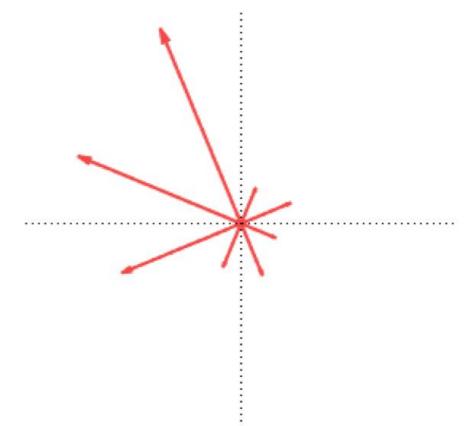
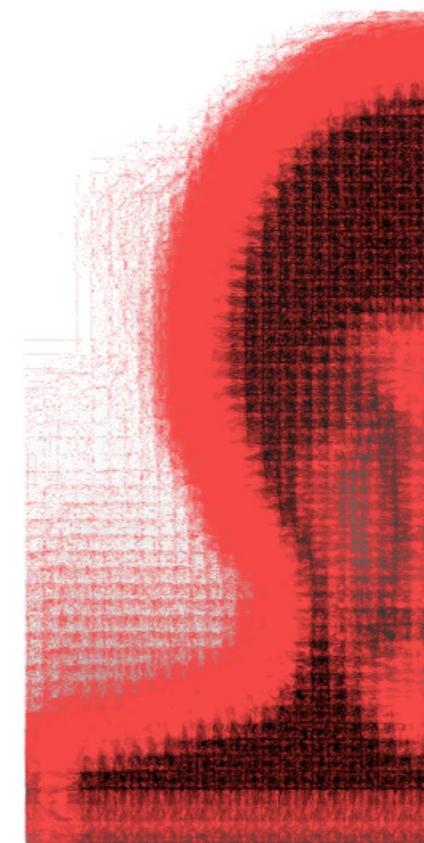
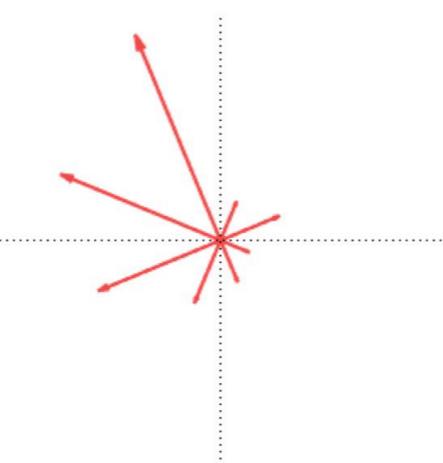
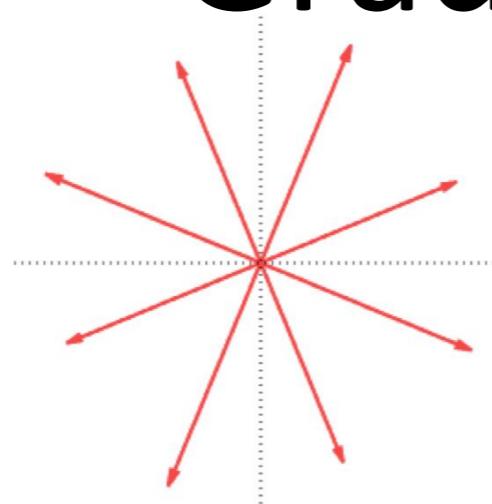


$L(x, y, 1^2)$

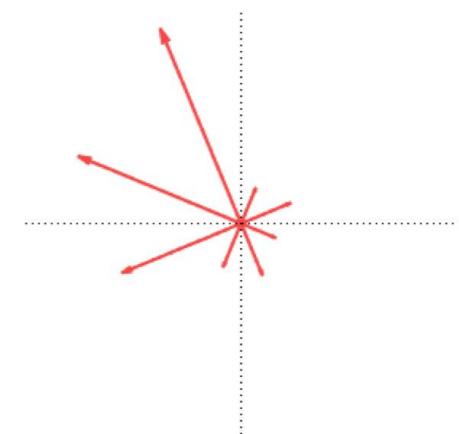
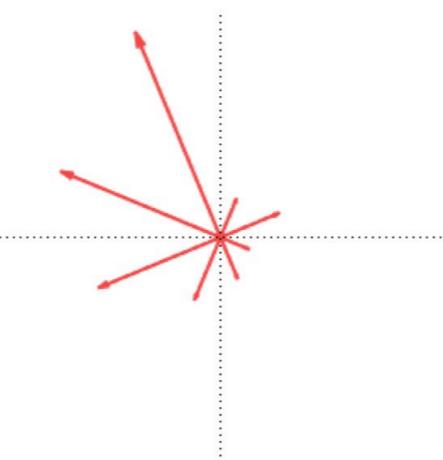
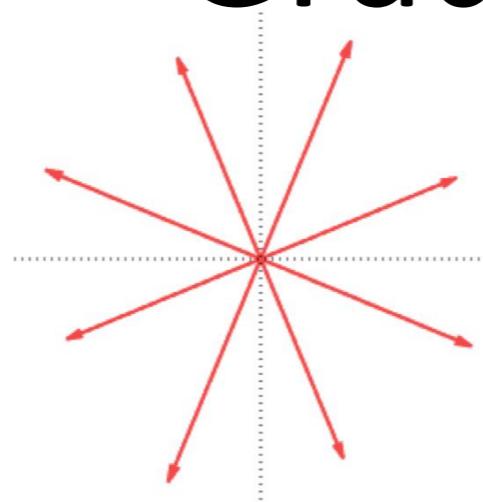
Gradient Histograms



Gradient Histograms

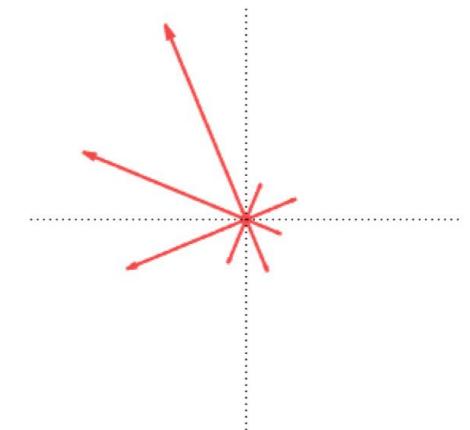
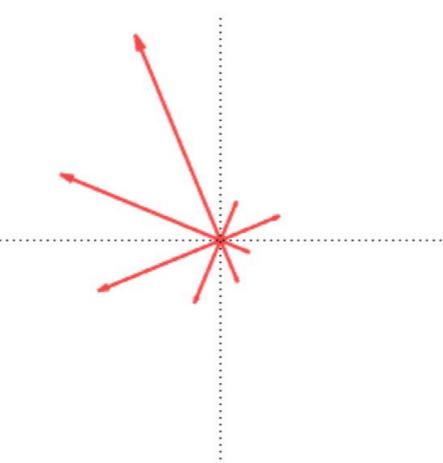
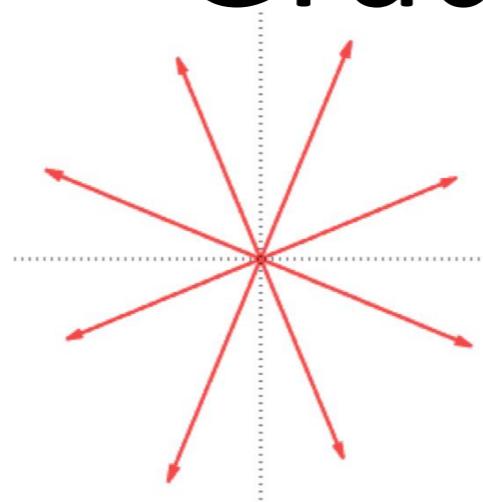


Gradient Histograms



Much more similar!

Gradient Histograms



Much more similar!

... but how do we determine the correct scale?

Estimating Scale

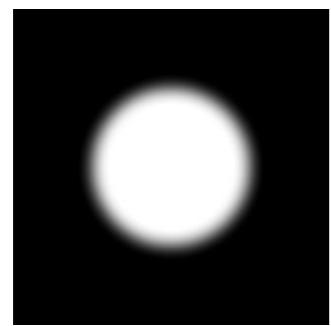


Query $r = 67$

Templates:

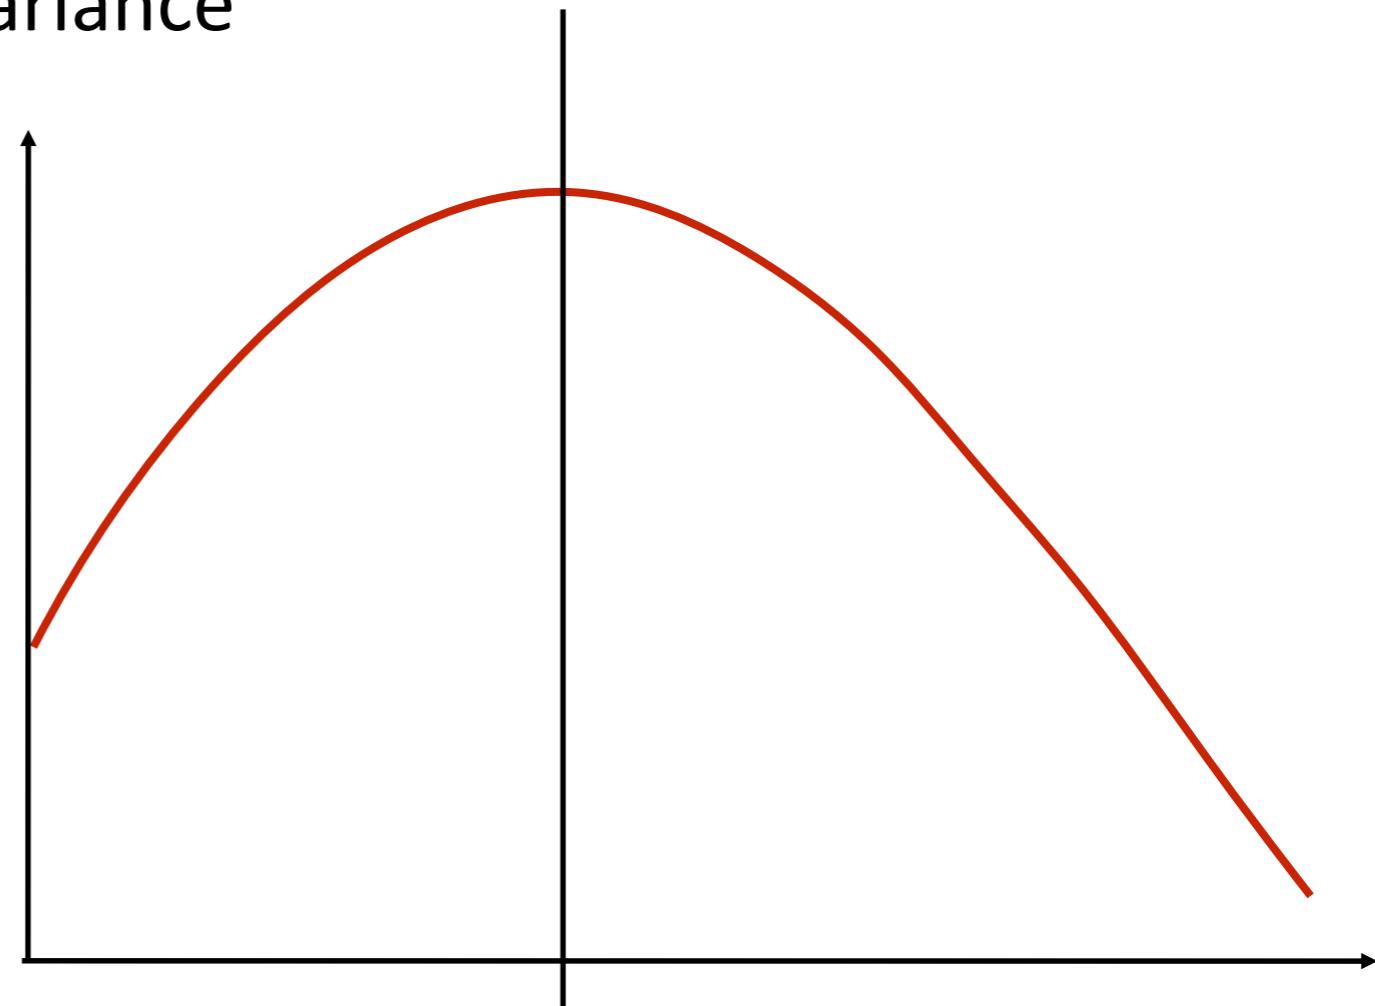


Estimating Scale



Query $r = 67$

Covariance

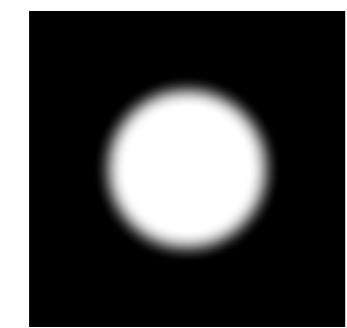


61

Templates:

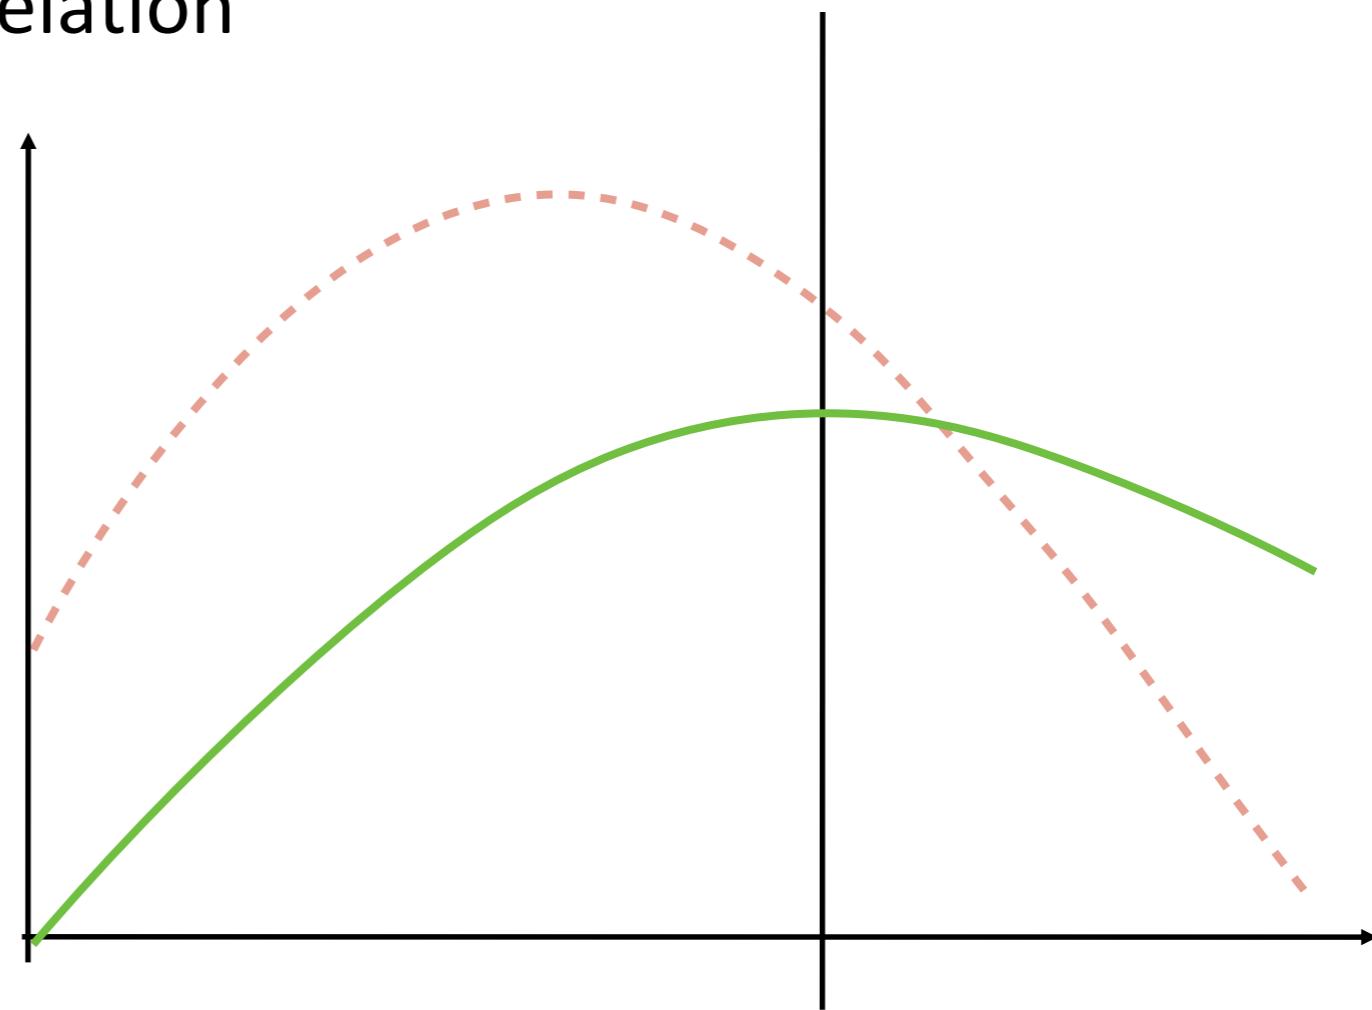


Estimating Scale

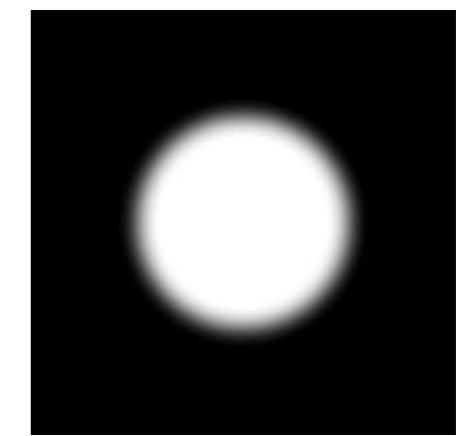
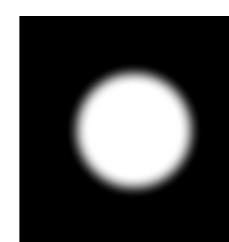


Query. $r = 67$

Correlation



Templates:



67

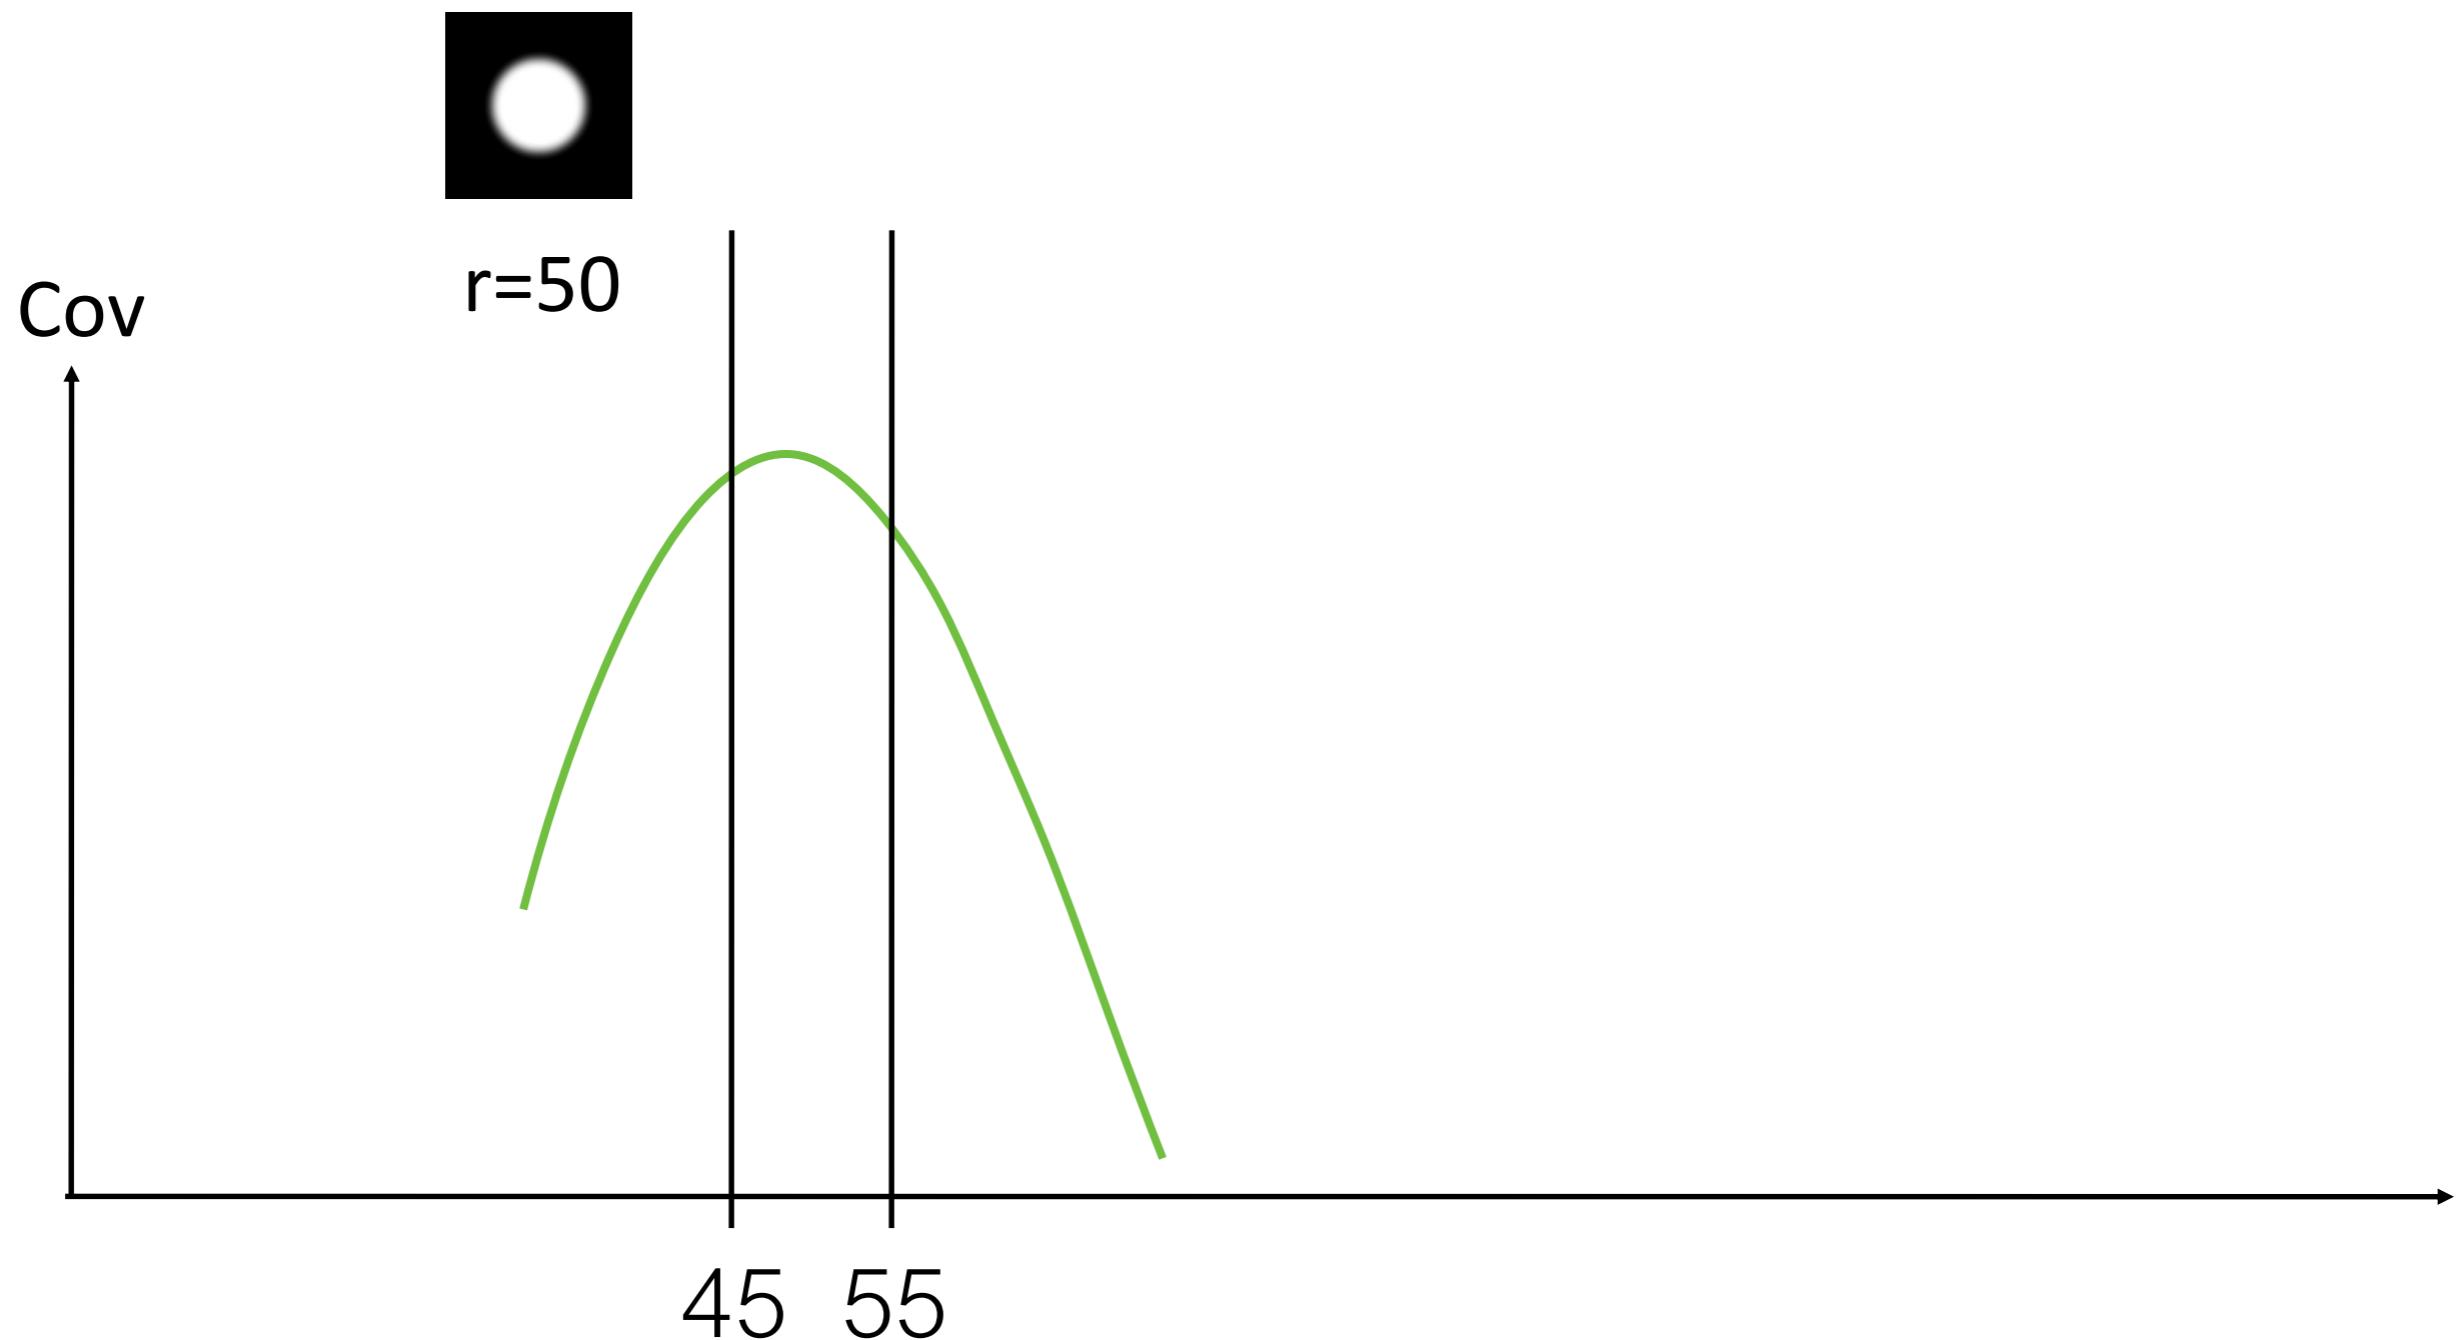
Estimating Relative Scale

In practice: Discretized set of templates



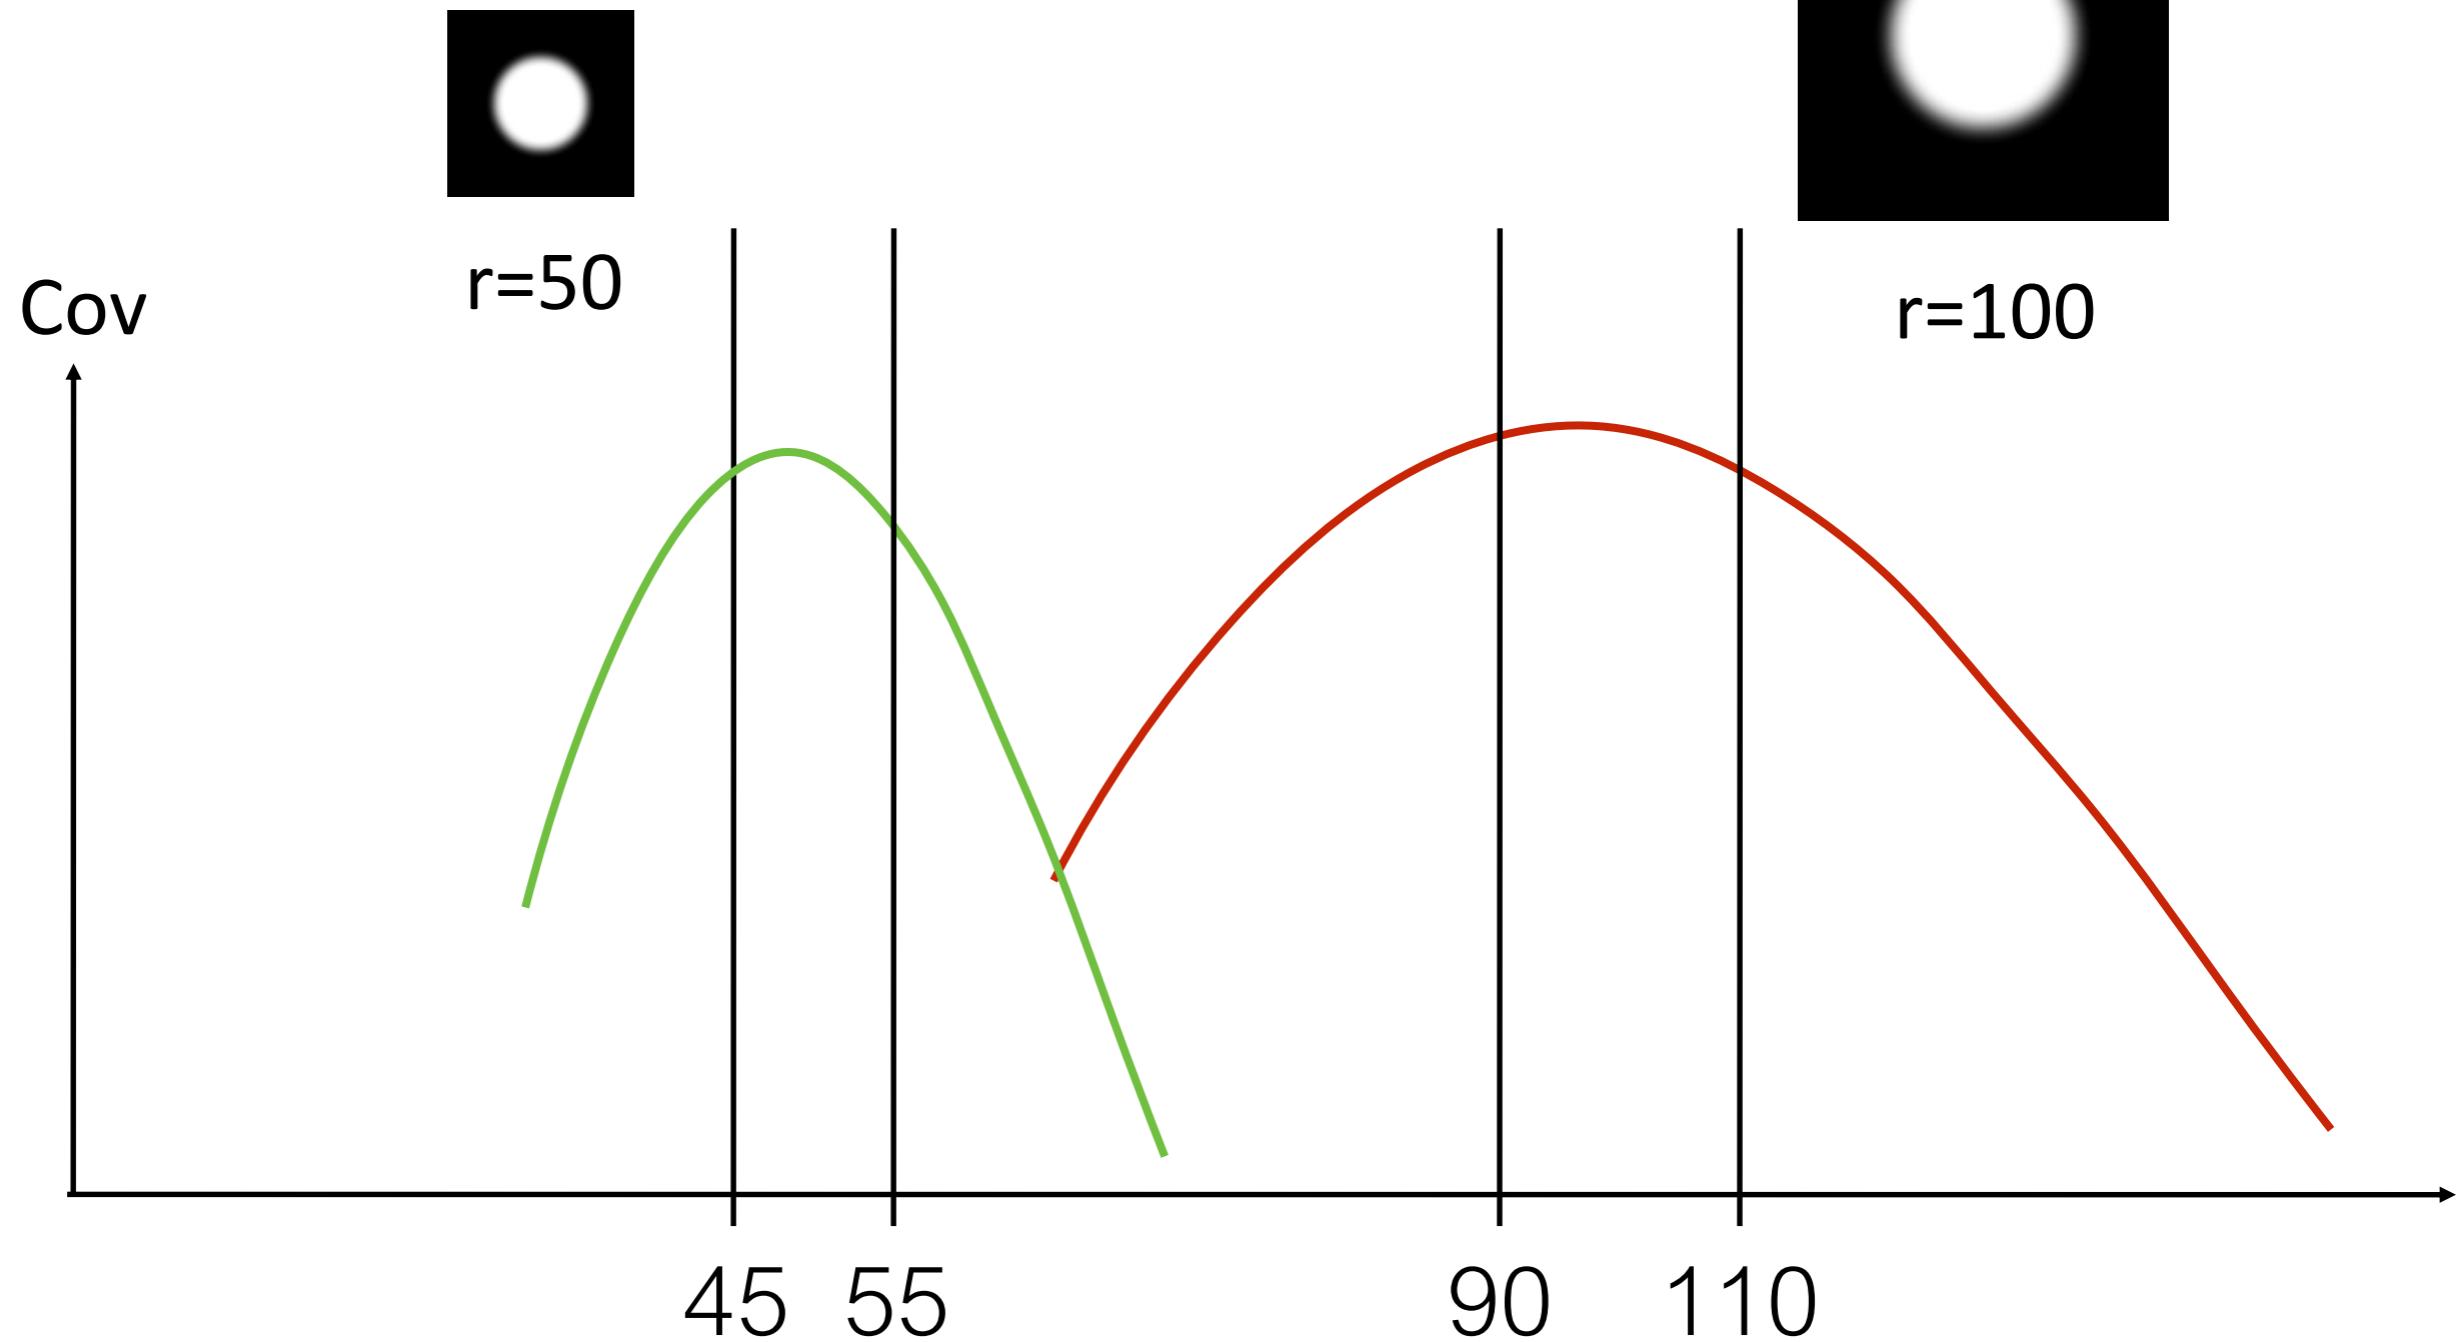
Estimating Relative Scale

In practice: Discretized set of templates



Estimating Relative Scale

In practice: Discretized set of templates

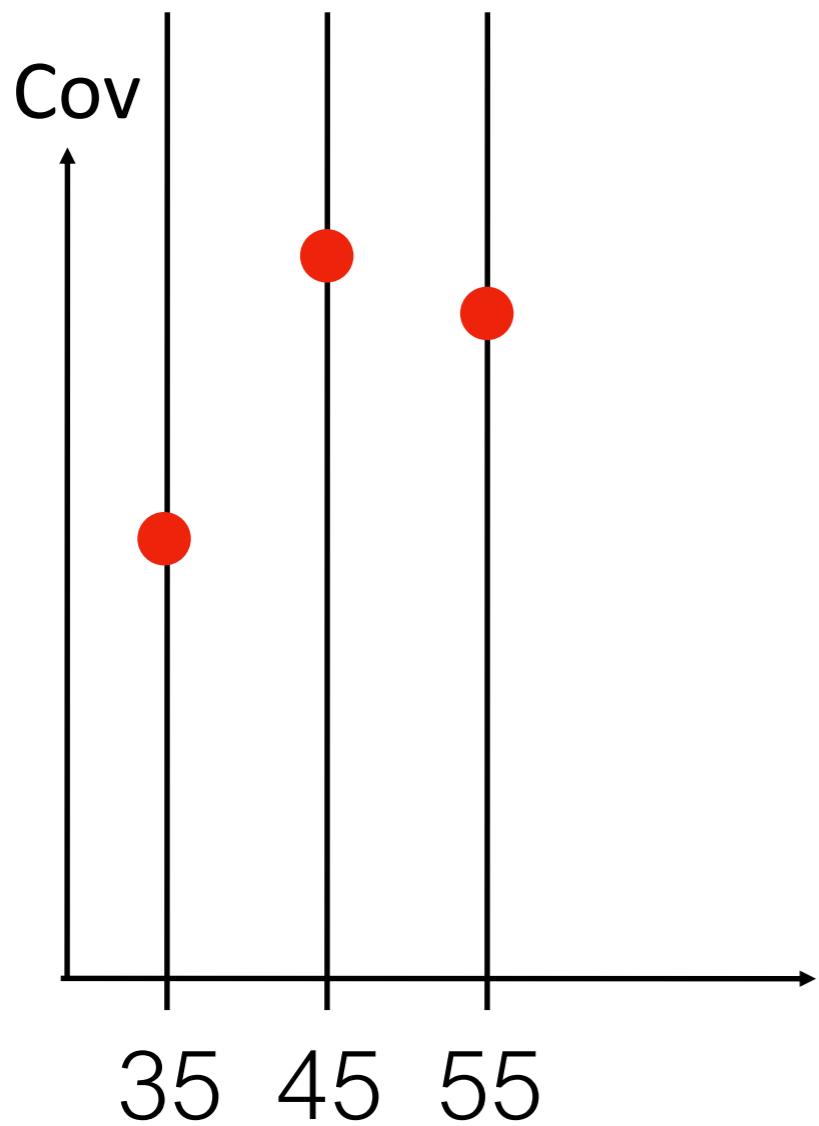


“Sub-Scale” Precision

How to determine scale more accurately?

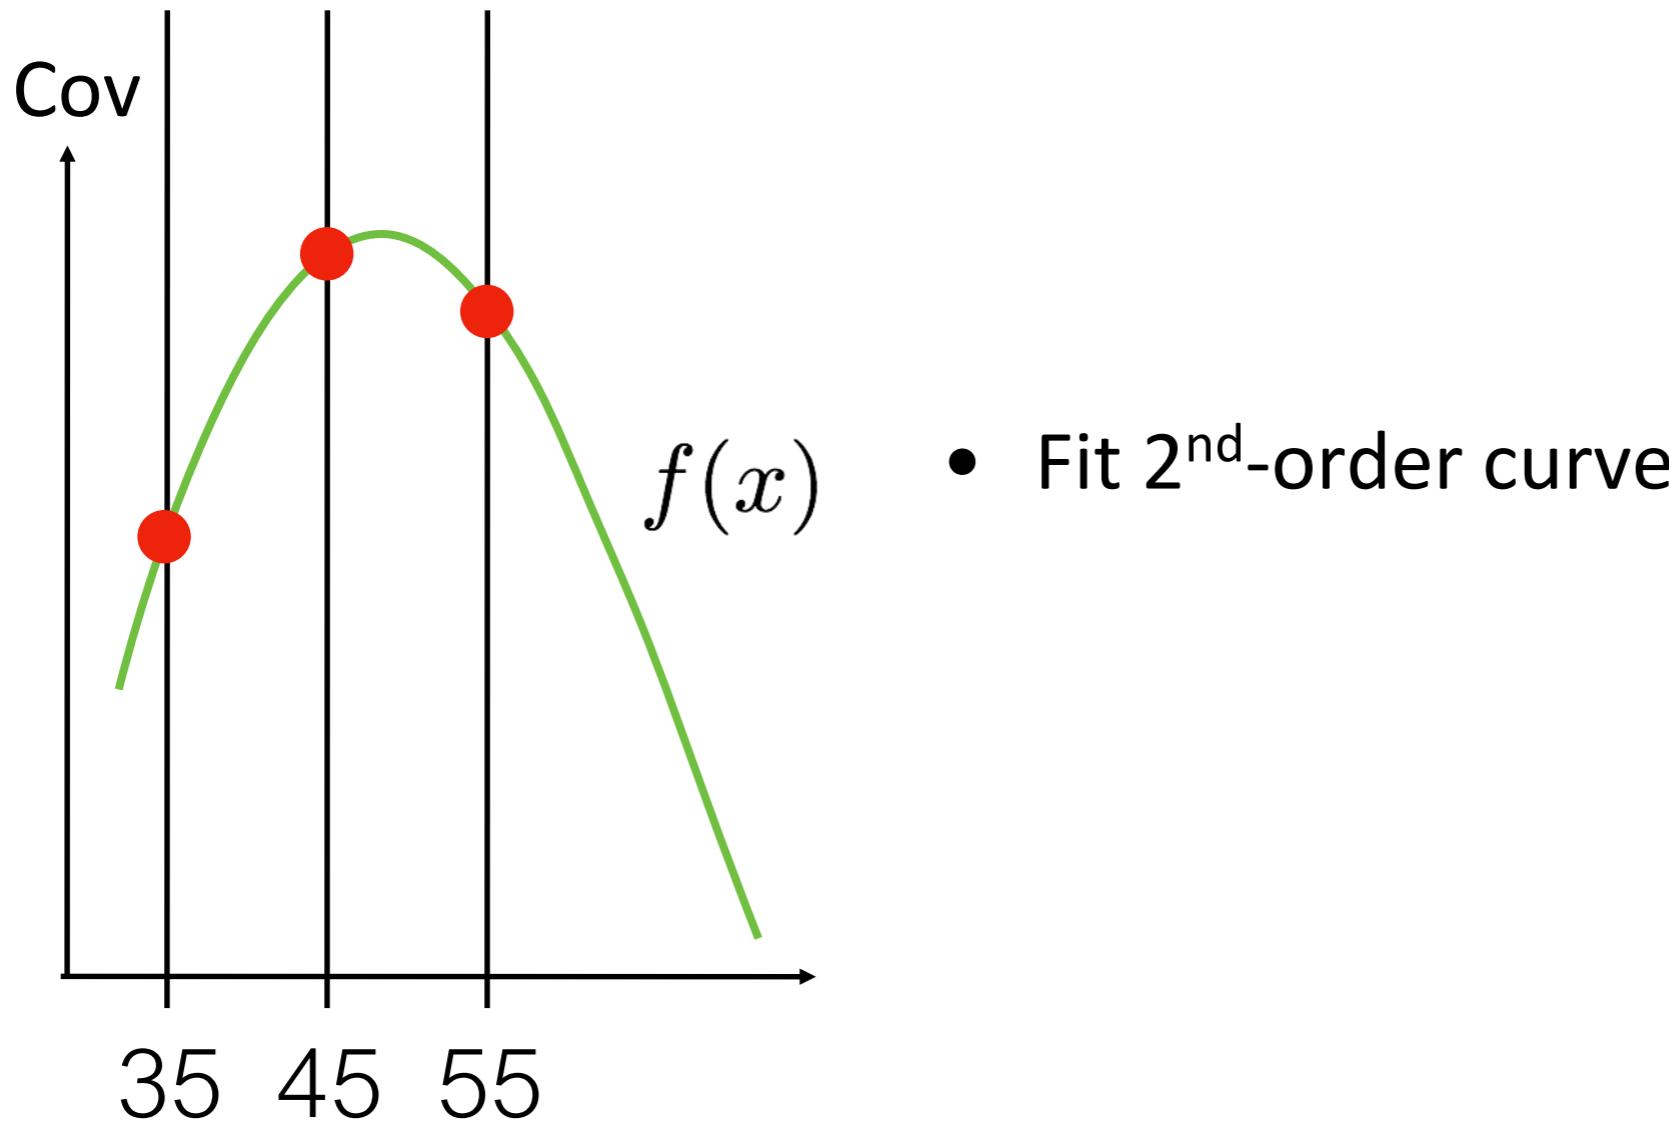
“Sub-Scale” Precision

How to determine scale more accurately?



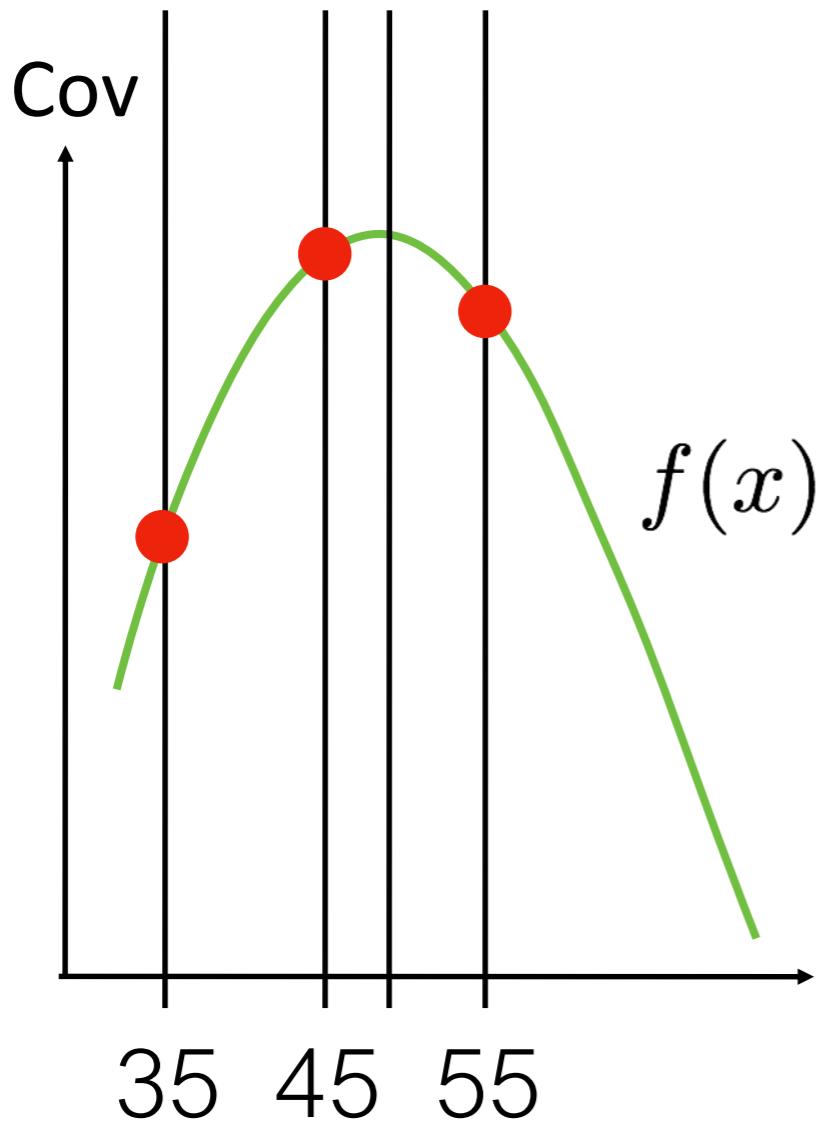
“Sub-Scale” Precision

How to determine scale more accurately?



“Sub-Scale” Precision

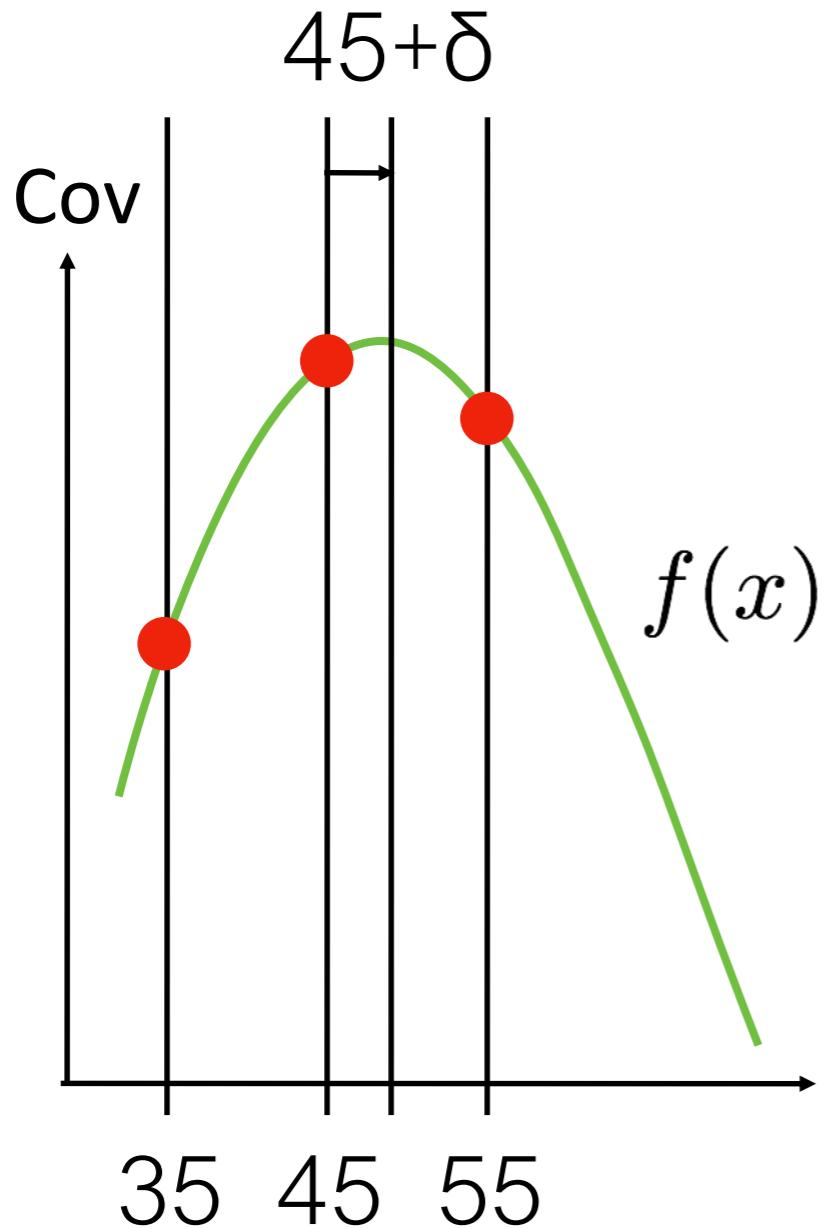
How to determine scale more accurately?



- Fit 2nd-order curve
- Determine position of maximum

“Sub-Scale” Precision

How to determine scale more accurately?



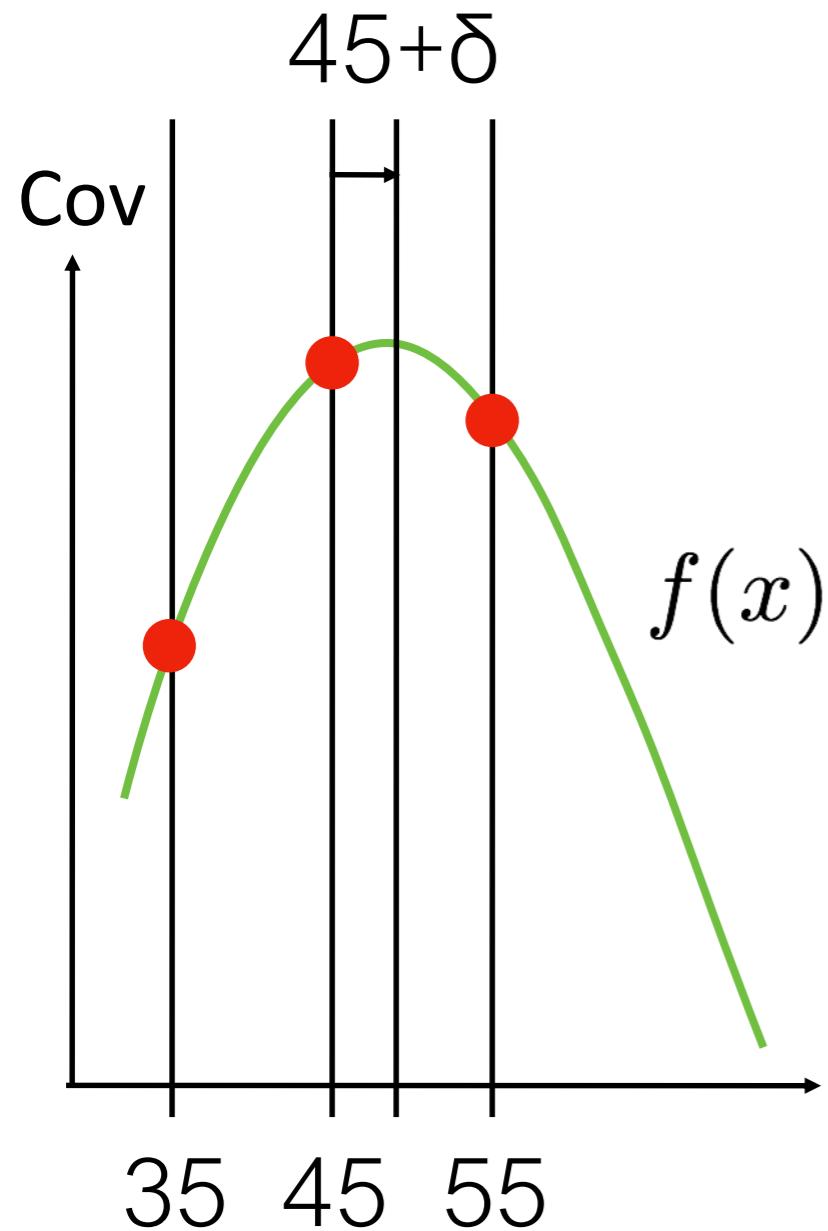
- Fit 2nd-order curve
- Determine position of maximum
- Obtain refinement of scale

“Sub-Scale” Precision

Curve fitting via 2nd-order Taylor expansion

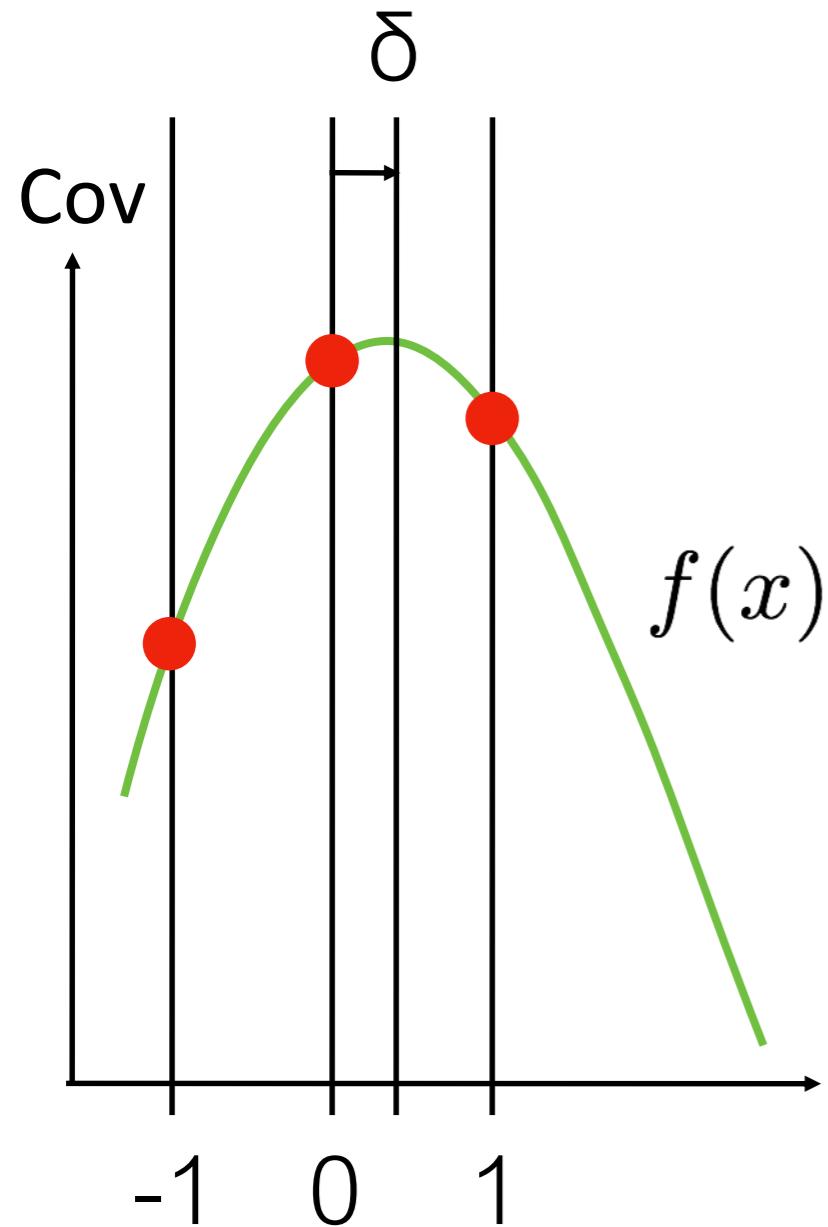
“Sub-Scale” Precision

Curve fitting via 2nd-order Taylor expansion



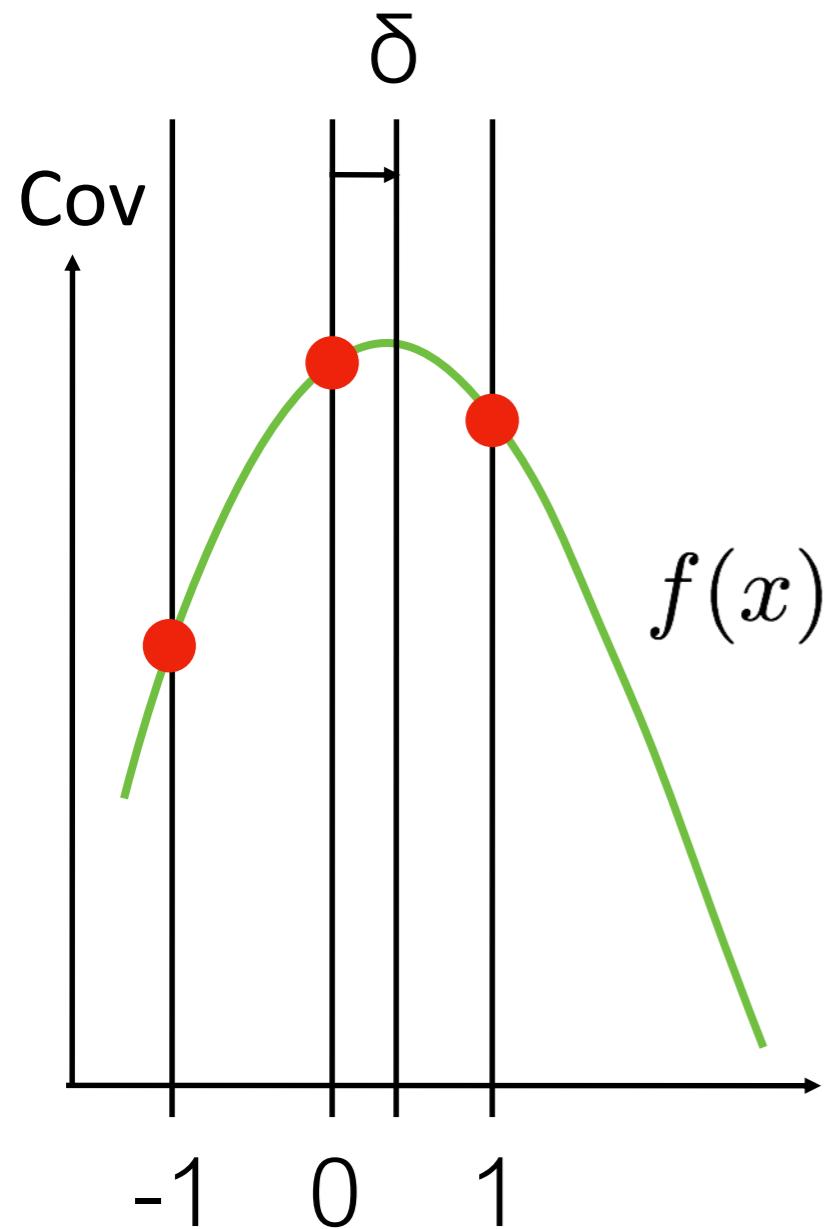
“Sub-Scale” Precision

Curve fitting via 2nd-order Taylor expansion



“Sub-Scale” Precision

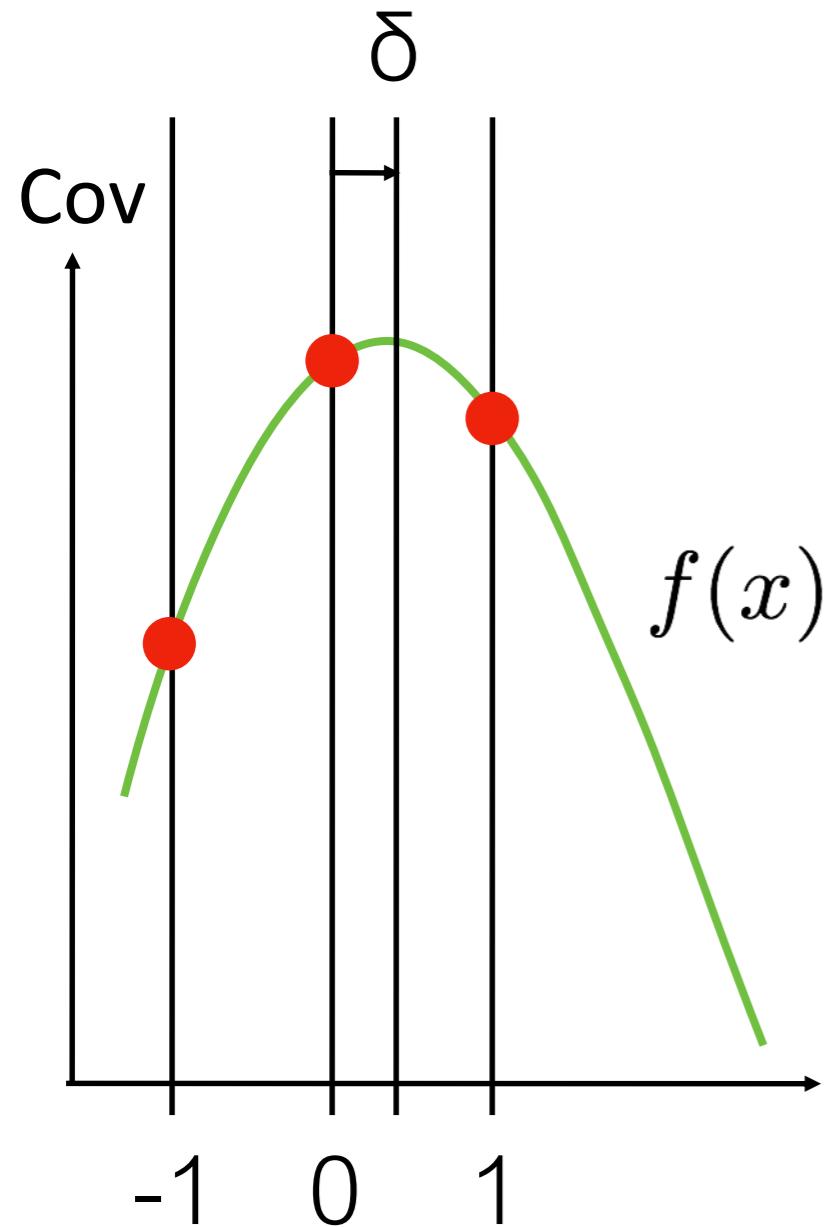
Curve fitting via 2nd-order Taylor expansion



$$f(\delta) = f(0) + \delta f'(0) + \frac{\delta^2}{2} f''(0)$$

“Sub-Scale” Precision

Curve fitting via 2nd-order Taylor expansion

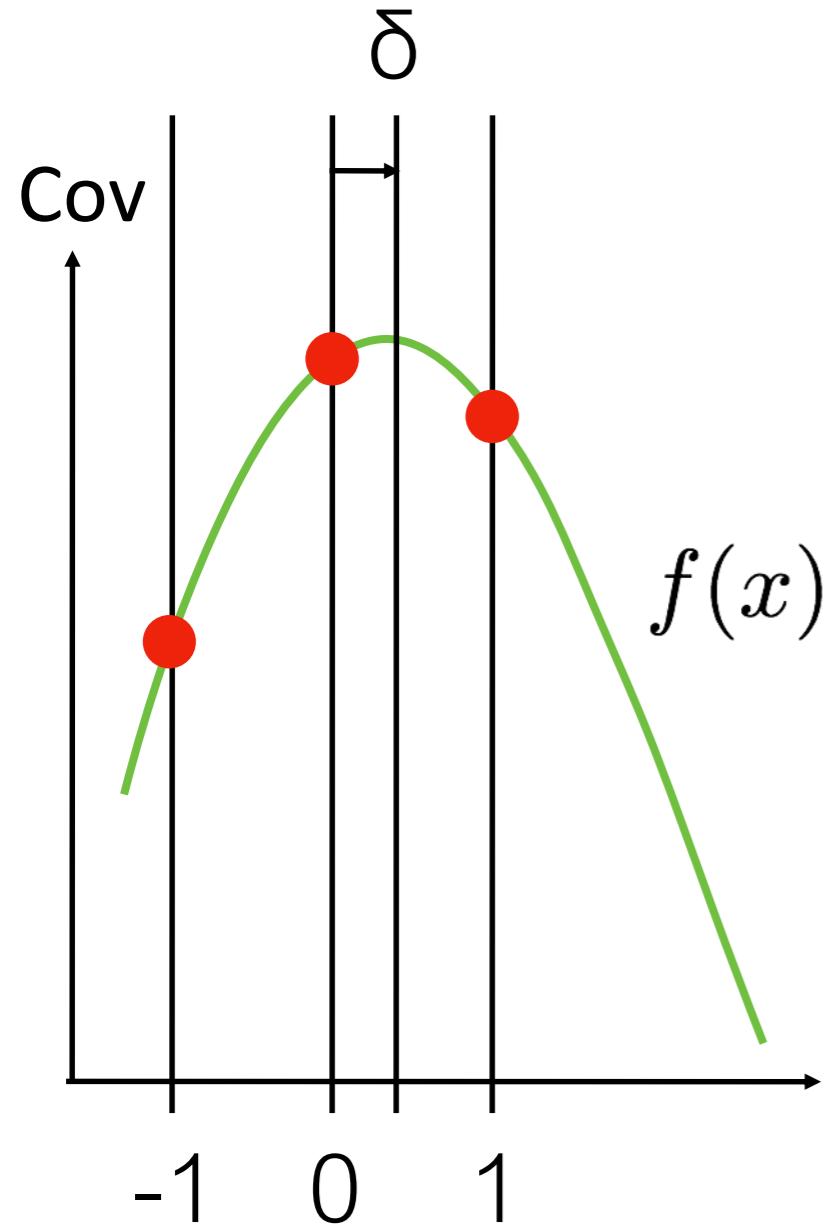


$$f(\delta) = f(0) + \delta f'(0) + \frac{\delta^2}{2} f''(0)$$

$$f'(0) = \frac{f(1) - f(-1)}{2} = a$$

“Sub-Scale” Precision

Curve fitting via 2nd-order Taylor expansion



$$f(\delta) = f(0) + \delta f'(0) + \frac{\delta^2}{2} f''(0)$$

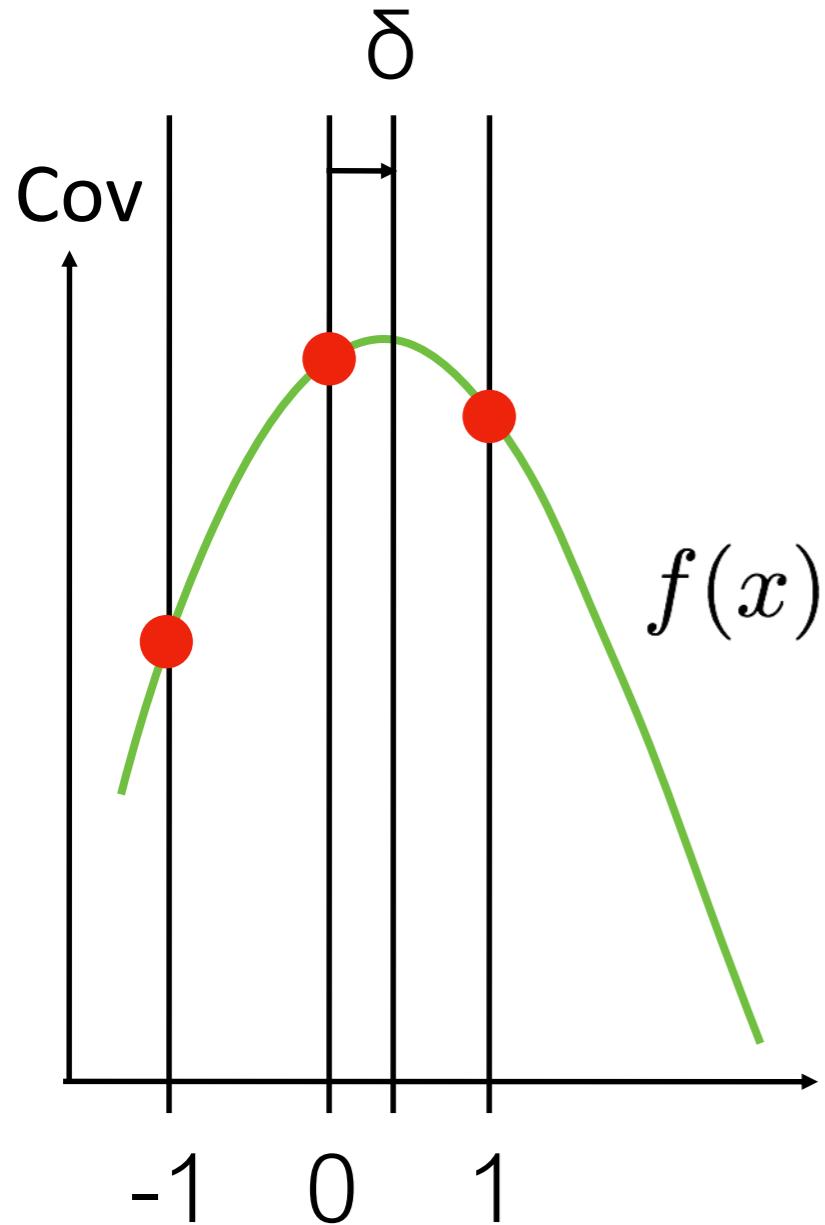
$$f'(0) = \frac{f(1) - f(-1)}{2} = a$$

$$f''(0) = f(1) - 2f(0) + f(-1) = b$$

discrete Laplacian

“Sub-Scale” Precision

Curve fitting via 2nd-order Taylor expansion



$$f(\delta) = f(0) + \delta f'(0) + \frac{\delta^2}{2} f''(0)$$

$$f'(0) = \frac{f(1) - f(-1)}{2} = a$$

$$f''(0) = f(1) - 2f(0) + f(-1) = b$$

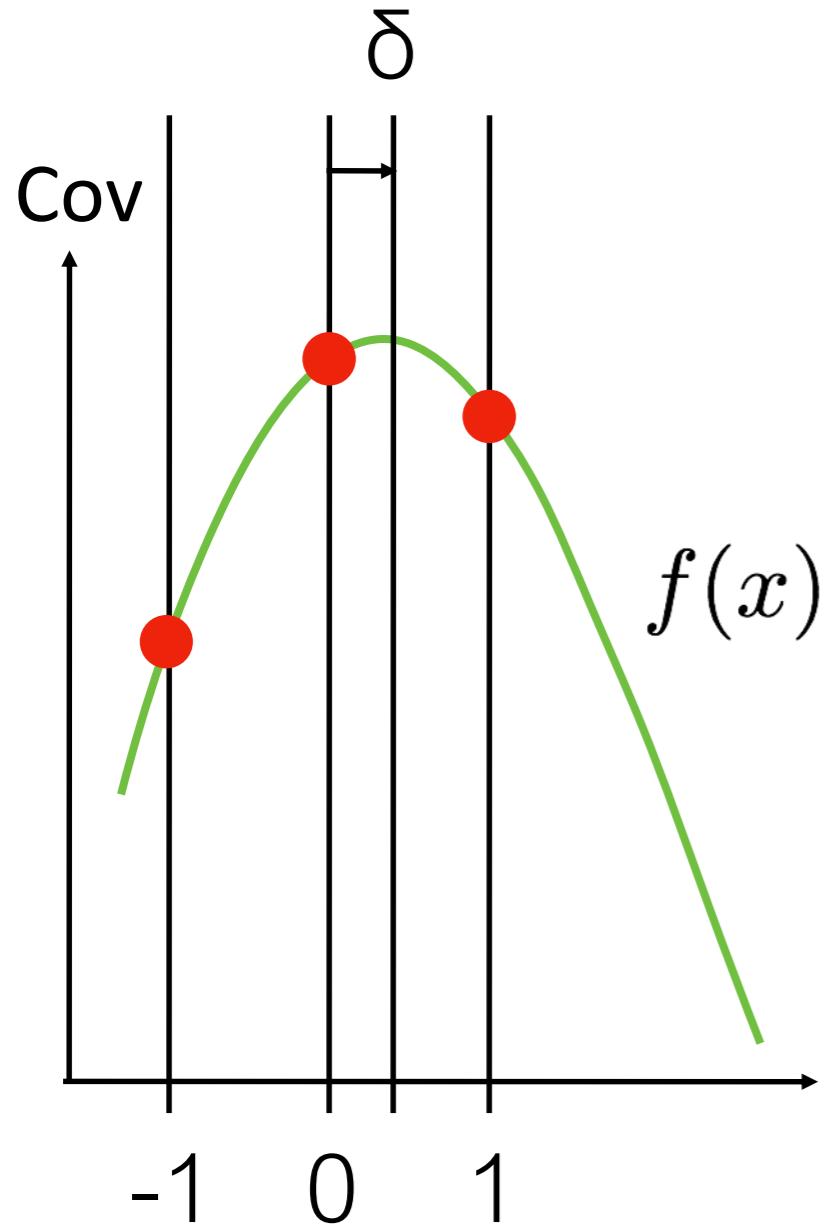
discrete Laplacian

Compute maximum:

$$f'(\delta) = a + b\delta$$

“Sub-Scale” Precision

Curve fitting via 2nd-order Taylor expansion



$$f(\delta) = f(0) + \delta f'(0) + \frac{\delta^2}{2} f''(0)$$

$$f'(0) = \frac{f(1) - f(-1)}{2} = a$$

$$f''(0) = f(1) - 2f(0) + f(-1) = b$$

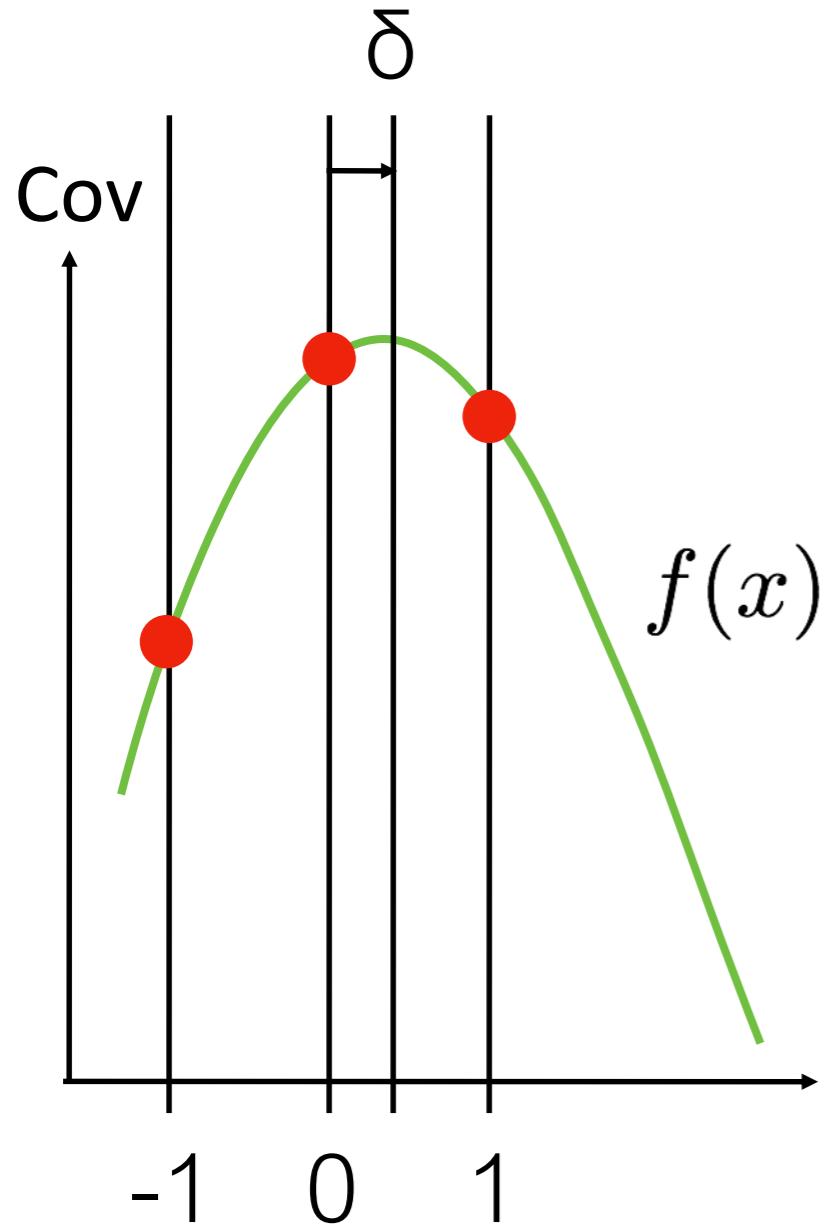
discrete Laplacian

Compute maximum:

$$f'(\delta) = a + b\delta \stackrel{!}{=} 0$$

“Sub-Scale” Precision

Curve fitting via 2nd-order Taylor expansion



$$f(\delta) = f(0) + \delta f'(0) + \frac{\delta^2}{2} f''(0)$$

$$f'(0) = \frac{f(1) - f(-1)}{2} = a$$

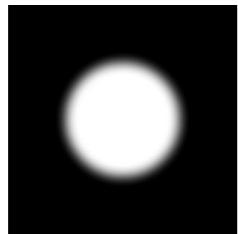
$$f''(0) = f(1) - 2f(0) + f(-1) = b$$

discrete Laplacian

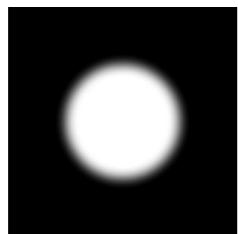
Compute maximum:

$$f'(\delta) = a + b\delta \stackrel{!}{=} 0 \Rightarrow \delta = \frac{-a}{b}$$

Detect Scale and Position



Detect Scale and Position

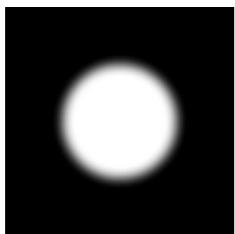


1	0	1	2	3	4	2
2	7	8	2	3	3	3
2	3	1	4	5	6	2
2	7	12	15	20	7	2
2	8	12	25	19	11	3
3	9	11	18	16	7	1
1	2	3	2	1	3	1

2	1	2	2	2	3	3
1	2	3	2	2	3	4
5	11	9	8	9	8	1
1	11	15	20	22	9	5
2	12	16	35	22	11	4
3	9	15	18	21	7	1
2	2	9	8	7	9	1

3	3	3	3	3	3	3
3	4	4	4	4	4	3
3	5	6	6	6	5	2
2	7	12	12	12	7	2
2	5	12	20	13	11	3
3	9	12	13	13	9	1
1	2	5	5	5	3	1

Detect Scale and Position



1	0	1	2	3	4	2
2	7	8	2	3	3	3
2	3	1	4	5	6	2
2	7	12	15	20	7	2
2	8	12	25	19	11	3
3	9	11	18	16	7	1
1	2	3	2	1	3	1

2	1	2	2	2	3	3
1	2	3	2	2	3	4
5	11	9	8	9	8	1
1	11	15	20	22	9	5
2	12	16	35	22	11	4
3	9	15	18	21	7	1
2	2	9	8	7	9	1

3	3	3	3	3	3	3
3	4	4	4	4	4	3
3	5	6	6	6	5	2
2	7	12	12	12	7	2
2	5	12	20	13	11	3
3	9	12	13	13	9	1
1	2	5	5	5	3	1

Larger than neighbors in image and scale!

Sub-Pixel Refinement

2nd order curve fitting (1D):

$$f(x) \approx f(0) + x \cdot f'(0) + \frac{x^2}{2} f''(0)$$

Sub-Pixel Refinement

2nd order curve fitting (1D):

$$f(x) \approx f(0) + x \cdot f'(0) + \frac{x^2}{2} f''(0)$$

2nd order surface fitting (3D):

$$f(x, y, \sigma) \approx f(0, 0, 0) + (x \quad y \quad \sigma) \cdot \nabla f(0, 0, 0) + \frac{1}{2} (x \quad y \quad \sigma) \mathbf{H}(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

Sub-Pixel Refinement

2nd order curve fitting (1D):

$$f(x) \approx f(0) + x \cdot f'(0) + \frac{x^2}{2} f''(0)$$

2nd order surface fitting (3D):

$$f(x, y, \sigma) \approx f(0, 0, 0) + (x \quad y \quad \sigma) \cdot \nabla f(0, 0, 0) + \frac{1}{2} (x \quad y \quad \sigma) \mathbb{H}(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

Hessian matrix:

$$\mathbb{H}(0, 0, 0) = \begin{pmatrix} f''_{xx}(0, 0, 0) & f''_{xy}(0, 0, 0) & f''_{x\sigma}(0, 0, 0) \\ f''_{xy}(0, 0, 0) & f''_{yy}(0, 0, 0) & f''_{y\sigma}(0, 0, 0) \\ f''_{x\sigma}(0, 0, 0) & f''_{y\sigma}(0, 0, 0) & f''_{\sigma\sigma}(0, 0, 0) \end{pmatrix}$$

Sub-Pixel Refinement

2nd order surface fitting (3D):

$$f(x, y, \sigma) \approx f(0, 0, 0) + (x \quad y \quad \sigma) \cdot \nabla f(0, 0, 0) + \frac{1}{2} (x \quad y \quad \sigma) H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

Sub-Pixel Refinement

2nd order surface fitting (3D):

$$f(x, y, \sigma) \approx f(0, 0, 0) + (x \quad y \quad \sigma) \cdot \nabla f(0, 0, 0) + \frac{1}{2} (x \quad y \quad \sigma) H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

Set derivative to 0:

$$\begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} = \nabla f(0, 0, 0) + H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

Sub-Pixel Refinement

2nd order surface fitting (3D):

$$f(x, y, \sigma) \approx f(0, 0, 0) + (x \quad y \quad \sigma) \cdot \nabla f(0, 0, 0) + \frac{1}{2} (x \quad y \quad \sigma) H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

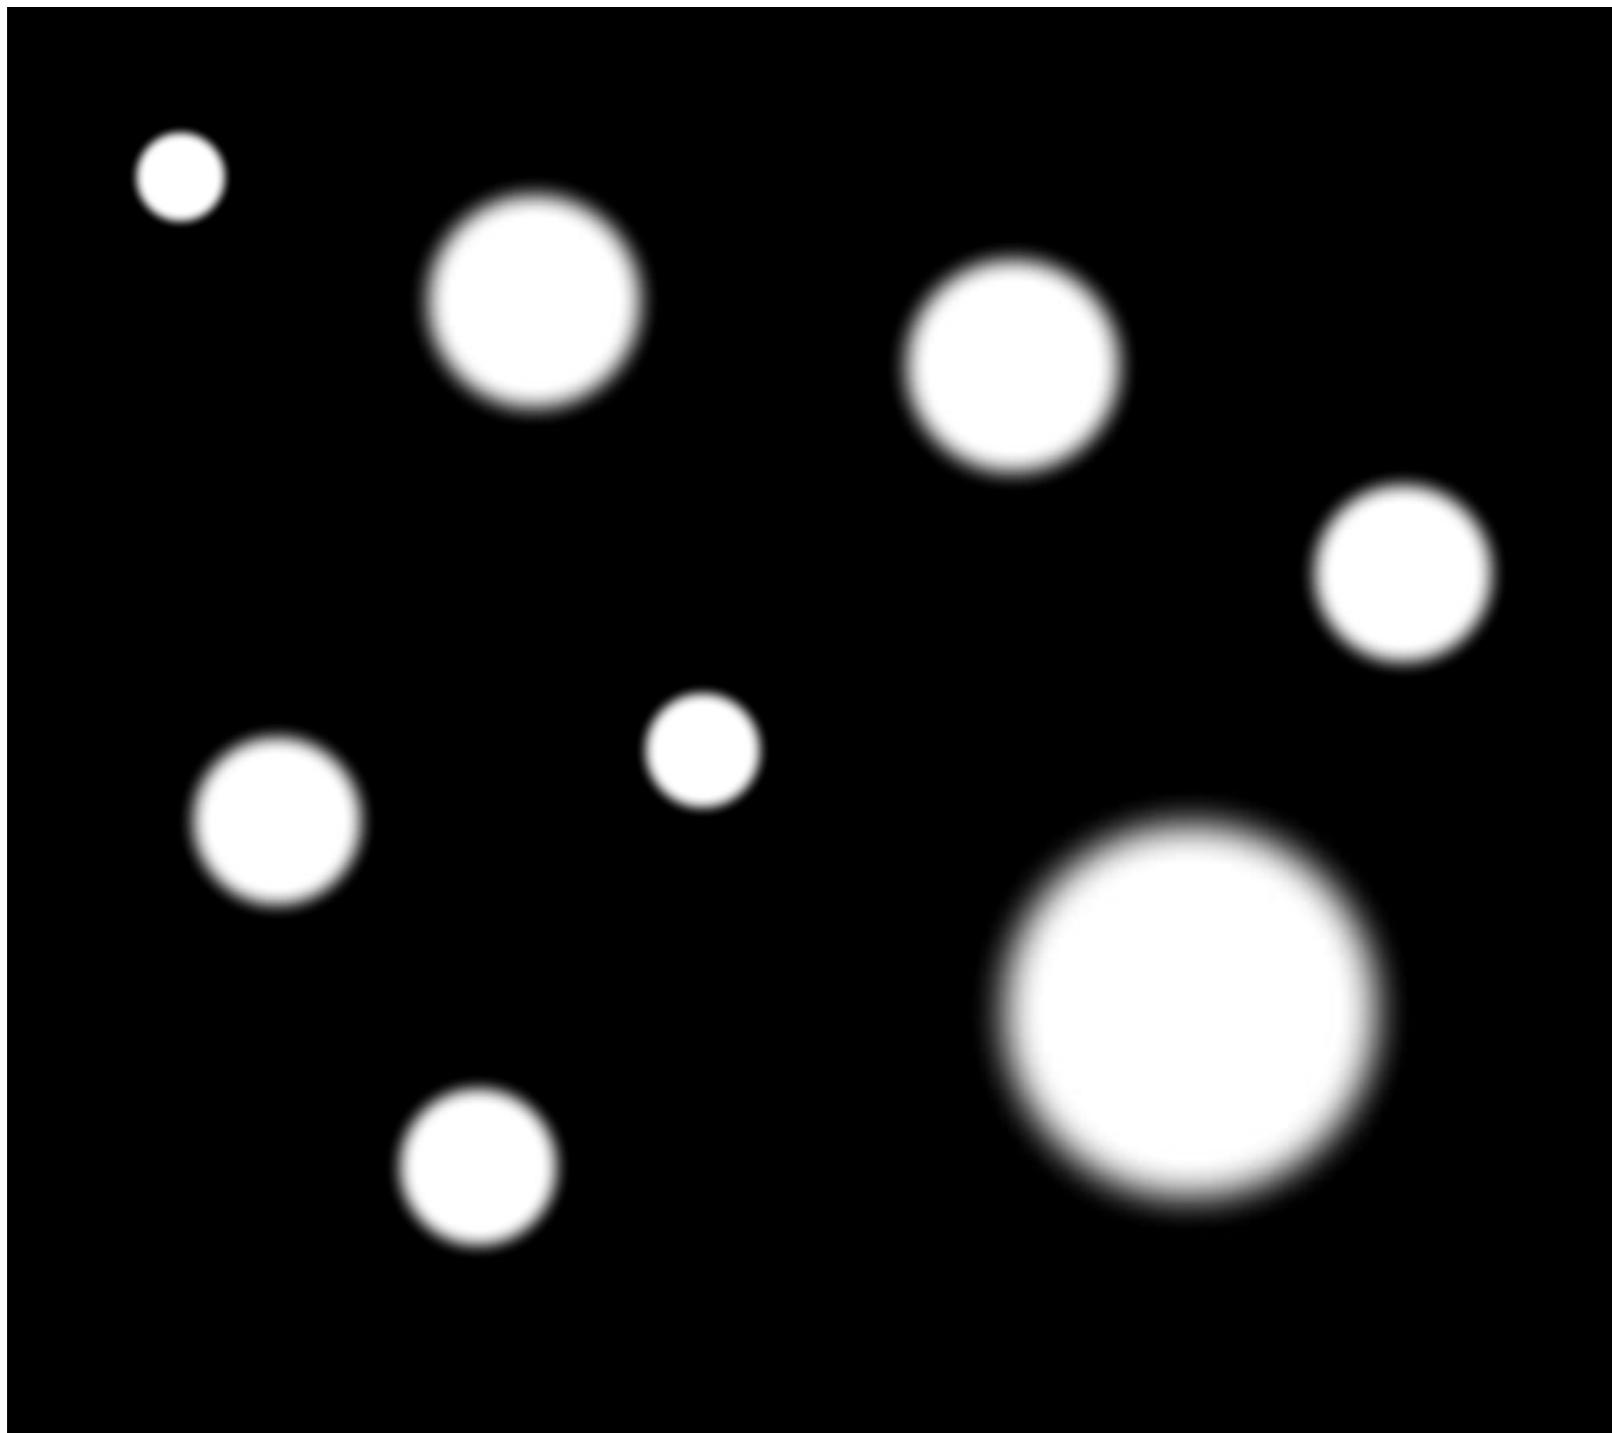
Set derivative to 0:

$$\begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} = \nabla f(0, 0, 0) + H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

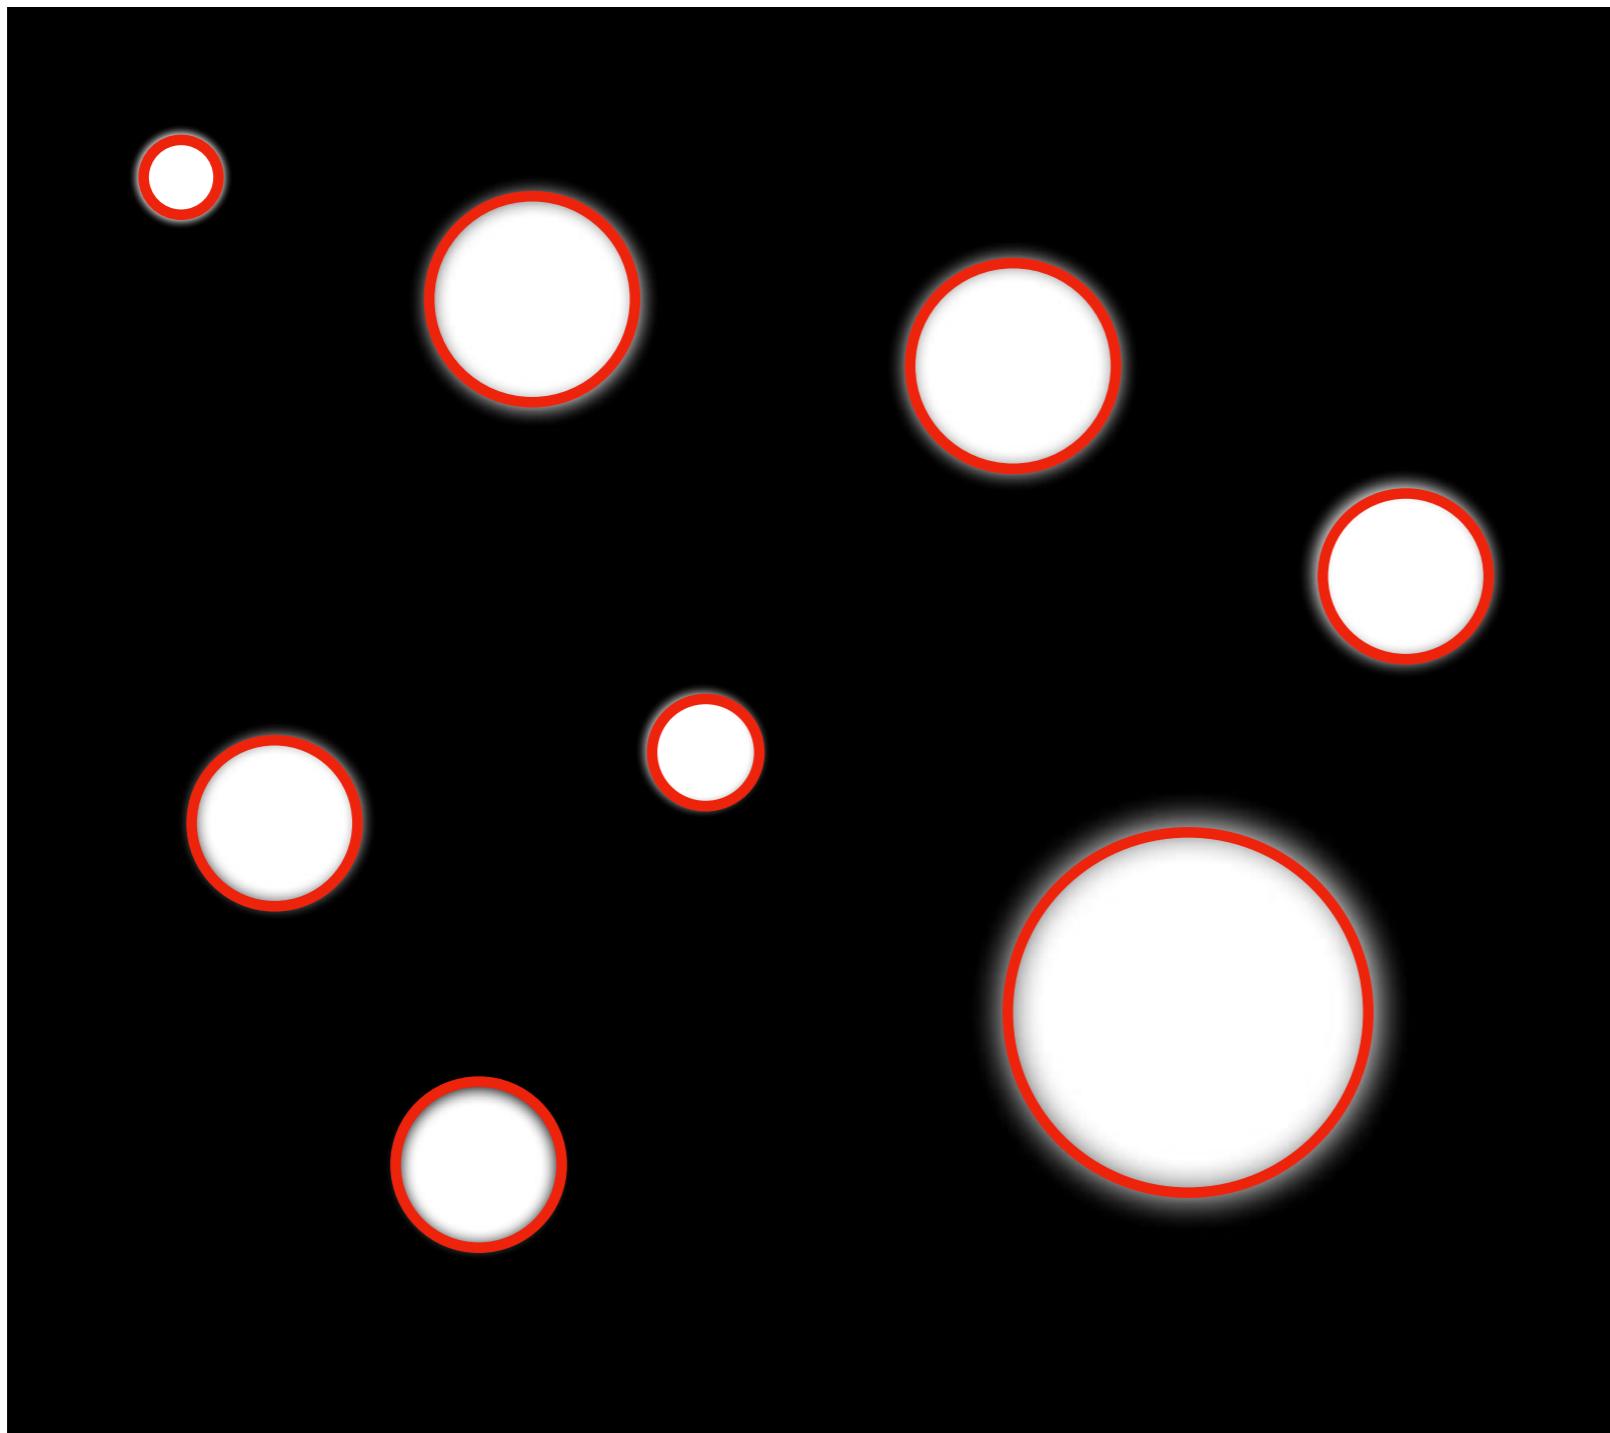
Solve small linear system:

$$\begin{pmatrix} x \\ y \\ \sigma \end{pmatrix} = -H(0, 0, 0)^{-1} \nabla f(0, 0, 0)$$

Application



Application



Lessons Learned

- Main lessons from this lecture
 - Image filters: Gaussian, gradients (edge detection)
 - Similarity measures: ZNCC, histogram of gradients
 - Scale space: Gaussian pyramid by applying Gaussian Filter iteratively
 - Sub-pixel refinement: Fitting 2nd order surfaces
- Next lecture: Local features
 - Will use Gaussian filters, scale space representation, vector bouquets

Next Lecture

Jan. 20	Introduction, Linear classifiers and filtering	
Jan. 23	Filtering, gradients, scale	Lab 1
Jan. 27	Local features	
Jan. 30	Learning a classifier	
Feb. 3	Convolutional neural networks	Lab 2
Feb. 6	More convolutional neural networks	
Feb. 10	Robust model fitting and RANSAC	
Feb. 13	Image registration	Lab 3
Feb. 17	Camera Geometry	
Feb. 20	More camera geometry	
Feb. 24	Generative neural networks	
Feb. 27	Generative neural networks	
Mar. 2	TBA	
Mar. 9	TBA	