

SSY098 - Image Analysis

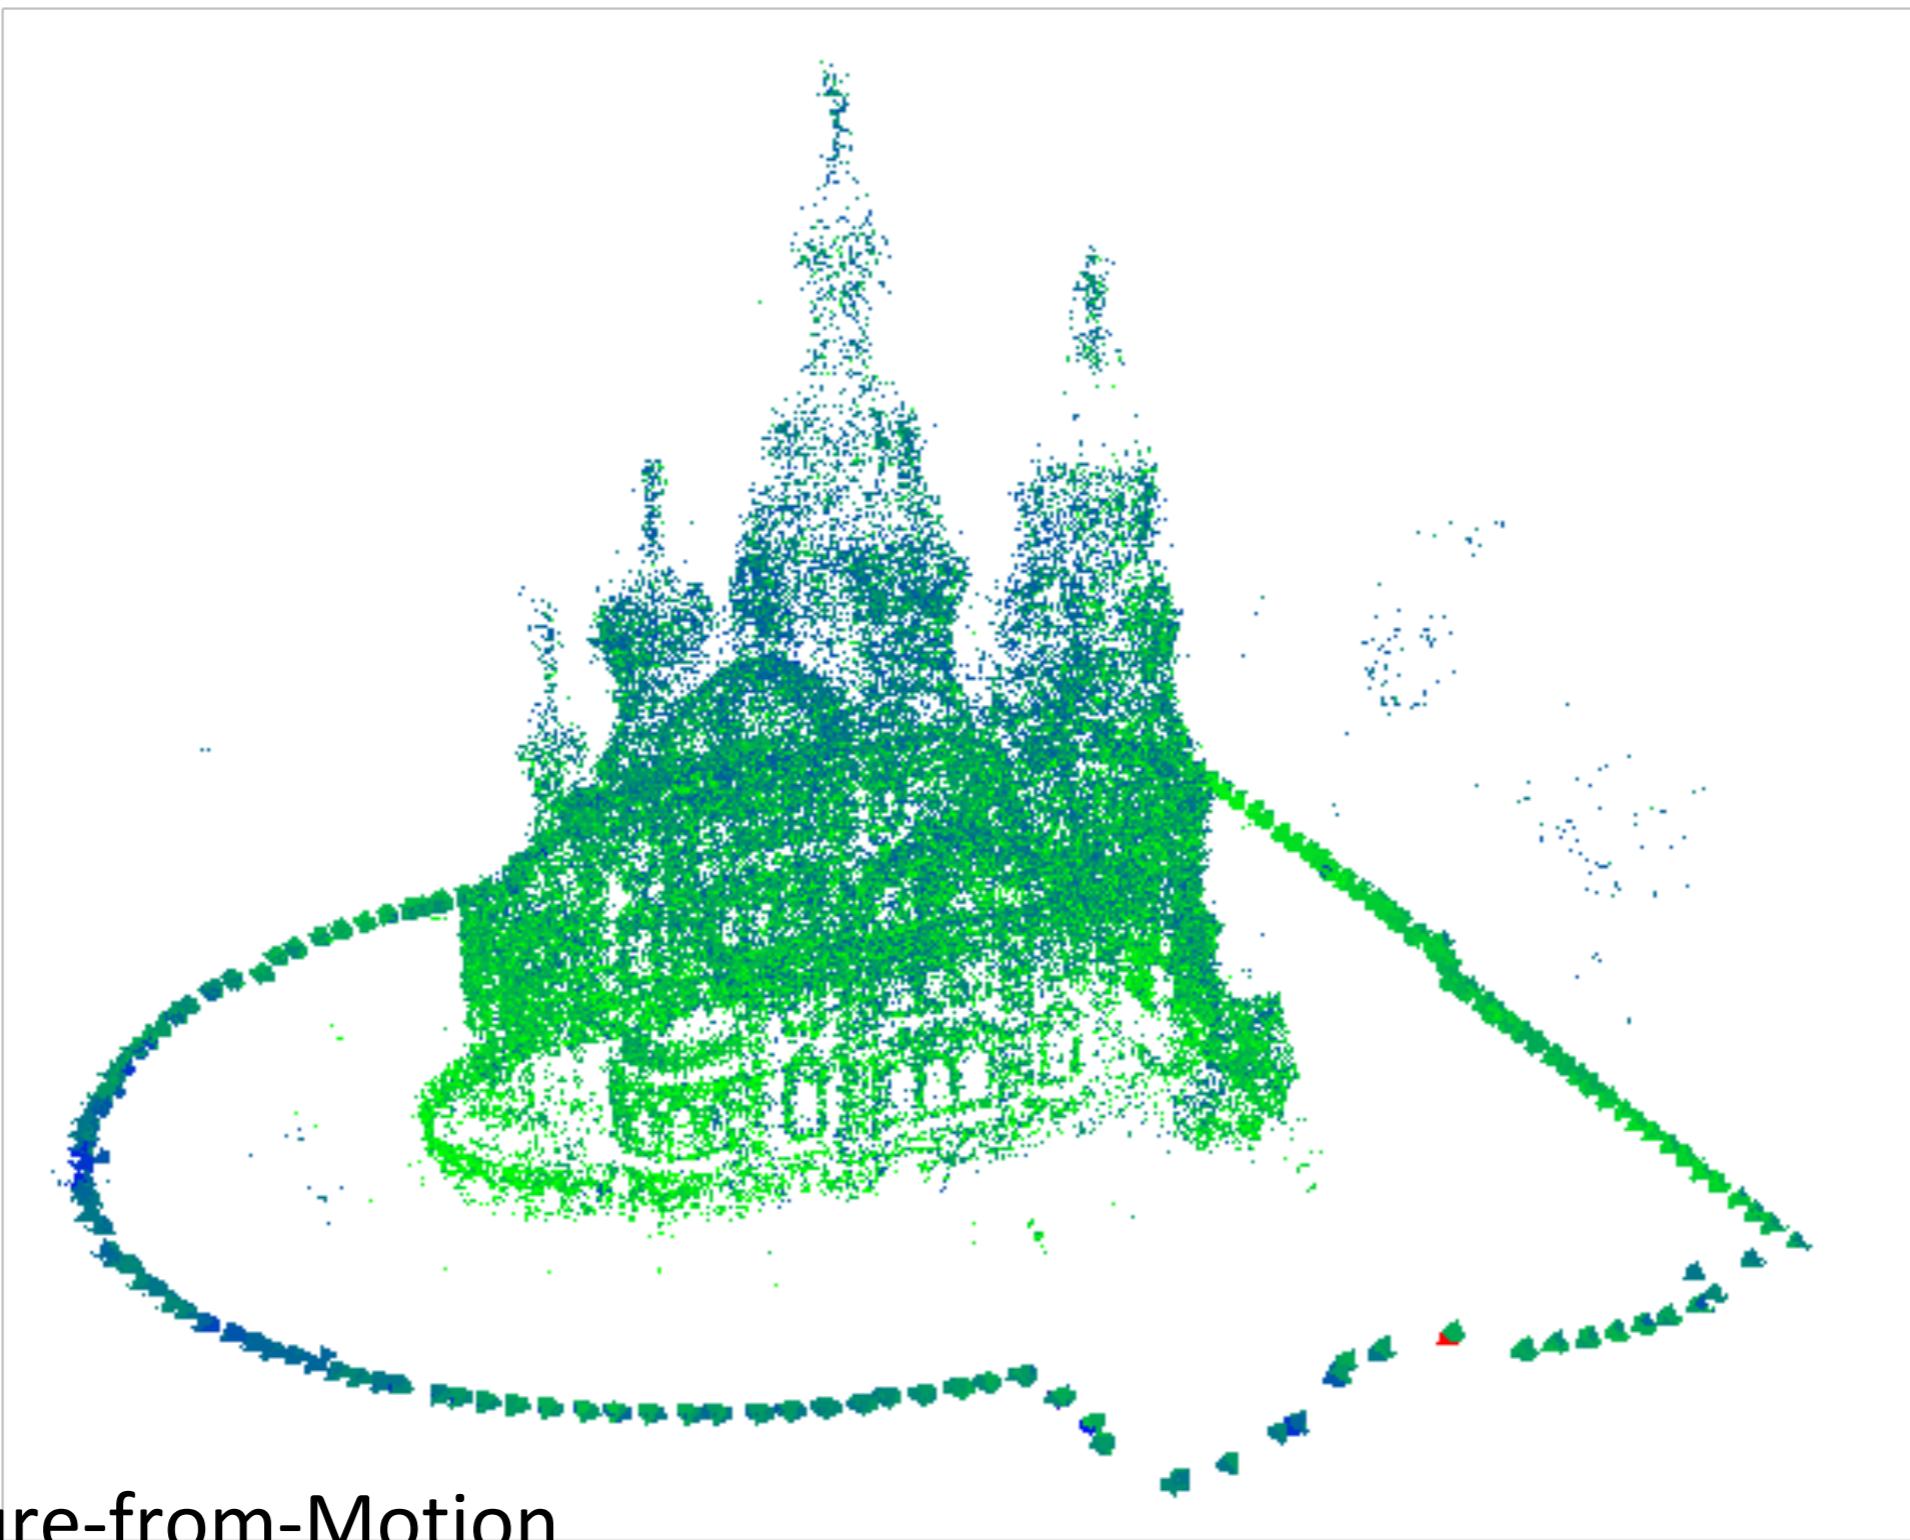
Lecture 11 - Generative Neural Networks

Torsten Sattler

Last Lecture

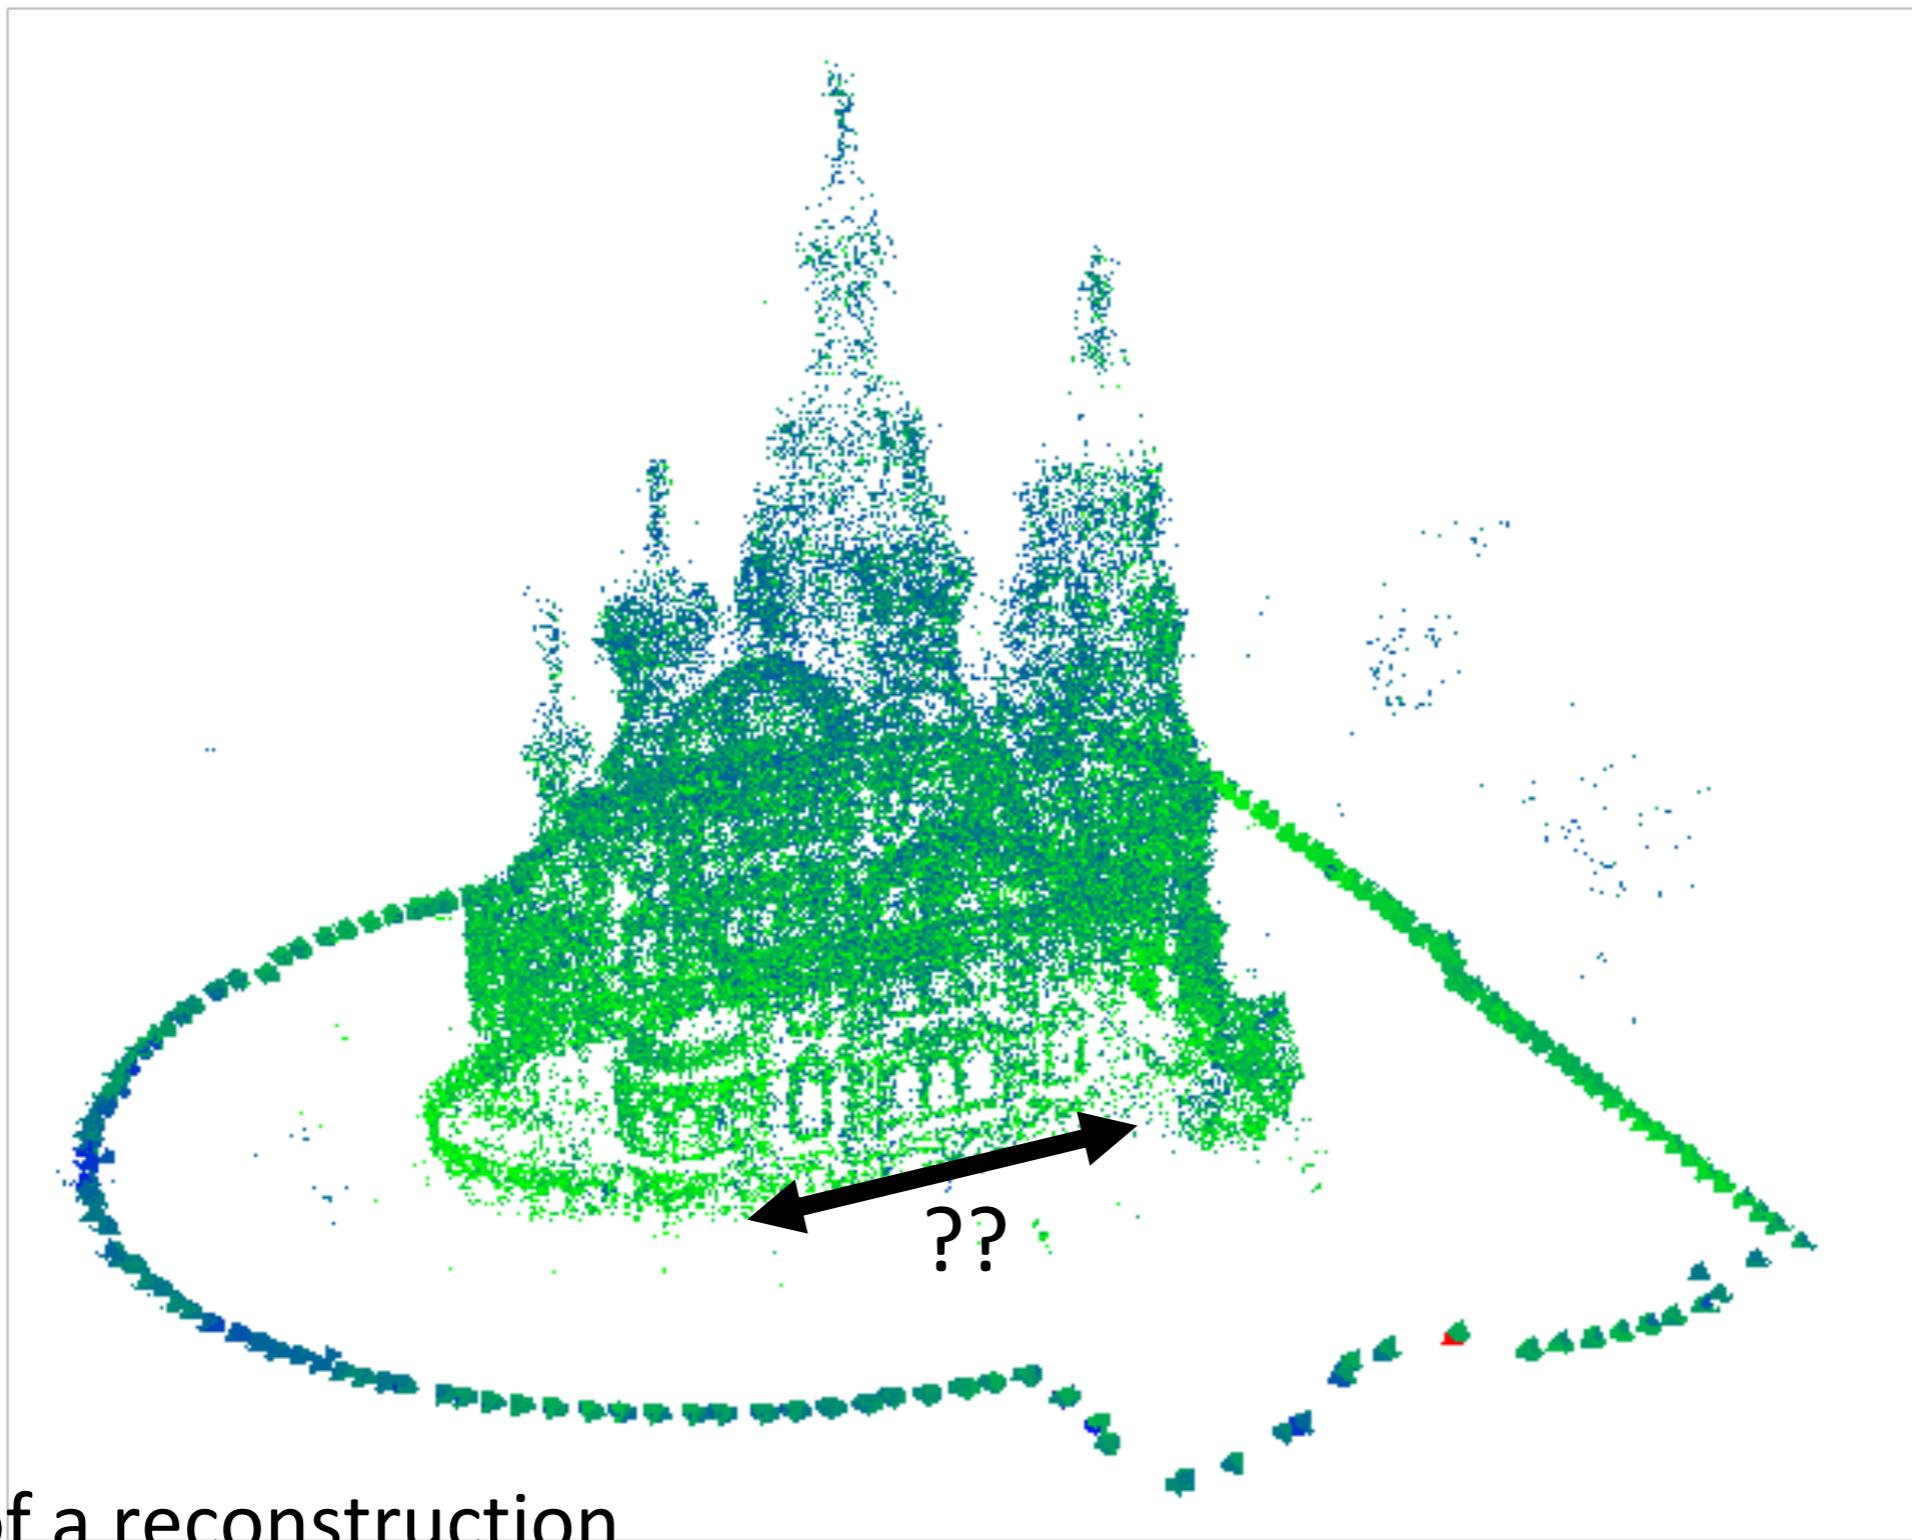
Jan. 20	Introduction, Linear classifiers and filtering	
Jan. 23	Filtering, gradients, scale	Lab 1
Jan. 27	Local features	
Jan. 30	Learning a classifier	
Feb. 3	Convolutional neural networks	Lab 2
Feb. 6	More convolutional neural networks	
Feb. 10	Robust model fitting and RANSAC	
Feb. 13	Image registration	Lab 3
Feb. 17	Camera Geometry	
Feb. 20	More camera geometry	
Feb. 24	Generative neural networks	Lab 4
Feb. 27	Generative neural networks	
Mar. 2	Visual Localization & Feature Learning	
Mar. 9	No lecture	

Last Lecture



Structure-from-Motion

Last Lecture



Scale of a reconstruction

Last Lecture

3D point \mathbf{X} seen from camera with pose \mathbf{R} , \mathbf{t} , intrinsics \mathbf{K}

$$\mathbf{K}(\mathbf{R}\mathbf{X} + \mathbf{t}) = \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \underline{x} \\ \underline{z} \\ \underline{y} \\ \underline{z} \end{pmatrix}$$

3D reconstructions defined up to scale

Last Lecture

3D point \mathbf{X} seen from camera with pose \mathbf{R} , \mathbf{t} , intrinsics \mathbf{K}

$$\mathbf{K}(\mathbf{R}\mathbf{X} + \mathbf{t}) = \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \underline{x} \\ \underline{z} \\ \underline{y} \\ \underline{z} \end{pmatrix}$$

Scale 3D scene by arbitrary factor $s \neq 0$

$$\mathbf{K}(\mathbf{R}(s\mathbf{X}) + s\mathbf{t}) = s\mathbf{K}(\mathbf{R}\mathbf{X} + \mathbf{t})$$

3D reconstructions defined up to scale

Last Lecture

3D point \mathbf{X} seen from camera with pose \mathbf{R} , \mathbf{t} , intrinsics \mathbf{K}

$$\mathbf{K}(\mathbf{R}\mathbf{X} + \mathbf{t}) = \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \underline{x} \\ \underline{z} \\ \underline{y} \\ z \end{pmatrix}$$

Scale 3D scene by arbitrary factor $s \neq 0$

$$\begin{aligned} \mathbf{K}(\mathbf{R}(s\mathbf{X}) + s\mathbf{t}) &= s\mathbf{K}(\mathbf{R}\mathbf{X} + \mathbf{t}) \\ &= s \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \underline{sx} \\ \underline{sz} \\ \underline{sy} \\ sz \end{pmatrix} \end{aligned}$$

3D reconstructions defined up to scale

Last Lecture

3D point \mathbf{X} seen from camera with pose \mathbf{R} , \mathbf{t} , intrinsics \mathbf{K}

$$\mathbf{K}(\mathbf{R}\mathbf{X} + \mathbf{t}) = \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \underline{x} \\ \underline{z} \\ \underline{y} \\ z \end{pmatrix}$$

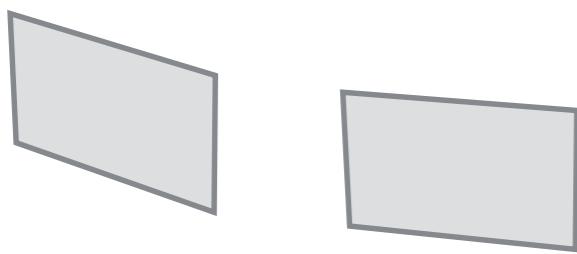
Scale 3D scene by arbitrary factor $s \neq 0$

$$\begin{aligned} \mathbf{K}(\mathbf{R}(s\mathbf{X}) + s\mathbf{t}) &= s\mathbf{K}(\mathbf{R}\mathbf{X} + \mathbf{t}) \\ &= s \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \underline{sx} \\ \underline{sz} \\ \underline{sy} \\ sz \end{pmatrix} = \begin{pmatrix} \underline{x} \\ \underline{z} \\ \underline{y} \\ z \end{pmatrix} \end{aligned}$$

3D reconstructions defined up to scale

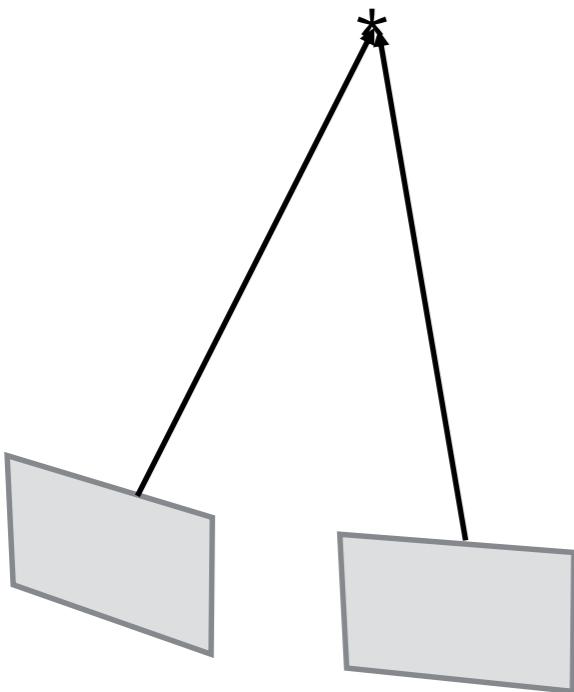
Last Lecture

Initialize motion from two views



Relative pose for two images

Last Lecture

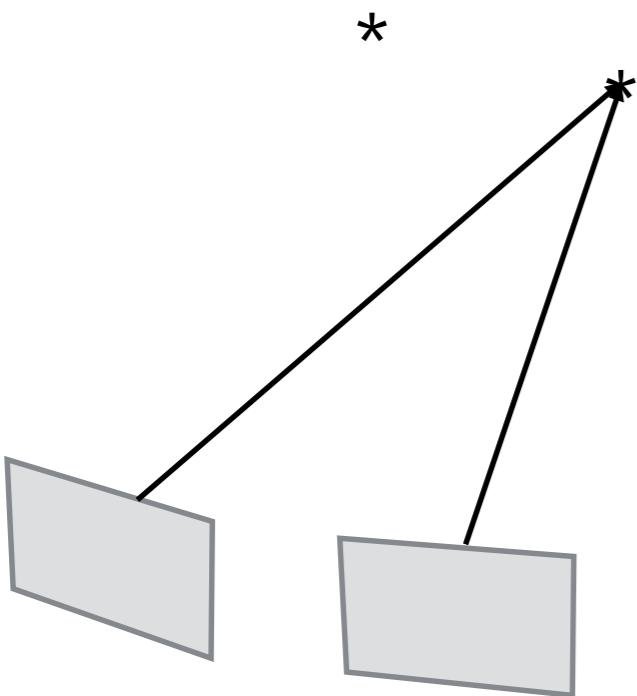


Initialize motion from two views

Initialize structure from two views

Triangulate 3D points

Last Lecture

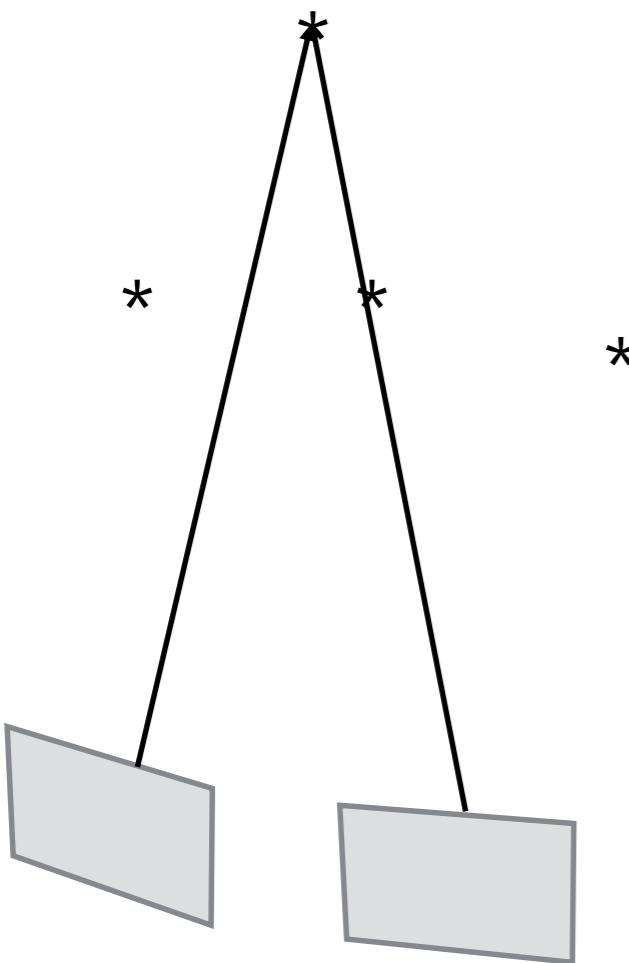


Initialize motion from two views

Initialize structure from two views

Triangulate 3D points

Last Lecture



Initialize motion from two views

Initialize structure from two views

Triangulate 3D points

Last Lecture

*

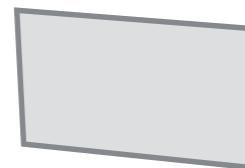
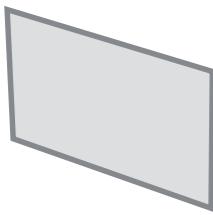
*

*

*

Initialize motion from two views

Initialize structure from two views



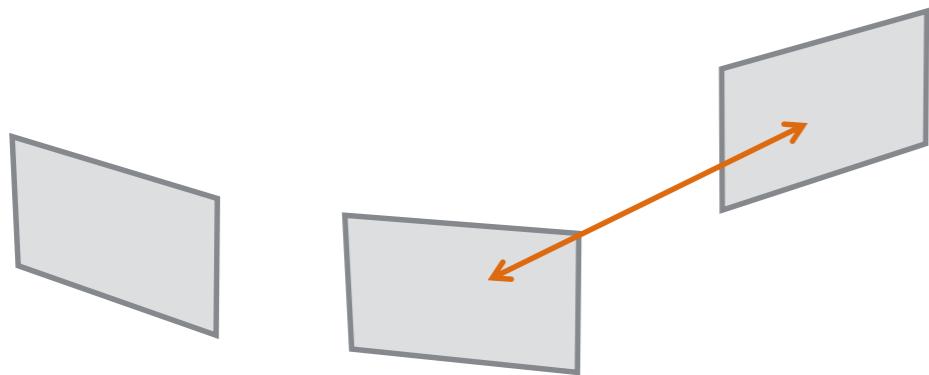
Last Lecture

*

*

*

*



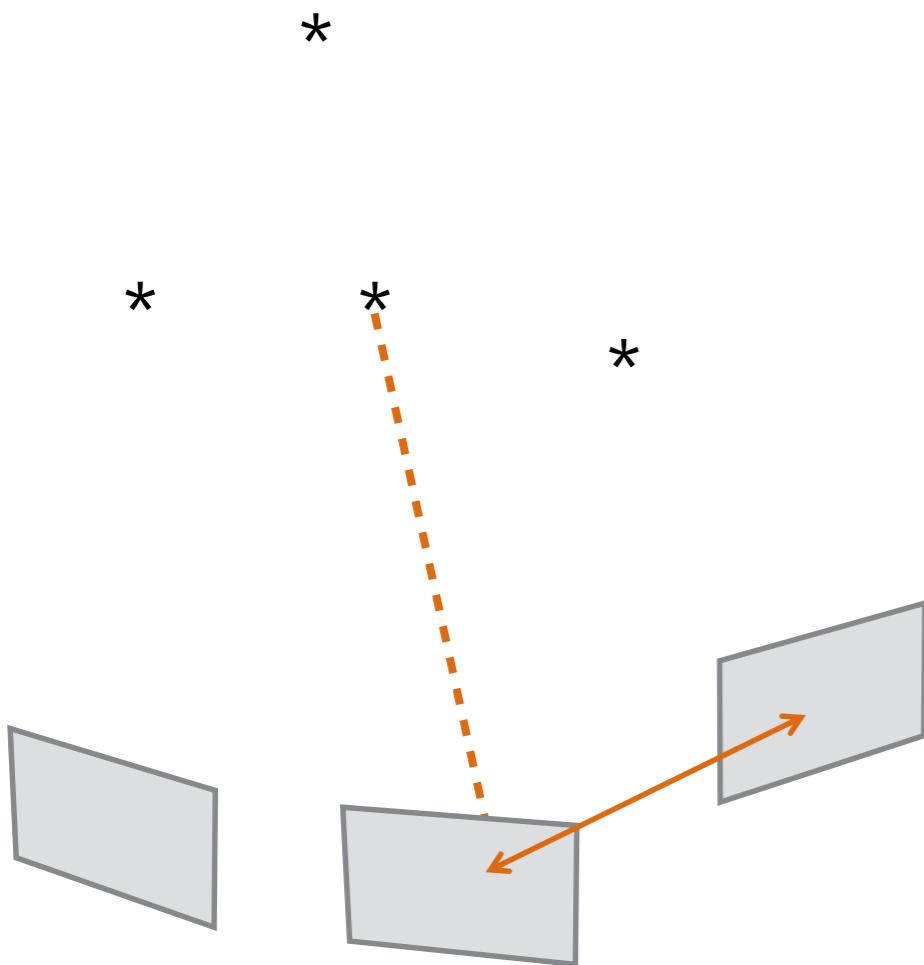
Match features

Initialize motion from two views

Initialize structure from two views

Extend motion

Last Lecture



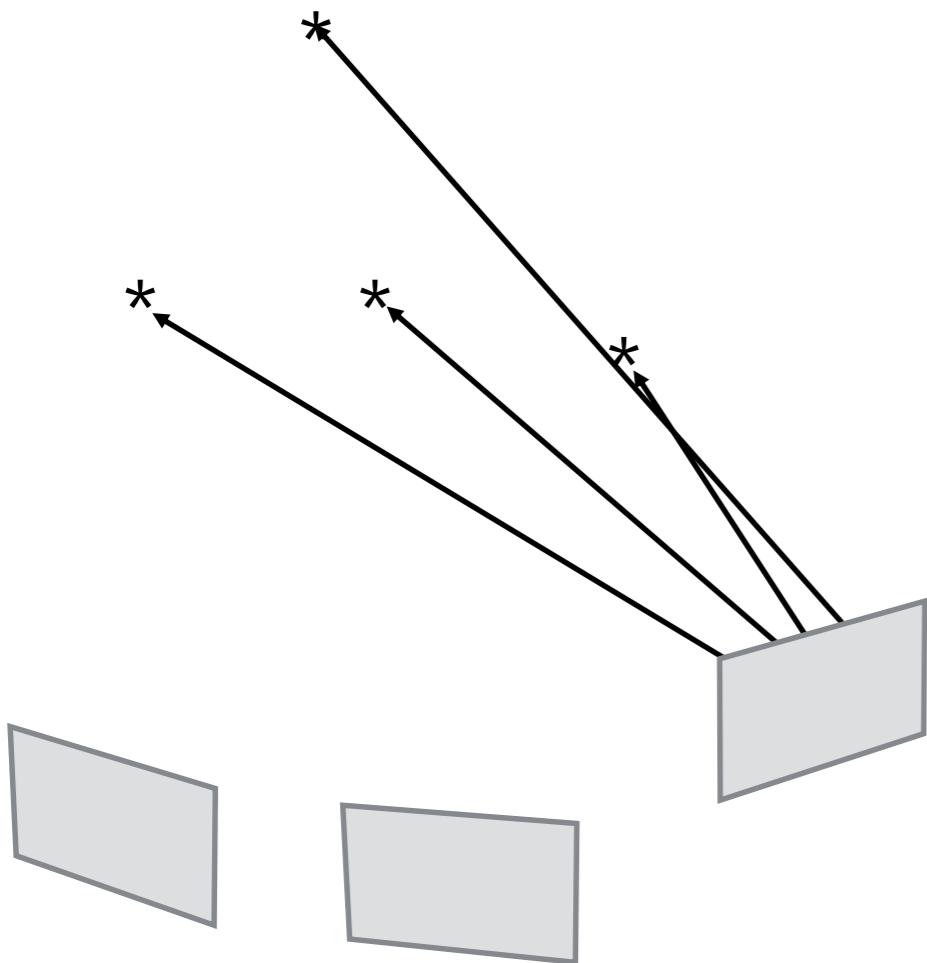
Initialize motion from two views

Initialize structure from two views

Extend motion

Transfer matches to 3D

Last Lecture



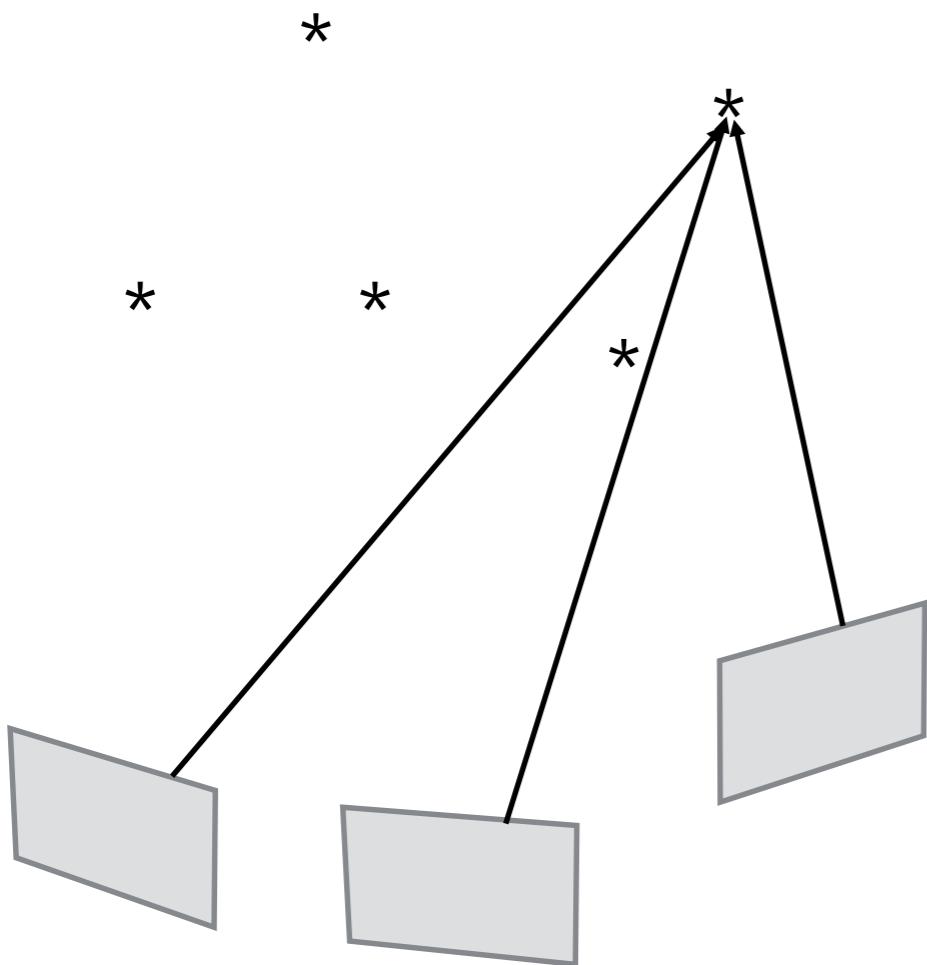
Initialize motion from two views

Initialize structure from two views

Extend motion

Camera pose for third camera

Last Lecture



Initialize motion from two views

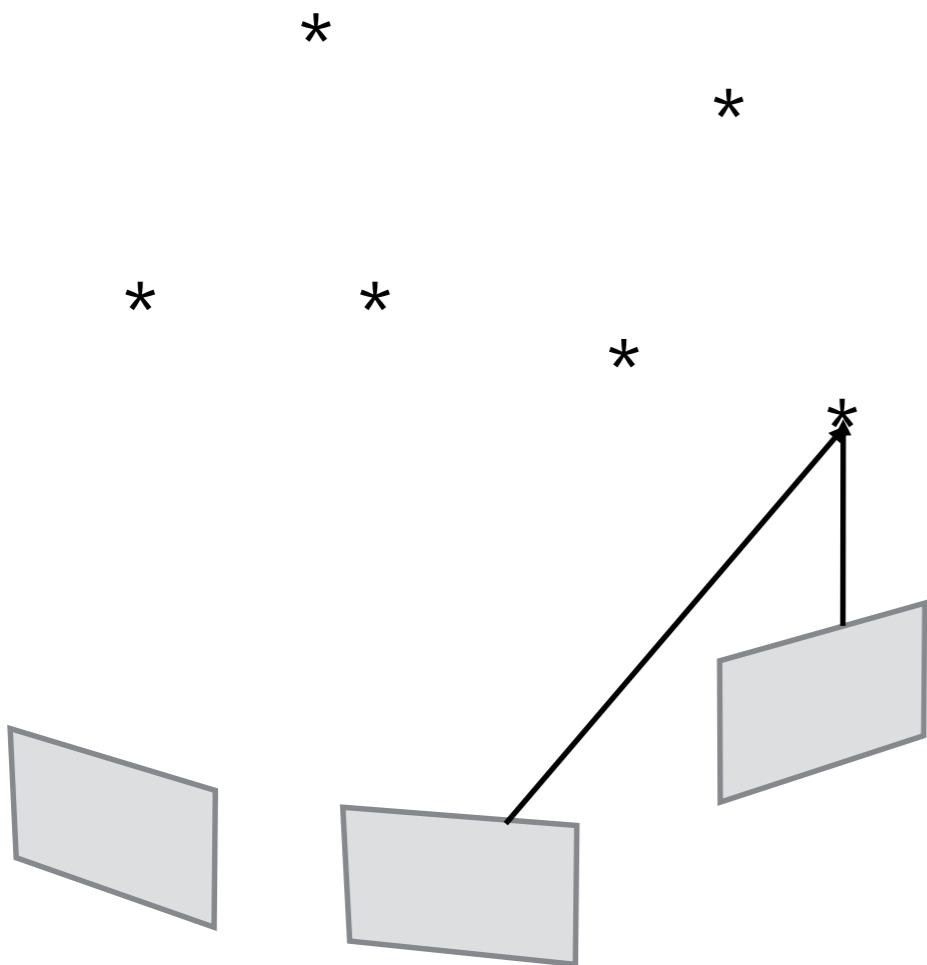
Initialize structure from two views

Extend motion

Extend structure

Triangulate points

Last Lecture



Initialize motion from two views

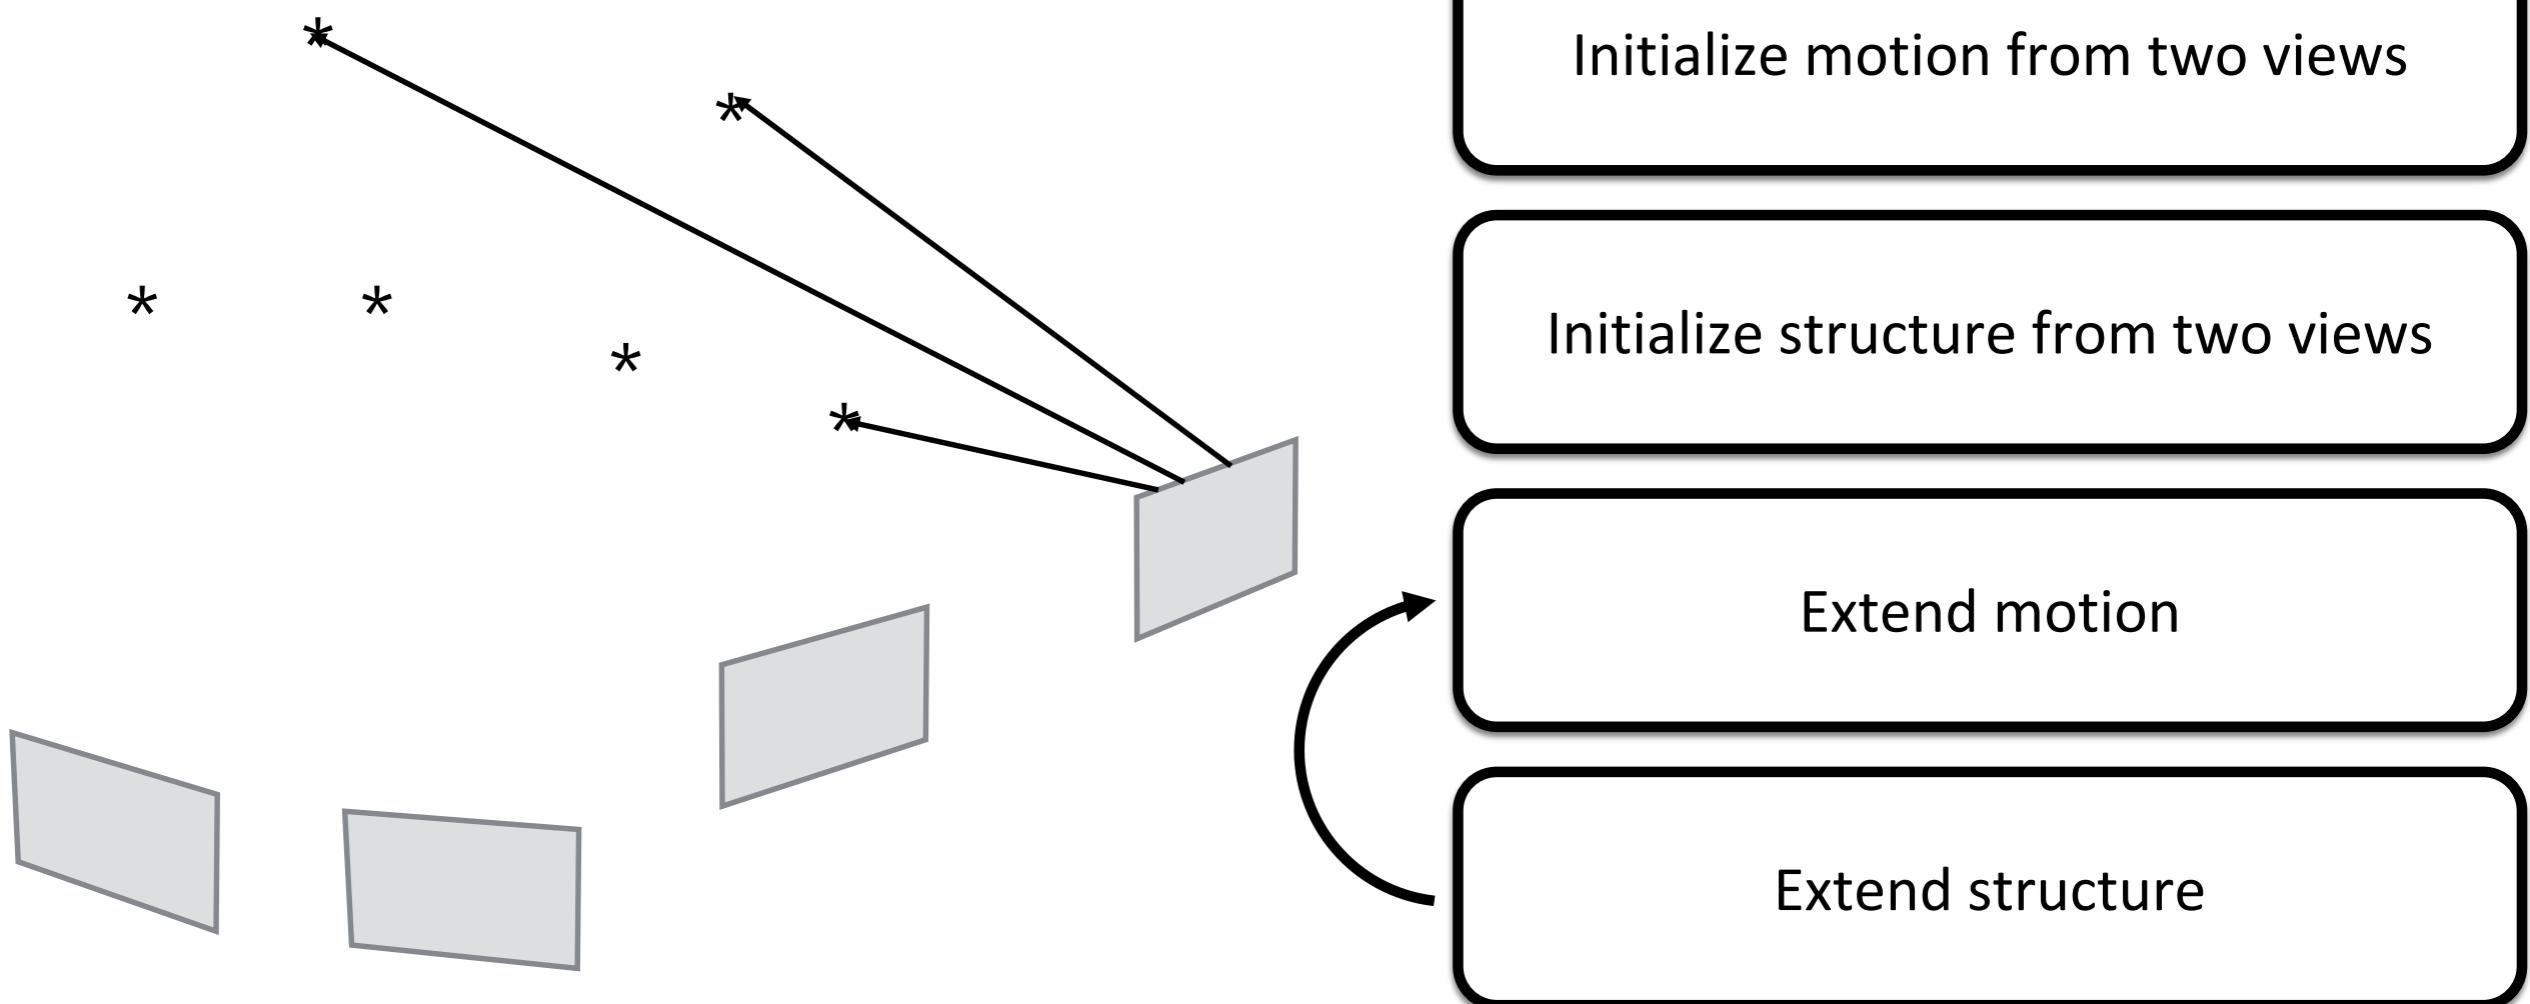
Initialize structure from two views

Extend motion

Extend structure

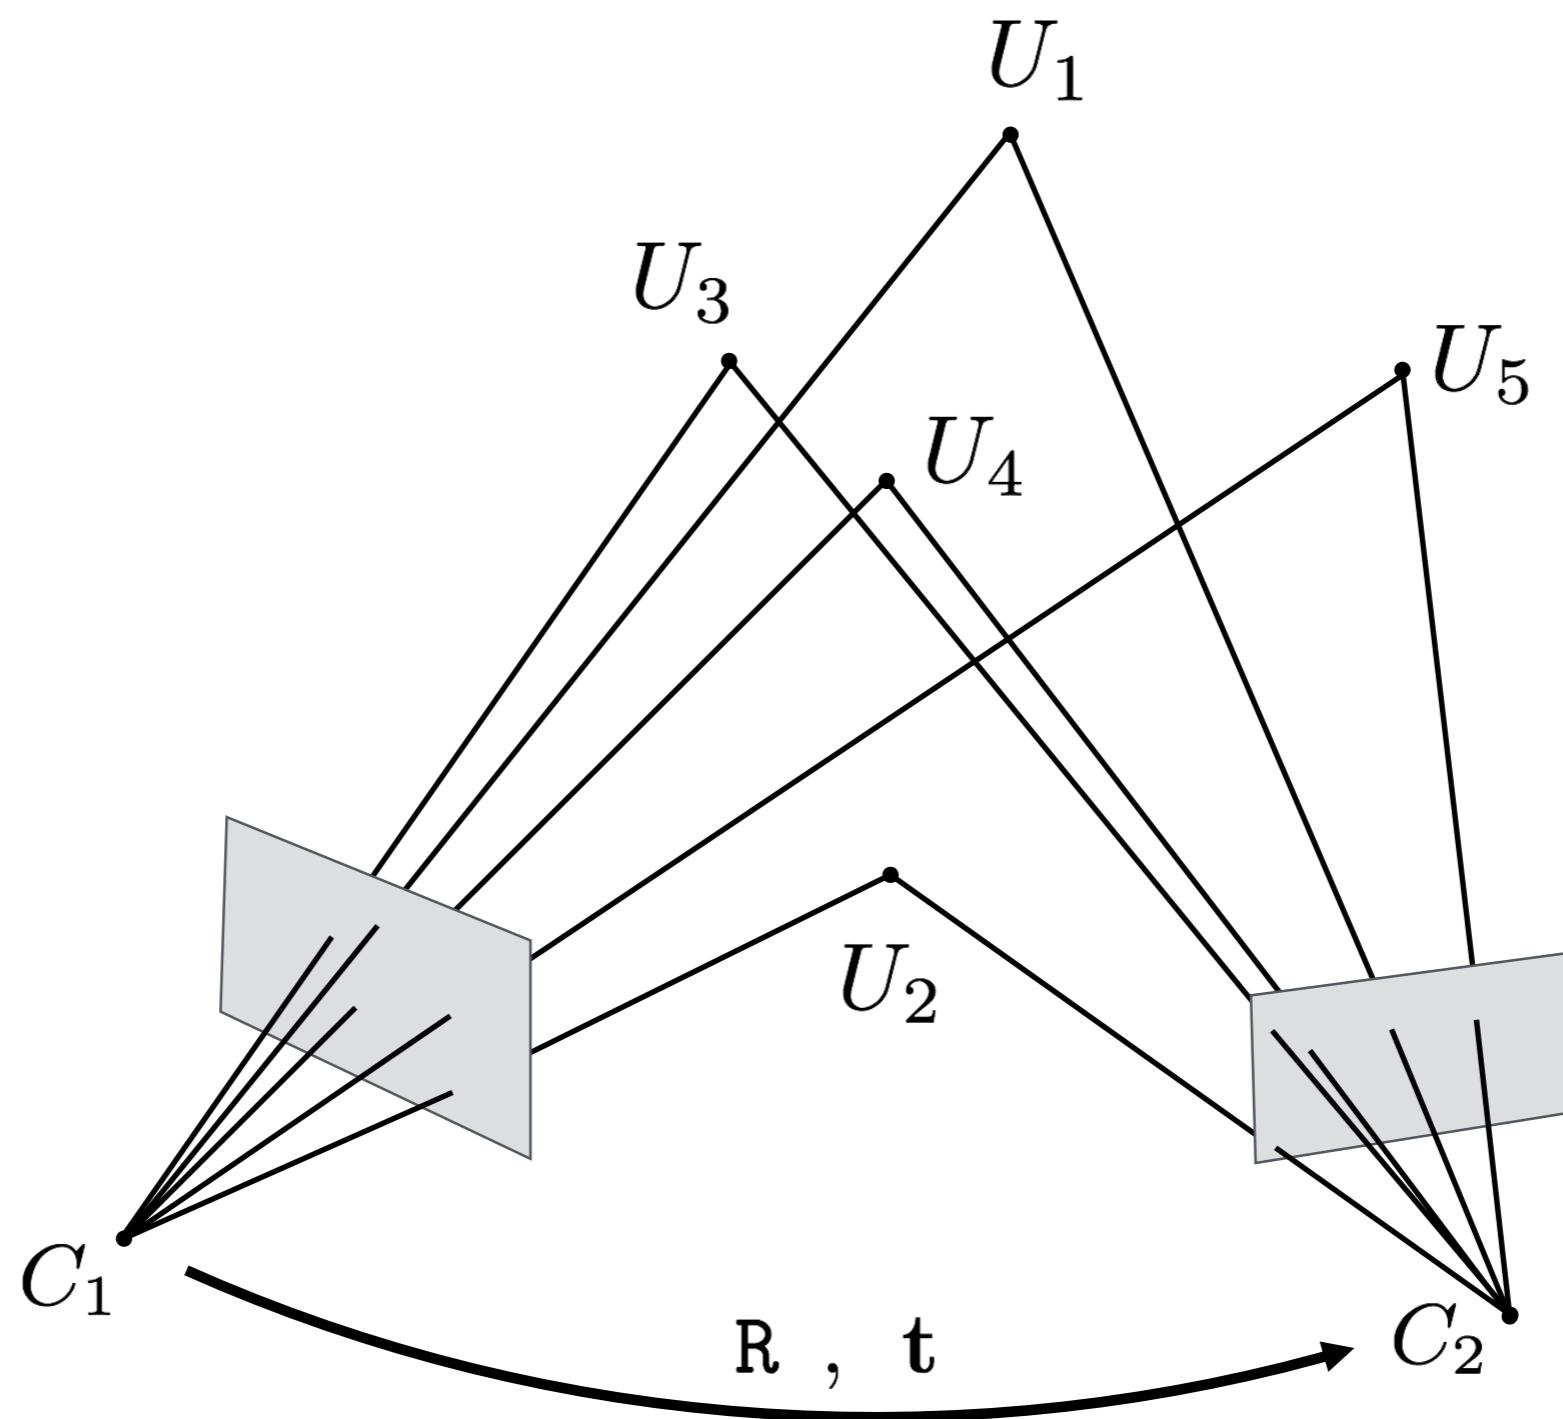
Triangulate points

Last Lecture



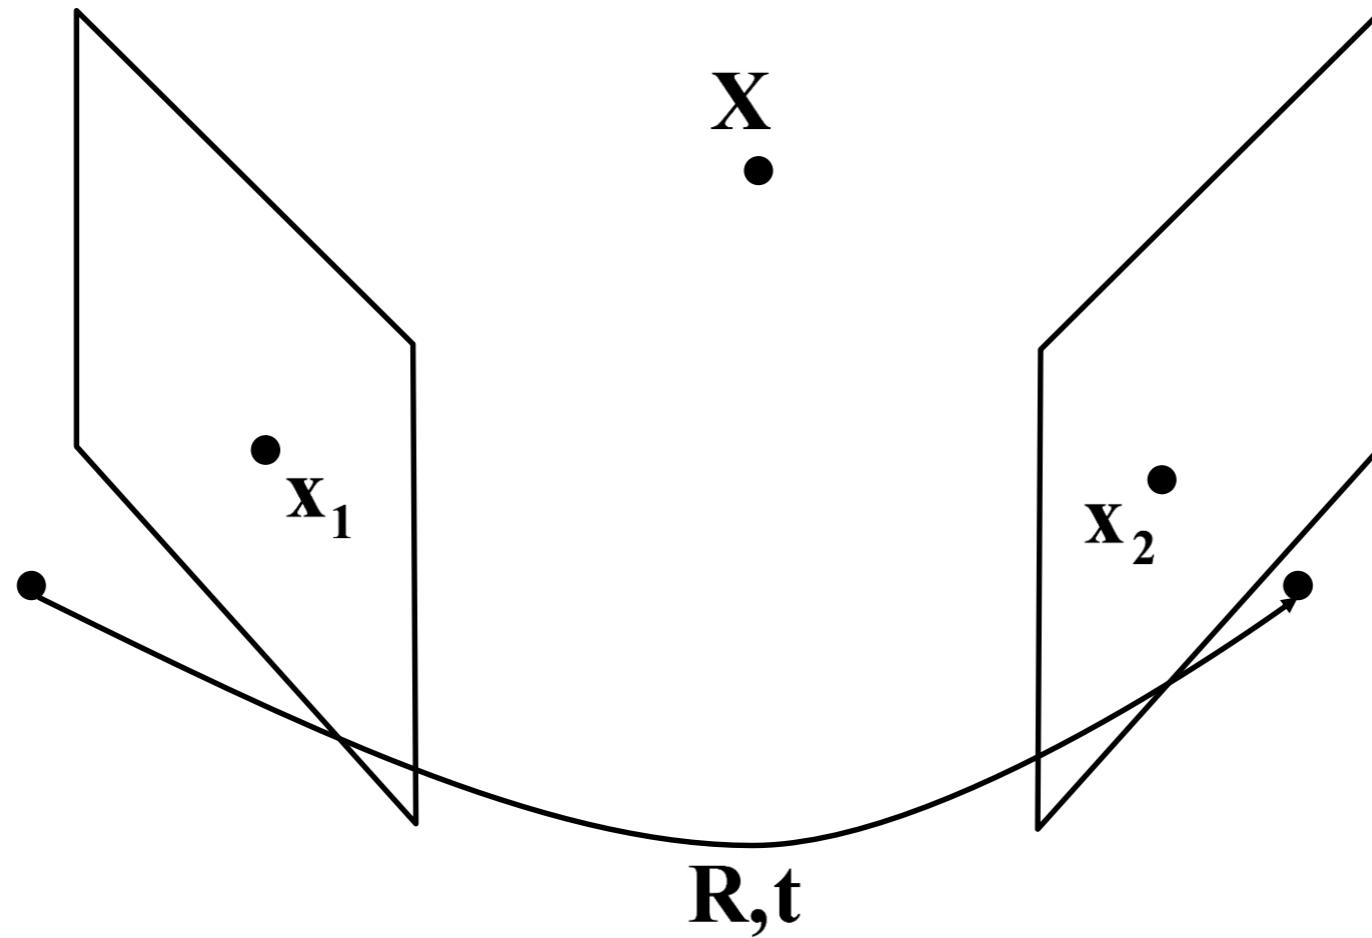
Camera pose for fourth image

Last Lecture



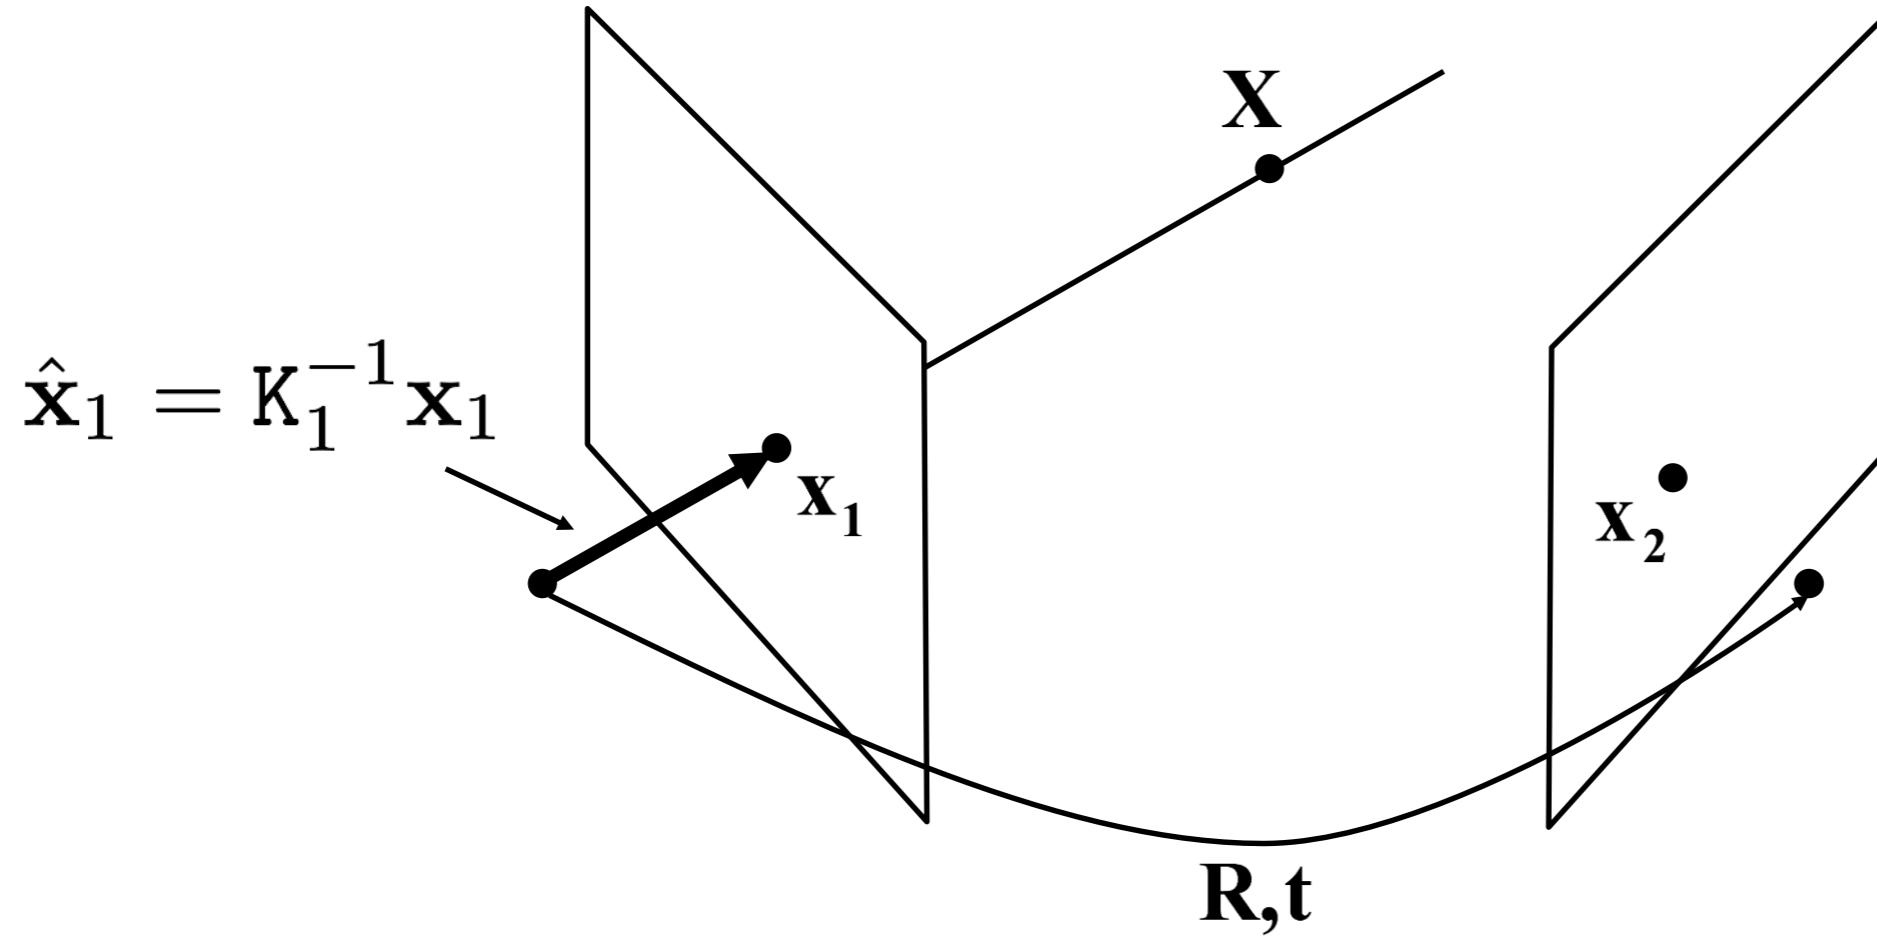
Relative pose estimation

Last Lecture



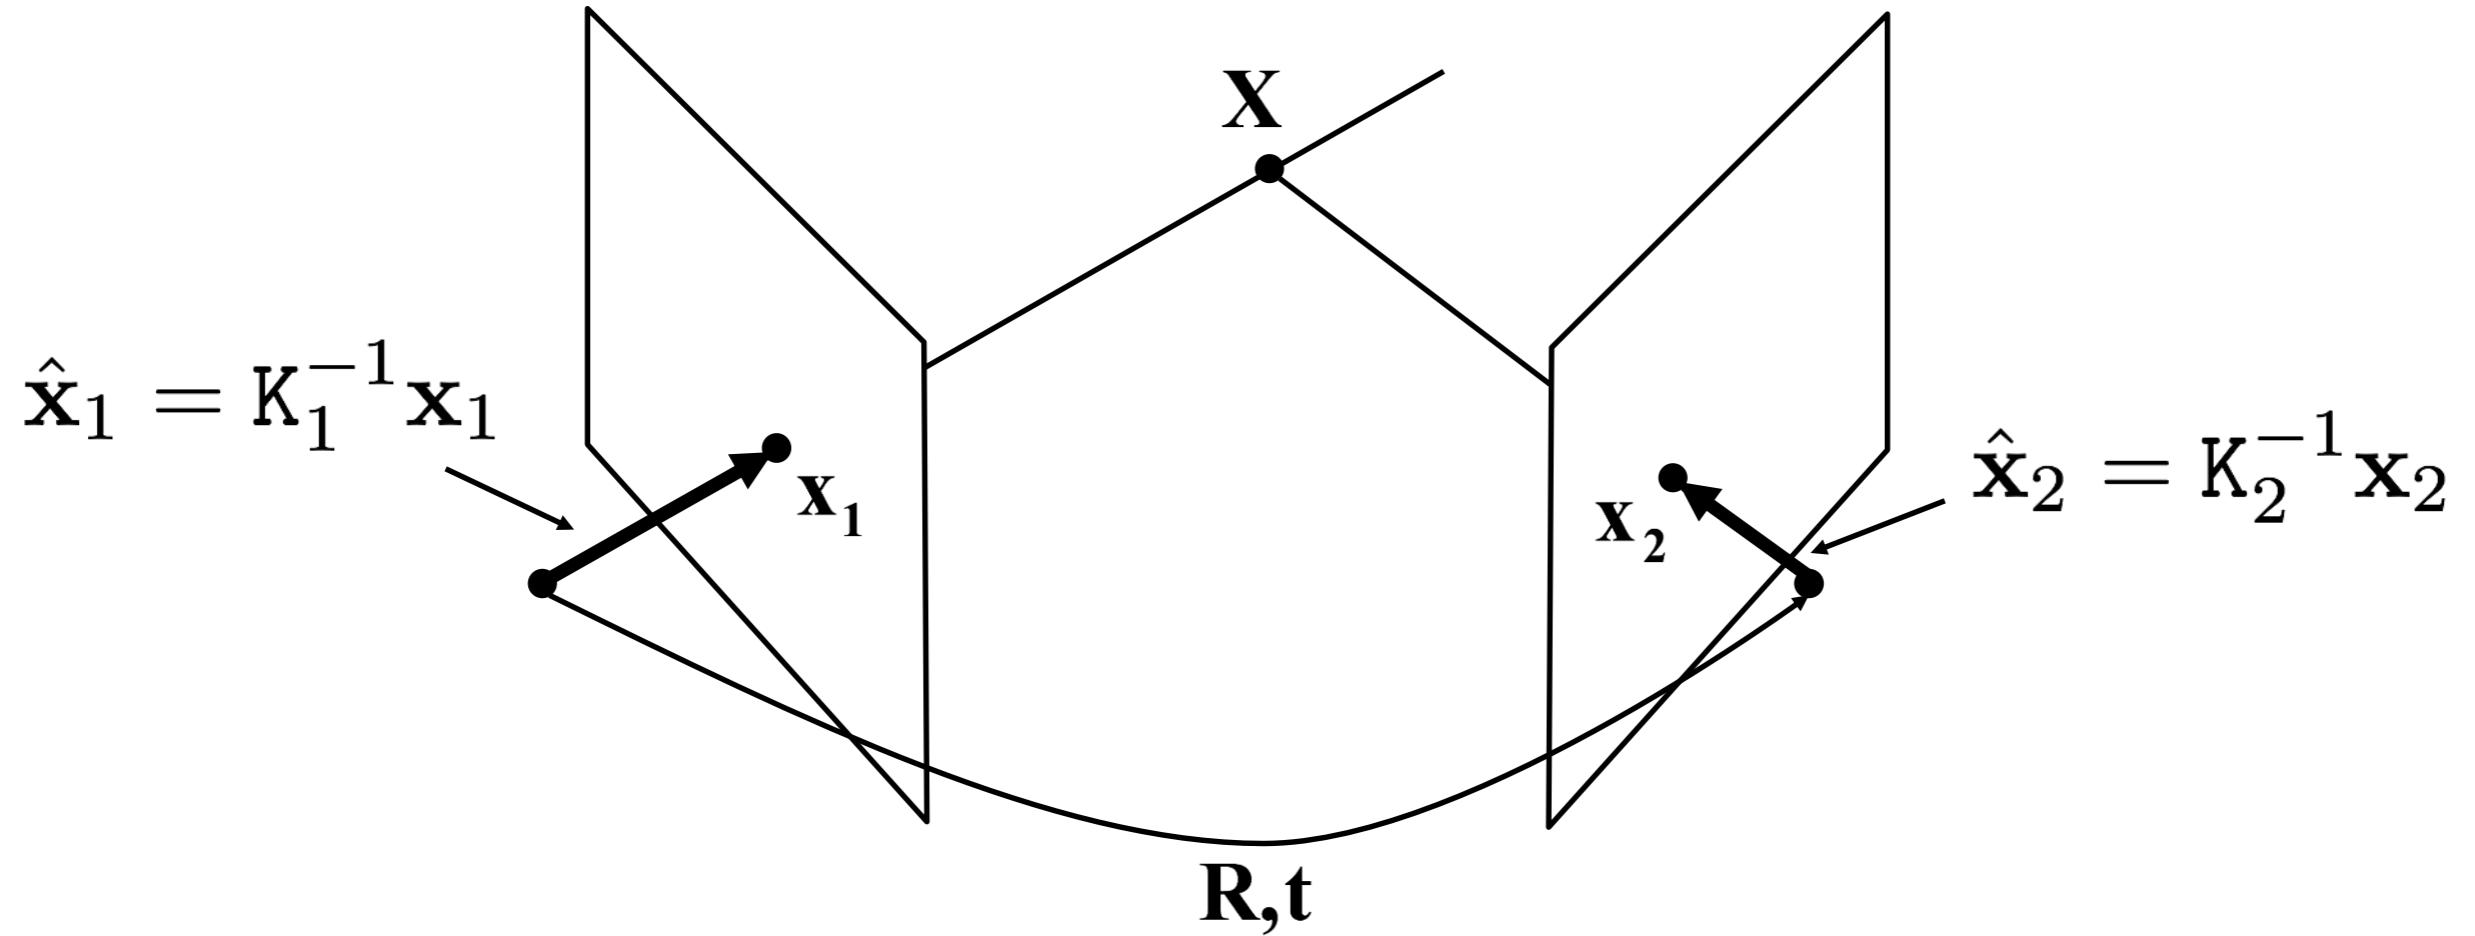
Epipolar geometry

Last Lecture



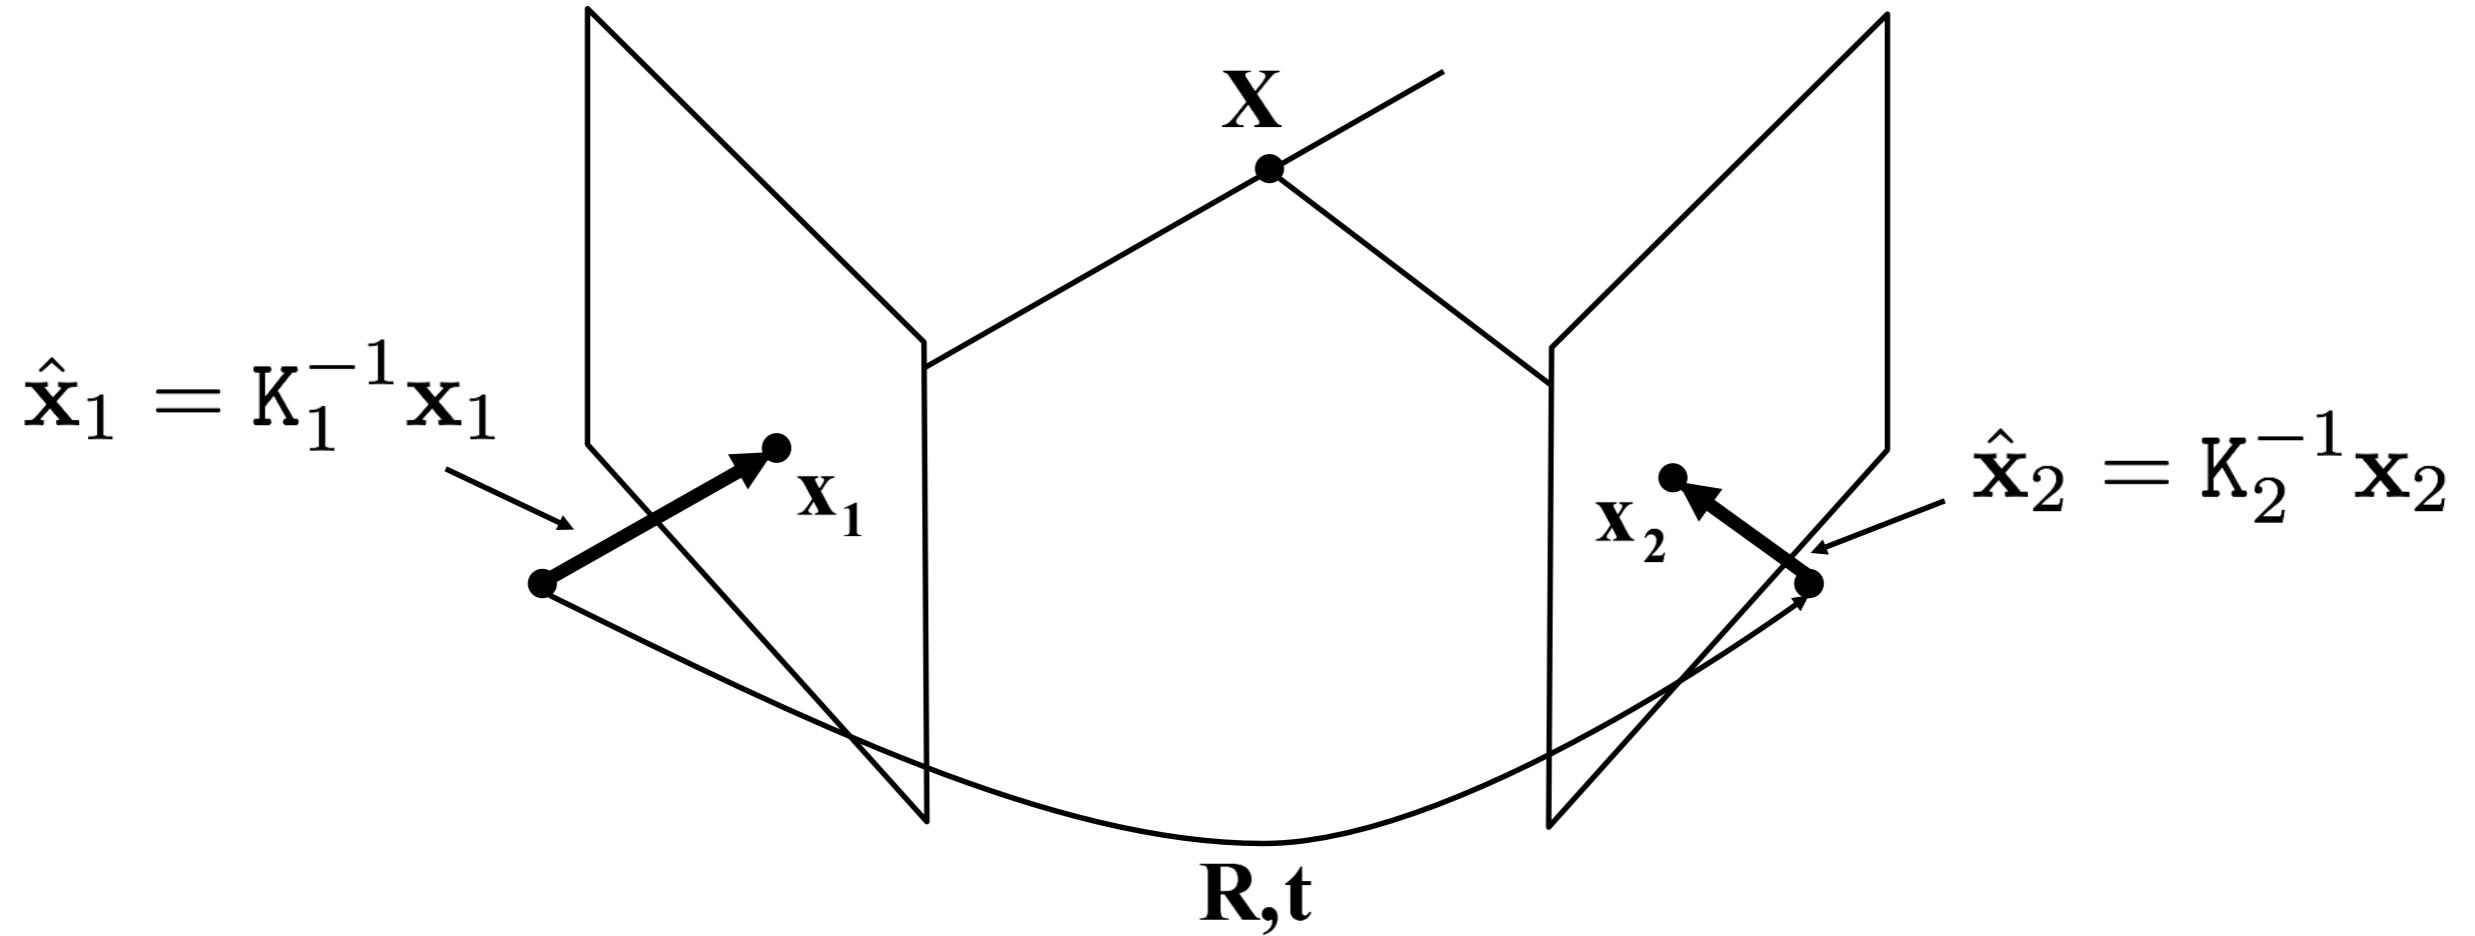
Epipolar geometry

Last Lecture



Epipolar geometry

Last Lecture



$$\text{epipolar constraint: } 0 = \hat{\mathbf{x}}_2^T [\mathbf{t}] \times \mathbf{R} \hat{\mathbf{x}}_1$$

Epipolar geometry

Last Lecture

epipolar constraint:

$$0 = \hat{\mathbf{x}}_2^T [\mathbf{t}]_{\times} \mathbf{R} \hat{\mathbf{x}}_1$$

Essential and Fundamental matrices

Last Lecture

epipolar constraint: $0 = \hat{\mathbf{x}}_2^T [\mathbf{t}]_{\times} \mathbf{R} \hat{\mathbf{x}}_1$

Essential matrix: $0 = \hat{\mathbf{x}}_2^T \mathbf{E} \hat{\mathbf{x}}_1$

Essential and Fundamental matrices

Last Lecture

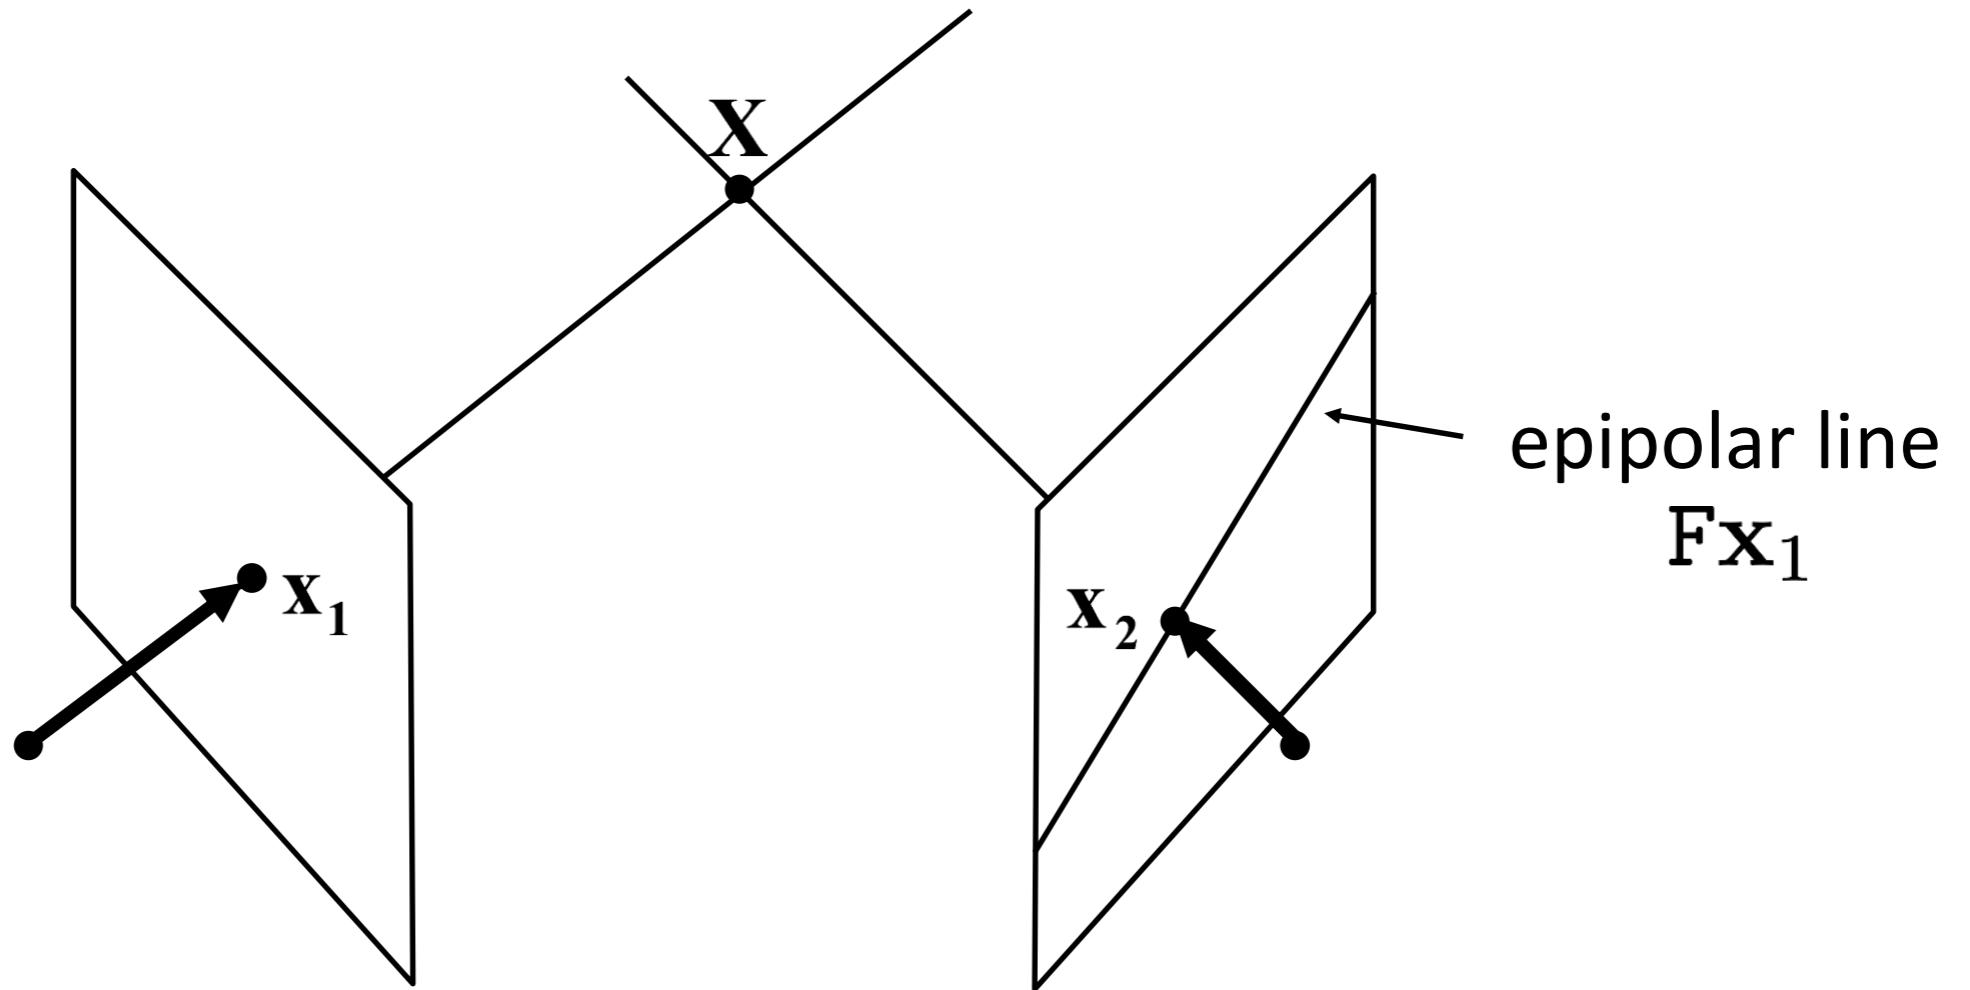
epipolar constraint: $0 = \hat{\mathbf{x}}_2^T [\mathbf{t}] \times \mathbf{R} \hat{\mathbf{x}}_1$

Essential matrix: $0 = \hat{\mathbf{x}}_2^T \mathbf{E} \hat{\mathbf{x}}_1$

Fundamental matrix: $0 = \mathbf{x}_2^T \mathbf{K}_2^{-T} \mathbf{E} \mathbf{K}_1^{-1} \mathbf{x}_1$
 $0 = \mathbf{x}_2^T \mathbf{F} \mathbf{x}_1$

Essential and Fundamental matrices

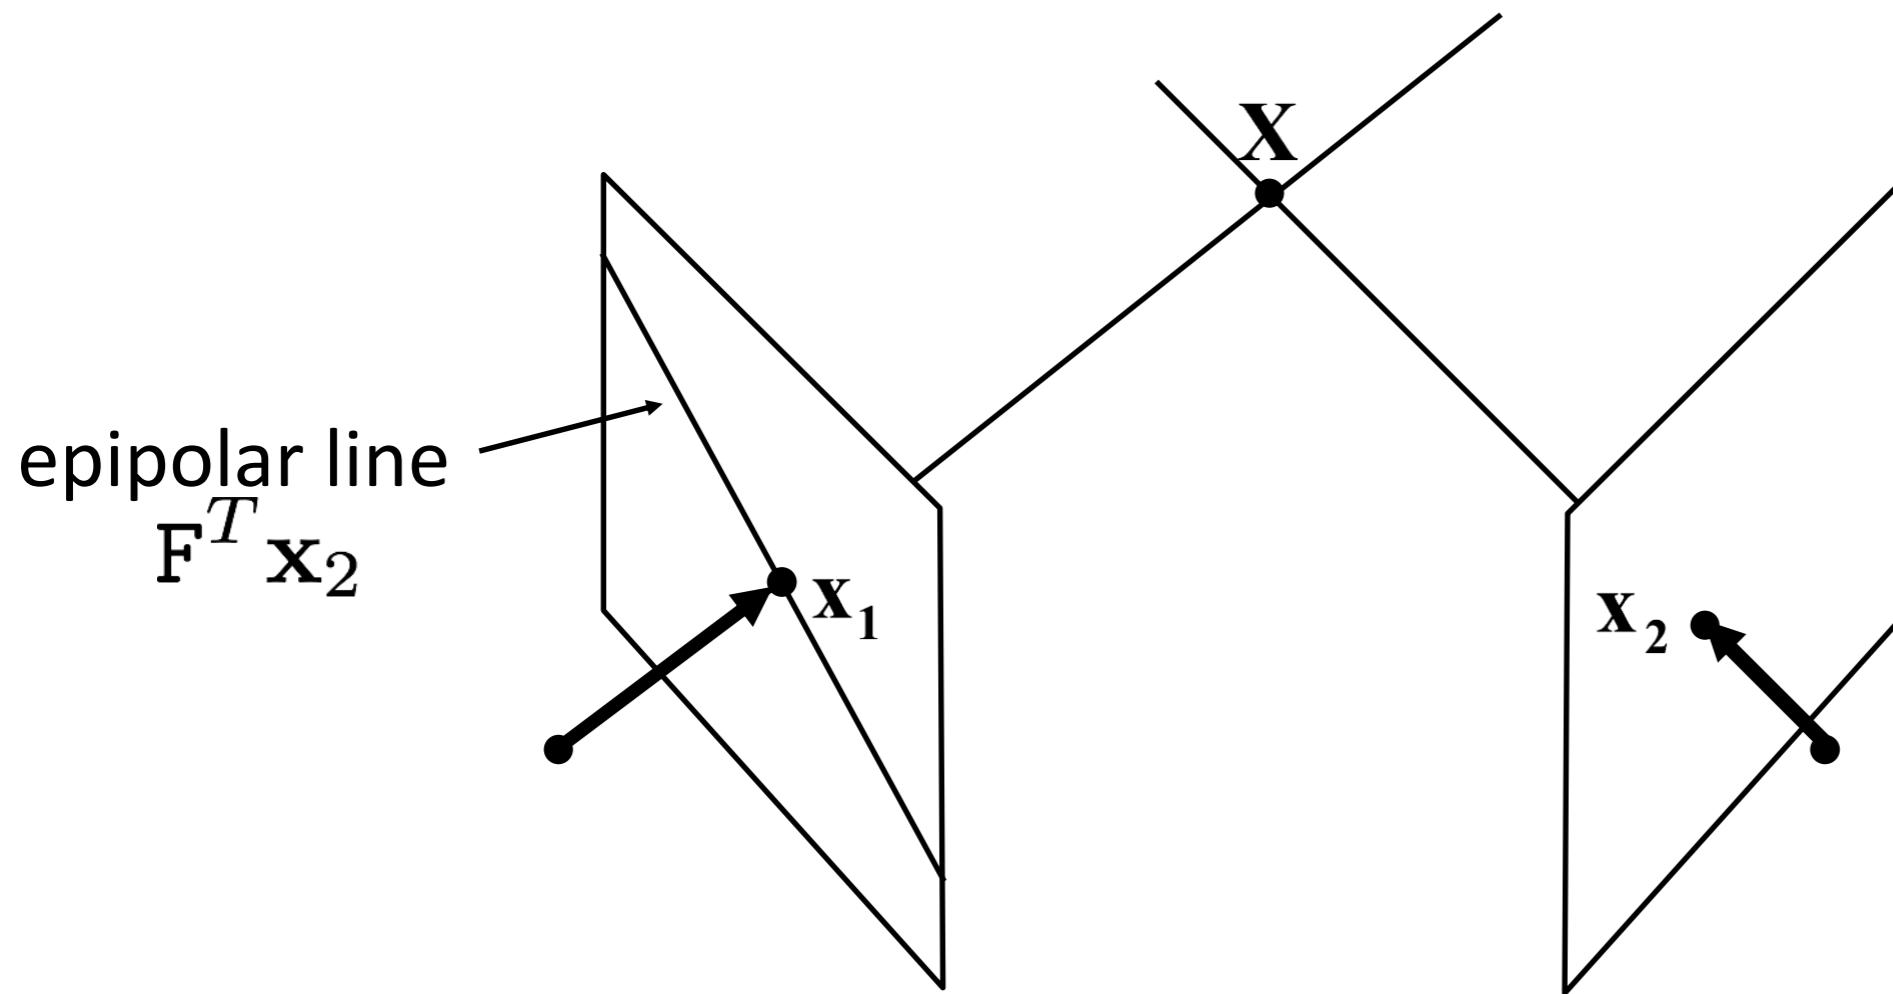
Last Lecture



$$0 = \mathbf{x}_2^T \mathbf{F} \mathbf{x}_1$$

Geometric Interpretation

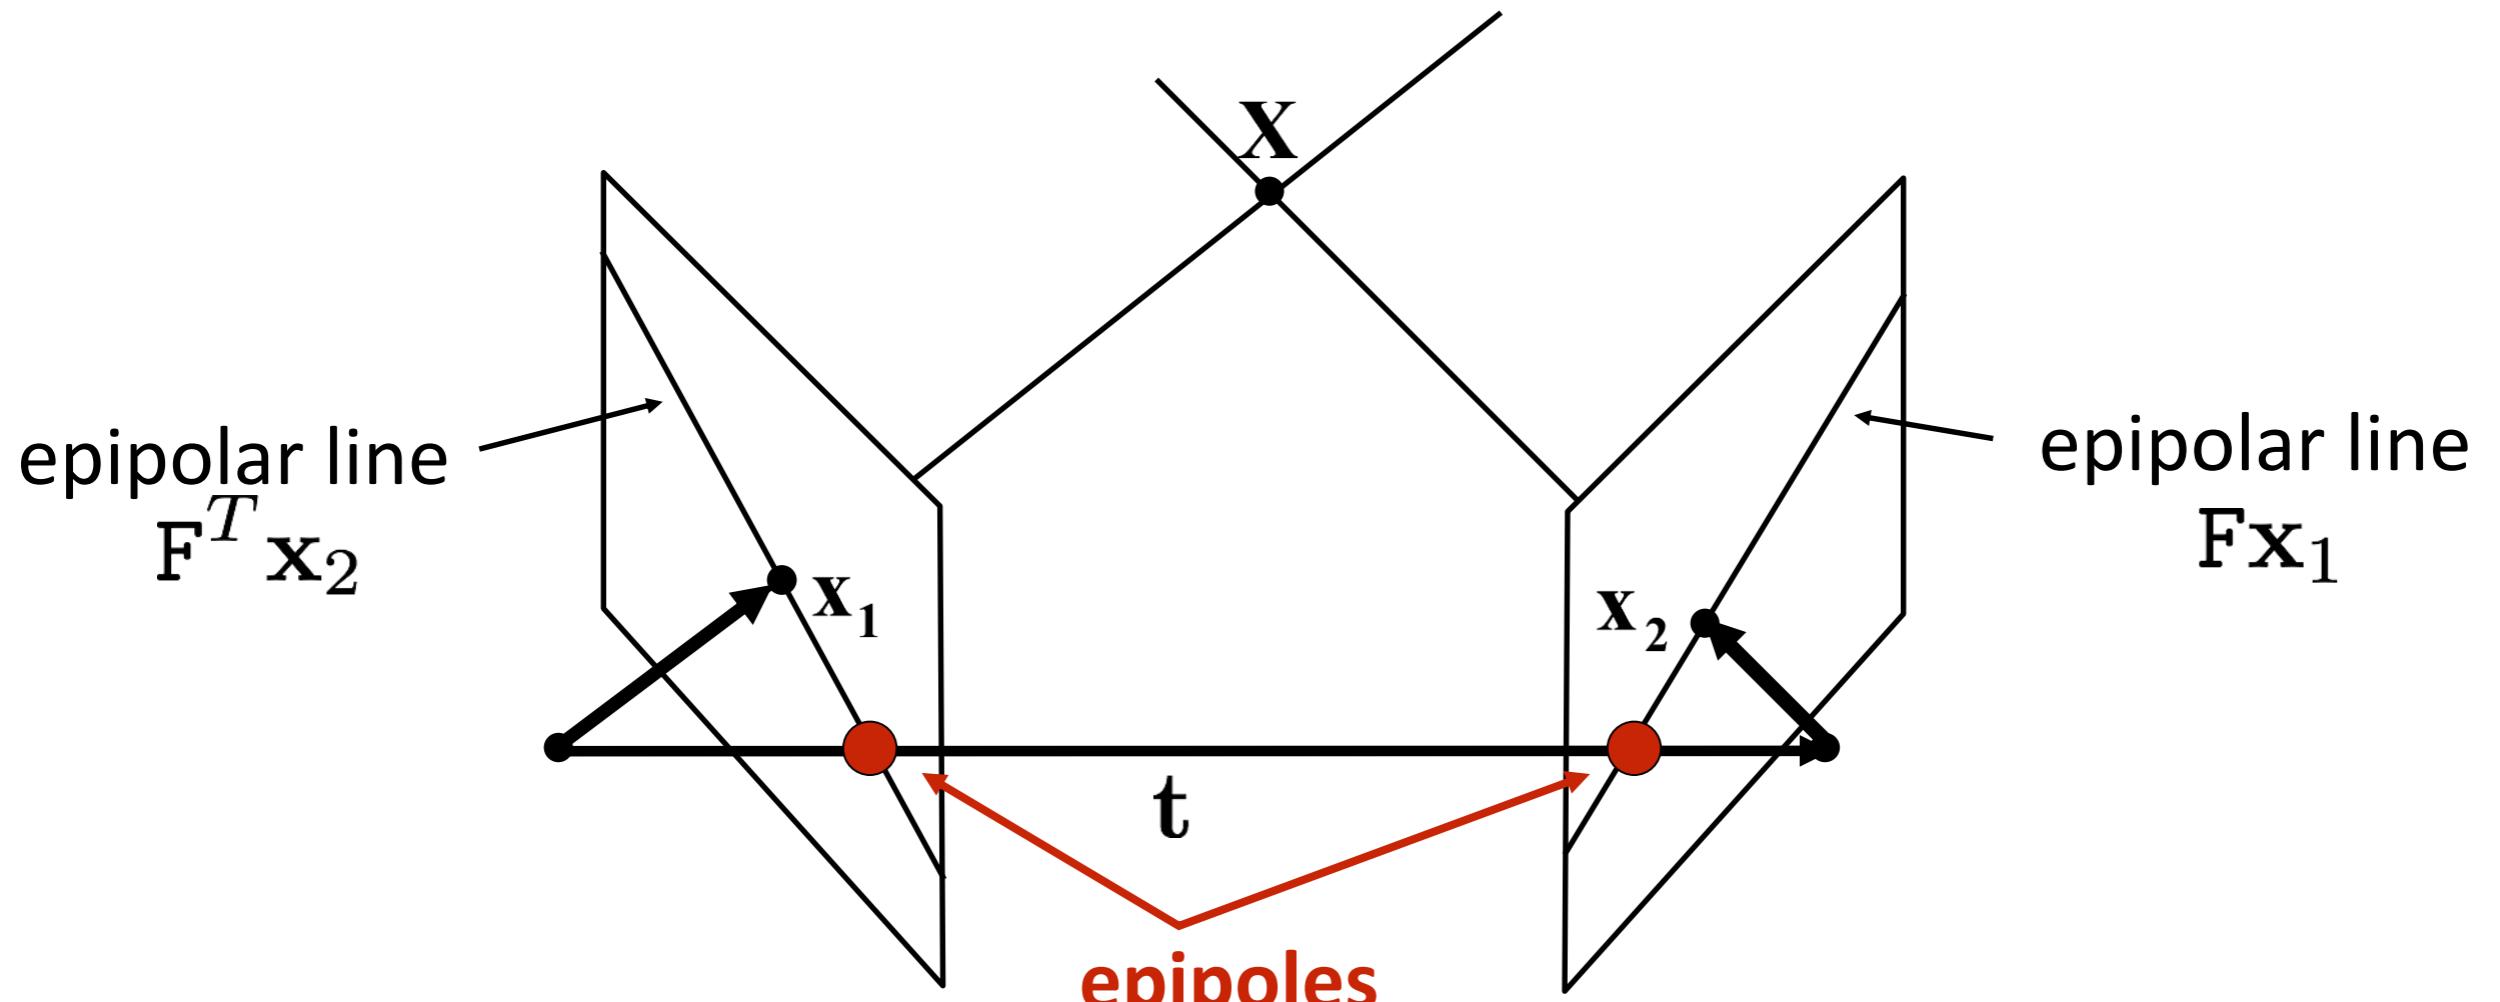
Last Lecture



$$0 = \mathbf{x}_2^T \mathbf{F} \mathbf{x}_1$$

Geometric Interpretation

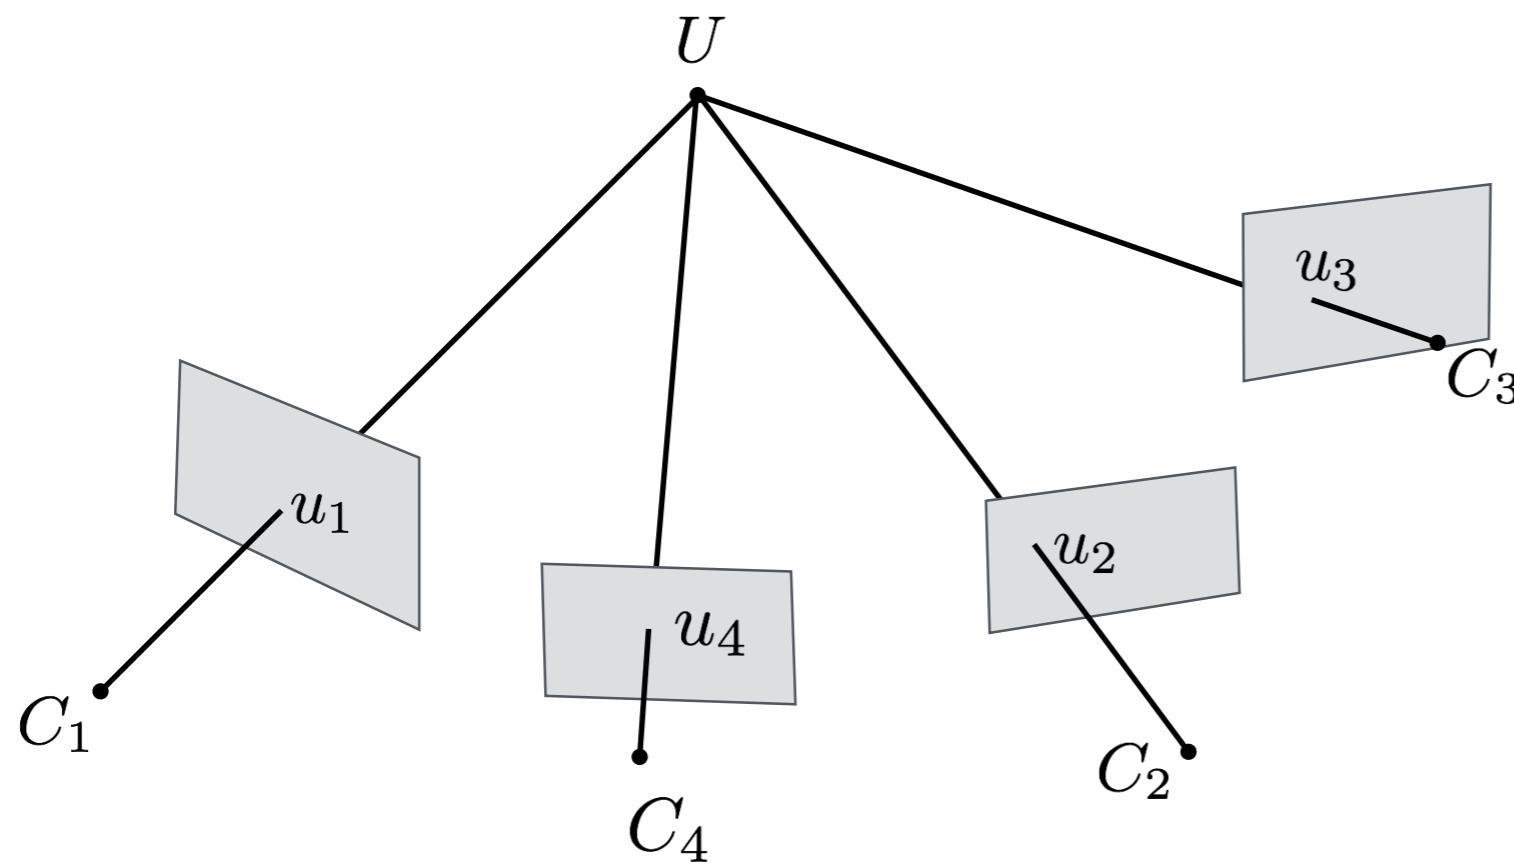
Last Lecture



$$0 = \mathbf{x}_2^T \mathbf{F} \mathbf{x}_1$$

Geometric Interpretation

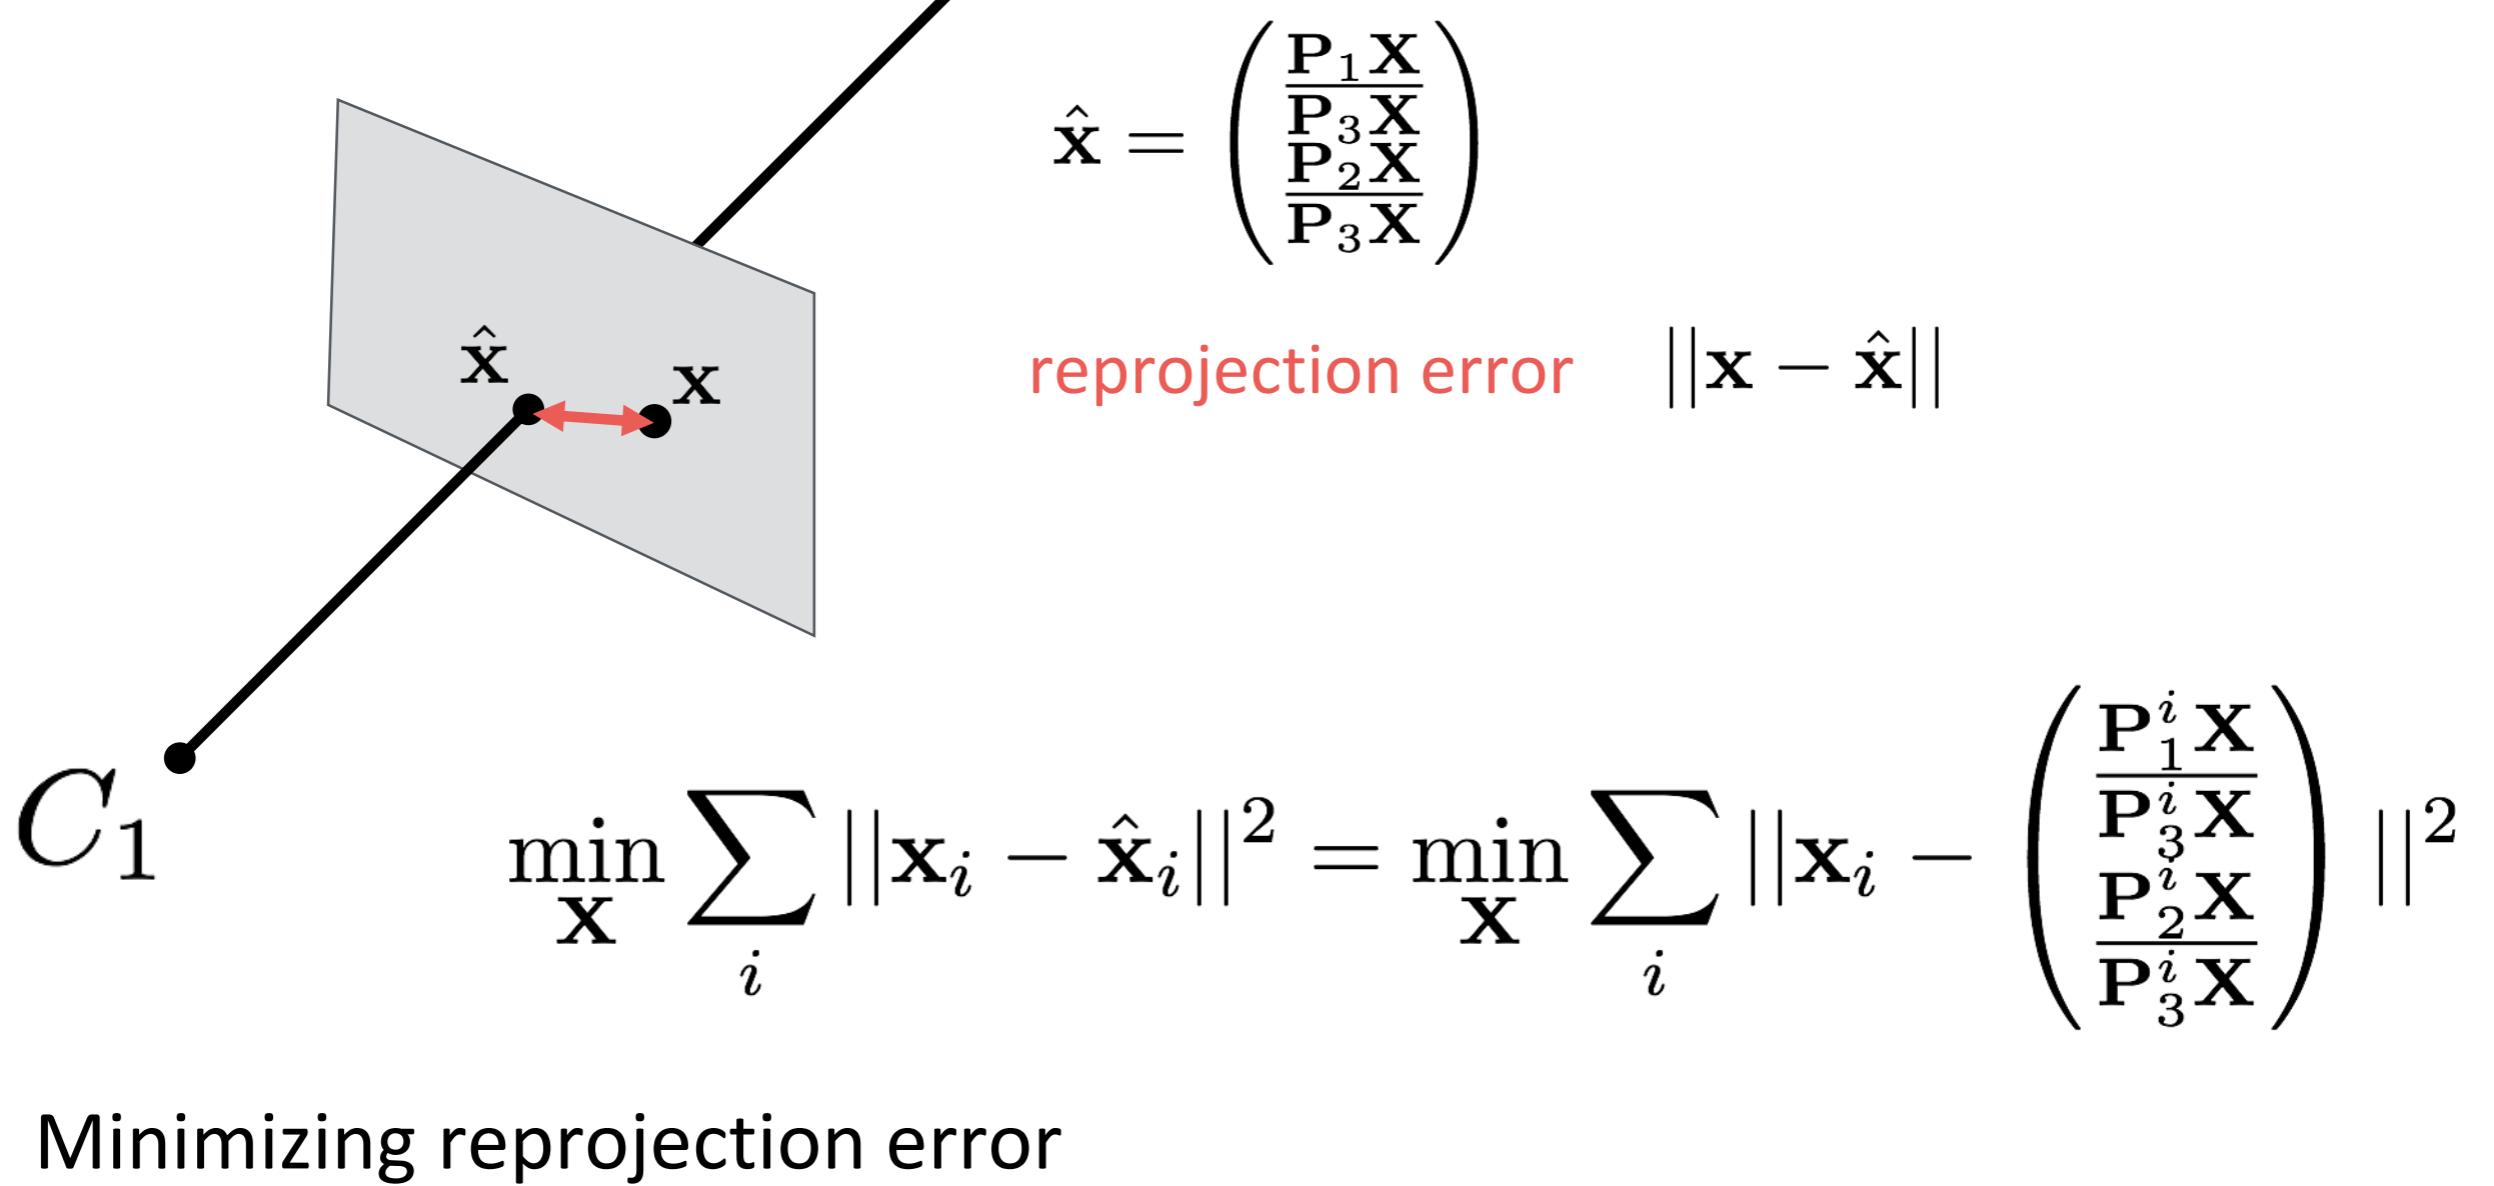
Last Lecture



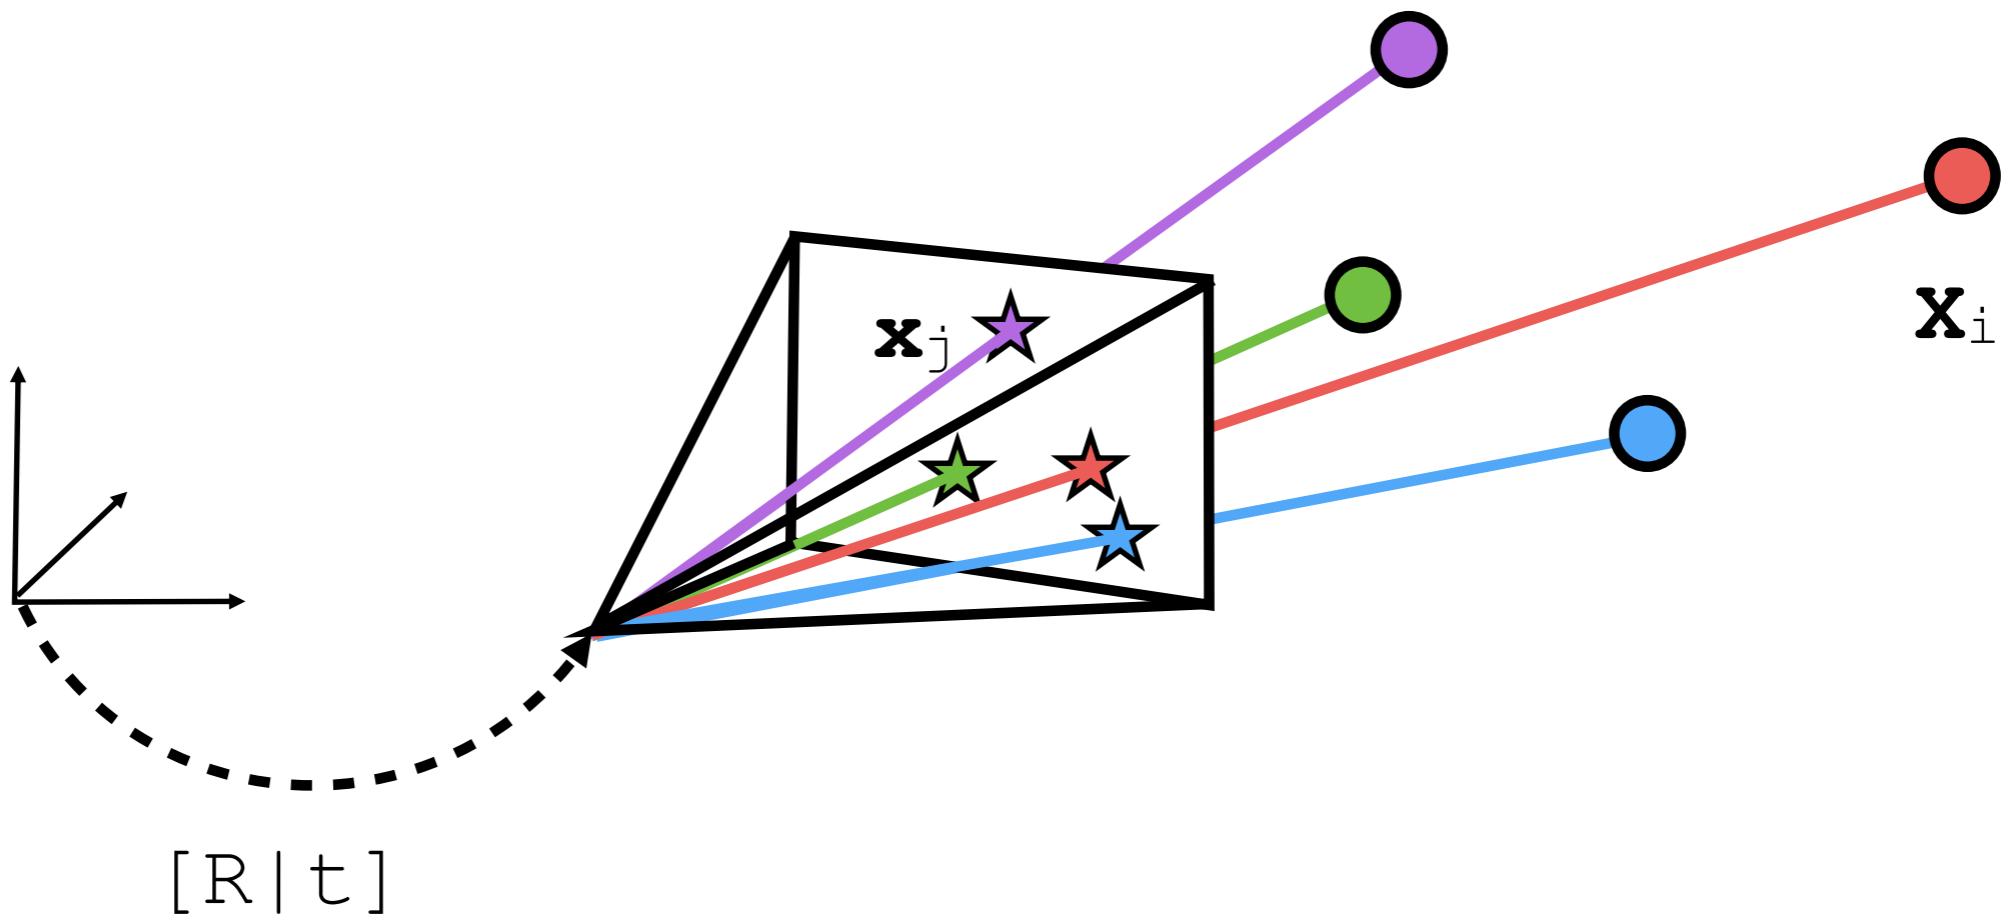
$$\begin{pmatrix} \mathbf{P}_3x - \mathbf{P}_1 \\ \mathbf{P}_3y - \mathbf{P}_2 \\ \mathbf{P}'_3x - \mathbf{P}'_1 \\ \mathbf{P}'_3y - \mathbf{P}'_2 \end{pmatrix} \mathbf{X} = 0$$

Triangulation

Last Lecture

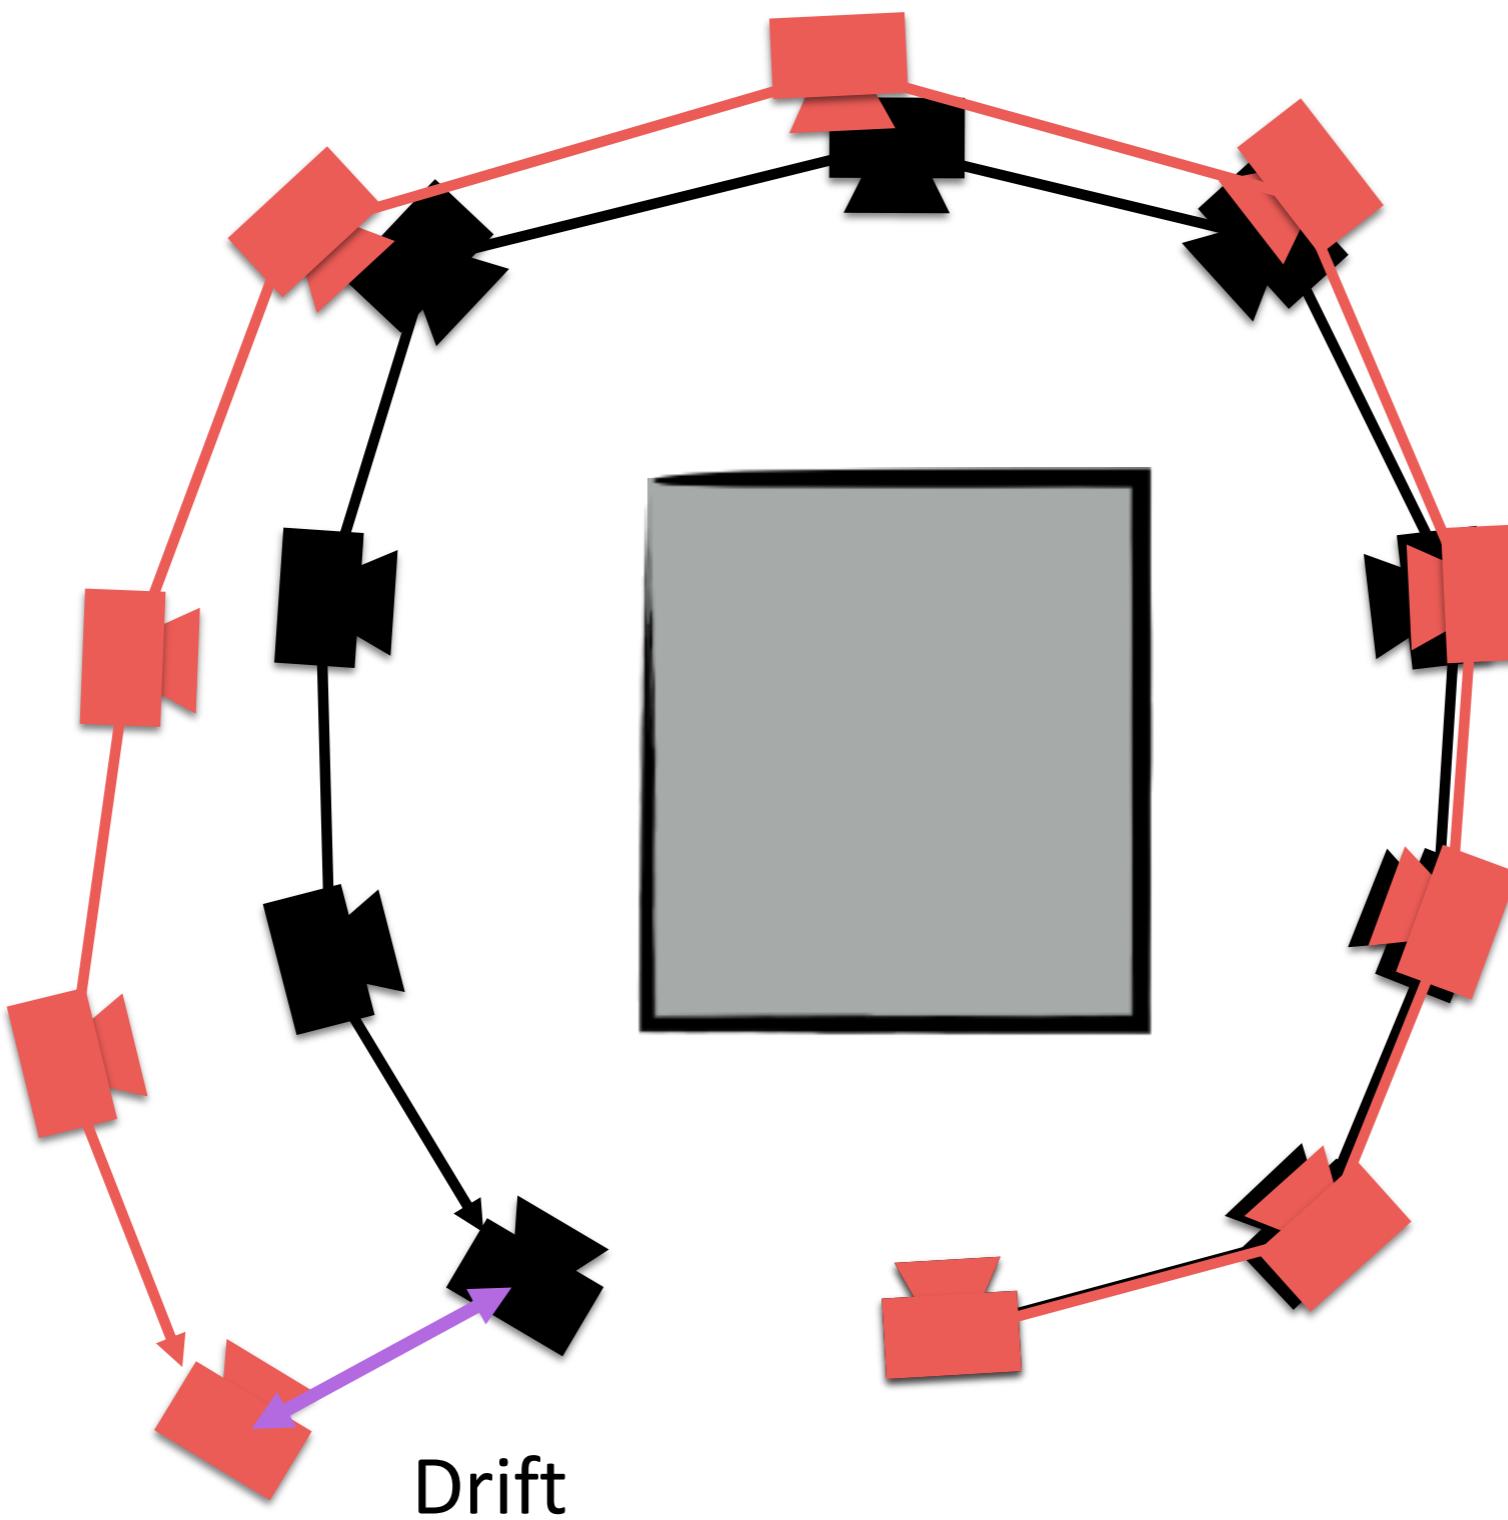


Last Lecture

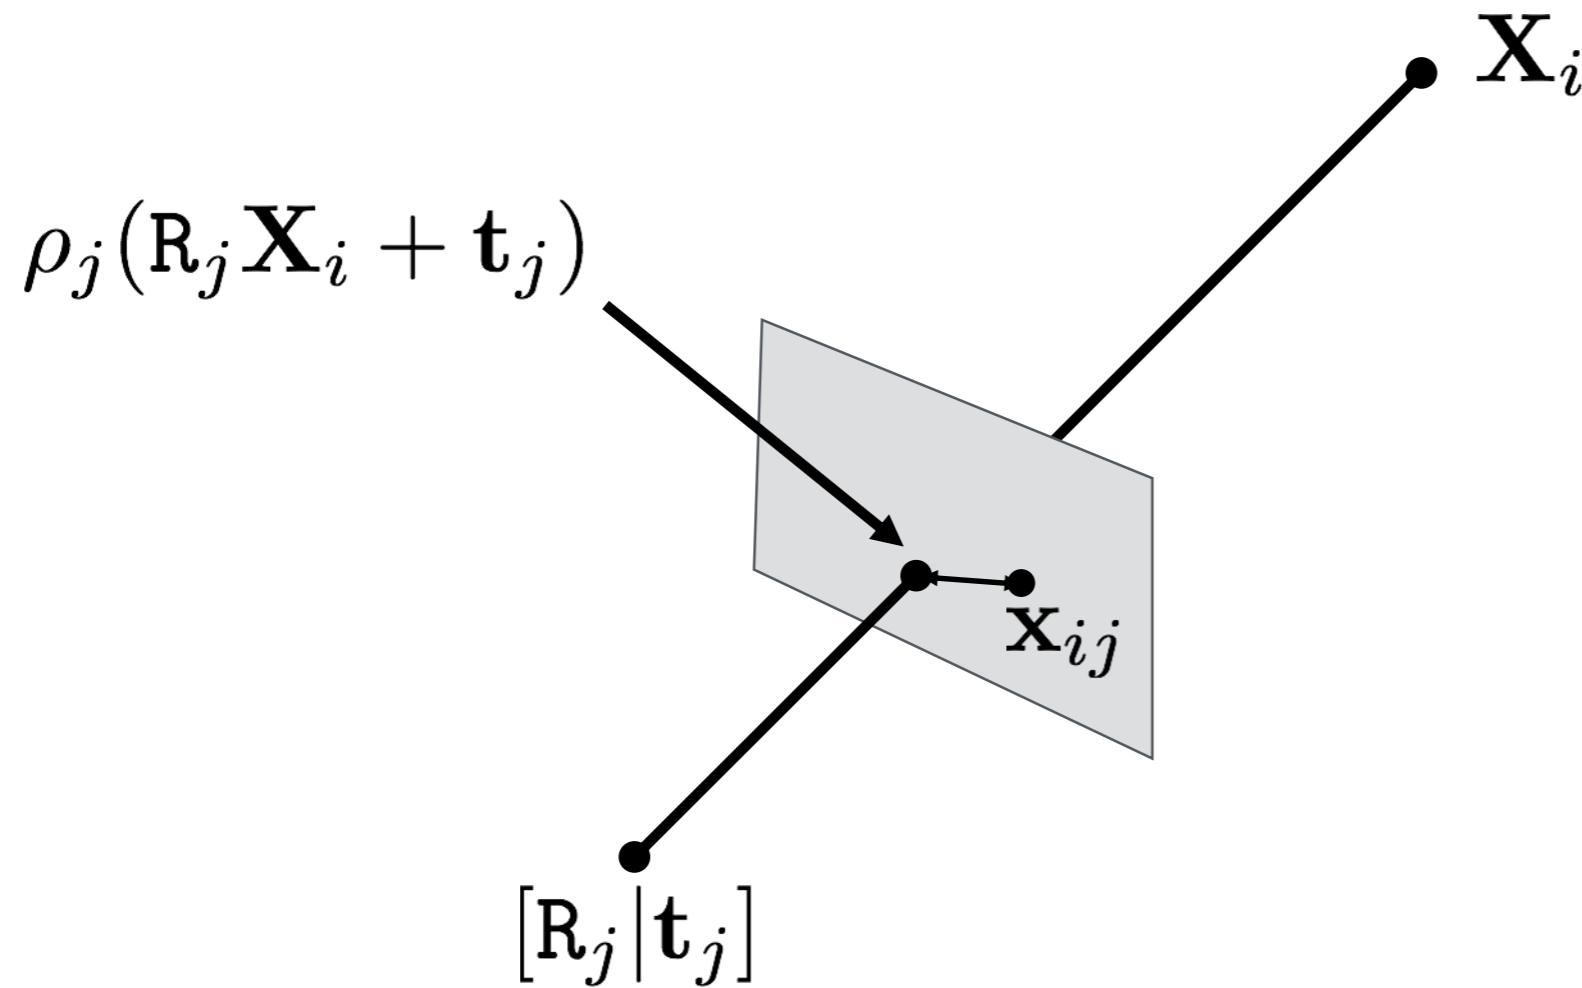


Absolute pose estimation

Last Lecture



Last Lecture

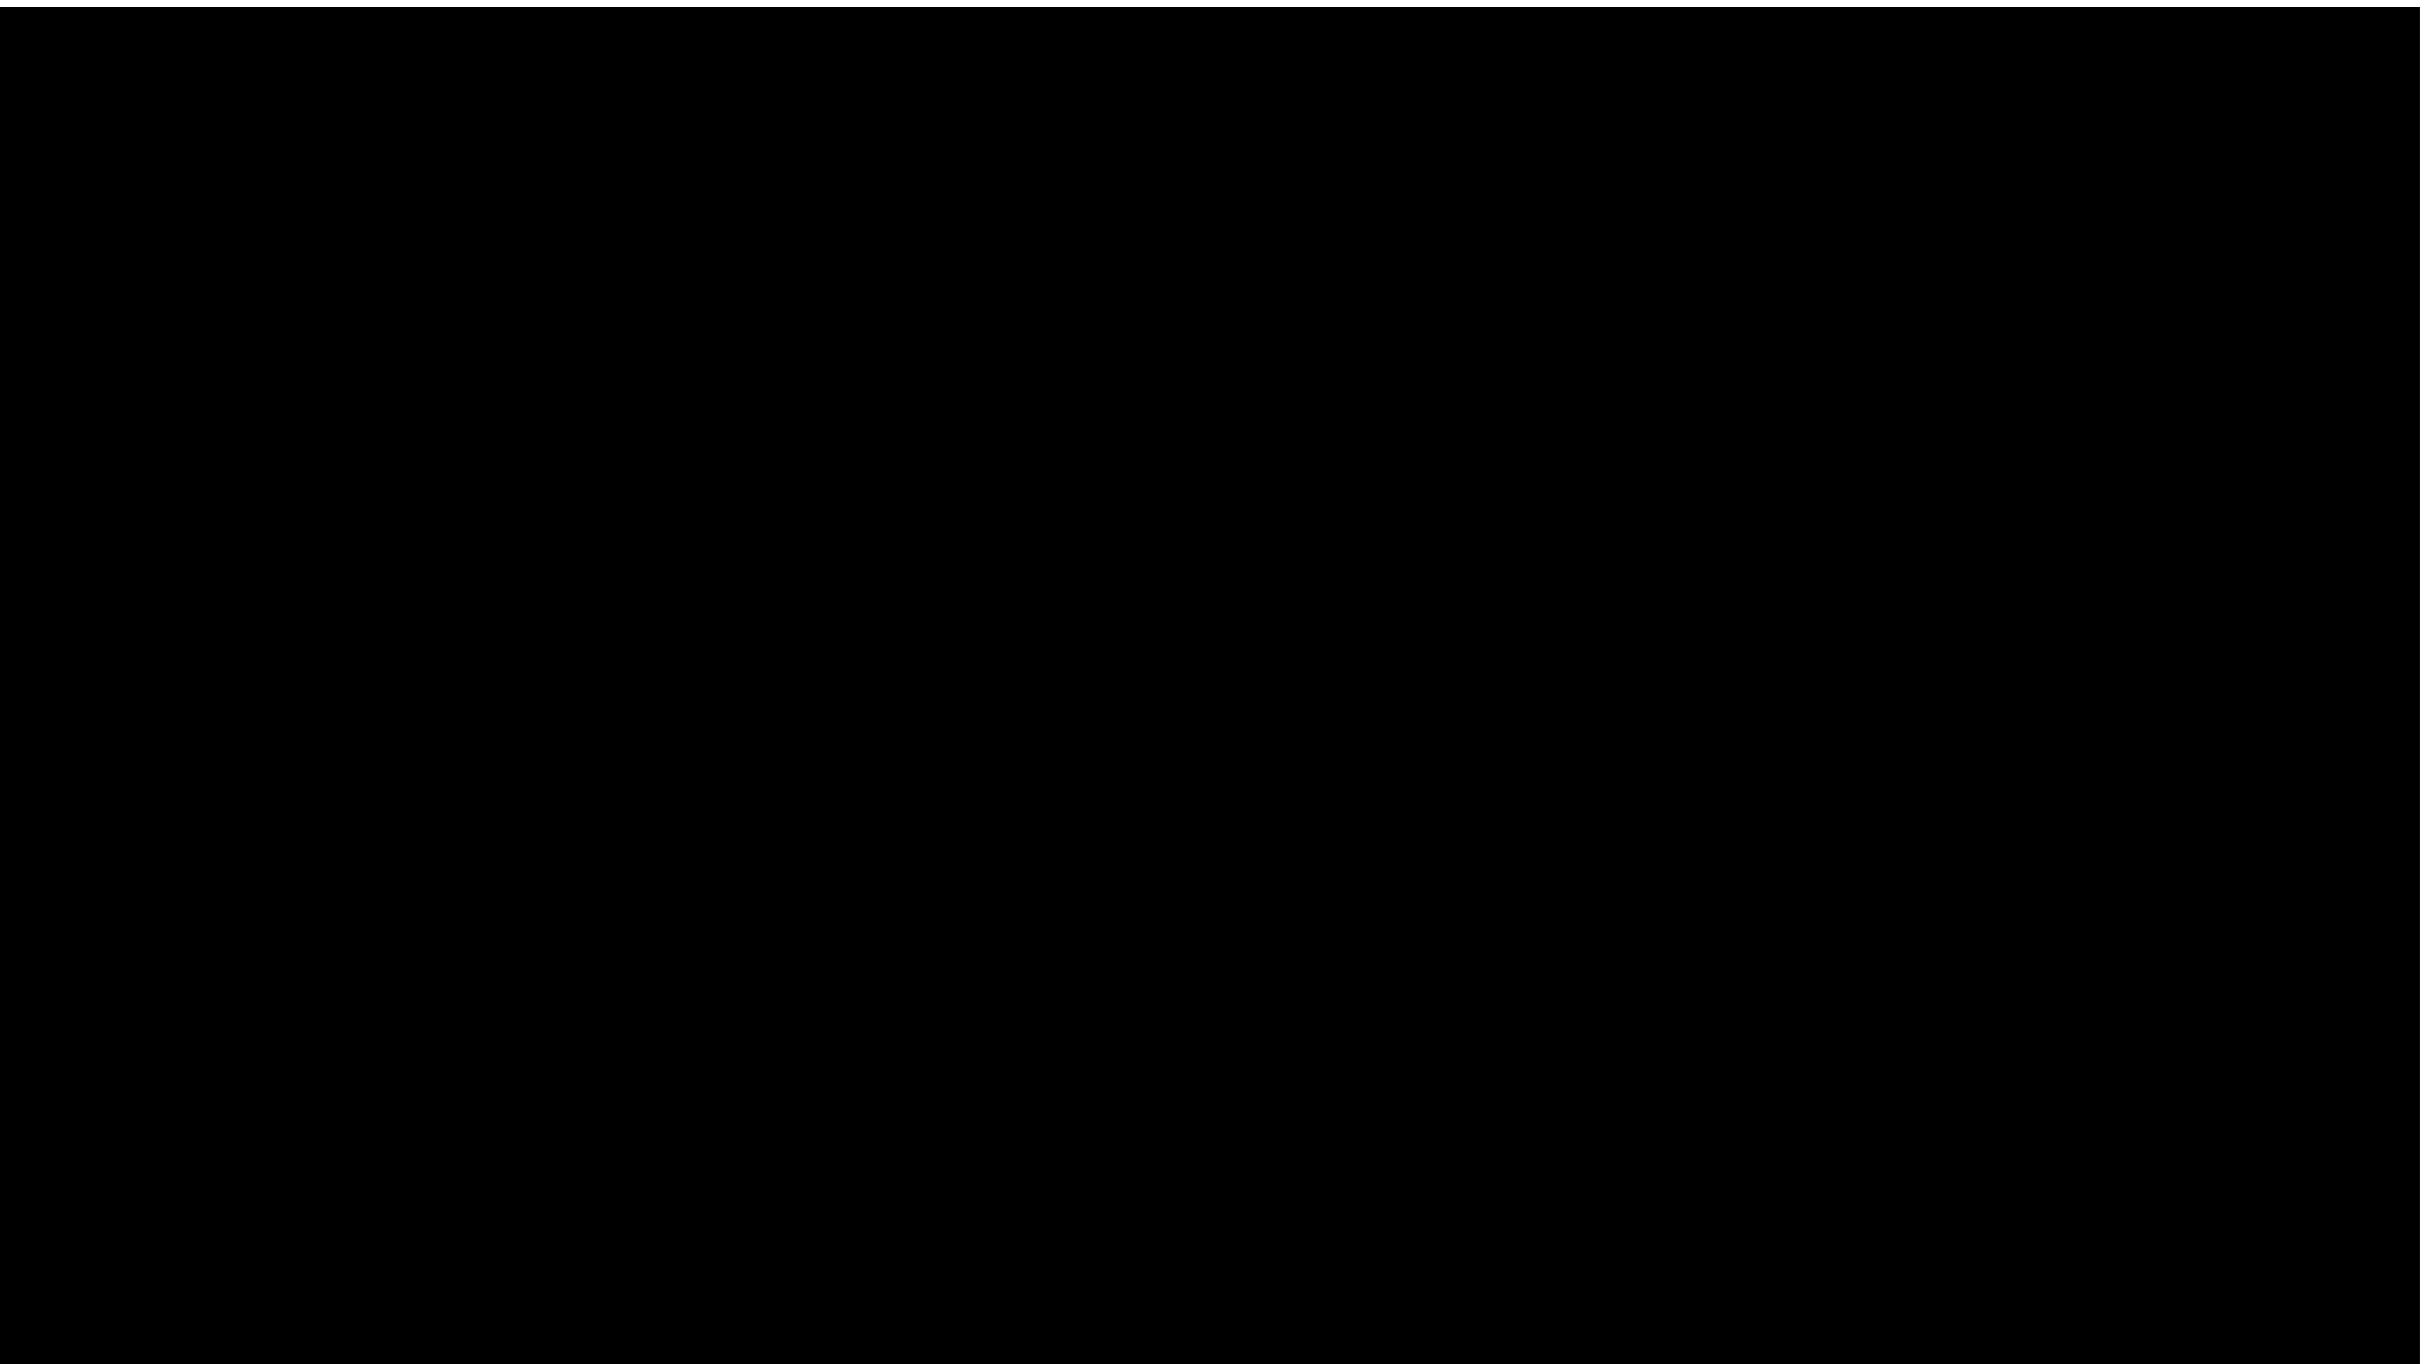


$\underset{\text{camera poses, 3D points}}{\operatorname{argmin}}$
Bundle adjustment

$$\sum_i \sum_j \Delta_{ij} \|\mathbf{x}_{ij} - \rho_j(\mathbf{R}_j \mathbf{X}_i + \mathbf{t}_j)\|^2$$

↑
point i visible in image j?

Last Lecture

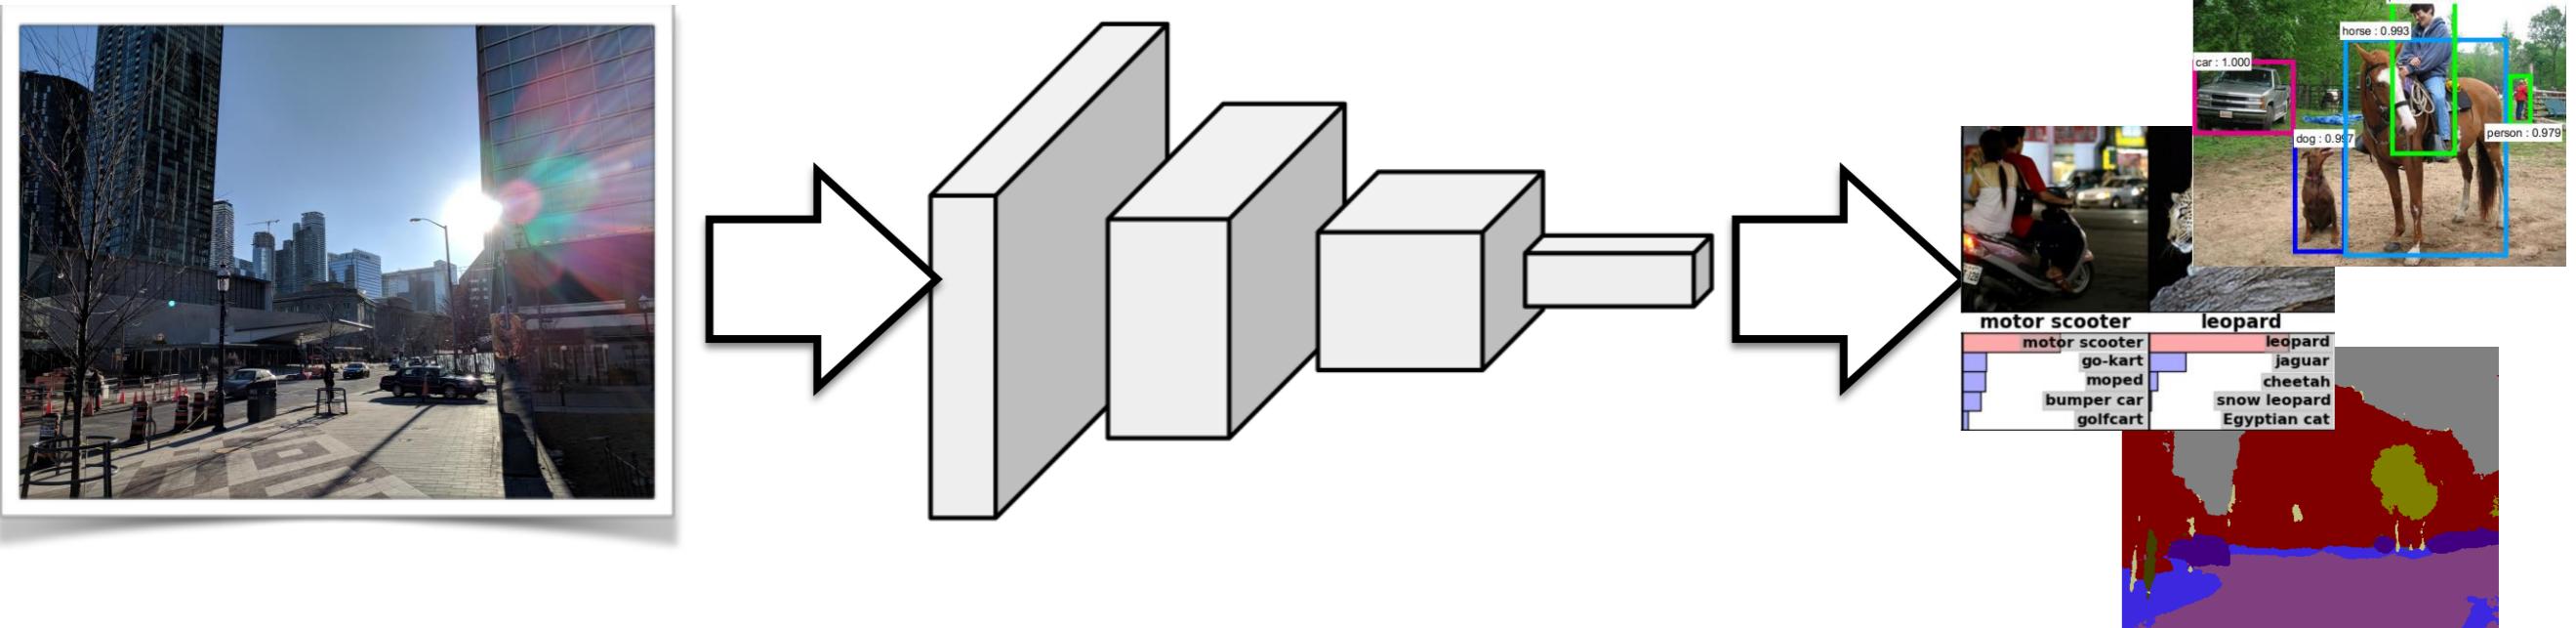


Bundle adjustment

Today

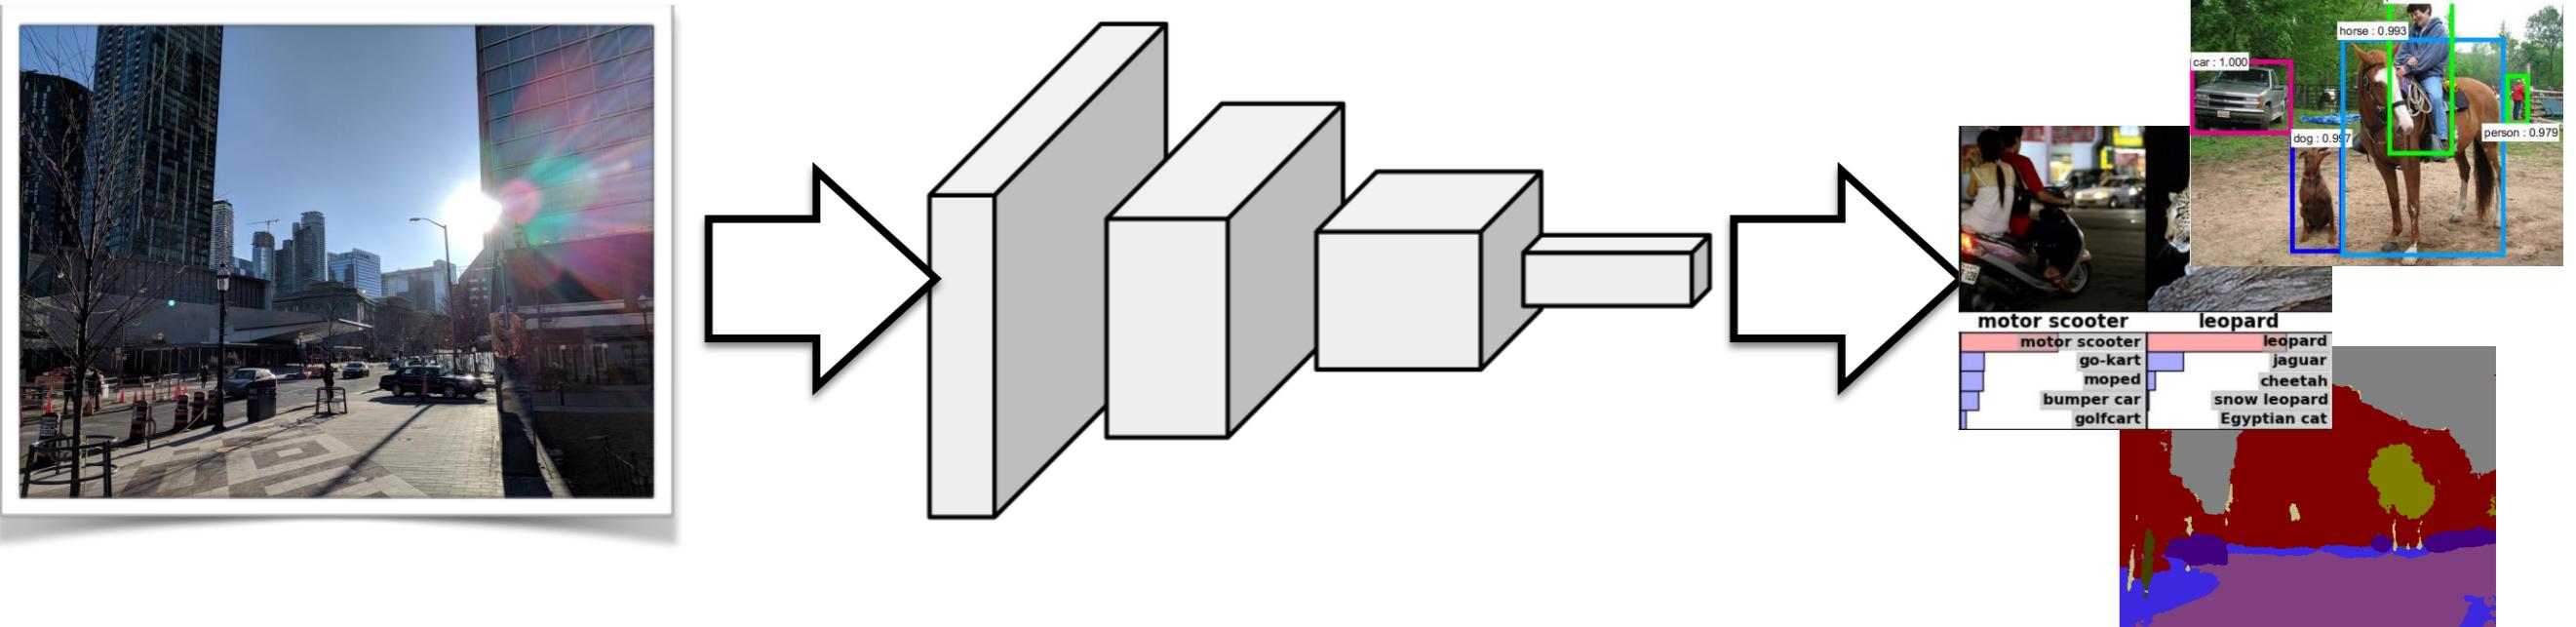
Generative Neural Networks

Recap: CNNs



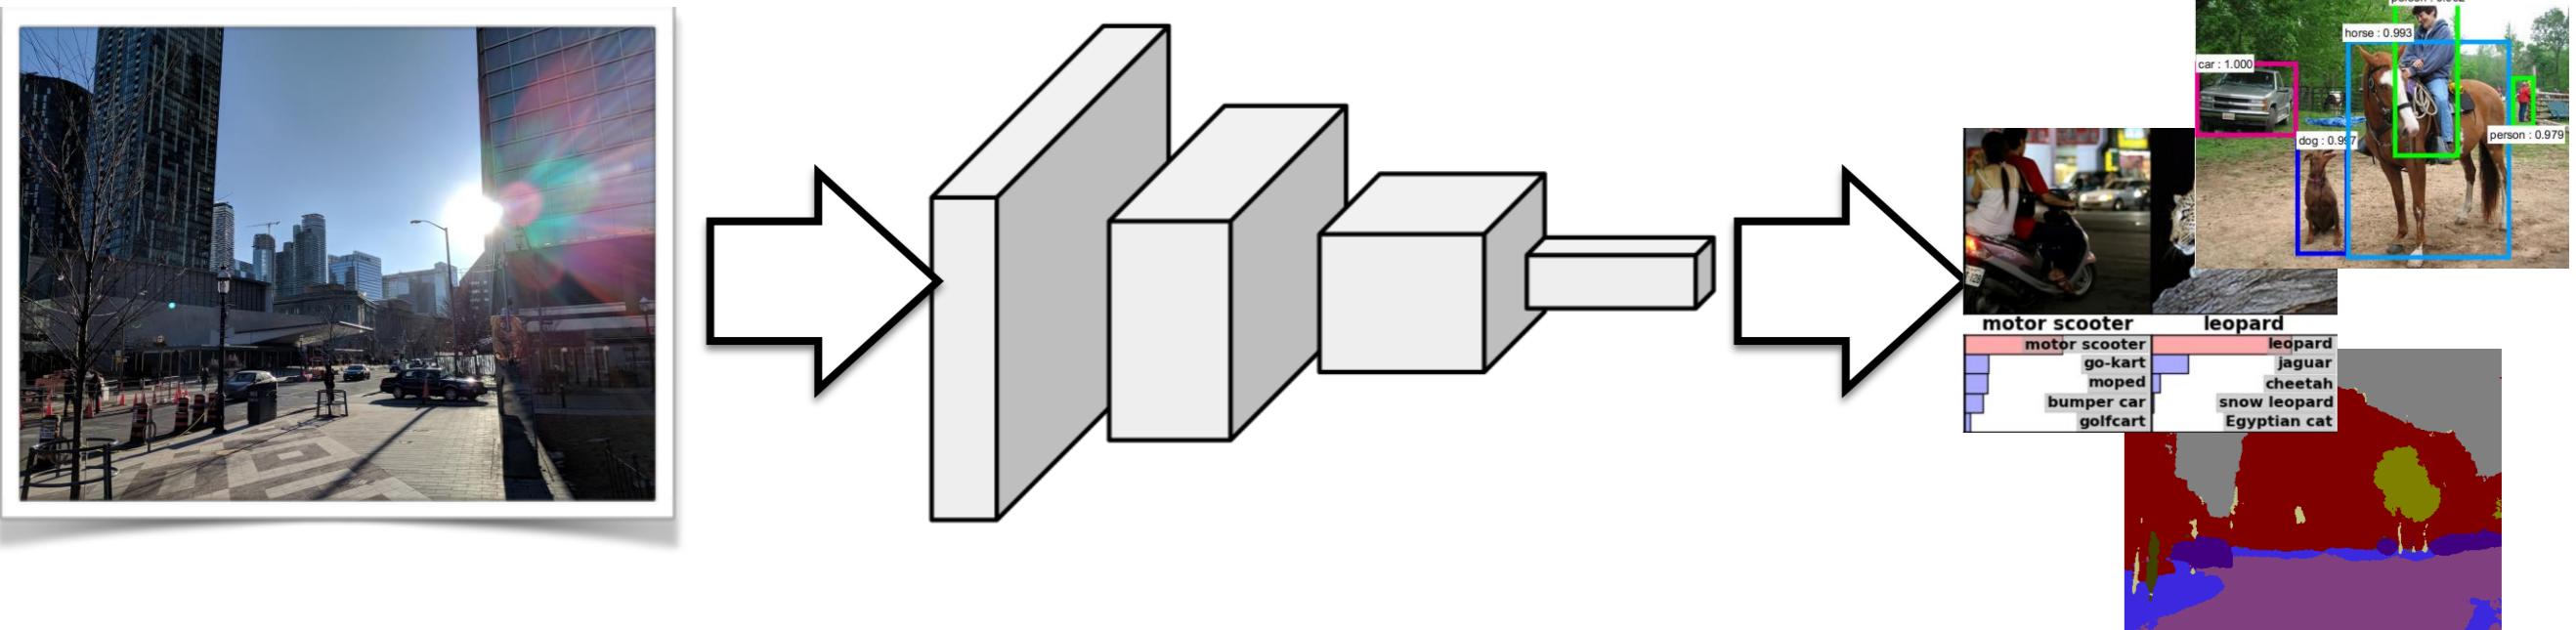
- Convolutional neural networks for classification, segmentation, regression, etc.

Recap: CNNs



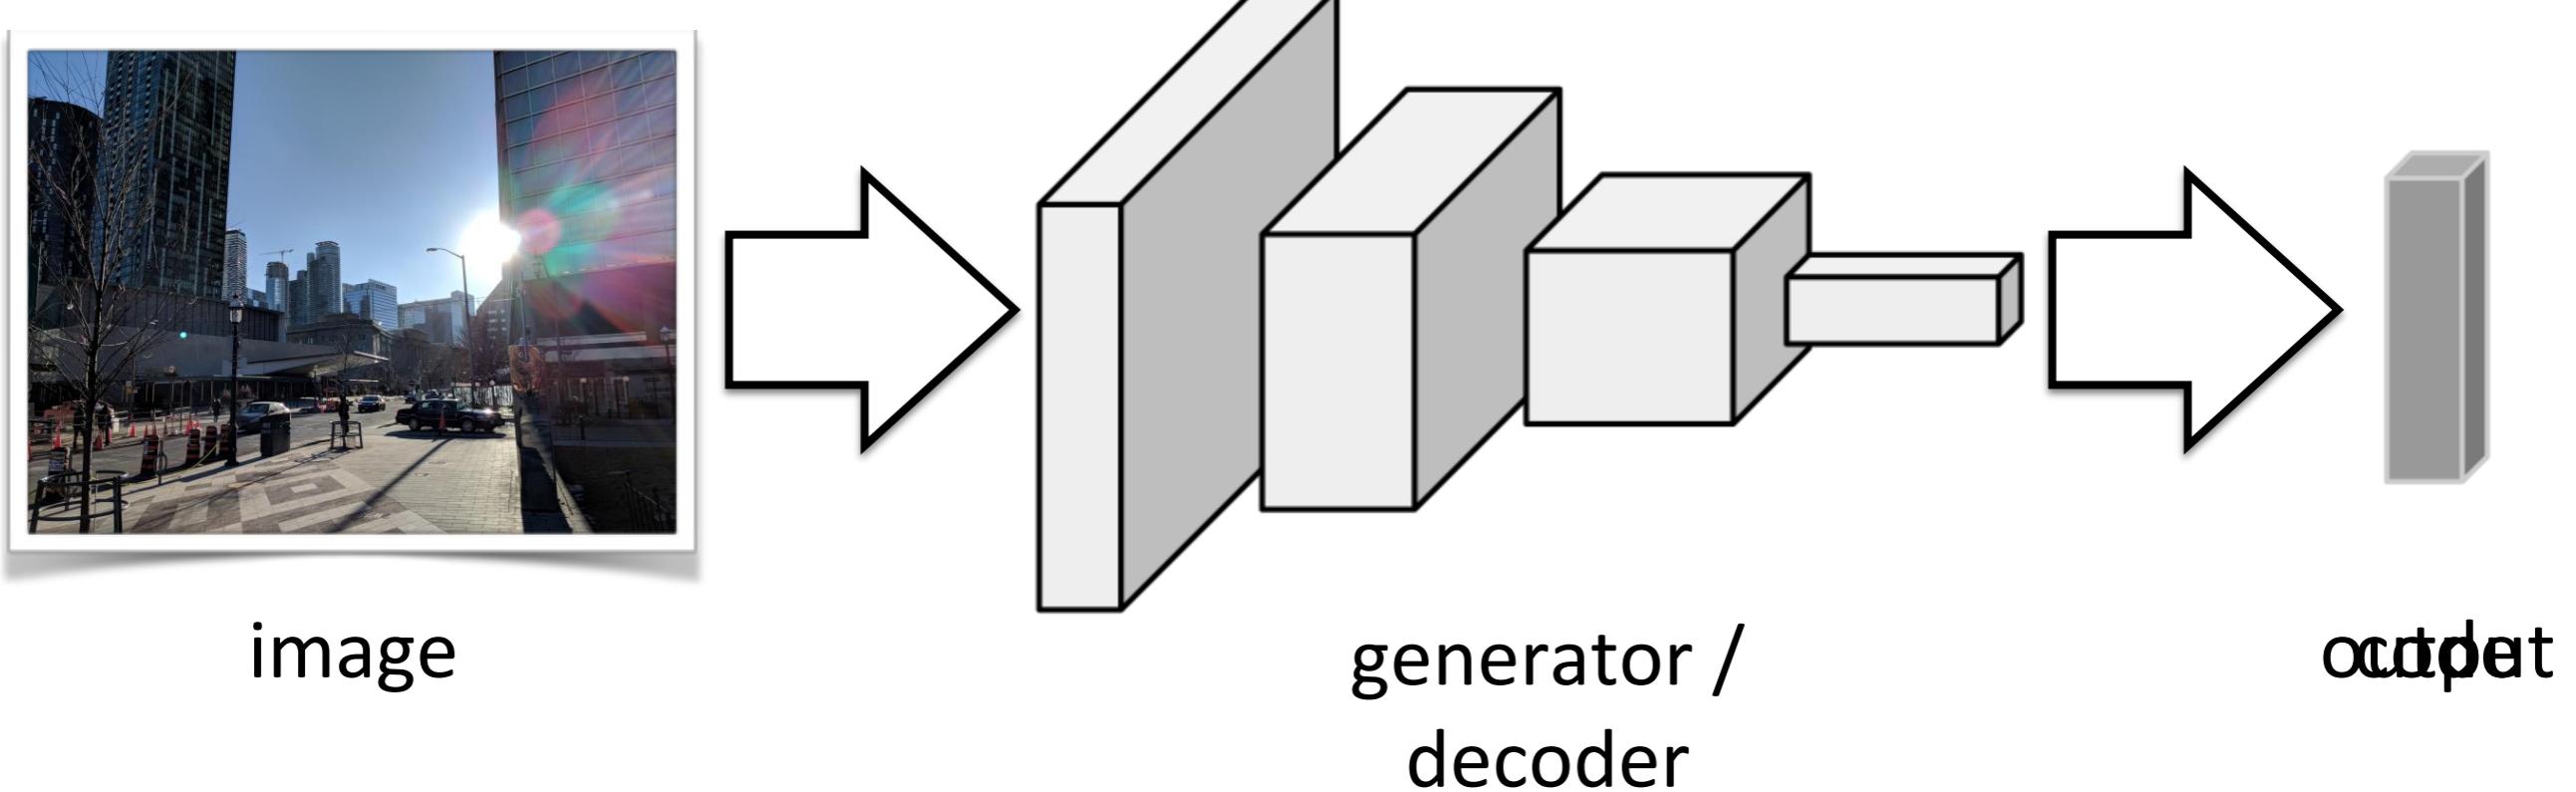
- Convolutional neural networks for classification, segmentation, regression, etc.
- Supervised learning: Training from (image, label) pairs

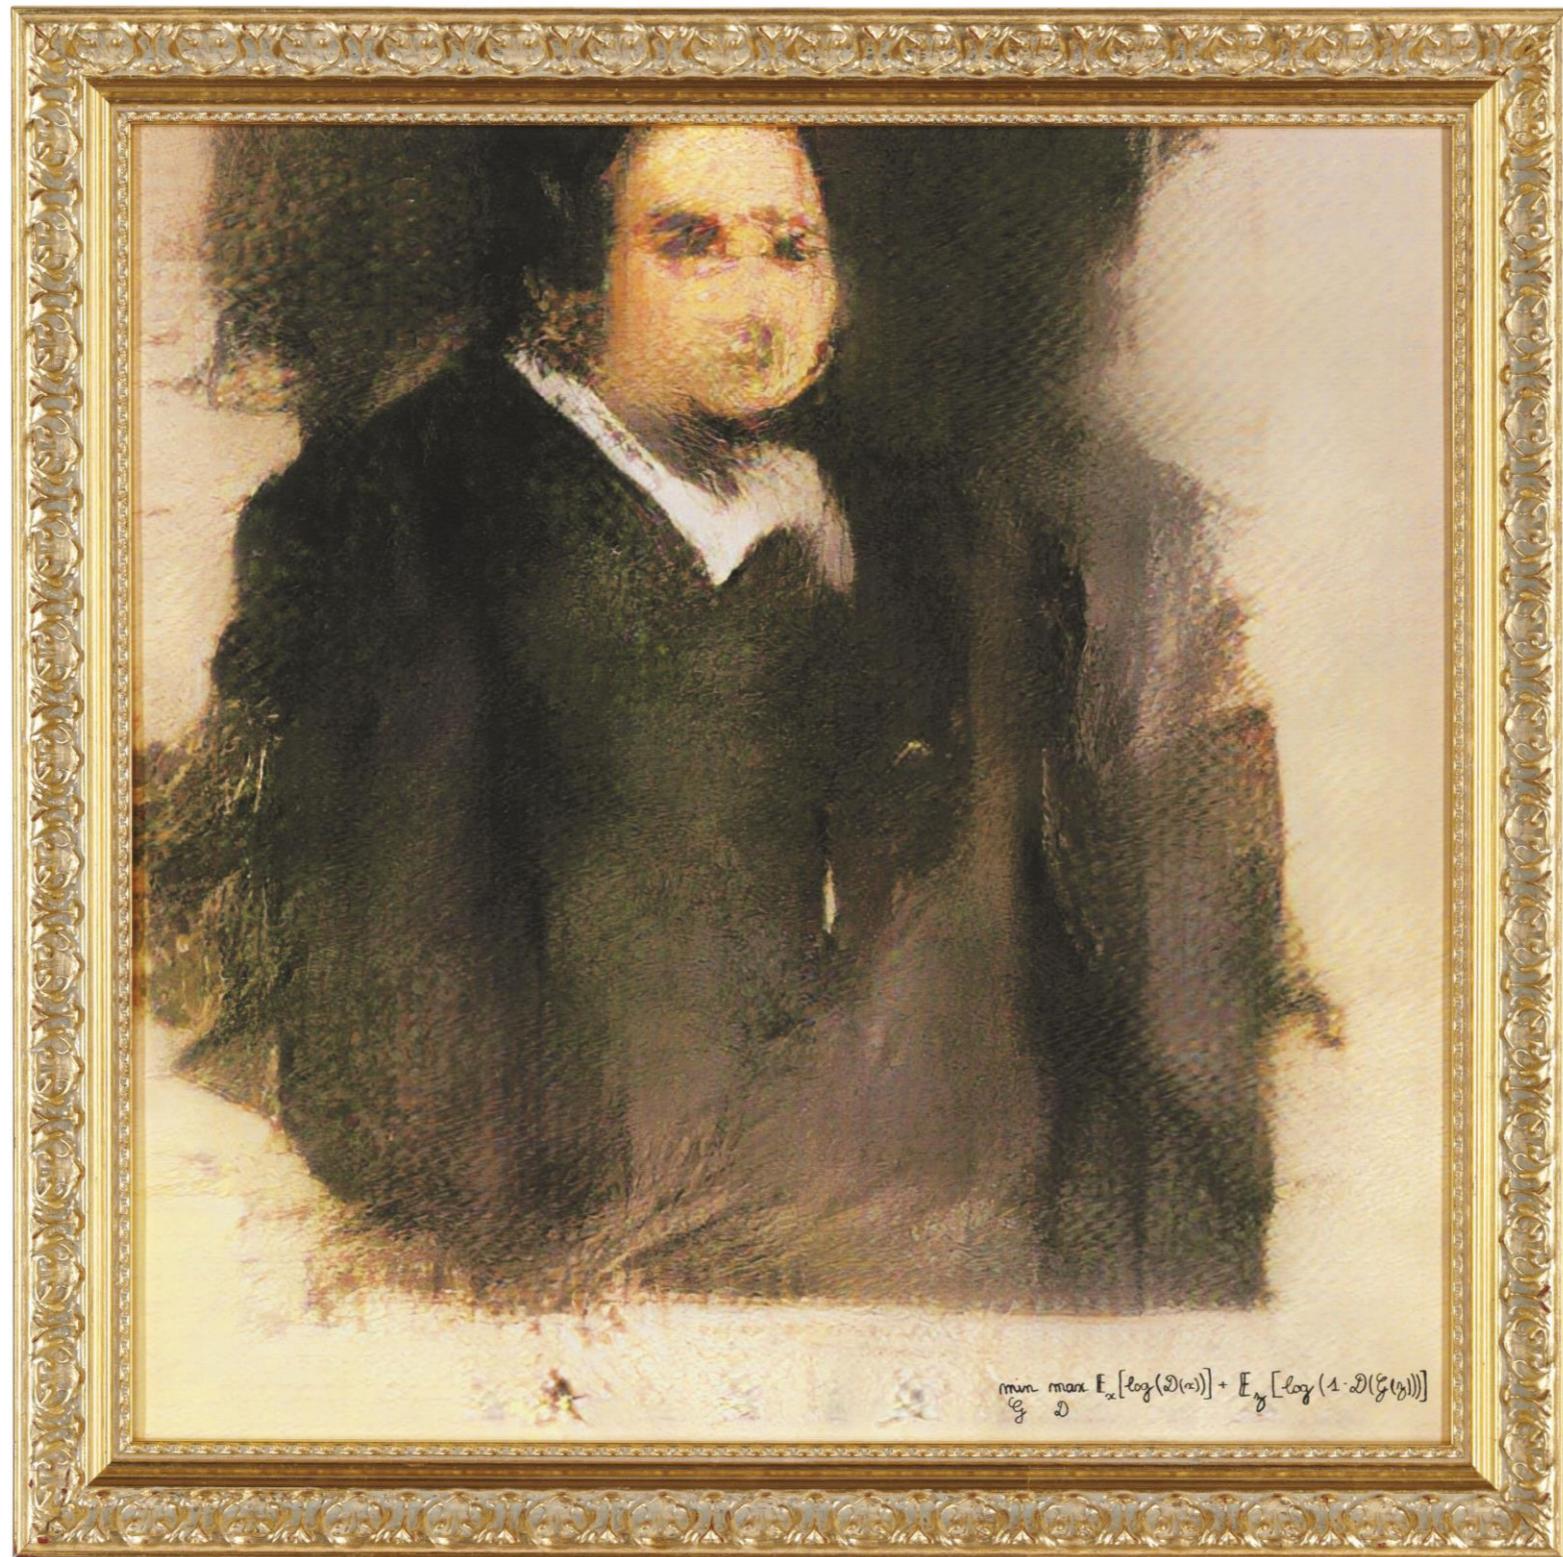
Recap: CNNs



- Convolutional neural networks for classification, segmentation, regression, etc.
- Supervised learning: Training from (image, label) pairs
- Training via stochastic gradient descent (or variants)

Generative Neural Networks





[Portrait of Edmond Belamy, 2018](#). Sold for \$432,500 on 25 October 2018 at

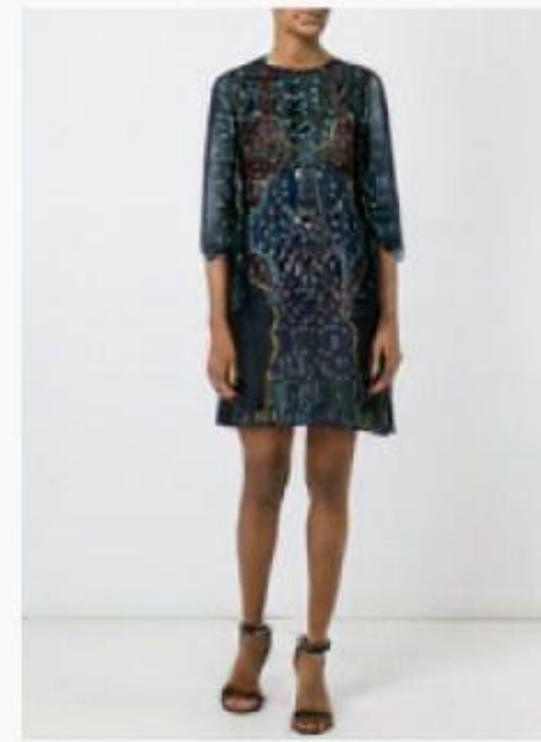
Christie's in New York. Image © Obvious

Automatic Image Generation



[Nguyen et al., Plug & Play Generative Networks: Conditional Iterative Generation of Images in Latent Space, CVPR 2017]

Commercial Applications



(vue.ai, covered in [Quartz](#))

[slide credit: Ian Goodfellow]

Commercial Applications



Image-to-Image Translation



[Zhu et al., Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, ICCV 2017]

Image-to-Image Translation



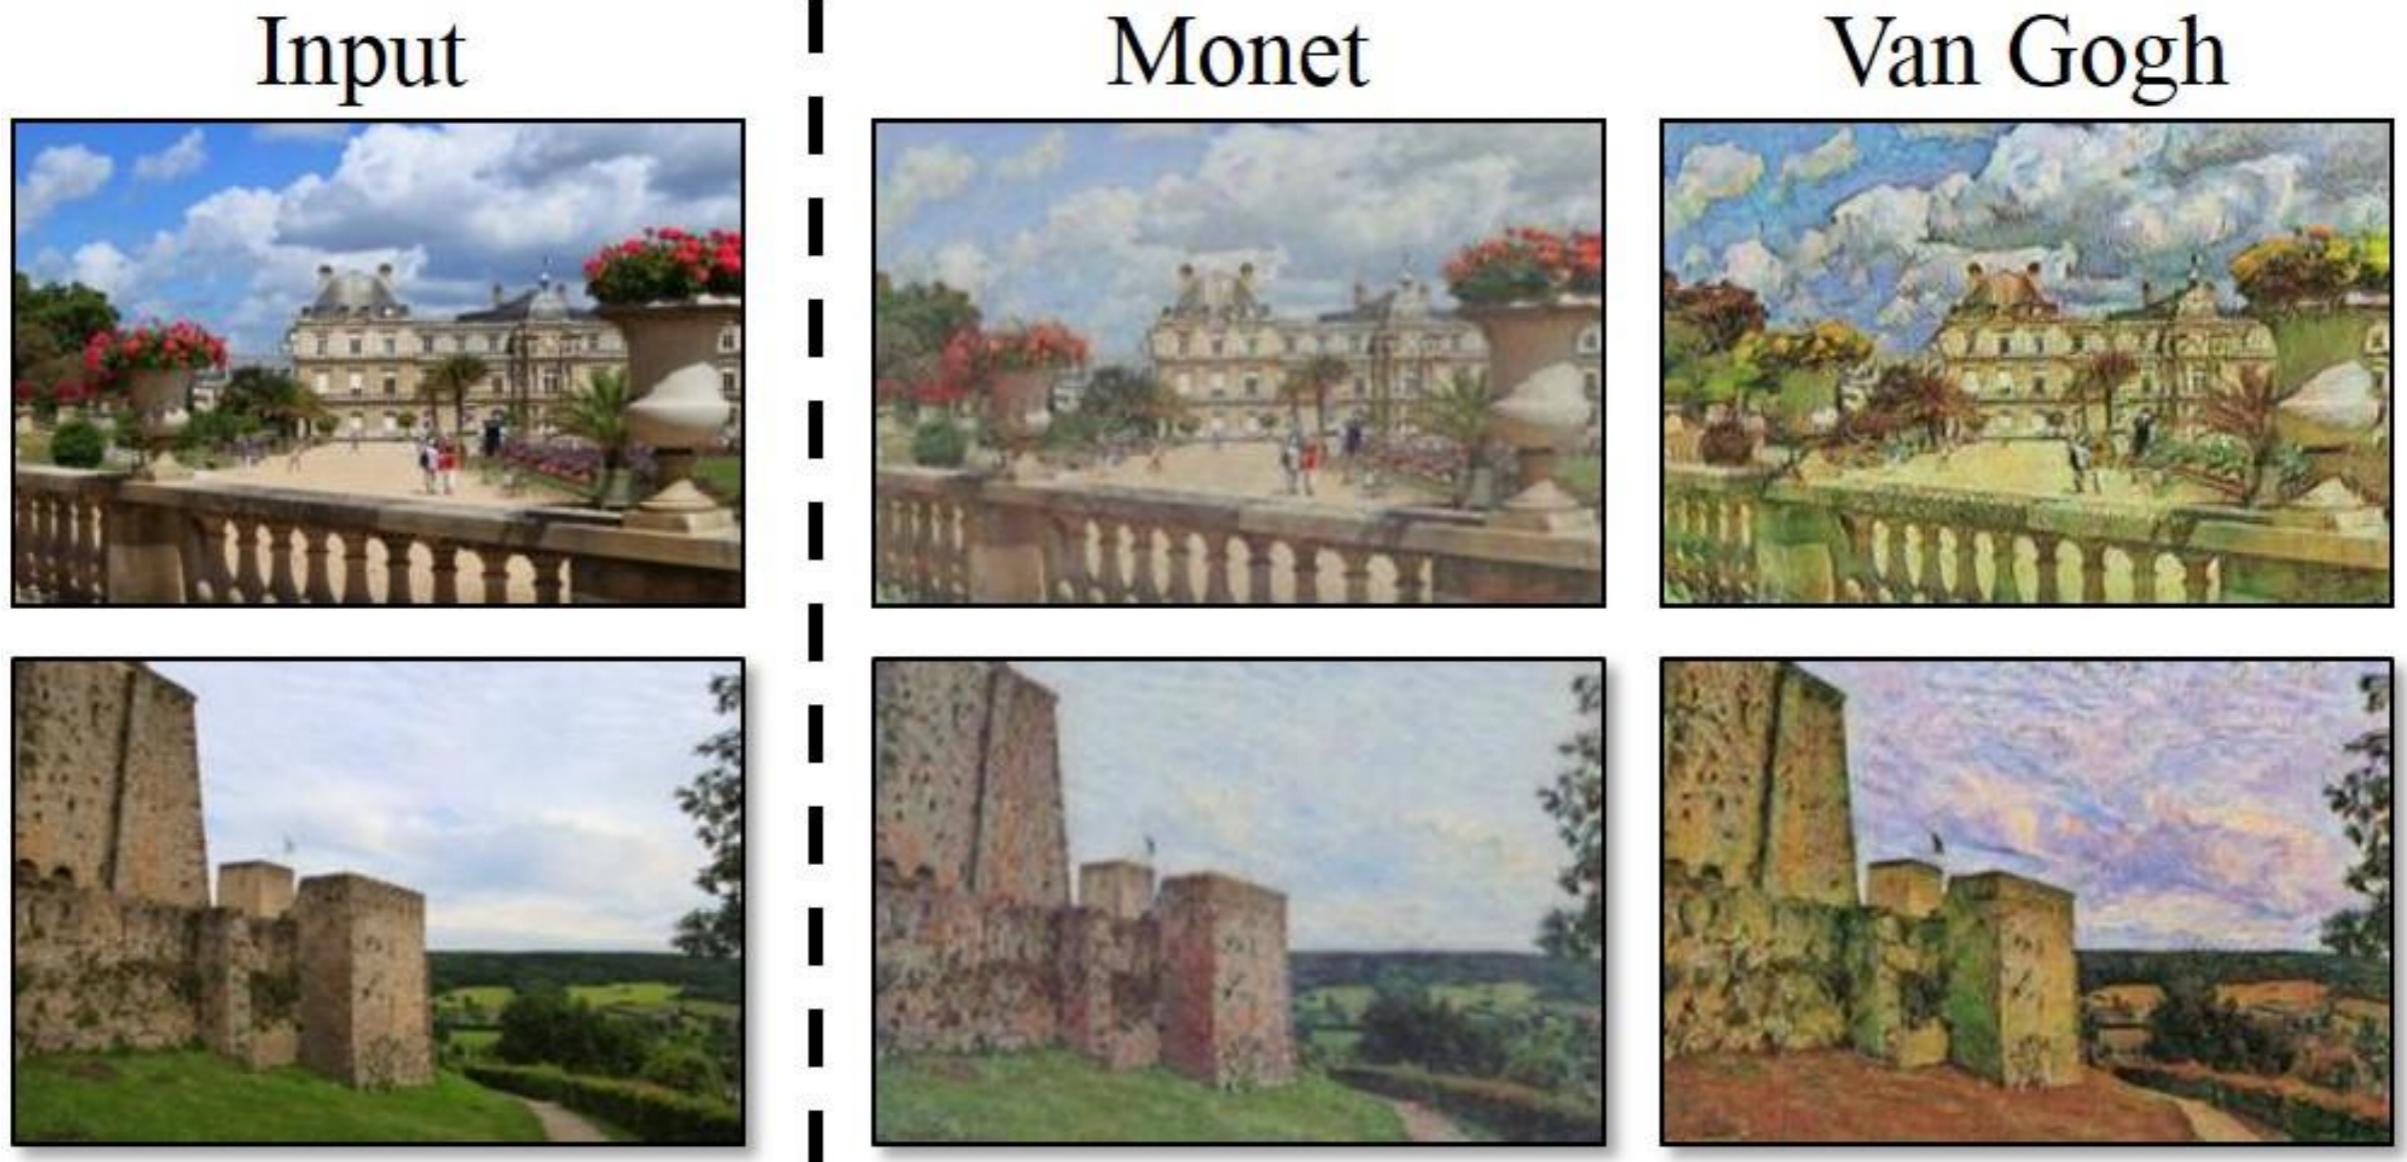
[Zhu et al., Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, ICCV 2017]

Image-to-Image Translation



[Zhu et al., Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, ICCV 2017]

Artistic Style Transfer



[Zhu et al., Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, ICCV 2017]

Training Data Generation

Input sunny image

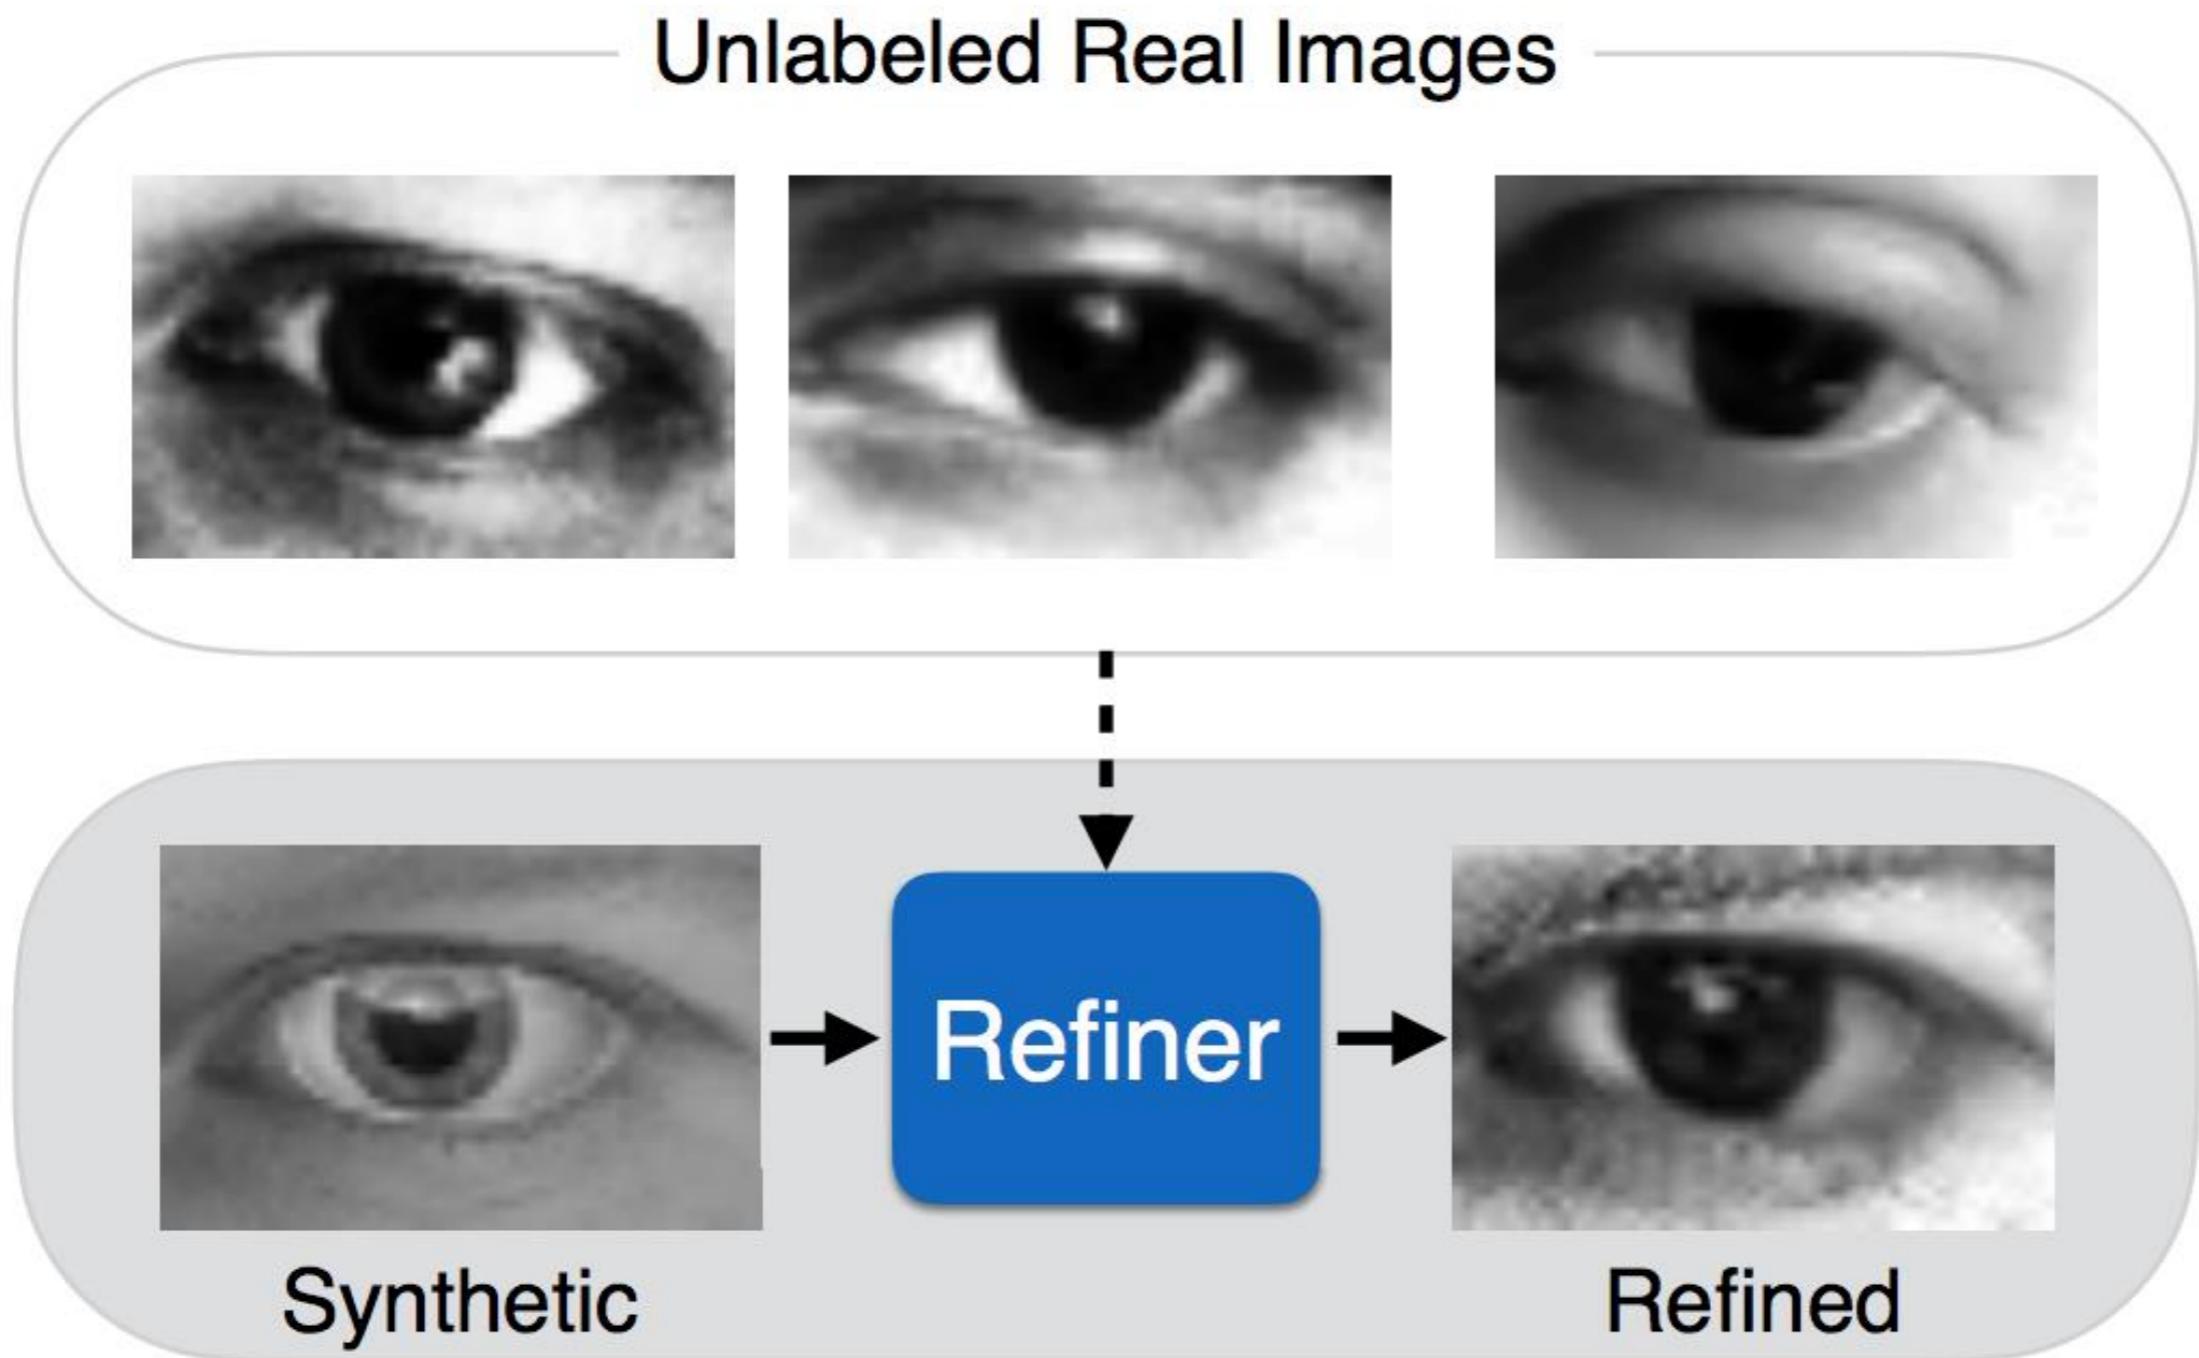


AI-generated rainy image



[Liu et al., Unsupervised Image-to-Image Translation Networks, NIPS 2017]

Training Data Generation



[Shrivastava et al., Learning from Simulated and Unsupervised Images through Adversarial Training, CVPR 2017]

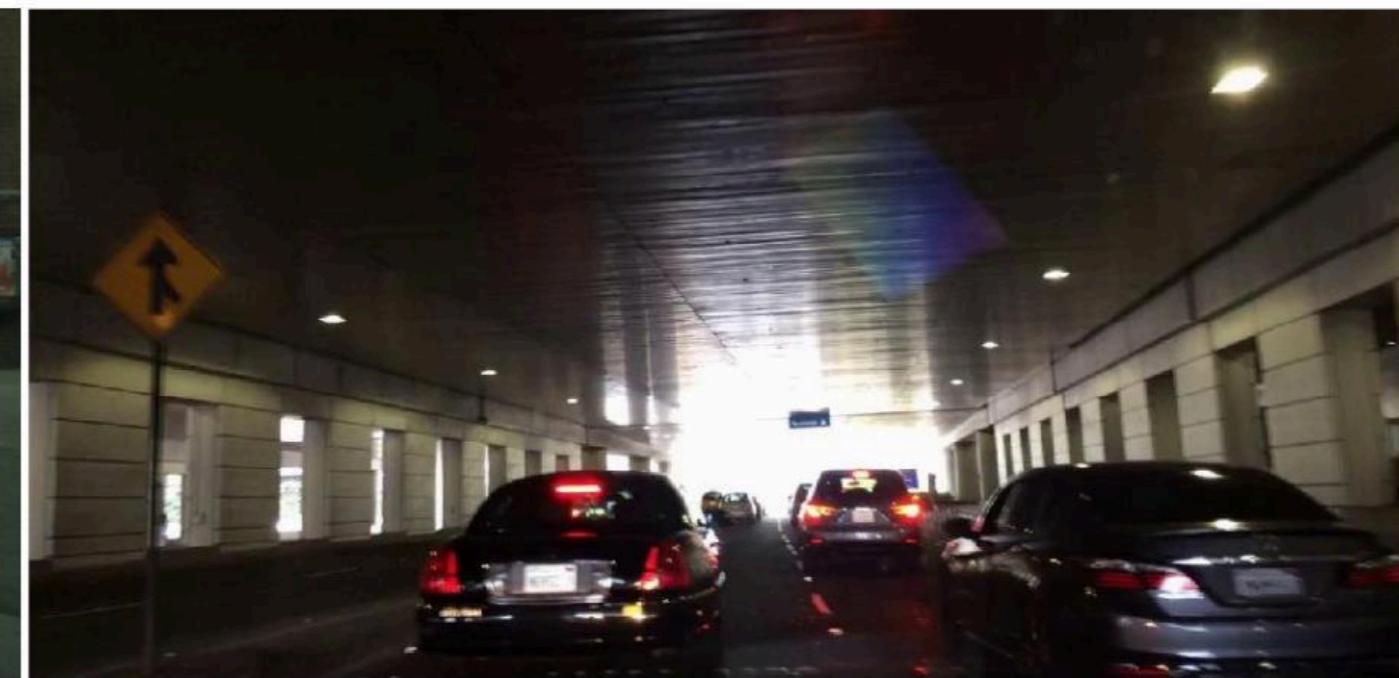
Domain Adaptation

Train



CityScapes (Germany)

Test



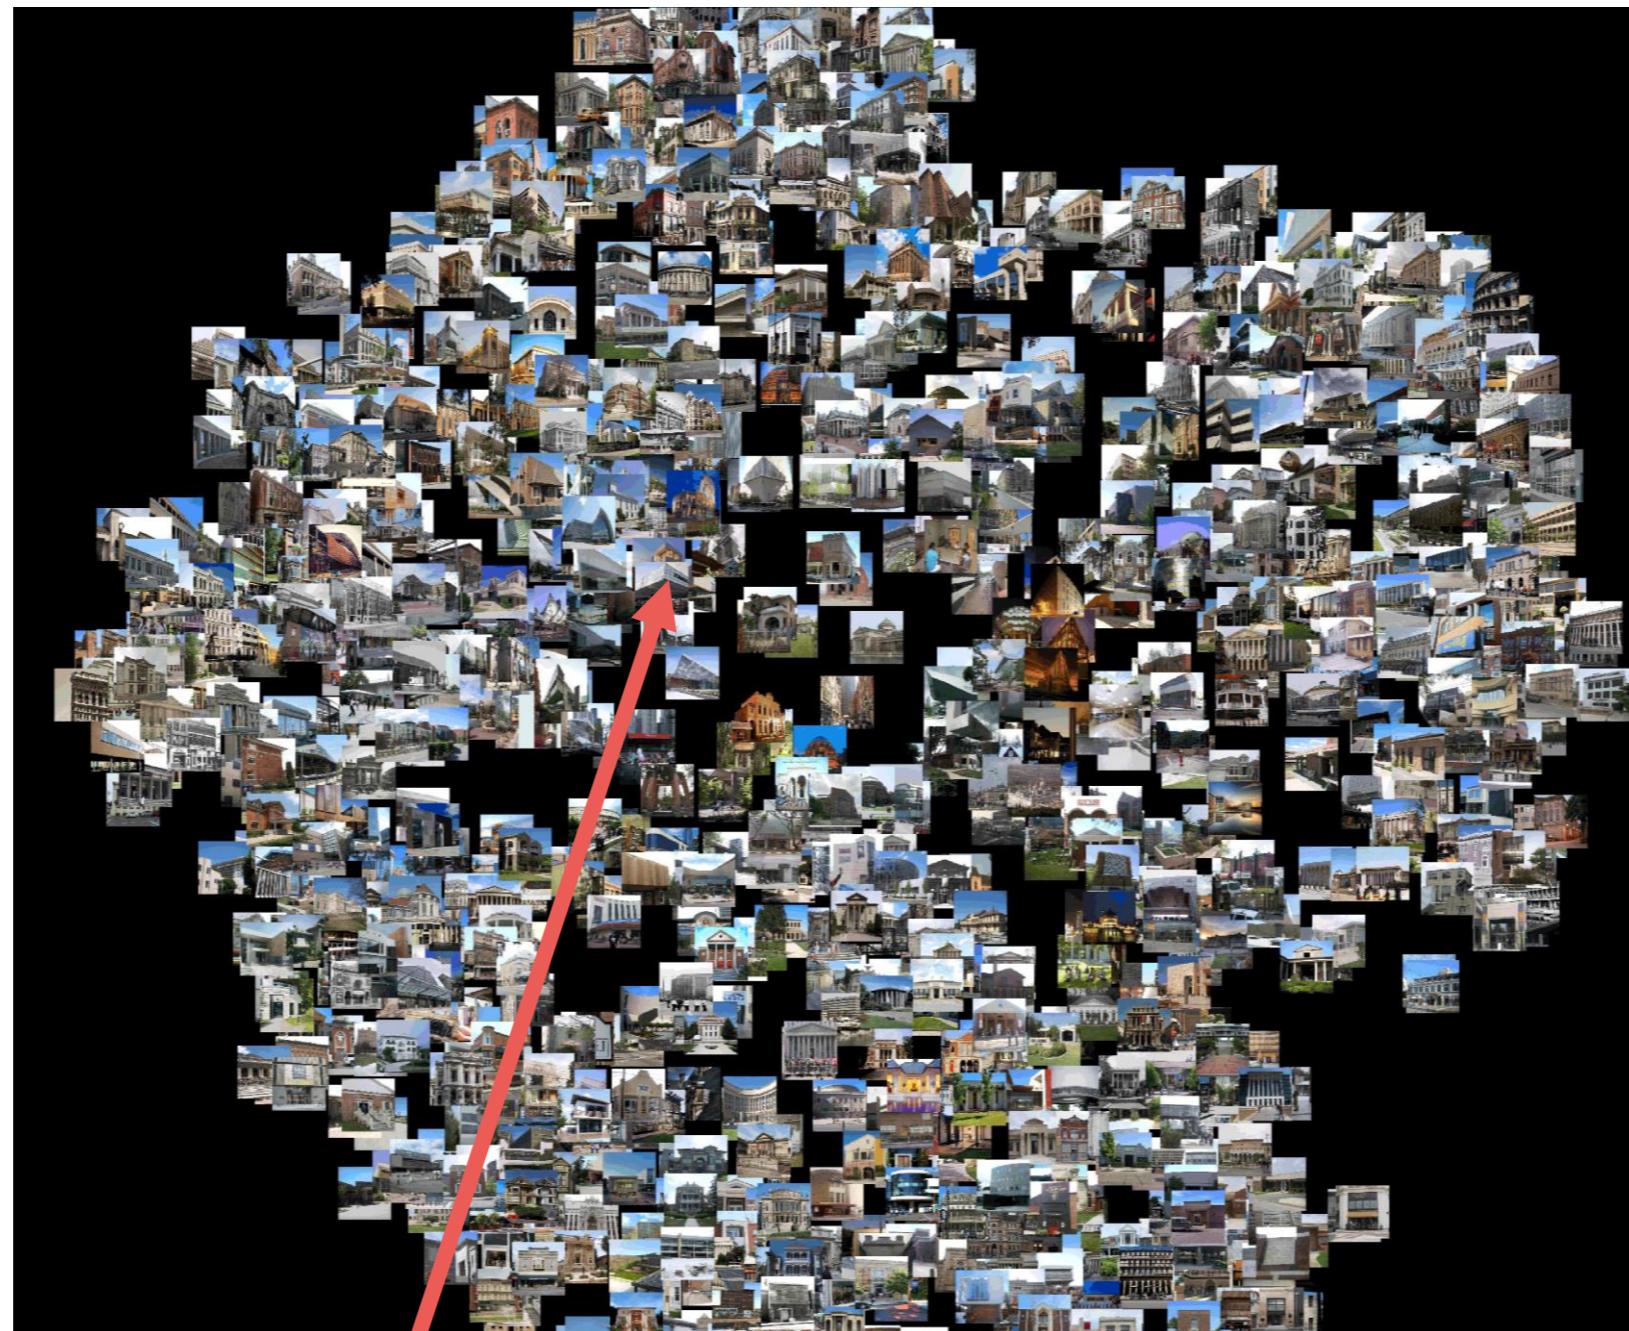
San Francisco

Source domain:
Plenty of labelled data

Target domain:
Little to no labelled data

slide credit: Judy Hoffman

Image Generation as Sampling Process



Generate image = pick image from distribution of real images

How to Model Distribution of Images?

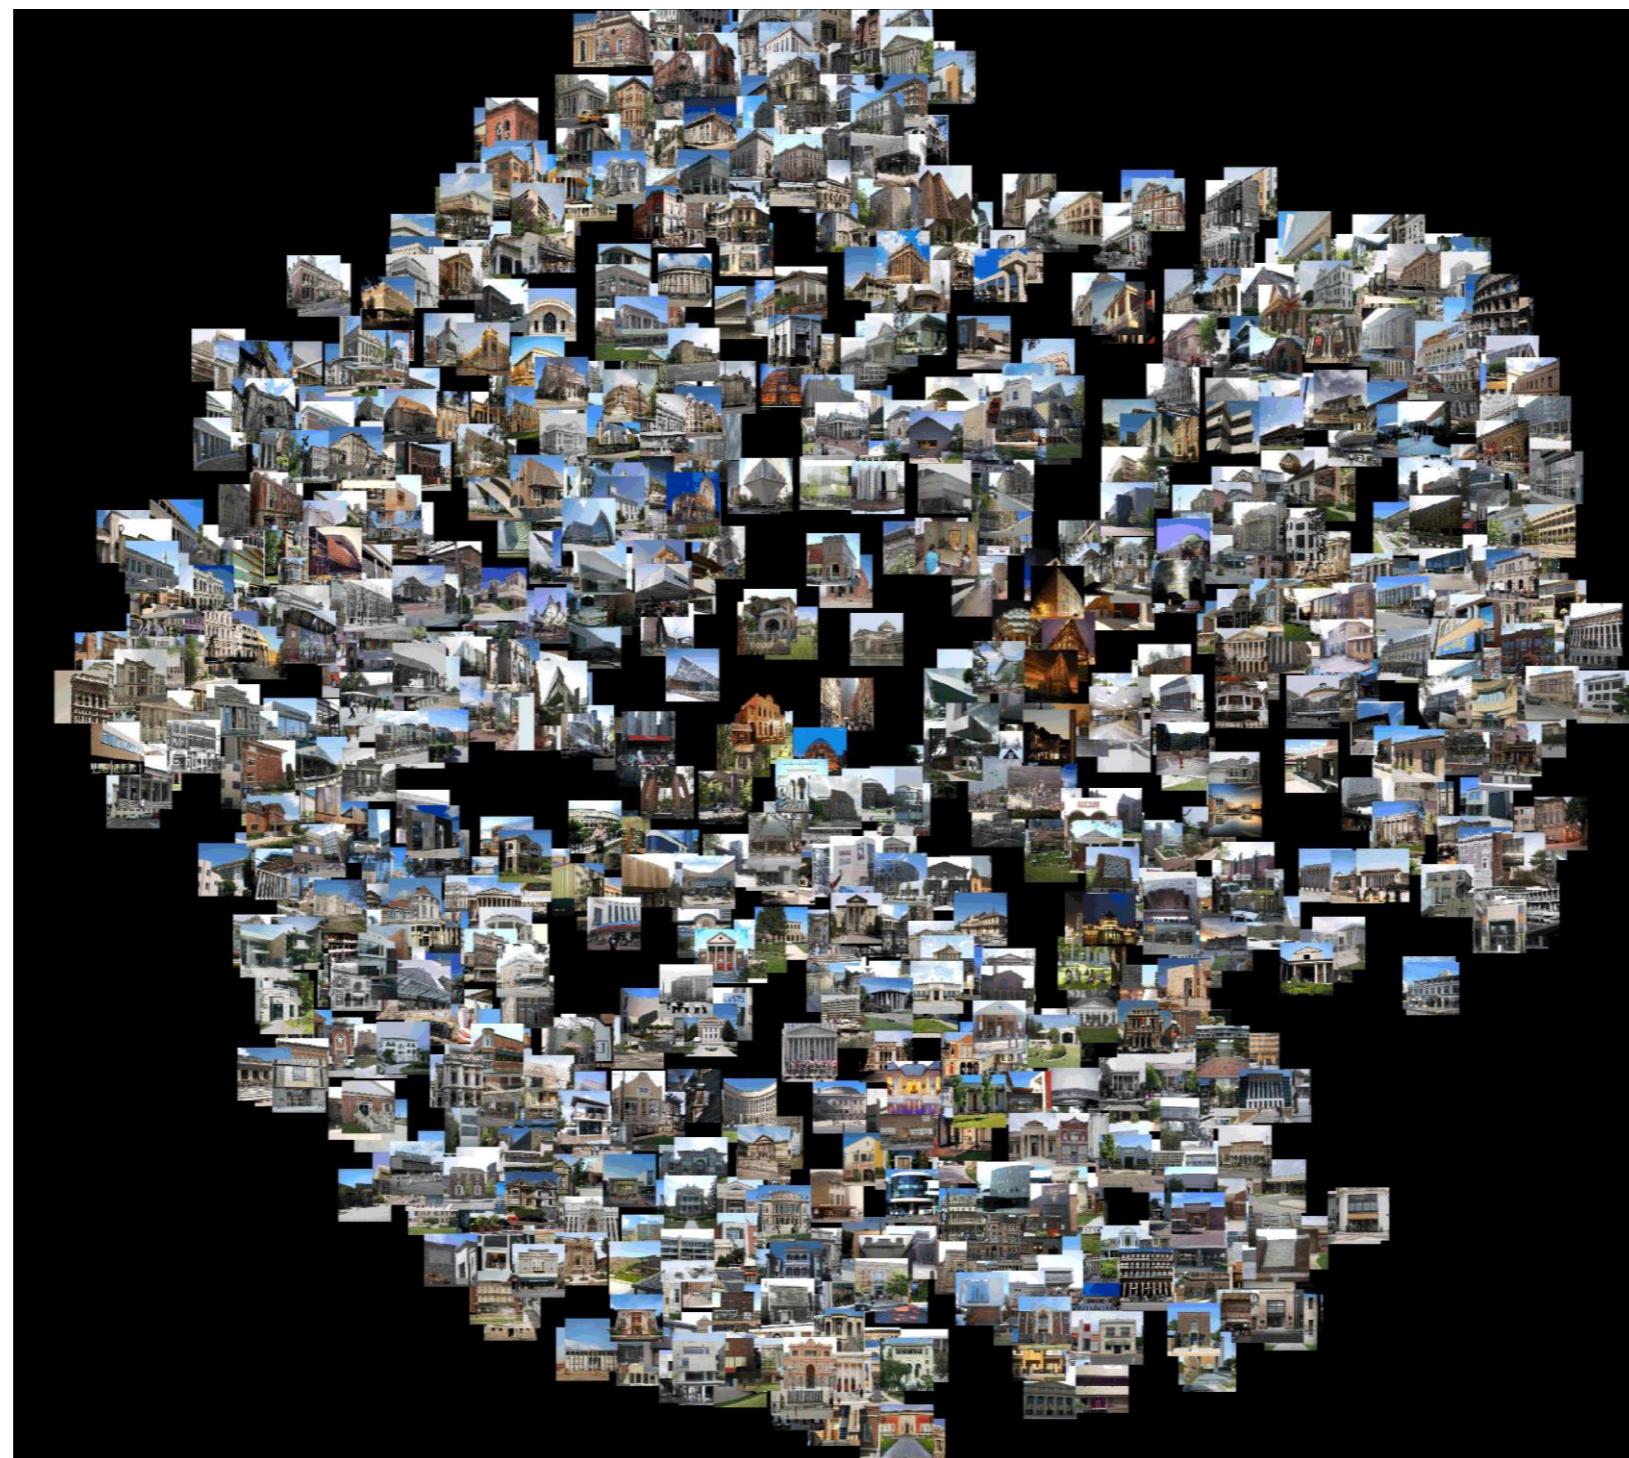
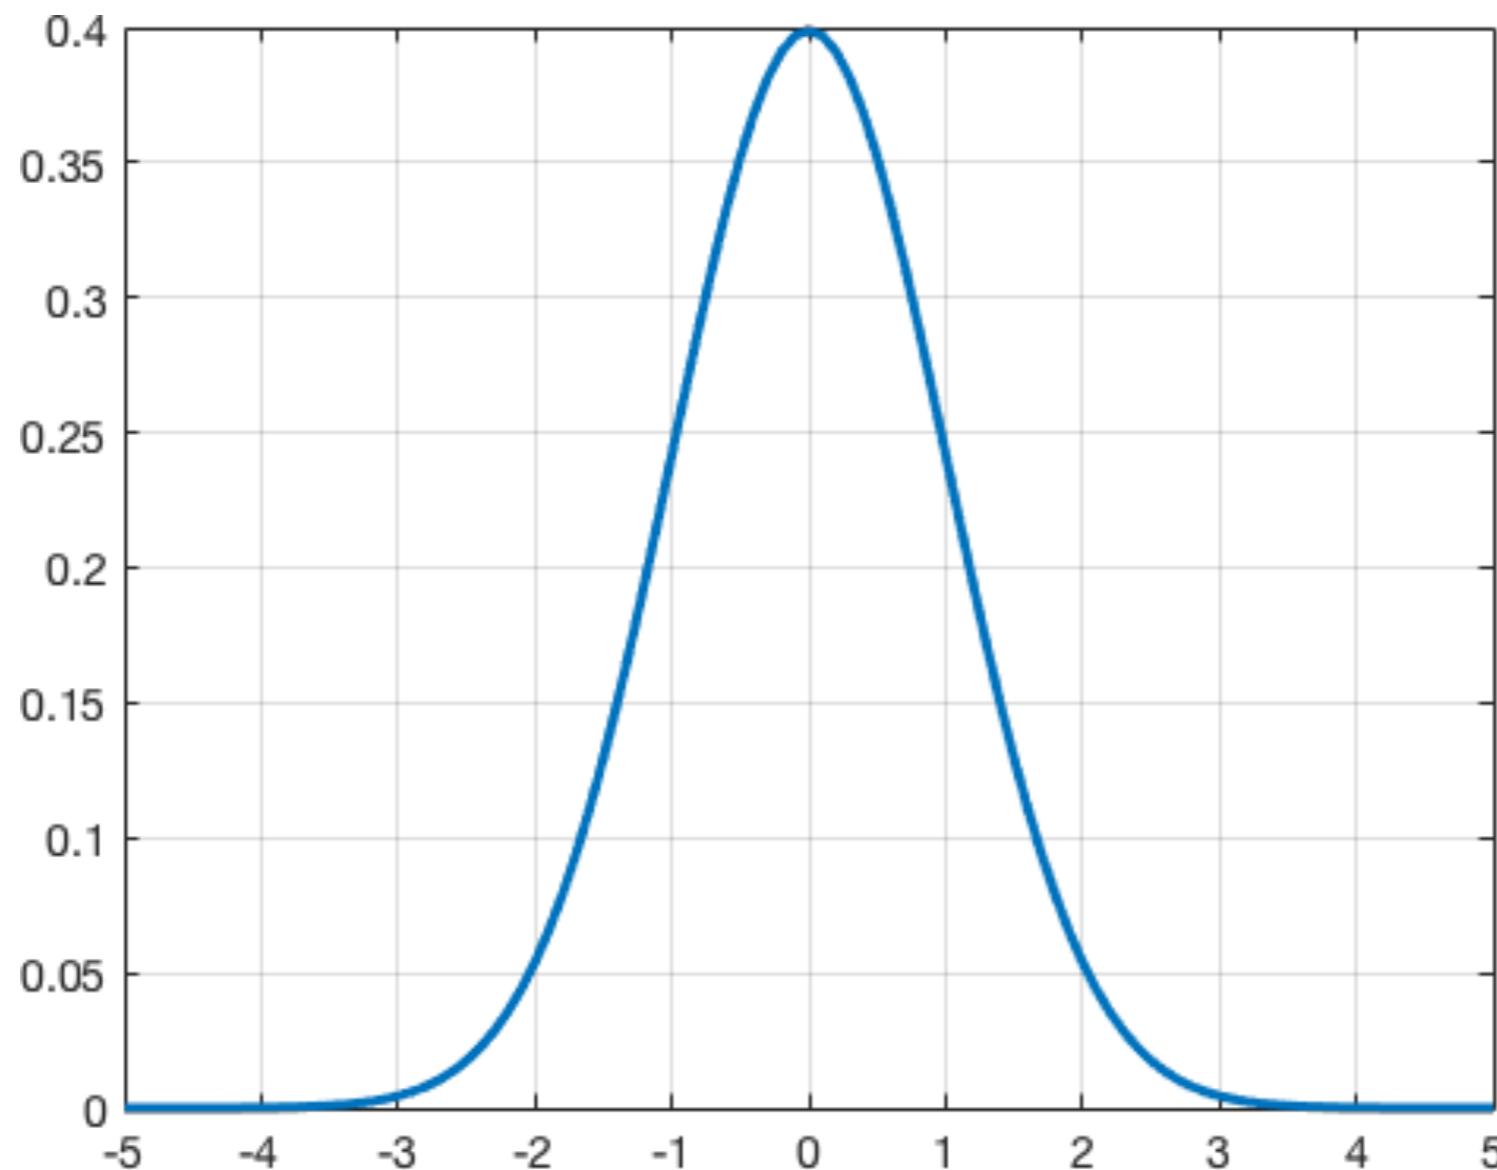


image from <http://3drepresentation.stanford.edu/>

Generation via Sampling

Example: Sampling from normal distribution



Generation via Sampling

Example: Sampling from normal distribution

Generation via Sampling

Example: Sampling from normal distribution

Box-Muller Transform:

Generation via Sampling

Example: Sampling from normal distribution

Box-Muller Transform:

- Input: U, V drawn uniformly from $(0, 1)$
- Output: Random numbers X, Y drawn from $\mathcal{N}(0, 1)$

Generation via Sampling

Example: Sampling from normal distribution

Box-Muller Transform:

- Input: U, V drawn uniformly from $(0, 1)$
- Output: Random numbers X, Y drawn from $\mathcal{N}(0, 1)$
- Algorithm:

$$X = \sqrt{-2 \ln U} \cos(2\pi V)$$

$$Y = \sqrt{-2 \ln U} \sin(2\pi V)$$

Generation via Sampling

Example: Sampling from normal distribution

Box-Muller Transform:

- Input: U, V drawn uniformly from $(0, 1)$
- Output: Random numbers X, Y drawn from $\mathcal{N}(0, 1)$
- Algorithm:

$$X = \sqrt{-2 \ln U} \cos(2\pi V)$$

$$Y = \sqrt{-2 \ln U} \sin(2\pi V)$$

Key idea: Sample from simple distribution, map to complex one

Generative Neural Networks

code



$$\mathbf{z} \sim p_{\text{model}}$$

Generative Neural Networks

code

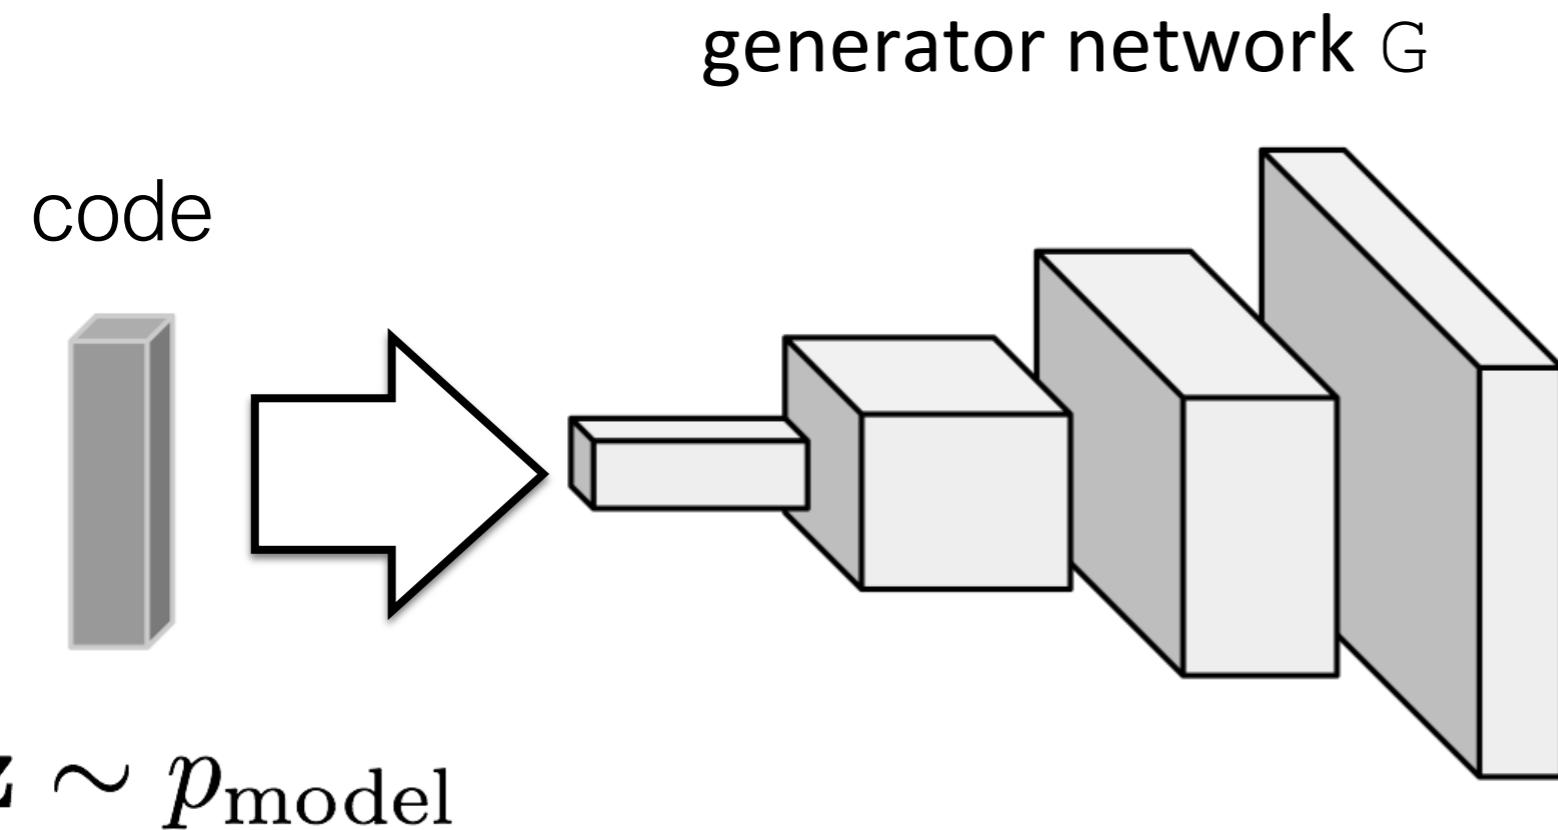


$$\mathbf{z} \sim p_{\text{model}}$$

normal distribution,
uniform distribution,

...

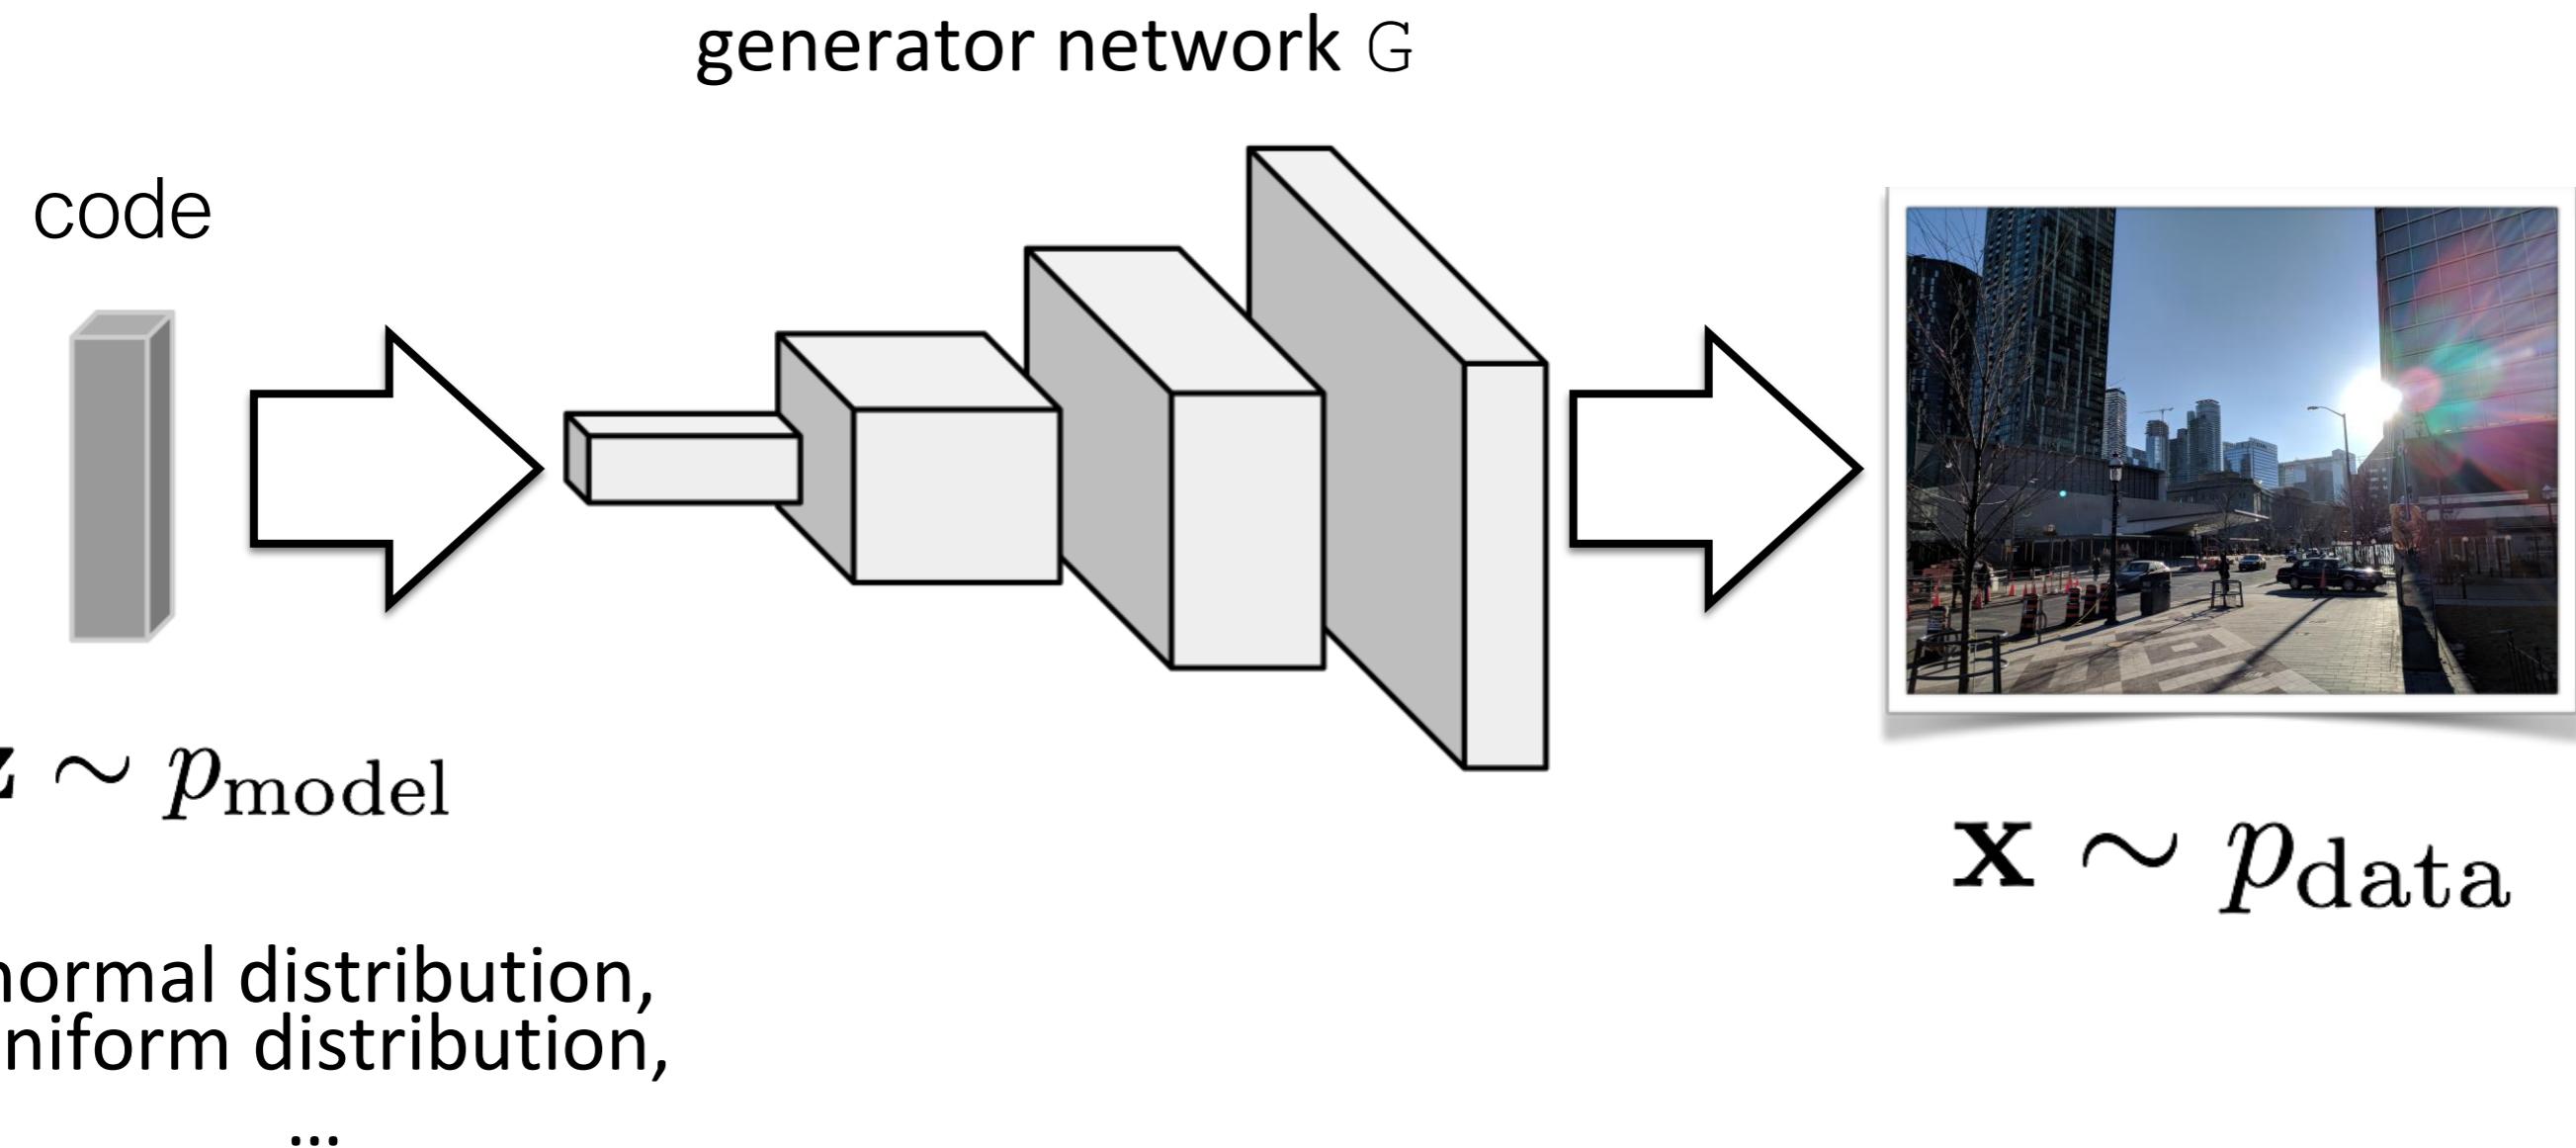
Generative Neural Networks



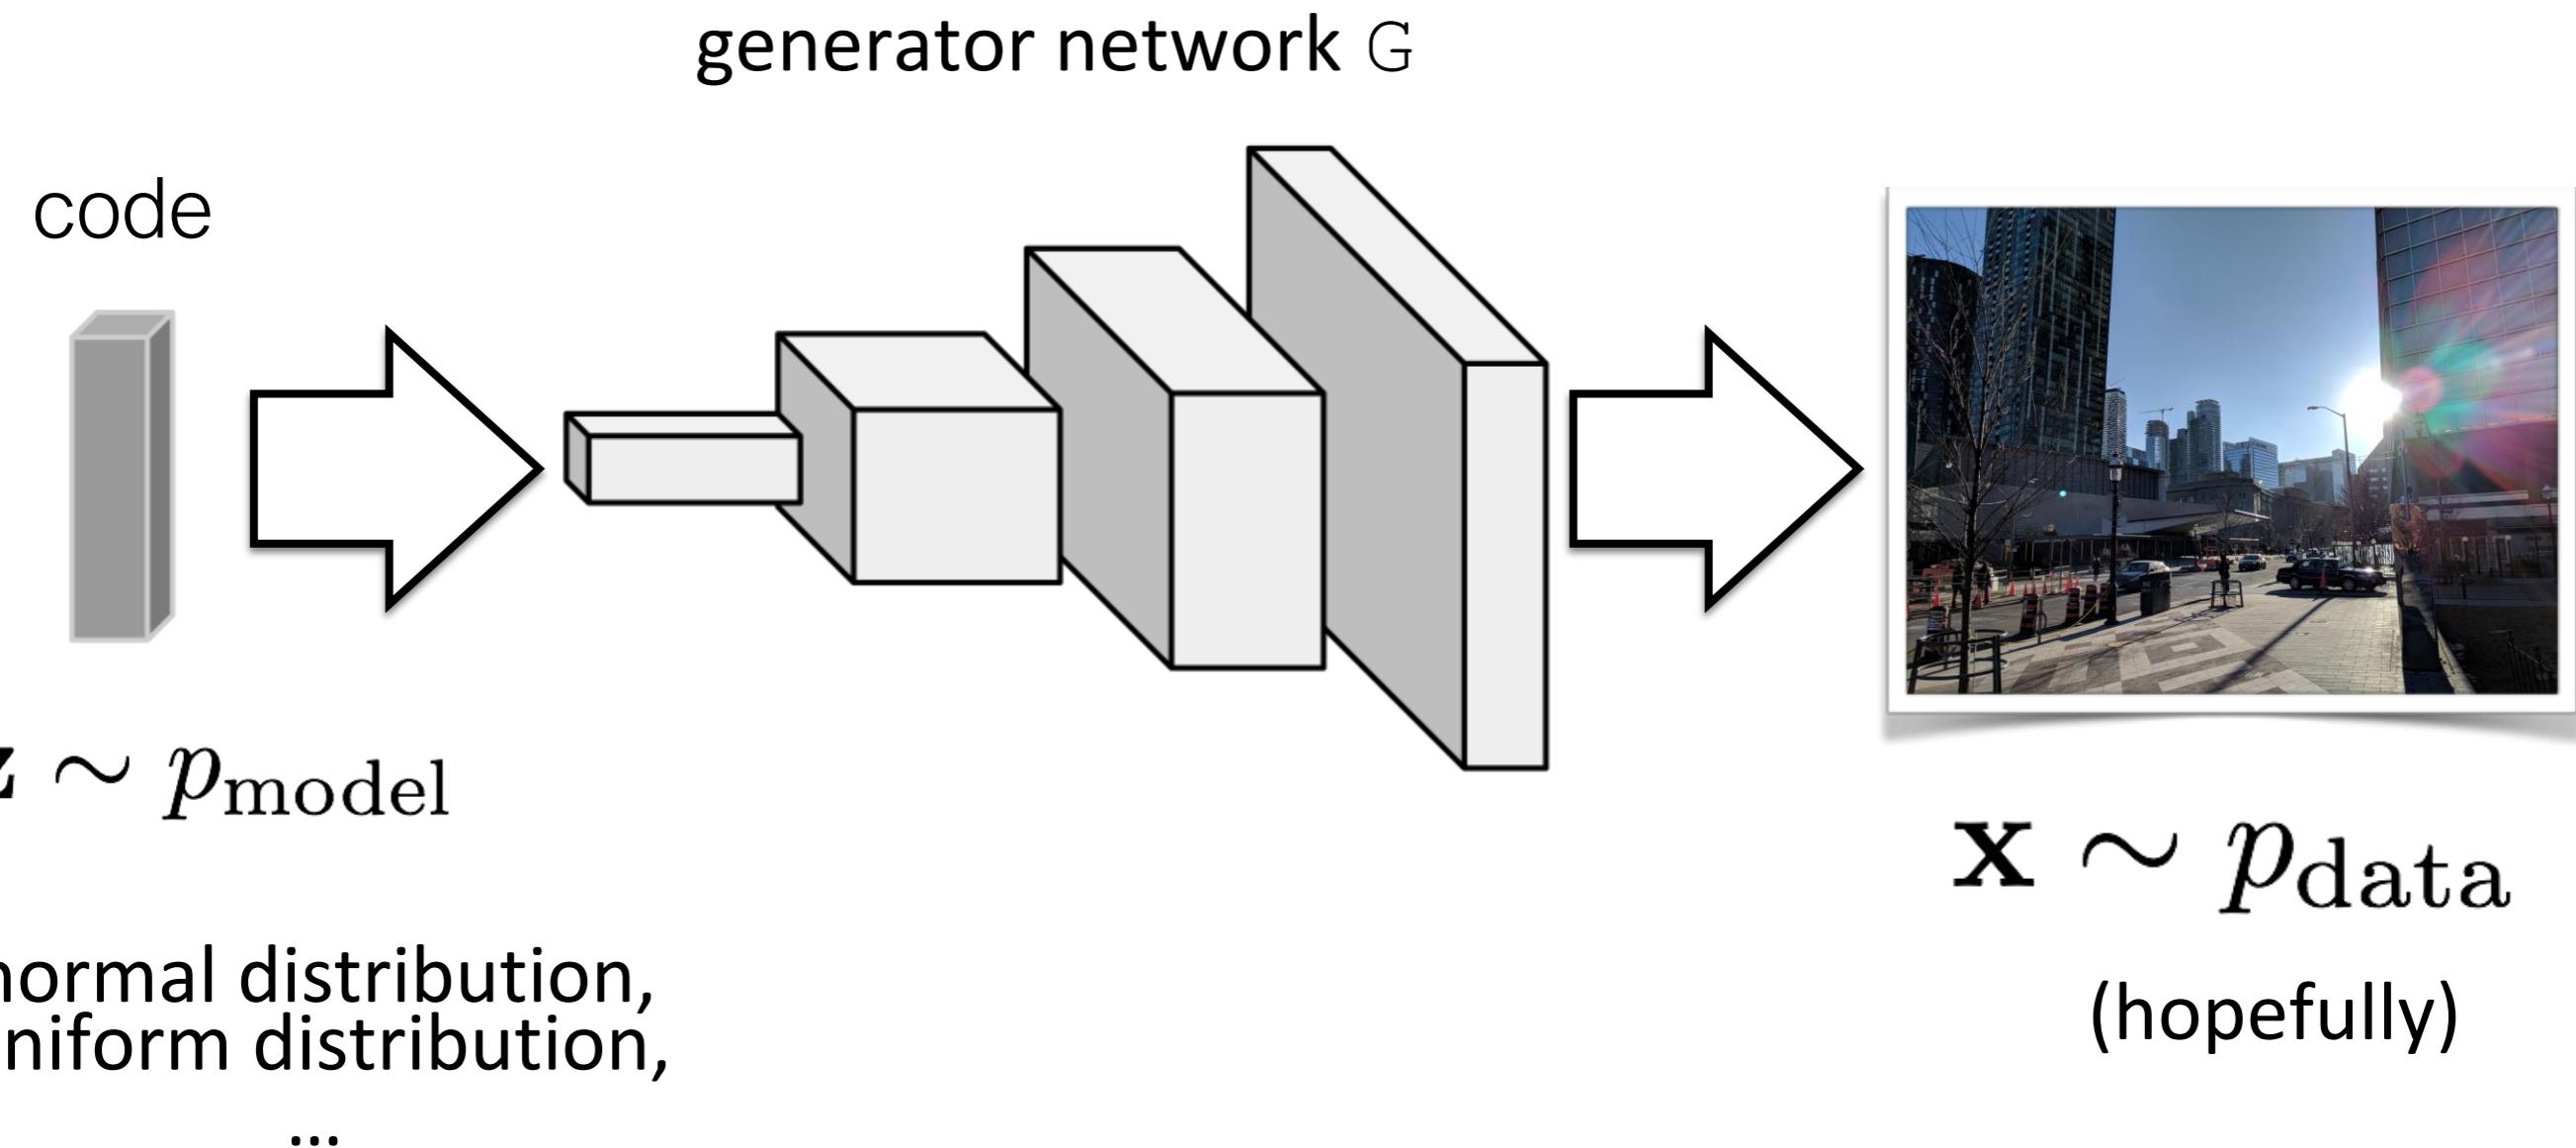
normal distribution,
uniform distribution,

...

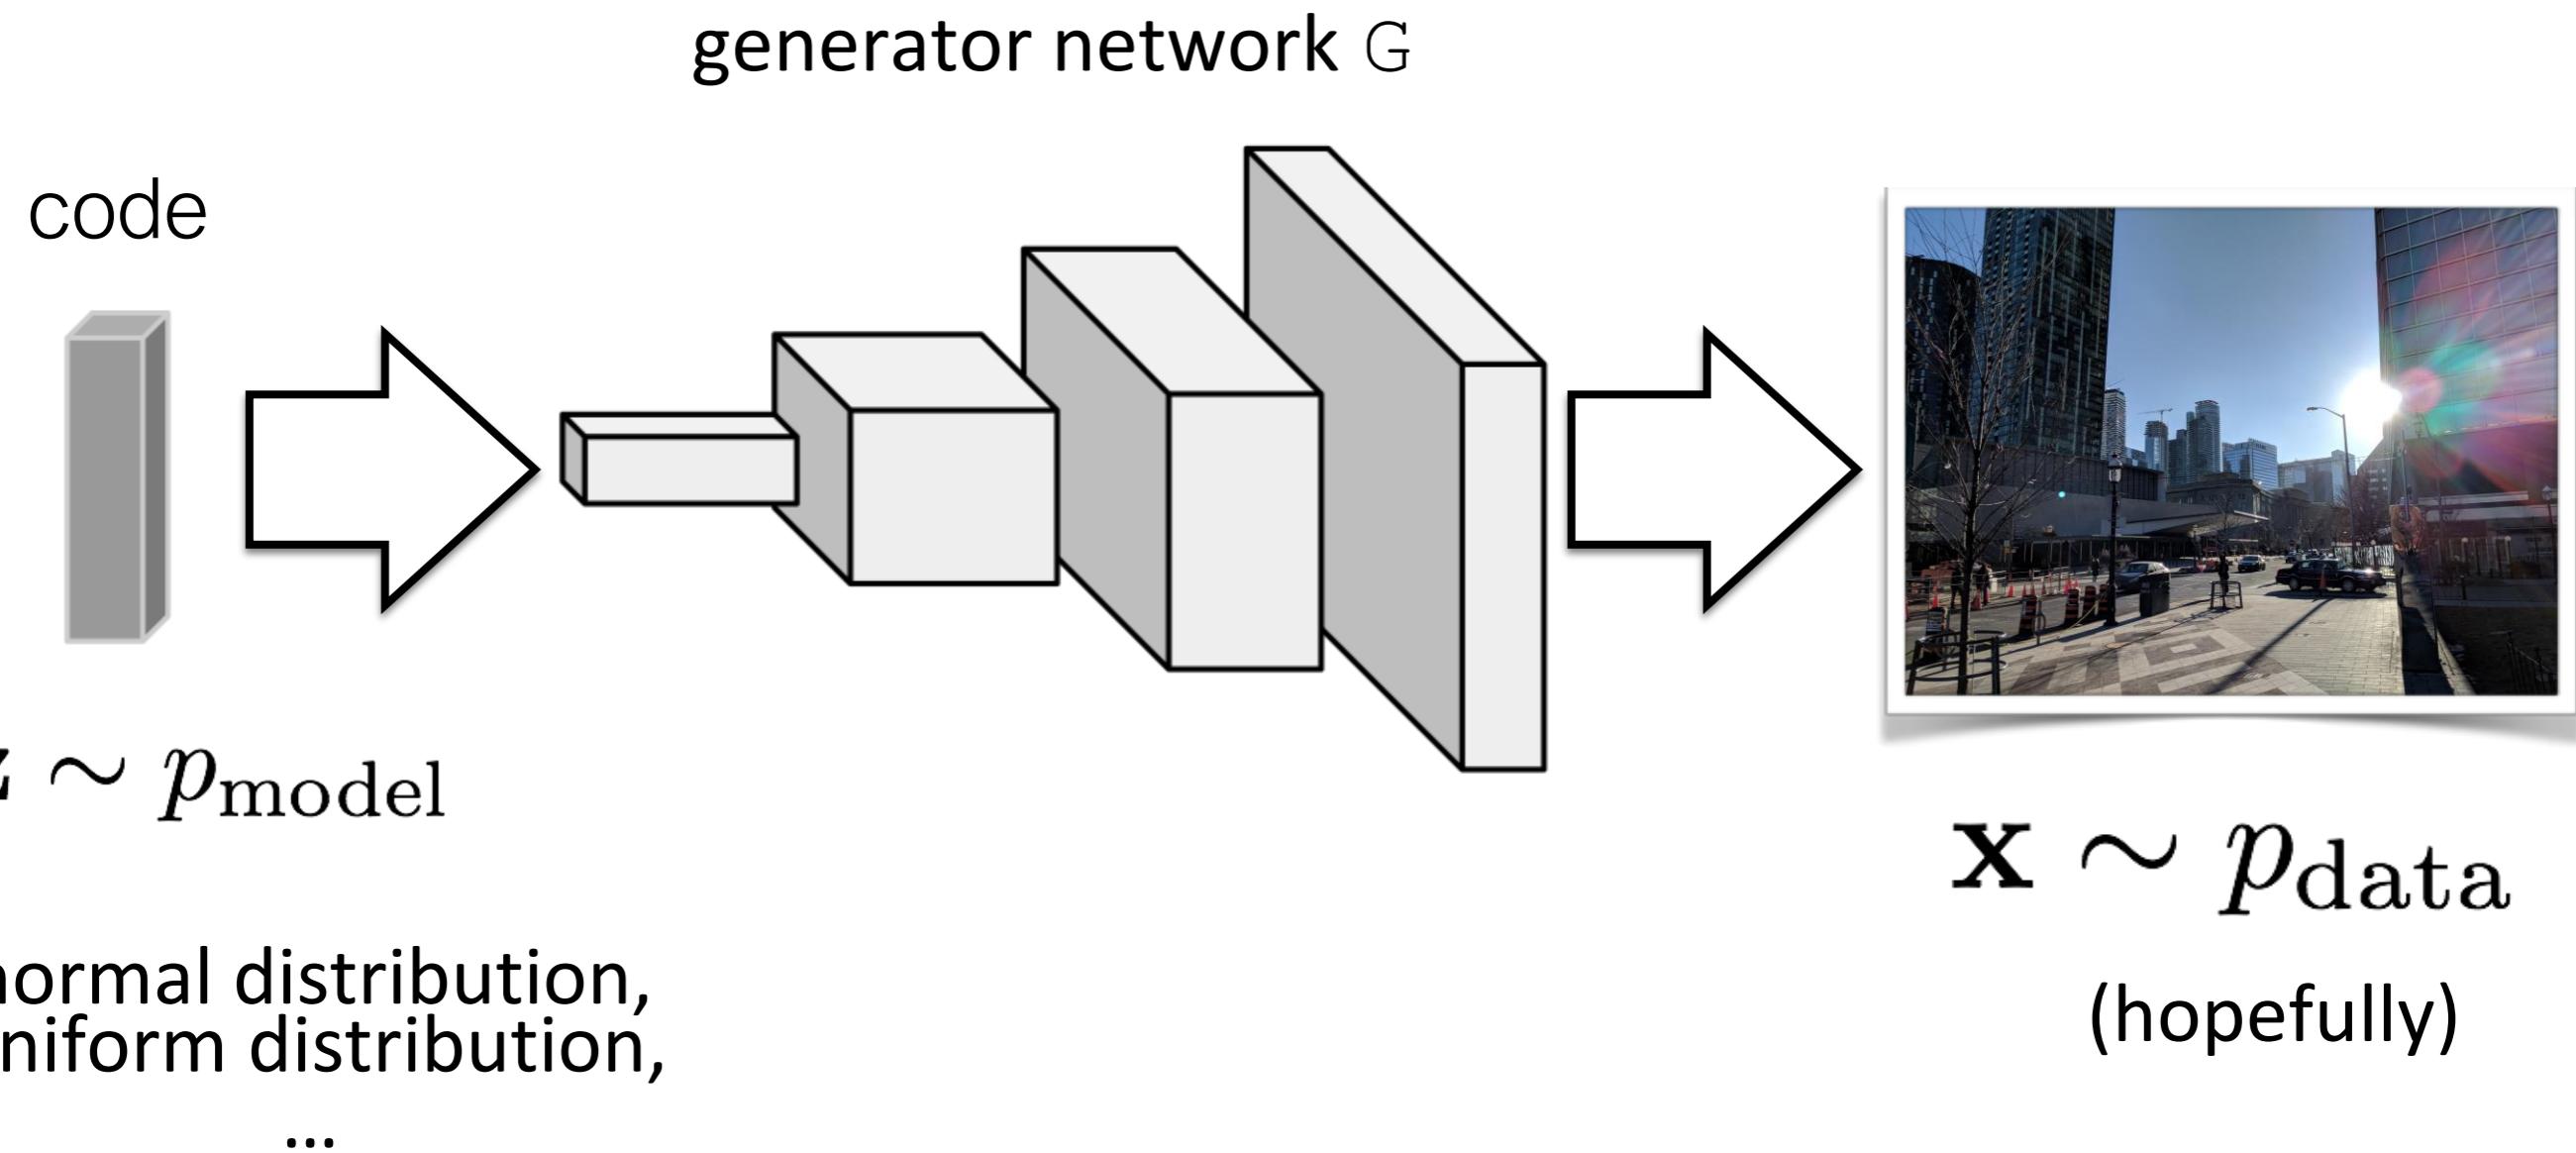
Generative Neural Networks



Generative Neural Networks



Generative Neural Networks



How to make sure that this mapping produces realistic images?

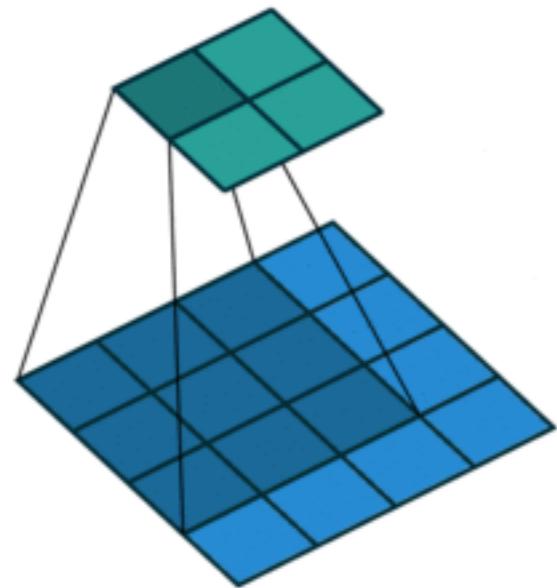
Today

- Variational Autoencoders (VAE)
- Generative Adversarial Networks (GANs)

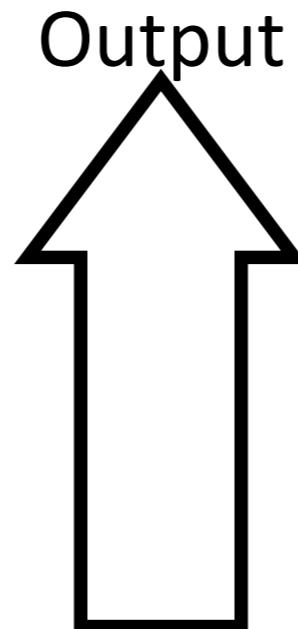
Today

- **Variational Autoencoders (VAE)**
- Generative Adversarial Networks (GANs)

Recap: (Transposed) Convolutions



convolution



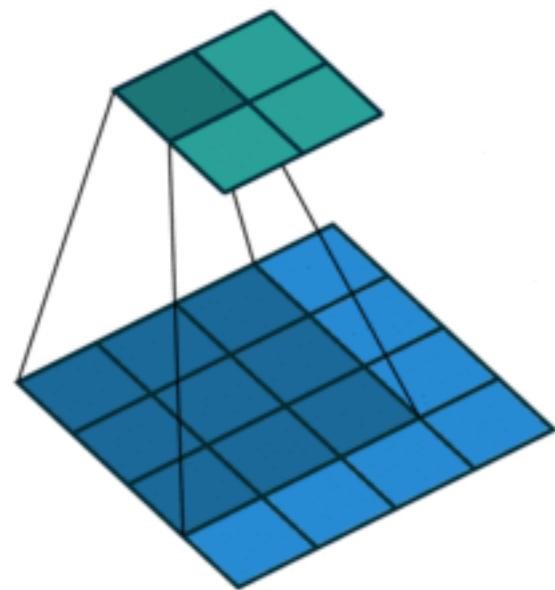
Input

Output

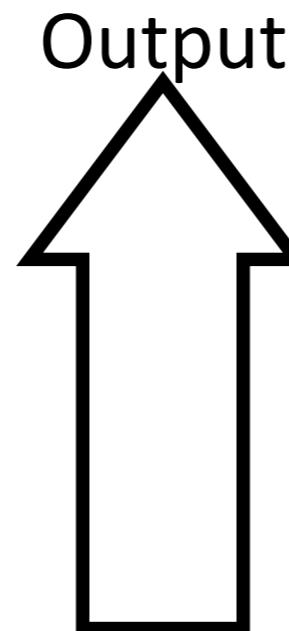
illustrations from https://github.com/vdumoulin/conv_arithmetic

slide credit: Laura Leal-Taixé

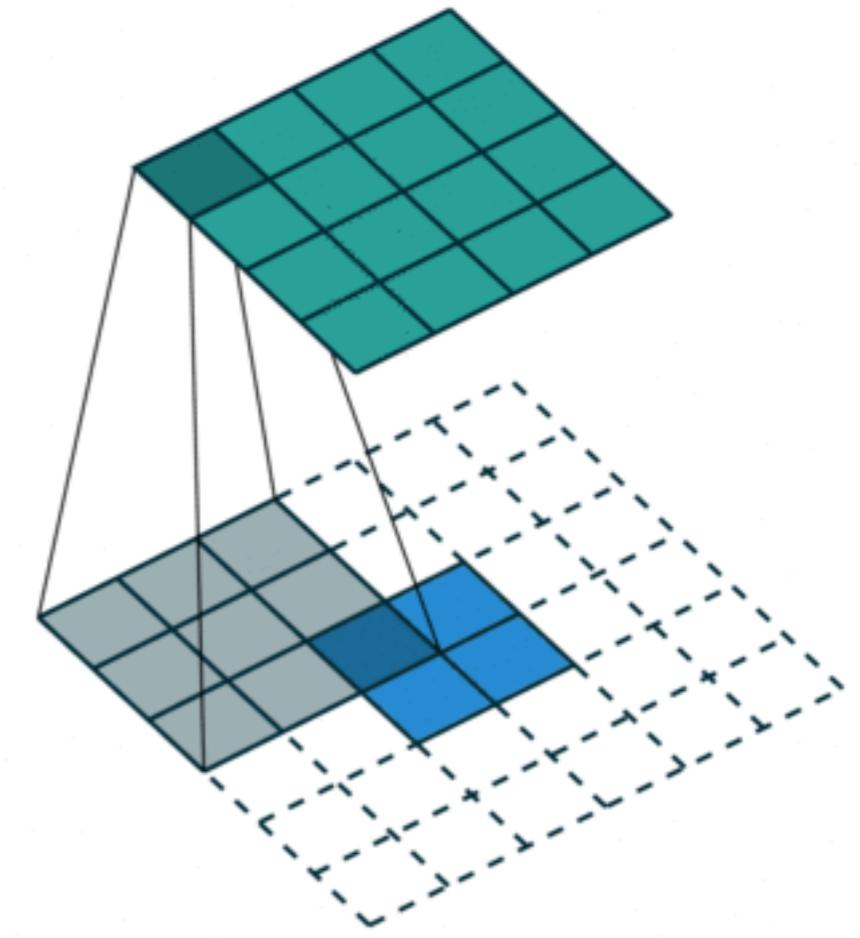
Recap: (Transposed) Convolutions



convolution



Input

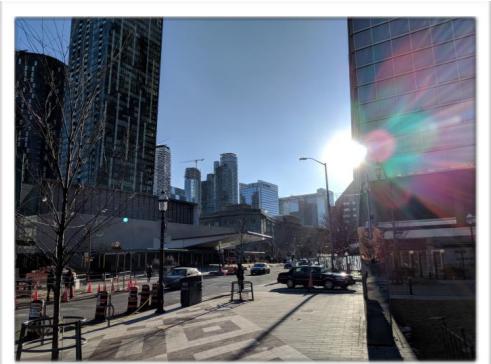


transposed convolution

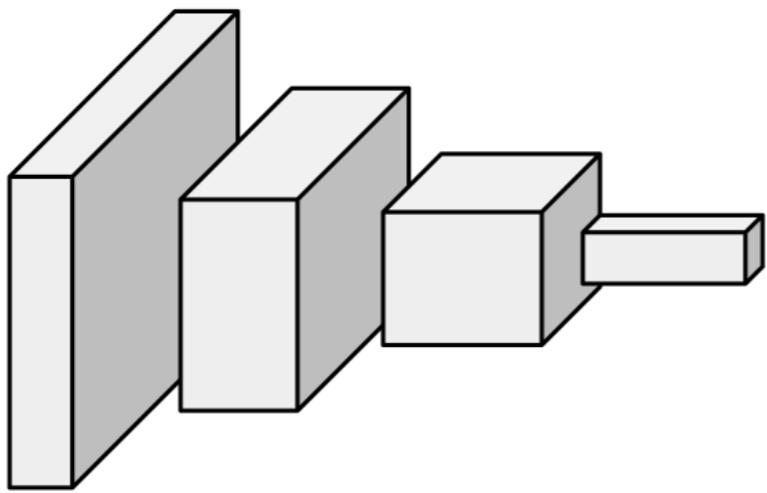
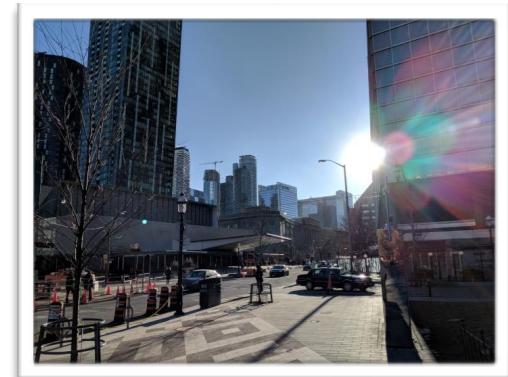
illustrations from https://github.com/vdumoulin/conv_arithmetic

slide credit: Laura Leal-Taixé

Autoencoder (AE)

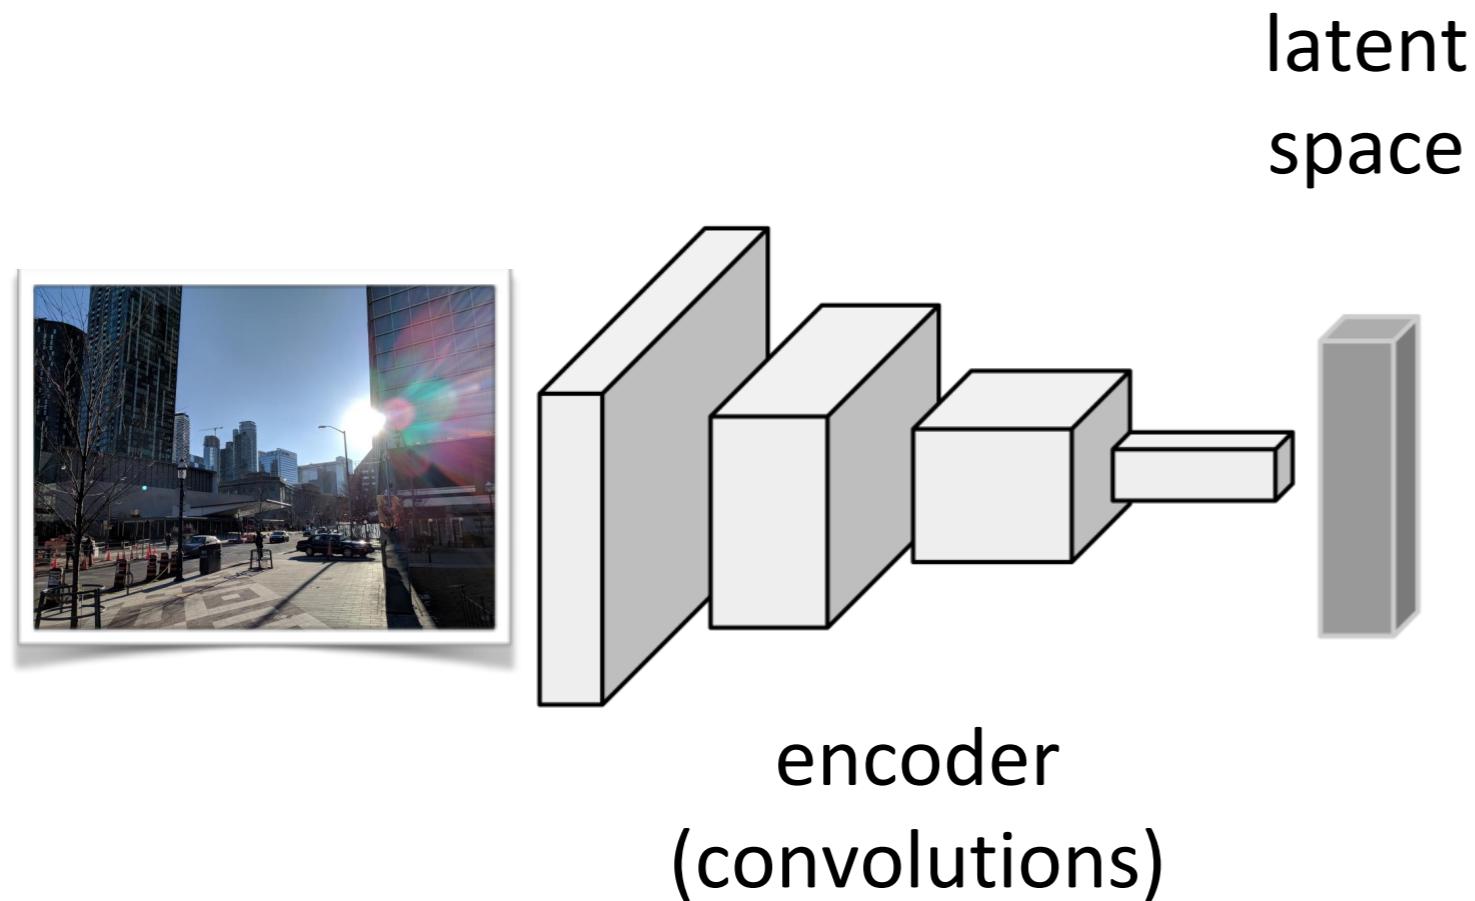


Autoencoder (AE)

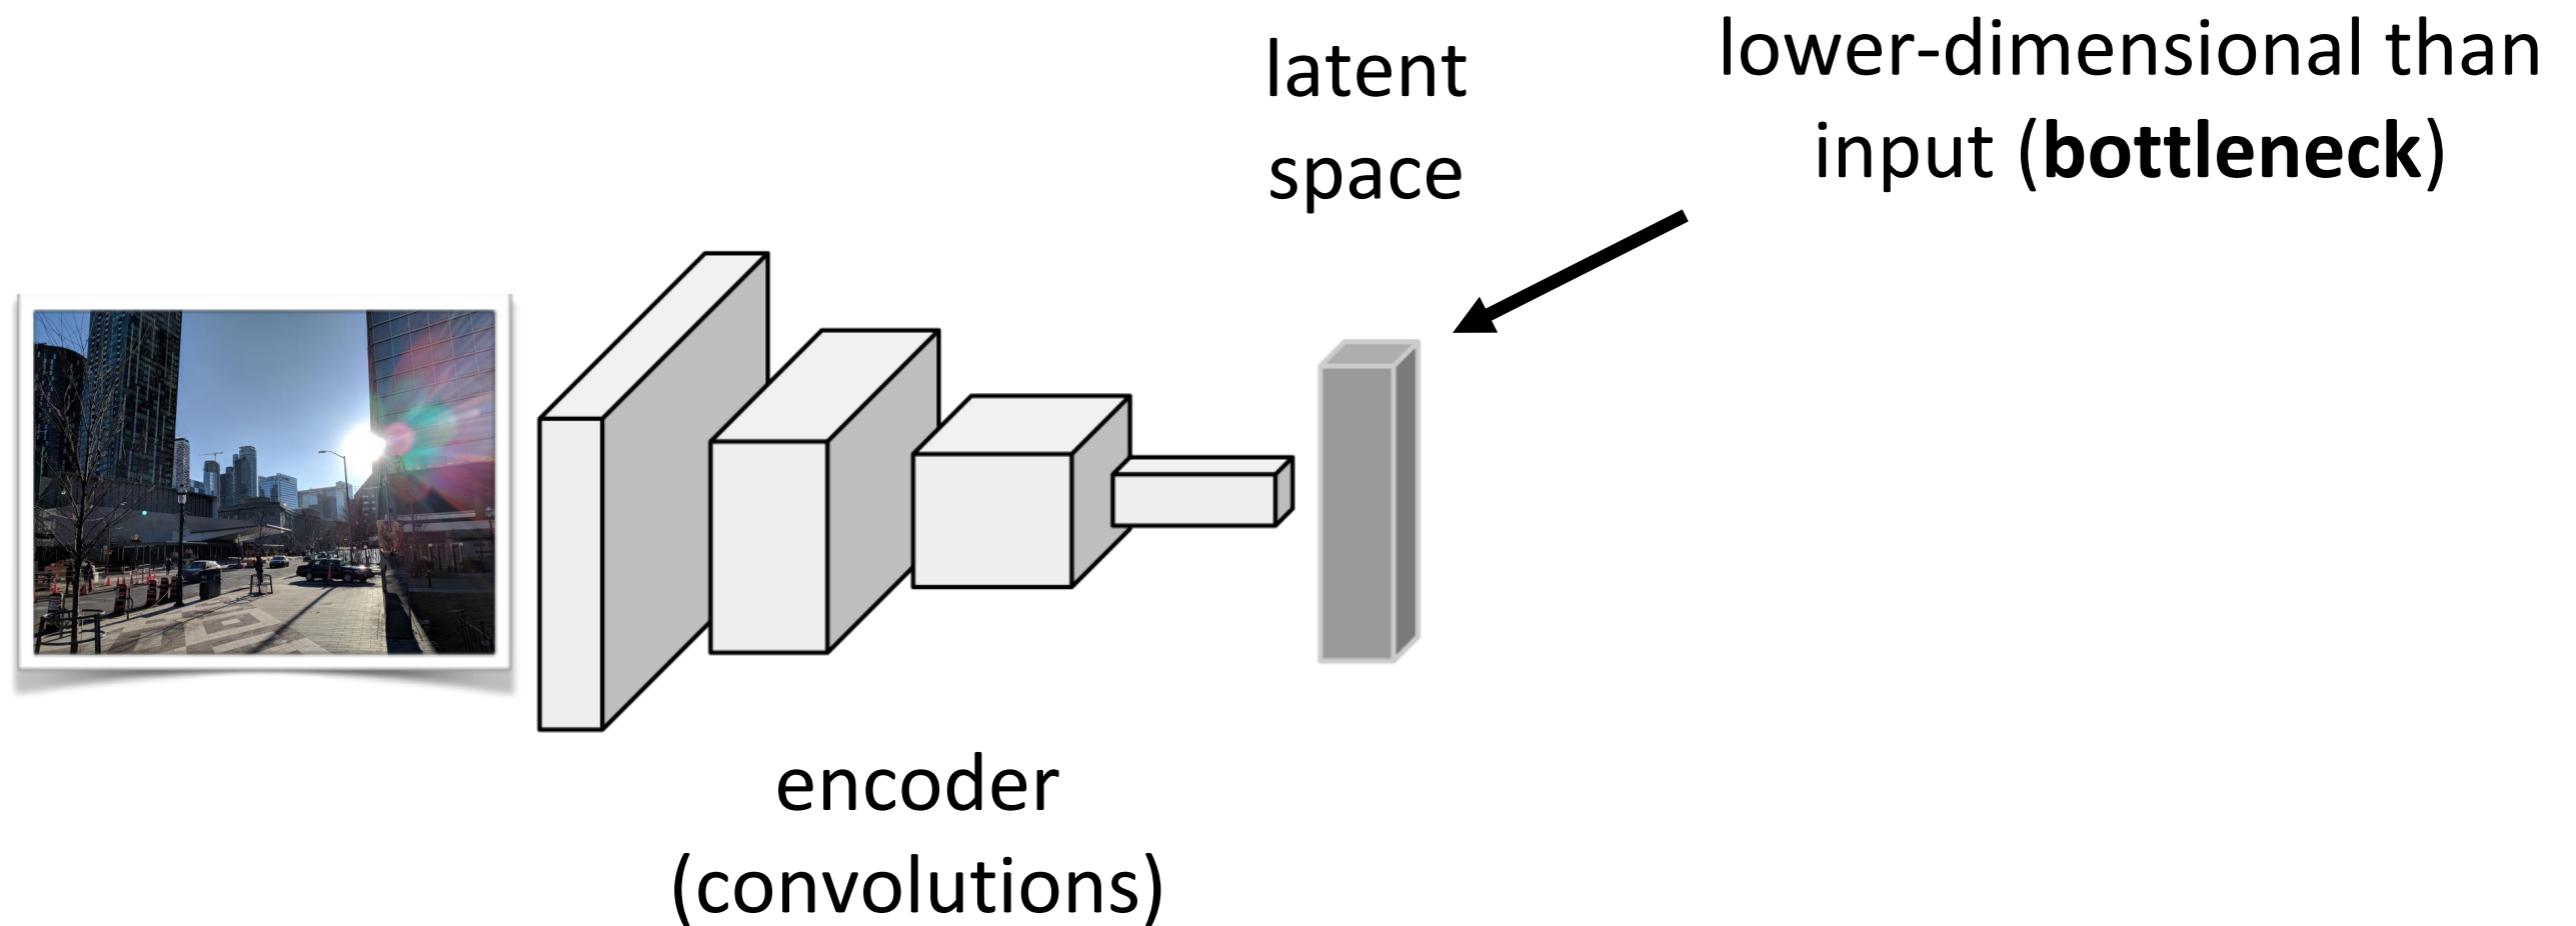


encoder
(convolutions)

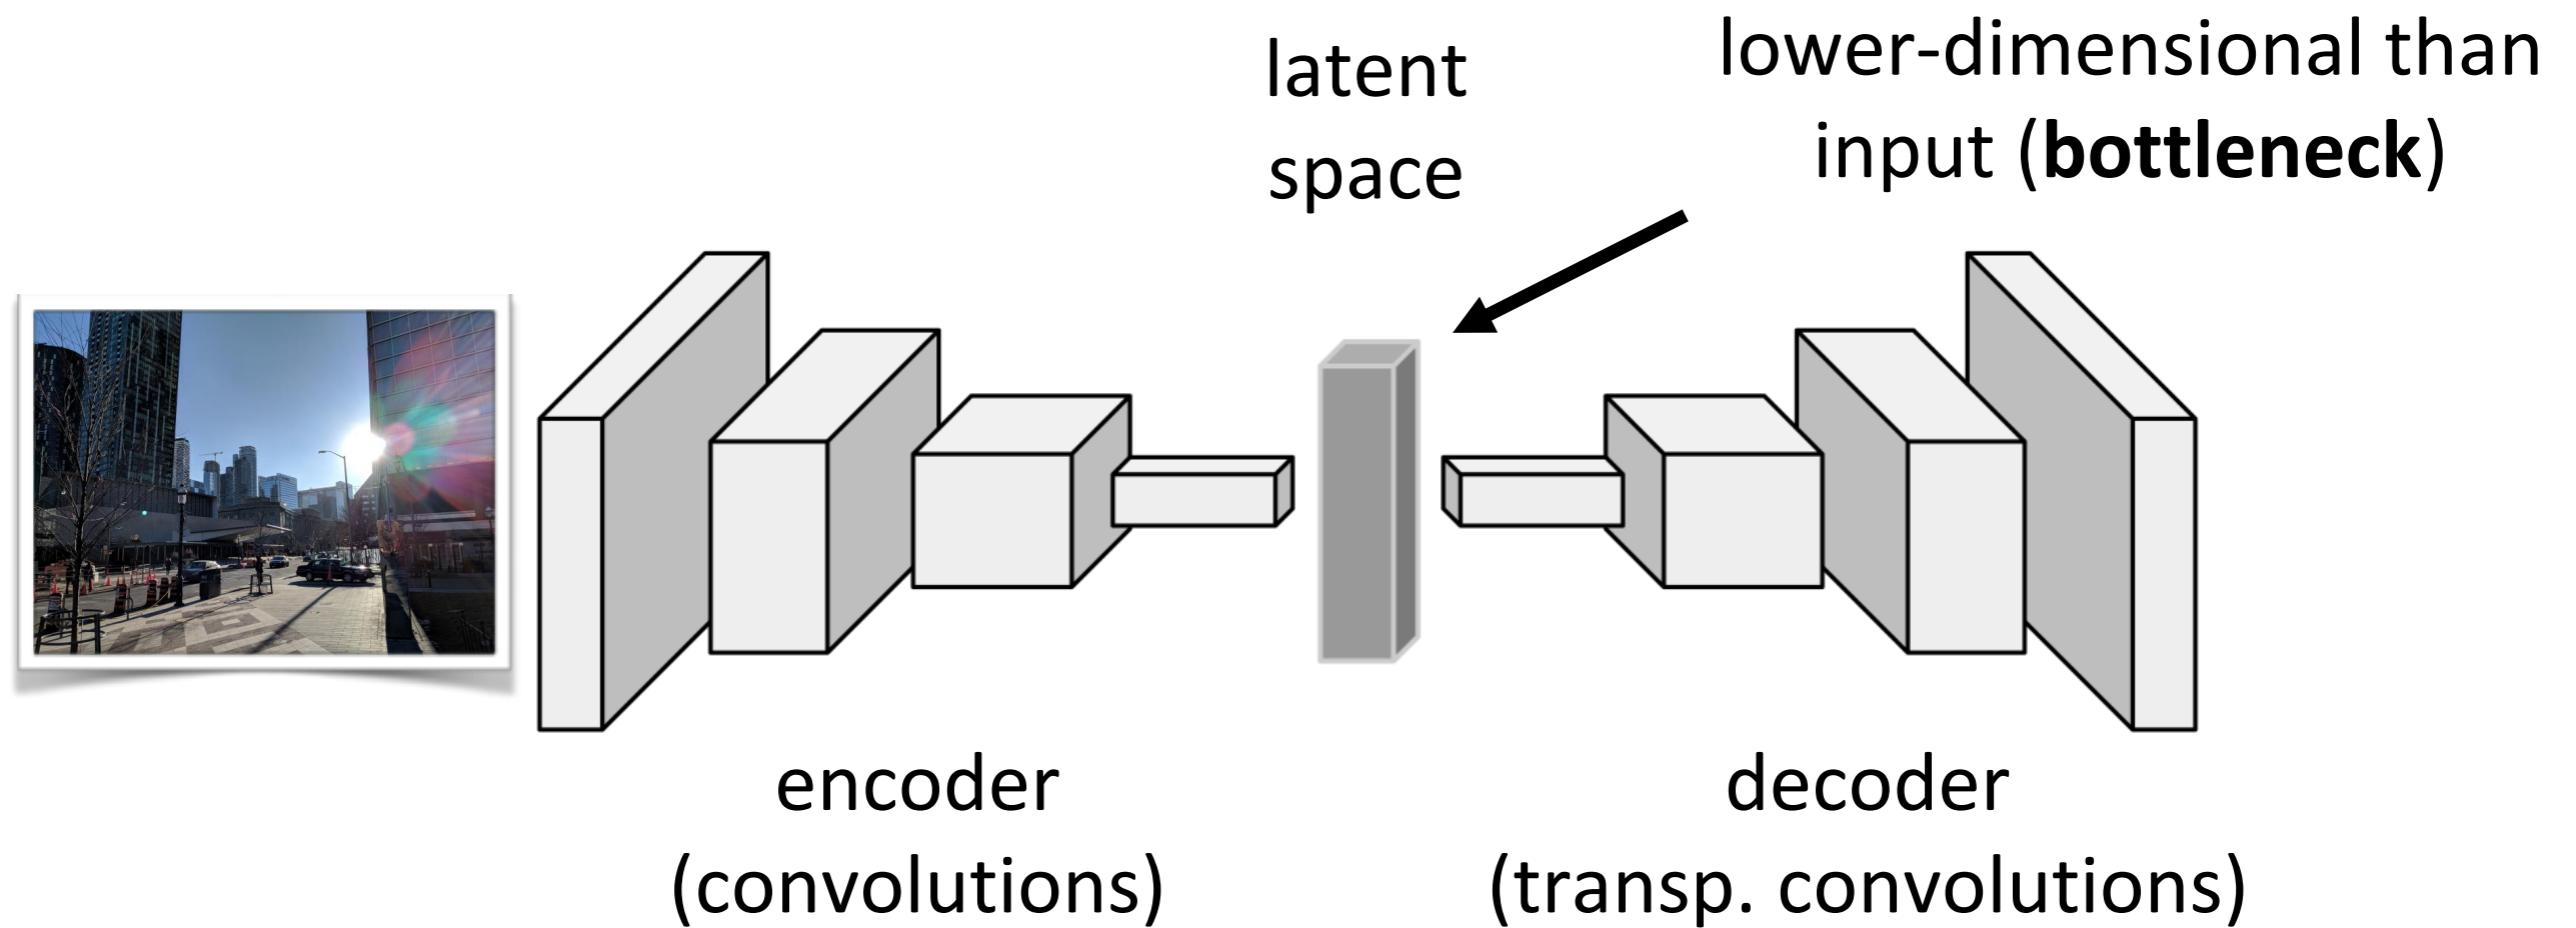
Autoencoder (AE)



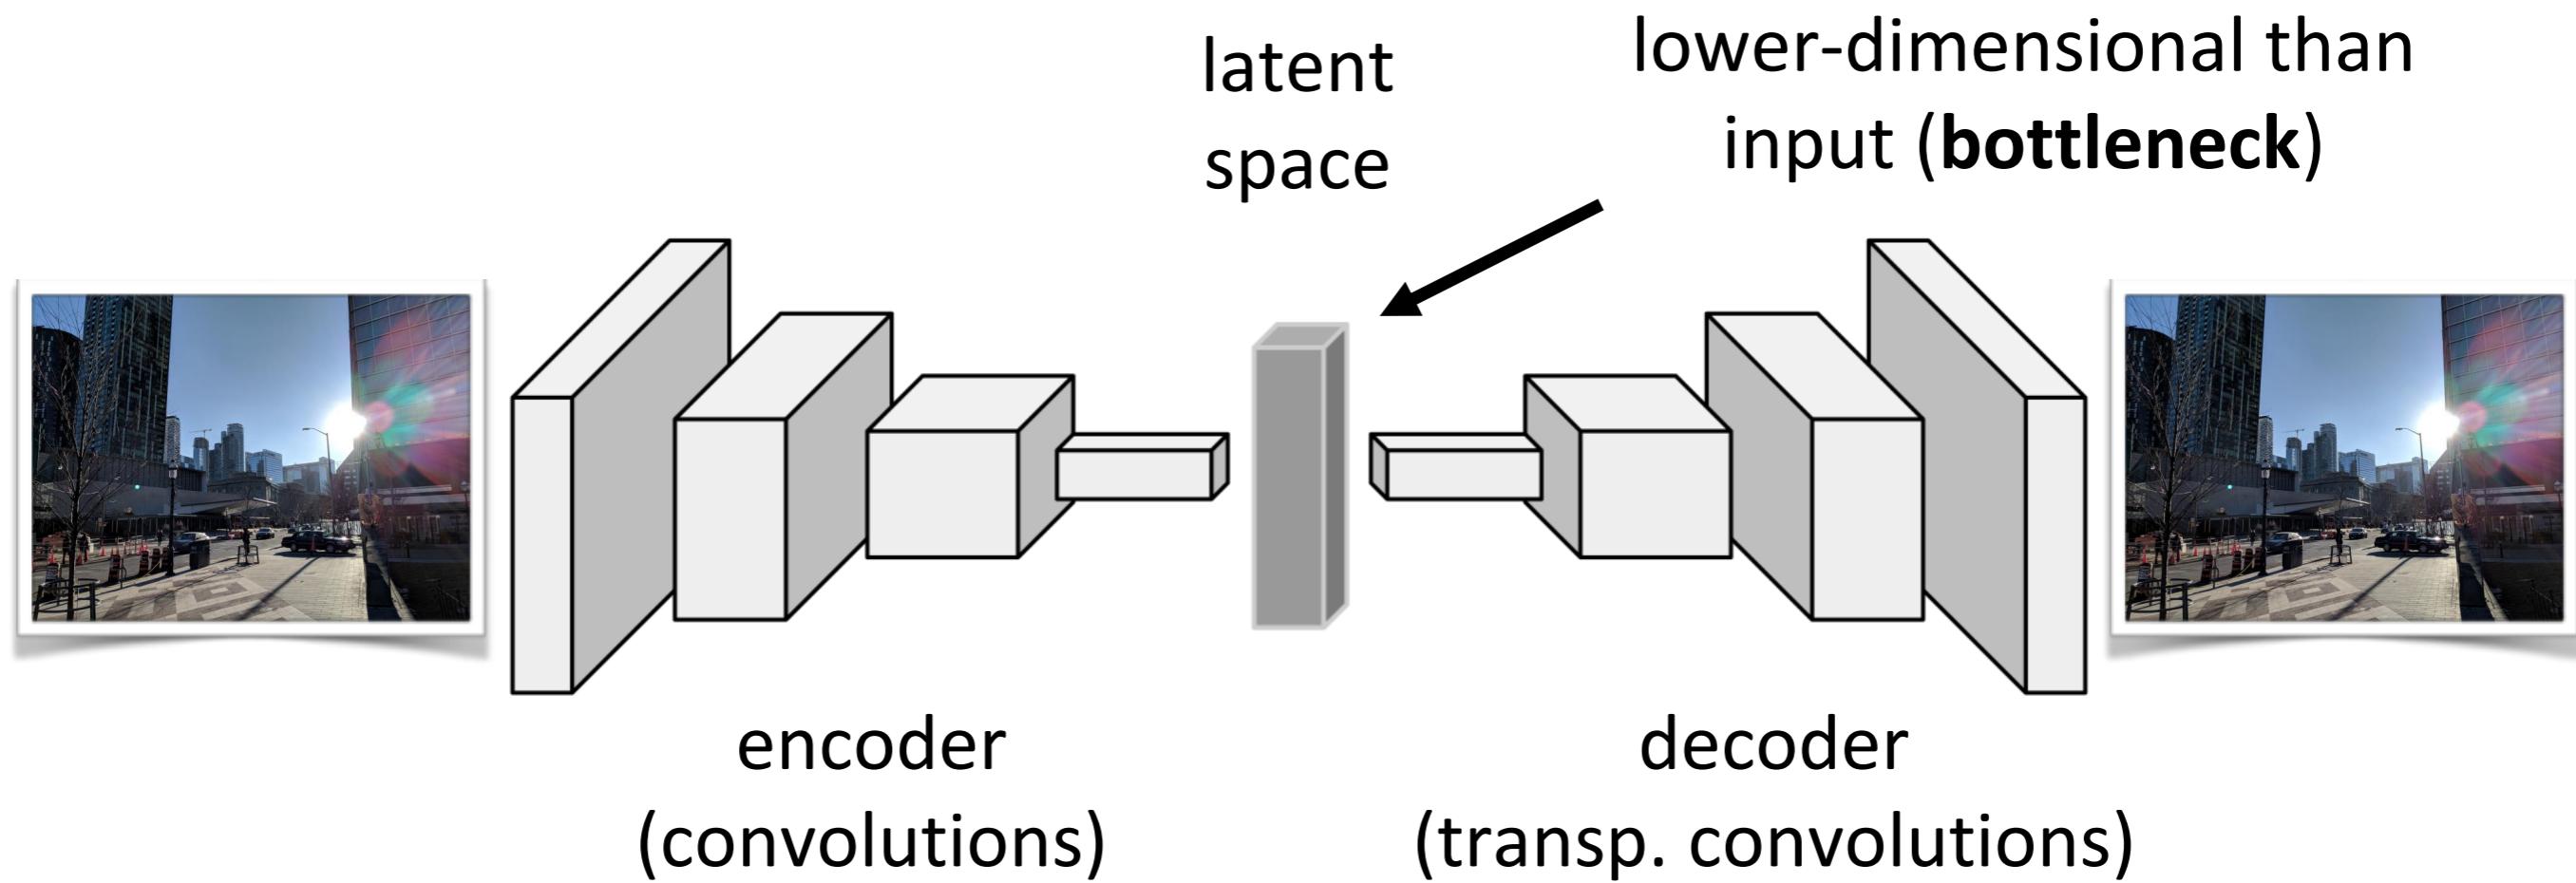
Autoencoder (AE)



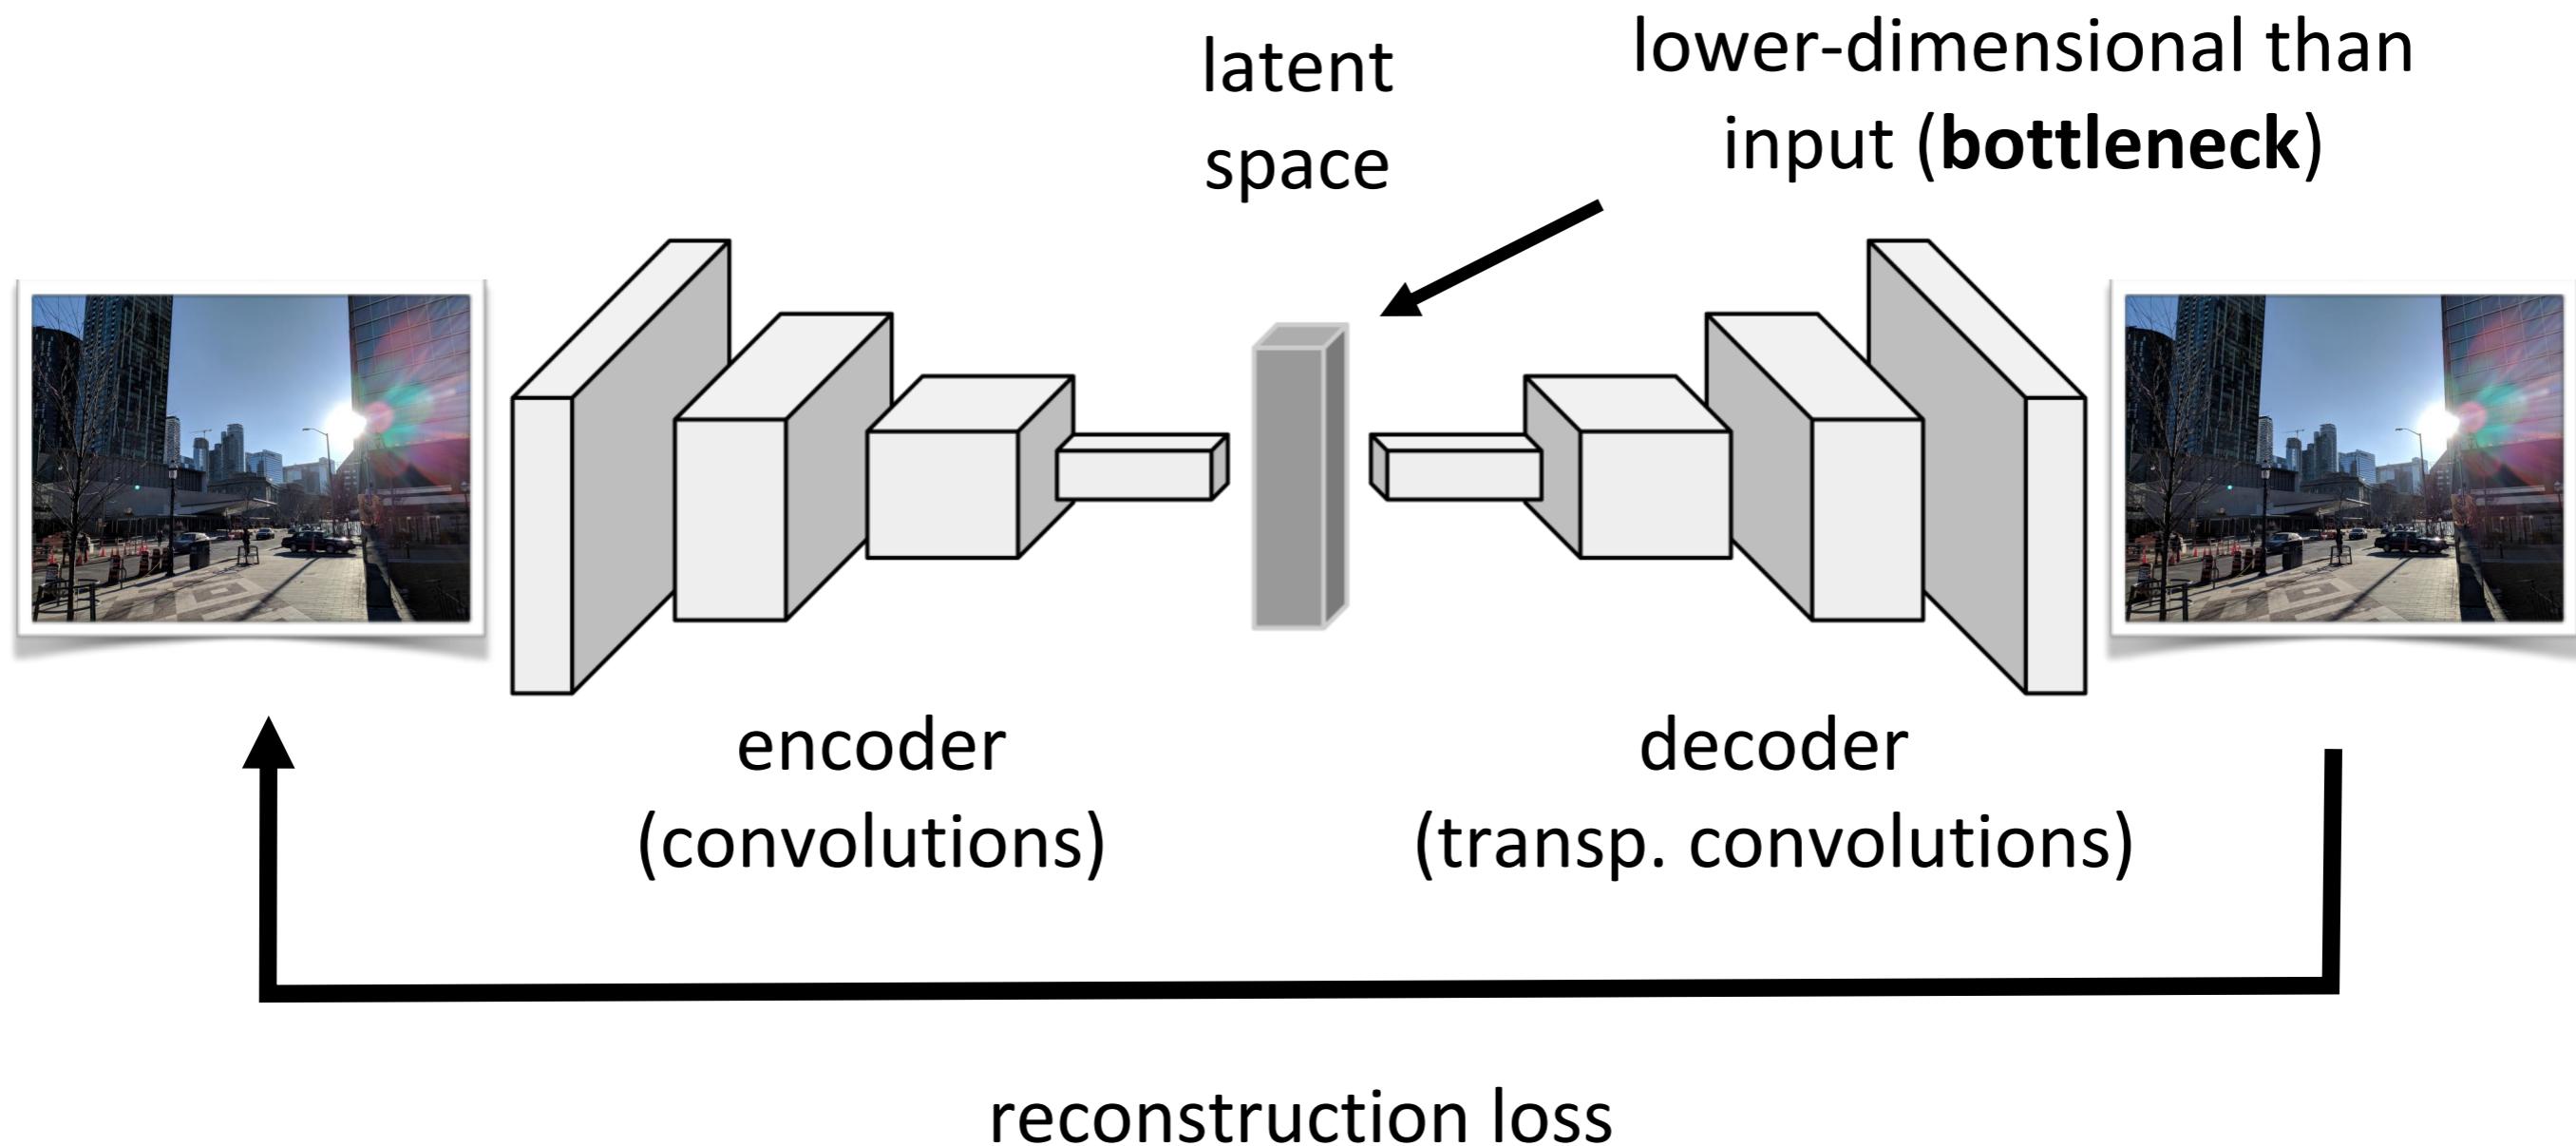
Autoencoder (AE)



Autoencoder (AE)



Autoencoder (AE)



Variational Autoencoder (VAE)

structured latent

space

$$\mathbf{z} \sim p_{\text{model}}$$

lower-dimensional than
input (**bottleneck**)

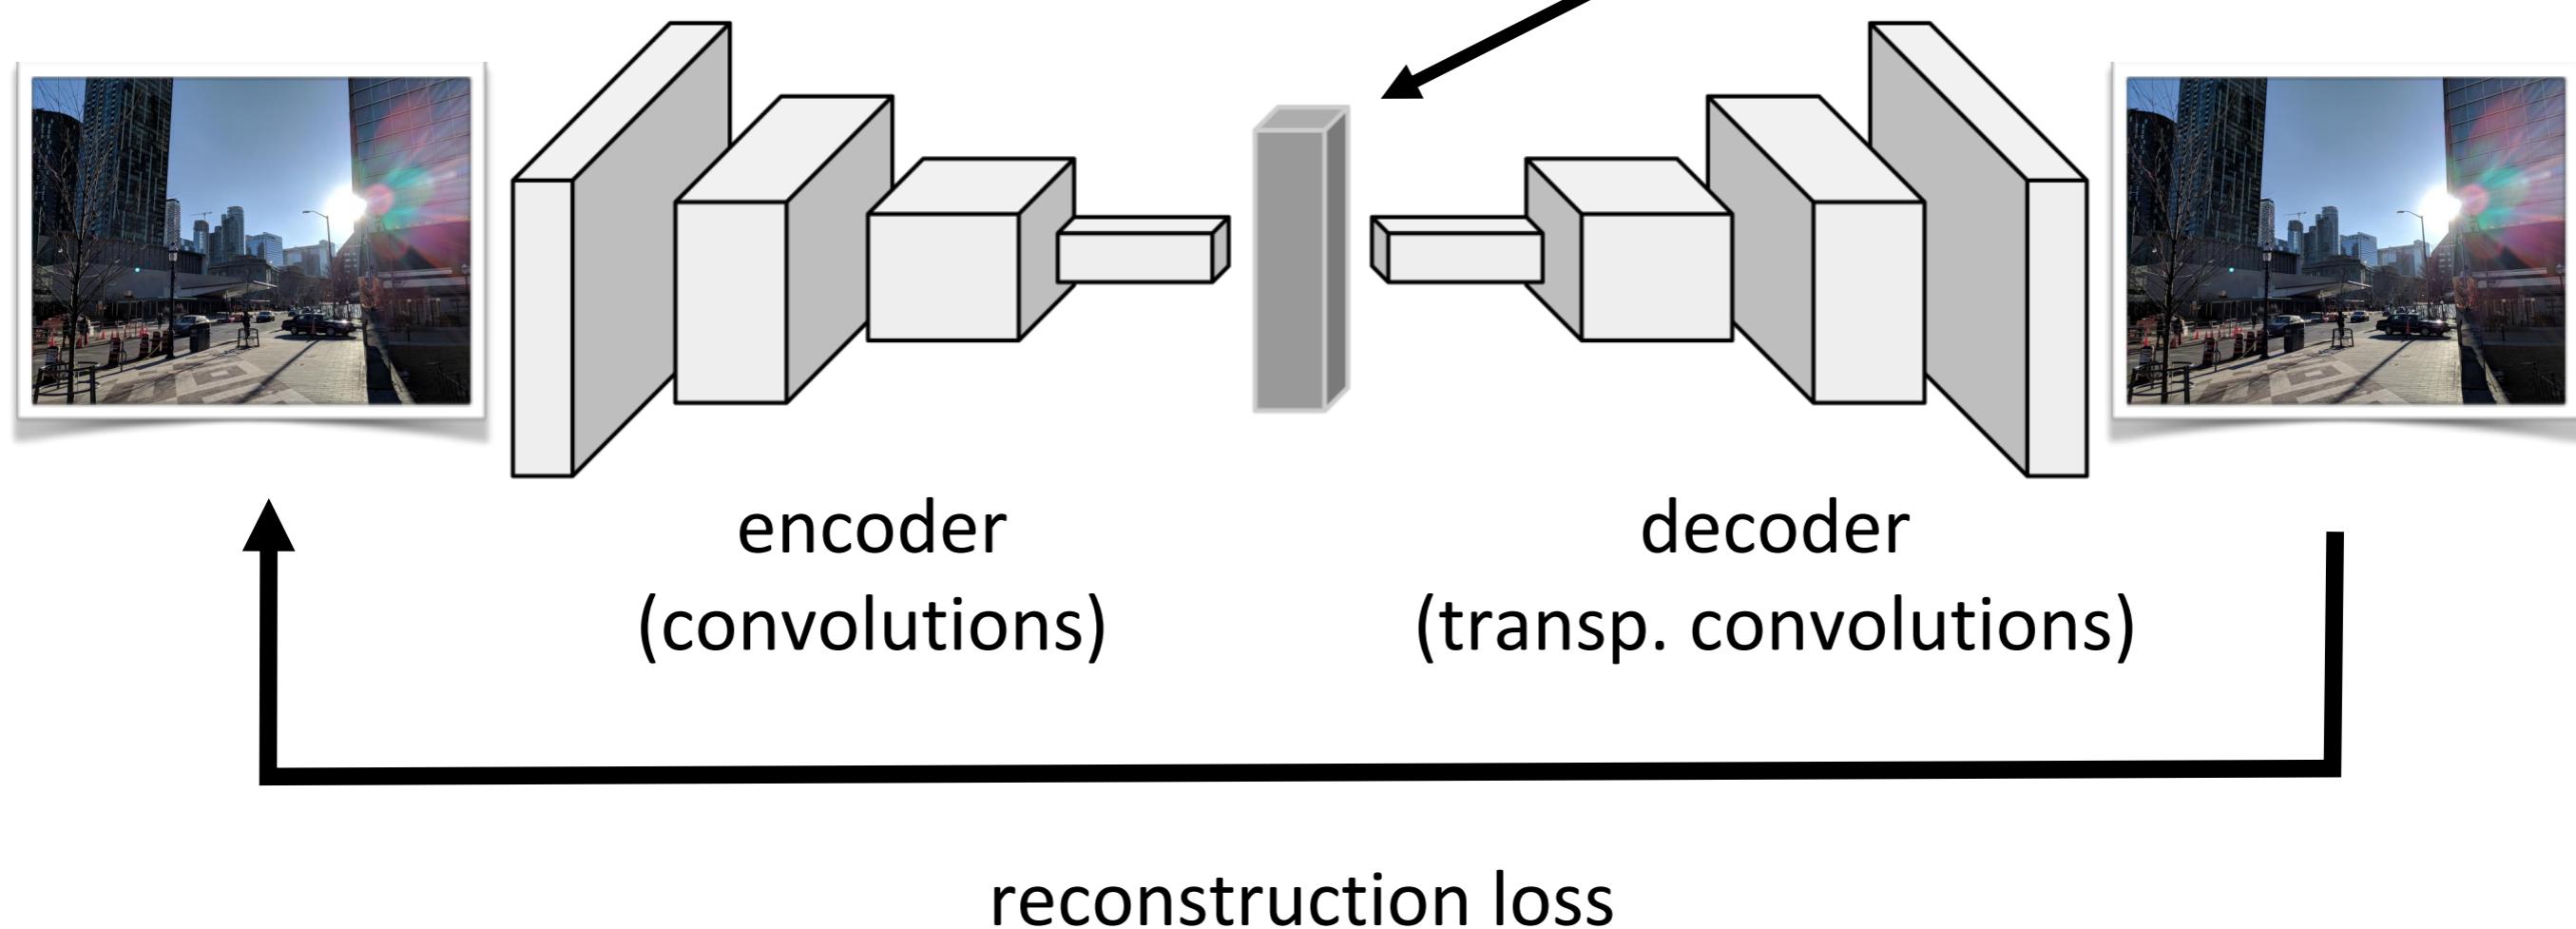


Image Generation via a VAE

sample point in latent space

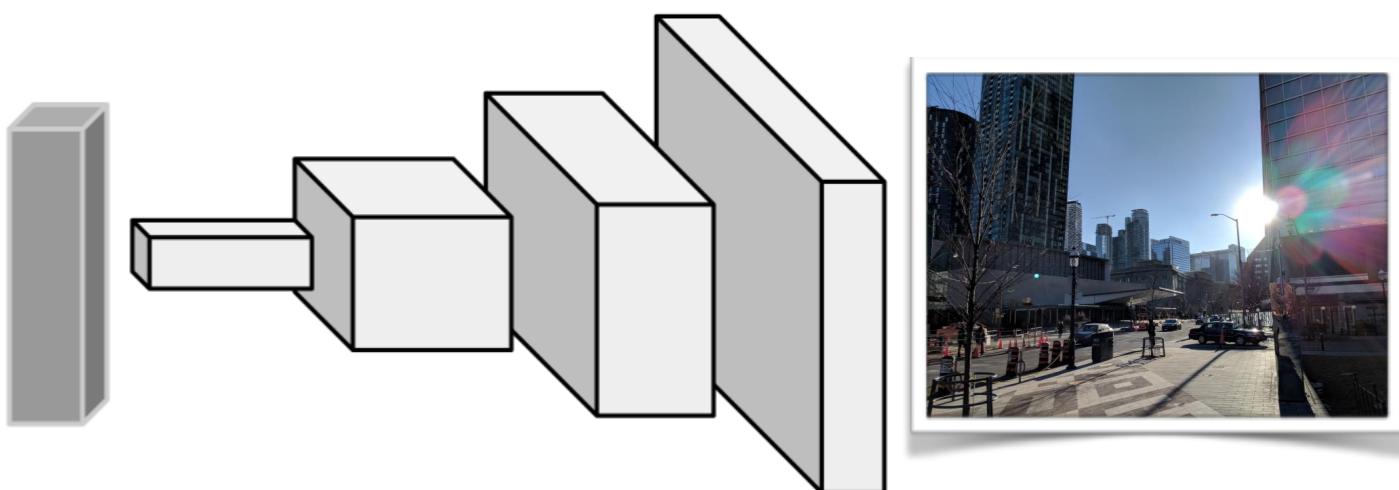
$$\mathbf{z} \sim p_{\text{model}}$$



Image Generation via a VAE

sample point in latent space

$$\mathbf{z} \sim p_{\text{model}}$$



decoder
(transp. convolutions)

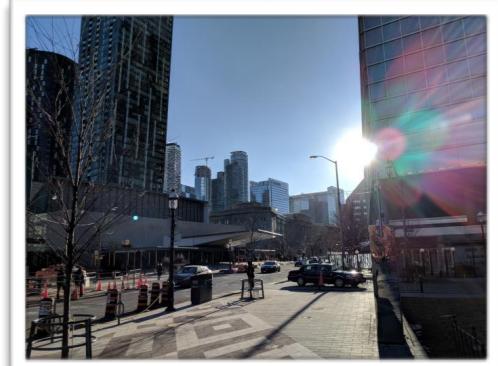
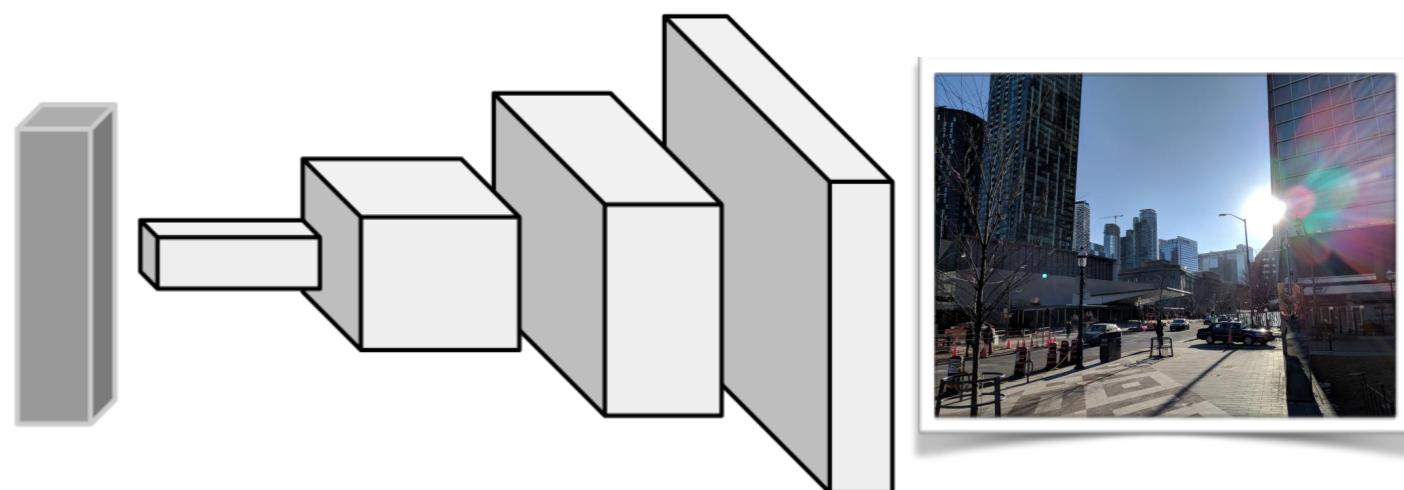


Image Generation via a VAE

sample point in latent space

$$\mathbf{z} \sim p_{\text{model}}$$



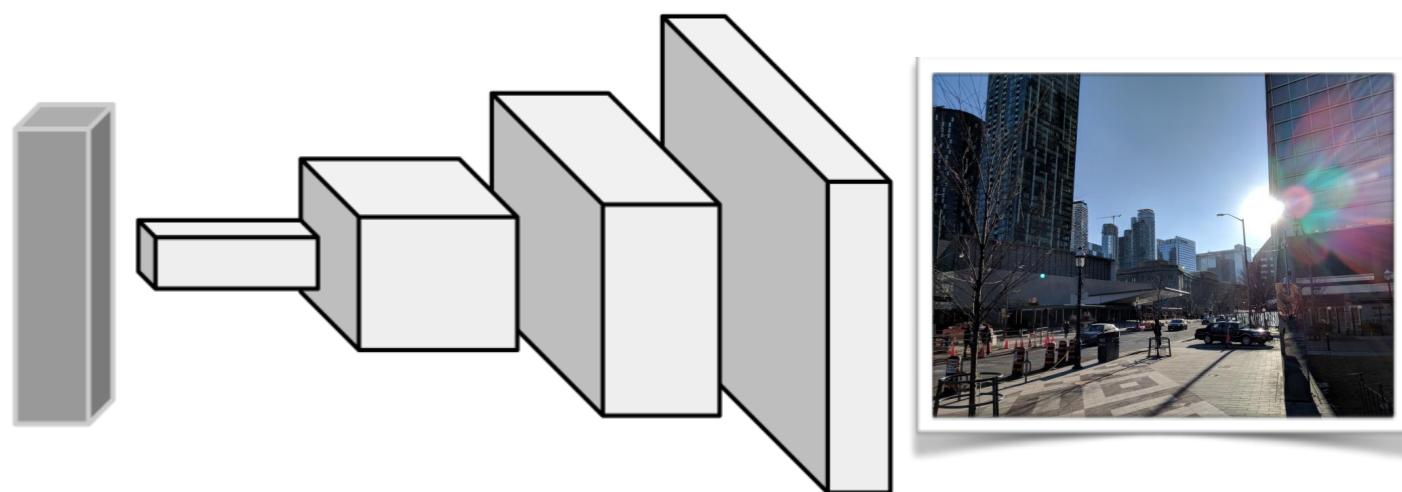
decoder
(transp. convolutions)

learn mapping from simple
to complex distribution

Image Generation via a VAE

sample point in latent space

$$\mathbf{z} \sim p_{\text{model}}$$



decoder
(transp. convolutions)

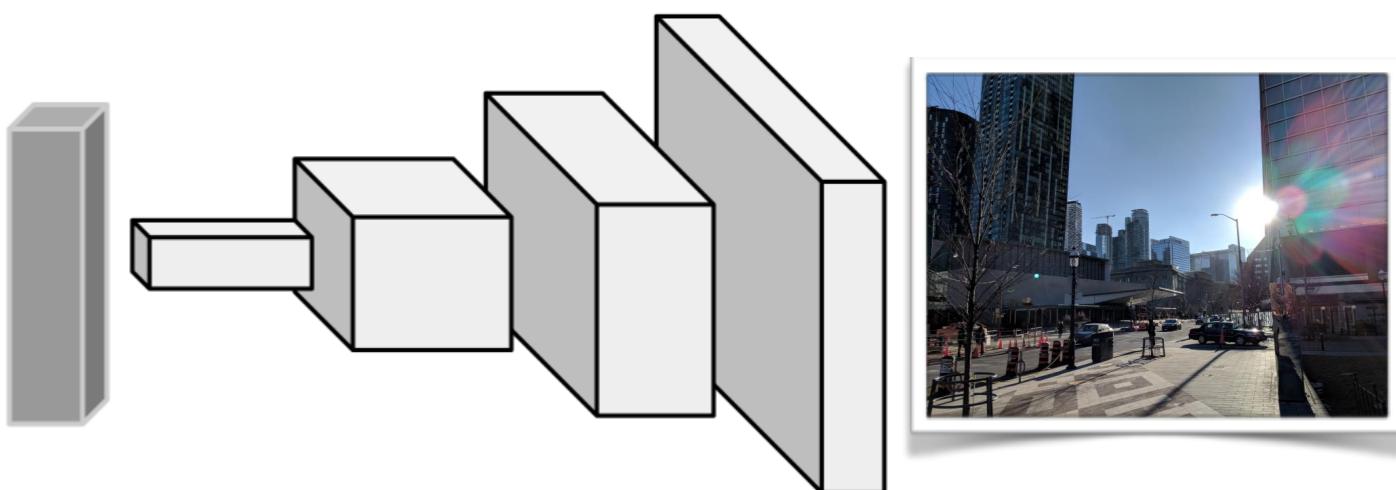
$$p(\mathbf{x}|\mathbf{z})$$

learn mapping from simple
to complex distribution

Image Generation via a VAE

sample point in latent space

$$\mathbf{z} \sim p_{\text{model}}$$



decoder
(transp. convolutions)

$$p(\mathbf{x}|\mathbf{z})$$

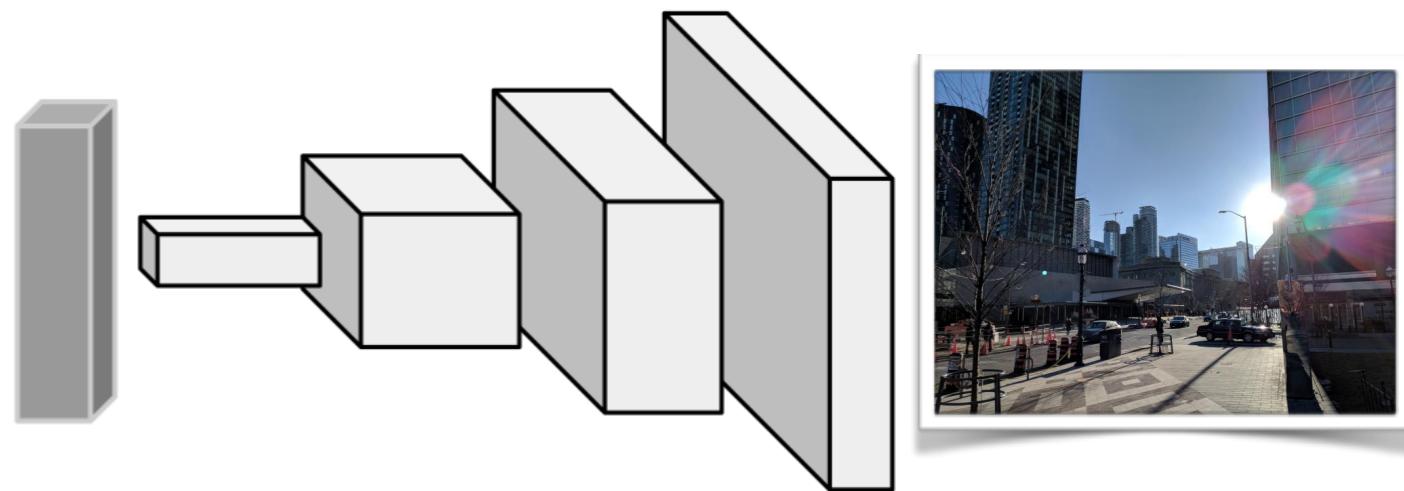
$$\mathbf{x} \sim p_{\text{data}} \\ (\text{ideally})$$

learn mapping from simple
to complex distribution

Image Generation via a VAE

sample point in latent space

$$\mathbf{z} \sim p_{\text{model}}$$



decoder

(transp. convolutions)

$$p(\mathbf{x}|\mathbf{z})$$

$$\mathbf{x} \sim p_{\text{data}}$$

(ideally)

How to ensure that
samples in latent space
lead to meaningful
images?

learn mapping from simple
to complex distribution

Theory

Want to maximize

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

Theory

Want to maximize

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

Intractable:

- Latent space is vector space
- Need to integrate over full space

Image Generation via a VAE

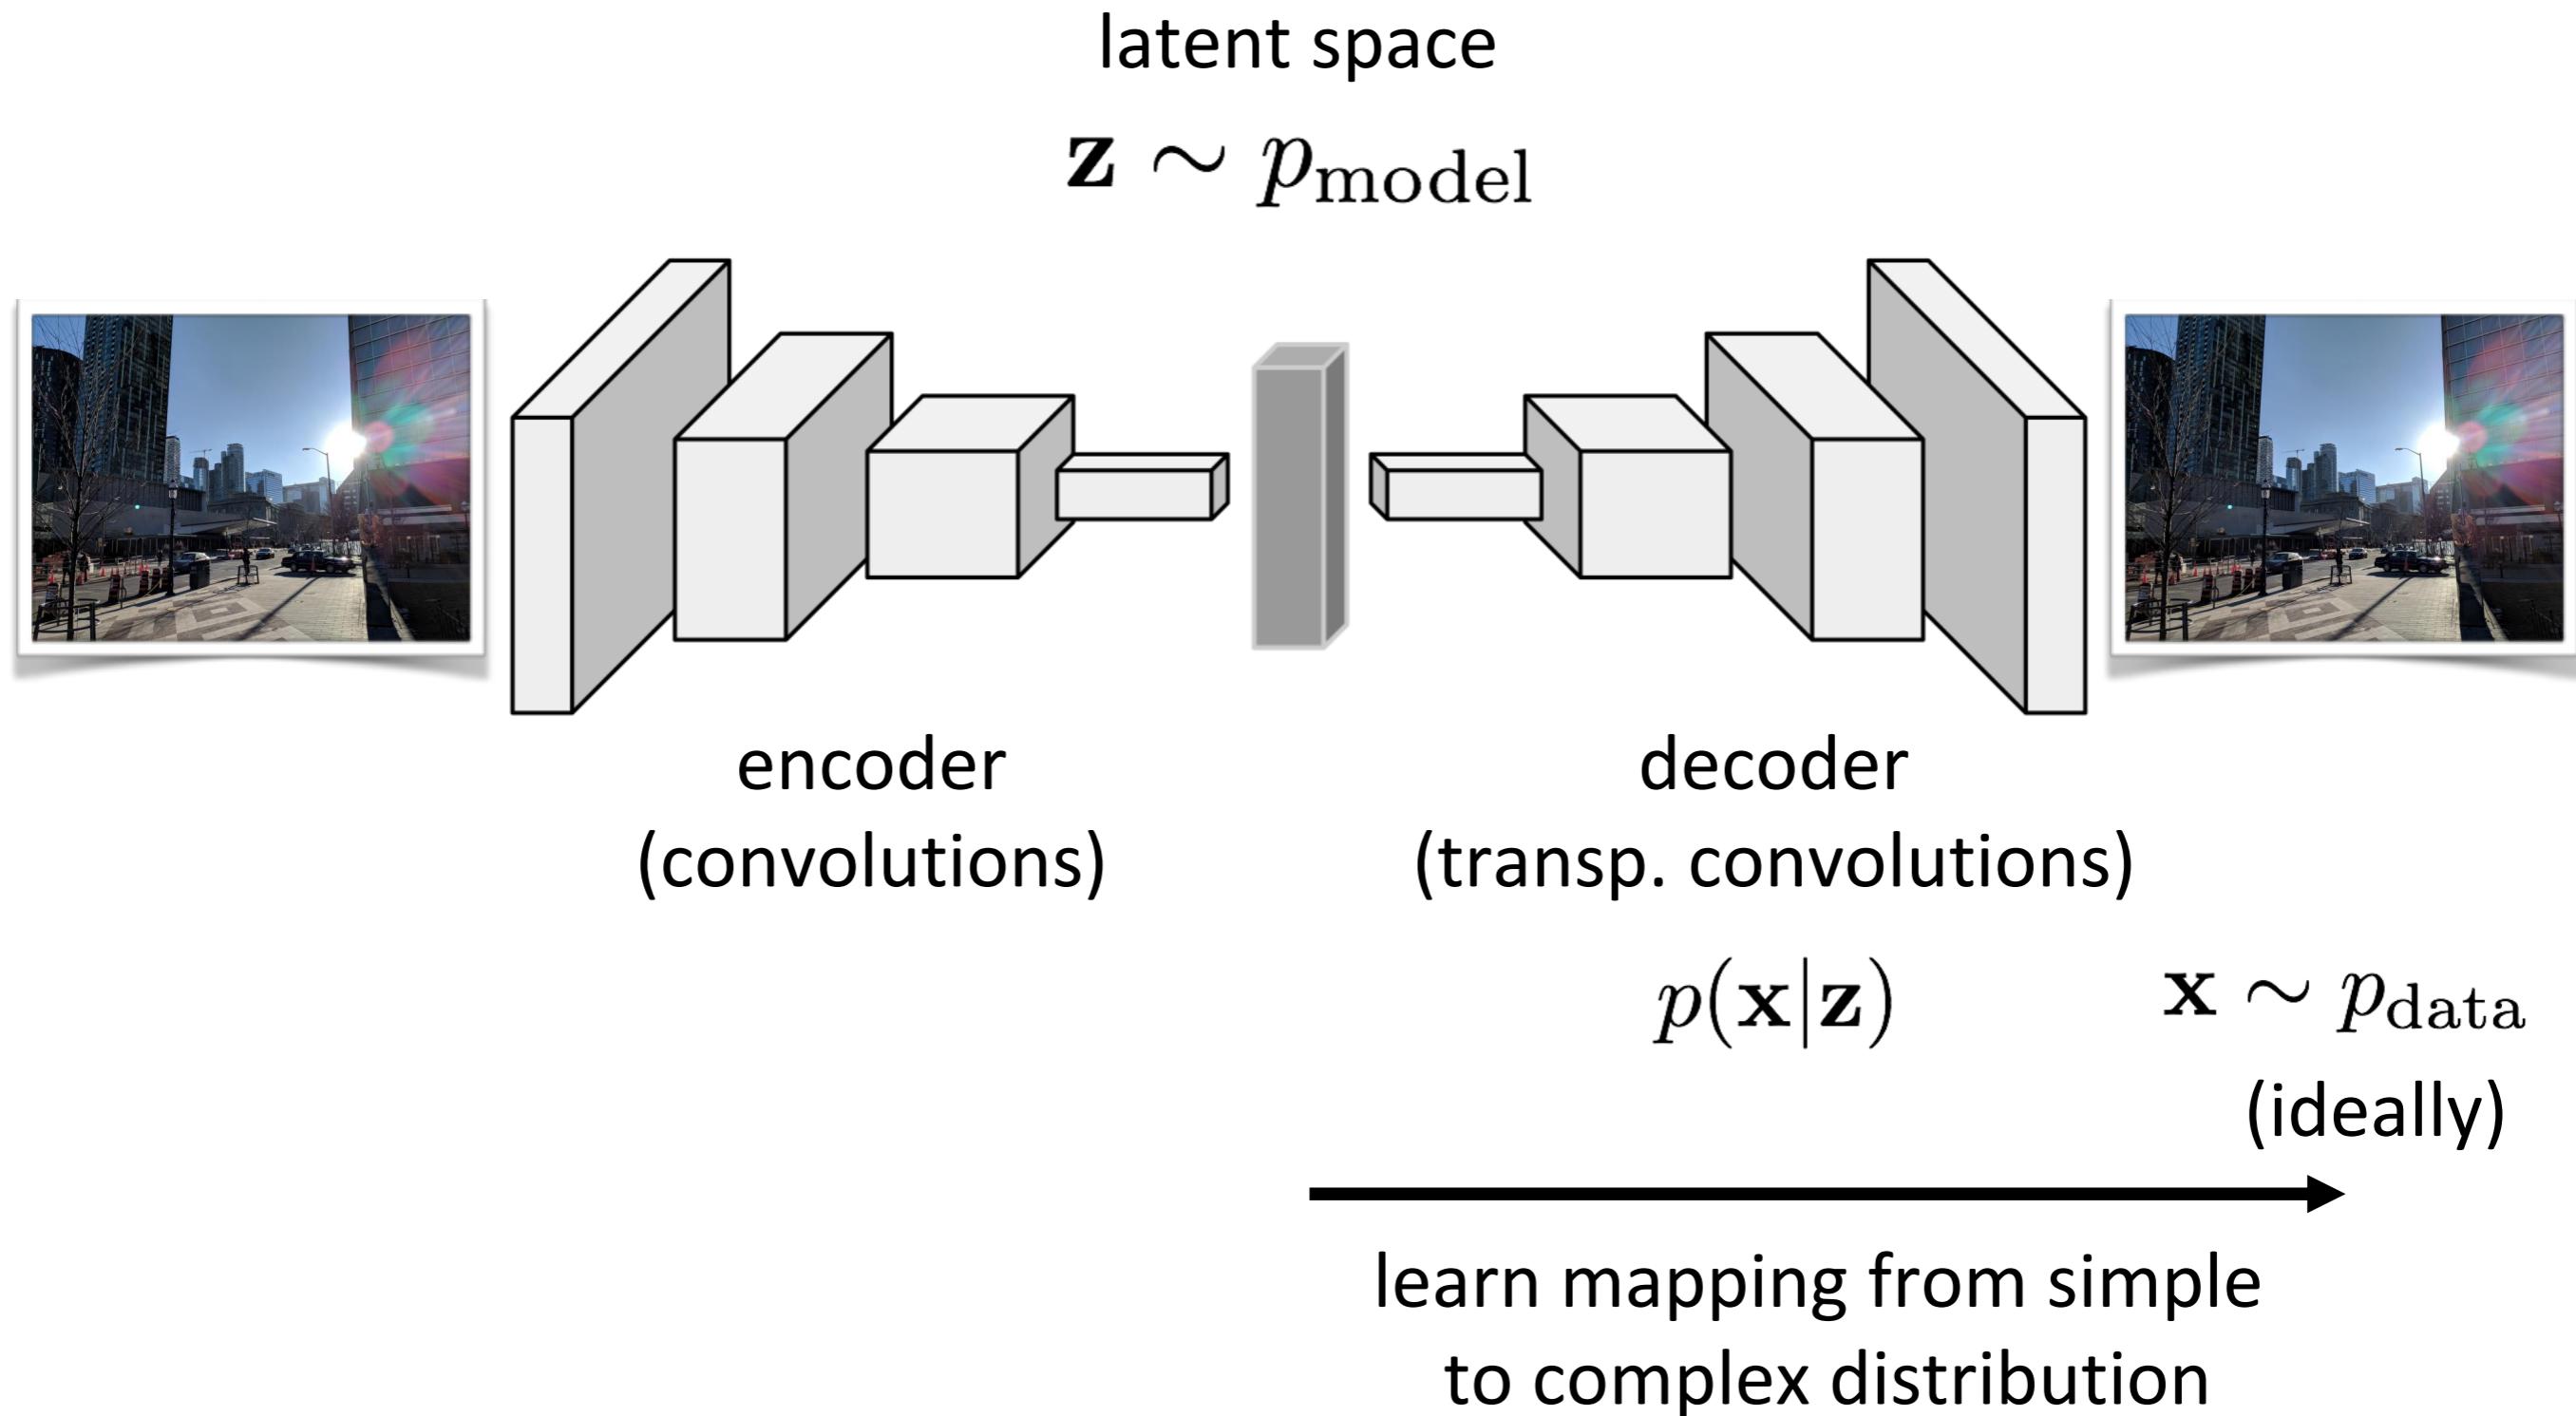


Image Generation via a VAE

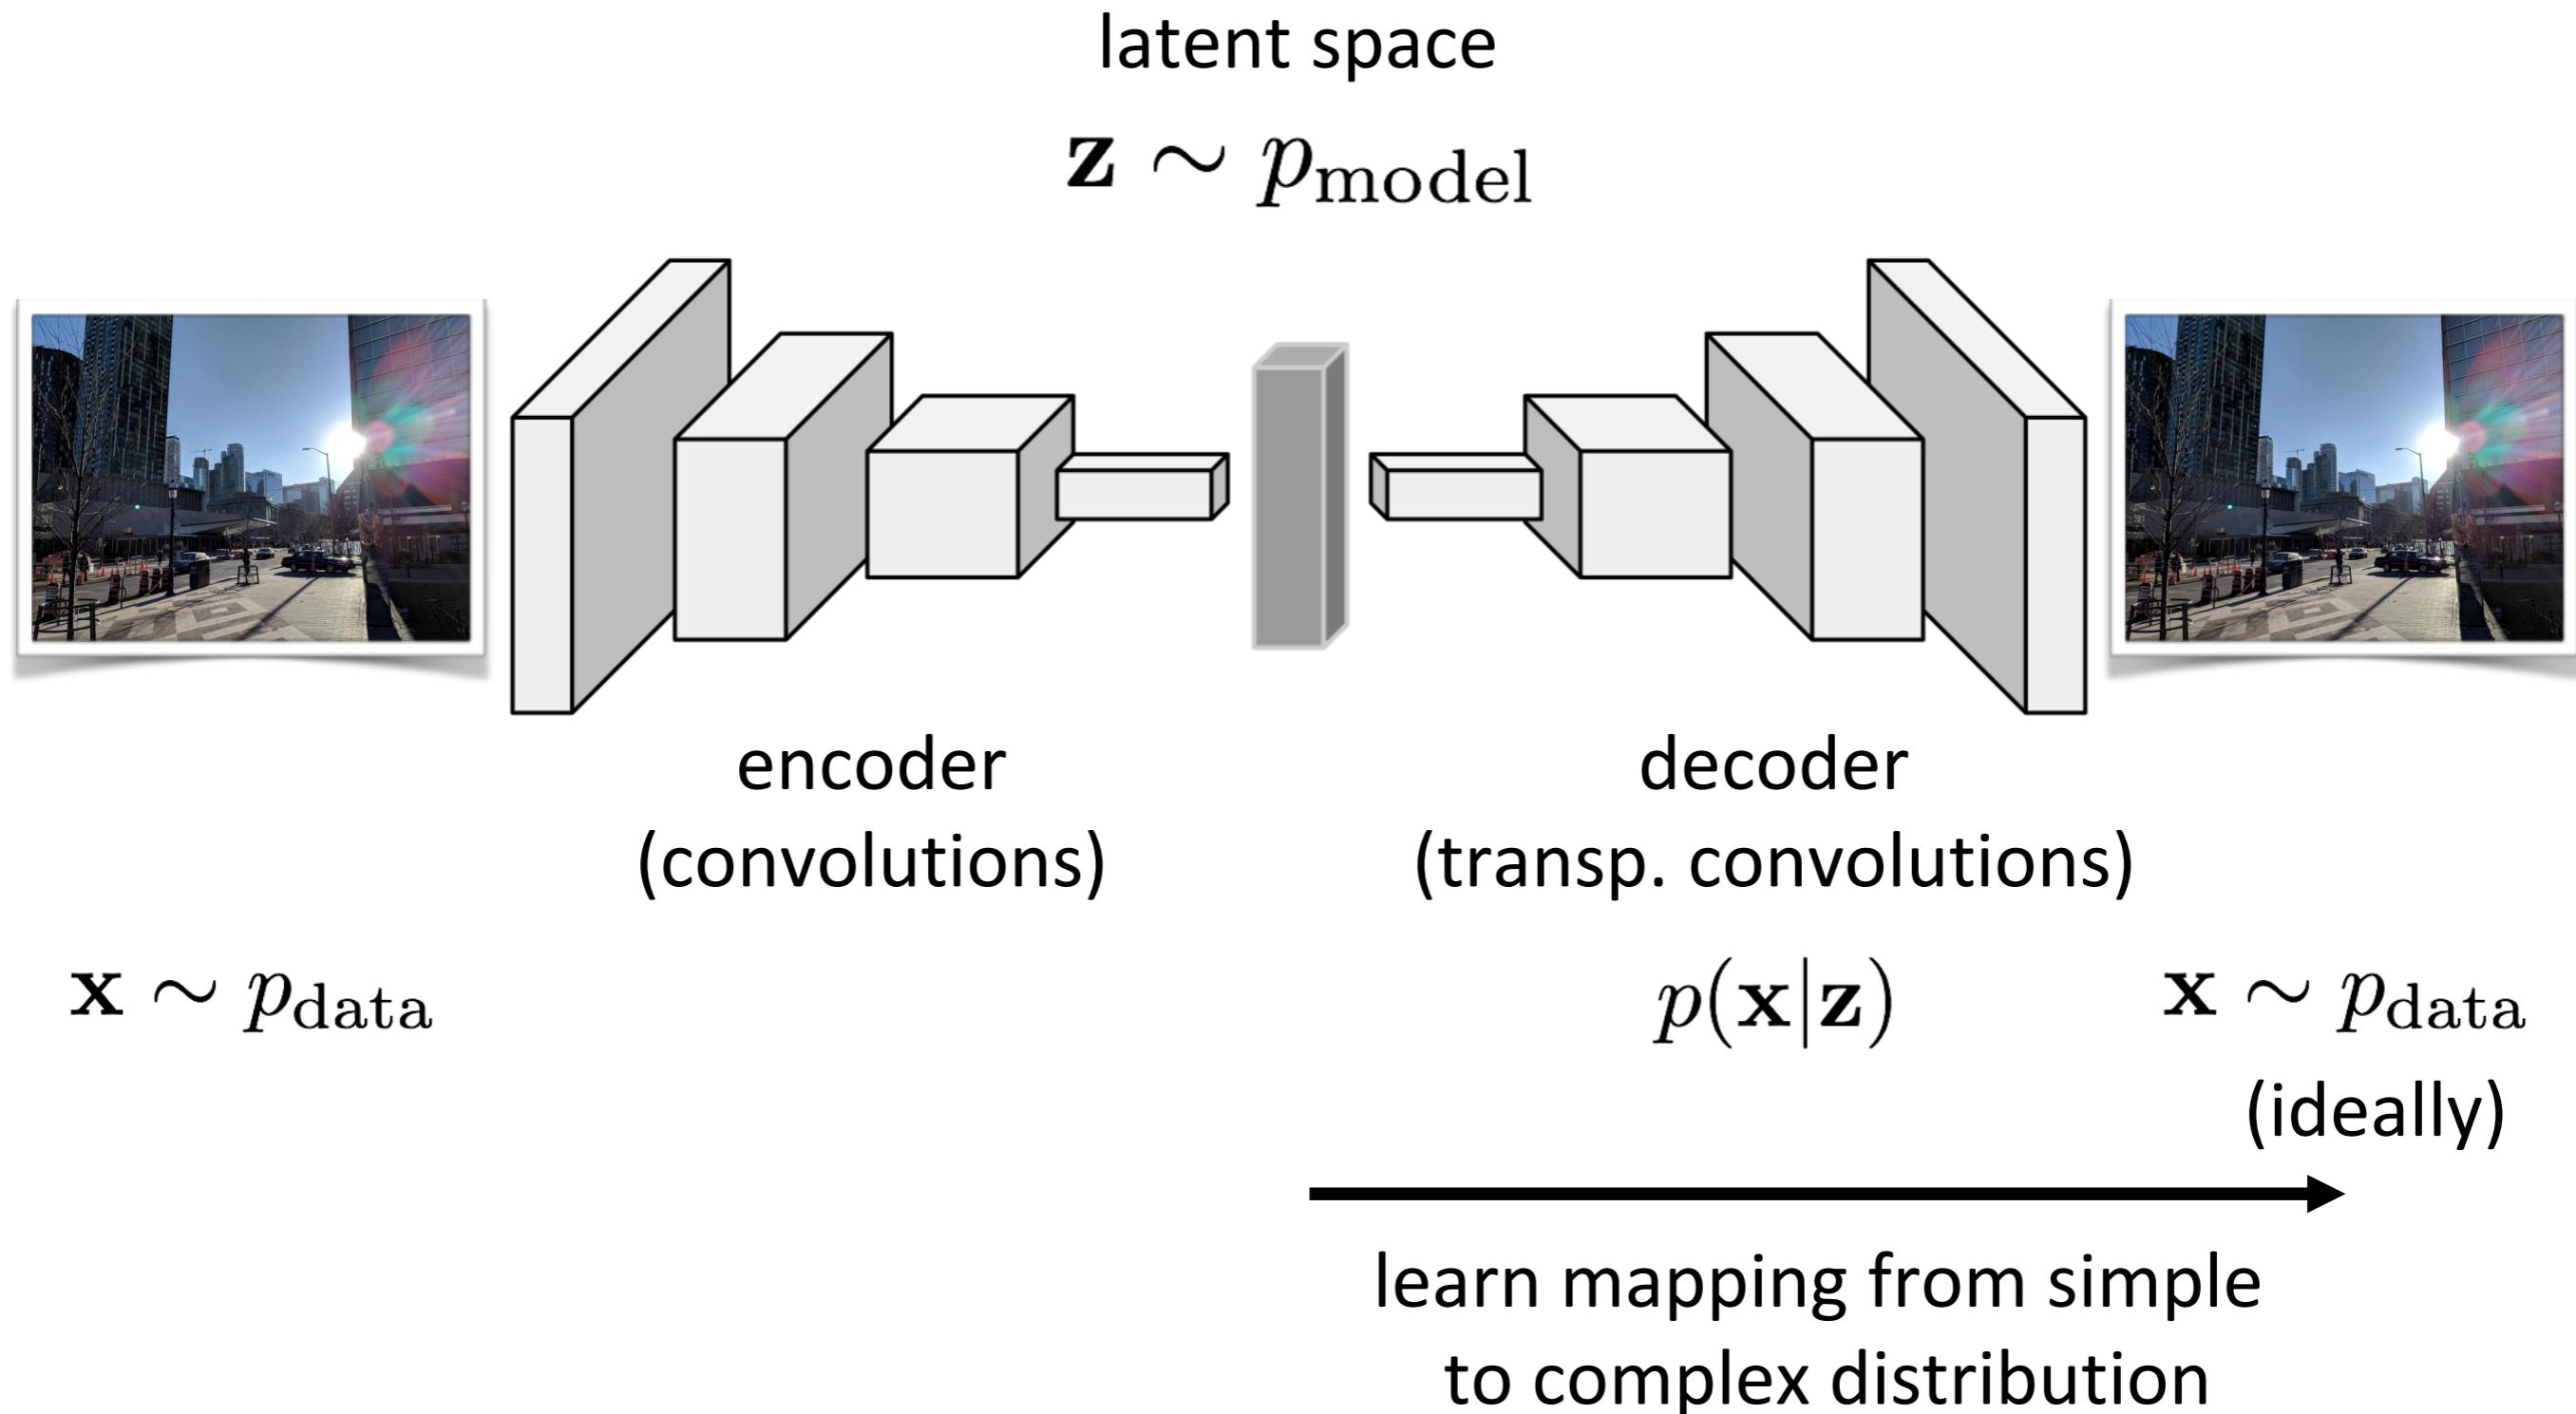


Image Generation via a VAE

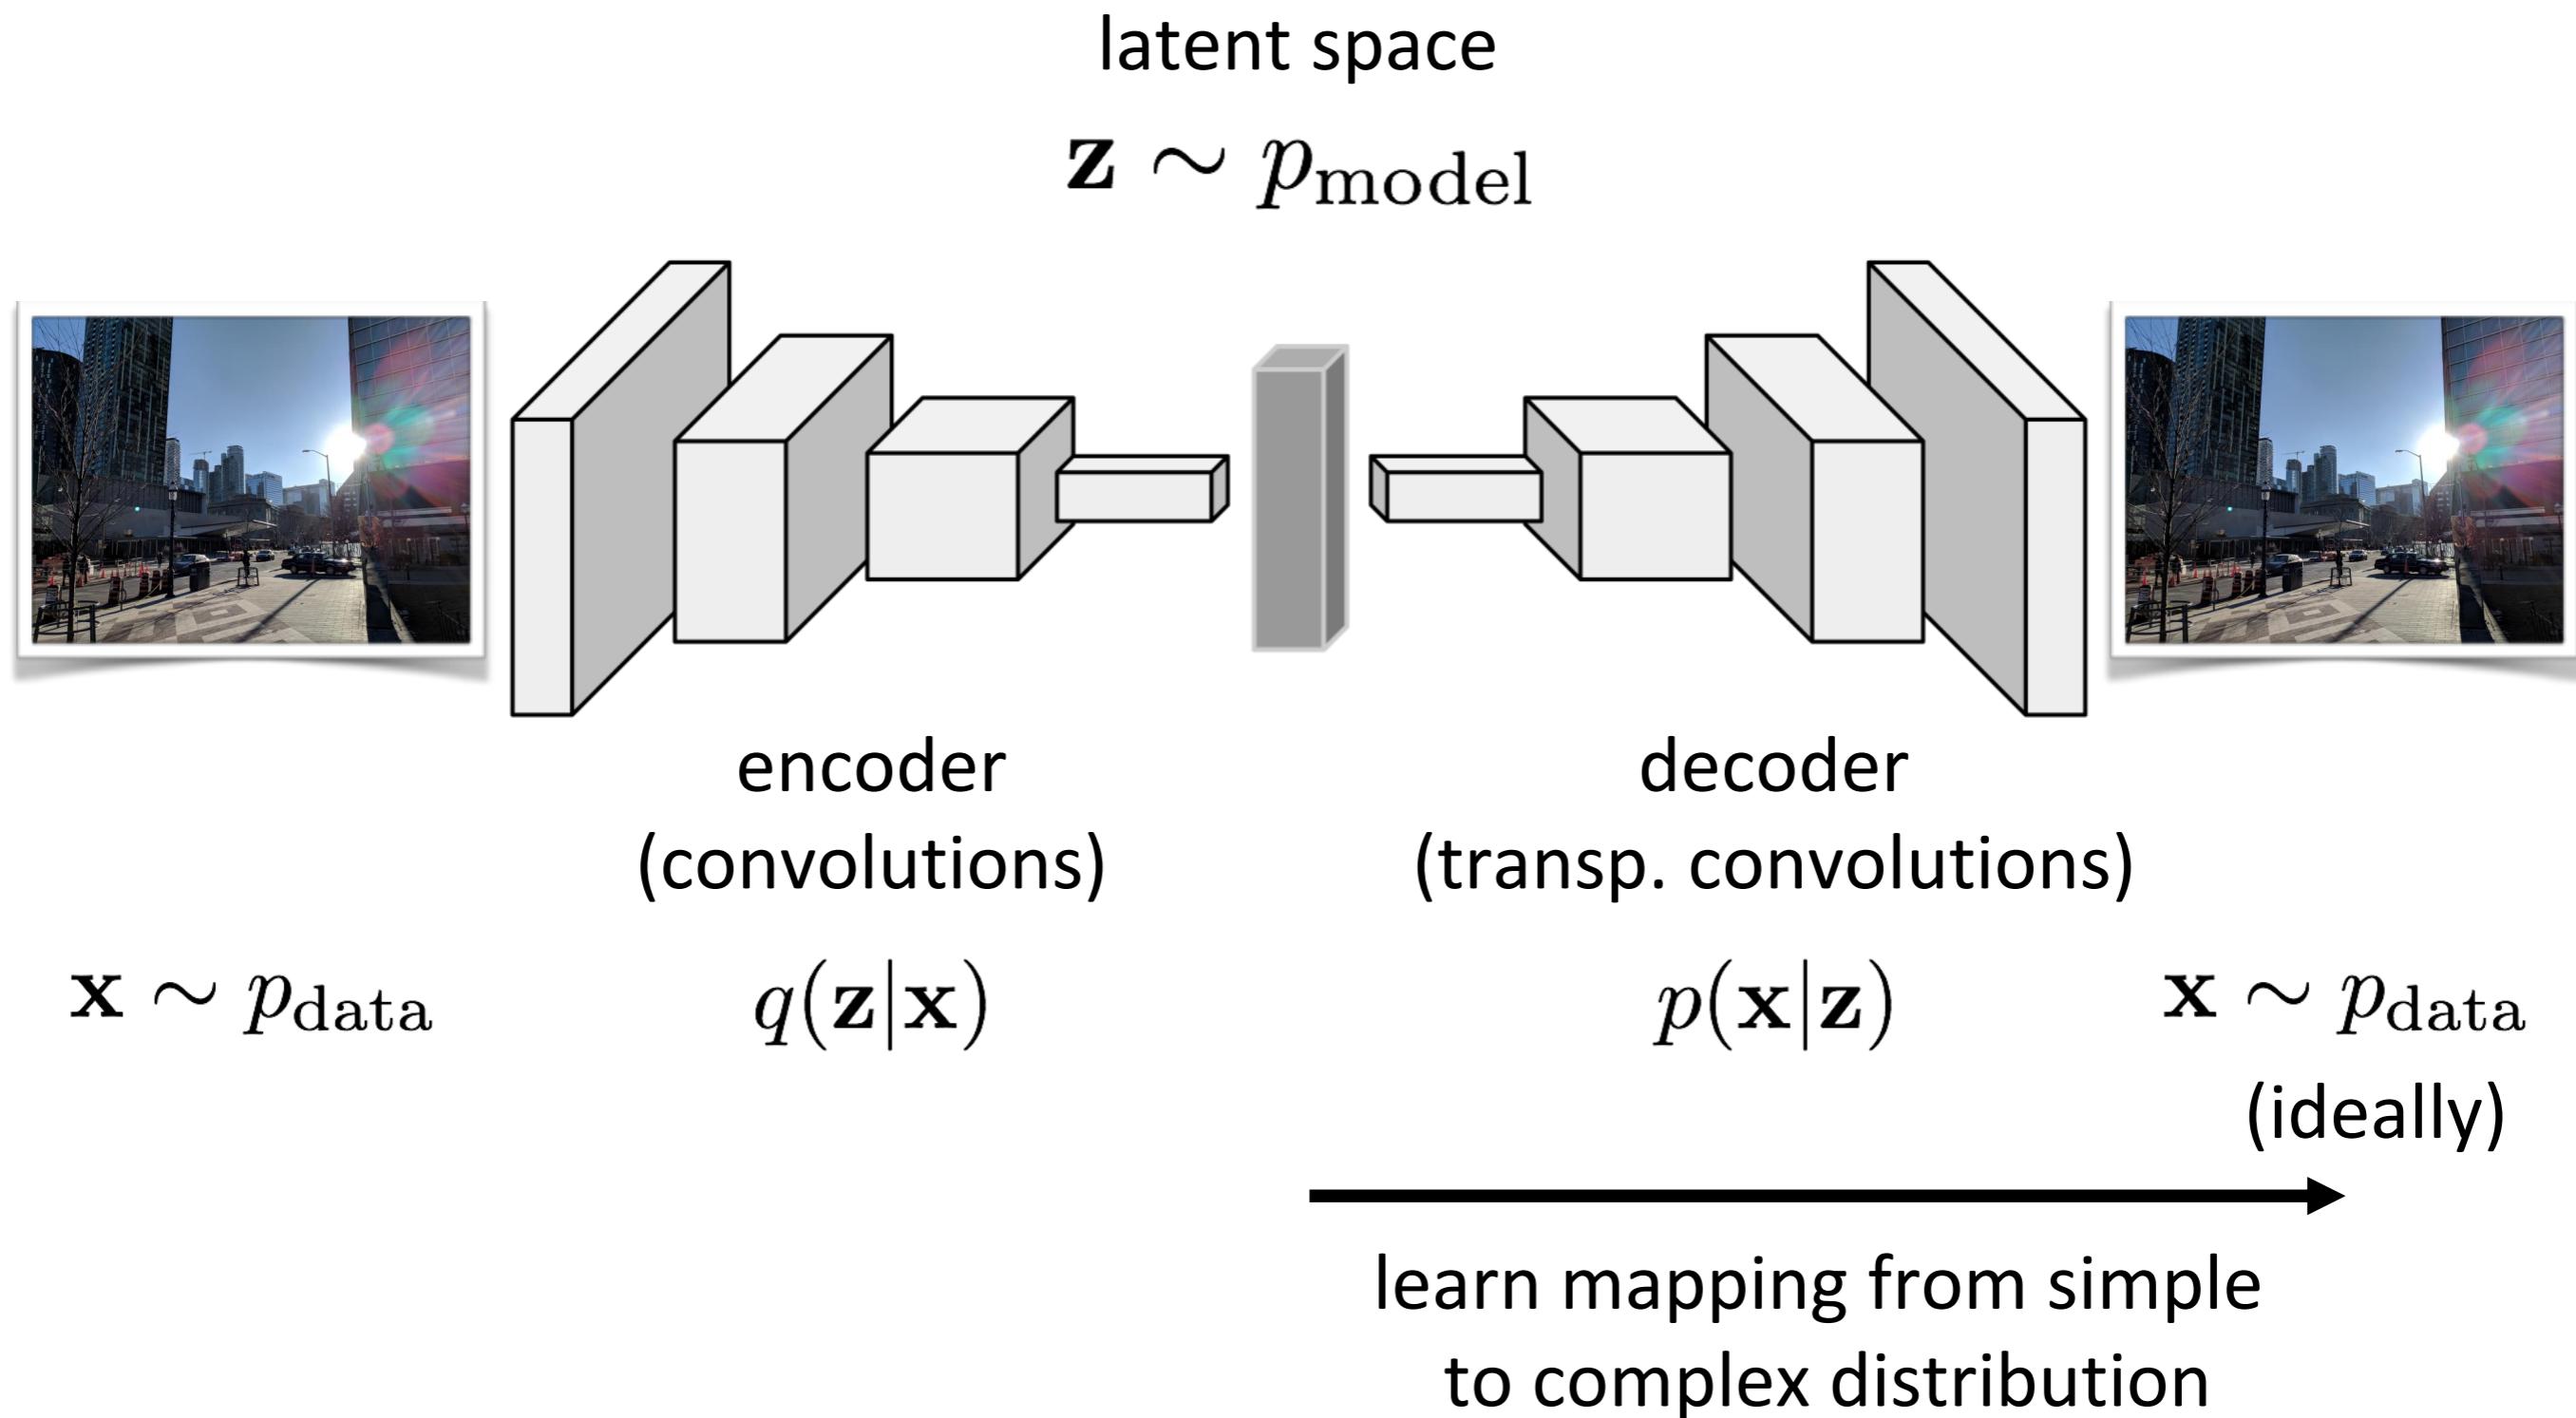
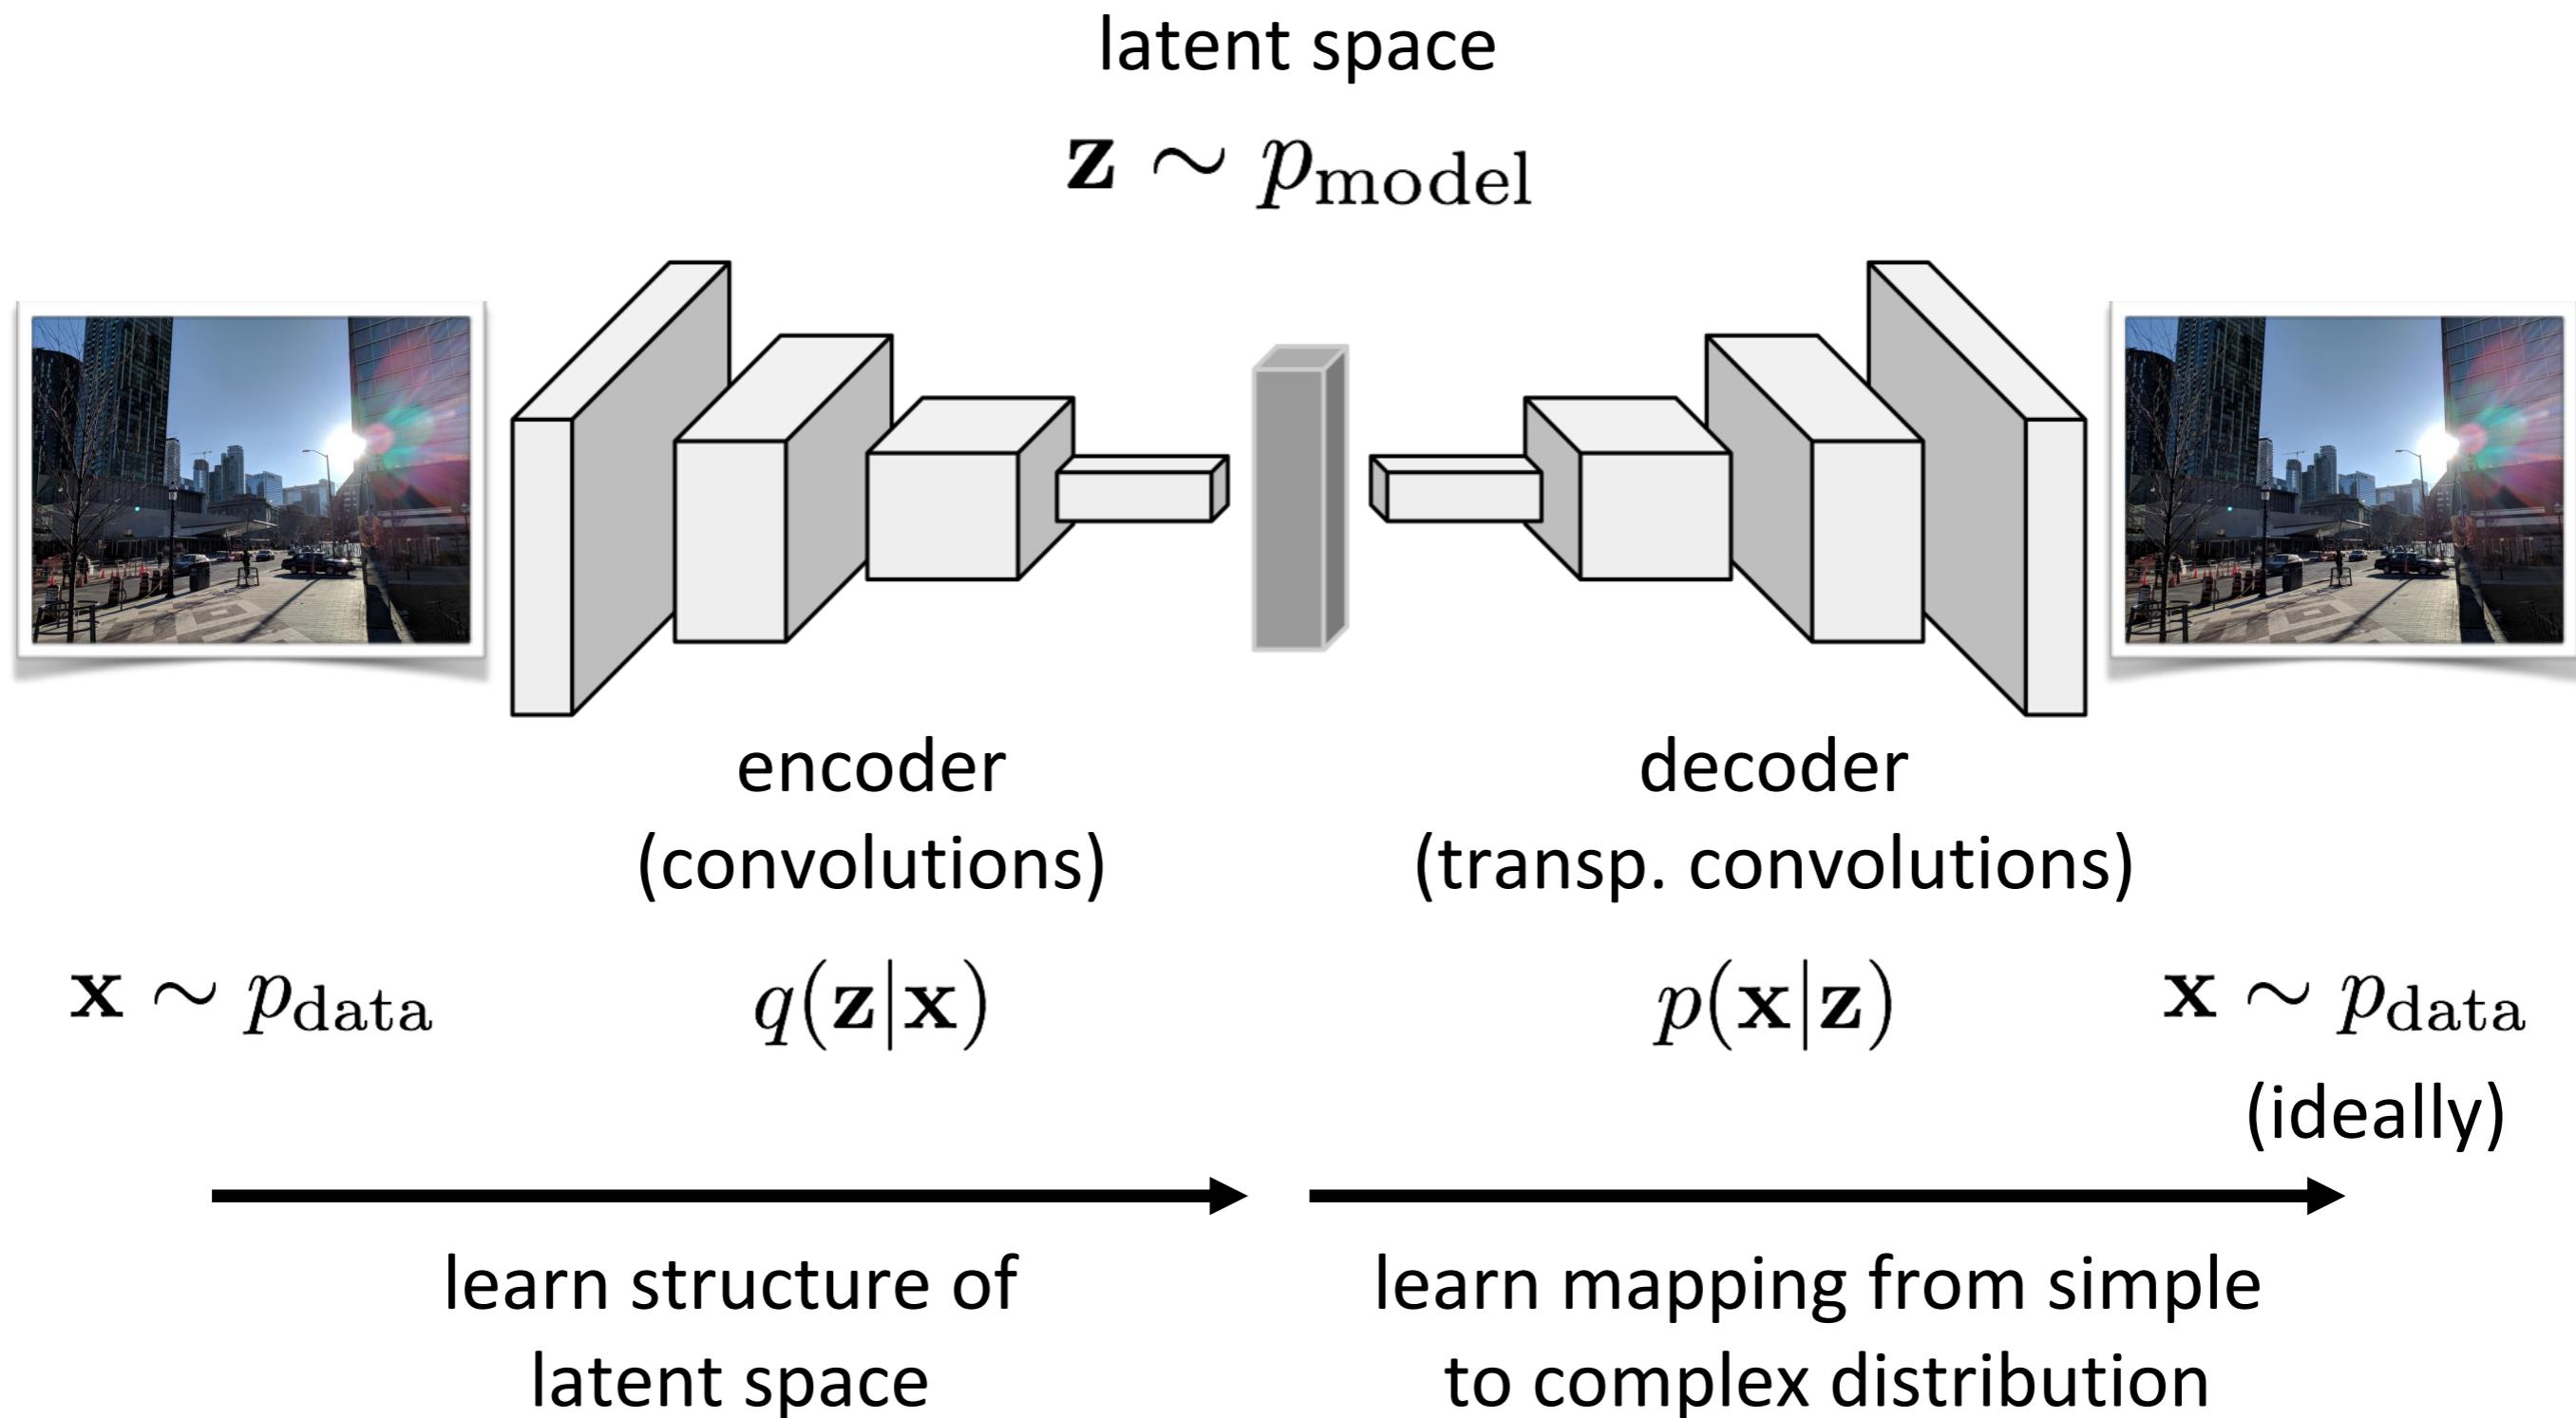


Image Generation via a VAE



Theory

Use encoder to generate probable samples, simplifying

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

Theory

Use encoder to generate probable samples, simplifying

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

to

$$p(\mathbf{x}) = E_{\mathbf{z} \sim q} p(\mathbf{x}|\mathbf{z})$$

Theory

Use encoder to generate probable samples, simplifying

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

to

$$p(\mathbf{x}) = E_{\mathbf{z} \sim q} p(\mathbf{x}|\mathbf{z})$$

Problem: Learned distribution q might not follow p_{model}

KL Divergence

Kullback-Leibler divergence: ``distance'' between probability distributions

$$D_{\text{KL}}(P||Q) = \int P(x) \log \left(\frac{P(x)}{Q(x)} \right) dx$$

KL Divergence

Kullback-Leibler divergence: ``distance'' between probability distributions

$$D_{\text{KL}}(P||Q) = \int P(x) \log \left(\frac{P(x)}{Q(x)} \right) dx$$

Properties:

KL Divergence

Kullback-Leibler divergence: ``distance'' between probability distributions

$$D_{\text{KL}}(P||Q) = \int P(x) \log \left(\frac{P(x)}{Q(x)} \right) dx$$

Properties:

- KL divergence is non-negative: $D_{\text{KL}}(P||Q) \geq 0$

KL Divergence

Kullback-Leibler divergence: “distance” between probability distributions

$$D_{\text{KL}}(P||Q) = \int P(x) \log \left(\frac{P(x)}{Q(x)} \right) dx$$

Properties:

- KL divergence is non-negative: $D_{\text{KL}}(P||Q) \geq 0$
- KL divergence is 0 exactly if: $D_{\text{KL}}(P||Q) = 0 \Leftrightarrow P(x) = Q(x) \forall x$

KL Divergence

Kullback-Leibler divergence: ``distance'' between probability distributions

$$D_{\text{KL}}(P||Q) = \int P(x) \log \left(\frac{P(x)}{Q(x)} \right) dx$$

Properties:

- KL divergence is non-negative: $D_{\text{KL}}(P||Q) \geq 0$
- KL divergence is 0 exactly if: $D_{\text{KL}}(P||Q) = 0 \Leftrightarrow P(x) = Q(x) \forall x$
- KL divergence is asymmetric: $D_{\text{KL}}(P||Q) \neq D_{\text{KL}}(Q||P)$

KL Divergence

Kullback-Leibler divergence: ``distance'' between probability distributions

$$D_{\text{KL}}(P||Q) = \int P(x) \log \left(\frac{P(x)}{Q(x)} \right) dx$$

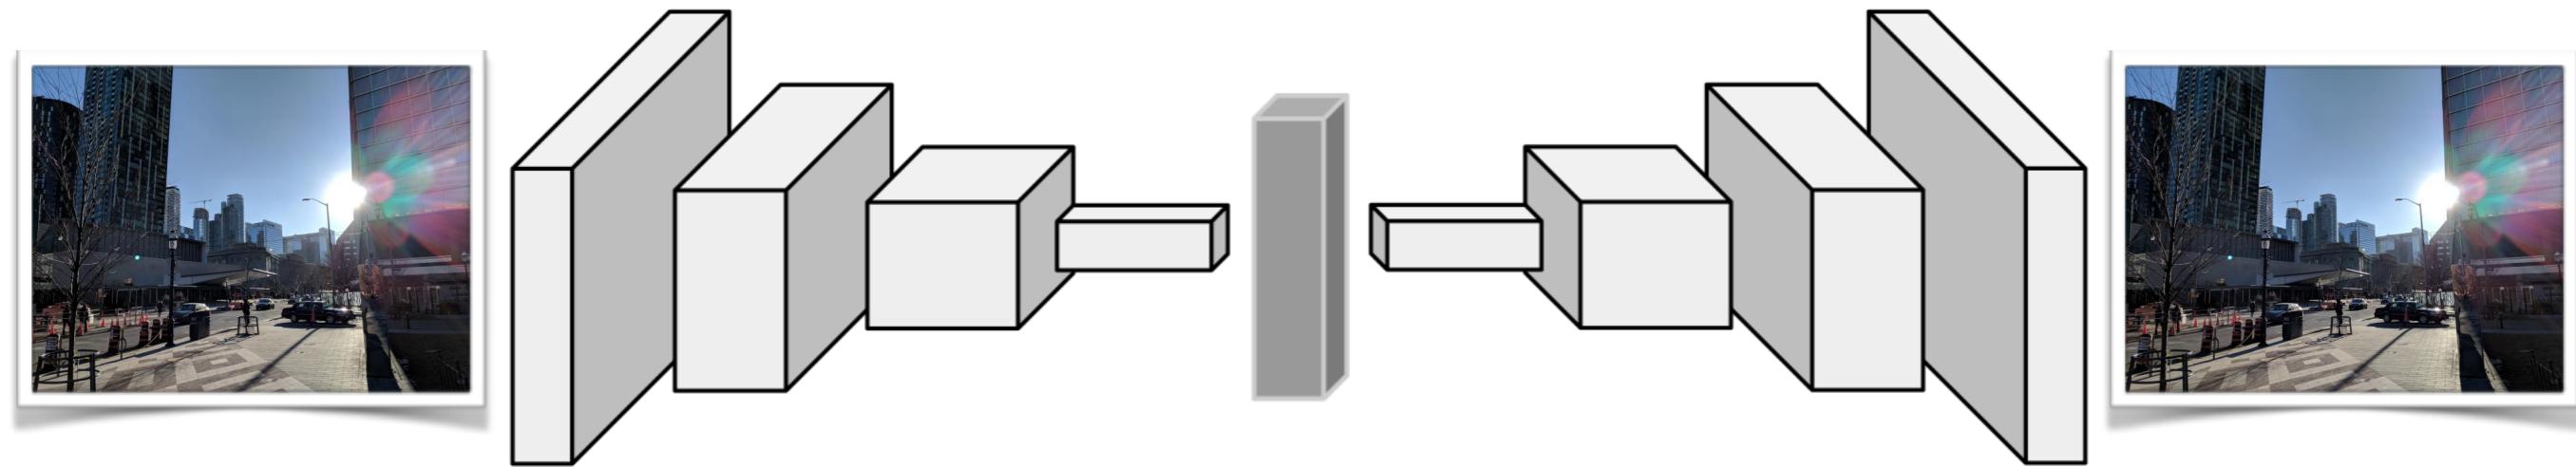
Equivalent to (when sampling):

$$D_{\text{KL}}(P||Q) = E_{x \sim P} [\log(P(x)) - \log(Q(x))]$$

Theory

latent space

$$\mathbf{z} \sim p_{\text{model}}$$



encoder

$$q(\mathbf{z}|\mathbf{x})$$

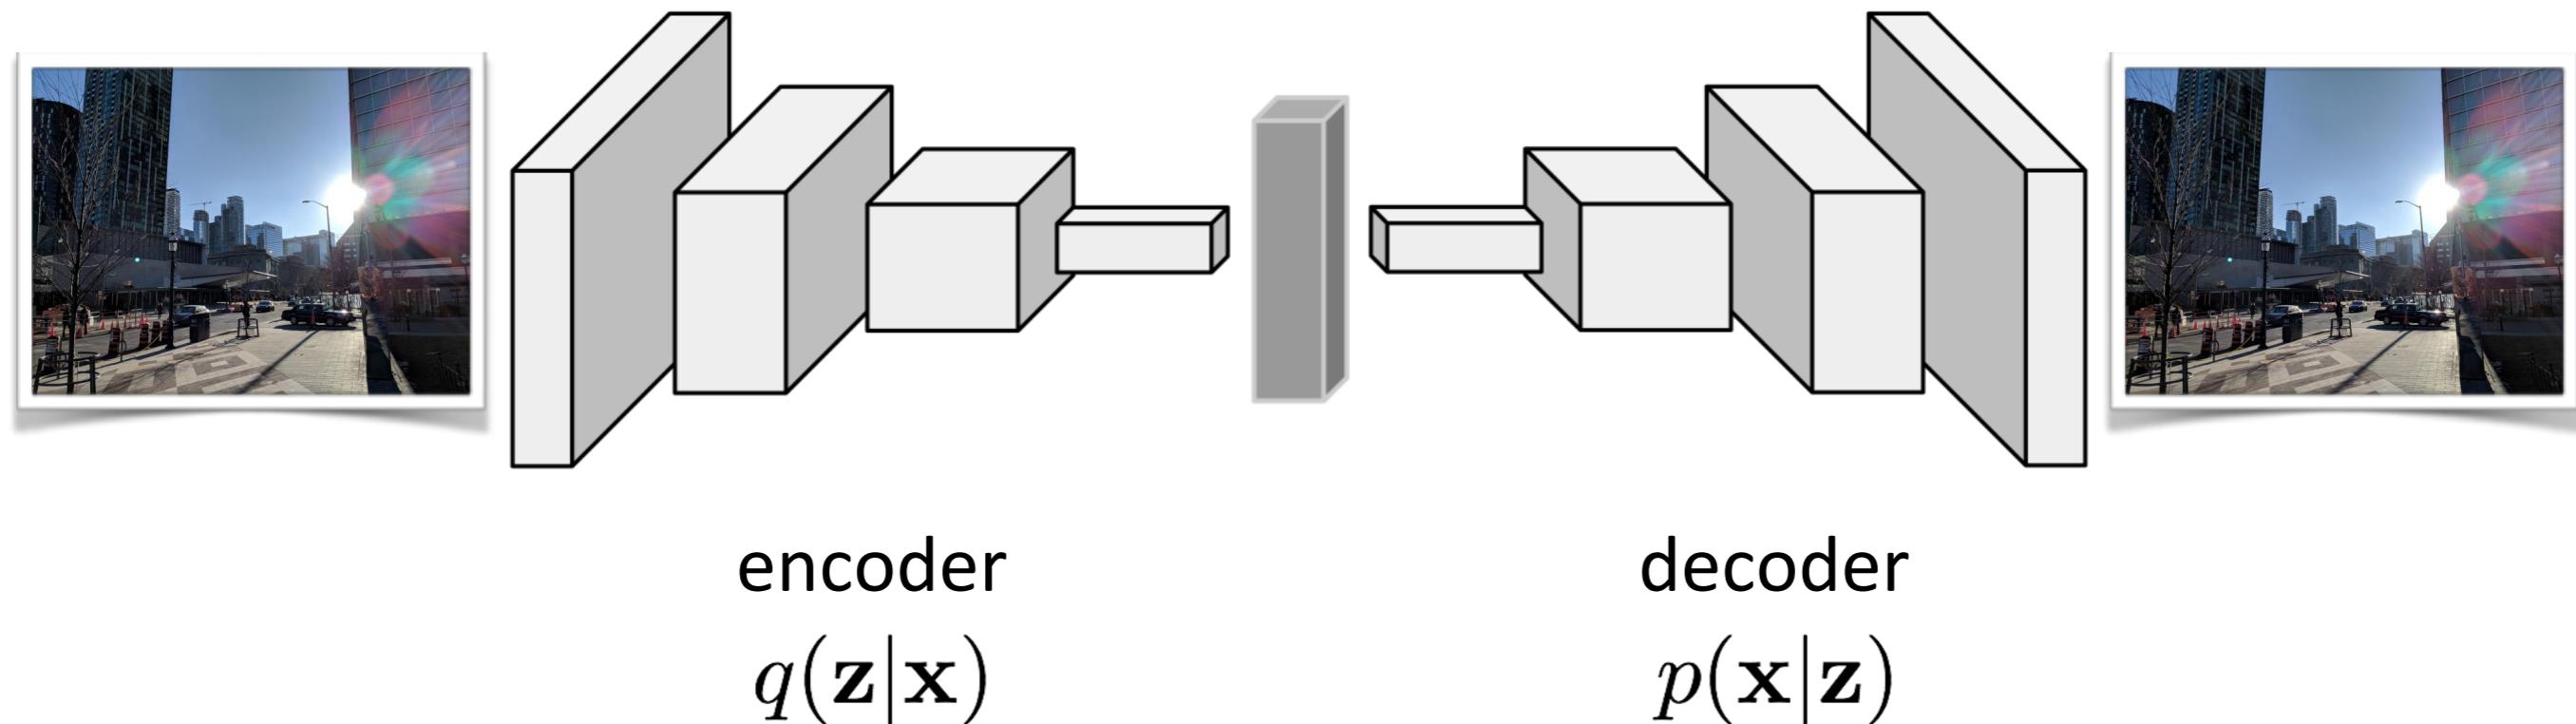
decoder

$$p(\mathbf{x}|\mathbf{z})$$

Theory

latent space

$$\mathbf{z} \sim p_{\text{model}}$$

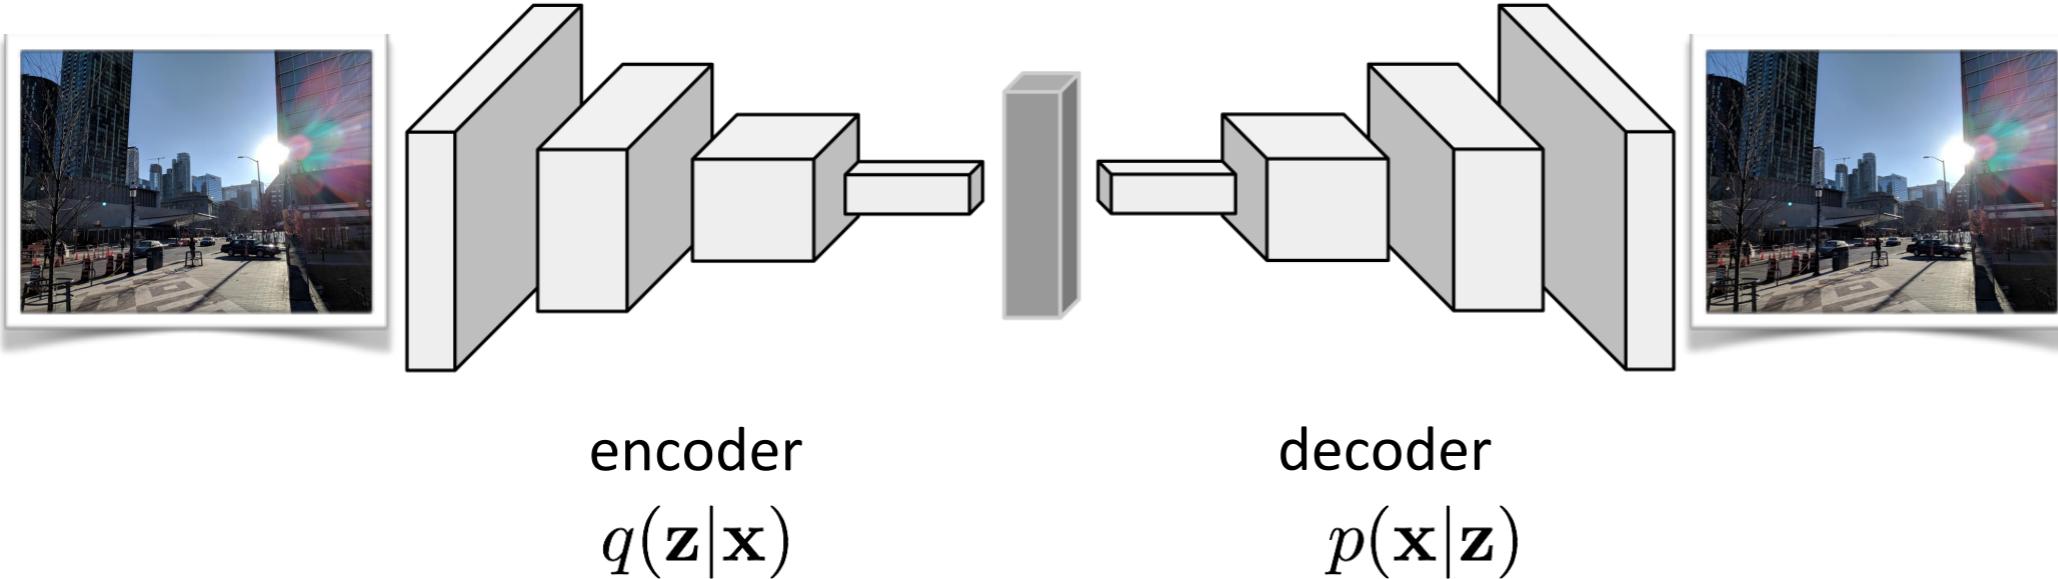


Want to ensure that $q(z|x)$ and $p(z|x)$ are the same:

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}|\mathbf{x}))]$$

Theory

$$\mathbf{z} \sim p_{\text{model}} \quad \text{latent space}$$

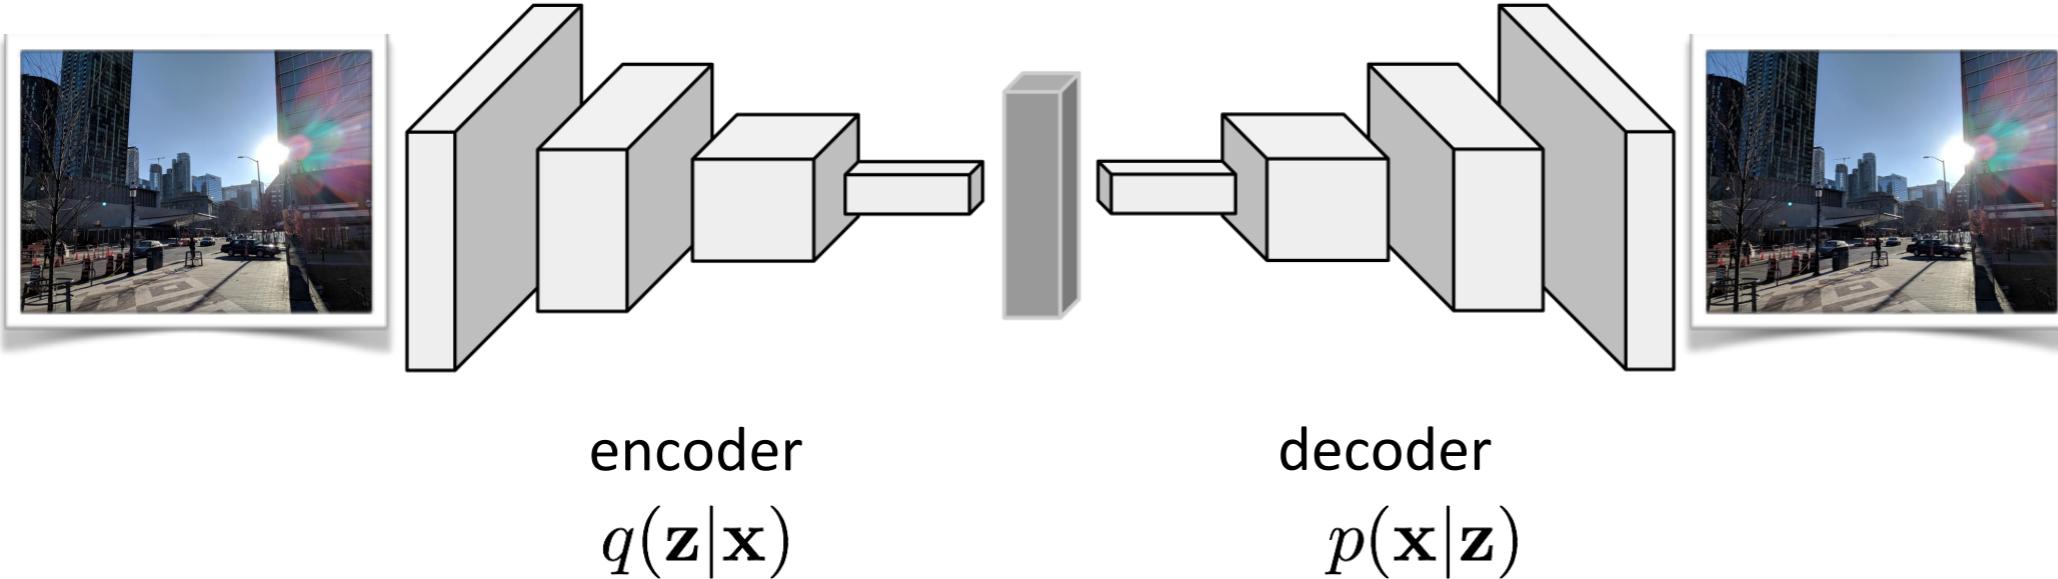


Use Bayes rule to include $p(\mathbf{x})$ and $p(\mathbf{x}|\mathbf{z})$:

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}|\mathbf{x}))]$$

Theory

$$\mathbf{z} \sim p_{\text{model}} \quad \text{latent space}$$

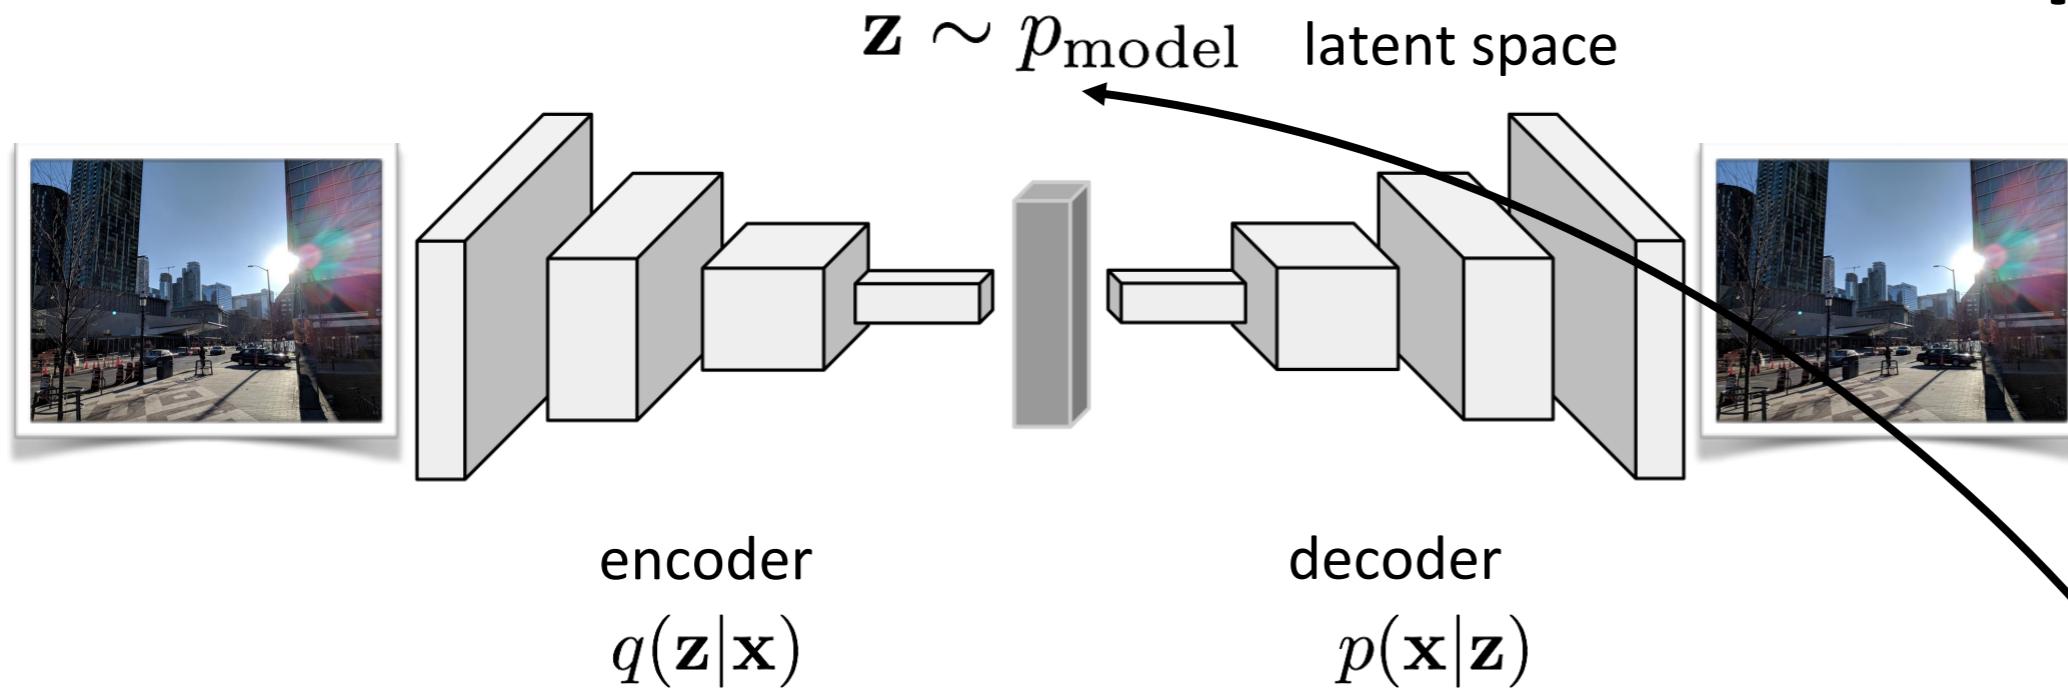


Use Bayes rule to include $p(\mathbf{x})$ and $p(\mathbf{x} | \mathbf{z})$:

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}|\mathbf{x}))]$$

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log\left(\frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}\right)]$$

Theory



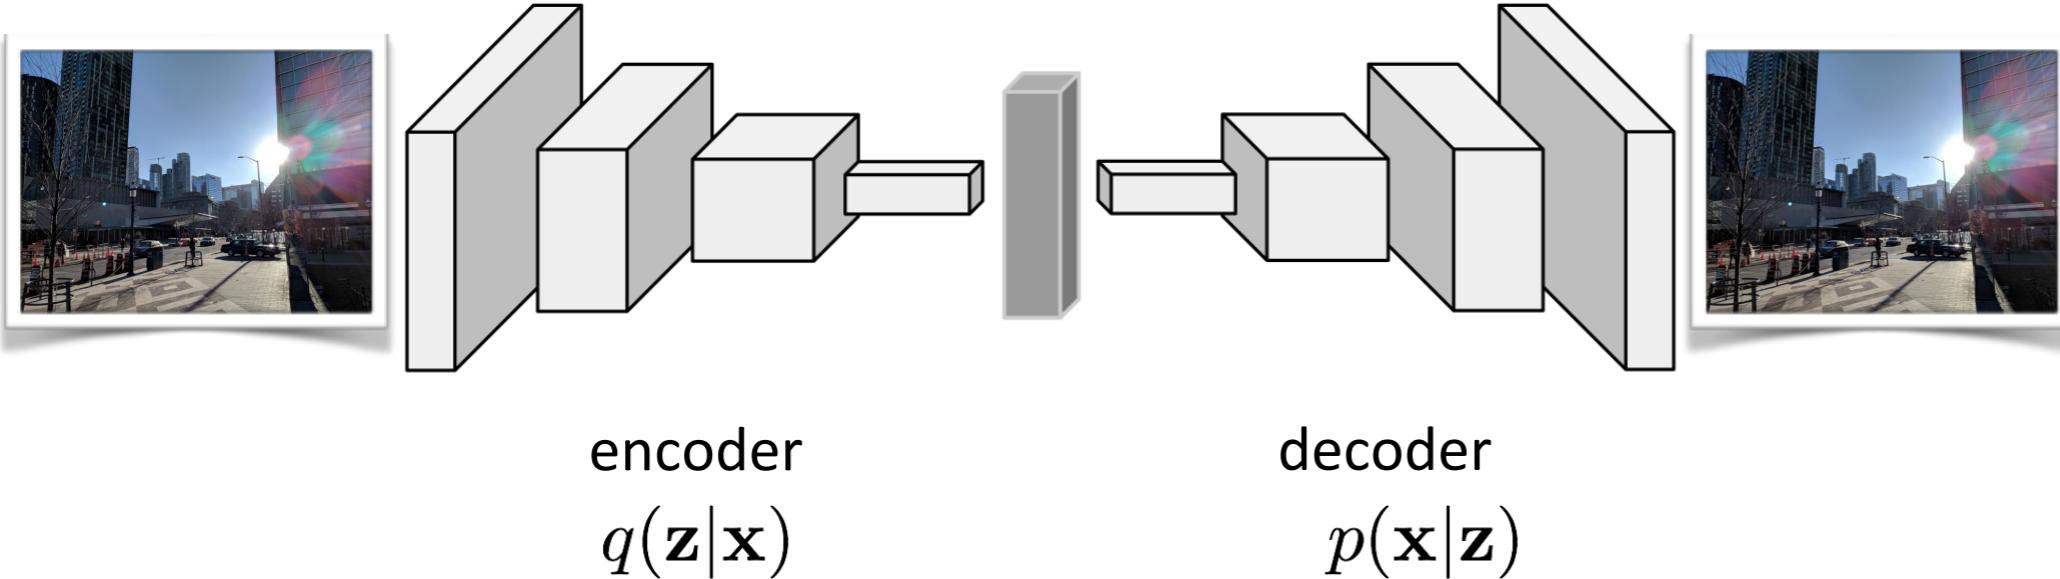
Use Bayes rule to include $p(\mathbf{x})$ and $p(\mathbf{x}|\mathbf{z})$:

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}|\mathbf{x}))]$$

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log\left(\frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}\right)]$$

Theory

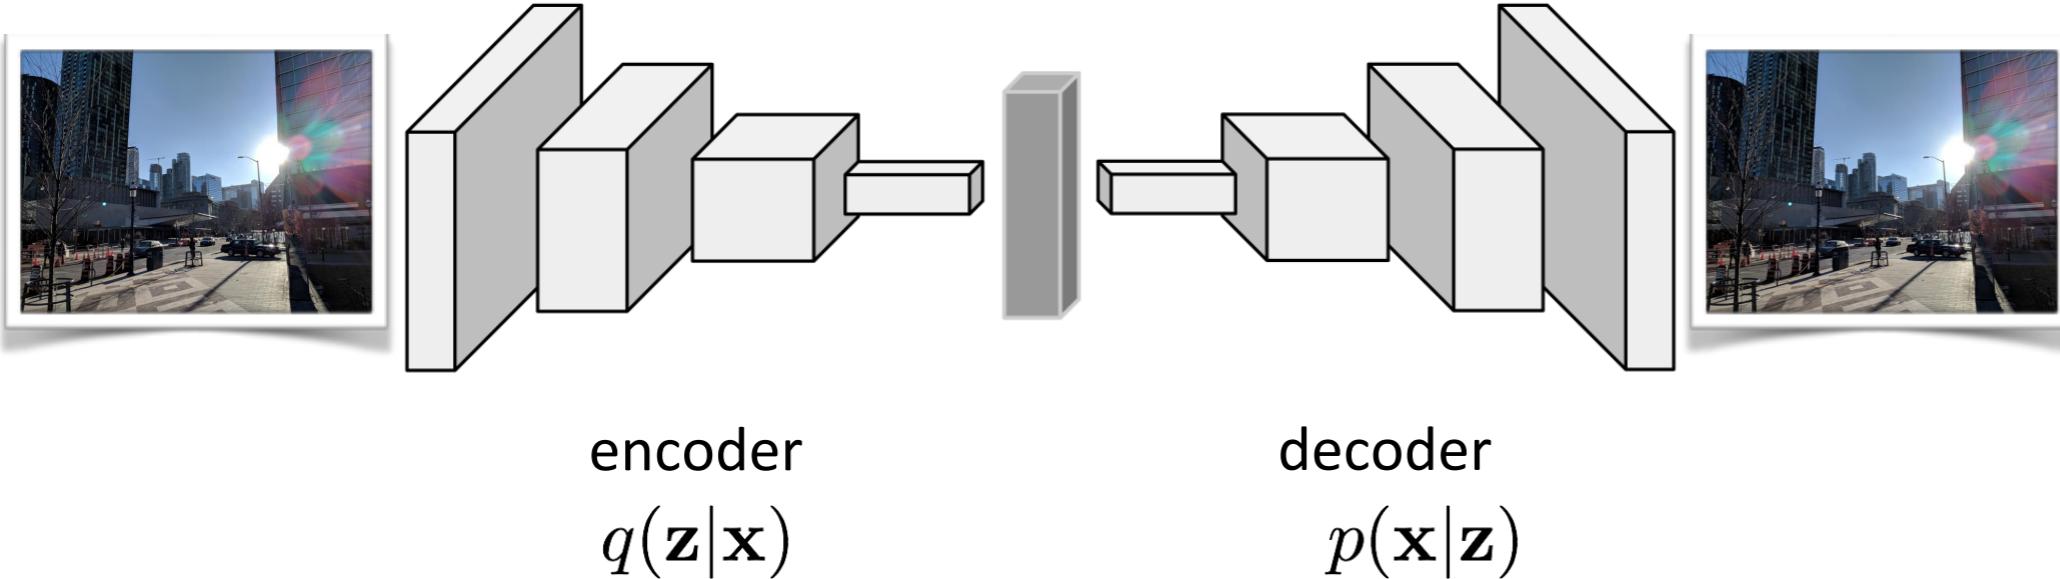
$\mathbf{z} \sim p_{\text{model}}$ latent space



$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log\left(\frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}\right)]$$

Theory

$\mathbf{z} \sim p_{\text{model}}$ latent space

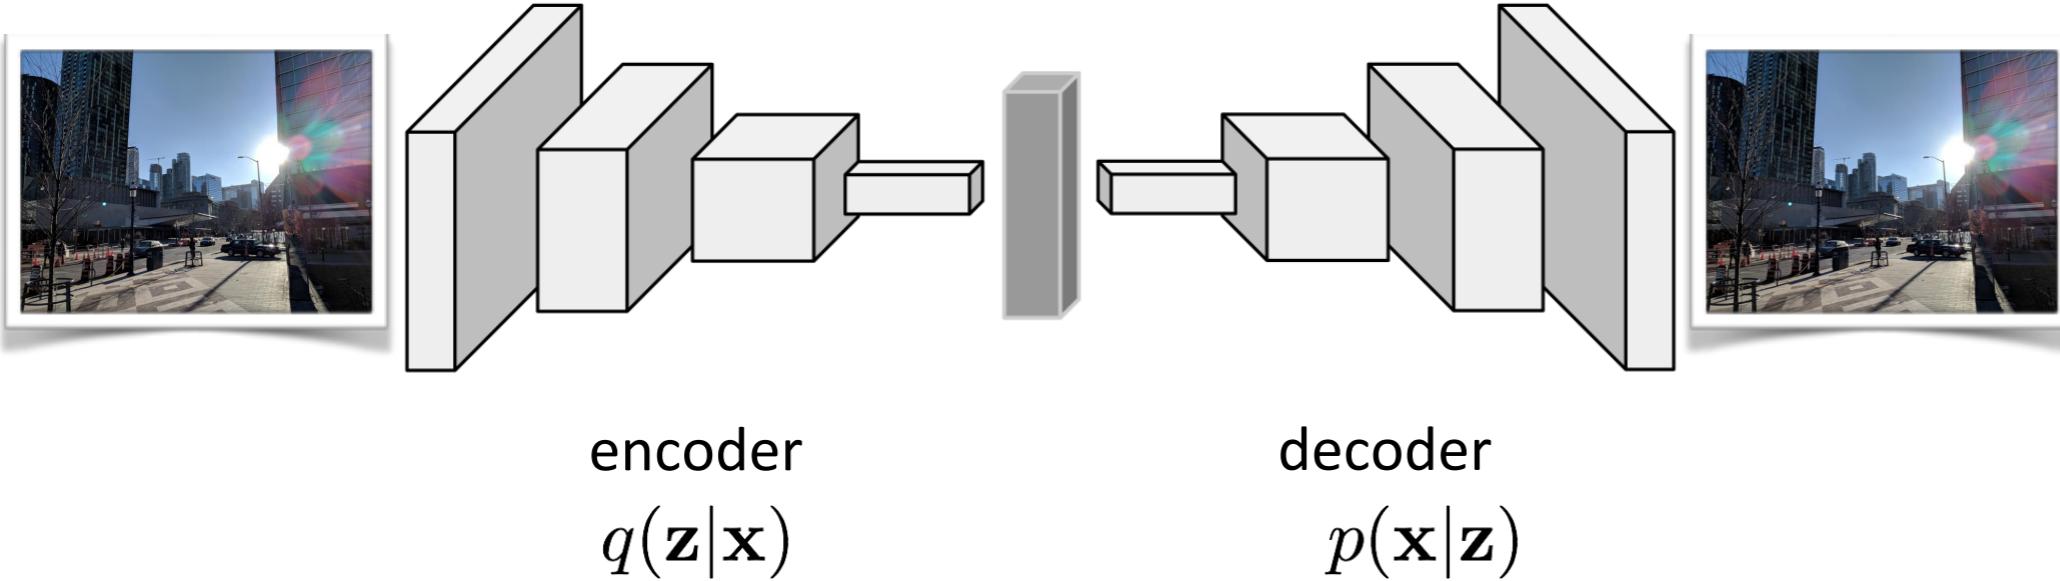


$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log\left(\frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}\right)]$$

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{x}|\mathbf{z})) - \log(p(\mathbf{z})) + \log(p(\mathbf{x}))]$$

Theory

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



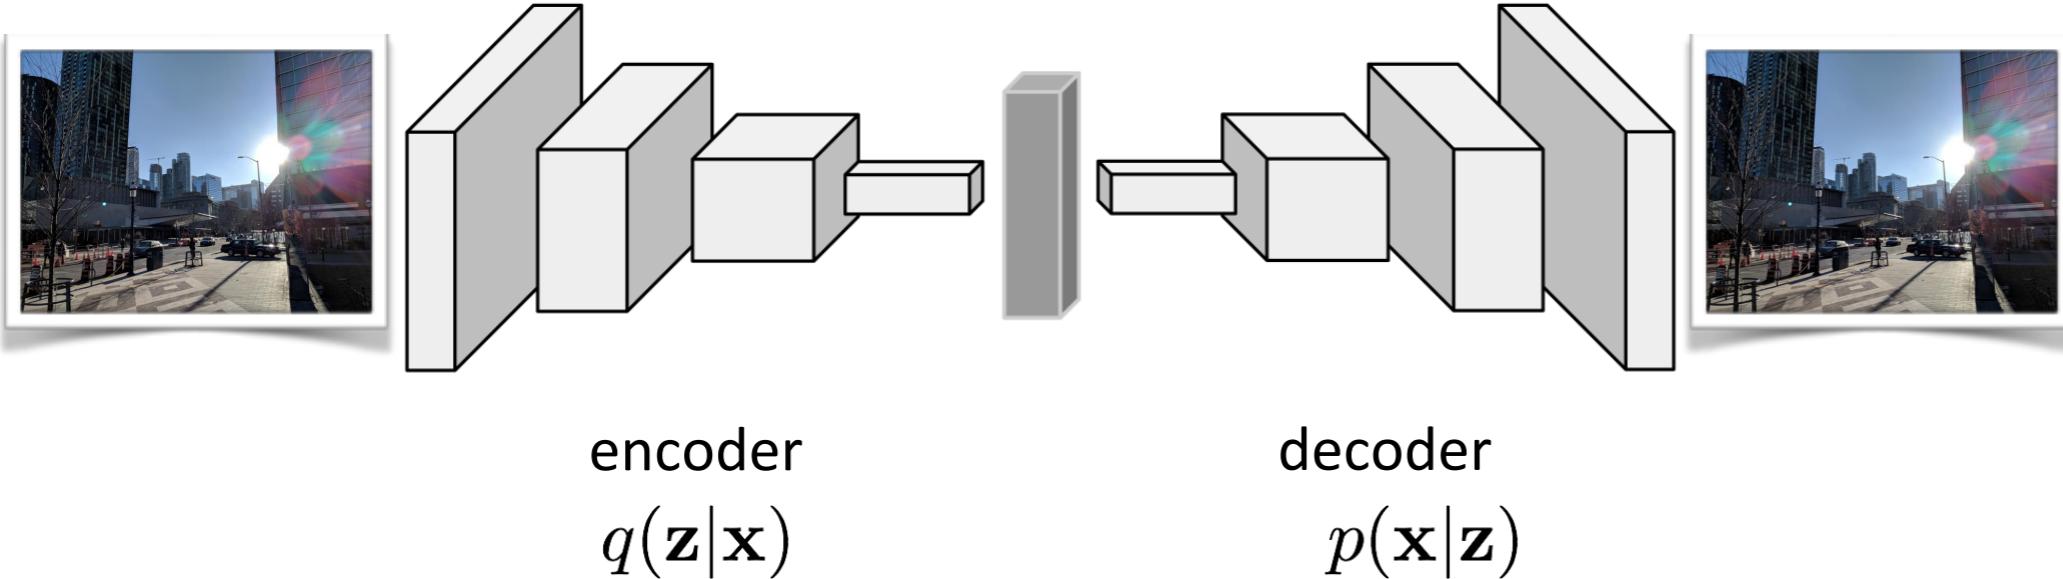
$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log\left(\frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}\right)]$$

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{x}|\mathbf{z})) - \log(p(\mathbf{z})) + \log(p(\mathbf{x}))]$$

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

Theory

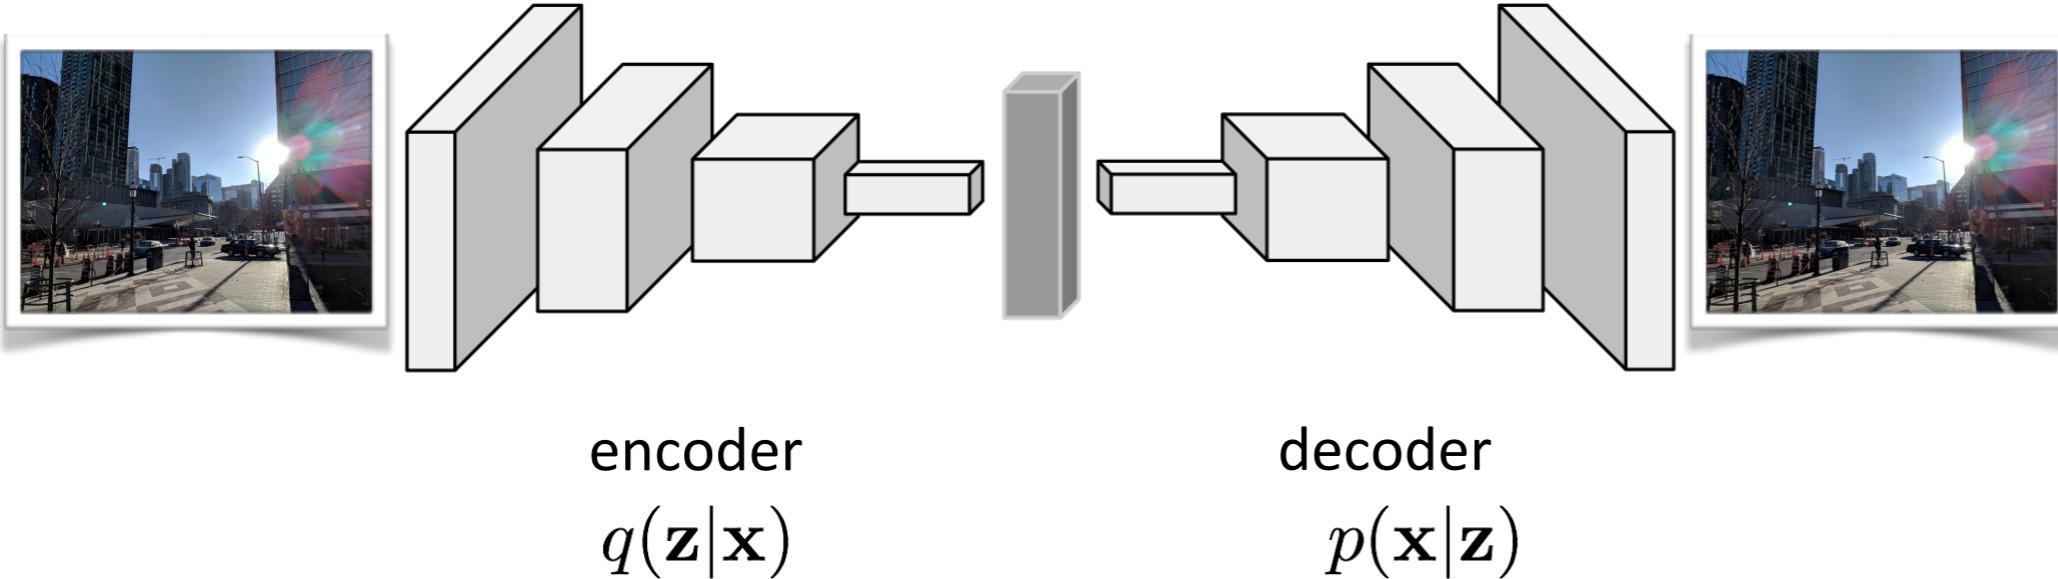
$\mathbf{z} \sim p_{\text{model}}$ latent space



$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

Theory

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



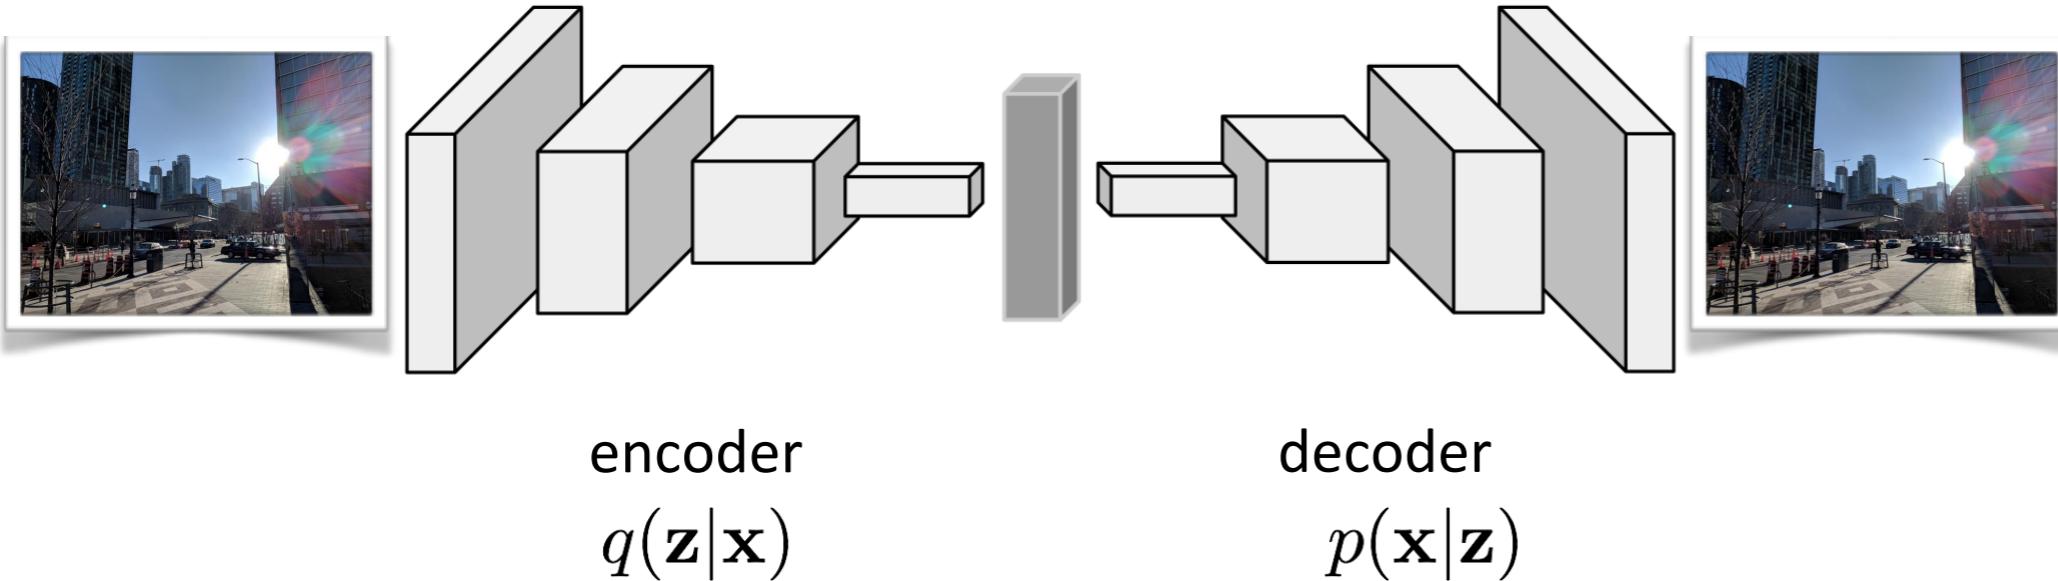
$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

Rearrange:

$$\log(p(\mathbf{x})) = D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) + E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

Theory

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



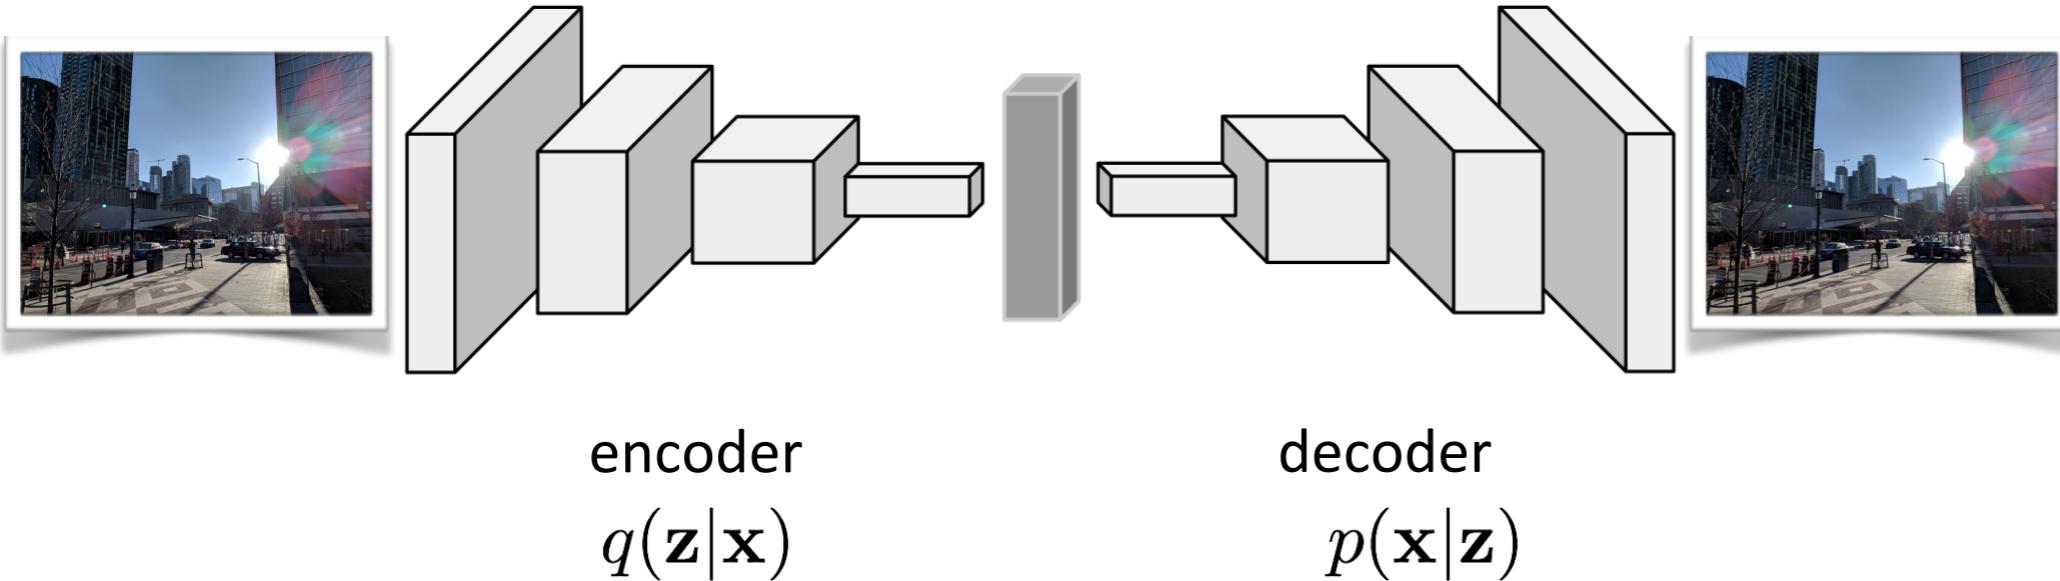
$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

Rearrange:

$$\log(p(\mathbf{x})) = D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) + E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

Theory

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



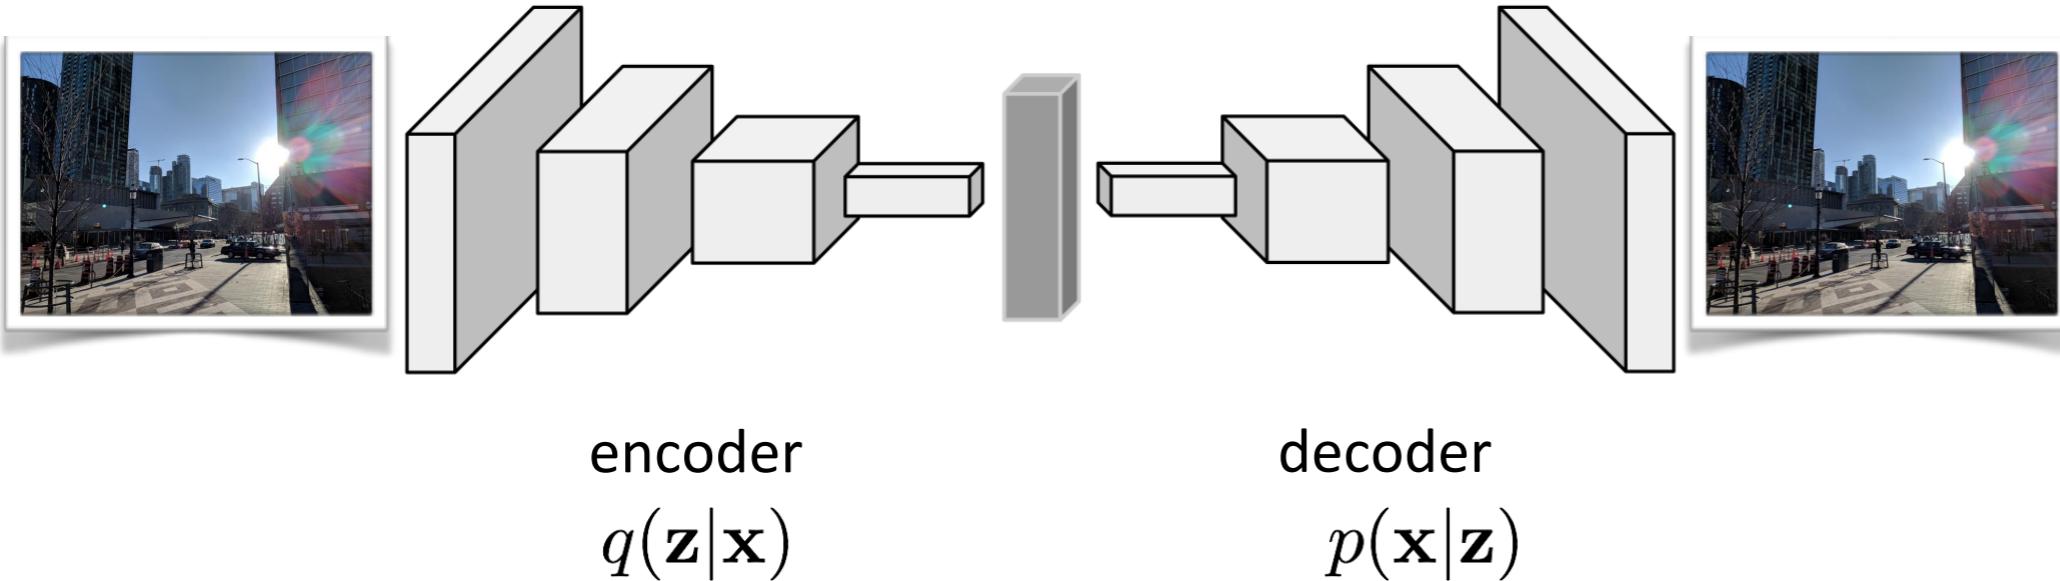
$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

Rearrange:

$$\log(p(\mathbf{x})) = D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) + \underbrace{E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))]}_{\text{decoder}} - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

Theory

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

Rearrange:

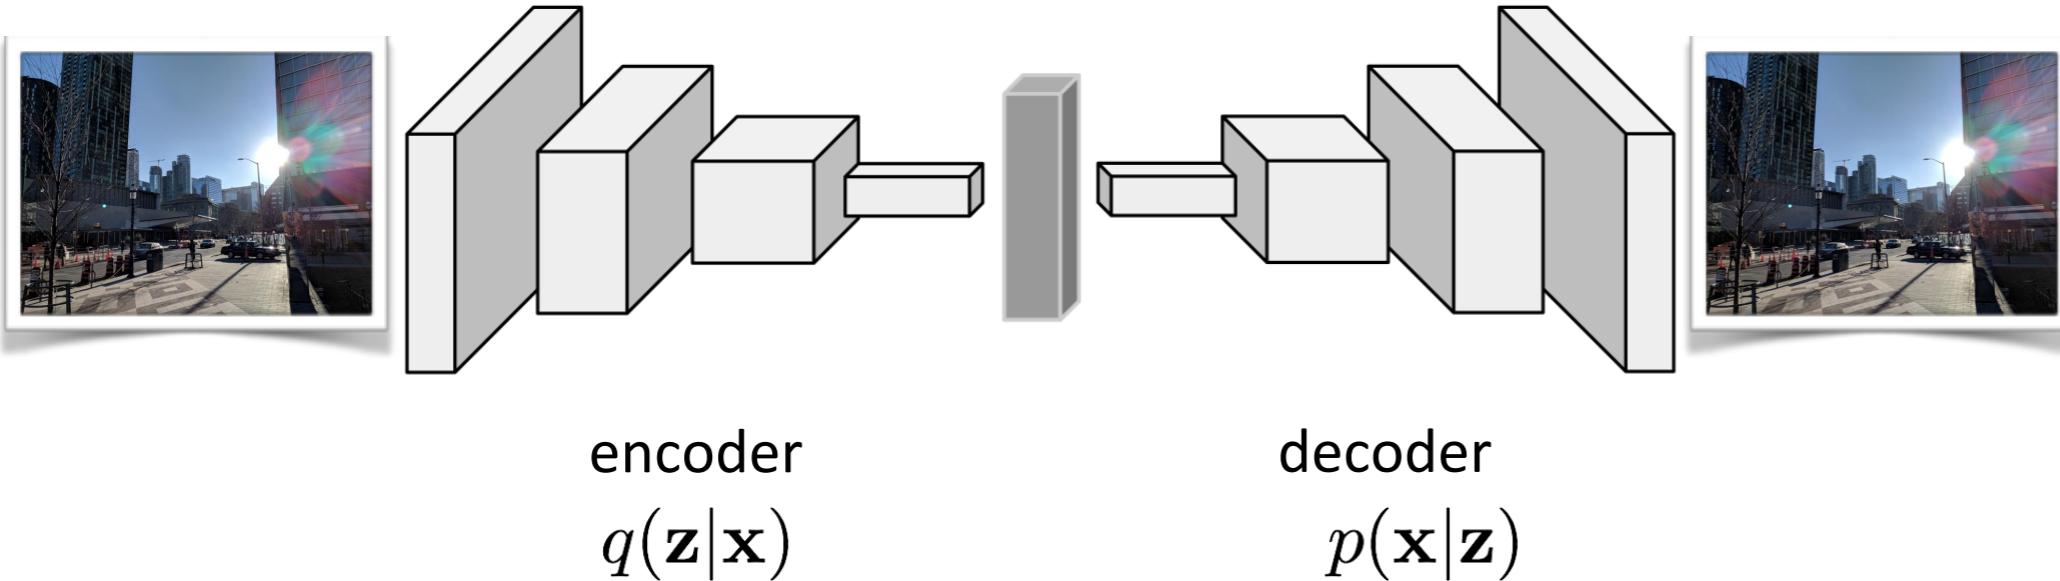
$$\log(p(\mathbf{x})) = D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) + E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

how to compute?

decoder

Theory

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

Rearrange:

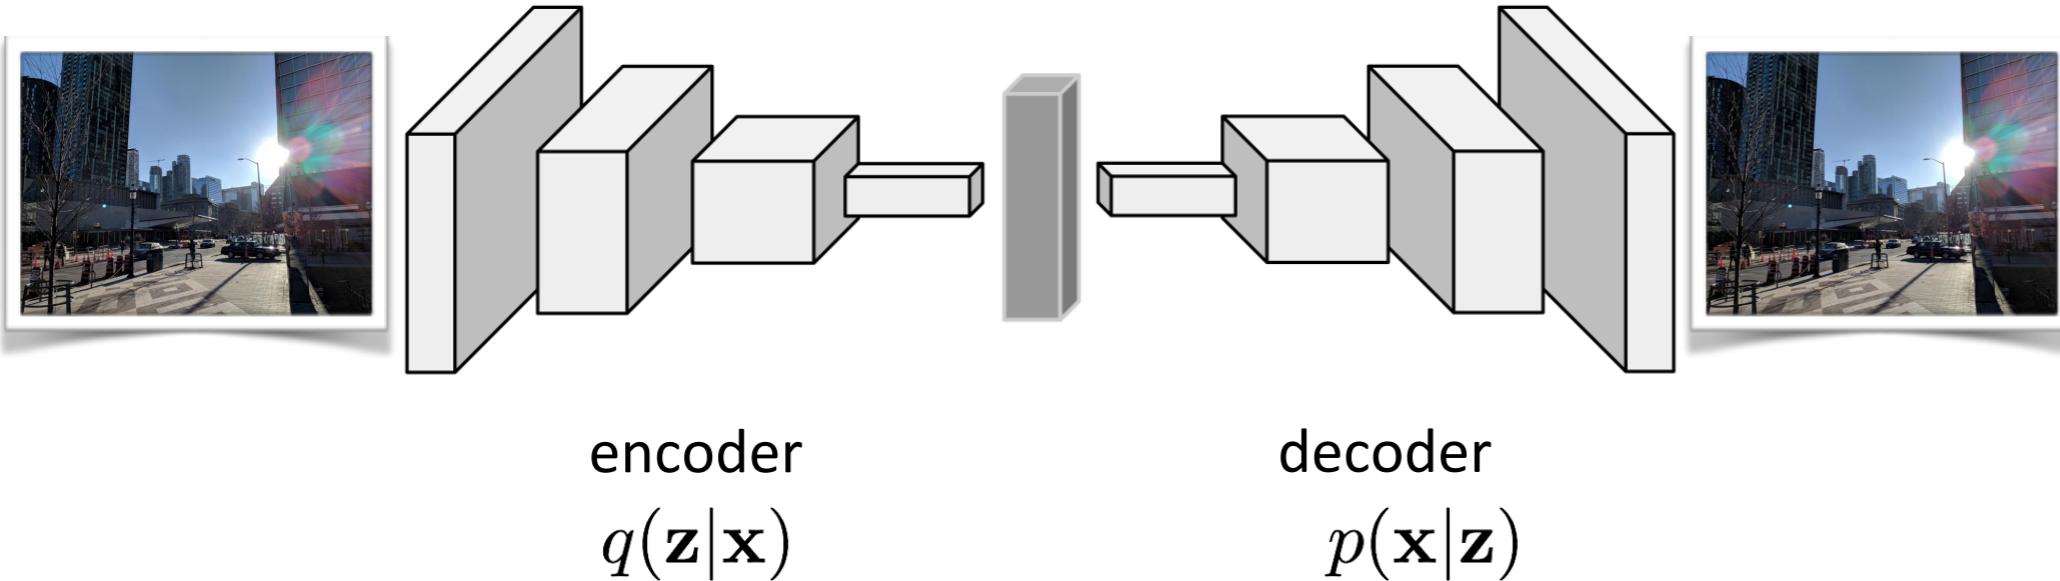
$$\log(p(\mathbf{x})) = D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) + E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

how to compute?
no idea!

decoder

Theory

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

Rearrange:

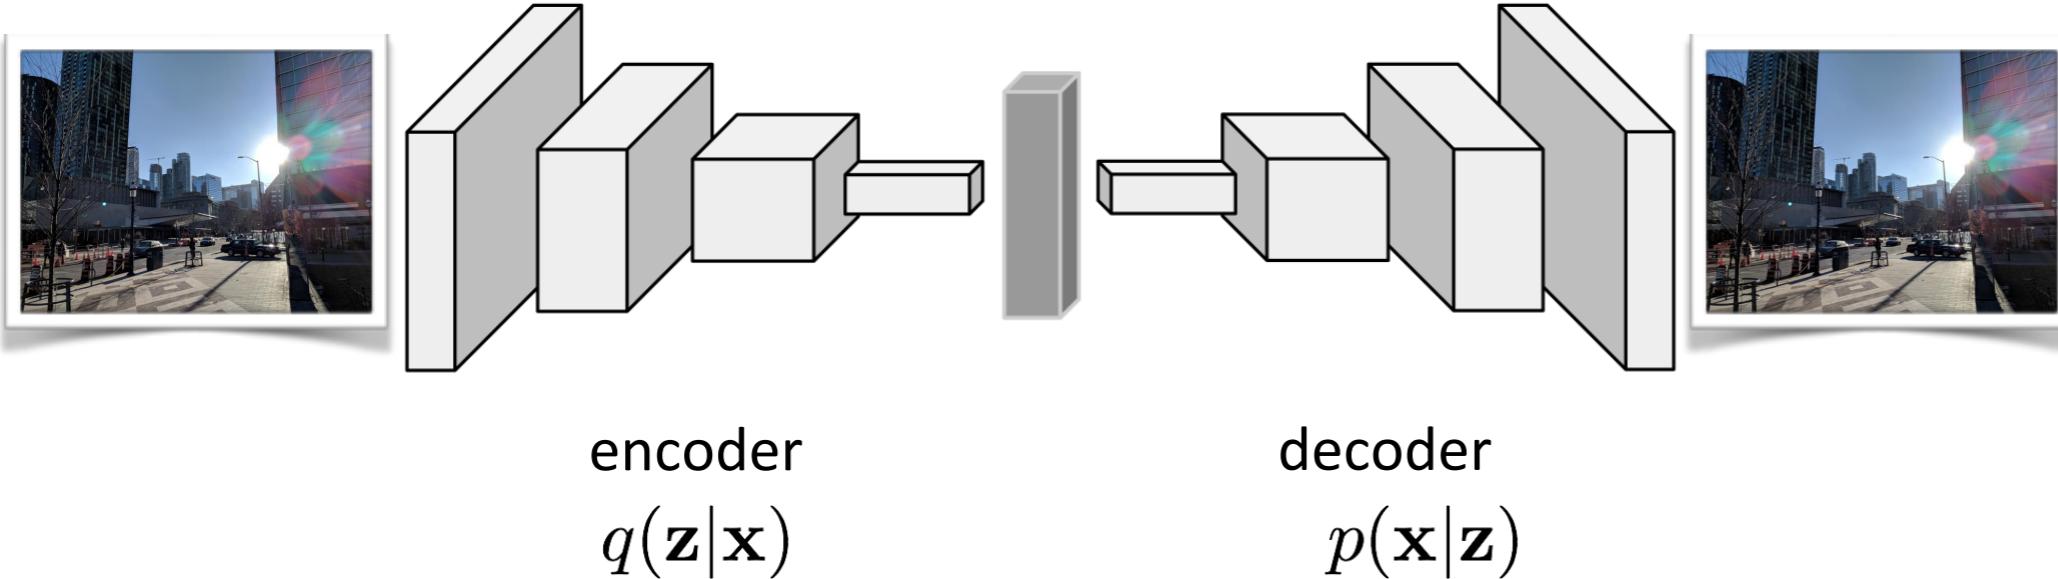
$$\log(p(\mathbf{x})) = D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) + E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

how to compute?
no idea! But ≥ 0

decoder

Theory

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

Rearrange:

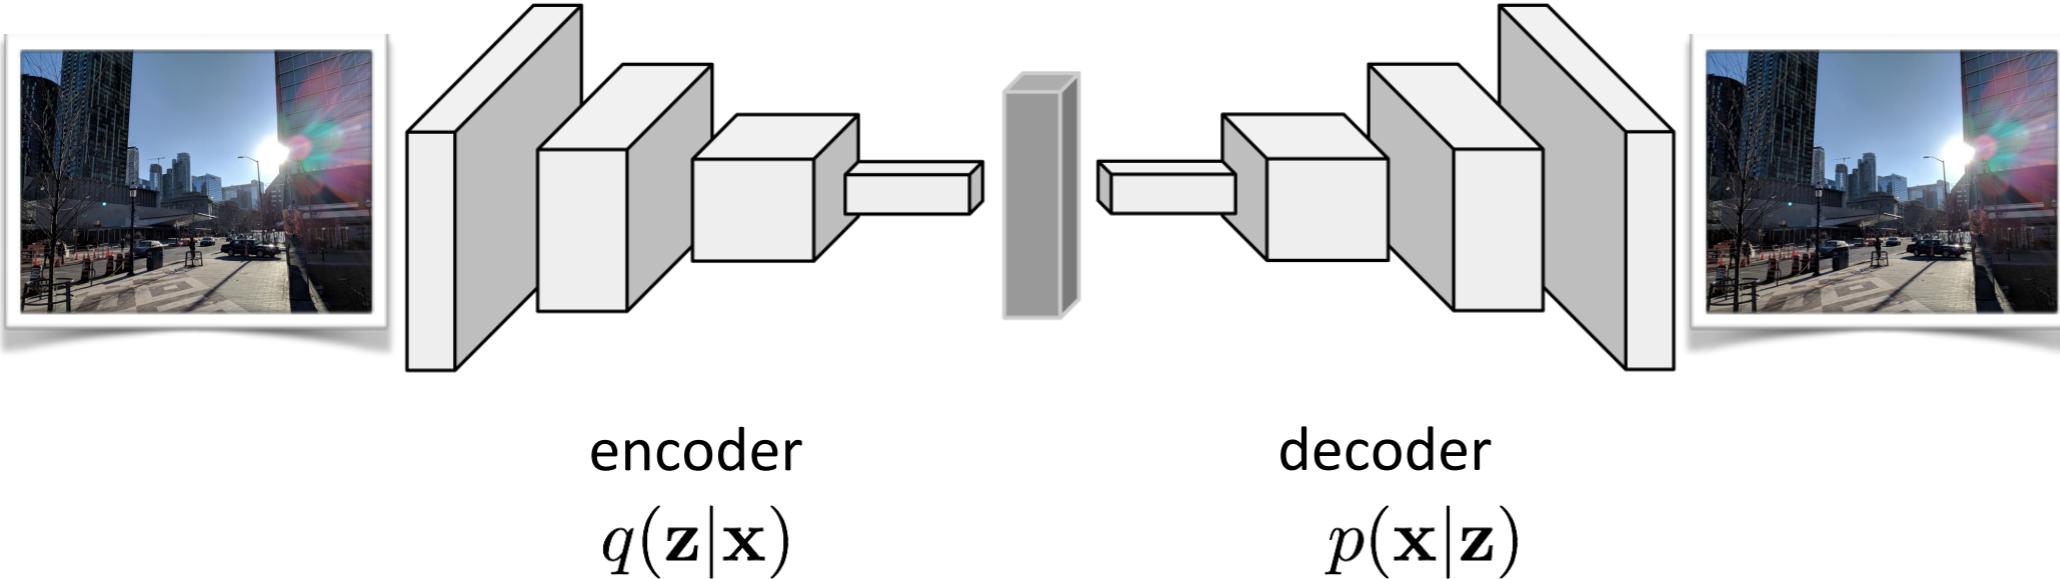
$$\log(p(\mathbf{x})) = \underbrace{D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x}))}_{\text{how to compute?}} + \underbrace{E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))]}_{\text{decoder}} - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

no idea! But ≥ 0

$$\geq E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

Theory

$$\mathbf{z} \sim p_{\text{model}} \quad \text{latent space}$$

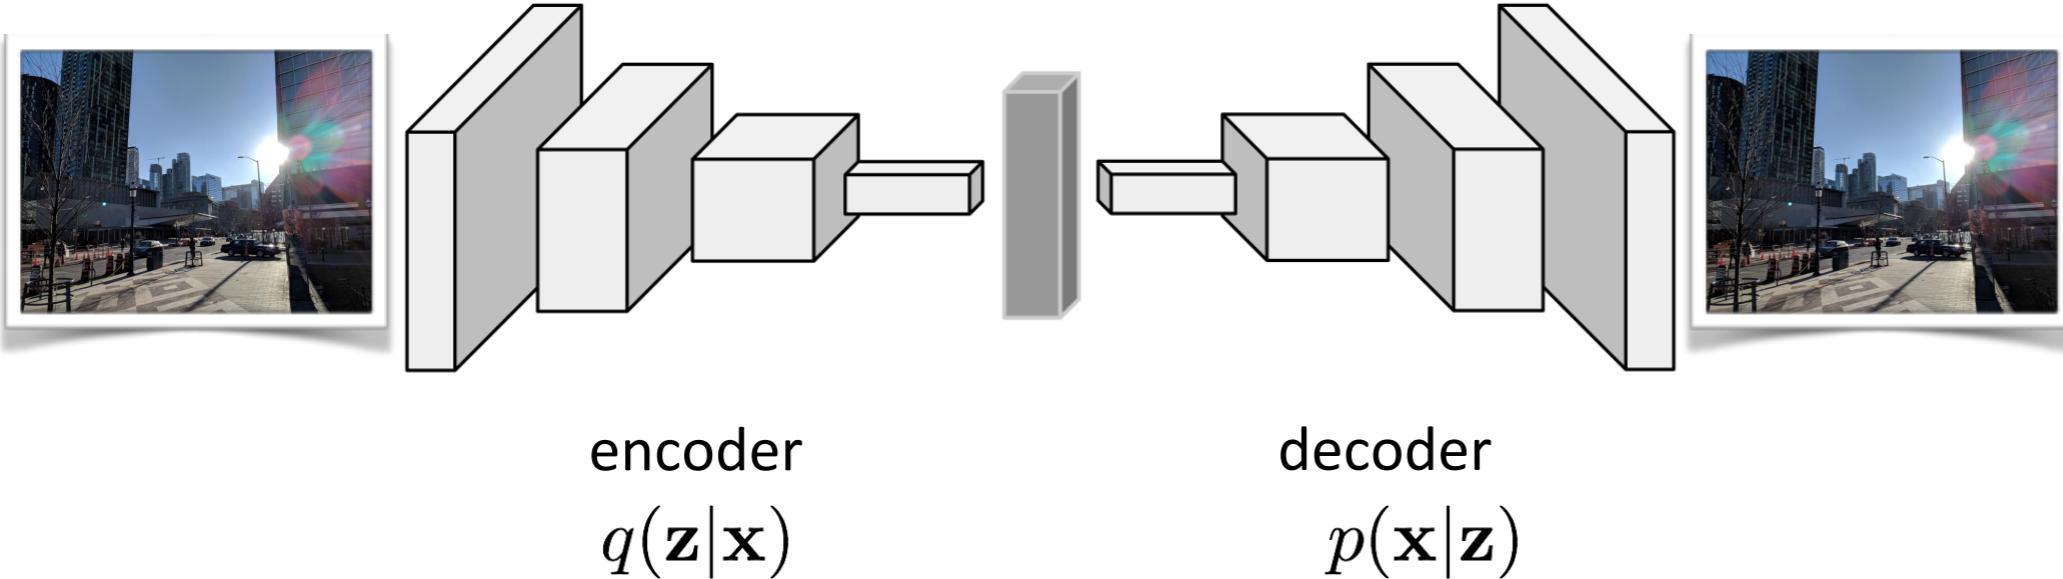


Evidence Lower BOund (ELBO):

$$\log(p(\mathbf{x})) \geq E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

Theory

$\mathbf{z} \sim p_{\text{model}}$ latent space



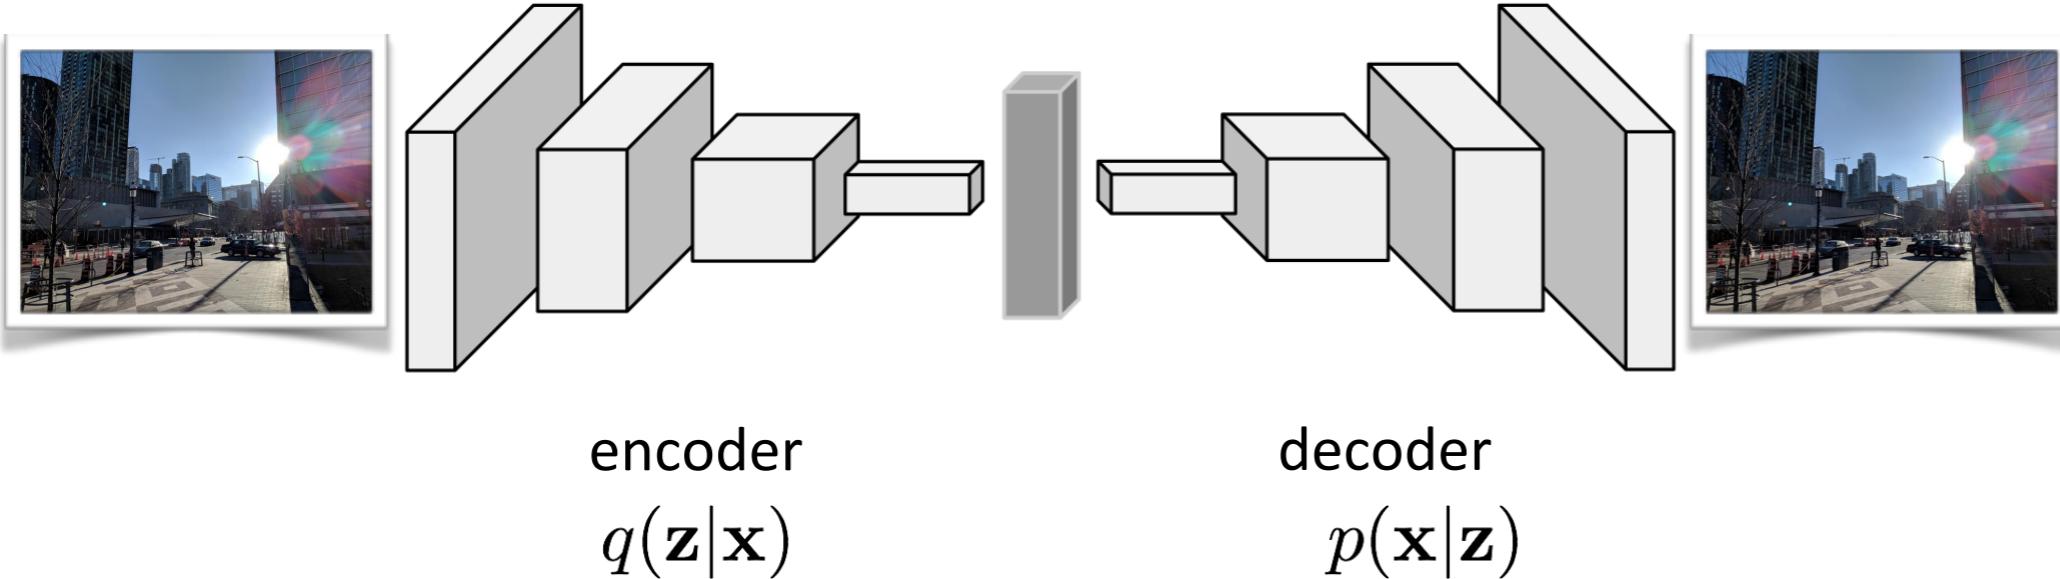
Evidence Lower BOund (ELBO):

$$\log(p(\mathbf{x})) \geq E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

likelihood

Theory

$\mathbf{z} \sim p_{\text{model}}$ latent space



Evidence Lower BOund (ELBO):

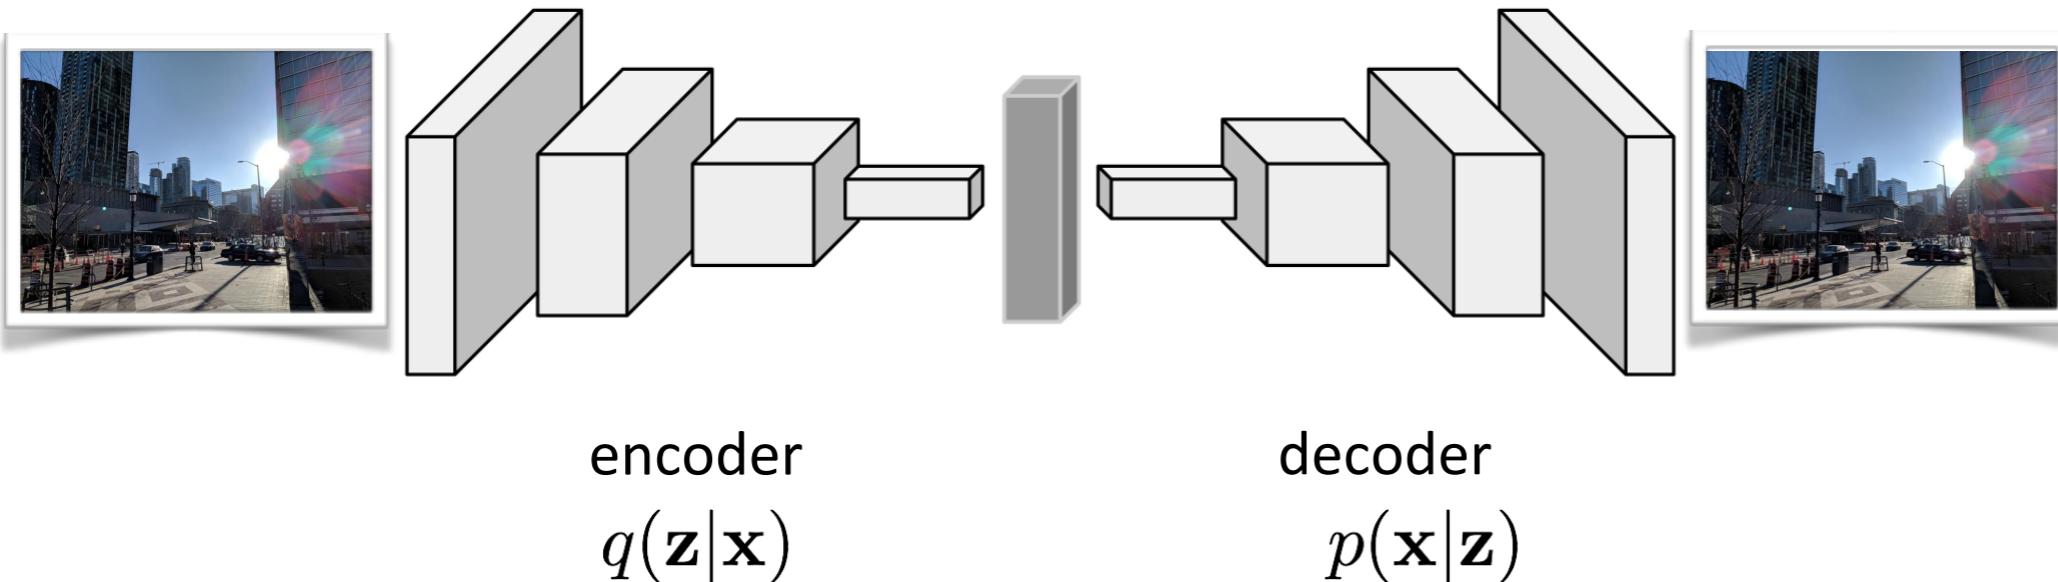
$$\log(p(\mathbf{x})) \geq E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

likelihood

KL divergence

Objective Function

$$\mathbf{z} \sim p_{\text{model}}$$

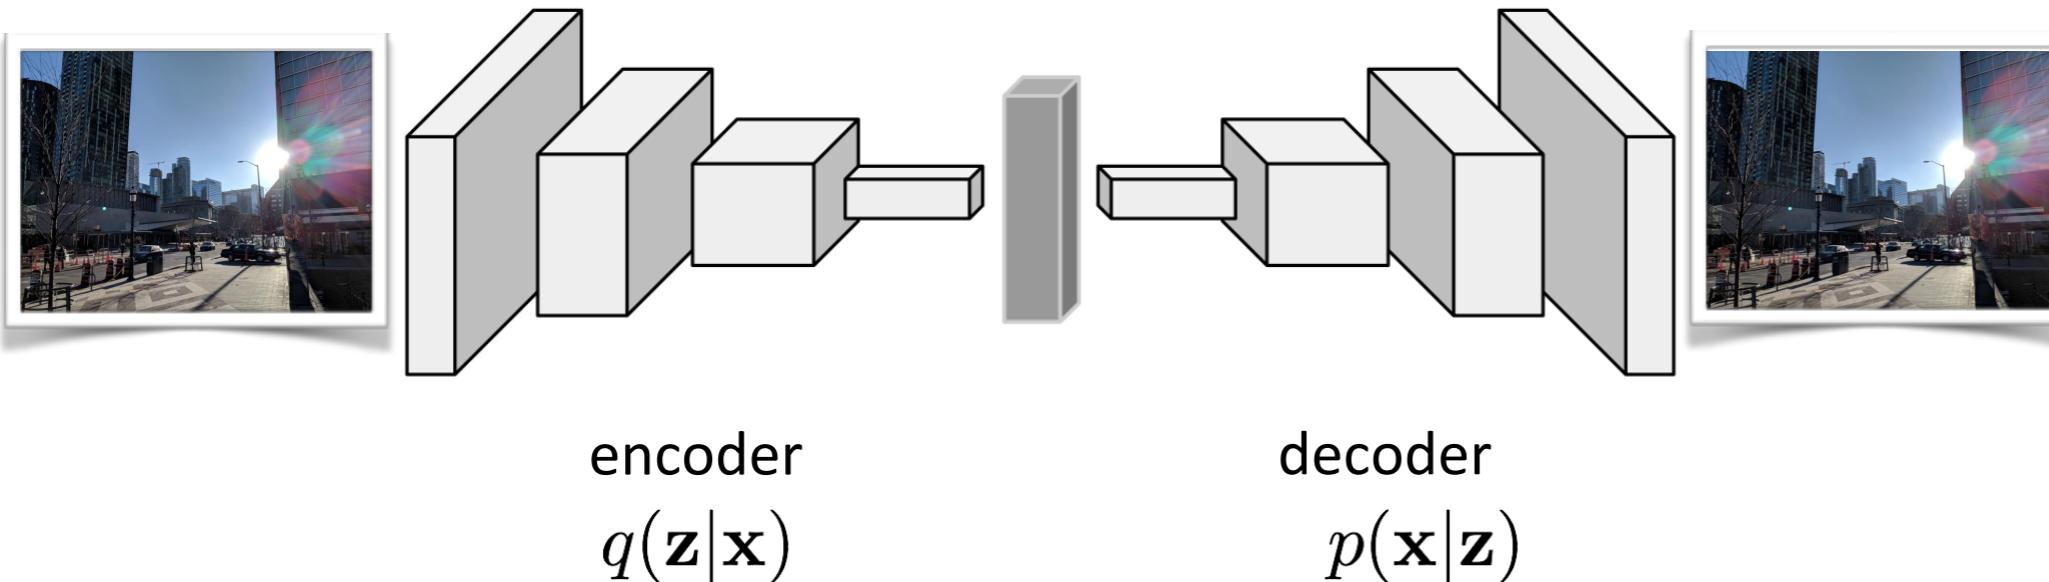


Maximize

$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}))$$

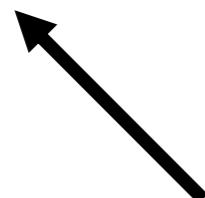
Objective Function

$$\mathbf{z} \sim p_{\text{model}}$$



Maximize

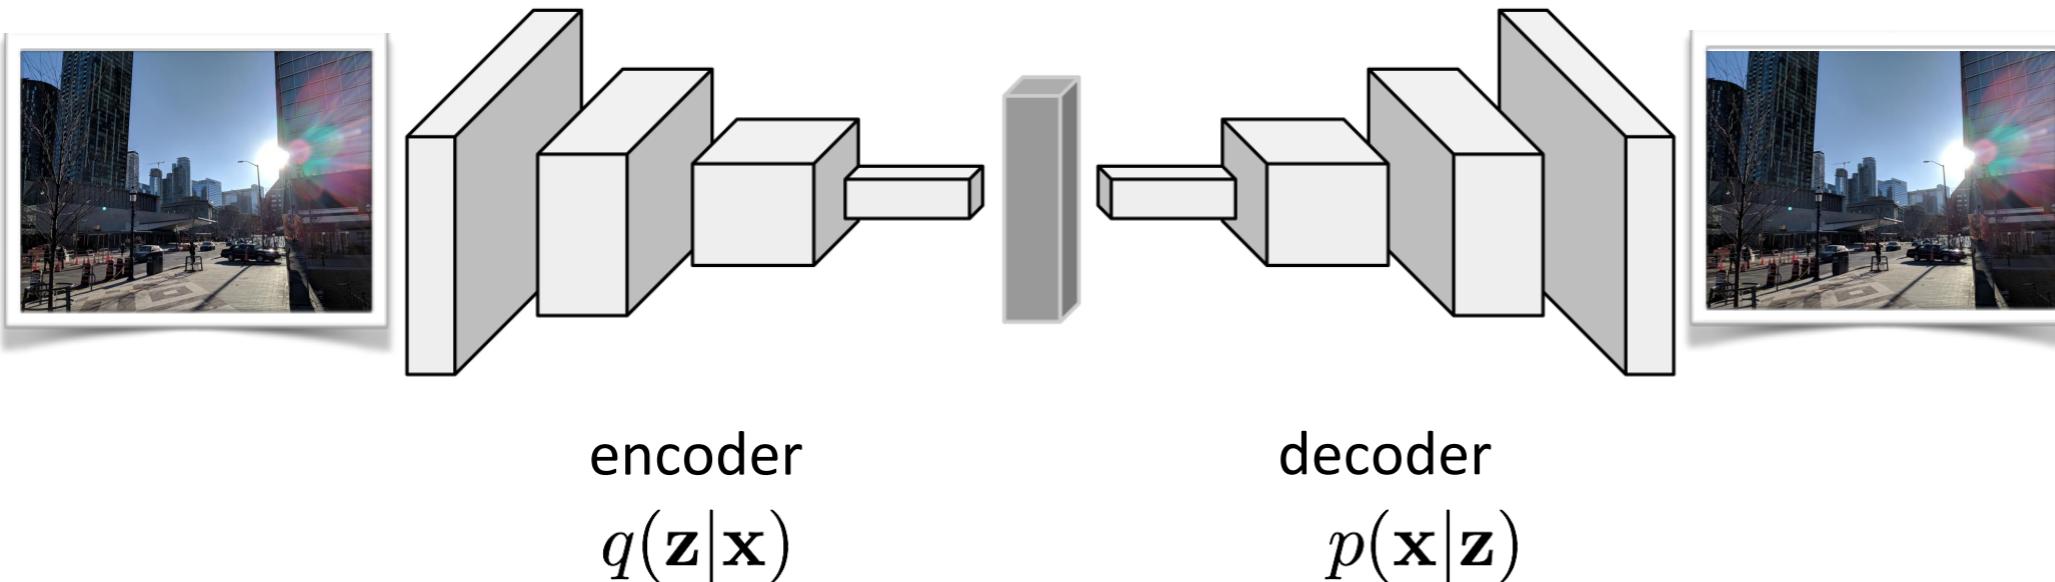
$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}))$$



reconstruction “loss”,
e.g., L2 norm (normal distribution)

Objective Function

$$\mathbf{z} \sim p_{\text{model}}$$



Maximize

$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}))$$

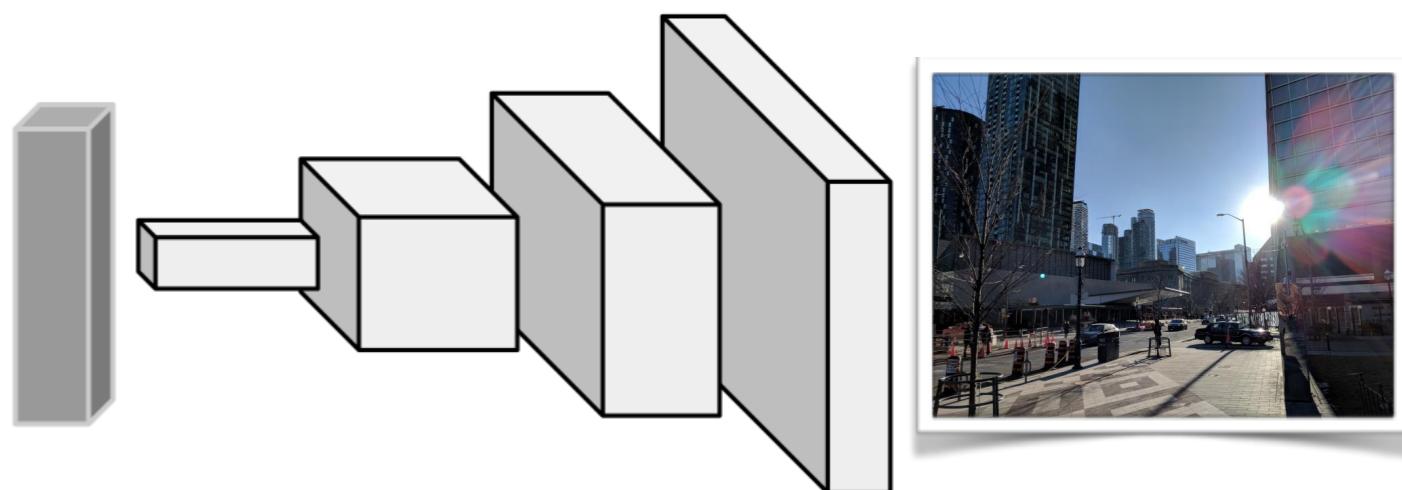
reconstruction “loss”,
e.g., L2 norm (normal distribution)

similarity between learned
and sampling distribution

At Test Time

sample point from latent space

$$\mathbf{z} \sim p_{\text{model}}$$



decoder
(transp. convolutions)

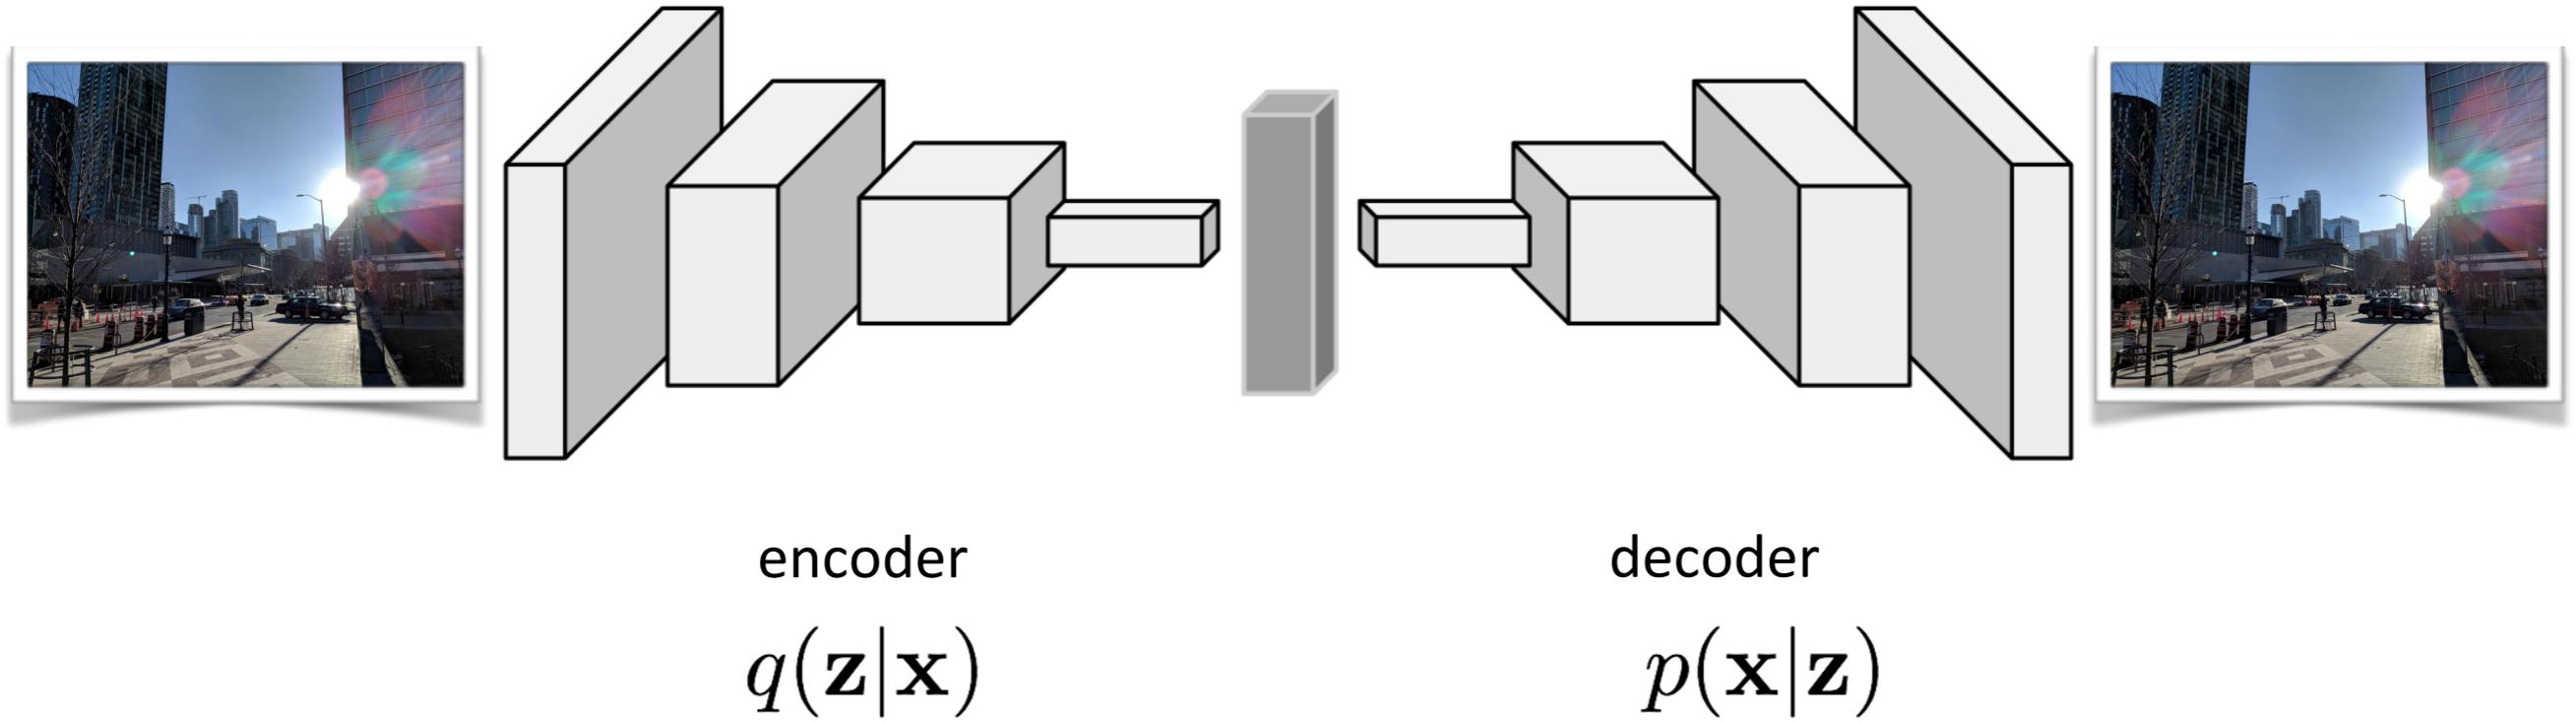
$$p(\mathbf{x}|\mathbf{z})$$

$$\mathbf{x} \sim p_{\text{data}}$$

use decoder to “translate”
to image

VAE - Example

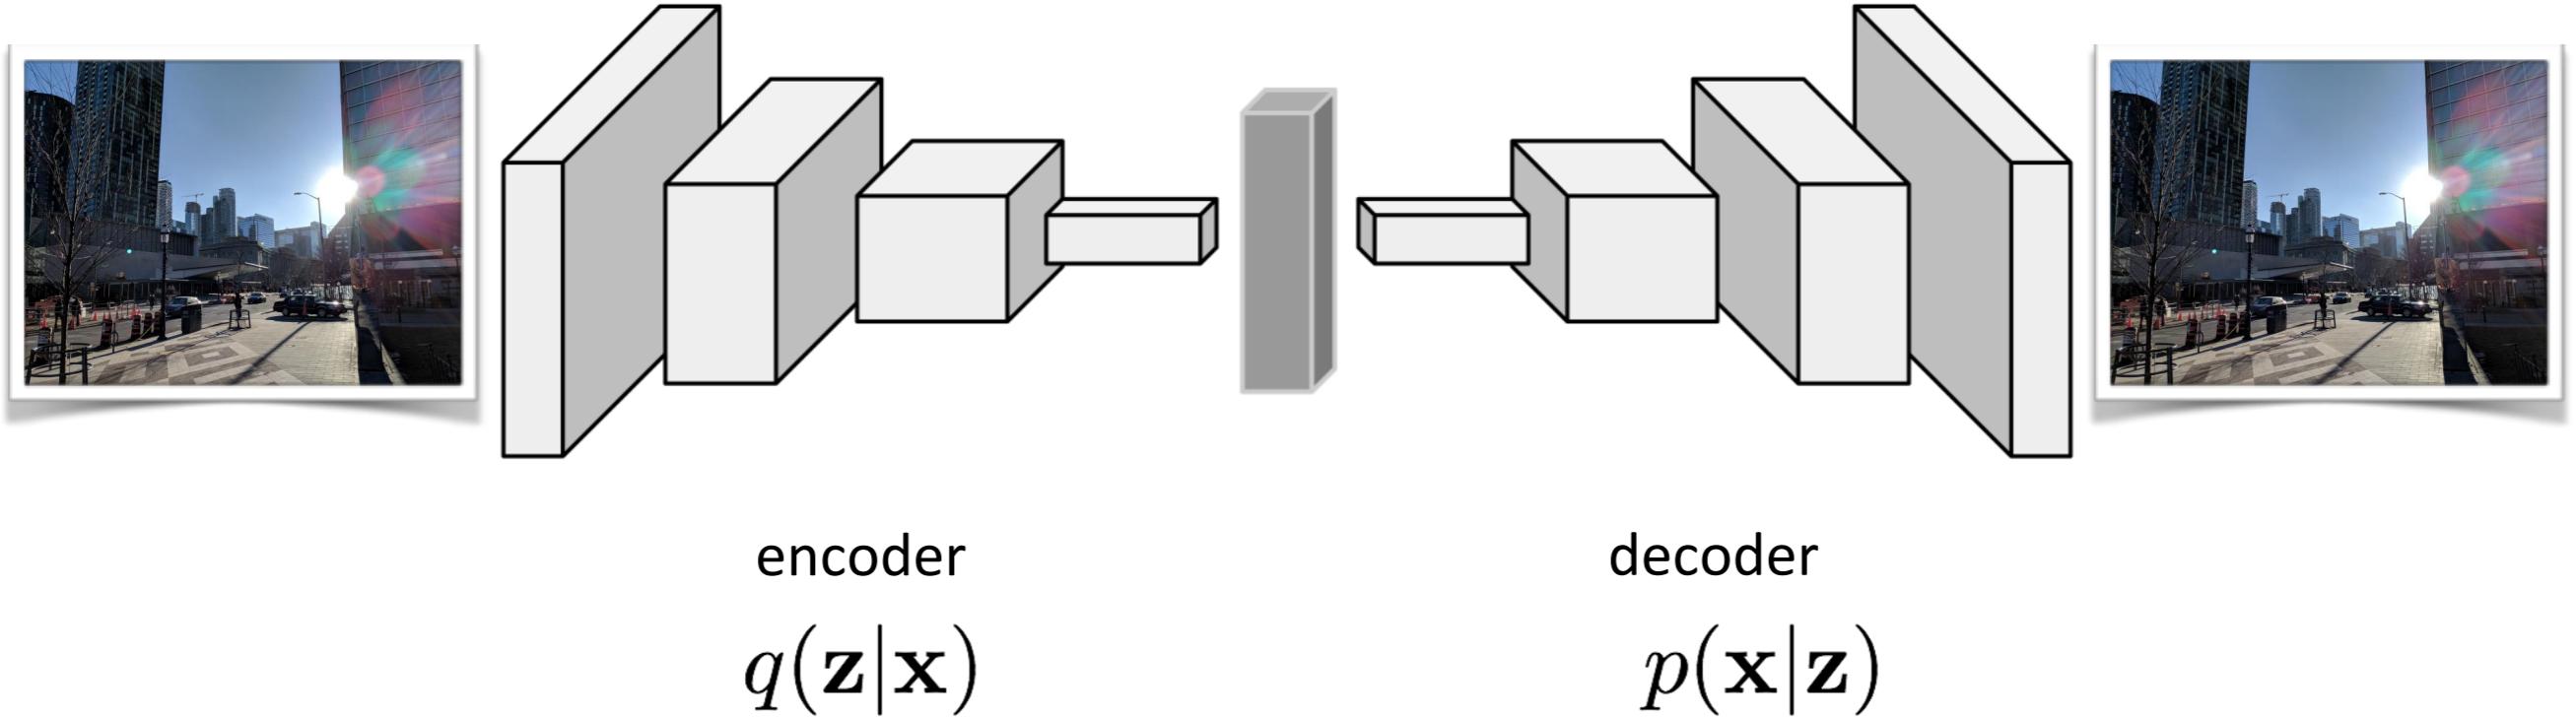
$$\mathbf{z} \sim p_{\text{model}}$$



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

VAE - Example

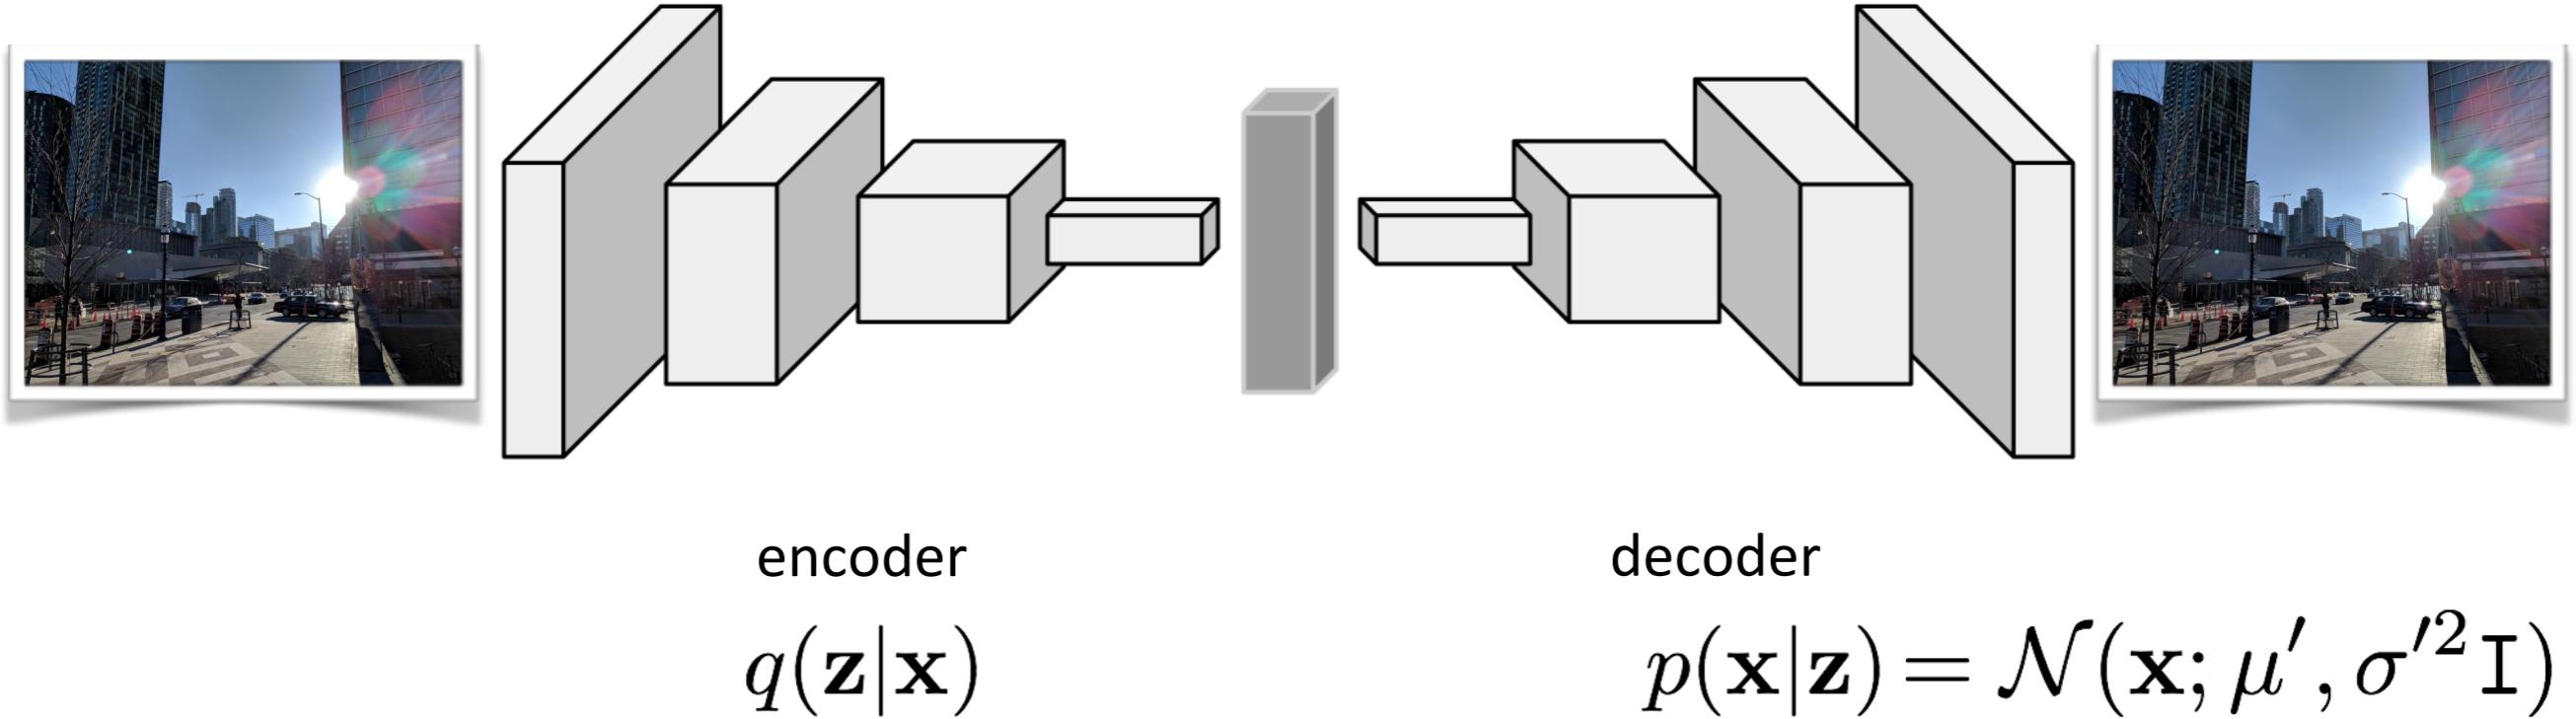
$$\mathbf{z} \sim p_{\text{model}} = \mathcal{N}(\mathbf{z}; \mathbf{0}, \mathbf{I})$$



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

VAE - Example

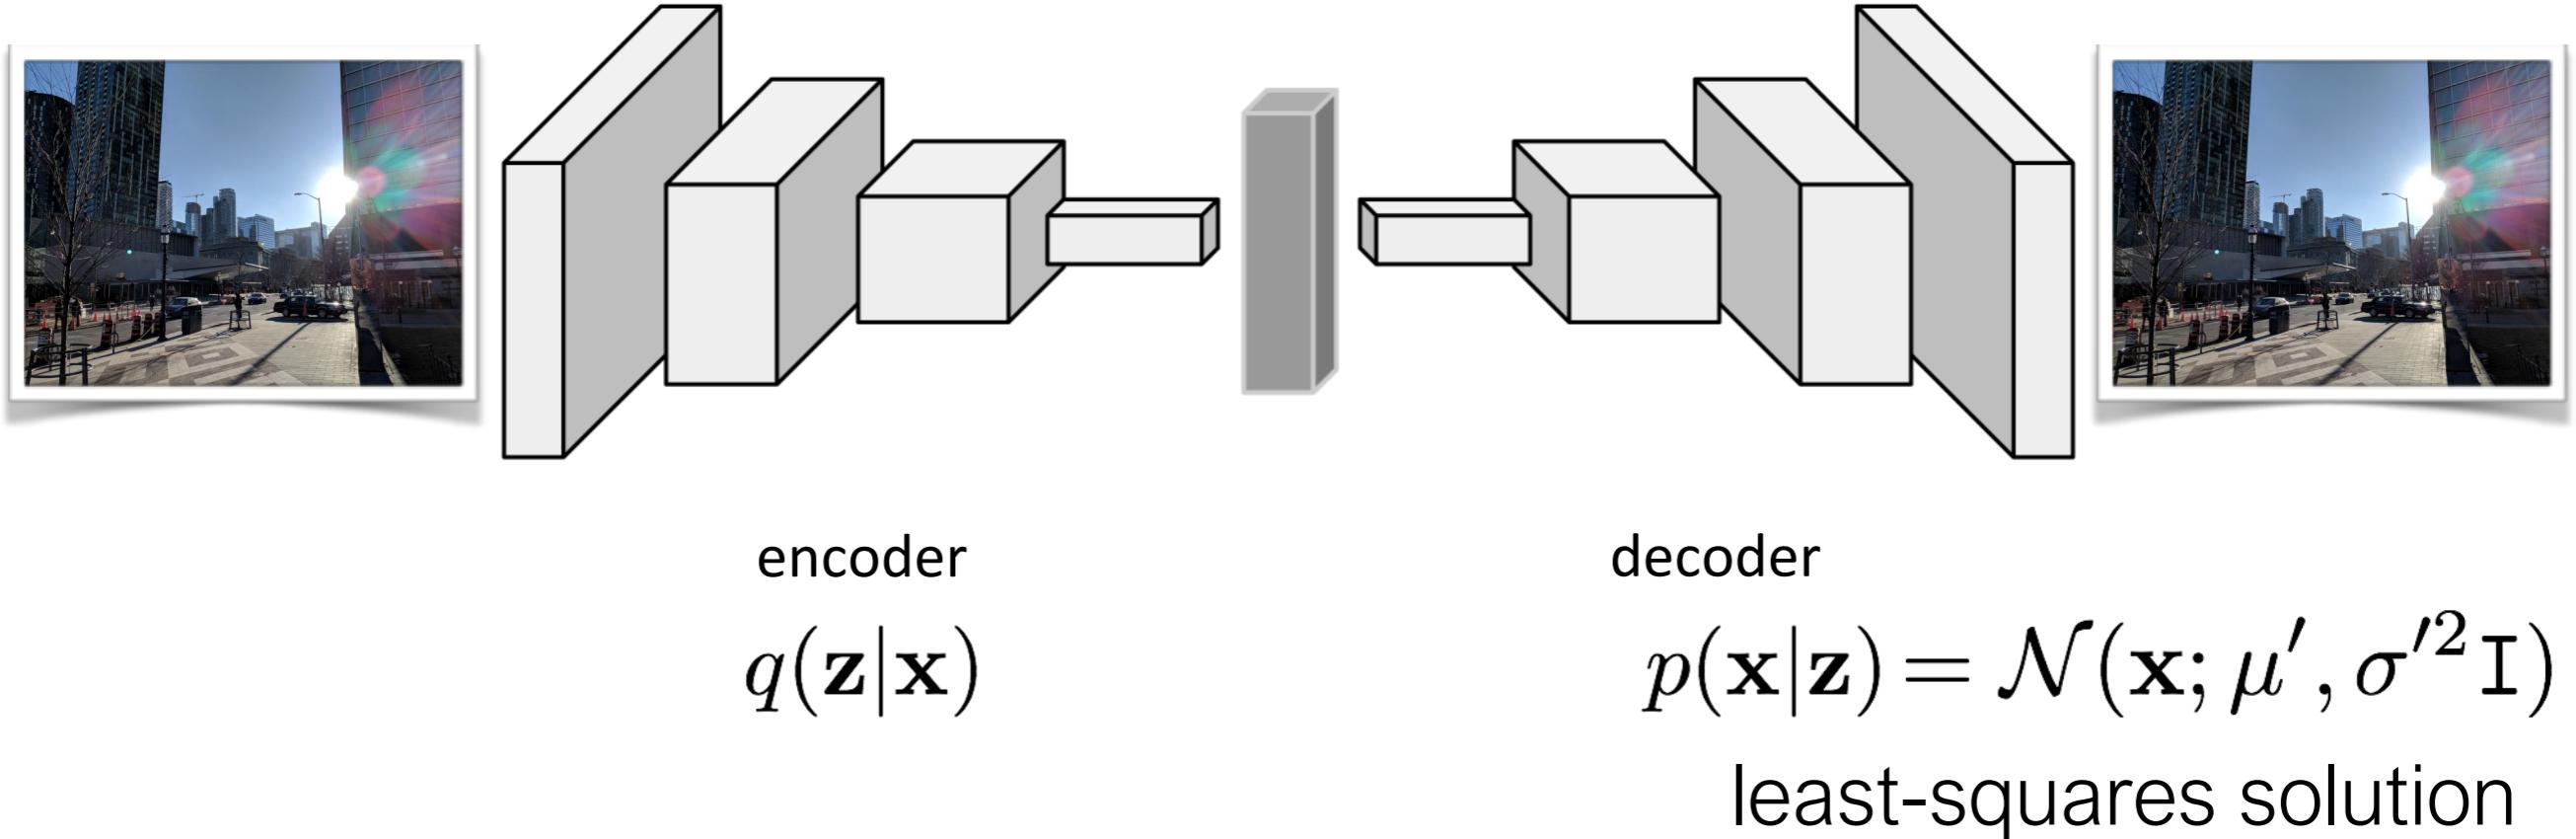
$$\mathbf{z} \sim p_{\text{model}} = \mathcal{N}(\mathbf{z}; \mathbf{0}, \mathbf{I})$$



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

VAE - Example

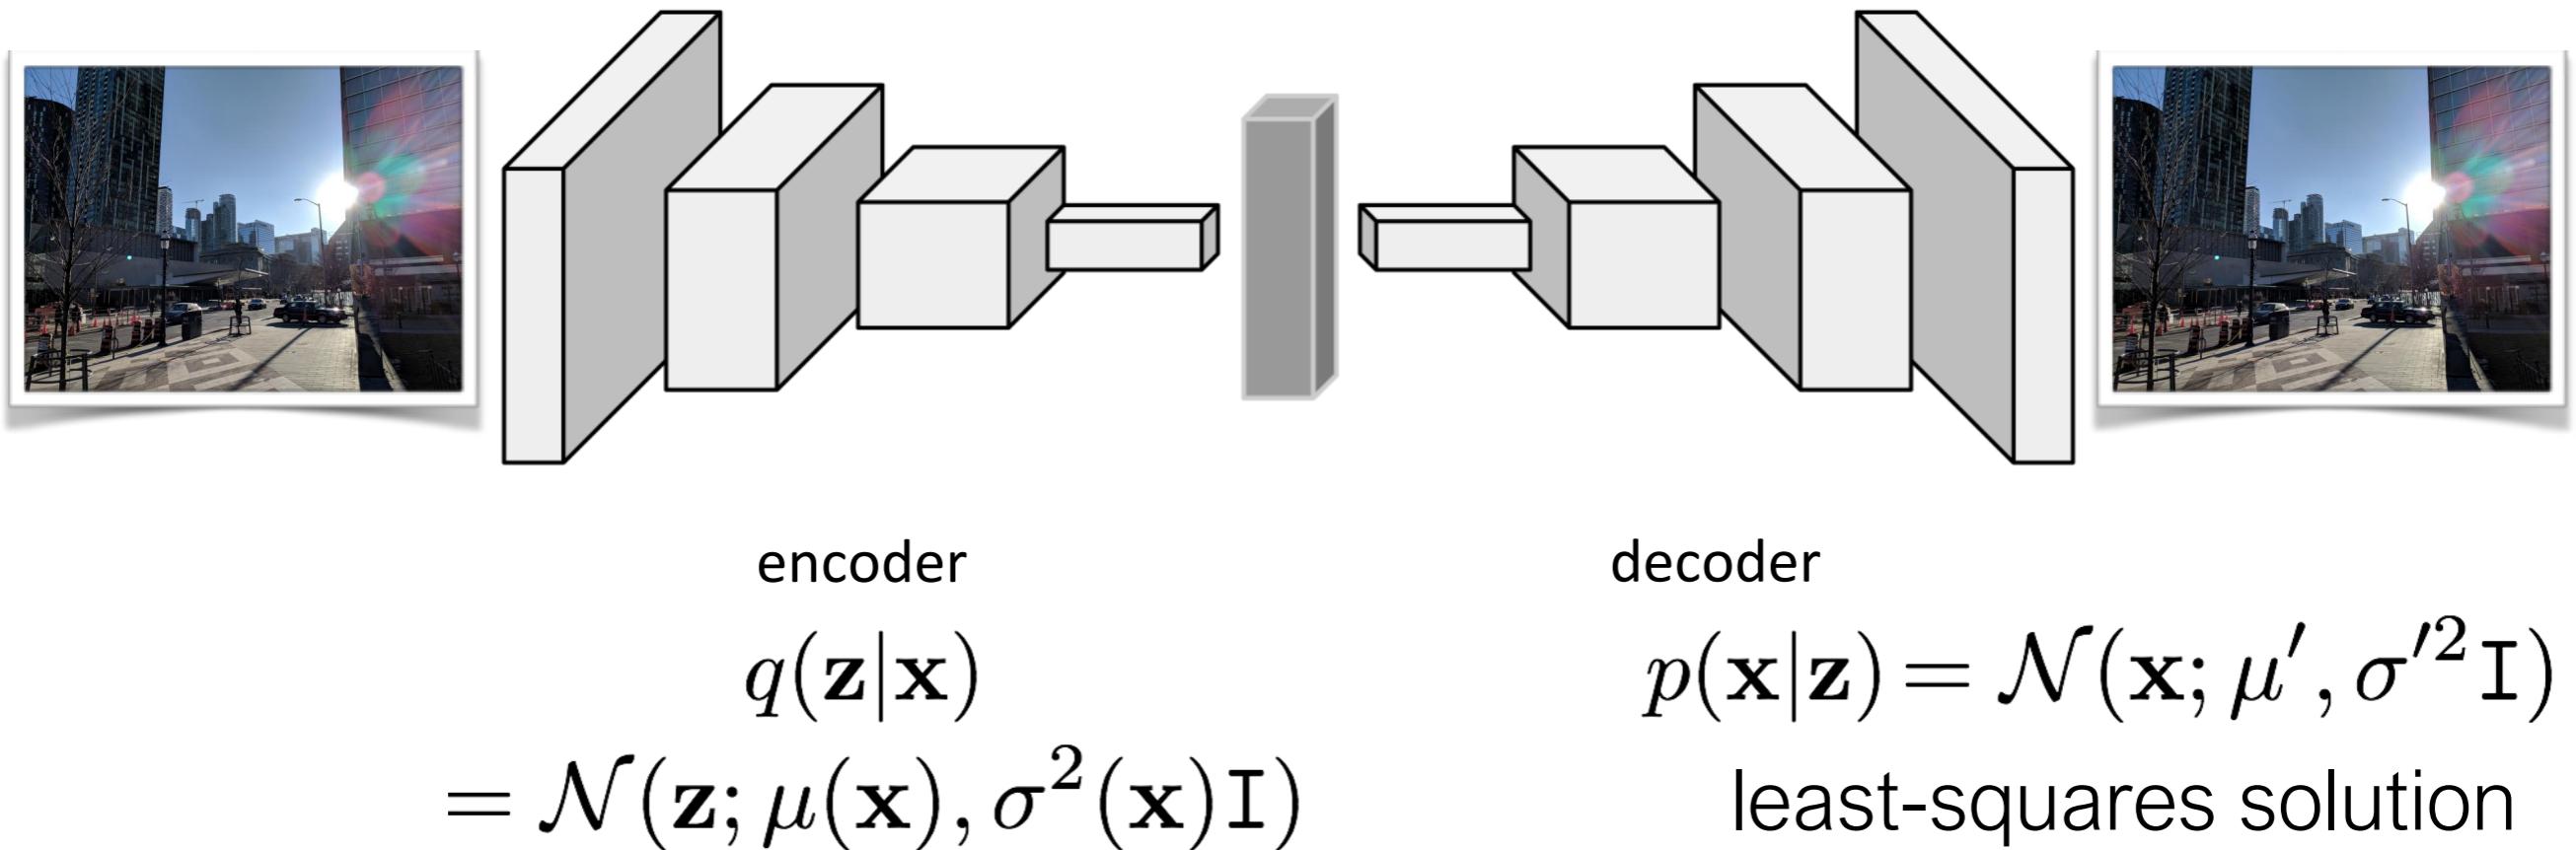
$$\mathbf{z} \sim p_{\text{model}} = \mathcal{N}(\mathbf{z}; \mathbf{0}, \mathbf{I})$$



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

VAE - Example

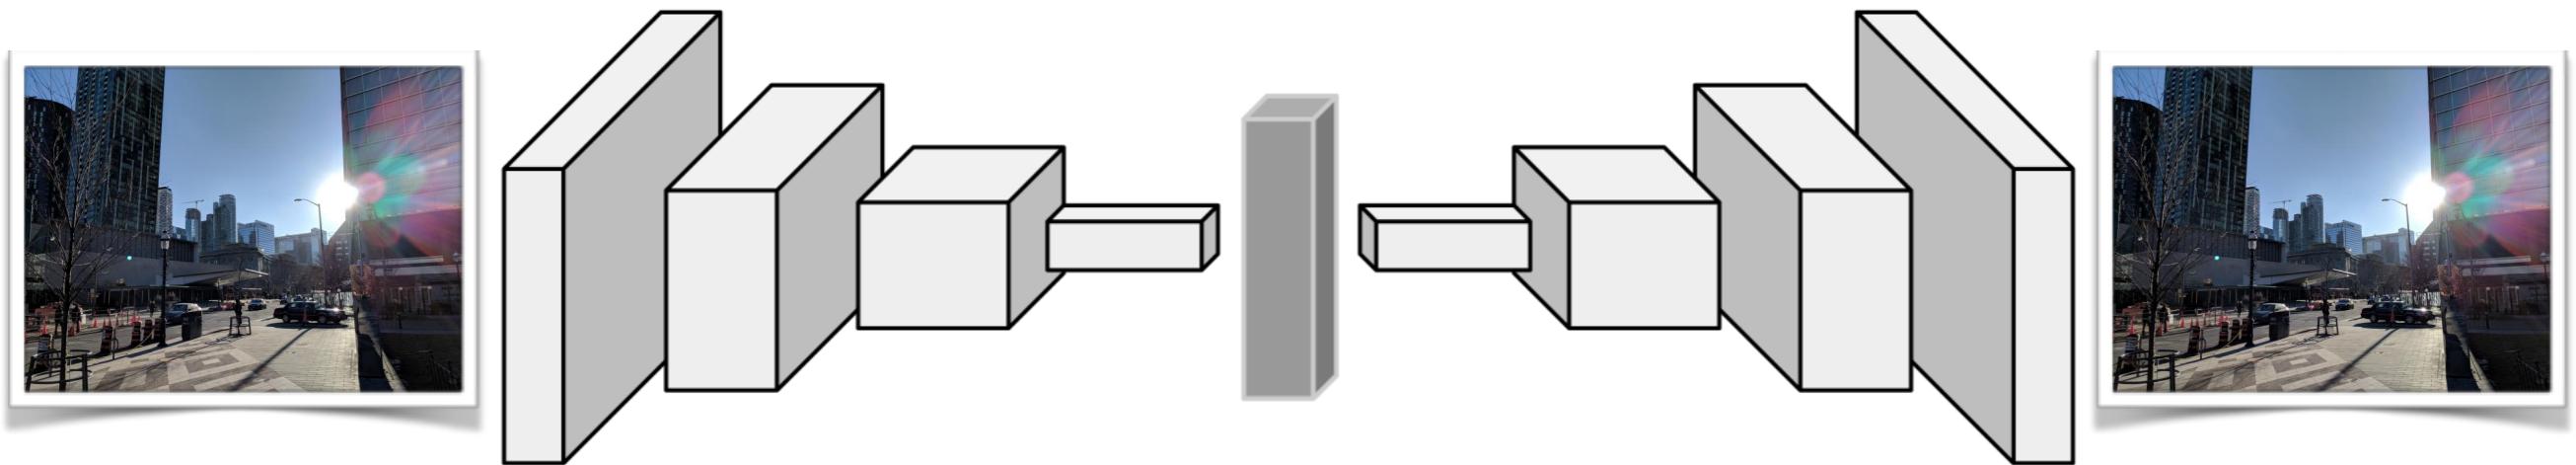
$$\mathbf{z} \sim p_{\text{model}} = \mathcal{N}(\mathbf{z}; \mathbf{0}, \mathbf{I})$$



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

VAE - Example

$$\mathbf{z} \sim p_{\text{model}} = \mathcal{N}(\mathbf{z}; \mathbf{0}, \mathbf{I})$$



encoder

$$q(\mathbf{z}|\mathbf{x}) \\ = \mathcal{N}(\mathbf{z}; \mu(\mathbf{x}), \sigma^2(\mathbf{x})\mathbf{I})$$

multi-variate Gaussian distribution

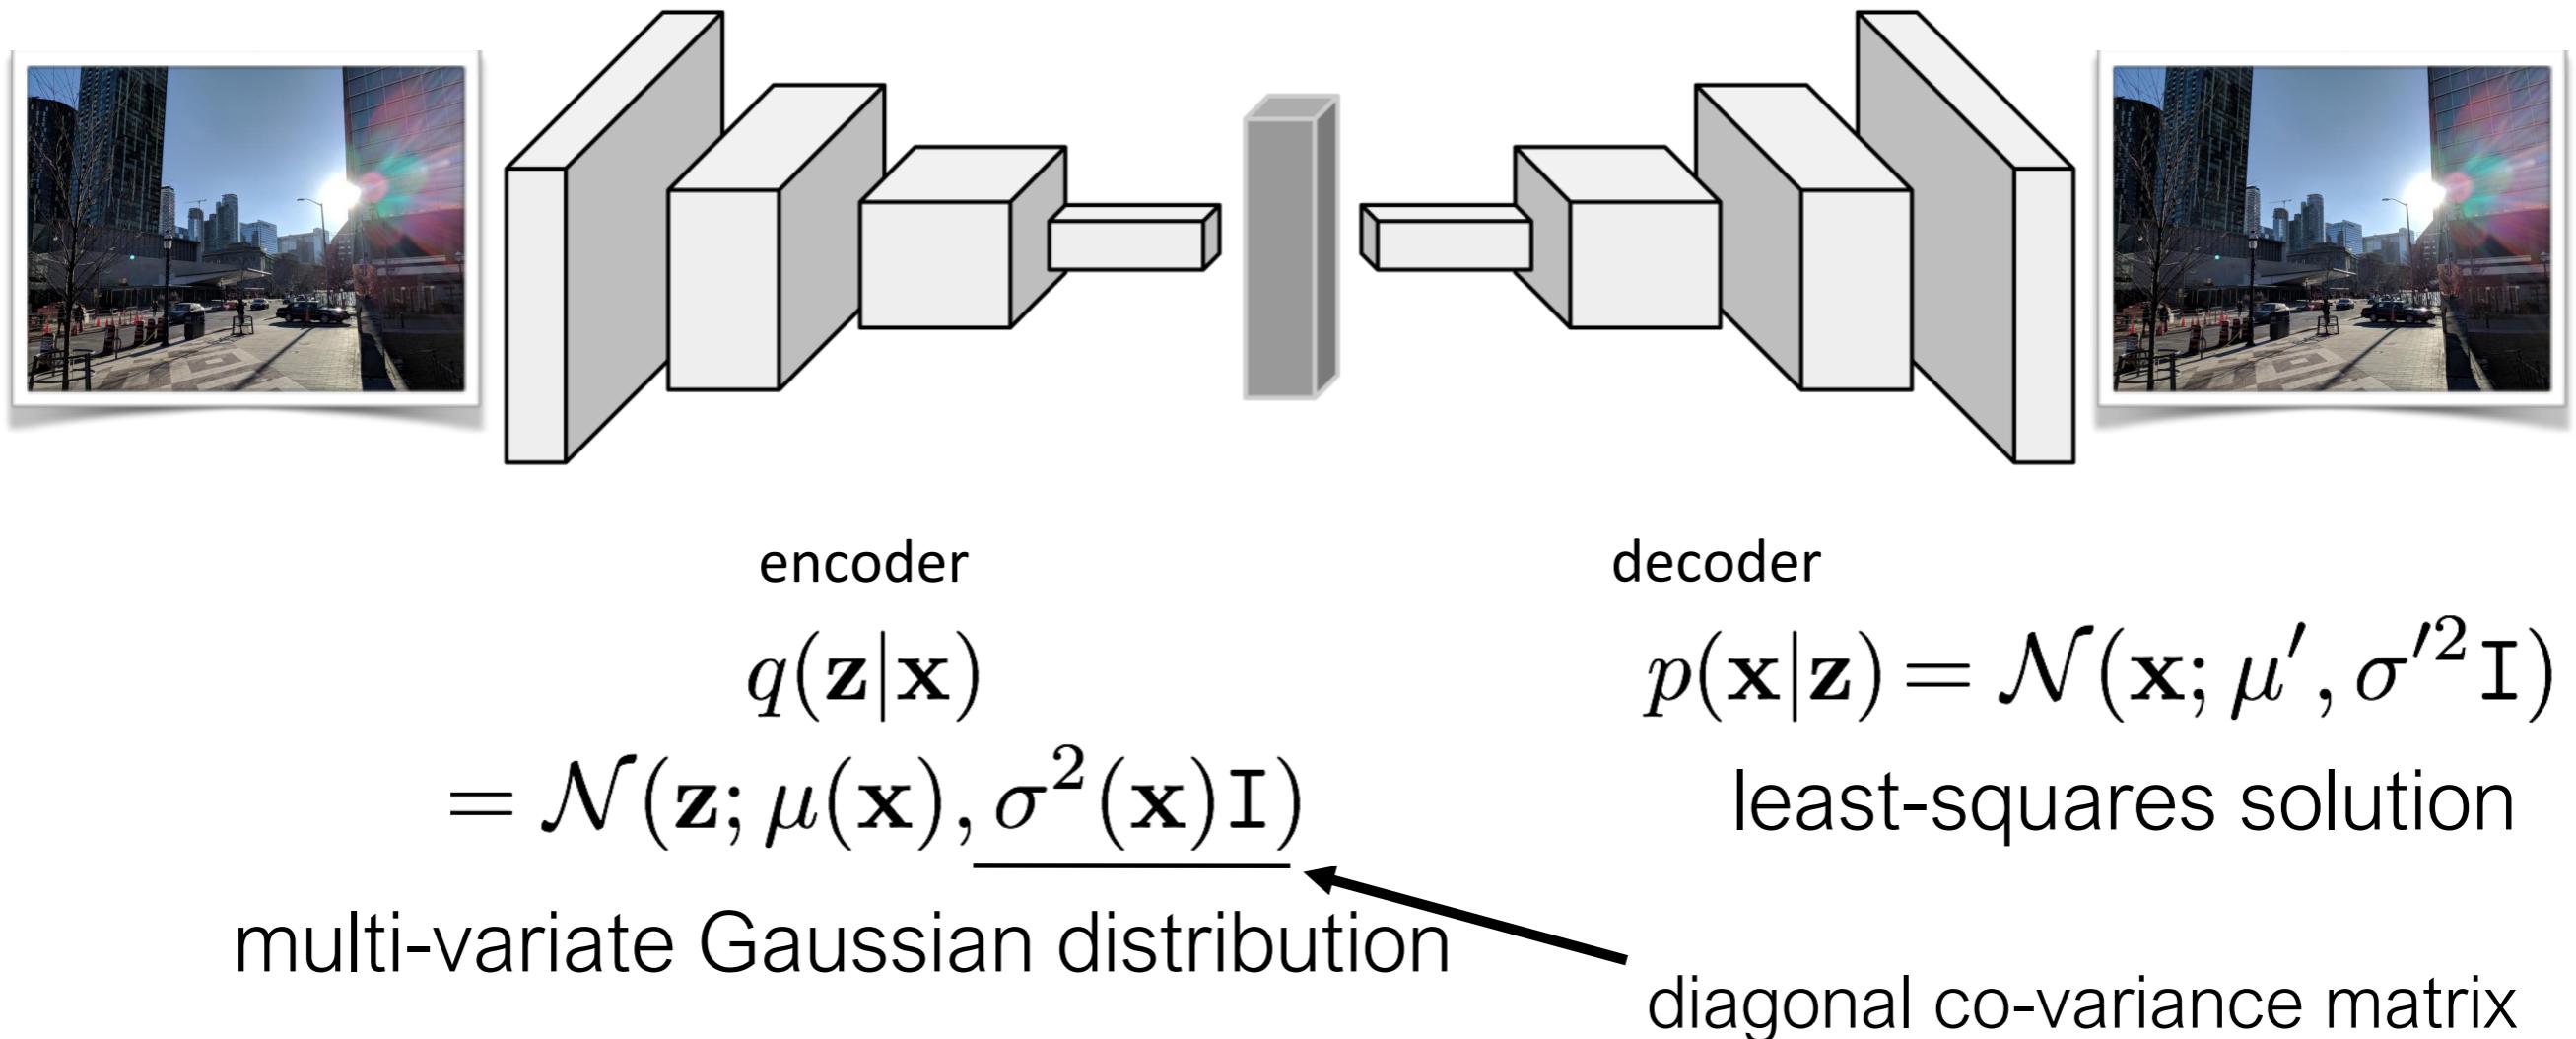
decoder

$$p(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{x}; \mu', \sigma'^2 \mathbf{I}) \\ \text{least-squares solution}$$

[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

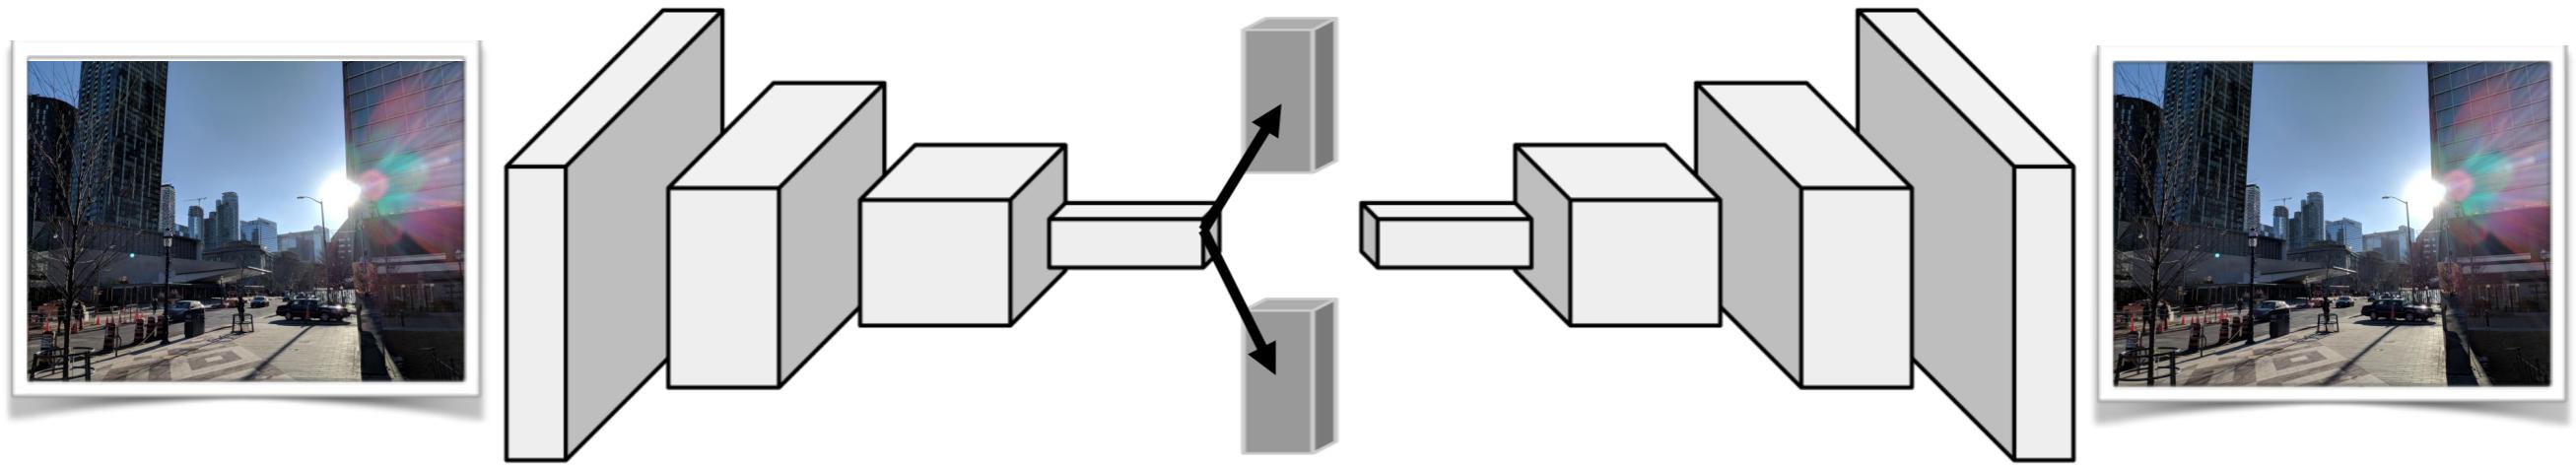
VAE - Example

$$\mathbf{z} \sim p_{\text{model}} = \mathcal{N}(\mathbf{z}; \mathbf{0}, \mathbf{I})$$



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

VAE - Example

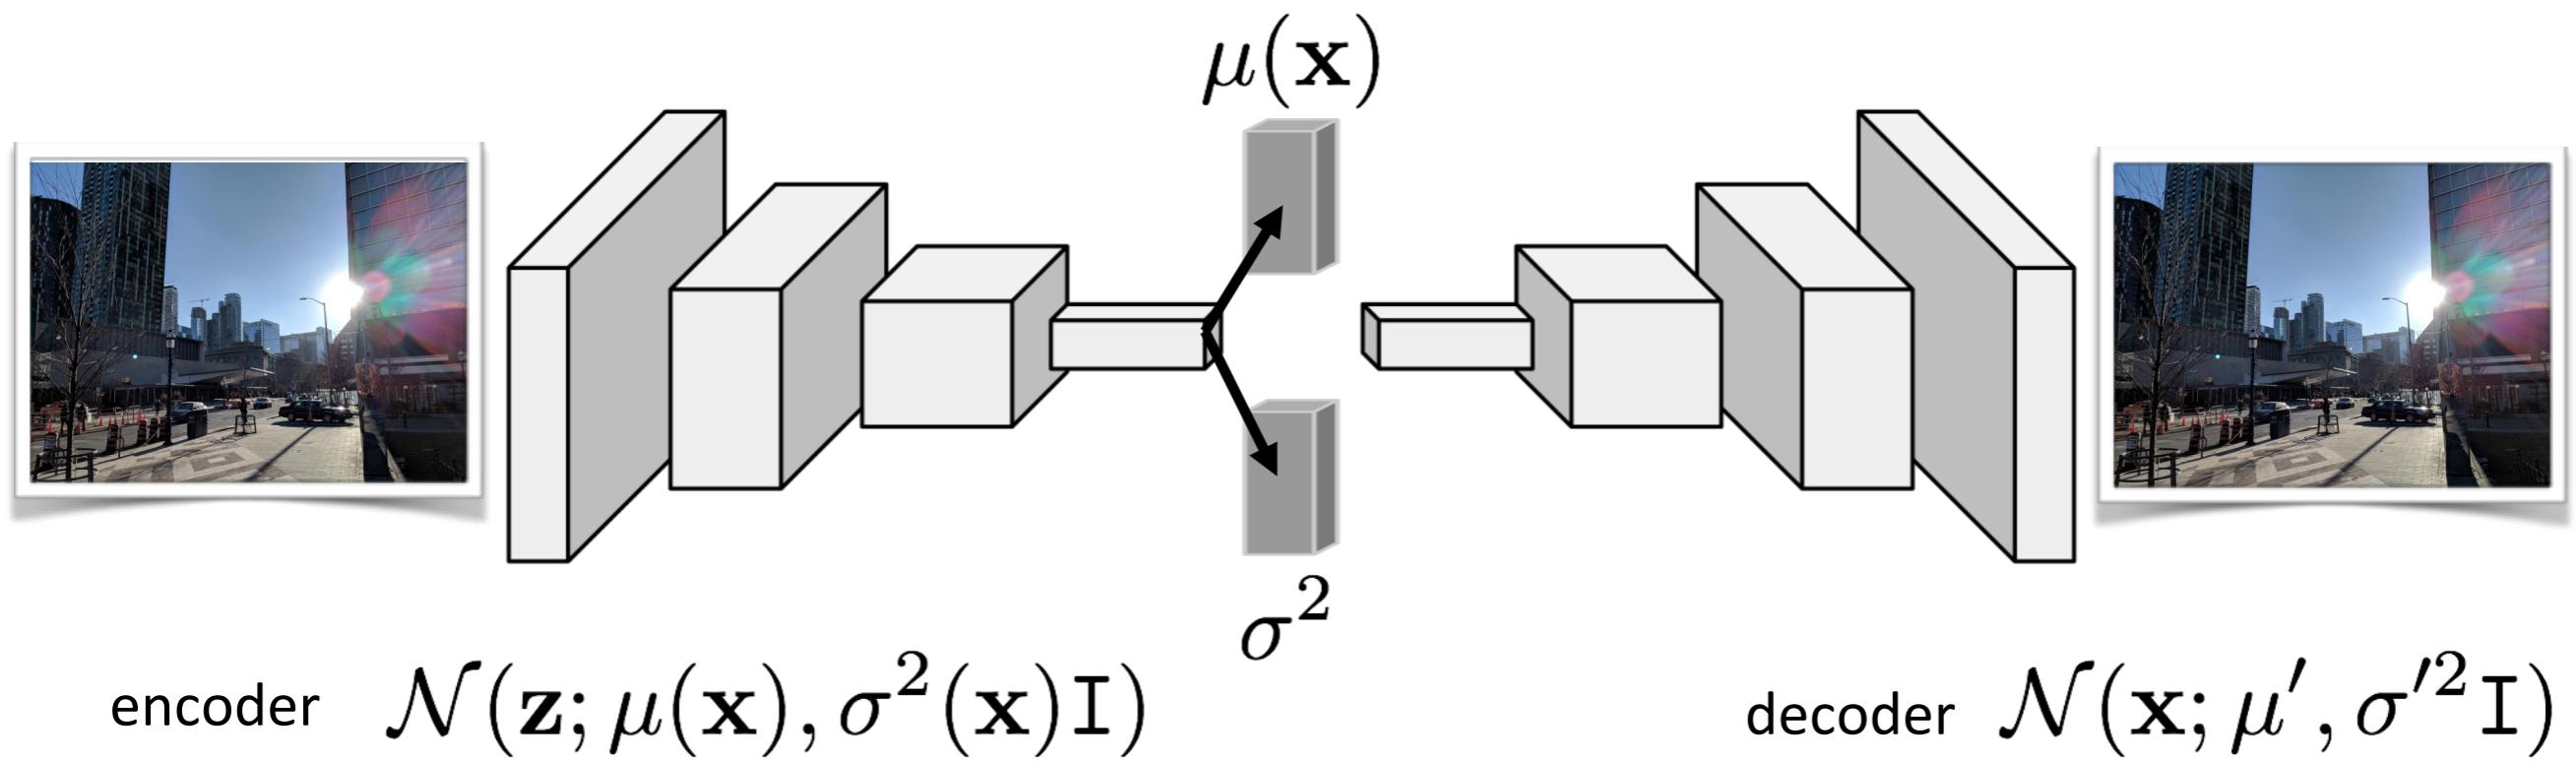


encoder $\mathcal{N}(\mathbf{z}; \mu(\mathbf{x}), \sigma^2(\mathbf{x})\mathbf{I})$

decoder $\mathcal{N}(\mathbf{x}; \mu', \sigma'^2\mathbf{I})$

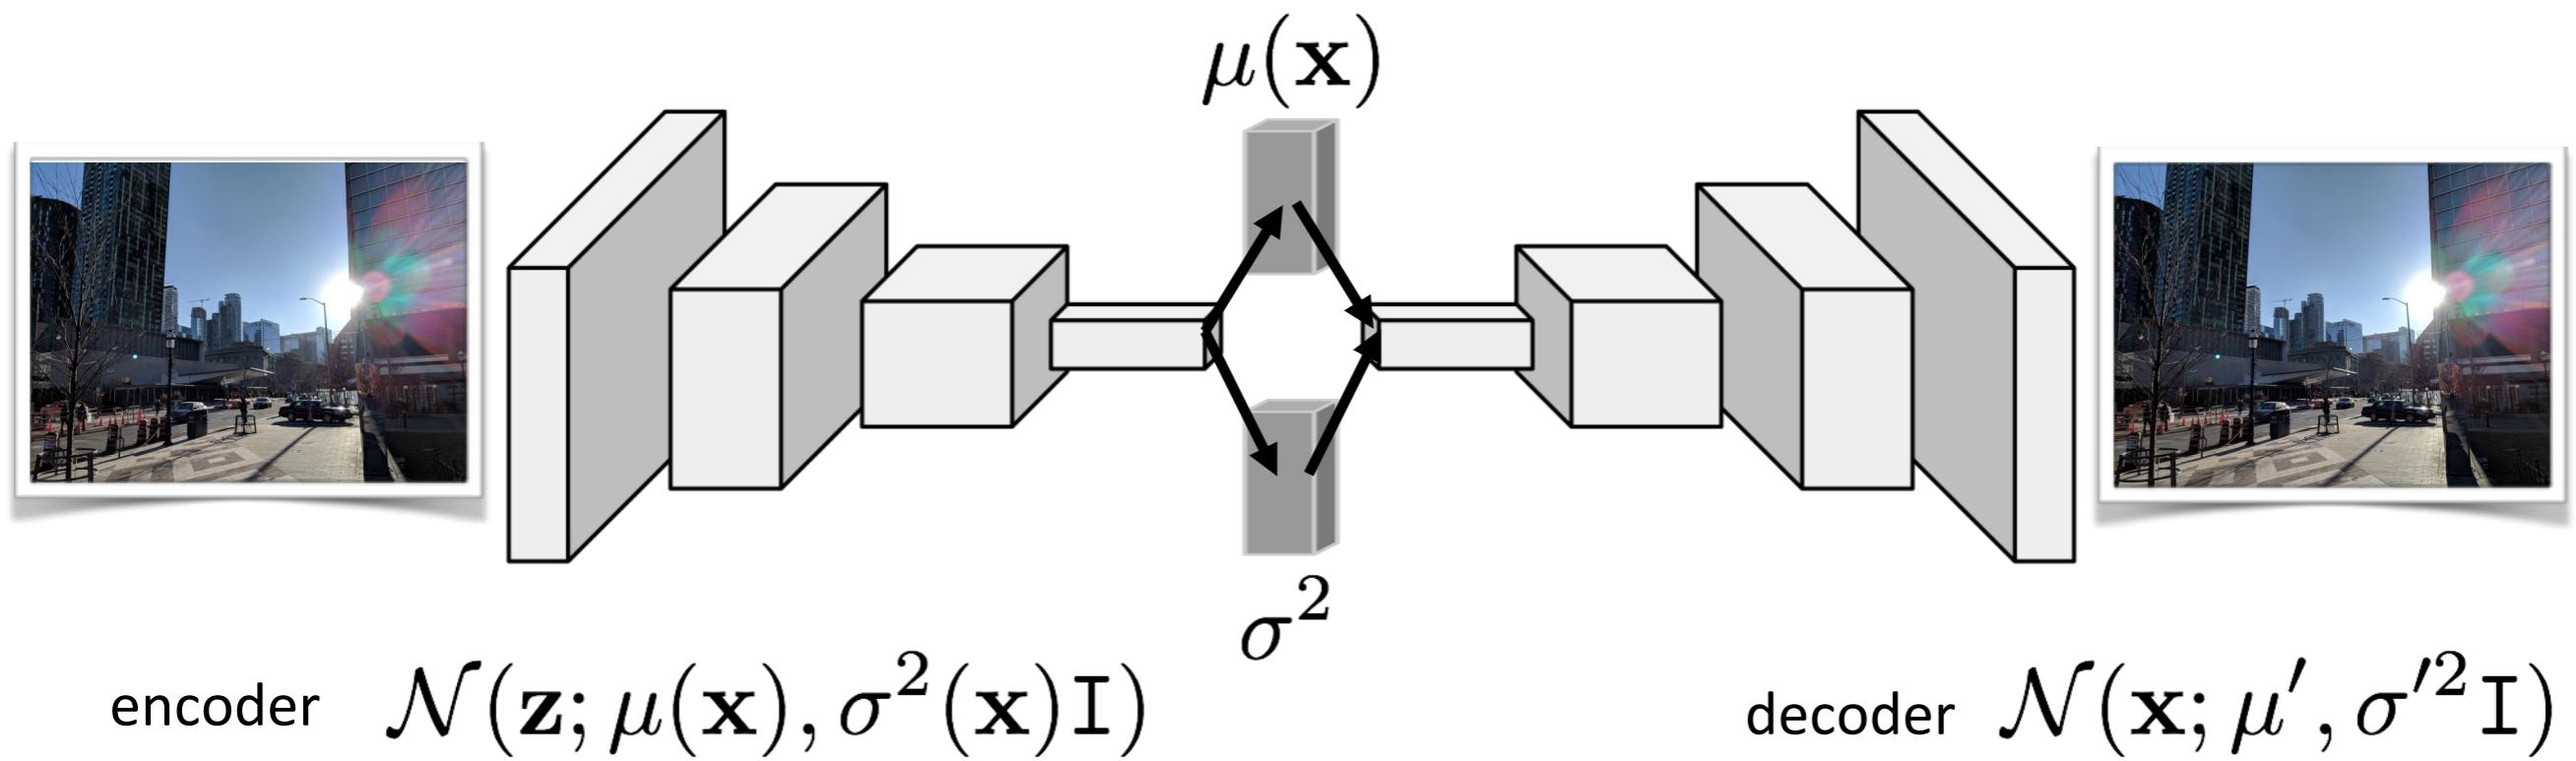
[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

VAE - Example



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

VAE - Example

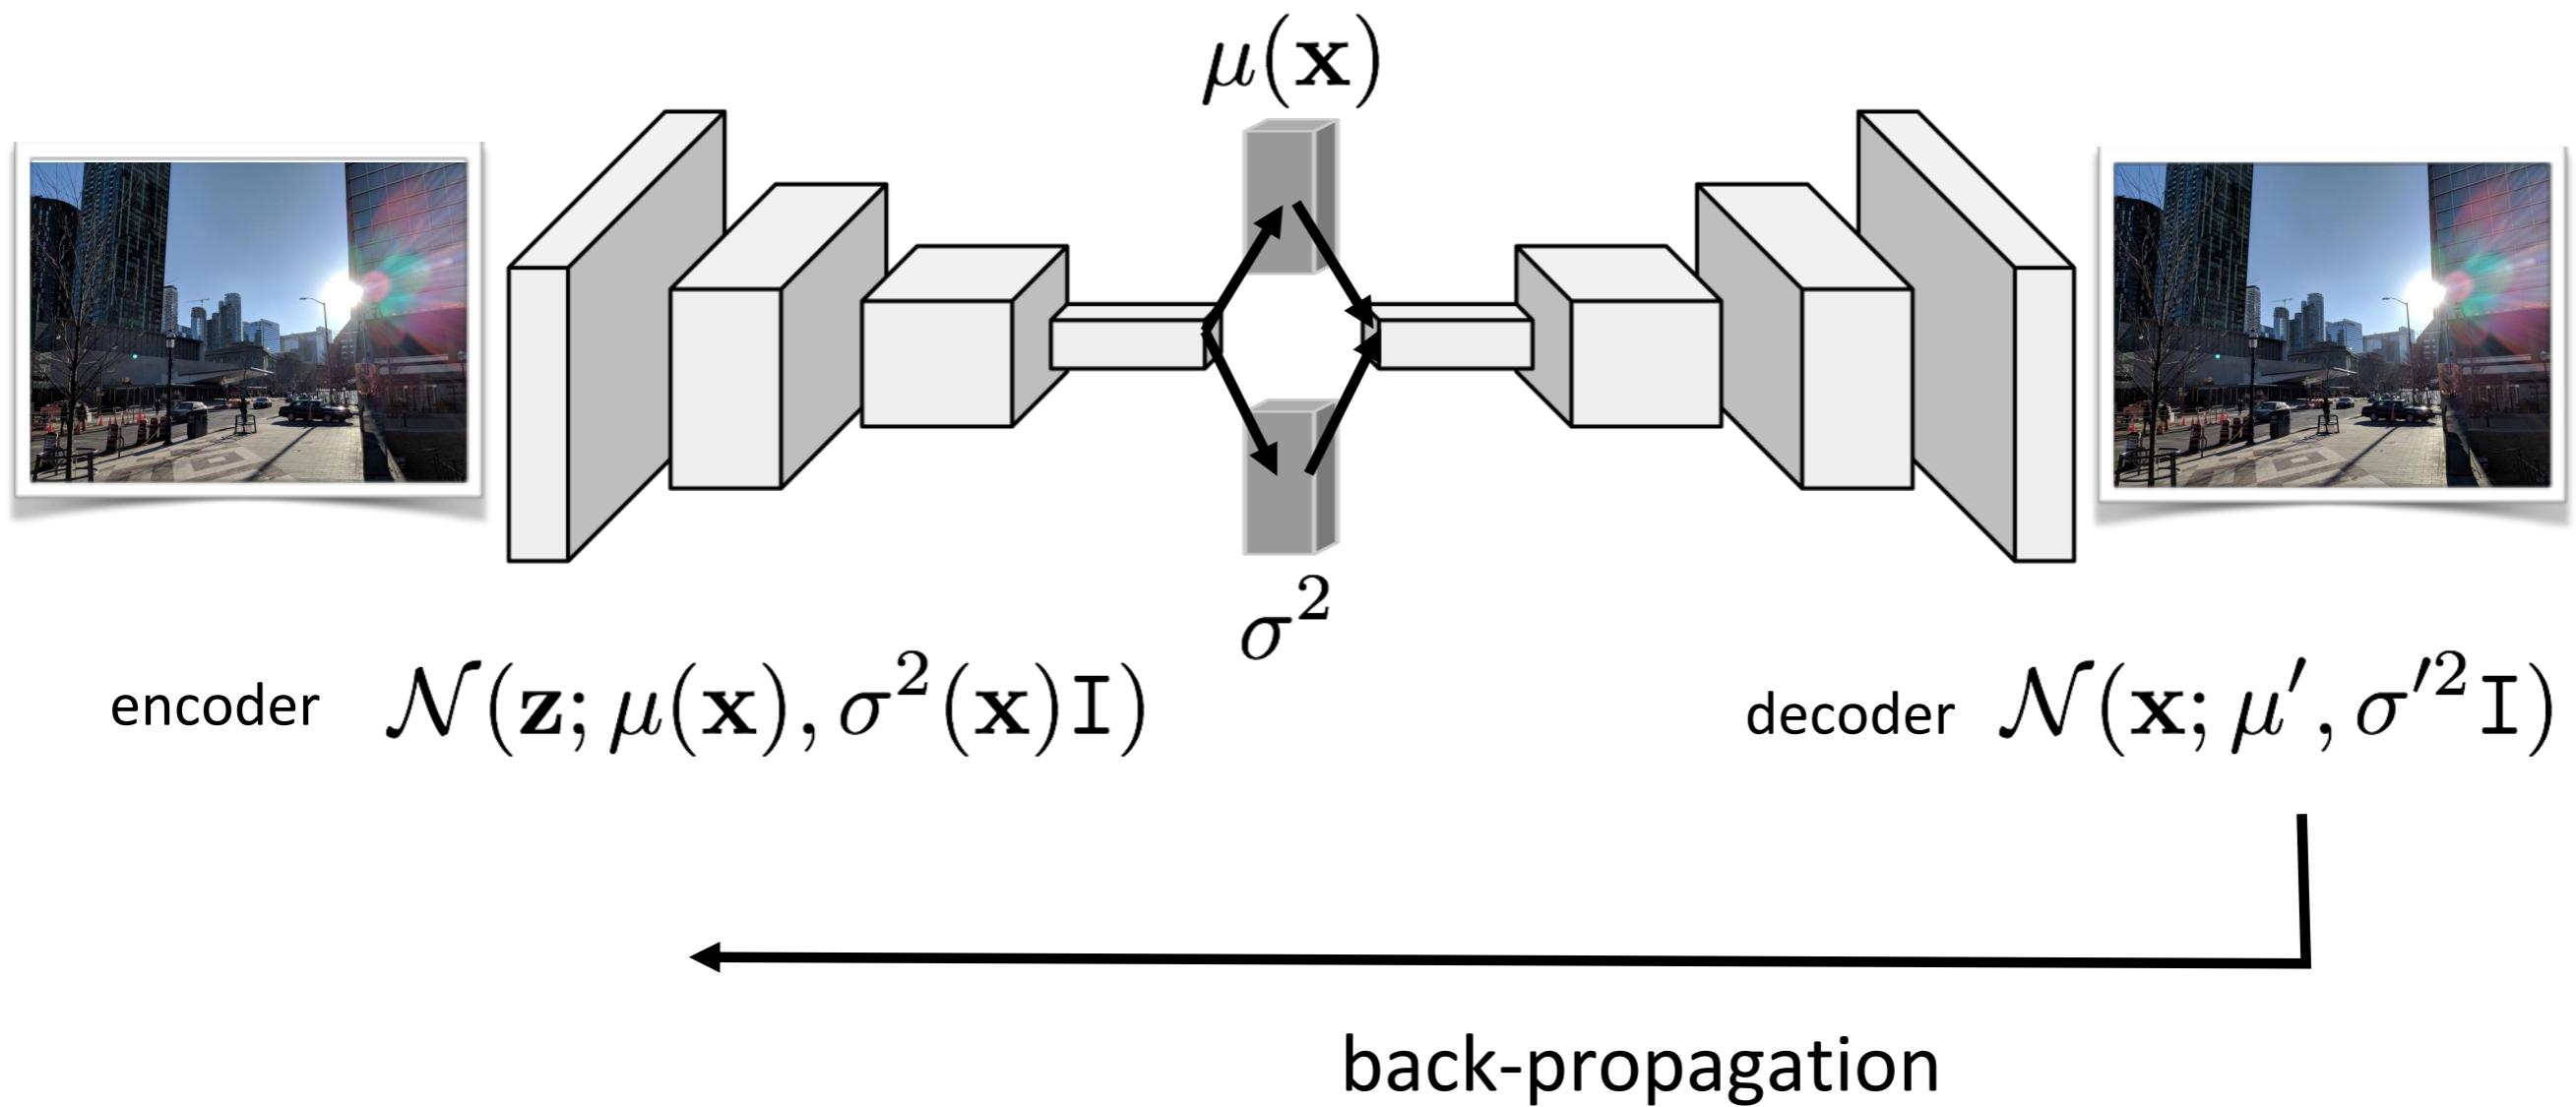


encoder $\mathcal{N}(\mathbf{z}; \mu(\mathbf{x}), \sigma^2(\mathbf{x})\mathbf{I})$

decoder $\mathcal{N}(\mathbf{x}; \mu', \sigma'^2\mathbf{I})$

[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

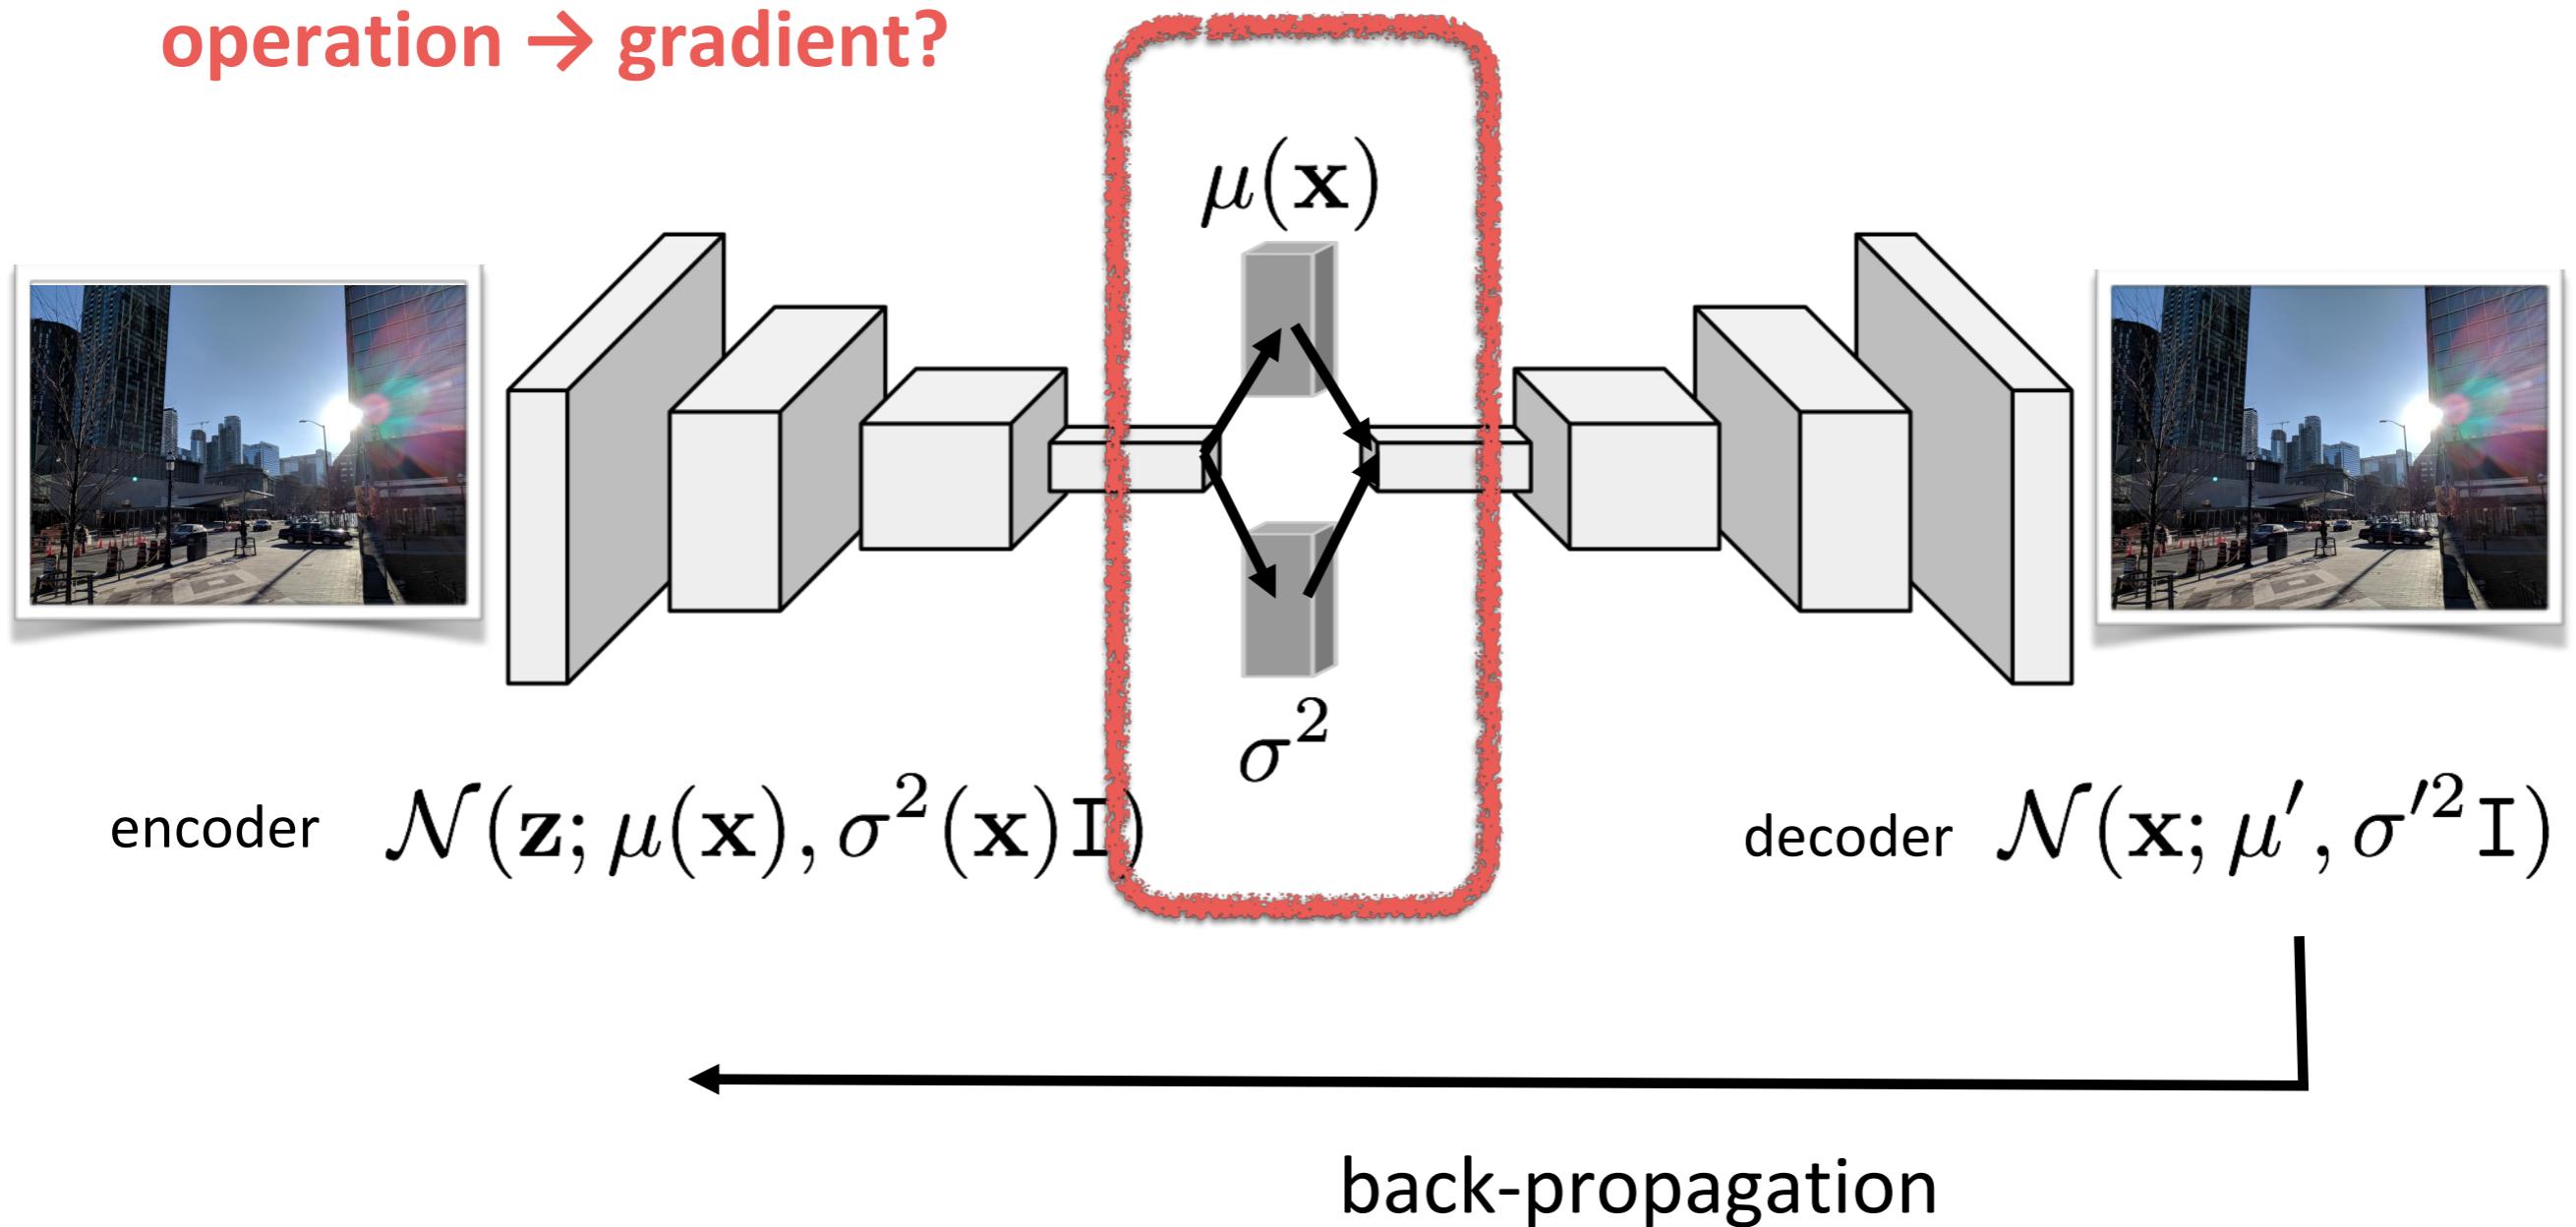
VAE - Example



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

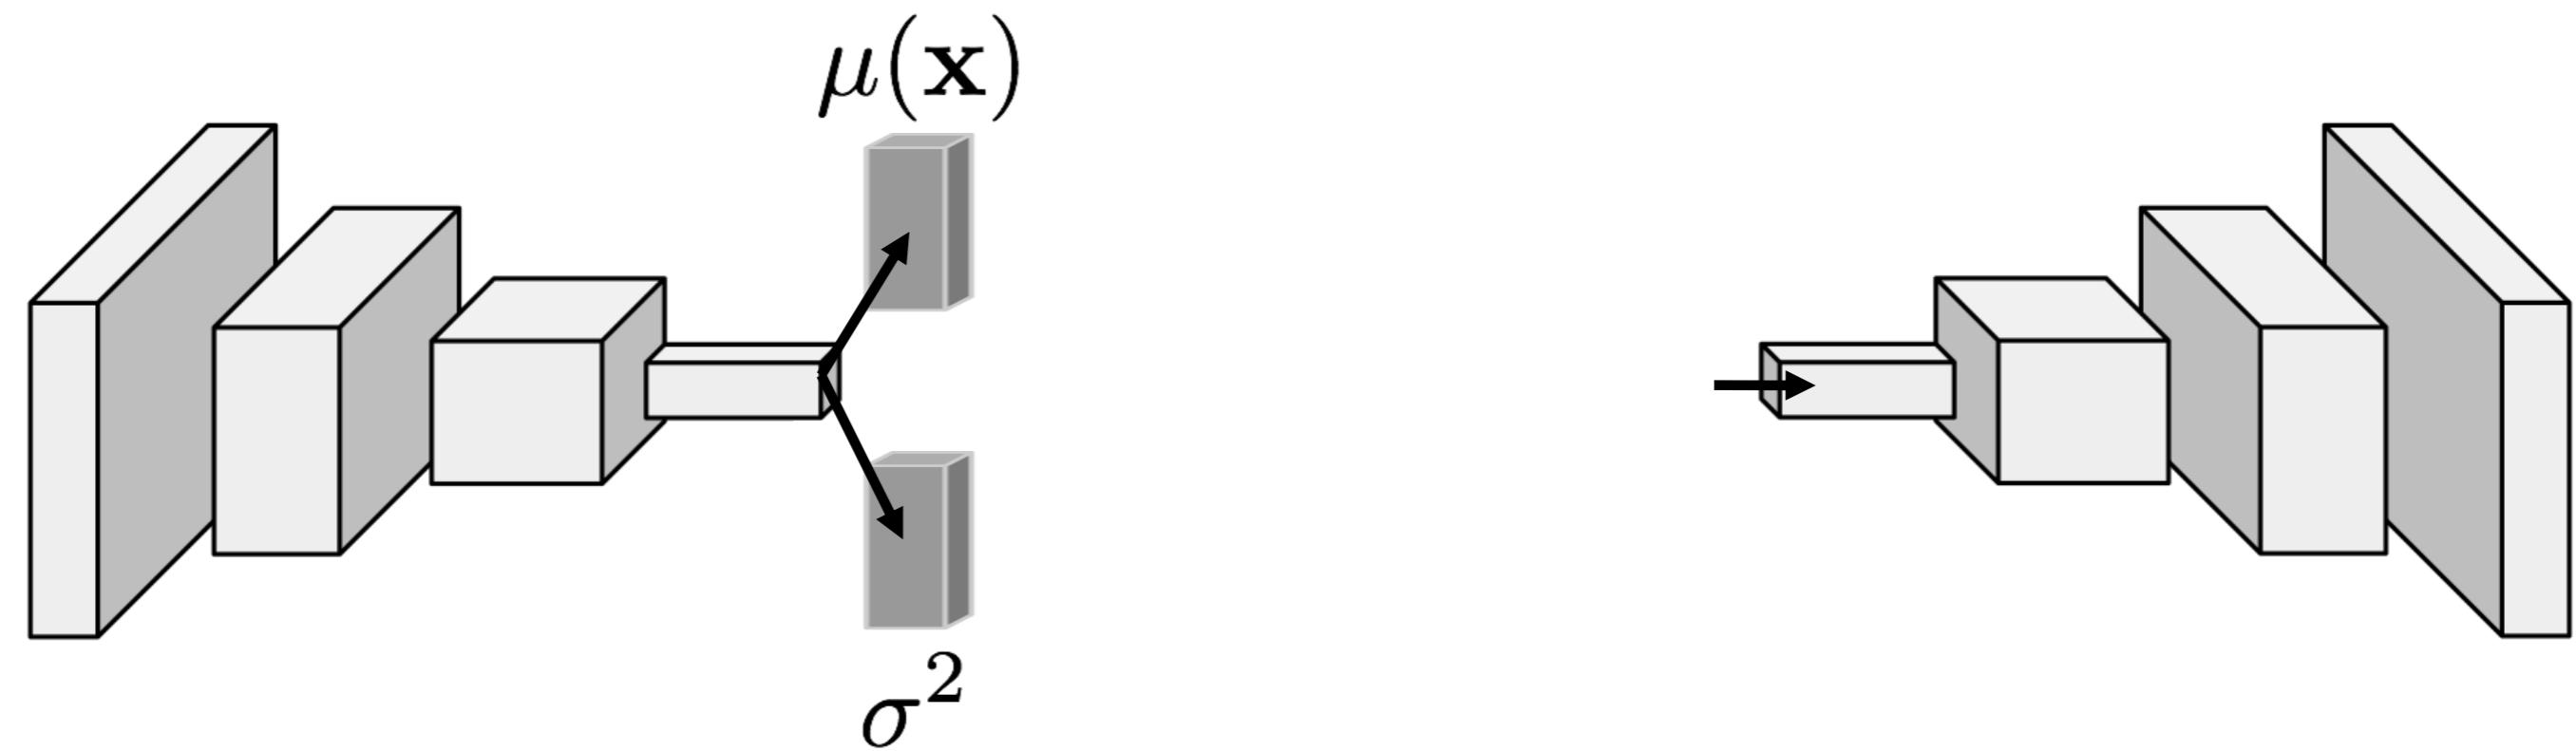
VAE - Example

sampling z is non-continuous
operation → gradient?



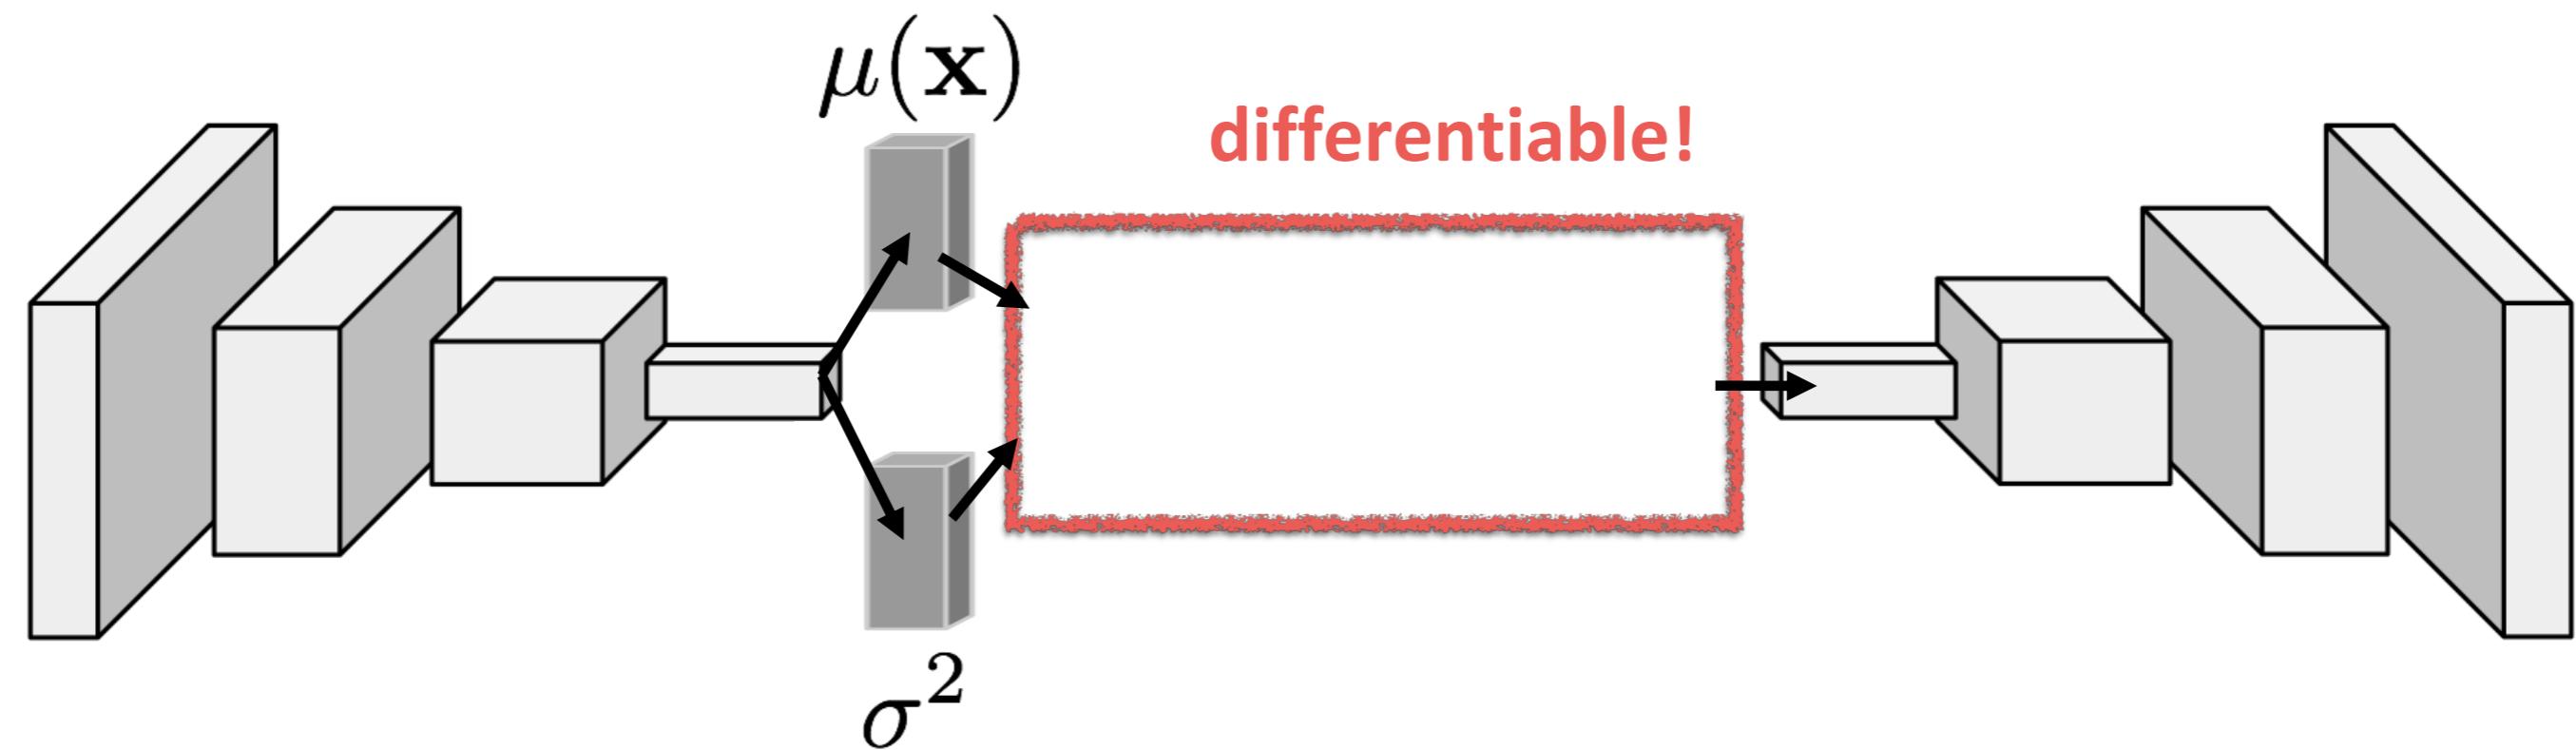
[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

The Reparametrization Trick



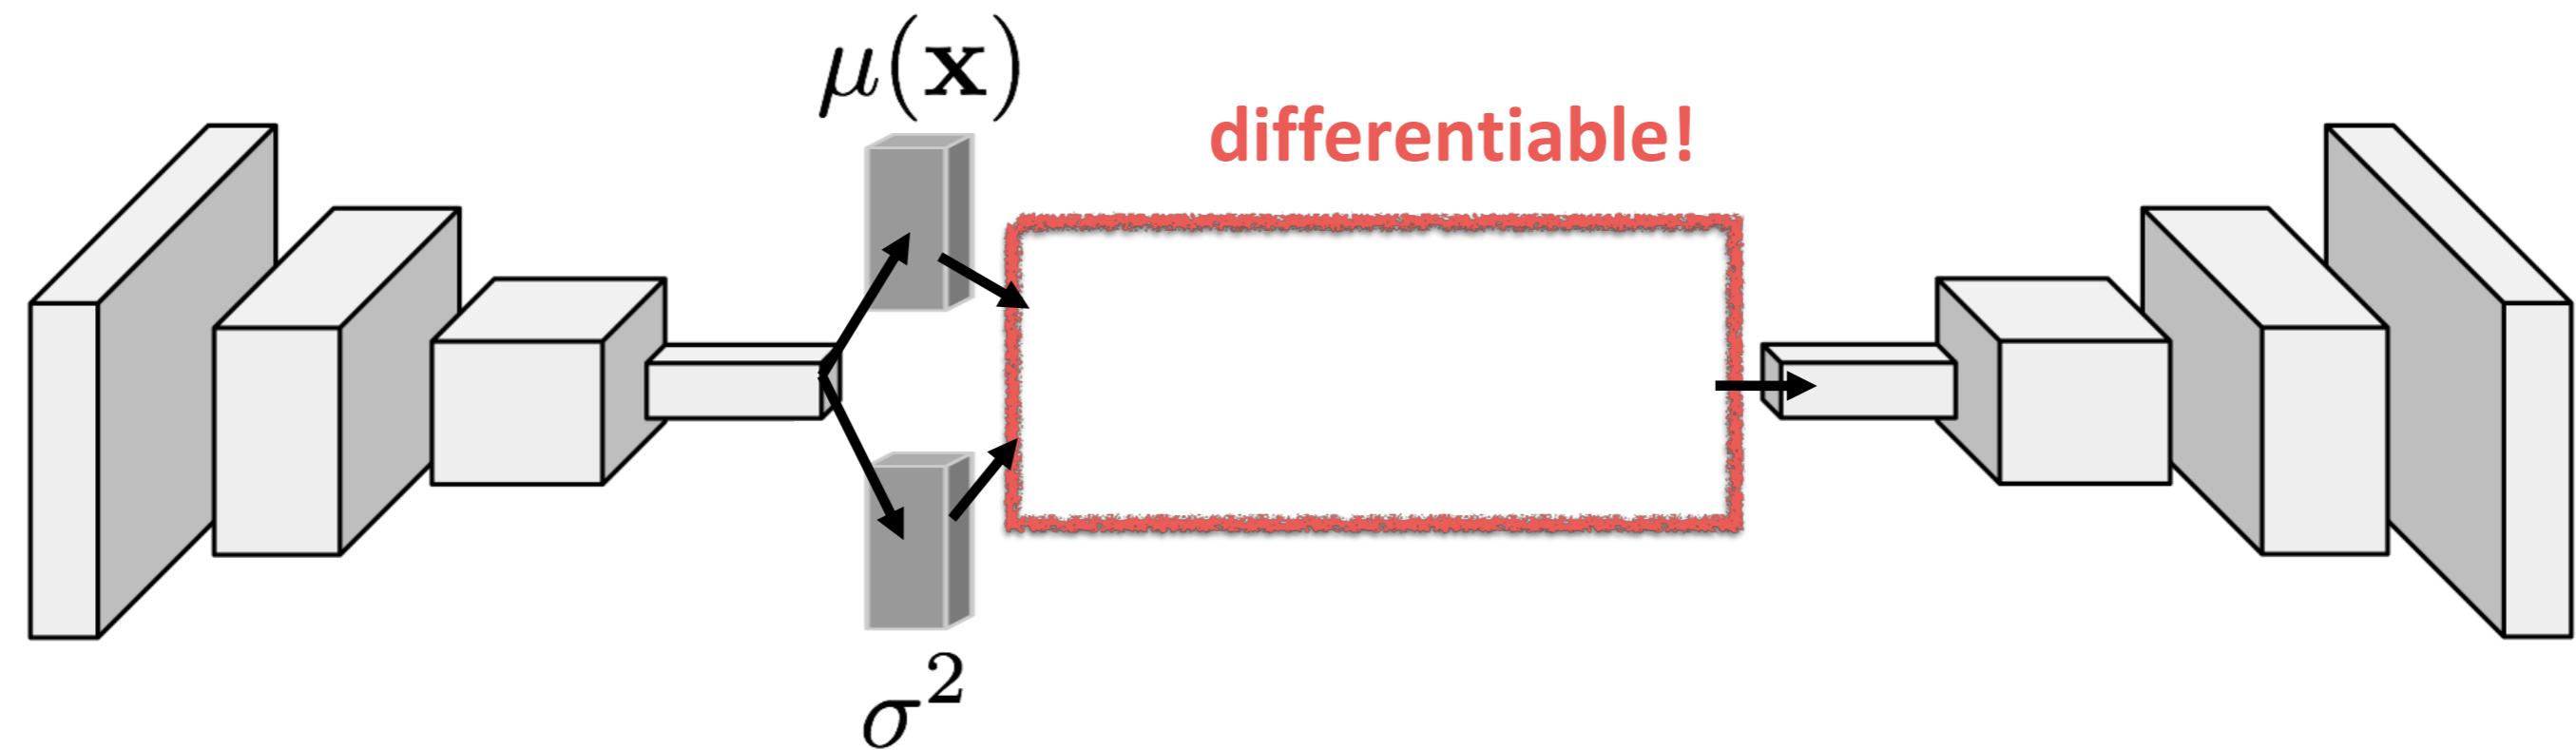
[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

The Reparametrization Trick



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

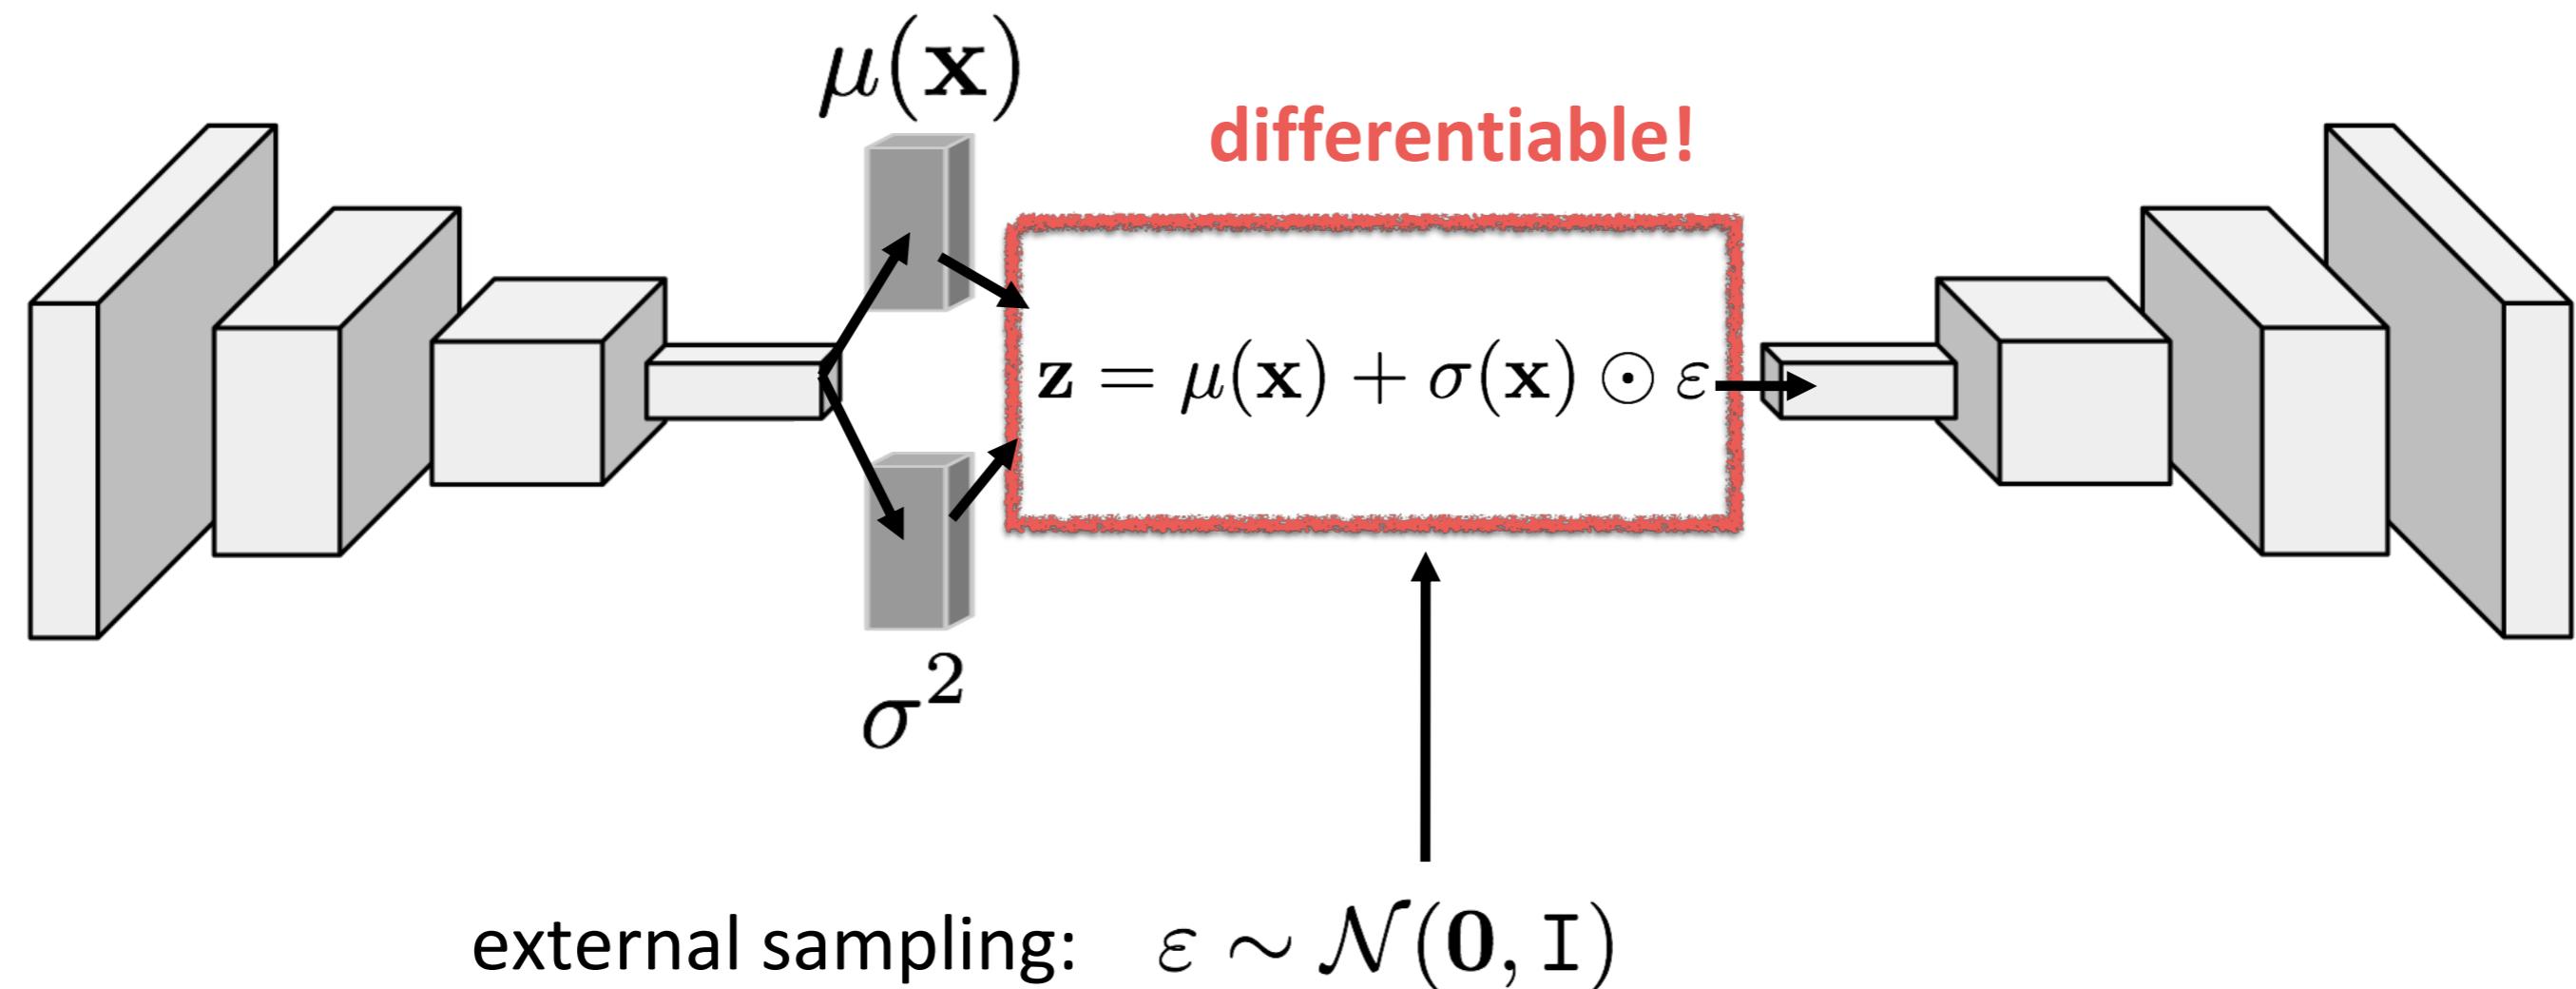
The Reparametrization Trick



external sampling: $\varepsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

The Reparametrization Trick



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

VAE Example - Results

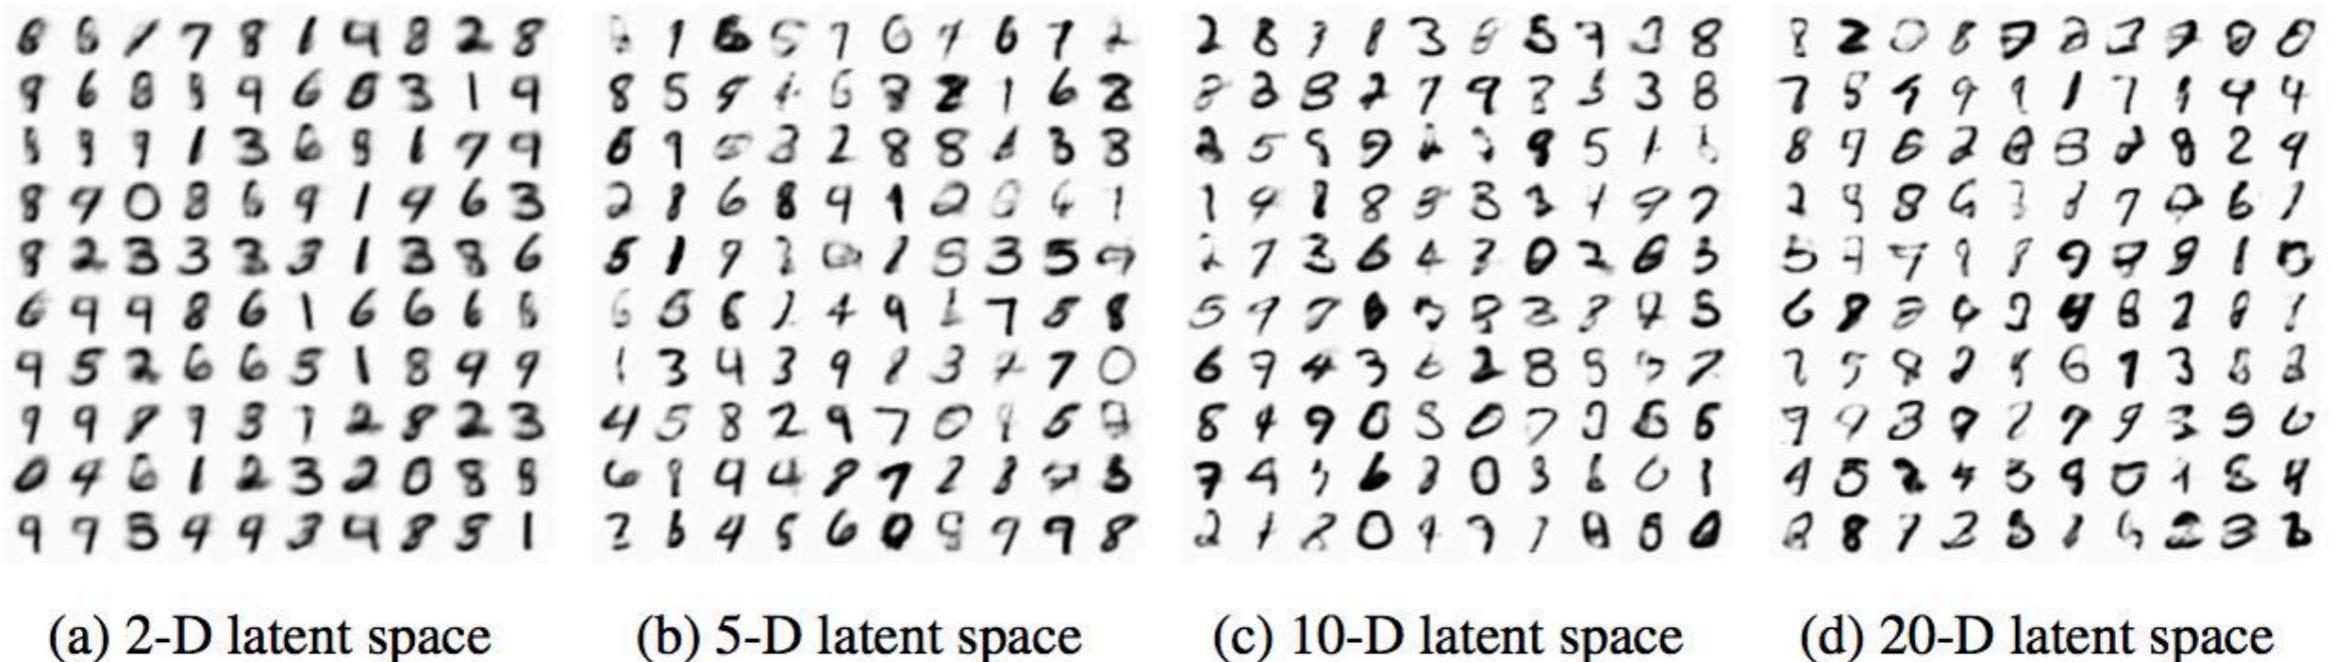


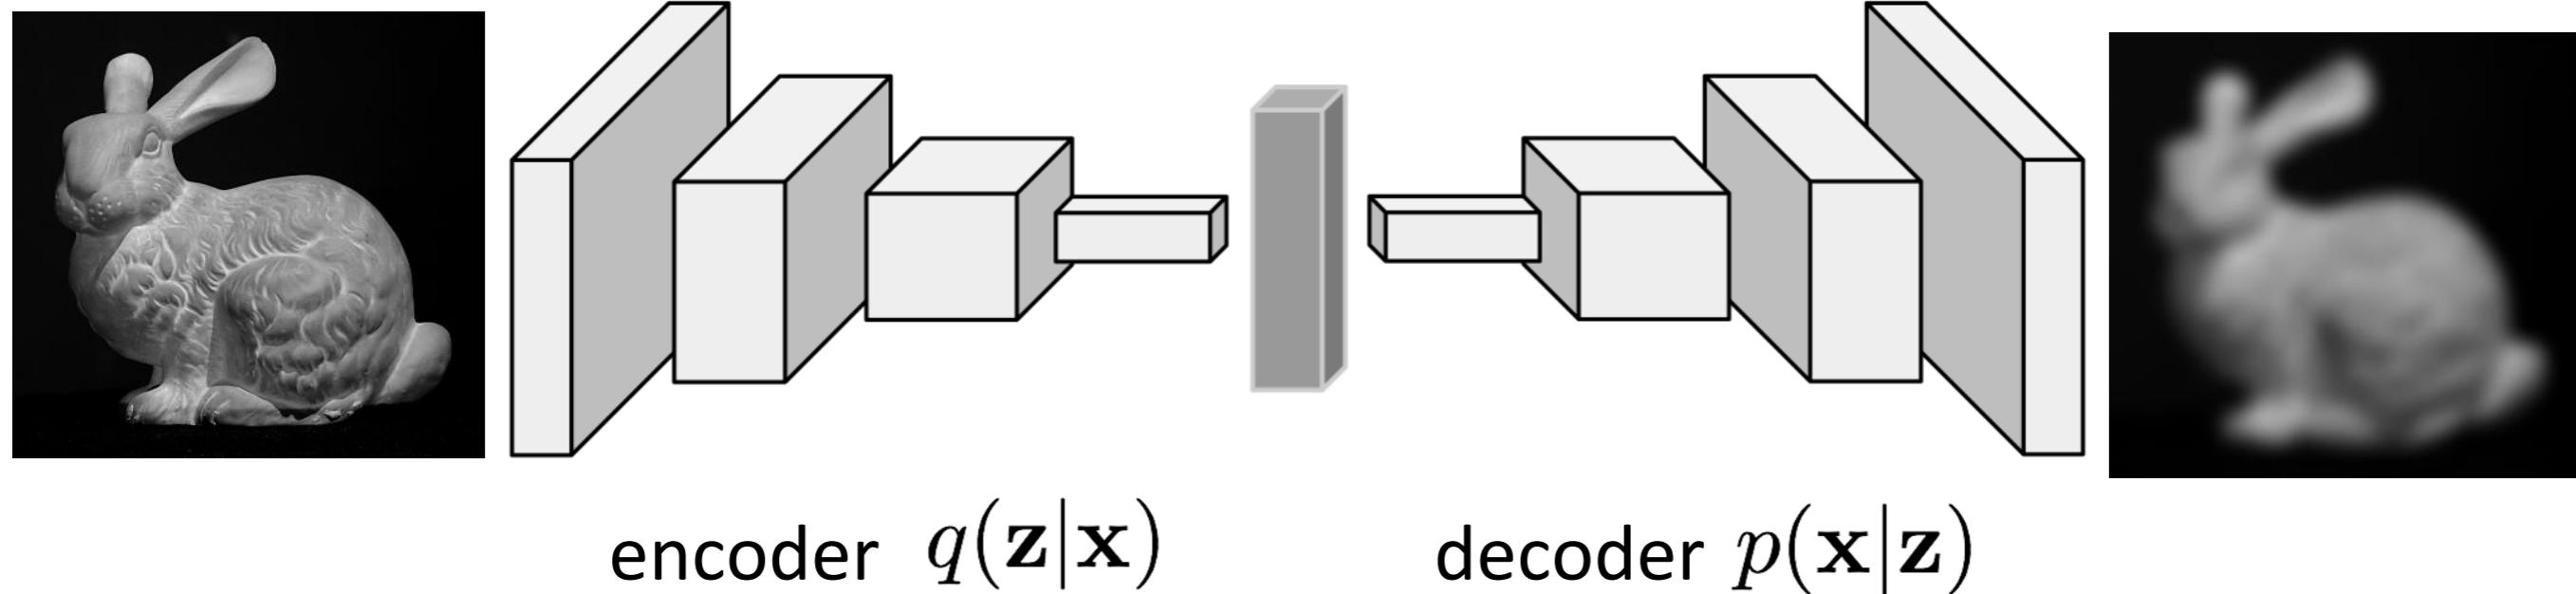
Figure 5: Random samples from learned generative models of MNIST for different dimensionalities of latent space.

Today

- Variational Autoencoders (VAE)
- **Generative Adversarial Networks (GANs)**

Variational Autoencoder

$$\mathbf{z} \sim p_{\text{model}} \quad \text{latent space}$$

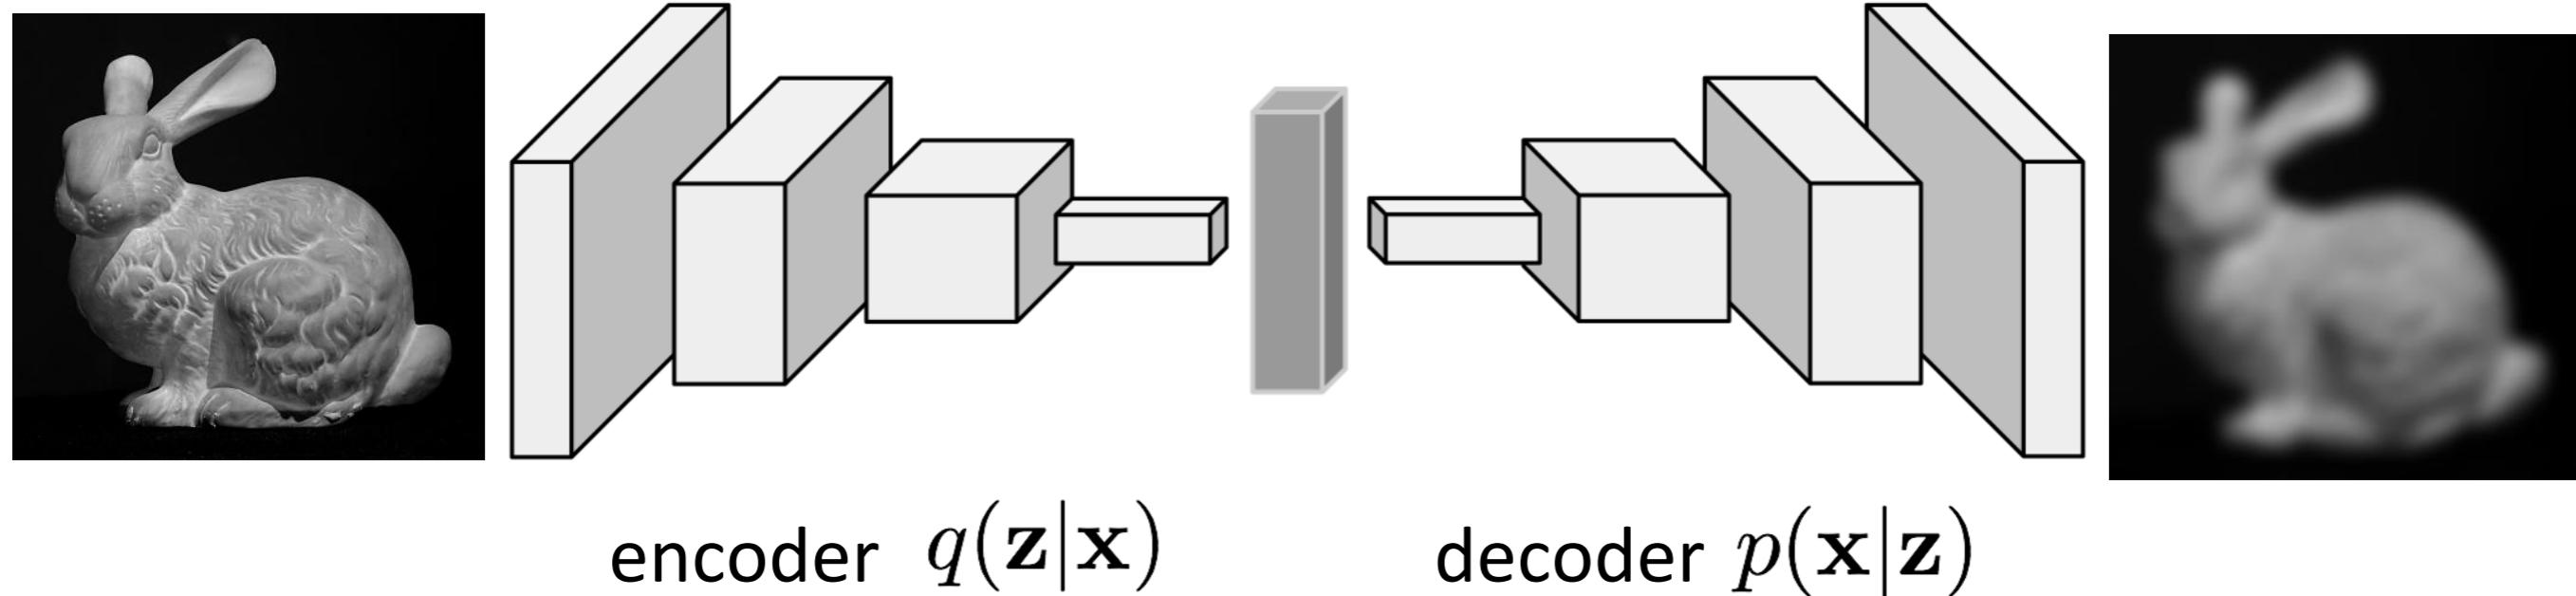


Maximize

$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$

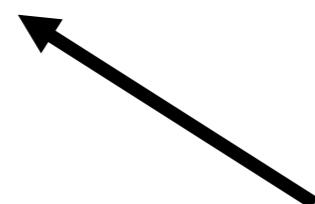
Variational Autoencoder

$$\mathbf{z} \sim p_{\text{model}} \quad \text{latent space}$$



Maximize

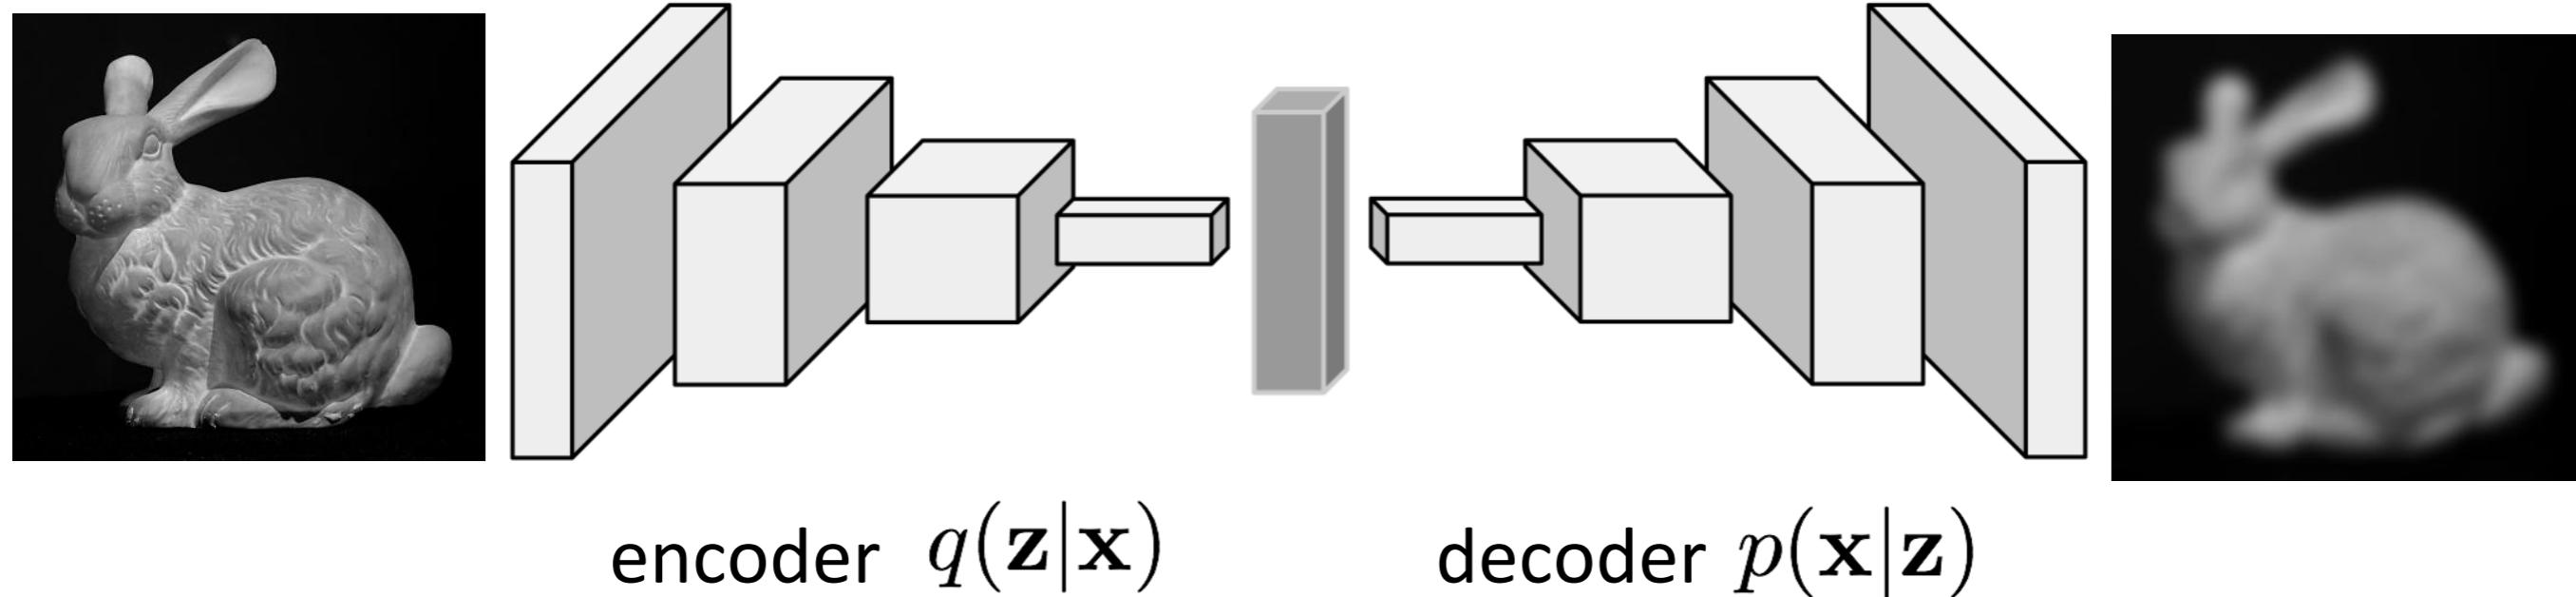
$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$



Typically a normal distribution \rightarrow blurry results

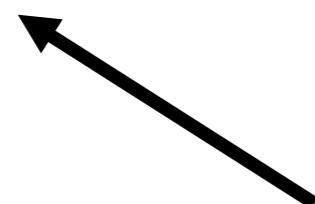
Variational Autoencoder

$$\mathbf{z} \sim p_{\text{model}} \quad \text{latent space}$$



Maximize

$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$



Better loss functions?

Typically a normal distribution → blurry results

The Need for a Better Loss Function



original



output 1



output 2

[Doersch, Tutorial on Variational Autoencoders, arXiv:1606.05908]

The Need for a Better Loss Function



original



output 1



output 2

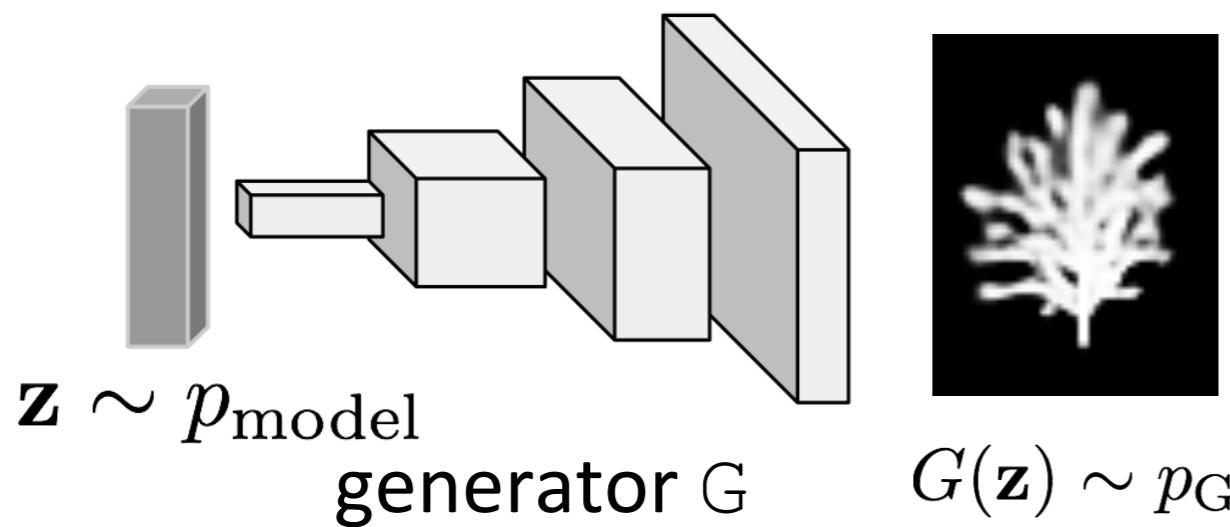
Output 2 better, but slightly shifted → larger loss under
Gaussian distribution

Generative Adversarial Networks (GANs)

Motivation: Learn loss function

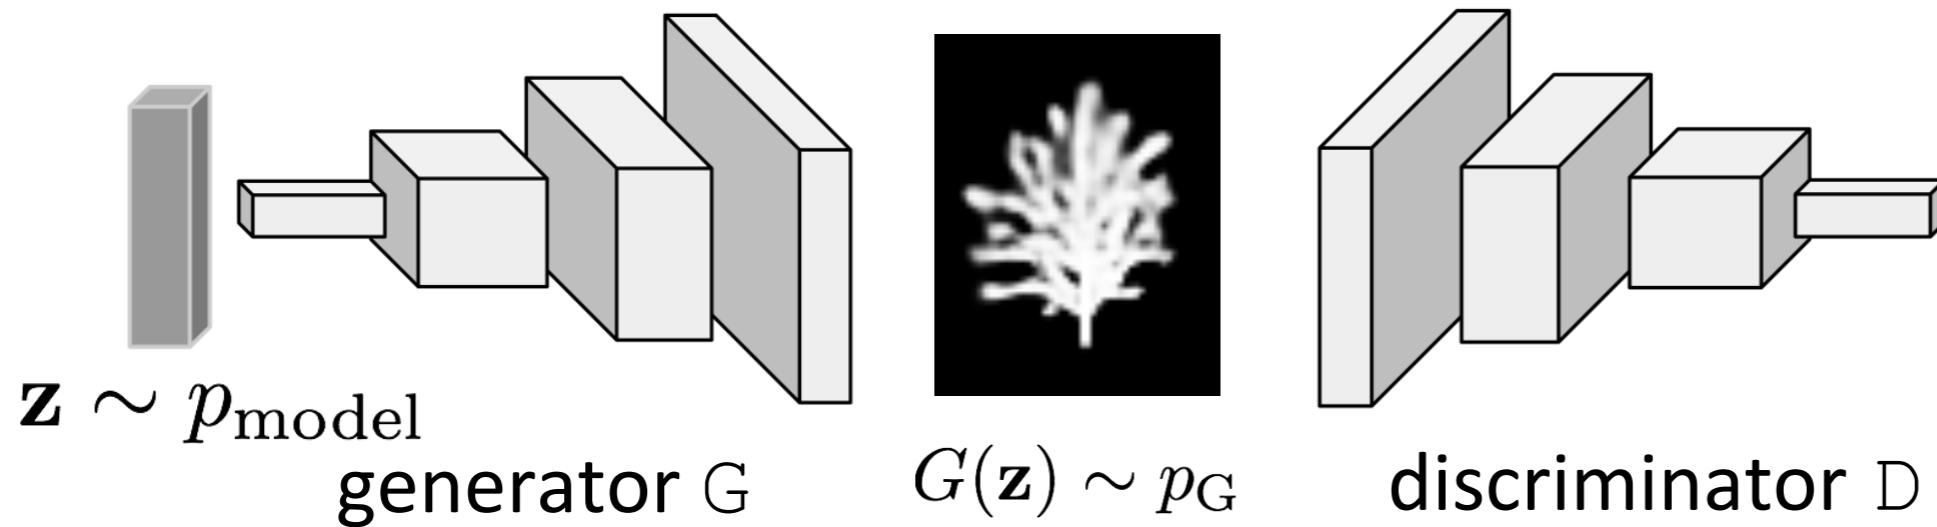
Generative Adversarial Networks (GANs)

Motivation: Learn loss function



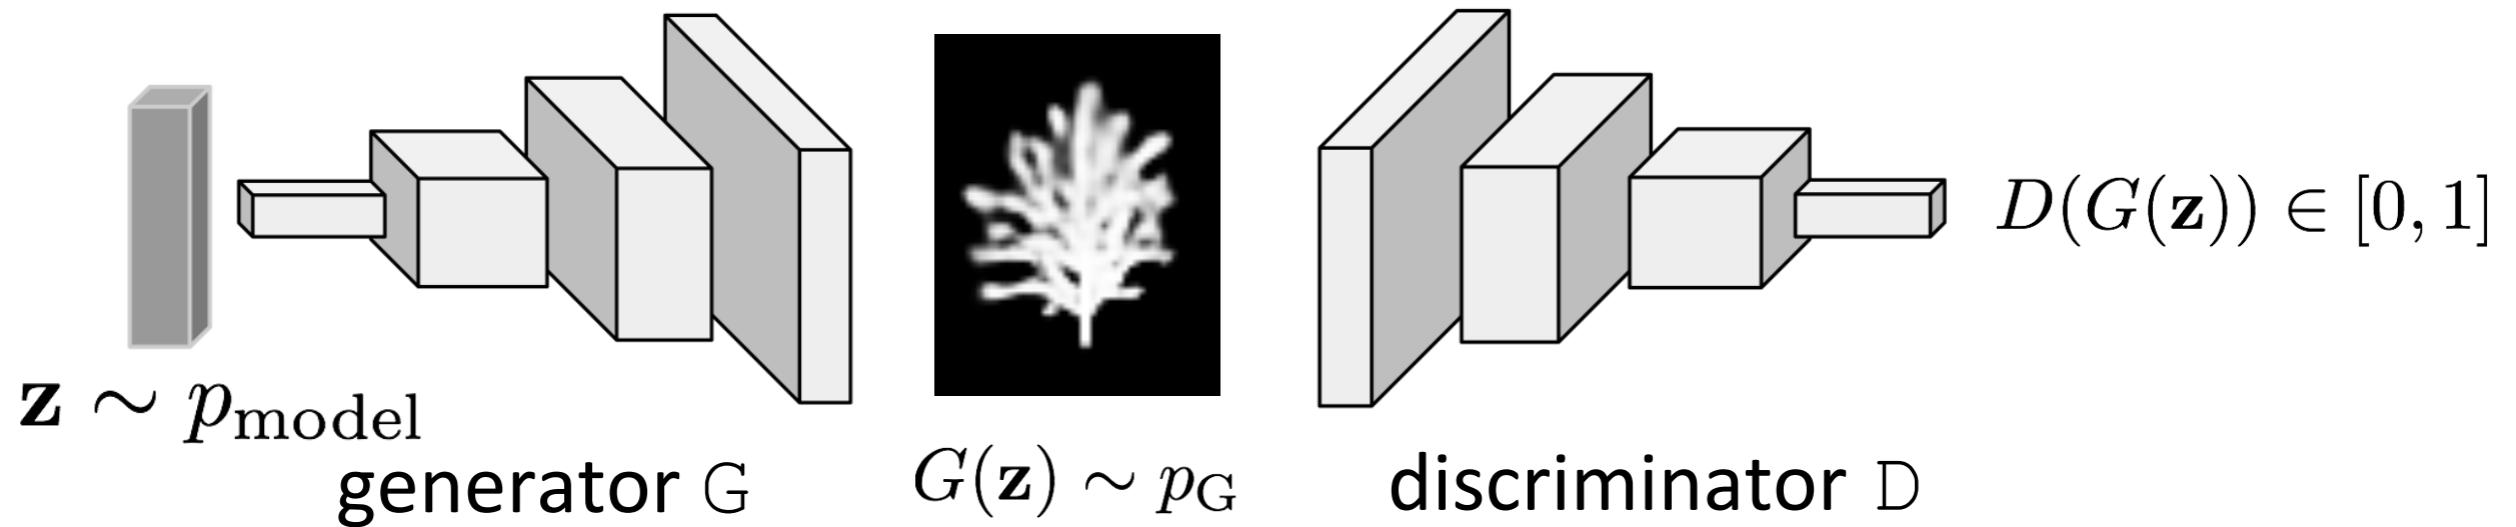
Generative Adversarial Networks (GANs)

Motivation: Learn loss function



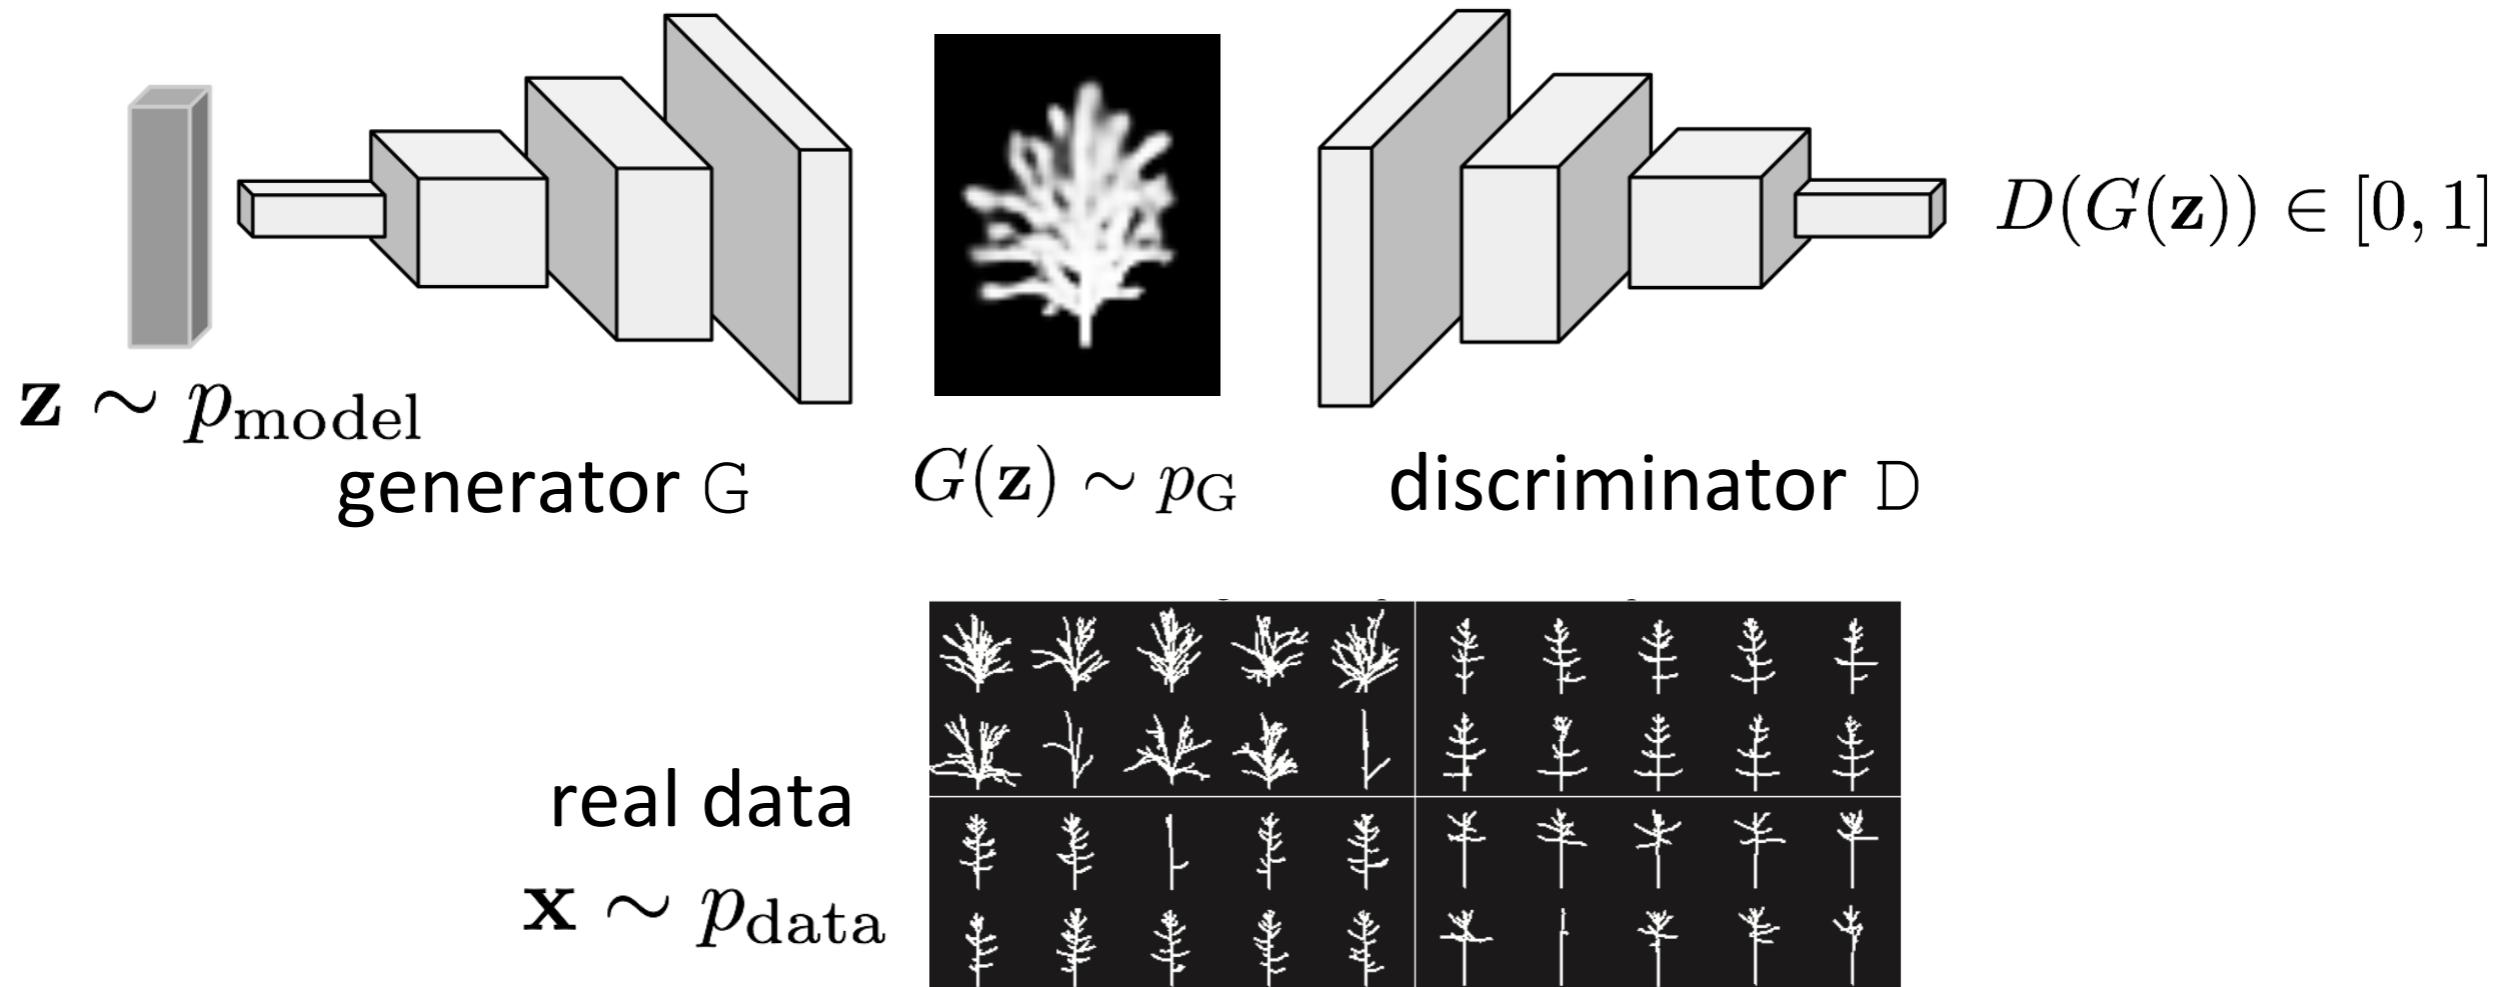
Generative Adversarial Networks (GANs)

Motivation: Learn loss function



Generative Adversarial Networks (GANs)

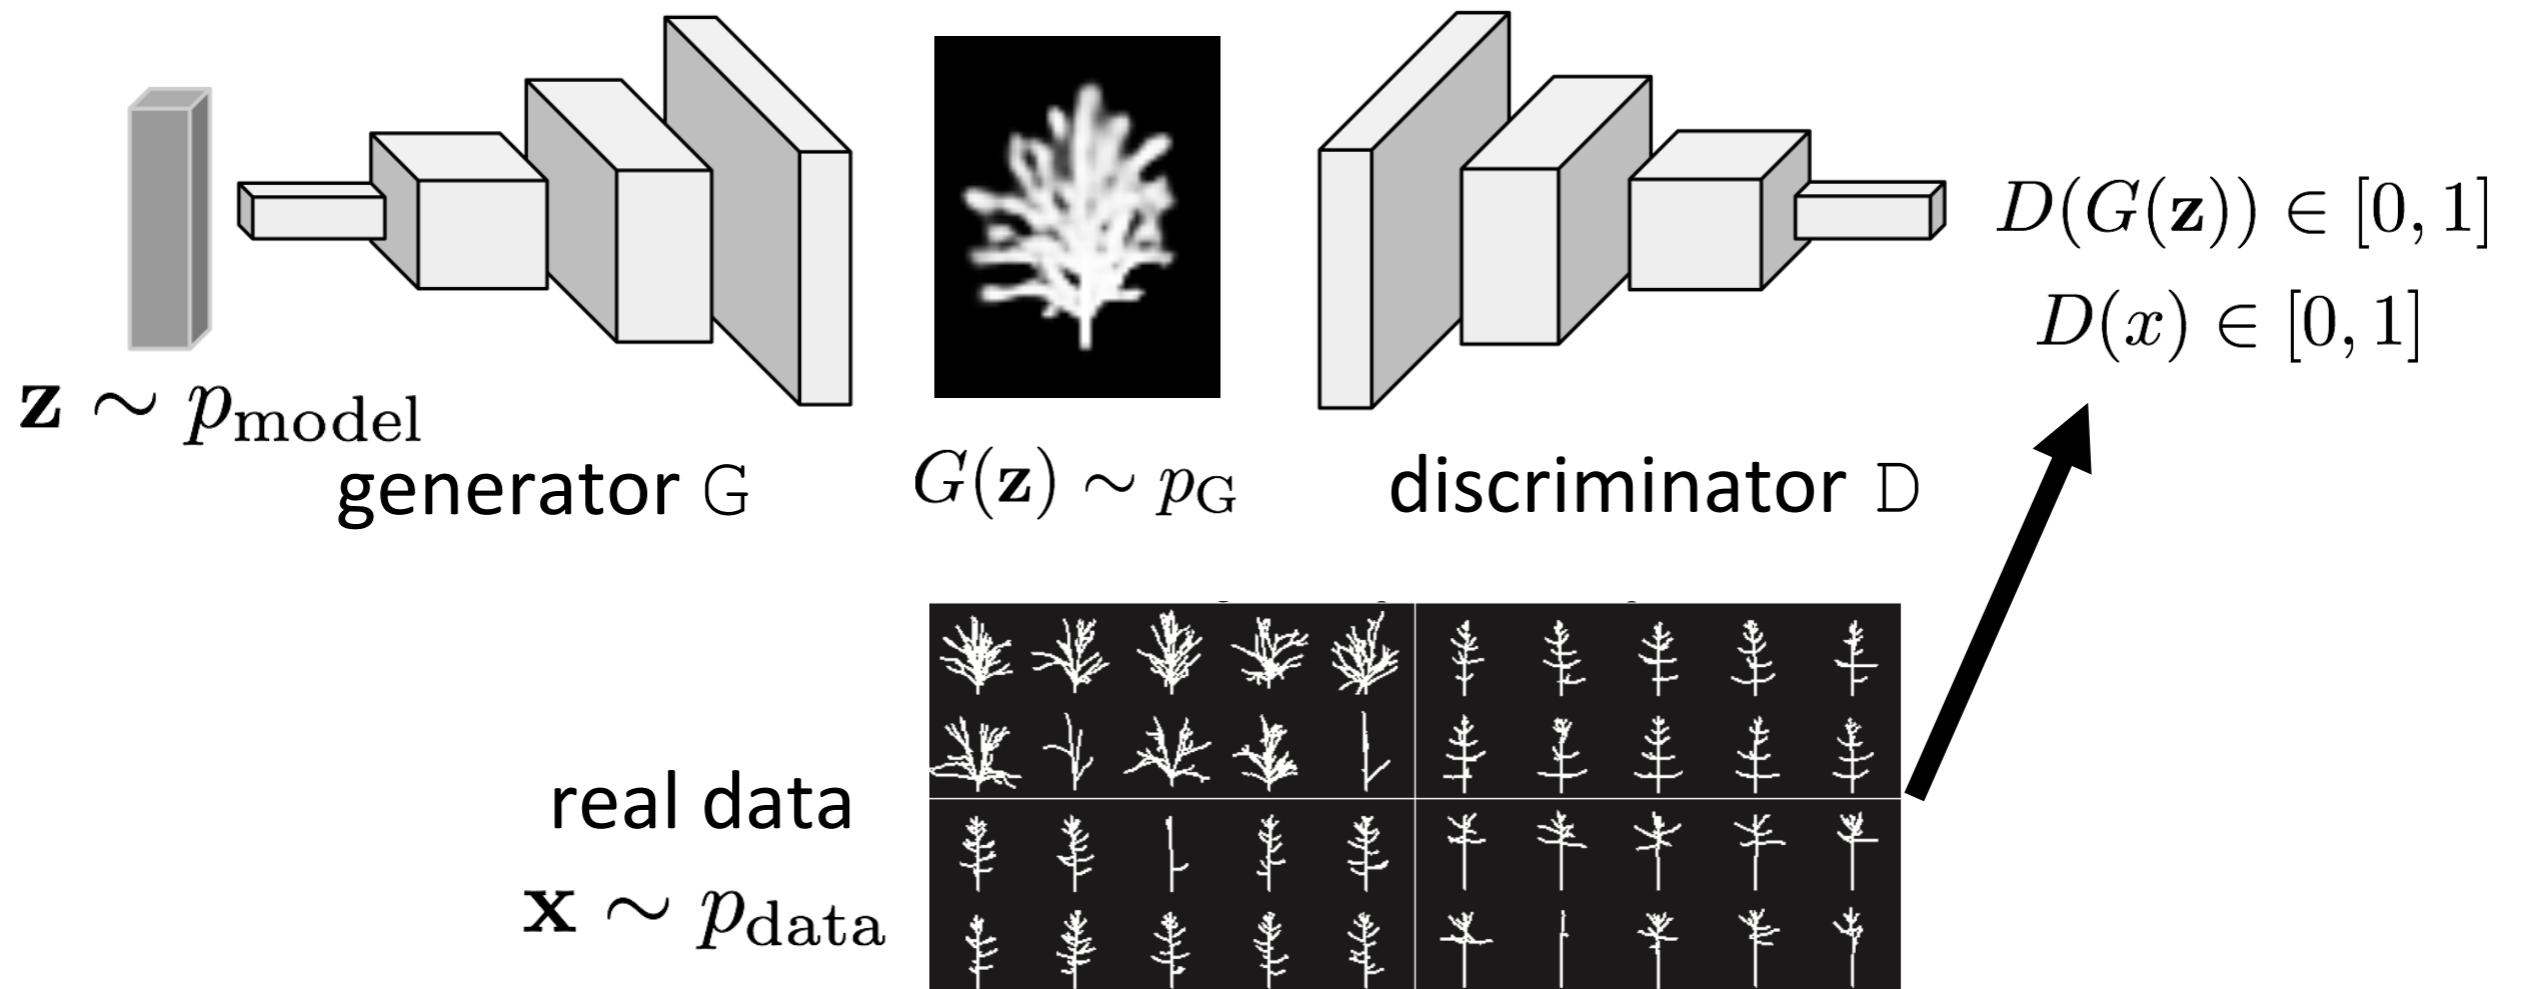
Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

Generative Adversarial Networks (GANs)

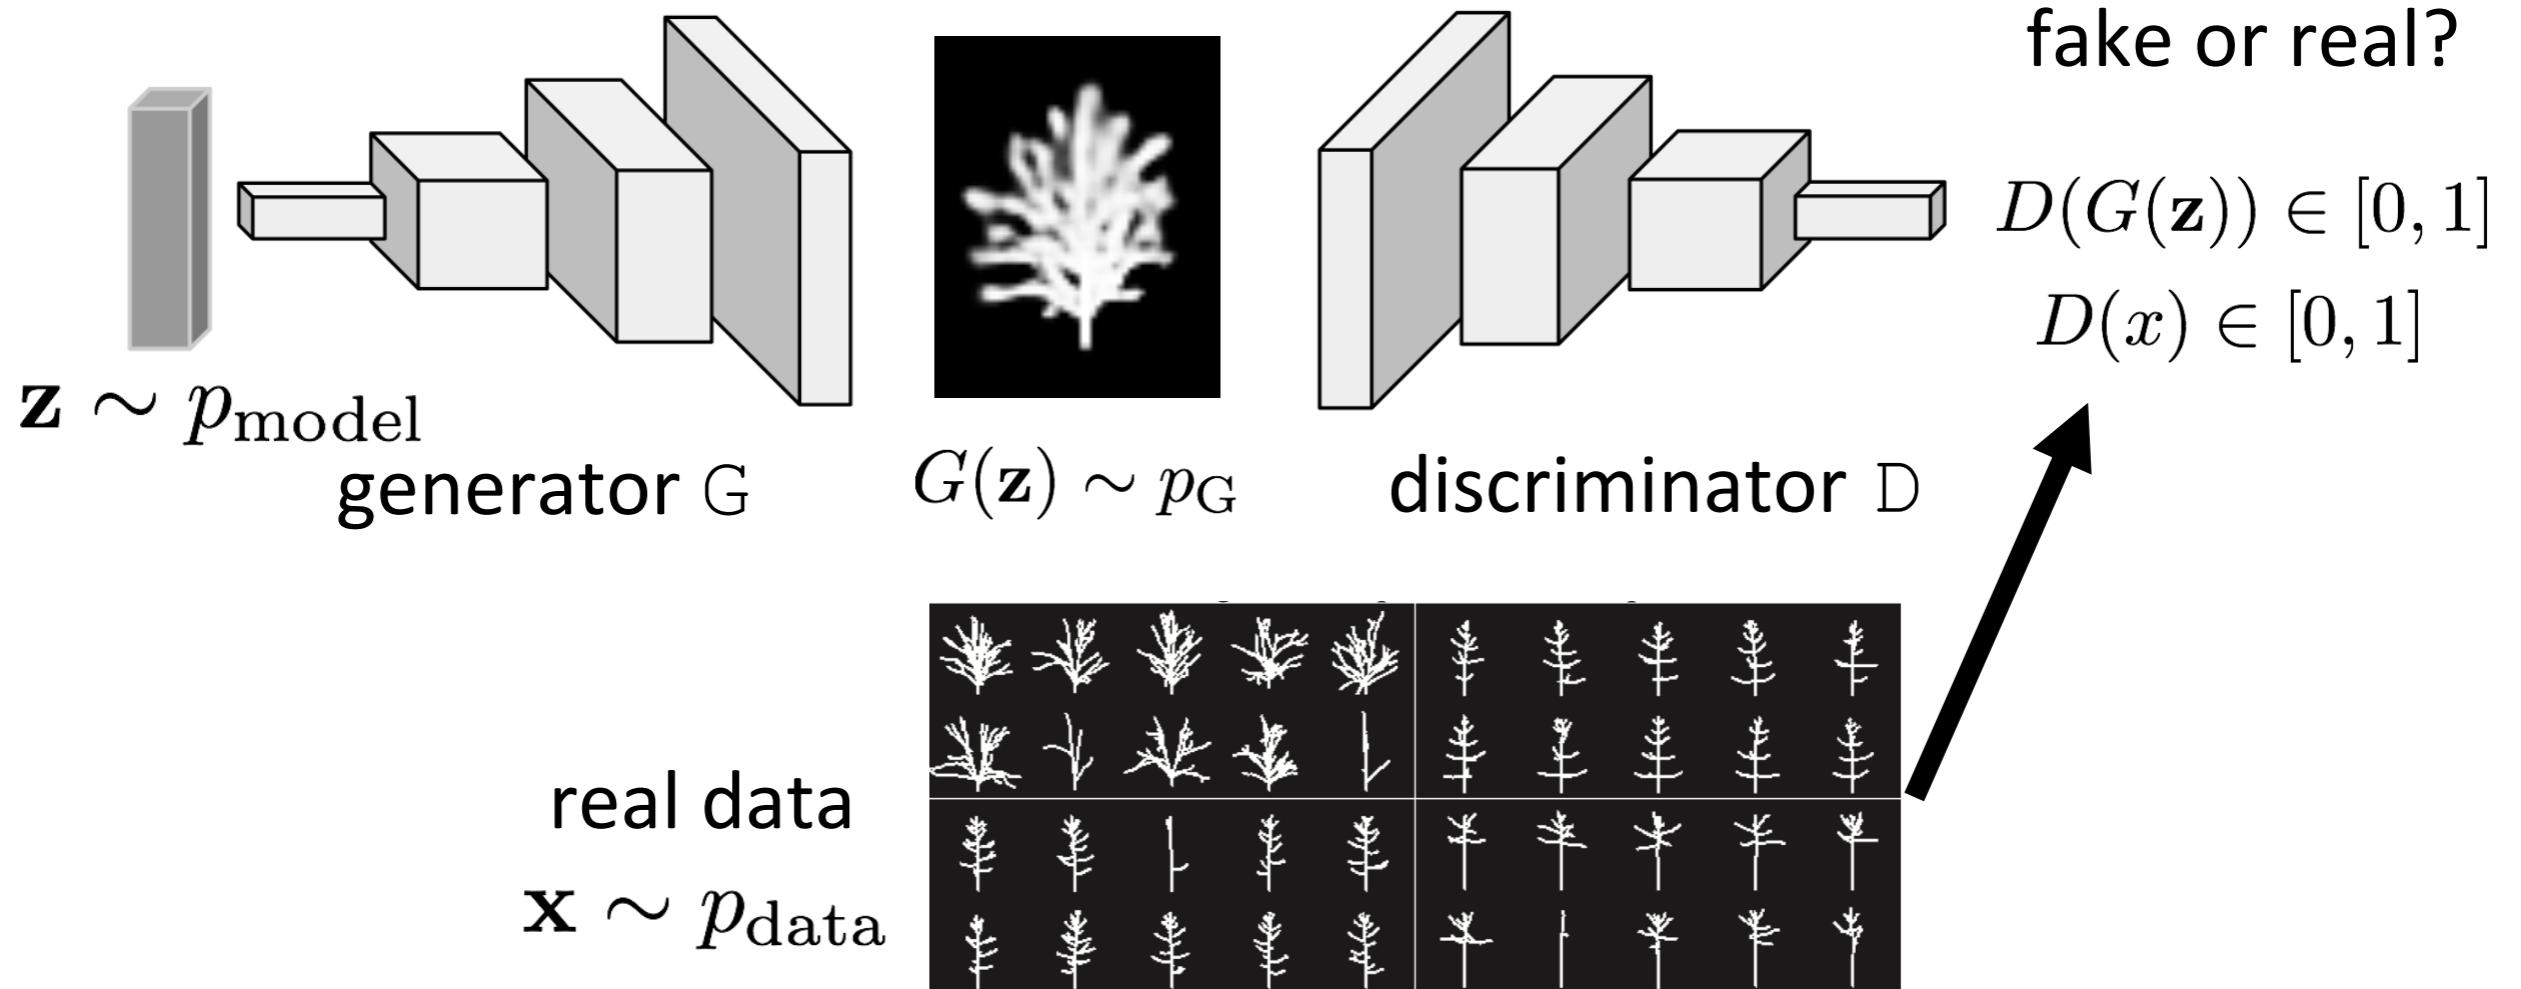
Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

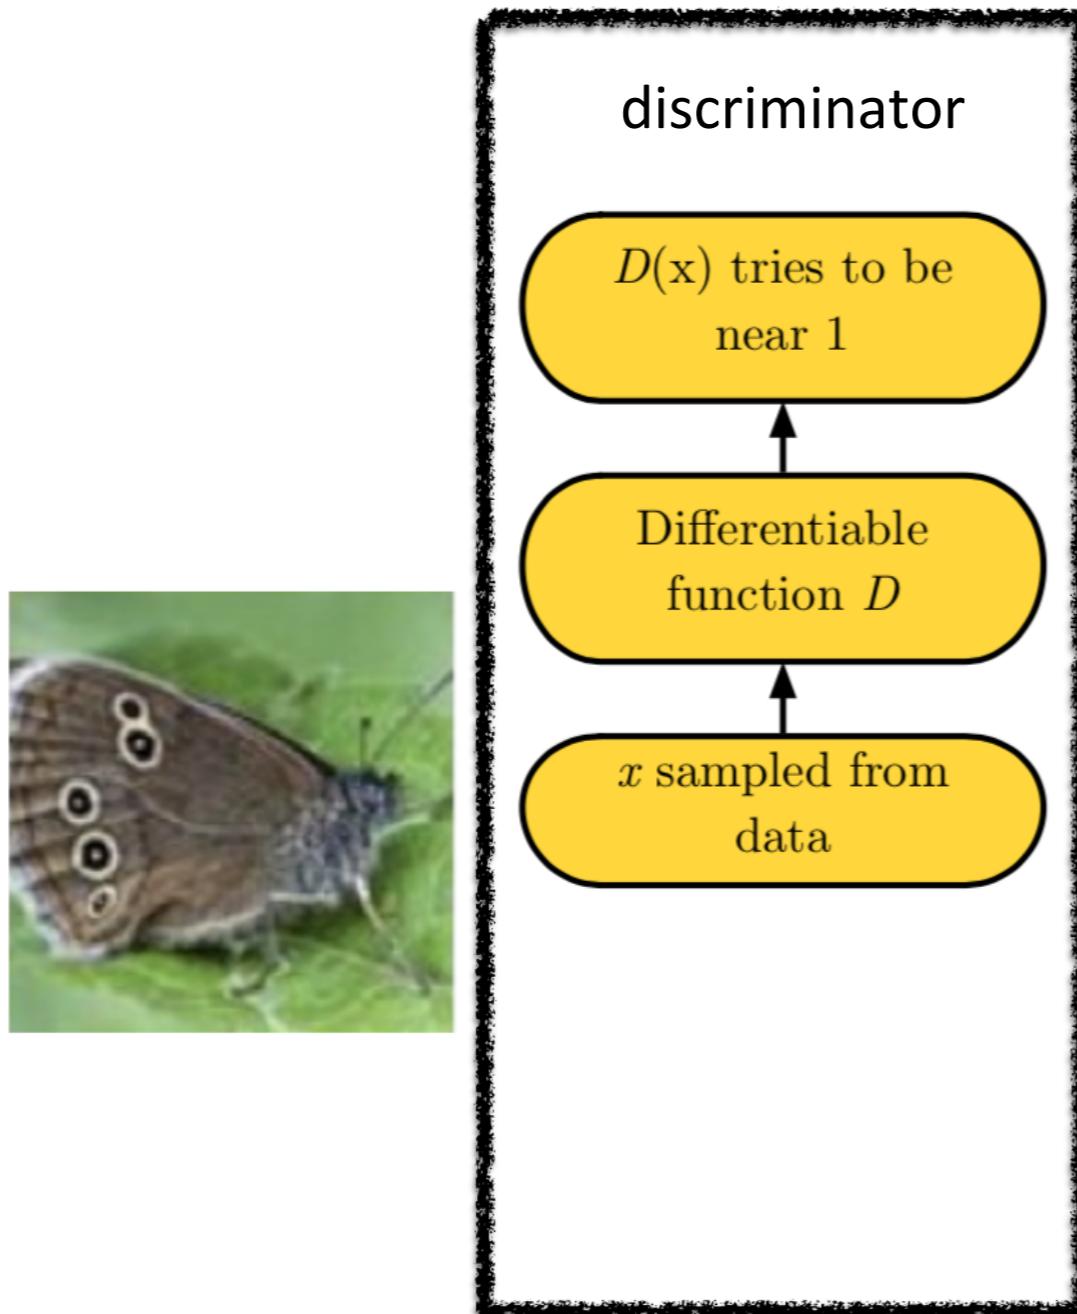
Generative Adversarial Networks (GANs)

Motivation: Learn loss function



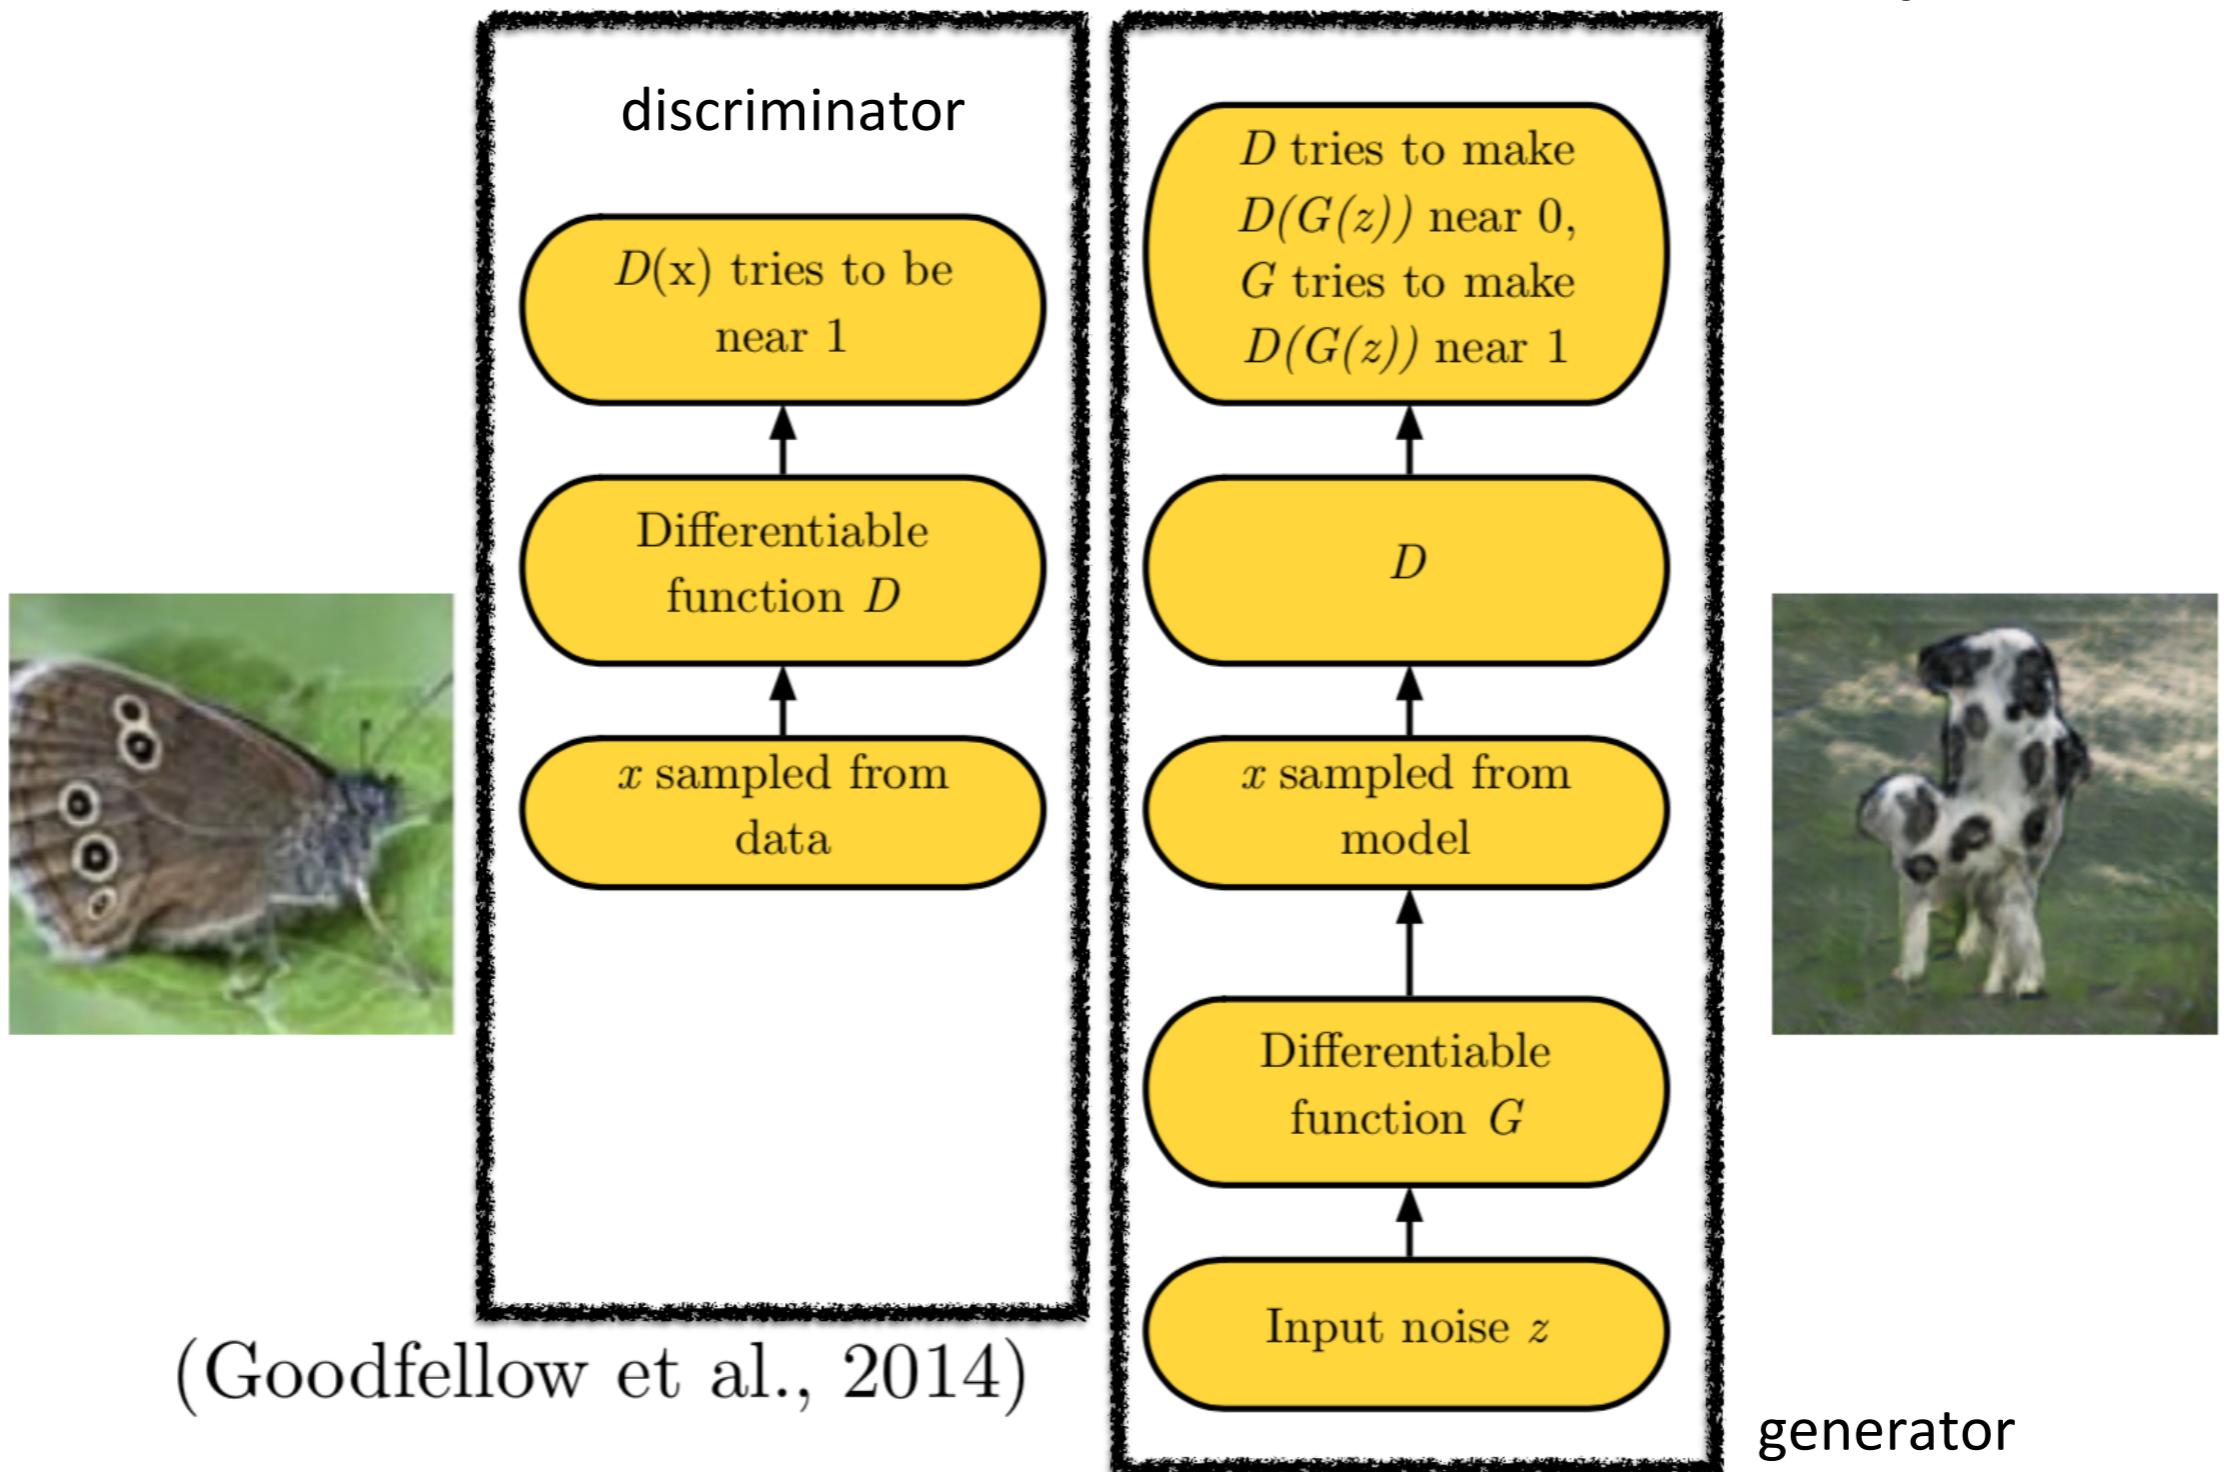
[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

Generative Adversarial Networks (GANs)



(Goodfellow et al., 2014)

Generative Adversarial Networks (GANs)



slide credit: Ian Goodfellow

Game Theoretic View

- Discriminator D tries to distinguish real and fake images

Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D , tries to make synthetic images look “as real as possible”

Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D , tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D , tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\max_D \quad \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D , tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\max_D \quad \underline{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))}$$

Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D , tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\min_G \max_D \quad \overbrace{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))}^{\text{red line}}$$

Game Theoretic View

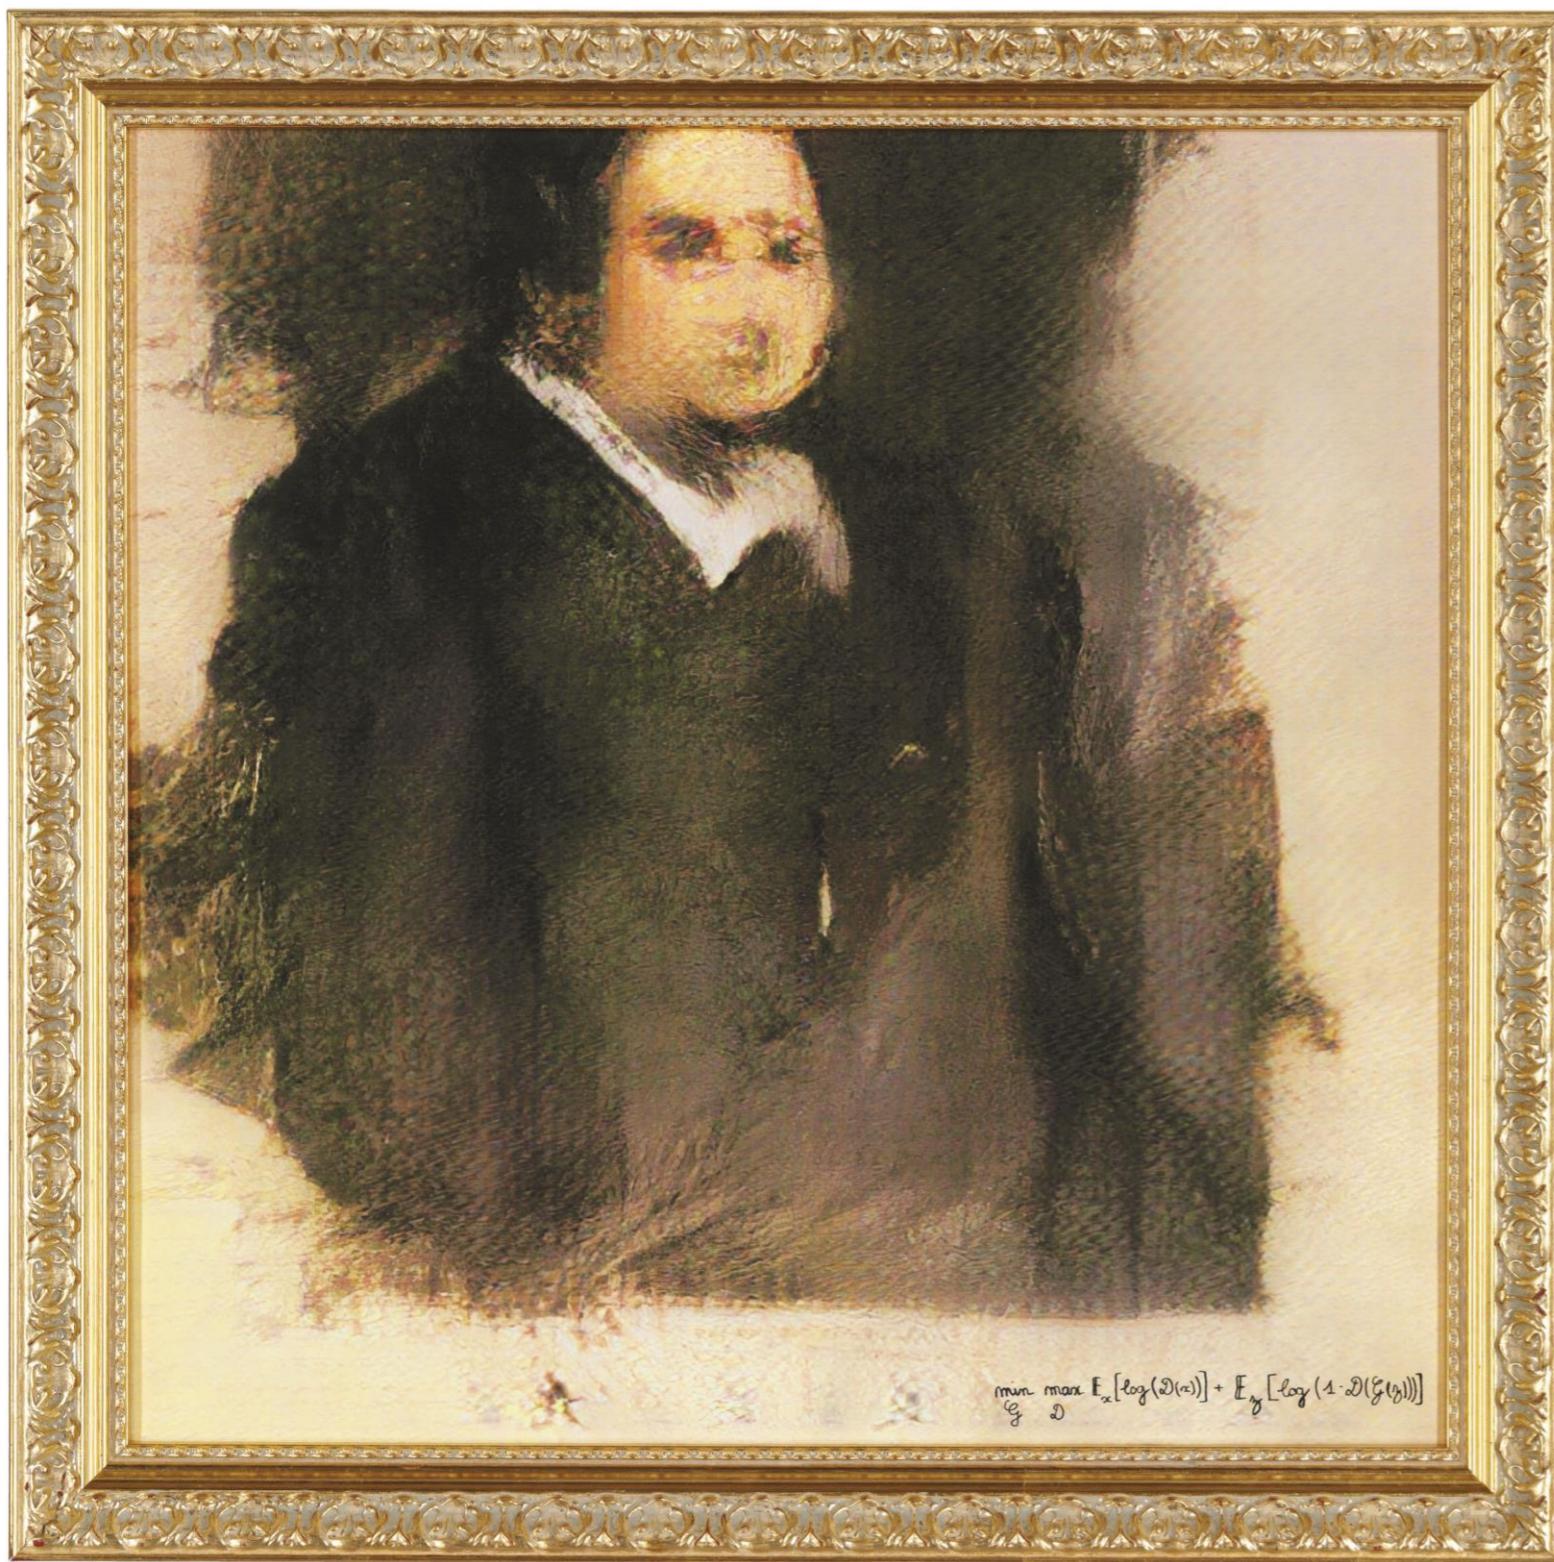
- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D , tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\min_G \max_D \quad \overbrace{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))}^{\text{red line}}$$

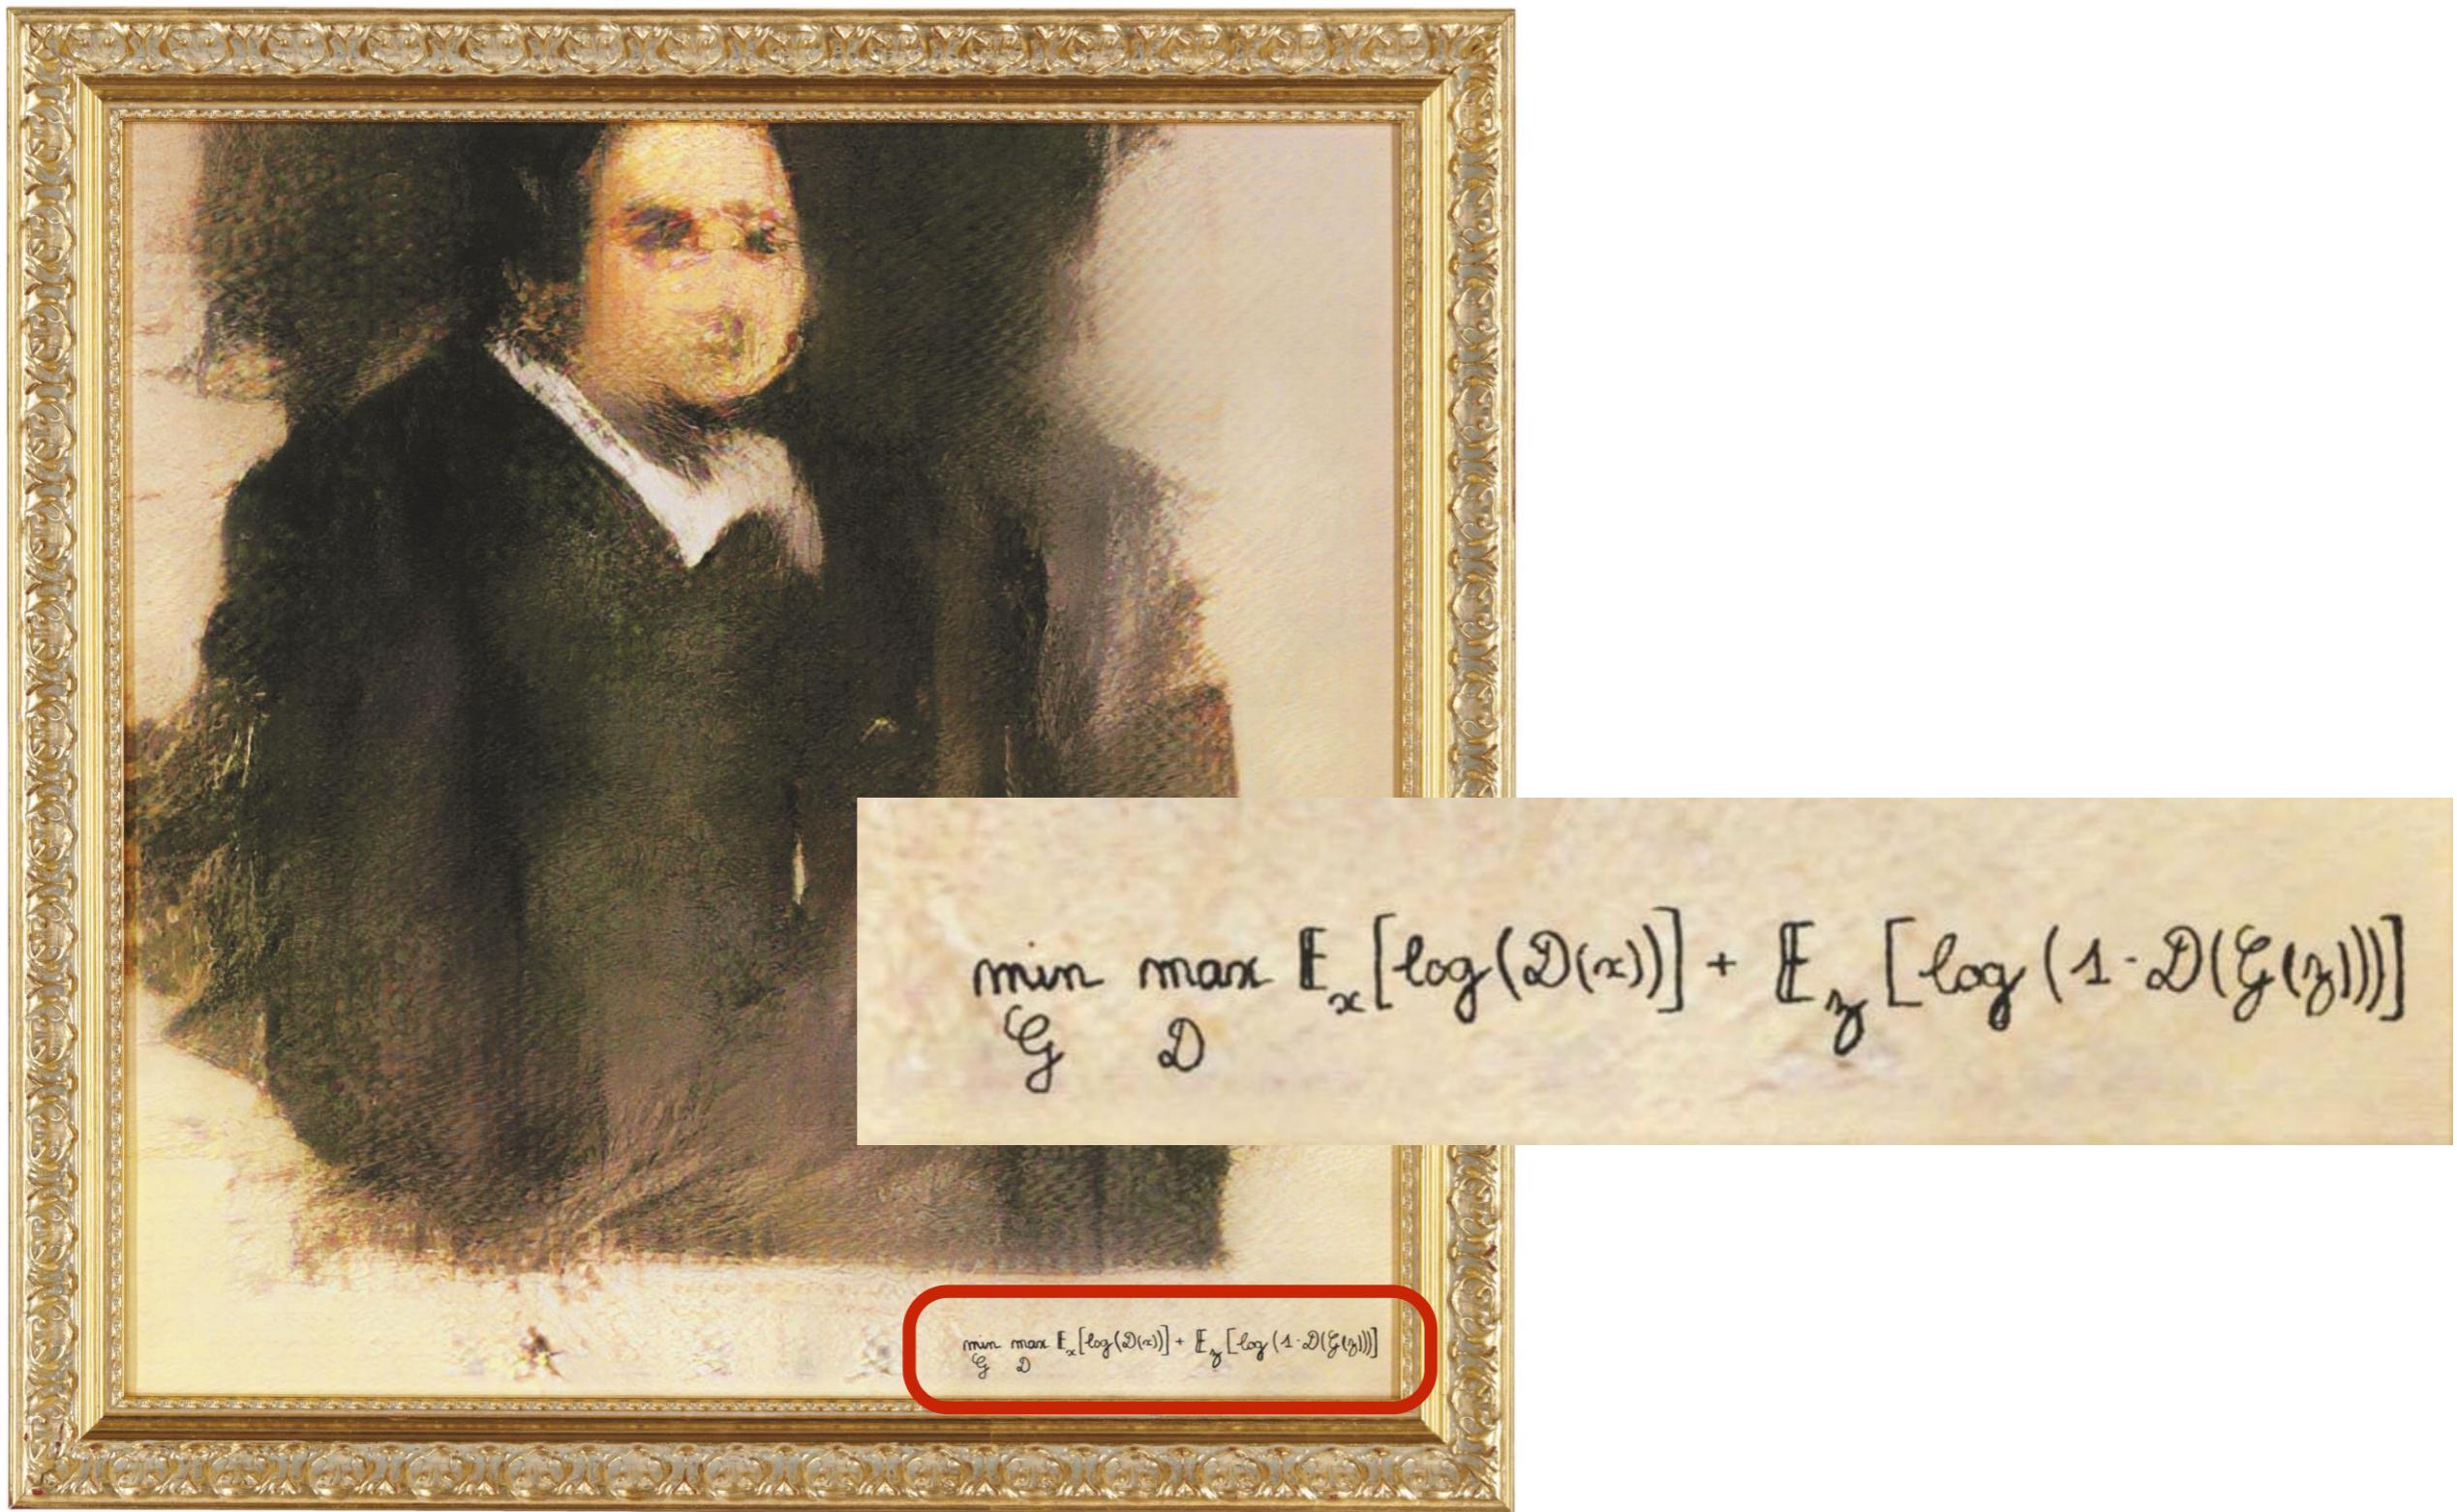
- Equilibrium / global optimum at $p_{\text{data}} = p_G$, i.e., when real and synthetic images are indistinguishable for D :

$$d(\mathbf{x}) = 1/2 \quad \forall \mathbf{x} \sim p_{\text{data}}, \mathbf{x} \sim p_G$$

Generating Art Revisited



Generating Art Revisited



Lessons Learned

- Main lessons from this lecture
 - Generative modeling via sampling from probability distributions
 - Generative networks learn mapping from simple probability distribution to target probability distribution
- Adversarial networks: Learn the loss function
- Next lecture: More generative modeling

Reading Materials

- Start: [Goodfellow, Bengio, Courville, Deep Learning, MIT press \(Chapter 20.10.4\)](#)
- Continue: Blog posts on [Wasserstein GANs](#), [Tutorial on VAEs](#)
- Advanced material:
 - ICCV 2017 GAN Tutorial ([website](#) with slides and videos)
 - CVPR 2018 GAN Tutorial ([website](#) with slides and videos)
 - Tensorflow implementations of [various GANs](#)
- GANs are an active research topic, be prepared to read a lot

Next Lecture

Jan. 20	Introduction, Linear classifiers and filtering	
Jan. 23	Filtering, gradients, scale	Lab 1
Jan. 27	Local features	
Jan. 30	Learning a classifier	
Feb. 3	Convolutional neural networks	Lab 2
Feb. 6	More convolutional neural networks	
Feb. 10	Robust model fitting and RANSAC	
Feb. 13	Image registration	Lab 3
Feb. 17	Camera Geometry	
Feb. 20	More camera geometry	
Feb. 24	Generative neural networks	Lab 4
Feb. 27	Generative neural networks	
Mar. 2	Visual Localization & Feature Learning	
Mar. 9	No lecture	