

SSY098 - Image Analysis

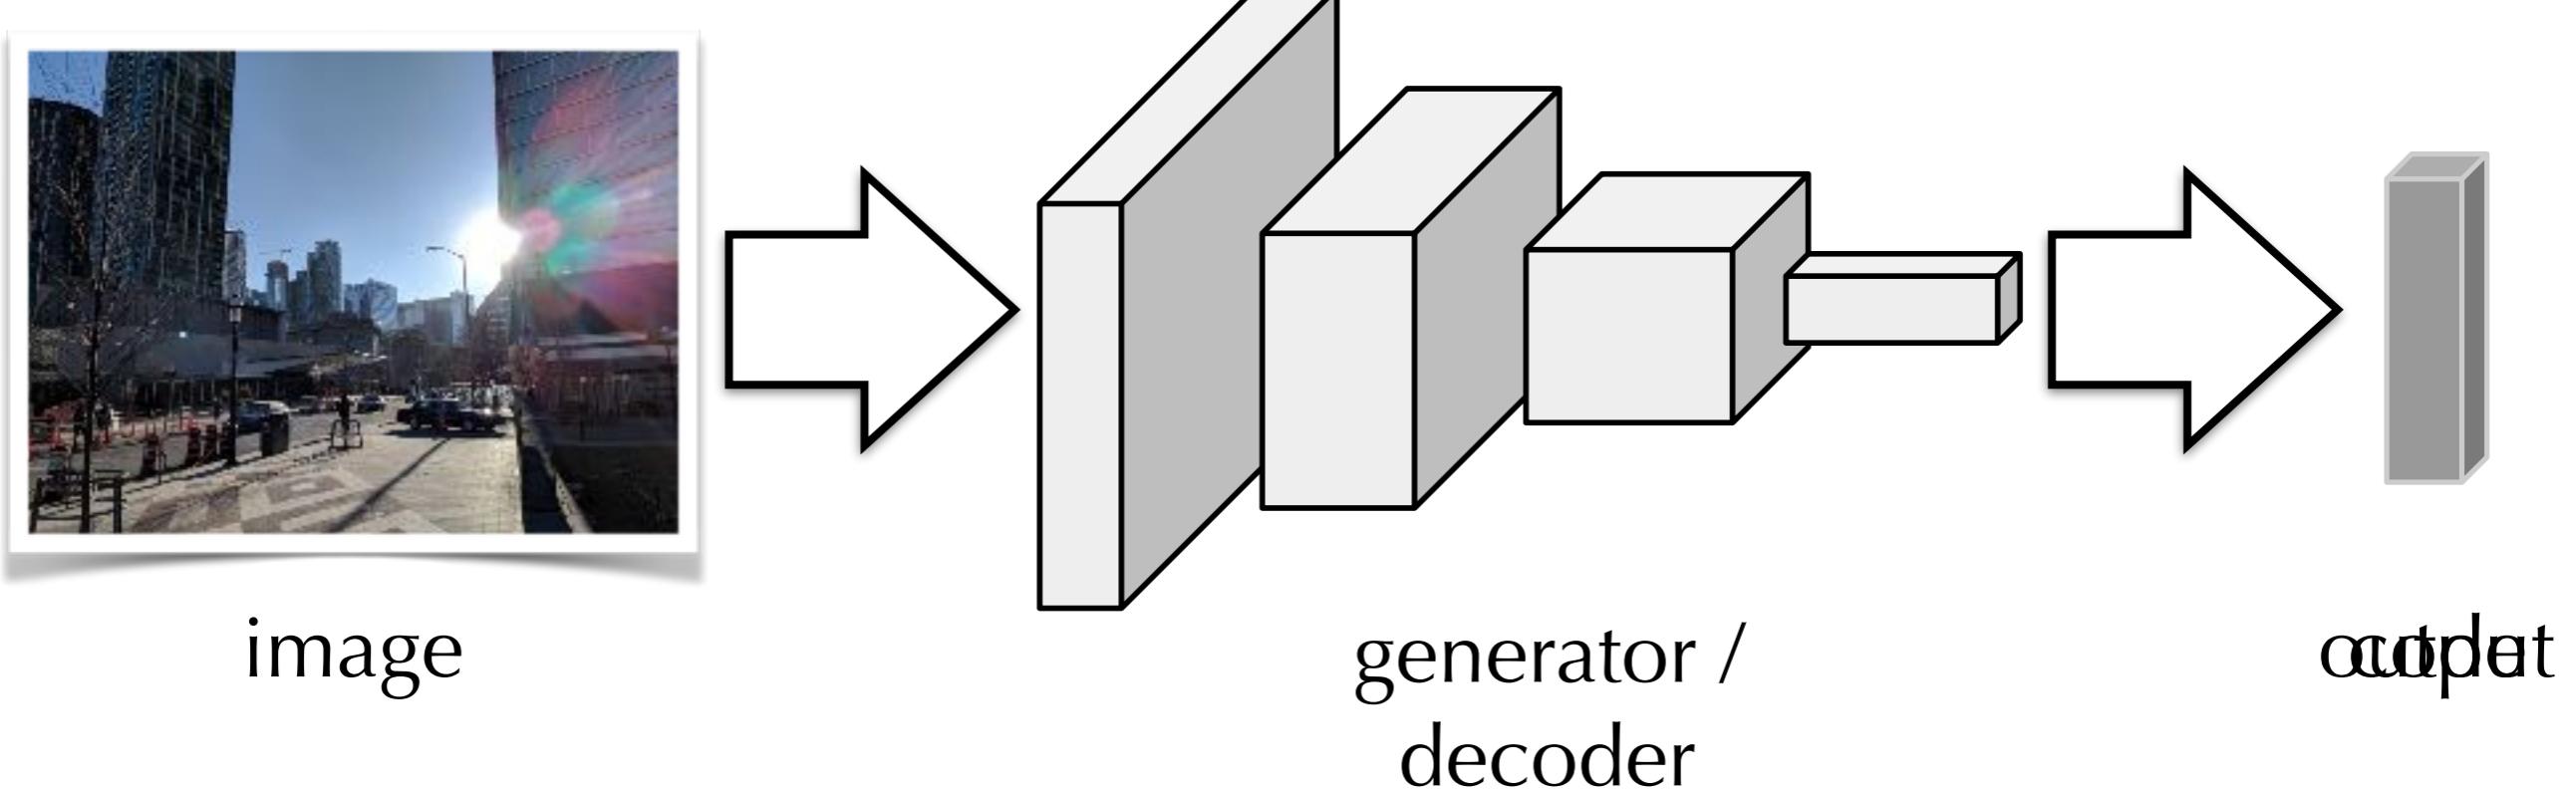
Lecture 12 - Generative Neural Networks

Torsten Sattler

Last Lecture

Jan. 20	Introduction, Linear classifiers and filtering	
Jan. 23	Filtering, gradients, scale	Lab 1
Jan. 27	Local features	
Jan. 30	Learning a classifier	
Feb. 3	Convolutional neural networks	Lab 2
Feb. 6	More convolutional neural networks	
Feb. 10	Robust model fitting and RANSAC	Lab 3
Feb. 13	Image registration	
Feb. 17	Camera Geometry	Lab 4
Feb. 20	More camera geometry	
Feb. 24	Generative neural networks	
Feb. 27	Generative neural networks	
Mar. 2	Visual Localization & Feature Learning	
Mar. 9	No lecture	

Generative Neural Networks





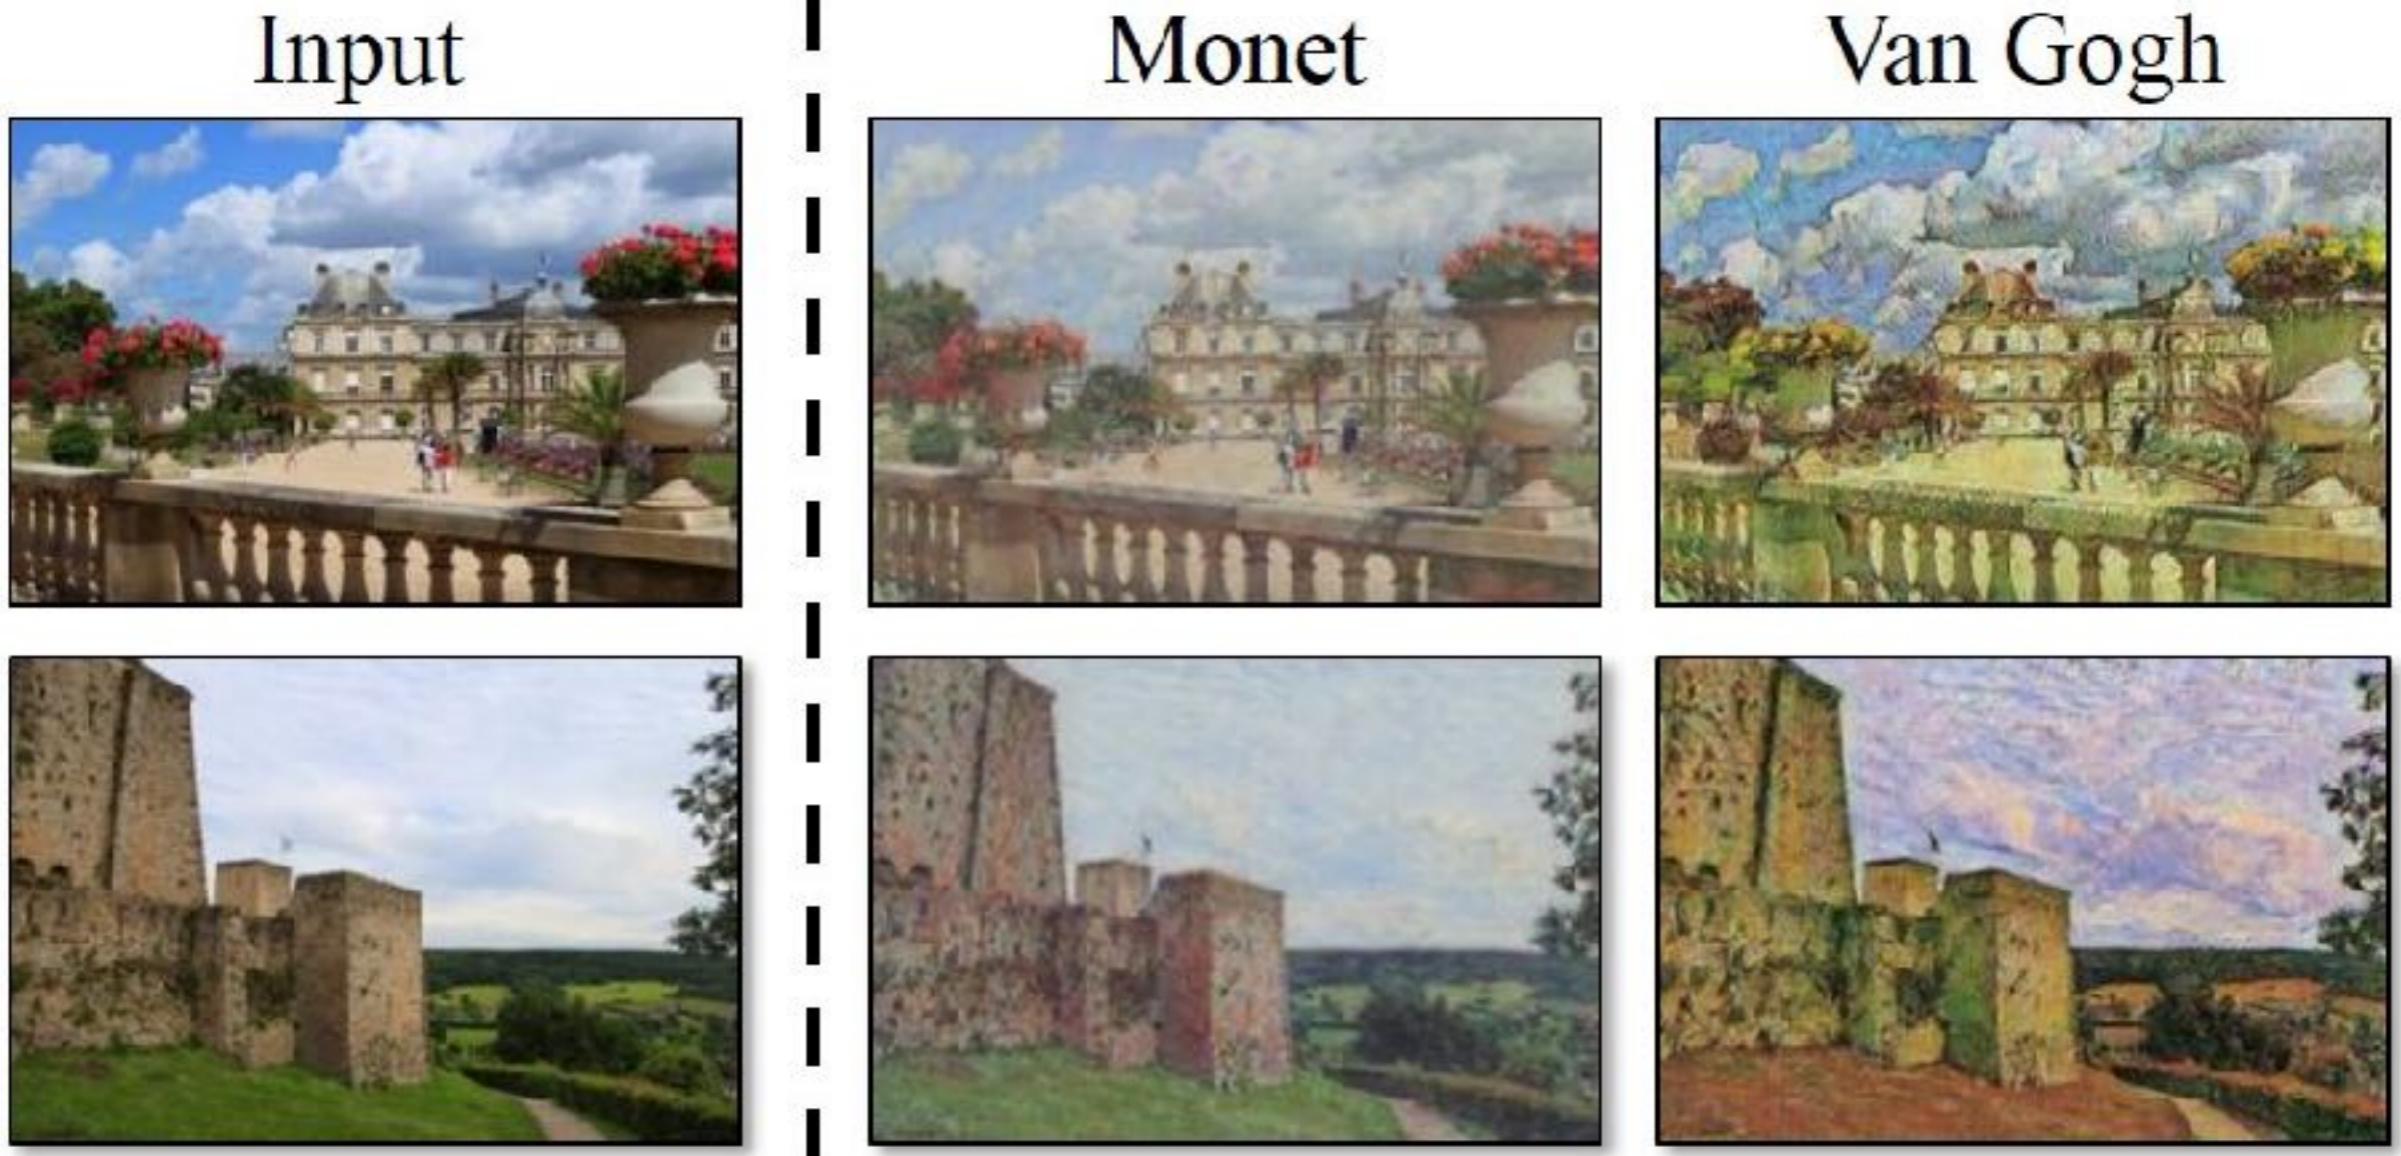
Portrait of Edmond Belamy, 2018. Sold for \$432,500 on 25 October at Christie's in New York. Image © Obvious

Image-to-Image Translation



[Zhu et al., Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks,
ICCV 2017]

Artistic Style Transfer



[Zhu et al., Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks,
ICCV 2017]

Training Data Generation

Input sunny image



AI-generated rainy image



[Liu et al., Unsupervised Image-to-Image Translation Networks, NIPS 2017]

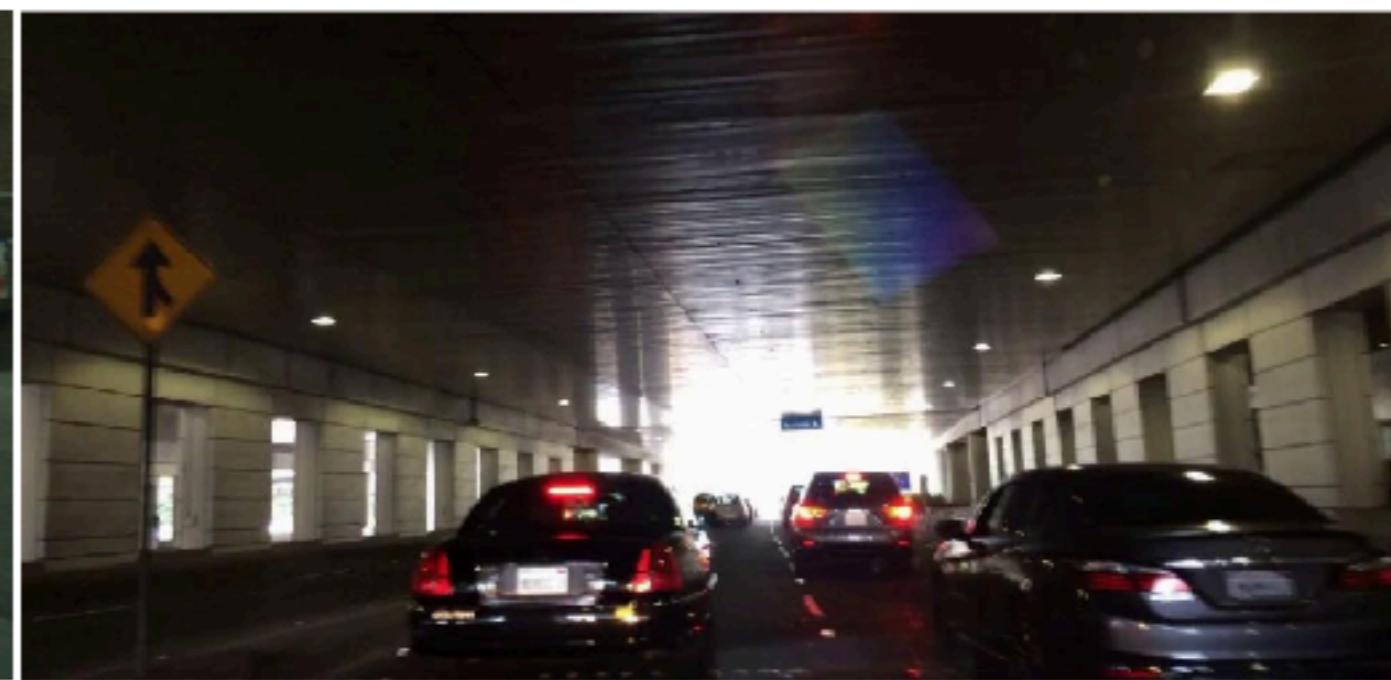
Domain Adaptation

Train



CityScapes (Germany)

Test



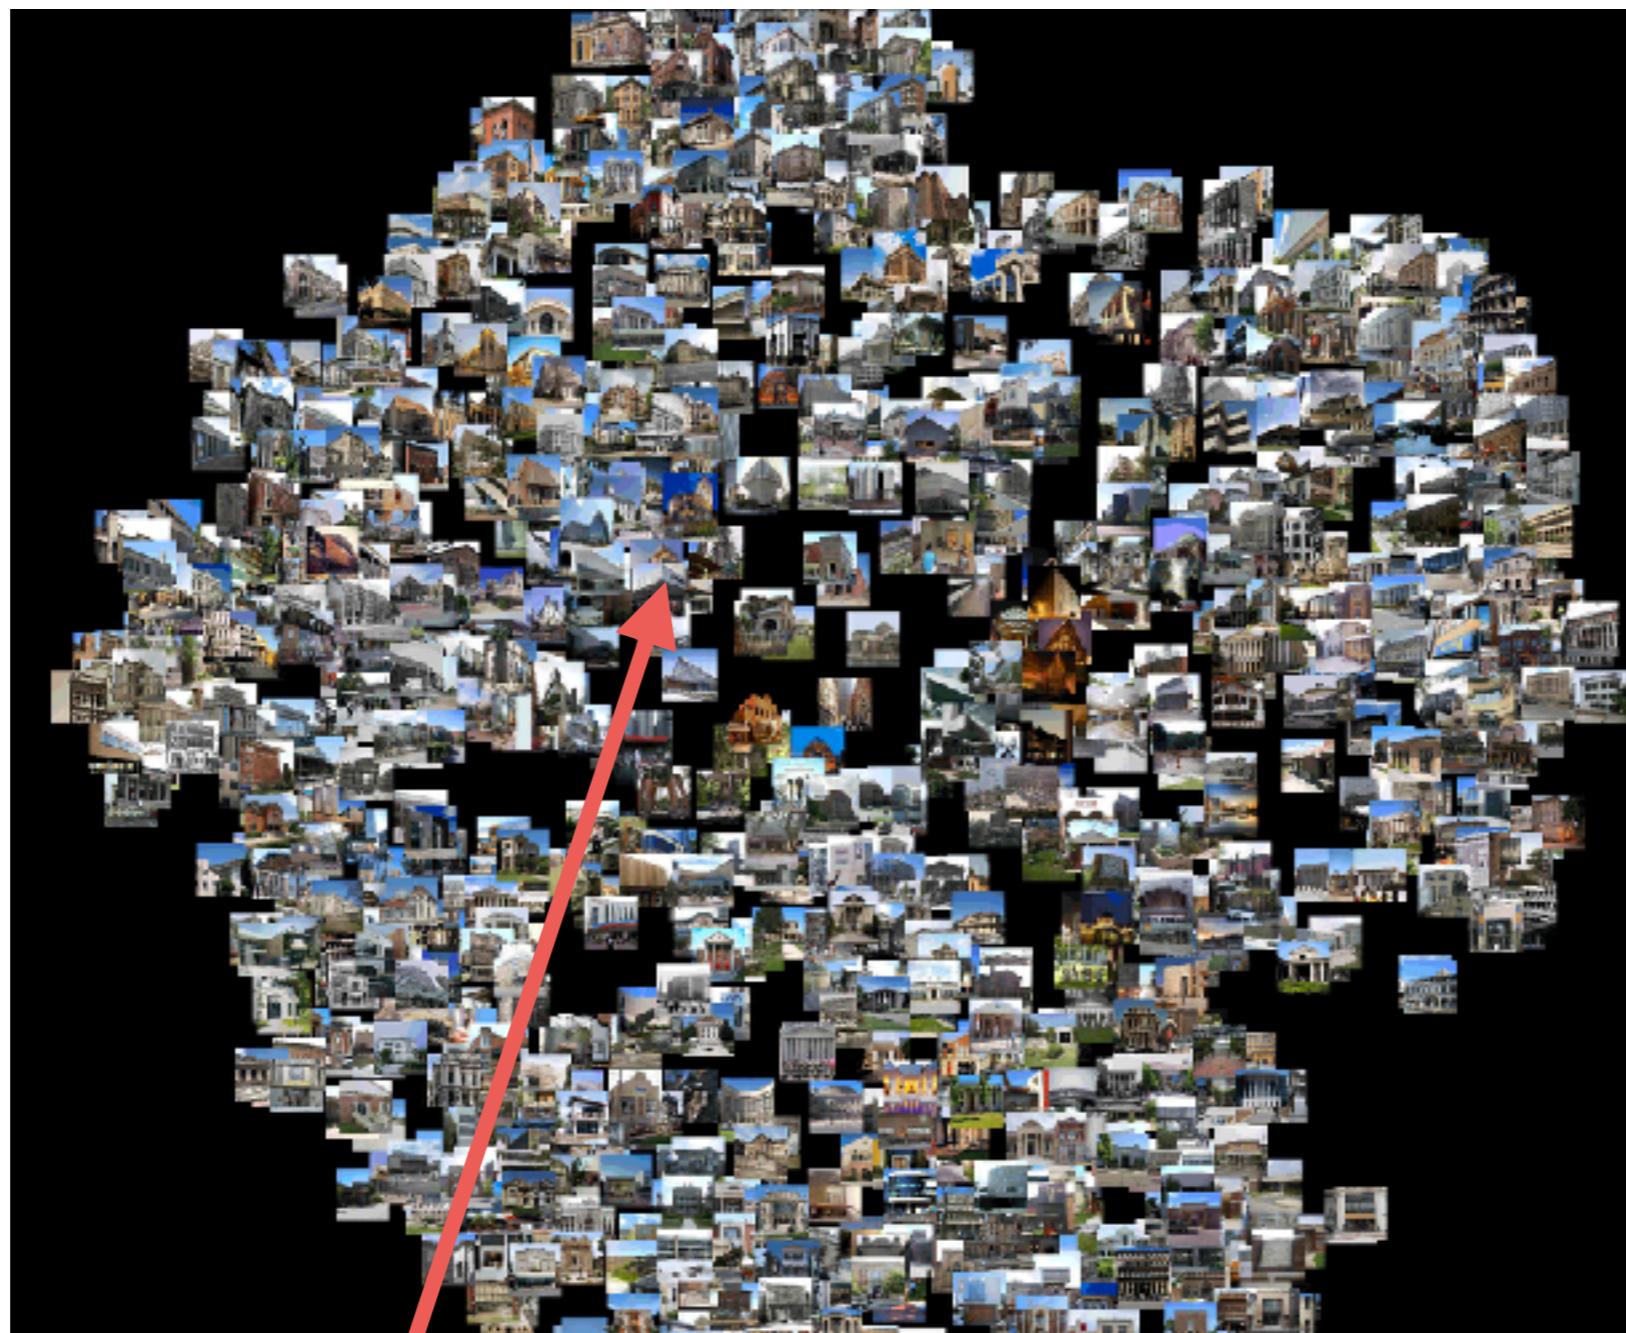
San Francisco

Source domain:
Plenty of labelled data

Target domain:
Little to no labelled data

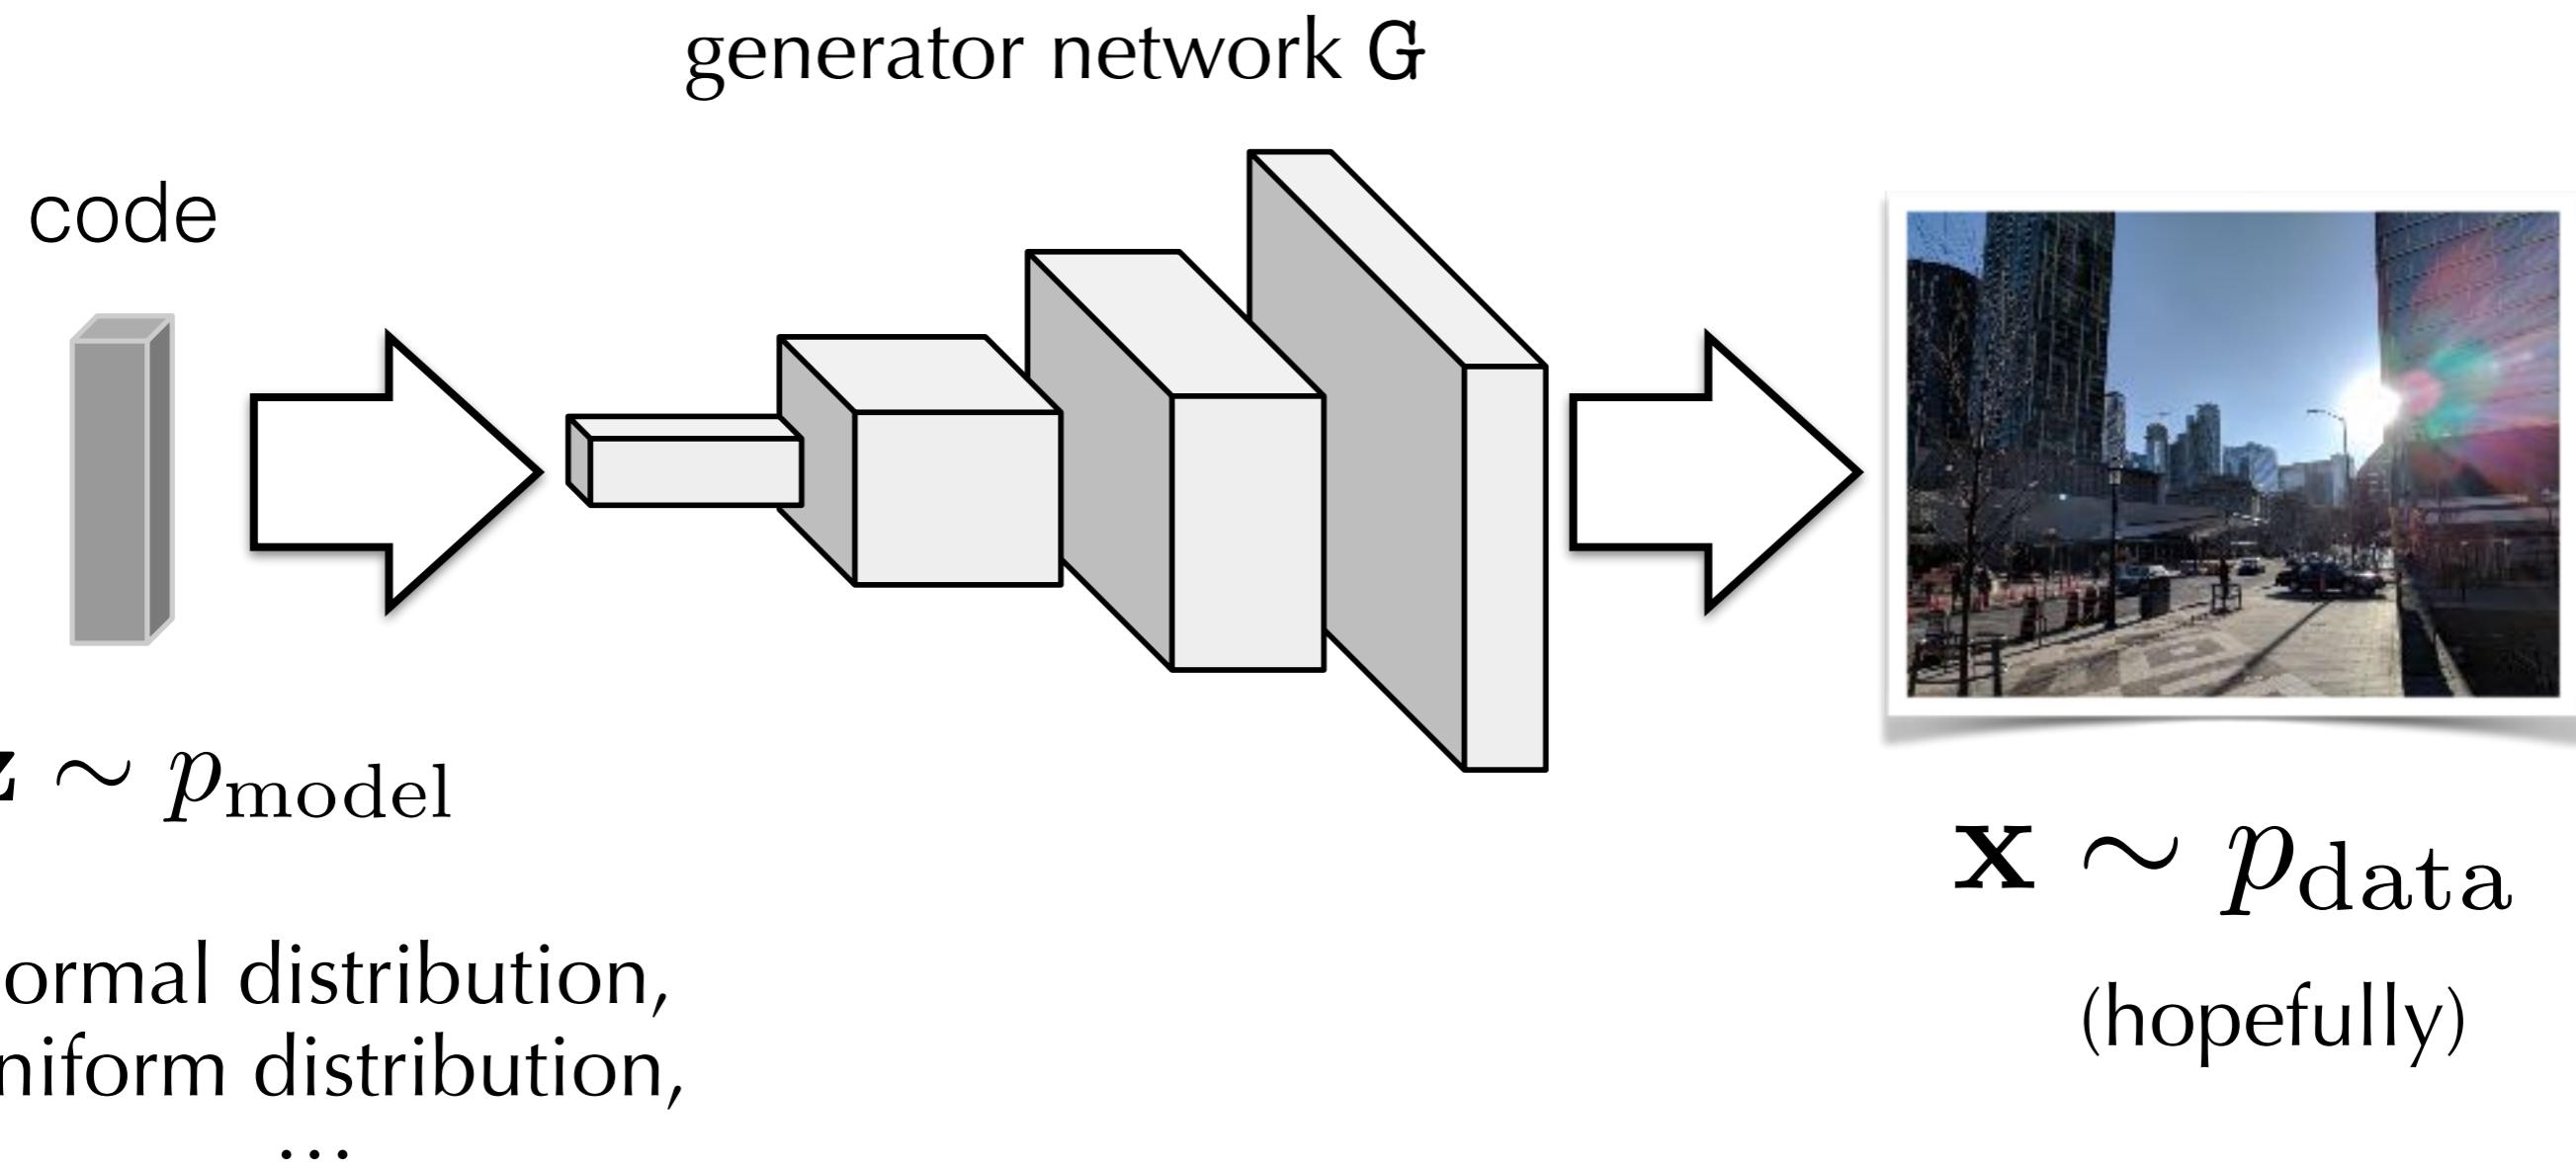
slide credit: Judy Hoffman

Image Generation as Sampling Process

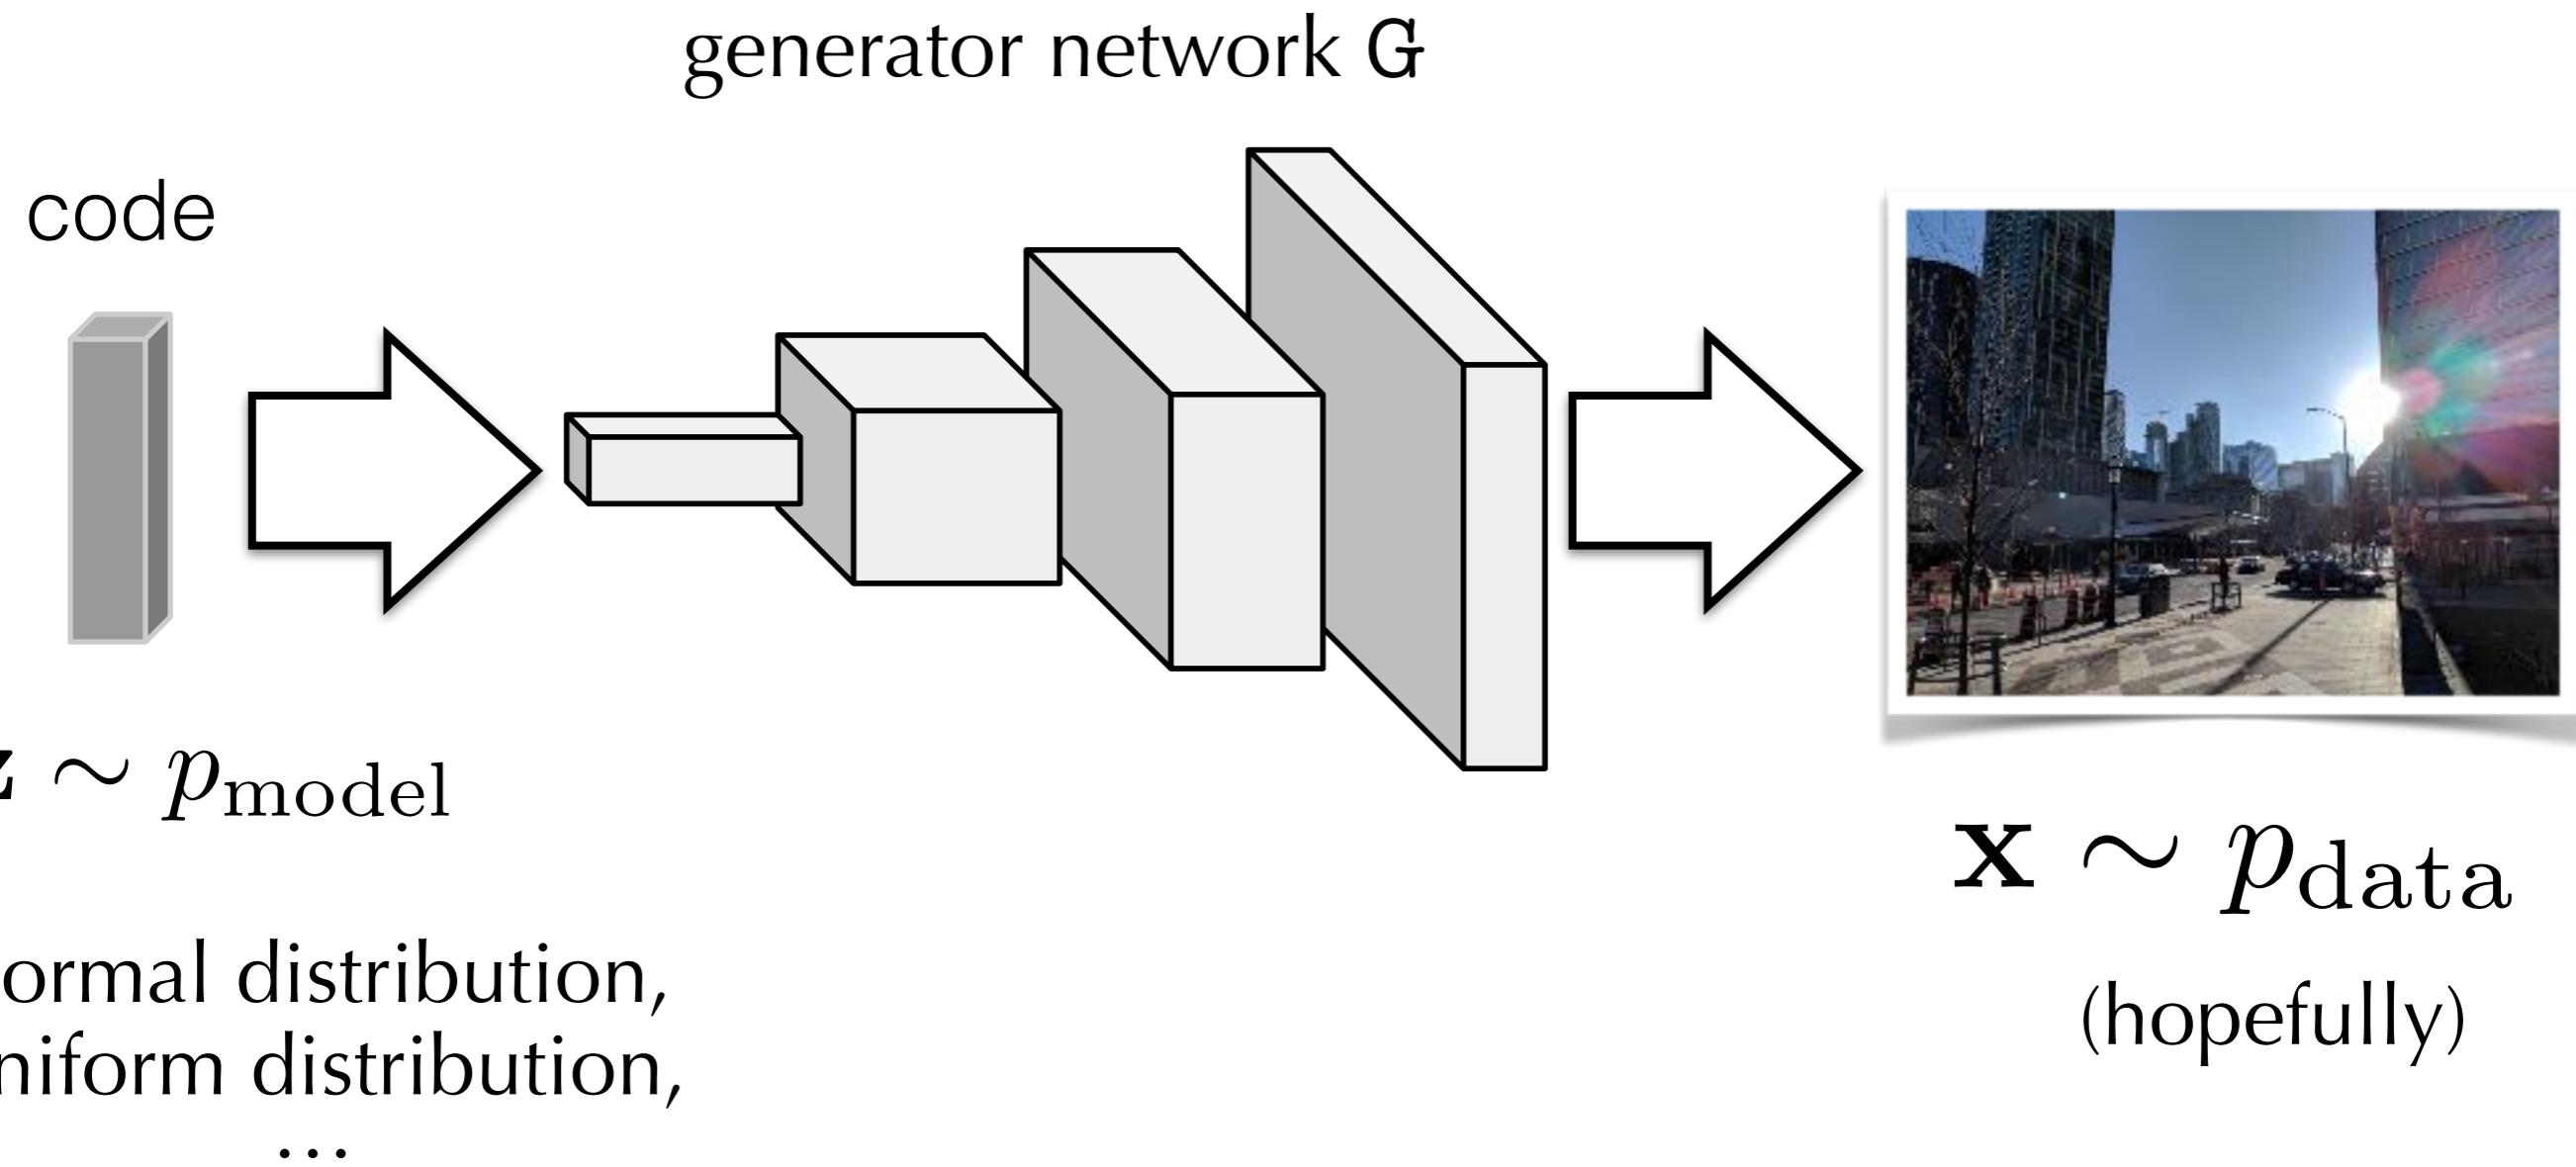


Generate image = pick image from distribution of real images

Generative Neural Networks



Generative Neural Networks



How to make sure that this mapping produces realistic images?

Variational Autoencoder (VAE)

structured latent

space

$$\mathbf{z} \sim p_{\text{model}}$$

lower-dimensional
than input (**bottleneck**)

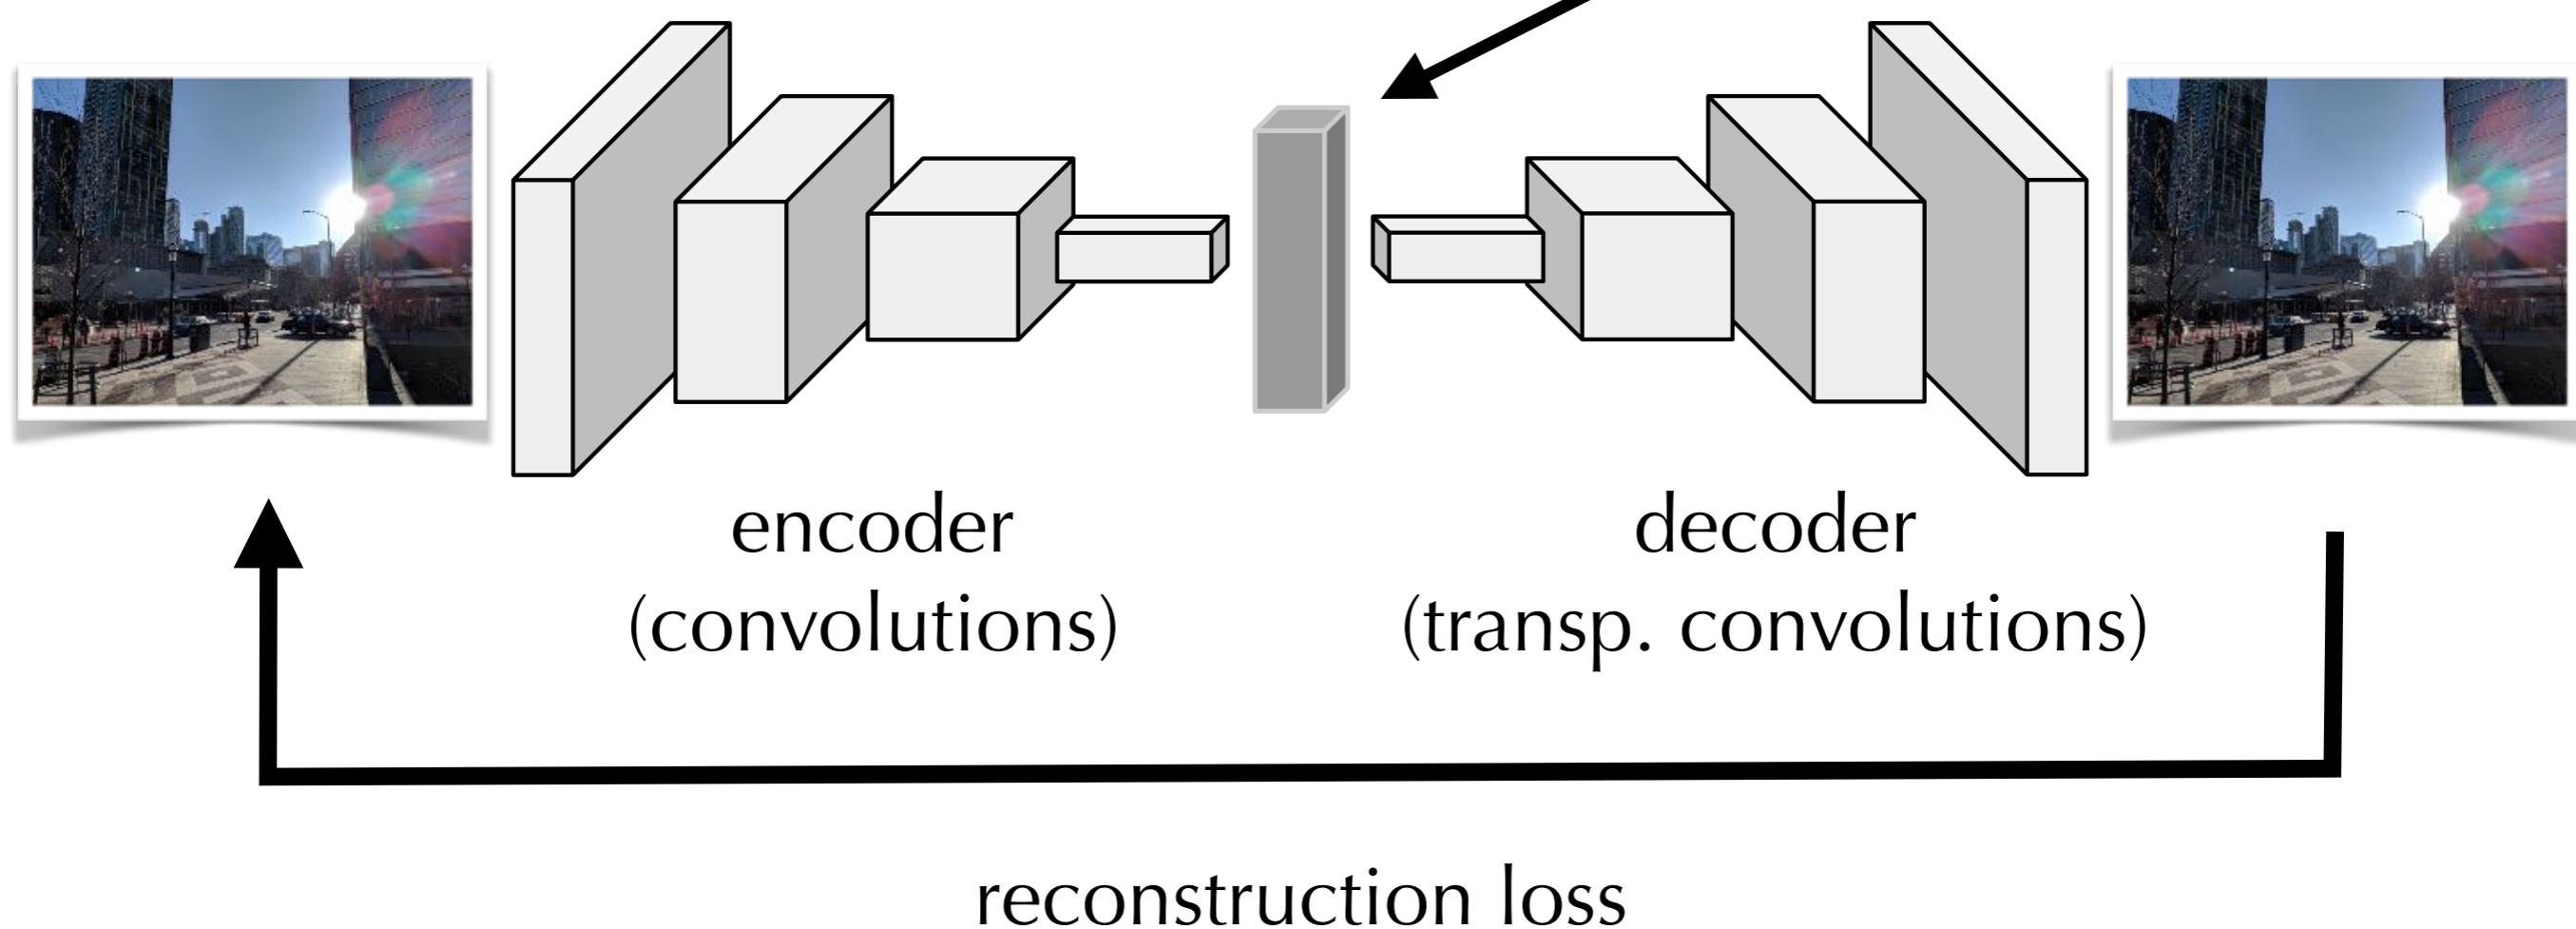
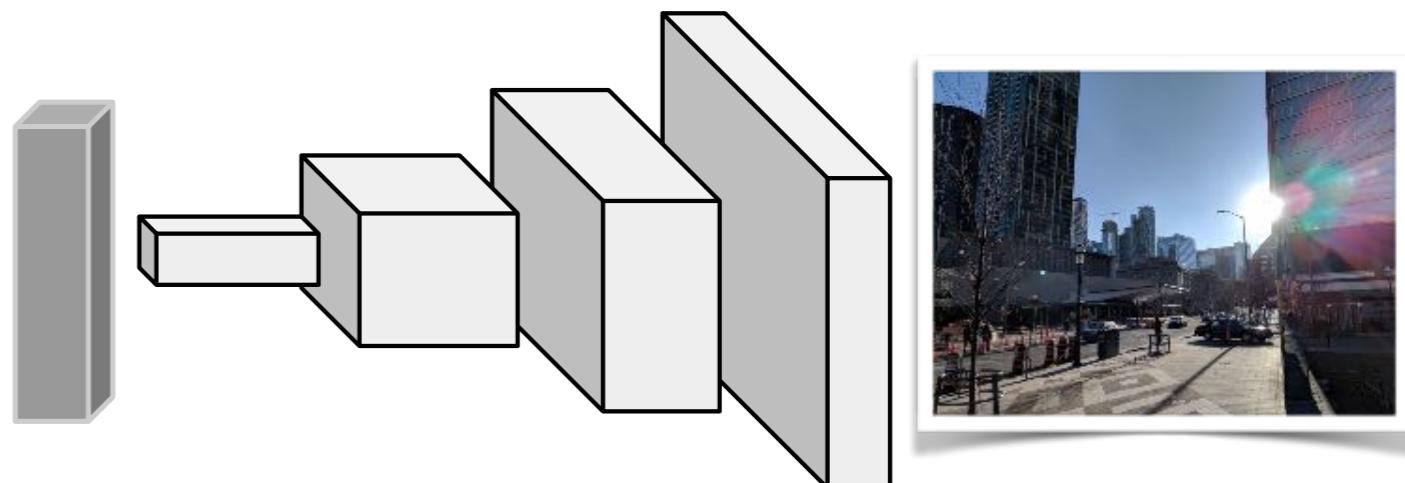


Image Generation via a VAE

sample point in latent space

$$\mathbf{z} \sim p_{\text{model}}$$

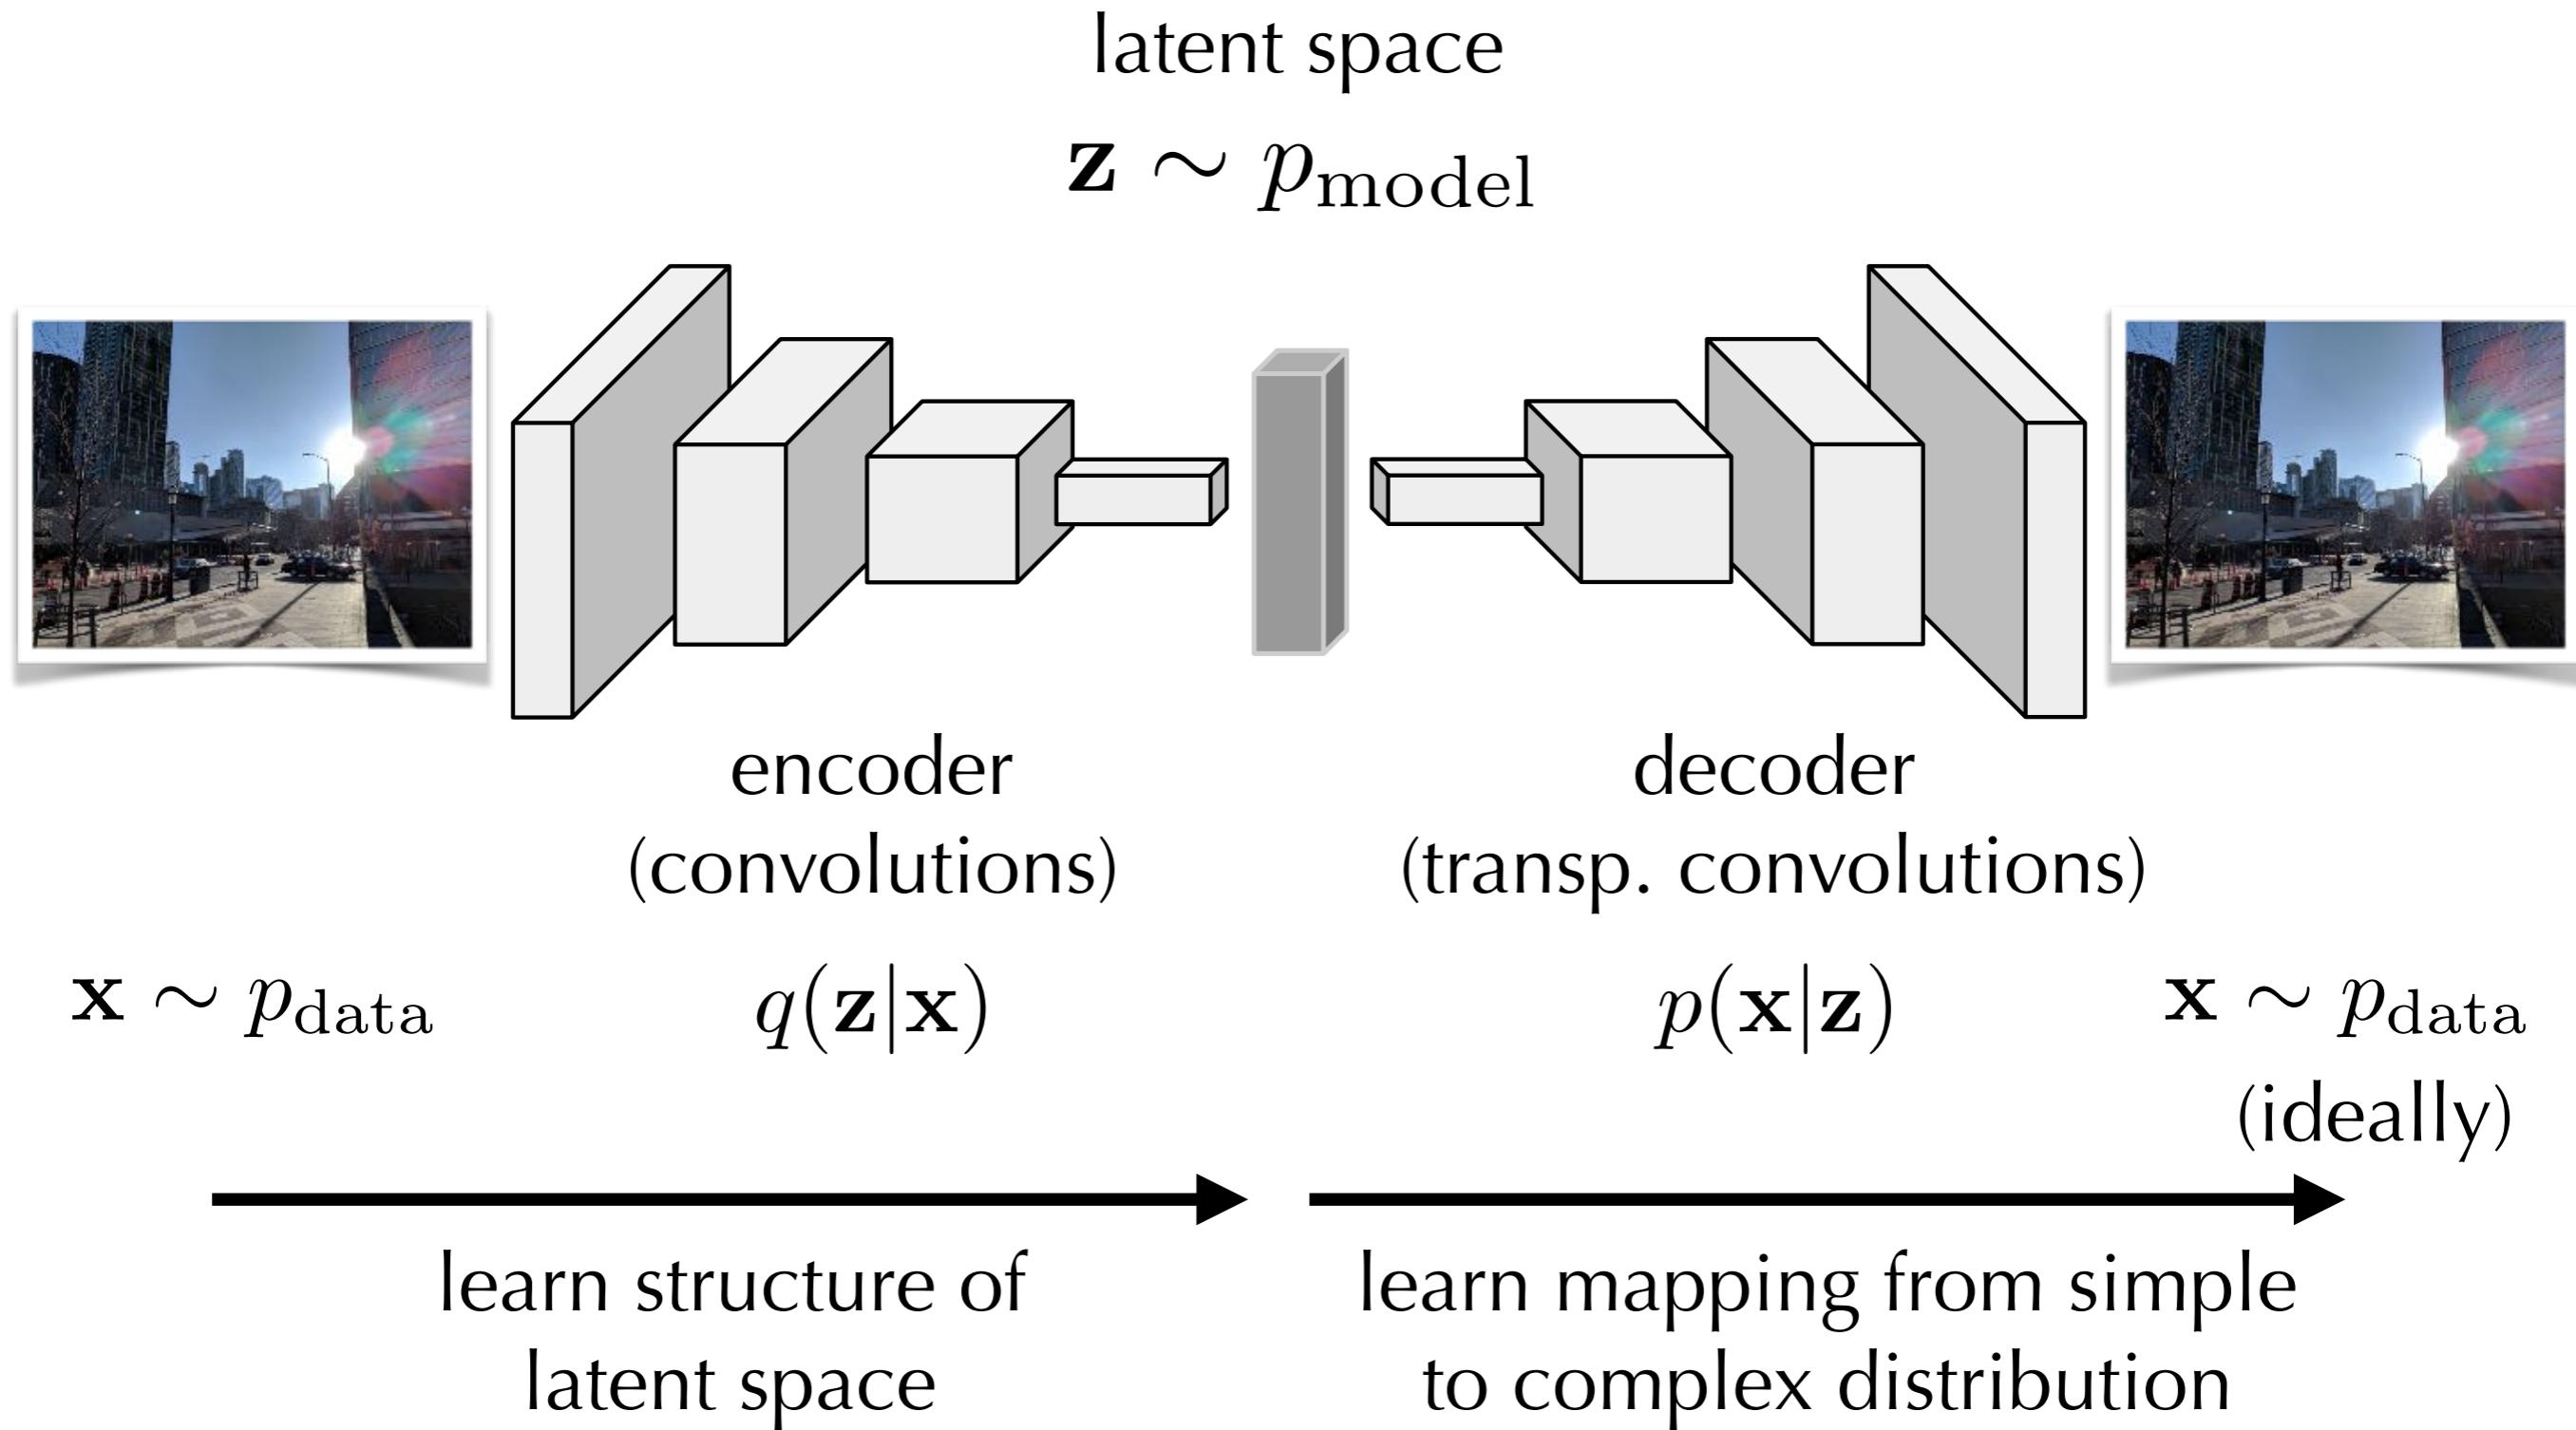


decoder
(transp. convolutions)

$$p(\mathbf{x}|\mathbf{z}) \qquad \mathbf{x} \sim p_{\text{data}} \\ (\text{ideally})$$

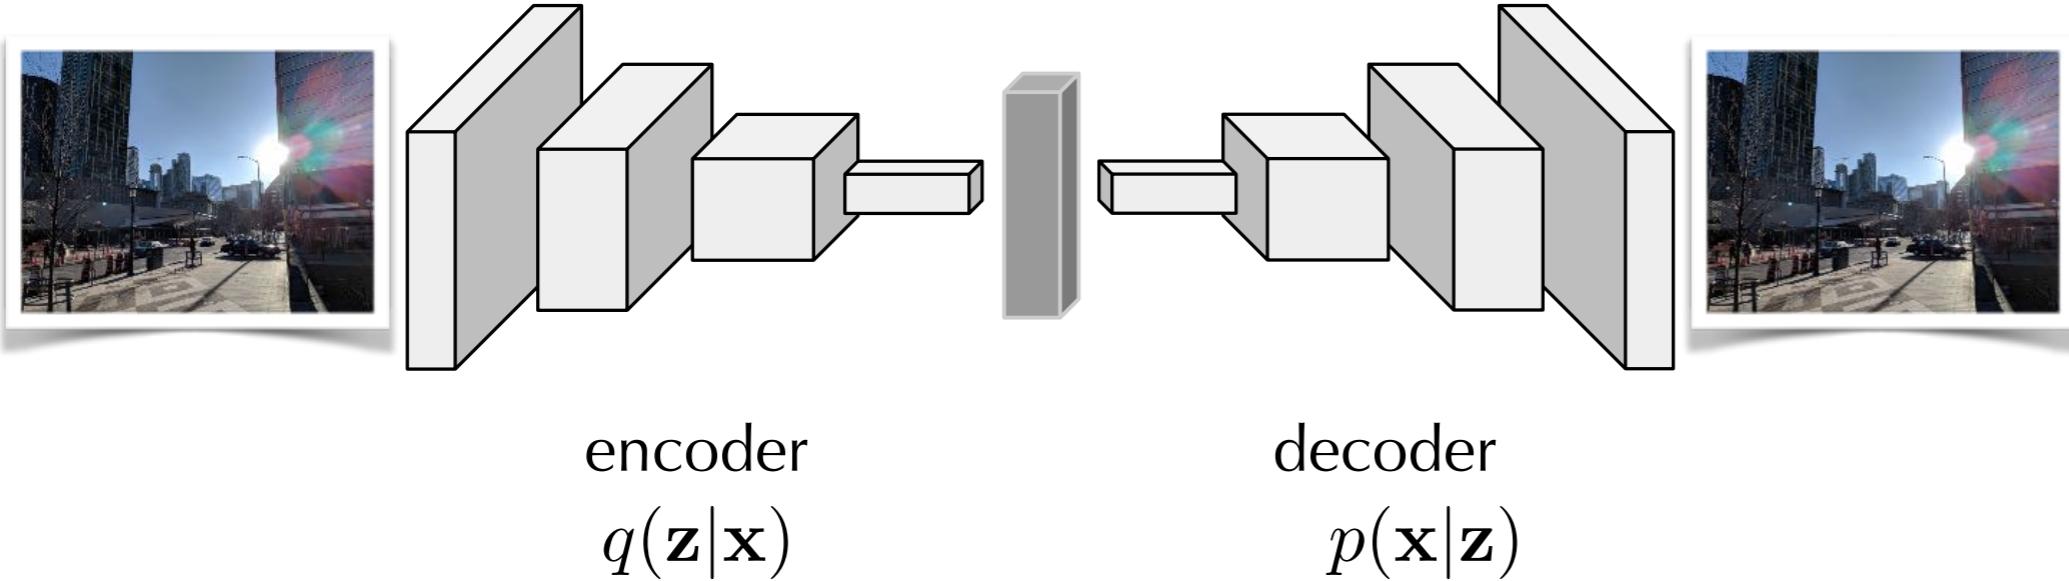
learn mapping from simple
to complex distribution

Image Generation via a VAE



Theory

latent space
 $\mathbf{z} \sim p_{\text{model}}$

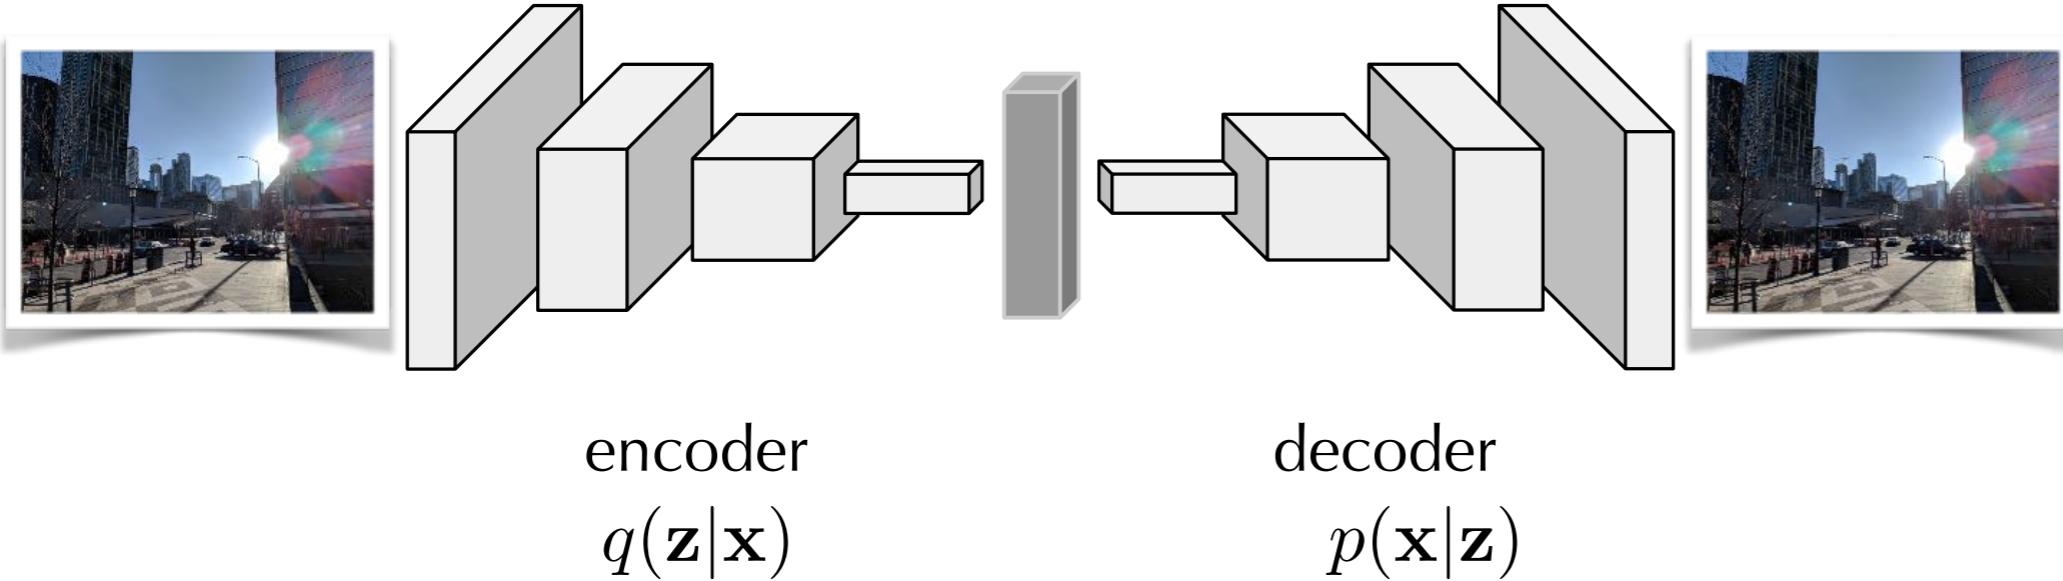


Evidence Lower BOund (ELBO):

$$\log(p(\mathbf{x})) \geq E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

Theory

latent space
 $\mathbf{z} \sim p_{\text{model}}$



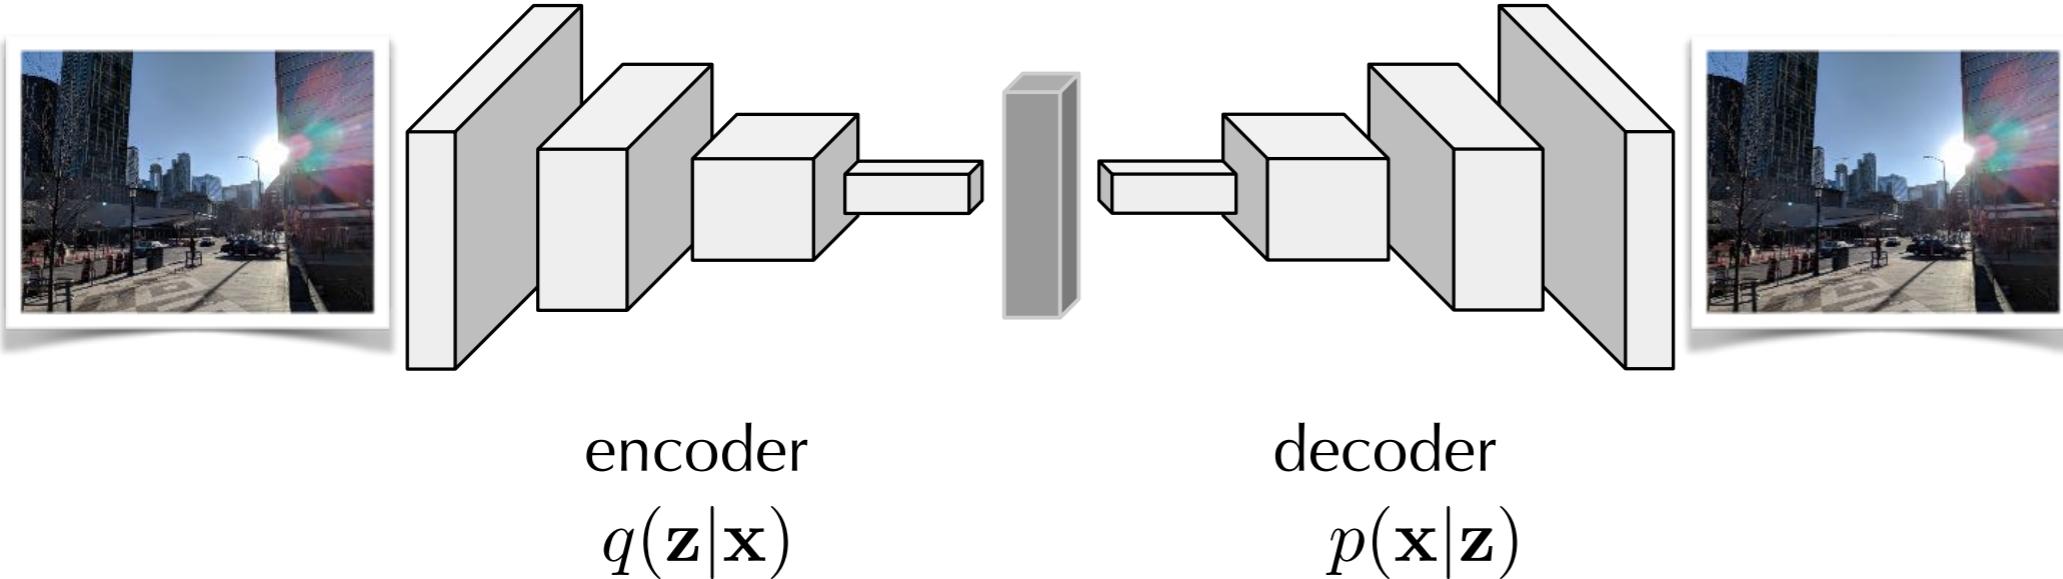
Evidence Lower BOund (ELBO):

$$\log(p(\mathbf{x})) \geq E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

likelihood

Theory

latent space
 $\mathbf{z} \sim p_{\text{model}}$



Evidence Lower BOund (ELBO):

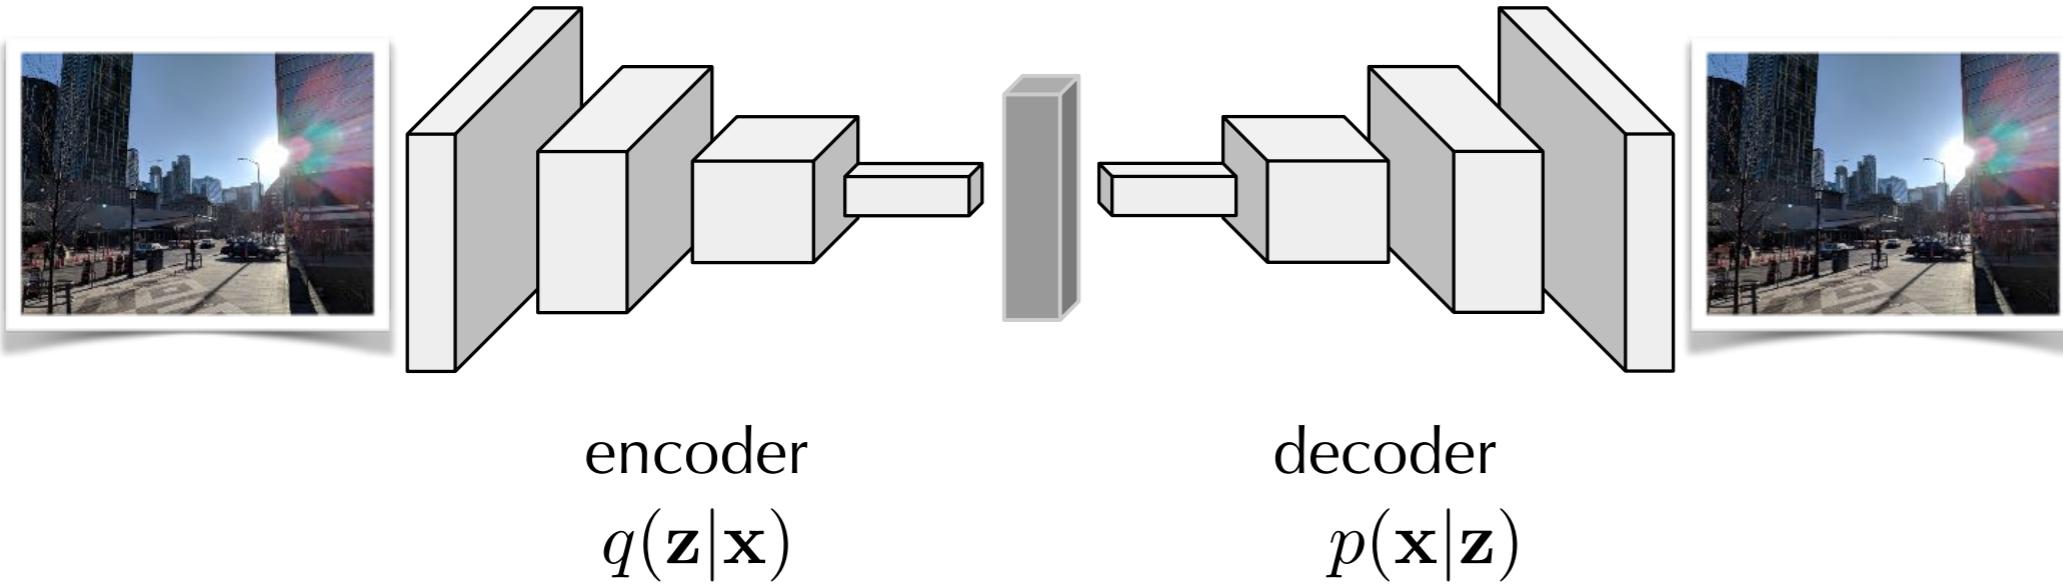
$$\log(p(\mathbf{x})) \geq E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

likelihood

KL divergence

Objective Function

$$\mathbf{z} \sim p_{\text{model}}$$



Maximize

$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}))$$

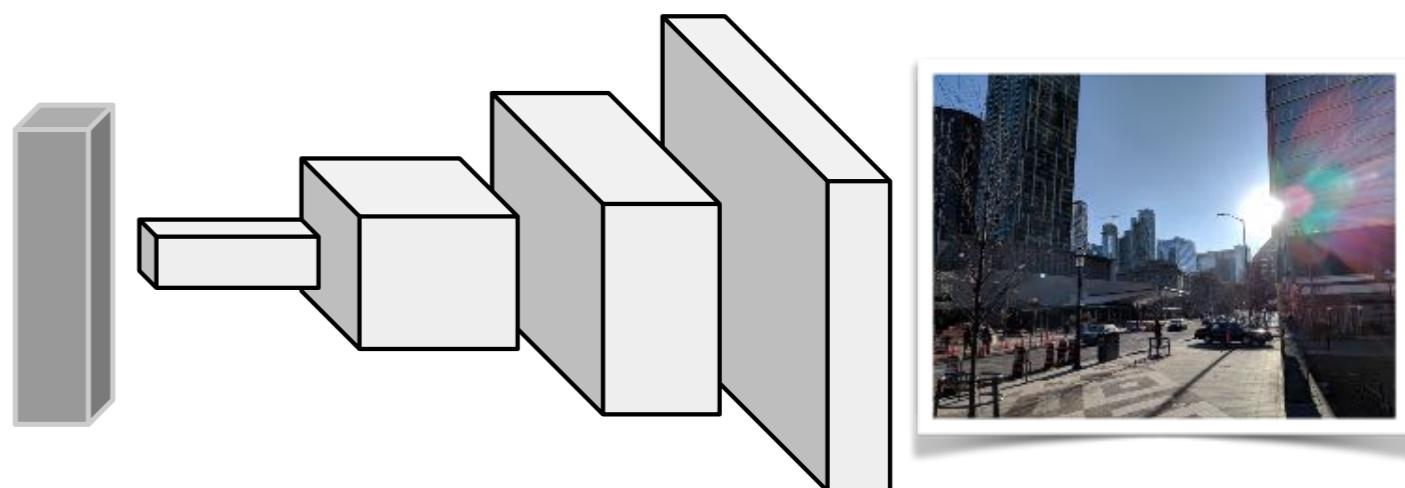
reconstruction ``loss'',
e.g., L2 norm (normal distribution)

similarity between learned
and sampling distribution

At Test Time

sample point from latent space

$$\mathbf{z} \sim p_{\text{model}}$$



decoder
(transp. convolutions)

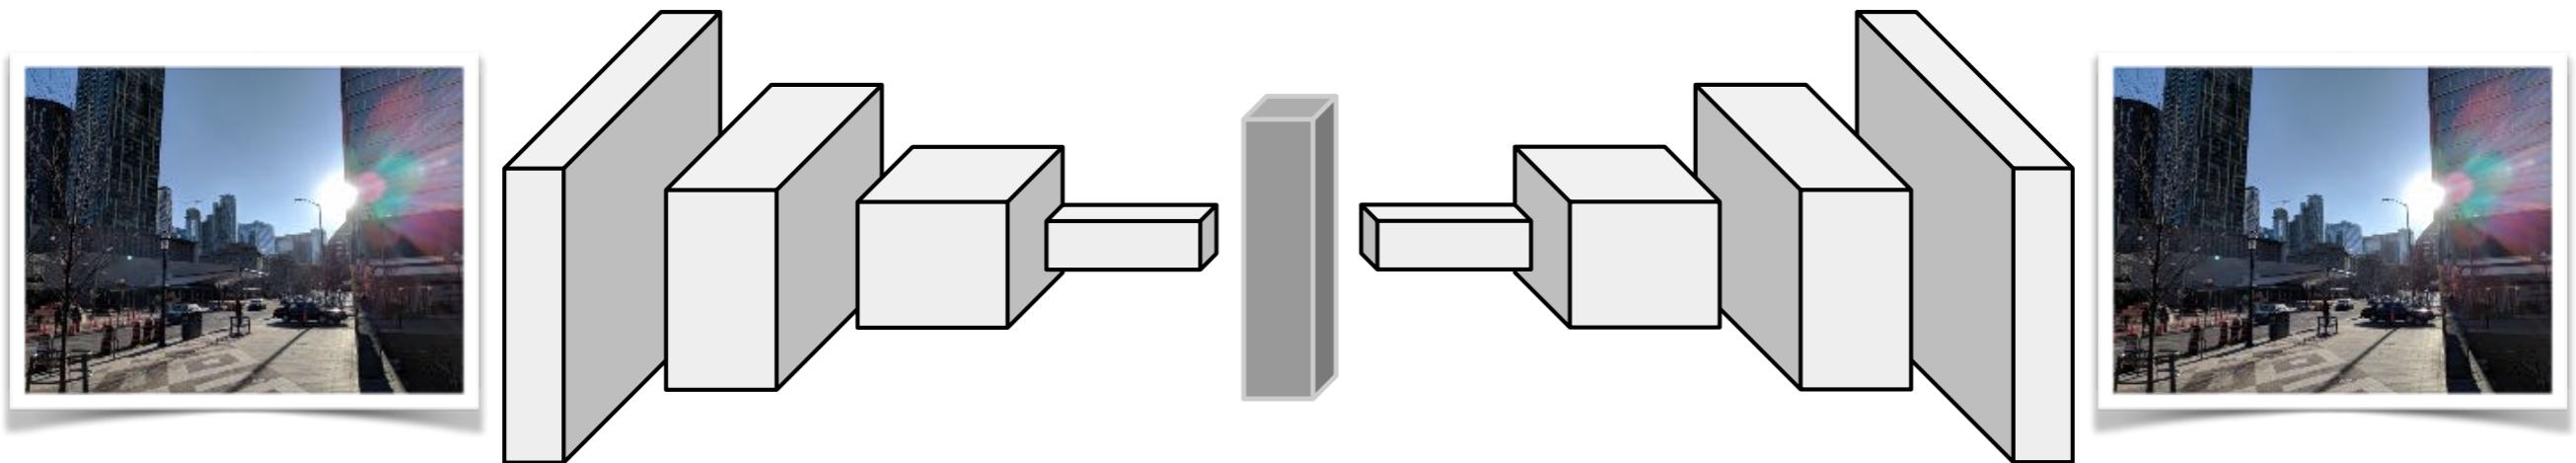
$$p(\mathbf{x}|\mathbf{z})$$

$$\mathbf{x} \sim p_{\text{data}}$$

use decoder to ``translate''
to image

VAE - Example

$$\mathbf{z} \sim p_{\text{model}} = \mathcal{N}(\mathbf{z}; \mathbf{0}, \mathbf{I})$$



encoder

$$q(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\mathbf{z}; \mu(\mathbf{x}), \sigma^2(\mathbf{x})\mathbf{I})$$

multi-variate Gaussian distribution

decoder

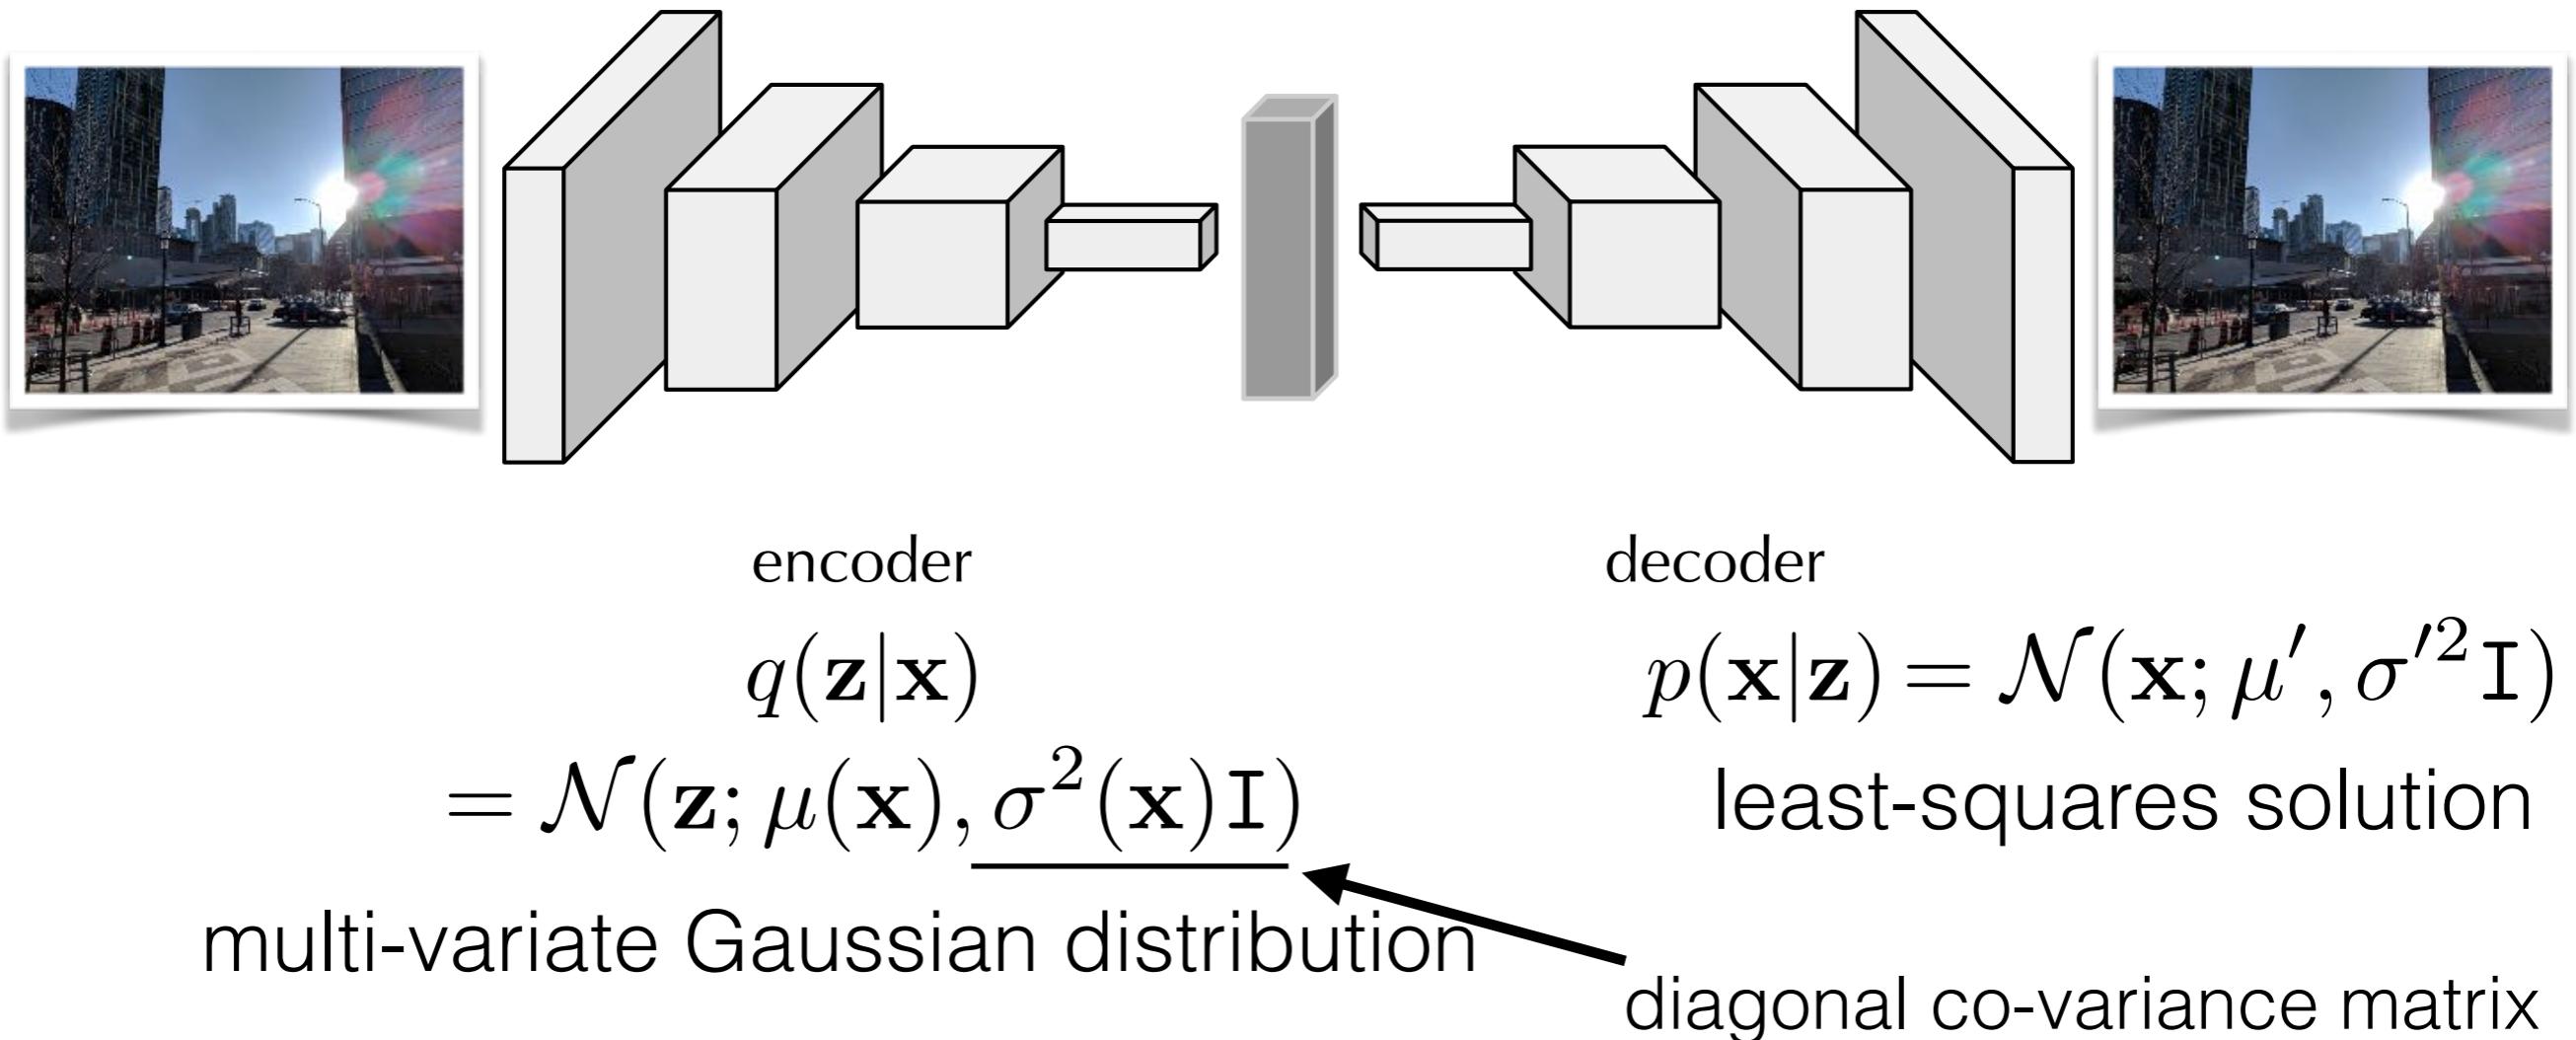
$$p(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{x}; \mu', \sigma'^2 \mathbf{I})$$

least-squares solution

[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

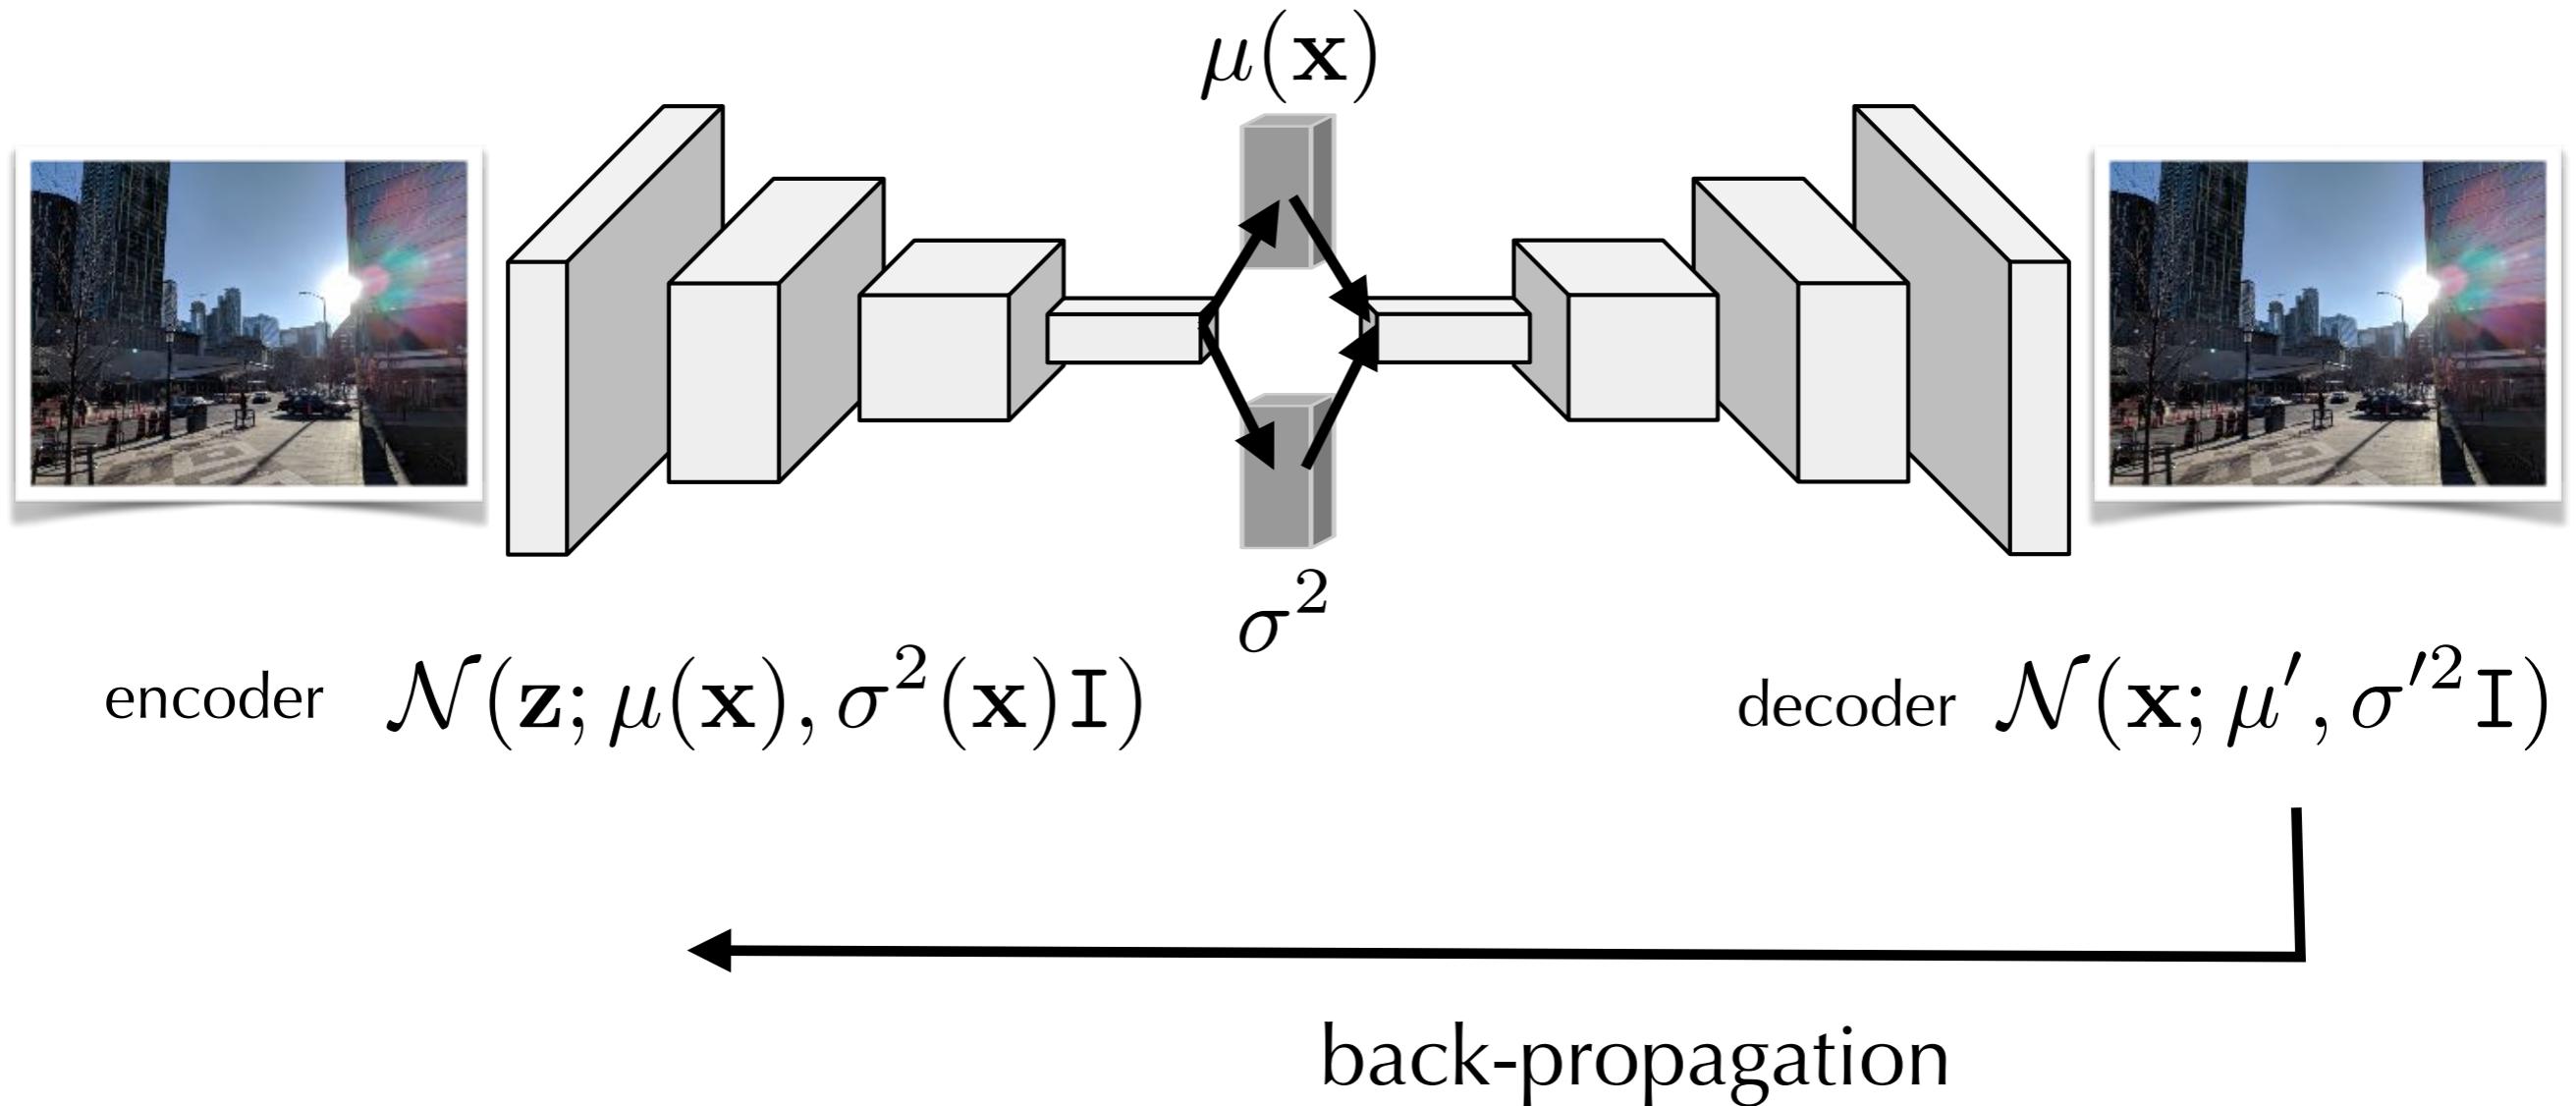
VAE - Example

$$\mathbf{z} \sim p_{\text{model}} = \mathcal{N}(\mathbf{z}; \mathbf{0}, \mathbf{I})$$



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

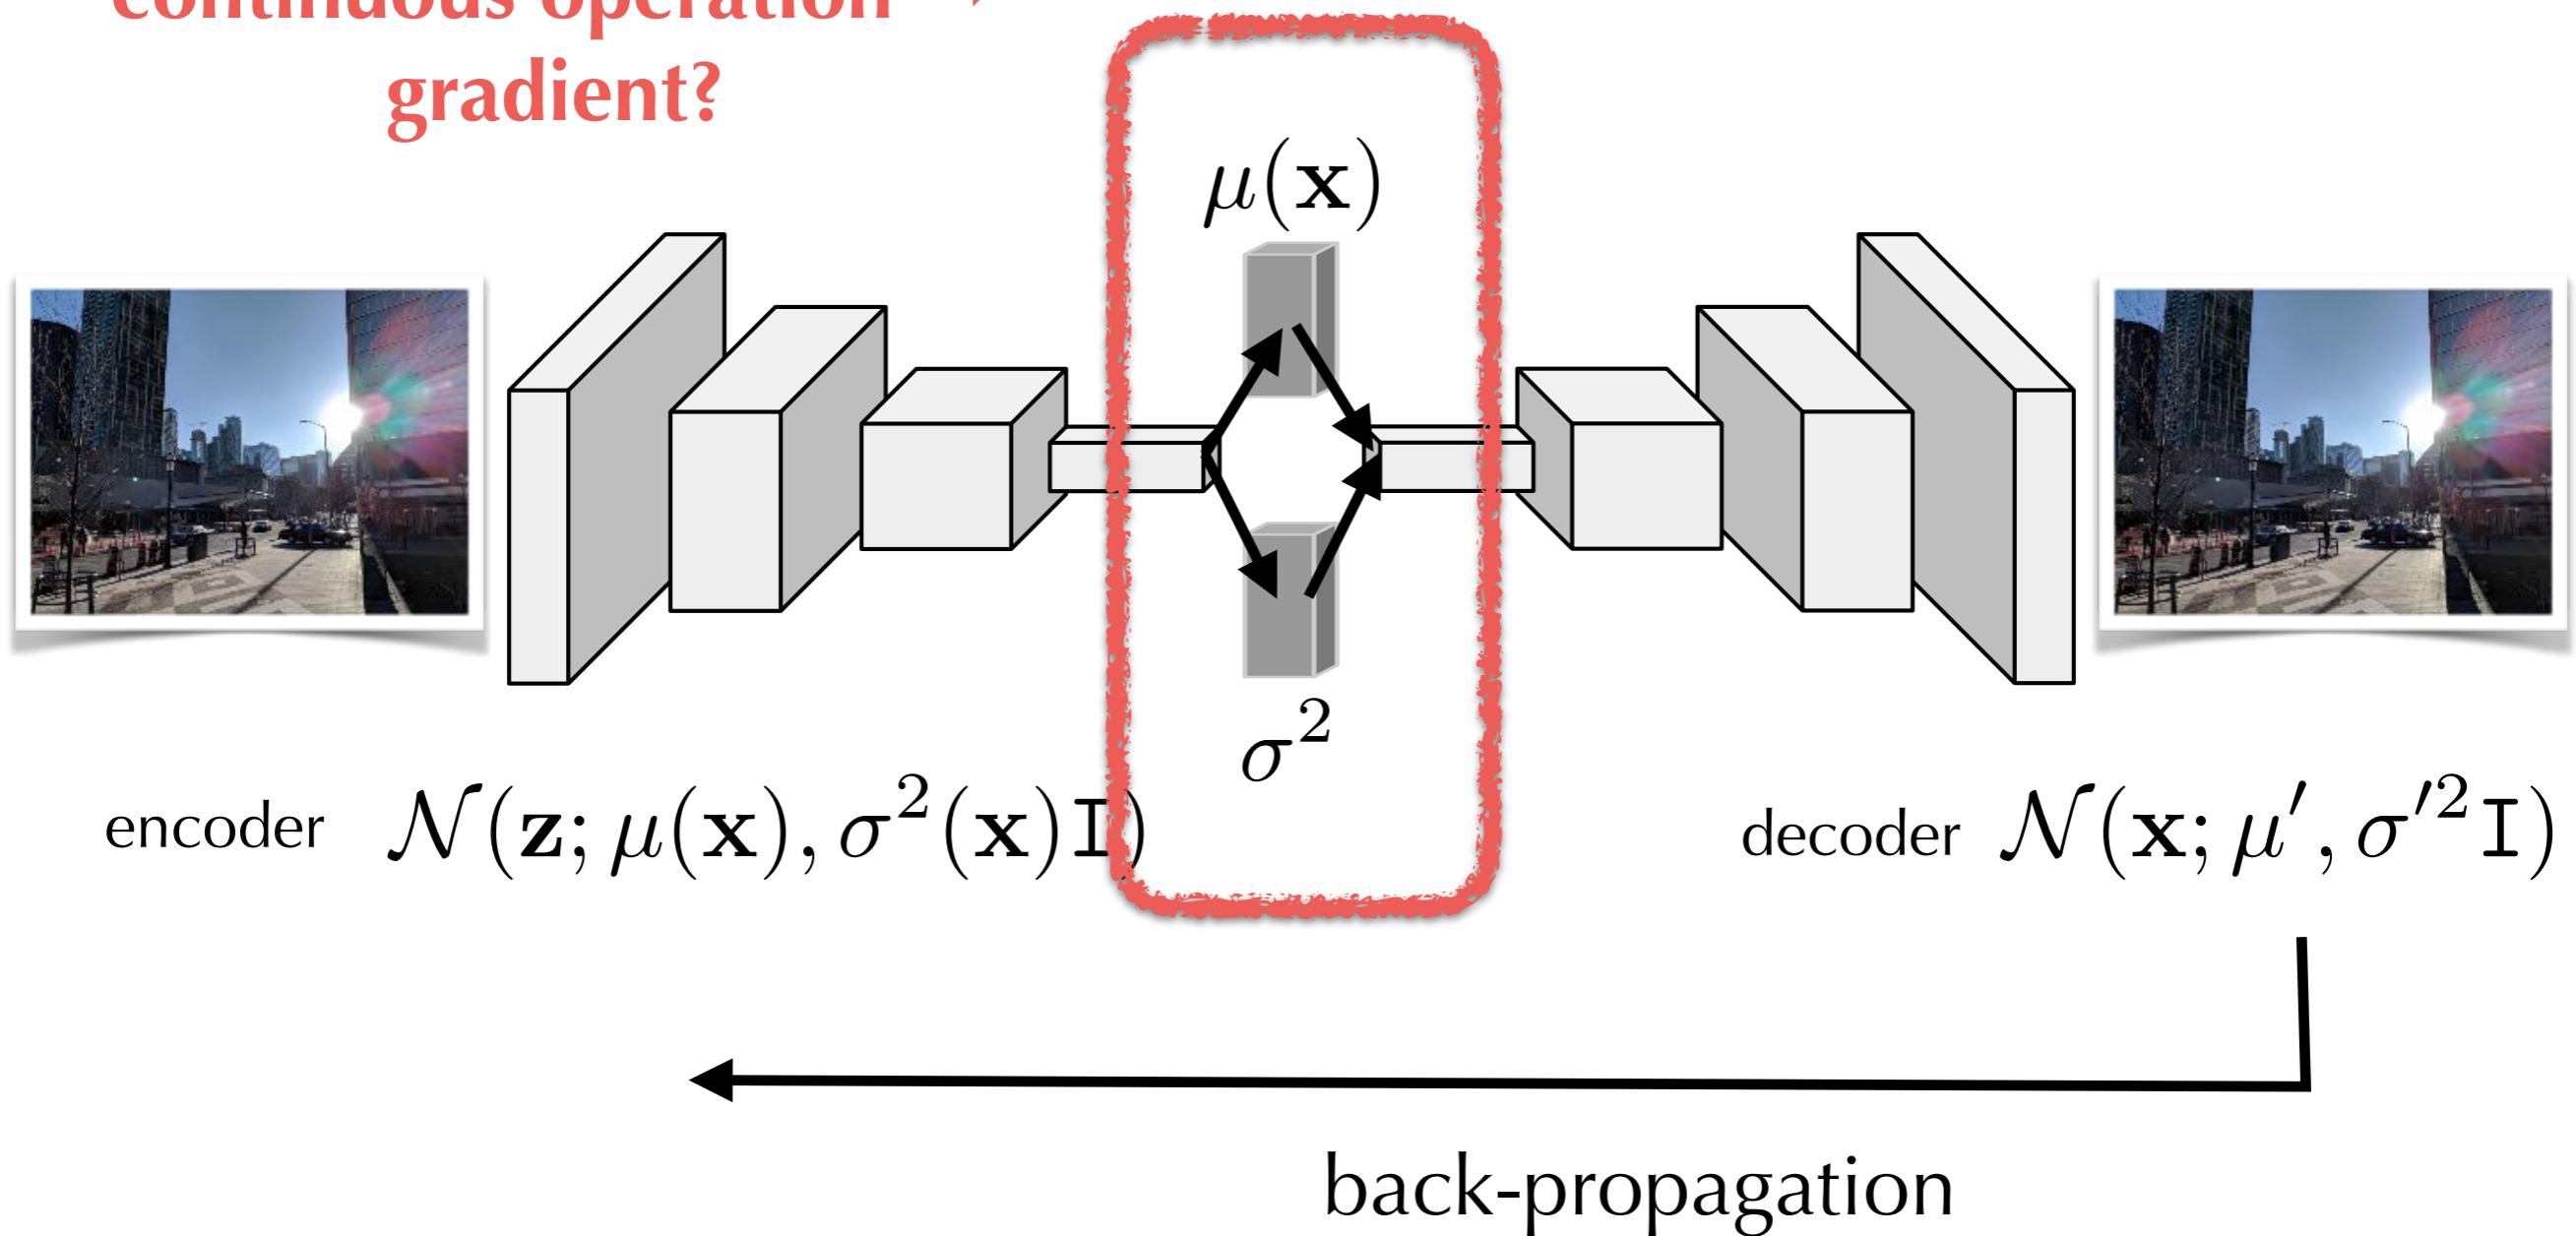
VAE - Example



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

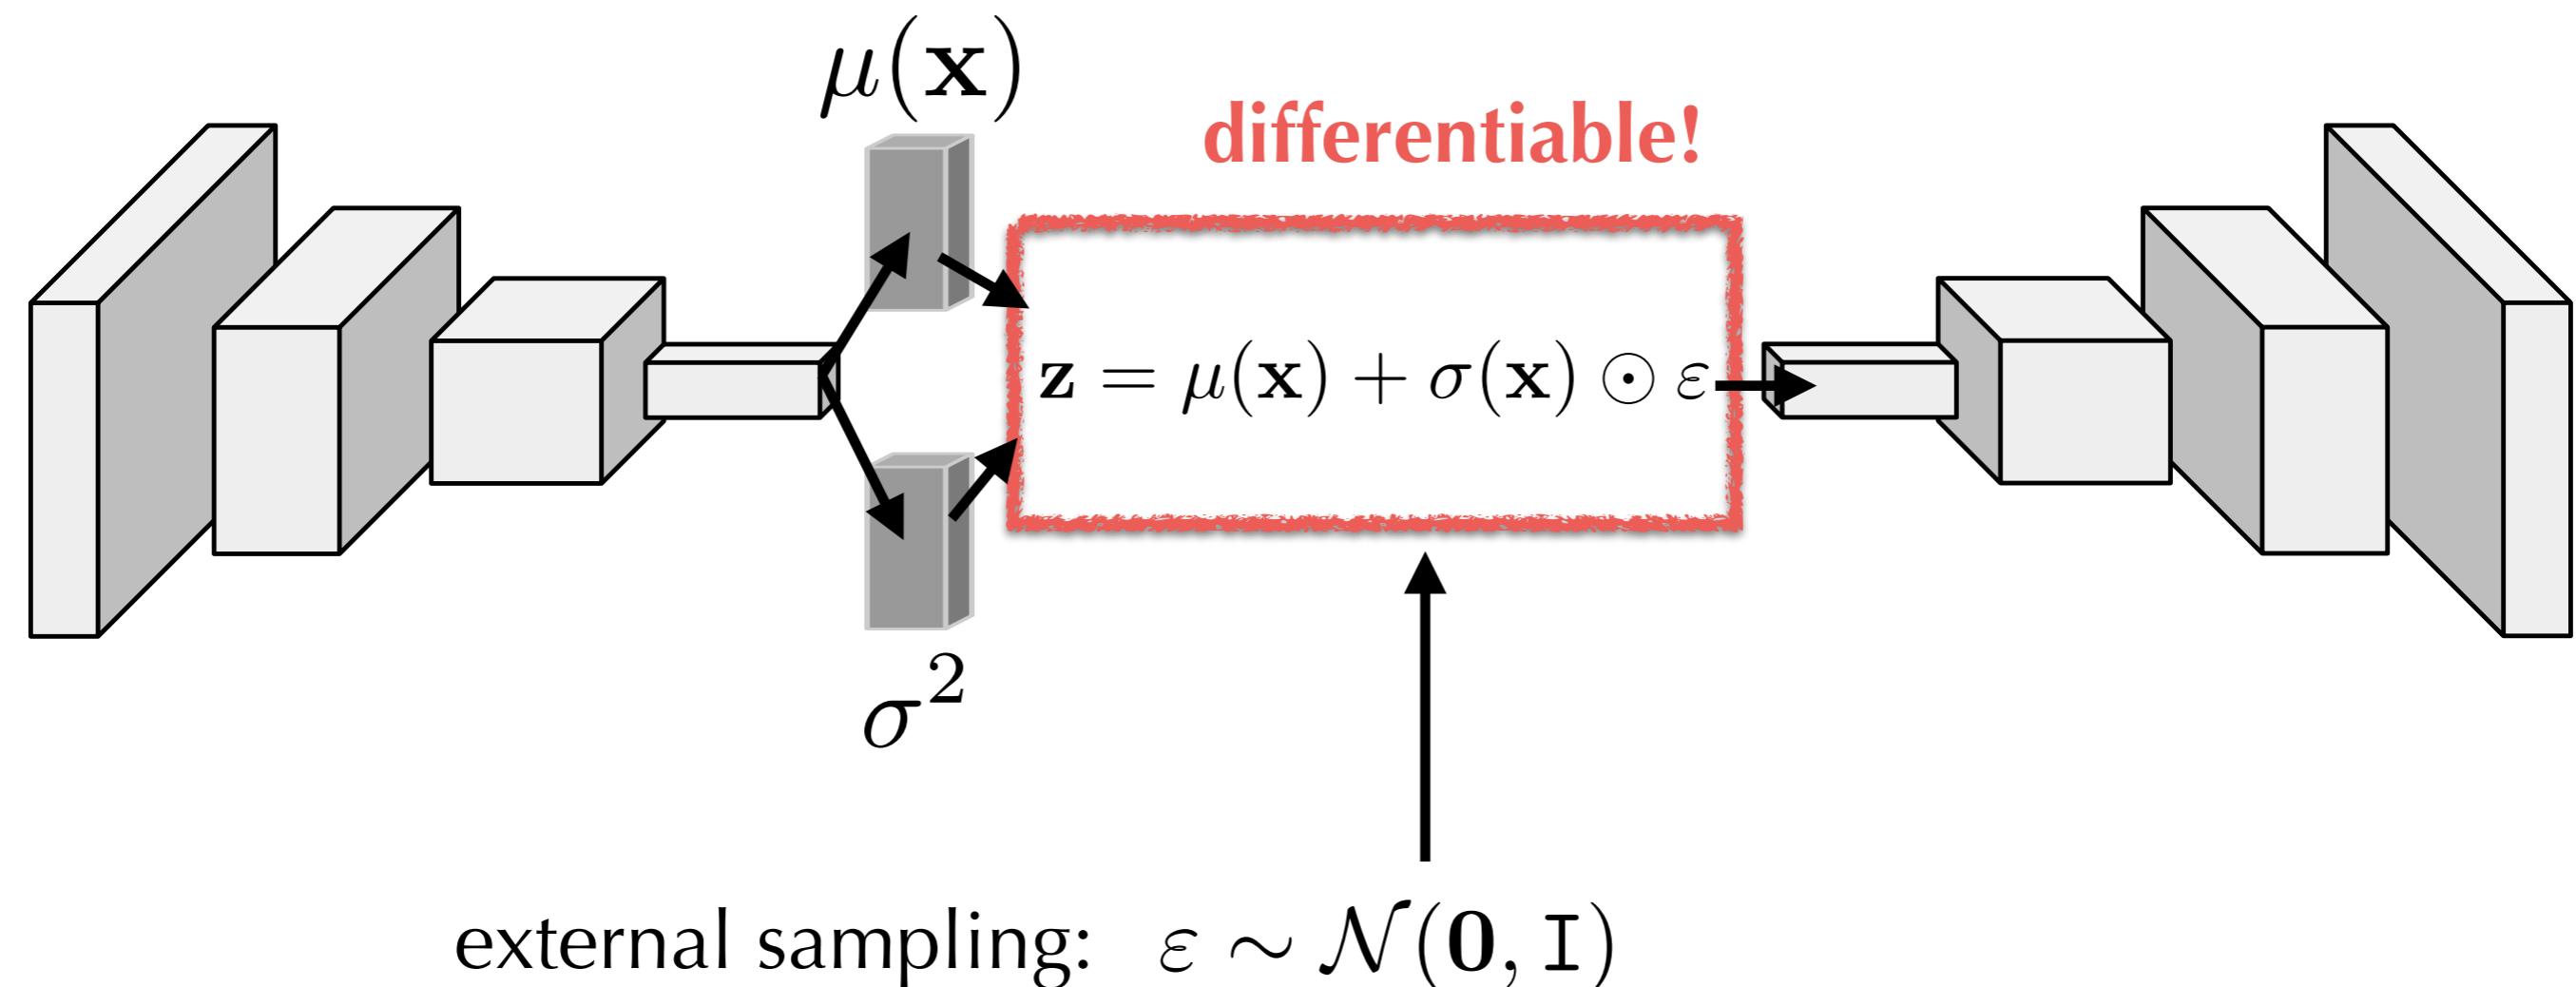
VAE - Example

sampling \mathbf{z} is non-continuous operation → gradient?



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

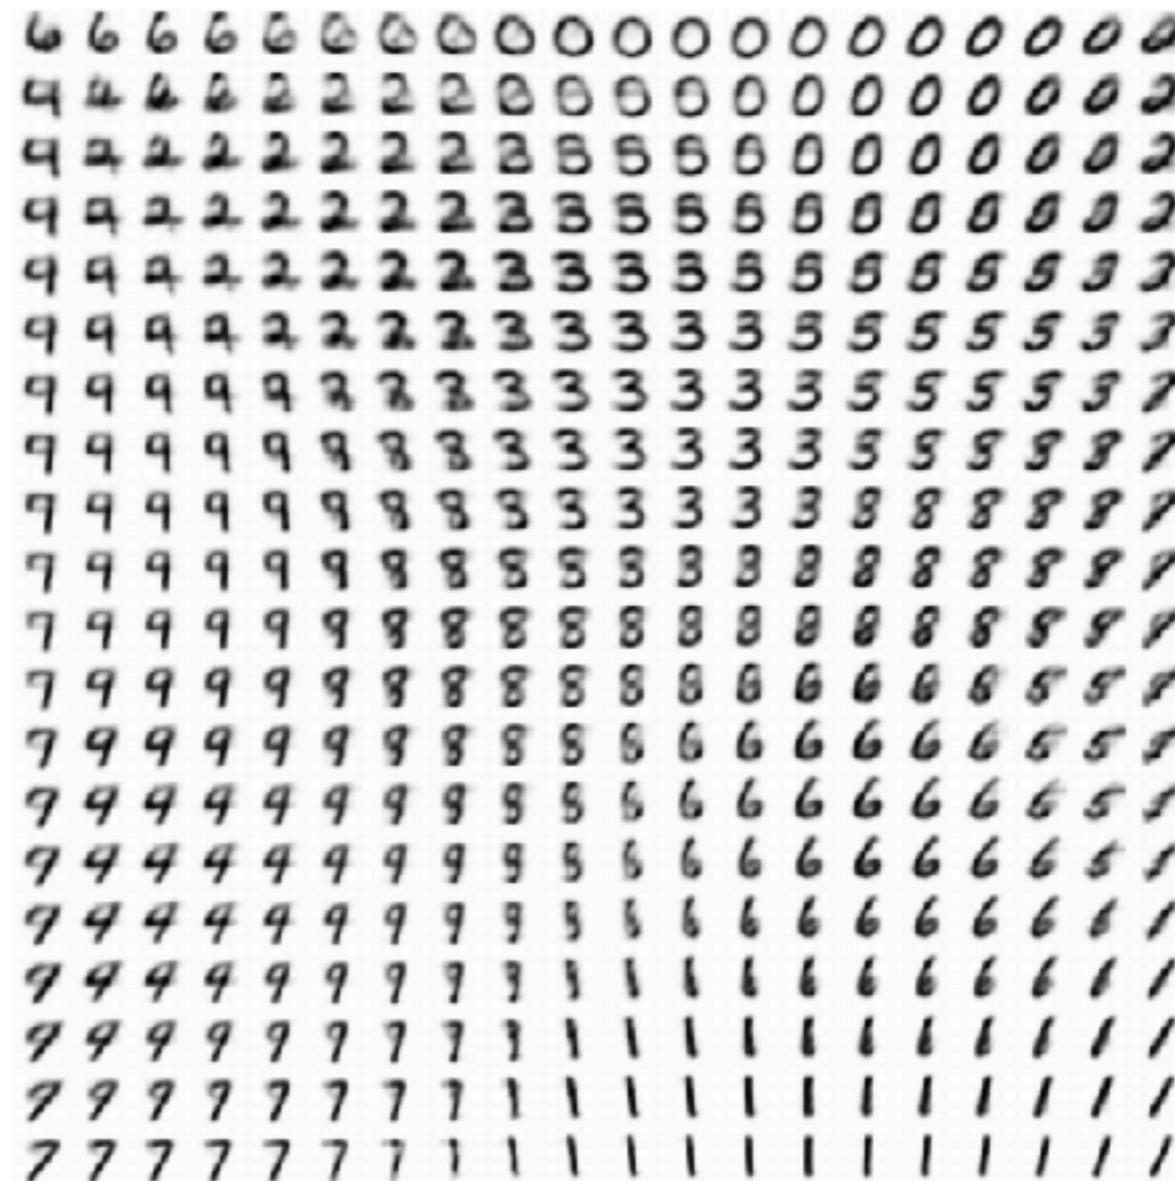
The Reparametrization Trick



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

VAE Example - Results

Learned manifold (2D latent space)



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

VAE Example - Results

Learned manifold (2D latent space)



[Kingma & Welling, Auto-Encoding Variational Bayes, ICLR 2014]

Today

Generative Neural Networks

Today

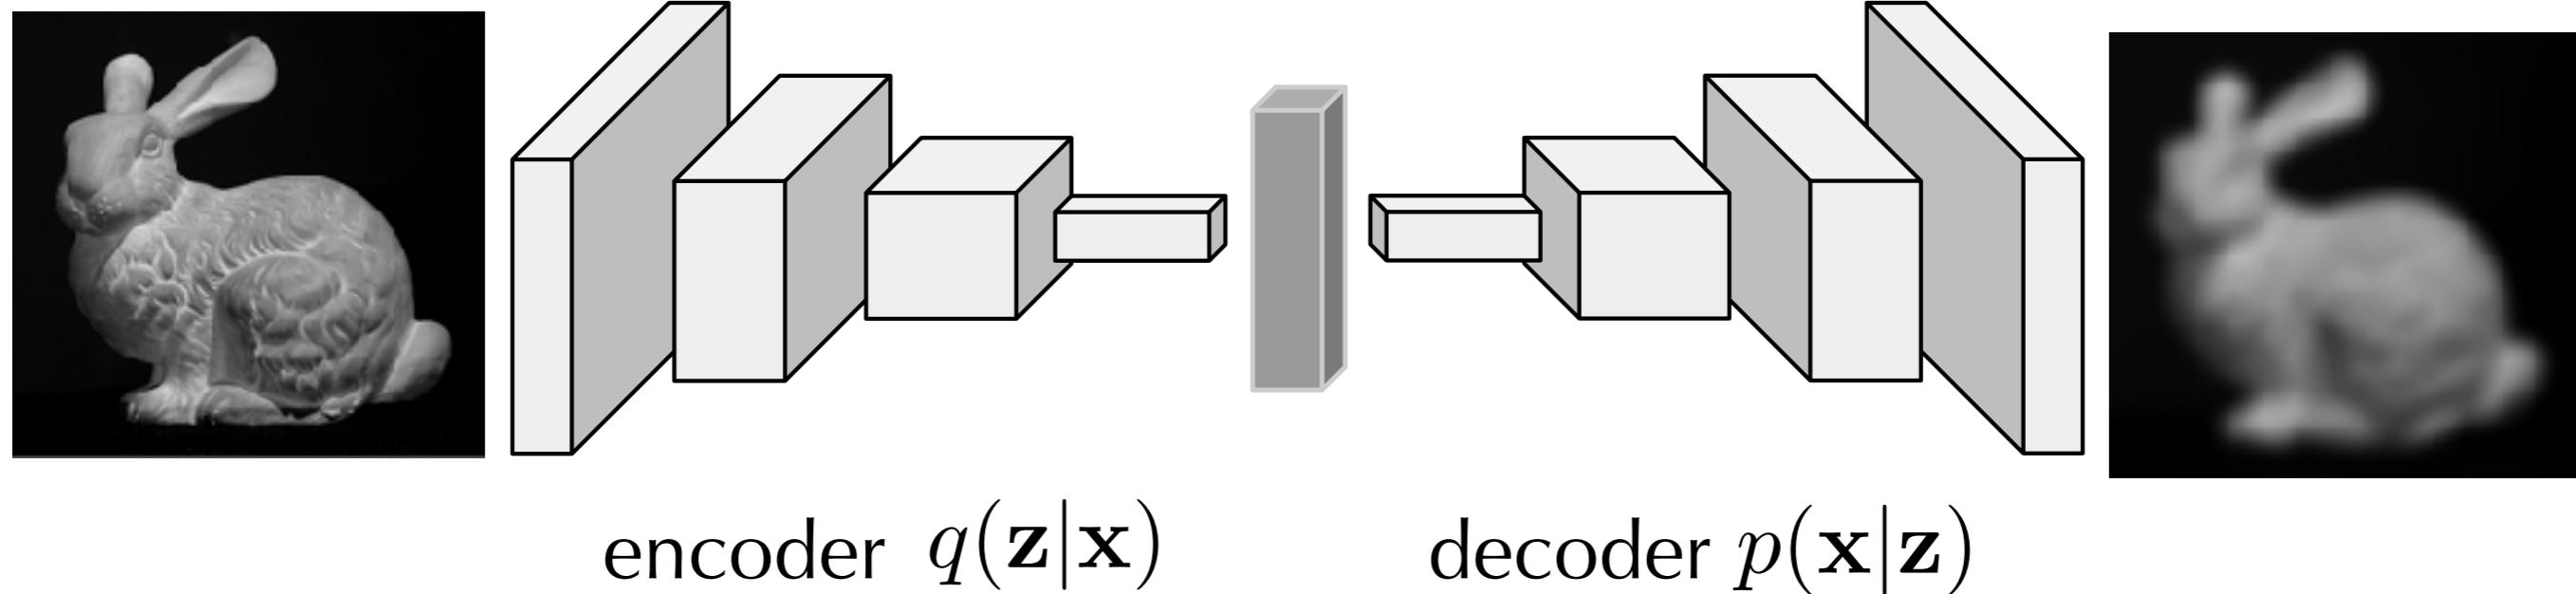
- Generative Adversarial Networks (GANs)
- (Unsupervised) Image Translation

Today

- **Generative Adversarial Networks (GANs)**
- (Unsupervised) Image Translation

Variational Autoencoder

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$

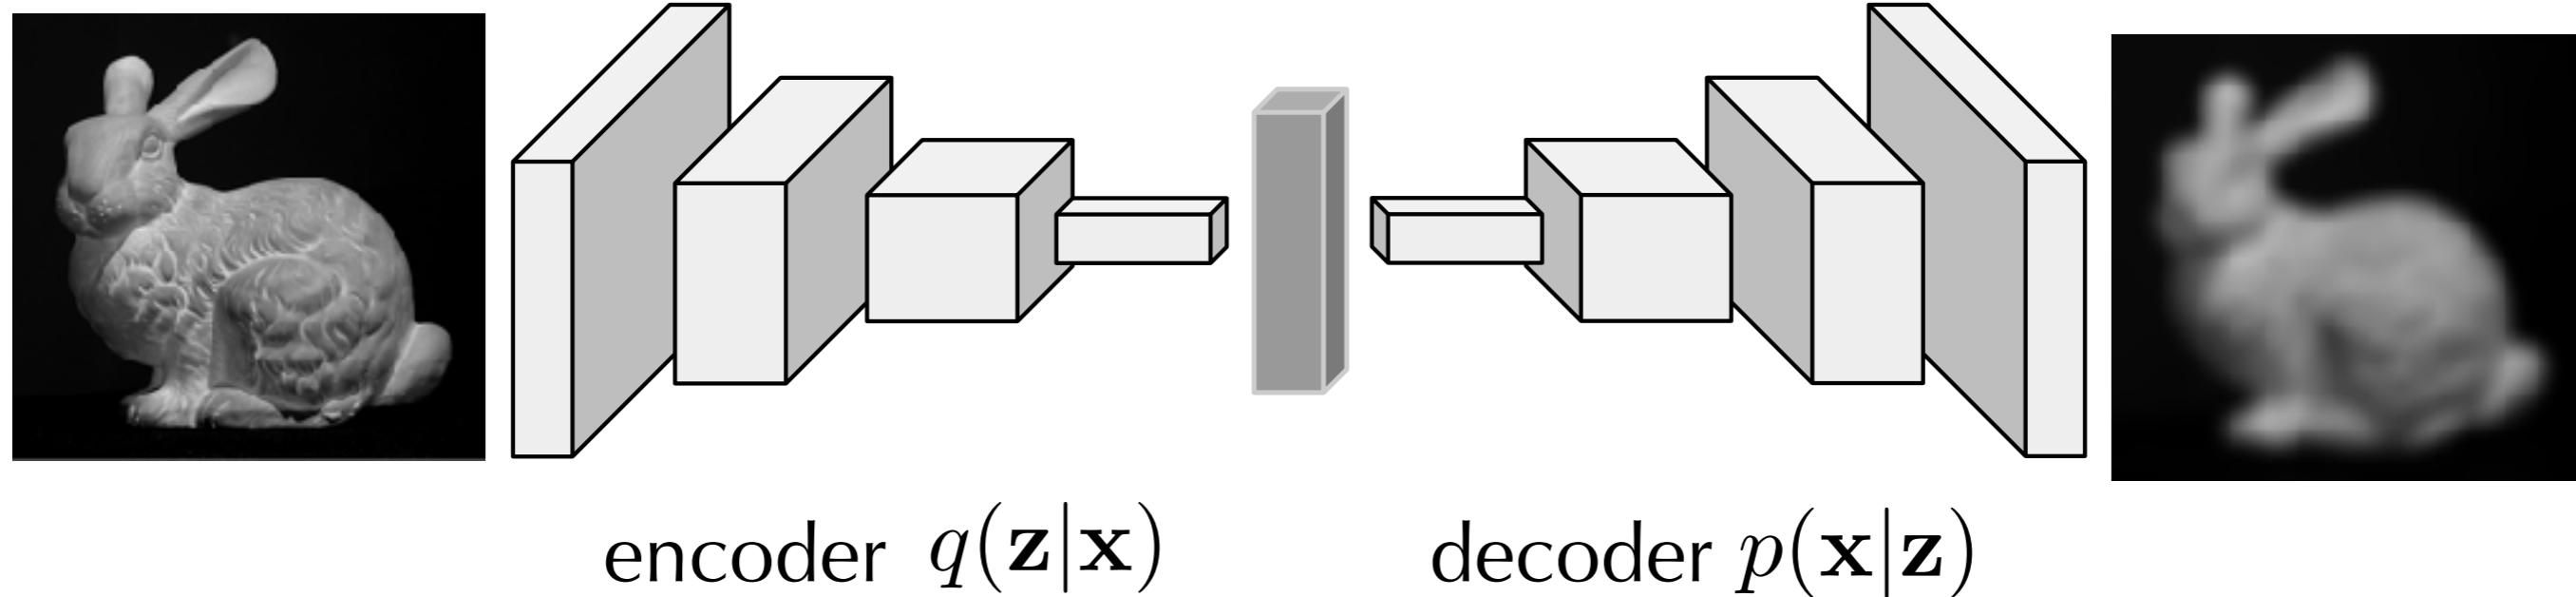


Maximize

$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$

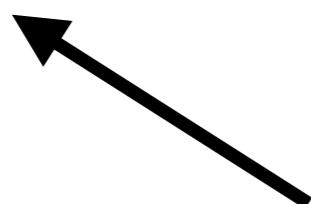
Variational Autoencoder

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



Maximize

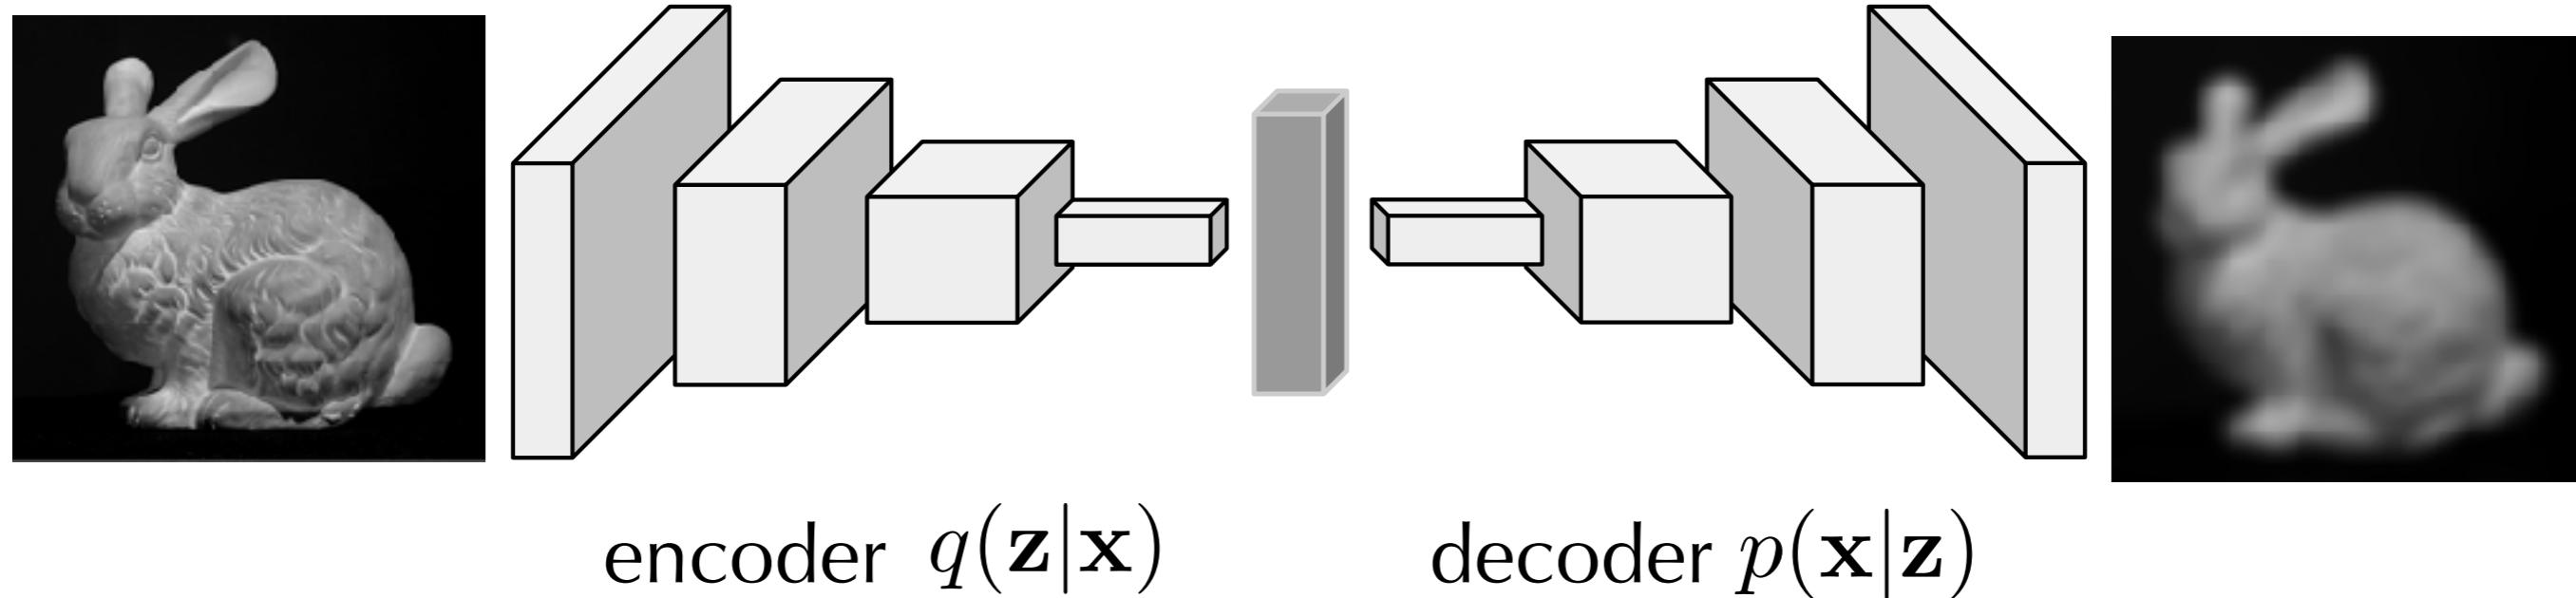
$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$



Typically a normal distribution \rightarrow blurry results

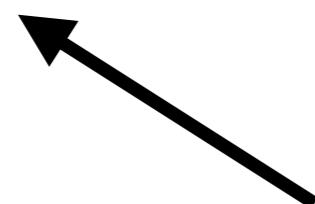
Variational Autoencoder

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



Maximize

$$E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$



Better loss functions?

Typically a normal distribution \rightarrow blurry results

The Need for a Better Loss Function



original



output 1



output 2

[Doersch, Tutorial on Variational Autoencoders, arXiv:1606.05908]

The Need for a Better Loss Function



original



output 1



output 2

Output 2 better, but slightly shifted → larger loss under Gaussian distribution

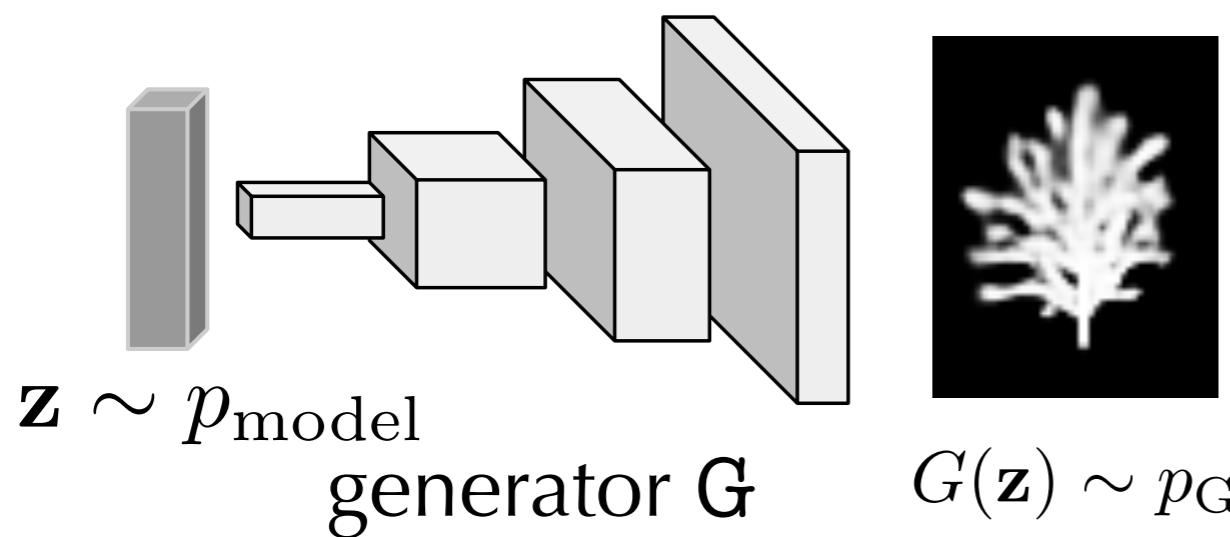
Generative Adversarial Networks (GANs)

Motivation: Learn loss function

[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

Generative Adversarial Networks (GANs)

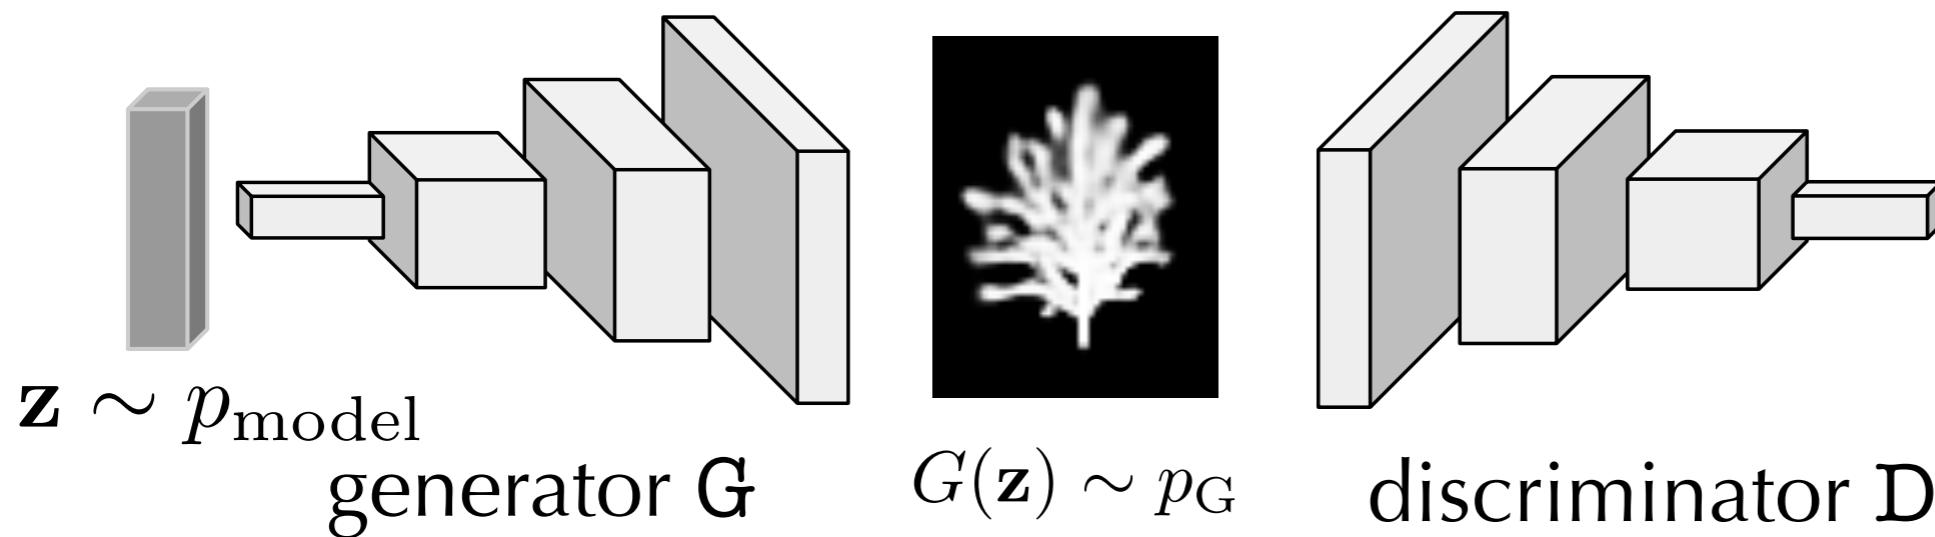
Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

Generative Adversarial Networks (GANs)

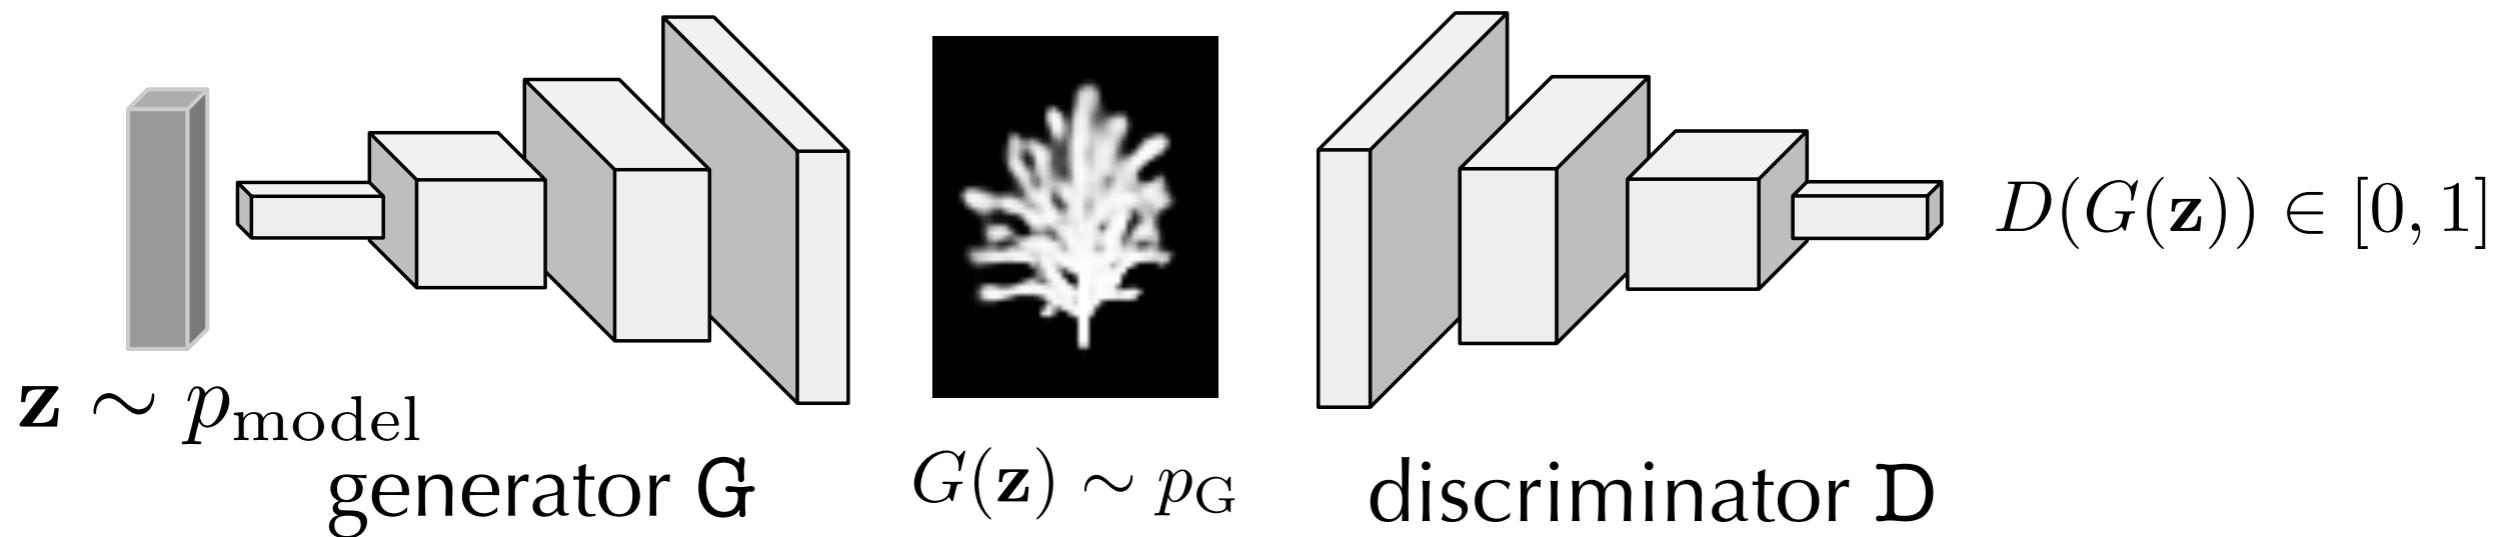
Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

Generative Adversarial Networks (GANs)

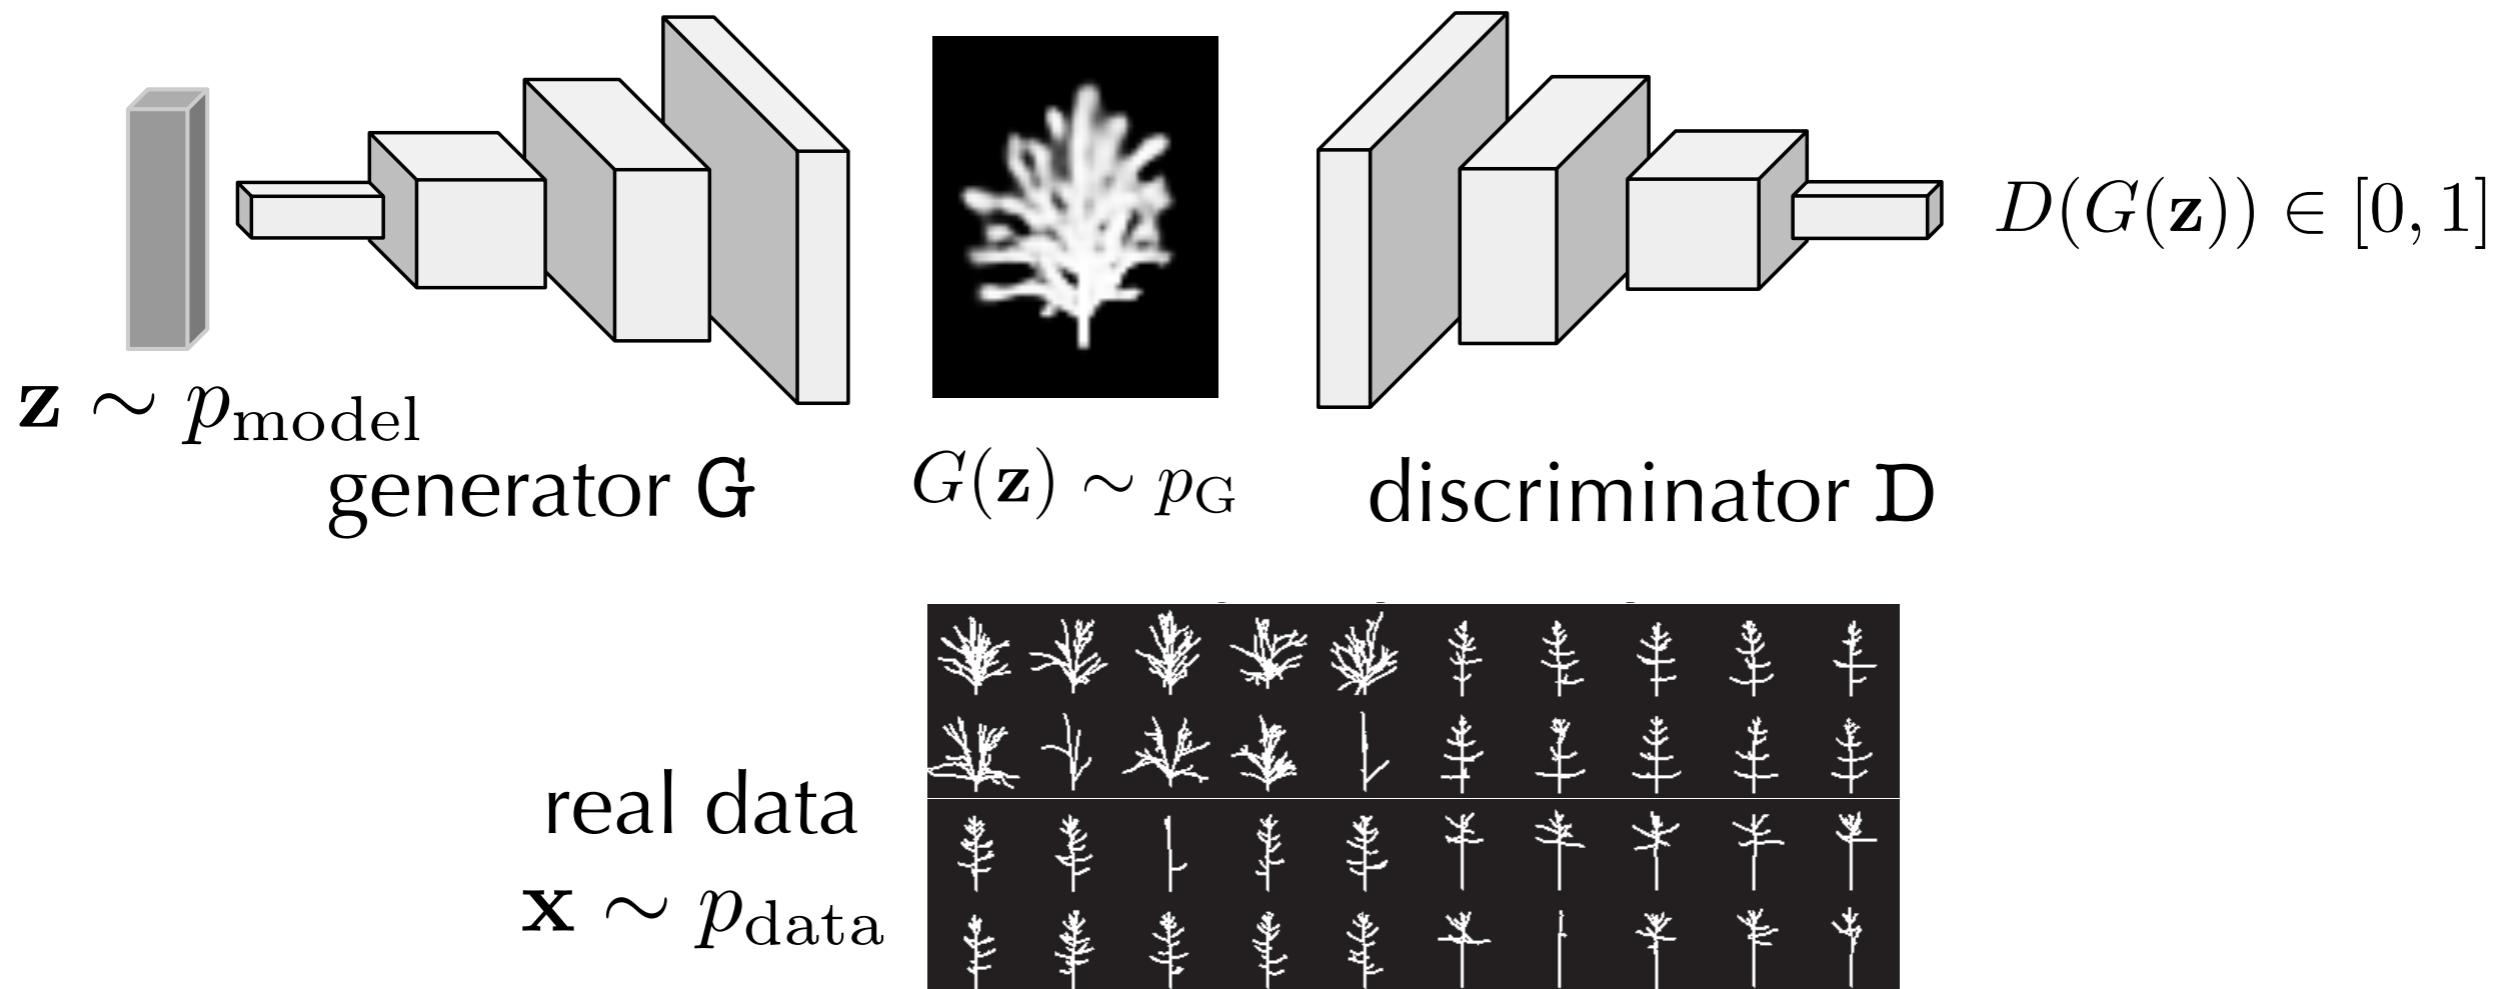
Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

Generative Adversarial Networks (GANs)

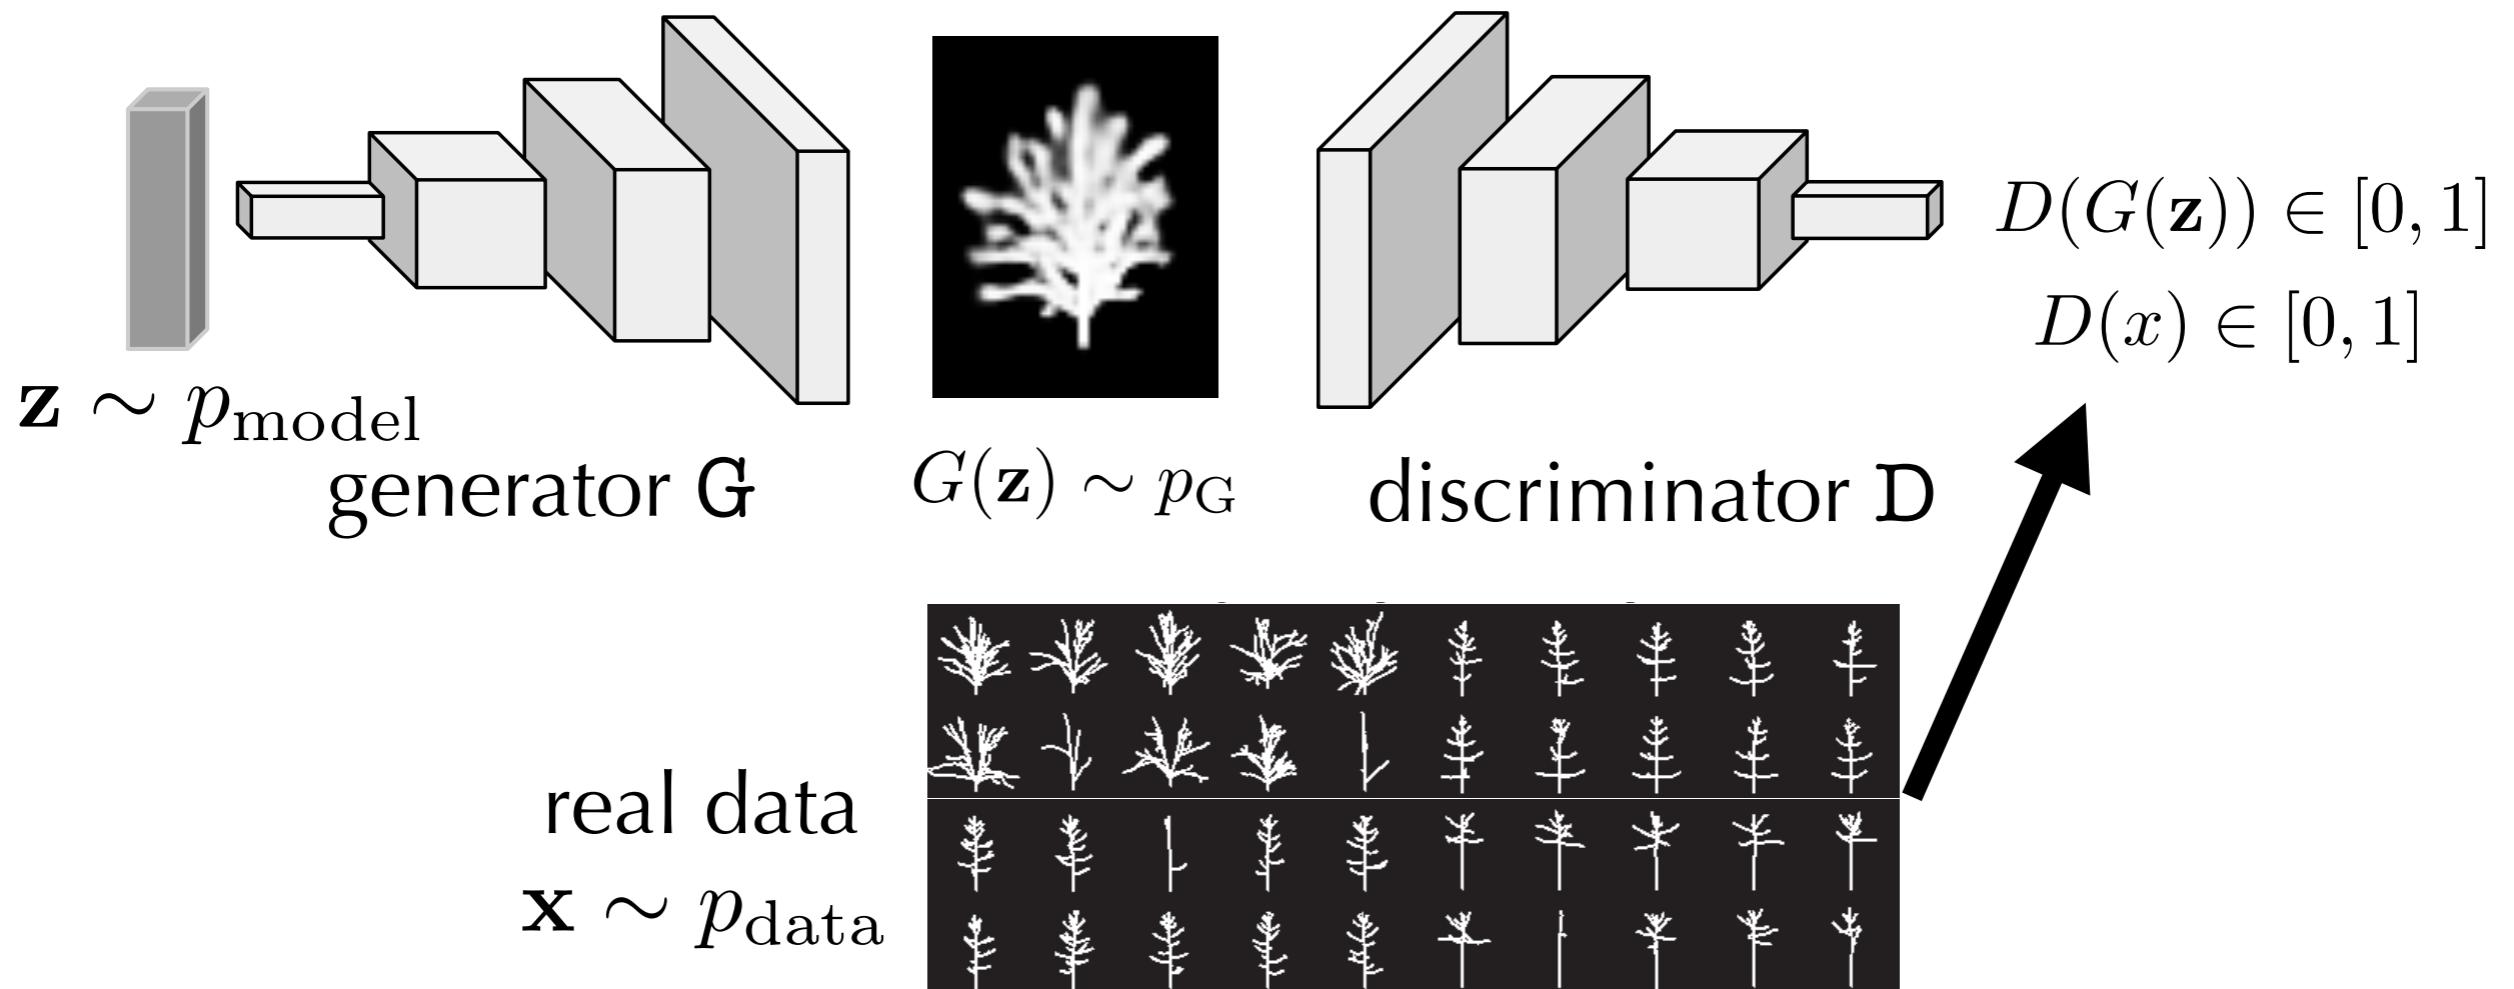
Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

Generative Adversarial Networks (GANs)

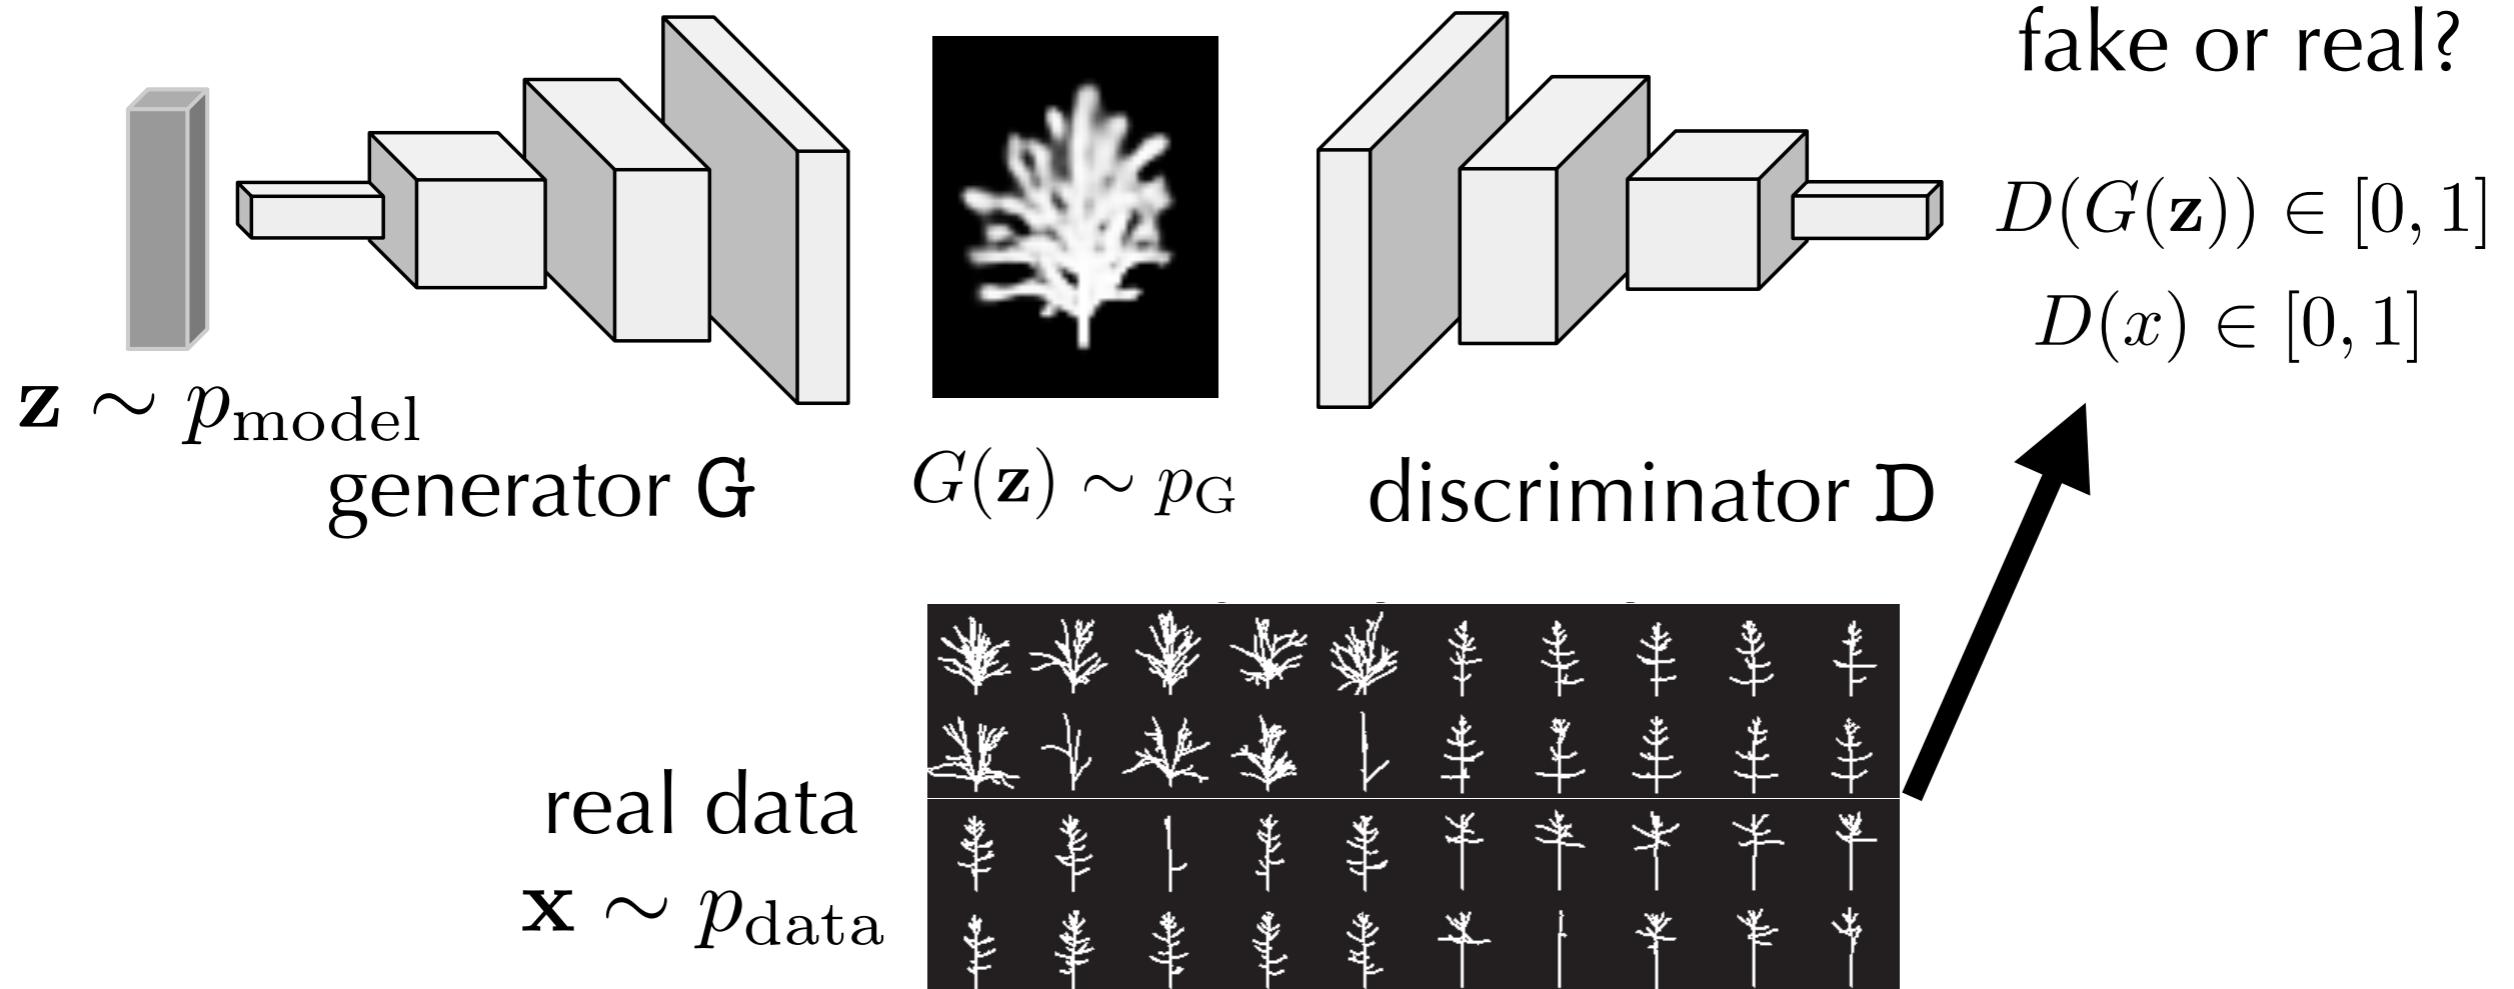
Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

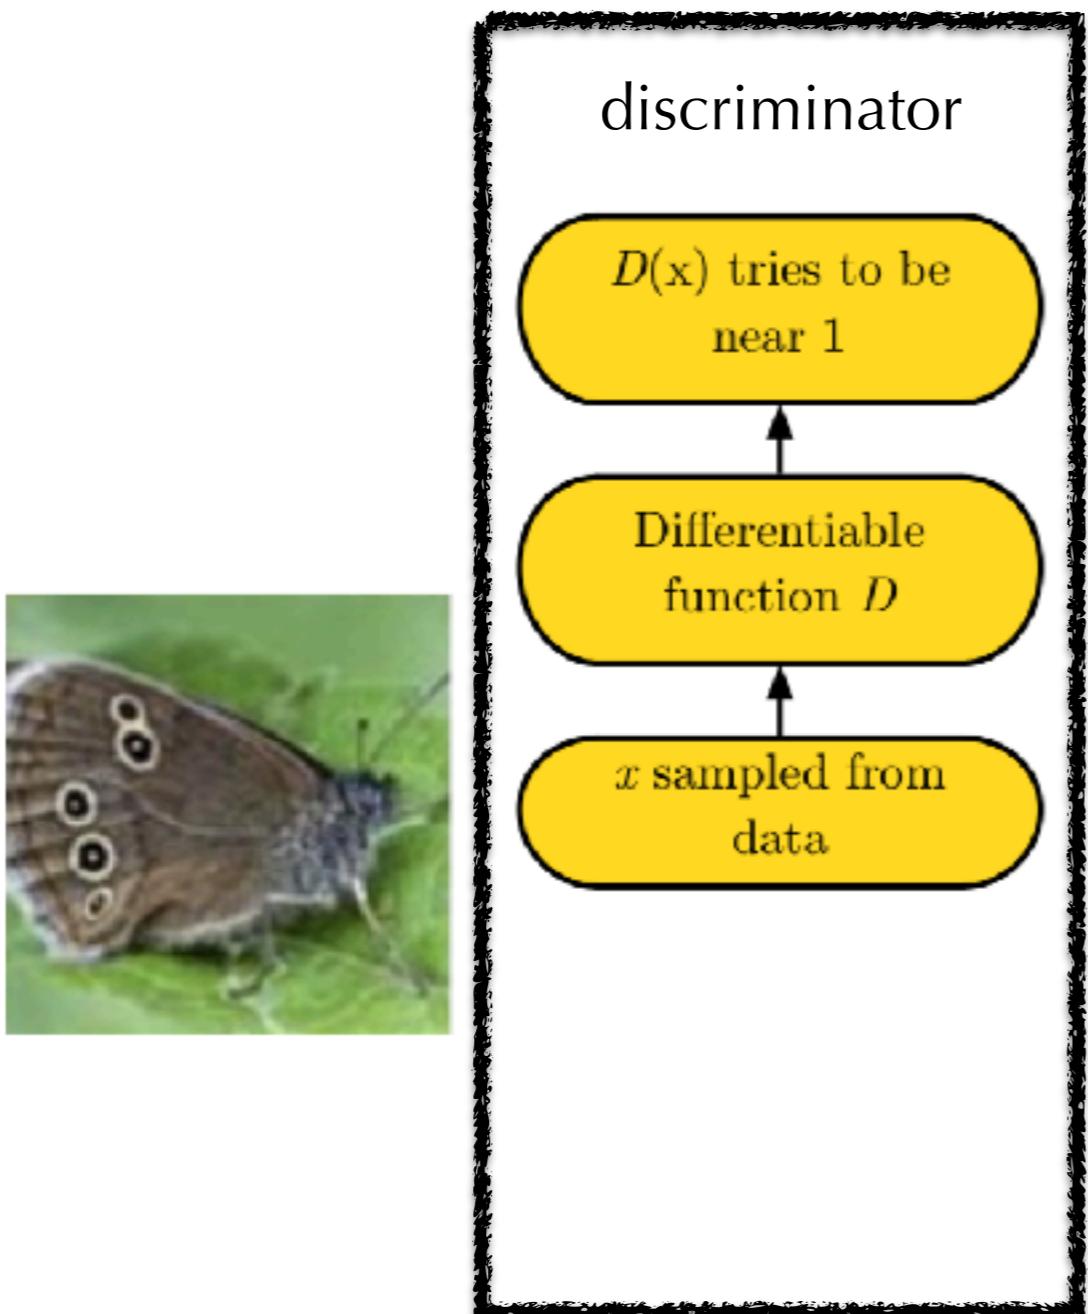
Generative Adversarial Networks (GANs)

Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

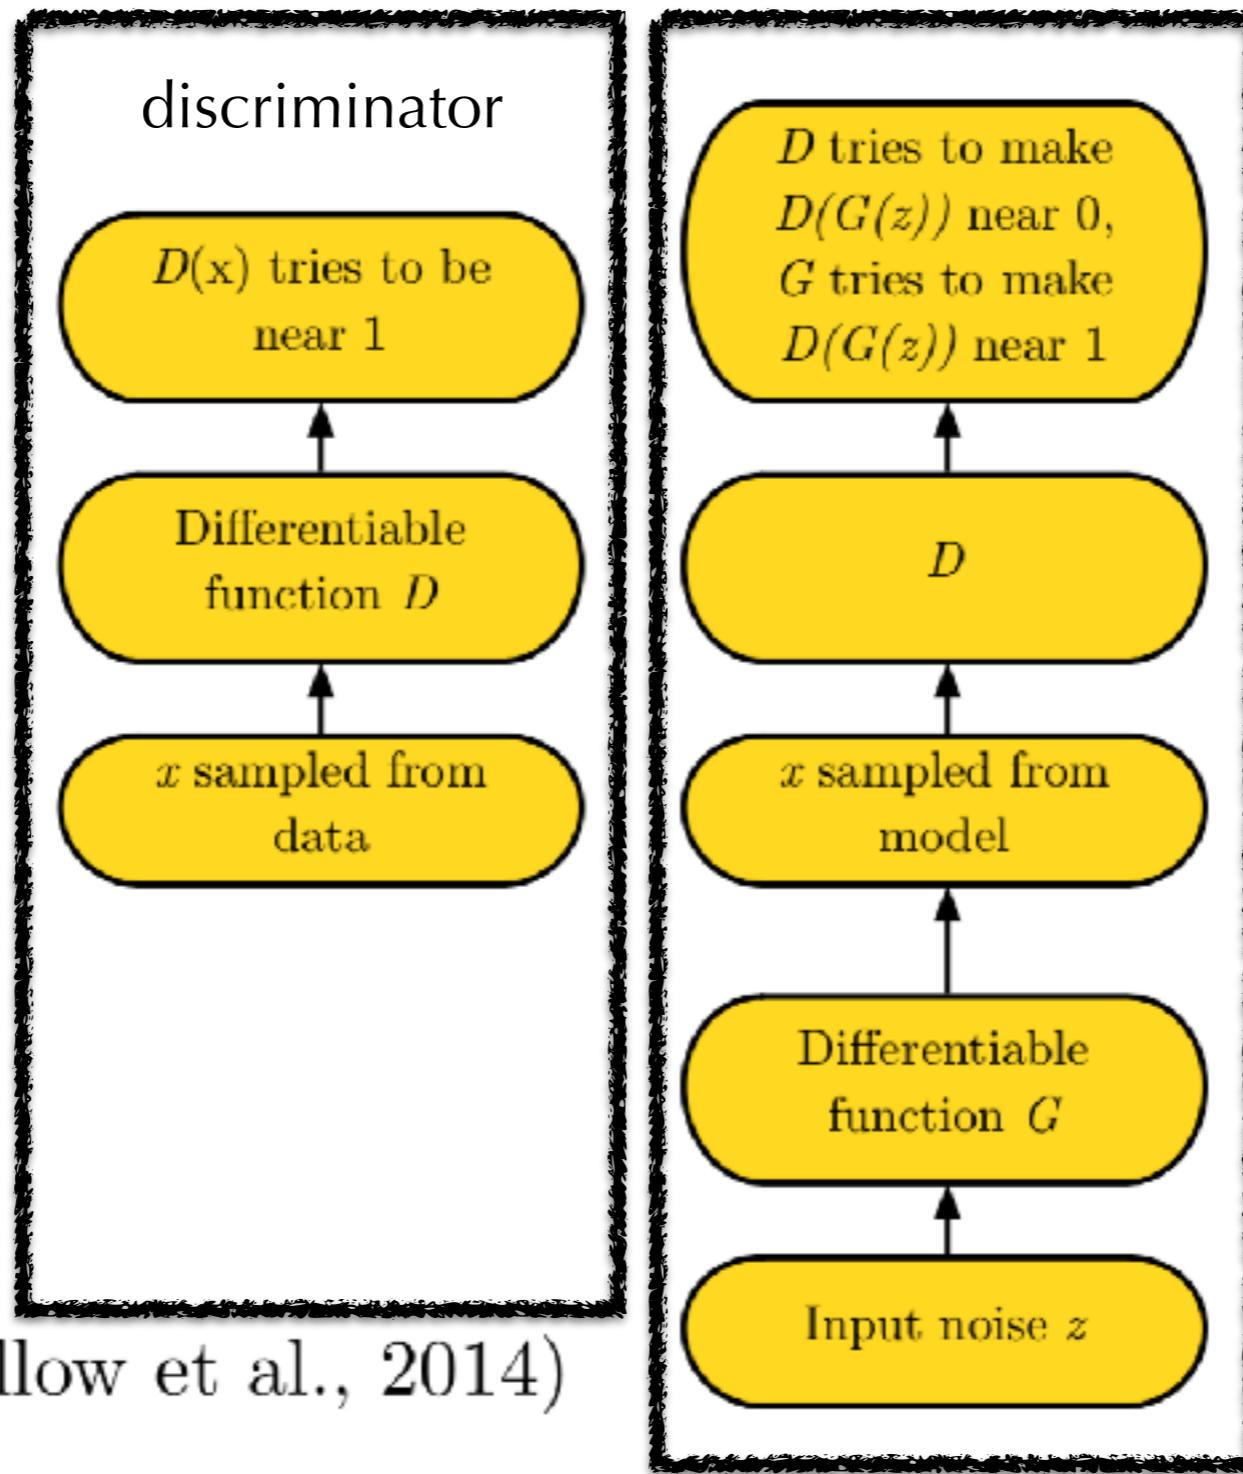
Generative Adversarial Networks (GANs)



(Goodfellow et al., 2014)

slide credit: Ian Goodfellow

Generative Adversarial Networks (GANs)



(Goodfellow et al., 2014)



generator

slide credit: Ian Goodfellow

Game Theoretic View

- Discriminator D tries to distinguish real and fake images

Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”

Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\max_D \quad \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\max_D \quad \underline{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))}$$

Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\min_G \max_D \quad \underline{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))}$$

Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\min_G \max_D \quad \underbrace{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))}_{\text{red line}}$$

Game Theoretic View

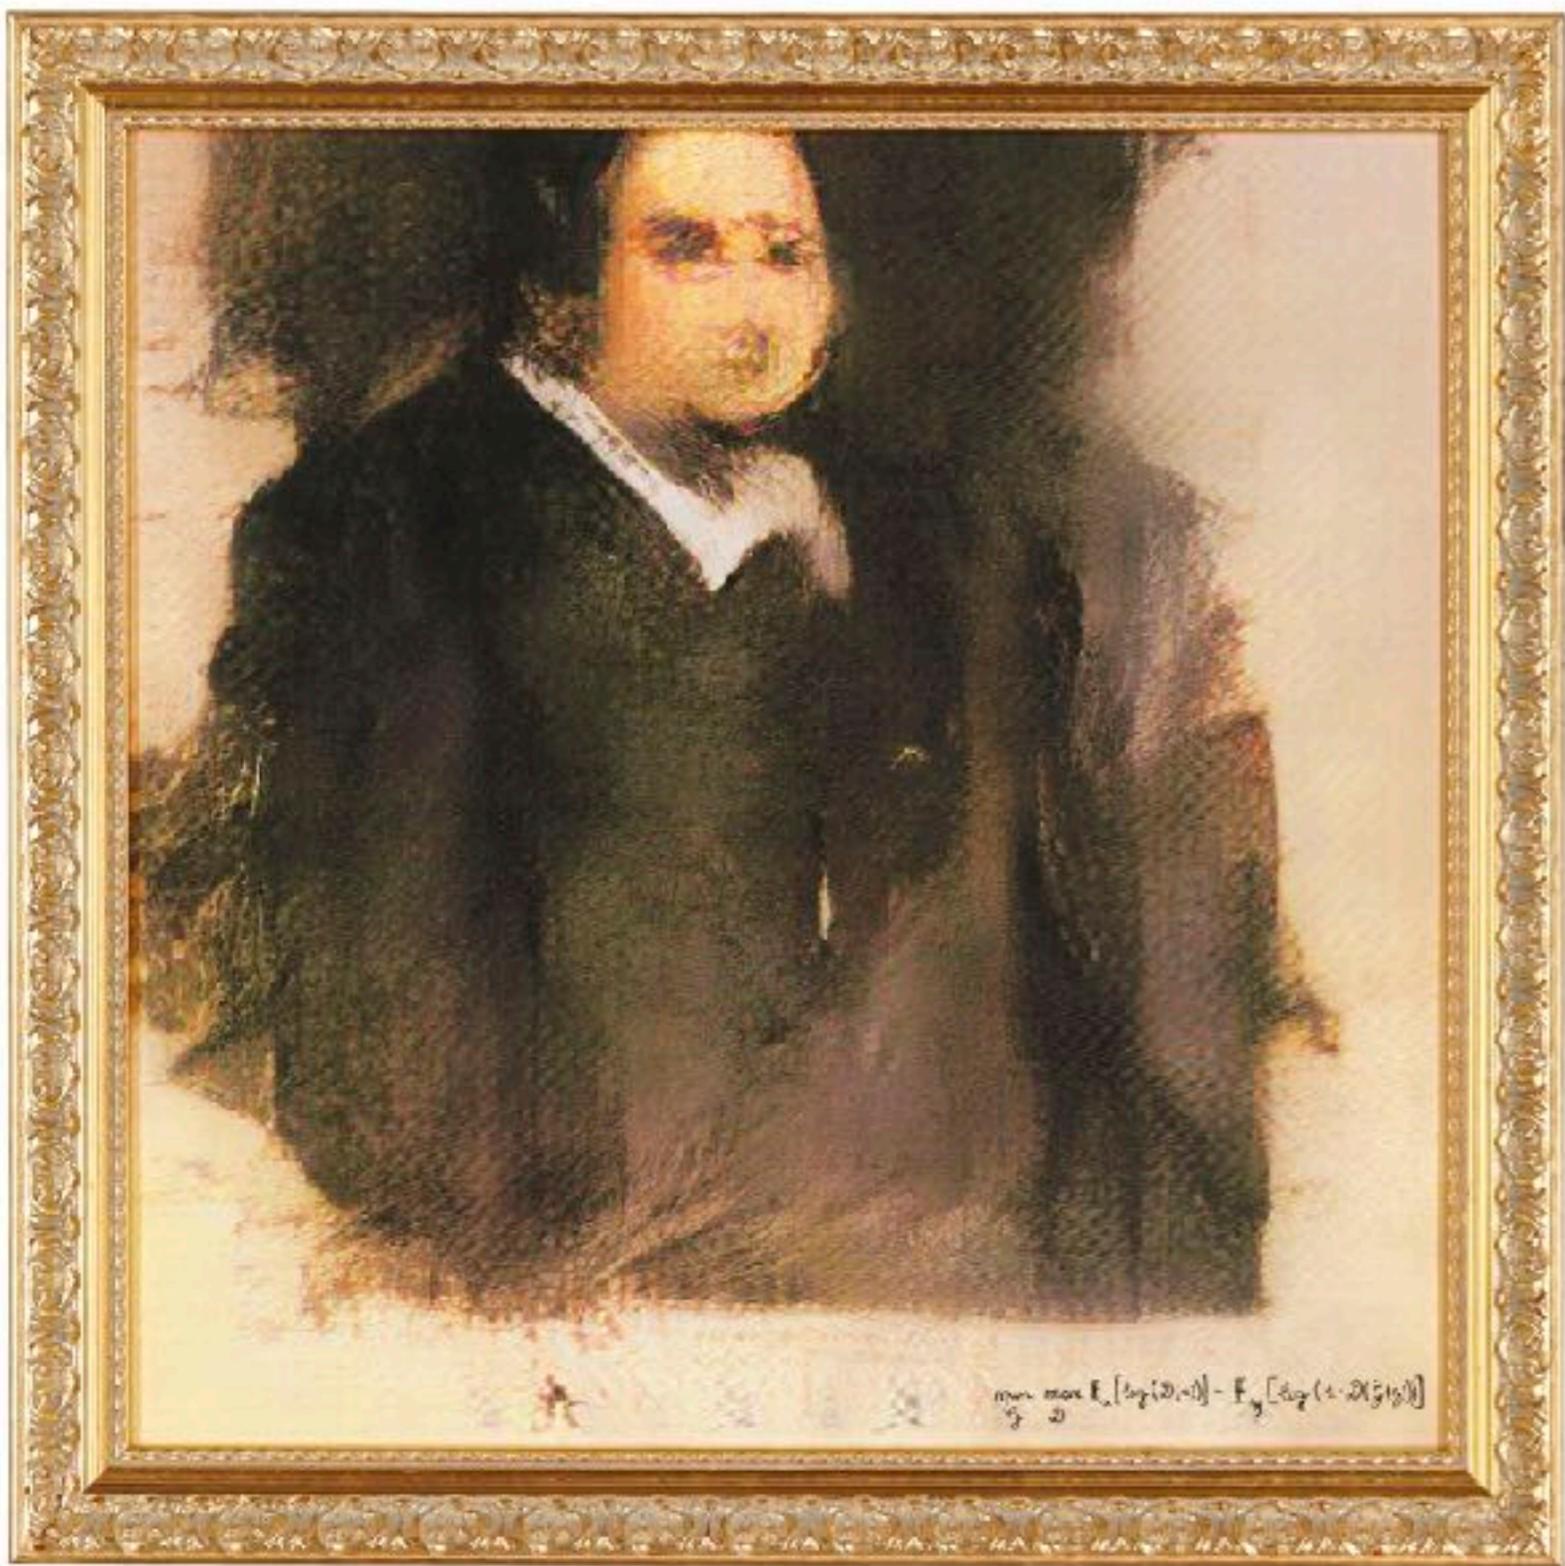
- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\min_G \max_D \quad \underbrace{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))}_{\text{red line}}$$

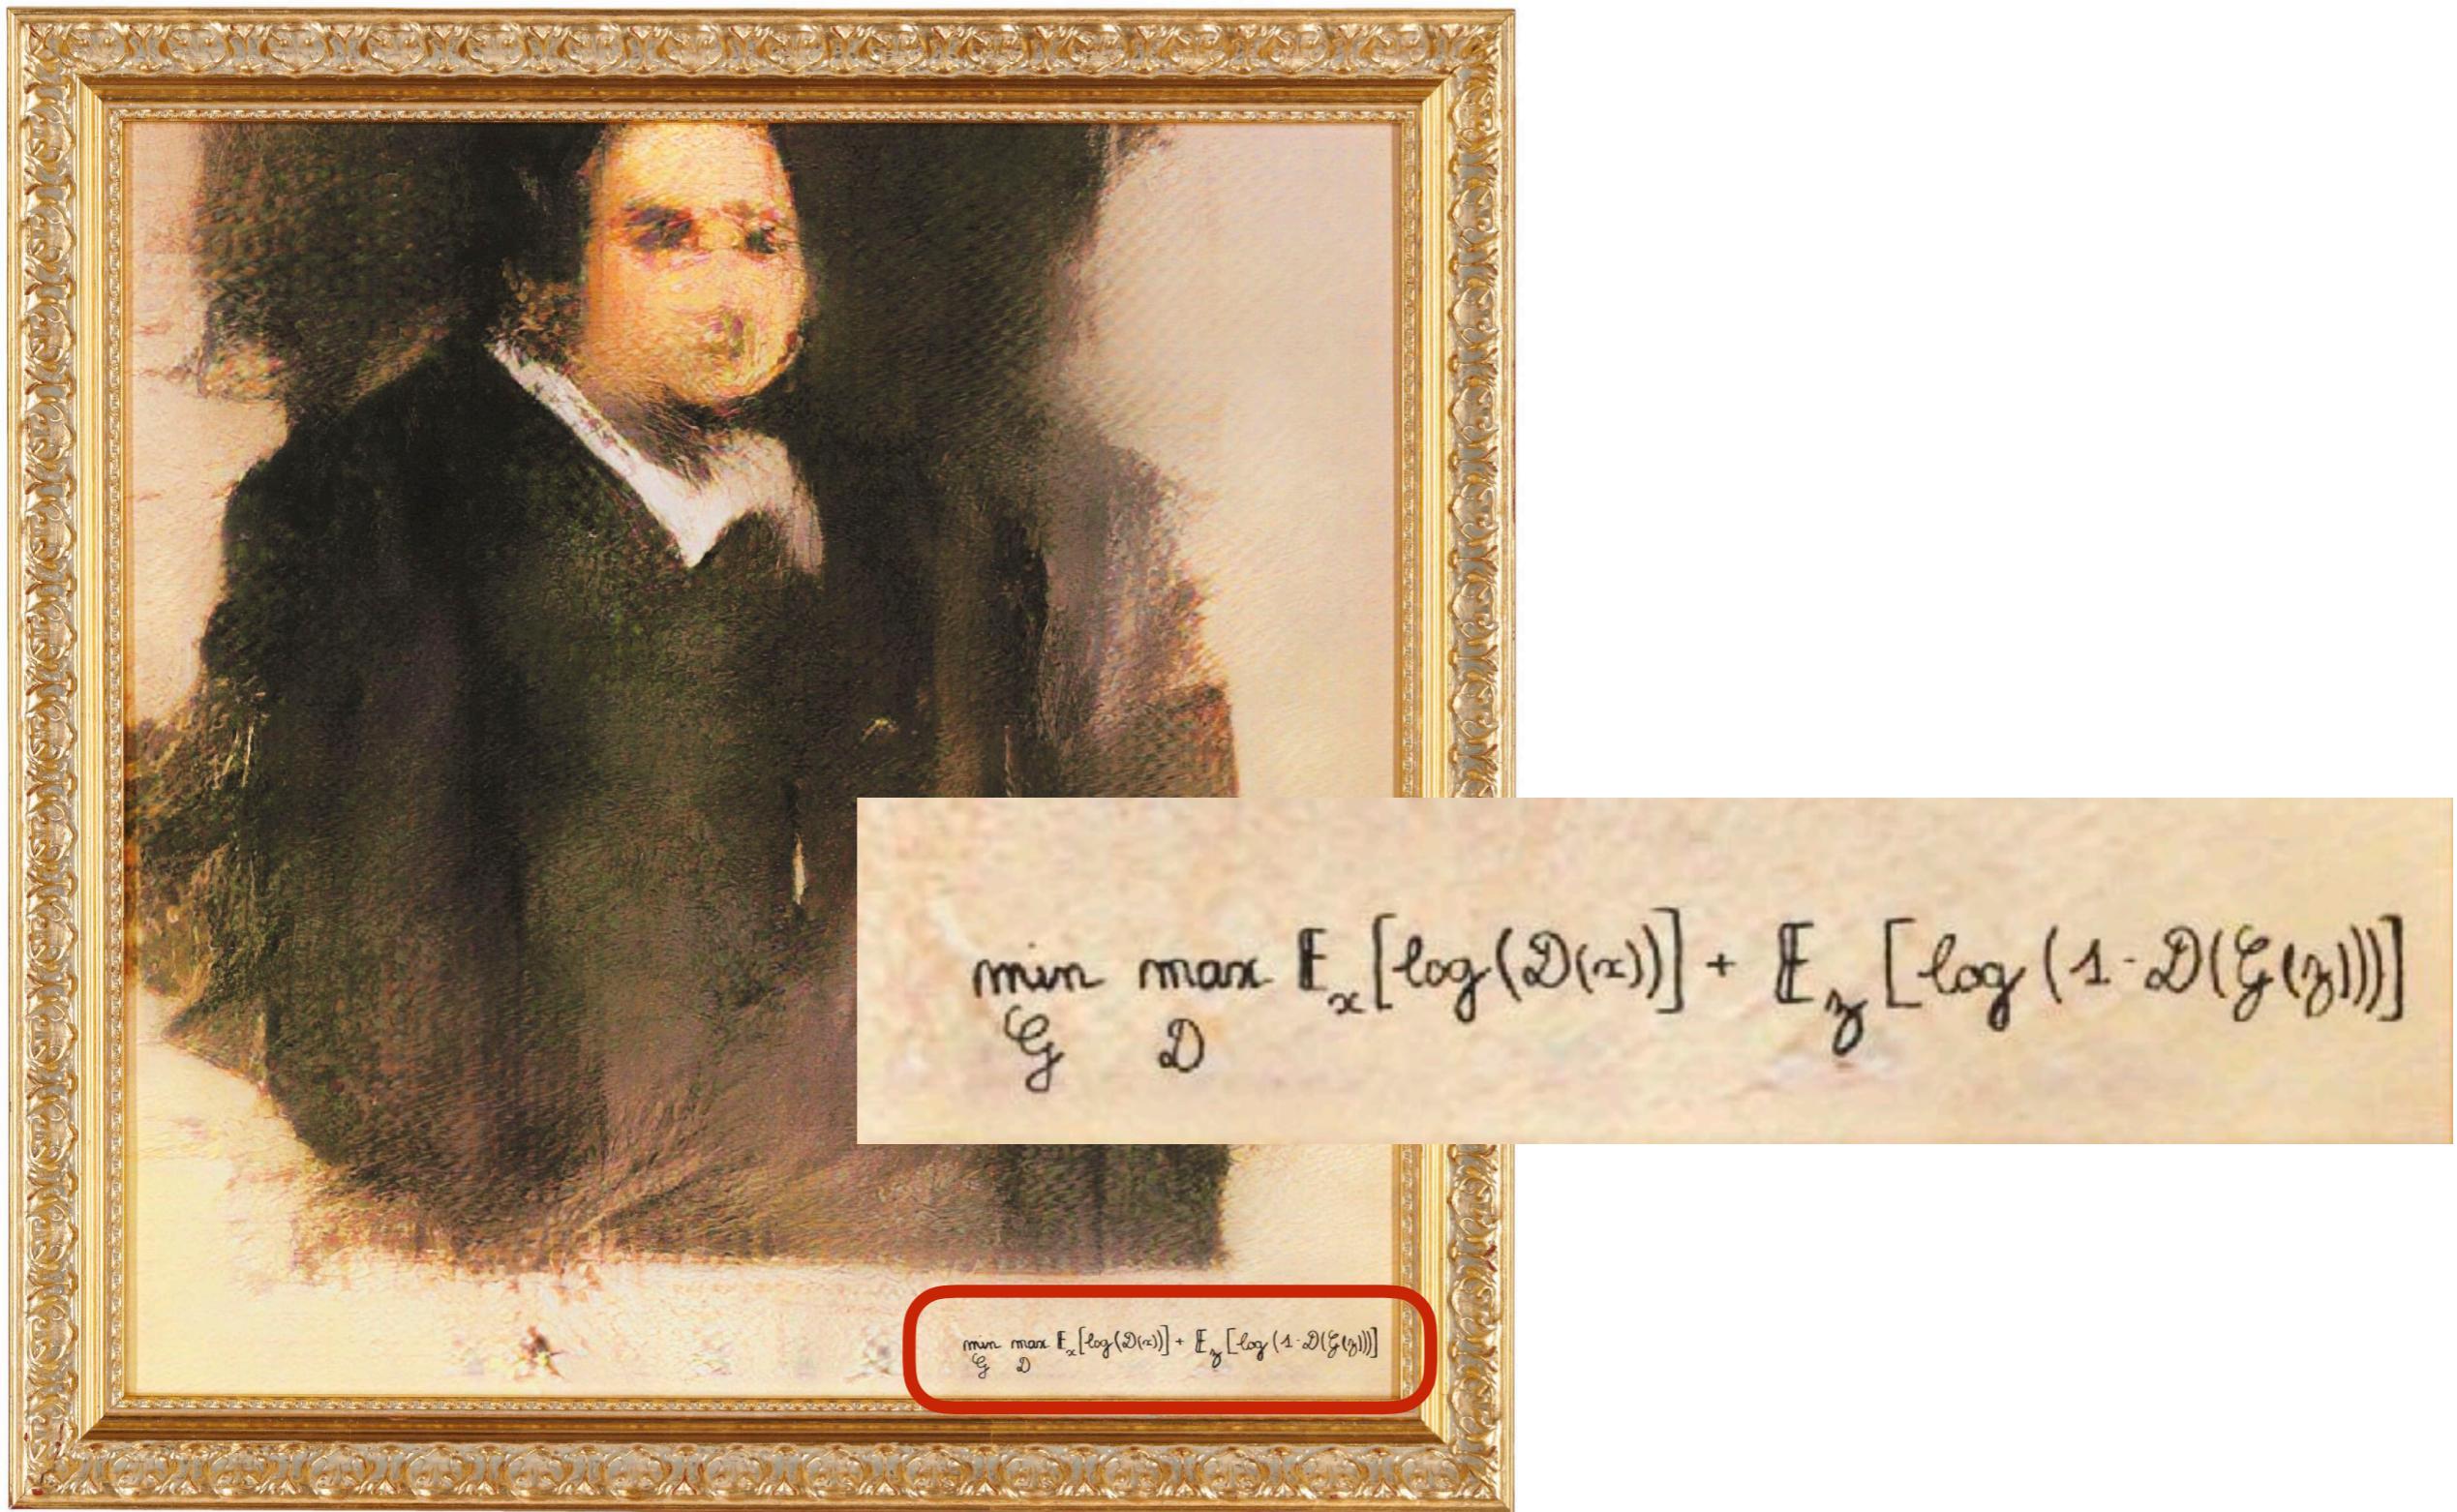
- Equilibrium / global optimum at $p_{\text{data}} = p_G$, i.e., when real and synthetic images are indistinguishable for D:

$$d(\mathbf{x}) = 1/2 \quad \forall \mathbf{x} \sim p_{\text{data}}, \mathbf{x} \sim p_G$$

Generating Art Revisited



Generating Art Revisited



Training GANs

```
for number of training iterations do  
    for  $k$  steps do
```

```
end for
```

Train discriminator D

Train generator G

```
end for
```

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Training GANs

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

end for

Train generator G

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Training GANs

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Training GANs is Hard

- Need to carefully **balance the generator and discriminator**

Training GANs is Hard

- Need to carefully **balance the generator and discriminator**
 - D becomes too powerful too soon → G cannot improve (gradient vanishes)

Training GANs is Hard

- Need to carefully **balance the generator and discriminator**
 - D becomes too powerful too soon → G cannot improve (gradient vanishes)
 - D not powerful enough → G does not improve

Training GANs is Hard

- Need to carefully **balance the generator and discriminator**
 - D becomes too powerful too soon → G cannot improve (gradient vanishes)
 - D not powerful enough → G does not improve
 - Black magic: Schedule training of generator and discriminator

Training GANs is Hard

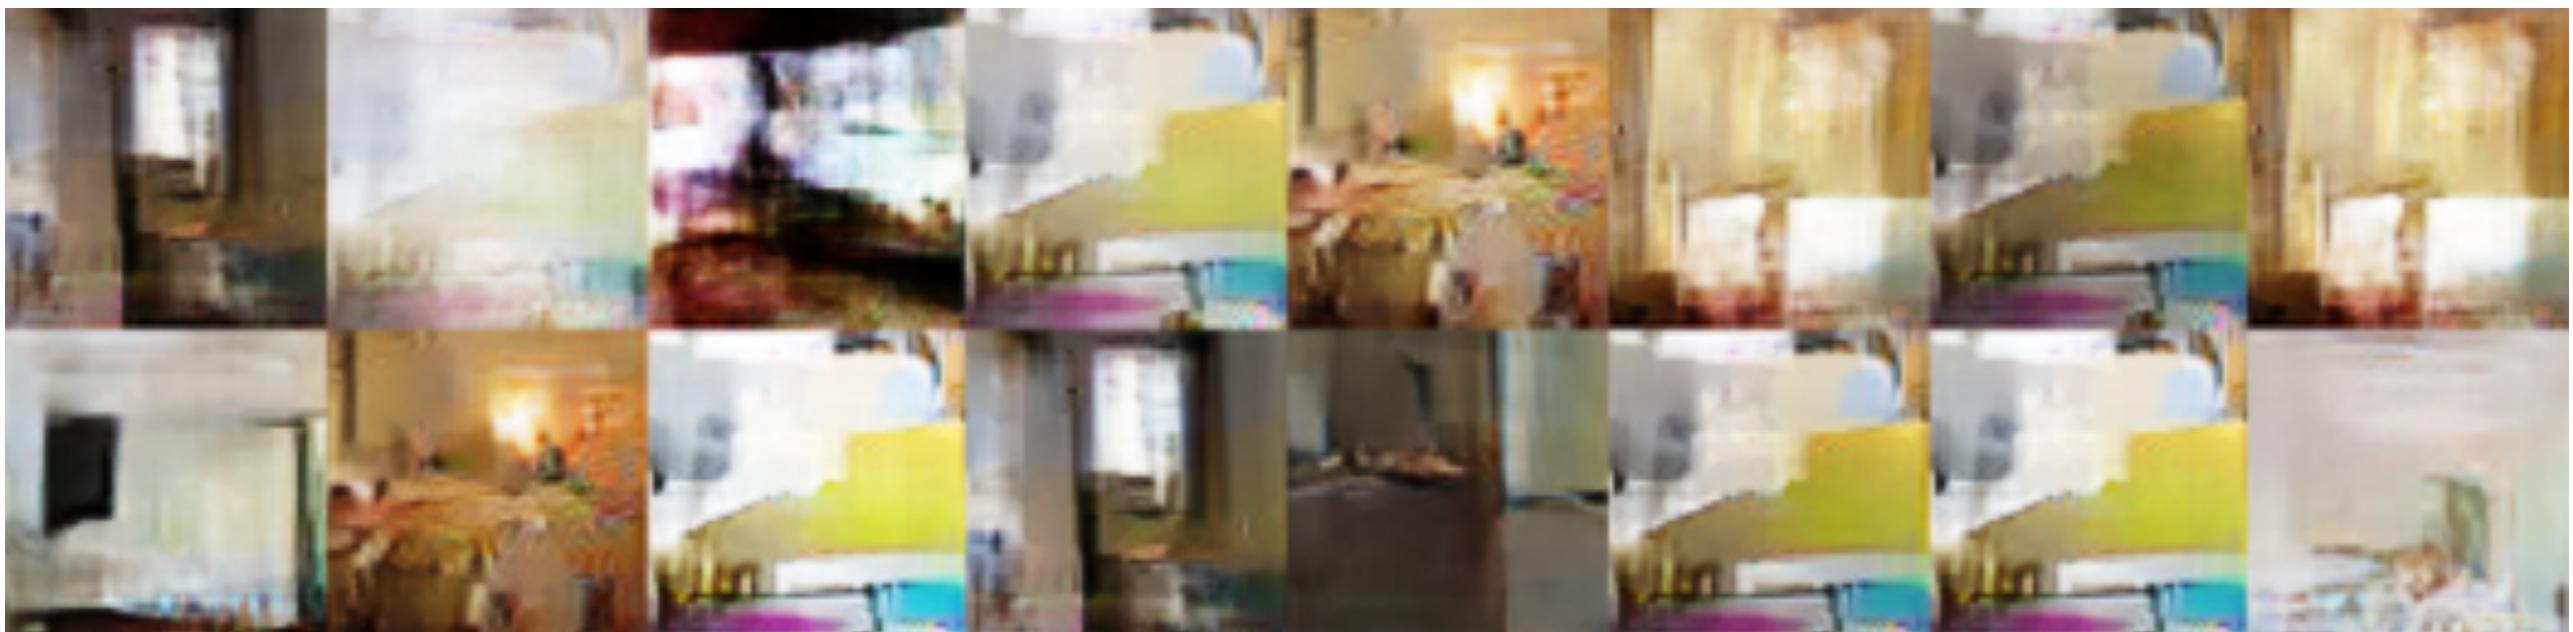
- Need to carefully **balance the generator and discriminator**
 - D becomes too powerful too soon → G cannot improve (gradient vanishes)
 - D not powerful enough → G does not improve
 - Black magic: Schedule training of generator and discriminator
- **Divergence:** GANs might not train at all

Training GANs is Hard

- Need to carefully **balance the generator and discriminator**
 - D becomes too powerful too soon → G cannot improve (gradient vanishes)
 - D not powerful enough → G does not improve
 - Black magic: Schedule training of generator and discriminator
- **Divergence:** GANs might not train at all
- **Mode collapse:** Generated images are not diverse

Mode Collapse

Generated images span only small subspace



[Arjovsky et al., Wasserstein GAN, arXiv:1701.07875]

Training GANs is Hard

- Need to carefully **balance the generator and discriminator**
 - D becomes too powerful too soon → G cannot improve (gradient vanishes)
 - D not powerful enough → G does not improve
 - Black magic: Schedule training of generator and discriminator
- **Divergence:** GANs might not train at all
- **Mode collapse:** Generated images are not diverse
- List of tips & tricks for training here: https://github.com/soumith/talks/blob/master/2017-ICCV_Venice/How_To_Train_a_GAN.pdf

Why is Training GANs so Hard?

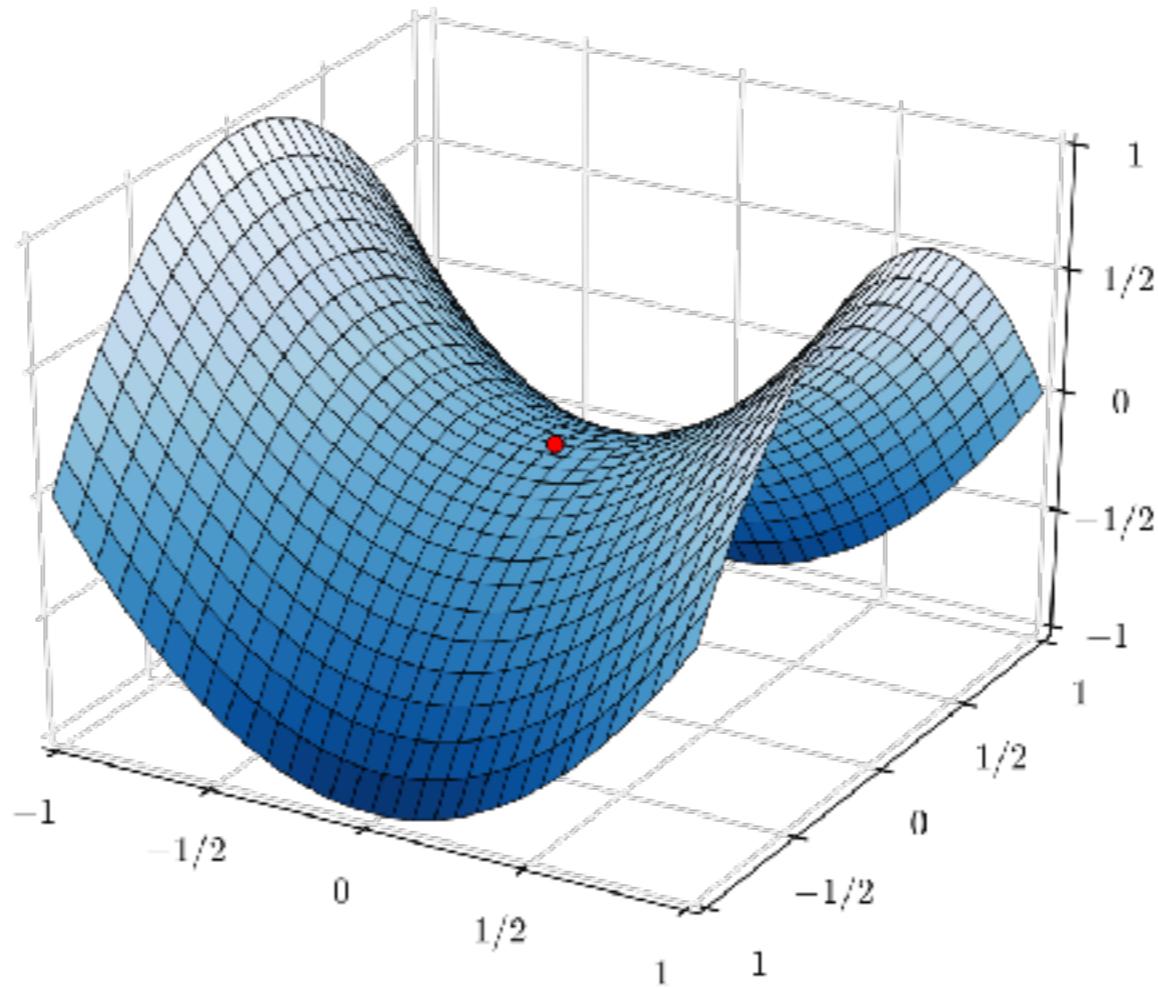


image: © Nicoguaro (CC Attribution Unported 3.0)

GAN equilibria are saddle points

$$\min_G \max_D \quad \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

- Optimal discriminator D^* for given G :

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$$

Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

- Optimal discriminator D^* for given G :

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$$

- GAN objective becomes

$$\min_G \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log \left(\frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right) + \mathbb{E}_{\mathbf{x} \sim p_G} \log \left(\frac{p_G(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right)$$

Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

- Optimal discriminator D^* for given G :

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$$

- GAN objective becomes

$$\min_G \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log \left(\frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right) + \mathbb{E}_{\mathbf{x} \sim p_G} \log \left(\frac{p_G(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right)$$

- Utilizing the Kullback-Leibler (KL) divergence

Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

- Optimal discriminator D^* for given G :

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$$

- GAN objective becomes

$$\min_G \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log \left(\frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right) + \mathbb{E}_{\mathbf{x} \sim p_G} \log \left(\frac{p_G(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right)$$

- Utilizing the Kullback-Leibler (KL) divergence

$$\min_G \text{KL} \left(p_{\text{data}} \parallel \frac{p_{\text{data}} + p_G}{2} \right) + \text{KL} \left(p_G \parallel \frac{p_{\text{data}} + p_G}{2} \right) - \log(4)$$

Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

- Optimal discriminator D^* for given G :

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$$

- GAN objective becomes

$$\min_G \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log \left(\frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right) + \mathbb{E}_{\mathbf{x} \sim p_G} \log \left(\frac{p_G(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right)$$

- Utilizing the Kullback-Leibler (KL) divergence

$$\min_G \text{KL} \left(p_{\text{data}} \parallel \frac{p_{\text{data}} + p_G}{2} \right) + \text{KL} \left(p_G \parallel \frac{p_{\text{data}} + p_G}{2} \right) - \log(4)$$

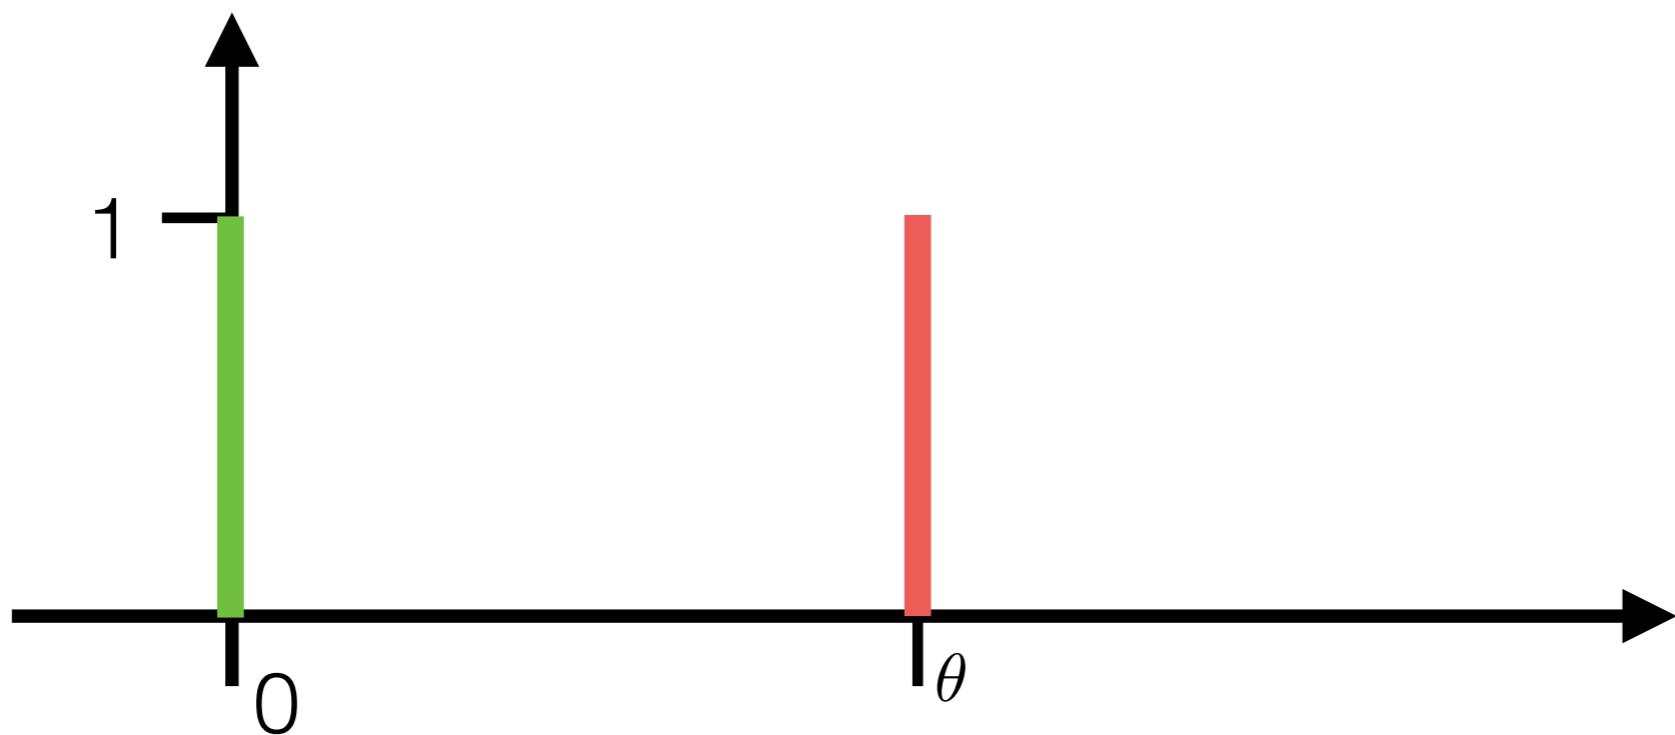

$2 \cdot \text{JSD}(p_{\text{data}} \parallel p_G)$ Jensen-Shannon divergence

Why is Training GANs so Hard?

Problem: Jensen-Shannon divergence gradient not always useful → problem when using gradient descent

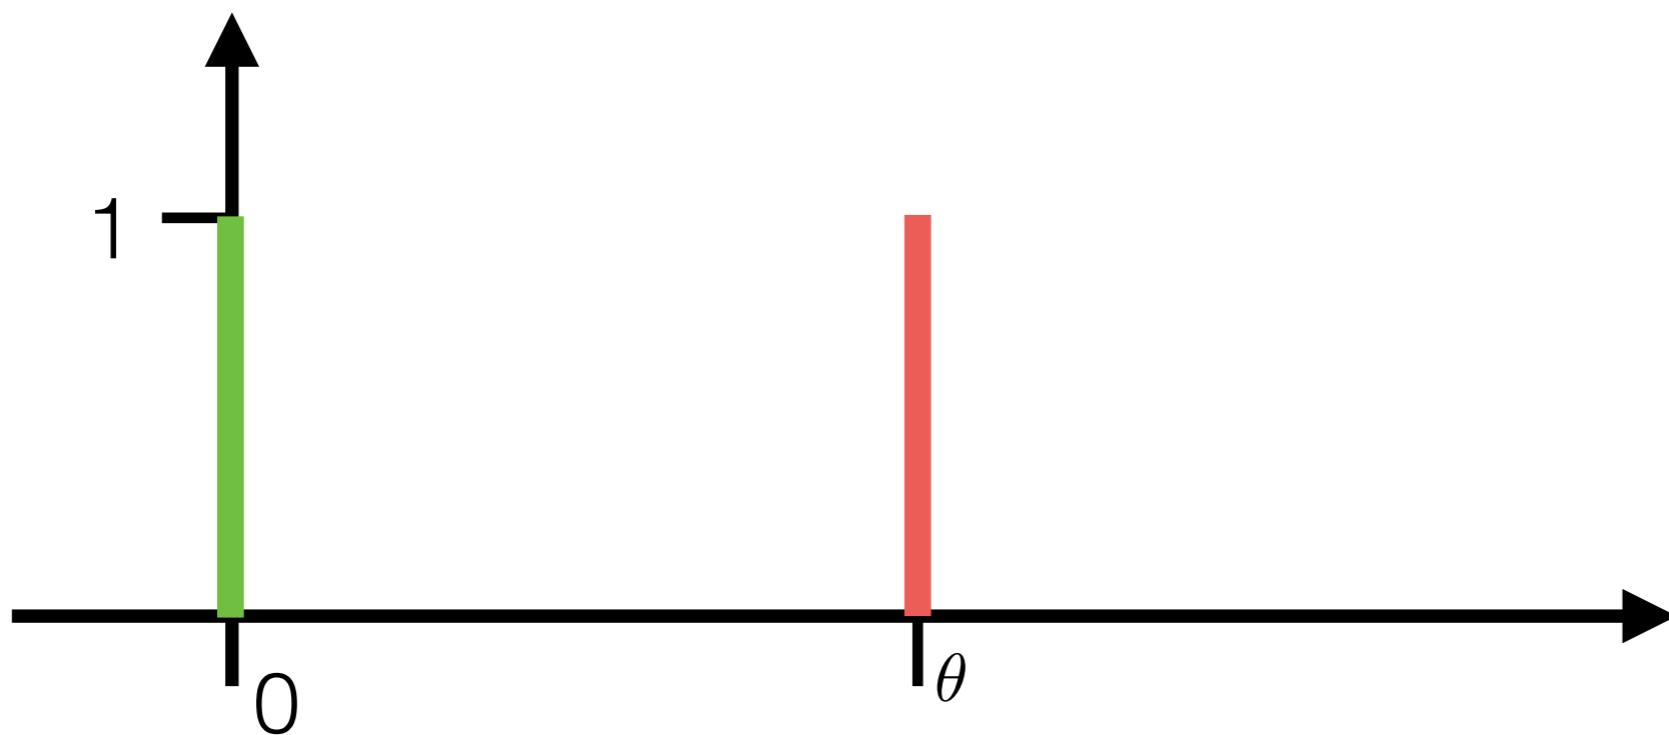
Why is Training GANs so Hard?

Problem: Jensen-Shannon divergence gradient not always useful → problem when using gradient descent



Why is Training GANs so Hard?

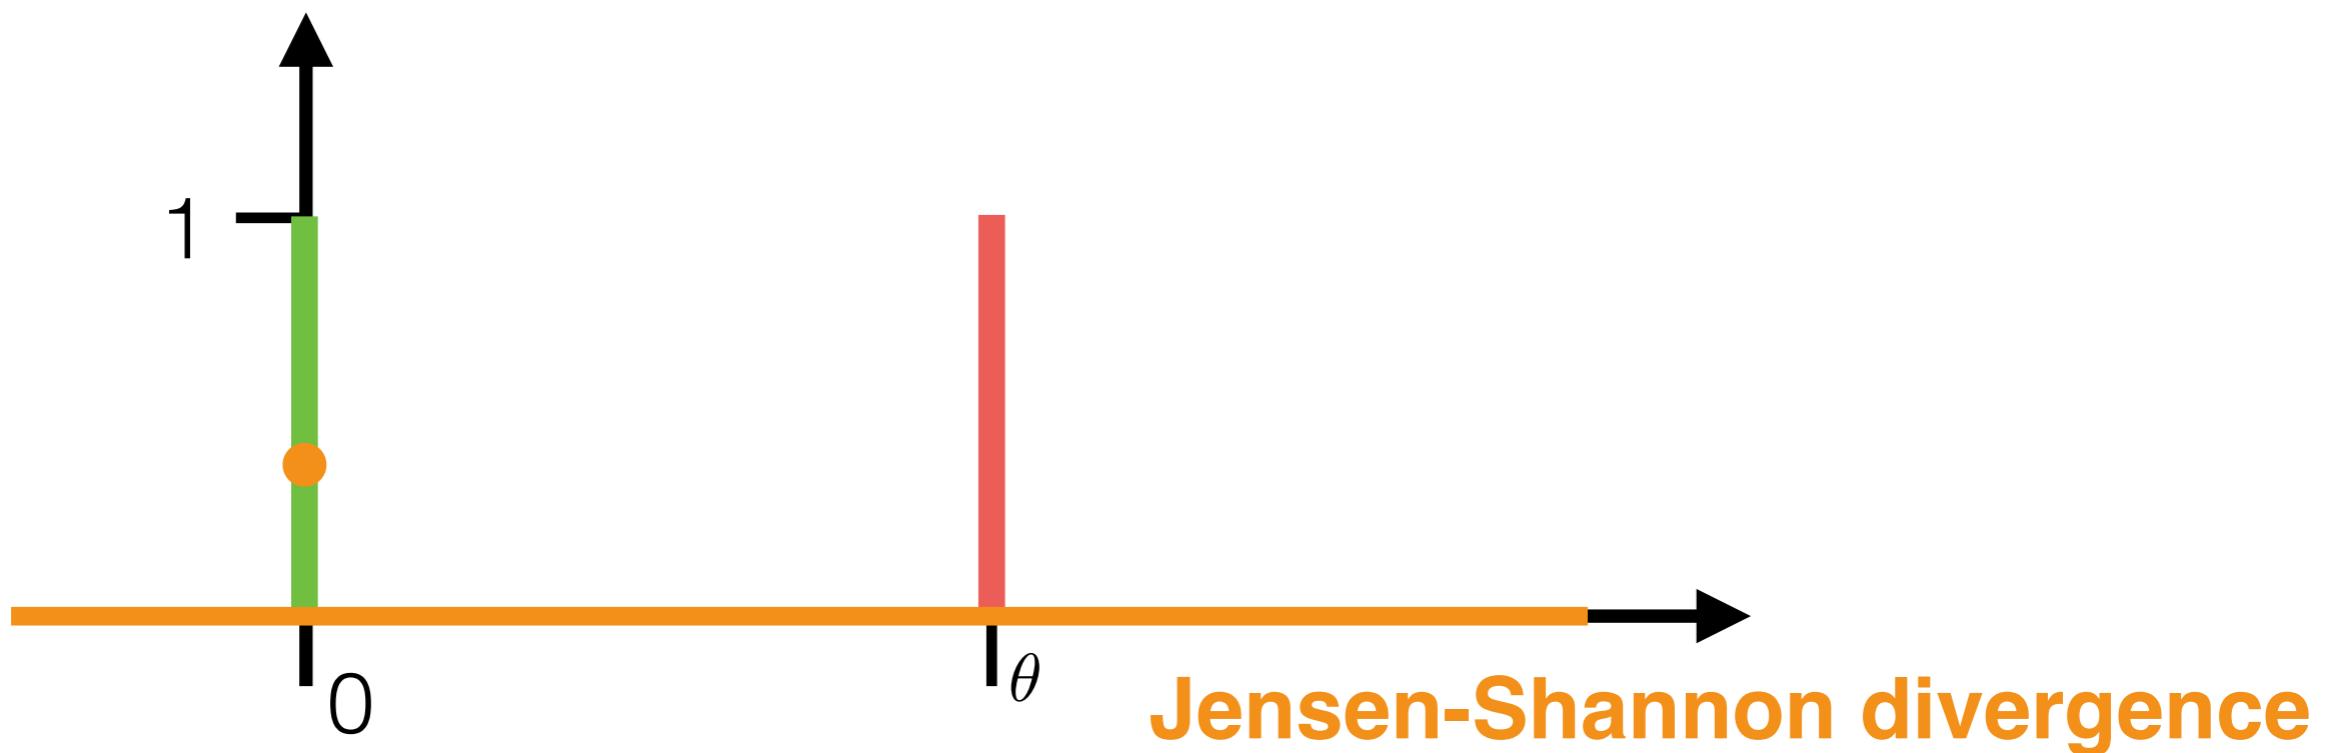
Problem: Jensen-Shannon divergence gradient not always useful → problem when using gradient descent



Simple one-parameter distribution $p_\theta(z) = (\theta, z)$, with $z \sim U[0, 1]$

Why is Training GANs so Hard?

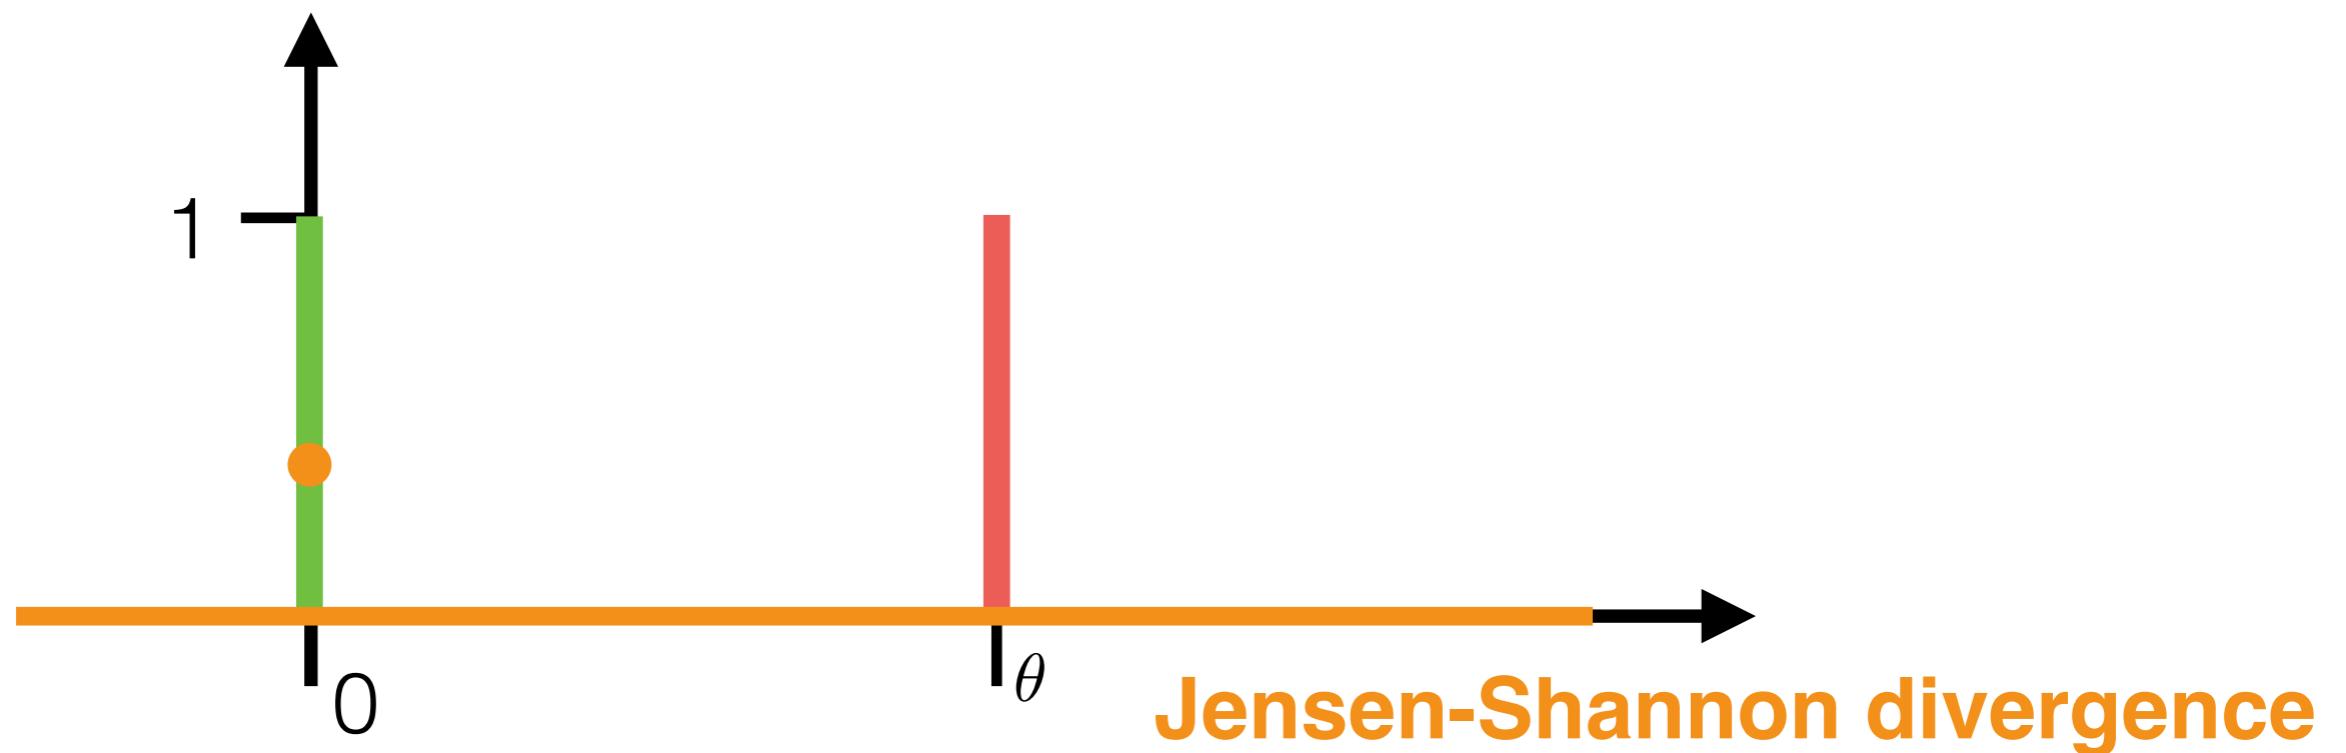
Problem: Jensen-Shannon divergence gradient not always useful → problem when using gradient descent



Simple one-parameter distribution $p_\theta(z) = (\theta, z)$, with $z \sim U[0, 1]$

Why is Training GANs so Hard?

Problem: Jensen-Shannon divergence gradient not always useful → problem when using gradient descent

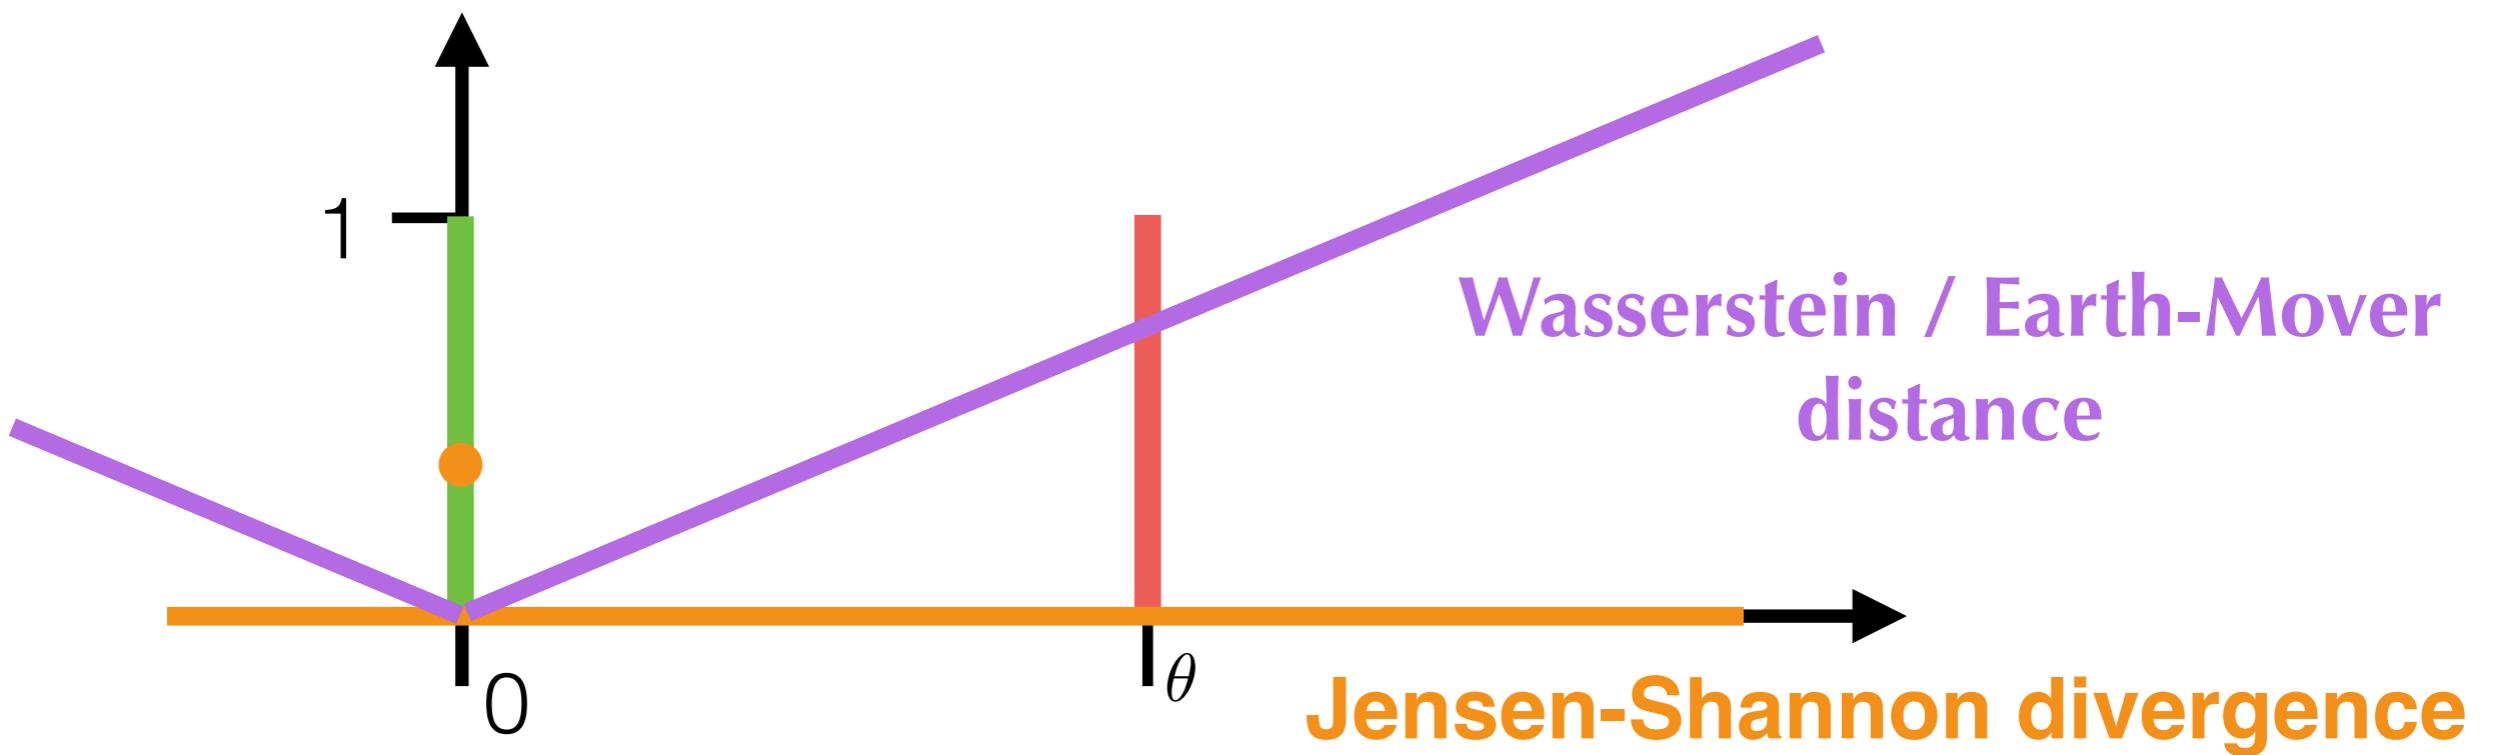


Simple one-parameter distribution $p_\theta(z) = (\theta, z)$, with $z \sim U[0, 1]$

Practical relevance: Operating on disjoint manifolds

Why is Training GANs so Hard?

Problem: Jensen-Shannon divergence gradient not always useful → problem when using gradient descent

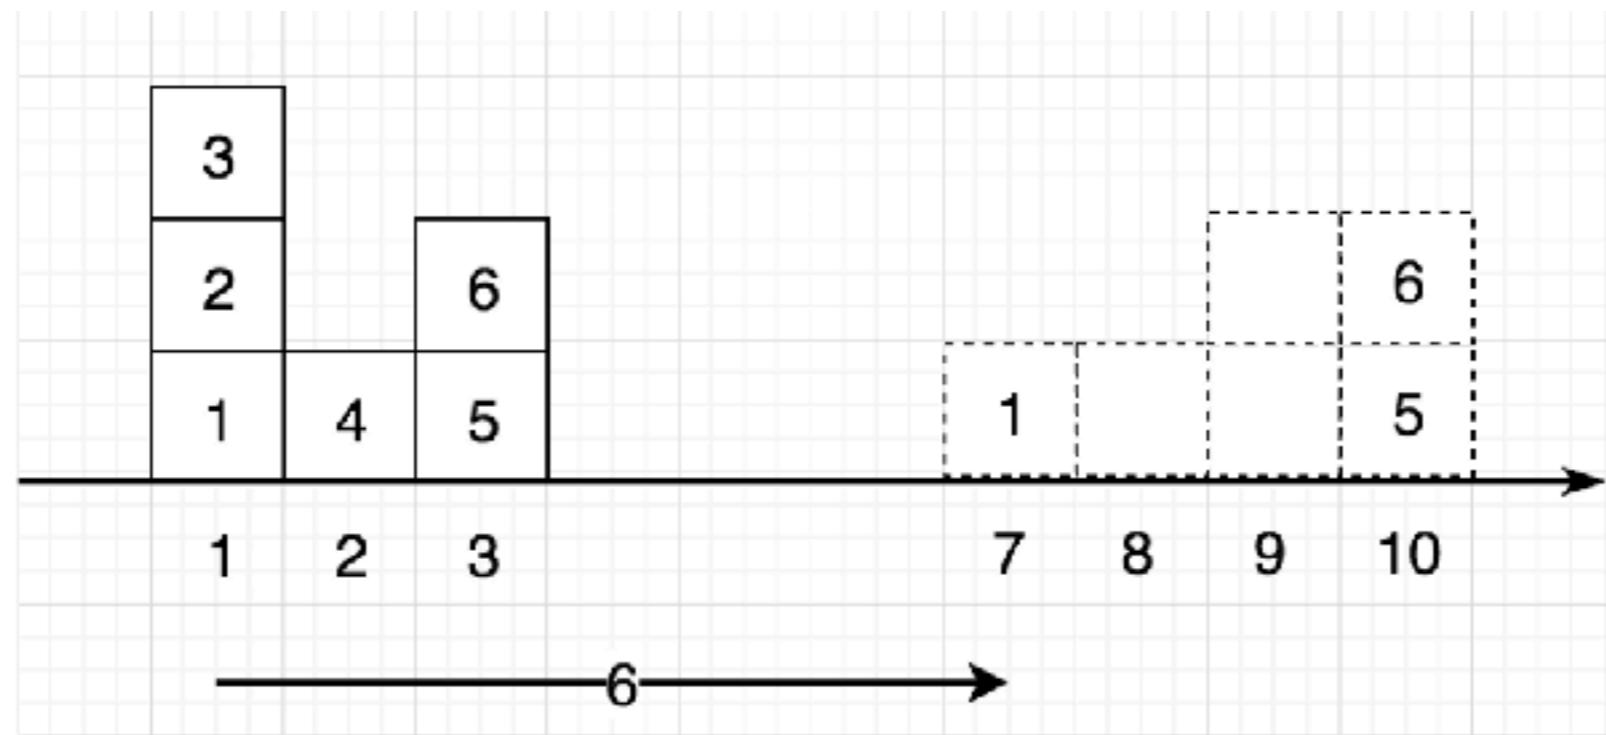


Simple one-parameter distribution $p_\theta(z) = (\theta, z)$, with $z \sim U[0, 1]$

Practical relevance: Operating on disjoint manifolds

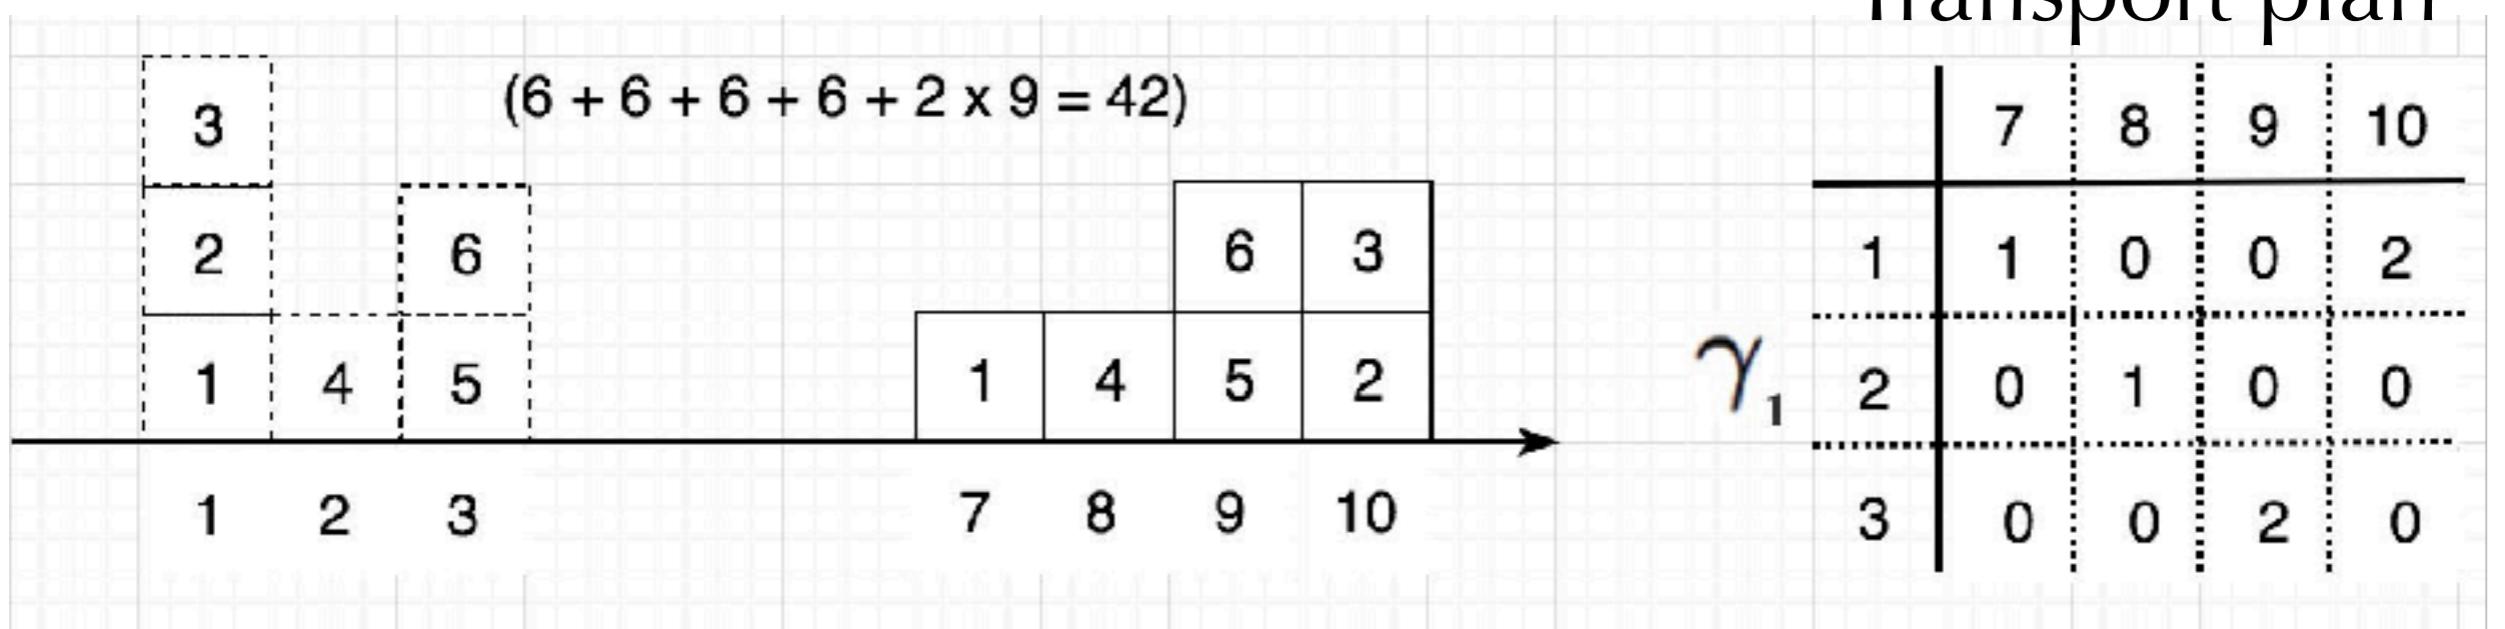
Wasserstein GANs (WGANS)

Earth-Mover distance: Minimum amount of work to move weight from $p(x)$ to $q(x)$



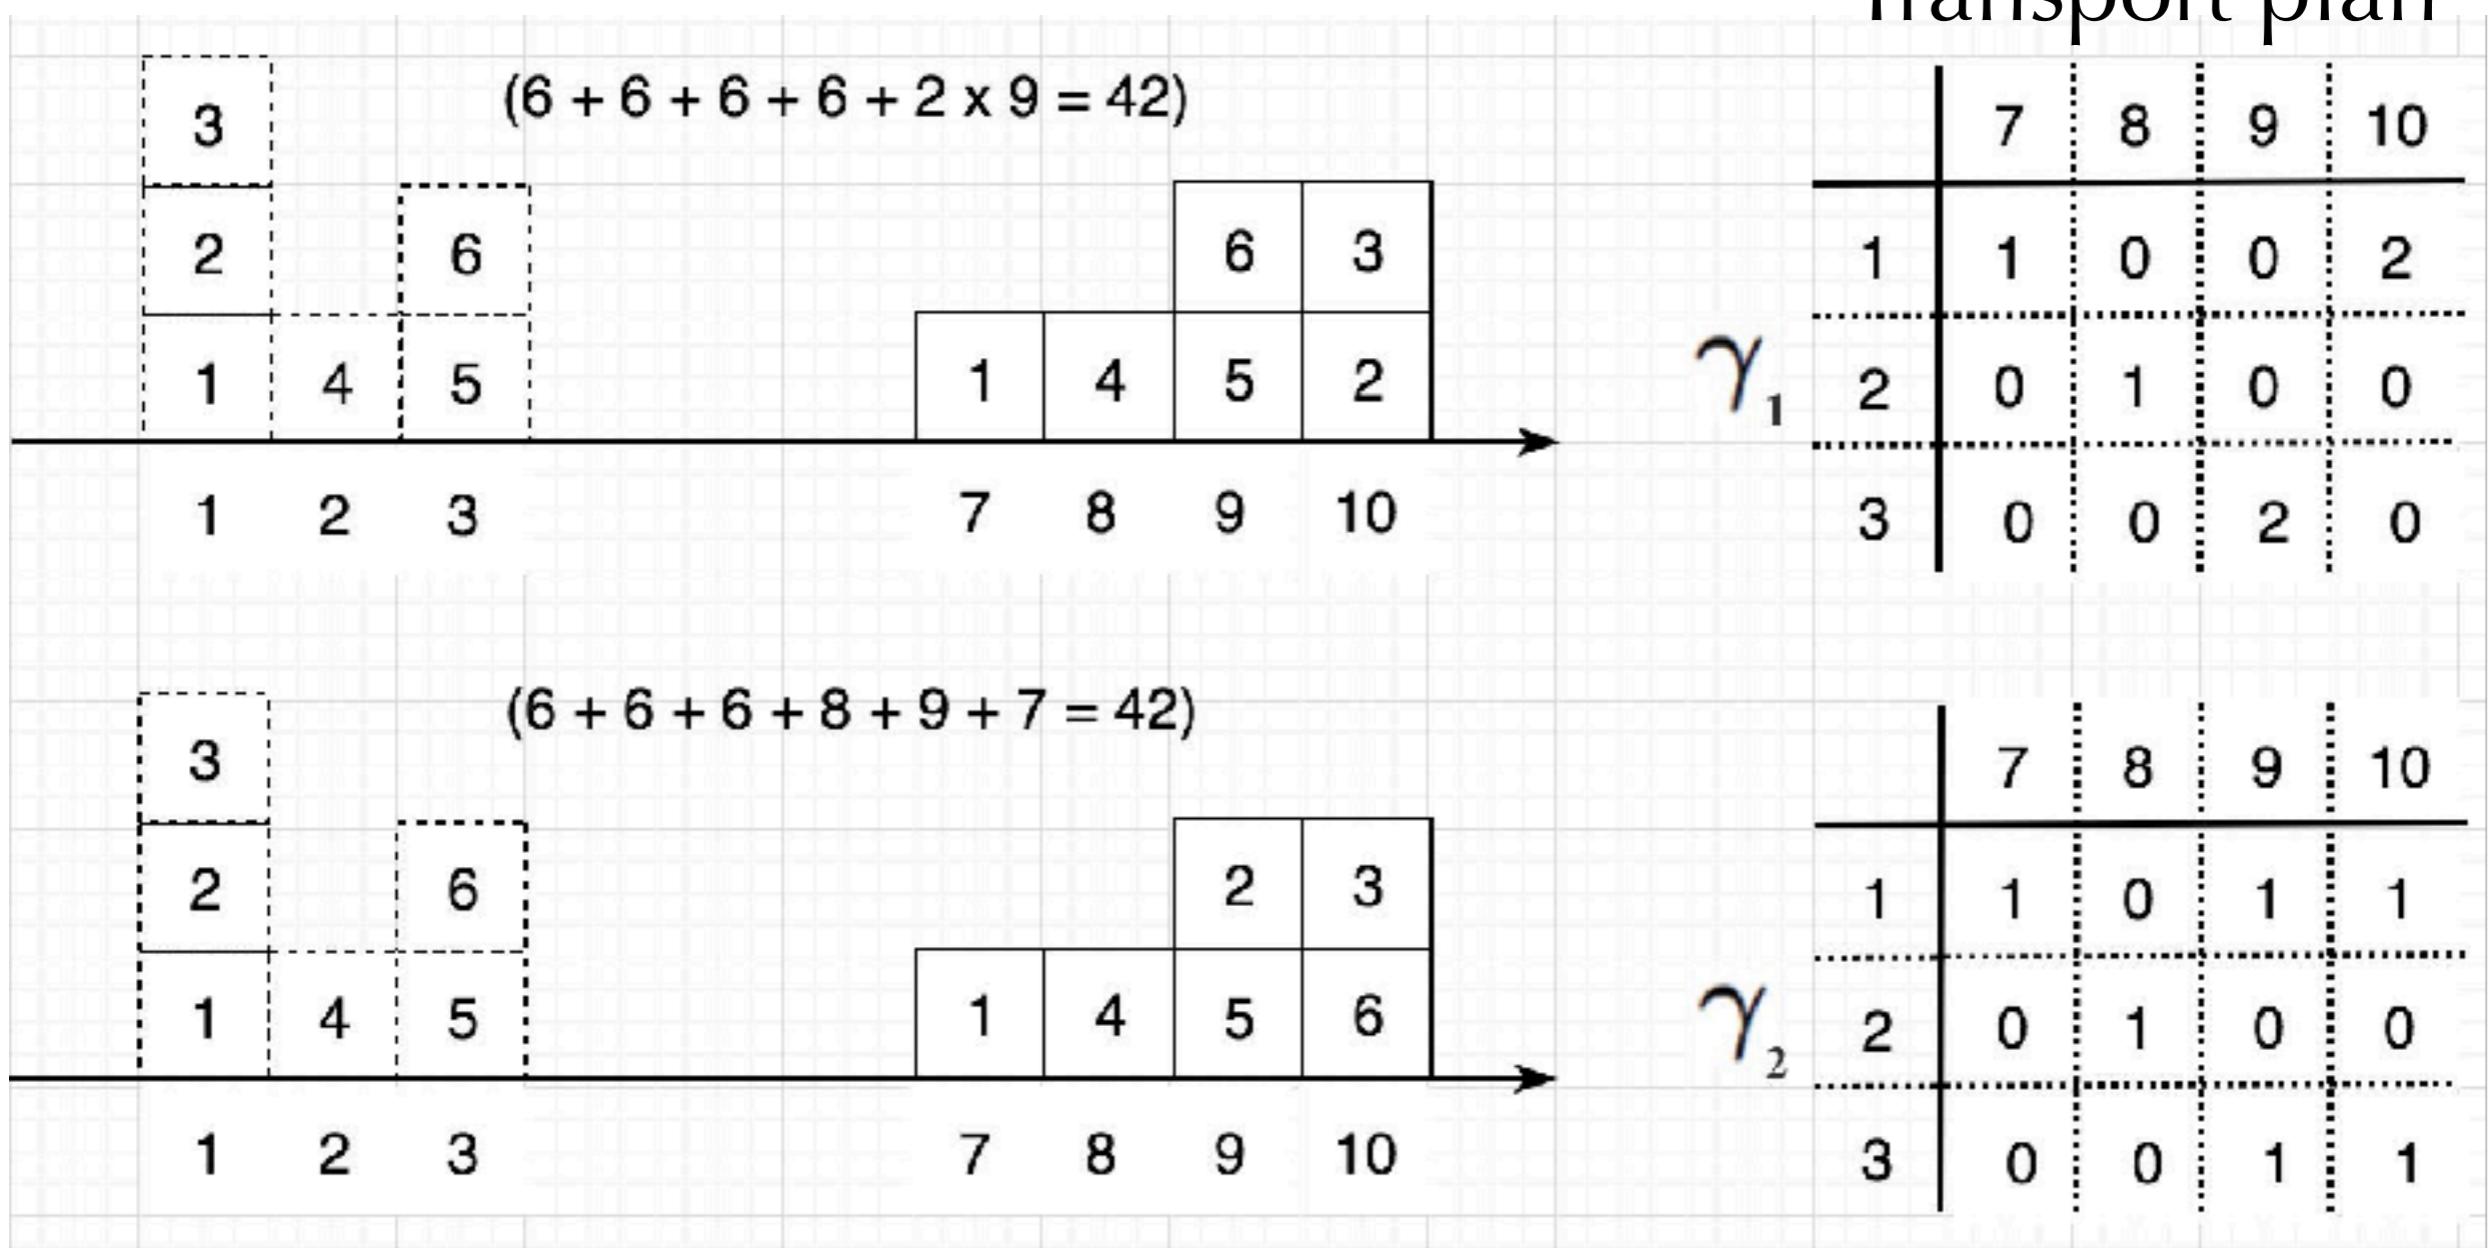
Wasserstein GANs (WGANS)

Earth-Mover distance: Minimum amount of work to move weight from $p(x)$ to $q(x)$



Wasserstein GANs (WGANS)

Earth-Mover distance: Minimum amount of work to move weight from $p(x)$ to $q(x)$



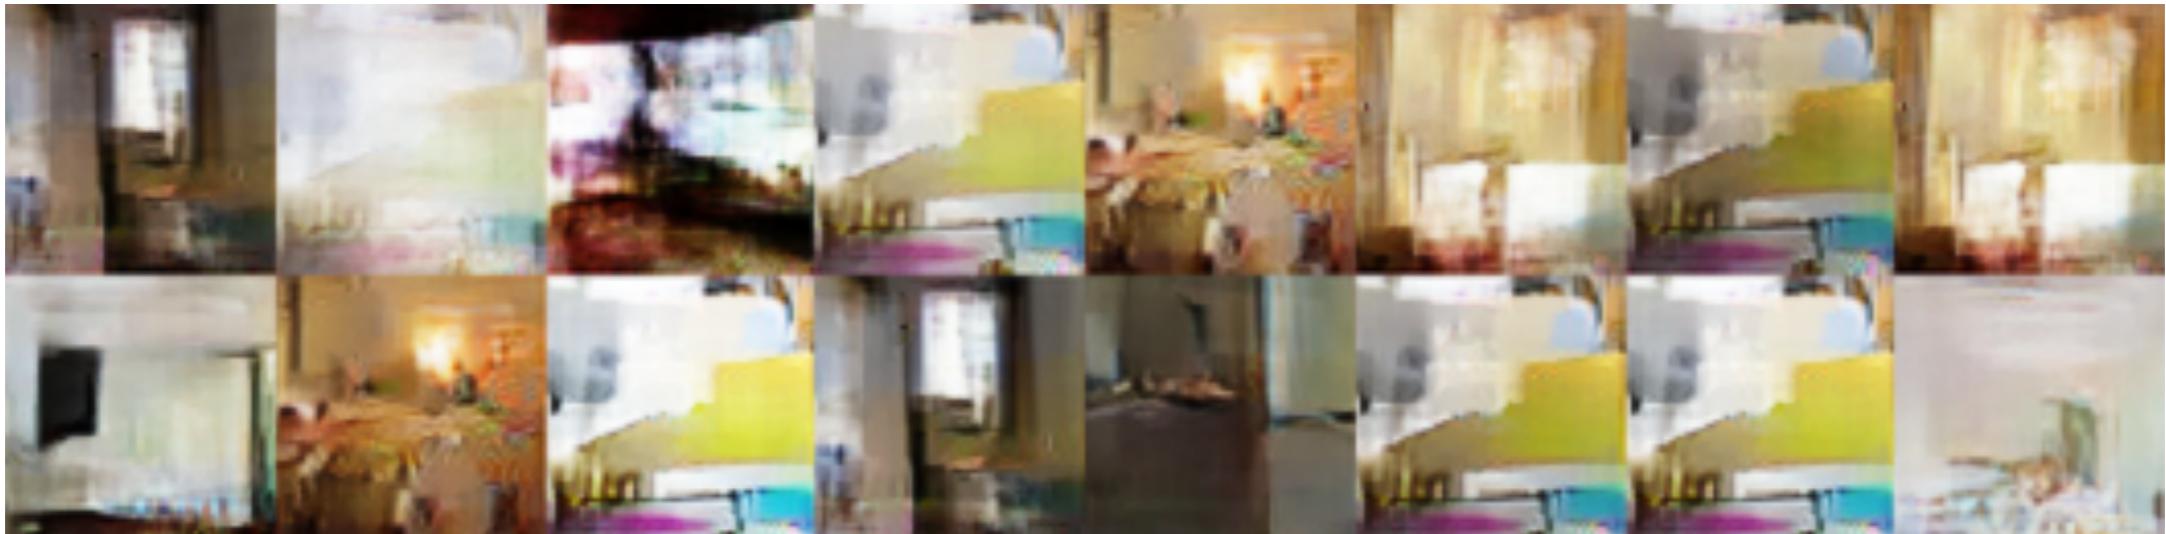
Wasserstein GANs (WGANS)

Earth-Mover distance: Minimum amount of work to move weight from $p(x)$ to $q(x)$

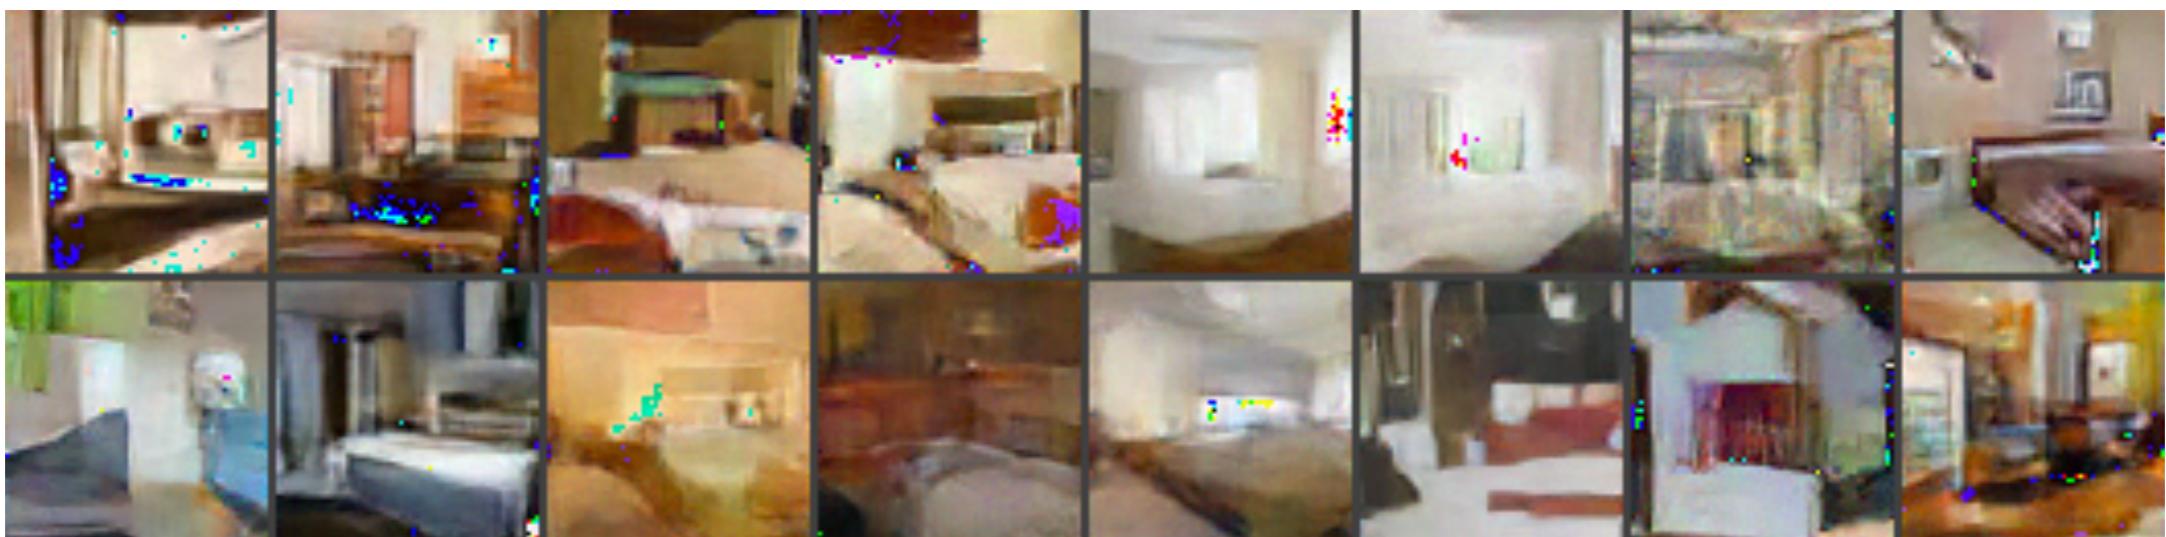
$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|],$$

WGANS vs. GANs

GAN



WGAN



[Arjovsky et al., Wasserstein GAN, arXiv:1701.07875]

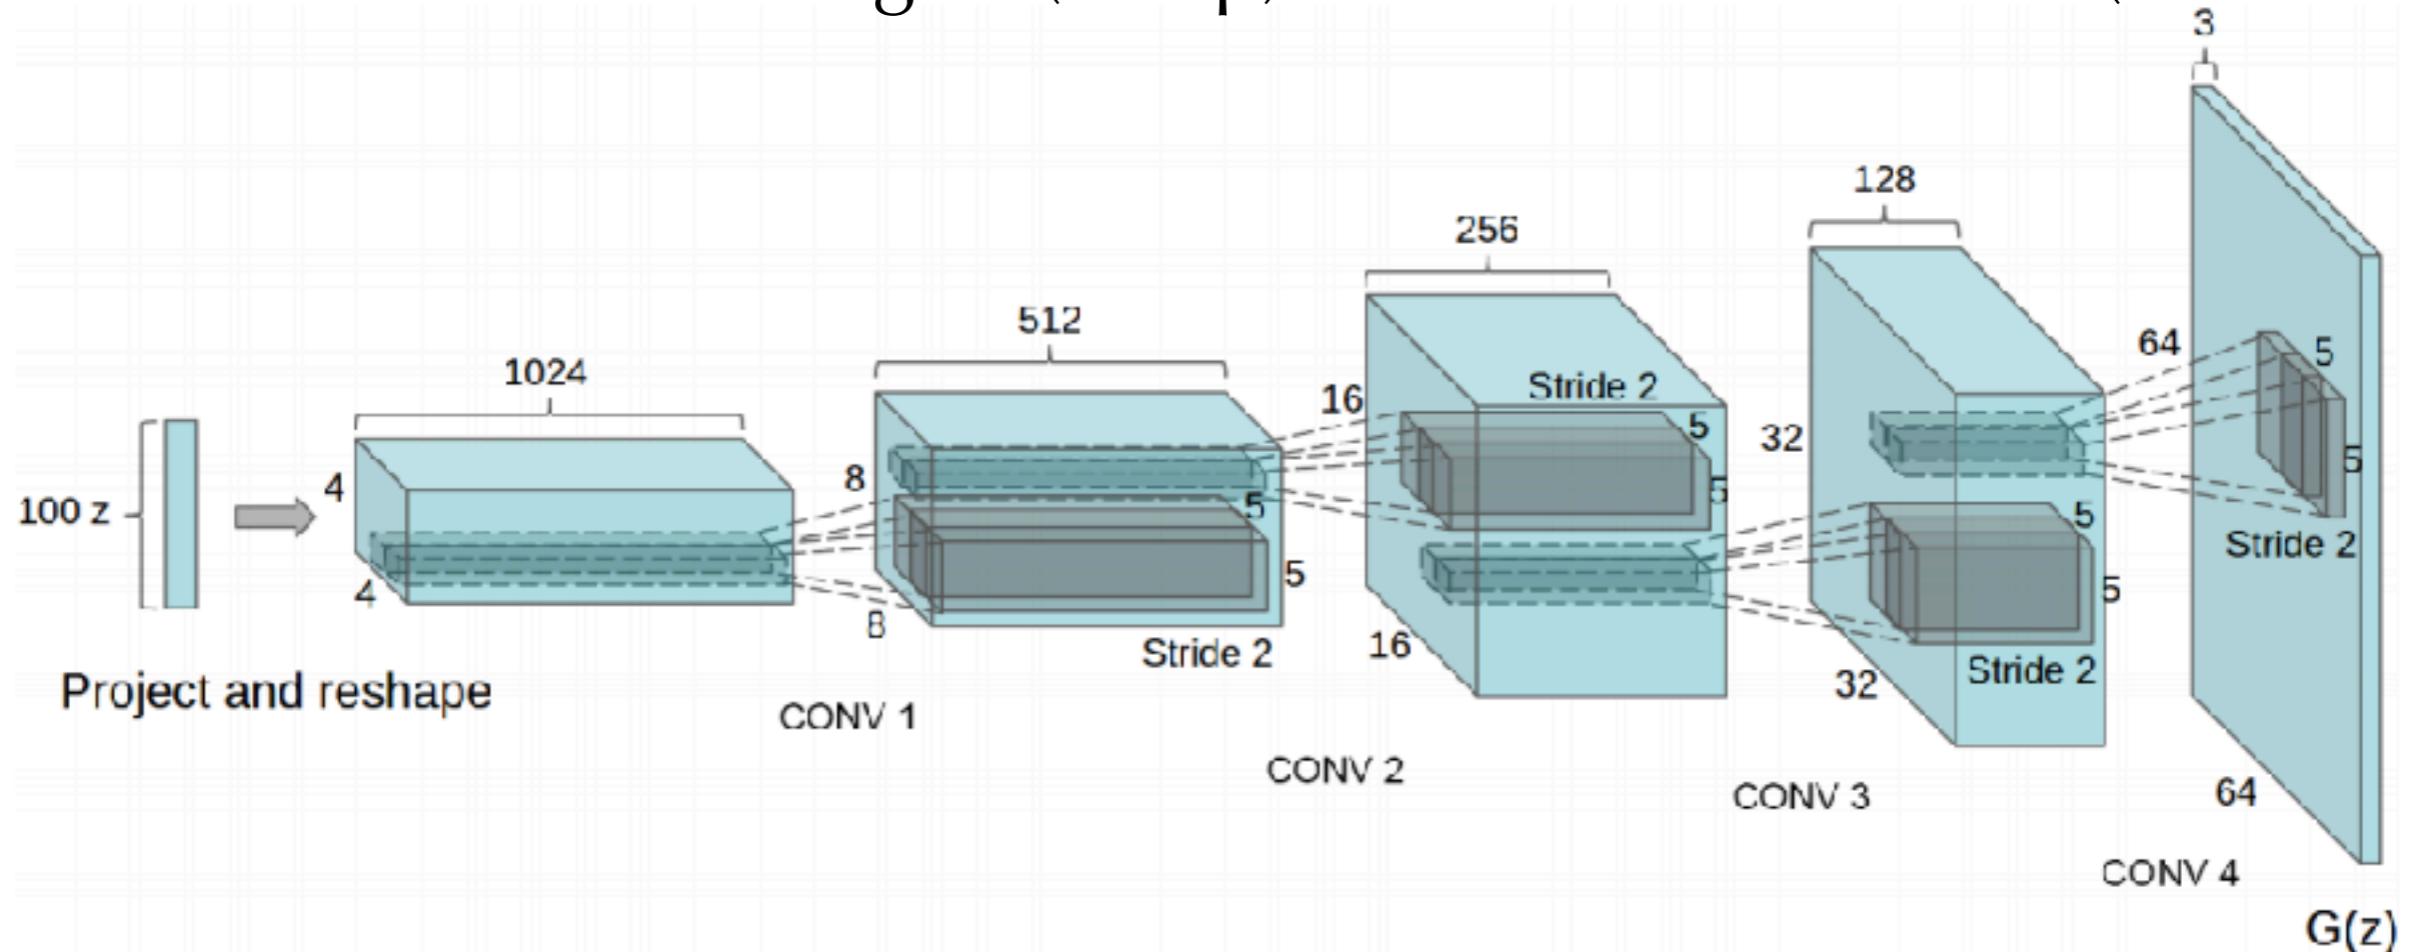
Deep Convolutional GANs (DCGANs)

- GANs are general technique, not restricted to images
- This lecture focus on images: (Deep) Convolutional GANs (DCGANs)

[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

Deep Convolutional GANs (DCGANs)

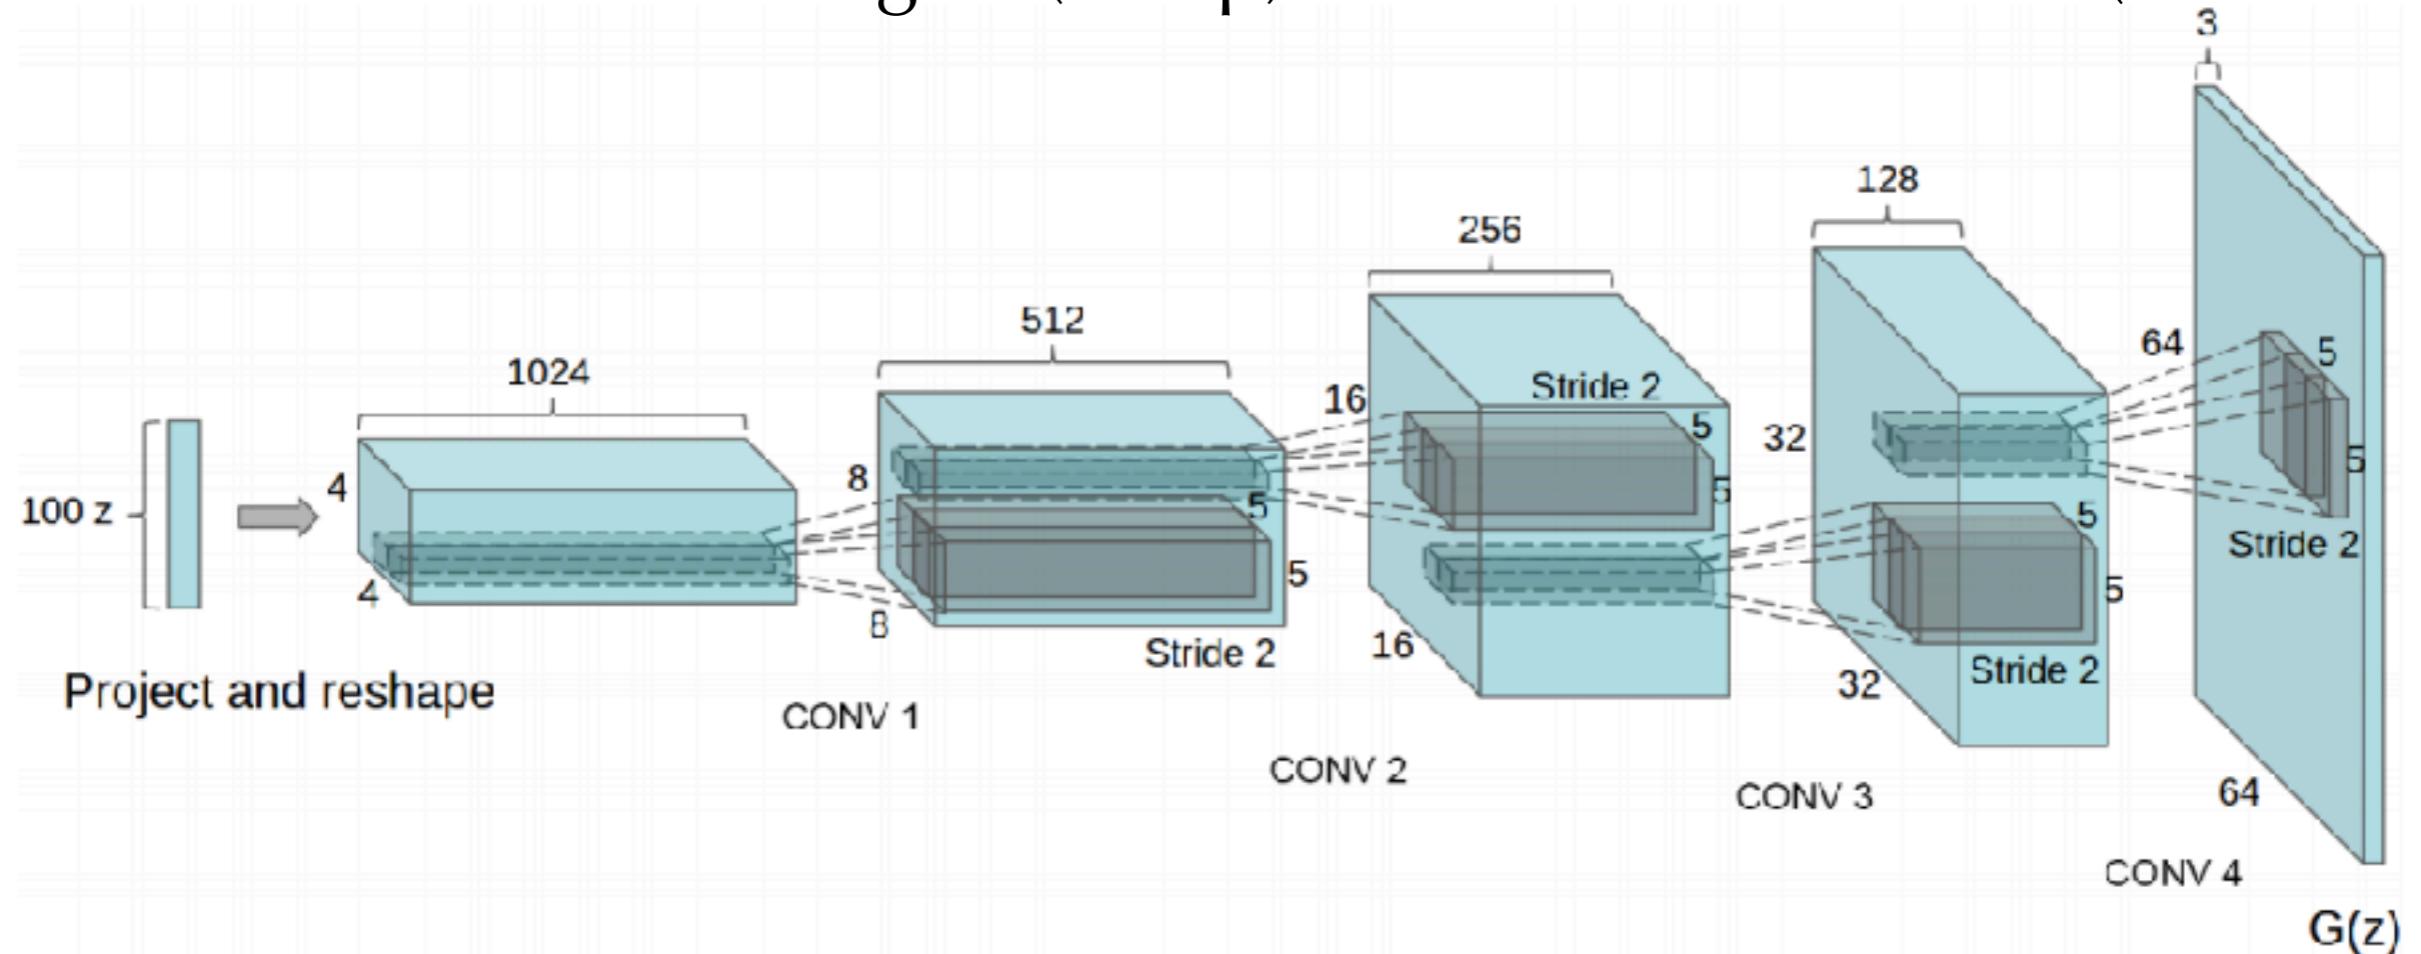
- GANs are general technique, not restricted to images
- This lecture focus on images: (Deep) Convolutional GANs (DCGANs)



[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

Deep Convolutional GANs (DCGANs)

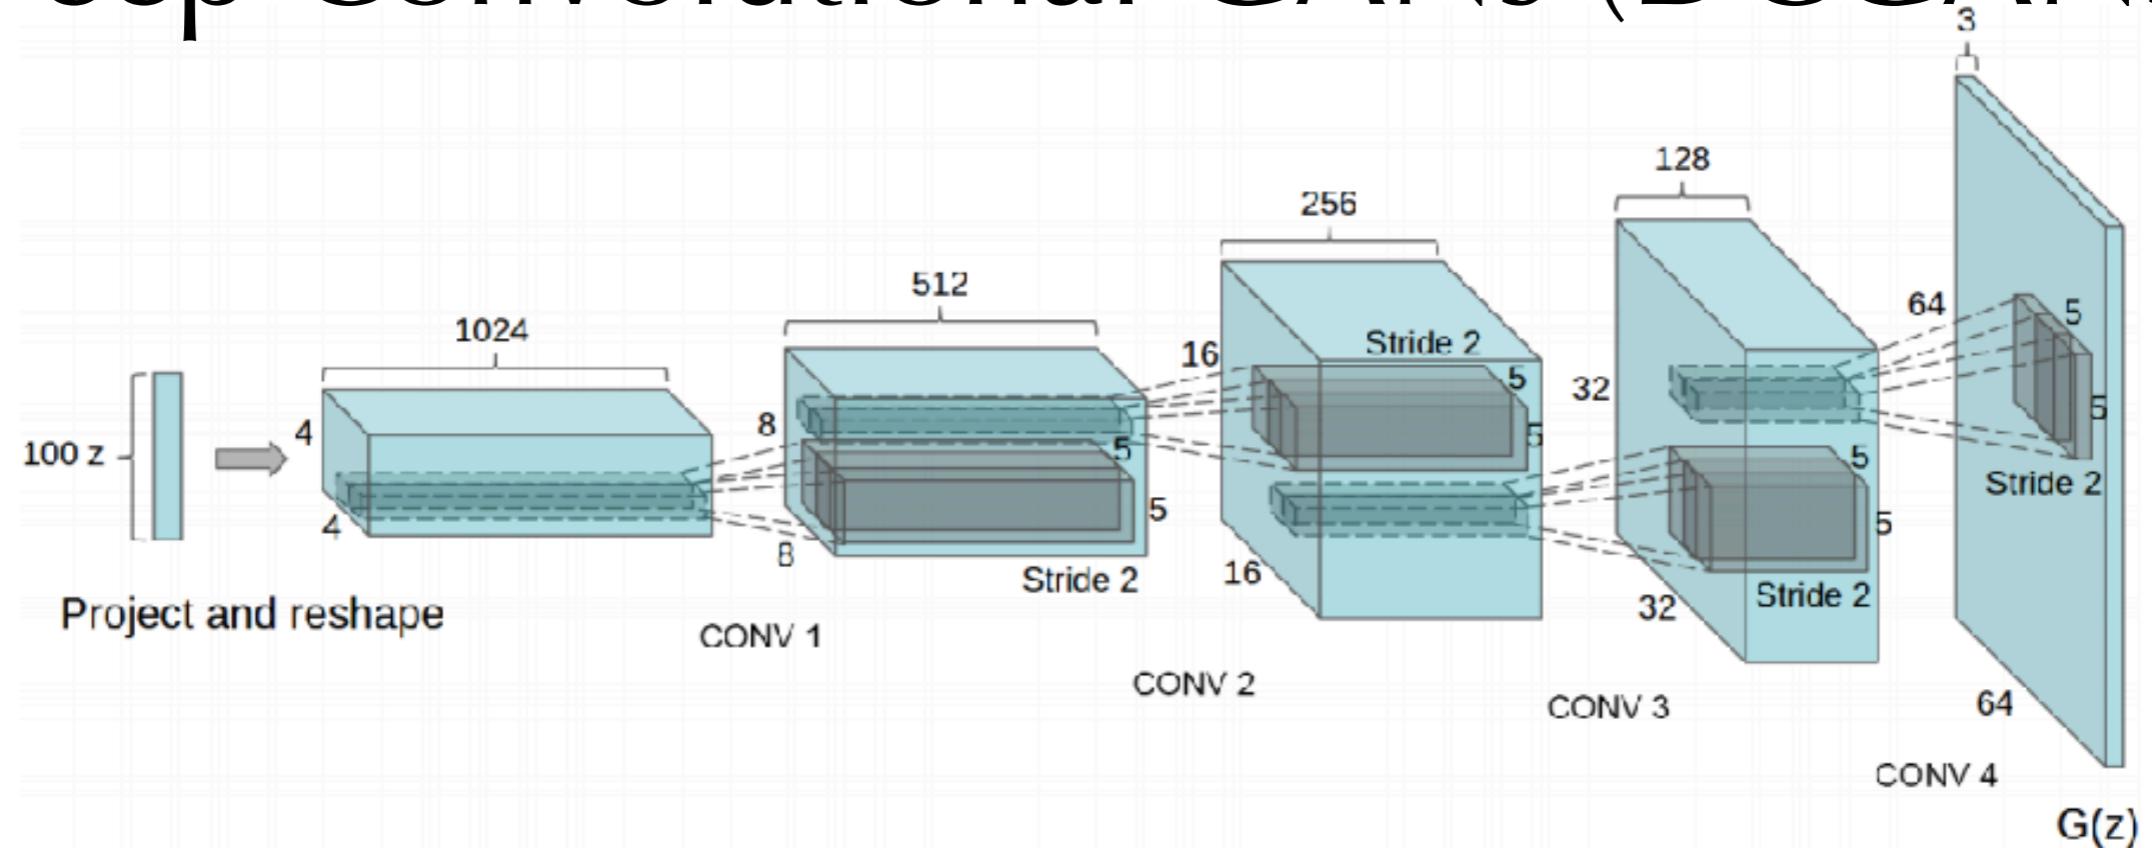
- GANs are general technique, not restricted to images
- This lecture focus on images: (Deep) Convolutional GANs (DCGANs)



- List of recommendations for stable DCGANs

[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

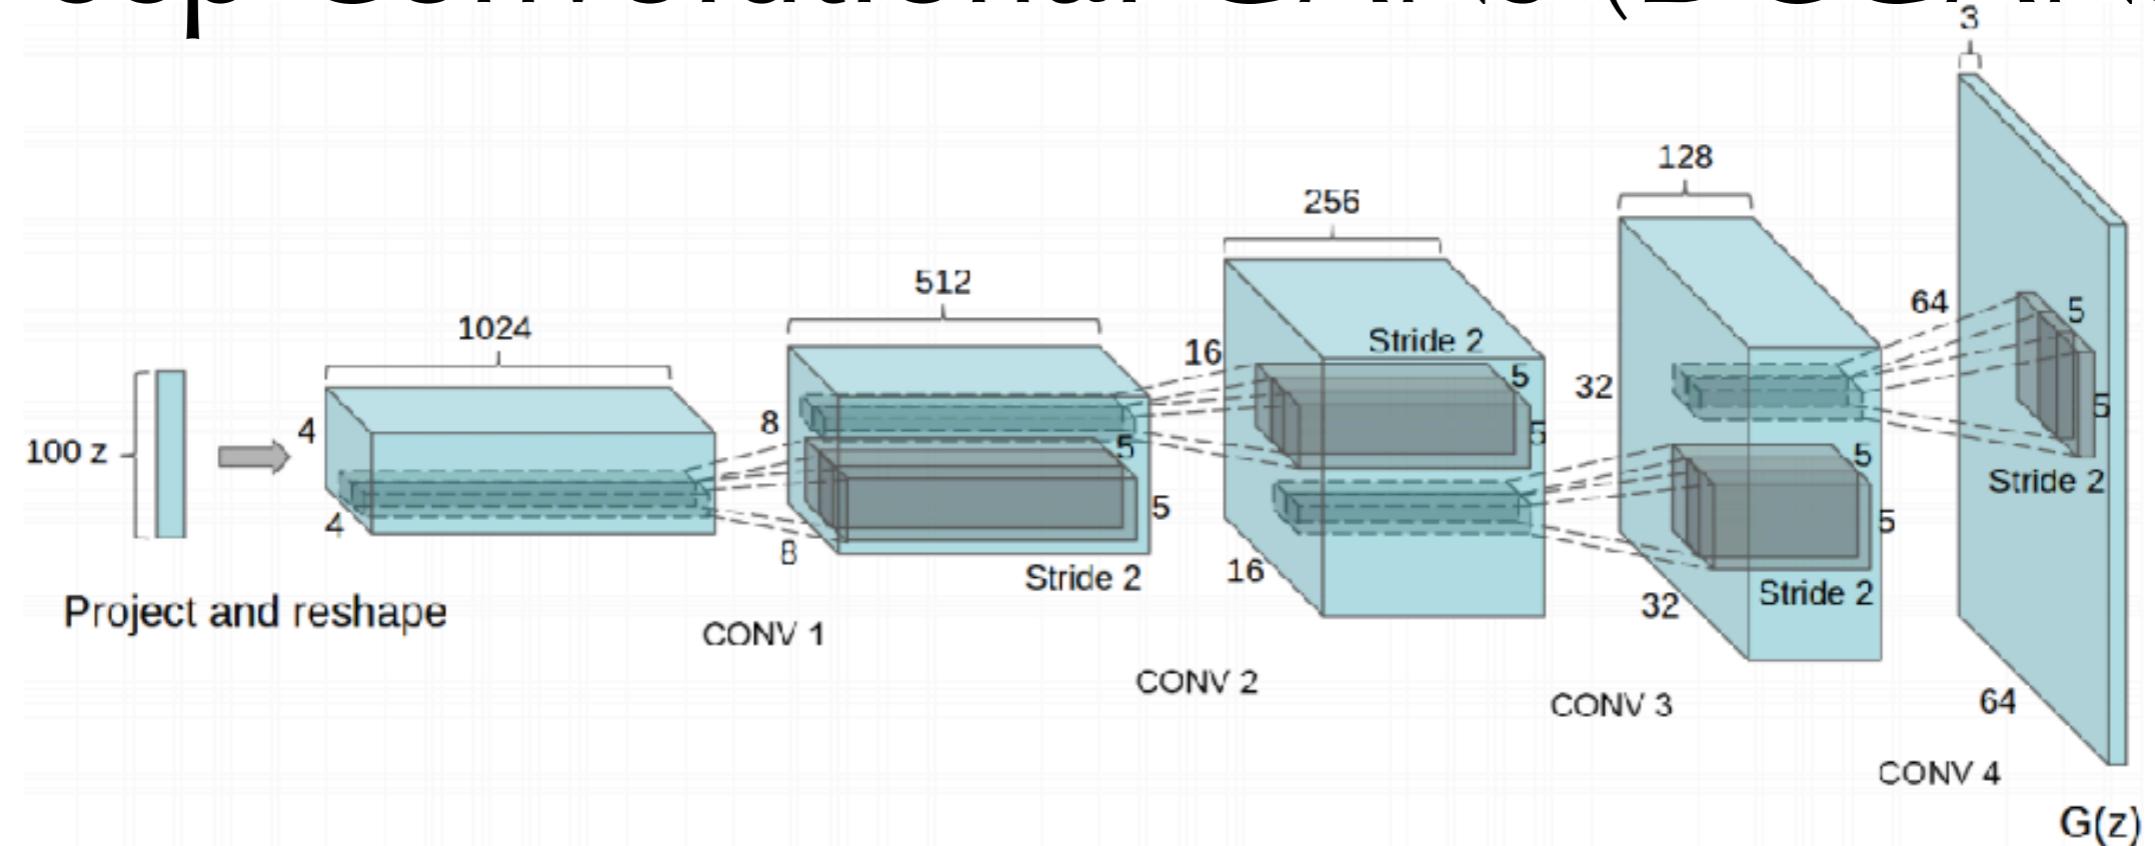
Deep Convolutional GANs (DCGANs)



- Discriminator: Use strided convolutions instead of max-pooling
- Generator: Use transposed convolutions
- Intuition: Learn up- / down-sampling

[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

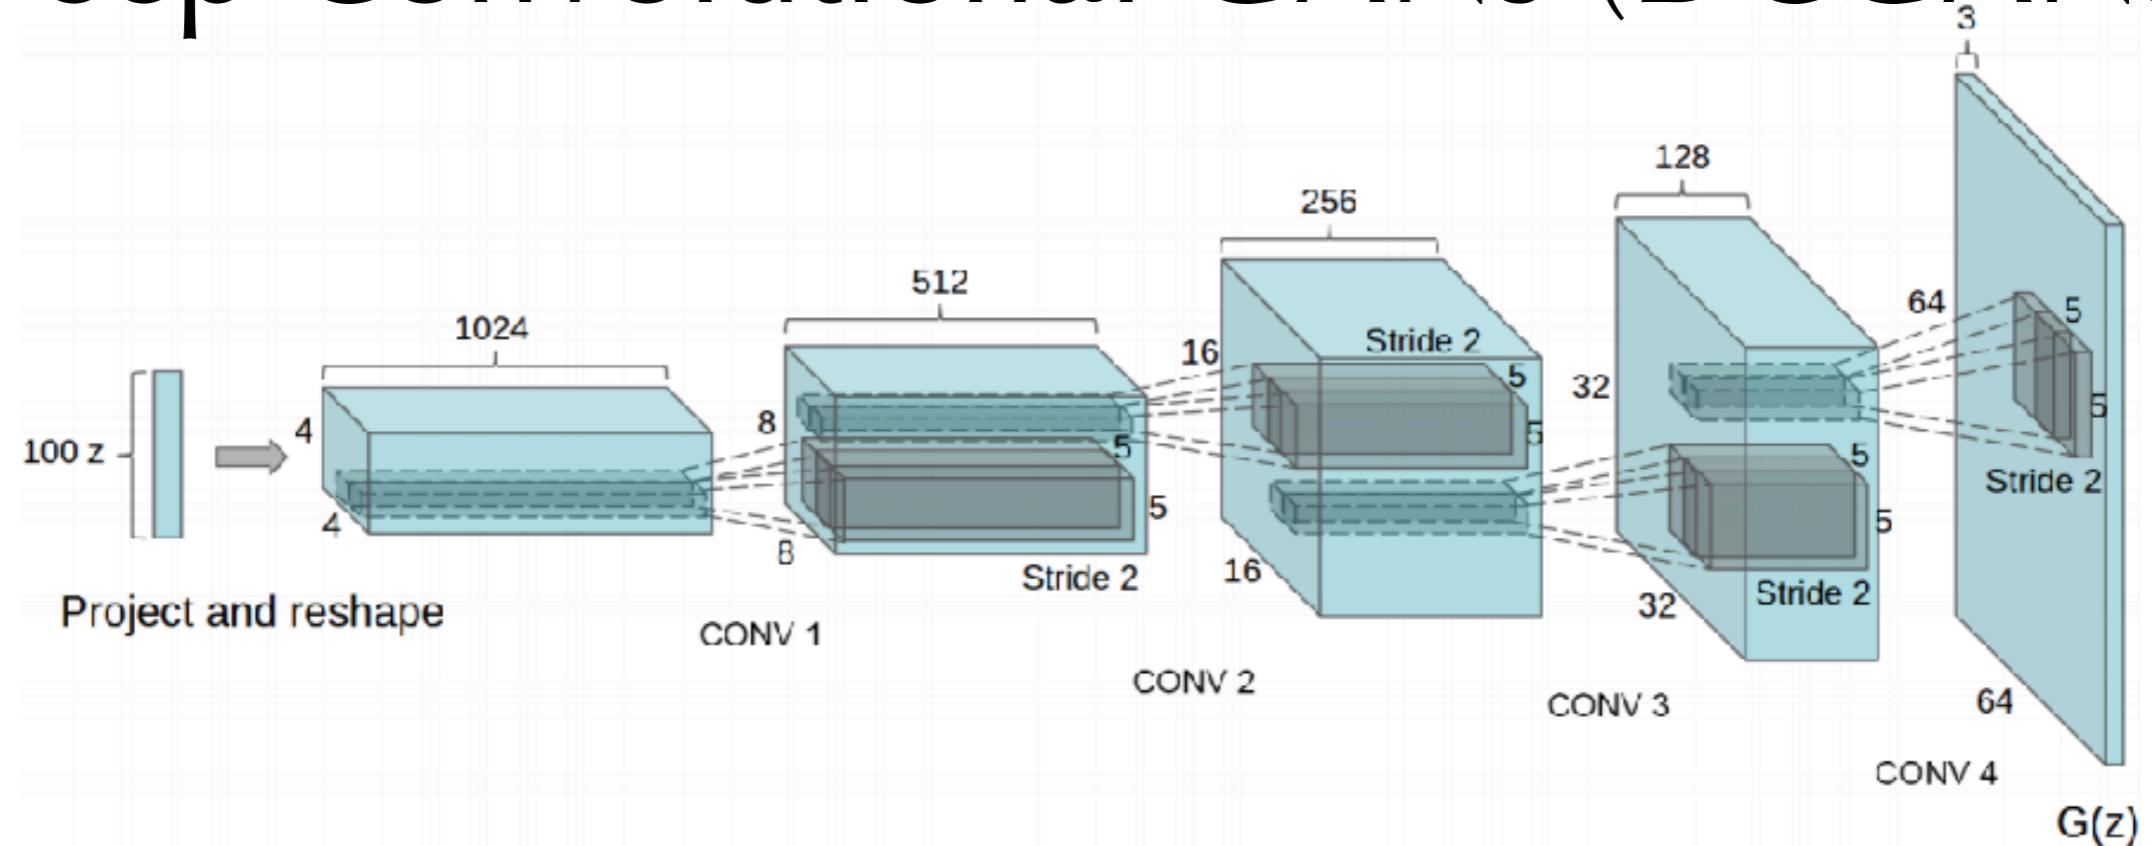
Deep Convolutional GANs (DCGANs)



- No fully connected hidden layers

[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

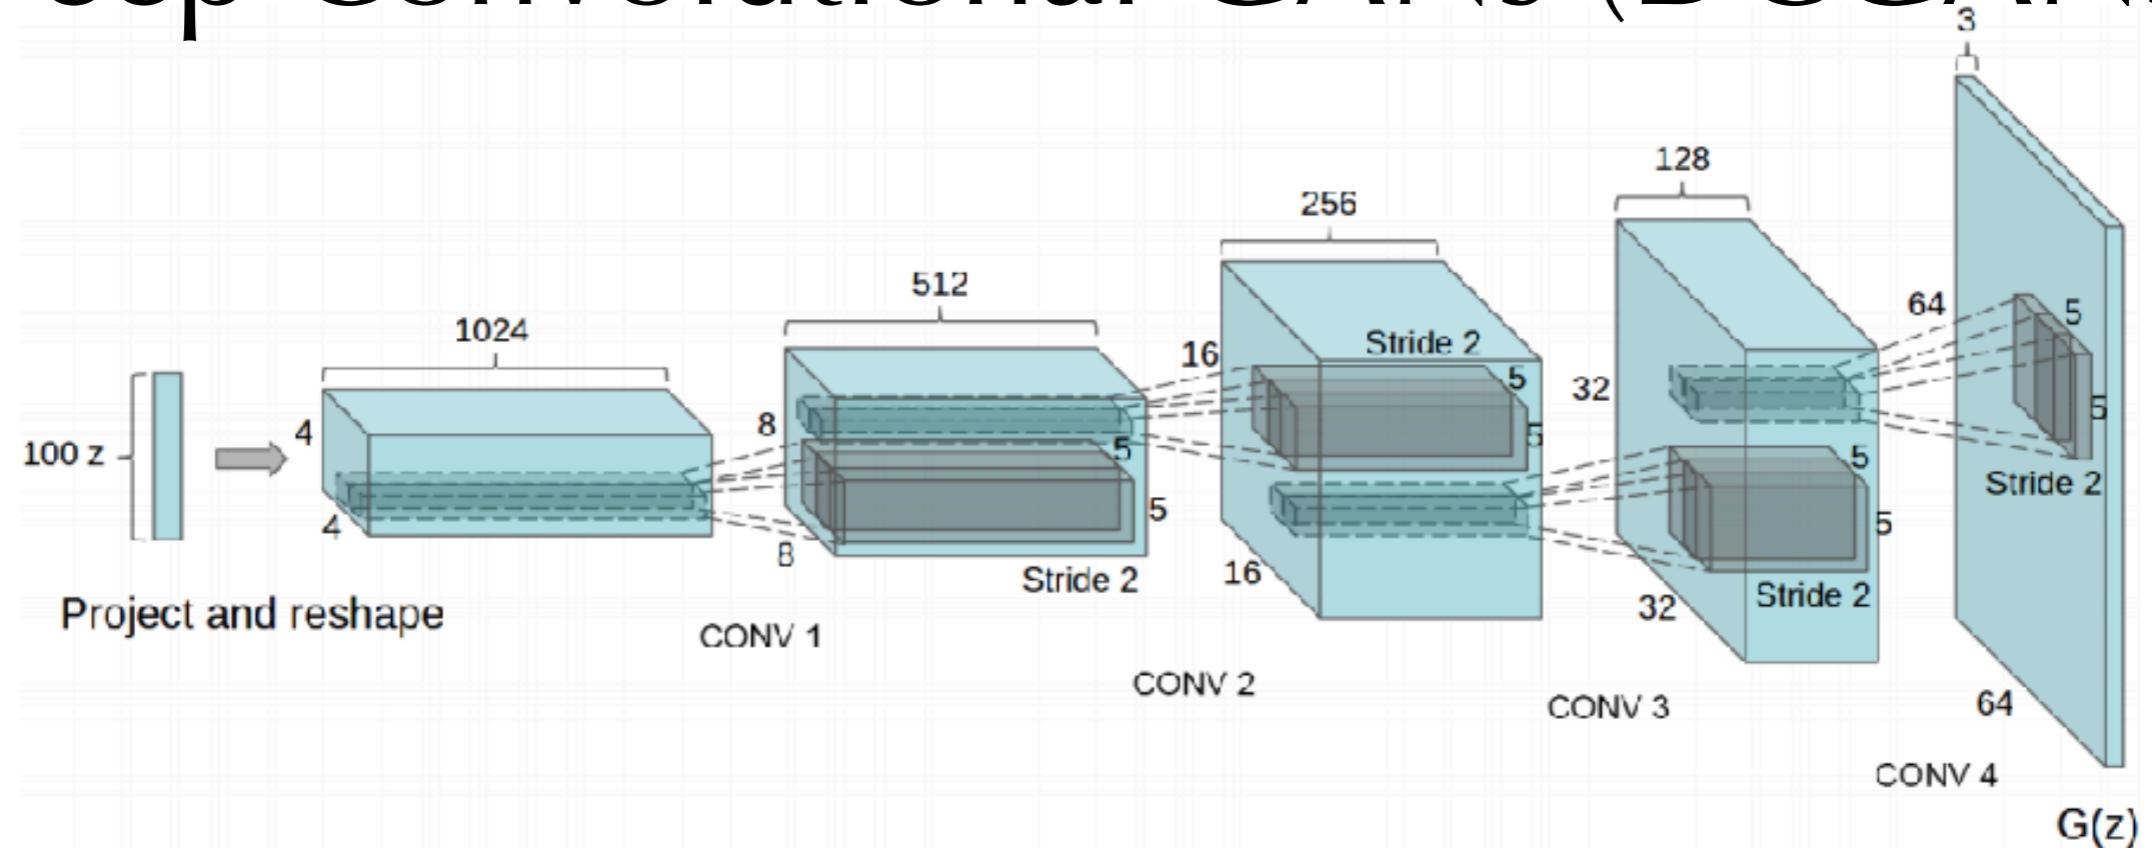
Deep Convolutional GANs (DCGANs)



- No fully connected hidden layers
- Observation: Improves convergence speed = training speed

[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

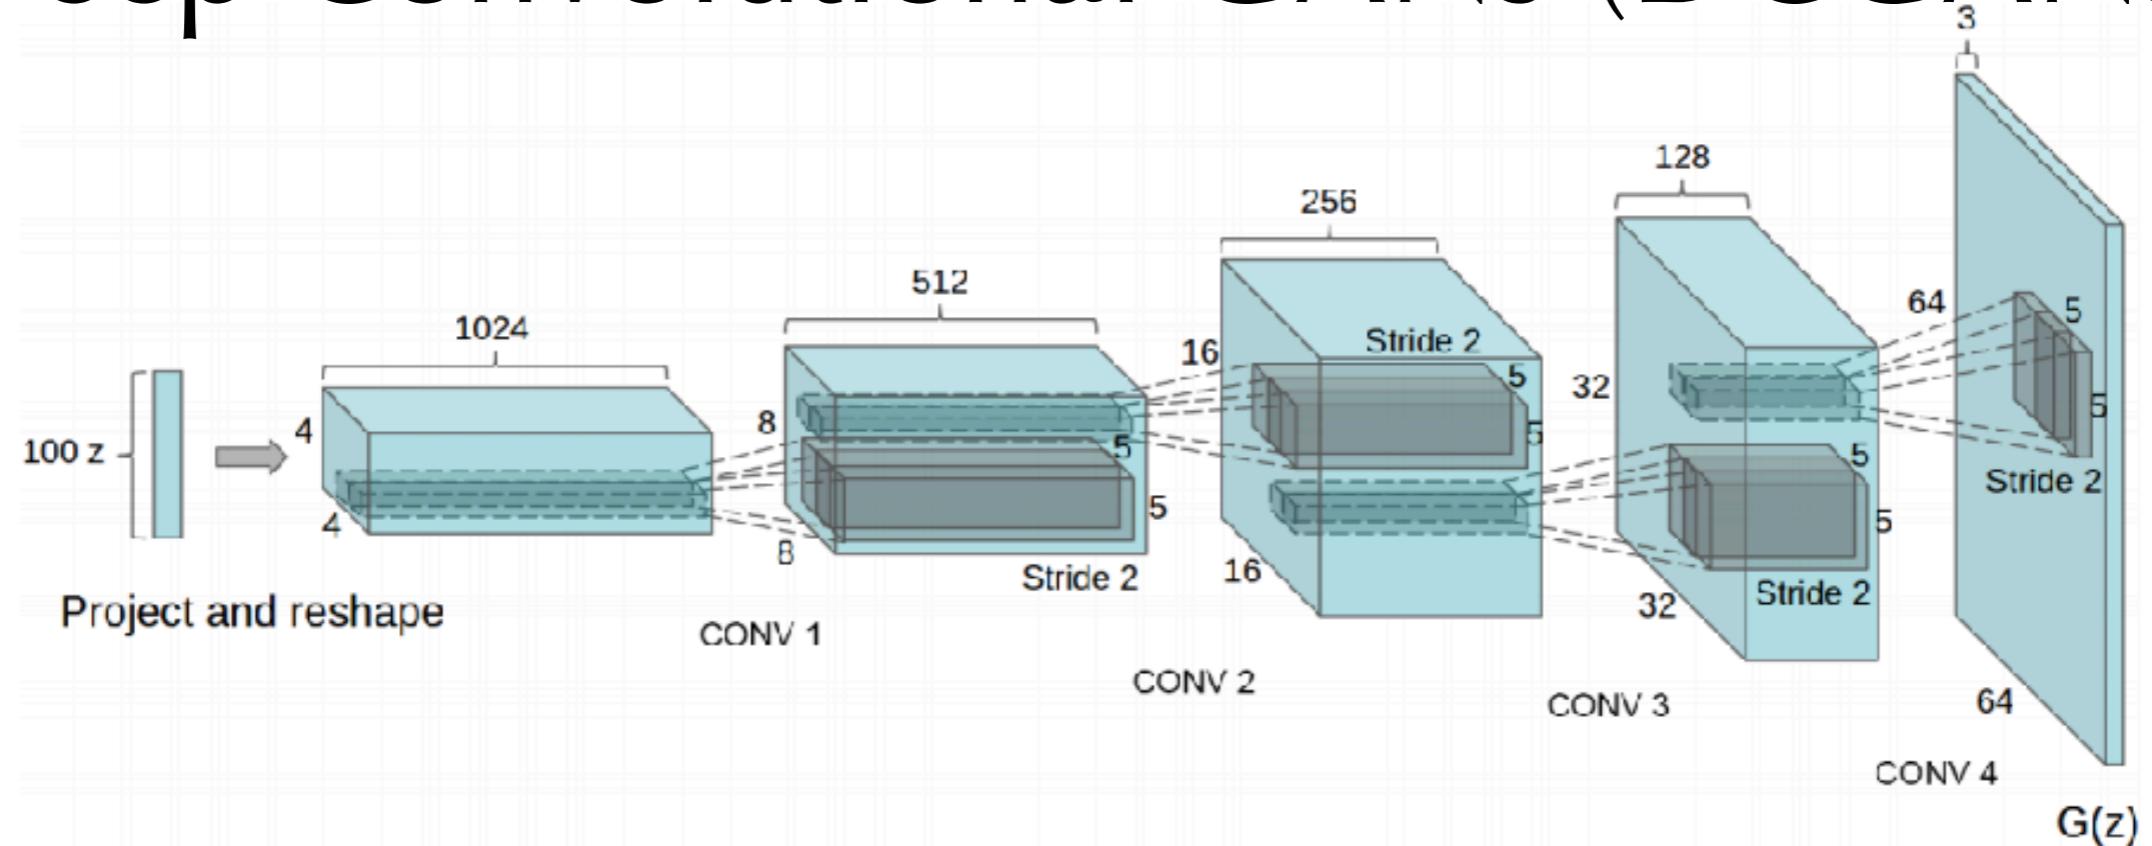
Deep Convolutional GANs (DCGANs)



- Use **batch normalization** (batch norm) in both generator and discriminator (but not in generator output layer and discriminator input layer)

[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

Deep Convolutional GANs (DCGANs)



- Use **batch normalization** (batch norm) in both generator and discriminator (but not in generator output layer and discriminator input layer)
- Observation: Helps generator to start training (prevents generator from collapsing all samples to single point)

[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

Batch Normalization

- Motivation:
 - Distribution of input of network layers changes over time
 - Distribution potentially changes with each mini batch
 - Later layers are affected by distributions of previous layers
 - Makes training slower and harder

Batch Normalization

- Motivation:
 - Distribution of input of network layers changes over time
 - Distribution potentially changes with each mini batch
 - Later layers are affected by distributions of previous layers
 - Makes training slower and harder
- Batch normalization:
 - Normalize the distribution to stabilize training
 - Normalization per mini-batch for efficiency

Batch Normalization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots m\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

Batch Normalization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots m\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

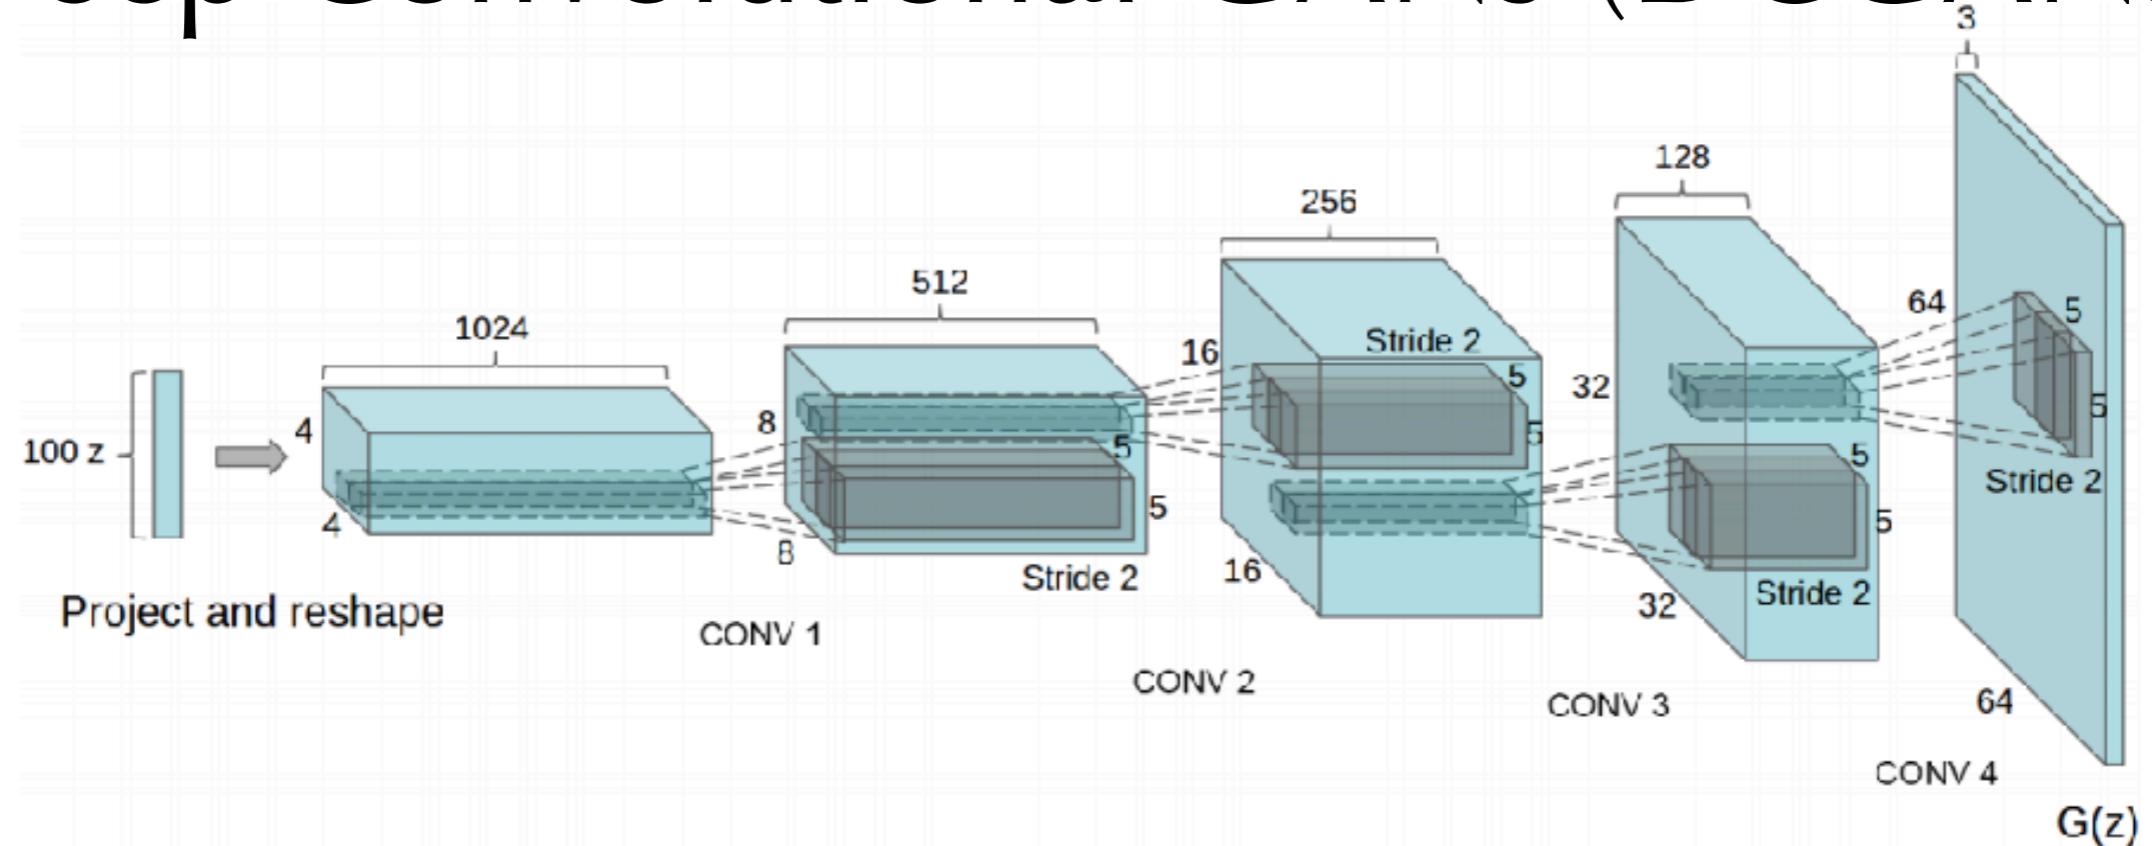
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

Observation: Stabilizes and accelerates the training process

Deep Convolutional GANs (DCGANs)

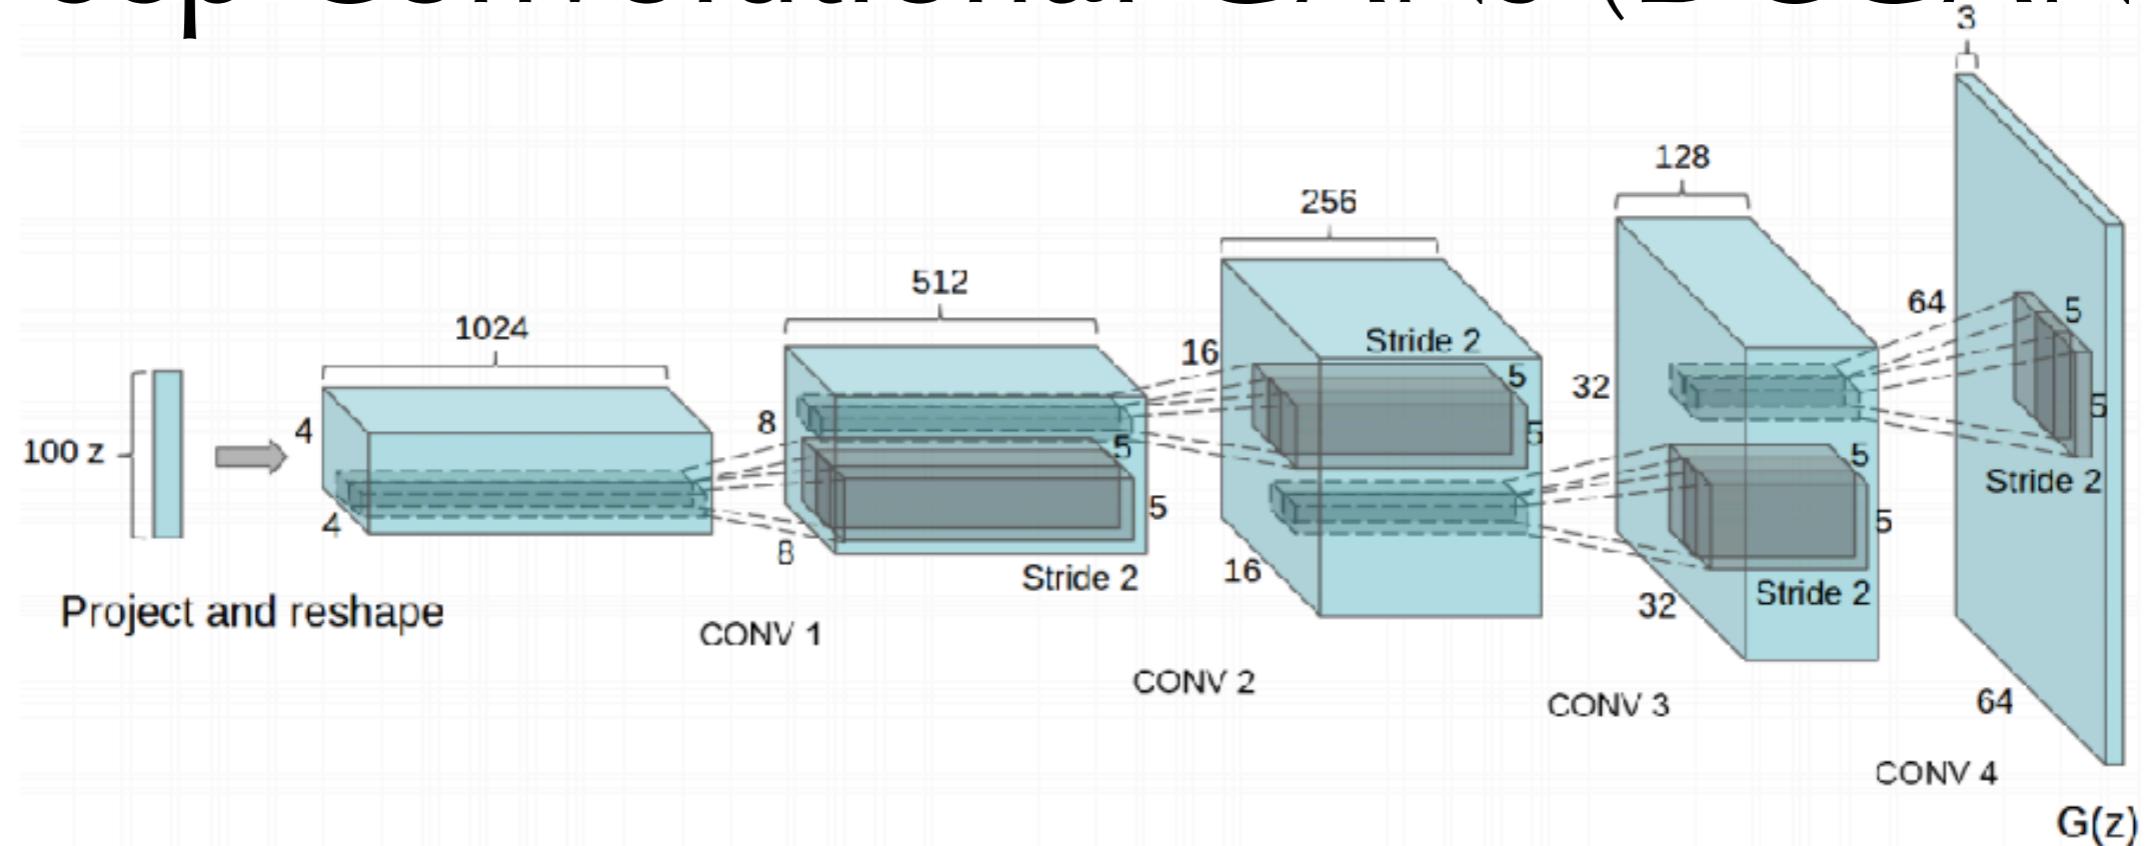


- Generator: Use ReLU in all layers, but tanh in output layer

[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

[Maas et al., Rectifier Nonlinearities Improve Neural Network Acoustic Models, ICML 2013]

Deep Convolutional GANs (DCGANs)



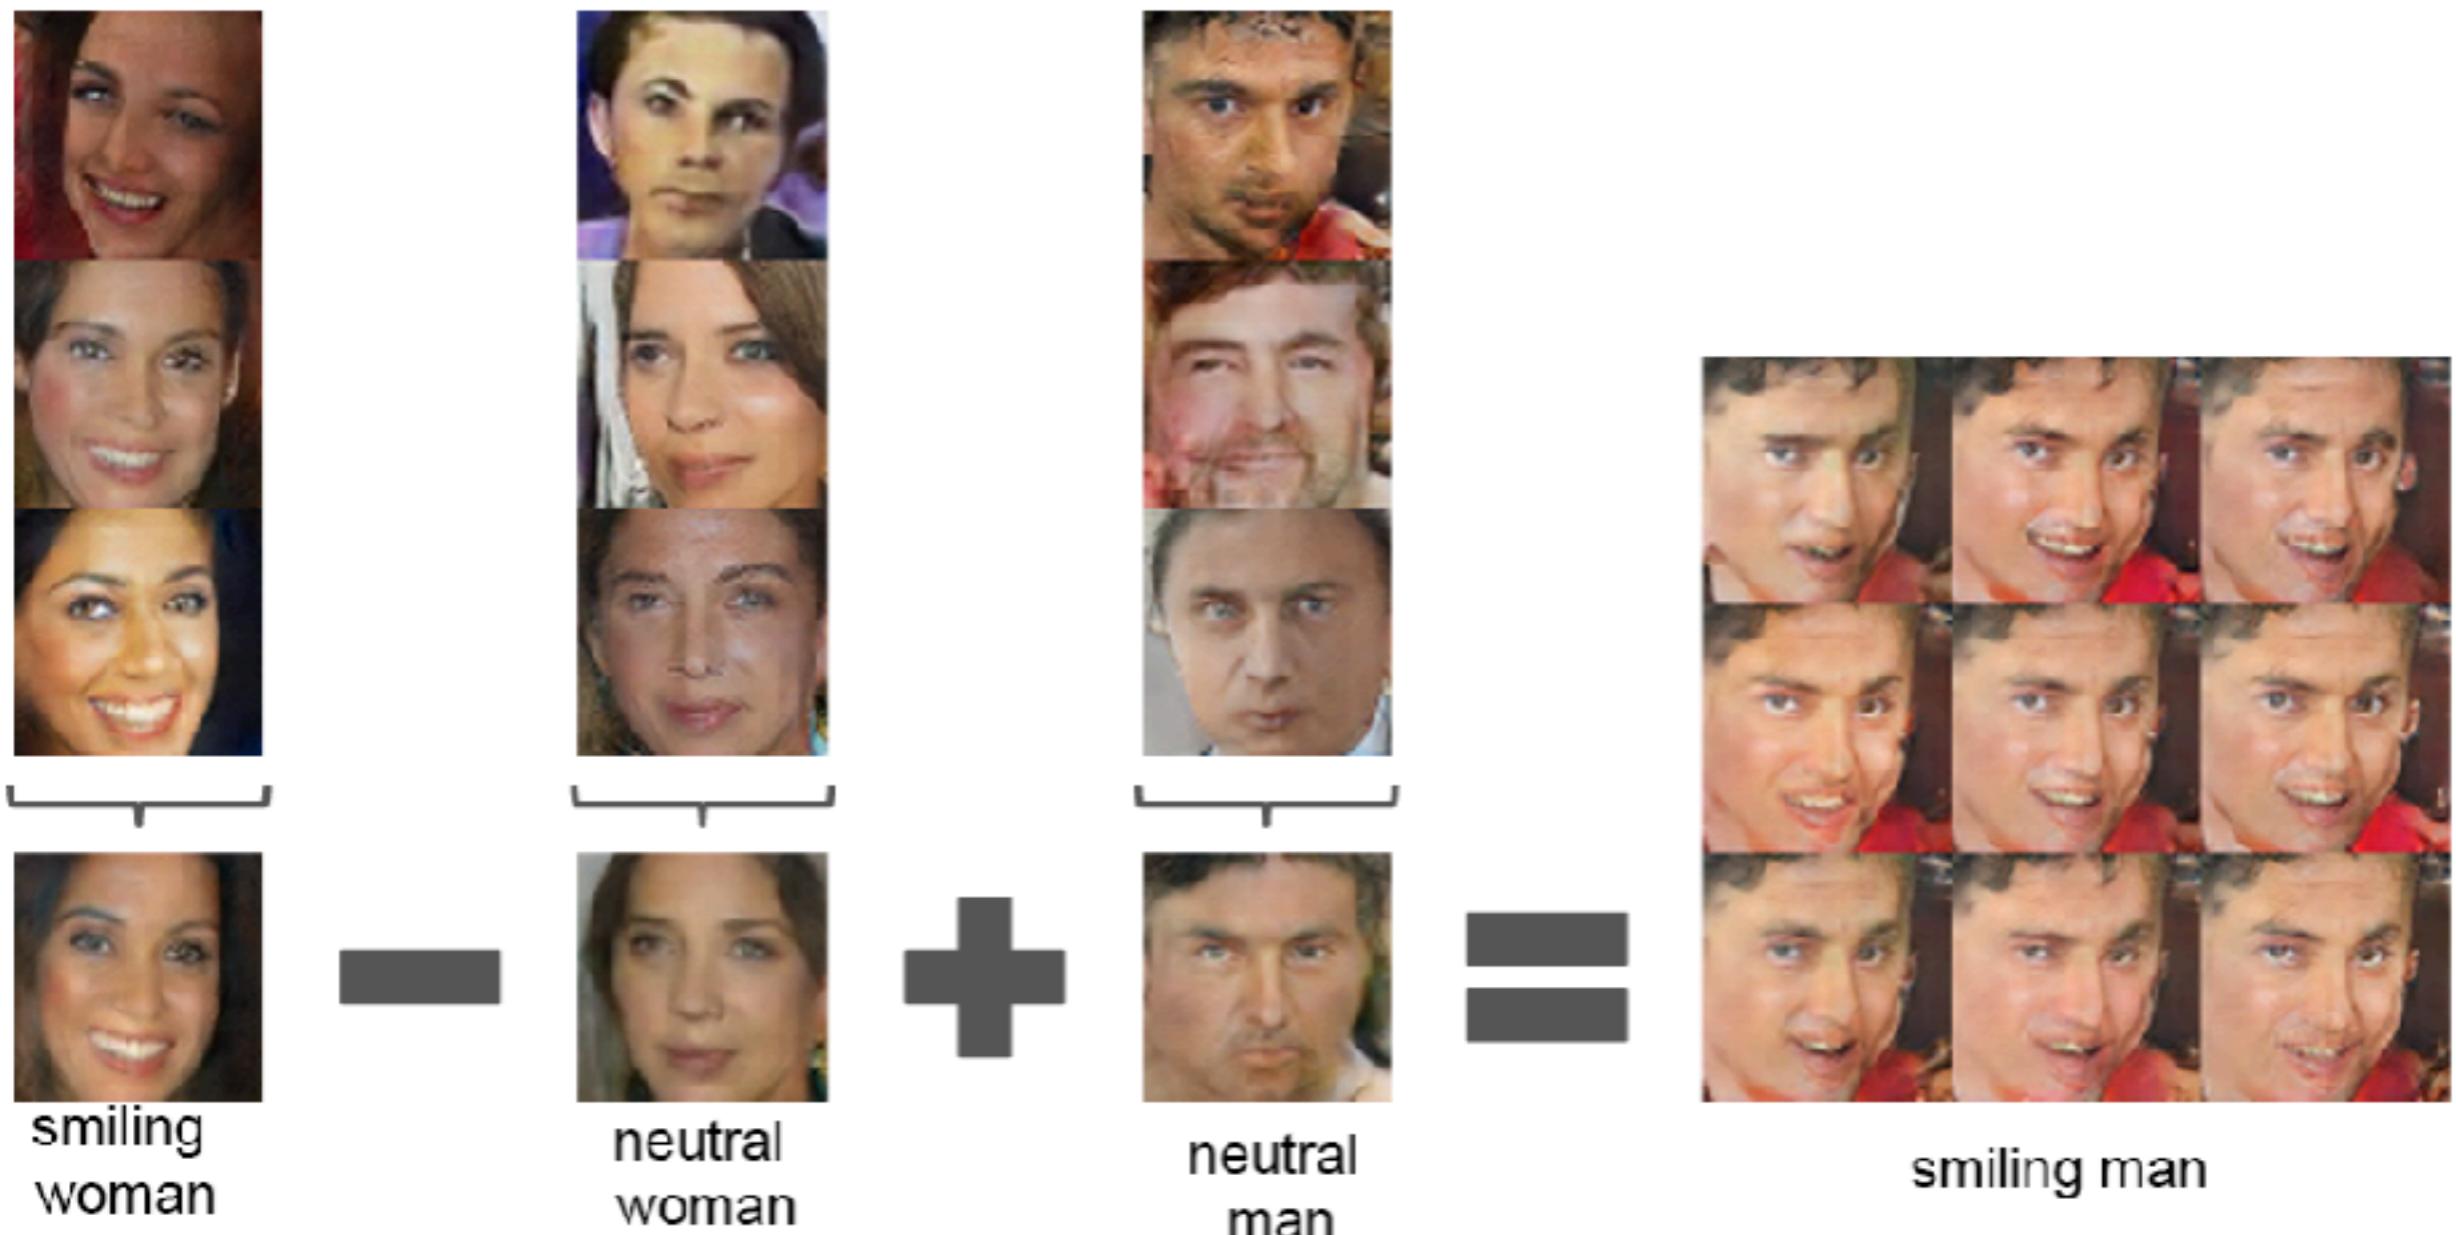
- Generator: Use ReLU in all layers, but tanh in output layer
- Discriminator: Use LeakyReLU

$$\text{LeakyReLU}(x) = \begin{cases} x & \text{if } x > 0 \\ 0.01x & \text{otherwise} \end{cases}$$

[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

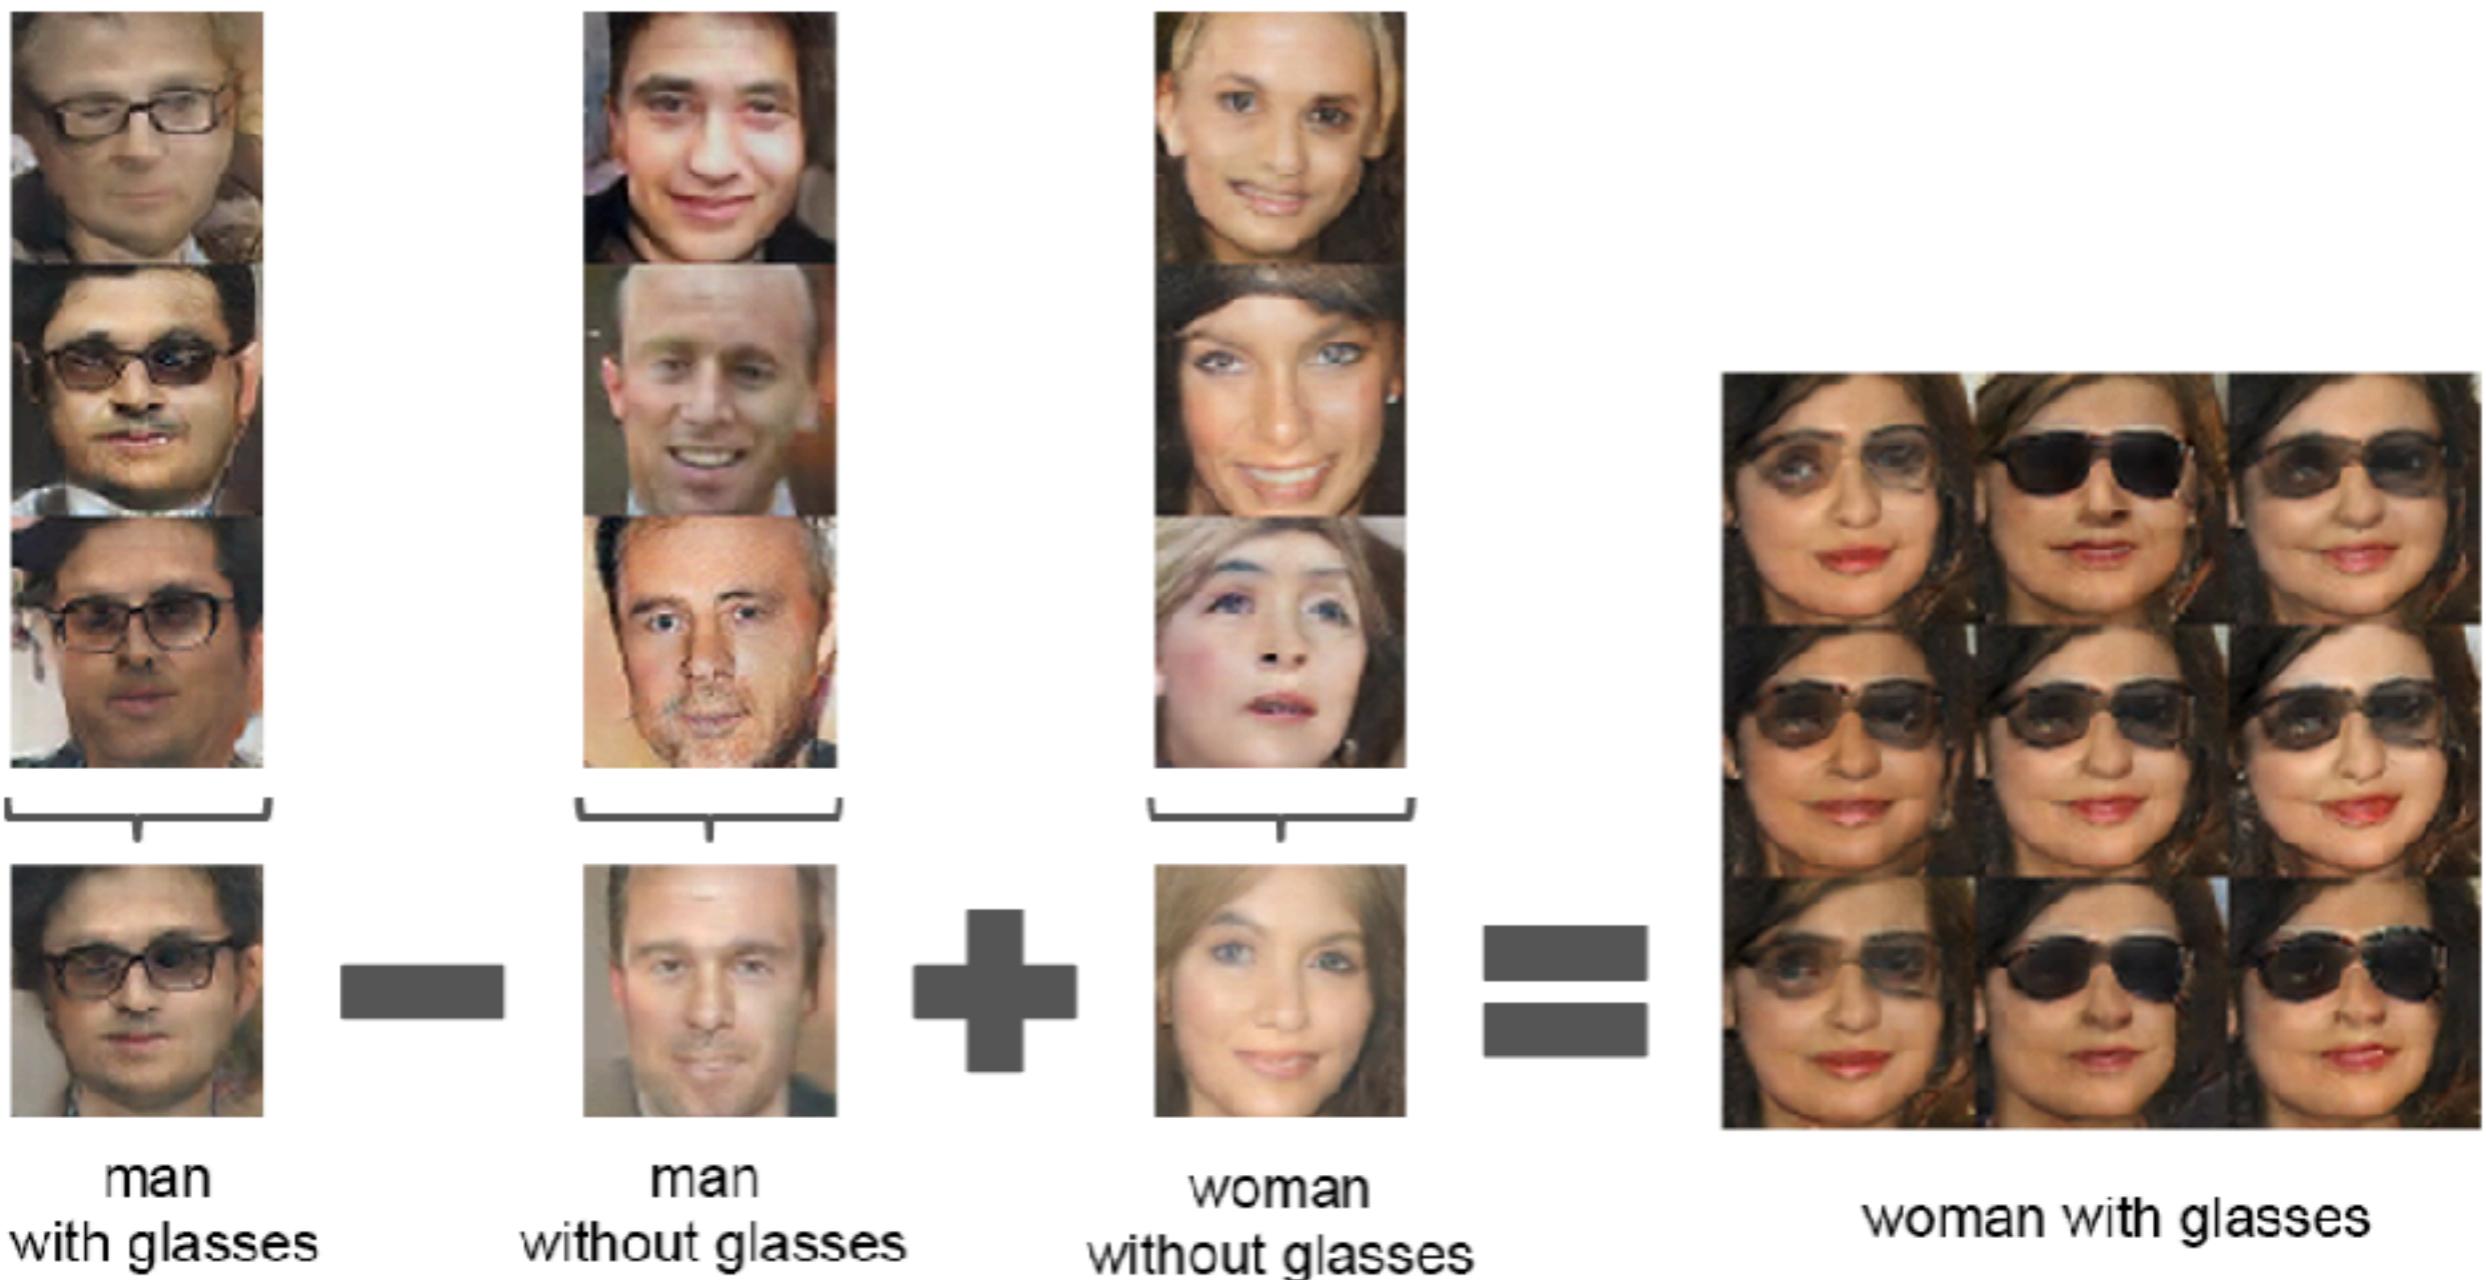
[Maas et al., Rectifier Nonlinearities Improve Neural Network Acoustic Models, ICML 2013]

Deep Convolutional GANs (DCGANs)



[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

Deep Convolutional GANs (DCGANs)



[Radford et al., Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks, ICLR 2016]

State-of-the-art Results

PROGRESSIVE GROWING OF GANs FOR IMPROVED
QUALITY, STABILITY, AND VARIATION

Tero Karras
NVIDIA

Timo Aila
NVIDIA

Samuli Laine
NVIDIA

Jaakko Lehtinen
NVIDIA
Aalto University

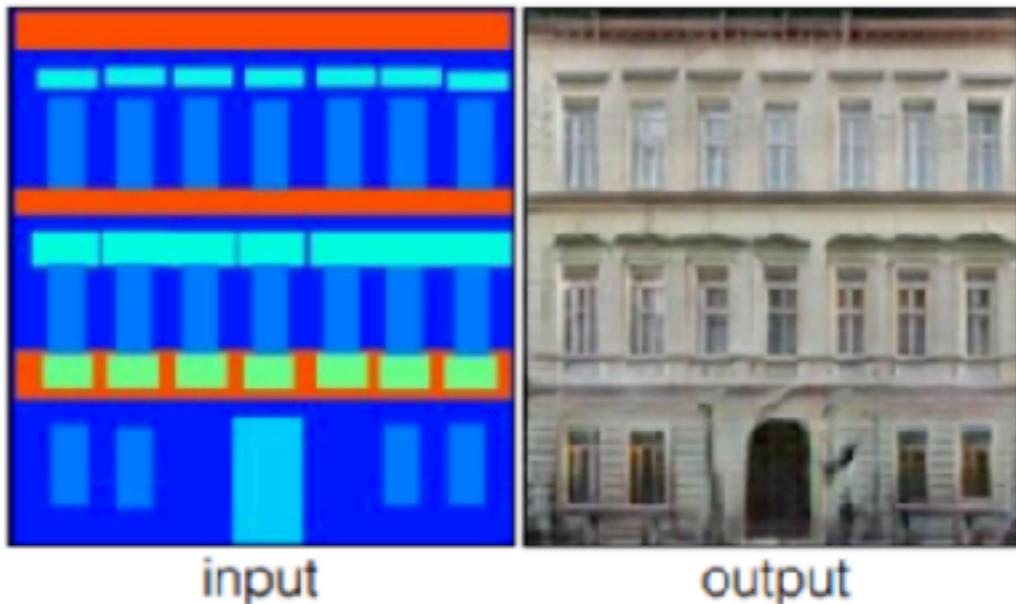


Today

- Generative Adversarial Networks (GANs)
- **(Unsupervised) Image Translation**

Image-To-Image Translation

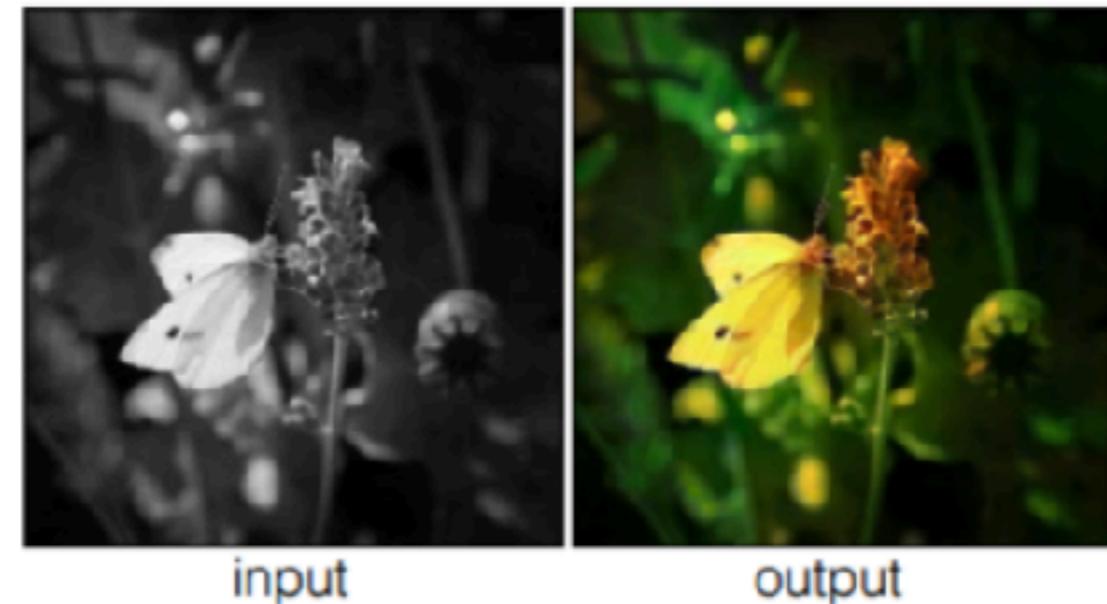
Labels to Facade



input

output

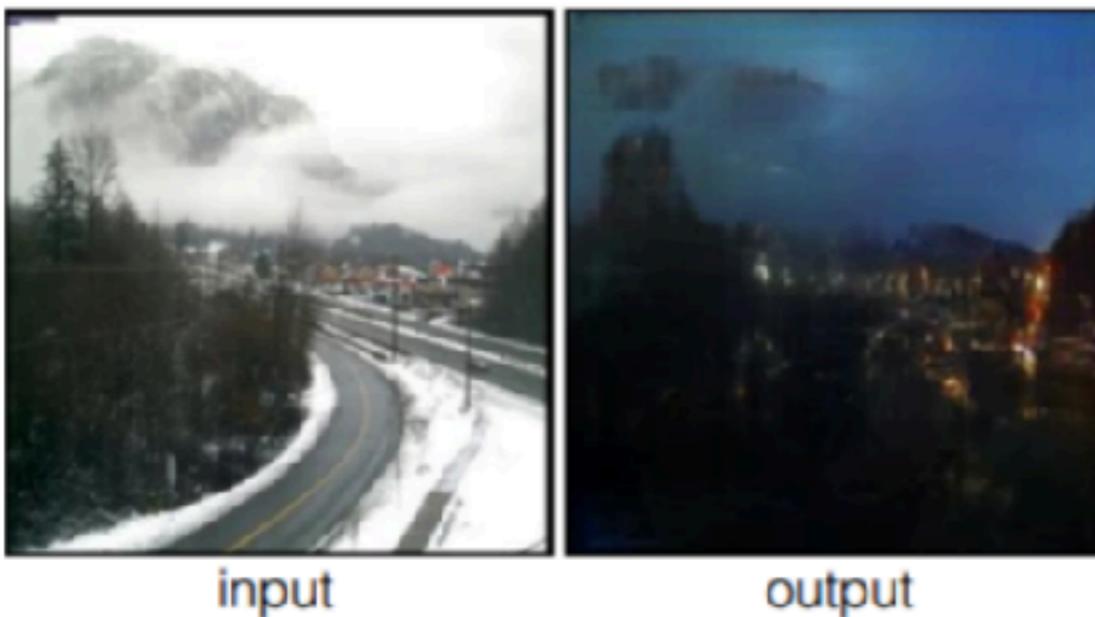
BW to Color



input

output

Day to Night



input

output

Edges to Photo

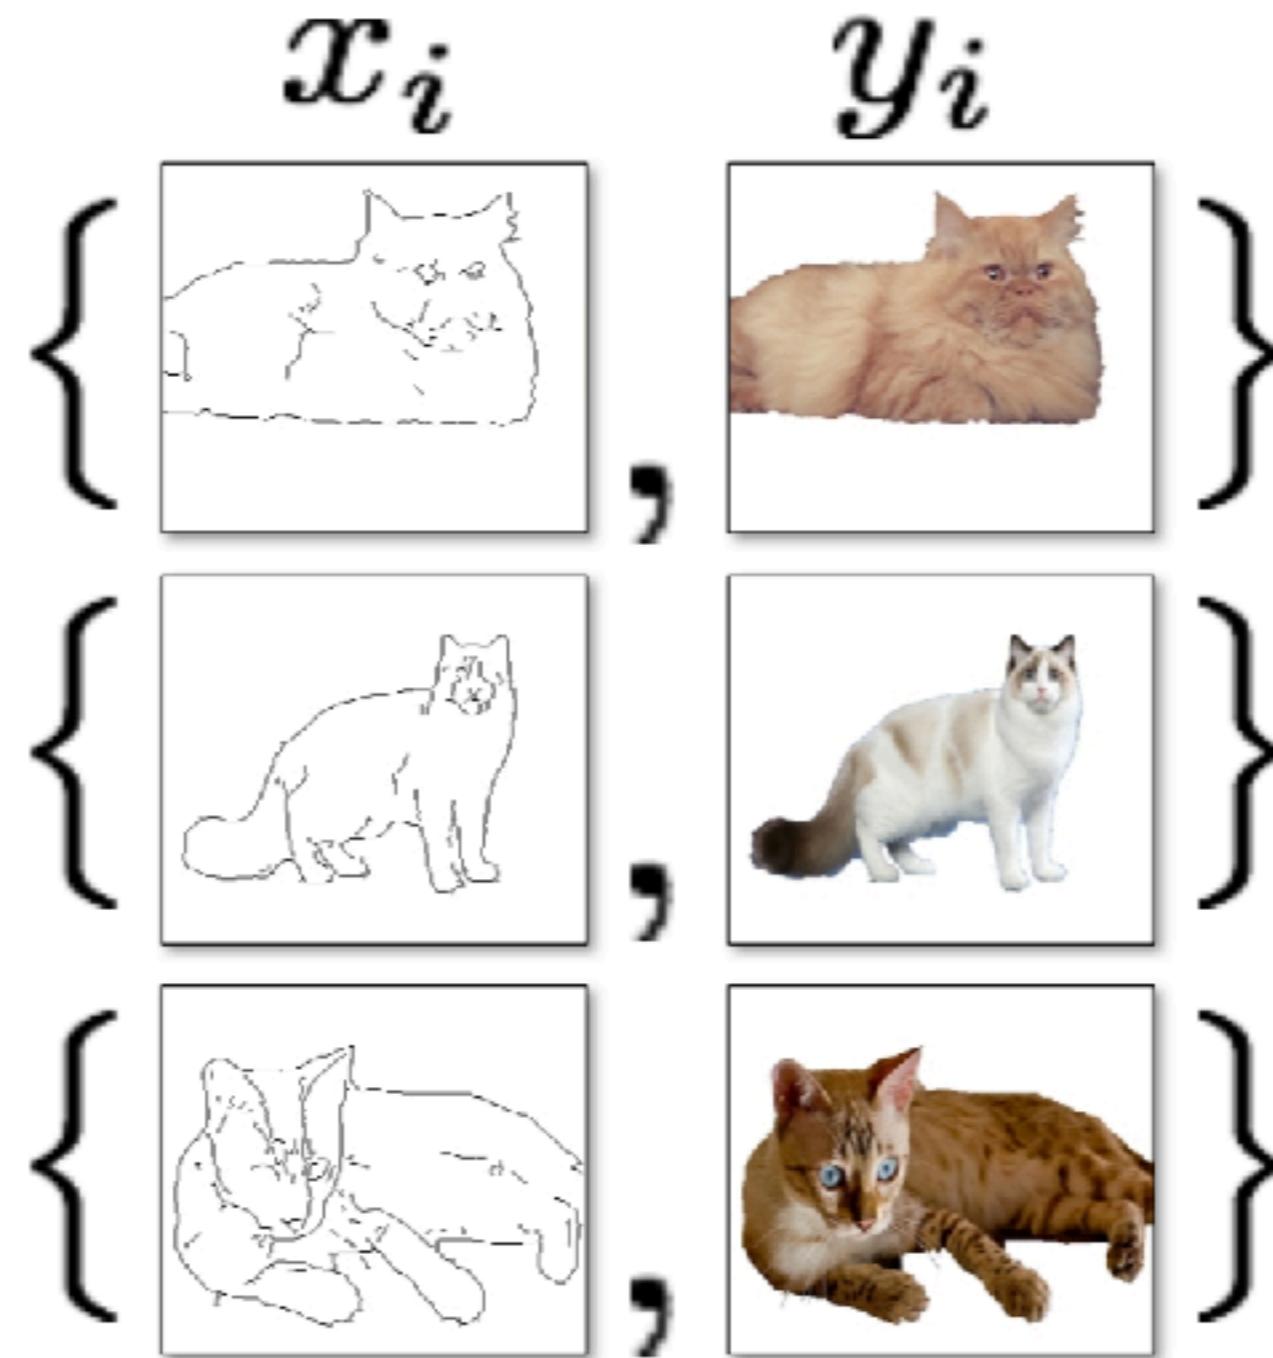


input

output

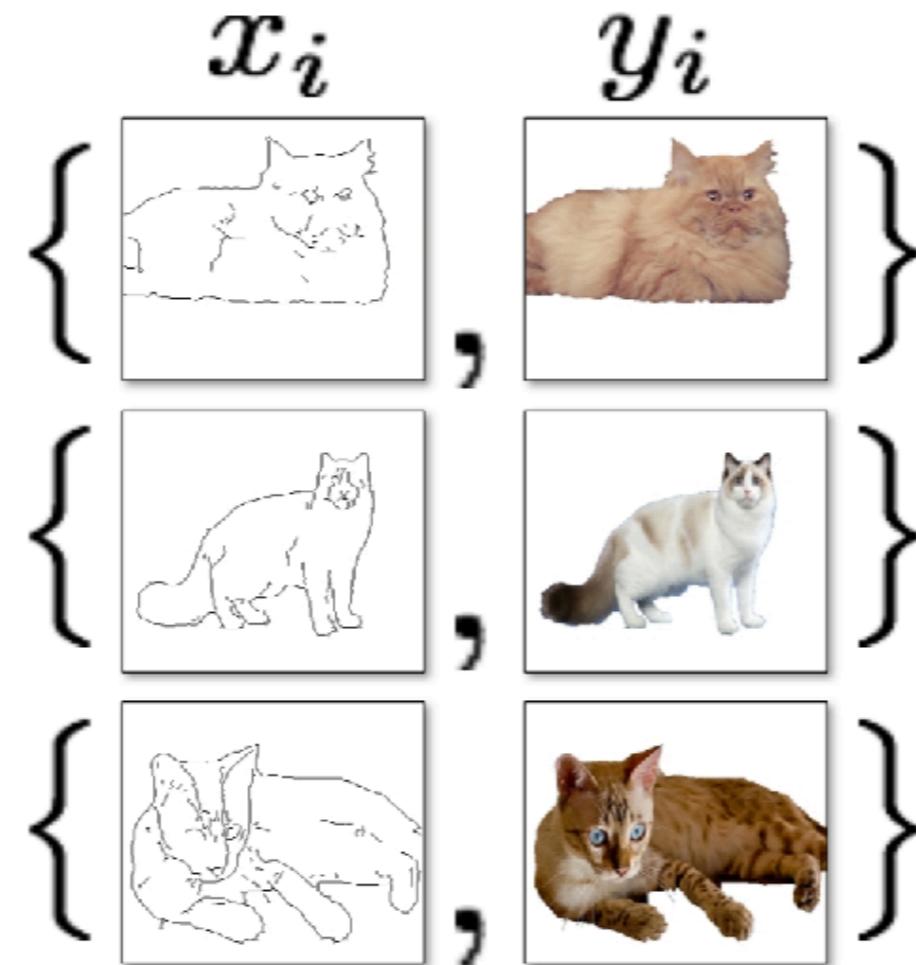
slide credit: Jun-Yan Zhu, Taesung Park

Supervised Image-To-Image Translation



slide credit: Jun-Yan Zhu, Taesung Park

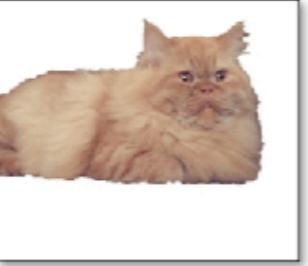
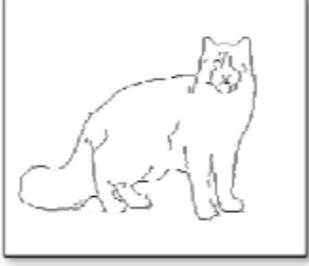
Supervised Image-To-Image Translation



$$\min_G \max_D E_{\mathbf{y}}[\log(D(\mathbf{y}))] + E_{\mathbf{x}}[\log(1 - D(G(\mathbf{x})))]$$

slide credit: Jun-Yan Zhu, Taesung Park

Supervised Image-To-Image Translation

x_i	y_i
{  }	, {  }
{  }	, {  }
{  }	, {  }

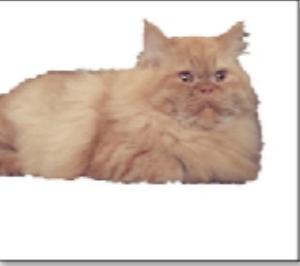
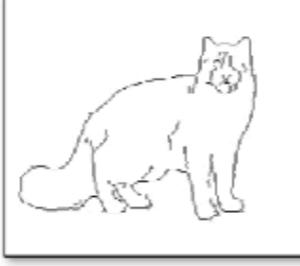
real cats



$$\min_G \max_D E_{\mathbf{y}}[\log(D(\mathbf{y}))] + E_{\mathbf{x}}[\log(1 - D(G(\mathbf{x})))]$$

slide credit: Jun-Yan Zhu, Taesung Park

Supervised Image-To-Image Translation

x_i	y_i
{  }	, 
{  }	, 
{  }	, 

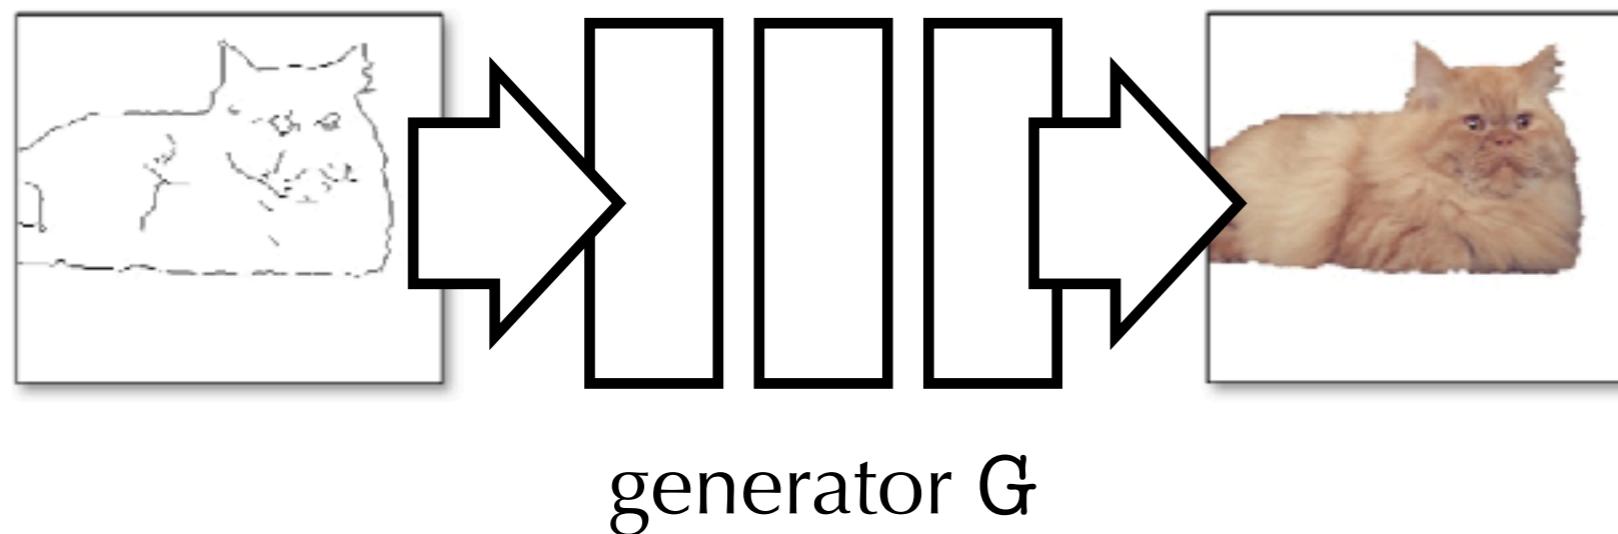
real cats

cat sketches

$$\min_G \max_D E_{\mathbf{y}}[\log(D(\mathbf{y}))] + E_{\mathbf{x}}[\log(1 - D(G(\mathbf{x})))]$$

slide credit: Jun-Yan Zhu, Taesung Park

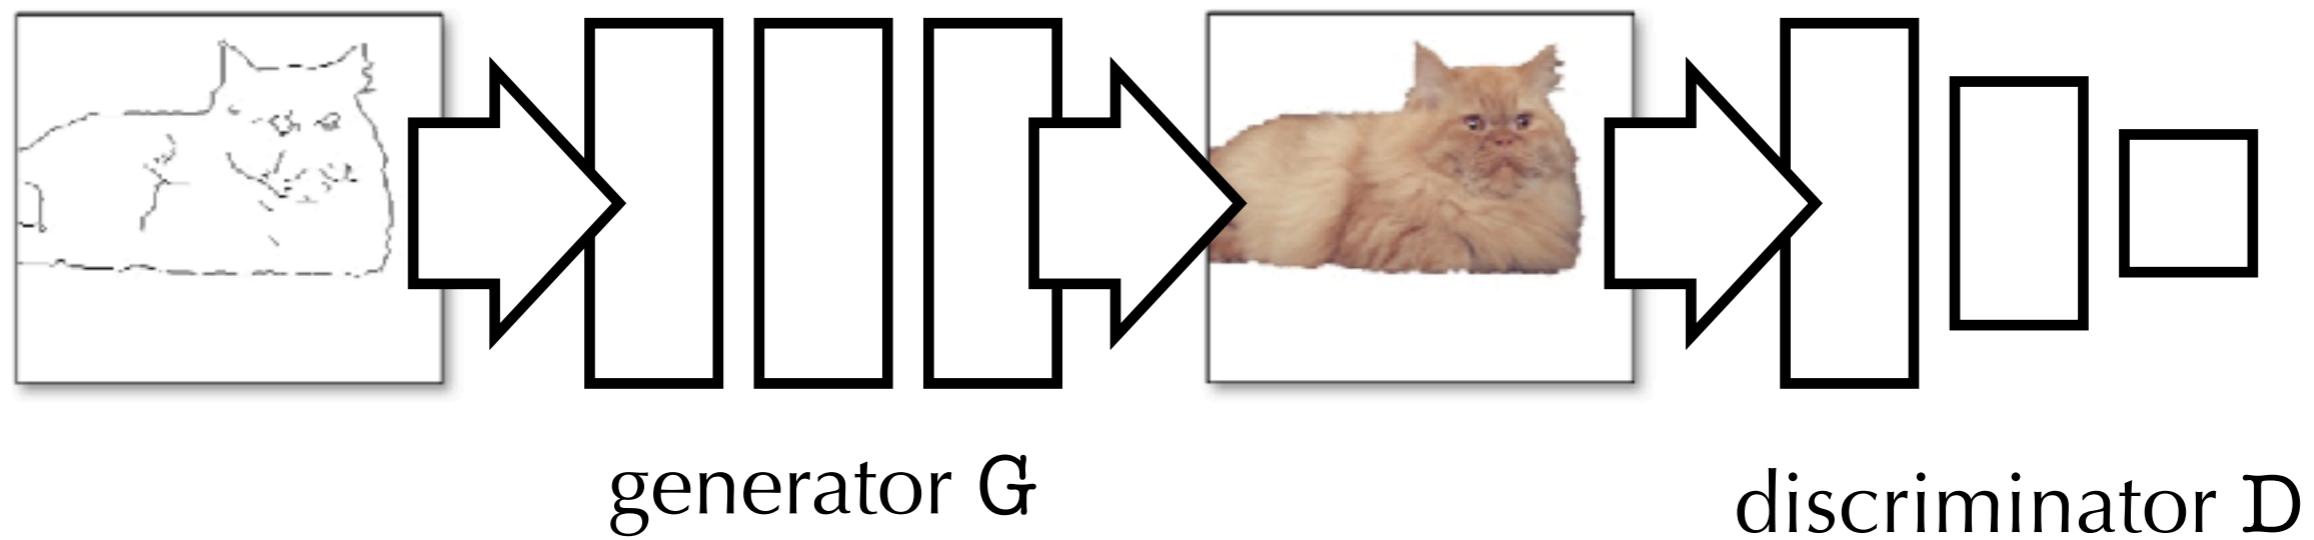
Supervised Image-To-Image Translation



$$\min_G \max_D E_{\mathbf{y}}[\log(D(\mathbf{y}))] + E_{\mathbf{x}}[\log(1 - D(G(\mathbf{x})))]$$

slide credit: Jun-Yan Zhu, Taesung Park

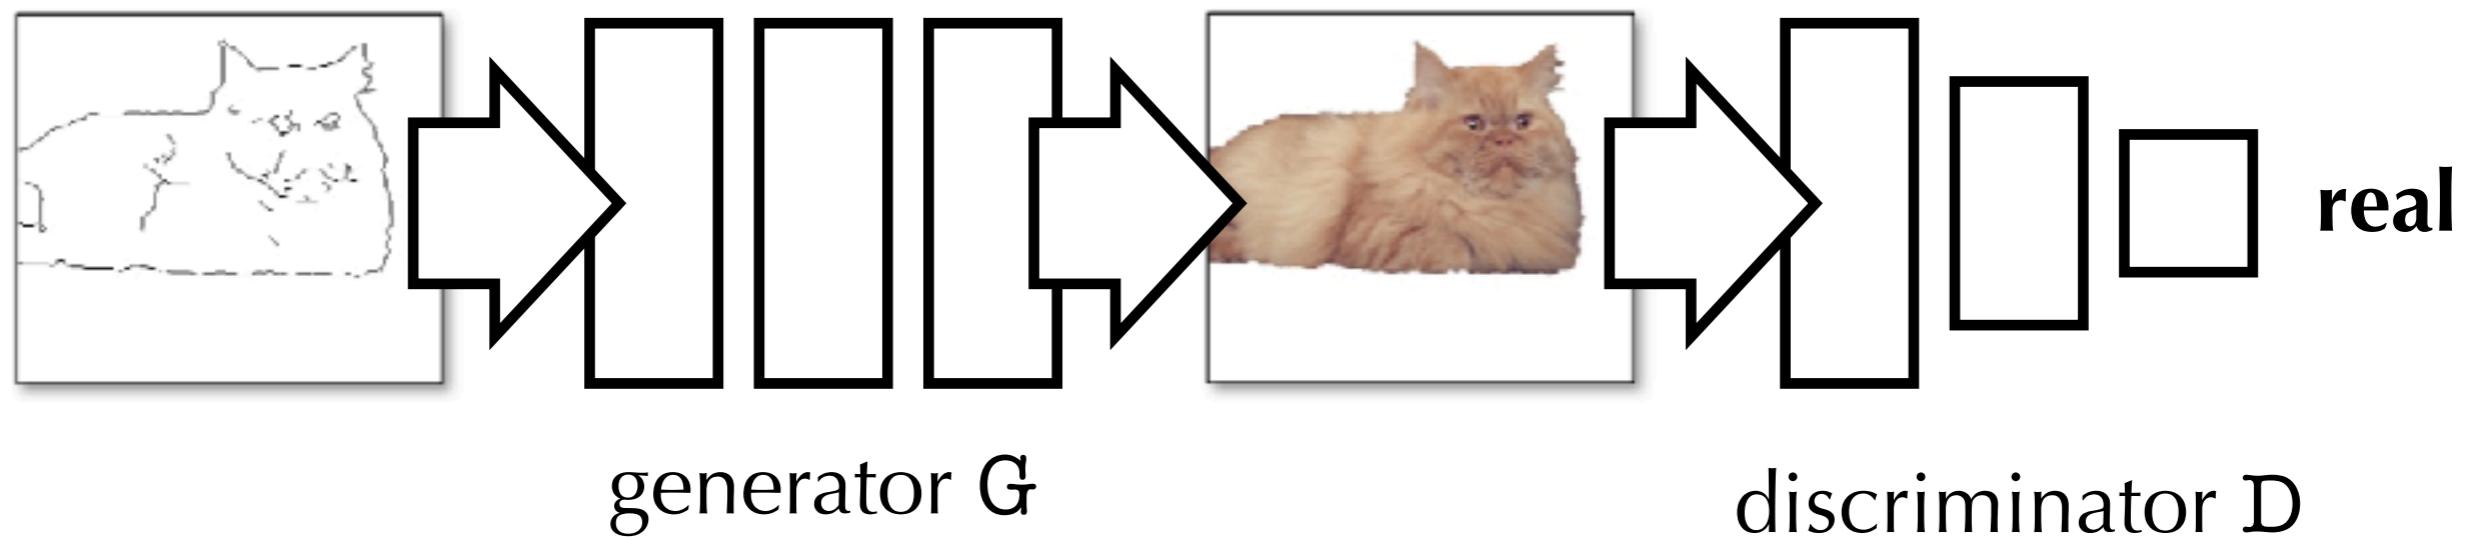
Supervised Image-To-Image Translation



$$\min_G \max_D E_{\mathbf{y}}[\log(D(\mathbf{y}))] + E_{\mathbf{x}}[\log(1 - D(G(\mathbf{x})))]$$

slide credit: Jun-Yan Zhu, Taesung Park

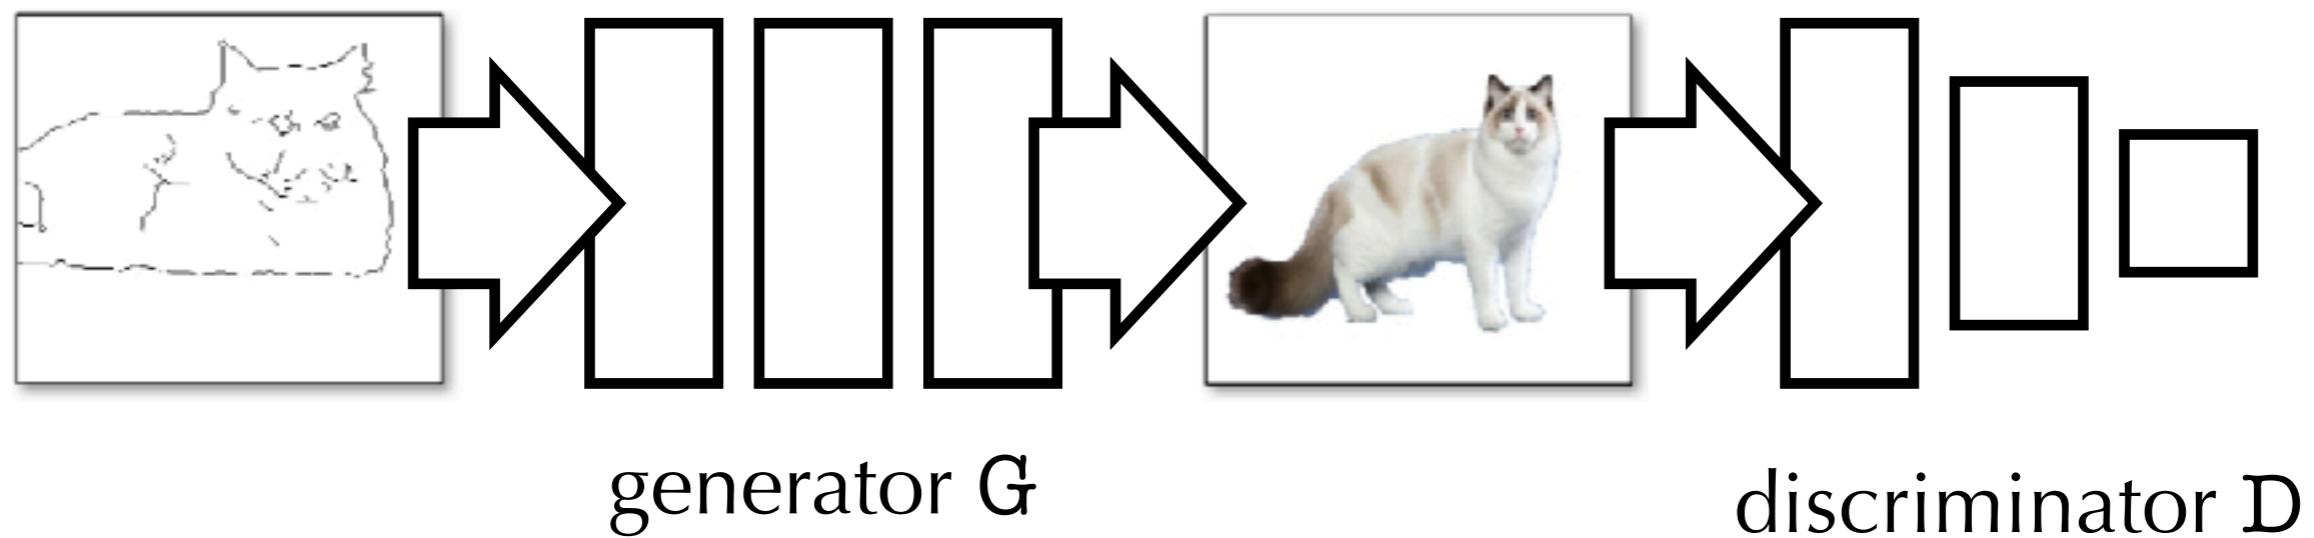
Supervised Image-To-Image Translation



$$\min_G \max_D E_{\mathbf{y}}[\log(D(\mathbf{y}))] + E_{\mathbf{x}}[\log(1 - D(G(\mathbf{x})))]$$

slide credit: Jun-Yan Zhu, Taesung Park

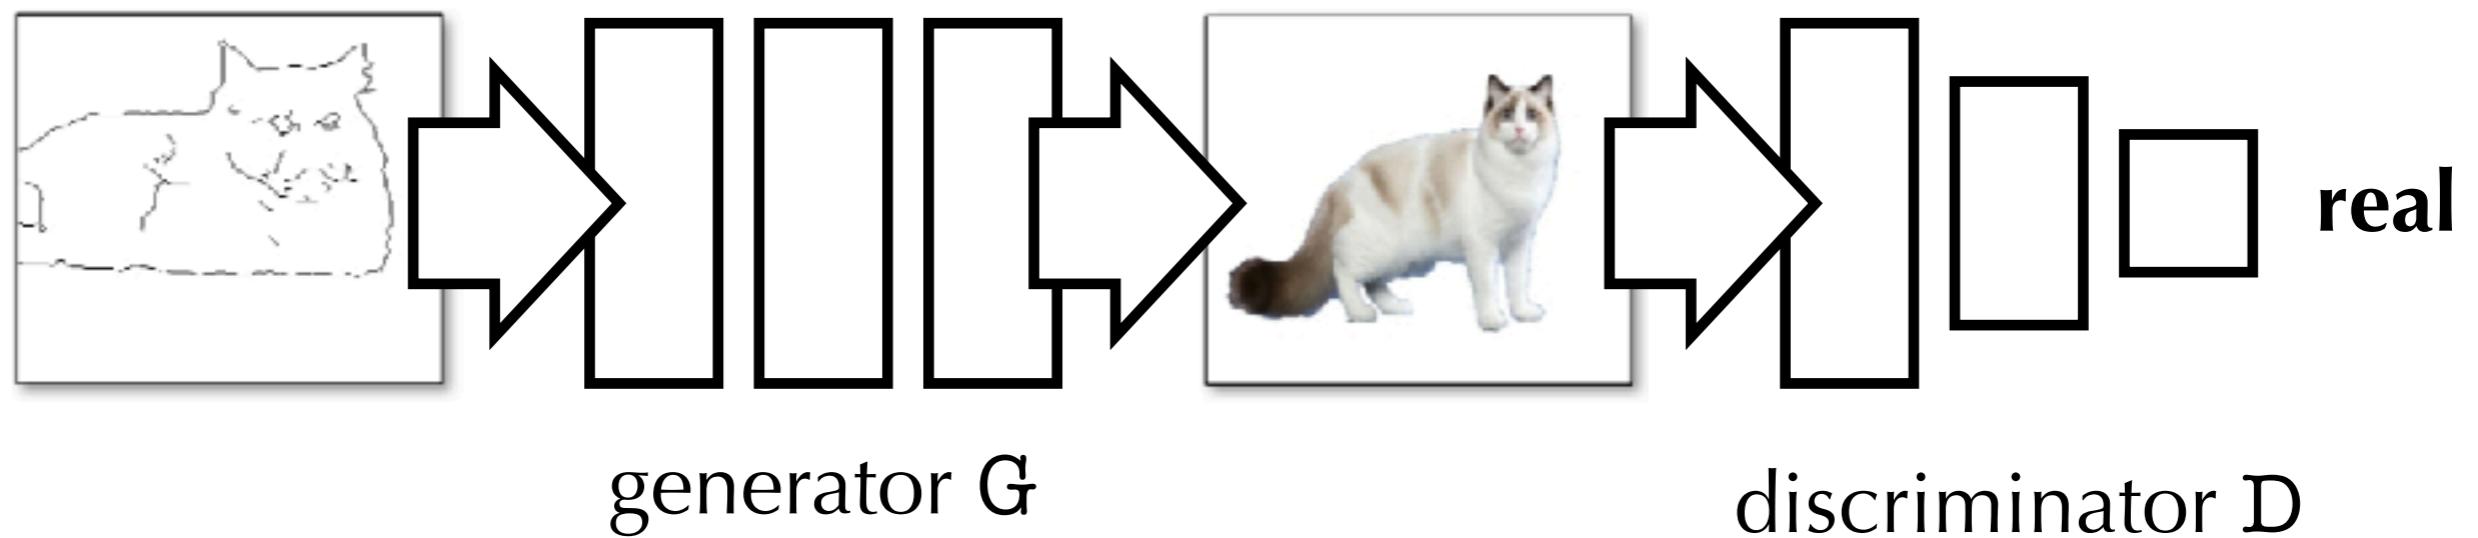
Supervised Image-To-Image Translation



$$\min_G \max_D E_{\mathbf{y}}[\log(D(\mathbf{y}))] + E_{\mathbf{x}}[\log(1 - D(G(\mathbf{x})))]$$

slide credit: Jun-Yan Zhu, Taesung Park

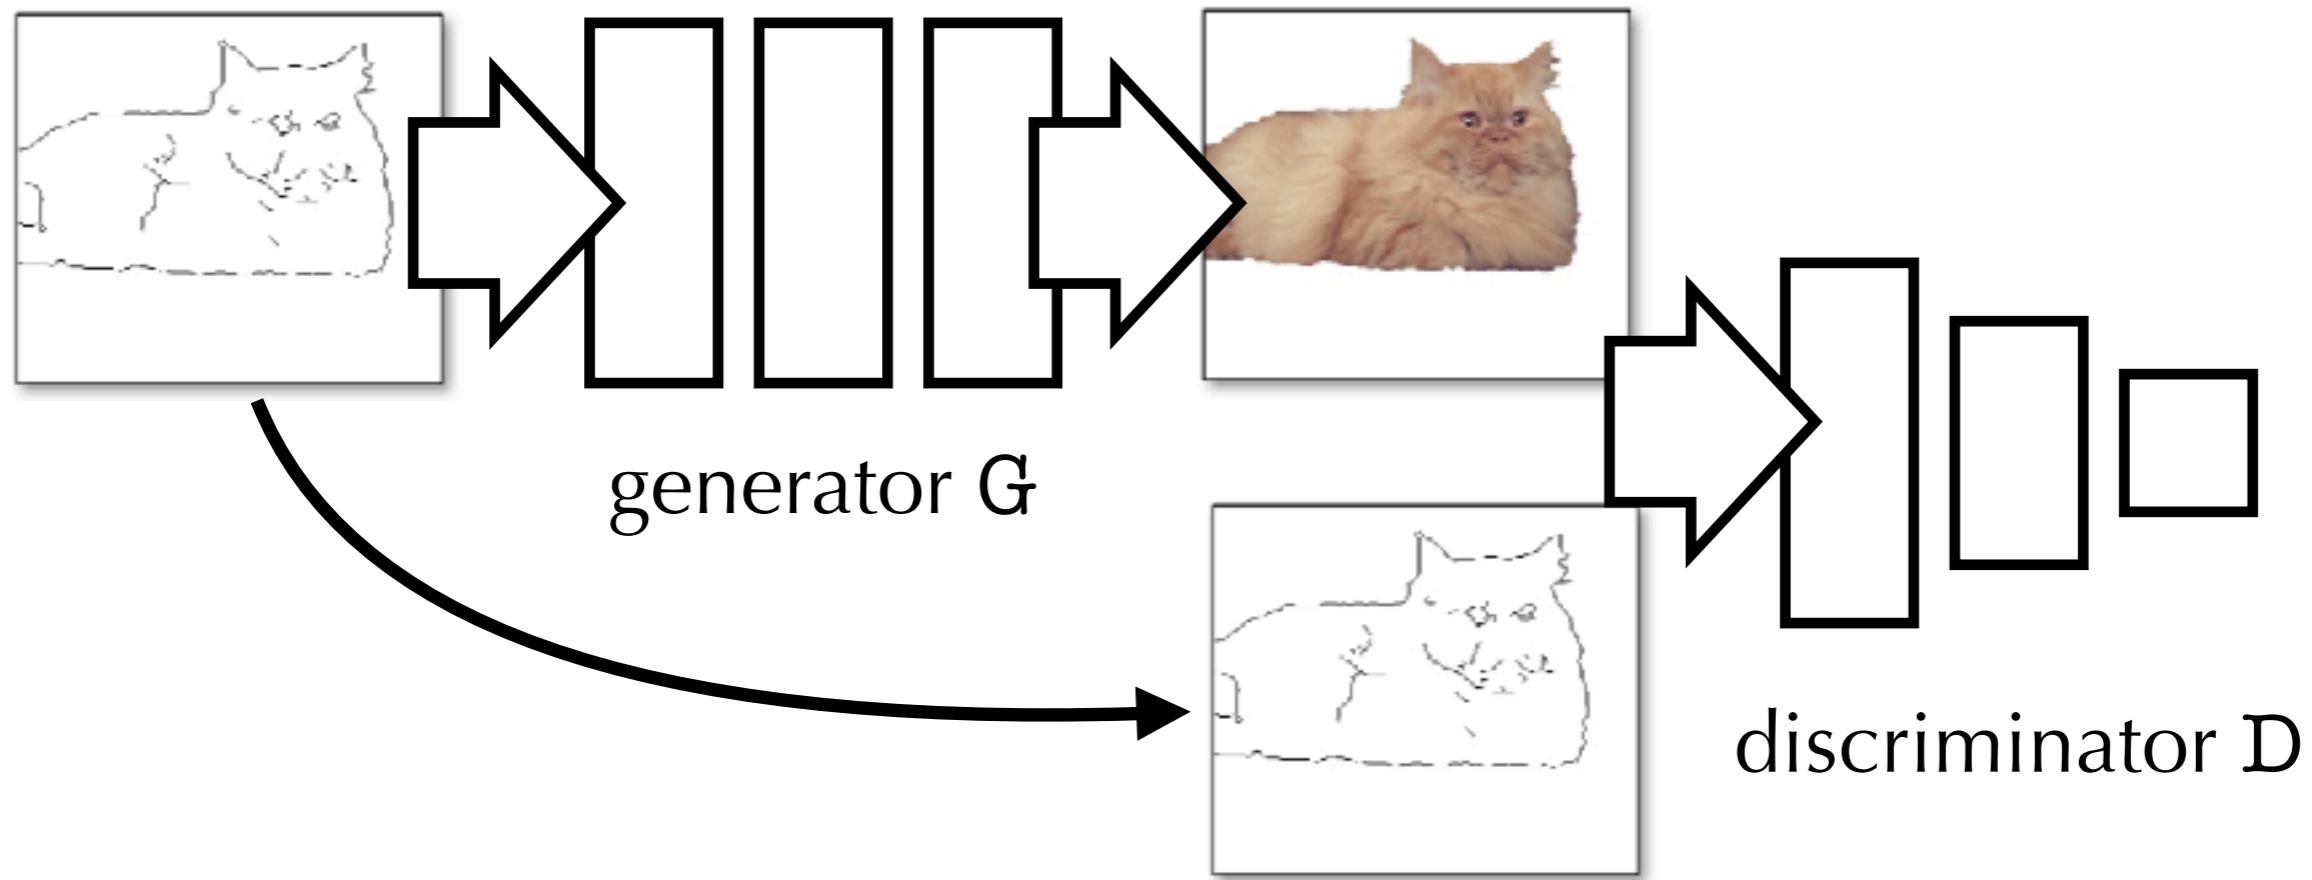
Supervised Image-To-Image Translation



$$\min_G \max_D E_{\mathbf{y}}[\log(D(\mathbf{y}))] + E_{\mathbf{x}}[\log(1 - D(G(\mathbf{x})))]$$

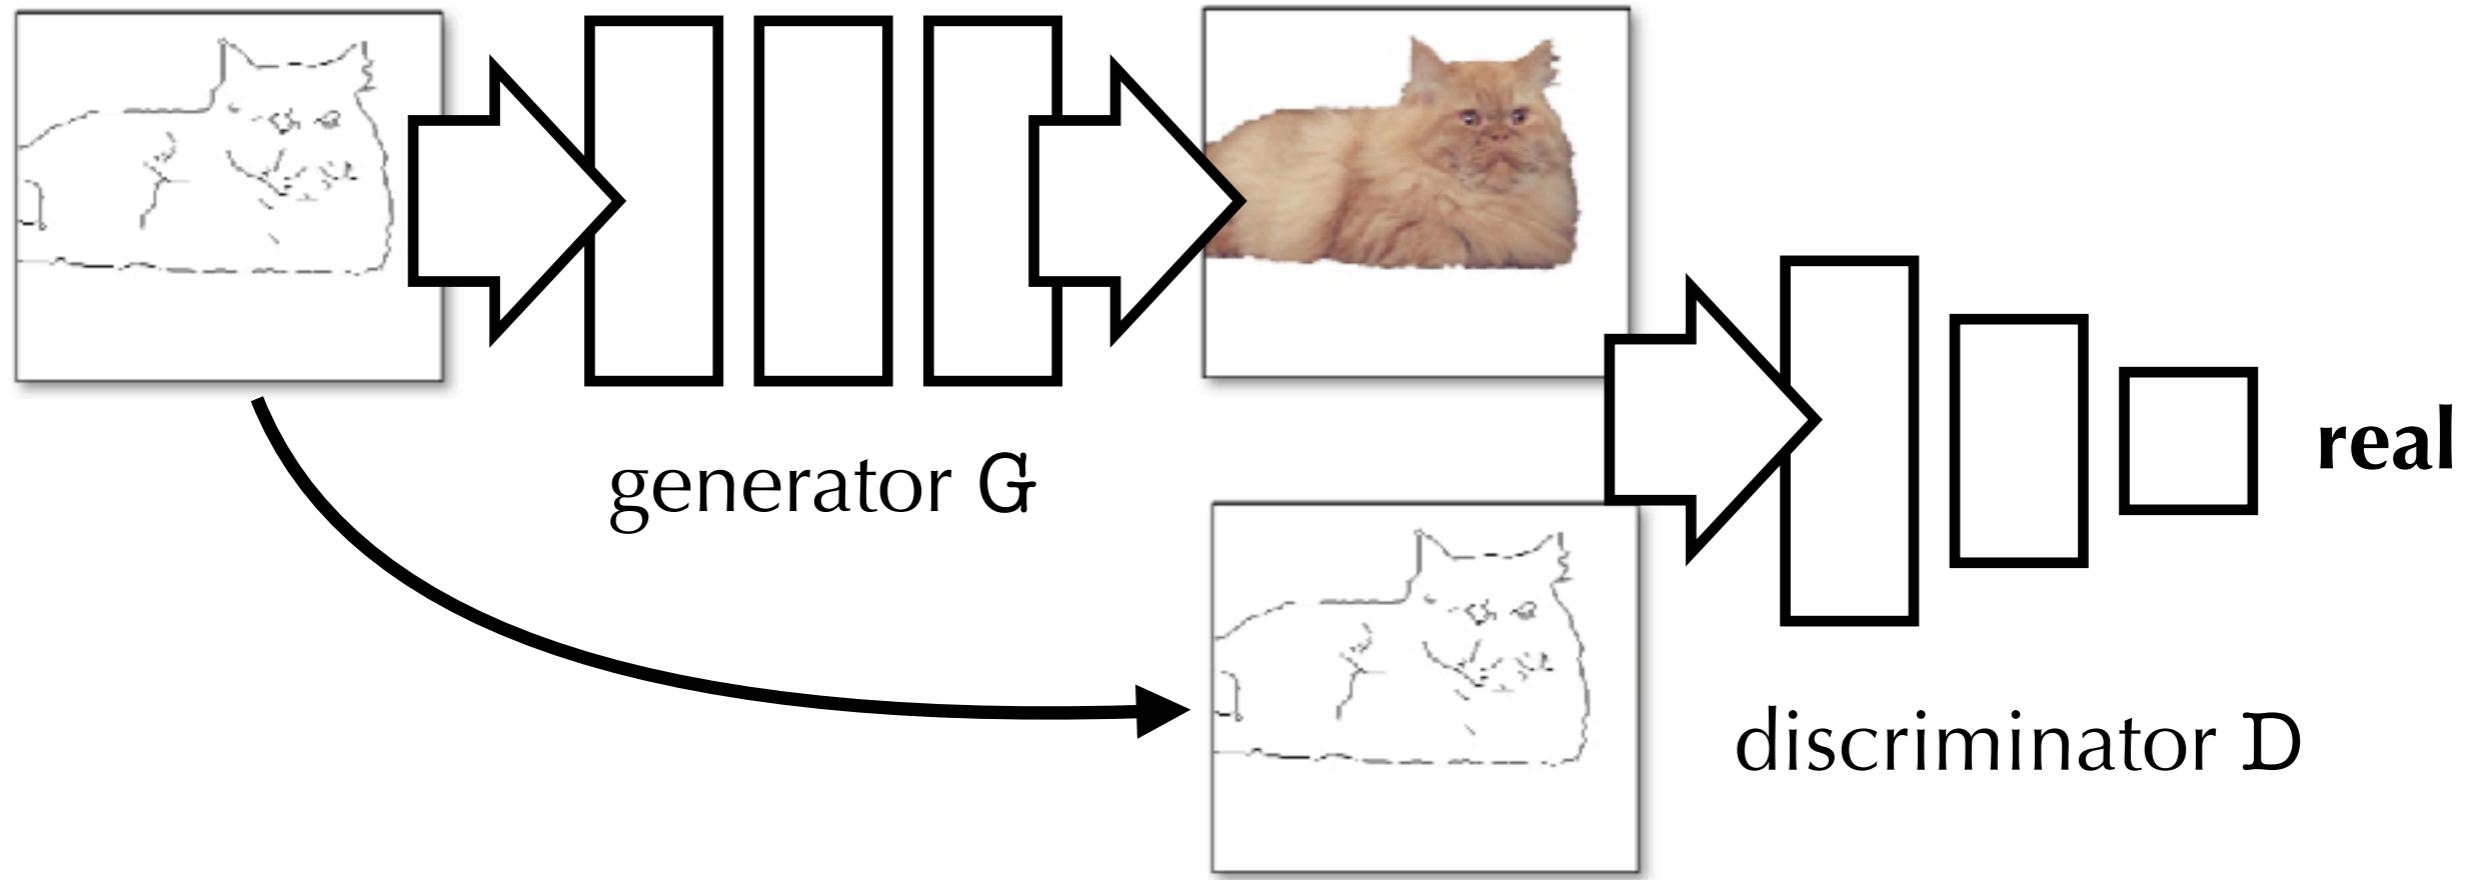
slide credit: Jun-Yan Zhu, Taesung Park

Conditional GANs



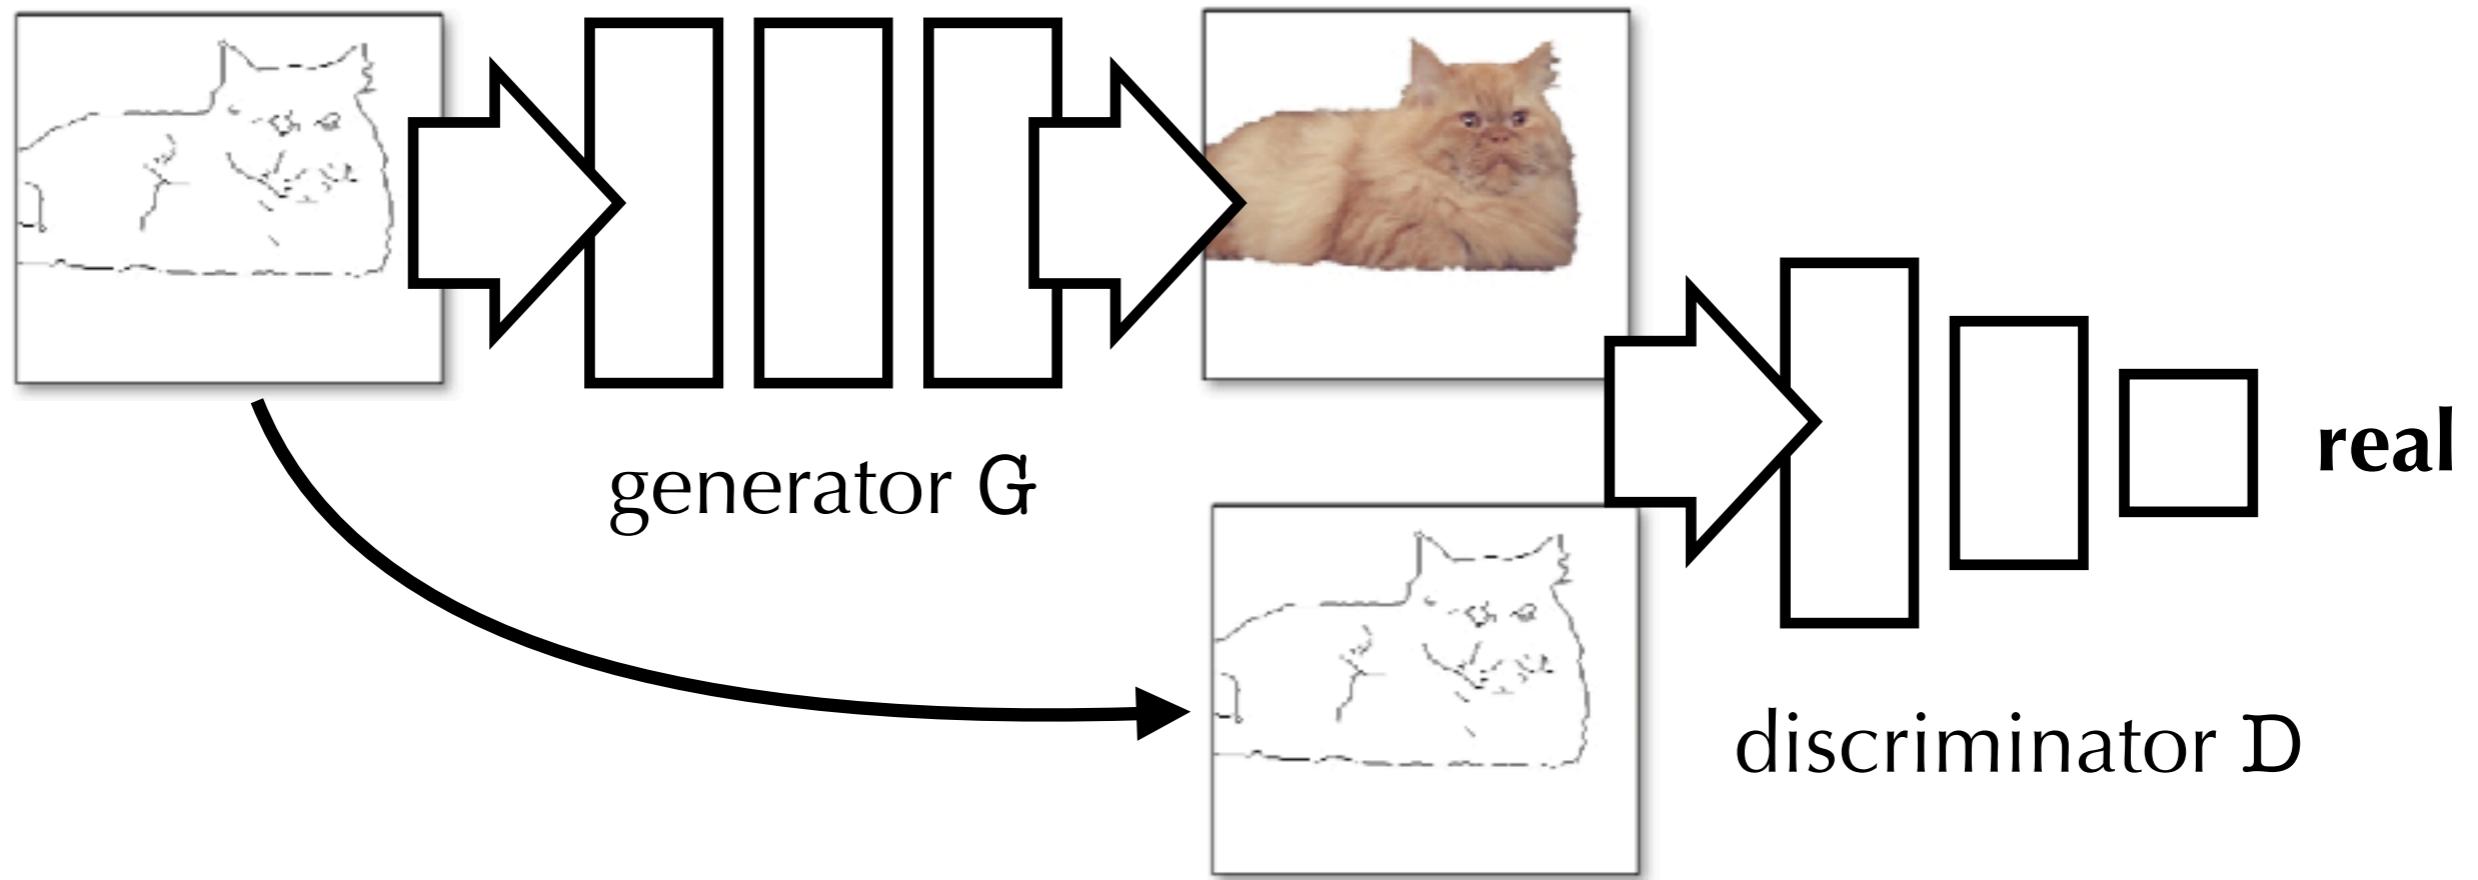
[Isola et al., Image-to-Image Translation with Conditional Adversarial Networks, CVPR 2017]

Conditional GANs



[Isola et al., Image-to-Image Translation with Conditional Adversarial Networks, CVPR 2017]

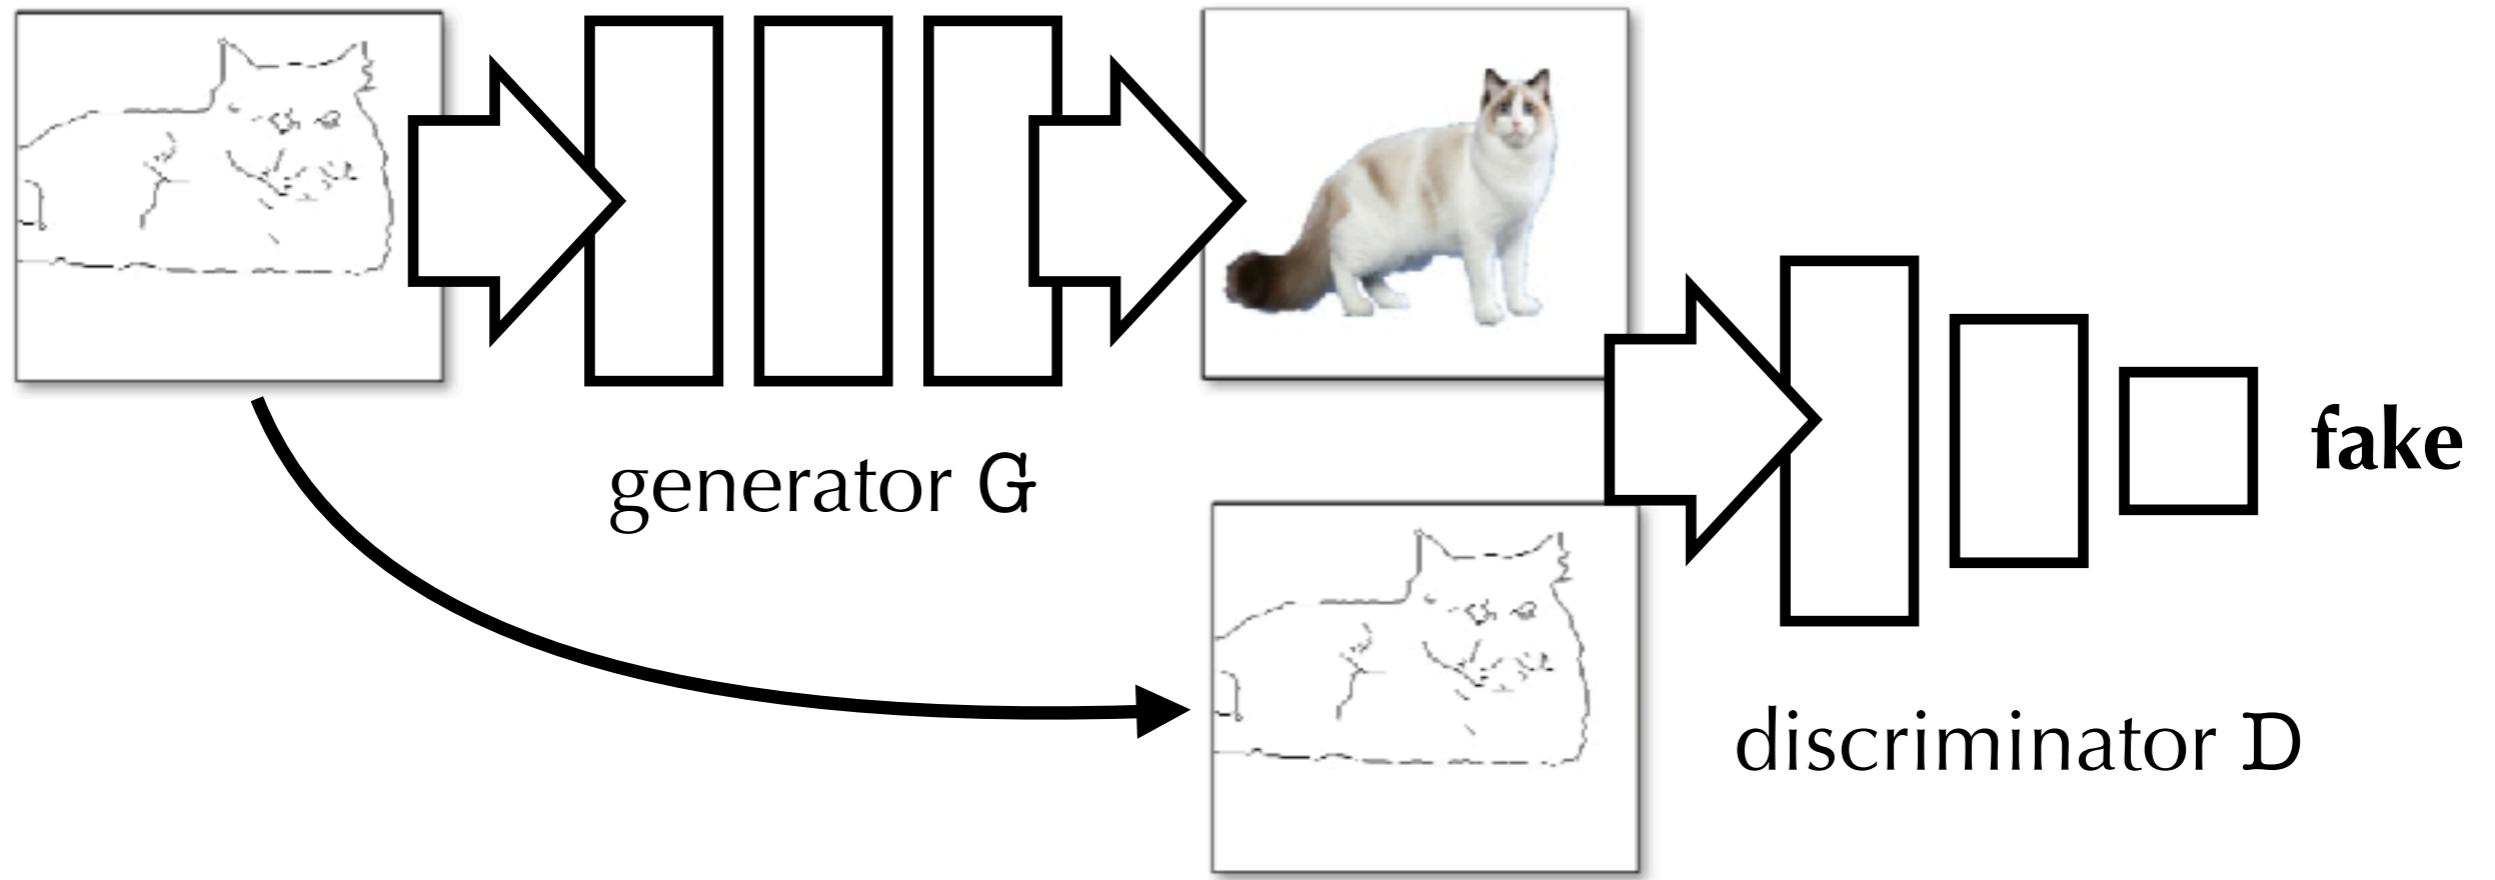
Conditional GANs



$$\min_G \max_D E_{\mathbf{x}, \mathbf{y}} [\log(D(\mathbf{x}, \mathbf{y}))] + E_{\mathbf{x}, \mathbf{y}} [\log(1 - D(\mathbf{x}, G(\mathbf{x})))]$$

[Isola et al., Image-to-Image Translation with Conditional Adversarial Networks, CVPR 2017]

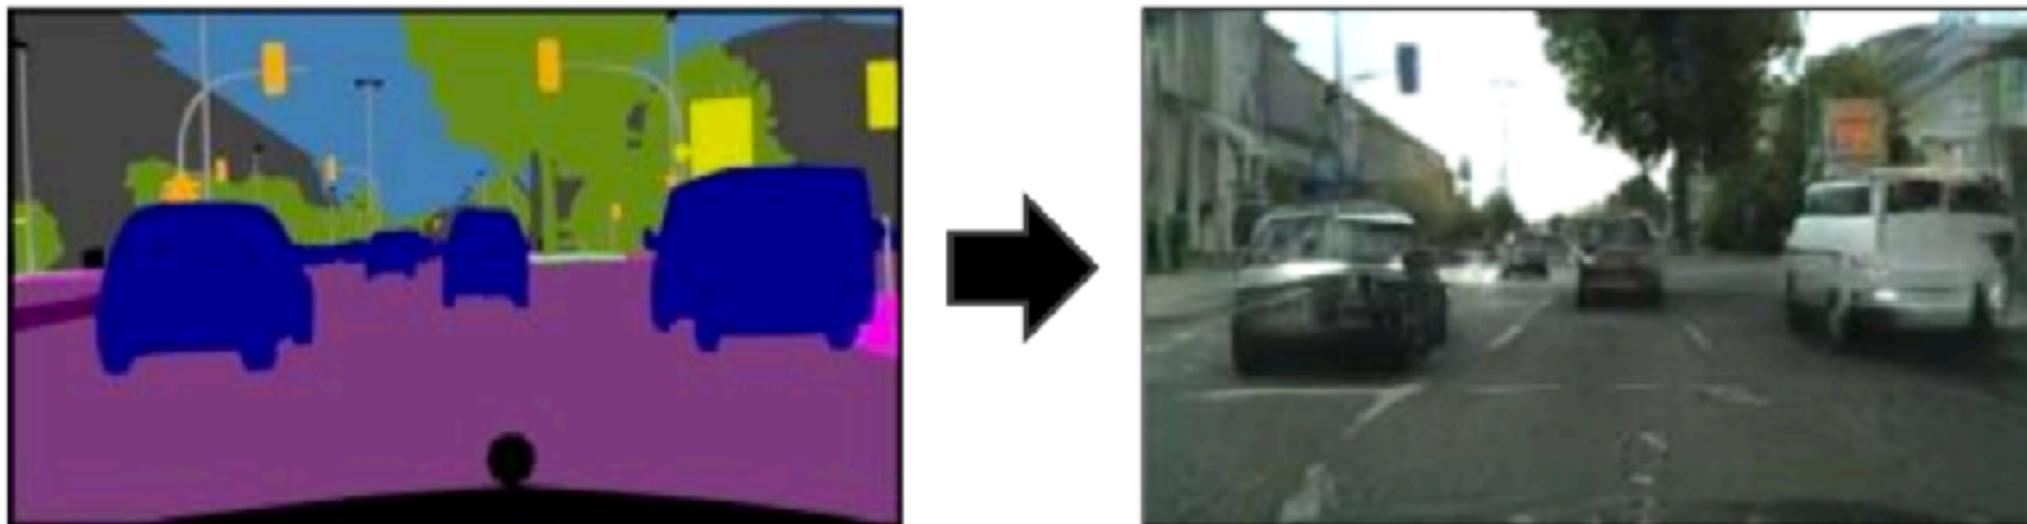
Conditional GANs



$$\min_G \max_D E_{\mathbf{x}, \mathbf{y}} [\log(D(\mathbf{x}, \mathbf{y}))] + E_{\mathbf{x}, \mathbf{y}} [\log(1 - D(\mathbf{x}, G(\mathbf{x})))]$$

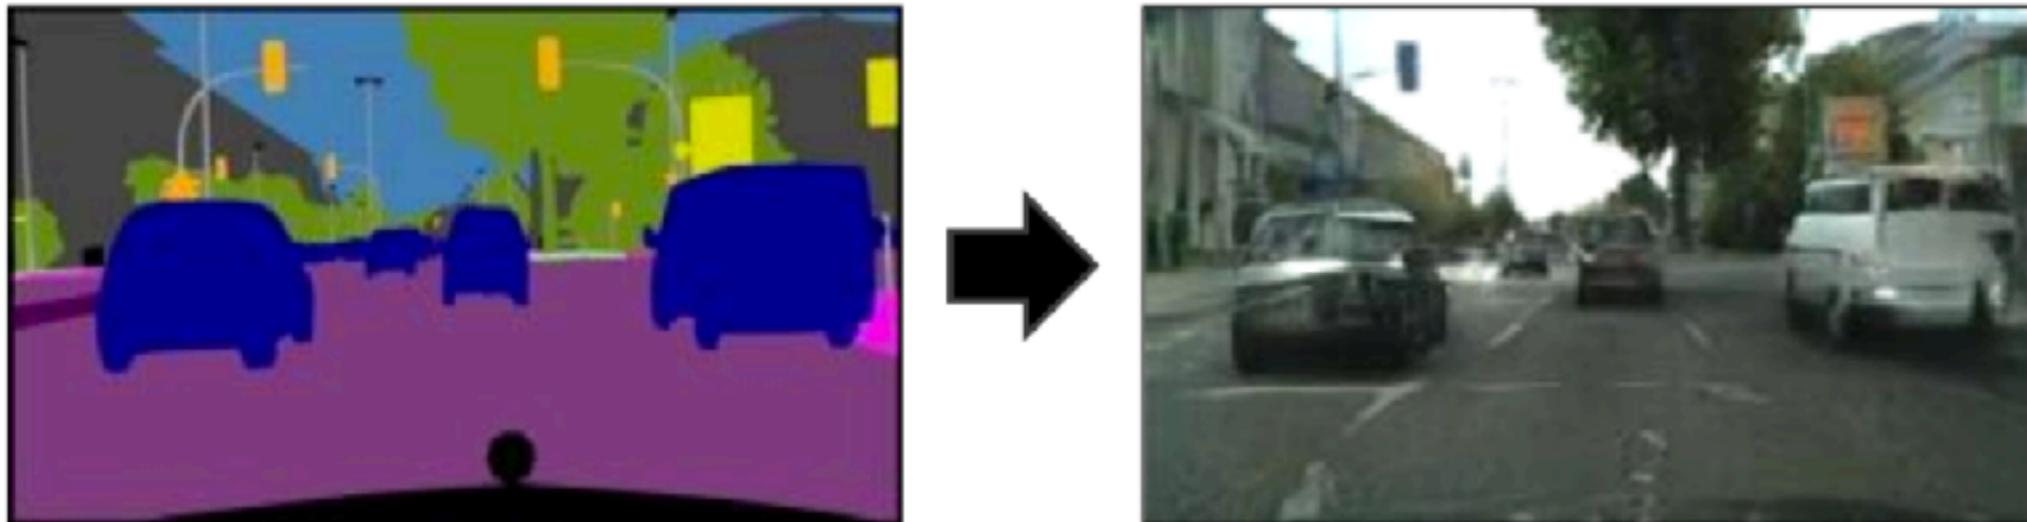
[Isola et al., Image-to-Image Translation with Conditional Adversarial Networks, CVPR 2017]

Supervised Image-To-Image Translation

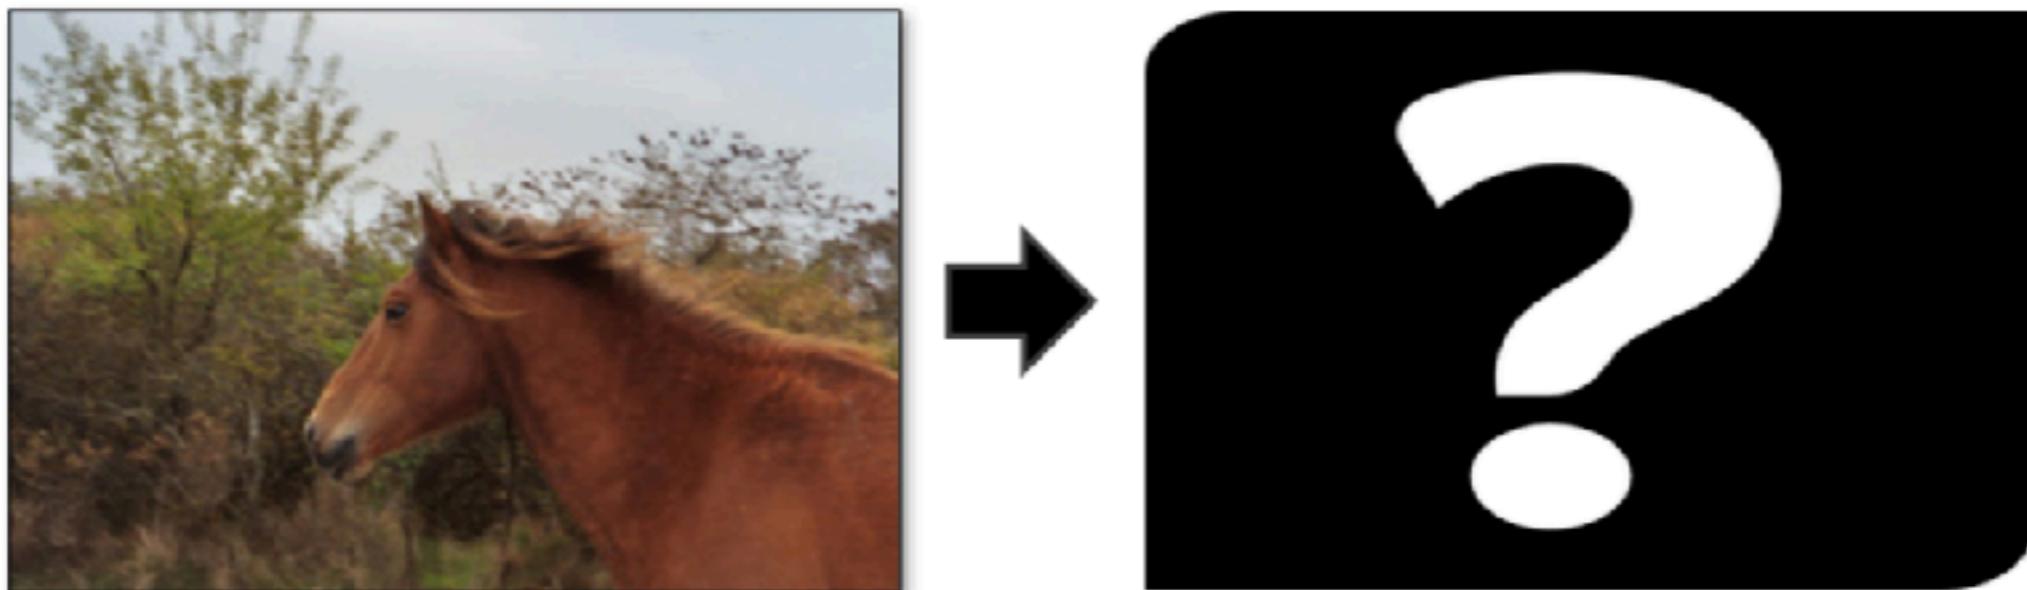


Label \leftrightarrow photo: per-pixel labeling

Supervised Image-To-Image Translation



Label \leftrightarrow photo: per-pixel labeling



Labelled data often hard or impossible to obtain

slide credit: Jun-Yan Zhu, Taesung Park

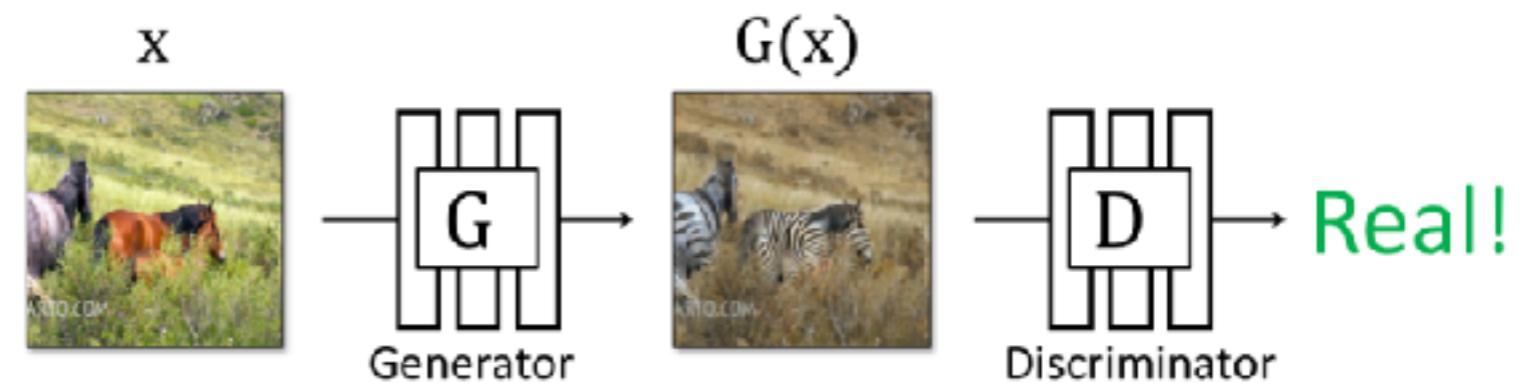
Unsupervised Image-To-Image Translation



independent sets,
no pairing!

slide credit: Jun-Yan Zhu, Taesung Park

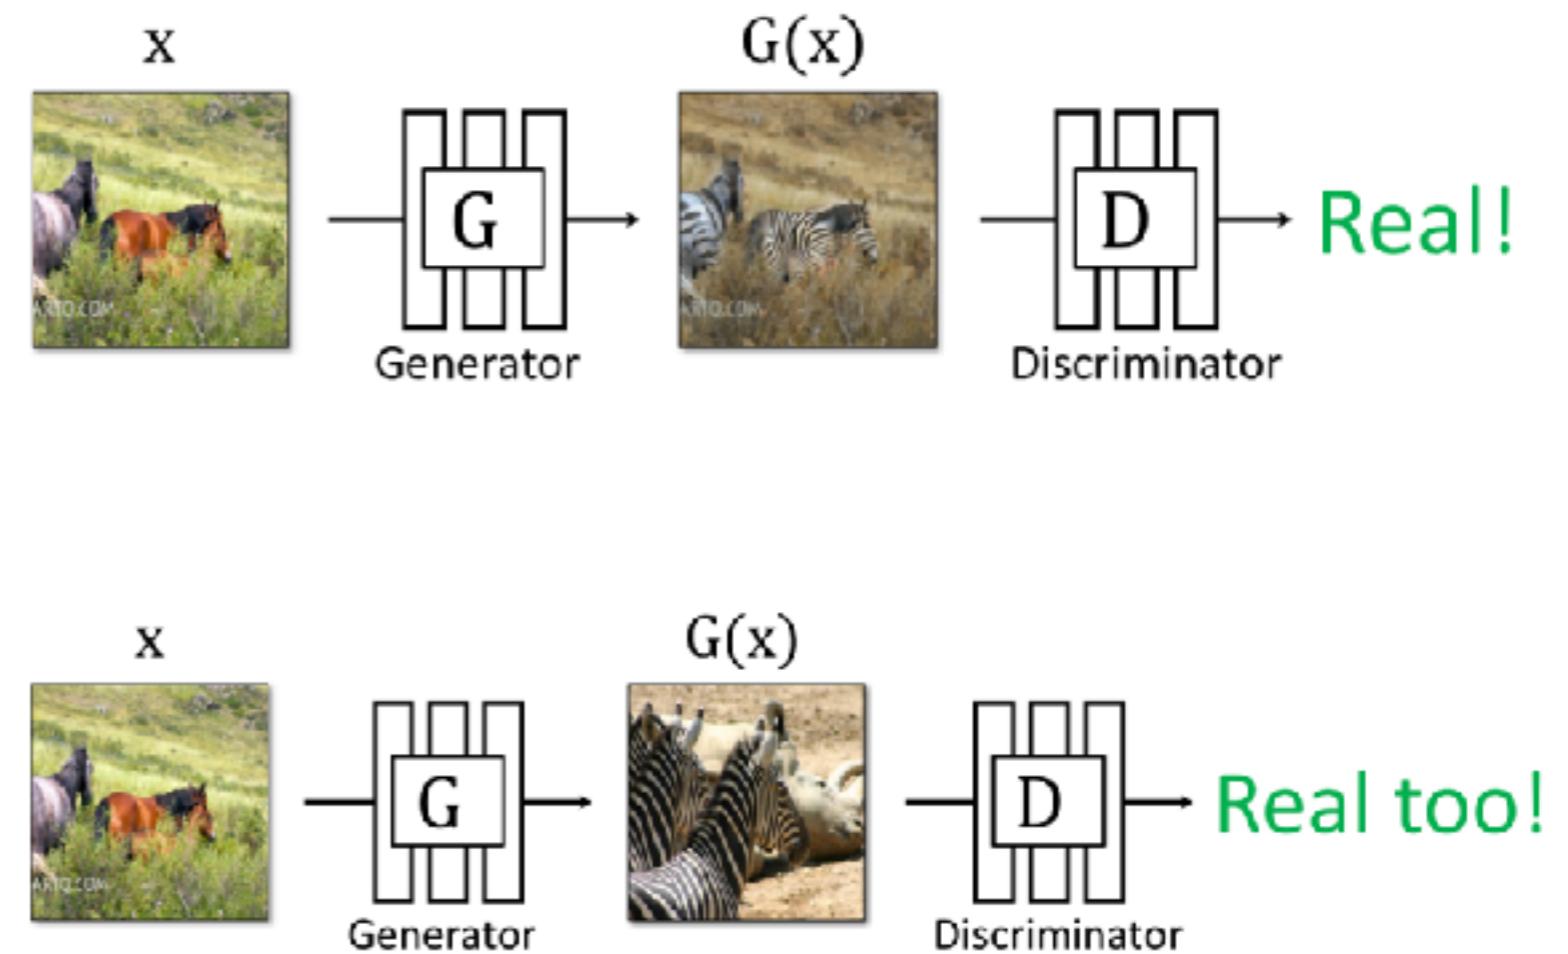
Unsupervised Image-To-Image Translation



independent sets,
no pairing!

slide credit: Jun-Yan Zhu, Taesung Park

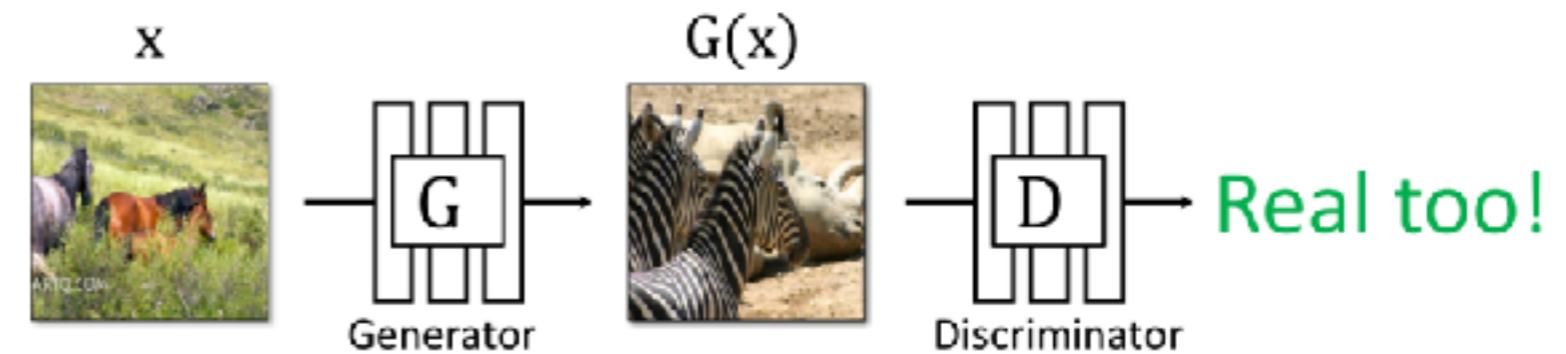
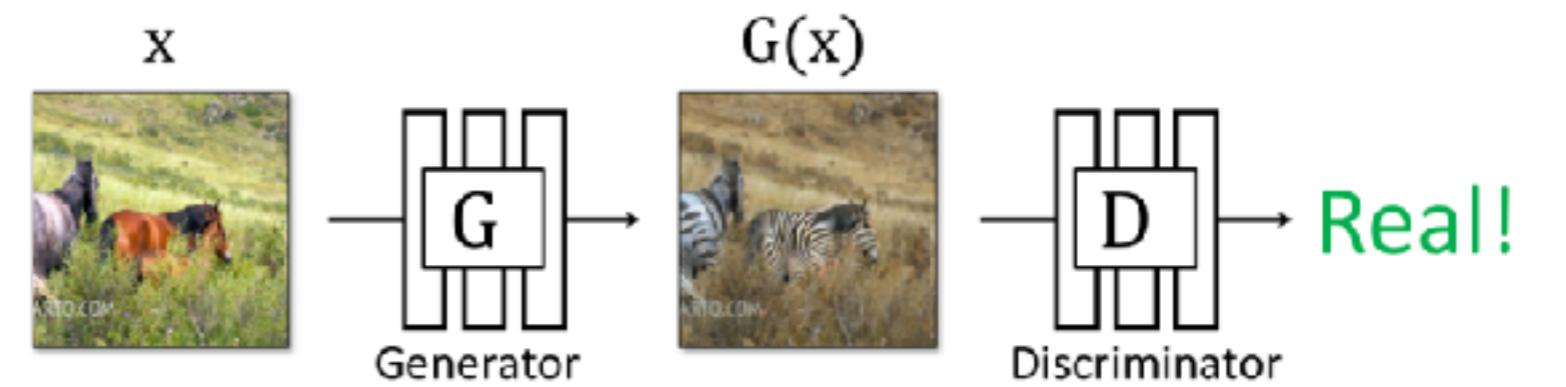
Unsupervised Image-To-Image Translation



independent sets,
no pairing!

slide credit: Jun-Yan Zhu, Taesung Park

Unsupervised Image-To-Image Translation

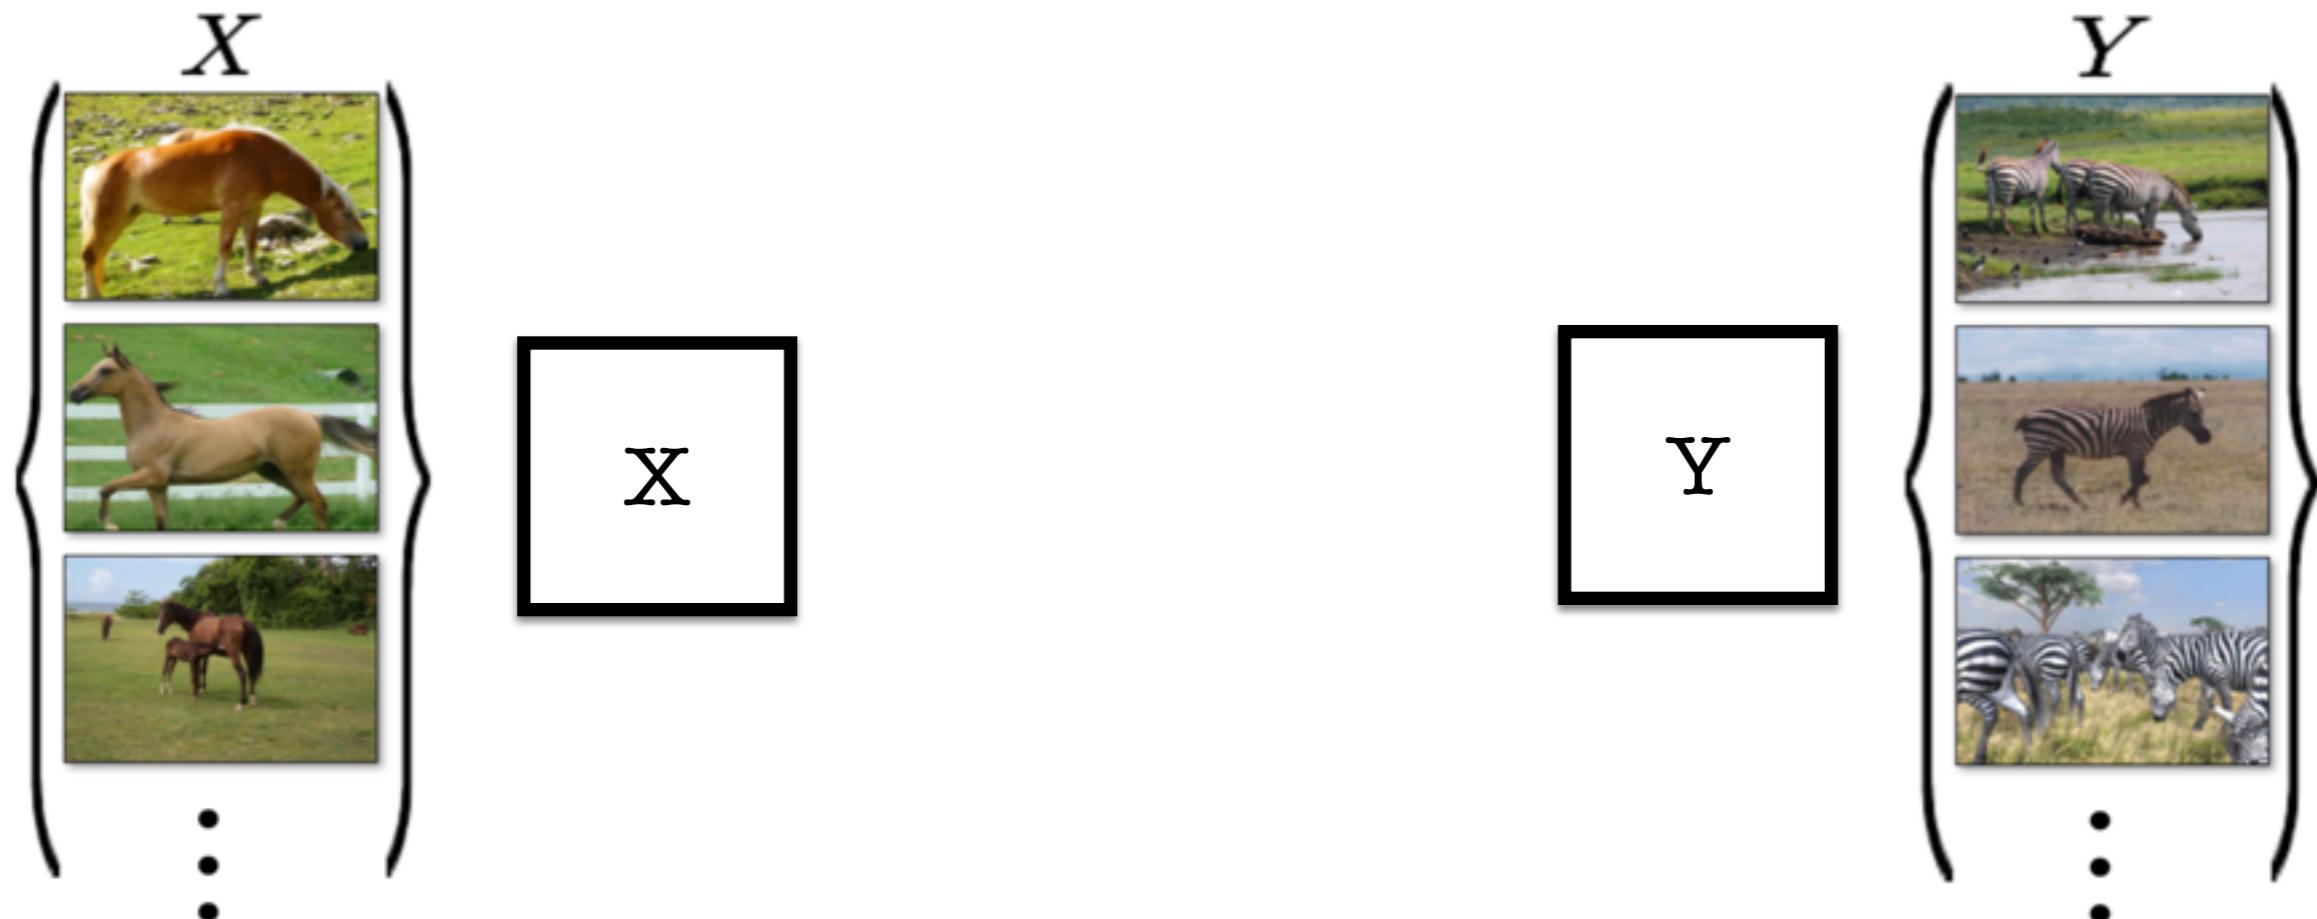


mode collapse

independent sets,
no pairing!

slide credit: Jun-Yan Zhu, Taesung Park

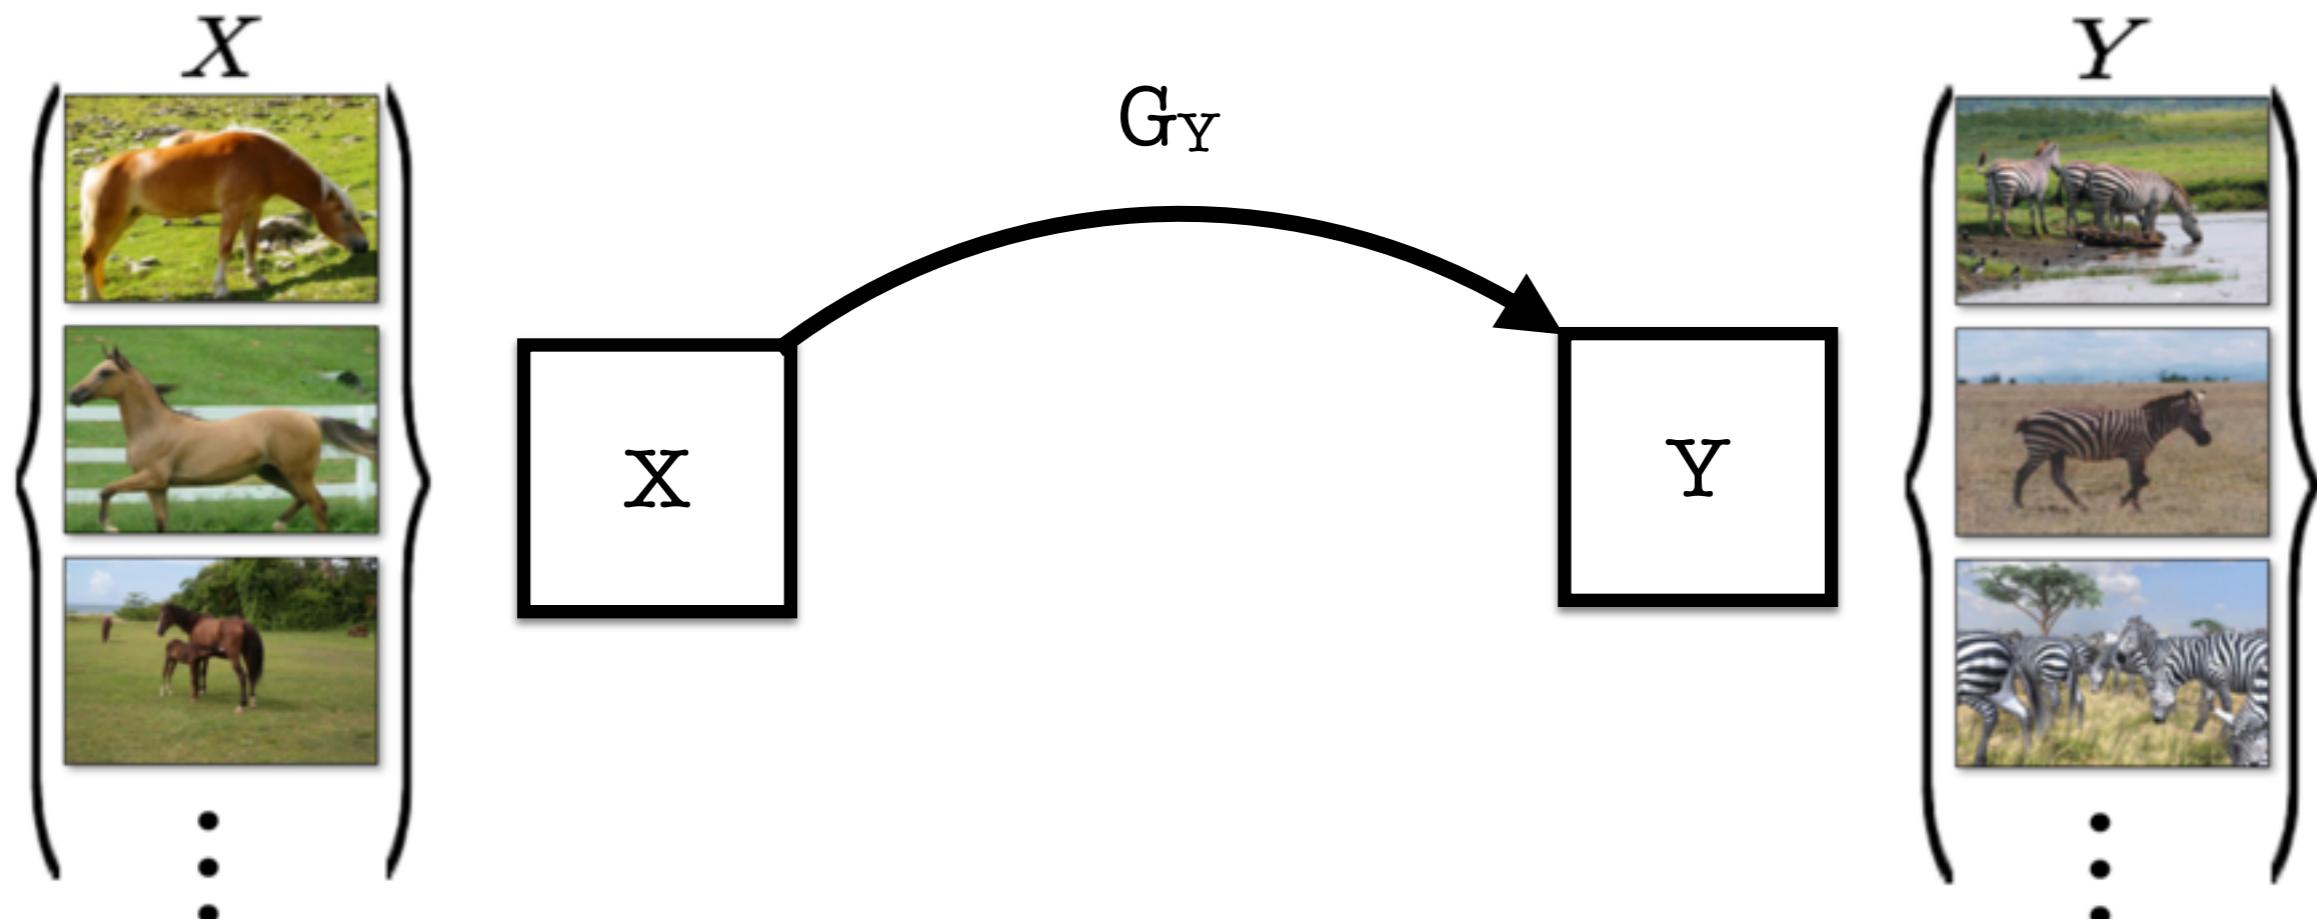
CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

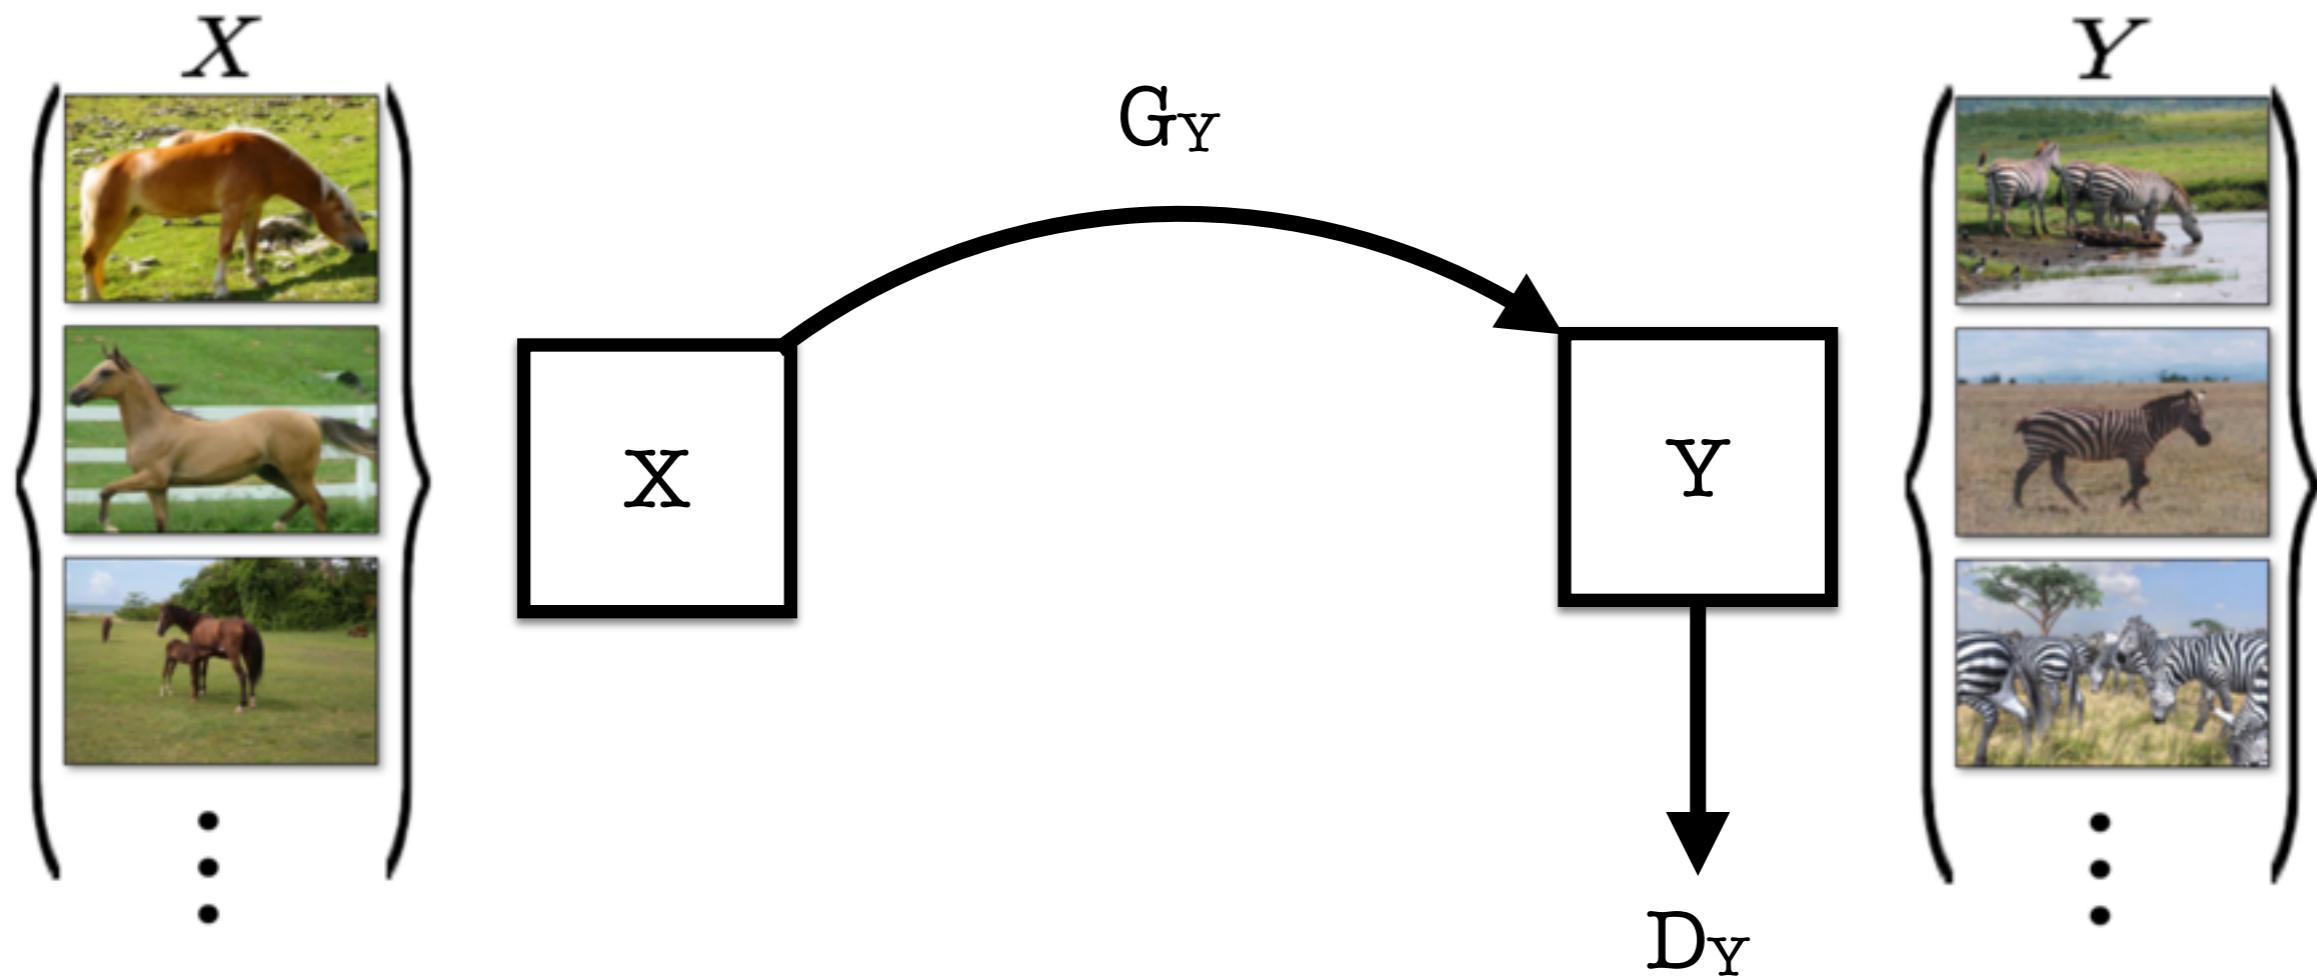
CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

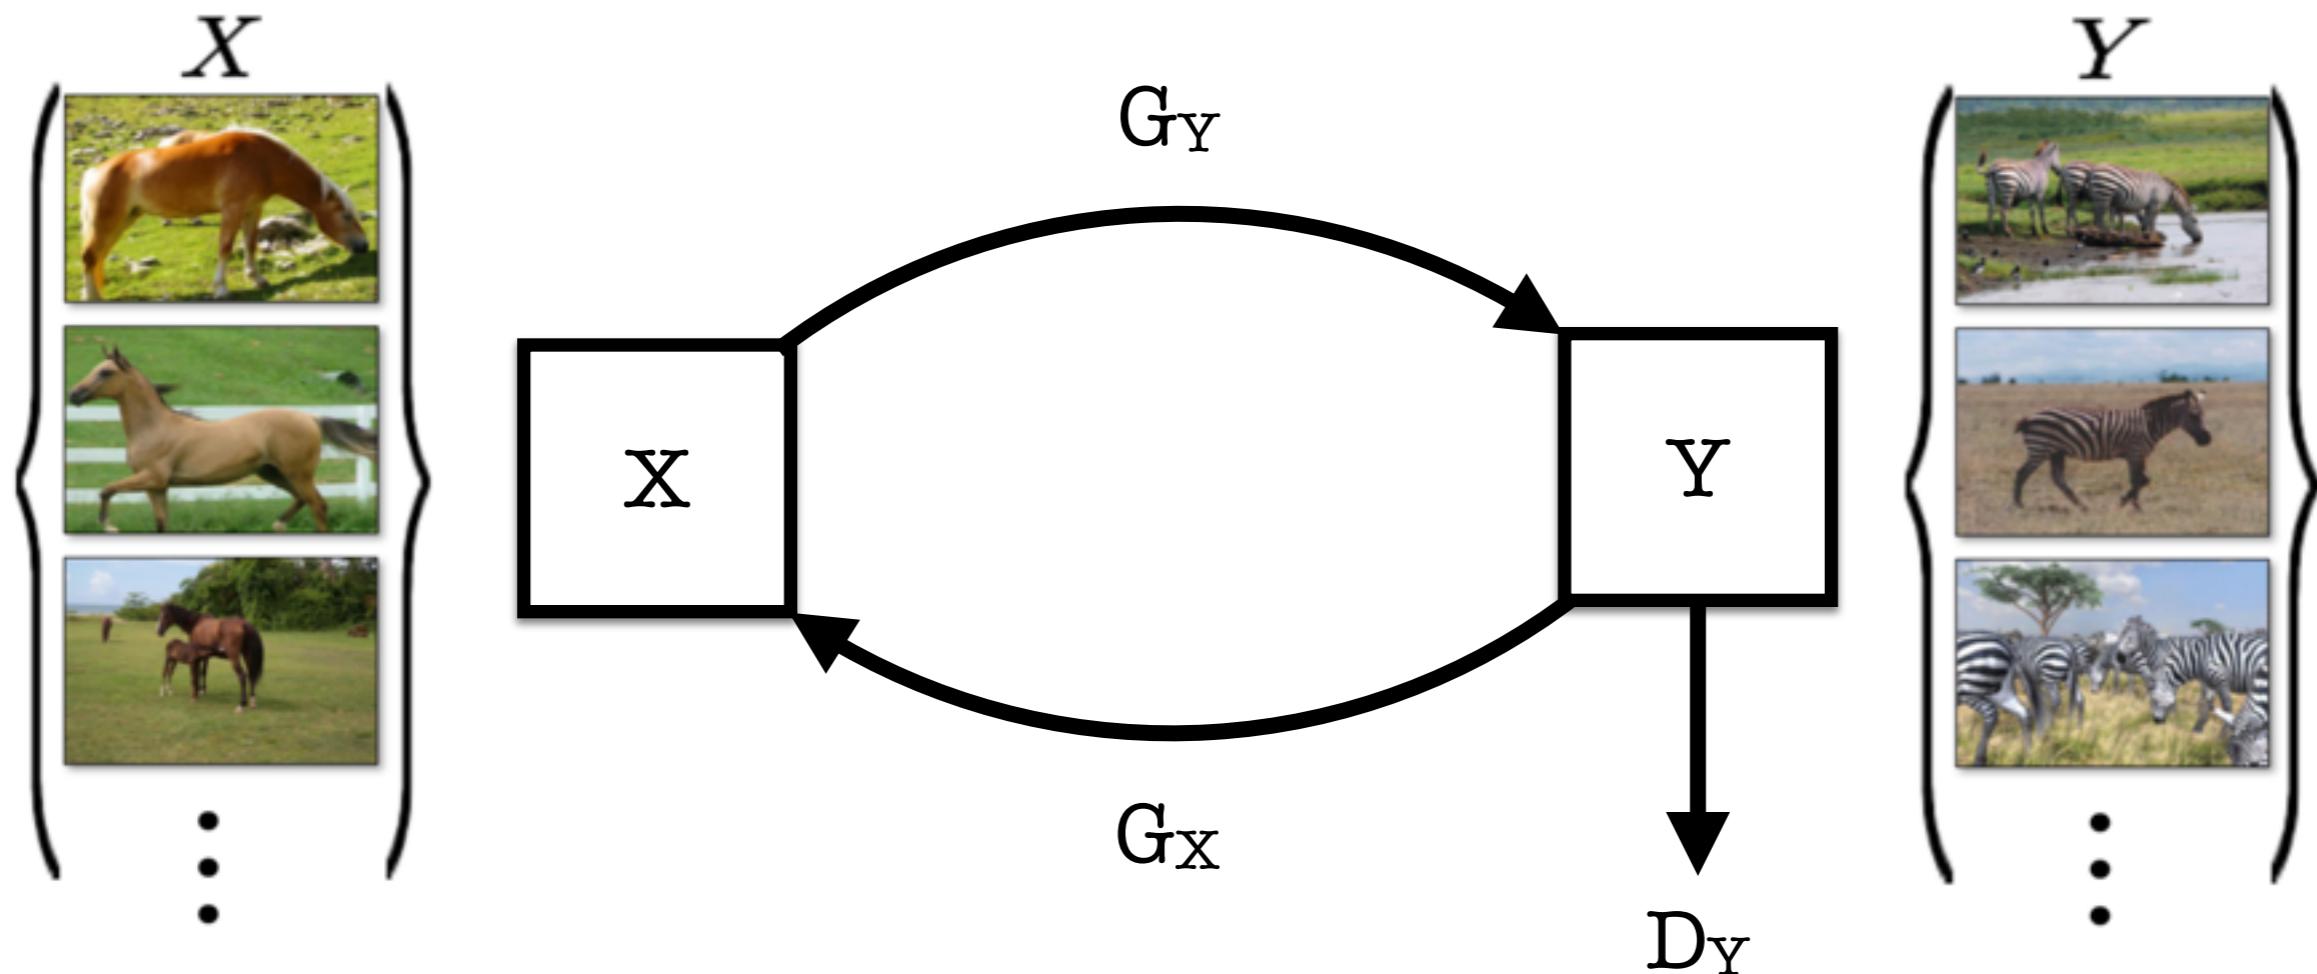
CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

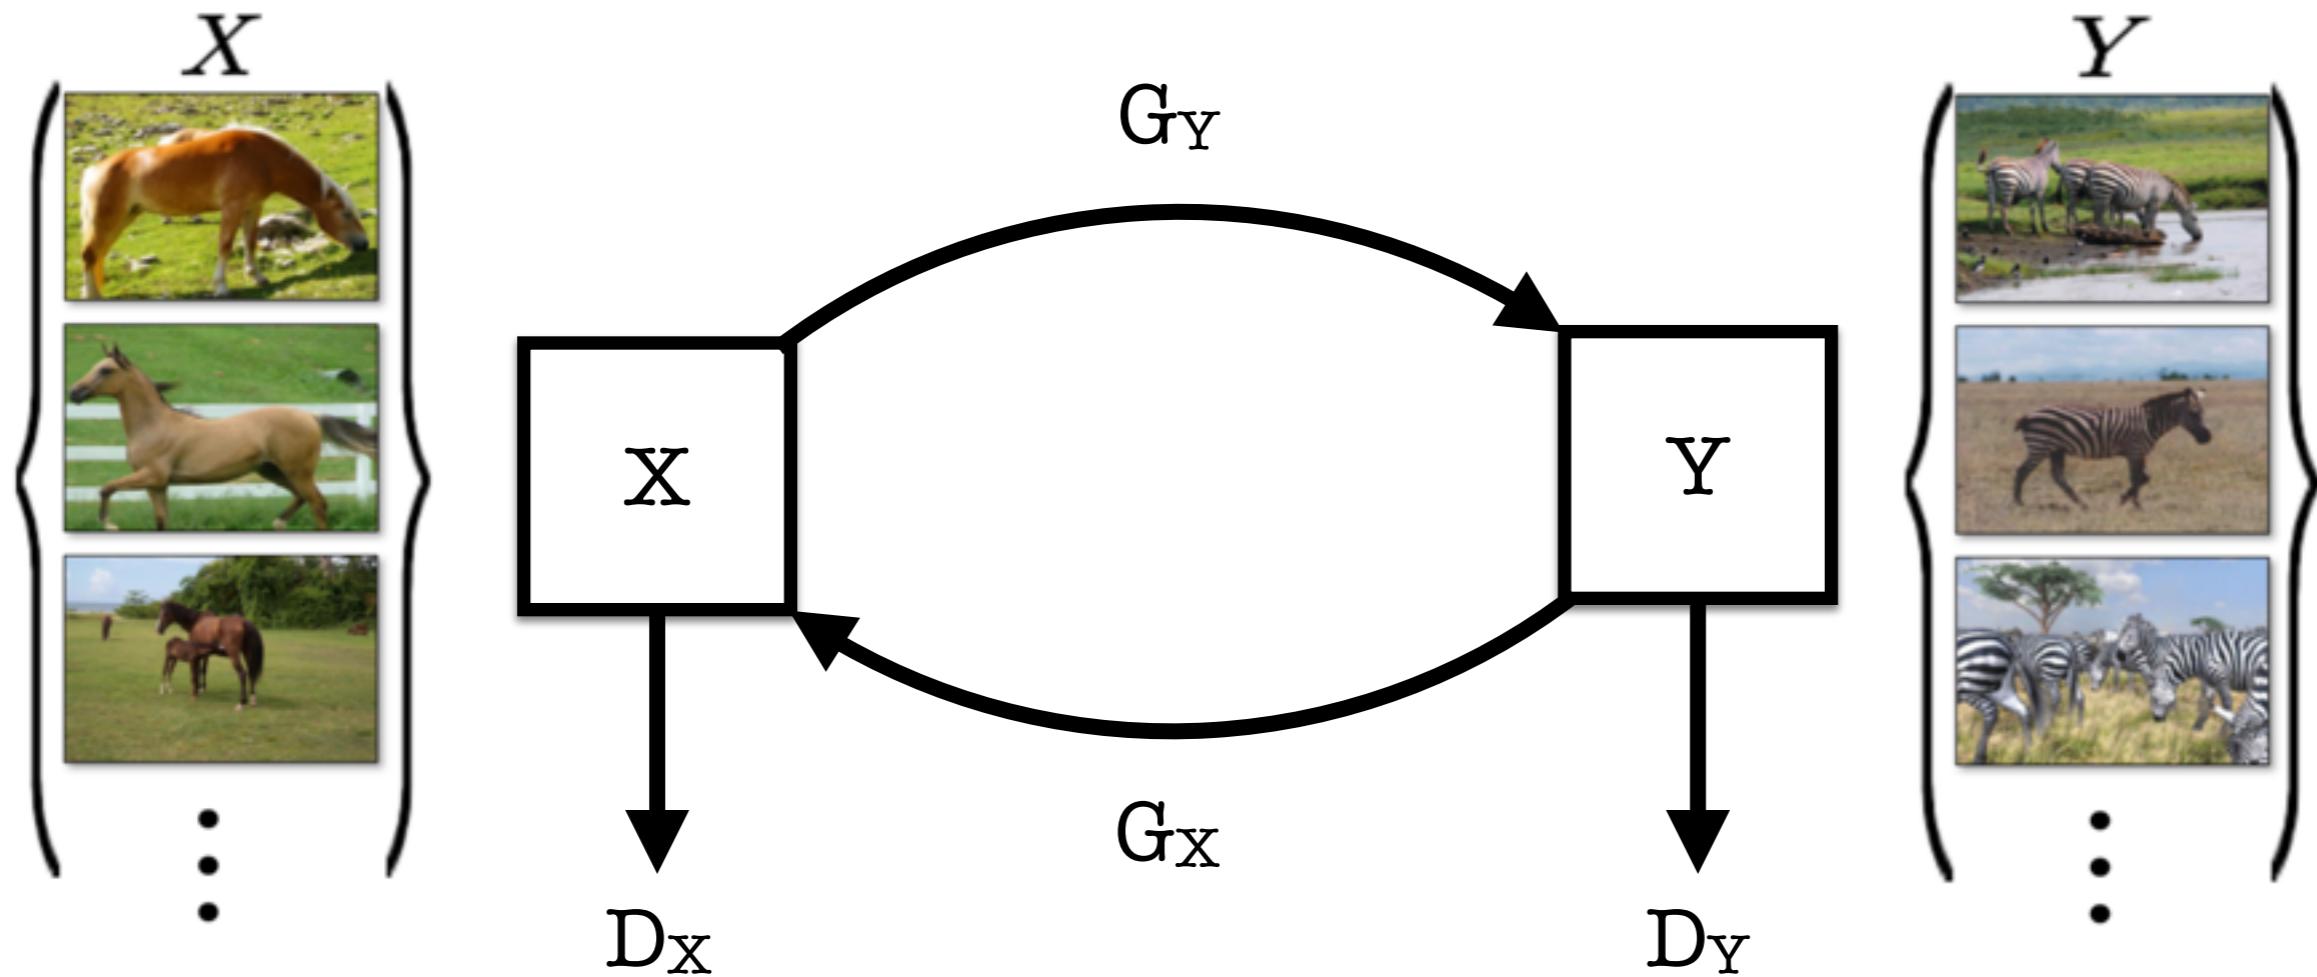
CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

CycleGan: Cycle -Consistent Adversarial Networks

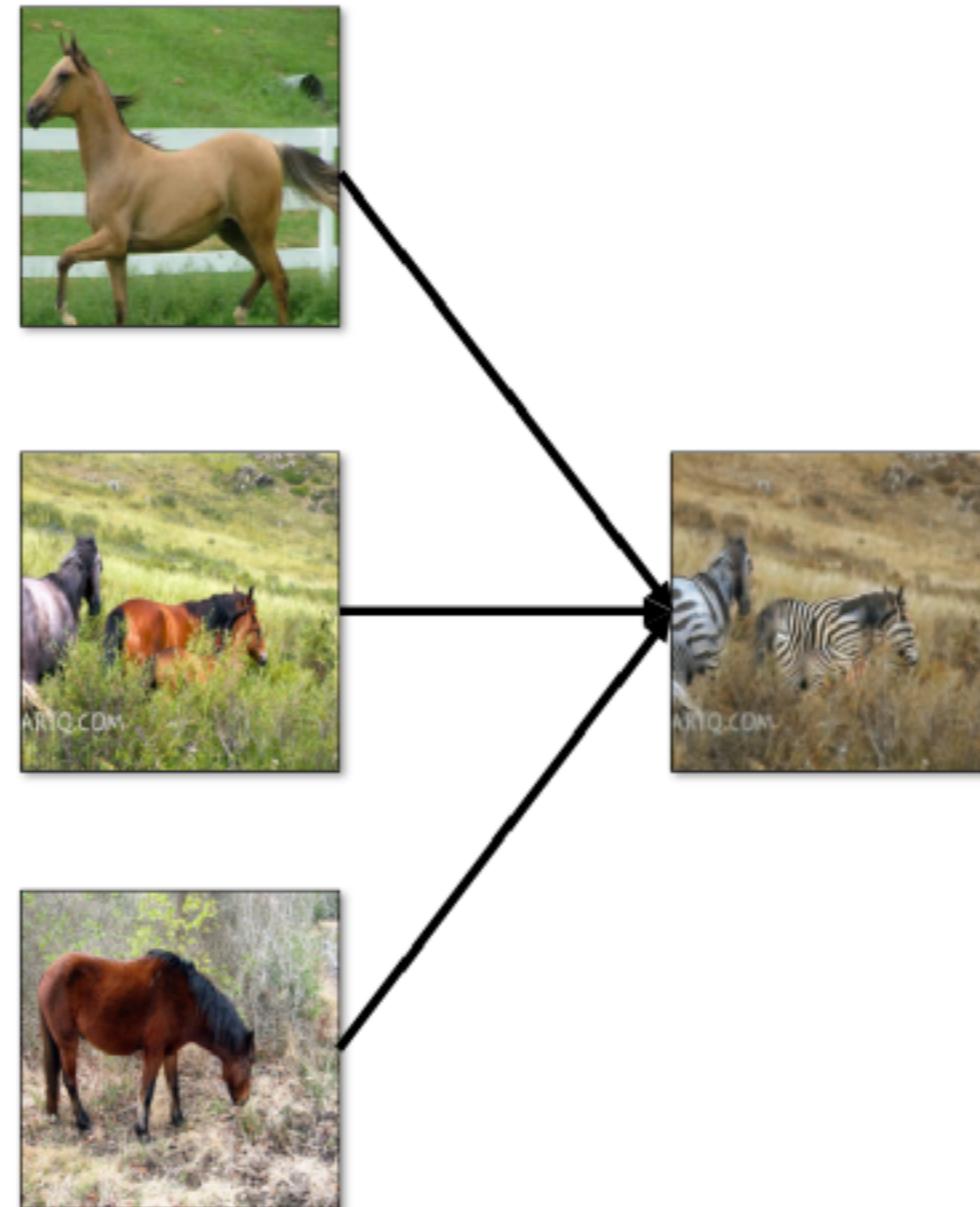


[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

CycleGan: Cycle -Consistent Adversarial Networks

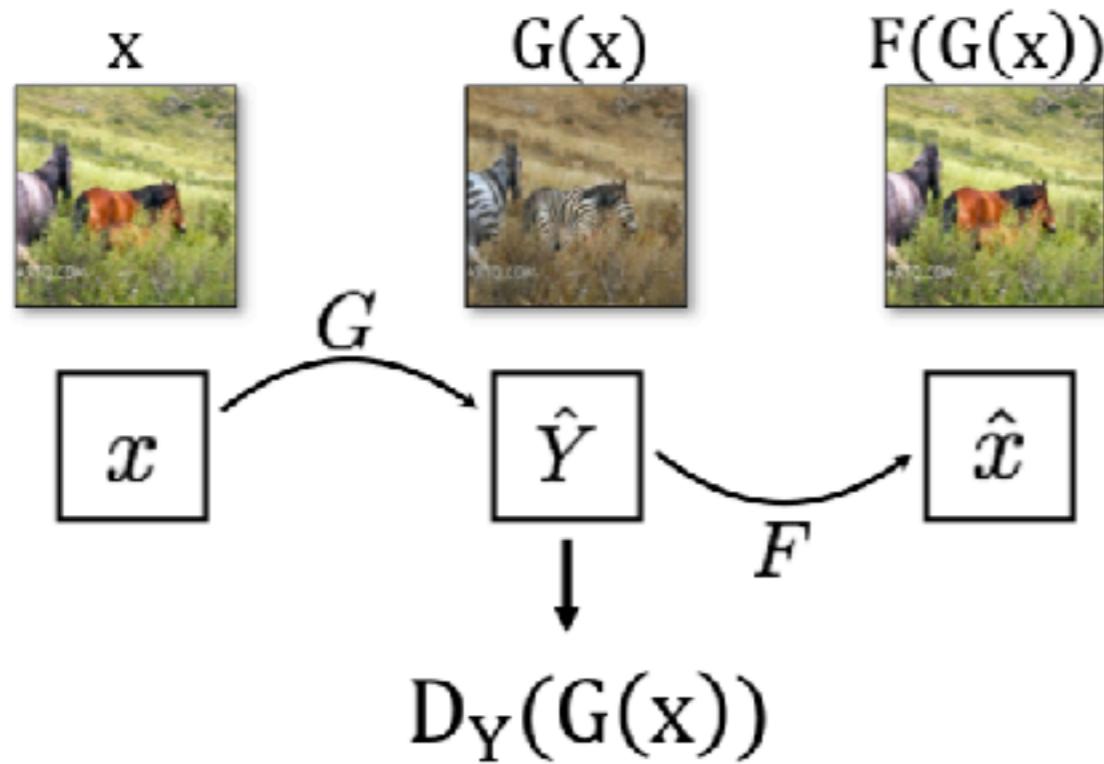
Does not prevent mode collapse



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

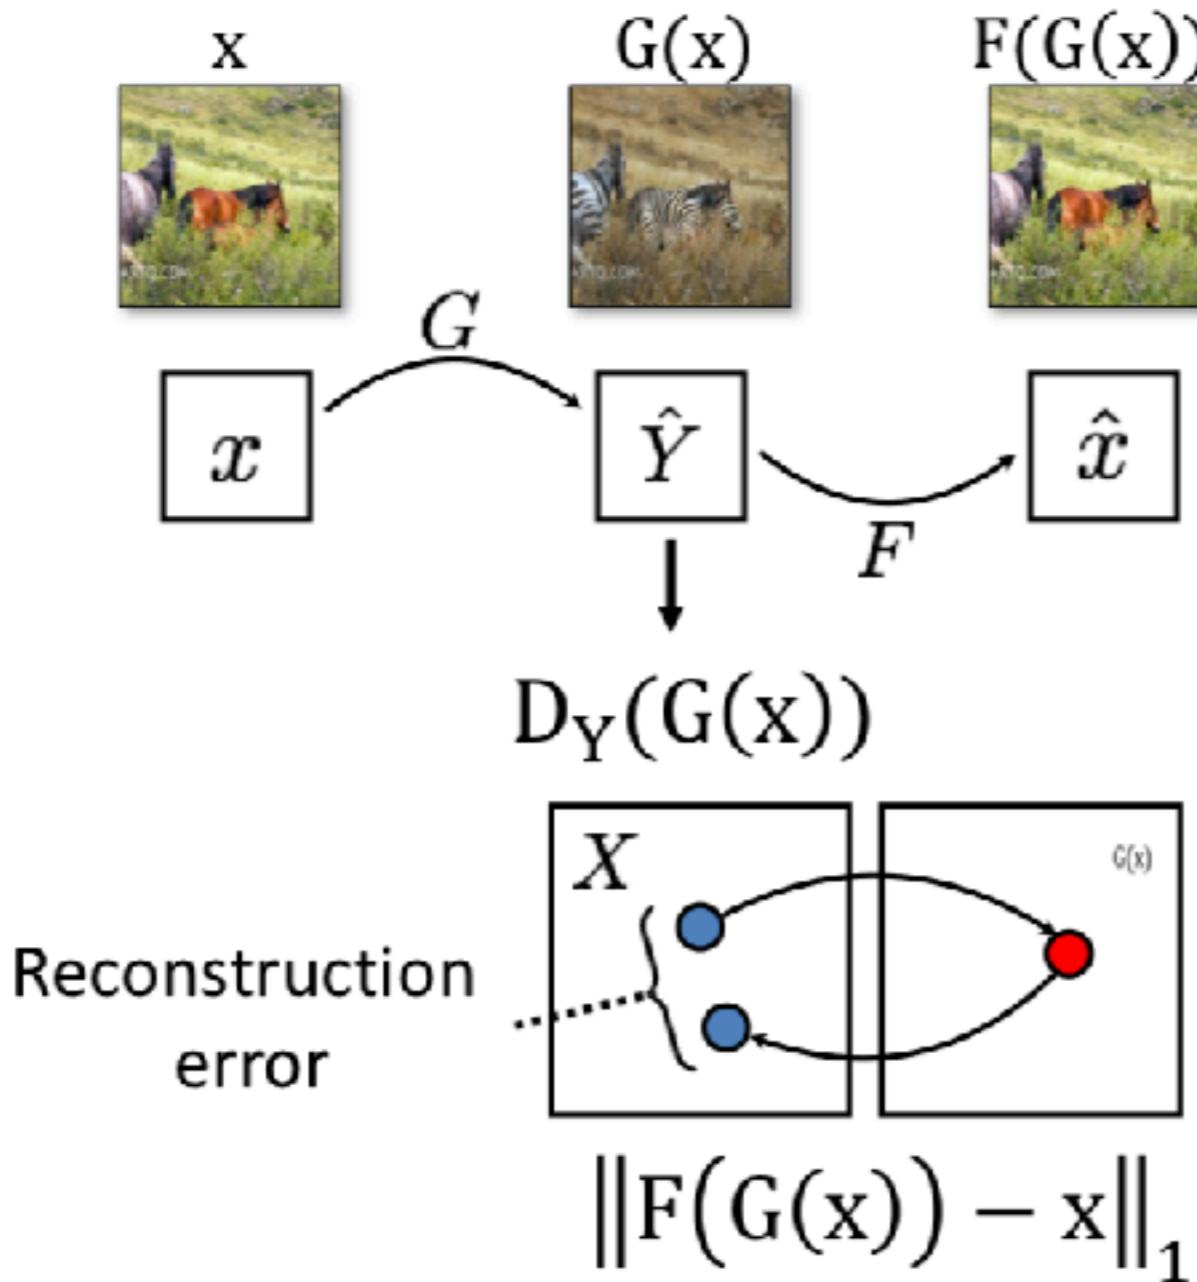
CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

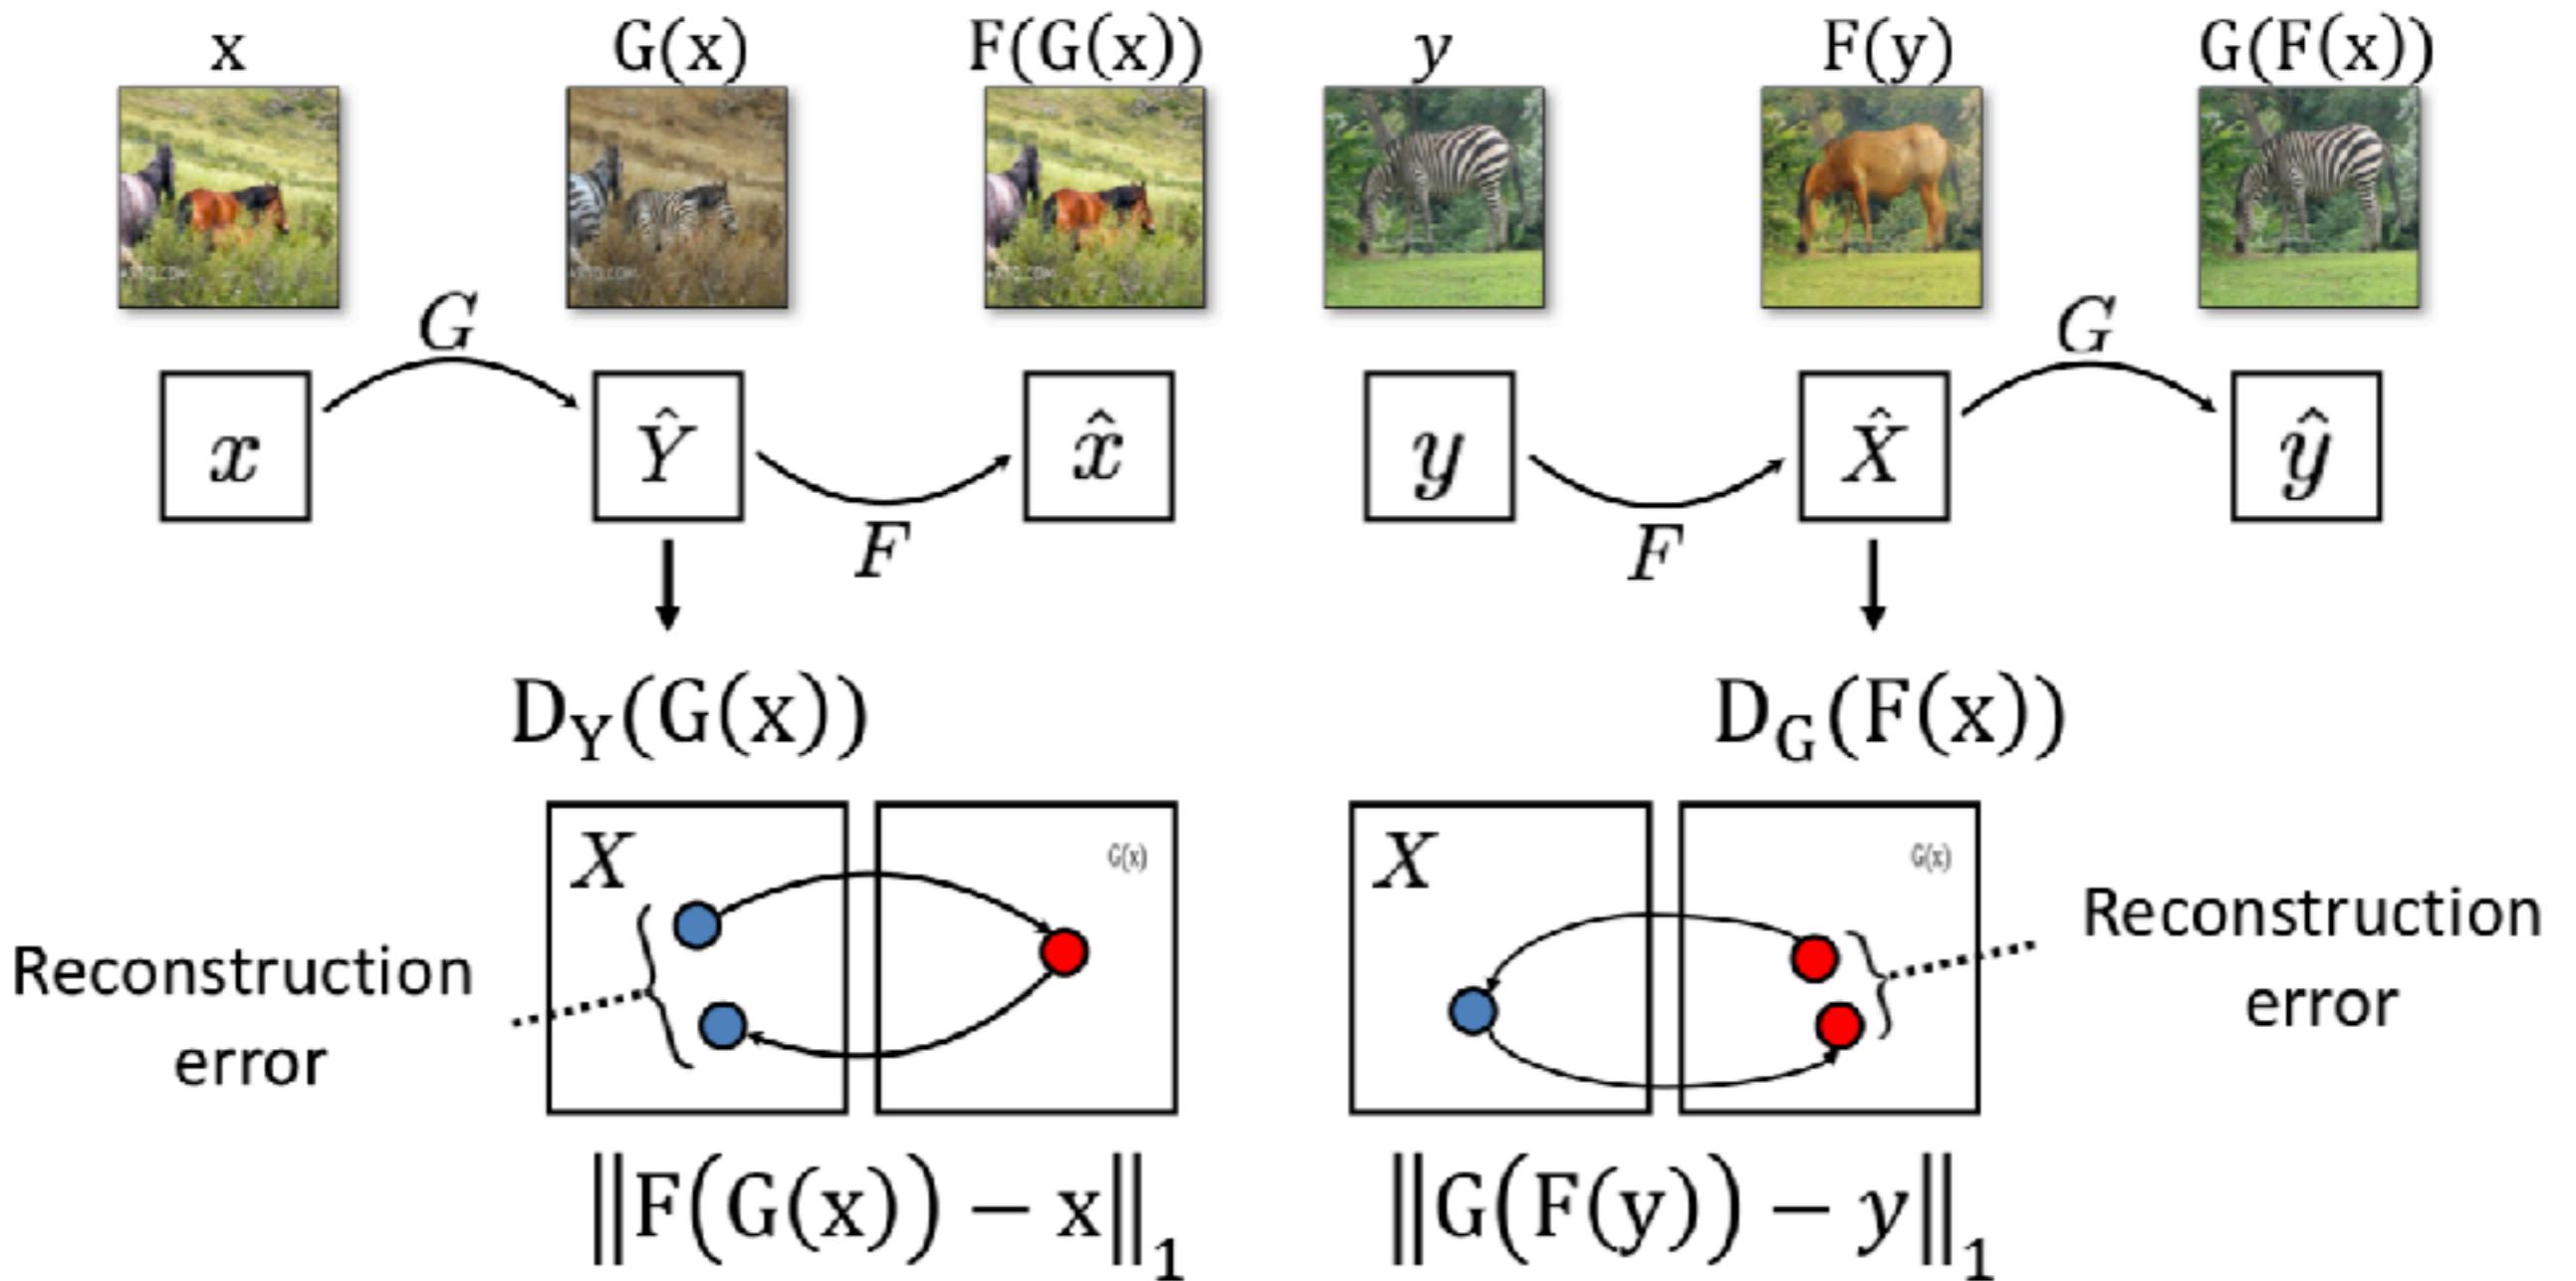
CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

CycleGan: Cycle -Consistent Adversarial Networks



winter Yosemite → summer Yosemite



summer Yosemite → winter Yosemite

CycleGan Application: Rendering to Image

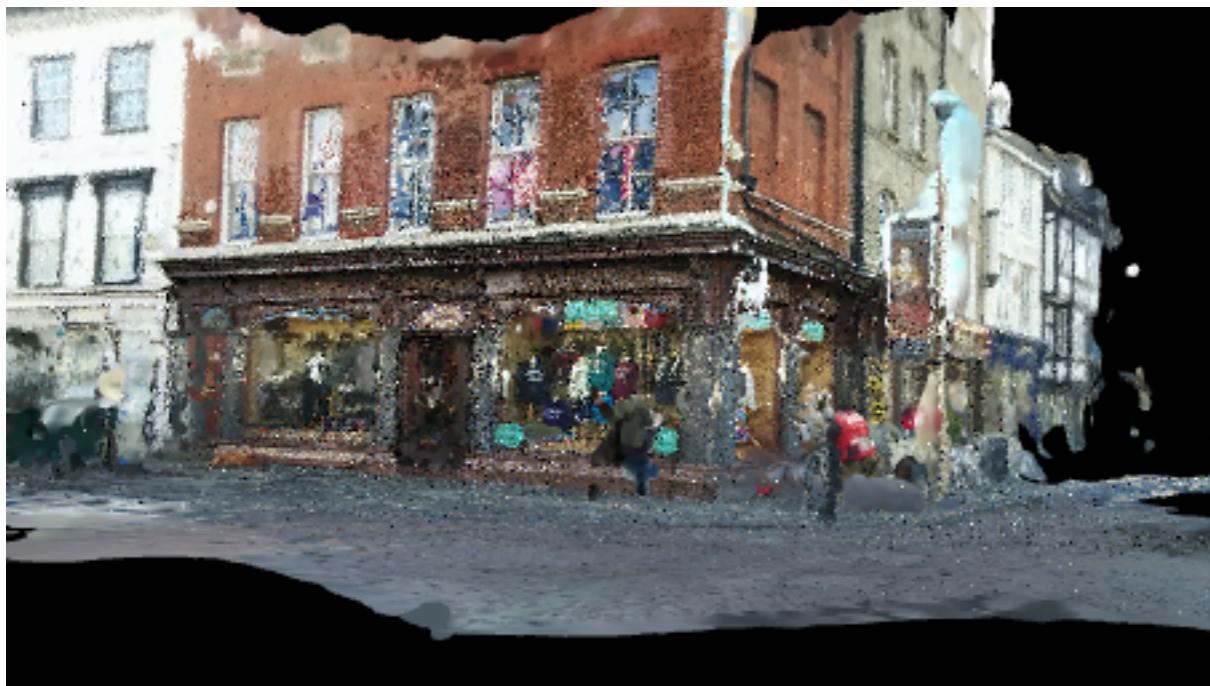


image credit: Markus Müller

CycleGan Application: Image Retrieval



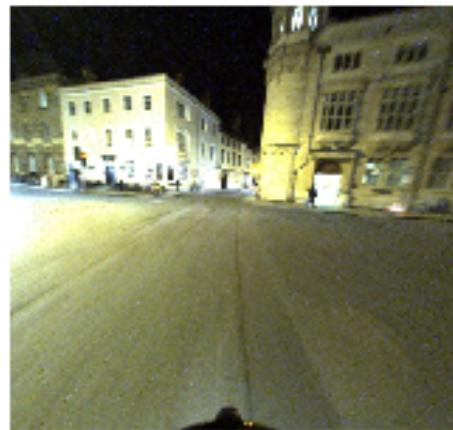
[Maddern, Pascoe, Linegar, Newman, 1 Year, 1000km: The Oxford RobotCar Dataset, IJRR 2017]



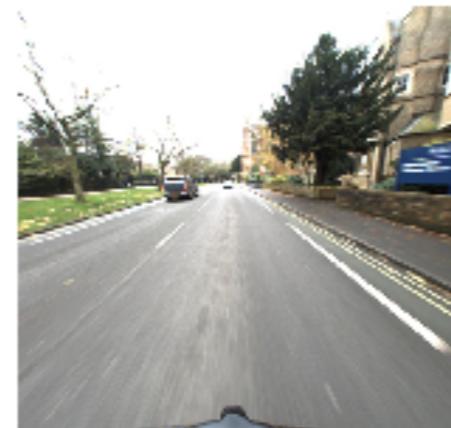
[Sattler et al., Benchmarking 6DOF Outdoor Visual Localization in Changing Conditions, CVPR 2018]

CycleGan Application: Image Retrieval

**Night-Time
Query**



**Retrieved Day
Image**

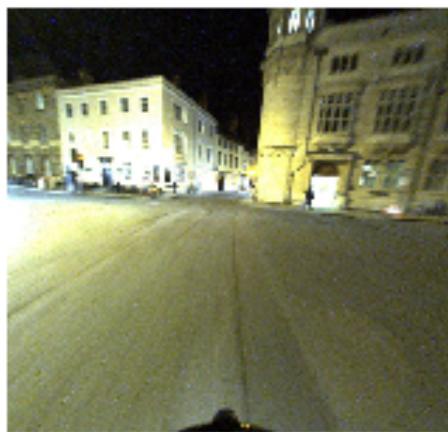


[Anoosheh et al., Night-to-Day Image Translation for Retrieval-based Localization, ICRA 2019]

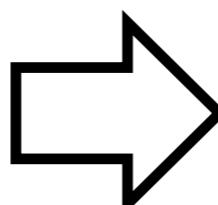
slide credit: Asha Anoosheh

CycleGan Application: Image Retrieval

**Night-Time
Query**



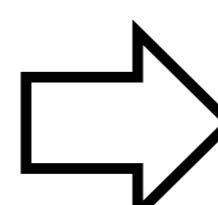
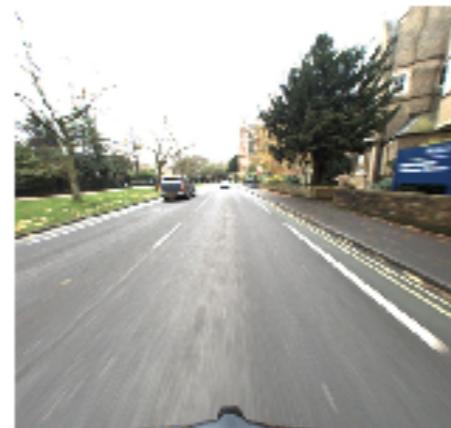
**Retrieved Day
Image**



**Translated
Query**



**Retrieved Day
Image**



[Anoosheh et al., Night-to-Day Image Translation for Retrieval-based Localization, ICRA 2019]

slide credit: Asha Anoosheh

Lessons Learned

- Main lessons from this lecture
 - Generative networks learn mapping from simple probability distribution to target probability distribution
 - Adversarial networks: Learn the loss function
 - Unsupervised image-to-image translation via cycle consistency
- Next lecture: Visual Localization & Feature Learning

Reading Materials

- Start: Goodfellow, Bengio, Courville, Deep Learning, MIT press (Chapter 20.10.4)
- Continue: Blog posts on Wasserstein GANs, Tutorial on VAEs
- Advanced material:
 - ICCV 2017 GAN Tutorial (website with slides and videos)
 - CVPR 2018 GAN Tutorial (website with slides and videos)
 - Tensorflow implementations of various GANs
 - Read the papers listed on the slides
- GANs are an active research topic, be prepared to read a lot

Next Lecture

Jan. 20	Introduction, Linear classifiers and filtering	
Jan. 23	Filtering, gradients, scale	Lab 1
Jan. 27	Local features	
Jan. 30	Learning a classifier	
Feb. 3	Convolutional neural networks	Lab 2
Feb. 6	More convolutional neural networks	
Feb. 10	Robust model fitting and RANSAC	Lab 3
Feb. 13	Image registration	
Feb. 17	Camera Geometry	Lab 4
Feb. 20	More camera geometry	
Feb. 24	Generative neural networks	
Feb. 27	Generative neural networks	
Mar. 2	Visual Localization & Feature Learning	
Mar. 9	No lecture	