

SSY097 - Image Analysis

Lecture 1 - Introduction and Linear Classifiers

*Torsten Sattler
(slides adapted from Olof Enqvist)*

Team



Jennifer Alvén
alven@chalmers.se



Carl Toft
carl.toft@chalmers.se



José Pedro Lopes Iglesias
jose.iglesias@chalmers.se



Torsten Sattler torsat@chalmers.se

Course Information

- Exercises (optional, non-graded)

Course Information

- Exercises (optional, non-graded)
- Four mandatory lab assignments
 - Hand in code and short report
 - Pass or fail
- Deadlines (23:59): February 2nd, February 13th, February 22nd, March 10th

Course Information

- Exercises (optional, non-graded)
- Four mandatory lab assignments
 - Hand in code and short report
 - Pass or fail
- Deadlines (23:59): February 2nd, February 13th, February 22nd, March 10th
- Mandatory mini project (two choices):
 - Hand in code, report, poster
 - Up to 3 bonus points for exams (do not help passing exam)

Course Information

- Exercises (optional, non-graded)
- Four mandatory lab assignments
 - Hand in code and short report
 - Pass or fail
- Deadlines (23:59): February 2nd, February 13th, February 22nd, March 10th
- Mandatory mini project (two choices):
 - Hand in code, report, poster
 - Up to 3 bonus points for exams (do not help passing exam)
- Recommend doing the labs and project in pairs

Course Information

- Exercises (optional, non-graded)
- Four mandatory lab assignments
 - Hand in code and short report
 - Pass or fail
- Deadlines (23:59): February 2nd, February 13th, February 22nd, March 10th
- Mandatory mini project (two choices):
 - Hand in code, report, poster
 - Up to 3 bonus points for exams (do not help passing exam)
- Recommend doing the labs and project in pairs
- Graded exam on March 23rd (pen and paper)

Check PingPong Course Page Regularly

SSY097 Bildanalys V19

Event

What's new

Overview

Content

Course information

- Schedule
- Projects
- Course Representatives
- Lab 1
- Lab 2
- Lab 3
- Lab 4
- Mini project

Documents

FAQ

Communication

Members

Message Board

Discussion

Project groups

Toolbox

Preferences

Answer questions

Project groups

Send message

Log book

Previous Next Search Maximize the content Save as PDF

Edit Print

Course information

Essential information

- There are four mandatory lab assignments.
After completing the lab you submit a mini report together with your code. Deadlines below.
- There is a mandatory mini project. You choose to do either Project A or Project B.
- The exam takes place on March 23 (time: TBA). Allowed materials: Pencil eraser.
You can get 3 bonus credits on the exam by completing extra tasks connected to the mini project.
These bonus credits are only valid for grades 4 and 5. They do not help in passing the exam.
The bonus credits are valid at all three exams in 2019 (March 23, June 10, Aug 19).
- We recommend that you do the labs and projects in pairs, but in special cases groups of 1-3 persons are ok.

Further information, such as aims and learning outcomes can be found on the course page at [student.portal.chalmers.se](#)

Staff

- Torsten Sattler, lecturer, torsat@chalmers.se
- Jennifer Alvén, teaching assistant, alven@chalmers.se
- Carl Toft, teaching assistant, carl.toft@chalmers.se
- José Pedro Lopes Iglesias, teaching assistant, jose.iglesias@chalmers.se
- Fredrik Karl, examiner, fredrik.karl@chalmers.se

Course material

The necessary course material will be provided during the course. This includes lab pm's and exercises.
If you like to read more about image analysis, you can use Szeliskis book.

- Richard Szeliski, Computer Vision: Algorithms and Applications, available at Cremona or as a [free pdf](#).

Deadlines

You need to submit your code and reports before the following dates. If there are minor errors, you will be given a chance to correct them later.
If you need more than one iteration to correct your mistakes, we do not promise to mark them before the re-exam.

- Lab 1: Sunday, February 2, 23:59
- Lab 2: Wednesday, February 13, 23:59
- Lab 3: Sunday, February 22, 23:59

Ask Questions

- Feel free to ask questions any time
- Please give us feedback (positive and negative) about the lecture, the labs, ...
 - Talk to me before or after the lecture
 - Write me an email
 - Come to my office
- Talk to the course representatives (see PingPong page)
- Happy to adapt based on your suggestions

Today

- What is Image Analysis
- Modeling Images
- Filtering and Linear Classifiers

What is Image Analysis?

Image Classification



Is there a robin in this image?

Object Detection



Where is the robin?

Object Detection



[YOLOv3](#)

Object Detection



[YOLOv3](#)

Image Classification

WORTHWILE GOALS FOR ROBOTS THAT CAN WORK WITH PEOPLE

- Object recognition capabilities of a two year old
- Language understanding capabilities of a four year old
- Manual dexterity of a six year old
- Social understanding of an eight year old
- The strength and tenderness of an 18 year old, holding an adult

ICRA 2018 plenary talk by Rodney Brooks

Image Segmentation

Full-Resolution Residual Networks for Semantic Segmentation in Street Scenes

Tobias Pohlen, Alexander Hermans,
Markus Mathias, Bastian Leibe

Visual Computing Institute, Computer Vision Group
RWTH Aachen University



[Full-Resolution Residual Networks \(FRRNs\) for
Semantic Image Segmentation in Street Scenes](#)

Image Segmentation

Full-Resolution Residual Networks for Semantic Segmentation in Street Scenes

Tobias Pohlen, Alexander Hermans,
Markus Mathias, Bastian Leibe

Visual Computing Institute, Computer Vision Group
RWTH Aachen University

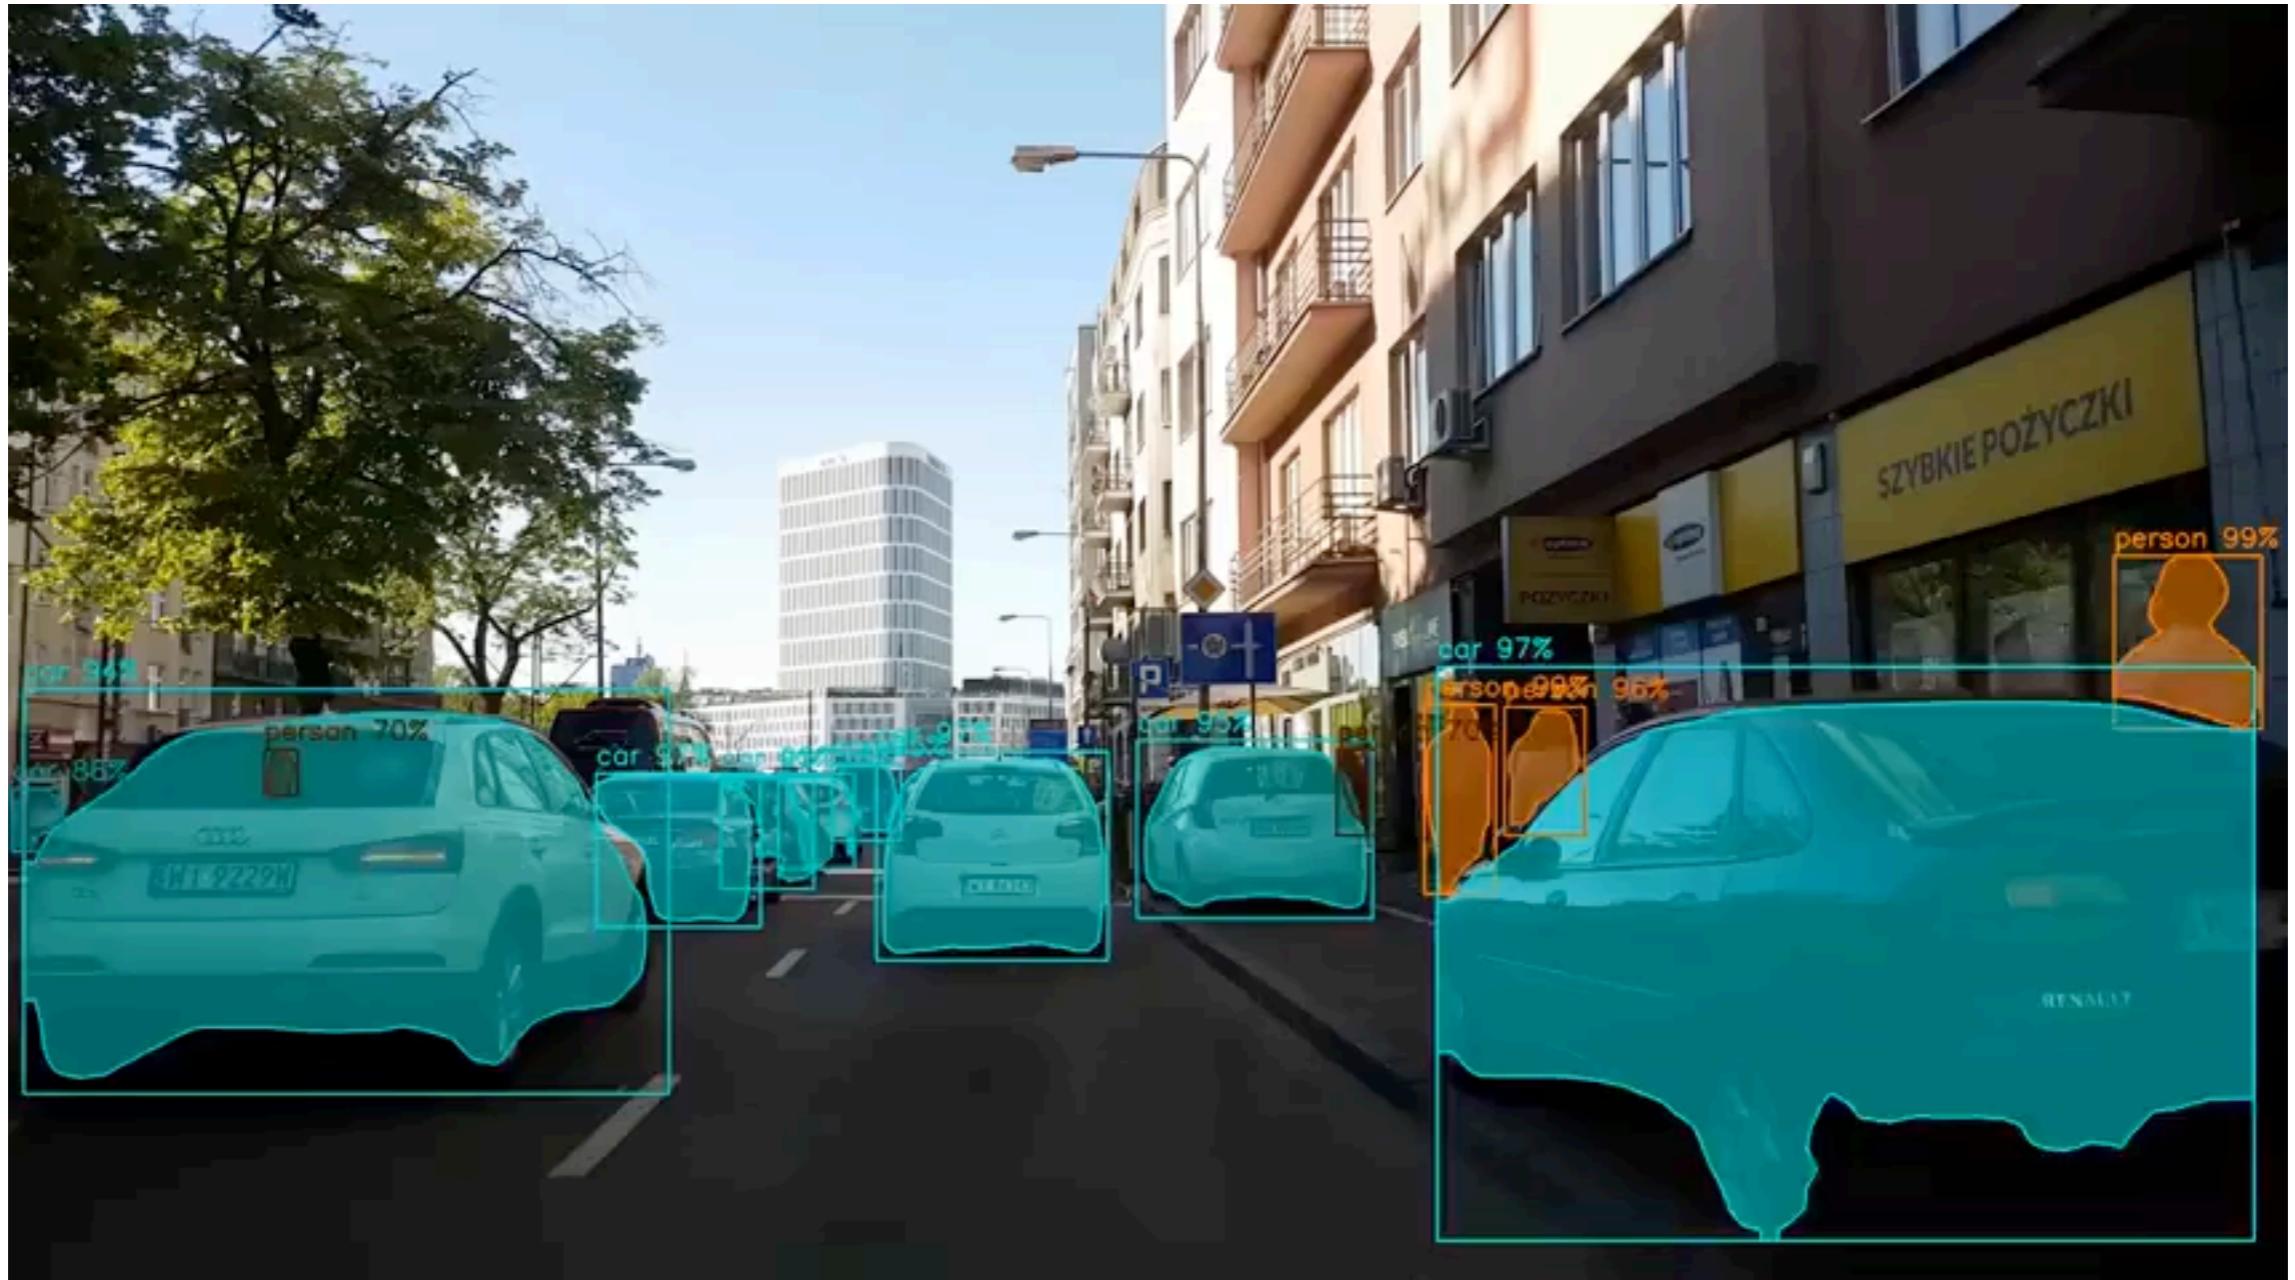


Visual Computing Institute
Computer Vision
Prof. Dr. Bastian Leibe

RWTHAACHEN
UNIVERSITY

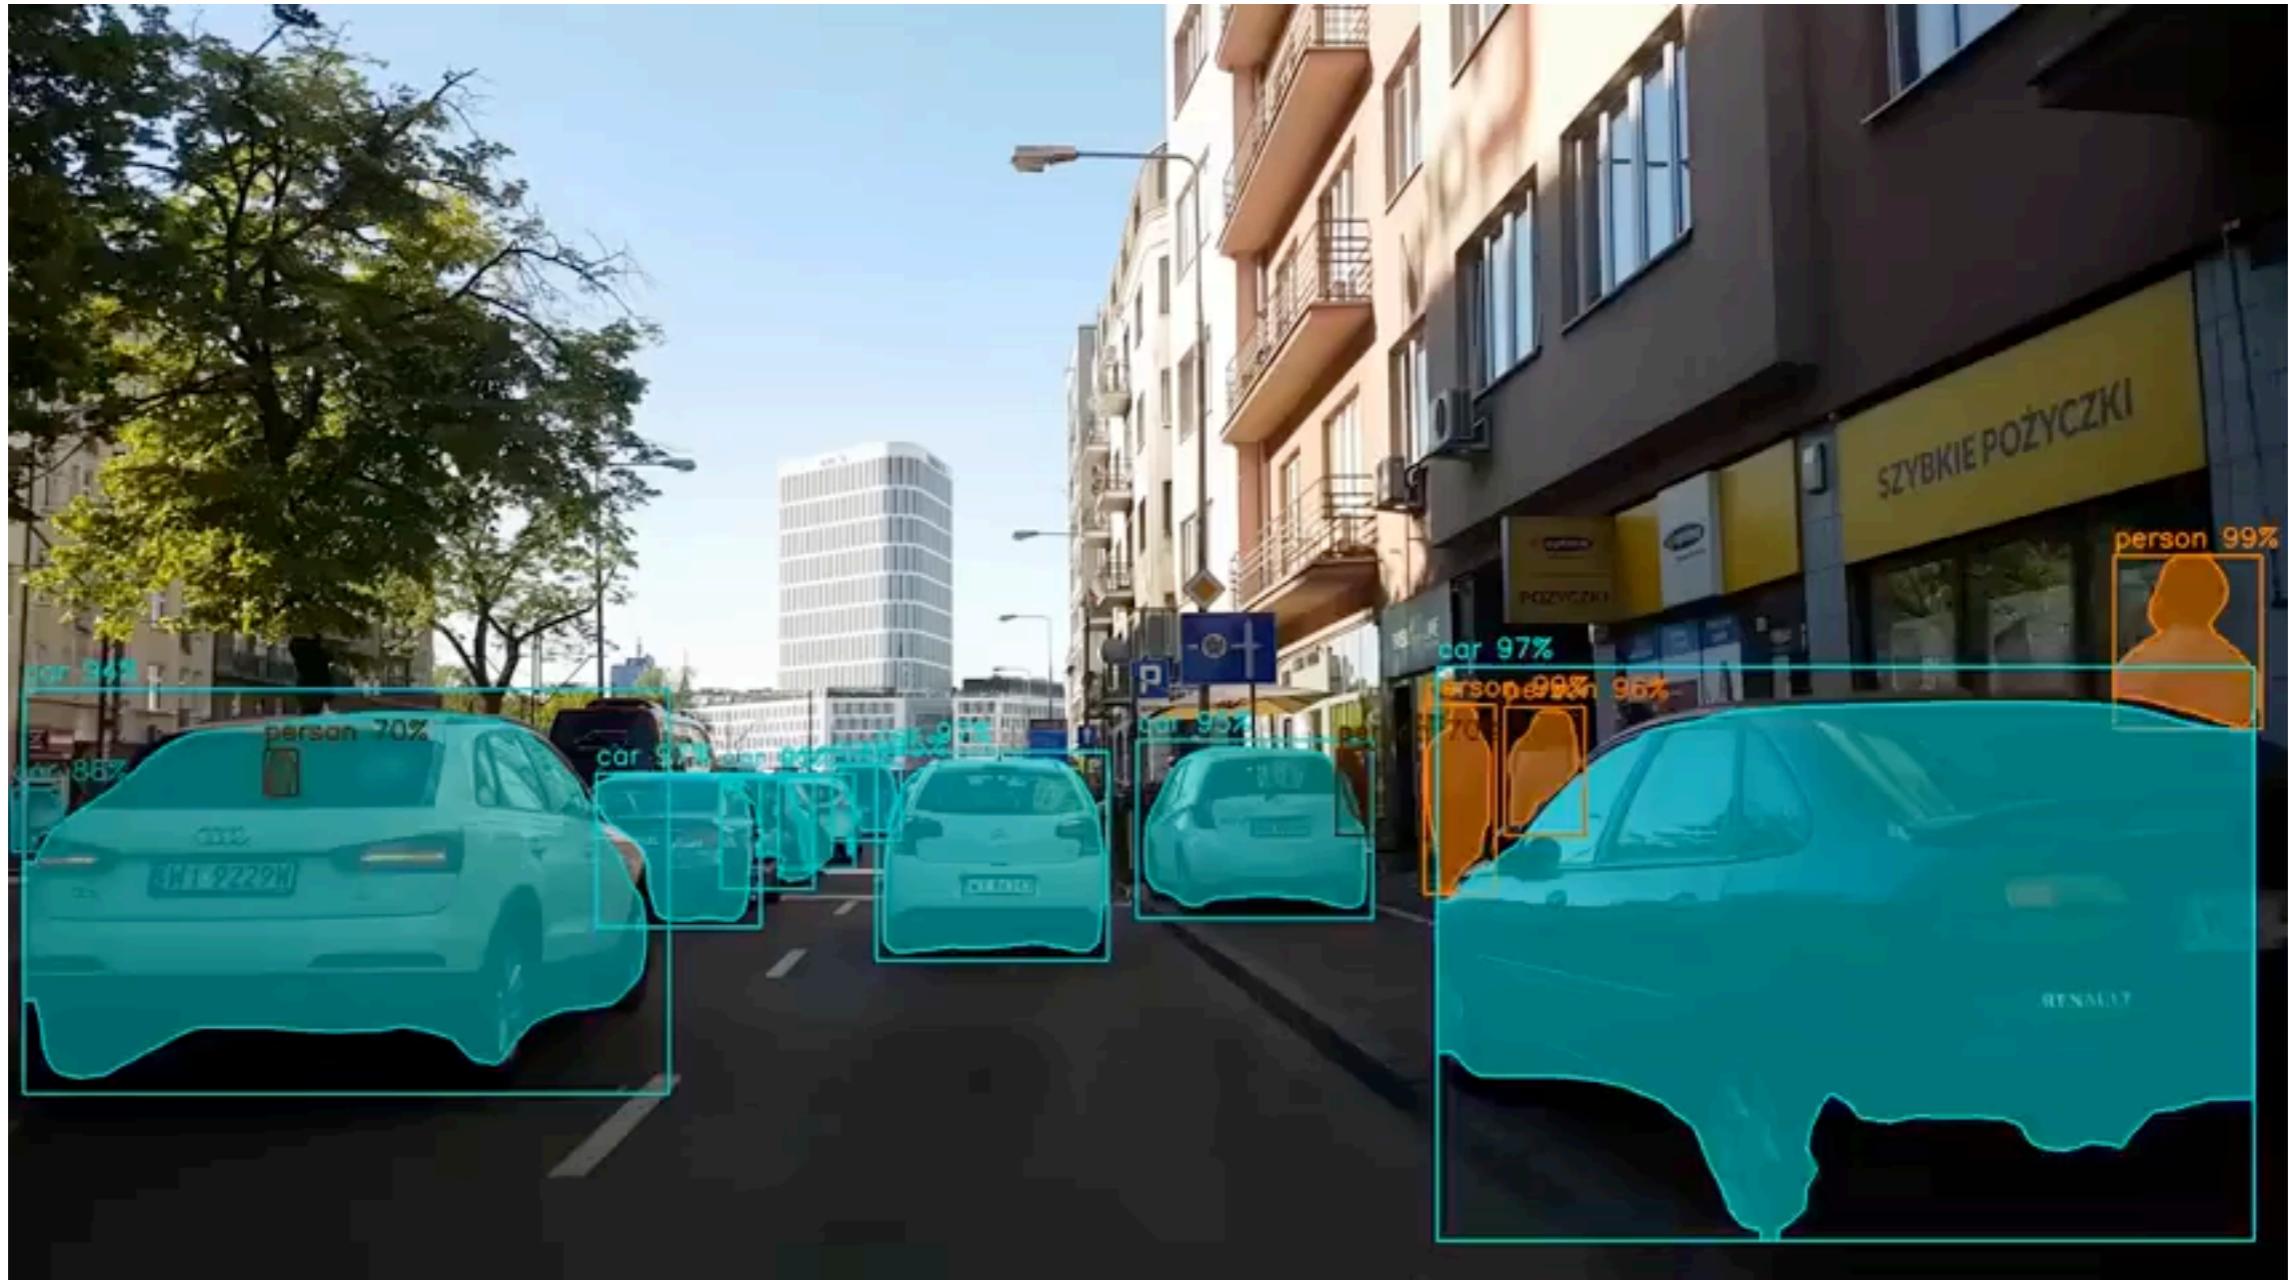
[Full-Resolution Residual Networks \(FRRNs\) for
Semantic Image Segmentation in Street Scenes](#)

Instance-Level Segmentation



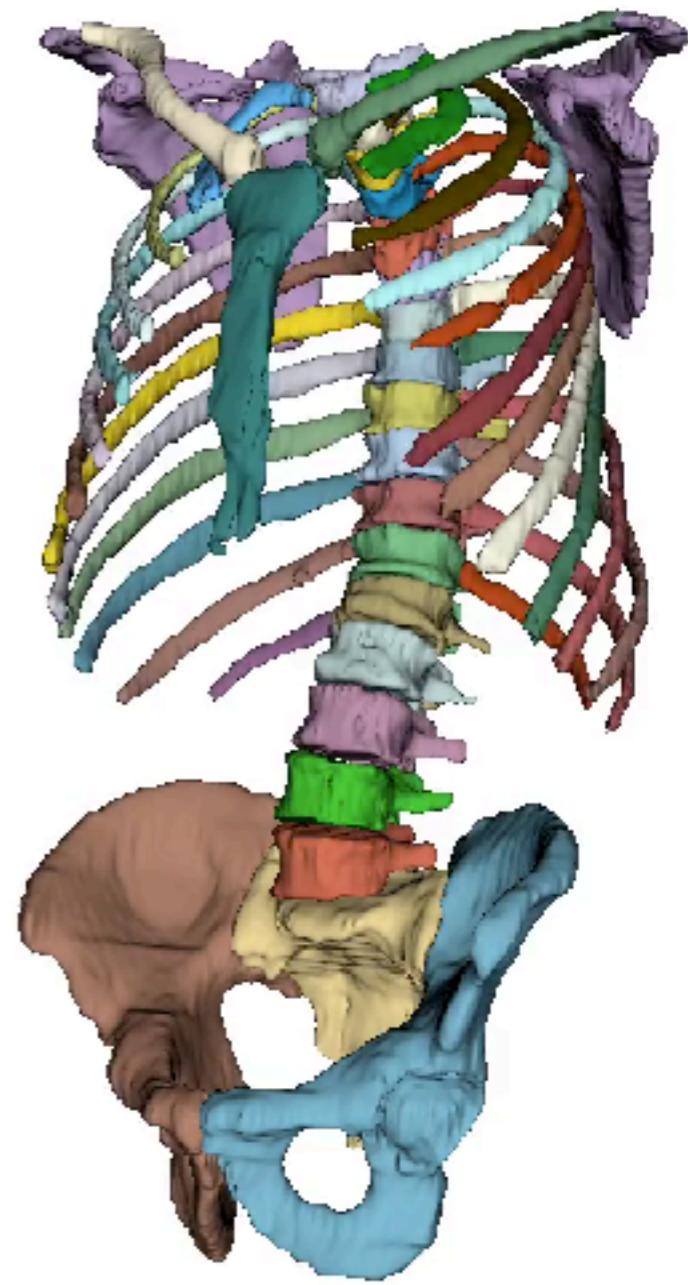
[Masked R-CNN](#)

Instance-Level Segmentation

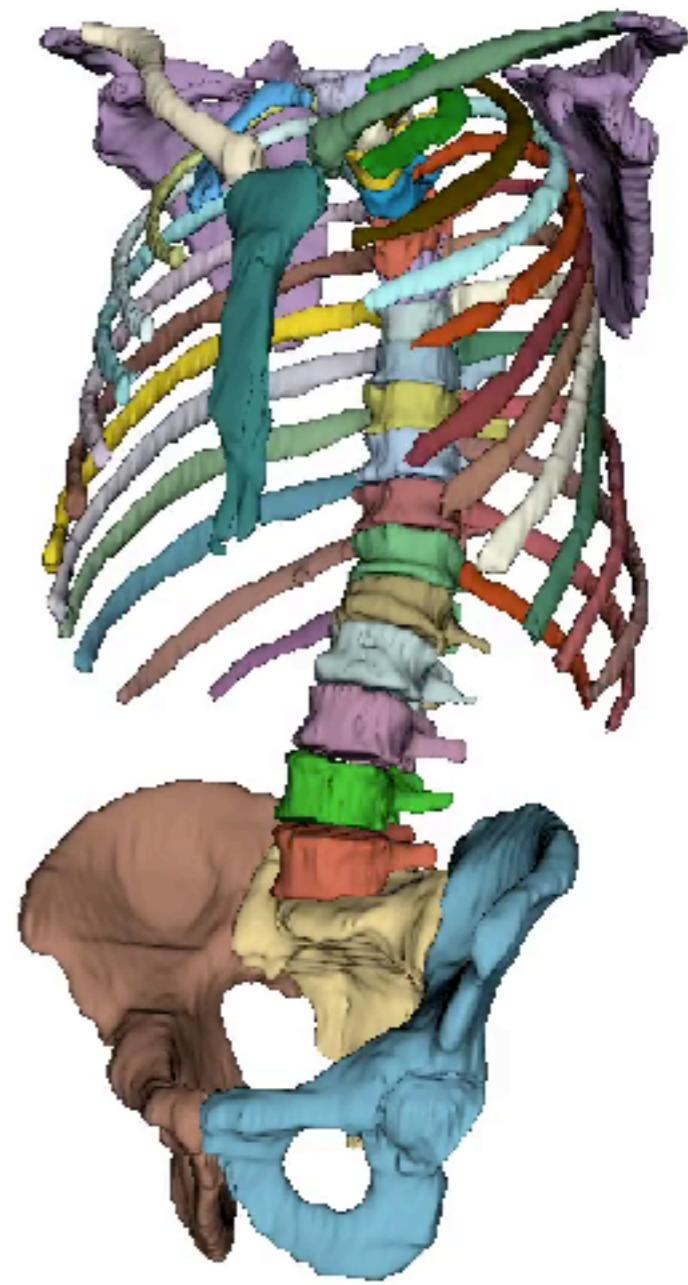


[Masked R-CNN](#)

3D Segmentation



3D Segmentation



3D Reconstruction

Structure-from-Motion Revisited

Johannes L. Schönberger, Jan-Michael Frahm

CVPR 2016

Code available at:
<https://github.com/colmap/colmap>

[Structure-from-Motion Revisited](#)

3D Reconstruction

Structure-from-Motion Revisited

Johannes L. Schönberger, Jan-Michael Frahm

CVPR 2016

Code available at:
<https://github.com/colmap/colmap>

[Structure-from-Motion Revisited](#)

3D Reconstruction

Pixelwise View Selection for Unstructured Multi-View Stereo

Johannes L. Schönberger, Enliang Zheng,
Marc Pollefeys, Jan-Michael Frahm

ECCV 2016

[Pixelwise View Selection for
Unstructured Multi-View Stereo](#)

3D Reconstruction

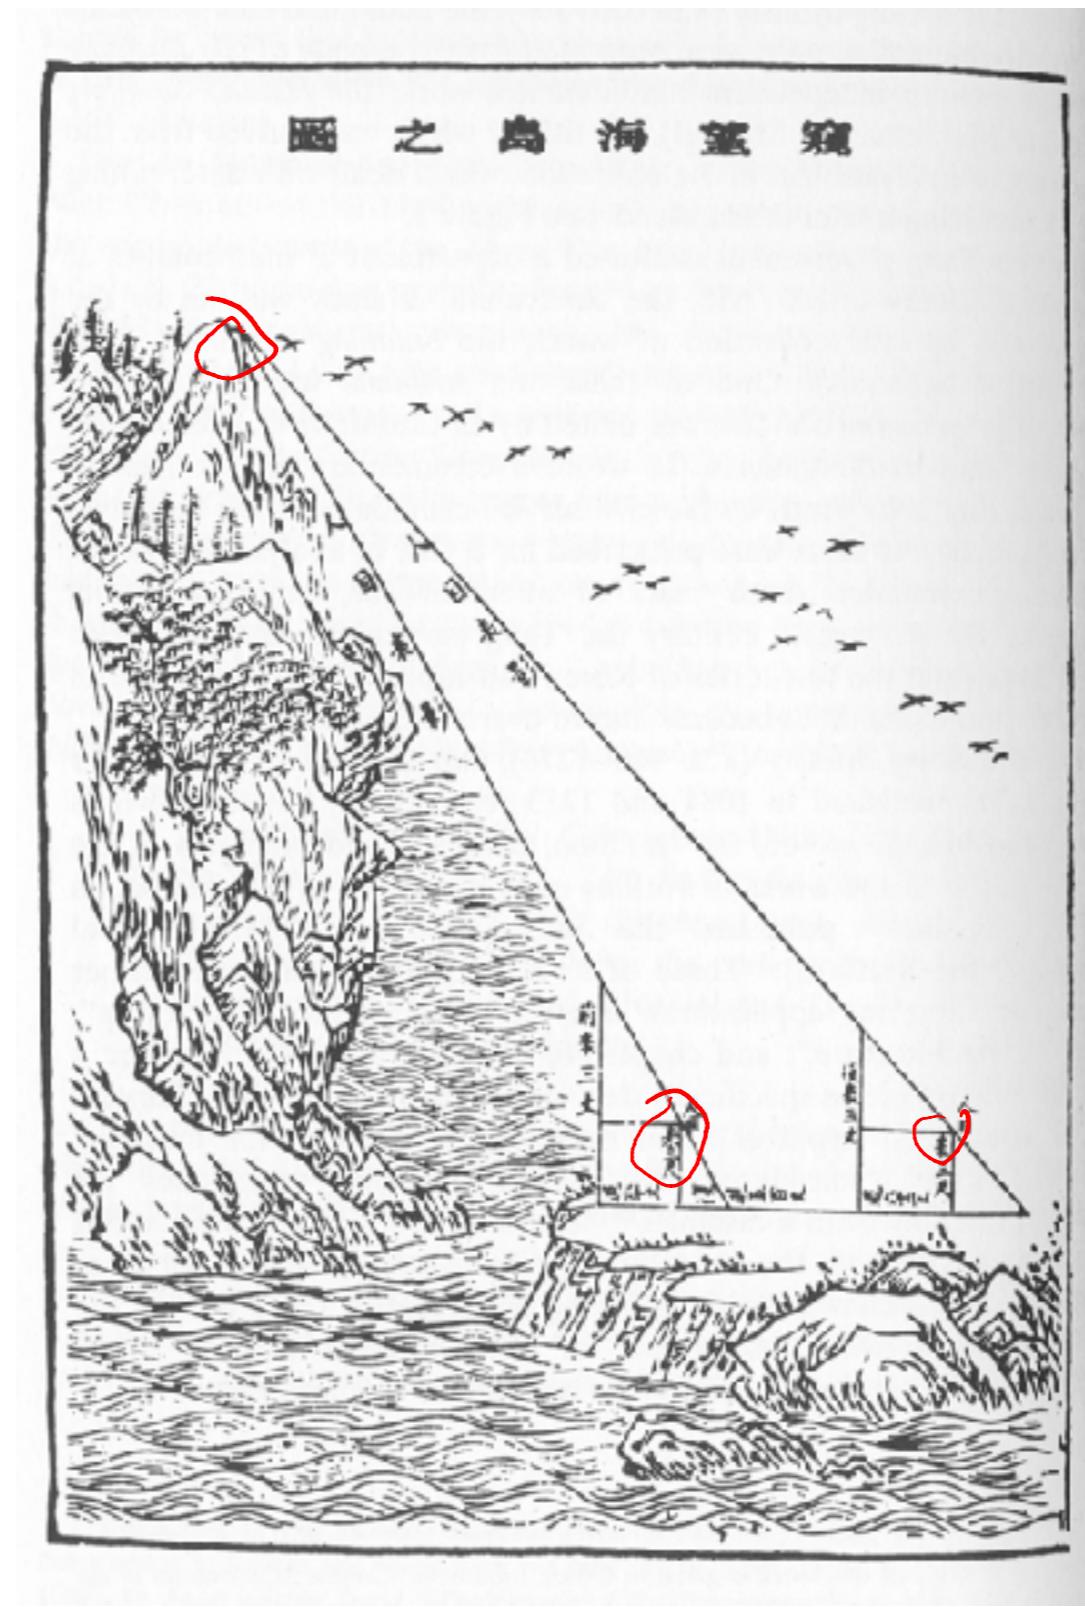
Pixelwise View Selection for Unstructured Multi-View Stereo

Johannes L. Schönberger, Enliang Zheng,
Marc Pollefeys, Jan-Michael Frahm

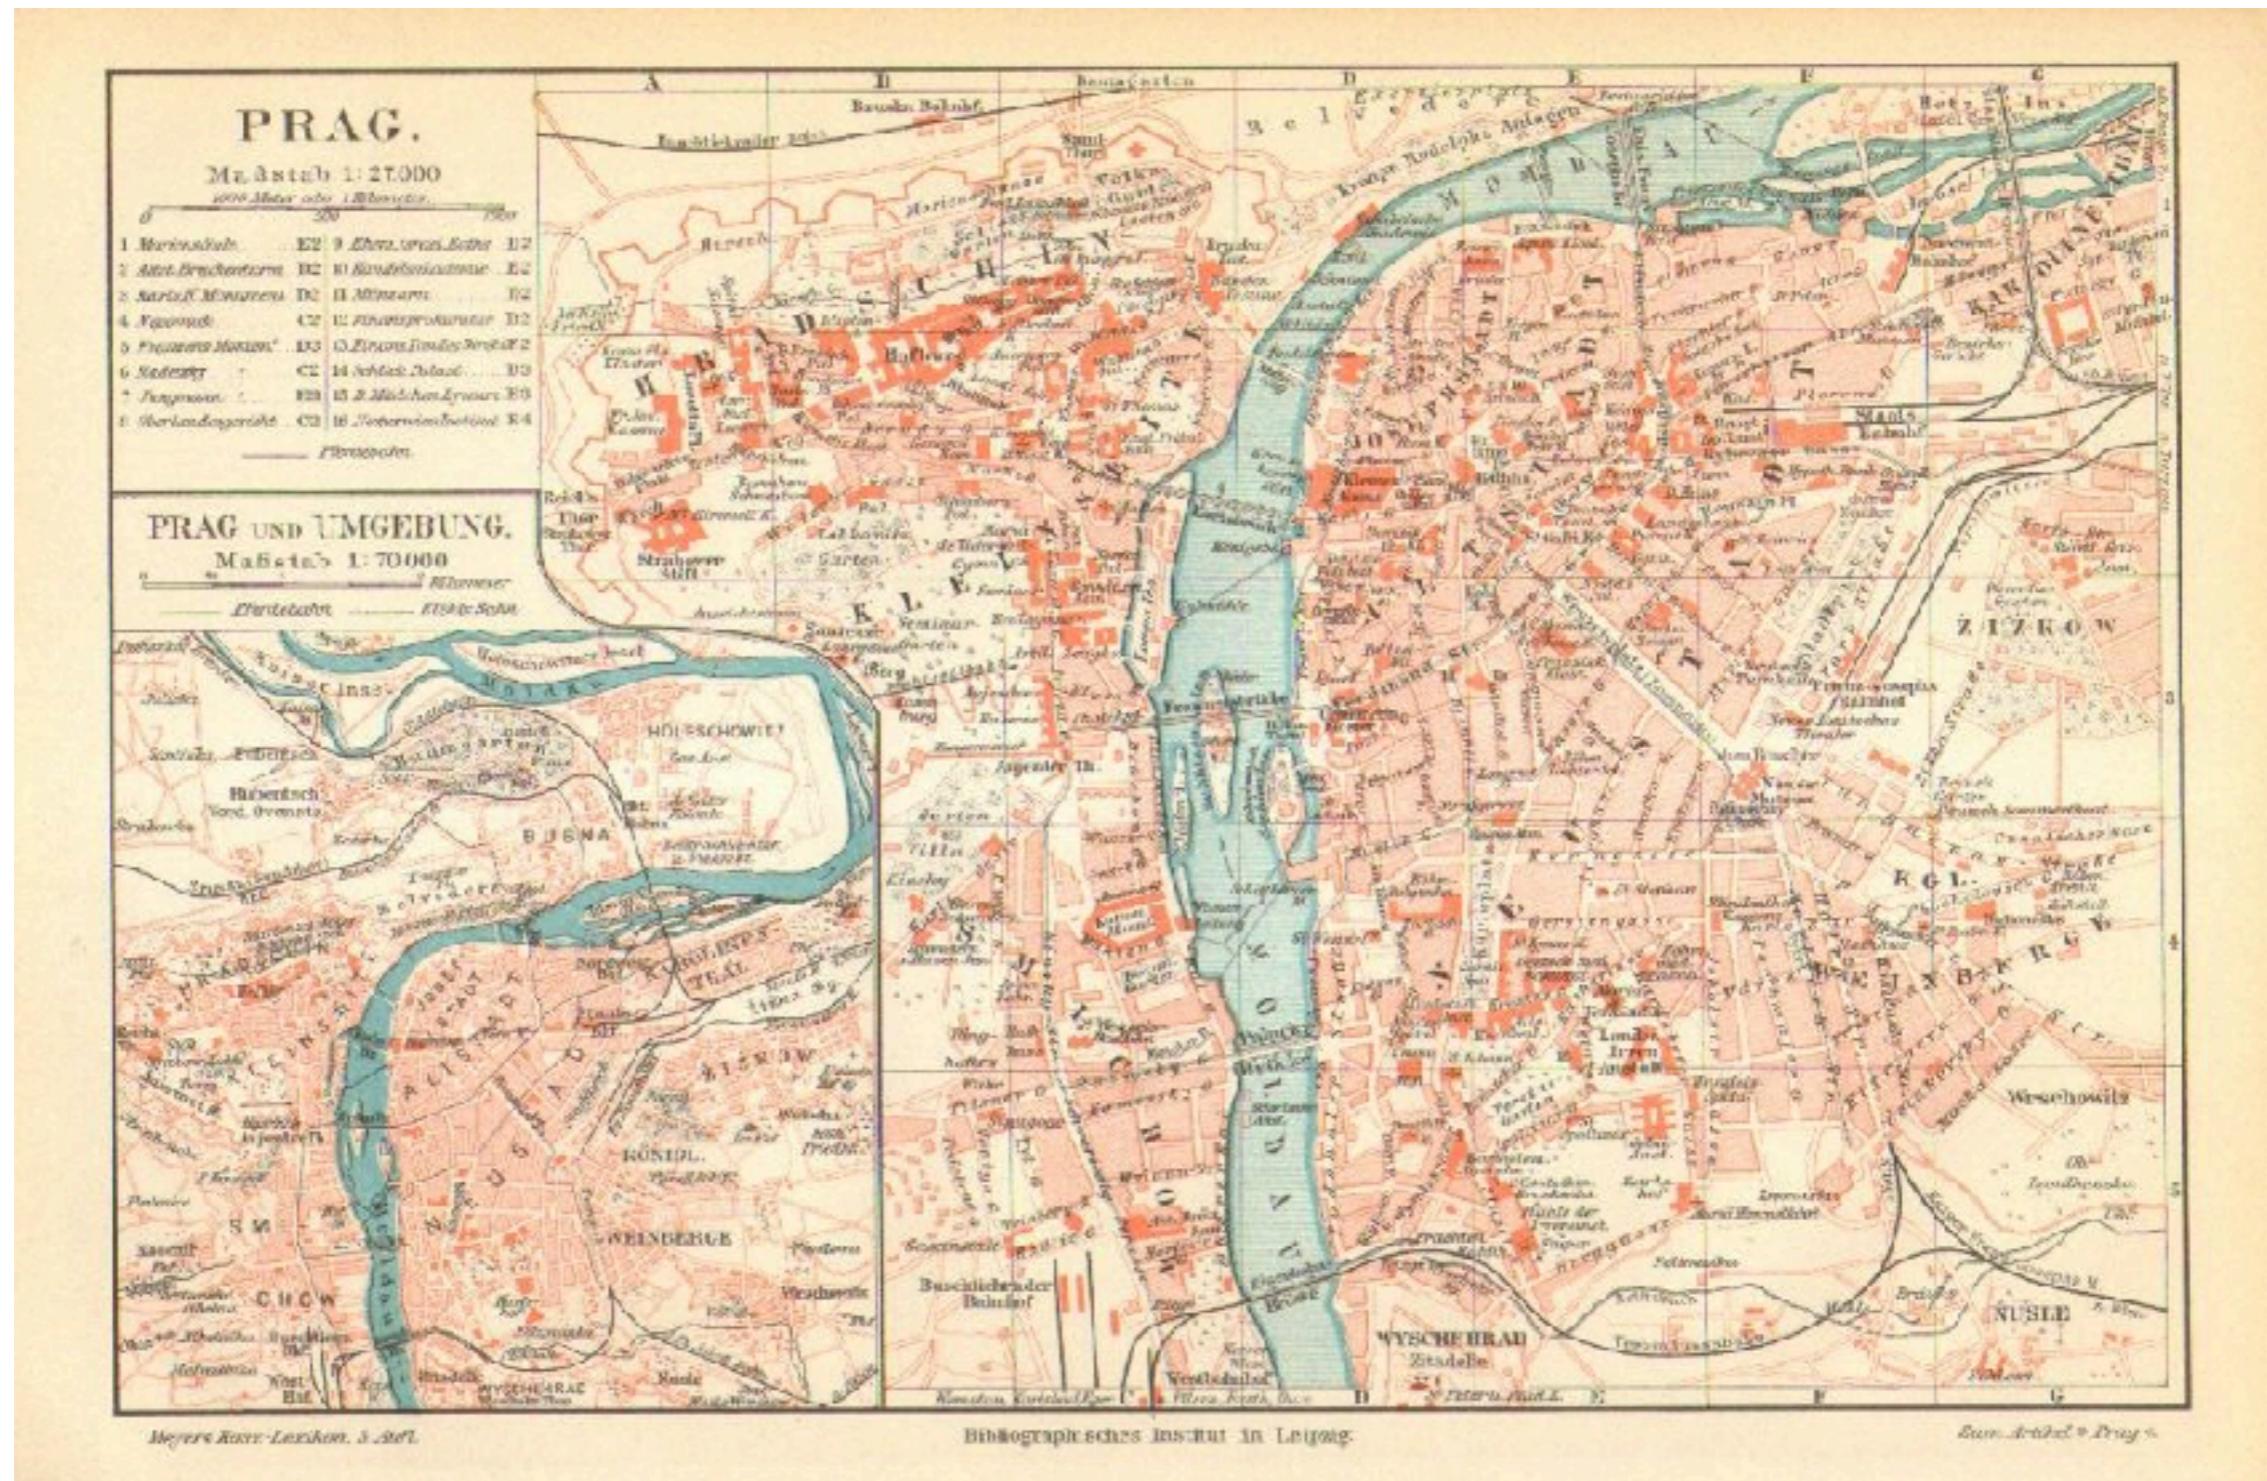
ECCV 2016

[Pixelwise View Selection for
Unstructured Multi-View Stereo](#)

3D Reconstruction



3D Reconstruction



Prague in the 19th century

Mixed Reality



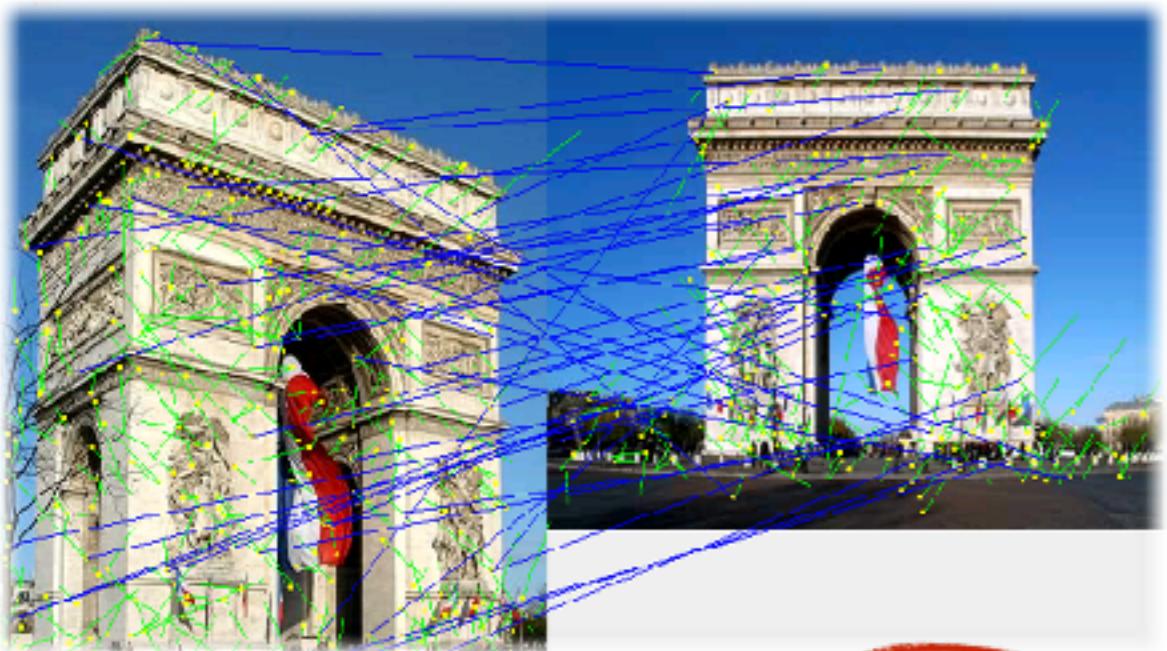
[Microsoft HoloLens](#)

Mixed Reality



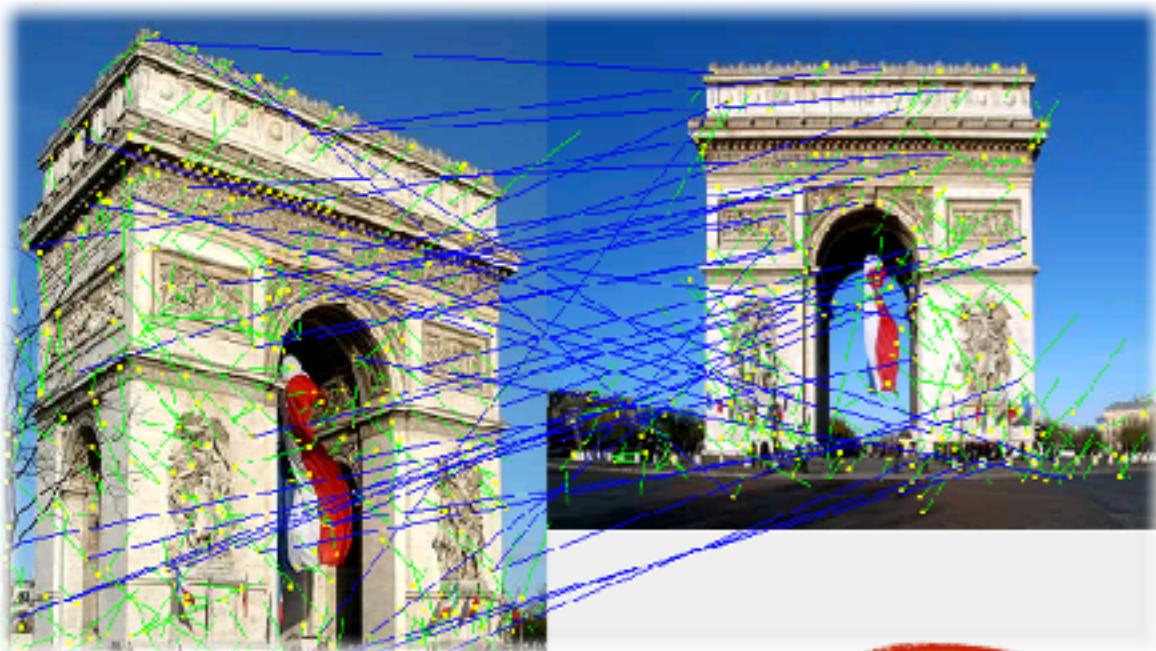
[Microsoft HoloLens](#)

Part 1. Basic image handling and sparse features



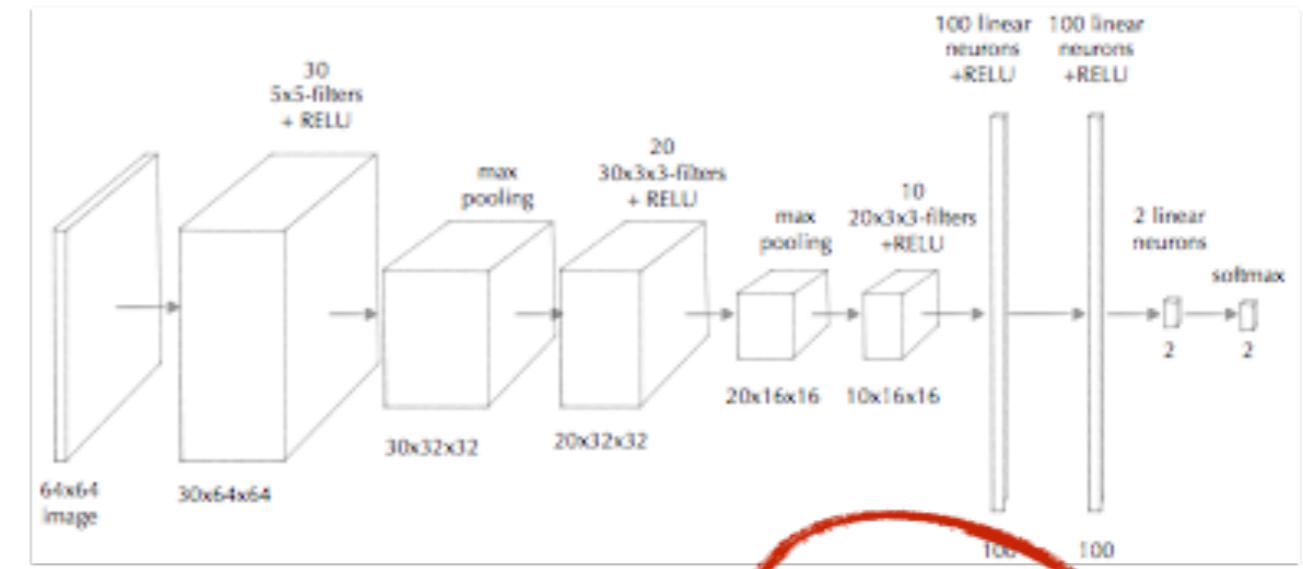
Lab 1

Part 1. Basic image handling and sparse features

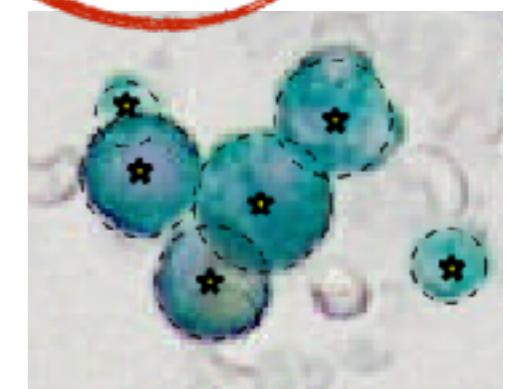
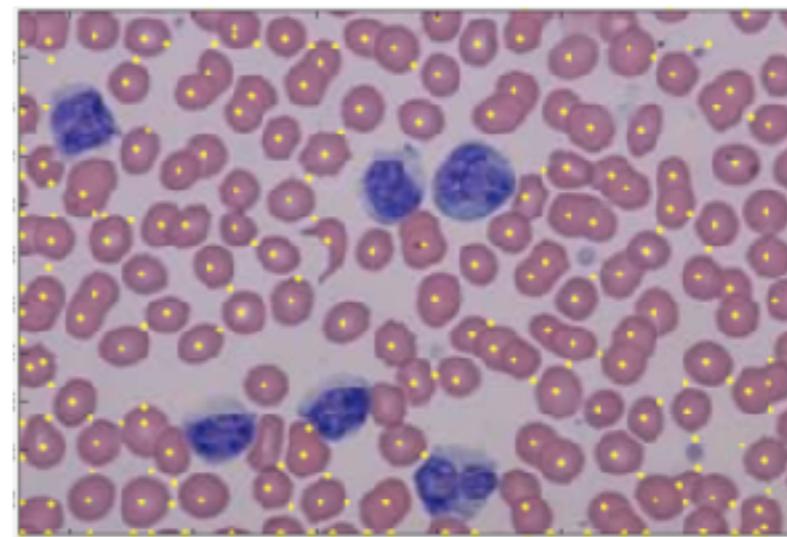


Lab 1

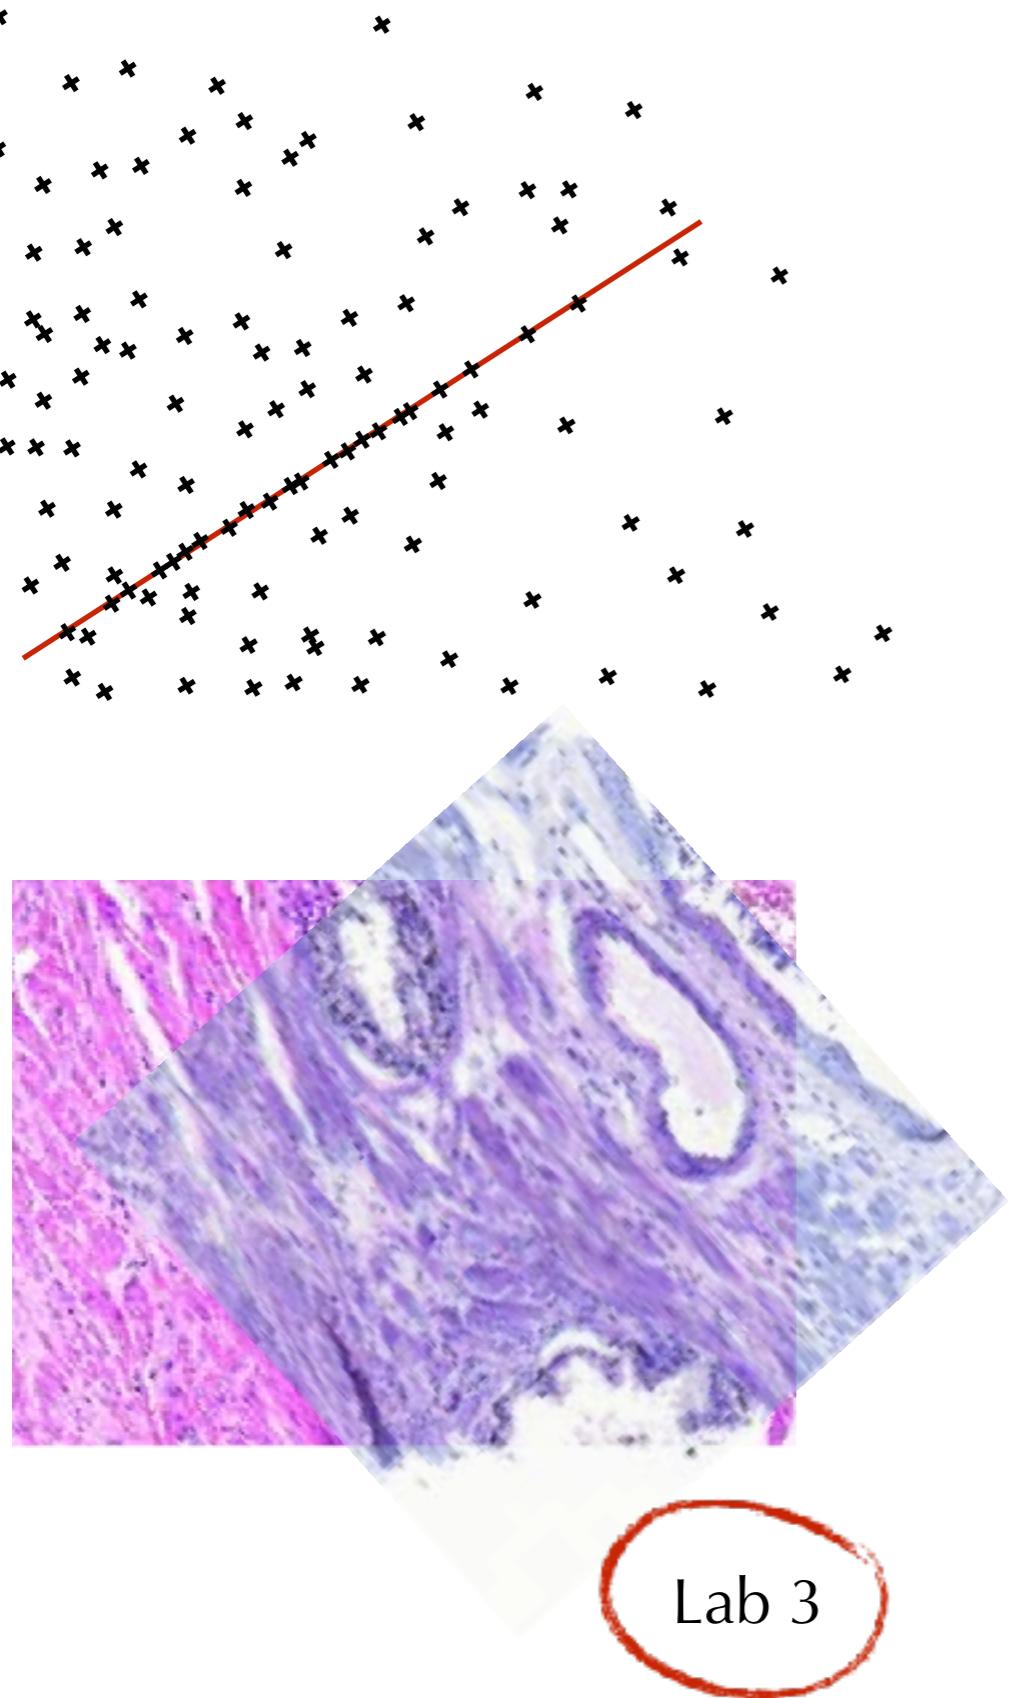
Part 2. Machine learning for classification, detection, and segmentation



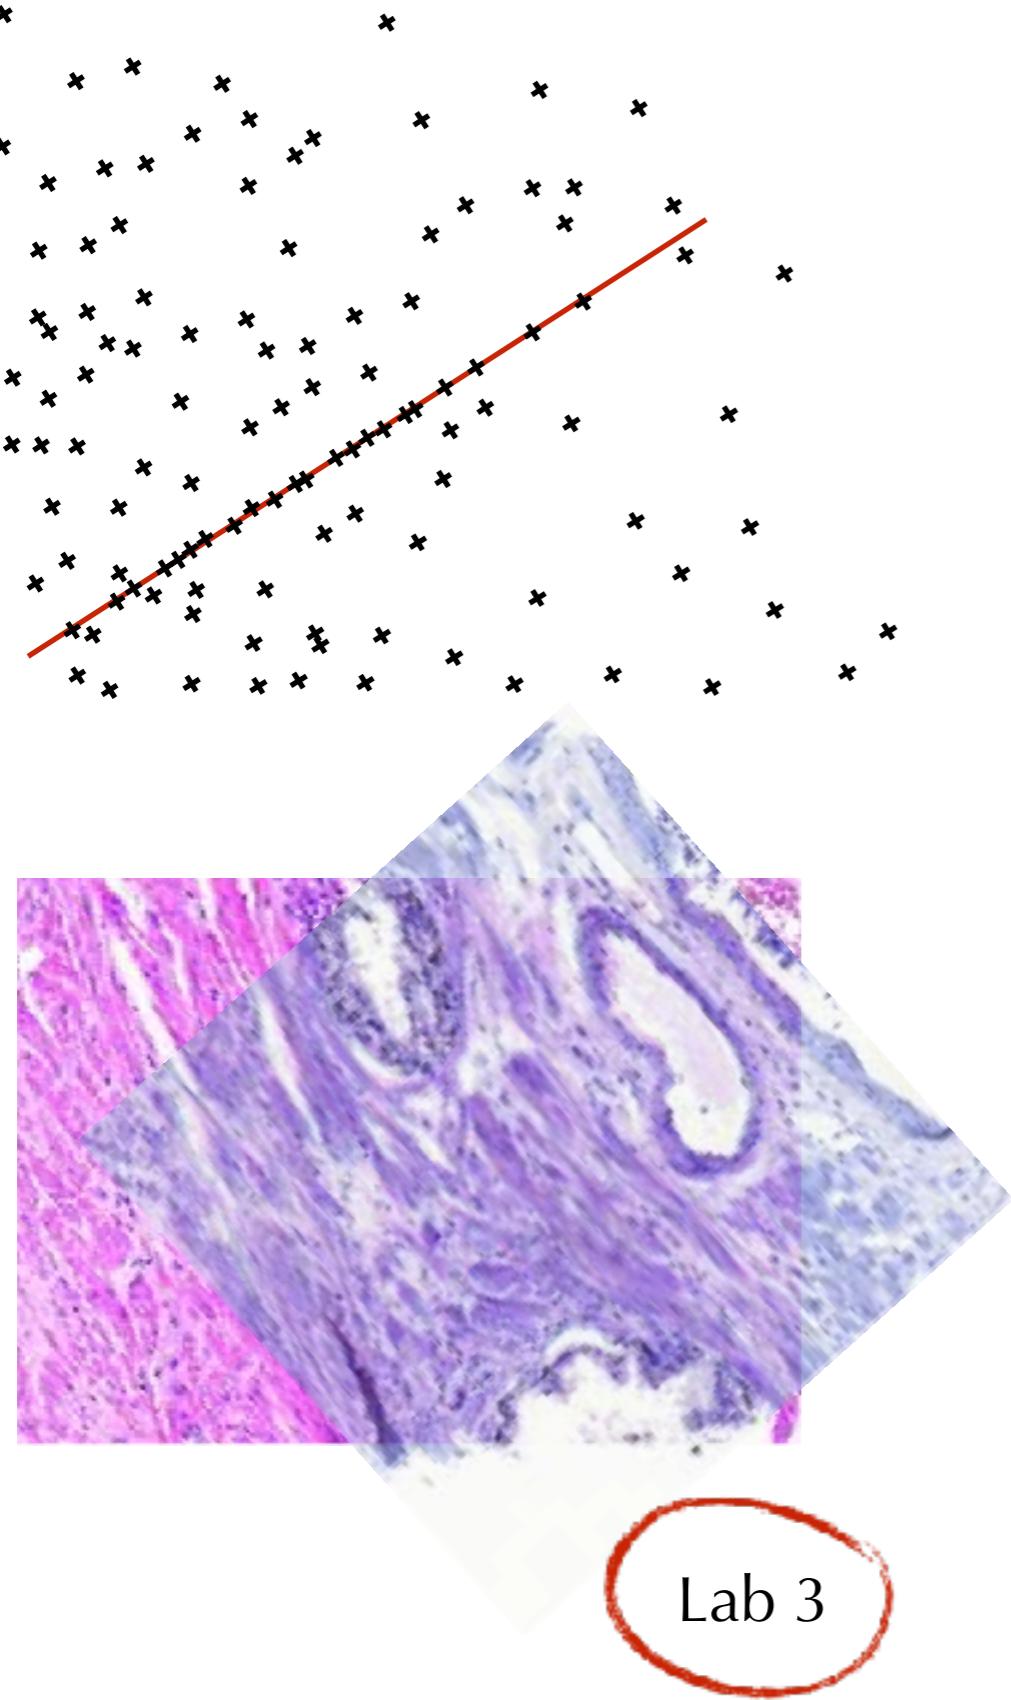
Lab 2
Project A



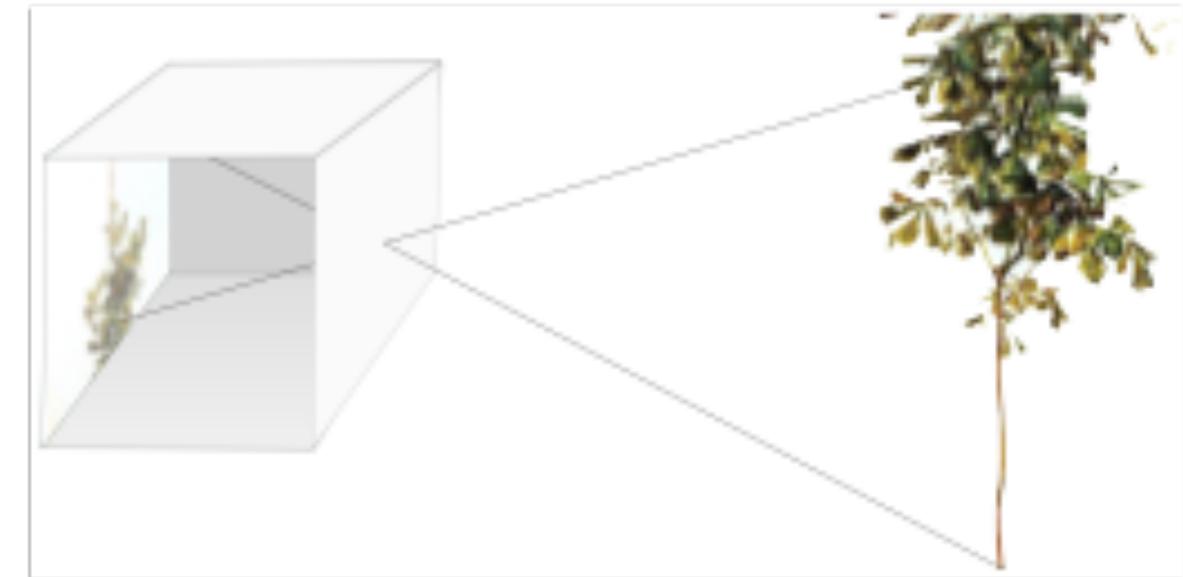
Part 3. Robust model estimation and image registration



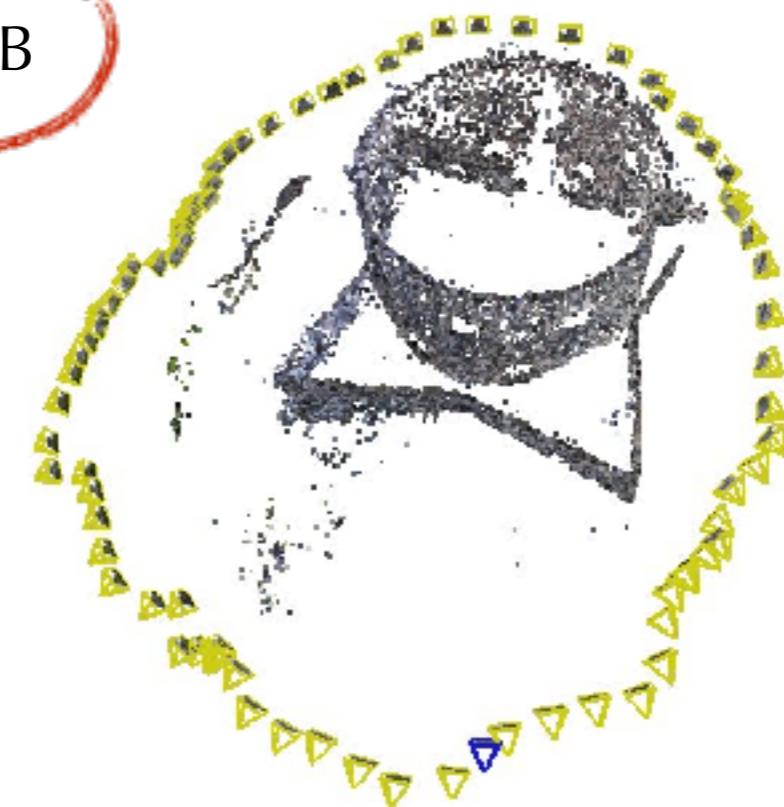
Part 3. Robust model estimation and image registration



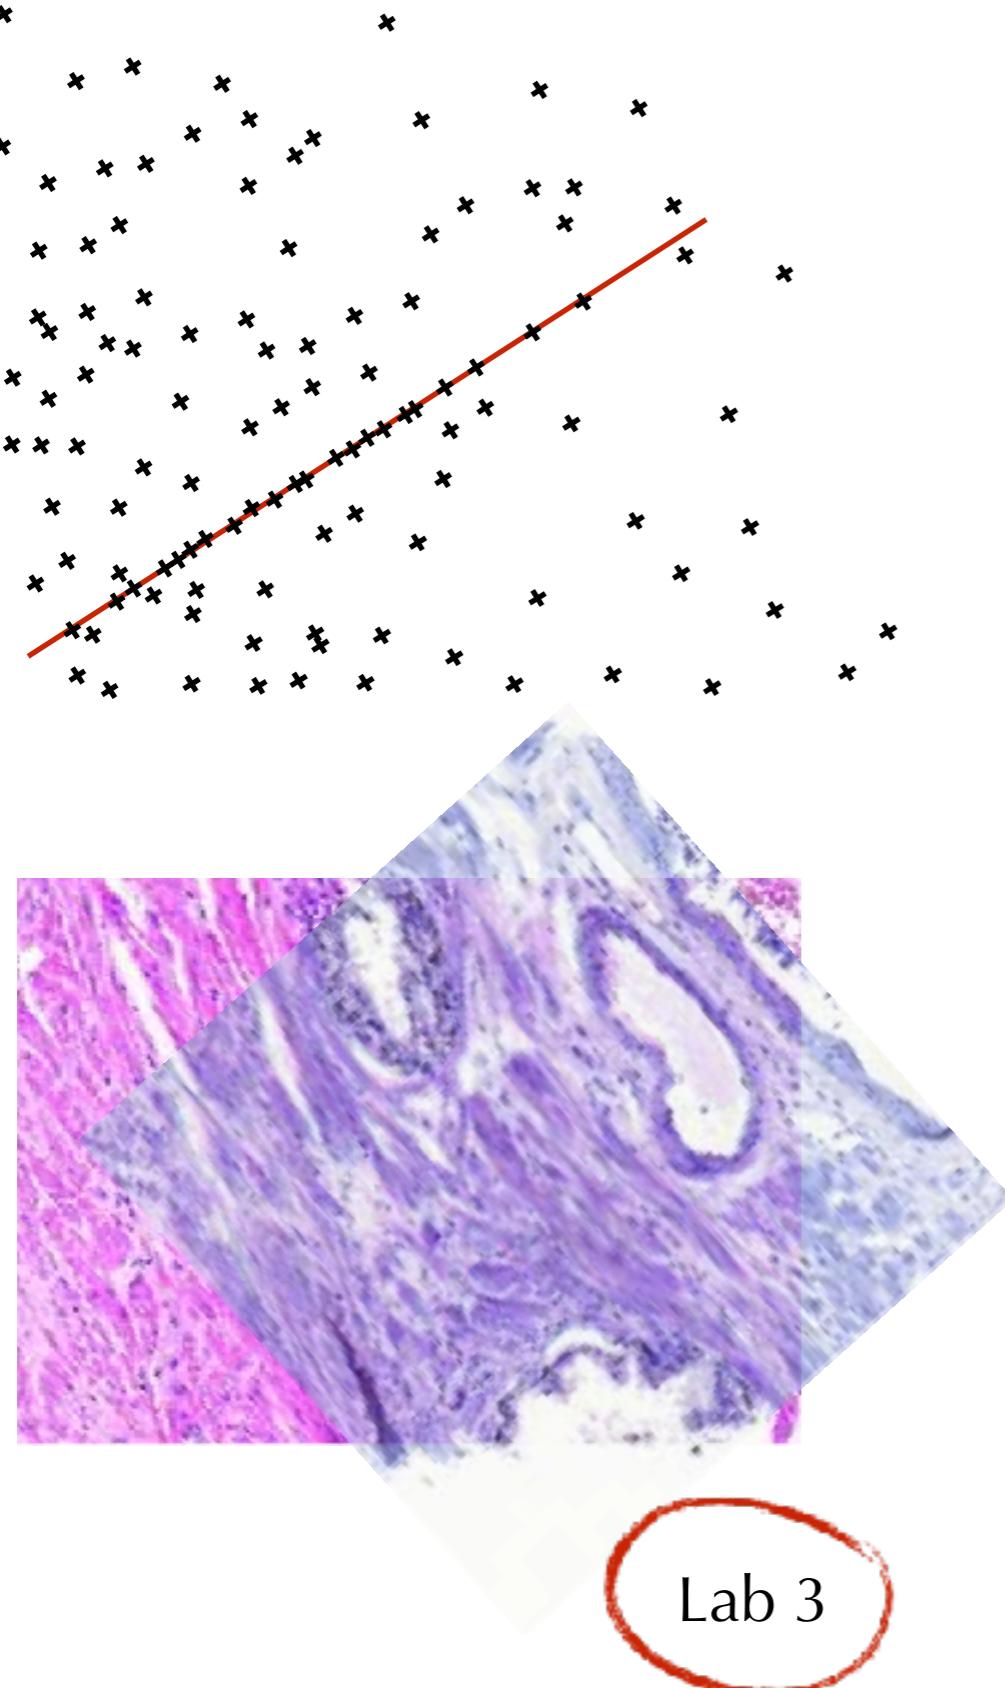
Part 4. Camera geometry and basic 3D reconstruction



Lab 4
Project B



Part 3. Robust model estimation and image registration



Part 4. Camera geometry and basic 3D reconstruction



Lab 4
Project B



EEN020 - Computer Vision

Part 5. Generative models for machine learning and machine learning in 3D vision

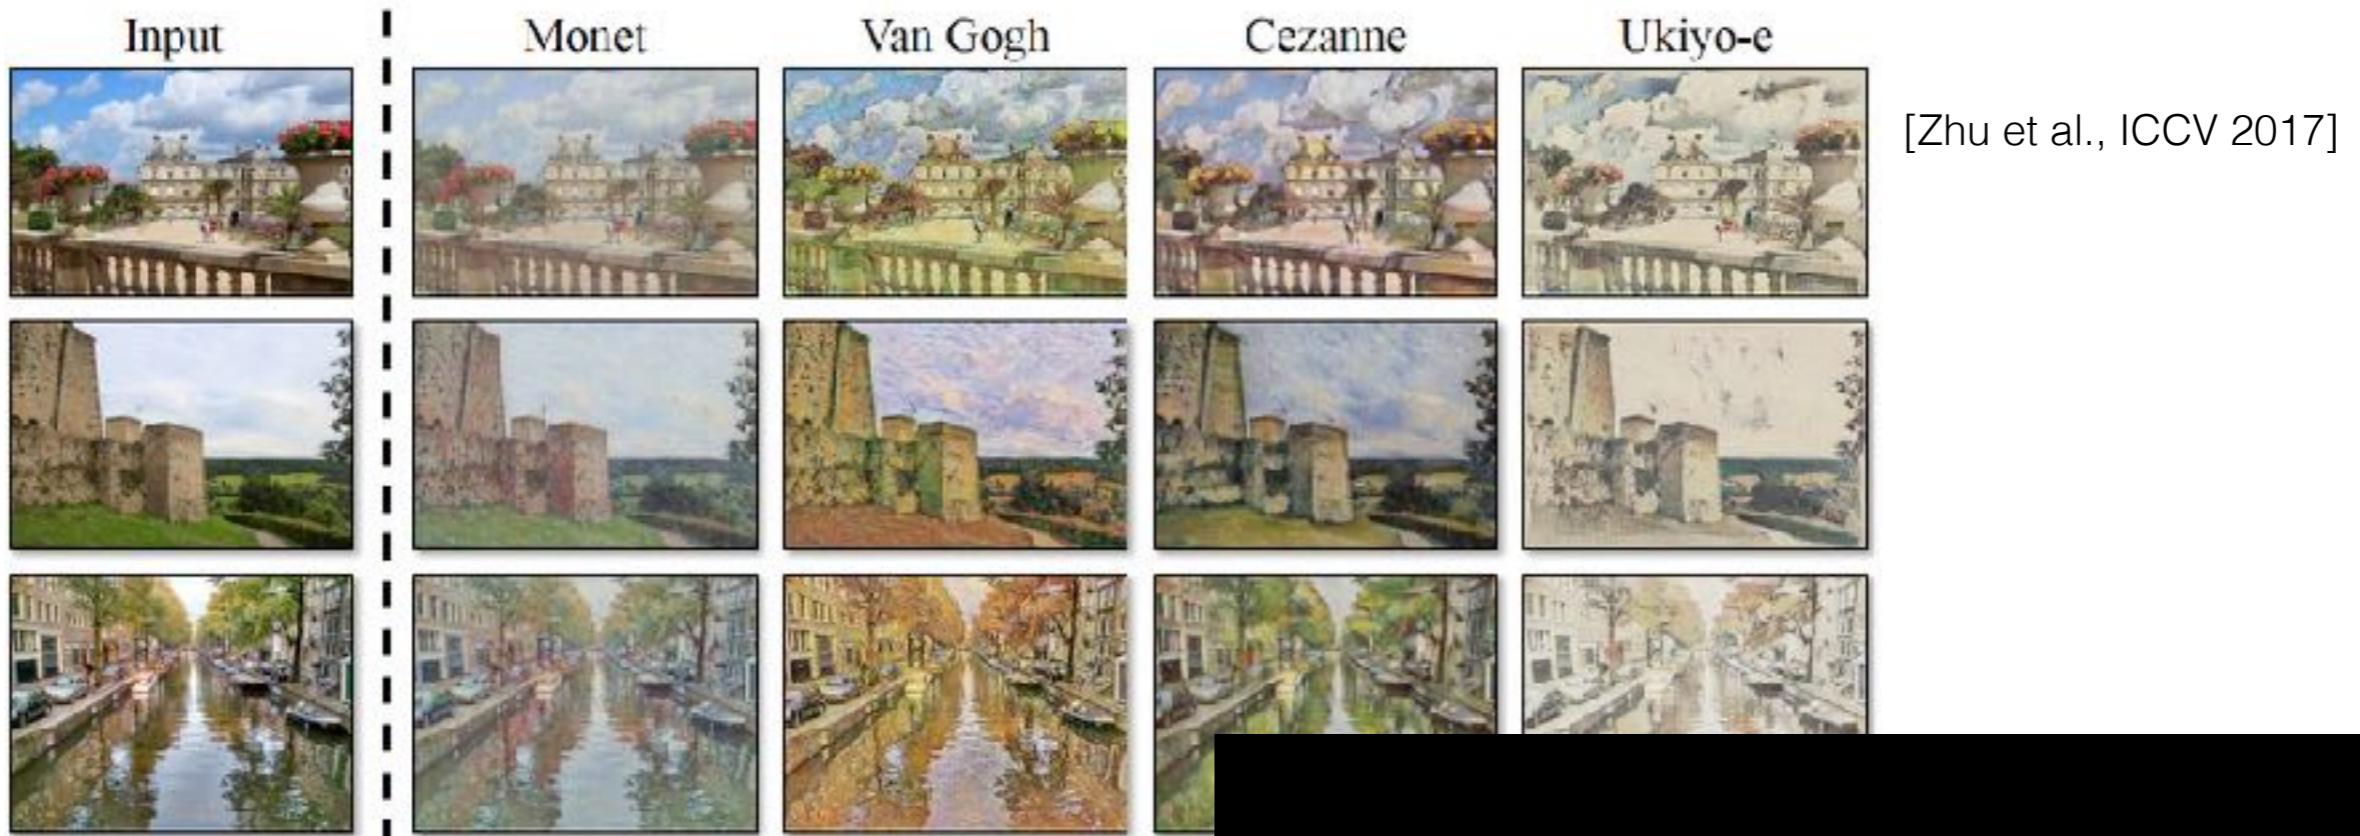


Learning Less is More -
6D Camera Localization
via 3D Surface Regression

Eric Brachmann Carsten Rother

Visual Learning Lab
University Heidelberg

Part 5. Generative models for machine learning and machine learning in 3D vision



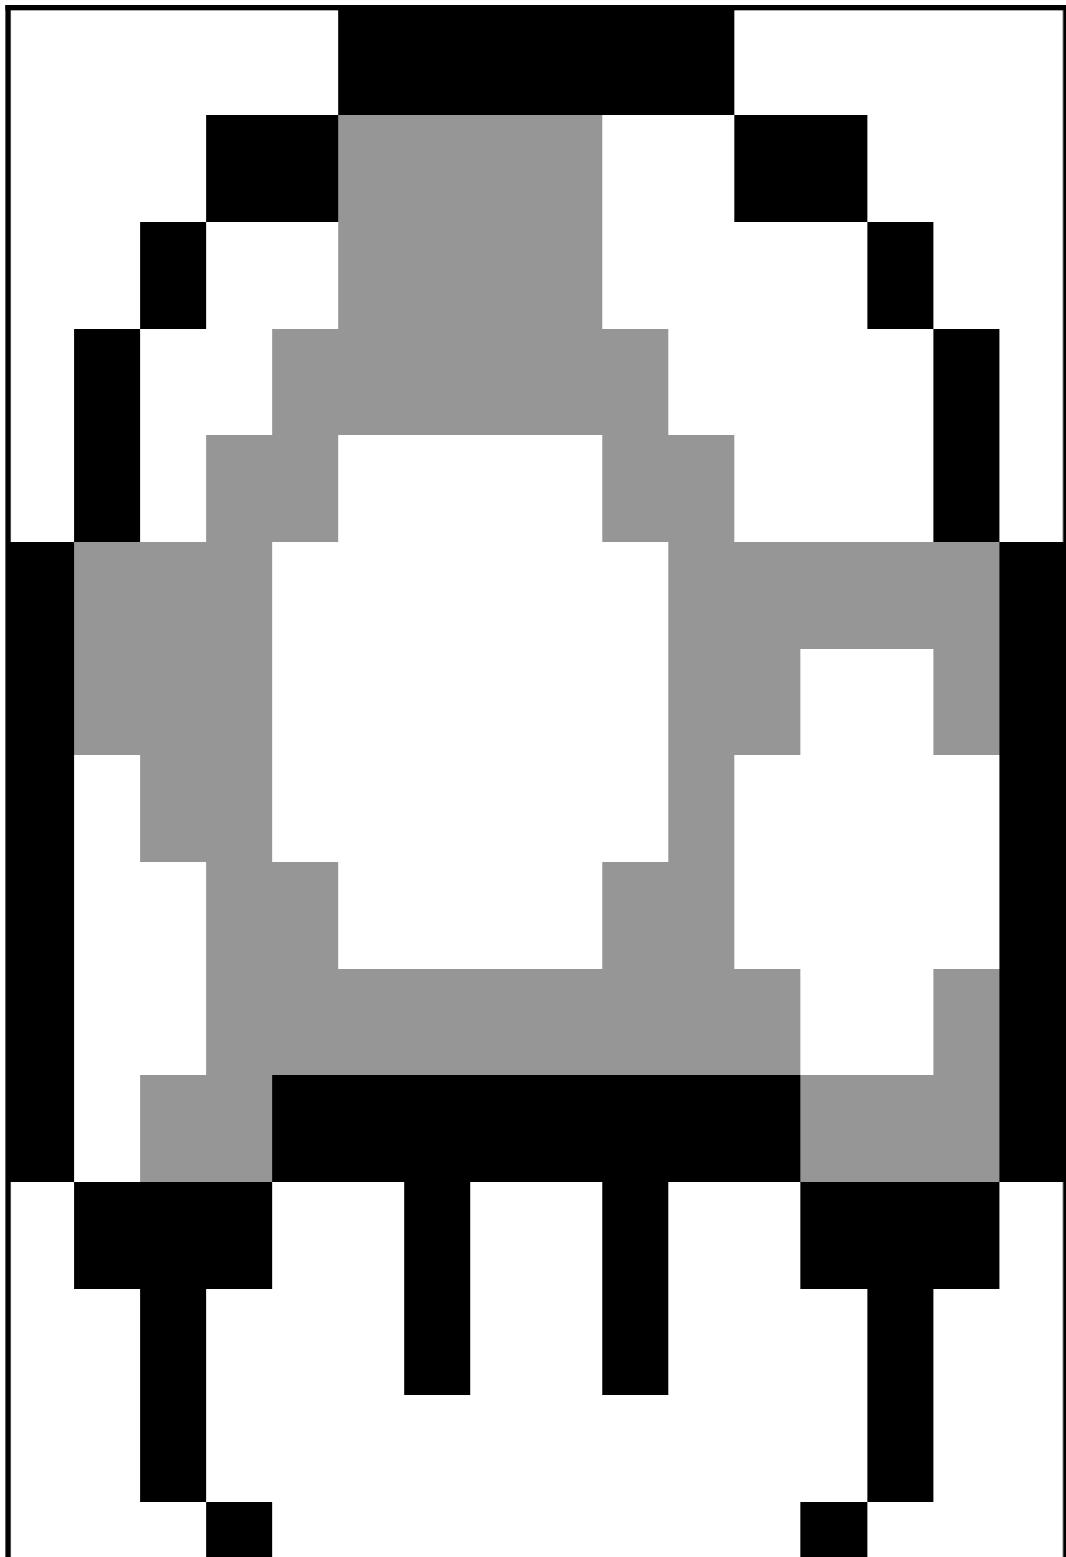
Learning Less is More -
6D Camera Localization
via 3D Surface Regression

Eric Brachmann Carsten Rother

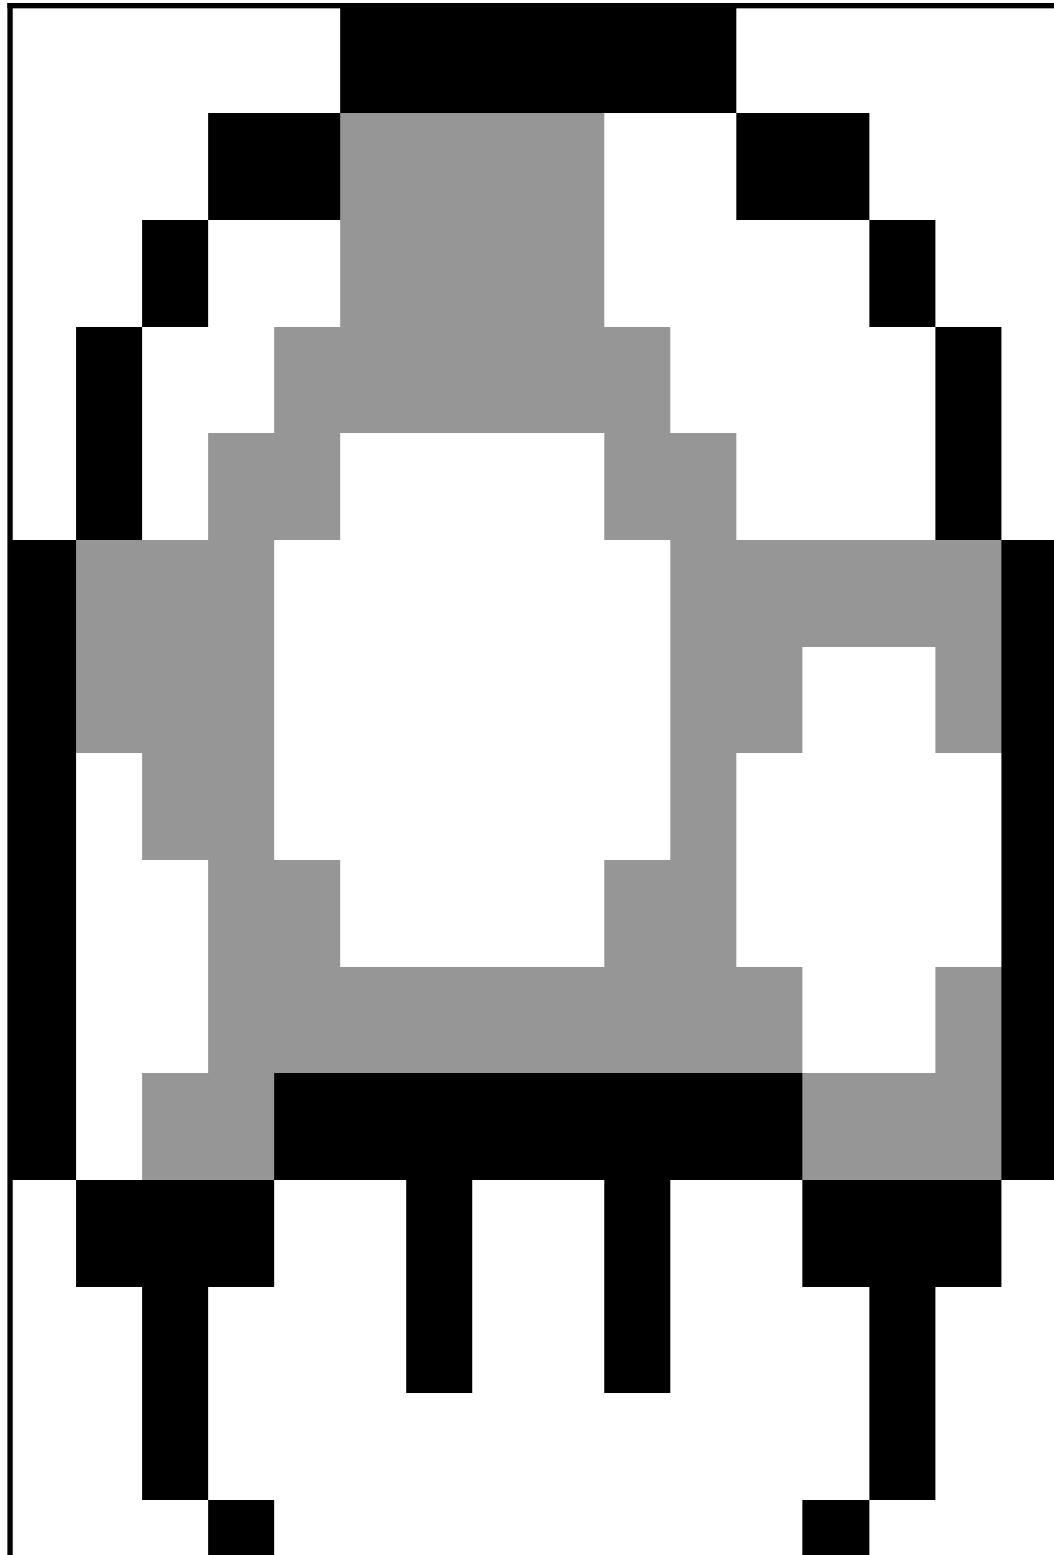
Visual Learning Lab
University Heidelberg

Modeling Images

What is an Image?



What is an Image?



25	25	25	25	25	25	0	0	0	0	0	25	25	25	25	25
5	5	5	5	5	5	0	0	0	0	0	5	5	5	5	5
25	25	25	0	0	0	15	15	15	15	25	25	0	0	25	25
5	5	5	0	0	0	0	0	0	0	0	0	0	0	25	25
25	25	0	25	25	25	15	15	15	15	25	25	25	25	0	25
5	5	0	25	25	25	15	15	15	15	25	25	25	25	0	25
25	0	25	25	15	15	15	0	15	15	25	25	25	25	25	0
5	0	25	25	15	15	15	0	25	25	25	25	25	25	25	0
0	15	15	15	25	25	25	25	25	25	15	15	15	15	15	0
0	0	0	0	5	5	5	5	5	5	0	0	0	0	0	0
0	15	15	15	25	25	25	25	25	25	15	15	15	15	15	0
0	0	0	0	5	5	5	5	5	5	0	0	0	0	0	0
0	25	15	15	25	25	25	25	25	25	15	25	25	25	25	0
0	5	0	0	5	5	5	5	5	5	0	25	25	25	25	0
0	25	25	15	0	15	25	25	25	25	15	0	25	25	25	25
0	5	0	0	0	0	15	15	15	15	0	0	25	25	15	0
0	25	15	15	0	0	0	0	0	0	0	0	15	15	15	0
25	0	0	0	25	25	0	25	25	0	25	25	0	0	0	25
5	5	0	25	25	25	0	25	25	0	25	25	0	0	25	25
25	5	0	25	25	25	0	25	25	0	25	25	0	0	25	25
5	5	0	25	25	25	0	25	25	0	25	25	0	0	25	25
25	5	0	25	25	25	0	25	25	0	25	25	0	0	25	25
5	5	0	25	25	25	0	25	25	0	25	25	0	0	25	25

Images are Vectors

Addition, Subtraction

$$\begin{array}{|c|c|c|c|} \hline 150 & 170 & \dots & \dots \\ \hline \vdots & 150 & & \\ \hline & 150 & & \\ \hline \end{array} + \begin{array}{|c|c|c|c|} \hline 0 & 0 & 0 & 0 \\ \hline 0 & 150 & 150 & 0 \\ \hline 0 & 170 & 130 & 0 \\ \hline 0 & 0 & 0 & 0 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 150 & 130 & \dots & \dots \\ \hline \vdots & 300 & 300 & \\ \hline & 300 & 360 & \\ \hline \end{array}$$

Images are Vectors

Scalar multiplication, dot product

$$2 \cdot \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} = \begin{bmatrix} 2 & 4 & 6 \\ 8 & 10 & 12 \\ 14 & 16 & 18 \end{bmatrix}$$

$$\begin{bmatrix} 6 & 2 & 5 \\ 1 & 0 & 3 \\ 0 & 1 & 0 \end{bmatrix} \cdot \begin{bmatrix} 10 & 20 & 30 \\ 40 & 50 & 60 \\ 90 & 80 & 70 \end{bmatrix} = 6 \cdot 10 + 2 \cdot 20 + 0 \cdot 30 + 1 \cdot 40 + 0 \cdot 30 + \dots$$

Images are Vectors



Average
speed limit sign

Images are Vectors



Average
speed limit sign



image source: [pinterest](#)

Average
president

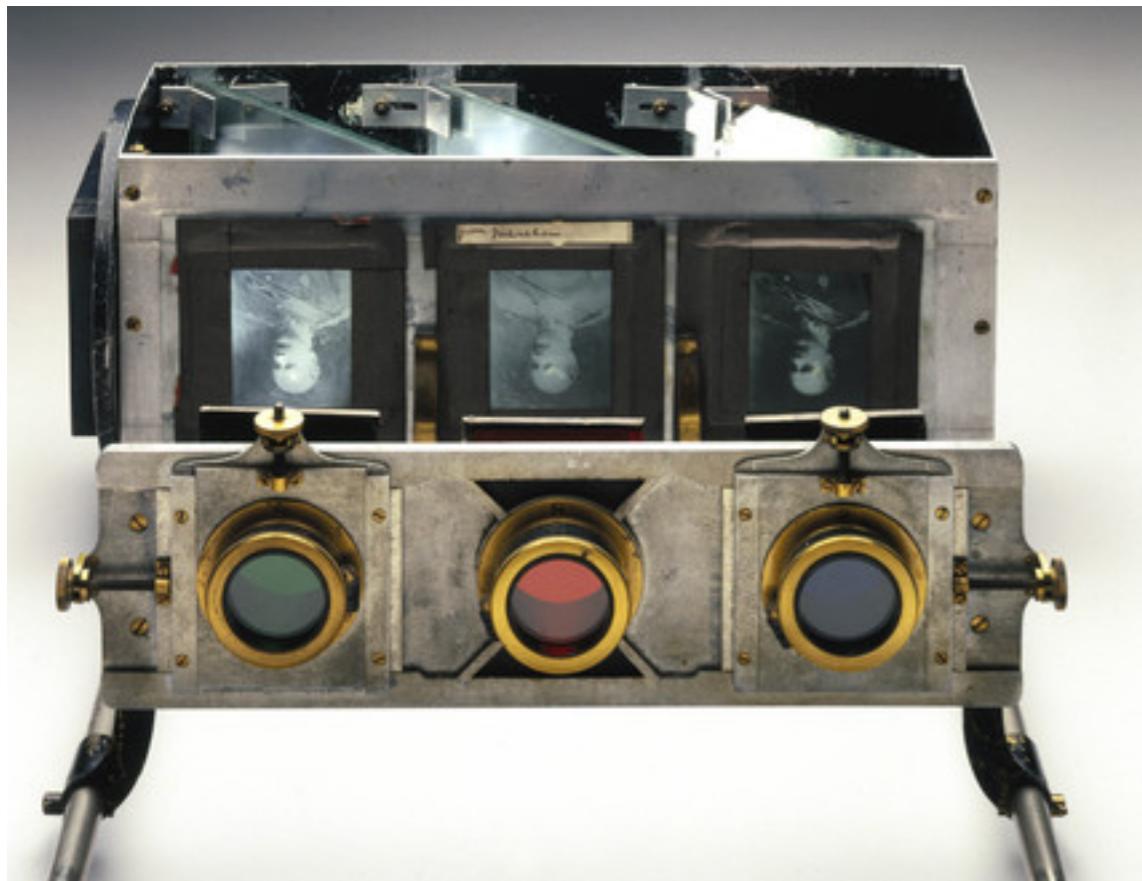
Color



Color



Color



RGB



122	124	145	153	100	123	109		
233	232	211	198	188	177	155		
54	121	121	121	78	13	87	76	
54	56	78	89	23	23	12	12	
255	78	18	24	25	53	23	54	23
255	255	23	34	24	35	54	23	34
255	244	34	53	35	35	52	43	63
255	34	56	34	87	76	56	87	
244	98	9	76	58	48	59	28	
67	28	38	28	28	38	38		
28	8	0	29	9	82	0		

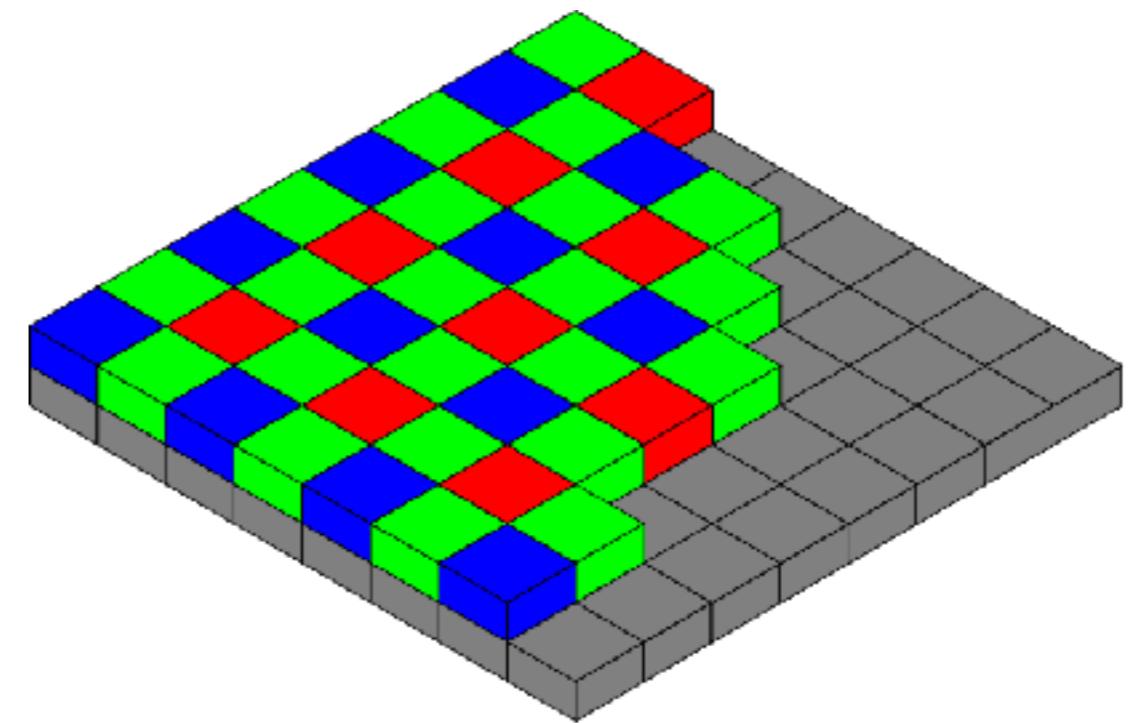
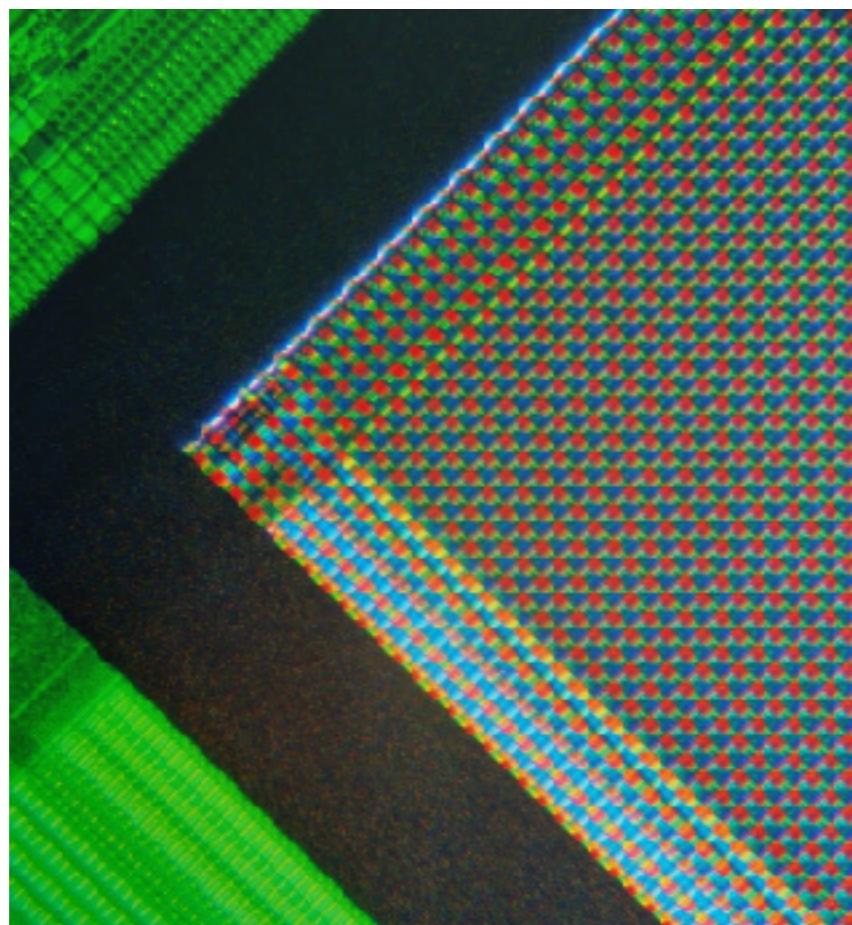
RGB



54	122	124	145	153	100	123	109
54	233	232	211	198	188	177	155
255	121	121	121	78	13	87	76
255	56	78	89	23	23	12	12
255	78	18	24	25	53	23	54
255	255	23	34	24	35	54	23
255	244	34	53	35	35	52	43
255	34	56	34	87	76	56	87
244	98	9	76	58	48	59	28
67	28	38	28	28	38	38	38
28	8	0	29	9	82	0	

Channel-wise operations

Digital Cameras

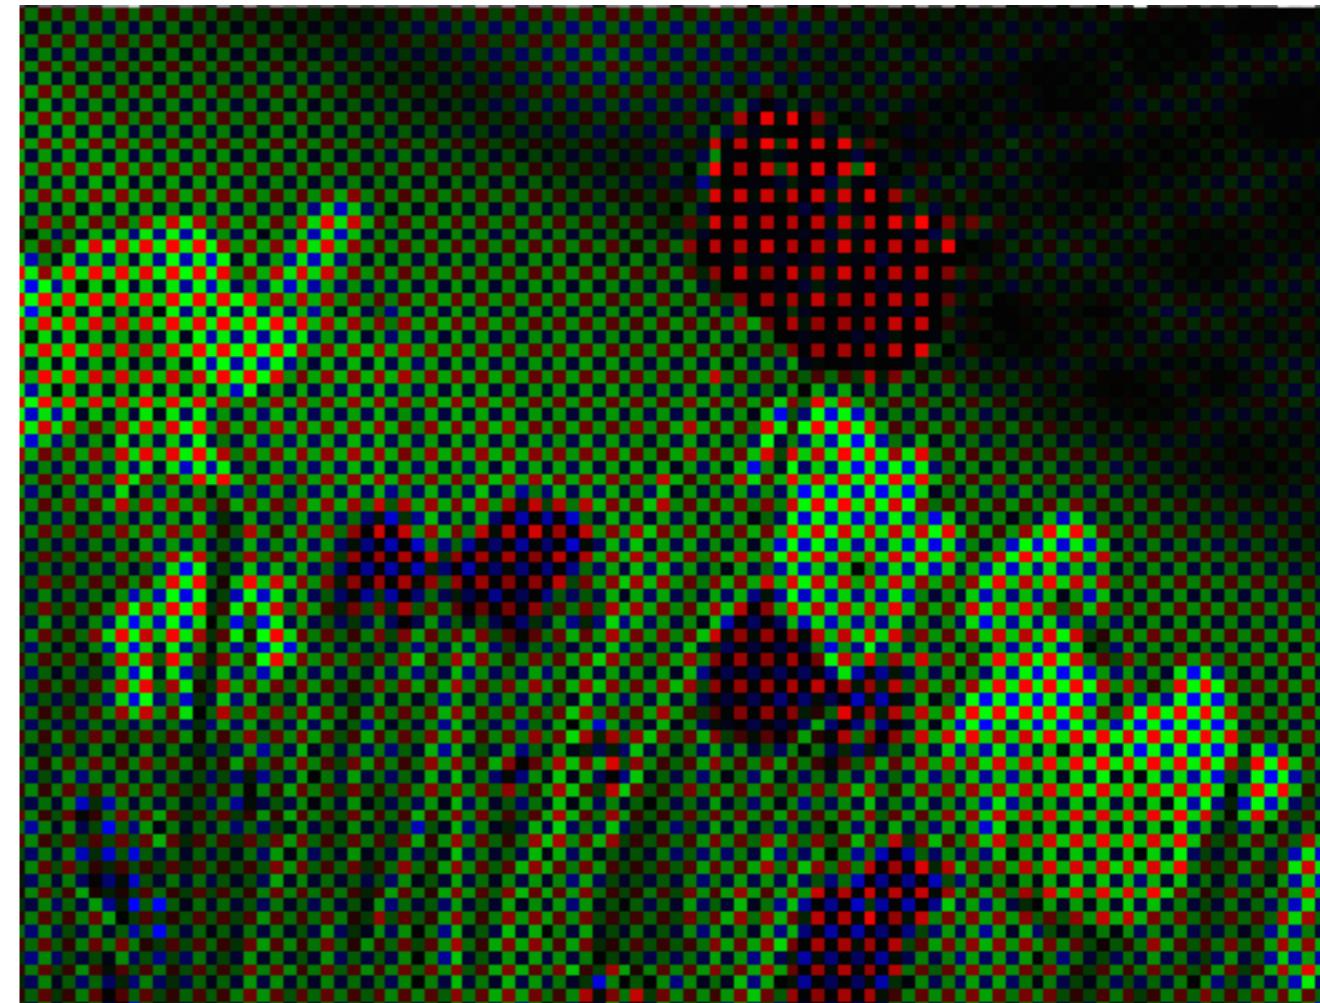


Bayer filter

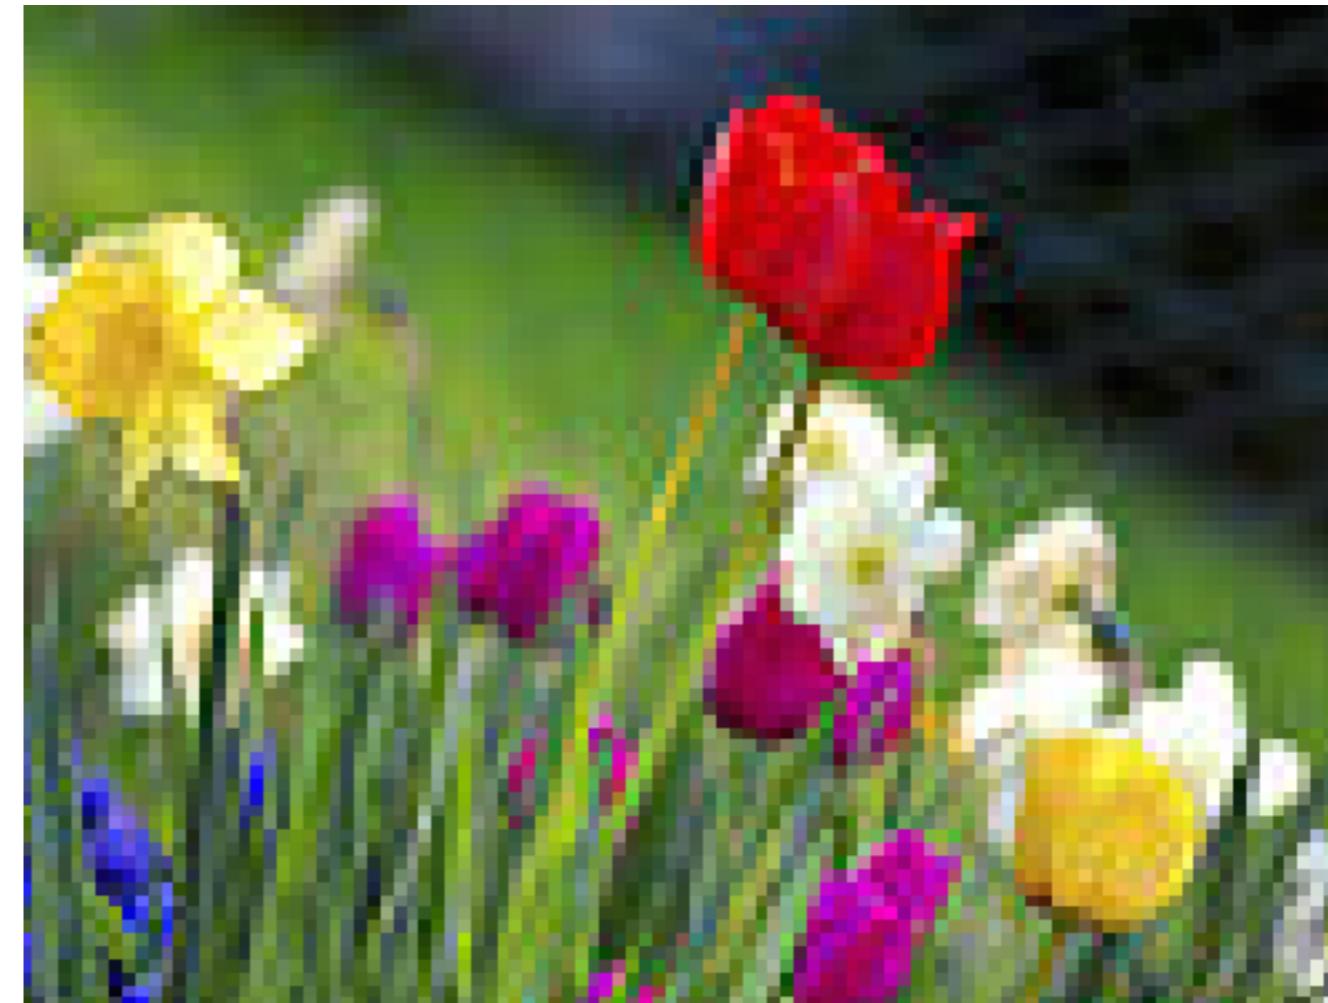
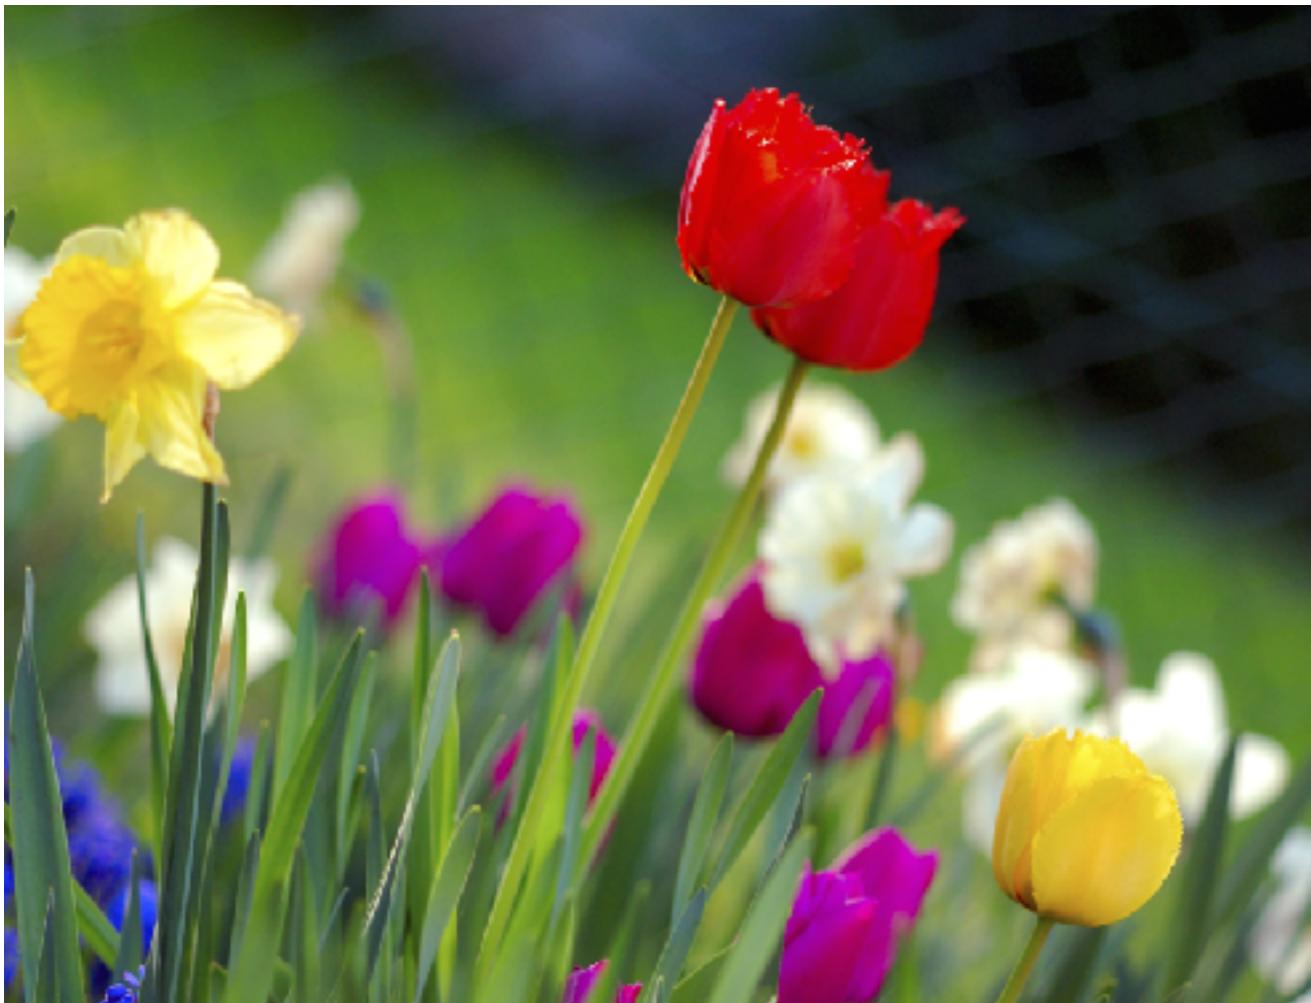
Digital Cameras



Digital Cameras



Digital Cameras

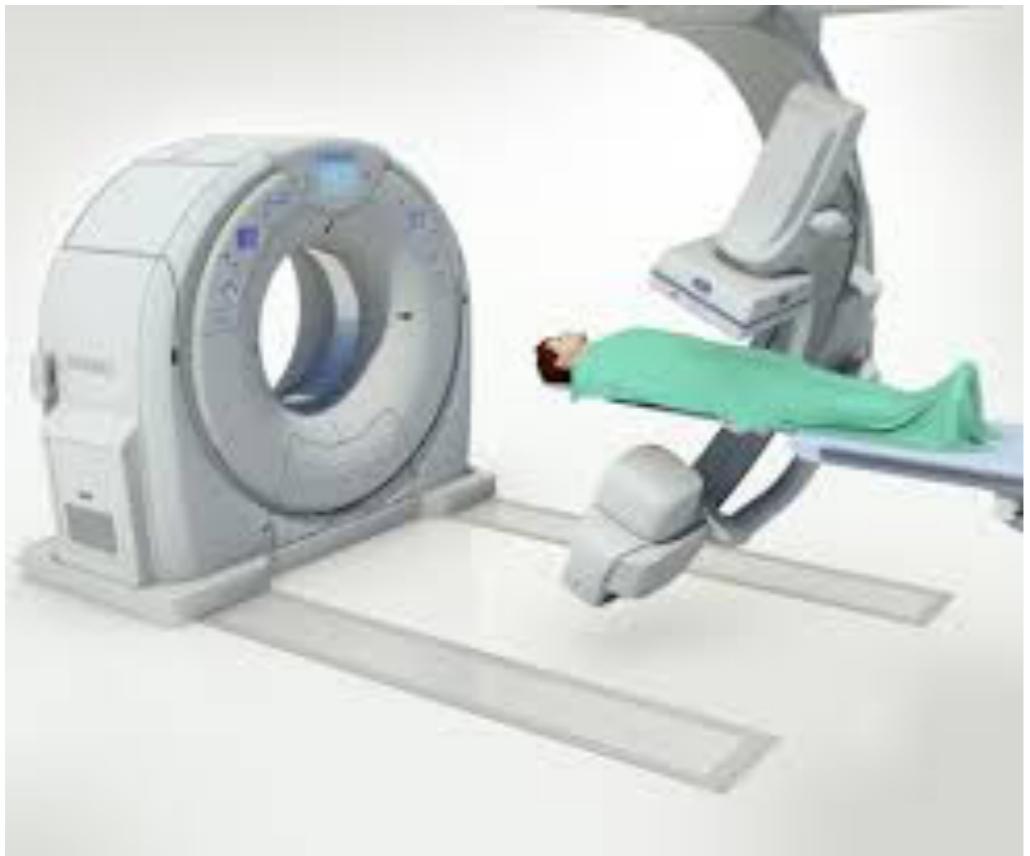


Demosaicing

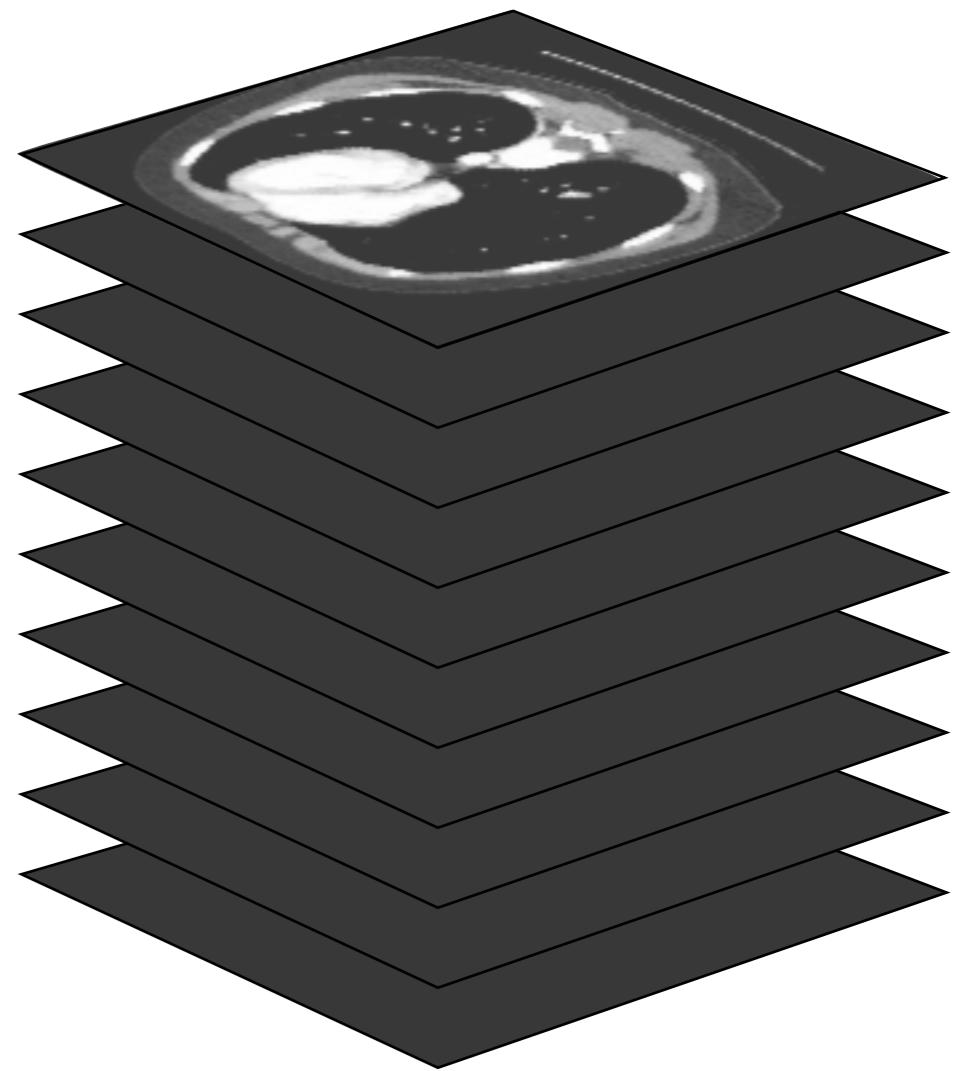
Lighting and White Balance



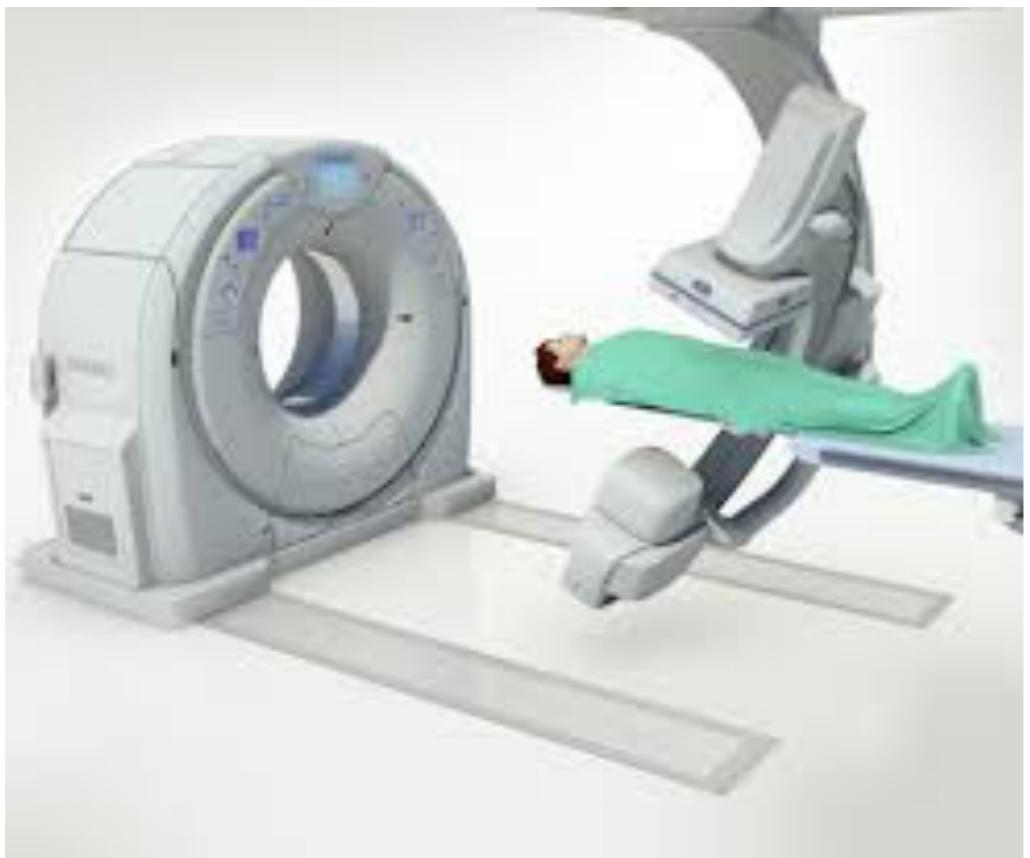
3D Images



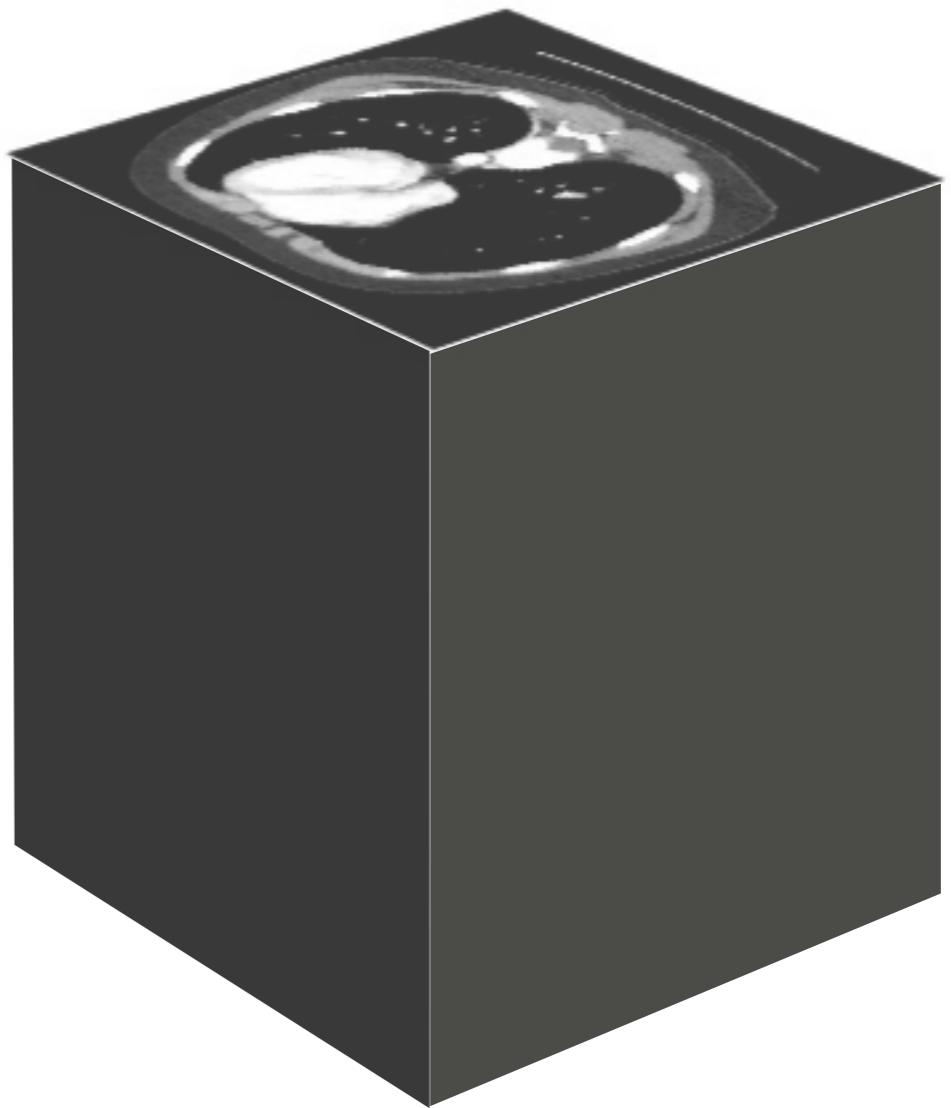
Computed tomography



3D Images

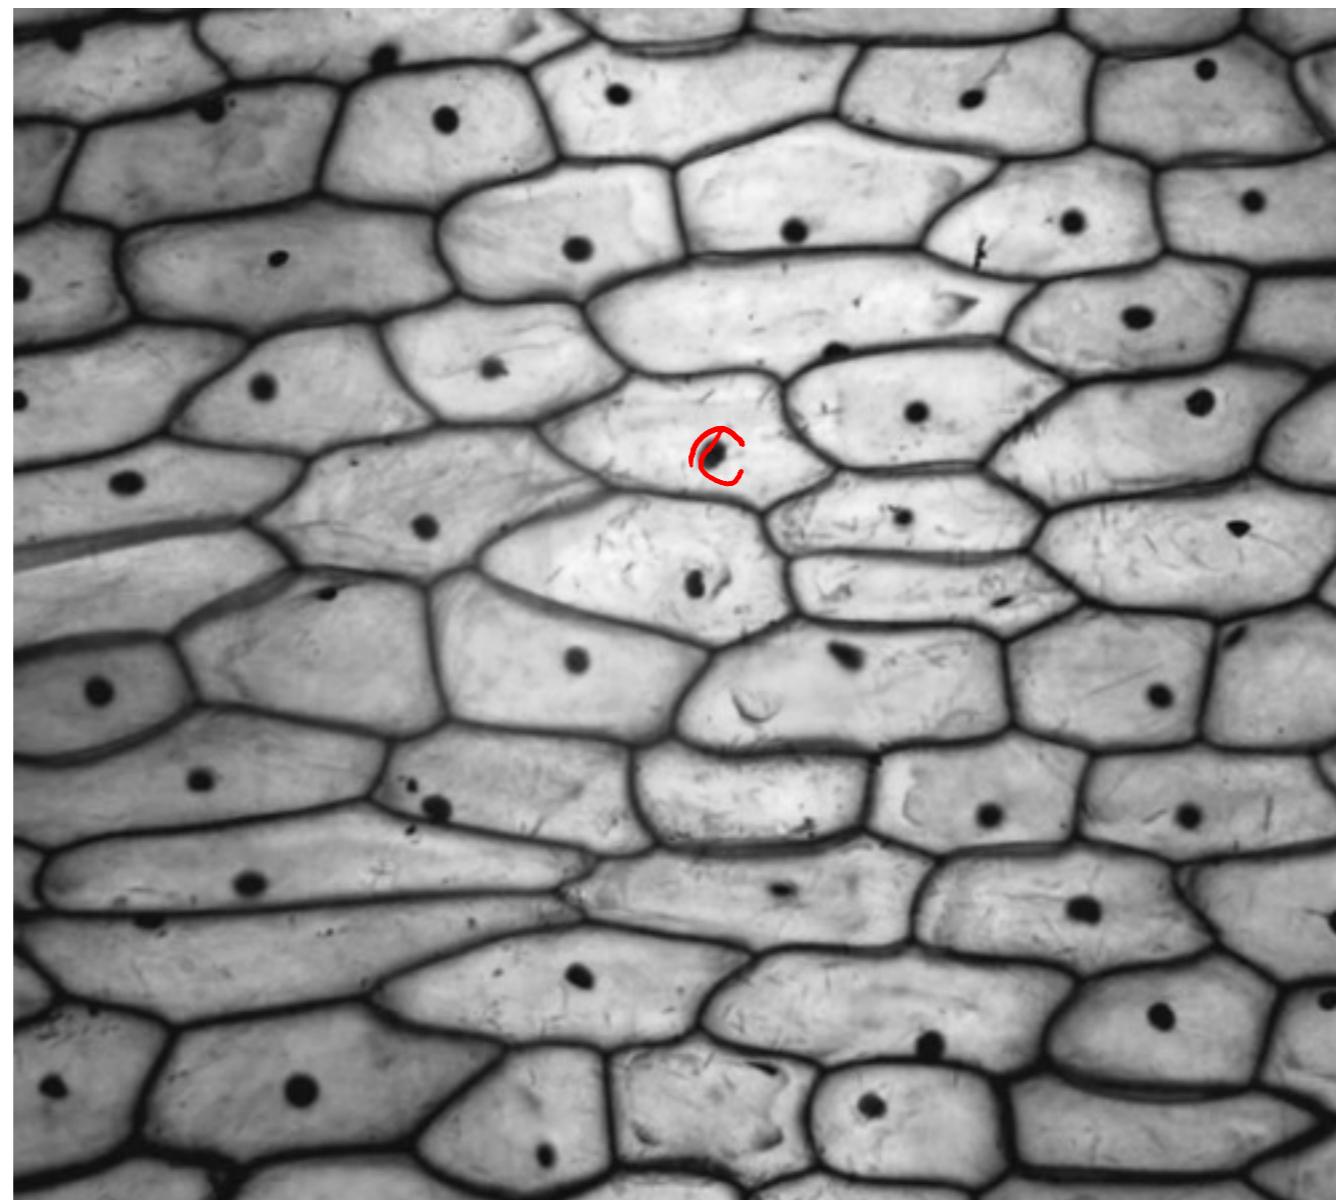


Computed tomography



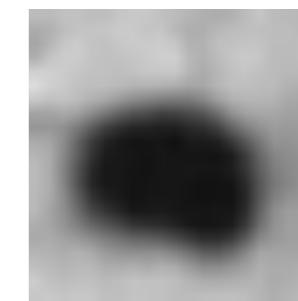
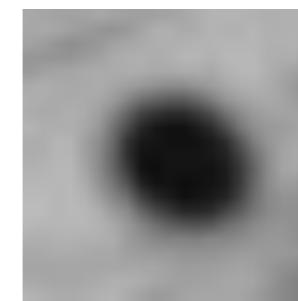
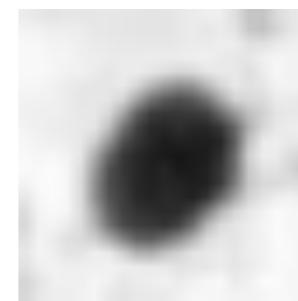
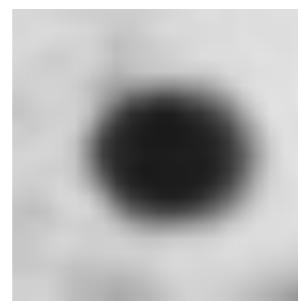
Filtering and Linear Classifiers

Count the Nuclei

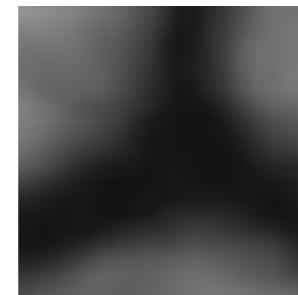


Classification

nuclei



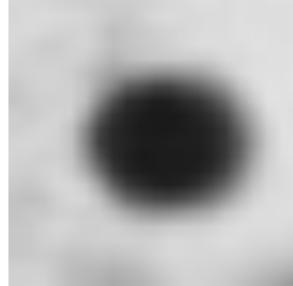
other

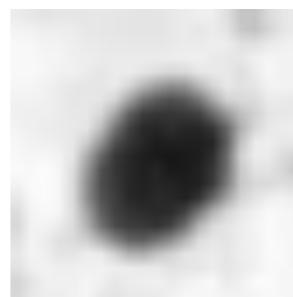


21 x 21 patches

A Linear Classifier

TARGET \rightarrow w


$$w \cdot$$
$$> \tau$$

$$w \cdot$$

$$> \tau$$

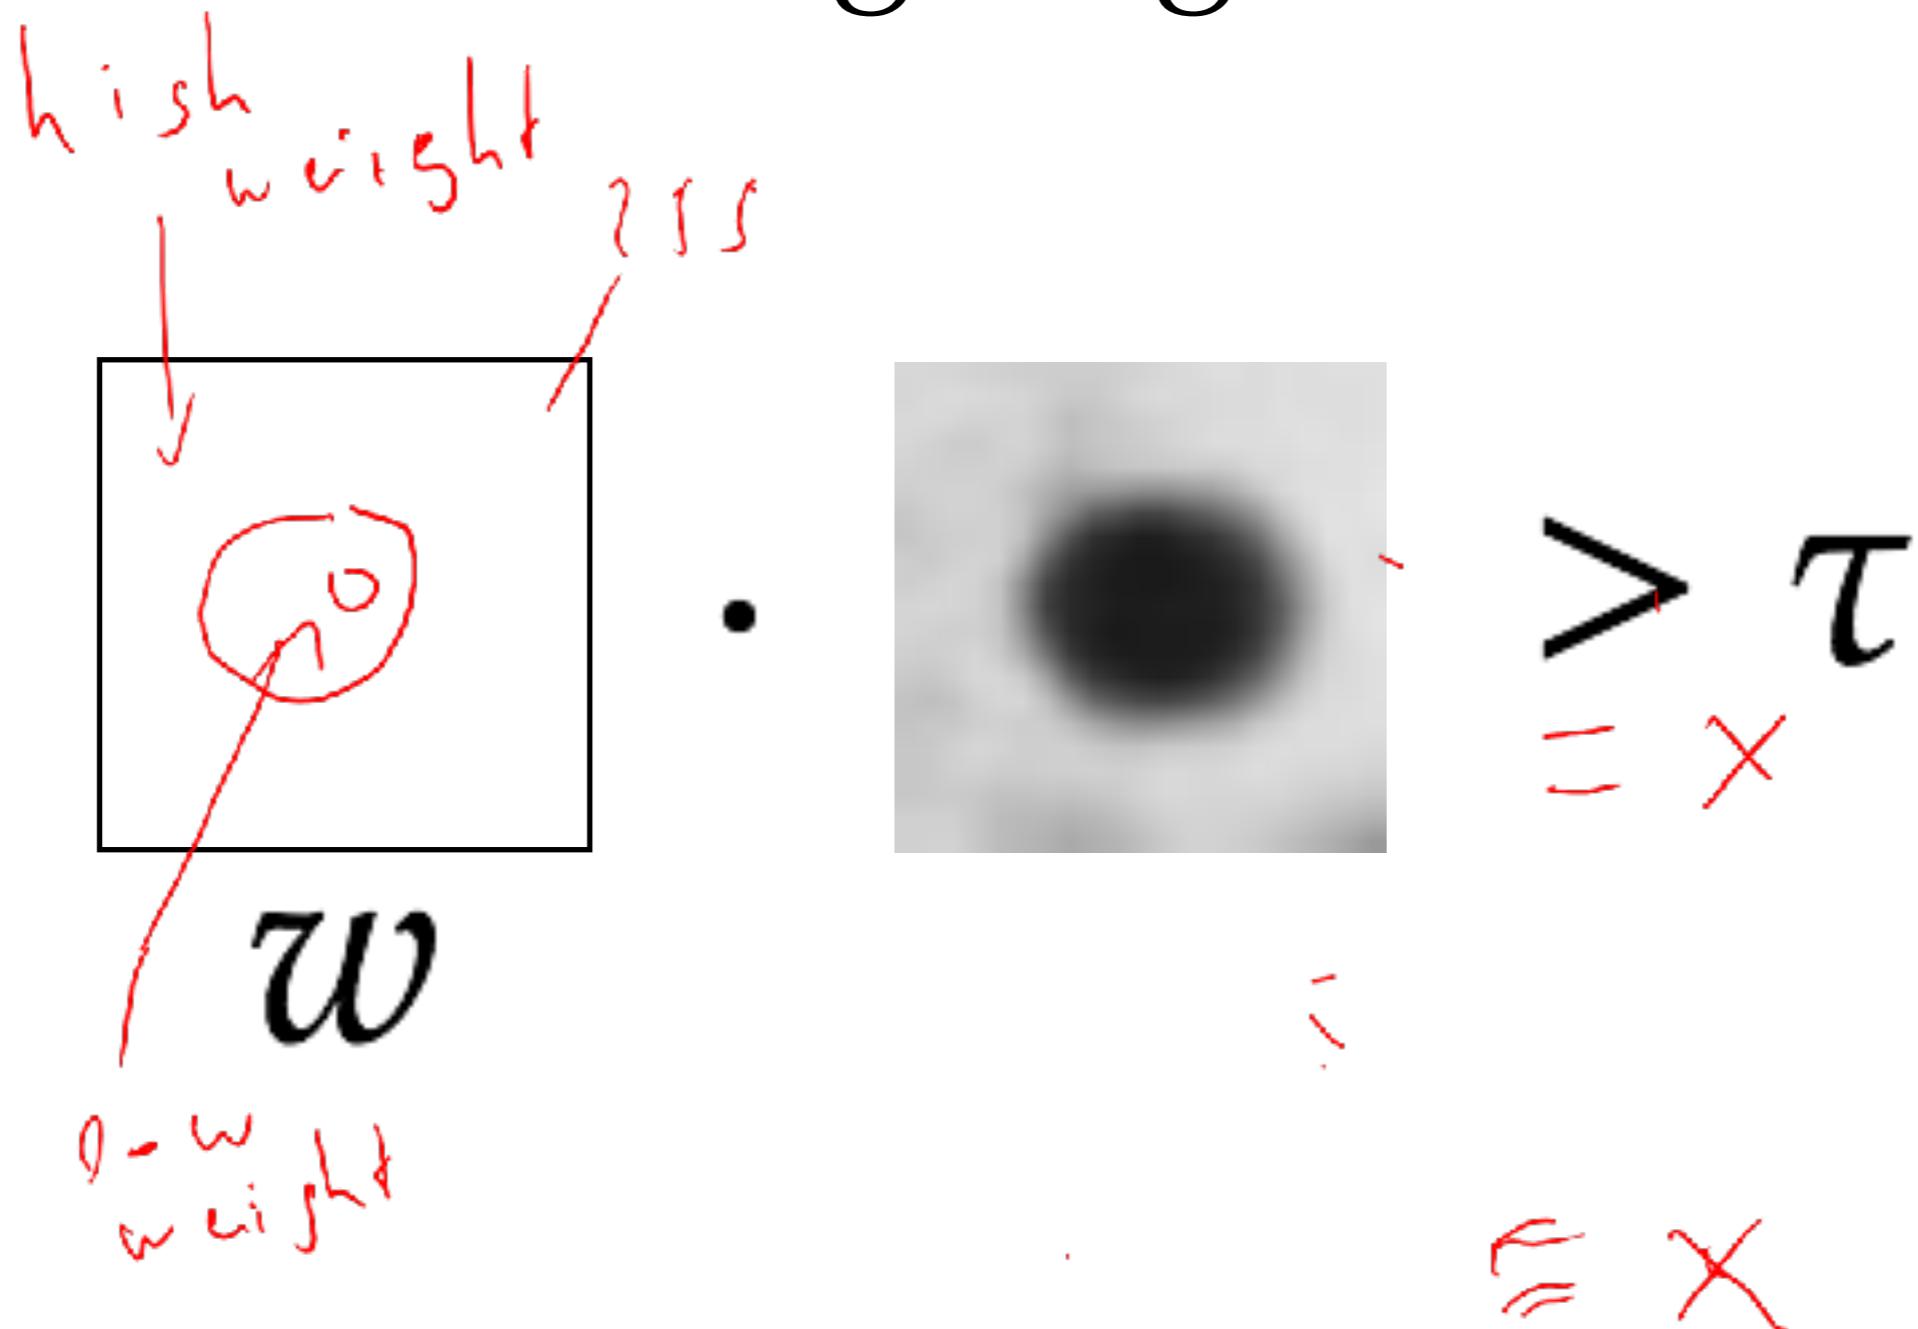
$$w \cdot$$

$$< \tau$$

$$w \cdot$$

$$< \tau$$

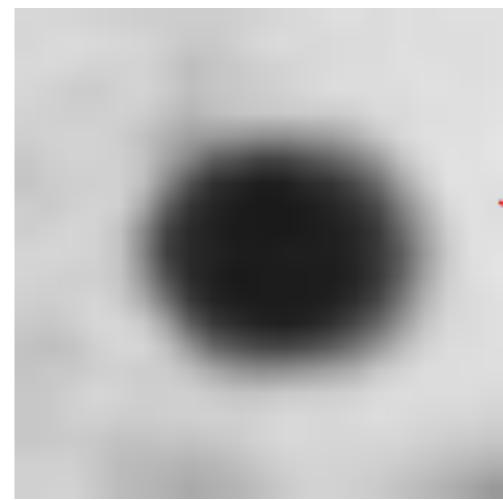
Designing a Classifier



Designing a Classifier



.



$> \tau$
 $\equiv X$



$\leq \tau$
 $\equiv X$

Designing a Better Classifier

$$w = \frac{1}{N} (w^l - \mu_{w^l} I)$$

$$w^l = \begin{array}{|c|c|c|}\hline x & x & x \\ \hline x & o & x \\ \hline x & x & x \\ \hline\end{array}$$

$$x = 252$$

$$w = \frac{1}{3} \left(\begin{array}{|c|c|c|}\hline x & x & x \\ \hline x & o & x \\ \hline x & x & x \\ \hline\end{array} - \begin{array}{|c|c|c|}\hline w^l & \dots \\ \hline \vdots & \vdots & \vdots \\ \hline\end{array} \right) = \frac{1}{3} \begin{array}{|c|c|c|}\hline 28 & \dots & \\ \hline \vdots & -224 & \\ \hline 28 & \dots & \\ \hline\end{array}$$

mean: $\mu_{w^l} = 224$

Designing a Better Classifier

1	28	28	28
2	21	-224	28
3	28	28	28

w

x	x	x
x	0	x
x	x	x

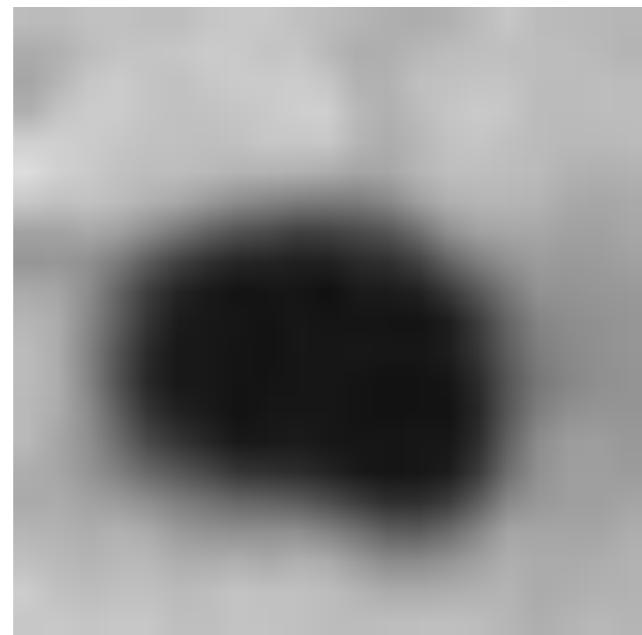
$$x = 252$$

y	y	y
y	y	y
y	y	y

$$y = 255$$

Choosing A Template

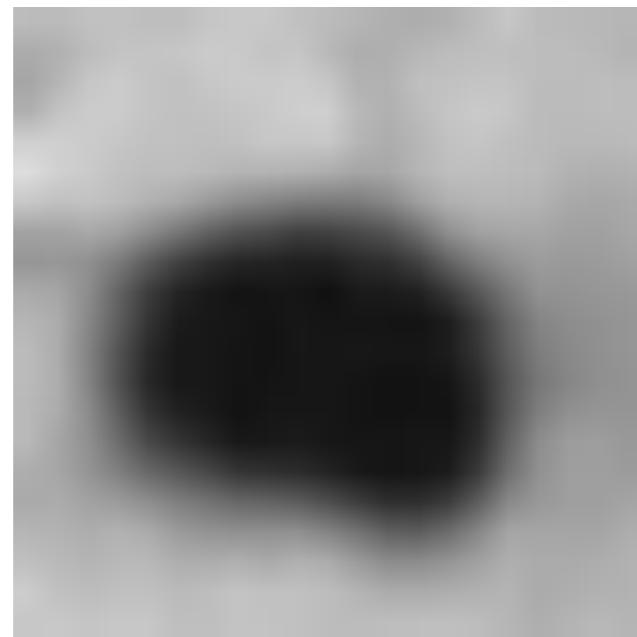
$w =$



some nucleus

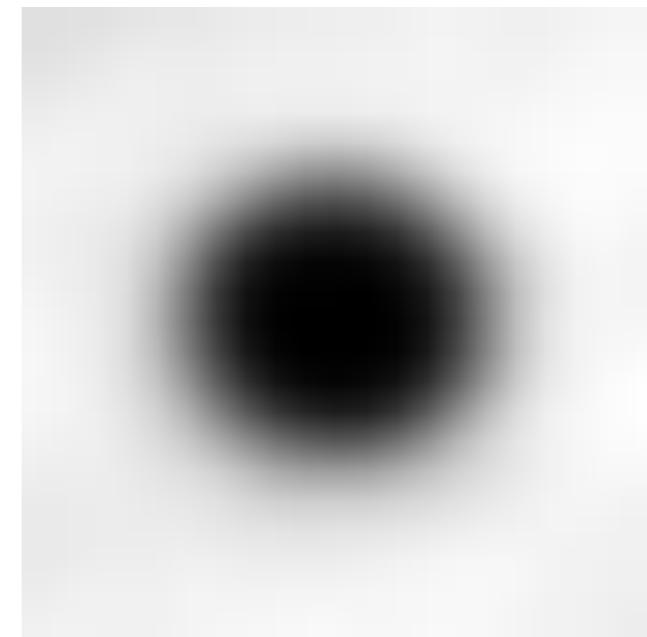
Choosing A Template

$w =$



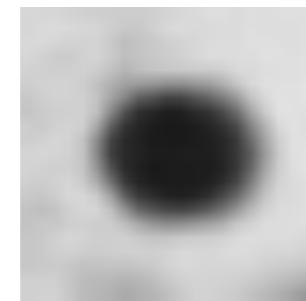
some nucleus

$w =$

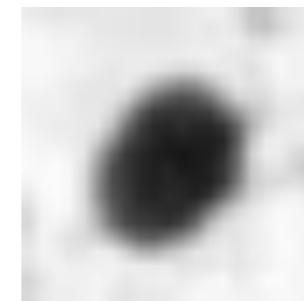


mean nucleus

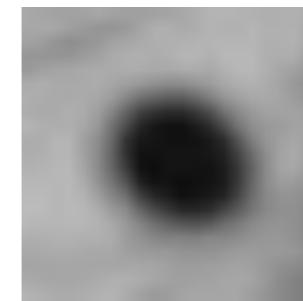
Choosing a Threshold



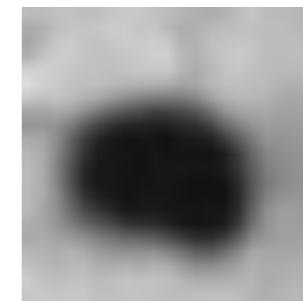
1.72



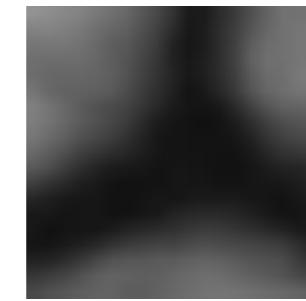
1.88



1.26



1.45



0.56



-0.07

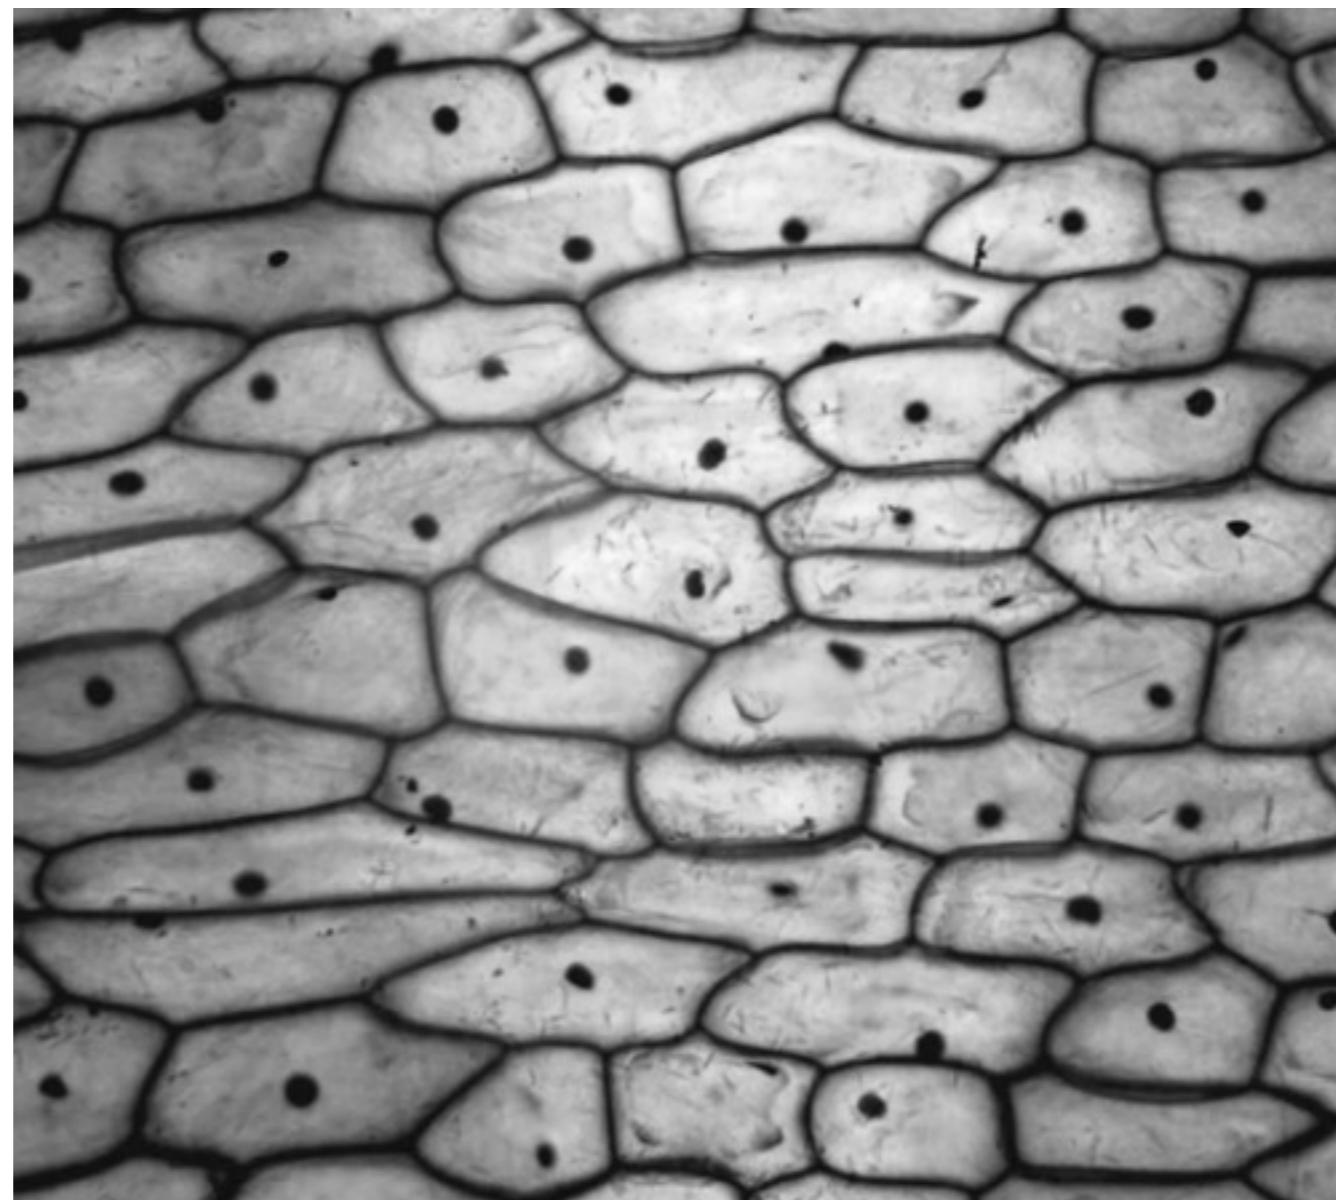


-0.06

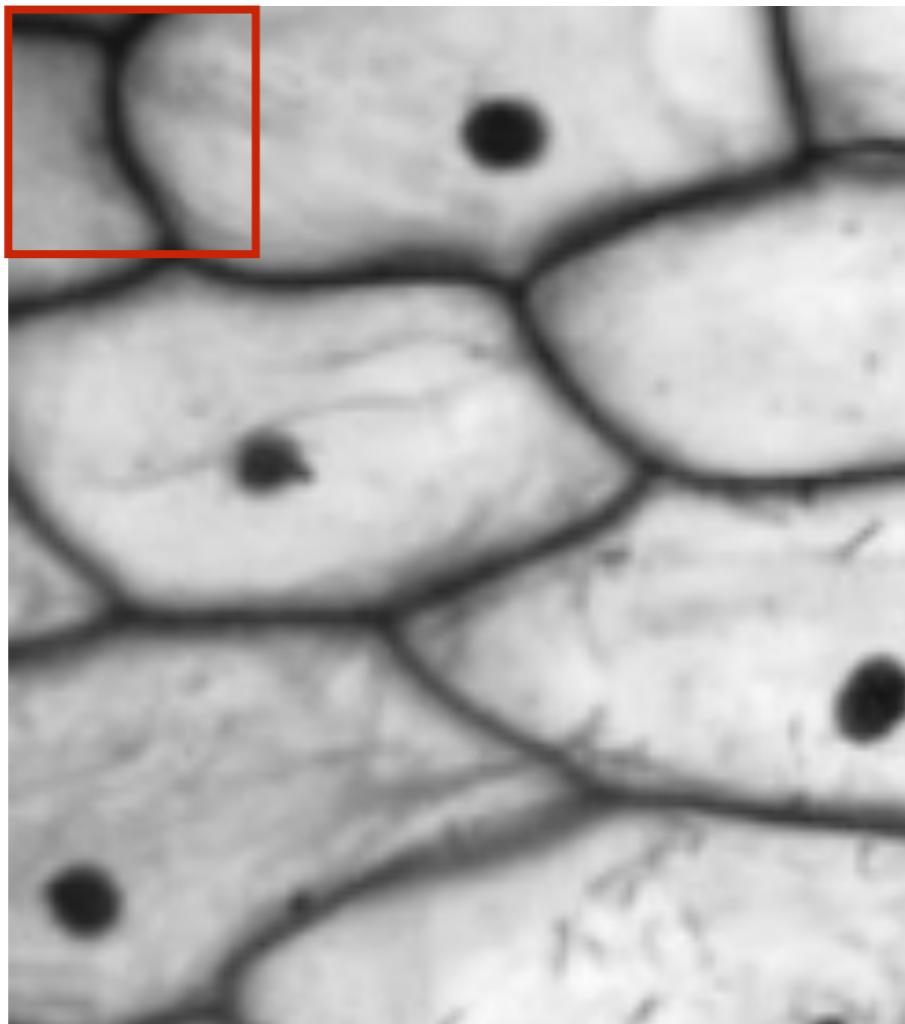


0.75

Count the Nuclei

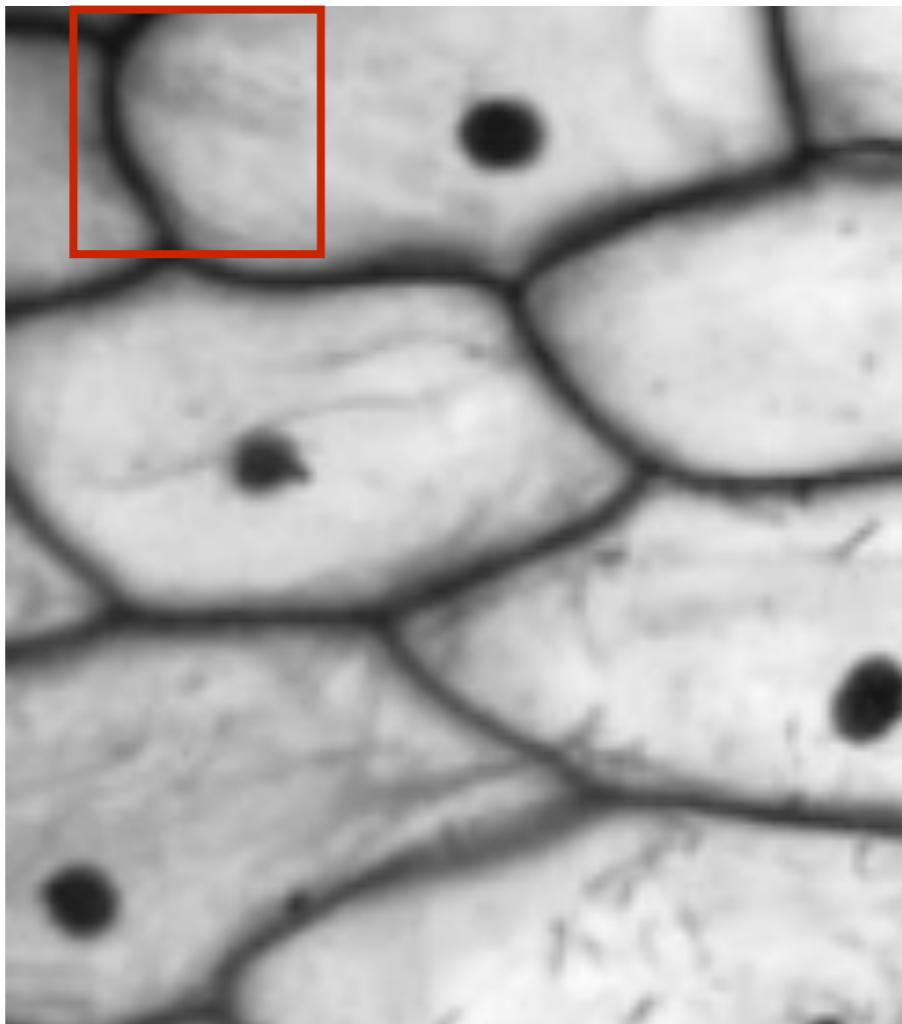


Sliding Window Classification



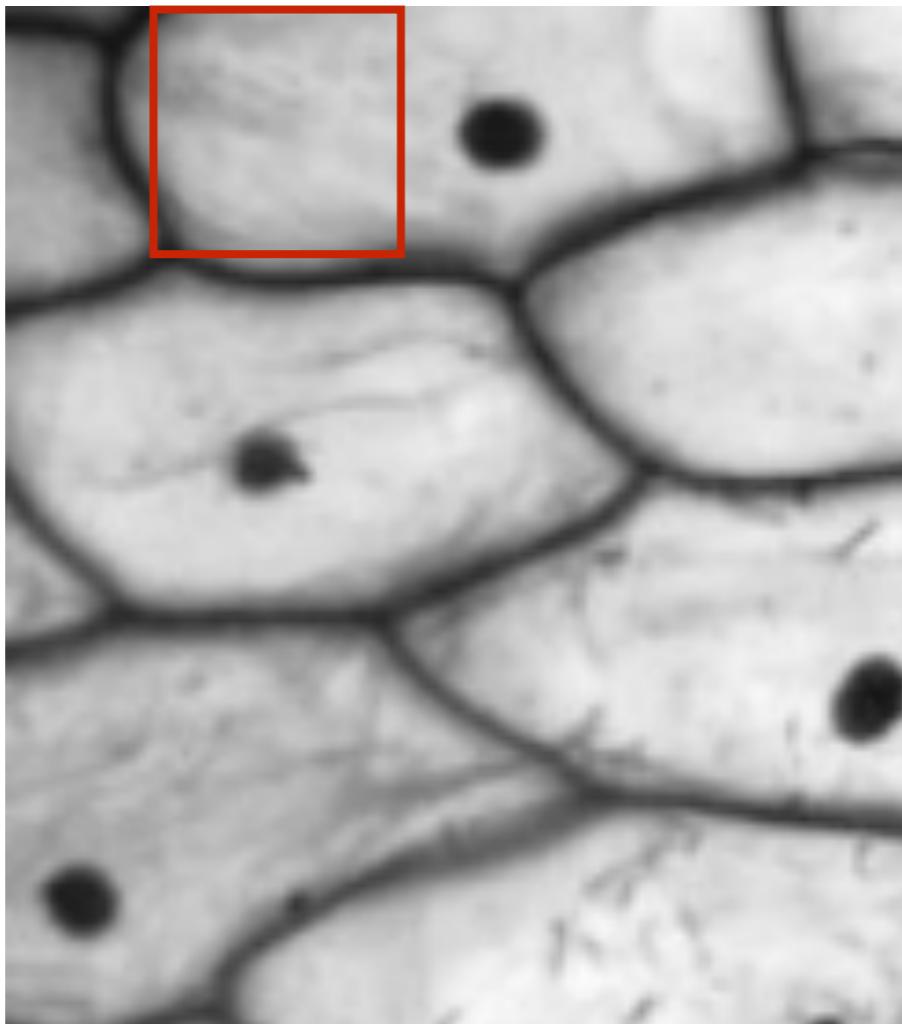
- $w < \tau$

Sliding Window Classification



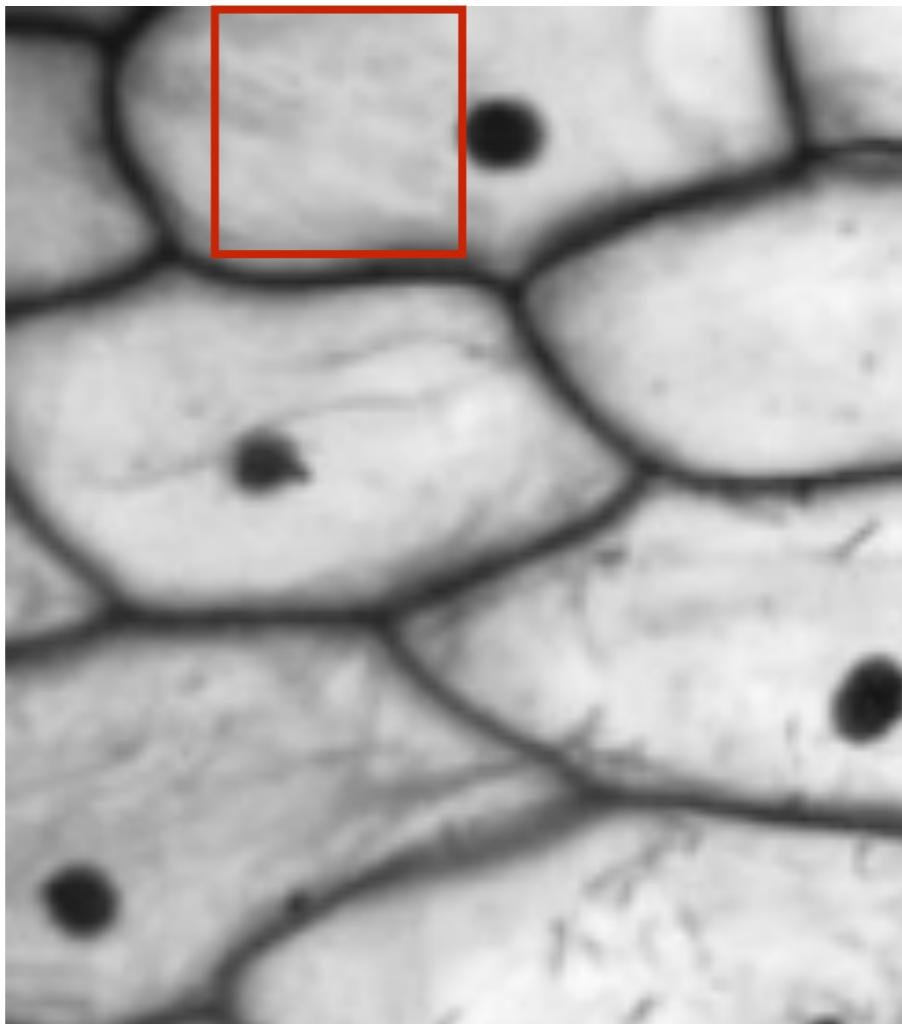
$$\cdot w < \tau$$

Sliding Window Classification



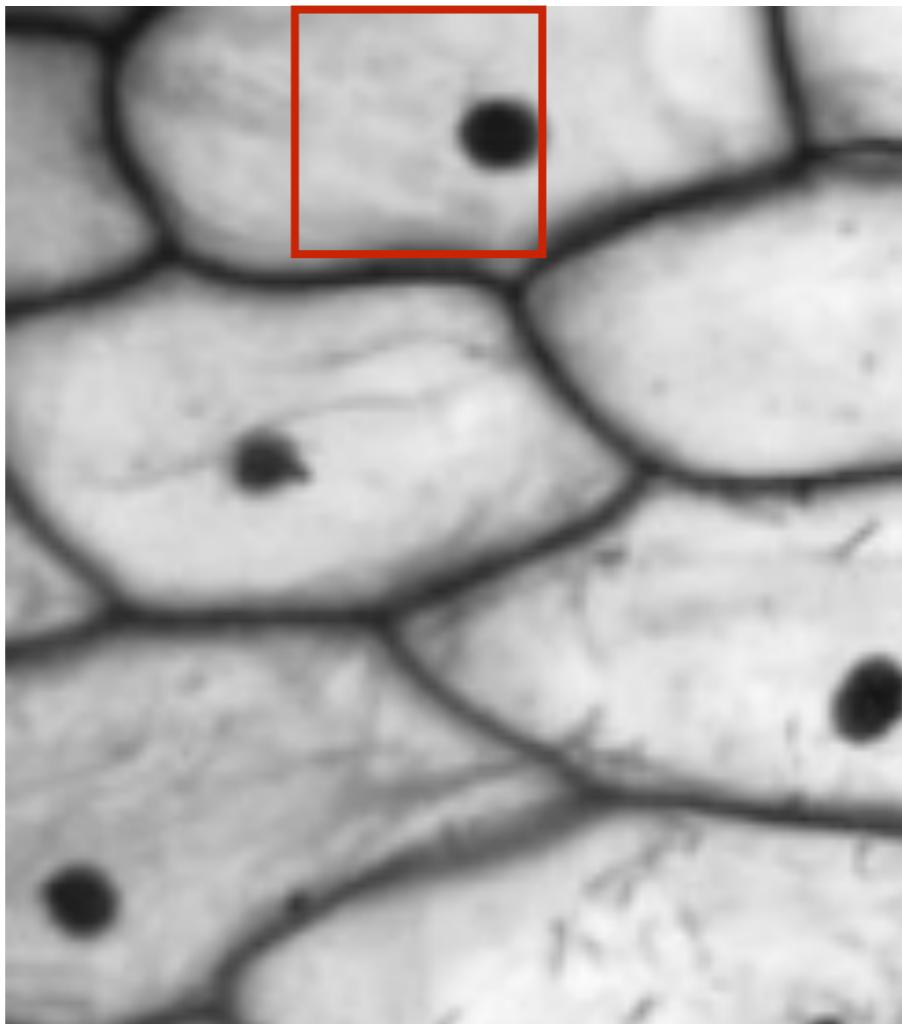
- $w < \tau$

Sliding Window Classification



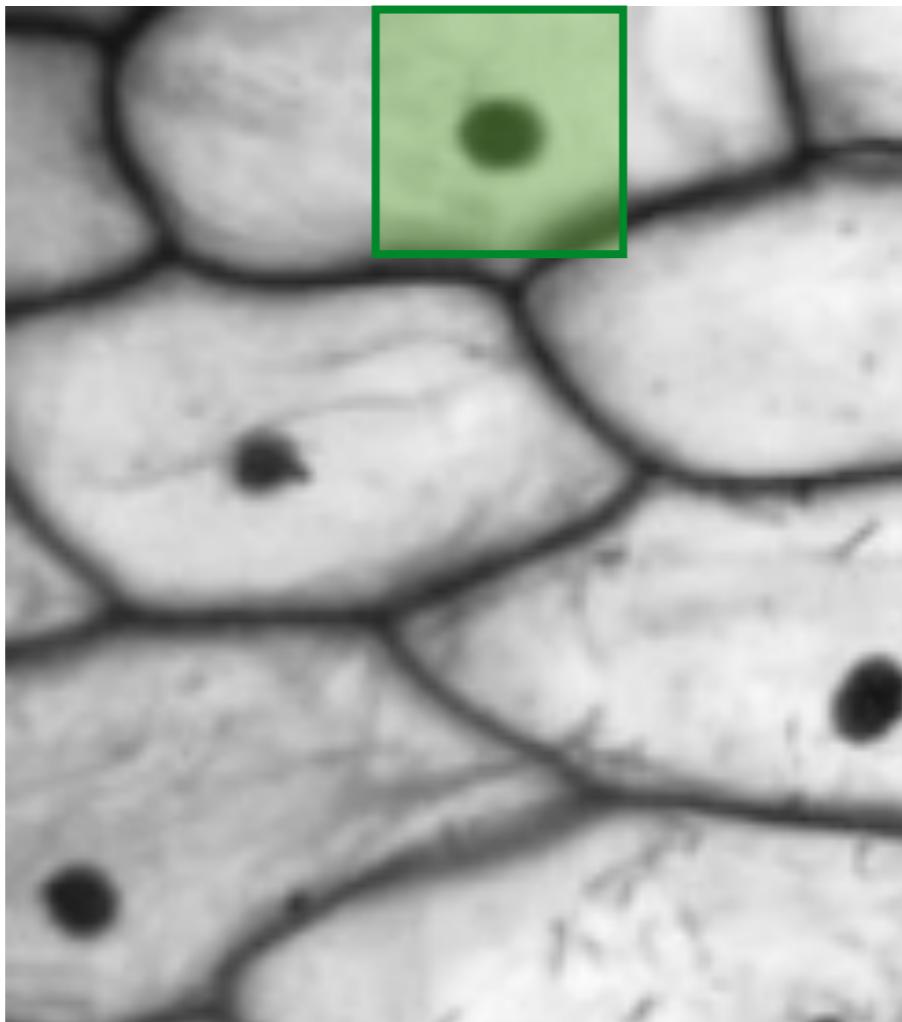
- $w < \tau$

Sliding Window Classification



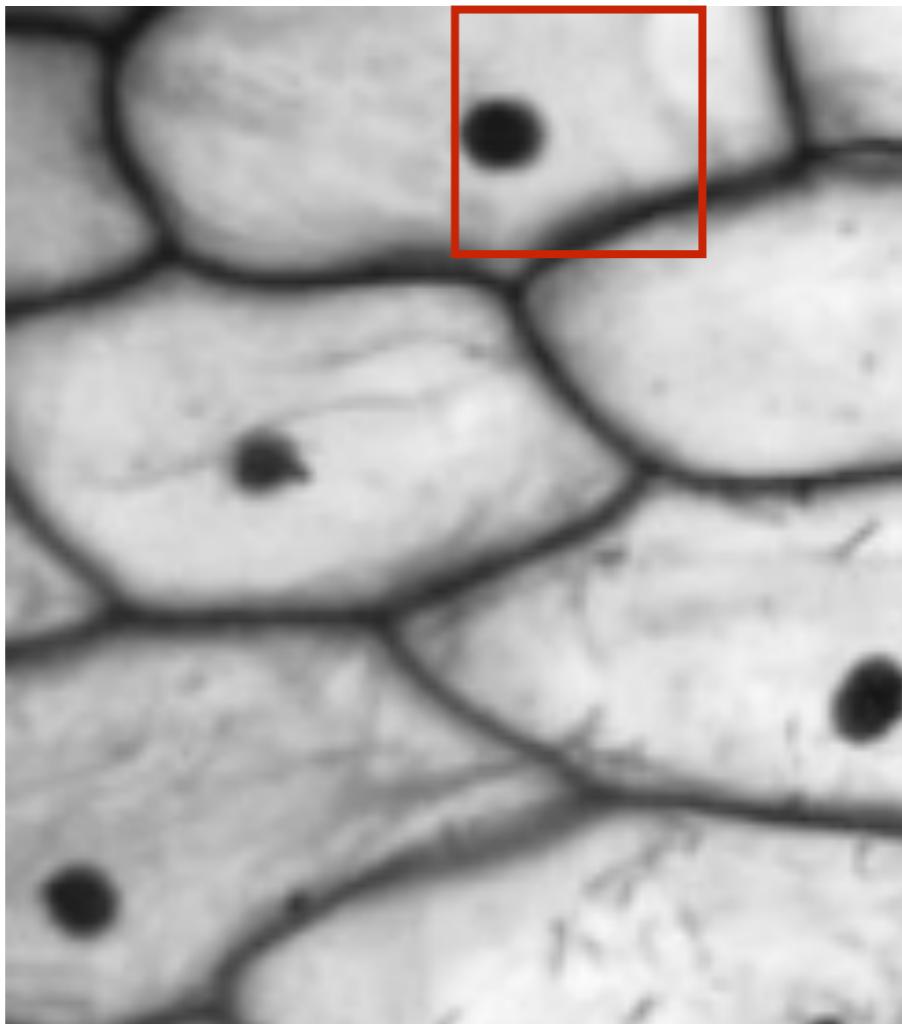
- $w < \tau$

Sliding Window Classification



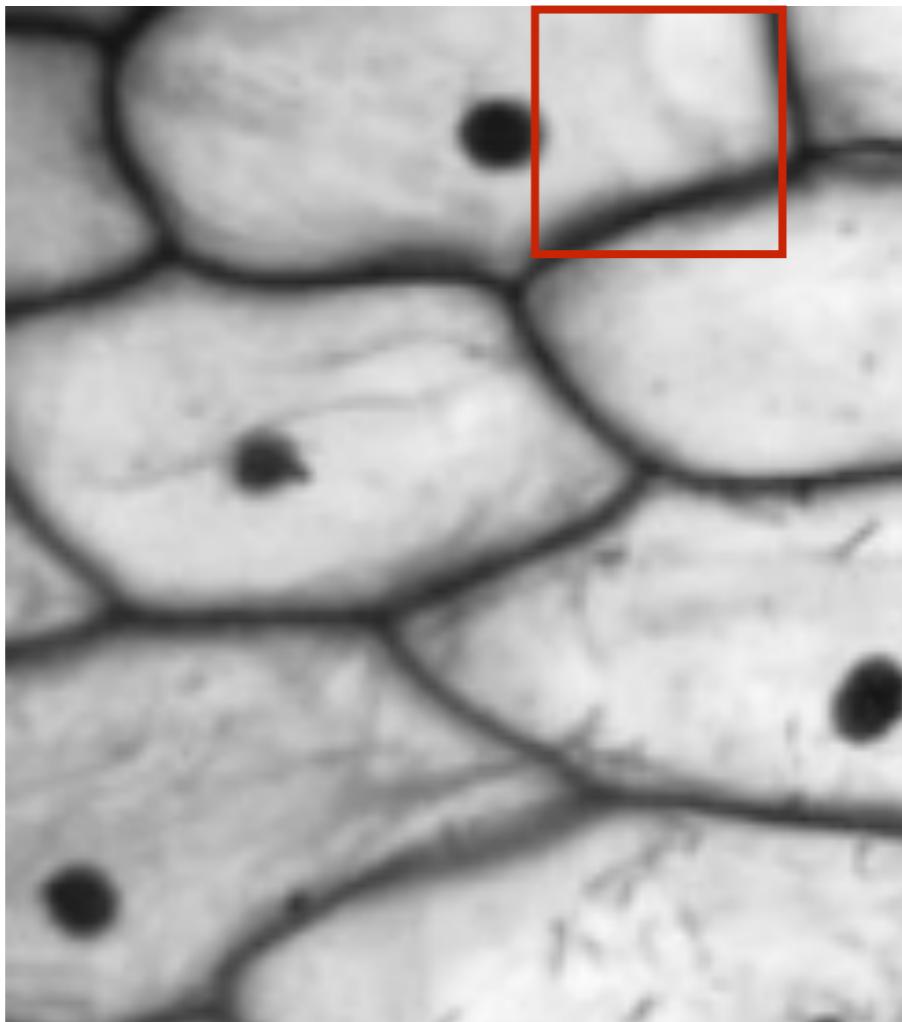
- $w > \tau$

Sliding Window Classification



- $w < \tau$

Sliding Window Classification



$$\cdot w < \tau$$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

			0			

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0						

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1				

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1				

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1	2			

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1	2	1	0	
	1	5	6	4	1	
	2	6				

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1	2	1	0	
	1	5	6	4	1	
	2	6	8			

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1	2	1	0	
	1	5	6	4	1	
	2	6	8	4	1	
	1	4	4	2	0	
	0	1	1	0	0	

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

	0	1	2	1	0	
	1	5	6	4	1	
	2	6	8	4	1	
	1	4	4	2	0	
	0	1	1	0	0	

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0							
0	0	0	0	0	0	0			
0	0	0	0	0	0	0			
0	0	5	10	5	0	0			
0	0	10	10	5	0	0			
0	0	5	5	0	0	0			
0	0	0	0	0	0	0			
0	0	0	0	0	0	0			

I

0									
	0	1	2	1	0				
	1	5	6	4	1				
	2	6	8	4	1				
	1	4	4	2	0				
	0	1	1	0	0				

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0						
0	0	0	0	0	0	0		
0	0	0	0	0	0	0		
0	0	5	10	5	0	0		
0	0	10	10	5	0	0		
0	0	5	5	0	0	0		
0	0	0	0	0	0	0		
0	0	0	0	0	0	0		

0	0							
	0	1	2	1	0			
1	5	6	4	1				
2	6	8	4	1				
1	4	4	2	0				
0	1	1	0	0				

I

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0 0 0

0	0	0	0	0	0	0
0	0	1	2	1	0	0
0	1	5	6	4	1	0
0	2	6	8	4	1	0
0	1	4	4	2	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0

$I \otimes w$

Sliding Dot Product

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

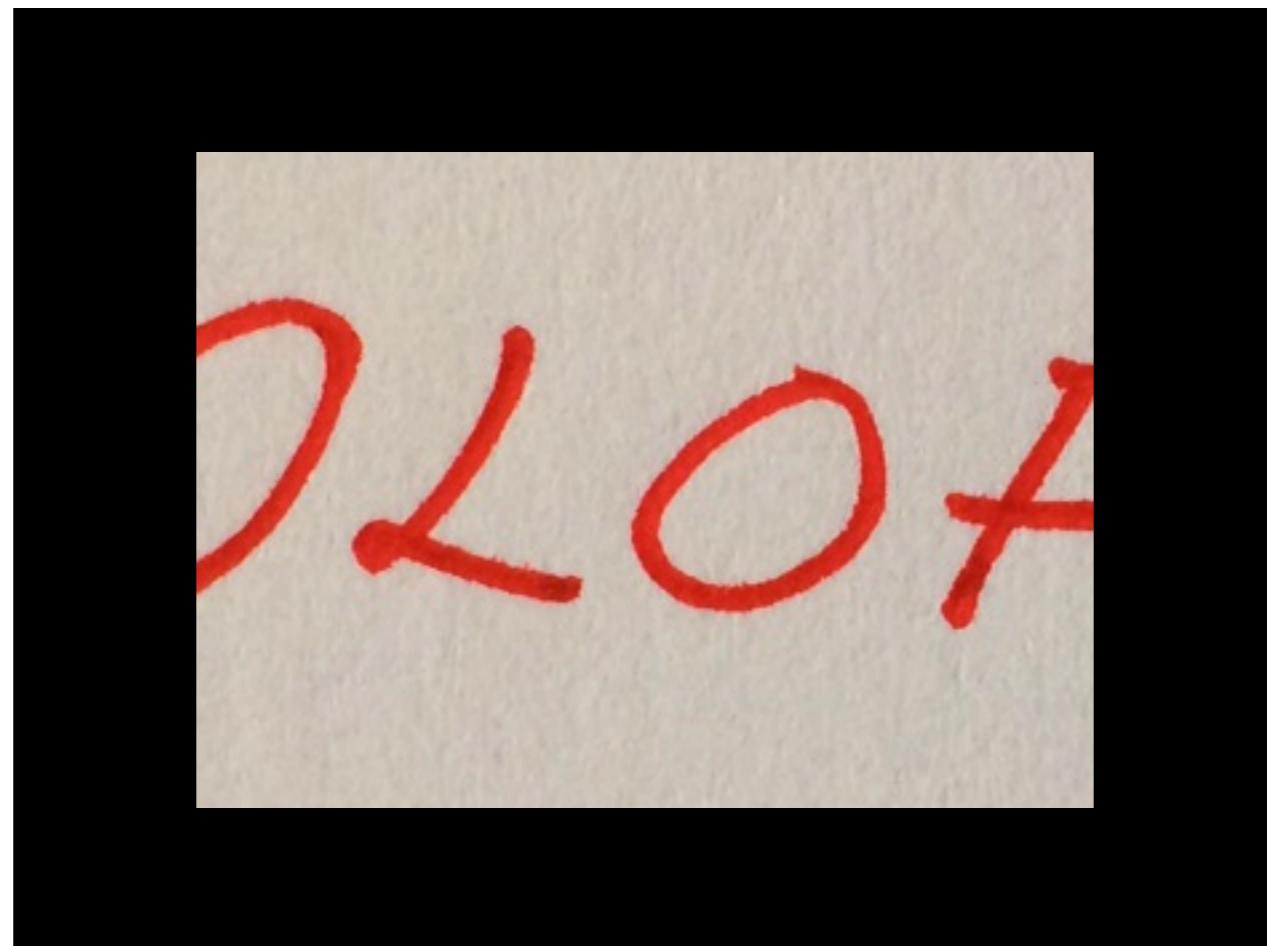
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	1	1	0	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$I \otimes w > 4$

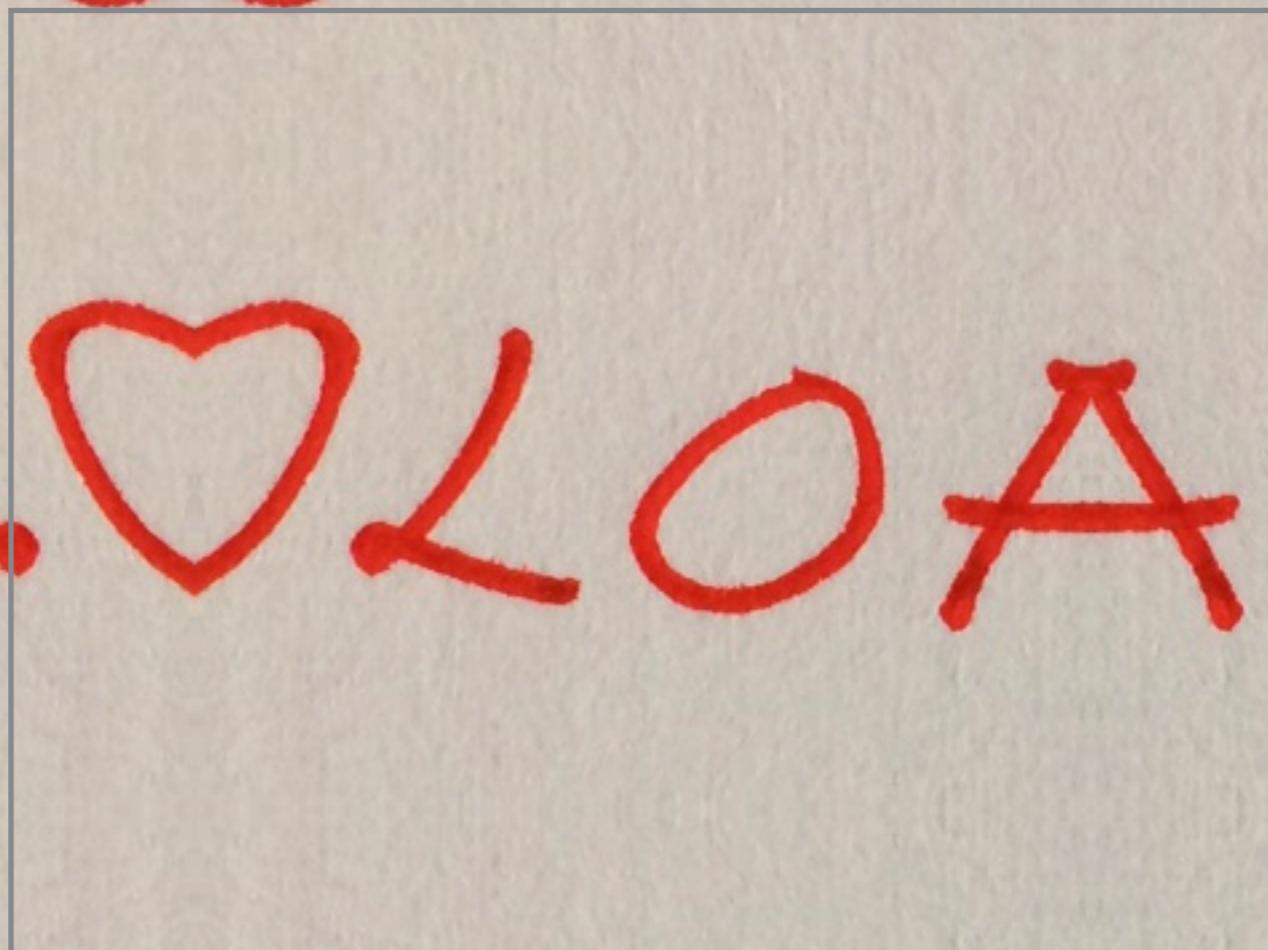
Boundary Padding



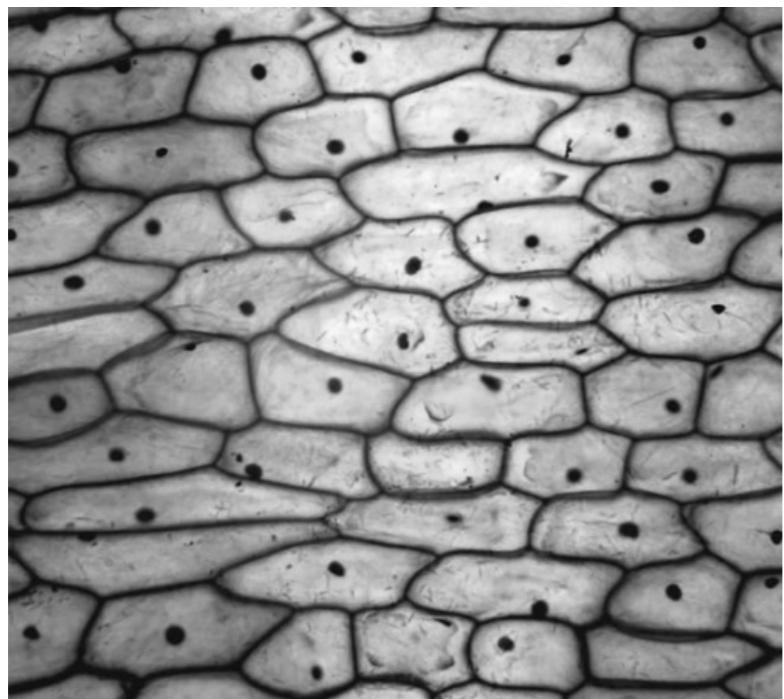
#OTTO#OTTO#OTTO
#OTTO#OTTO#OTTO

#OTTO#OTTO#OTTO

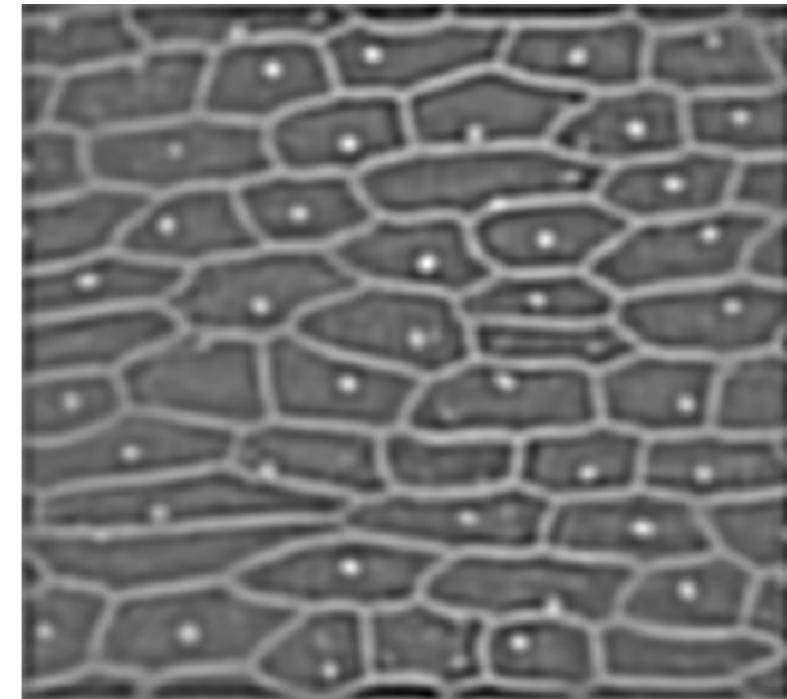
#OTTO#OTTO#OTTO



Sliding Window Classification



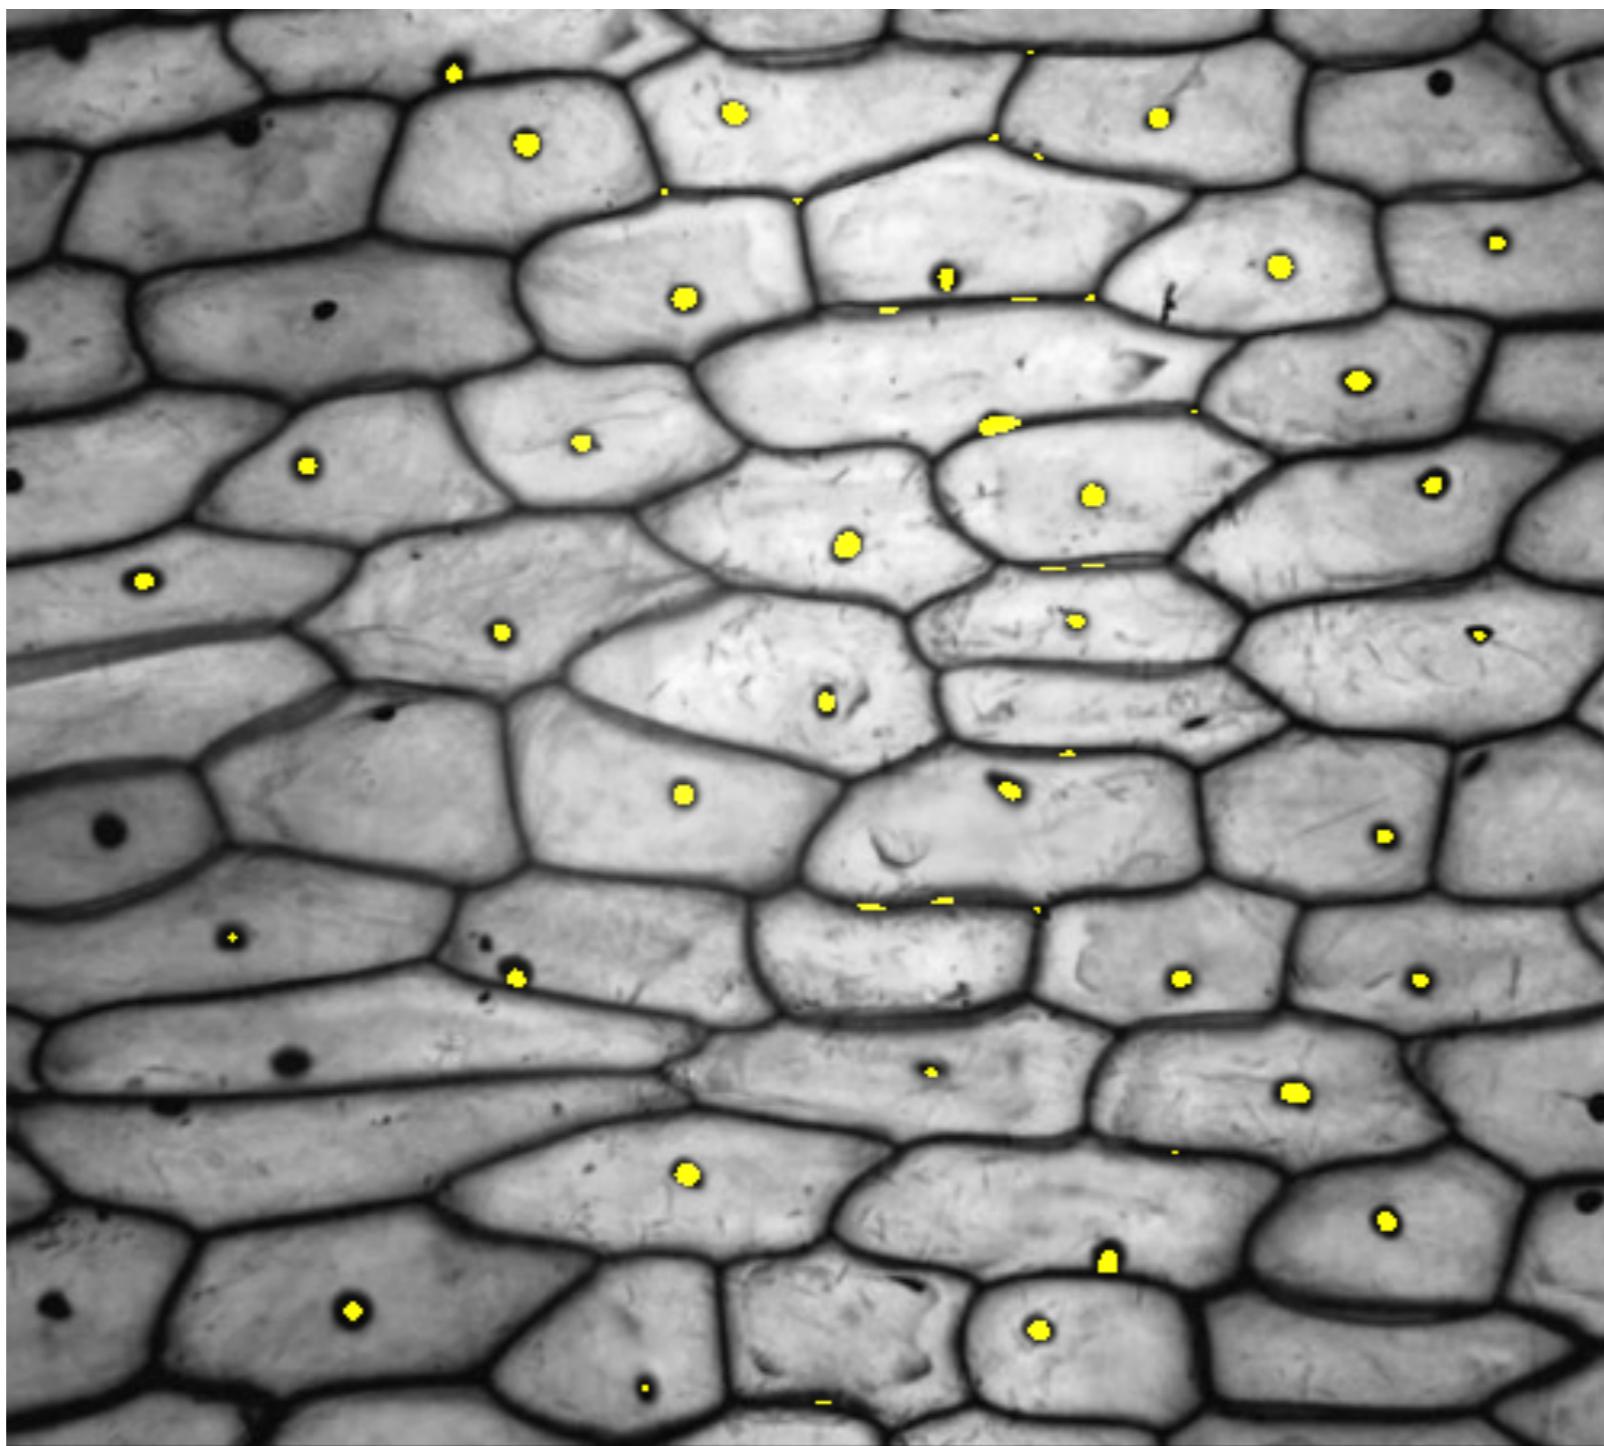
$$\otimes \quad \begin{matrix} & \bullet \\ & \circ \end{matrix} \quad =$$



Per-Pixel Classification Results



Per-Pixel Classification



Maximum Filter / Non-Maximum Suppression

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

result

Nonlinear Filters

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

0	0					

result

Nonlinear Filters

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

	0	0	0	0	0	0
0	1					

result

Nonlinear Filters

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

	0	0	0	0	0	0
	0	1	0			

result

Nonlinear Filters

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

	0	0	0	0	0	
	0	1	0	0	0	
	0	0	0	0	0	
	0	0	0	1		

result

Nonlinear Filters

Is this pixel larger than its neighbors?

1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

		0	0	0	0	0
		0	1	0	0	0
		0	0	0	0	0
		0	0	0	1	0
		0	0	0	0	0

result

Nonlinear Filters

Is this pixel larger than its neighbors?

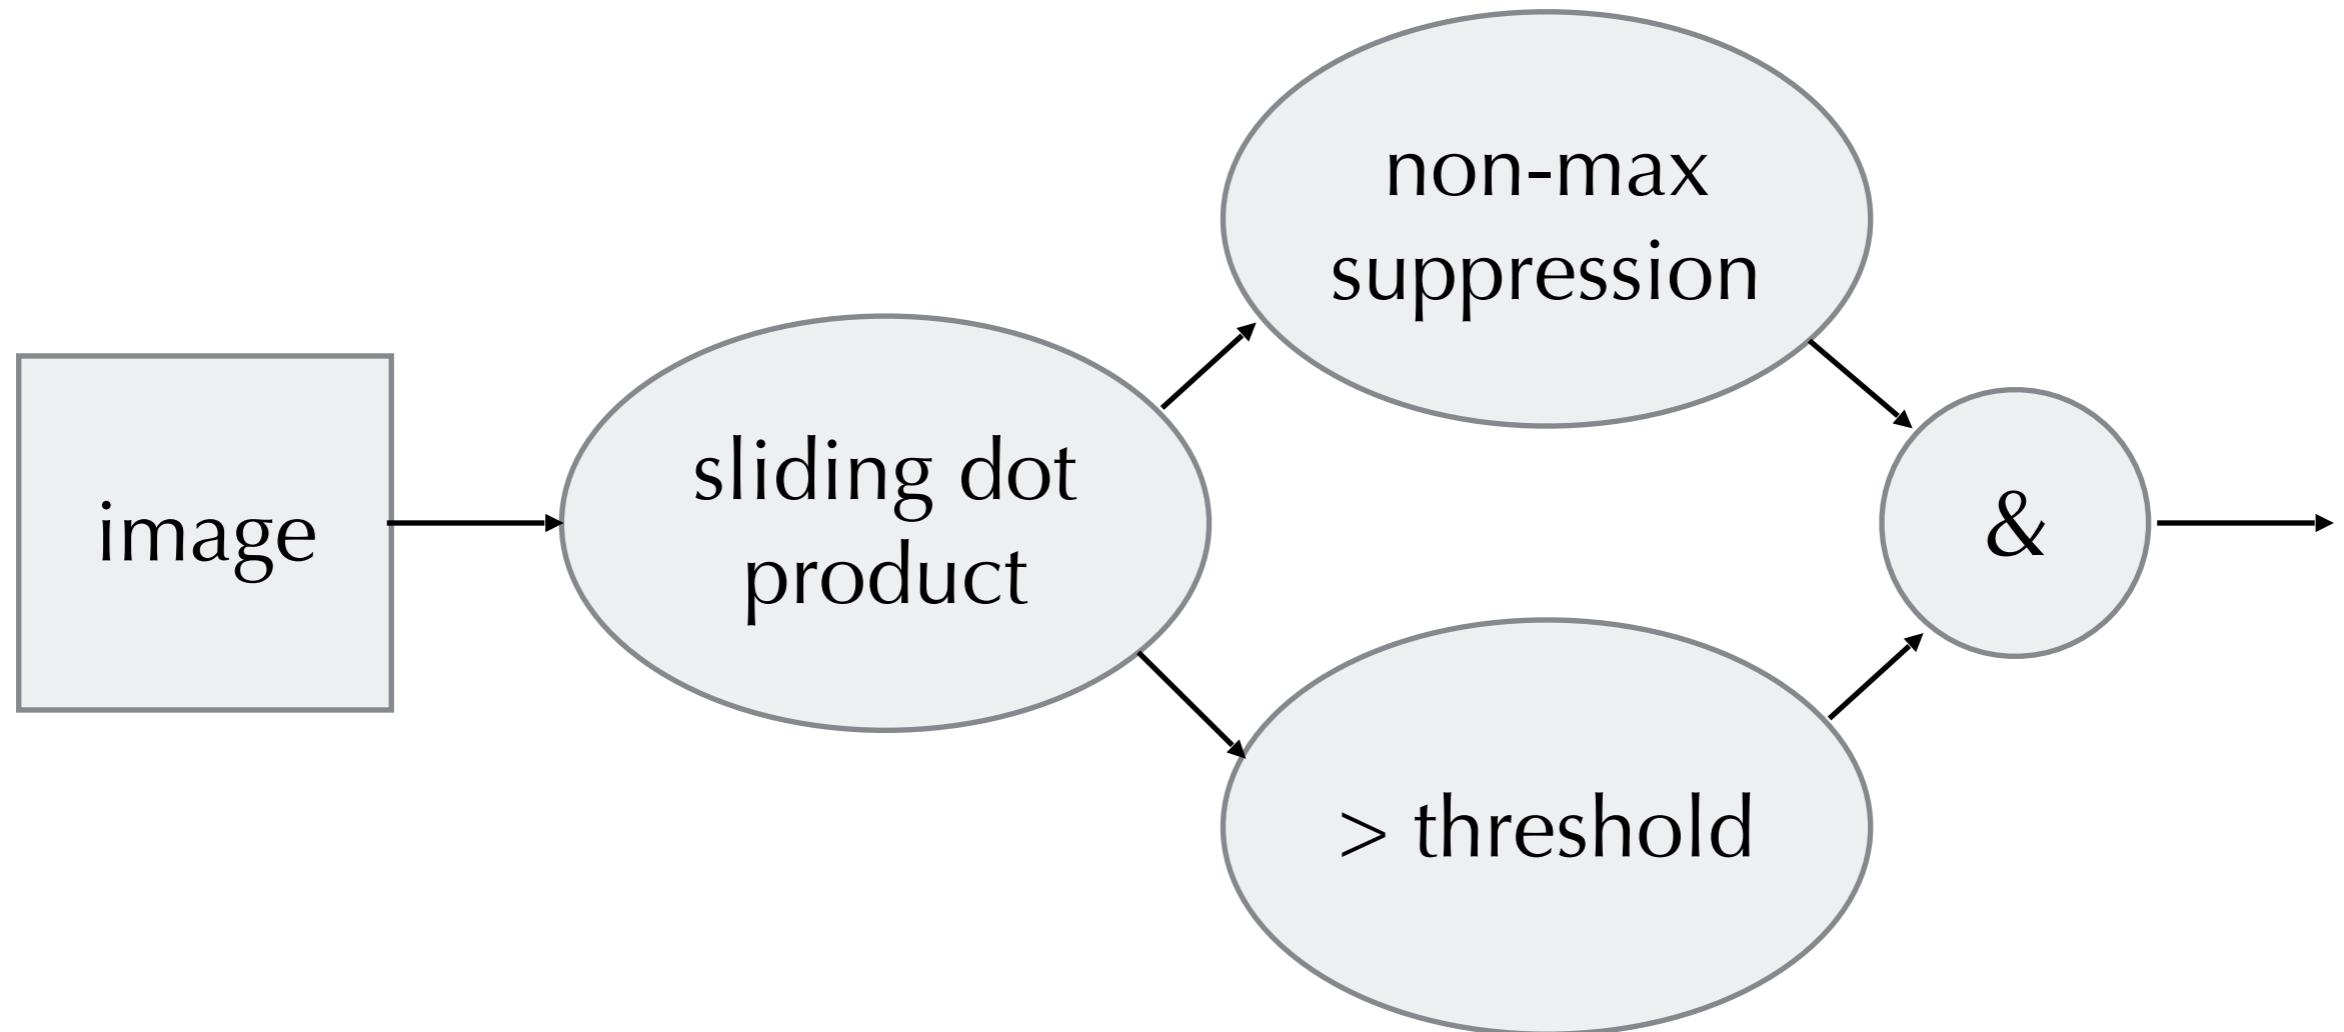
1	2	1	2	3	4	2
11	21	22	21	14	6	7
12	20	45	32	21	12	11
11	12	11	16	21	12	21
21	22	23	25	35	22	20
12	11	16	17	16	6	0
0	7	0	21	12	11	0

input

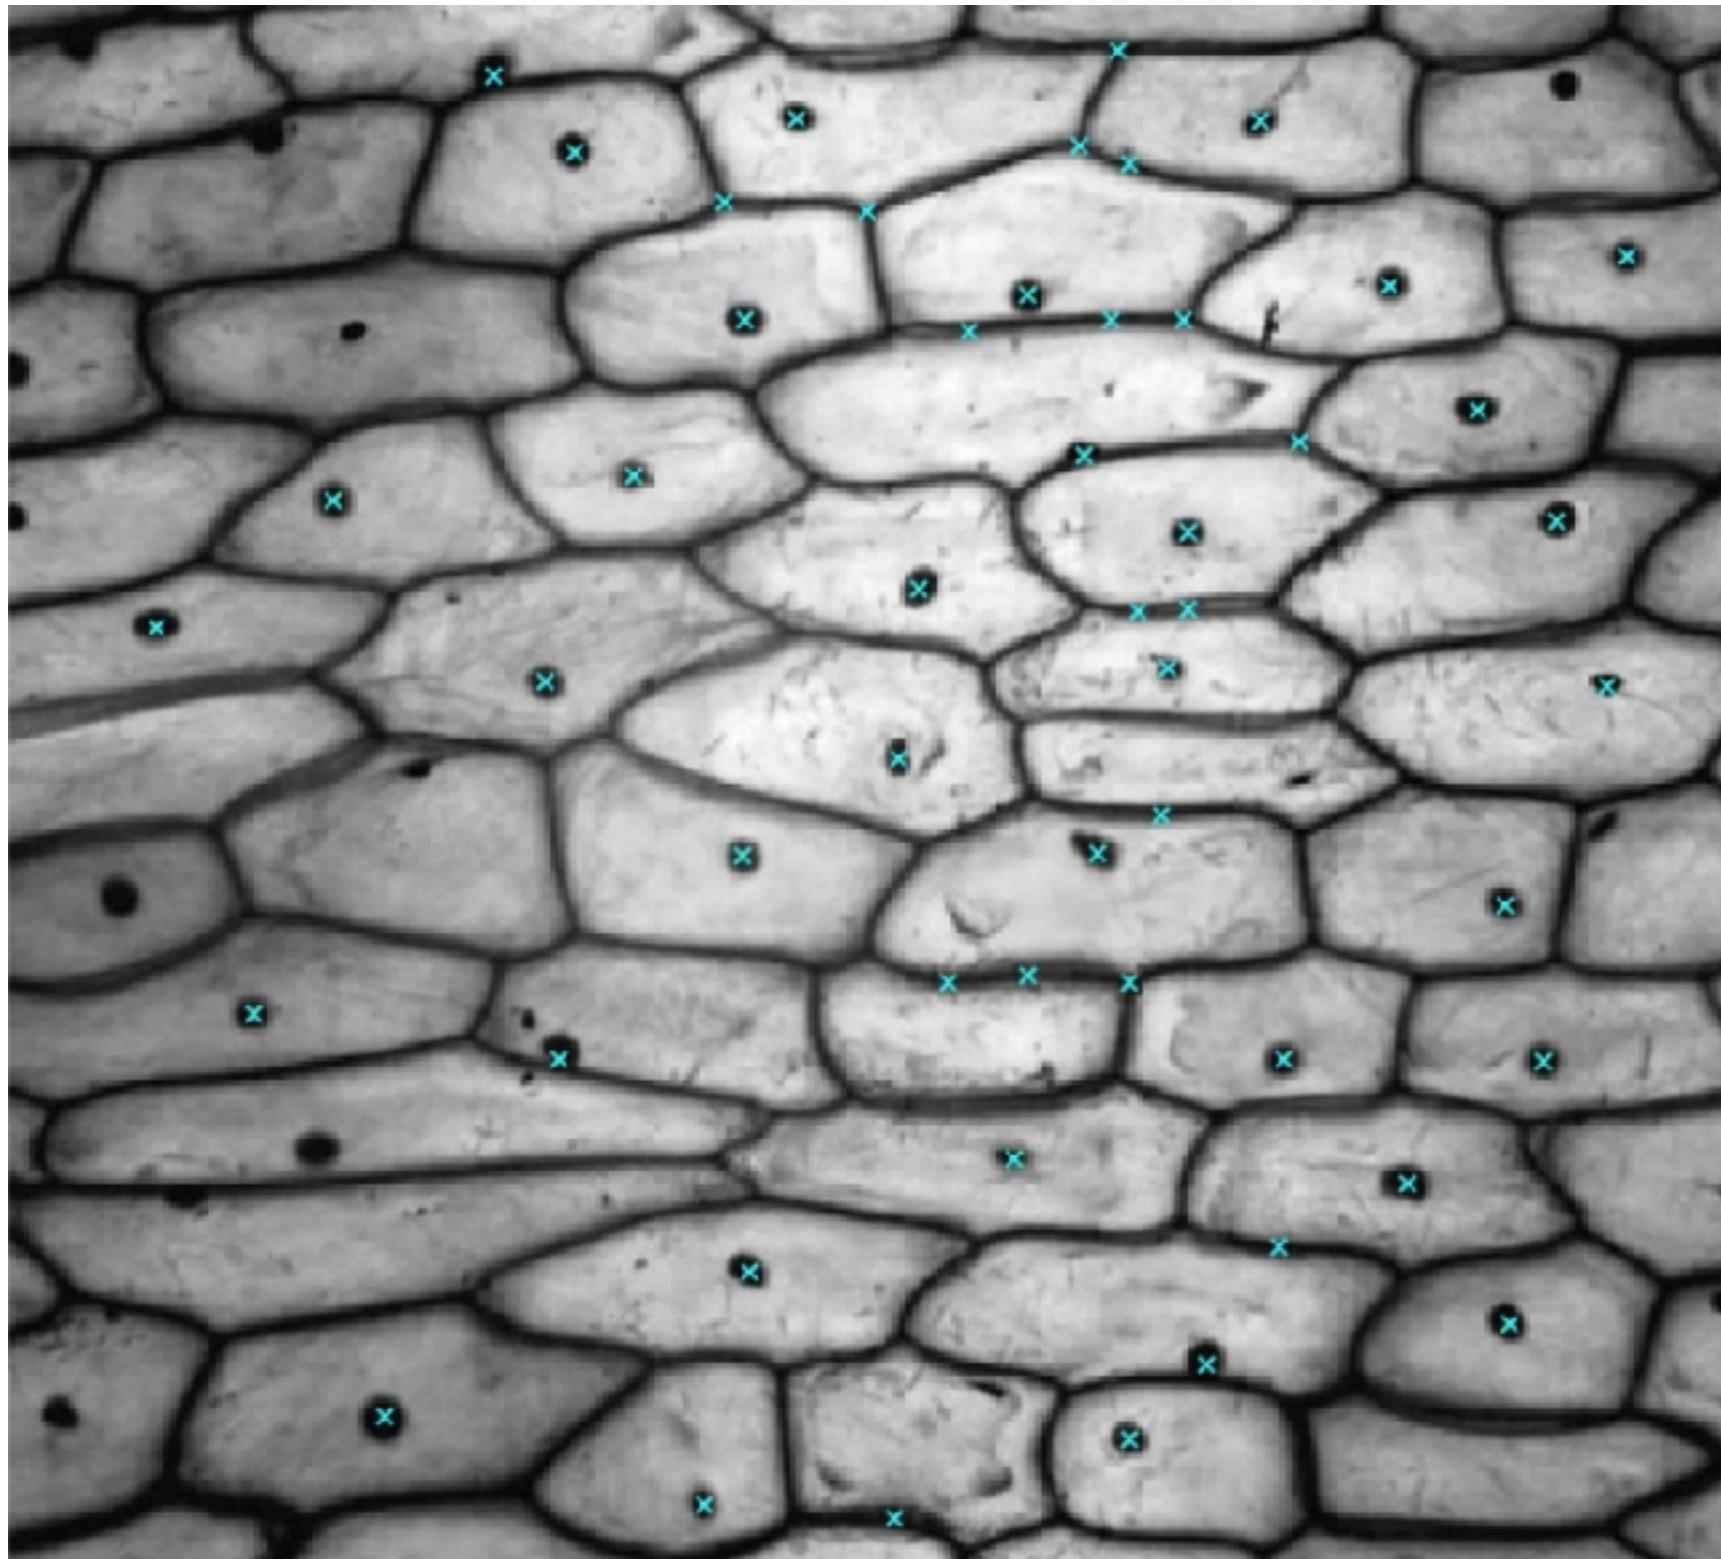
0	0	0	0	0	0	
0	1	0	0	0	0	
0	0	0	0	0	0	
0	0	0	1	0		
0	0	0	0	0	0	

result

A Sliding-Window Detector



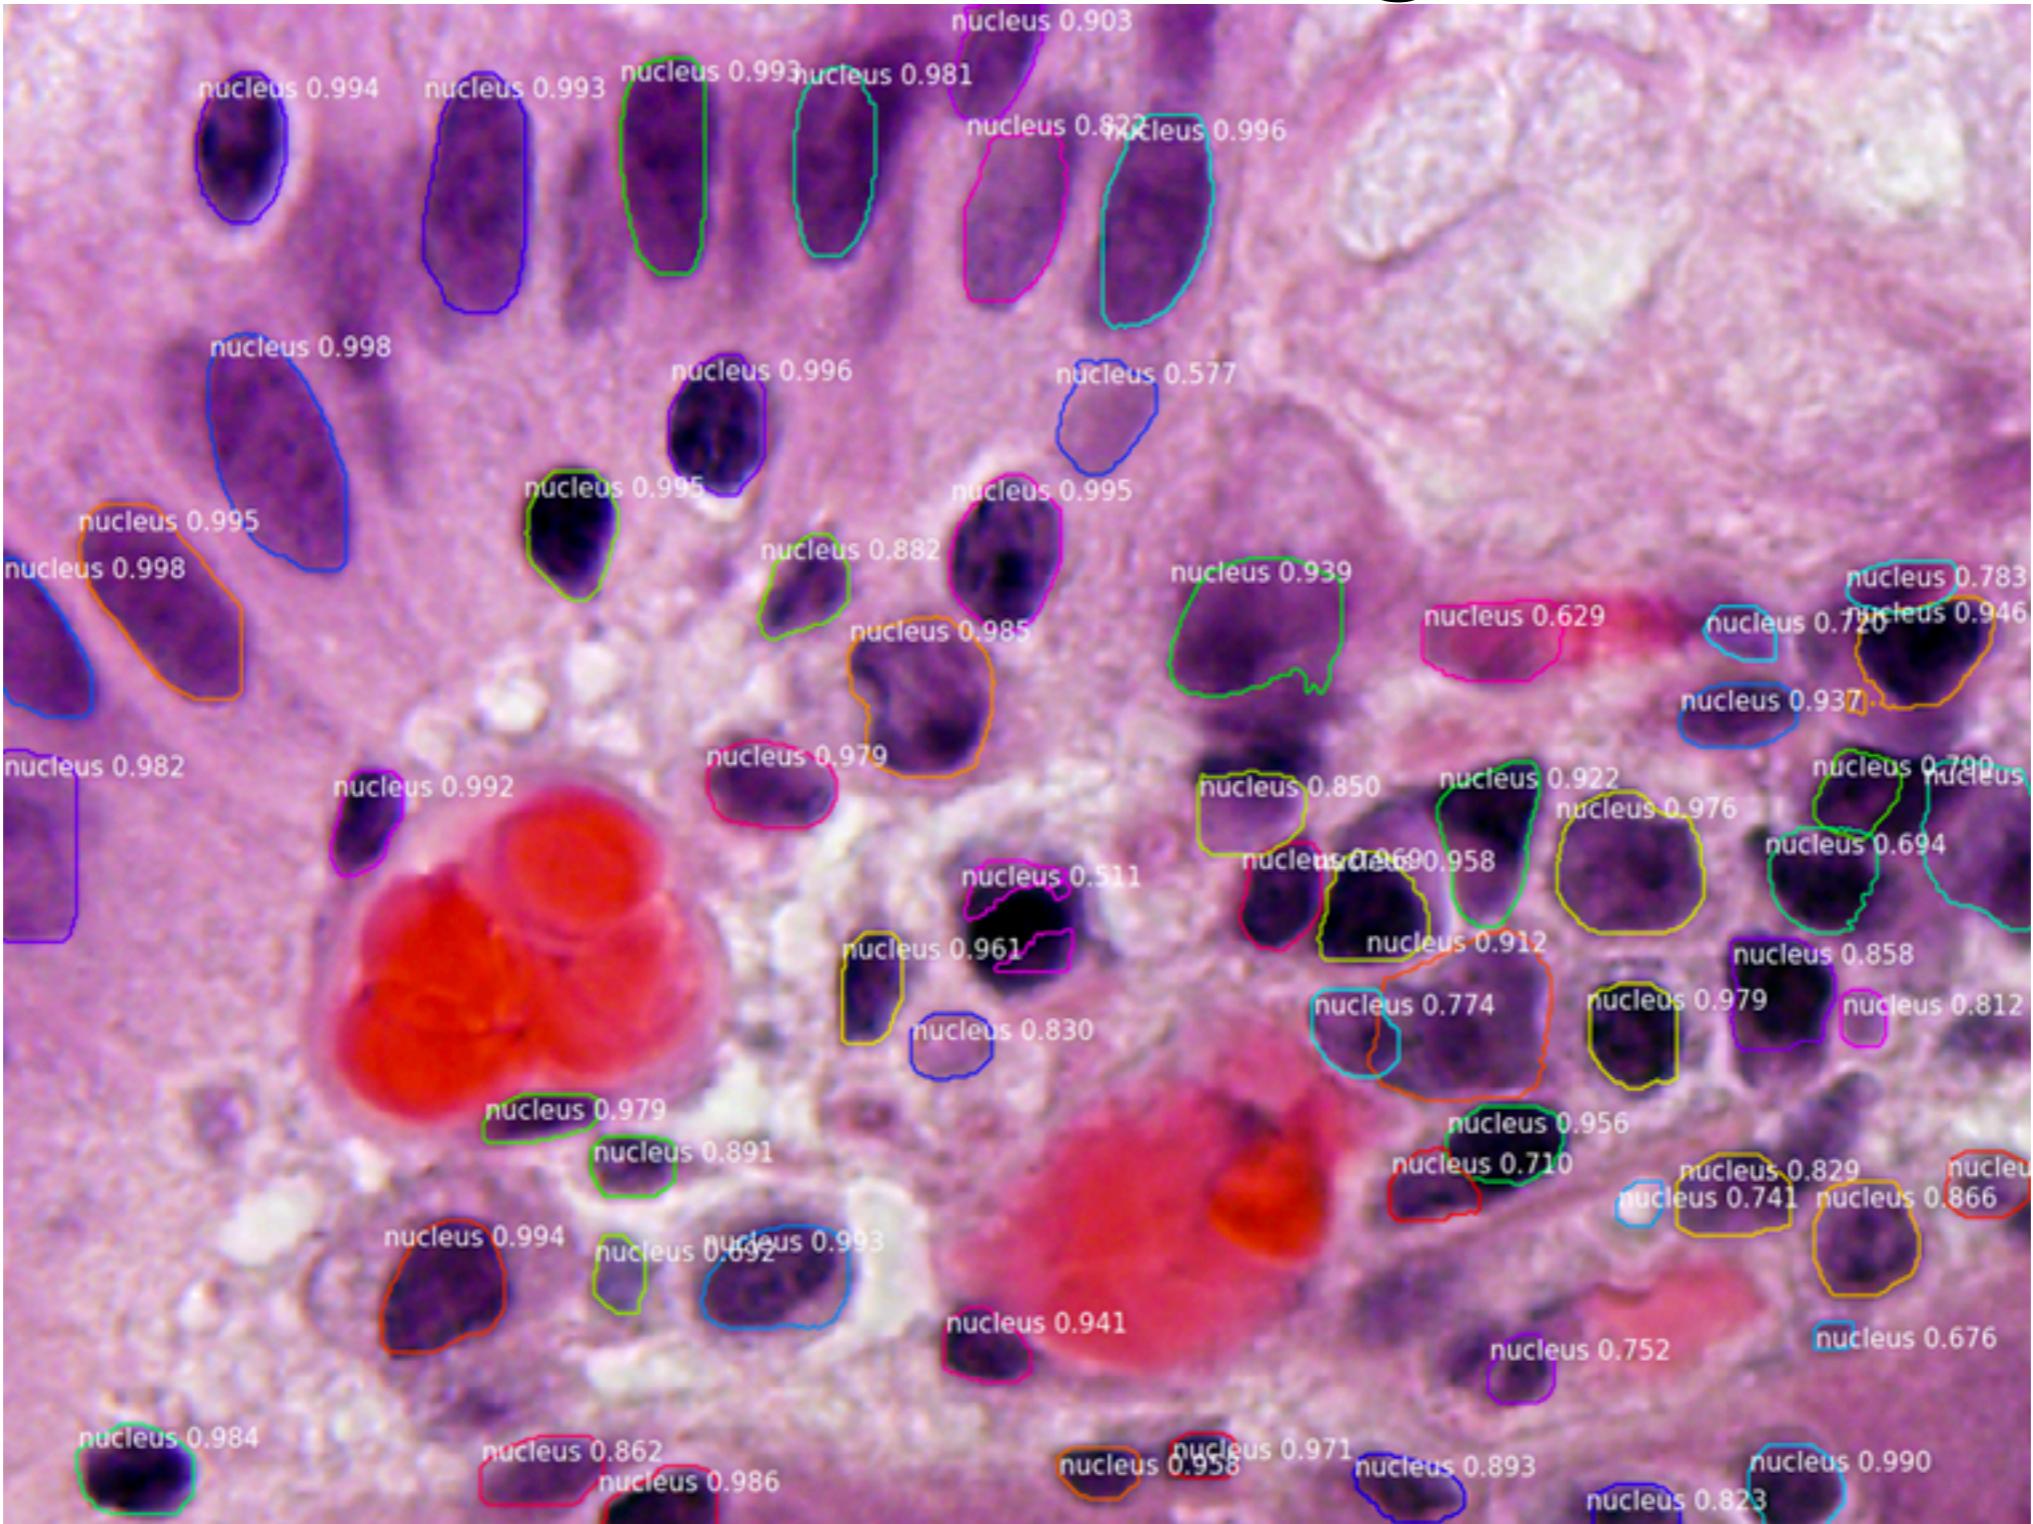
Result



Count: 52

Manual: 48

Instance Level Segmentation



Result from apply Masked R-CNN

Lessons Learned

- Main lessons from this lecture
 - Images are vectors: Addition, subtraction, dot product
 - Linear filters: Mean, dot product-based
 - Nonlinear filters: Non-maximum suppression
 - Sliding window classification

Lessons Learned

- Main lessons from this lecture
 - Images are vectors: Addition, subtraction, dot product
 - Linear filters: Mean, dot product-based
 - Nonlinear filters: Non-maximum suppression
 - Sliding window classification
- Next lecture: More filtering, gradients, multi-scale processing