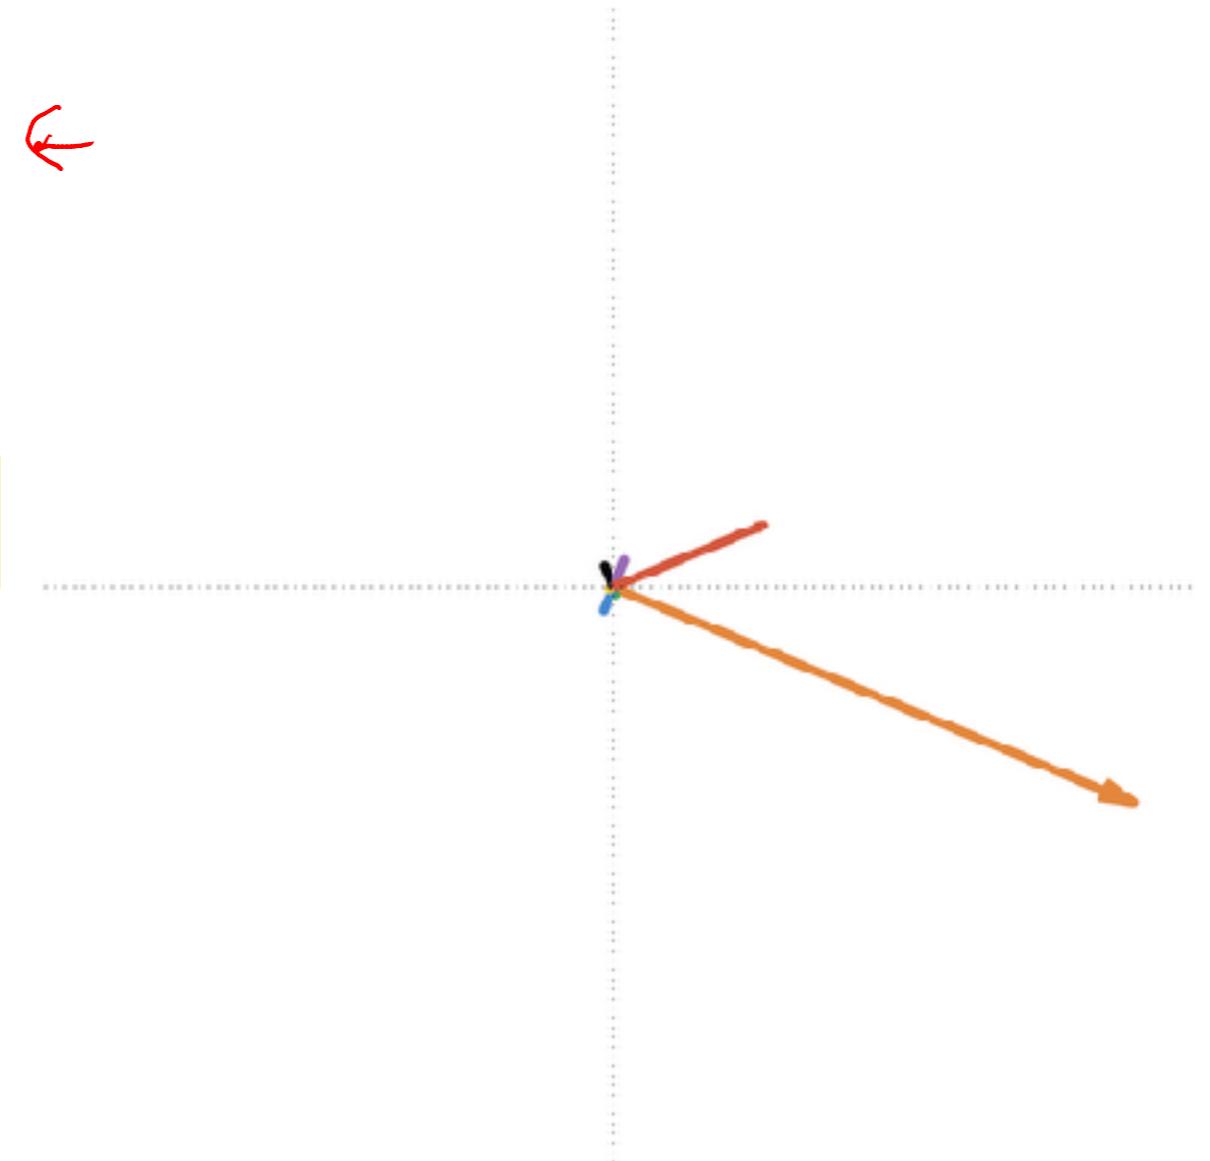
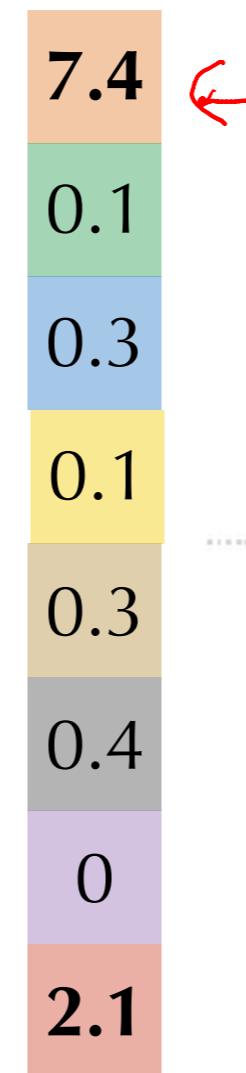
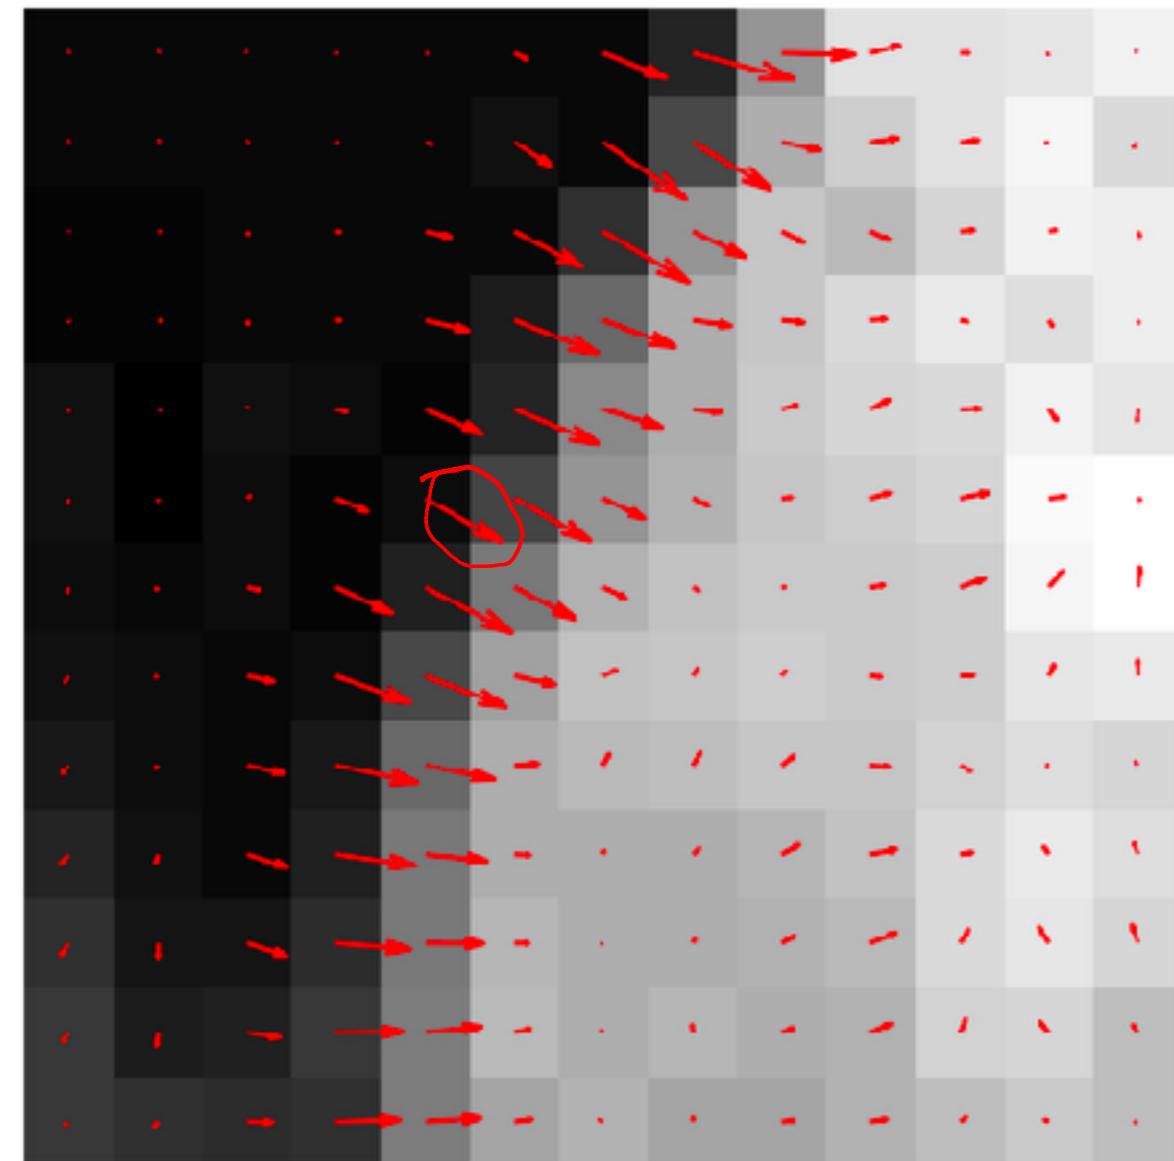


# **SSY097 - Image Analysis**

## Lecture 3 - Scale-Invariant Local Features

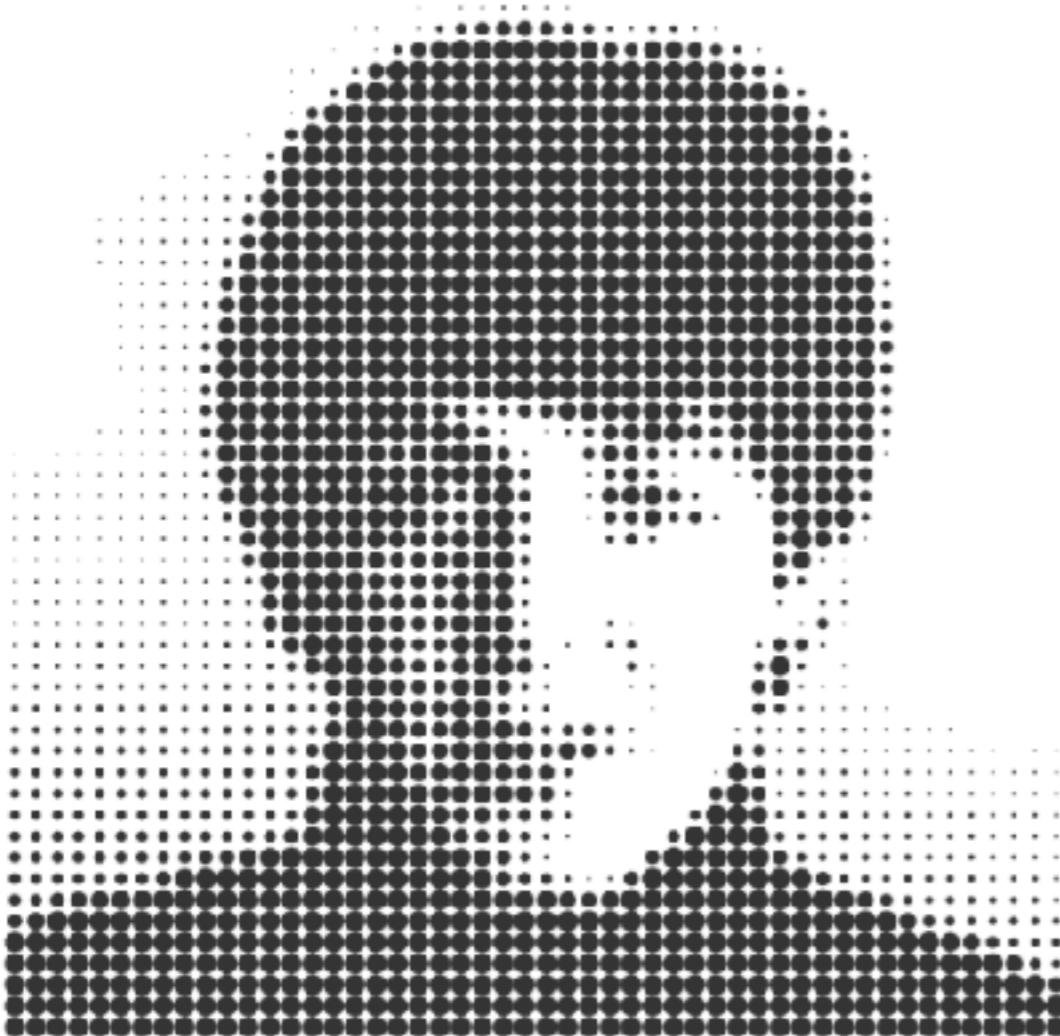
*Torsten Sattler  
(slides adapted from Olof Enqvist)*

# Last Lecture



Gradient histograms

# Last Lecture



Object detection at different scales

# Last Lecture

Scale space

$\sigma$



$L(x, y, 1^2)$



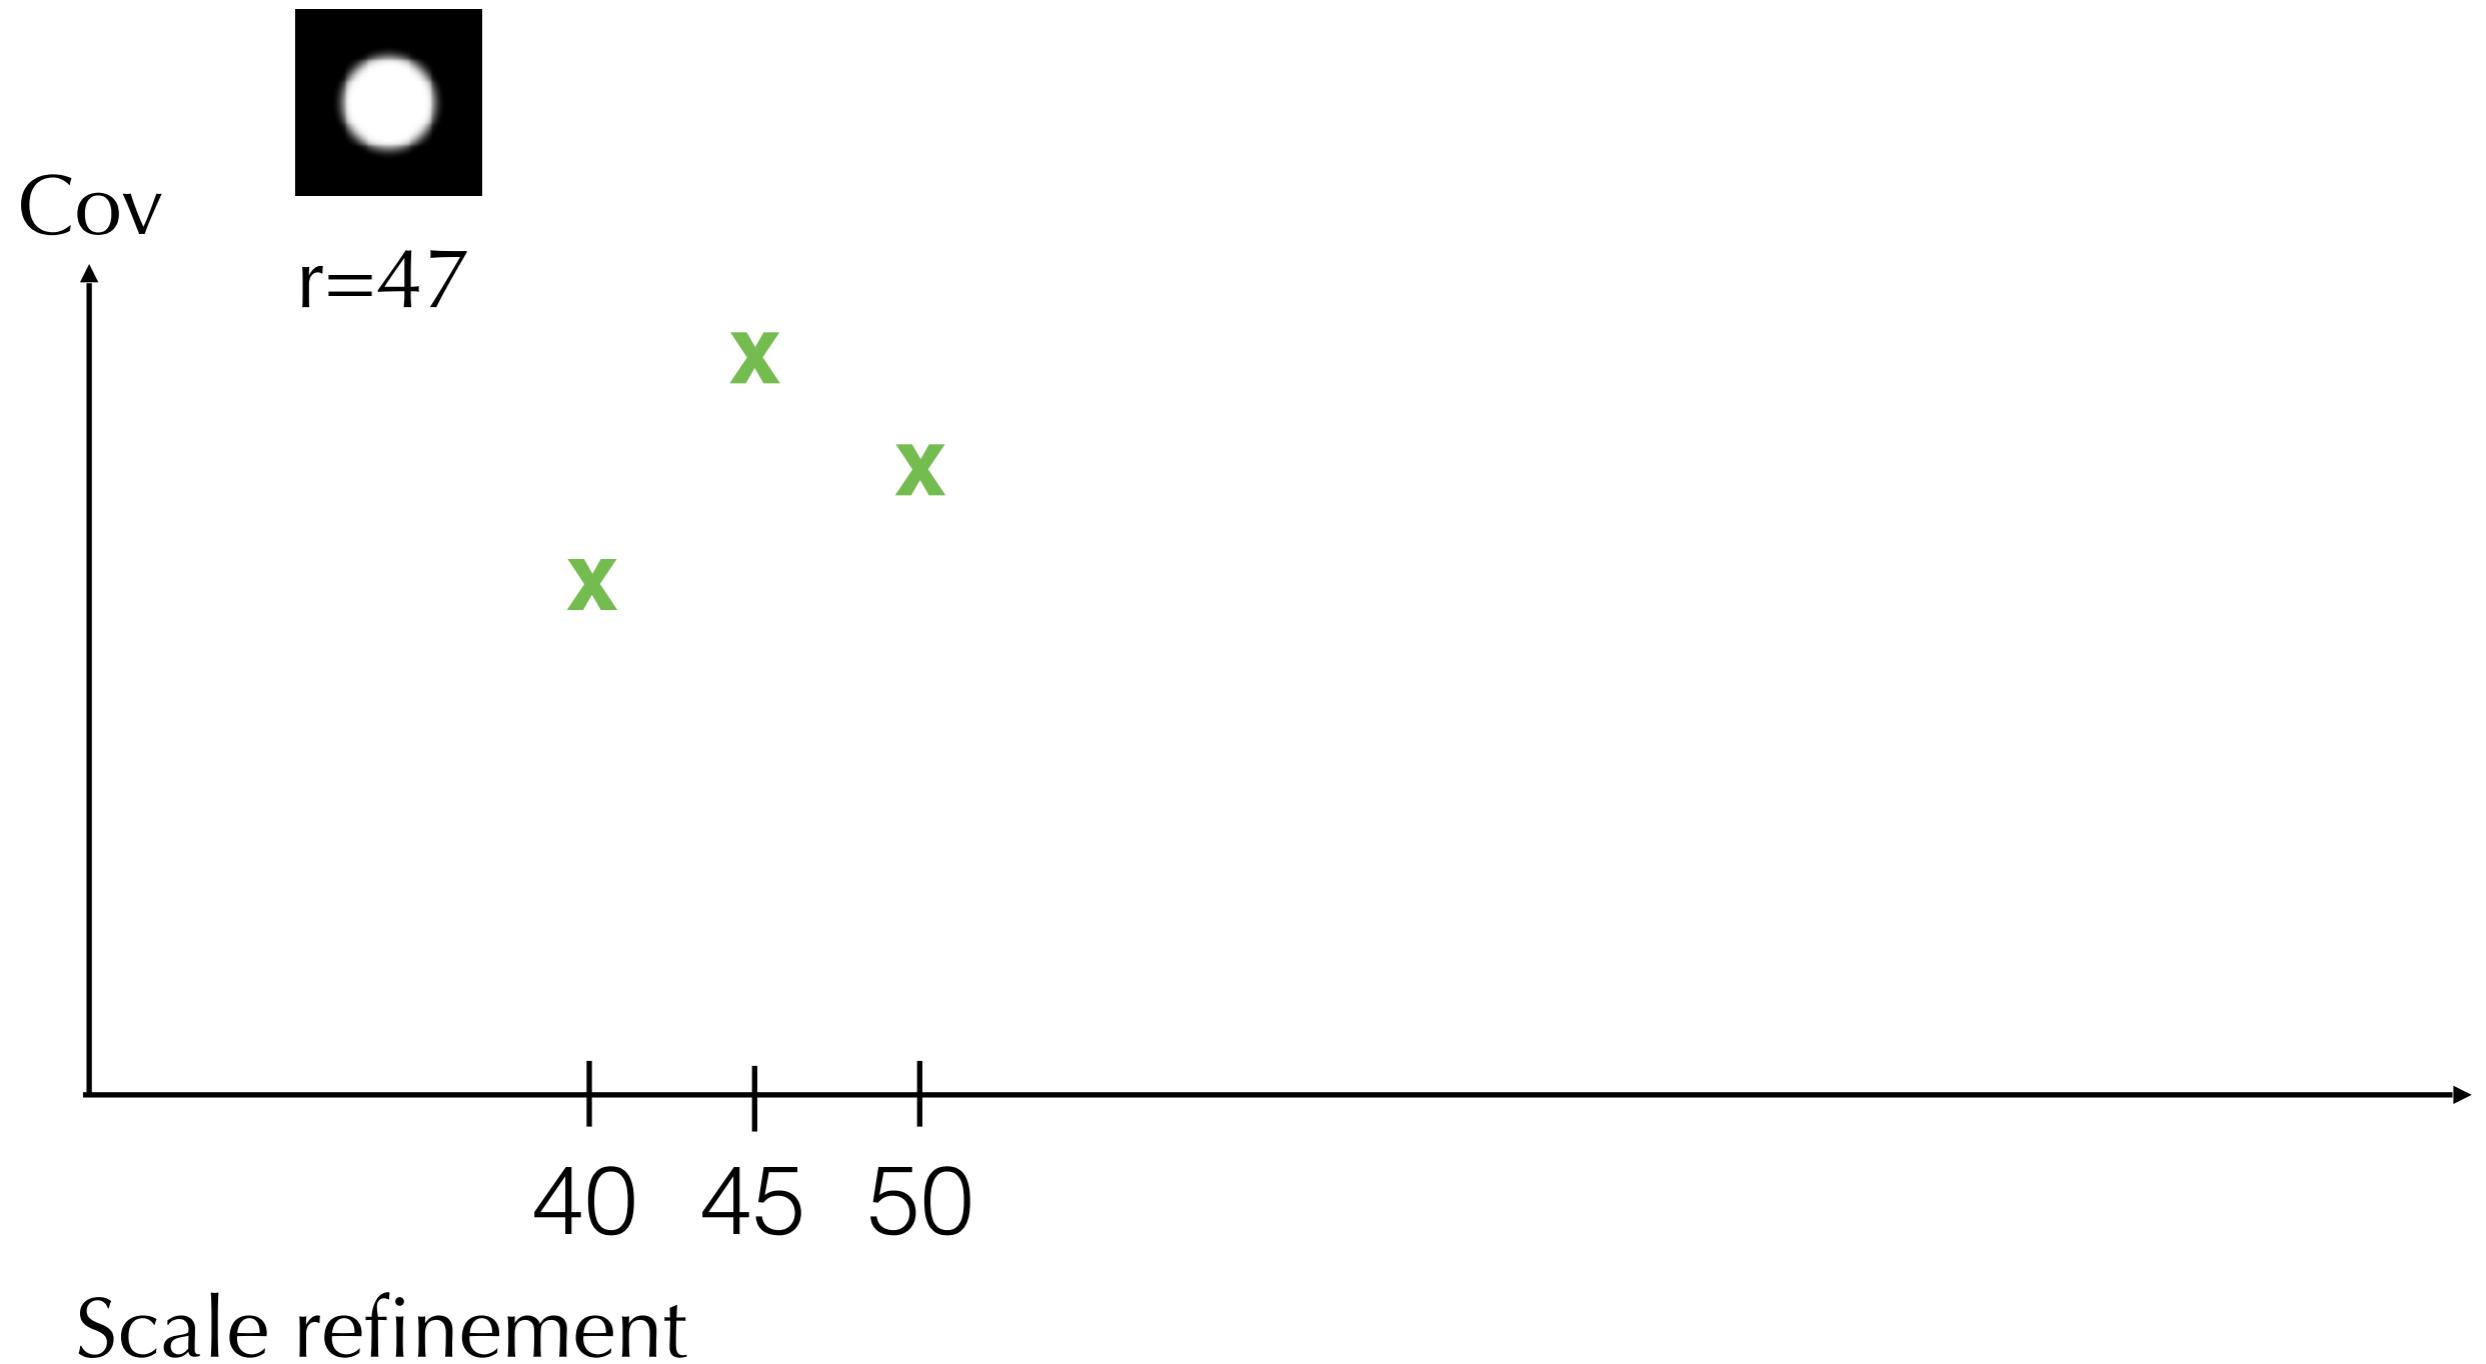
$L(x, y, 4^2)$



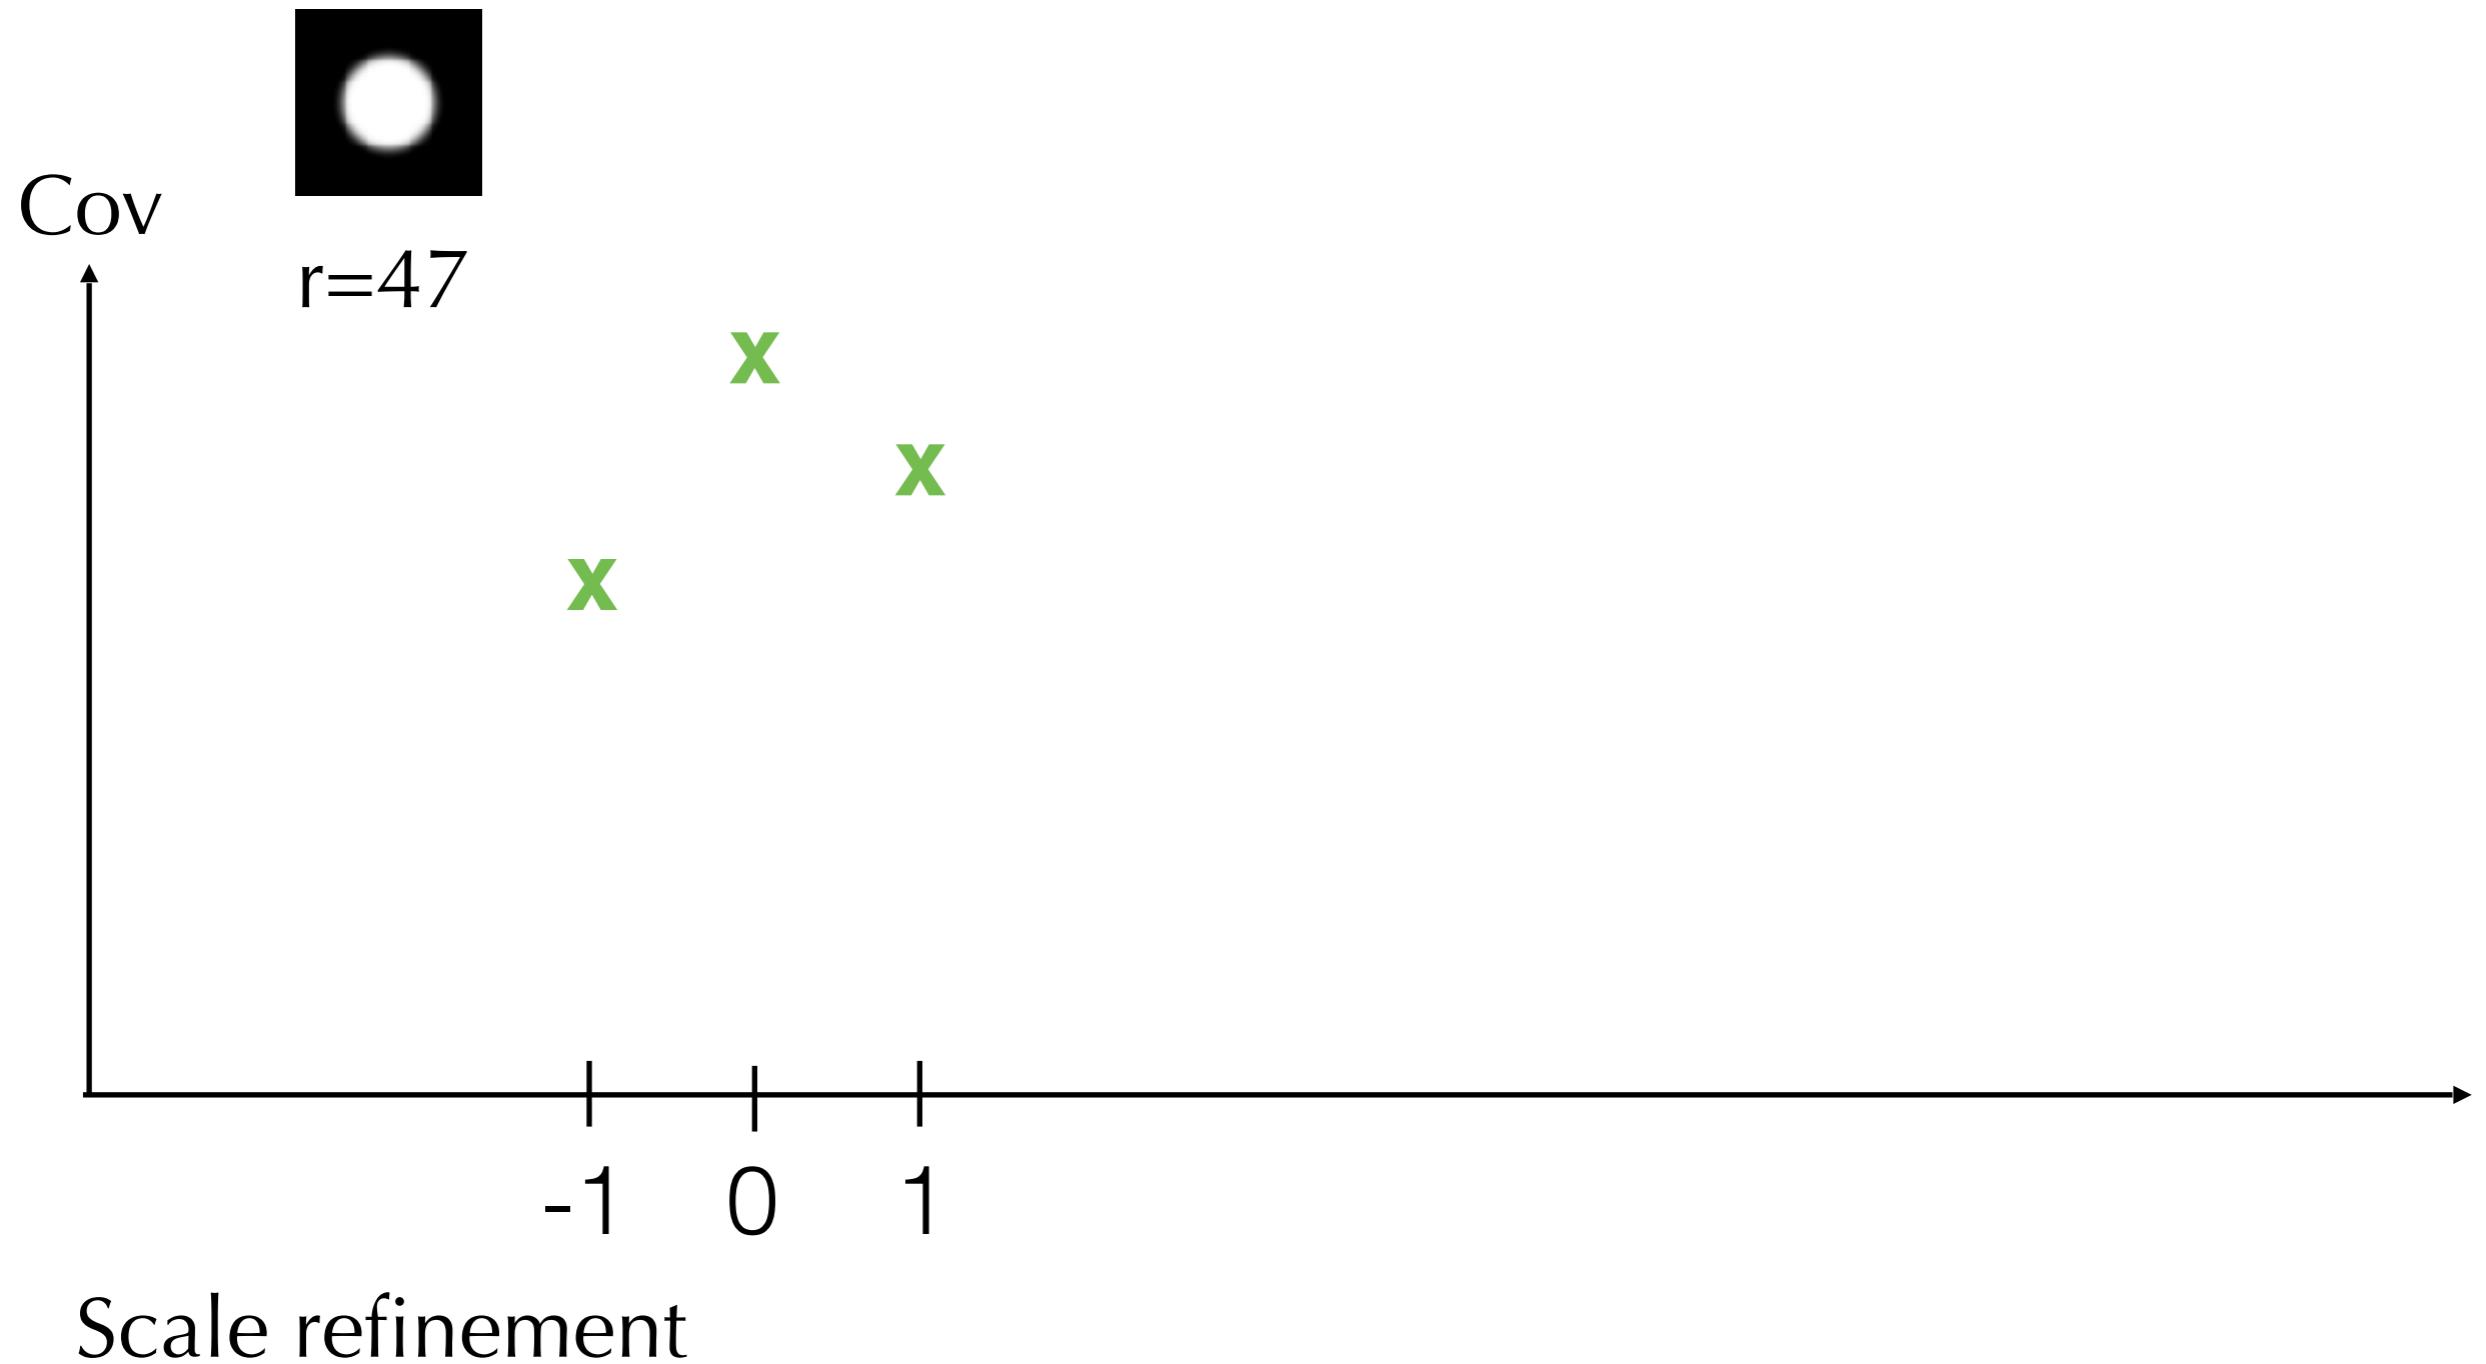
$L(x, y, 2^2)$



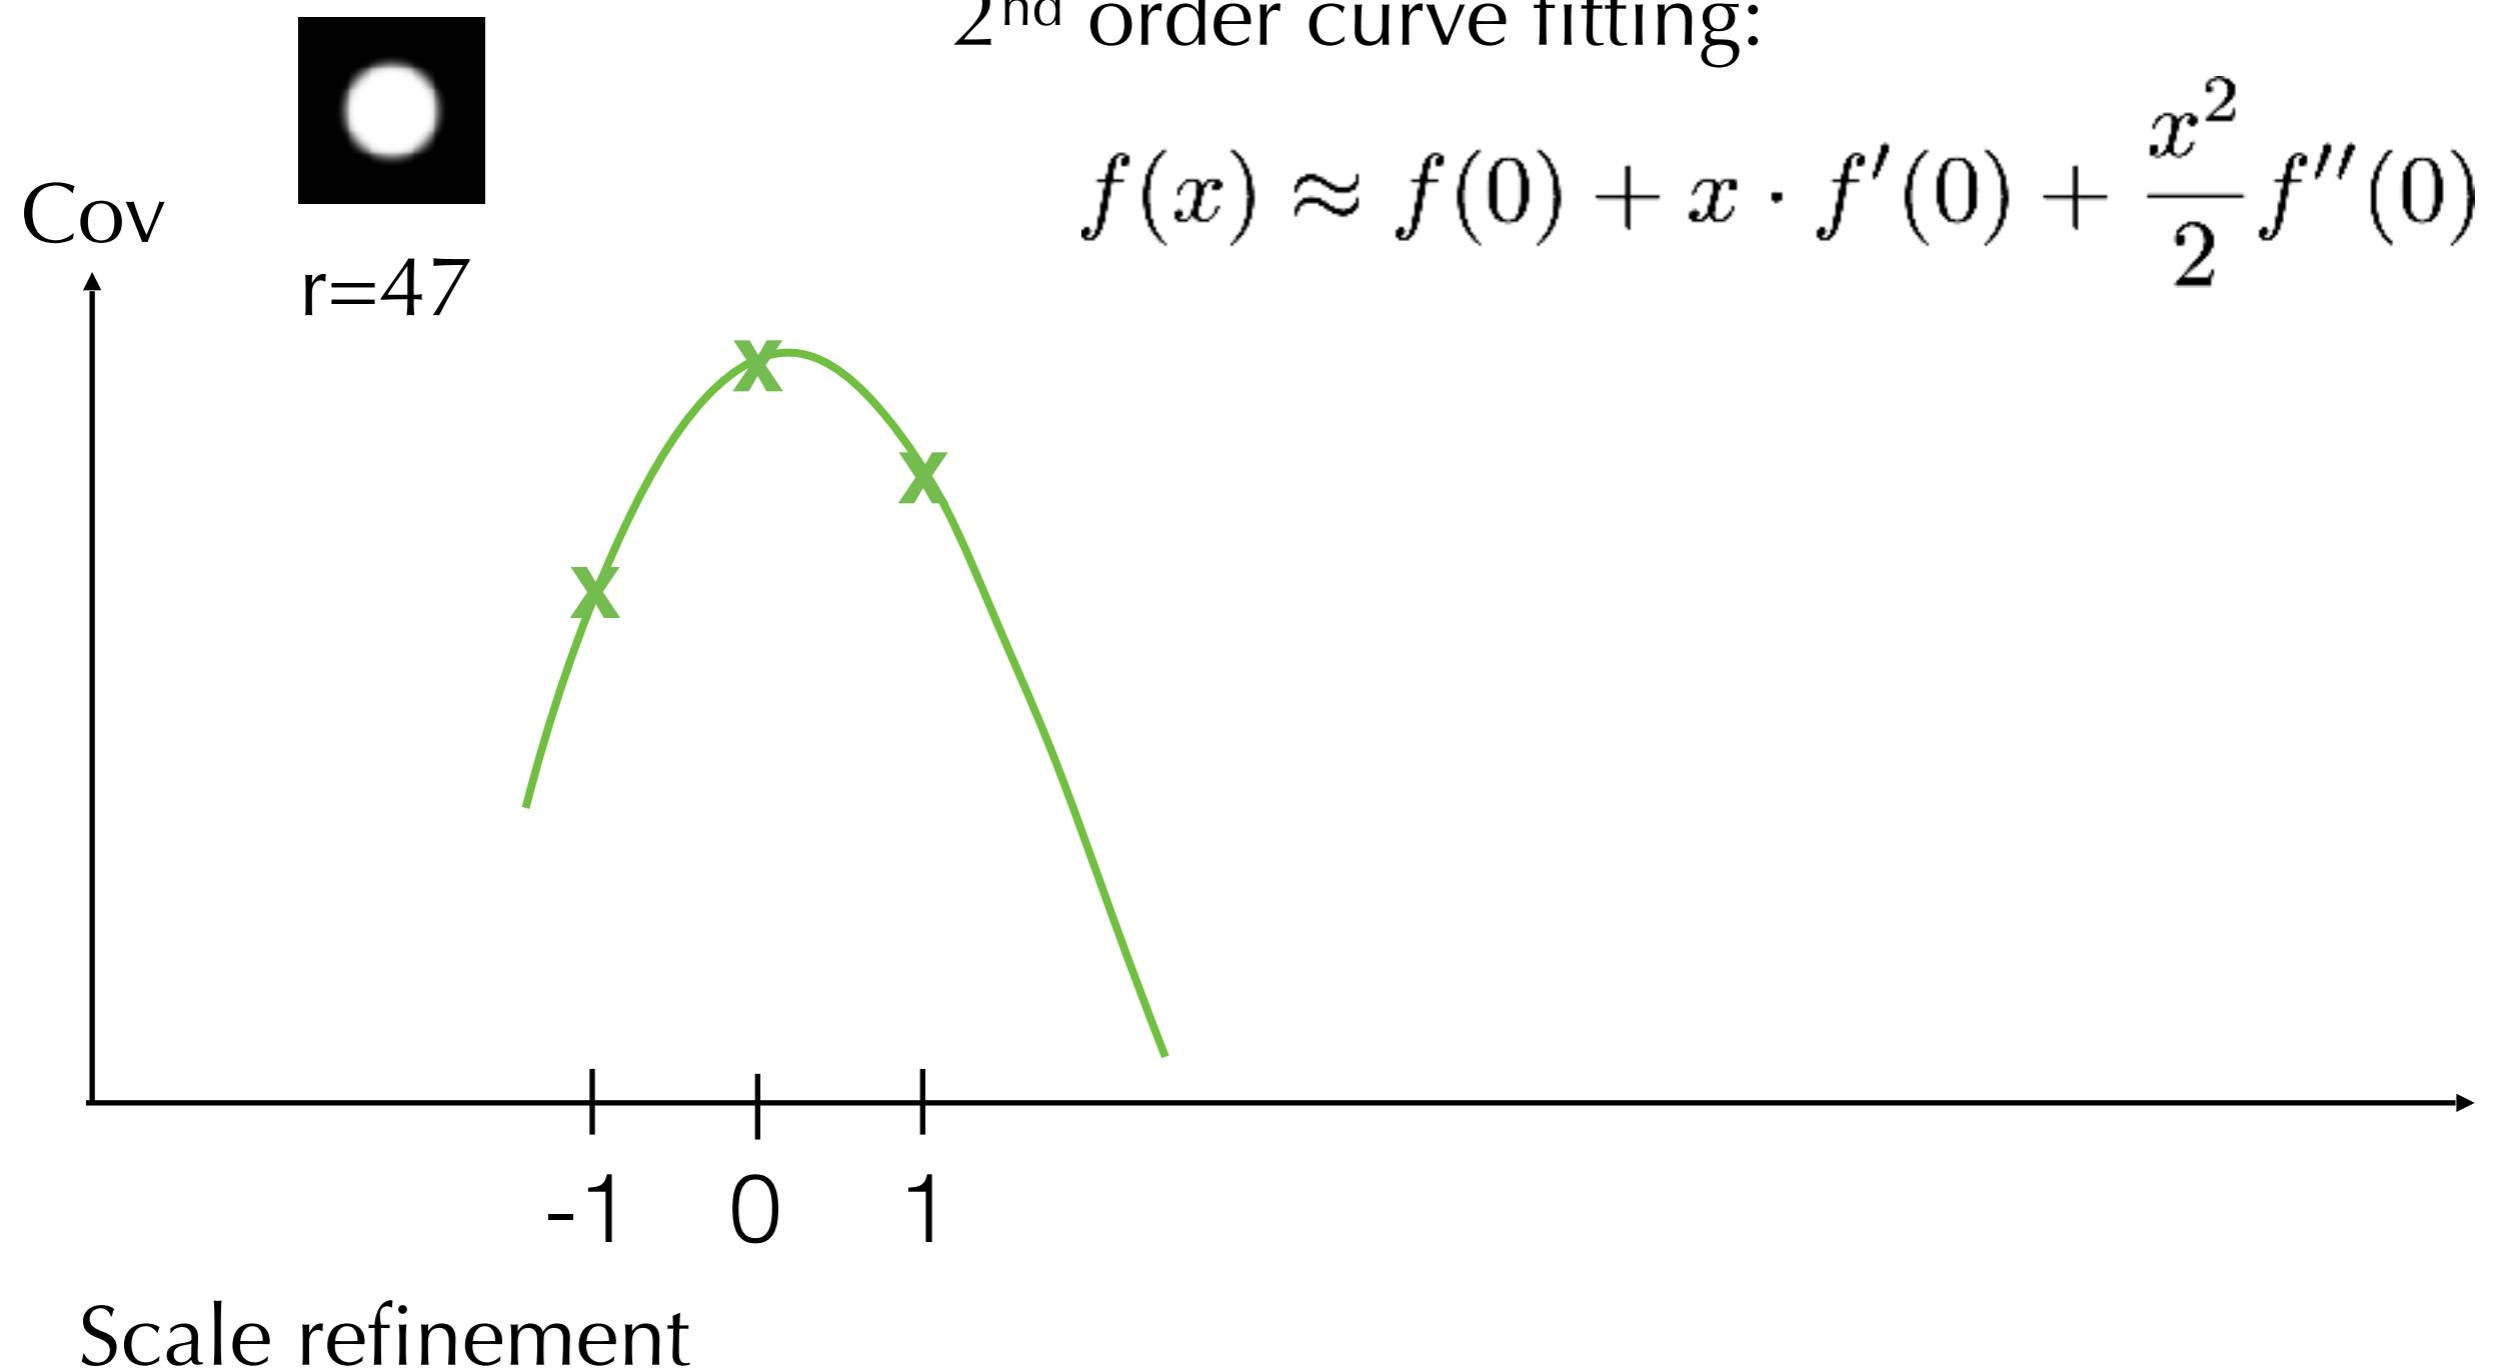
# Last Lecture



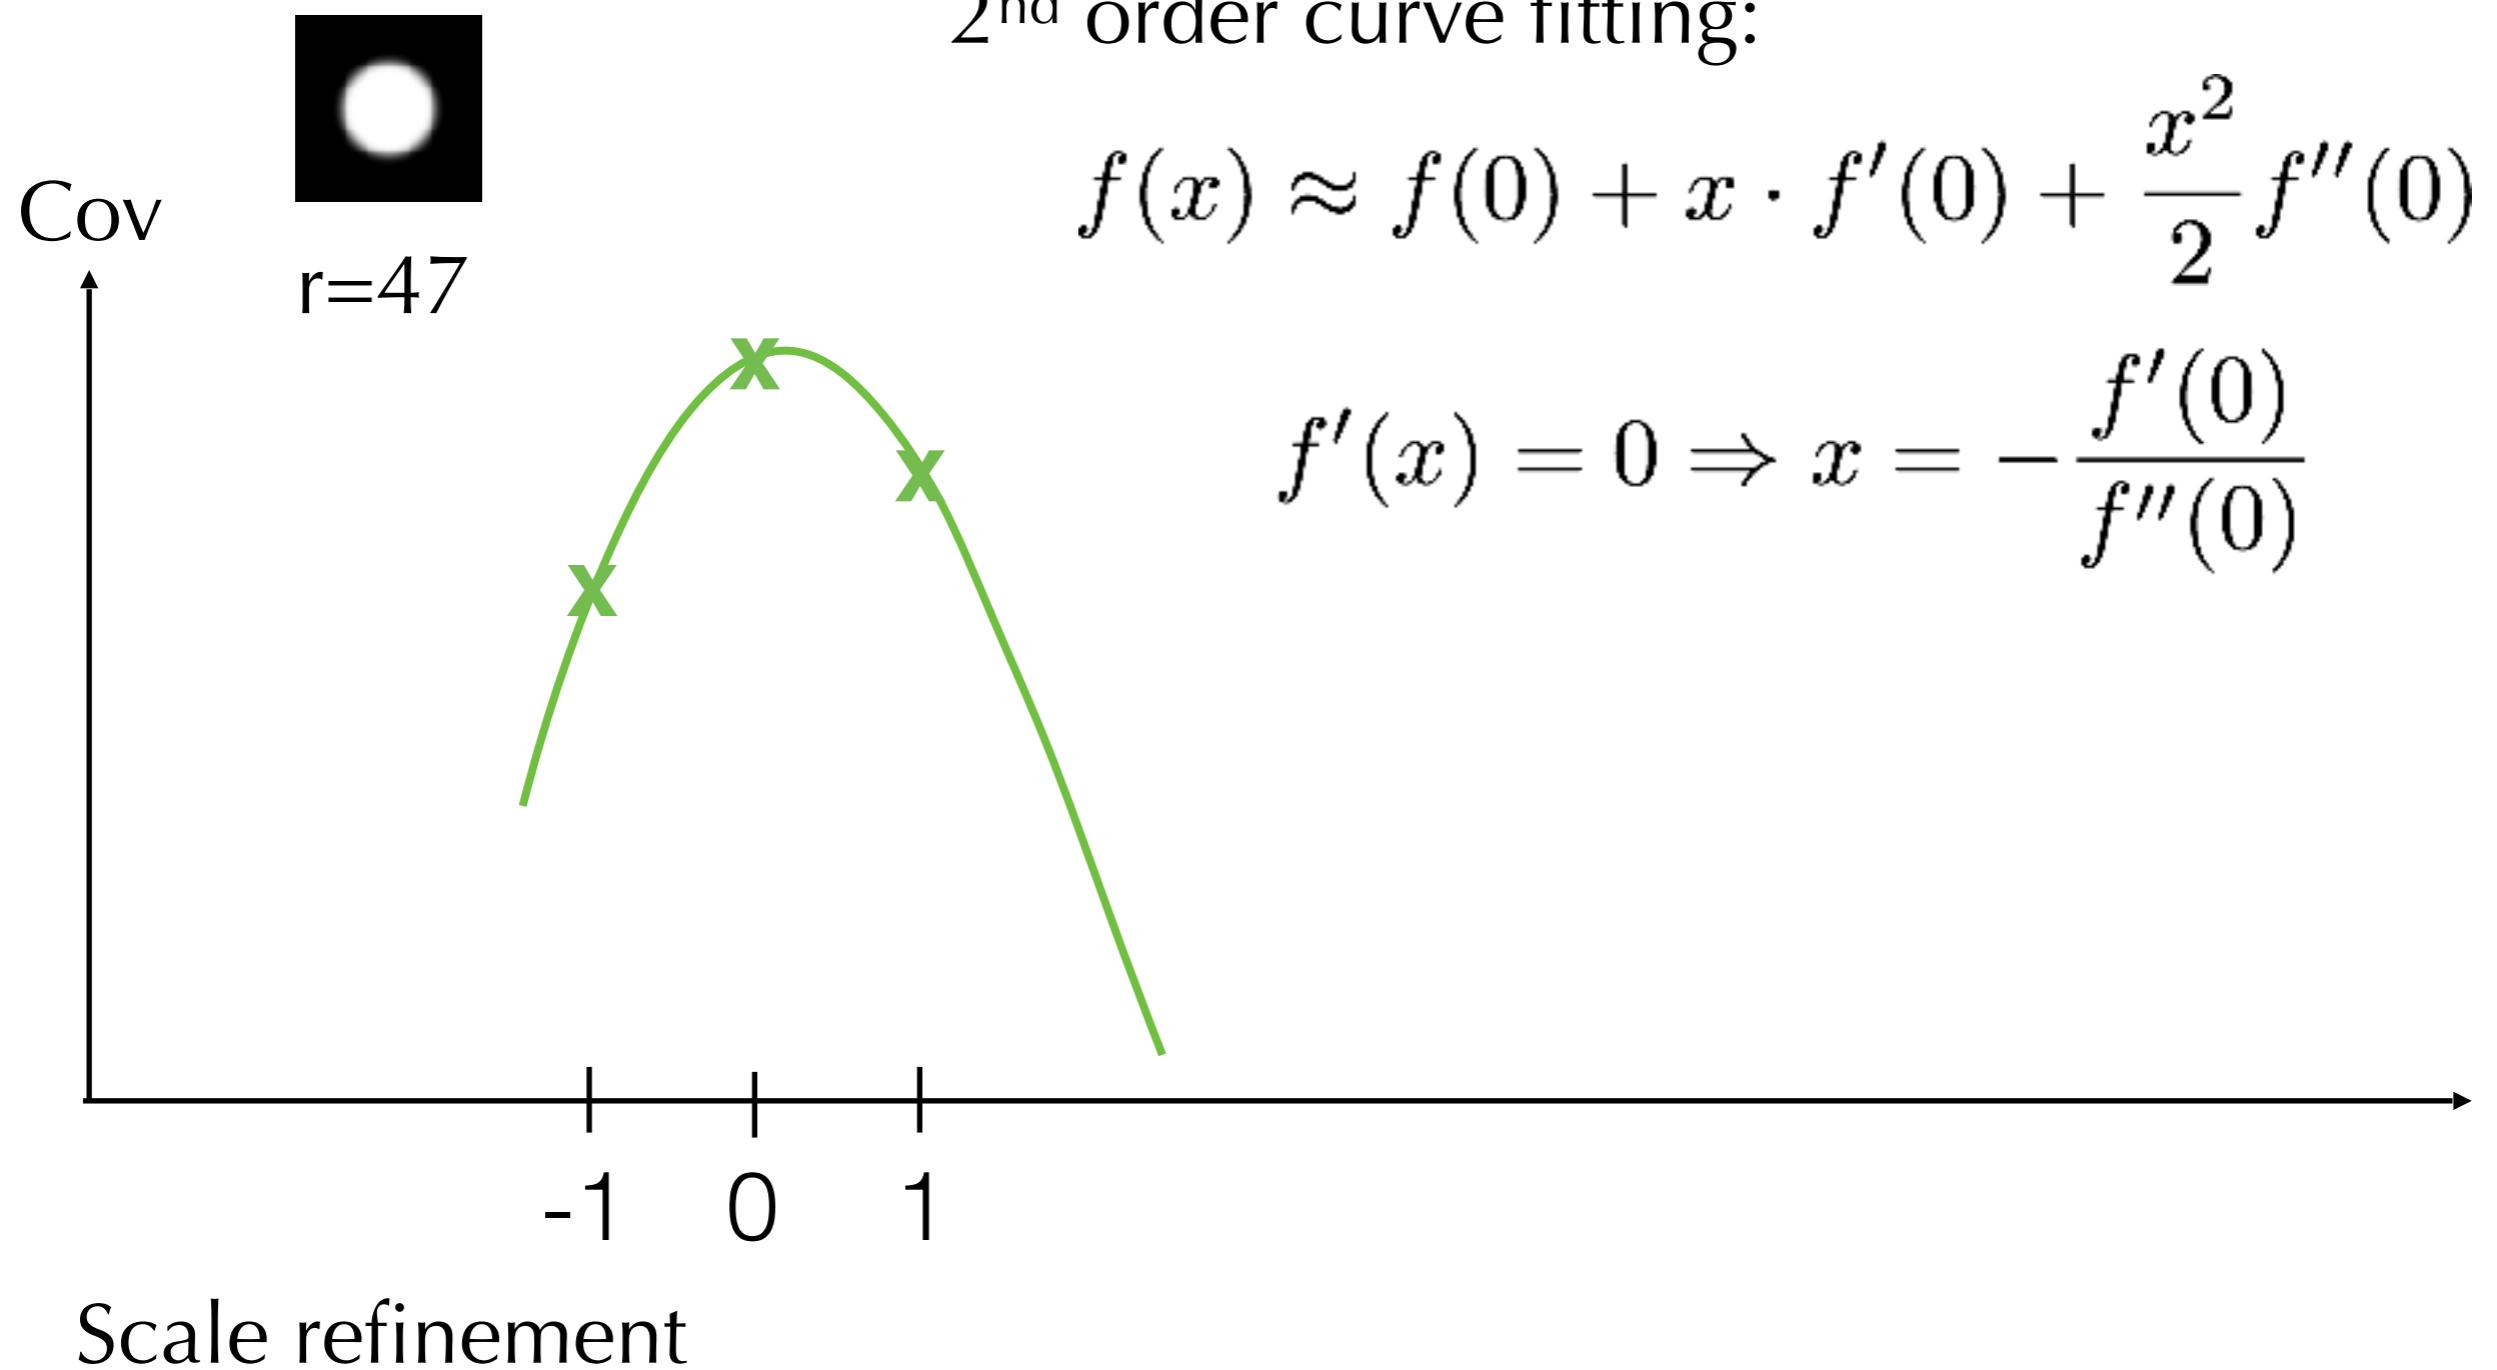
# Last Lecture



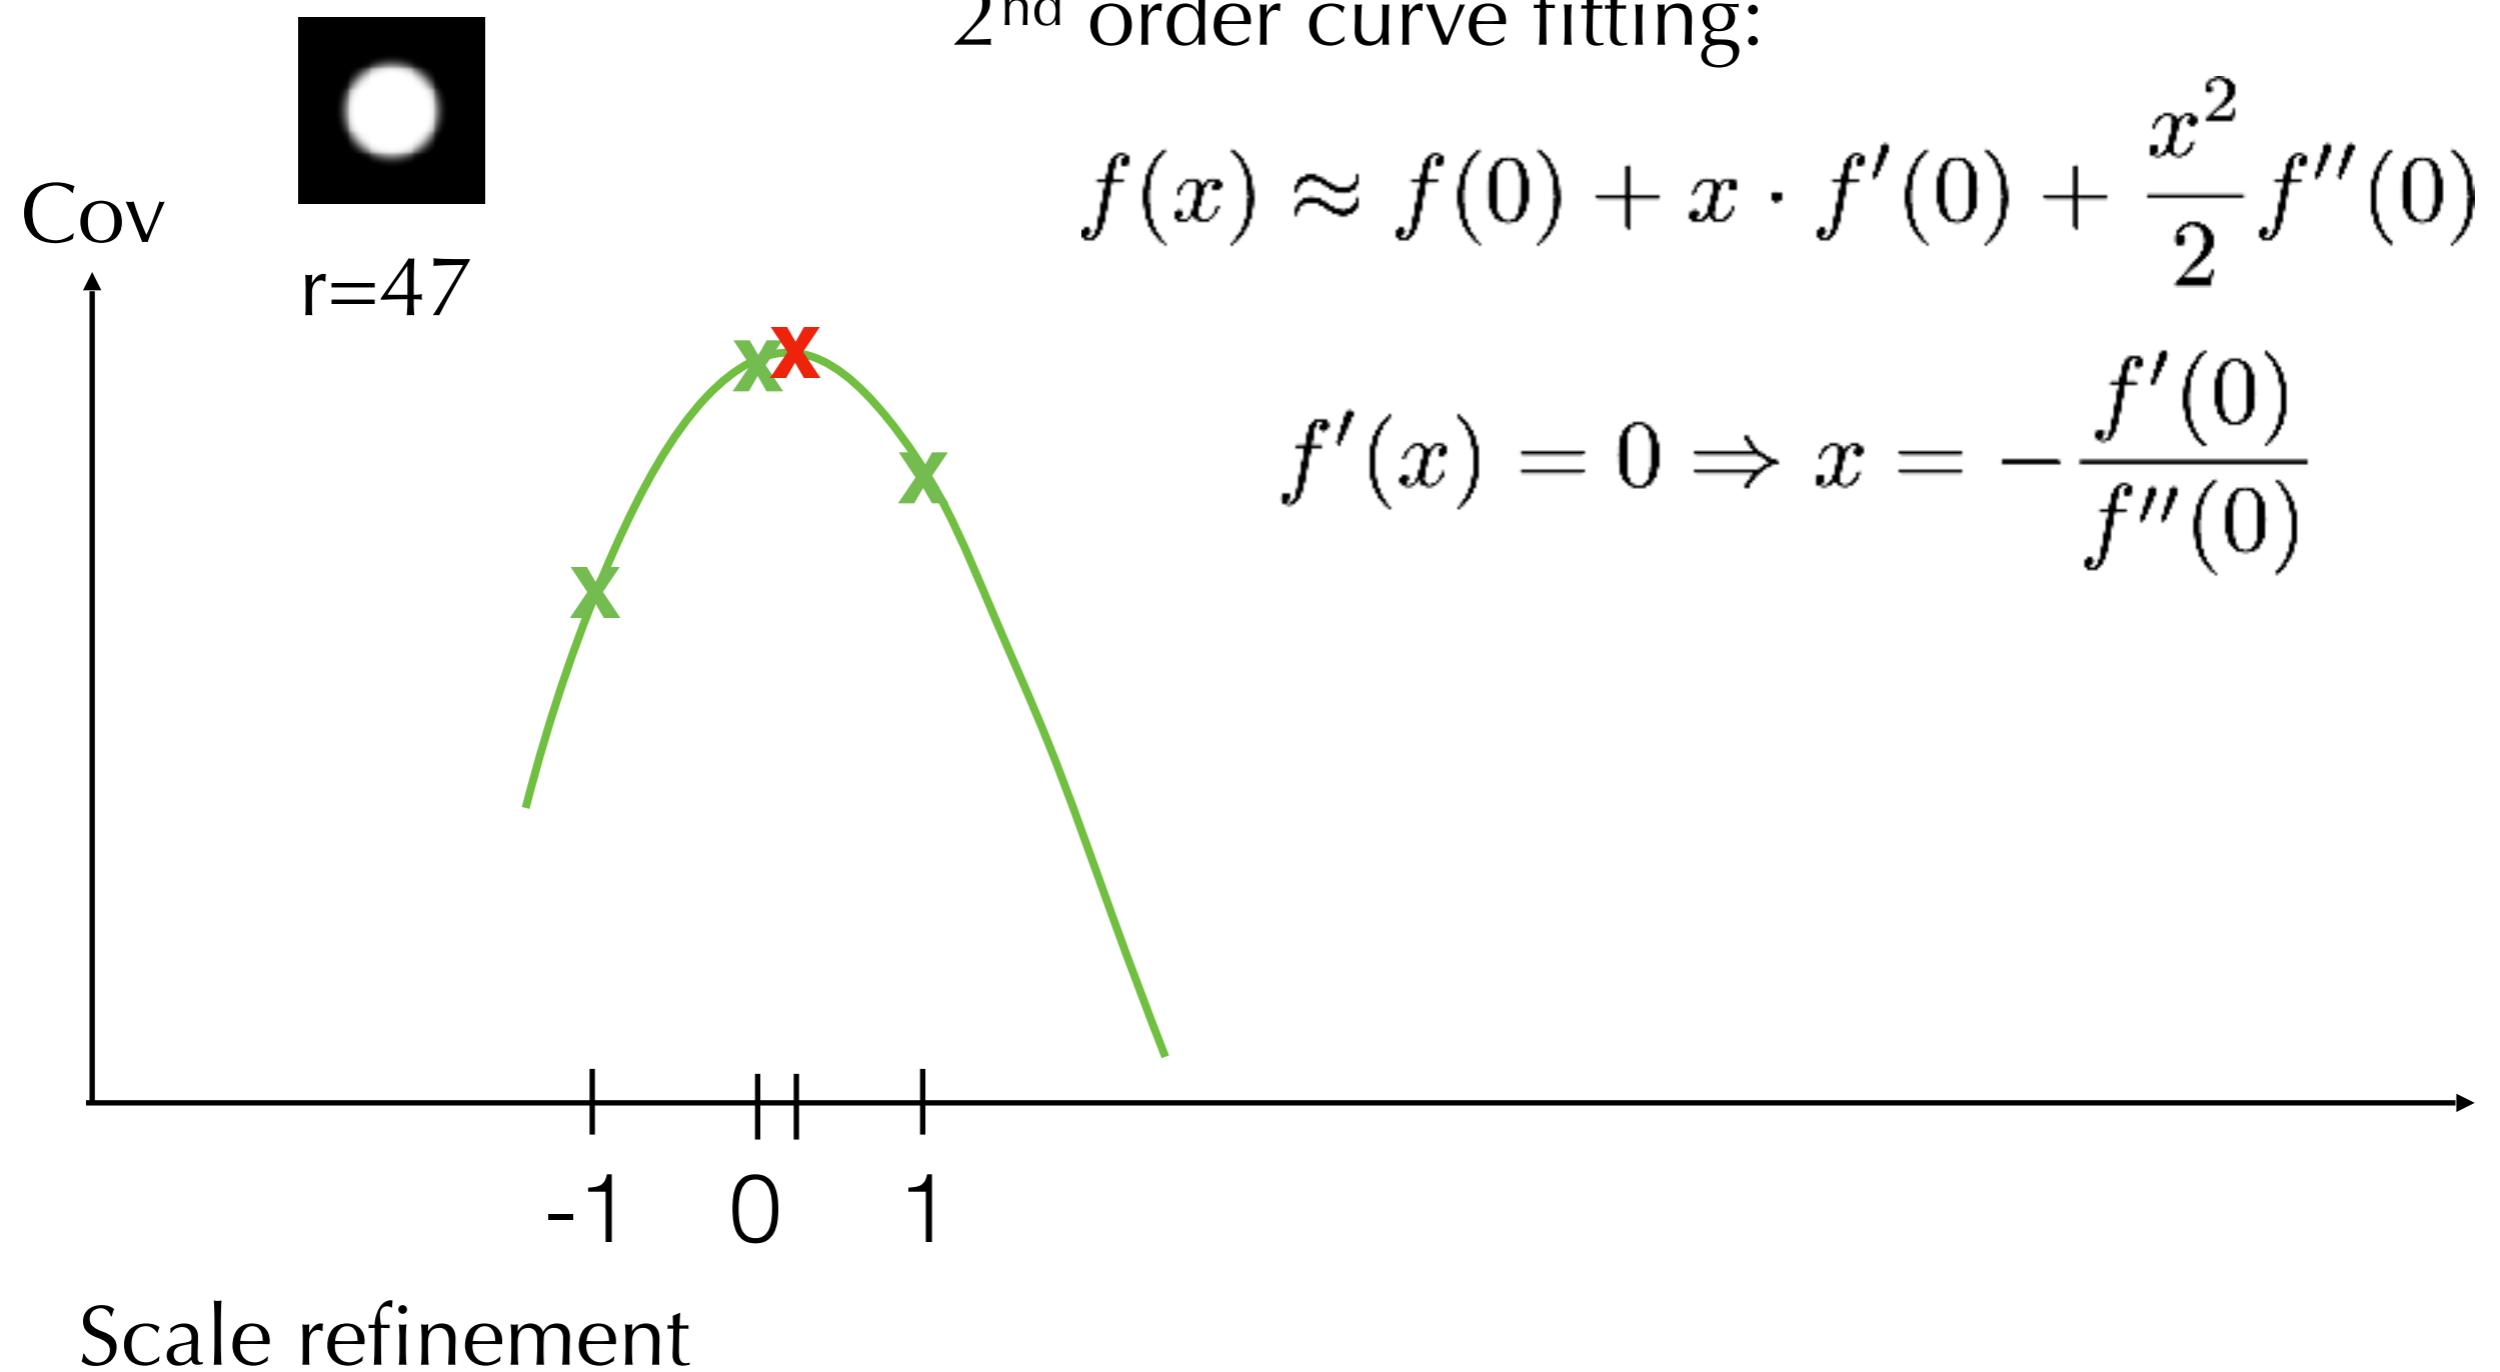
# Last Lecture



# Last Lecture



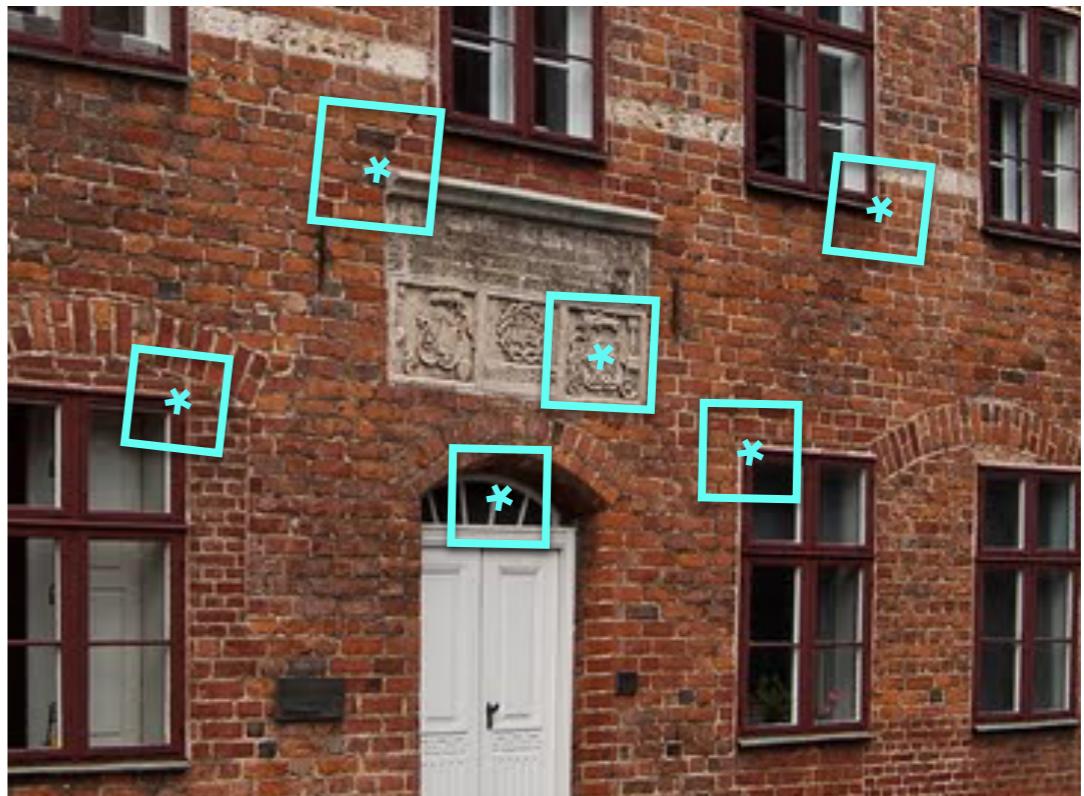
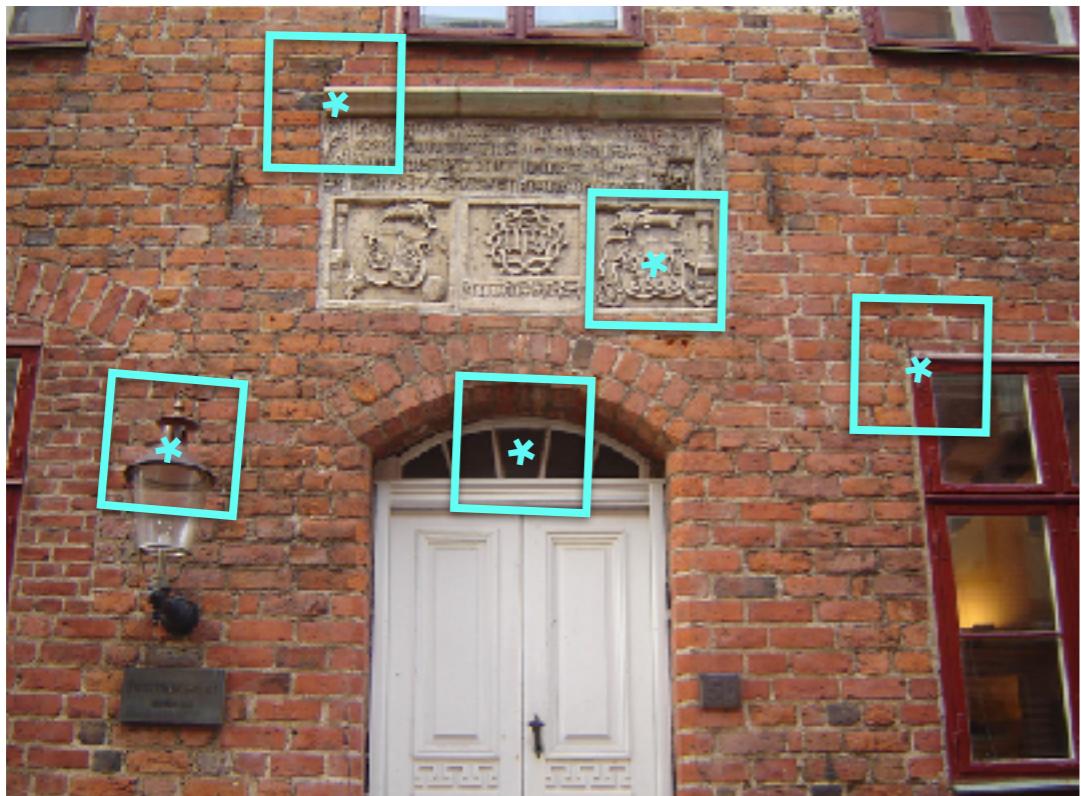
# Last Lecture



# Today

- Scale-Invariant Local Features
- Invariances\*

# Why?

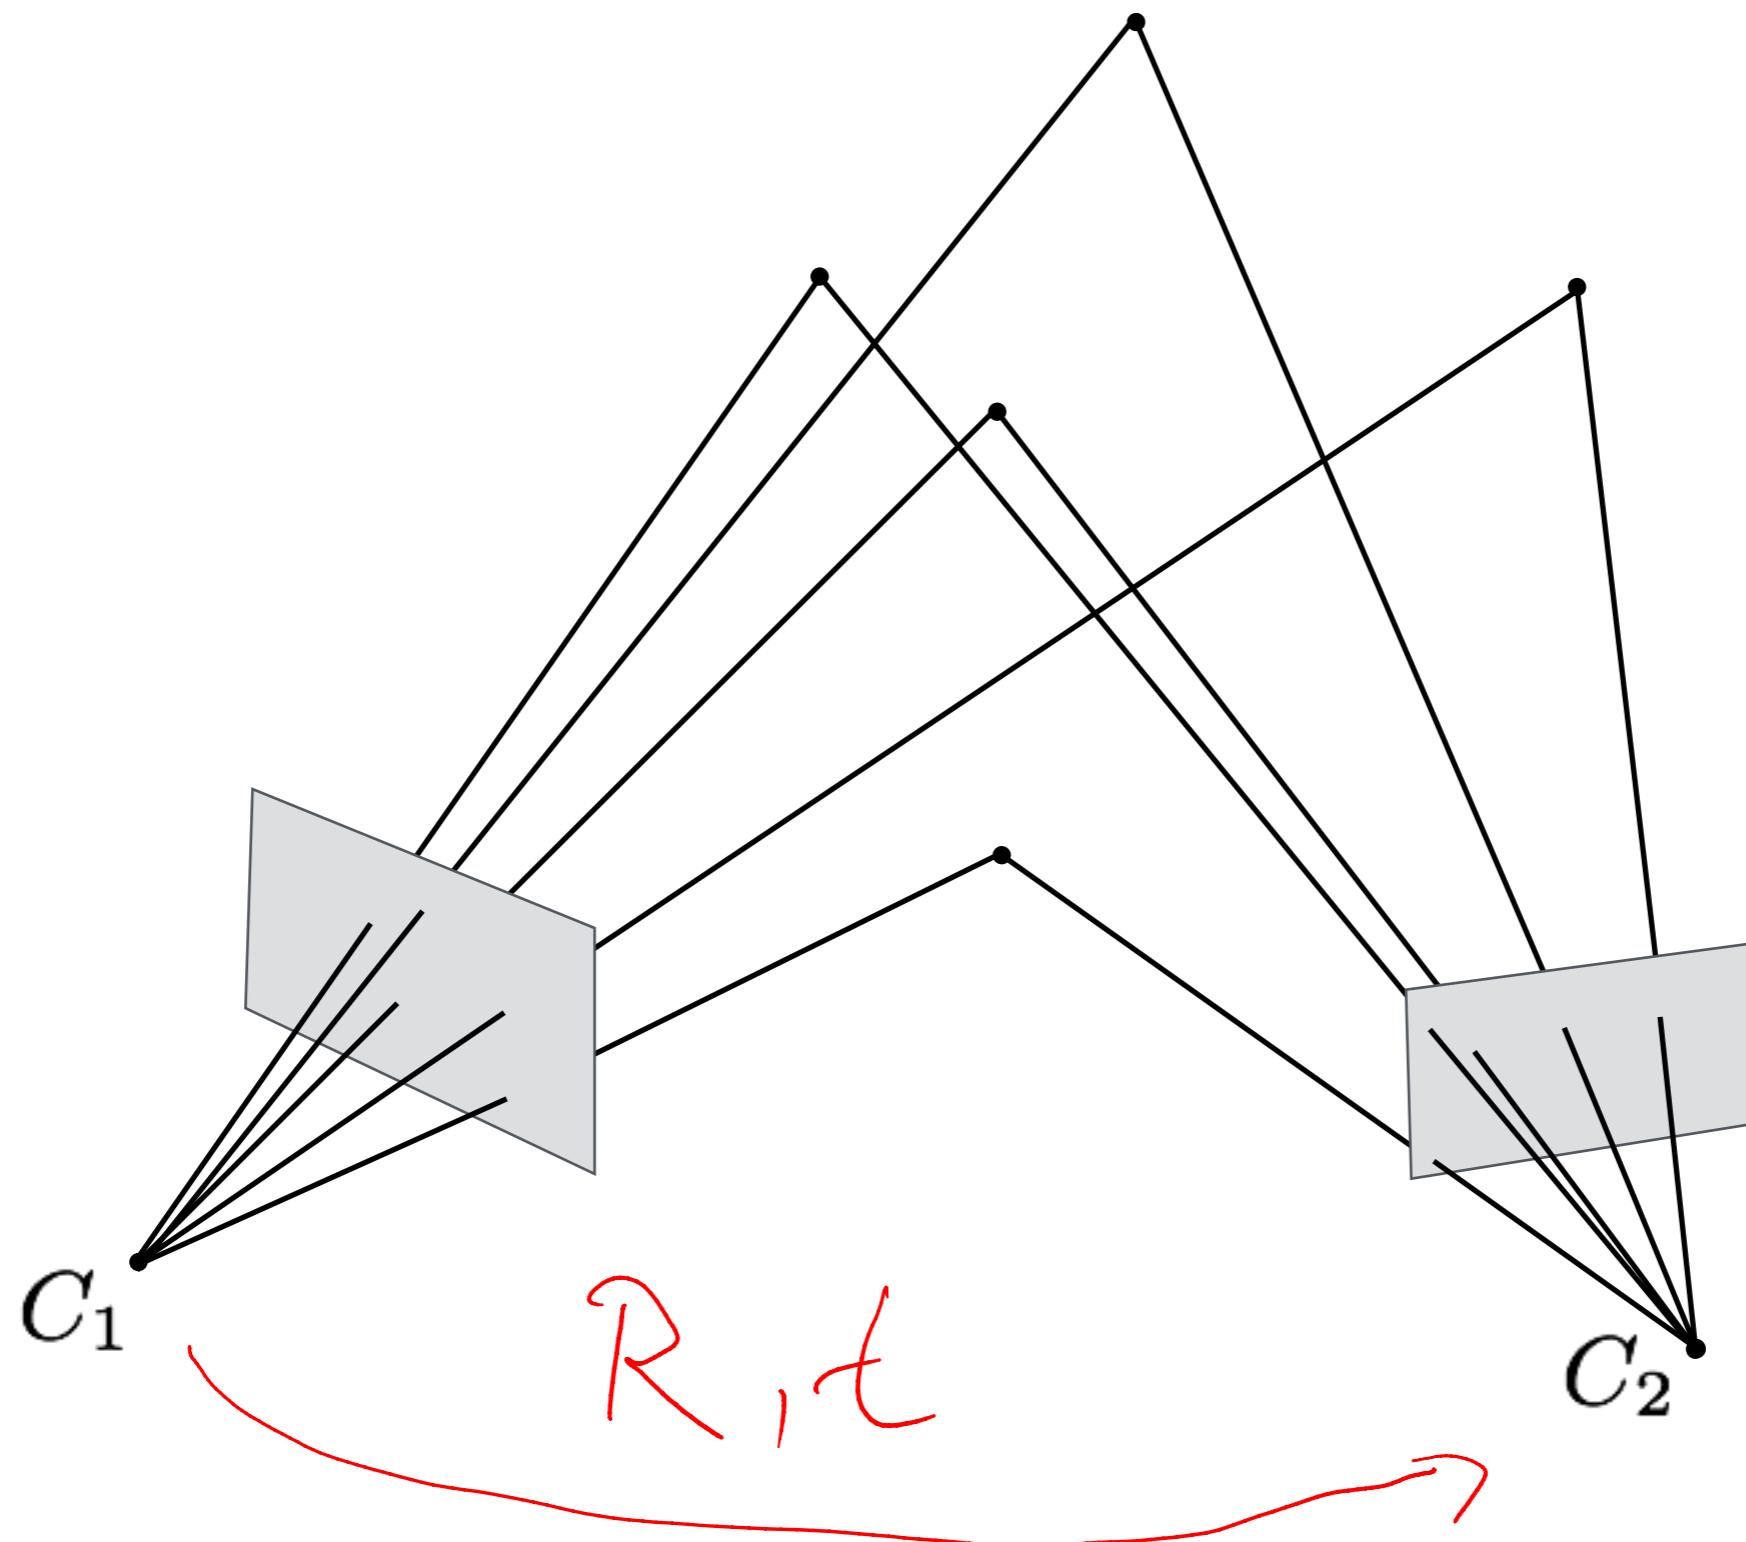


Patch matching

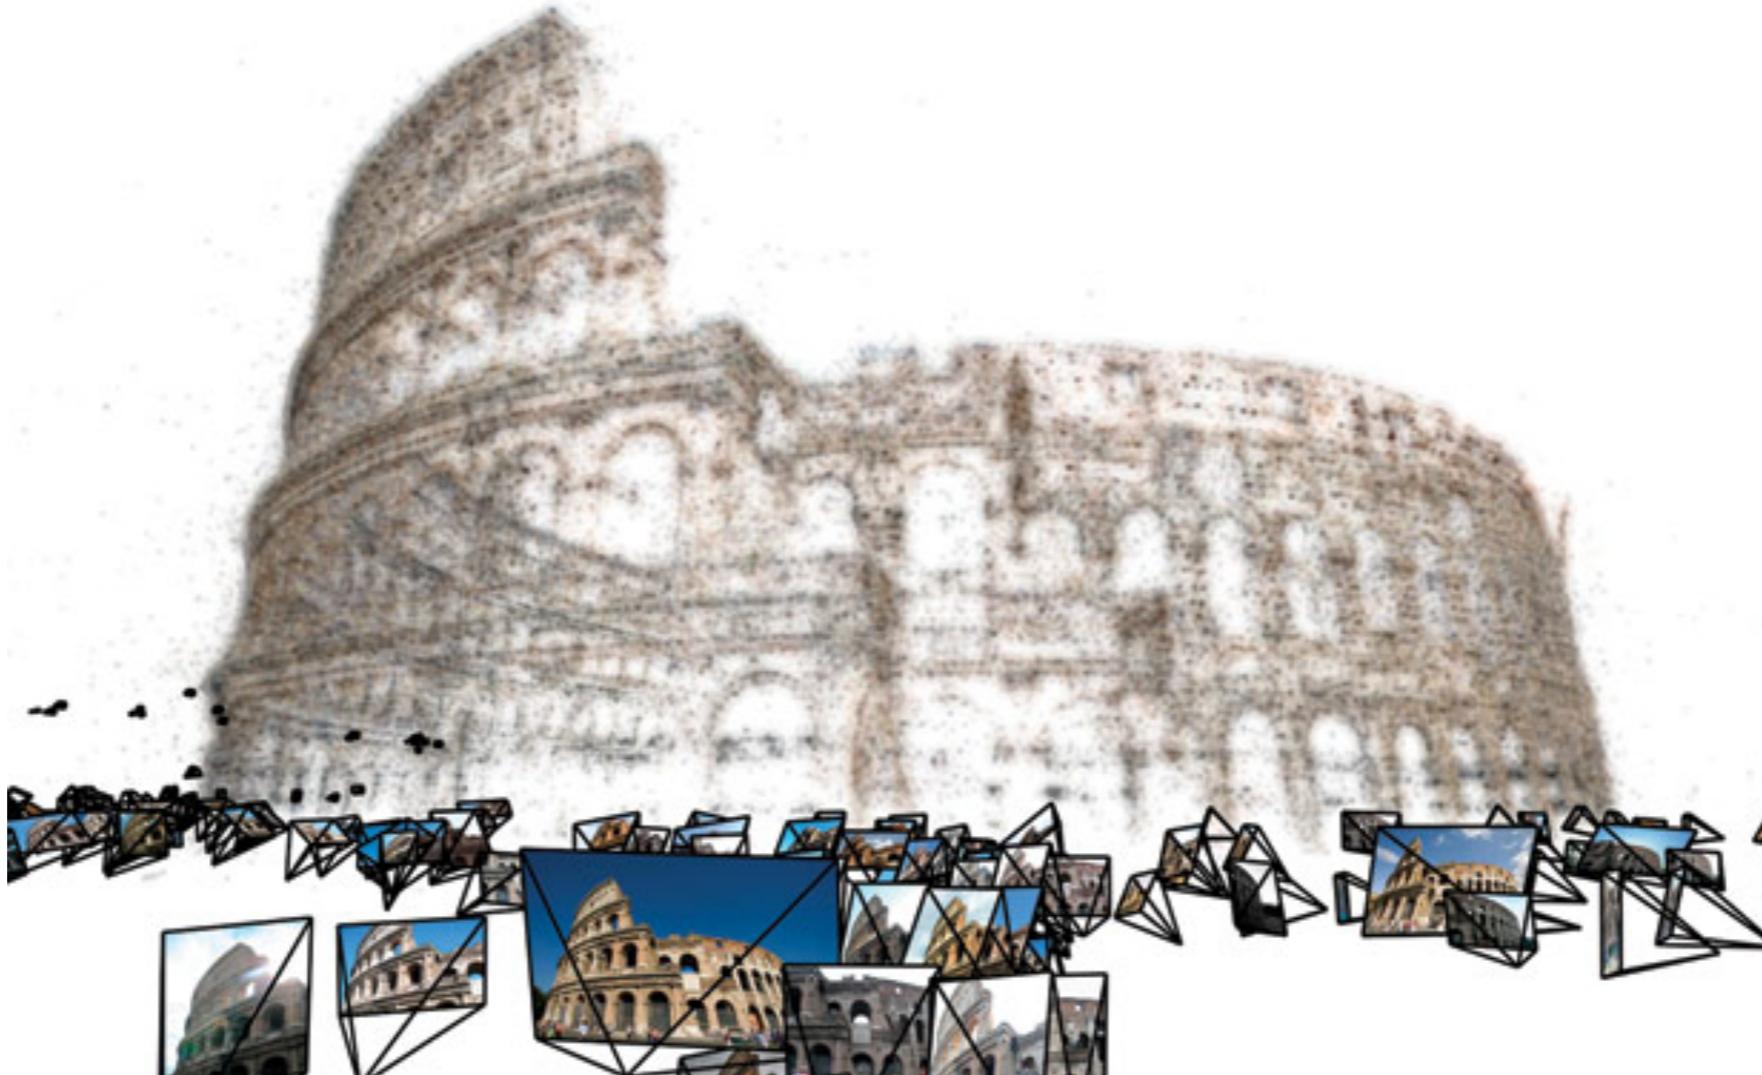
# Applications: Image Stitching



# Relative Pose Estimation



# Applications: 3D Reconstruction



Snavely et al.

# Scale-Invariant Feature Transform (SIFT) Features



[Lowe, Distinctive Image Features from Scale-Invariant Keypoints, IJCV 2004]

# Scale-Invariant Feature Transform (SIFT) Features



- Find a set of *interest* points

[Lowe, Distinctive Image Features from Scale-Invariant Keypoints, IJCV 2004]

# Scale-Invariant Feature Transform (SIFT) Features



- Find a set of *interest* points
- Extract oriented patch around each interest point

[Lowe, Distinctive Image Features from Scale-Invariant Keypoints, IJCV 2004]

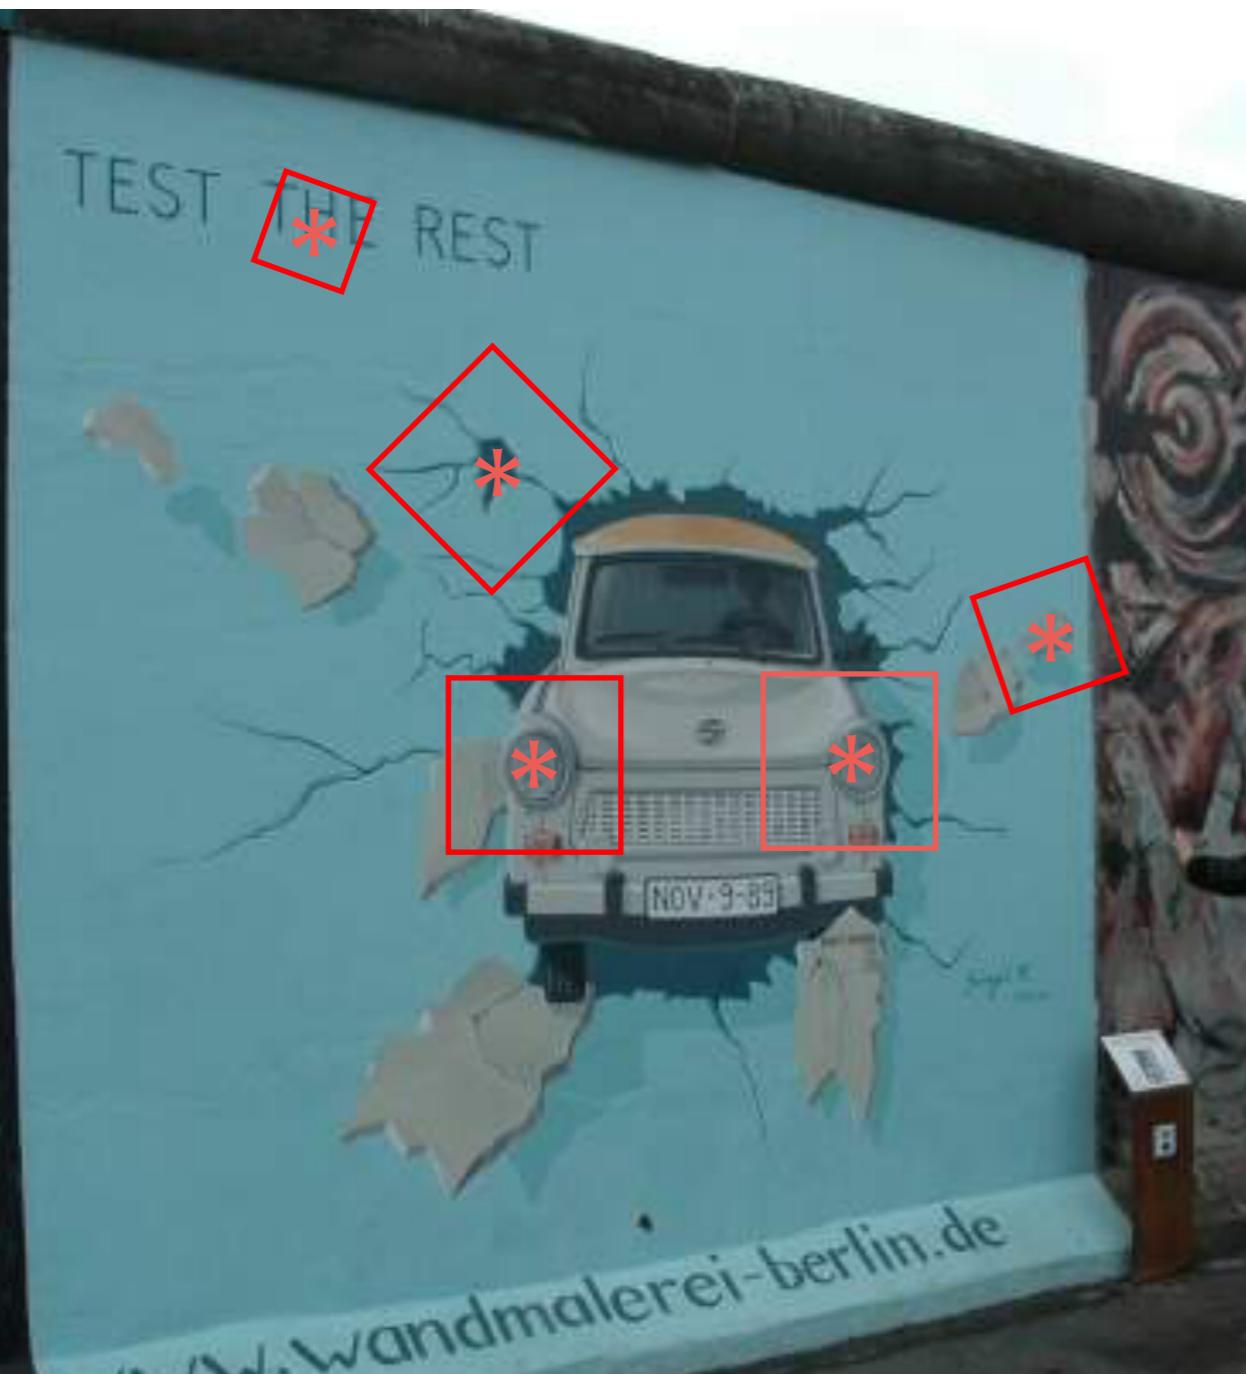
# Scale-Invariant Feature Transform (SIFT) Features



- Find a set of *interest* points
- Extract oriented patch around each interest point
- Compute SIFT descriptor for each patch

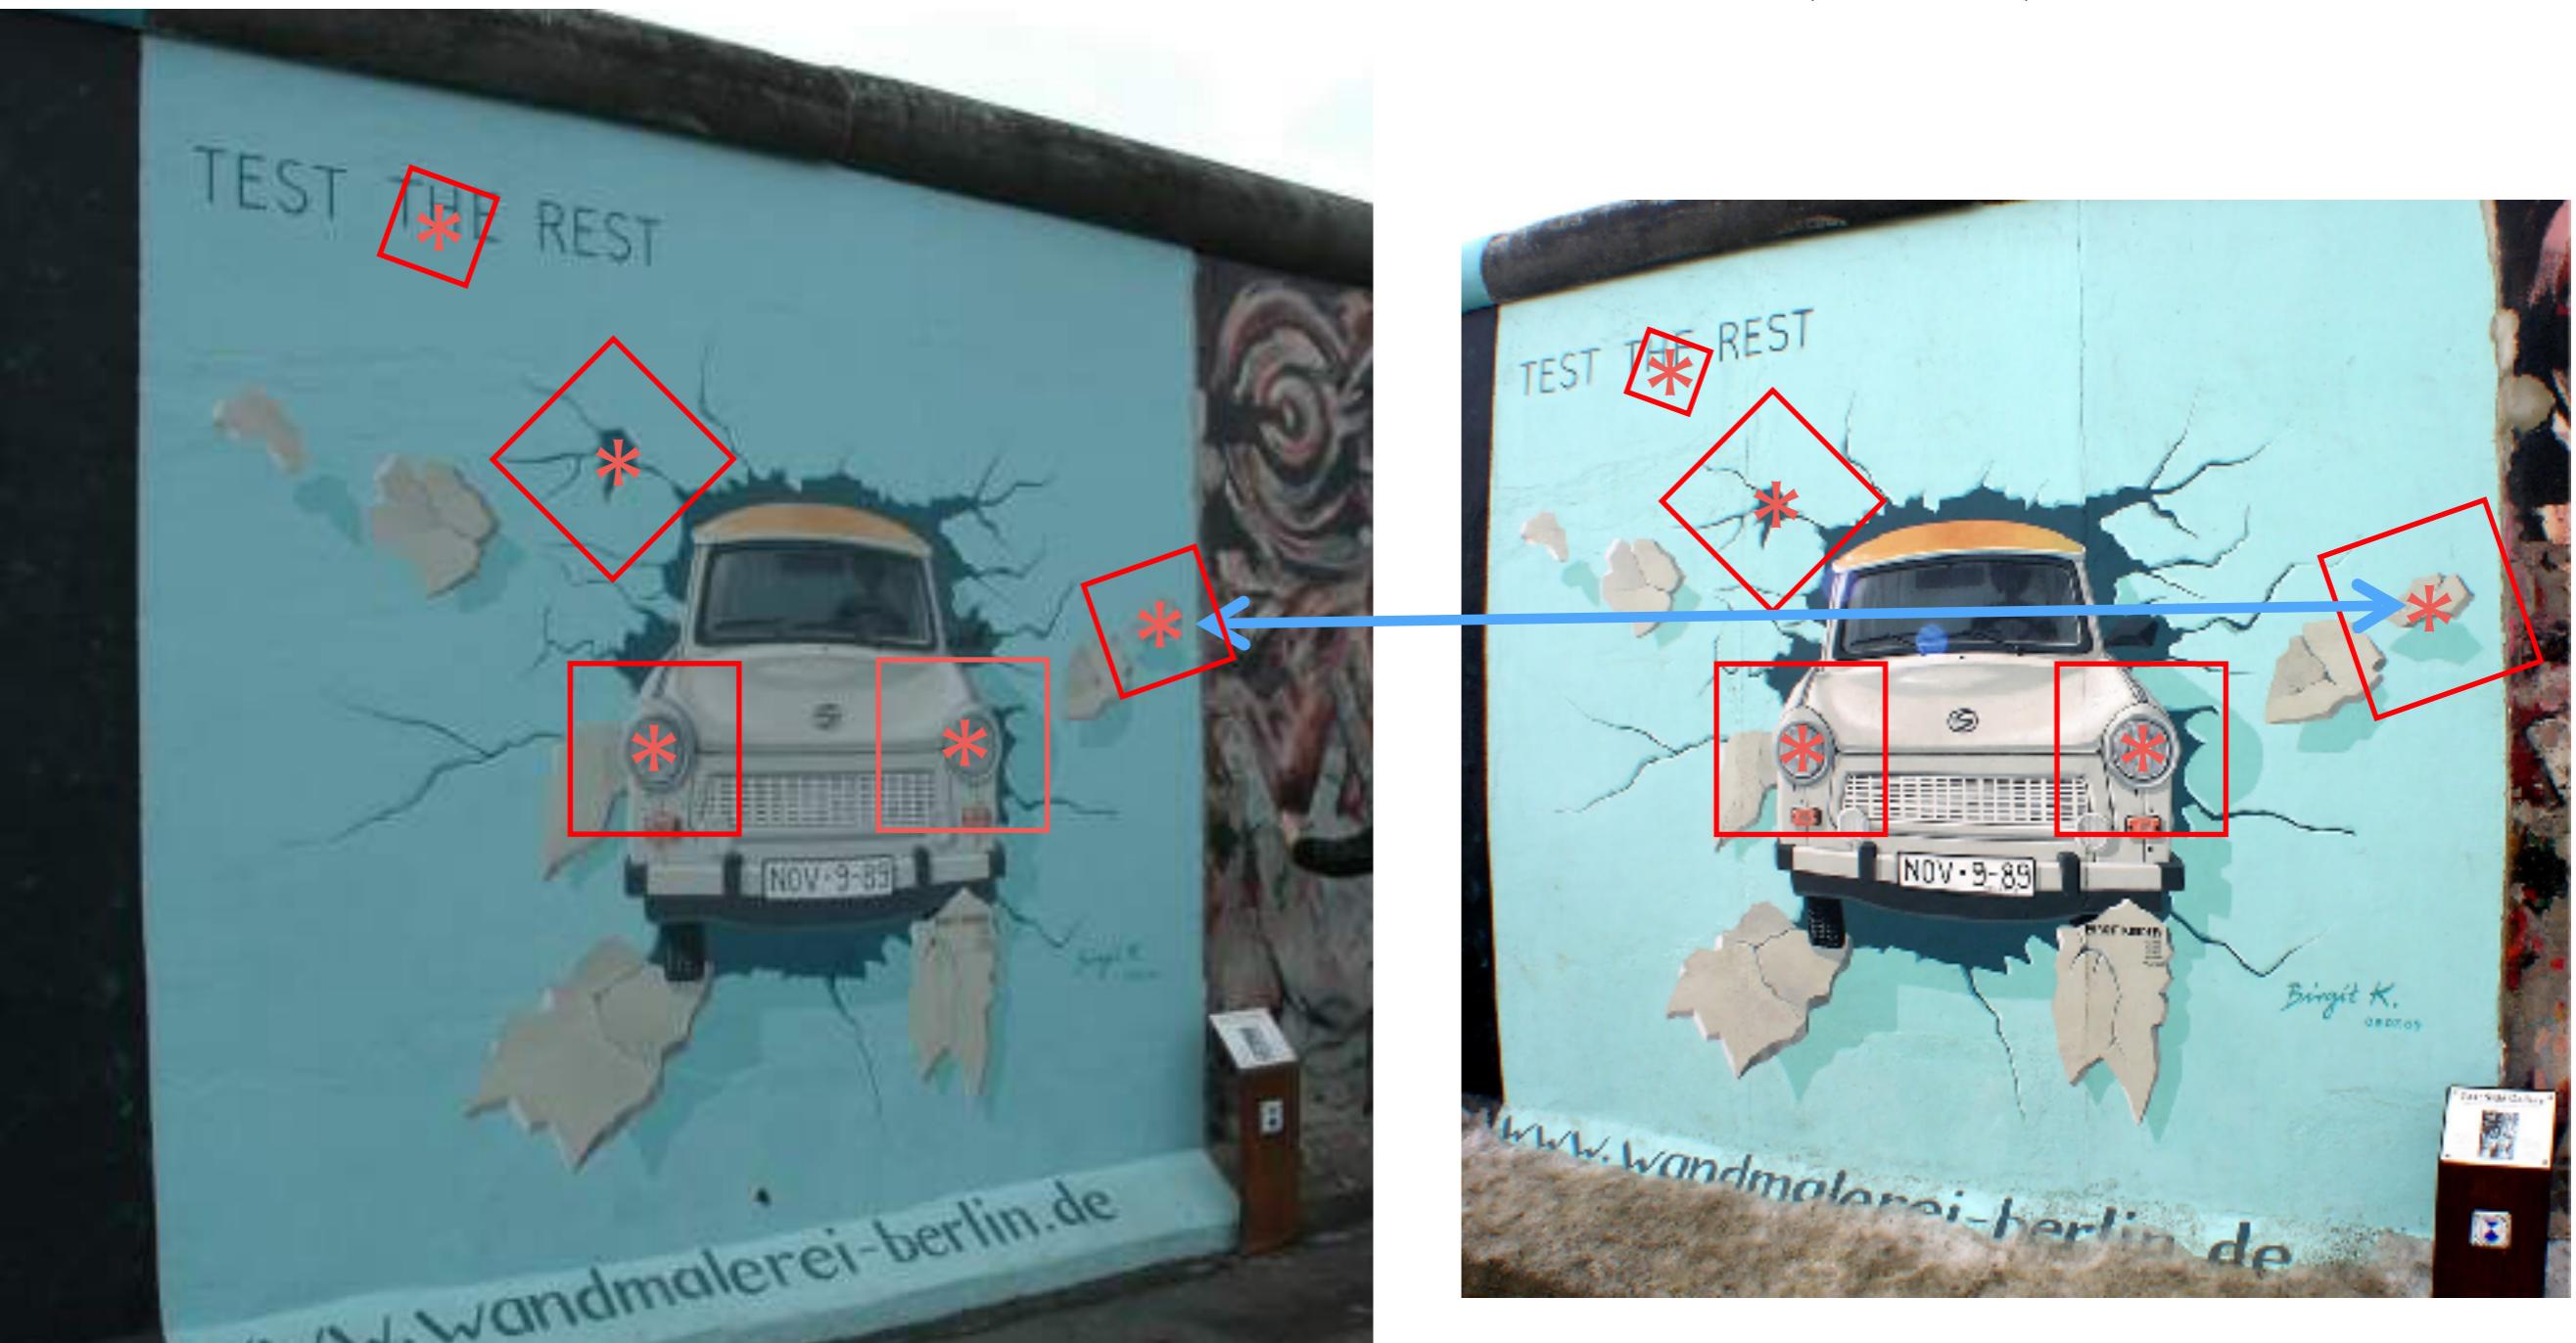
[Lowe, Distinctive Image Features from Scale-Invariant Keypoints, IJCV 2004]

# Scale-Invariant Feature Transform (SIFT) Features



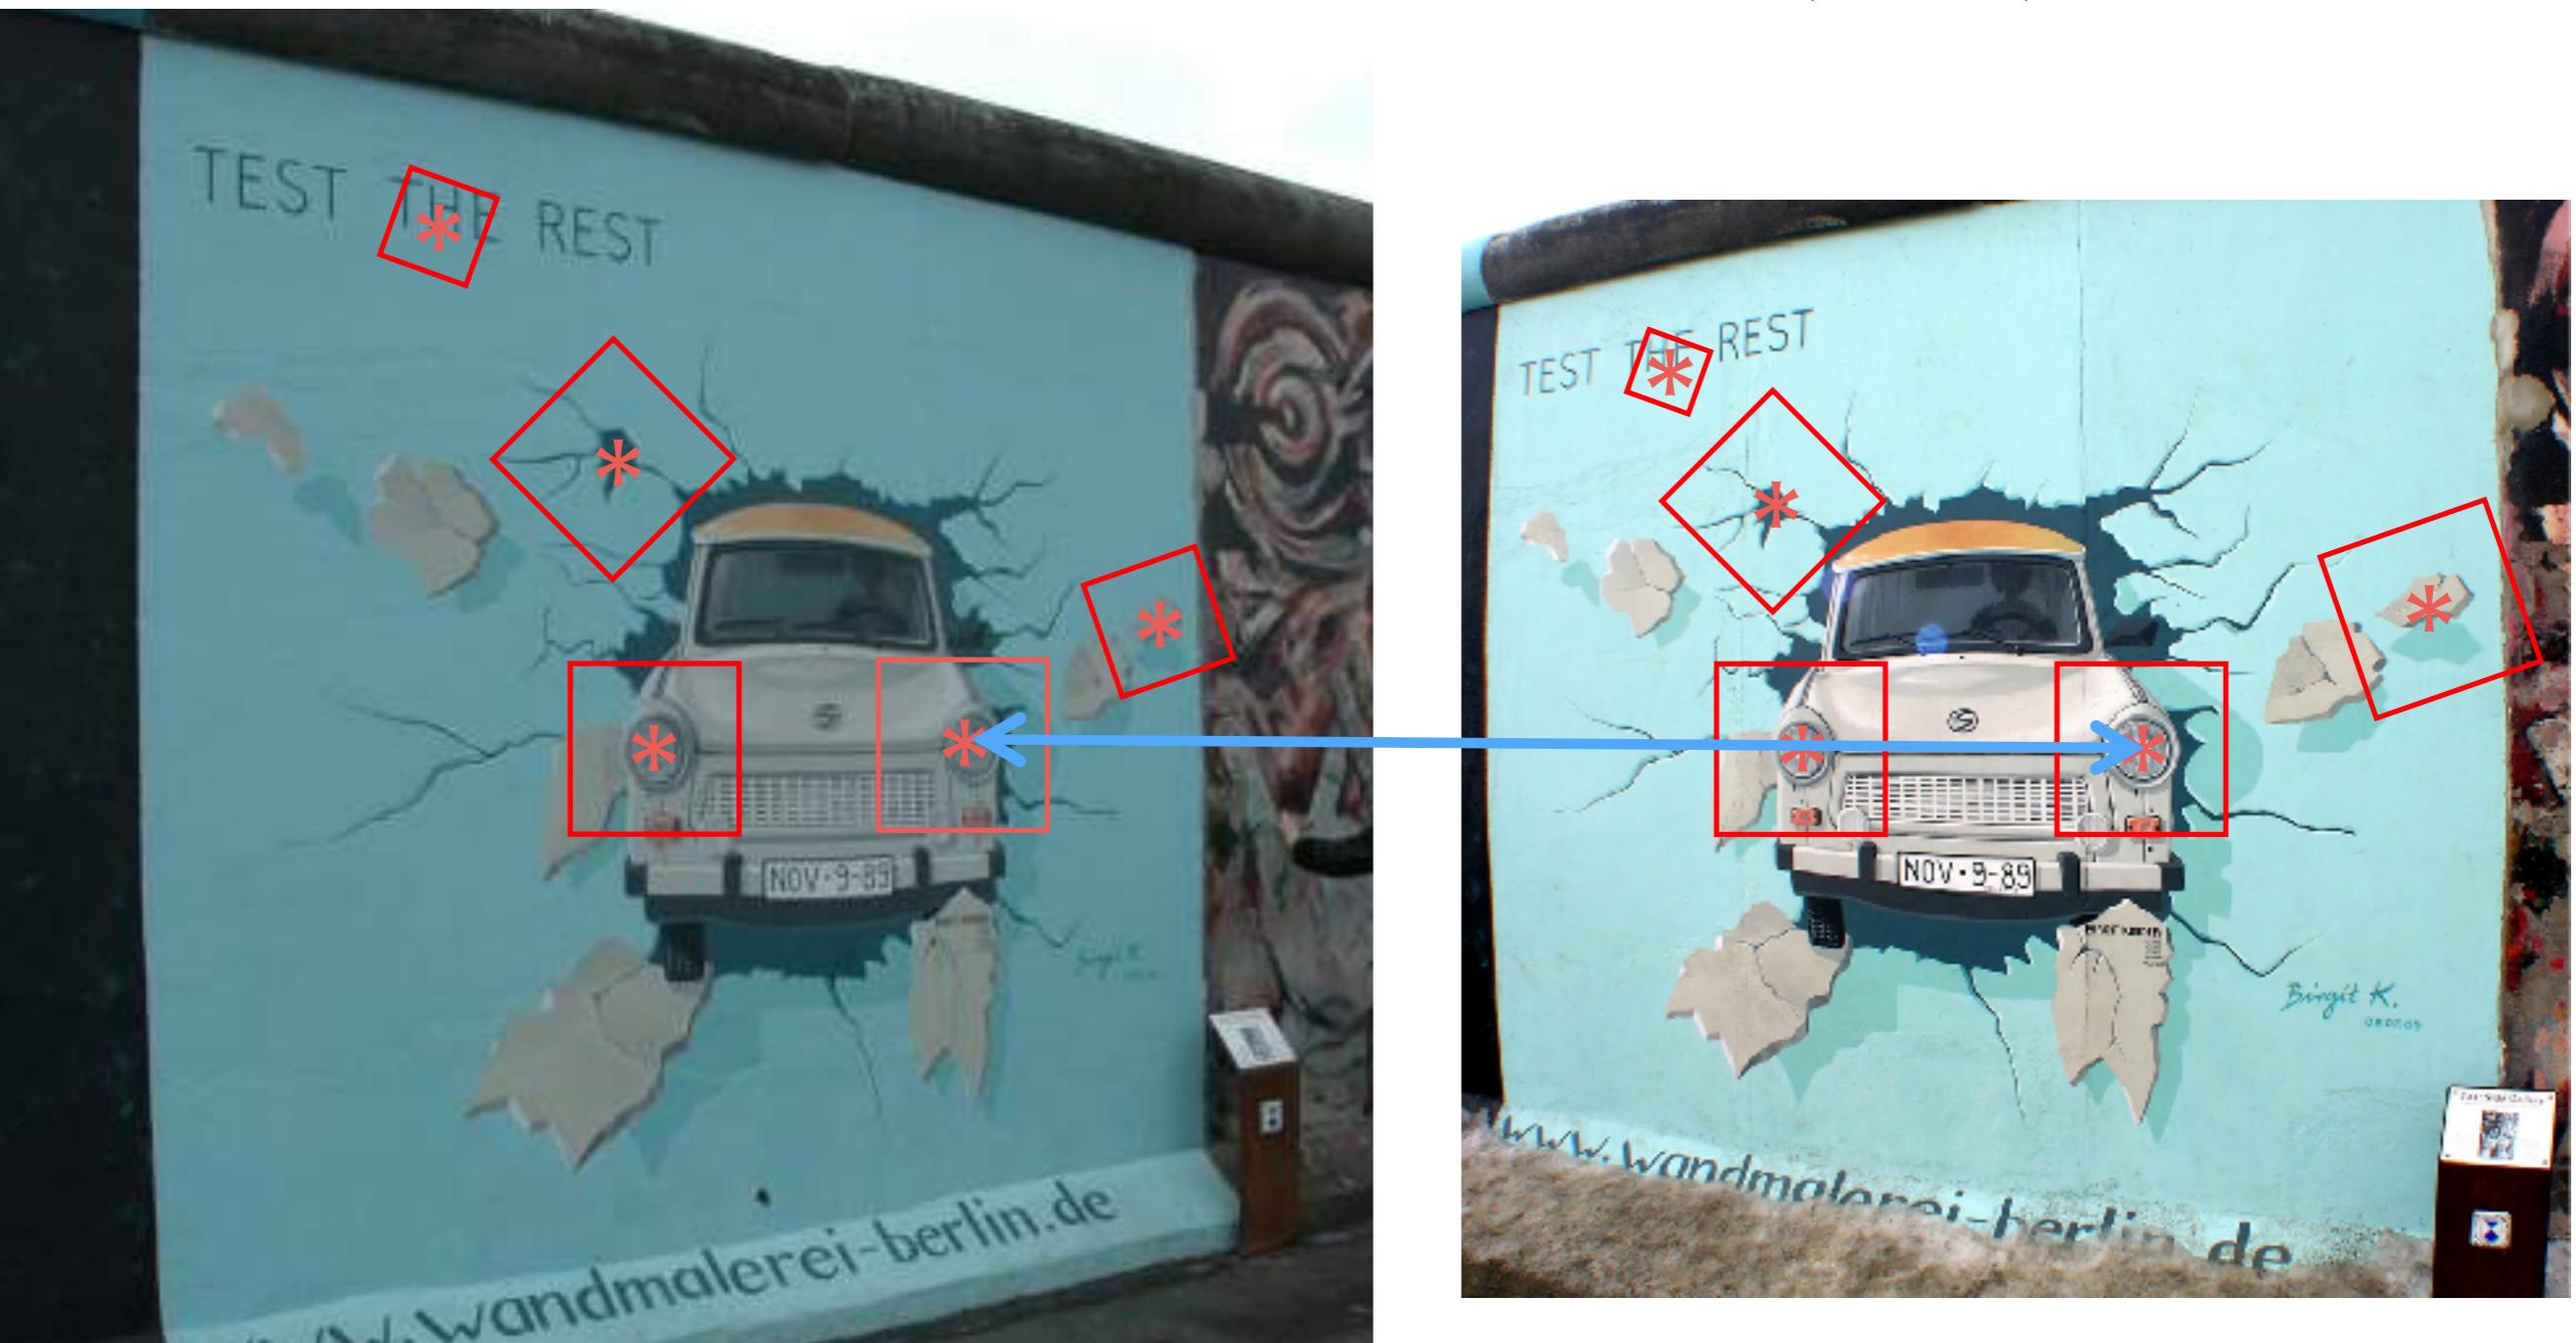
Use descriptor similarity ( $L_2$  distance) to find corresponding points in different images

# Scale-Invariant Feature Transform (SIFT) Features



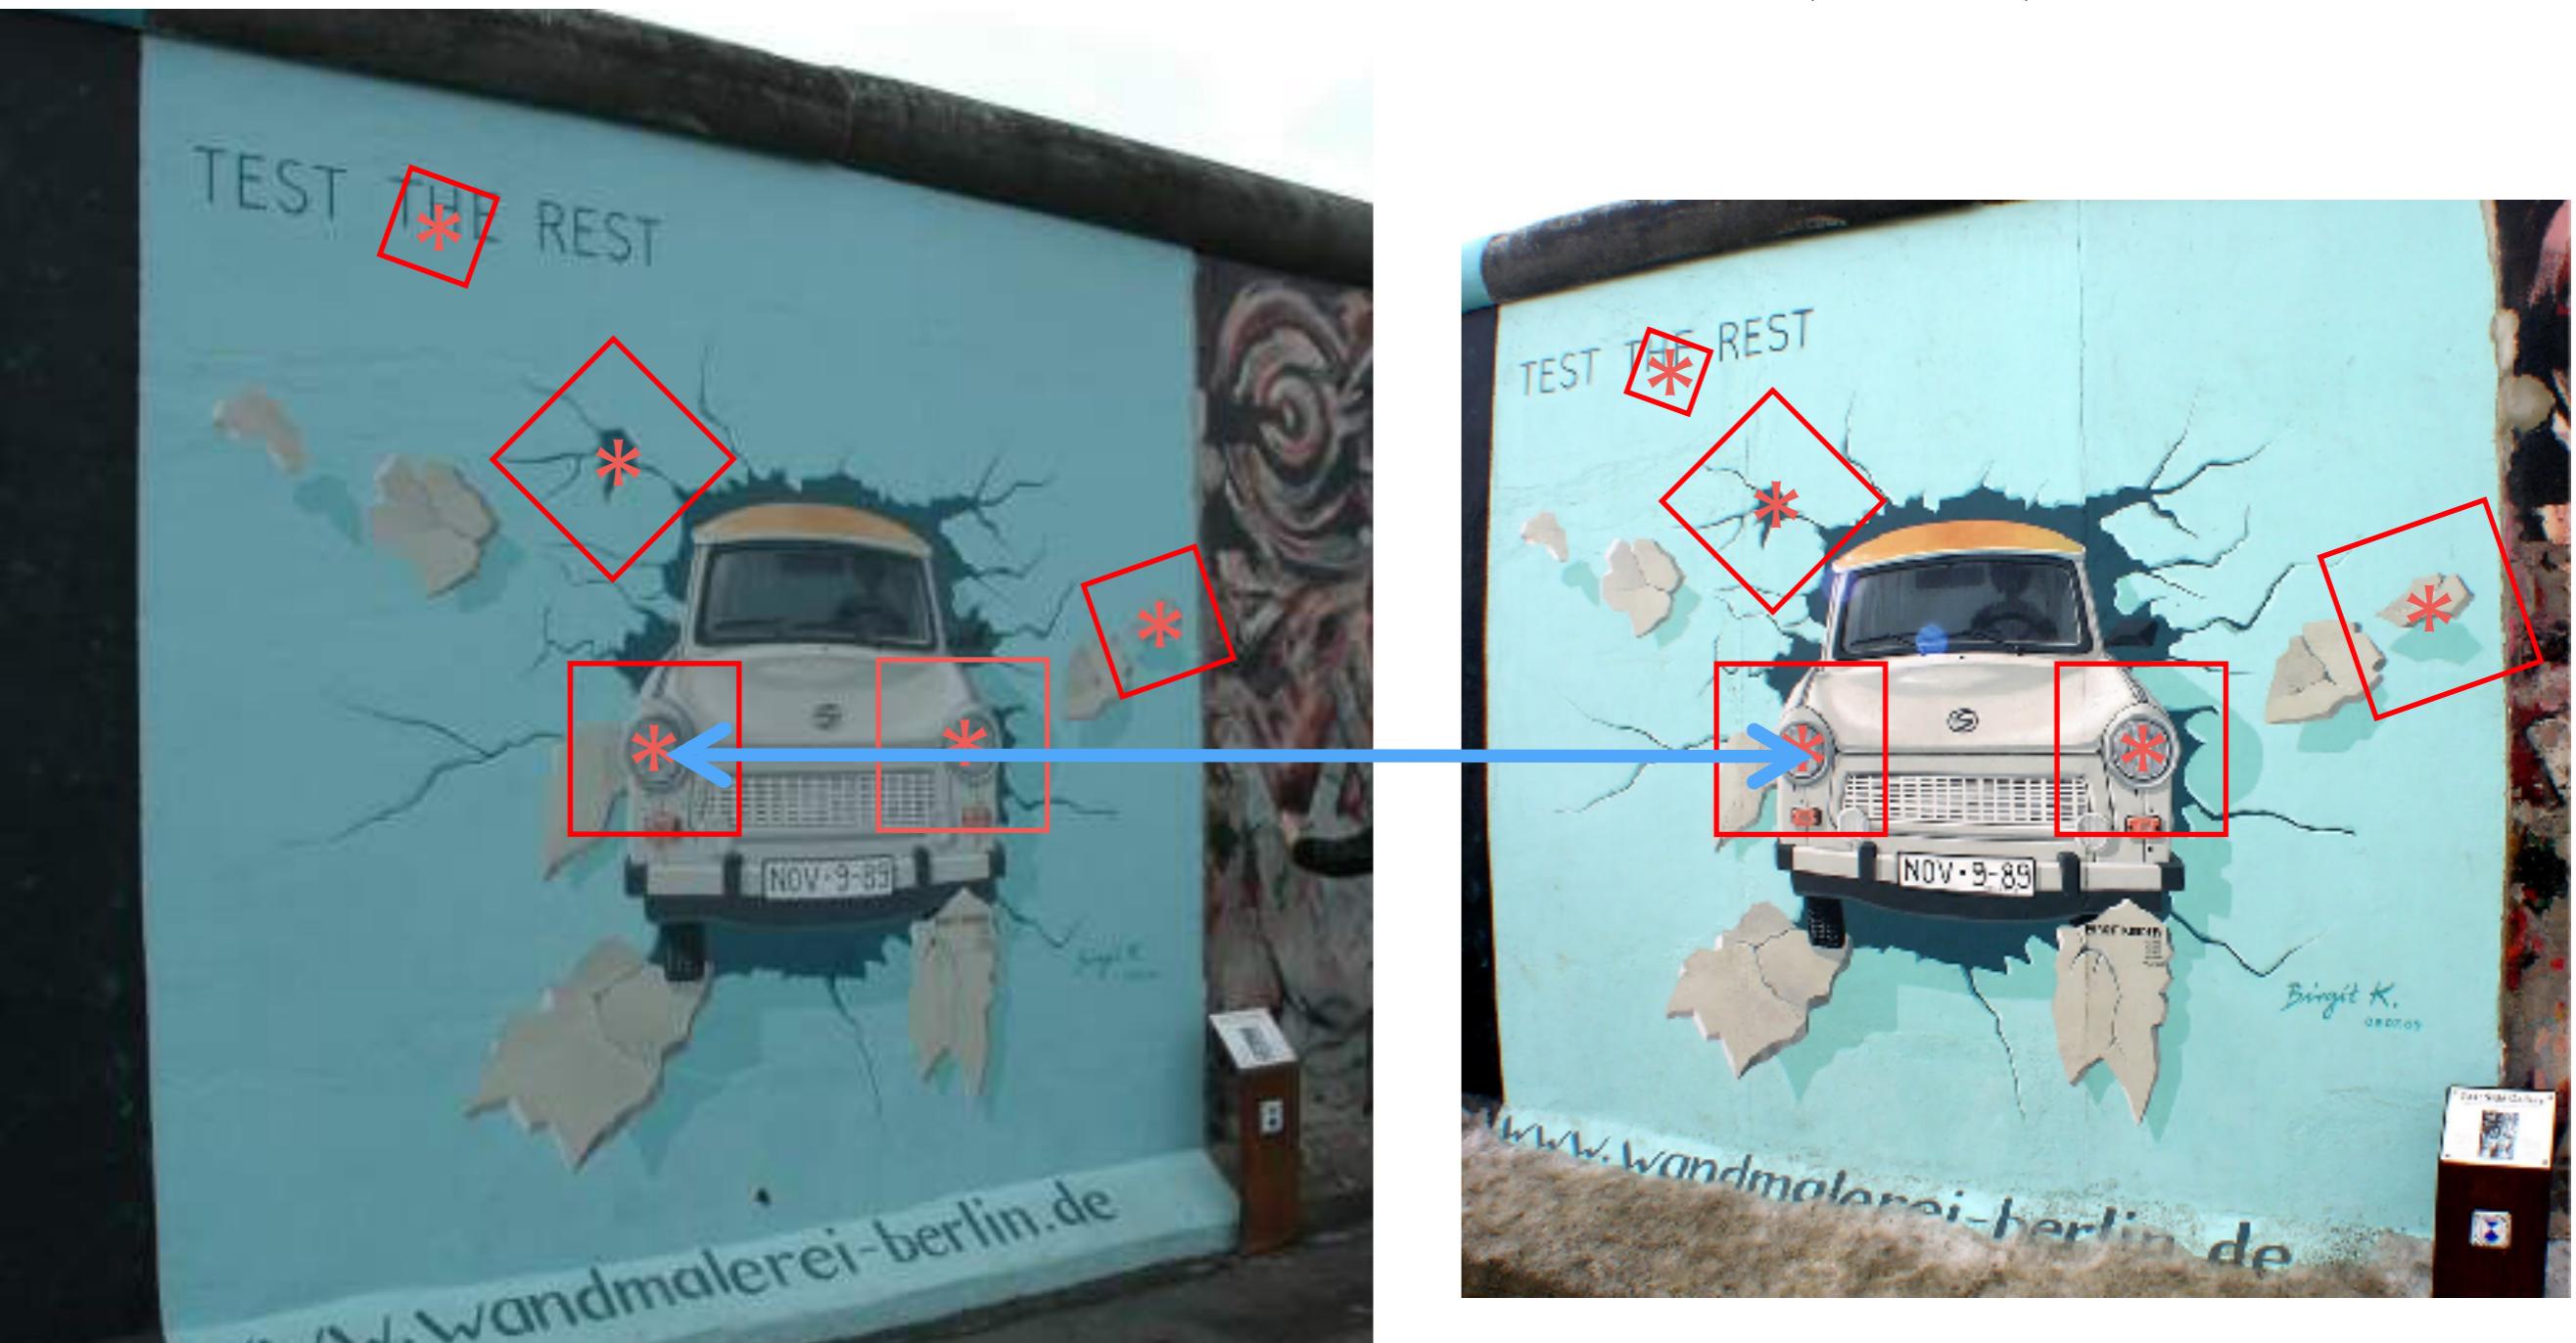
Use descriptor similarity ( $L_2$  distance) to find corresponding points in different images

# Scale-Invariant Feature Transform (SIFT) Features



Use descriptor similarity ( $L_2$  distance) to find corresponding points in different images

# Scale-Invariant Feature Transform (SIFT) Features



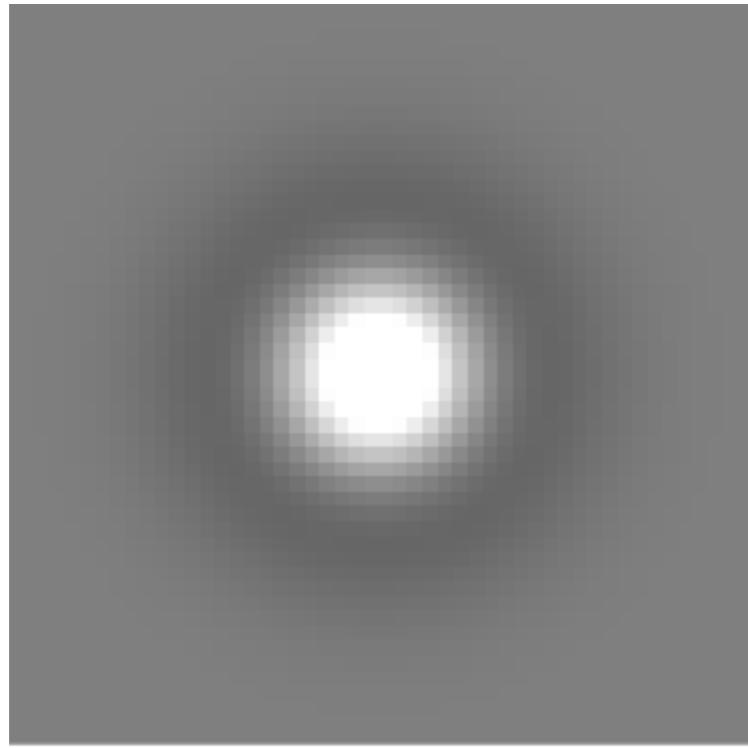
Use descriptor similarity ( $L_2$  distance) to find corresponding points in different images

# Scale-Invariant Feature Transform (SIFT) Features

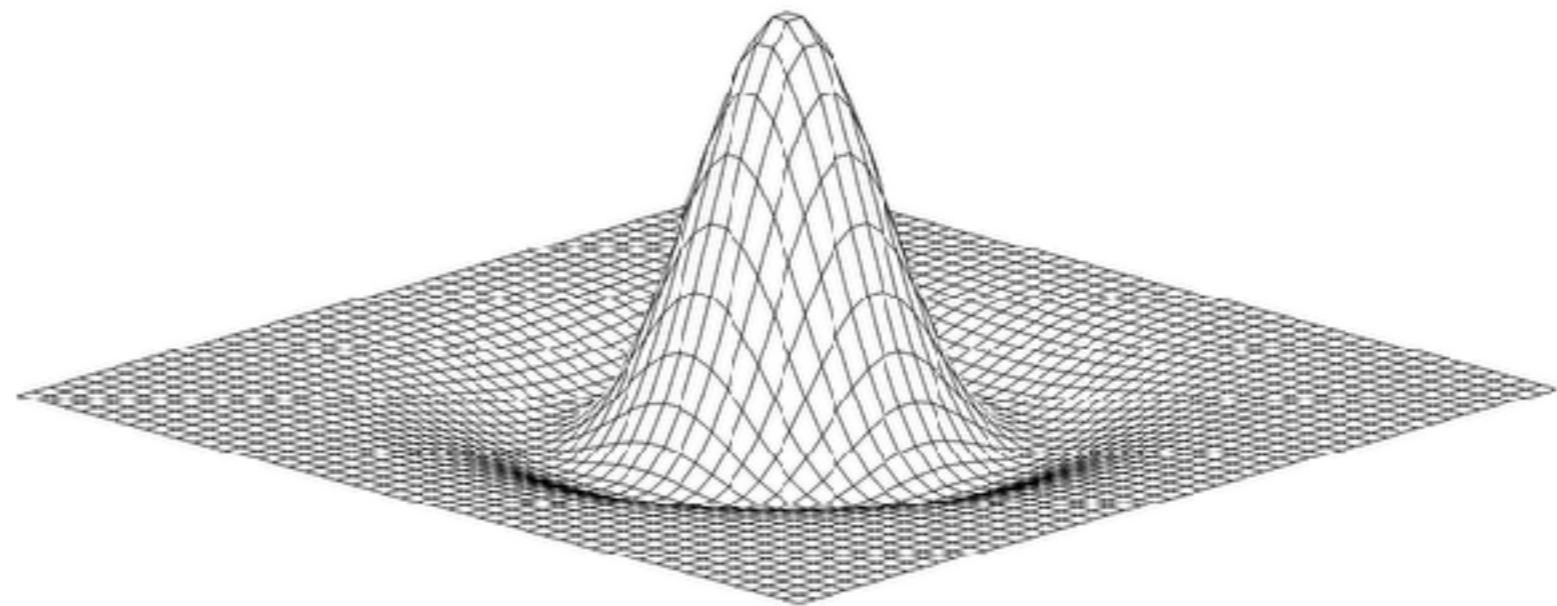
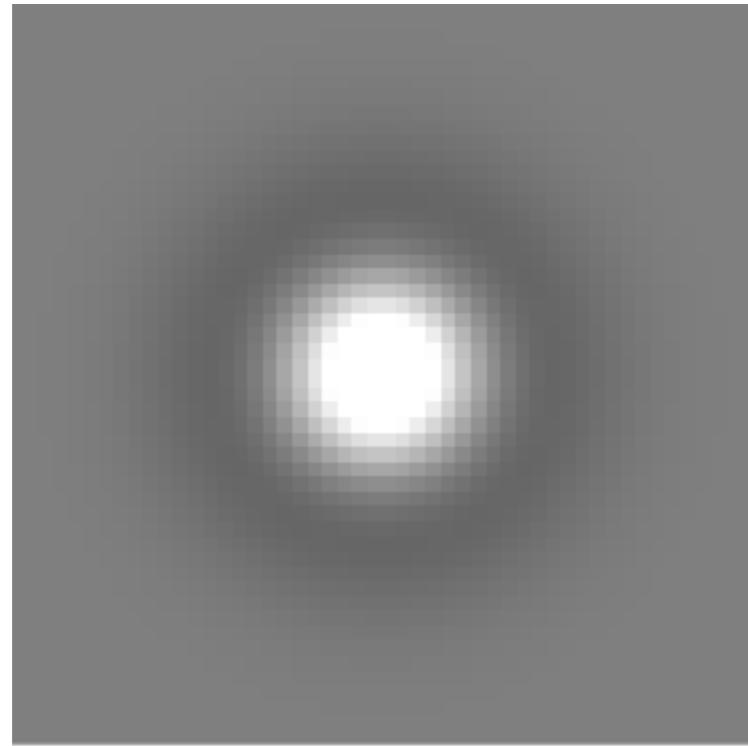


- Find a set of *interest* points
- Extract oriented patch around each interest point
- Compute SIFT descriptor for each patch

# Difference of Gaussians (DoG)



# Difference of Gaussians (DoG)



Difference of two Gaussian filters

- Local maxima at bright blobs
- Local minima at dark blobs

# The Need for Multi-Scale Detections

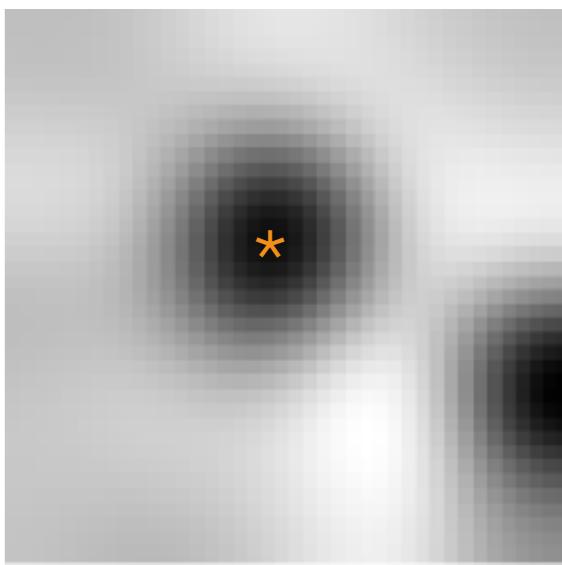
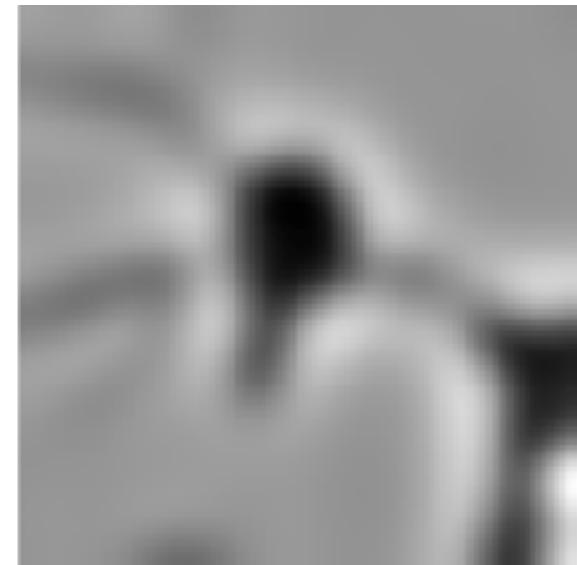
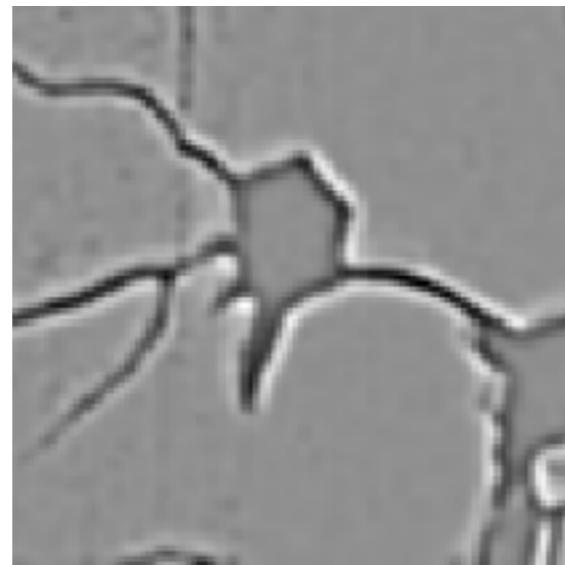
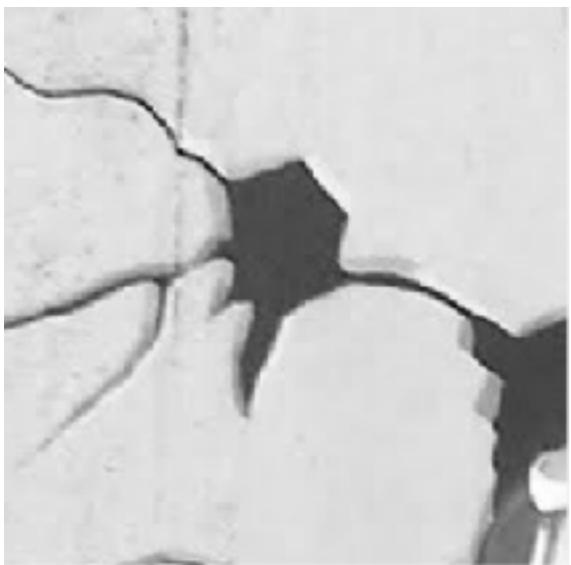
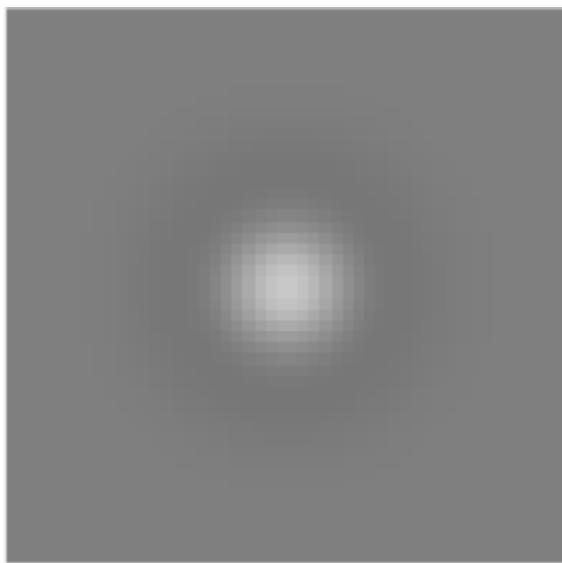
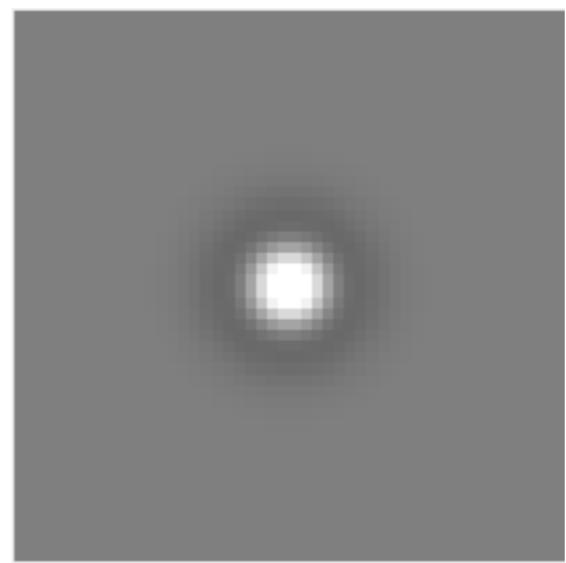
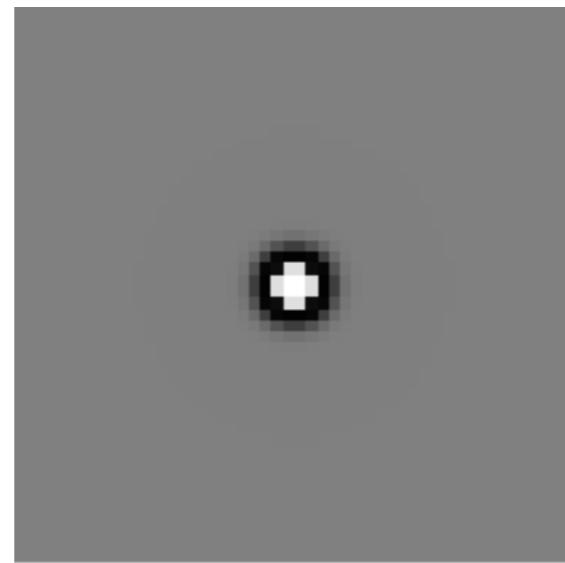


# The Need for Multi-Scale Detections

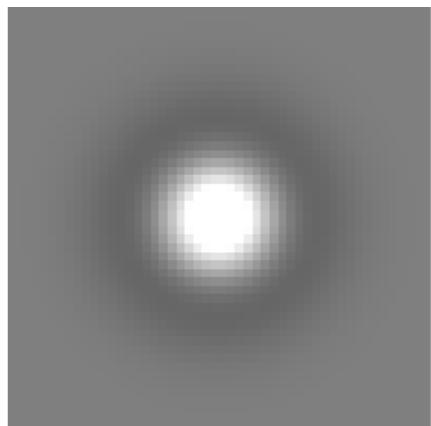
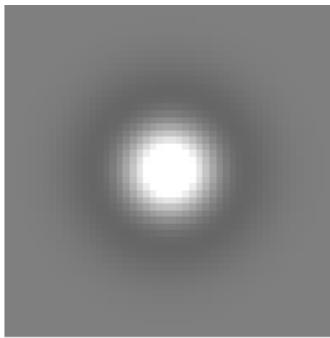
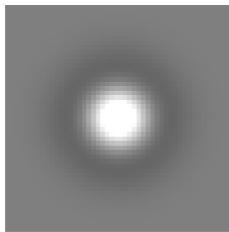


# Difference of Gaussians (DoG)

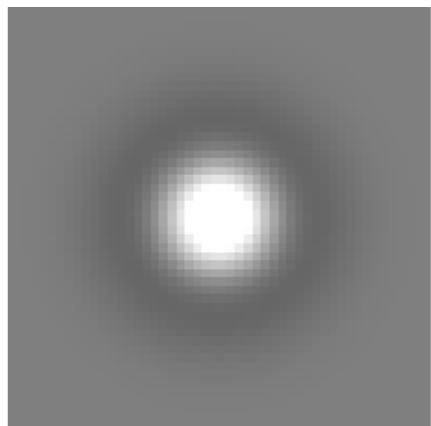
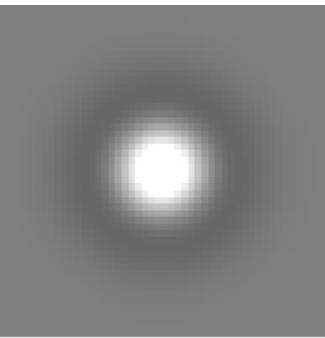
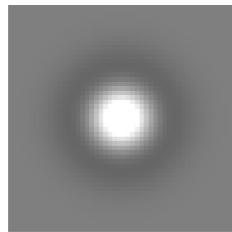
Use filters of different size



# Detect Scale and Position



# Detect Scale and Position

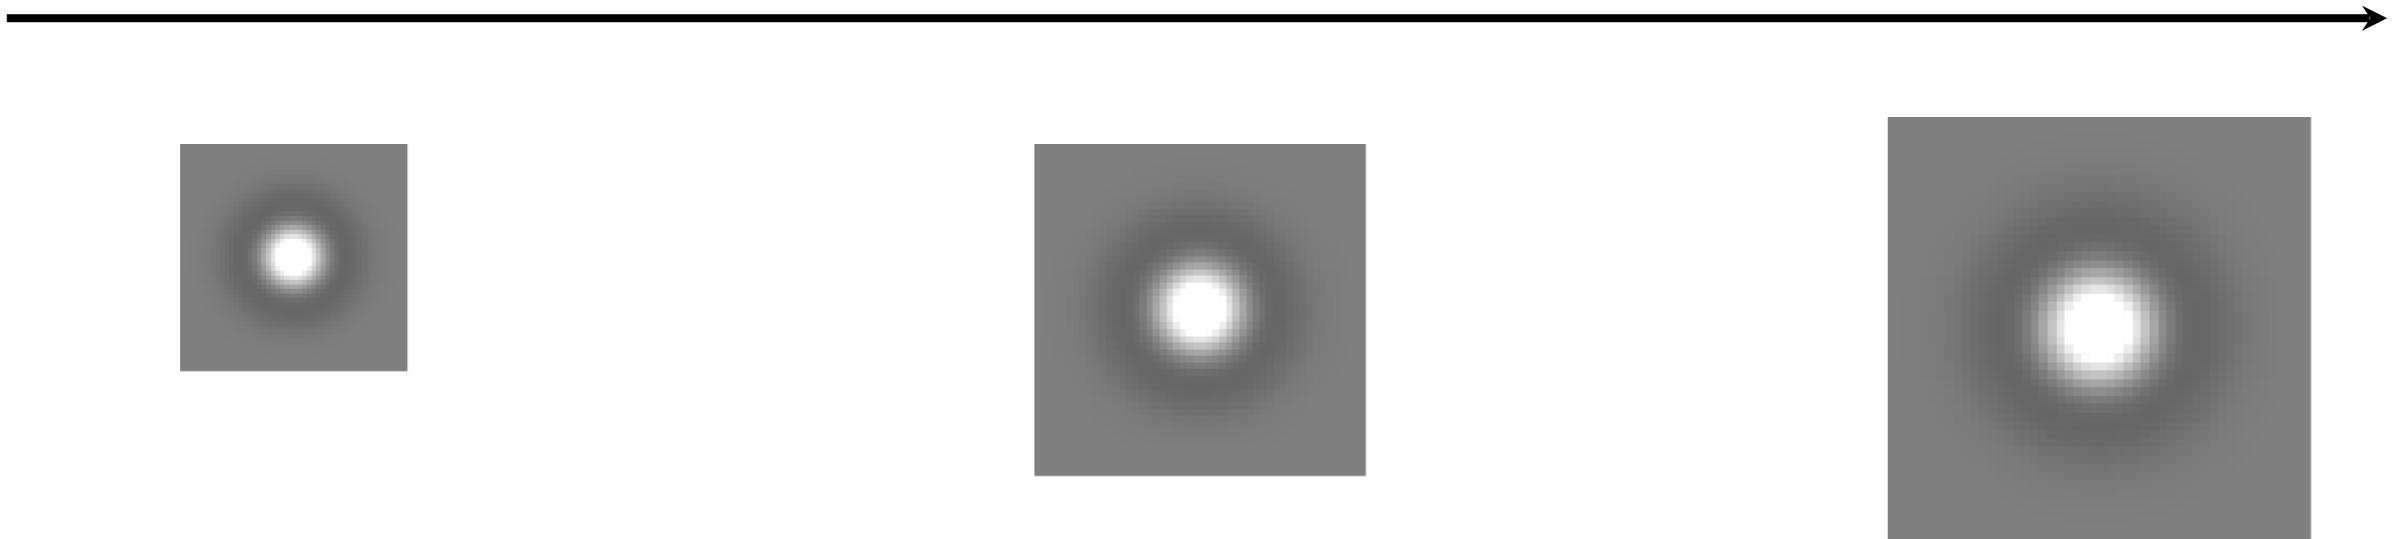


1	0	1	2	3	4	2
2	7	8	2	3	3	3
2	3	1	4	5	6	2
2	7	12	15	20	7	2
2	8	12	25	19	11	3
3	9	11	18	16	7	1
1	2	3	2	1	3	1

2	1	2	2	2	3	3
1	2	3	2	2	3	4
5	11	9	8	9	8	1
1	11	15	20	22	9	5
2	12	16	35	22	11	4
3	9	15	18	21	7	1
2	2	9	8	7	9	1

3	3	3	3	3	3	3	3
3	4	4	4	4	4	4	3
3	5	6	6	6	5	2	
2	7	12	12	12	7	2	
2	5	12	20	13	11	3	
3	9	12	13	13	9	1	
1	2	5	5	5	3	1	

# Detect Scale and Position



1	0	1	2	3	4	2
2	7	8	2	3	3	3
2	3	1	4	5	6	2
2	7	12	15	20	7	2
2	8	12	25	19	11	3
3	9	11	18	16	7	1
1	2	3	2	1	3	1

2	1	2	2	2	3	3
1	2	3	2	2	3	4
5	11	9	8	9	8	1
1	11	15	20	22	9	5
2	12	16	35	22	11	4
3	9	15	18	21	7	1
2	2	9	8	7	9	1

3	3	3	3	3	3	3
3	4	4	4	4	4	3
3	5	6	6	6	5	2
2	7	12	12	12	7	2
2	5	12	20	13	11	3
3	9	12	13	13	9	1
1	2	5	5	5	3	1

Non-maximum suppression in image **and scale!**

# Sub-Pixel Refinement

2<sup>nd</sup> order curve fitting (1D):

$$f(x) \approx f(0) + x \cdot f'(0) + \frac{x^2}{2} f''(0)$$

# Sub-Pixel Refinement

2<sup>nd</sup> order curve fitting (1D):

$$f(x) \approx f(0) + x \cdot f'(0) + \frac{x^2}{2} f''(0)$$

2<sup>nd</sup> order surface fitting (3D):

$$f(x, y, \sigma) \approx f(0, 0, 0) + (x \quad y \quad \sigma) \cdot \nabla f(0, 0, 0) + \frac{1}{2} (x \quad y \quad \sigma) H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

# Sub-Pixel Refinement

2<sup>nd</sup> order curve fitting (1D):

$$f(x) \approx f(0) + x \cdot f'(0) + \frac{x^2}{2} f''(0)$$

2<sup>nd</sup> order surface fitting (3D):

$$f(x, y, \sigma) \approx f(0, 0, 0) + (x \quad y \quad \sigma) \cdot \nabla f(0, 0, 0) + \frac{1}{2} (x \quad y \quad \sigma) \mathbf{H}(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

Hessian matrix:

$$\mathbf{H}(0, 0, 0) = \begin{pmatrix} f''_{xx}(0, 0, 0) & f''_{xy}(0, 0, 0) & f''_{x\sigma}(0, 0, 0) \\ f''_{xy}(0, 0, 0) & f''_{yy}(0, 0, 0) & f''_{y\sigma}(0, 0, 0) \\ f''_{x\sigma}(0, 0, 0) & f''_{y\sigma}(0, 0, 0) & f''_{\sigma\sigma}(0, 0, 0) \end{pmatrix}$$

# Sub-Pixel Refinement

2<sup>nd</sup> order surface fitting (3D):

$$f(x, y, \sigma) \approx f(0, 0, 0) + (x \quad y \quad \sigma) \cdot \nabla f(0, 0, 0) + \frac{1}{2} (x \quad y \quad \sigma) H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

# Sub-Pixel Refinement

2<sup>nd</sup> order surface fitting (3D):

$$f(x, y, \sigma) \approx f(0, 0, 0) + (x \ y \ \sigma) \cdot \nabla f(0, 0, 0) + \frac{1}{2} (x \ y \ \sigma) H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

Set derivative to 0:

$$\begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} = \nabla f(0, 0, 0) + H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

offset

# Sub-Pixel Refinement

2<sup>nd</sup> order surface fitting (3D):

$$f(x, y, \sigma) \approx f(0, 0, 0) + (x \ y \ \sigma) \cdot \nabla f(0, 0, 0) + \frac{1}{2} (x \ y \ \sigma) H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

Set derivative to 0:

$$\begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} = \nabla f(0, 0, 0) + H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

Solve small linear system:

$$\begin{pmatrix} x \\ y \\ \sigma \end{pmatrix} = -H(0, 0, 0)^{-1} \nabla f(0, 0, 0)$$

# Sub-Pixel Refinement

2<sup>nd</sup> order surface fitting (3D):

$$f(\underbrace{x, y, \sigma}_{\text{underlined}}) \approx f(0, 0, 0) + (x \ y \ \sigma) \cdot \nabla f(0, 0, 0) + \frac{1}{2} (x \ y \ \sigma) H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

Set derivative to 0:

$$\begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} = \nabla f(0, 0, 0) + H(0, 0, 0) \begin{pmatrix} x \\ y \\ \sigma \end{pmatrix}$$

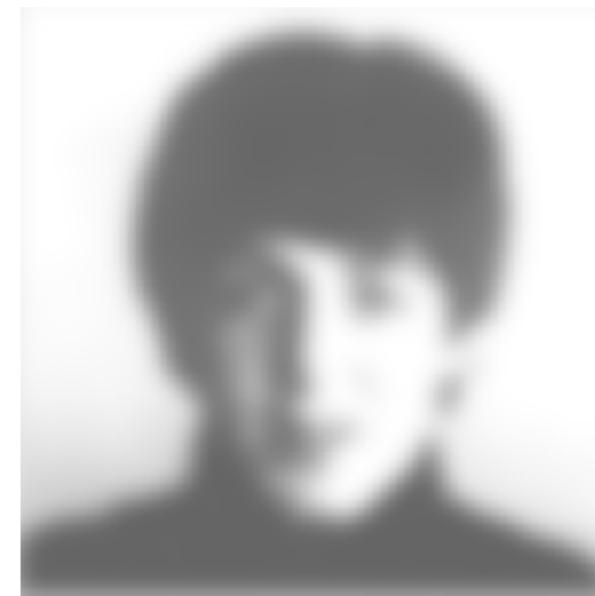
Solve small linear system:

$$\begin{pmatrix} x \\ y \\ \sigma \end{pmatrix} = -H(0, 0, 0)^{-1} \nabla f(0, 0, 0)$$

Potentially repeat for a few iterations

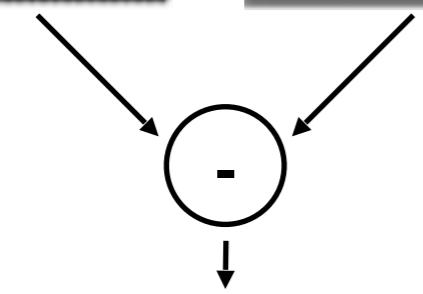
# Efficient implementation

Gaussian filtering

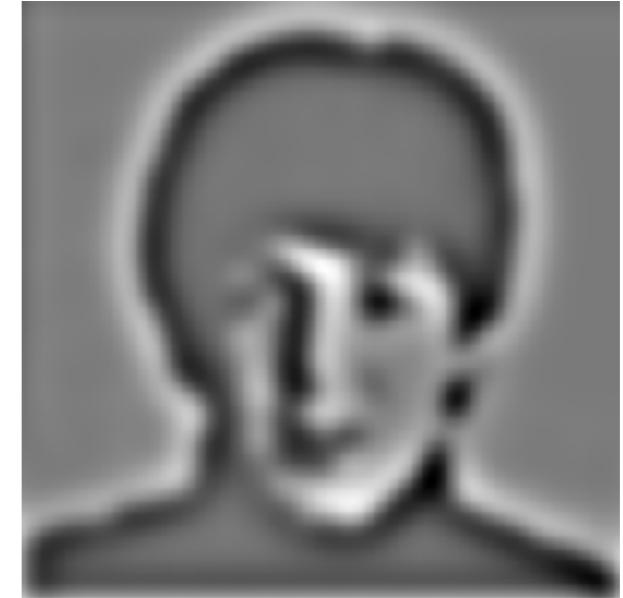
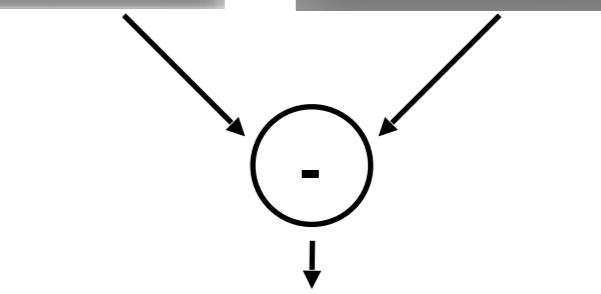
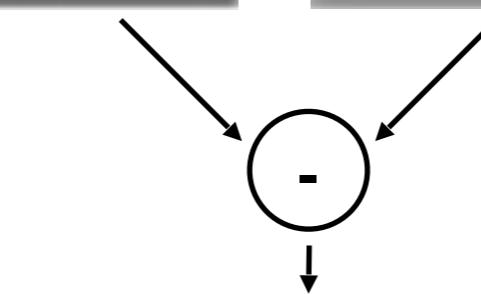


# Efficient implementation

Gaussian filtering



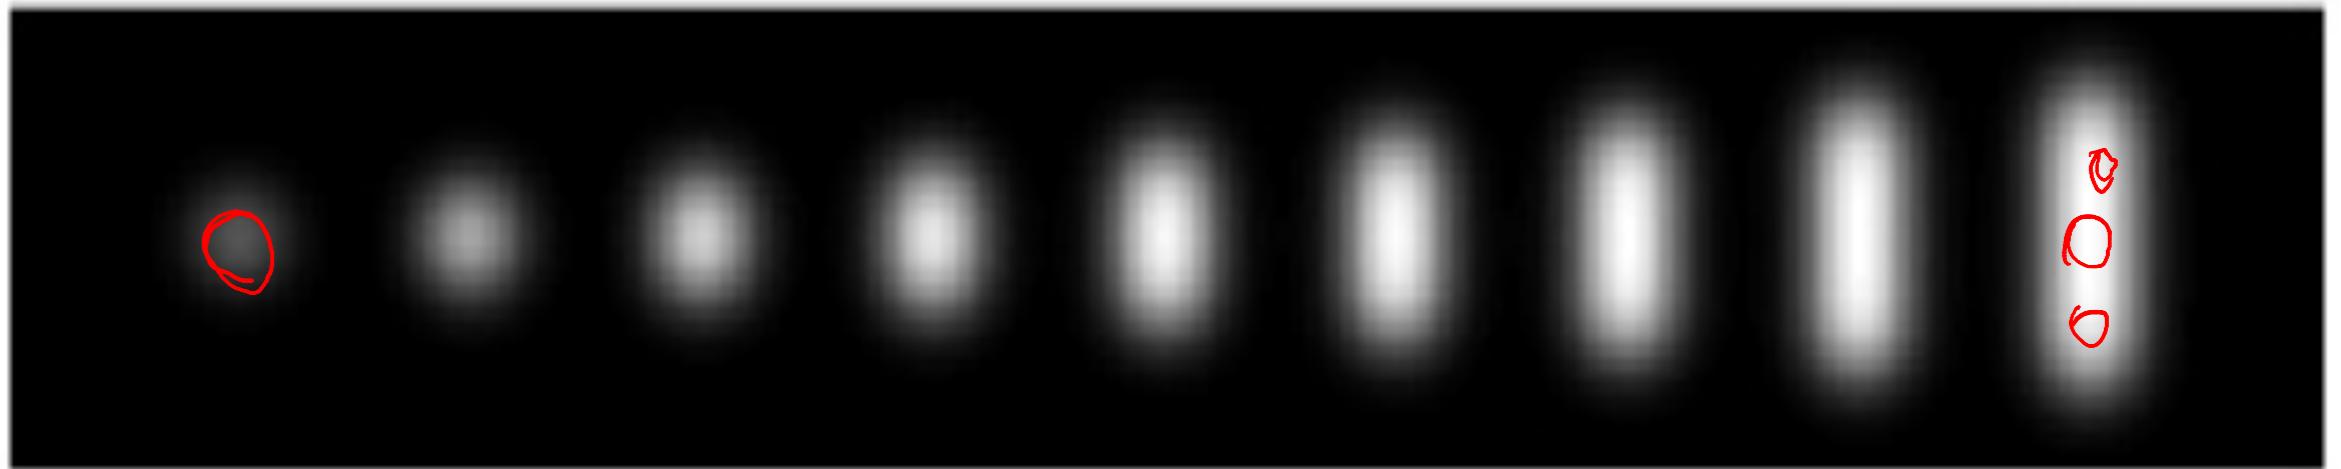
Difference  
of  
Gaussians



# Avoiding Edges



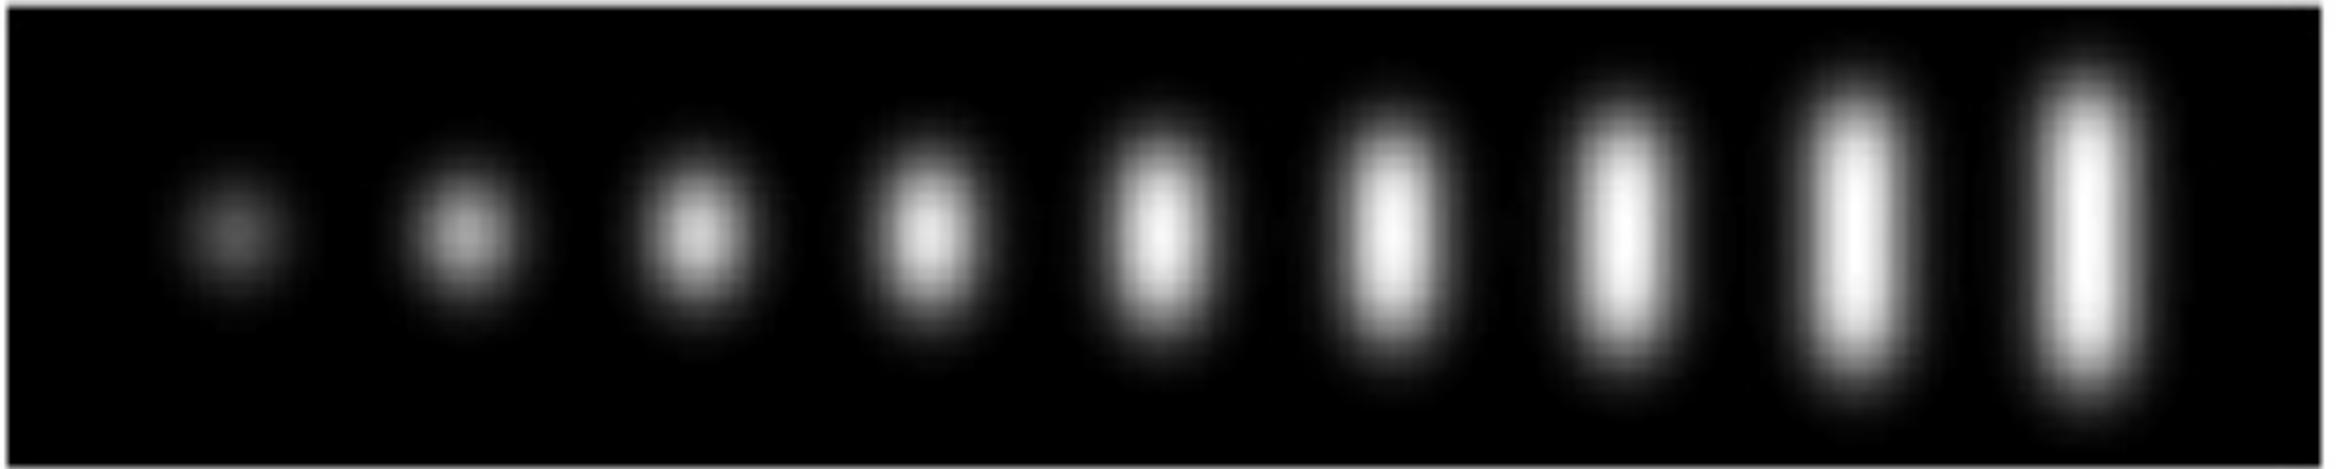
# Avoiding Edges



well-localized blob

~~edge~~

# Avoiding Edges



well-localized blob

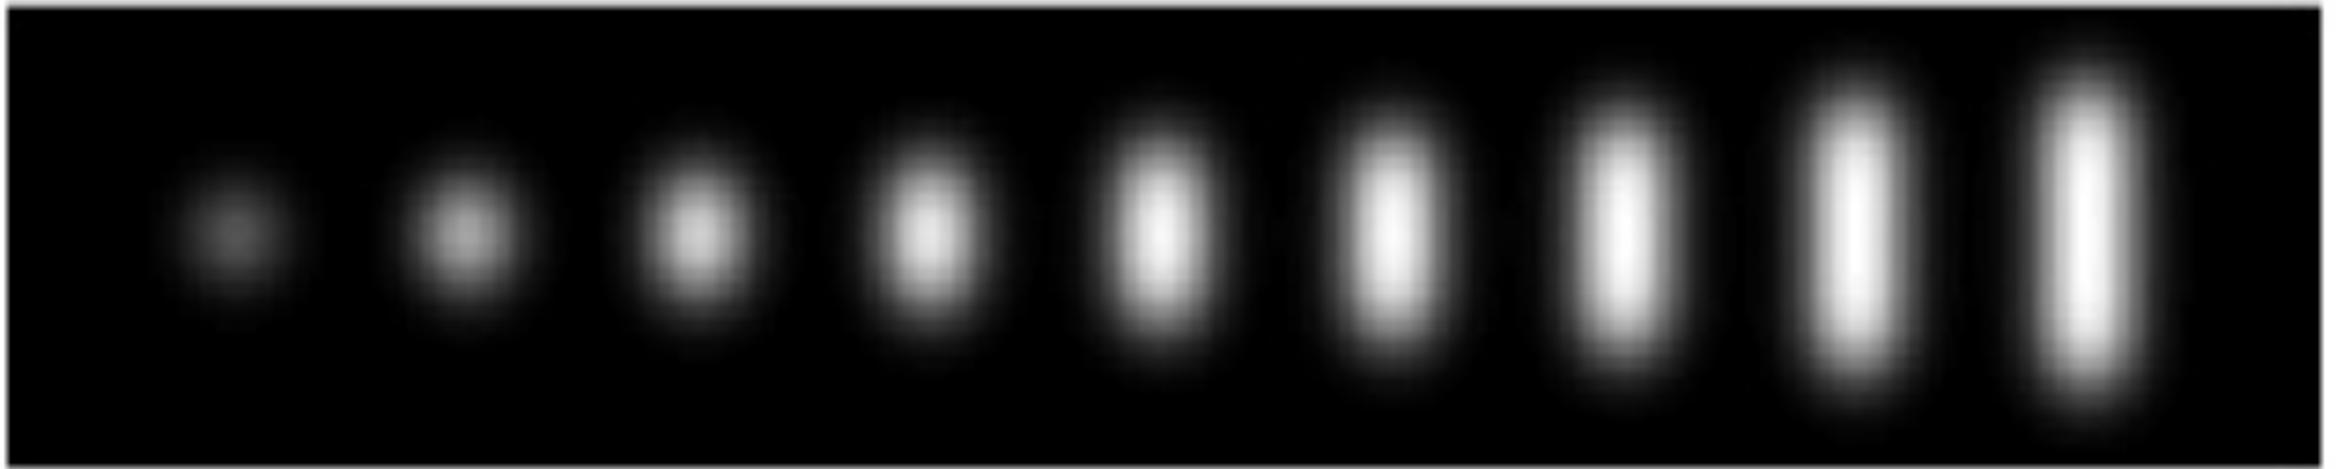
edge

Consider 2D Hessian around detected keypoint (after refinement):

$$\mathbf{H} = \begin{pmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{pmatrix}$$

*2nd order derivatives*

# Avoiding Edges



well-localized blob

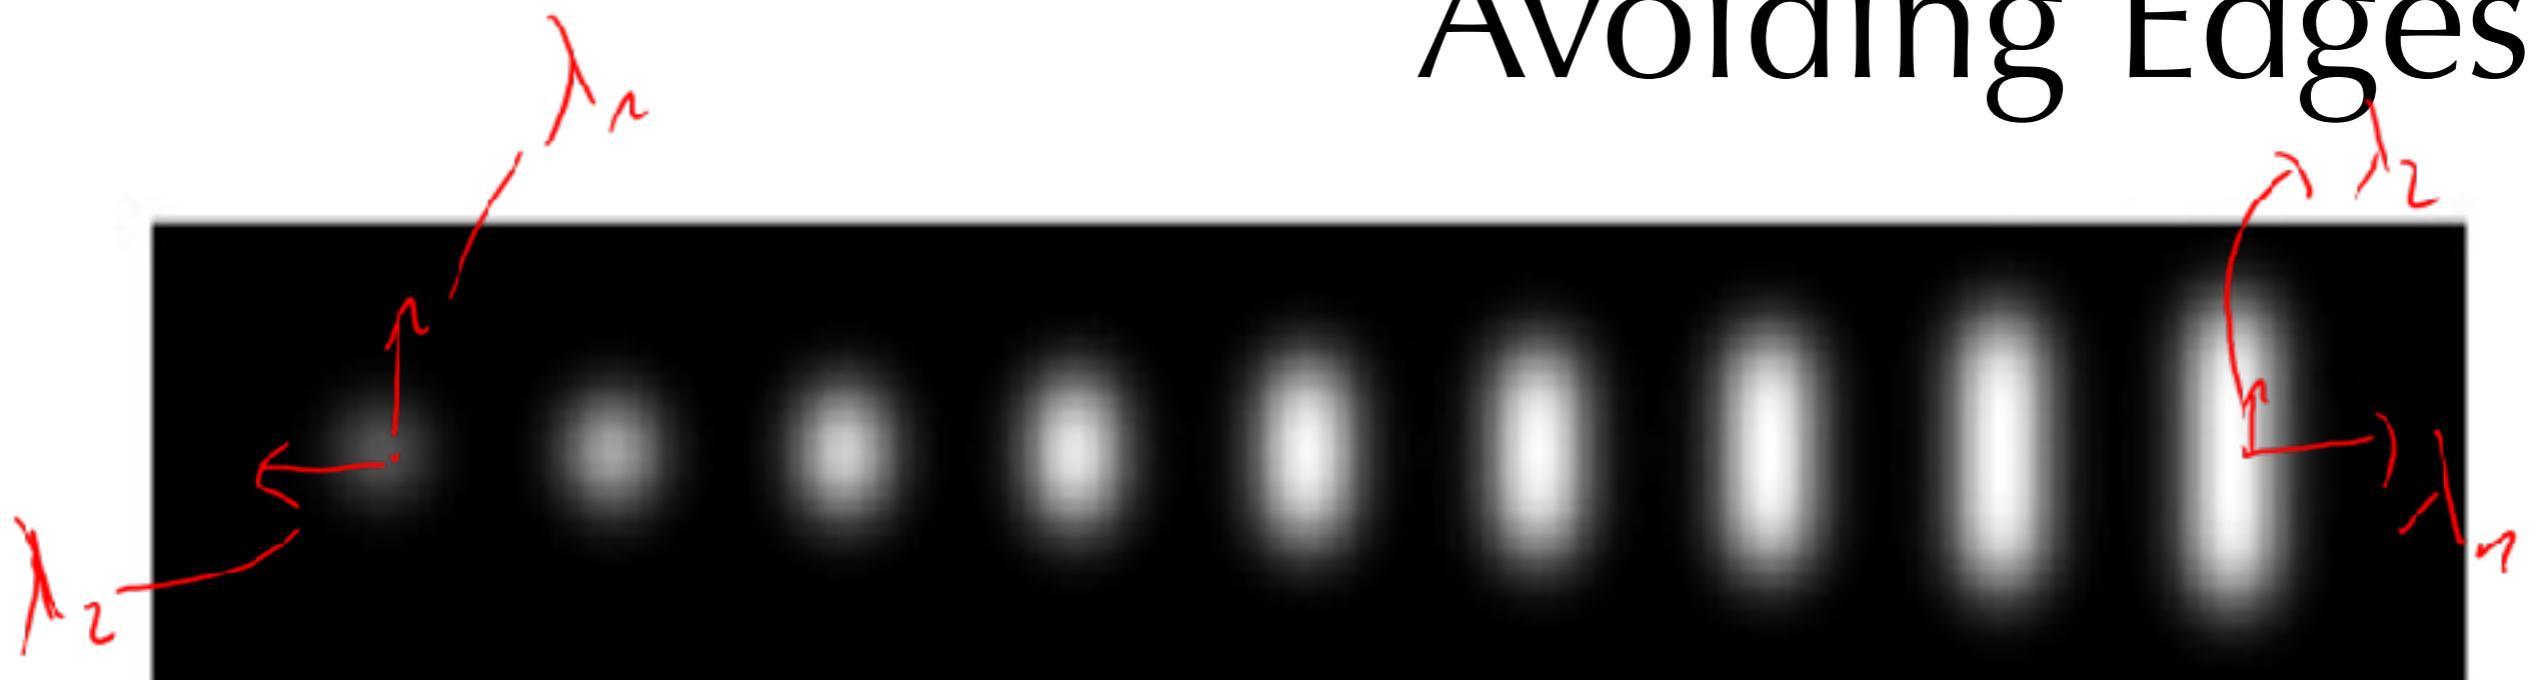
edge

Consider 2D Hessian around detected keypoint (after refinement):

$$\mathbf{H} = \begin{pmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{pmatrix}$$

Eigenvalues  $\lambda_1 \geq \lambda_2$  of  $\mathbf{H}$  proportional to *principal curvature*

# Avoiding Edges



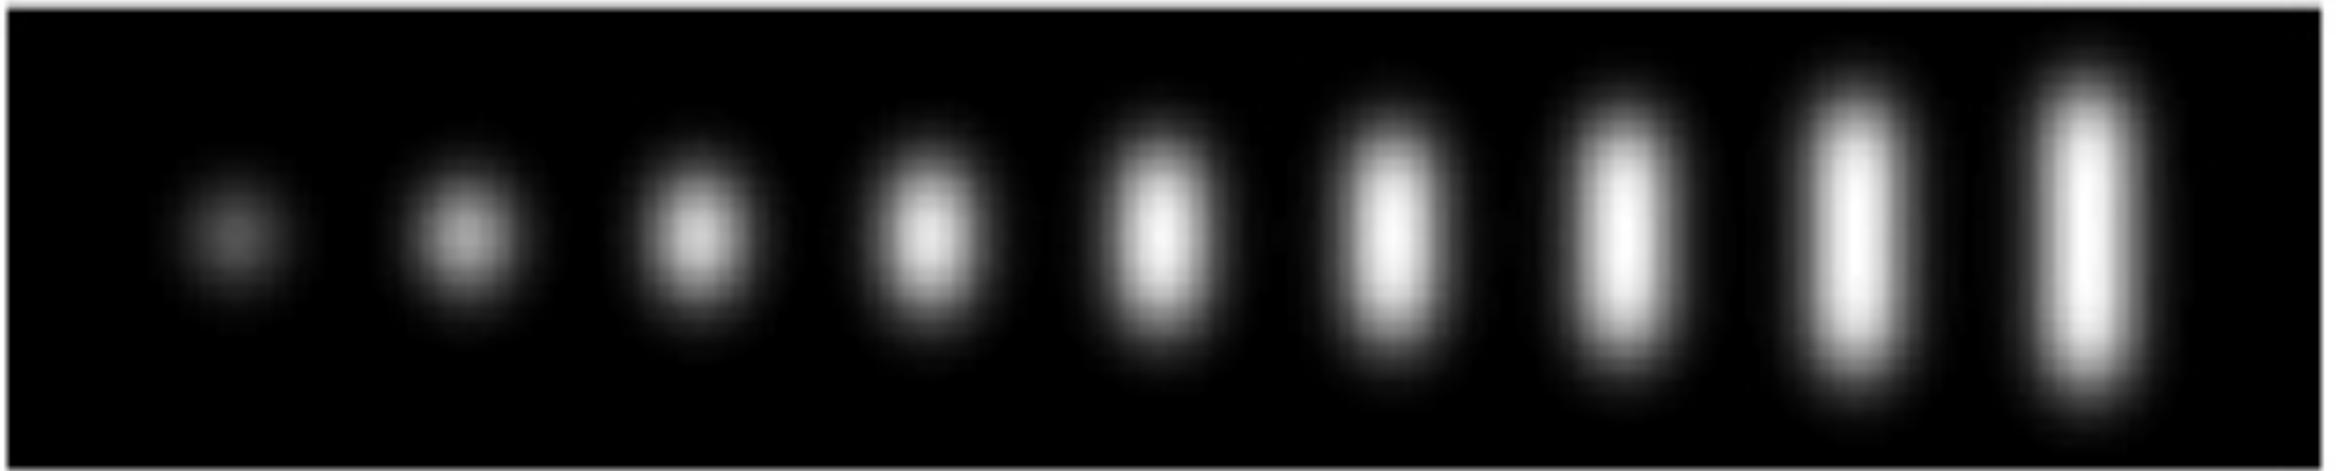
well-localized blob

$$\lambda_n \approx \lambda_z \gg 0$$

edge

$$\lambda_n \gg \lambda_z$$

# Avoiding Edges



well-localized blob

$$\lambda_1 \approx \lambda_2$$

Detecting edges: Look at ratio of Eigenvalues ( $\lambda_1 \geq \lambda_2$ )

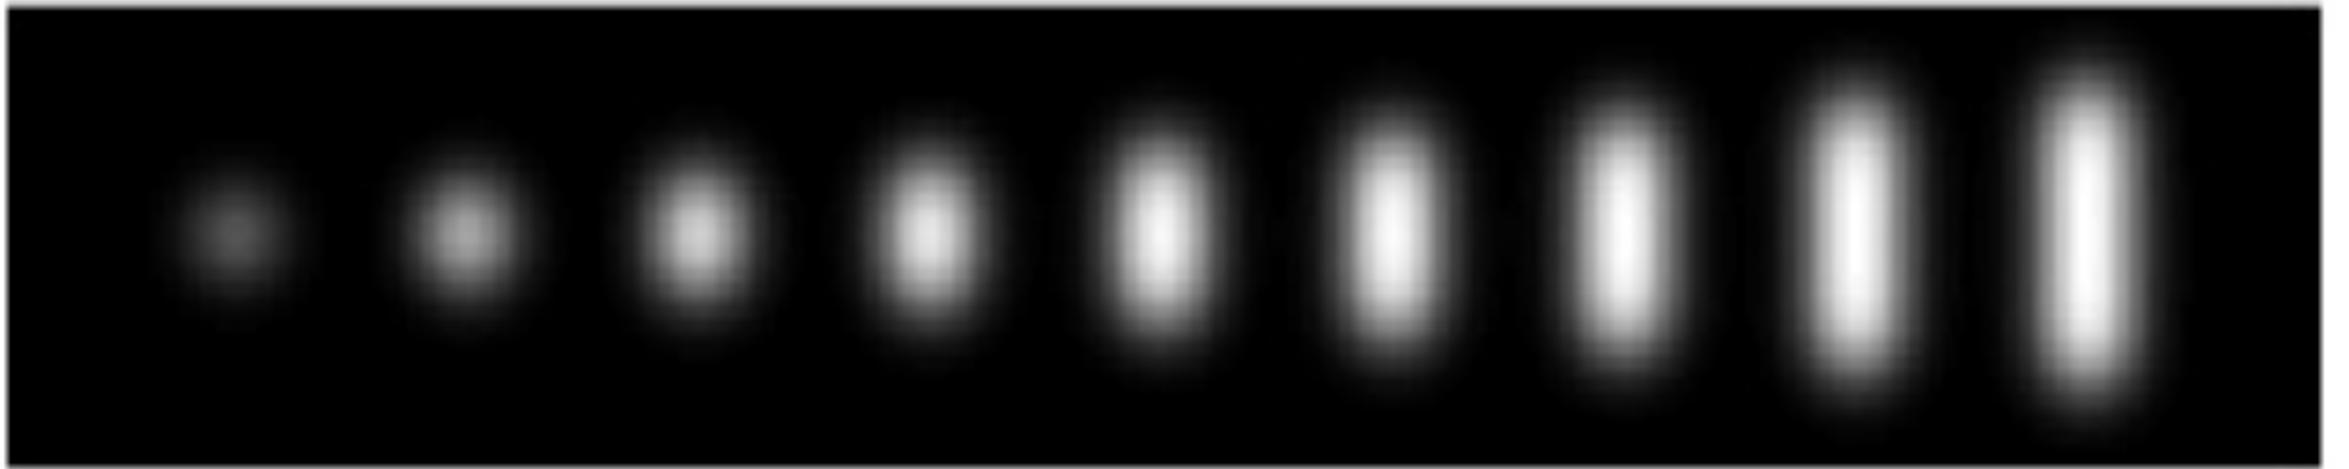
$$\frac{\lambda_1}{\lambda_2} < r$$

edge

$$\lambda_1 \gg \lambda_2$$

$$\gg \lambda$$

# Avoiding Edges



well-localized blob

edge

Detecting edges: Look at ratio of Eigenvalues ( $\lambda_1 \geq \lambda_2$ )

$$\frac{\lambda_1}{\lambda_2} < r$$

No need to explicitly compute Eigenvalues (see Lowe's paper)

# Avoiding Edges



# Avoiding Edges

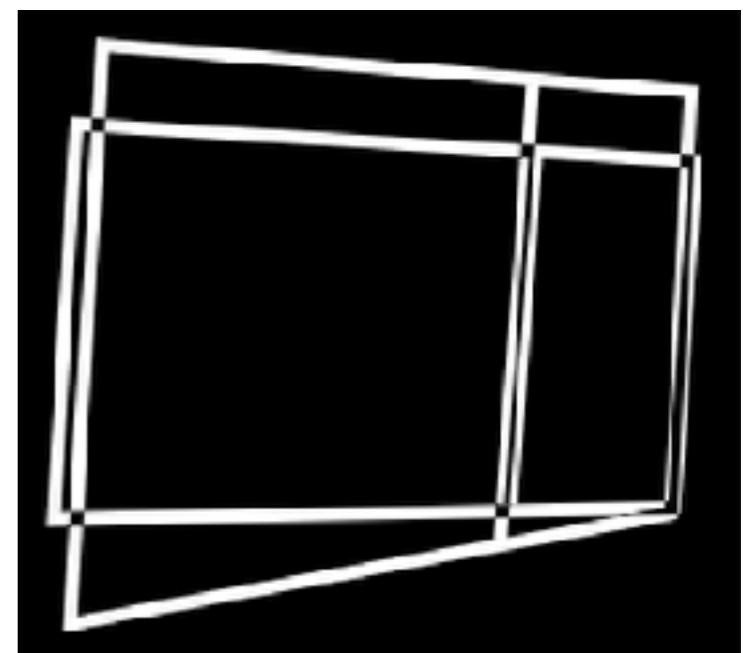
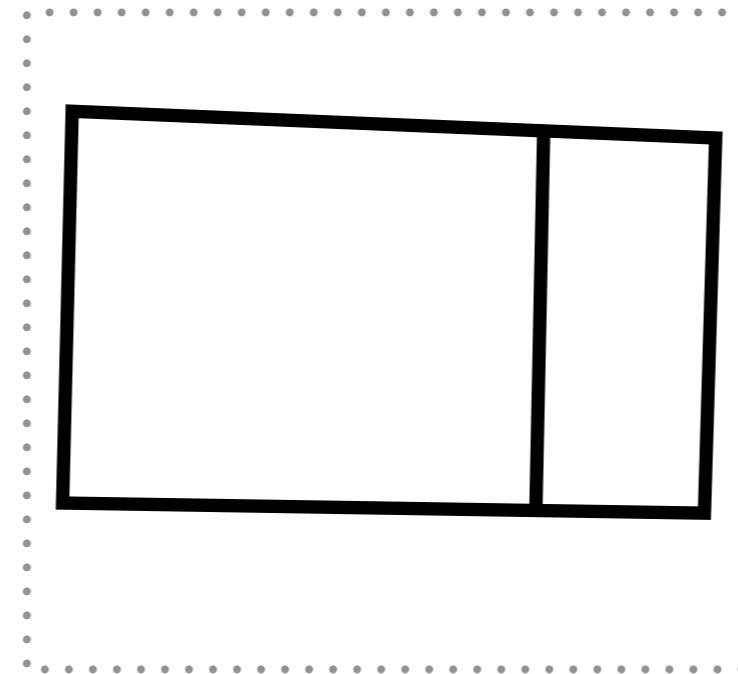
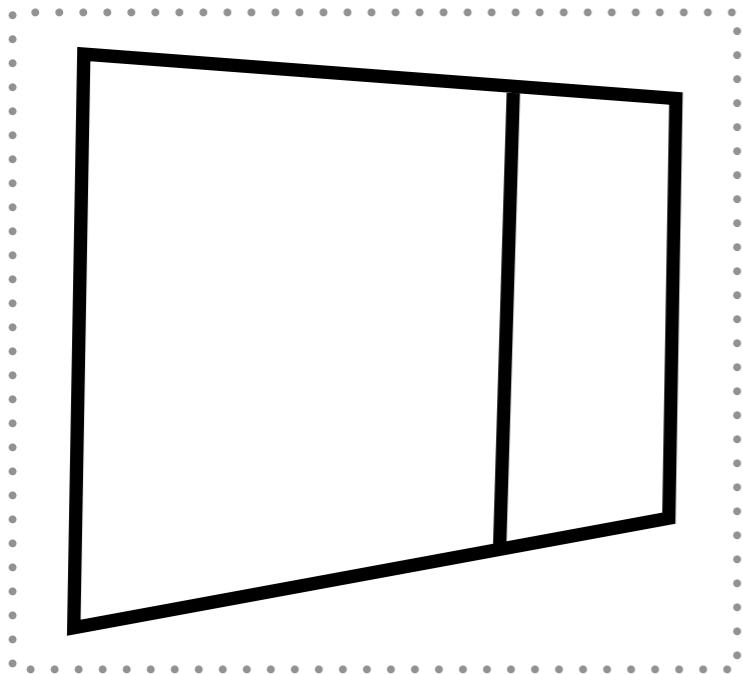


# Scale-Invariant Feature Transform (SIFT) Features

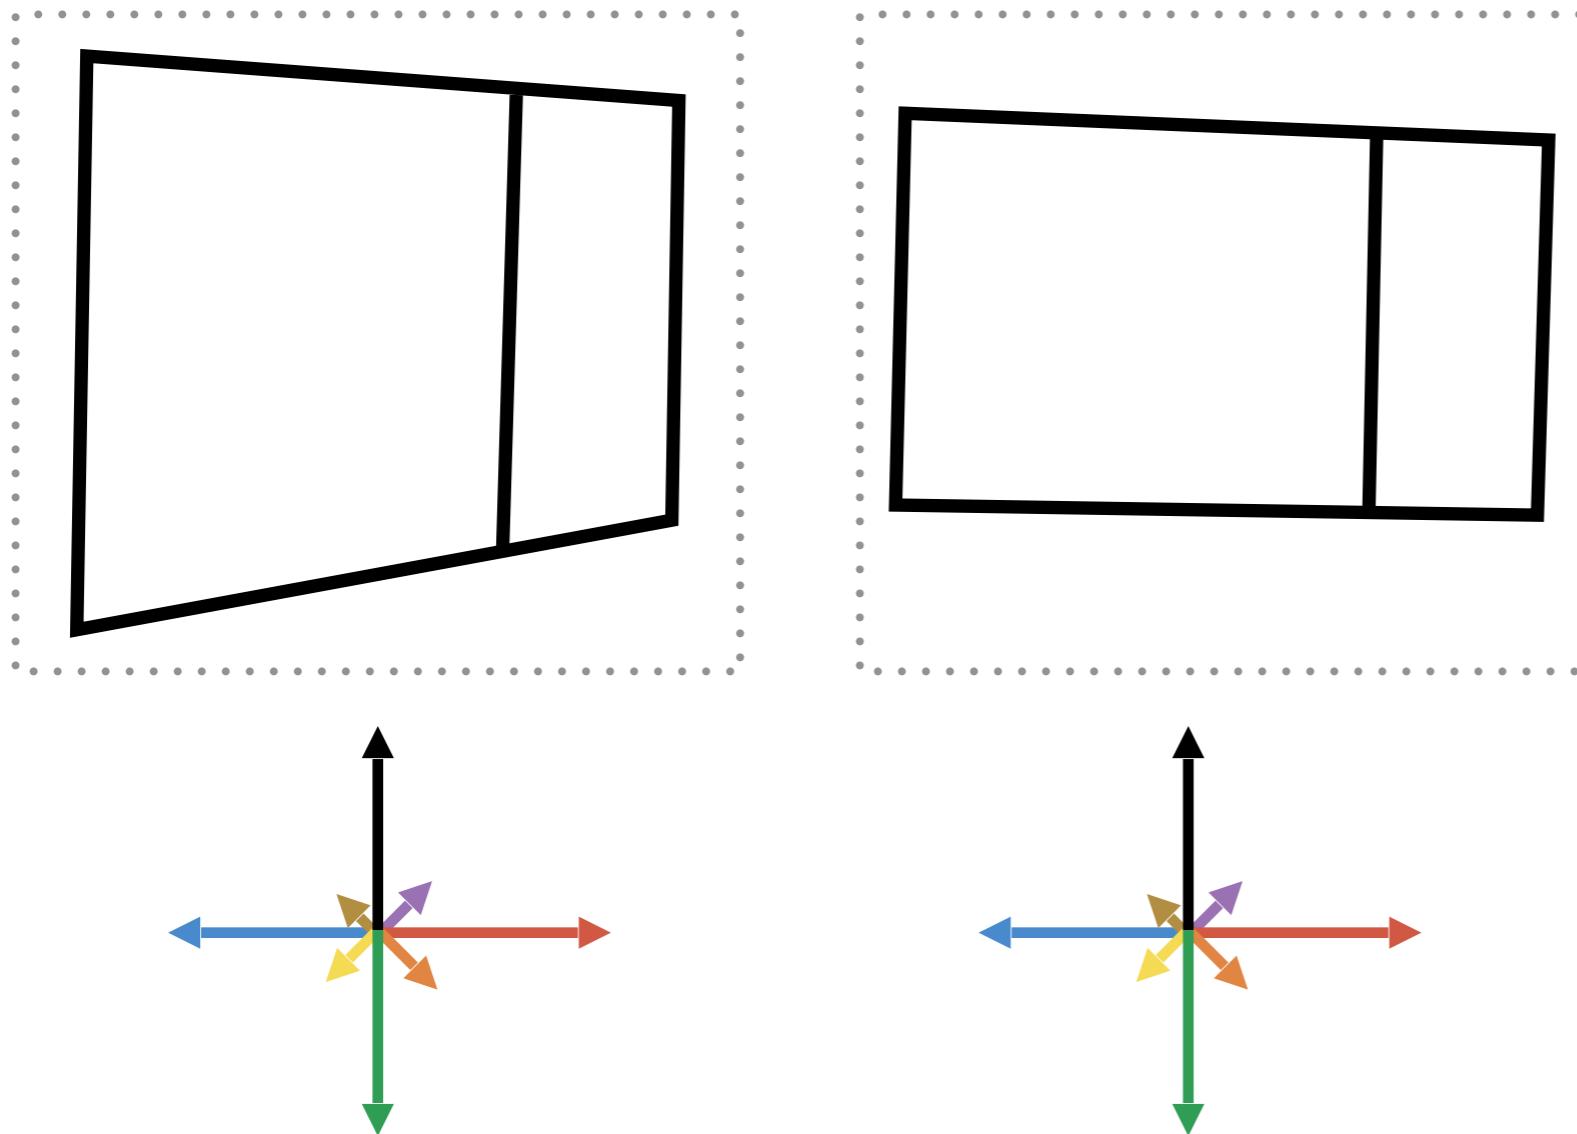


- Find a set of *interest* points
- Extract oriented patch around each interest point
- **Compute SIFT descriptor for each patch**

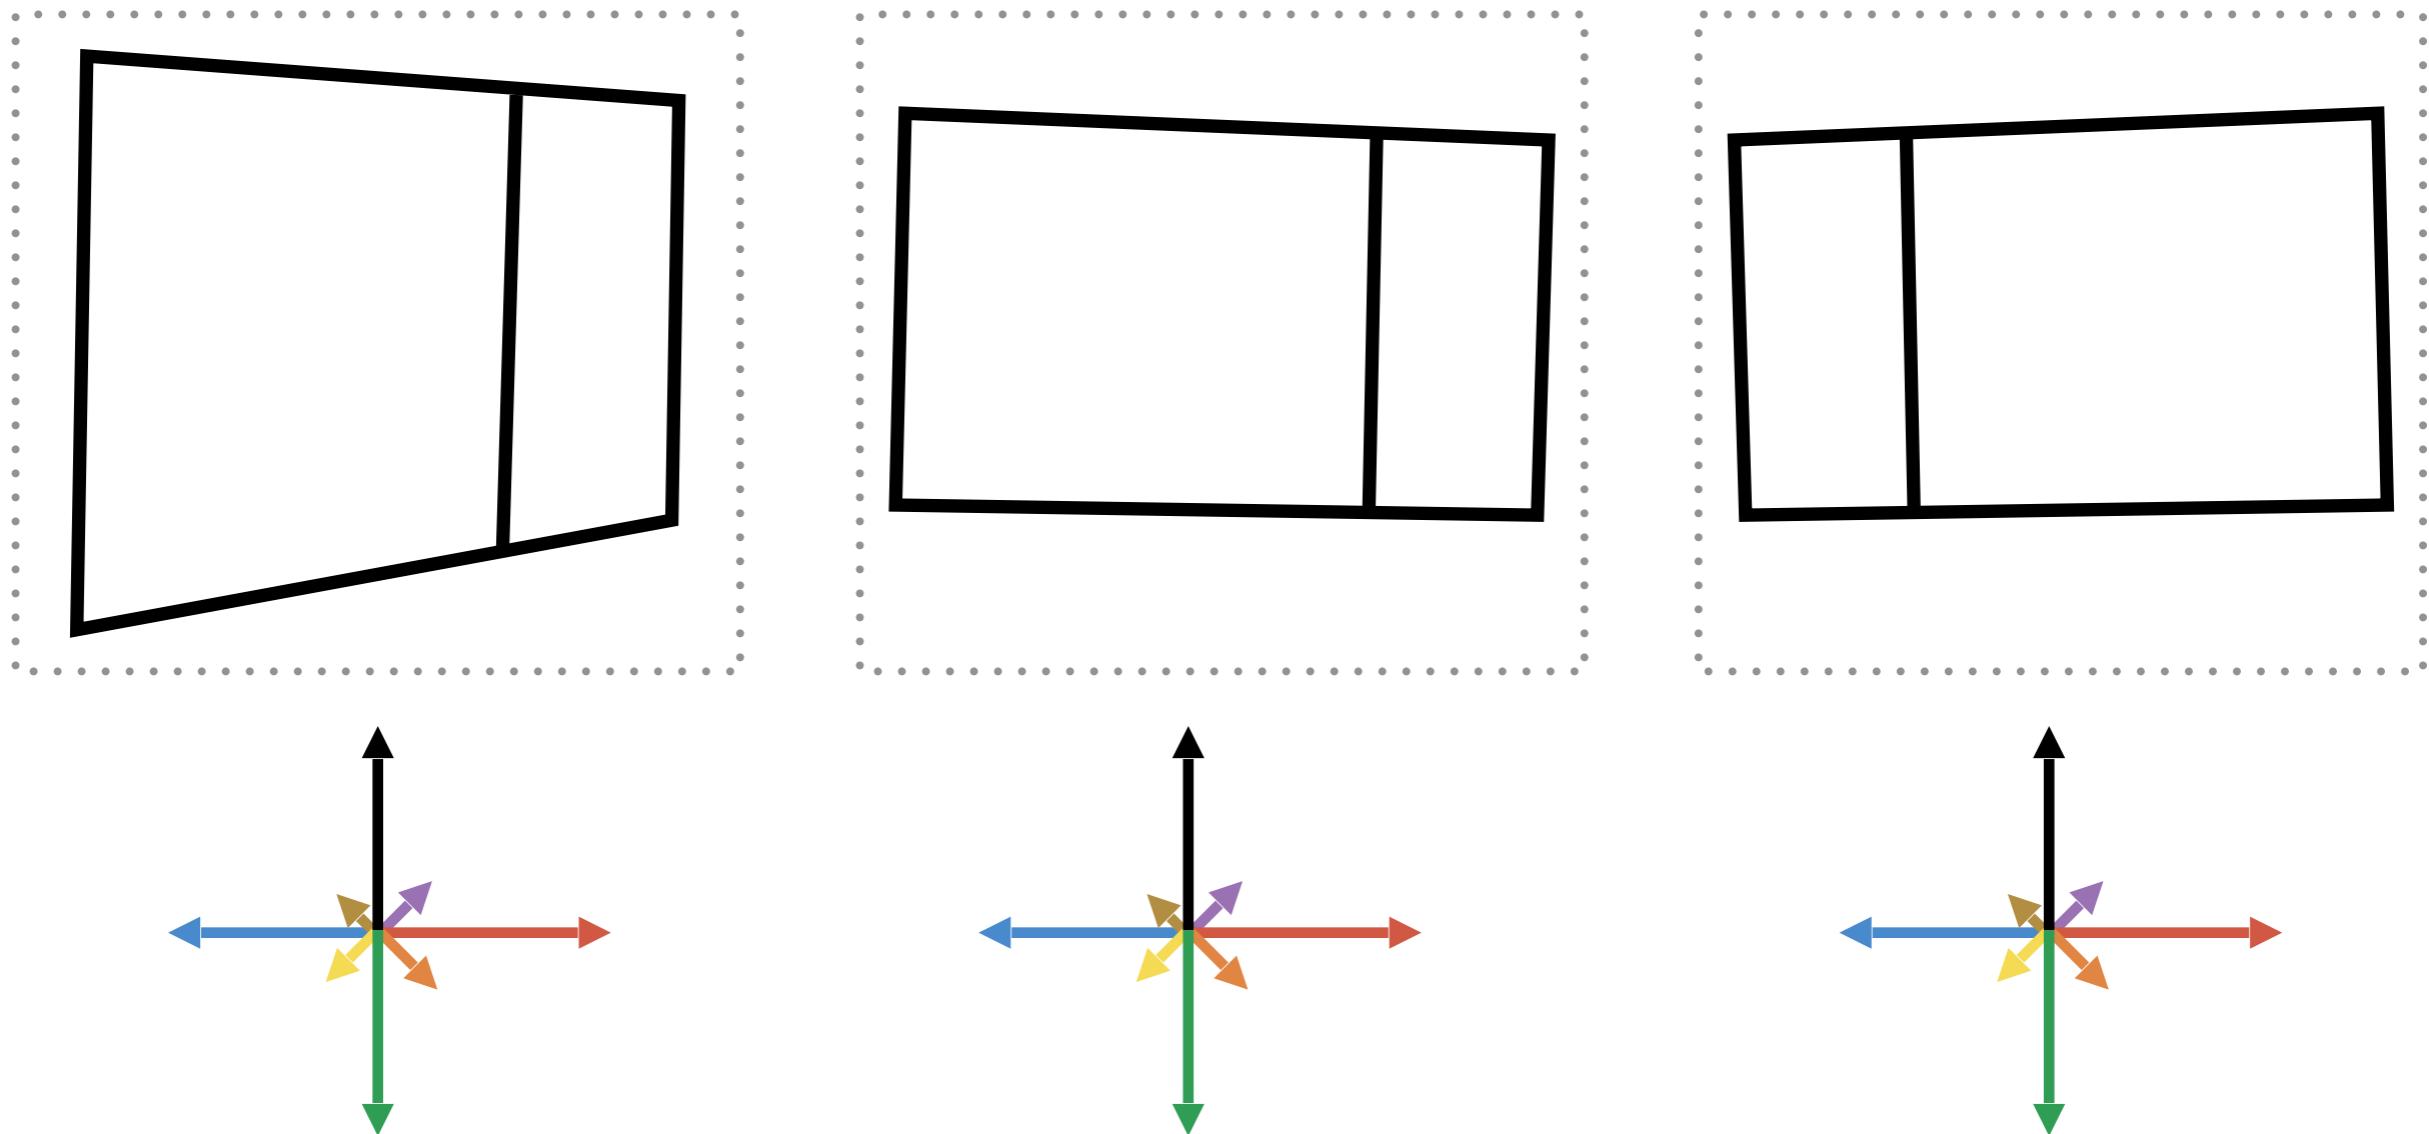
# Perspective Distortion



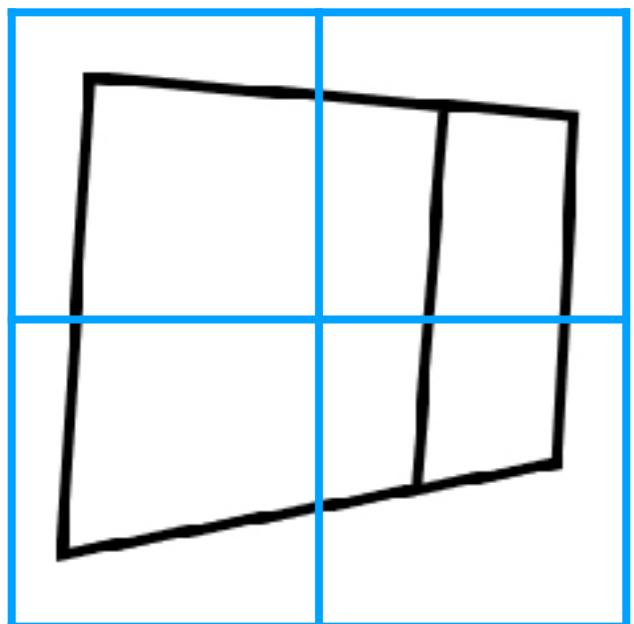
# Perspective Distortion



# Perspective Distortion

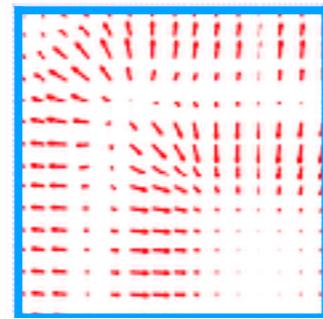
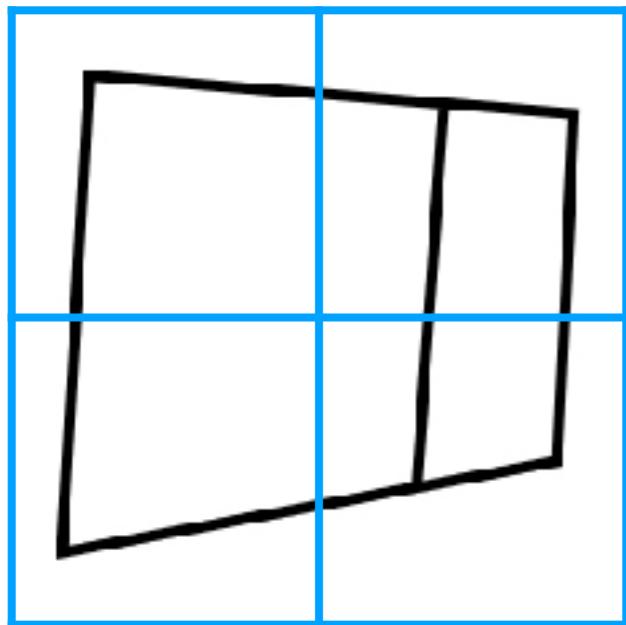


# SIFT Descriptor (Idea)



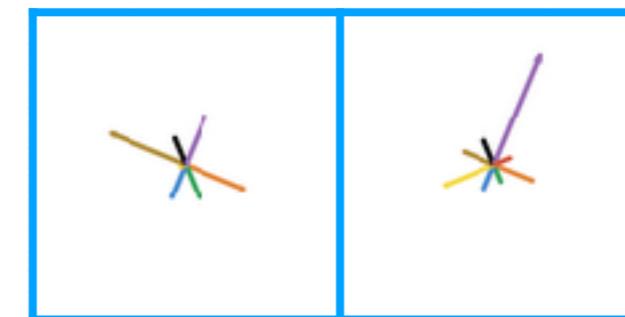
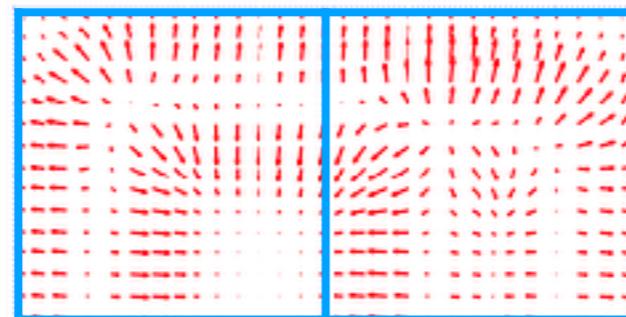
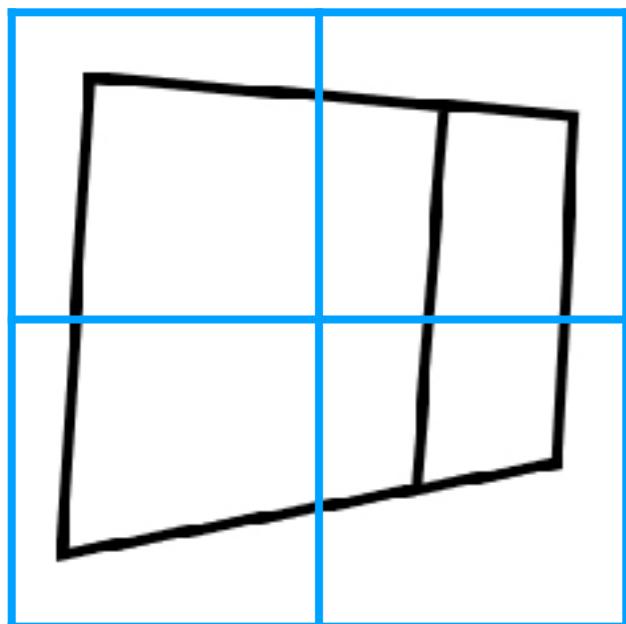
Split patch in  $2 \times 2$  regions  
(SIFT actually uses  $4 \times 4$  regions)

# SIFT Descriptor (Idea)



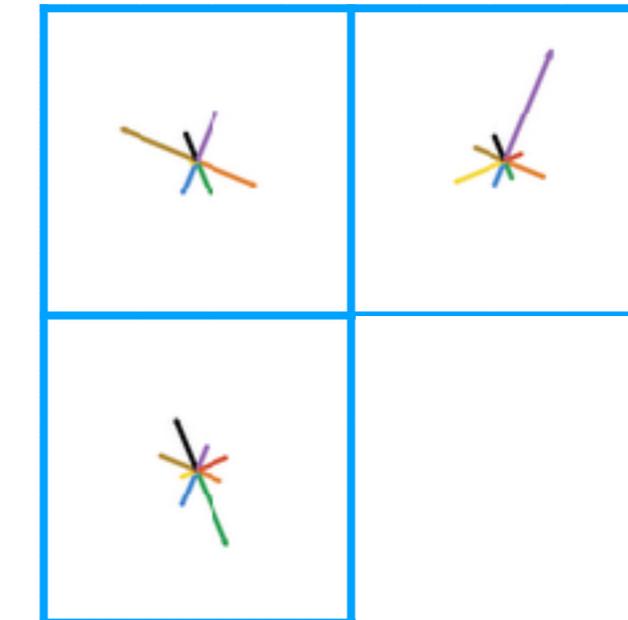
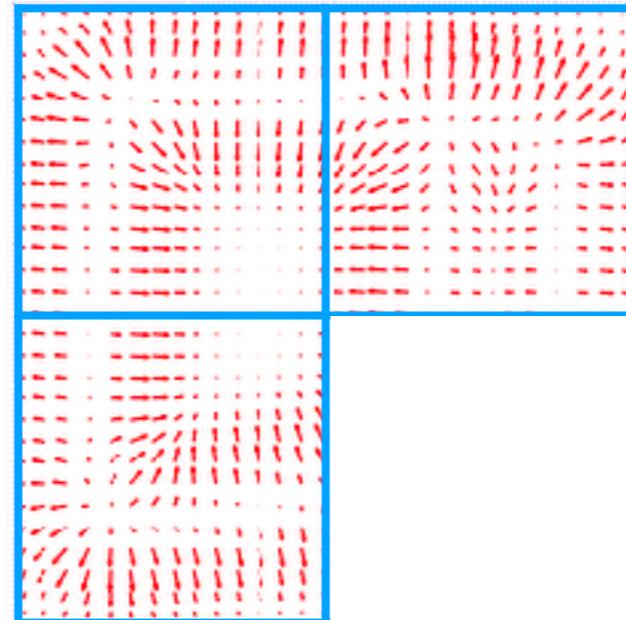
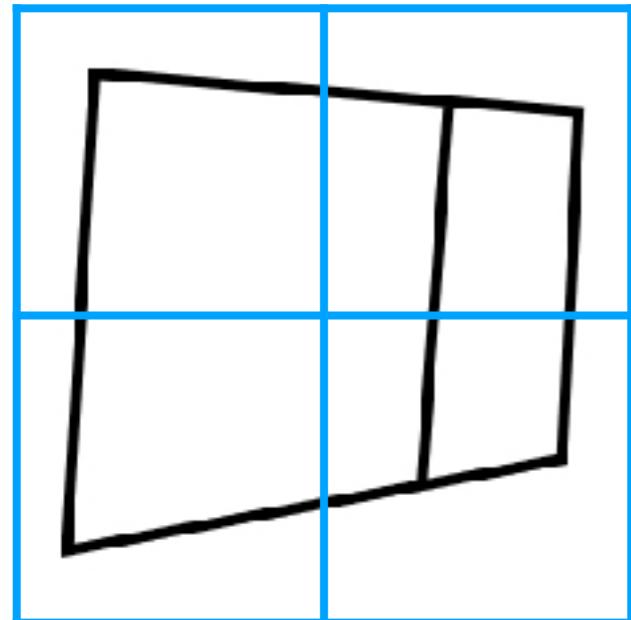
Compute a 8-bin gradient histogram for each region

# SIFT Descriptor (Idea)



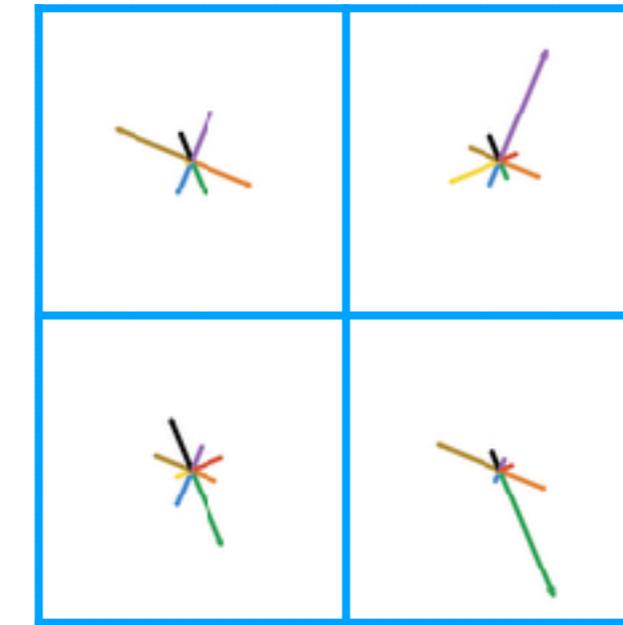
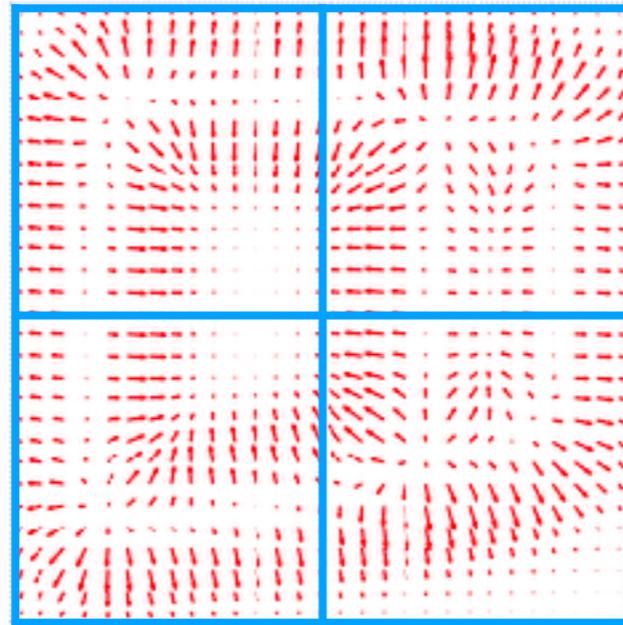
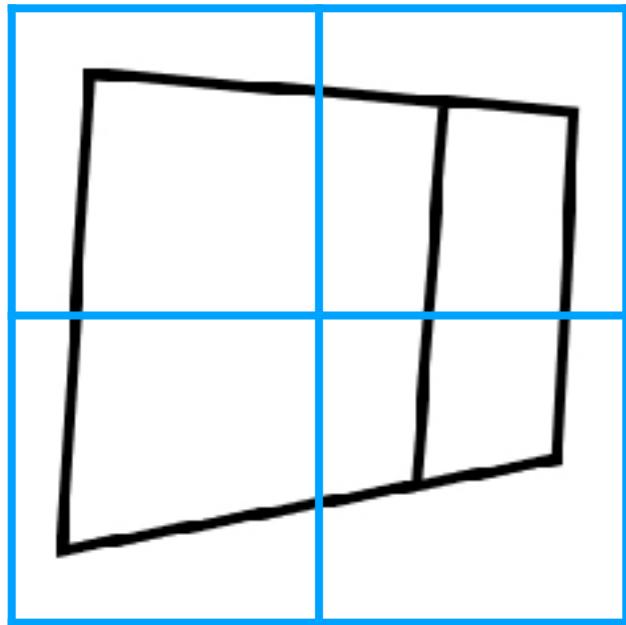
Compute a 8-bin gradient histogram for each region

# SIFT Descriptor (Idea)

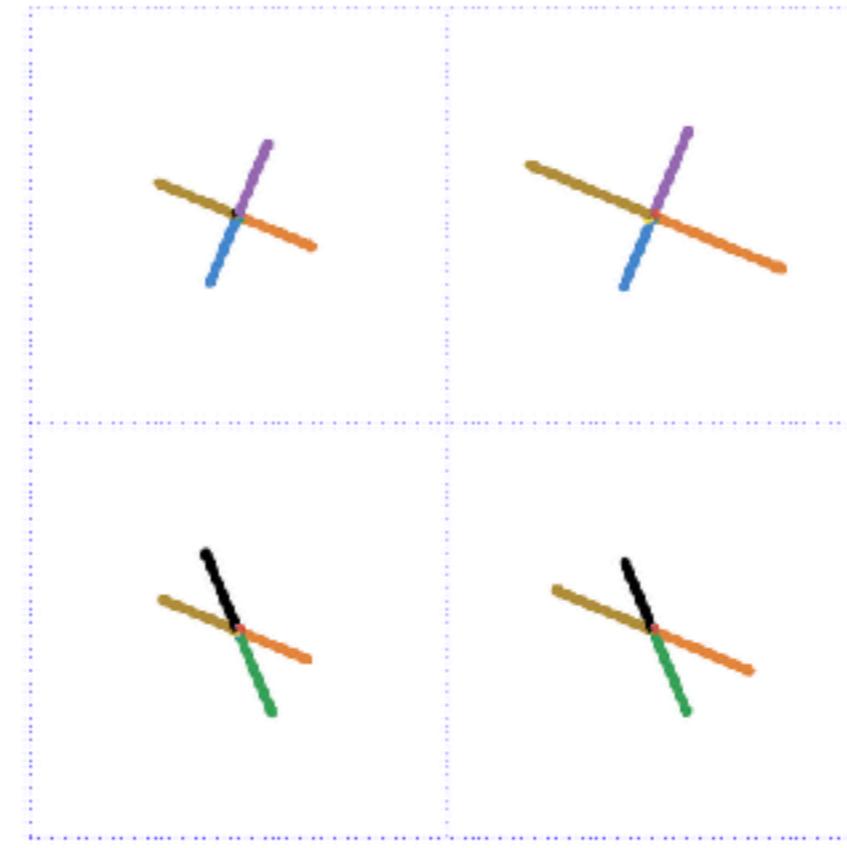
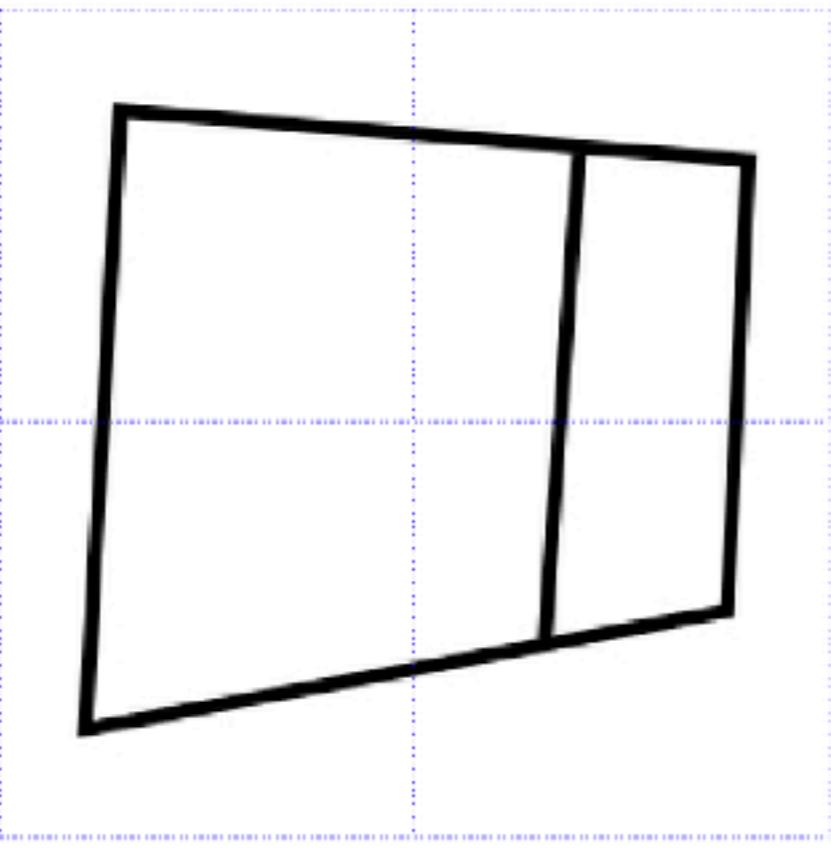
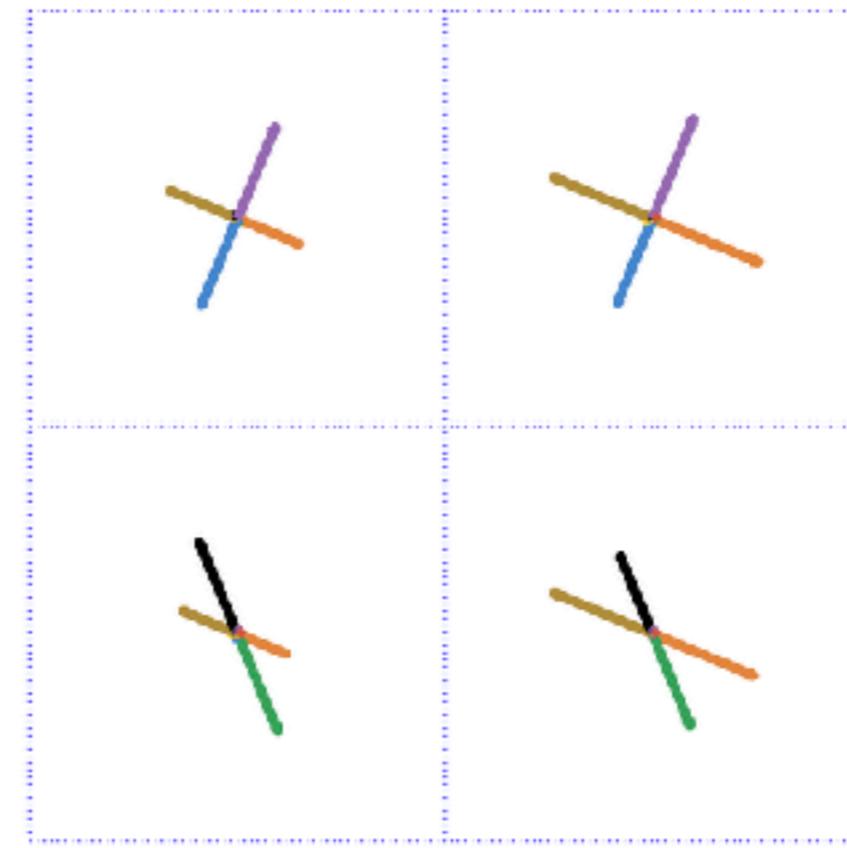
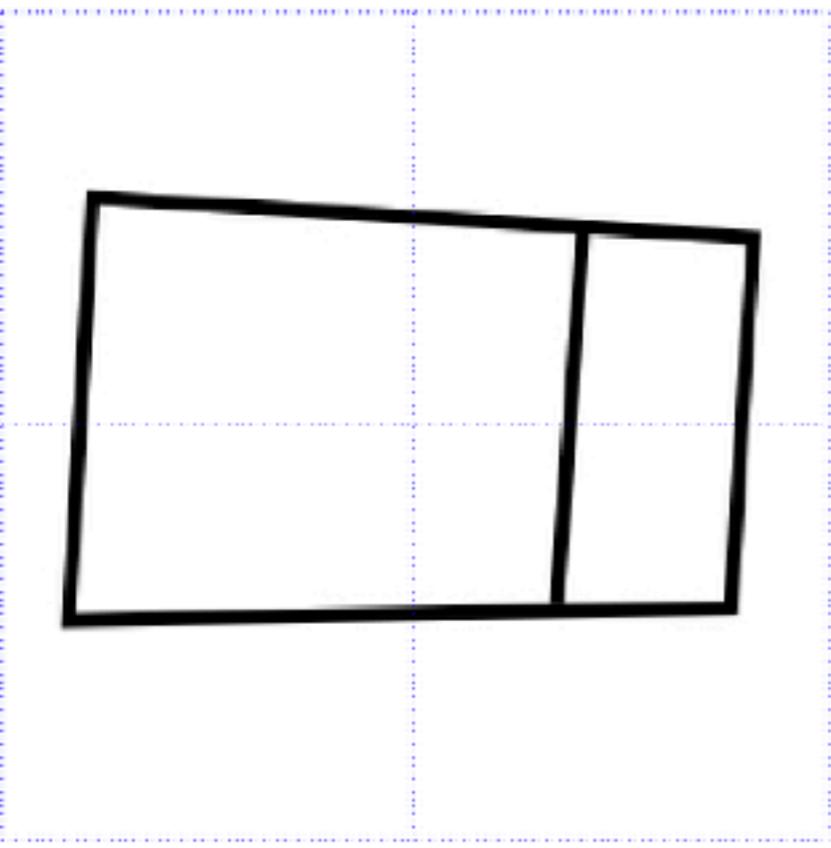


Compute a 8-bin gradient histogram for each region

# SIFT Descriptor (Idea)



$L_2$  Normalize the descriptor



# Scale-Invariant Feature Transform (SIFT) Features



- Find a set of *interest* points
- Extract oriented patch around each interest point
- Compute SIFT descriptor for each patch

# Scale Invariance



Detect a size $s$ , for each point

# Scale Invariance



Detect a size  $s$ , for each point

# Scale Invariance



Detect a size  $s$ , for each point  
Scale patch size to fit.  
Compute gradients in  $L(x, y, s^2)$

# Scale Invariance



size=5



size=0.5

# Scale Invariance



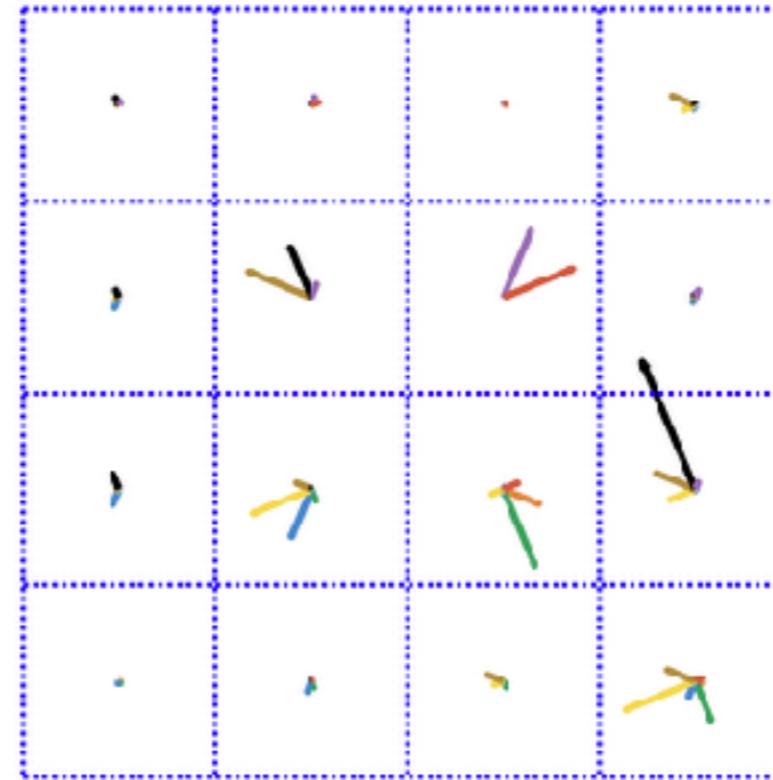
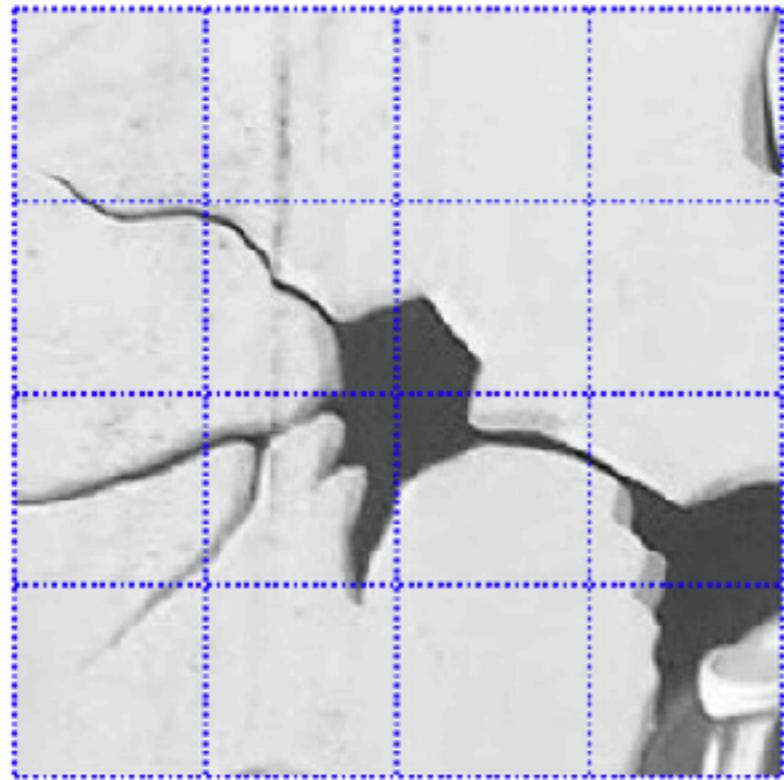
size=5

Map to canonical patch before descriptor extraction

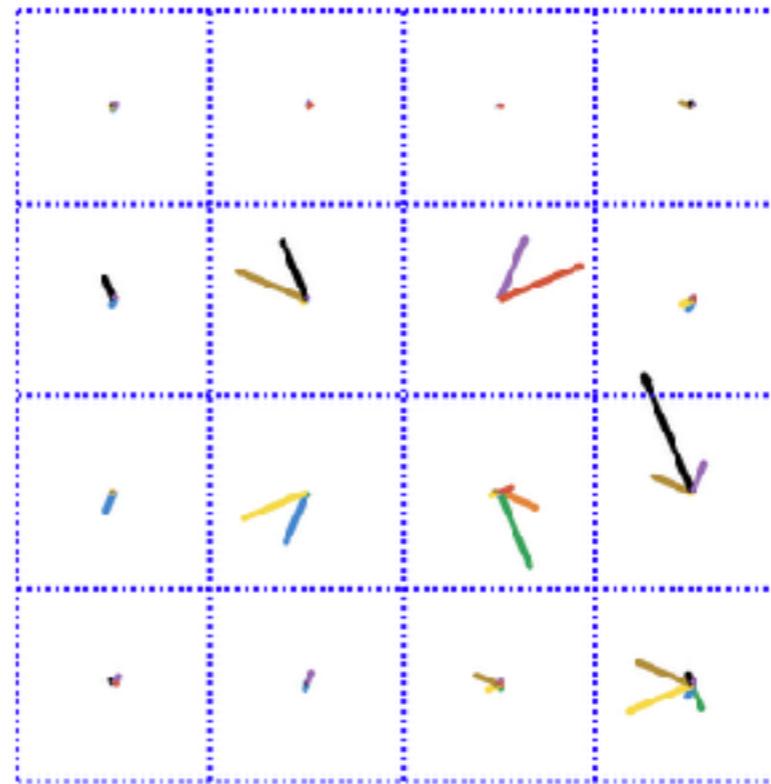
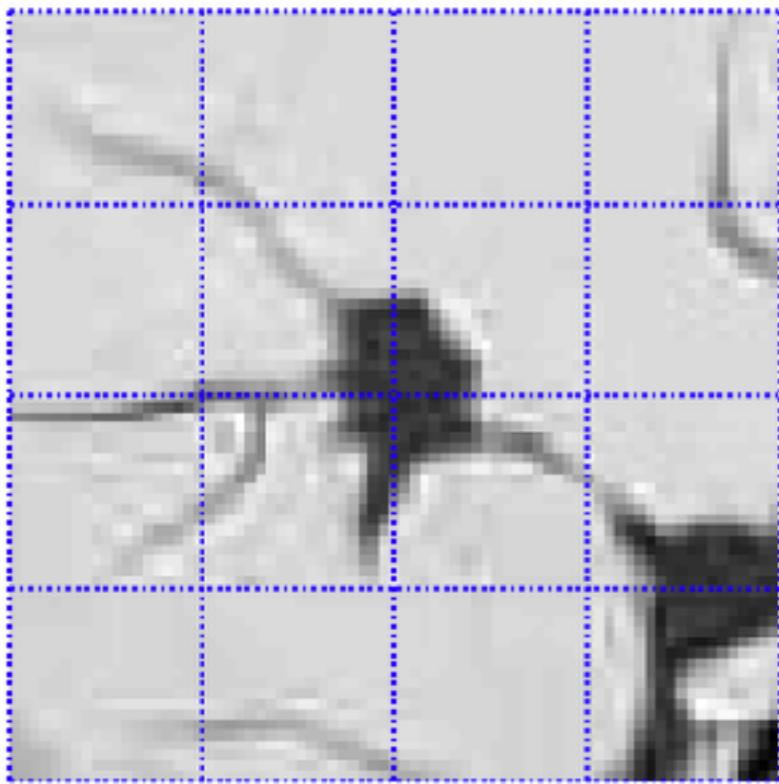
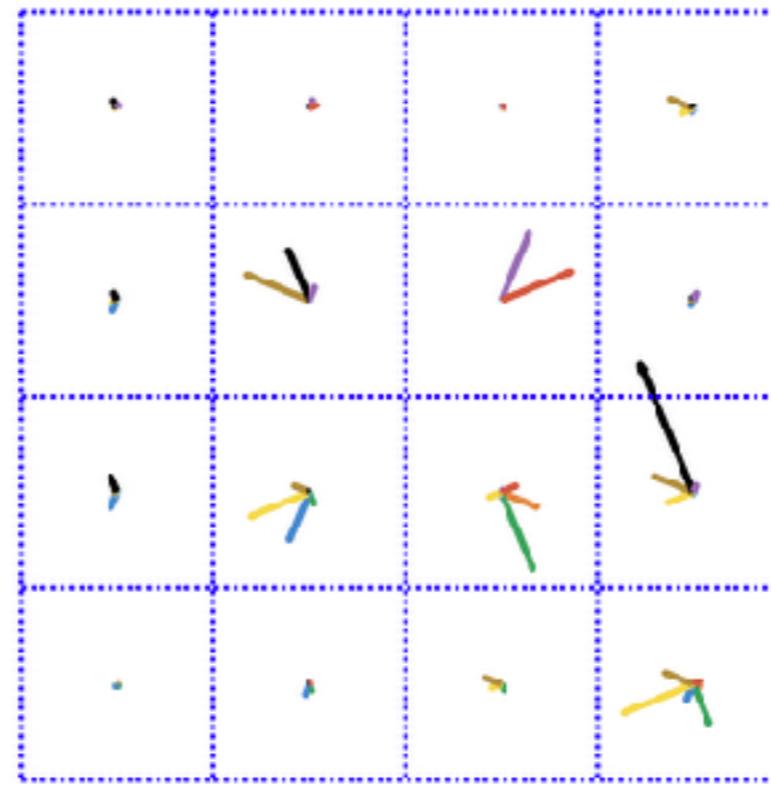
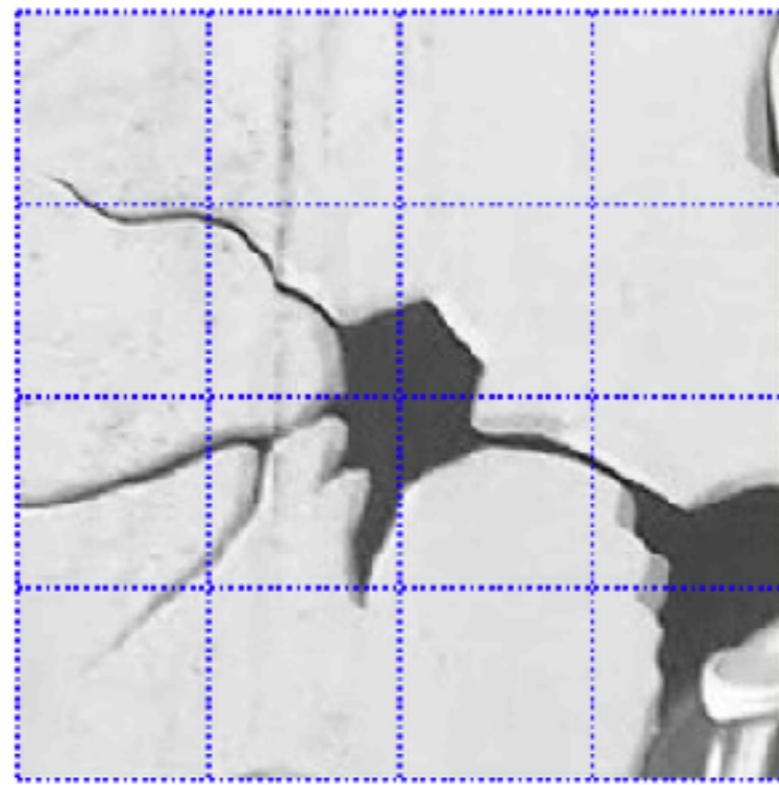


size=0.5

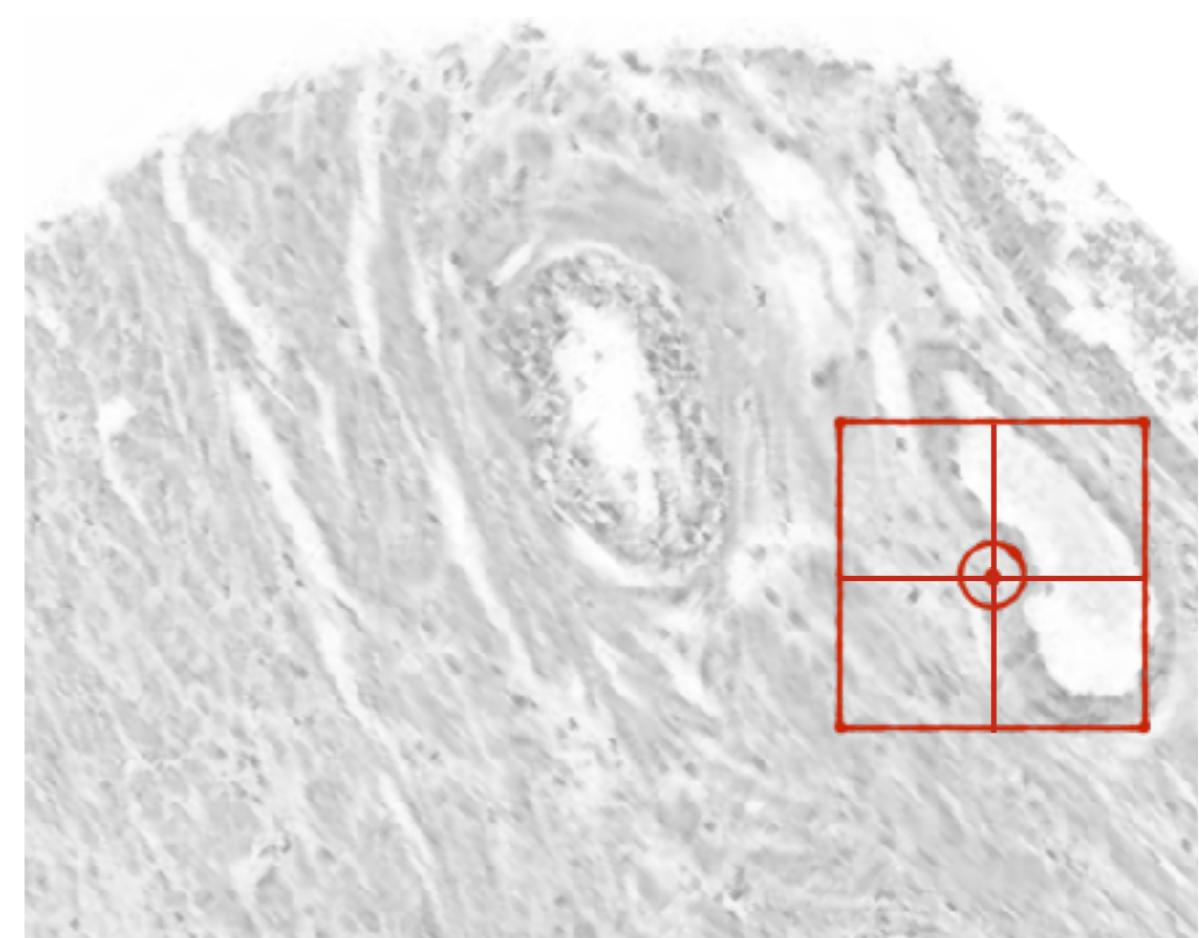
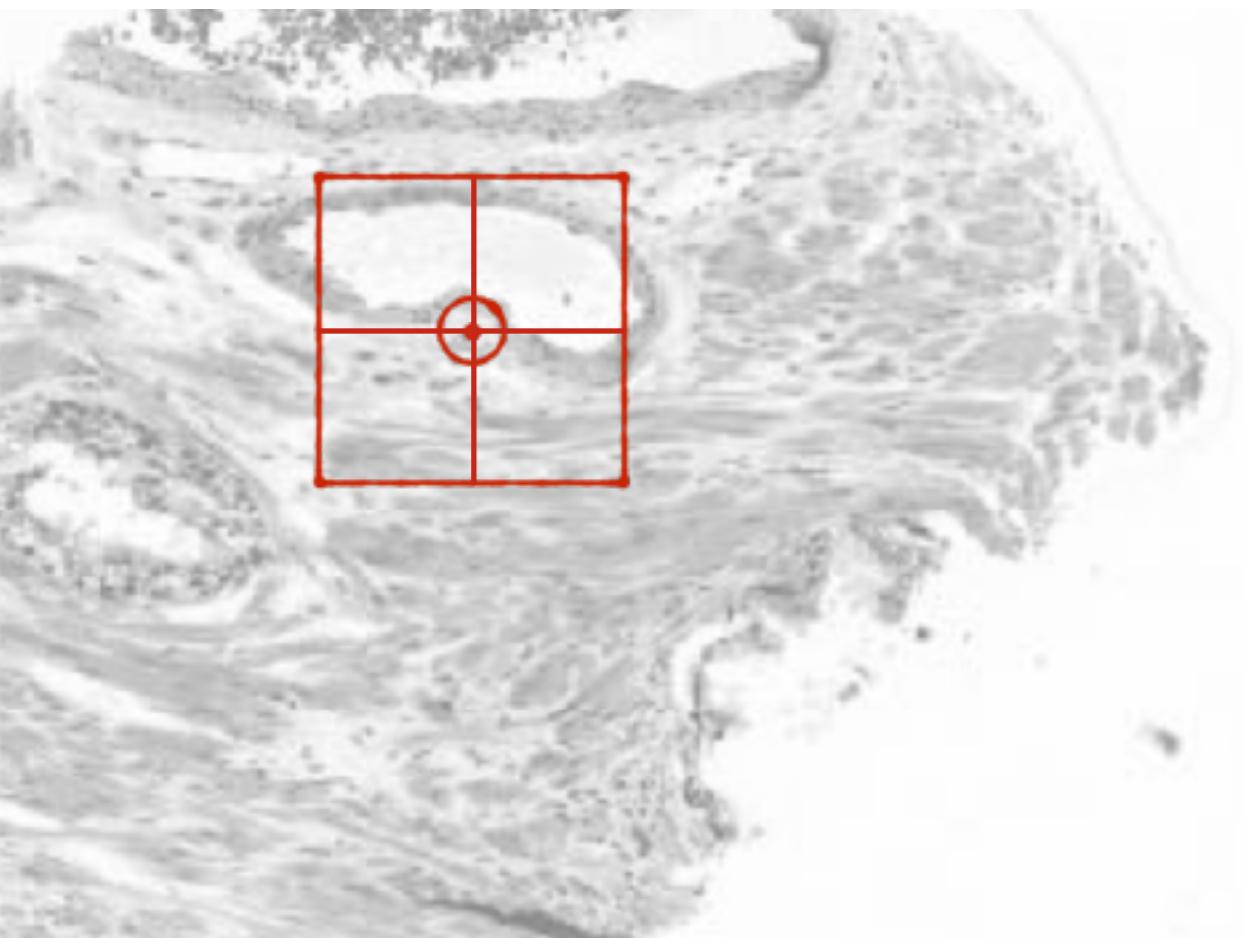
# Canonical Patches



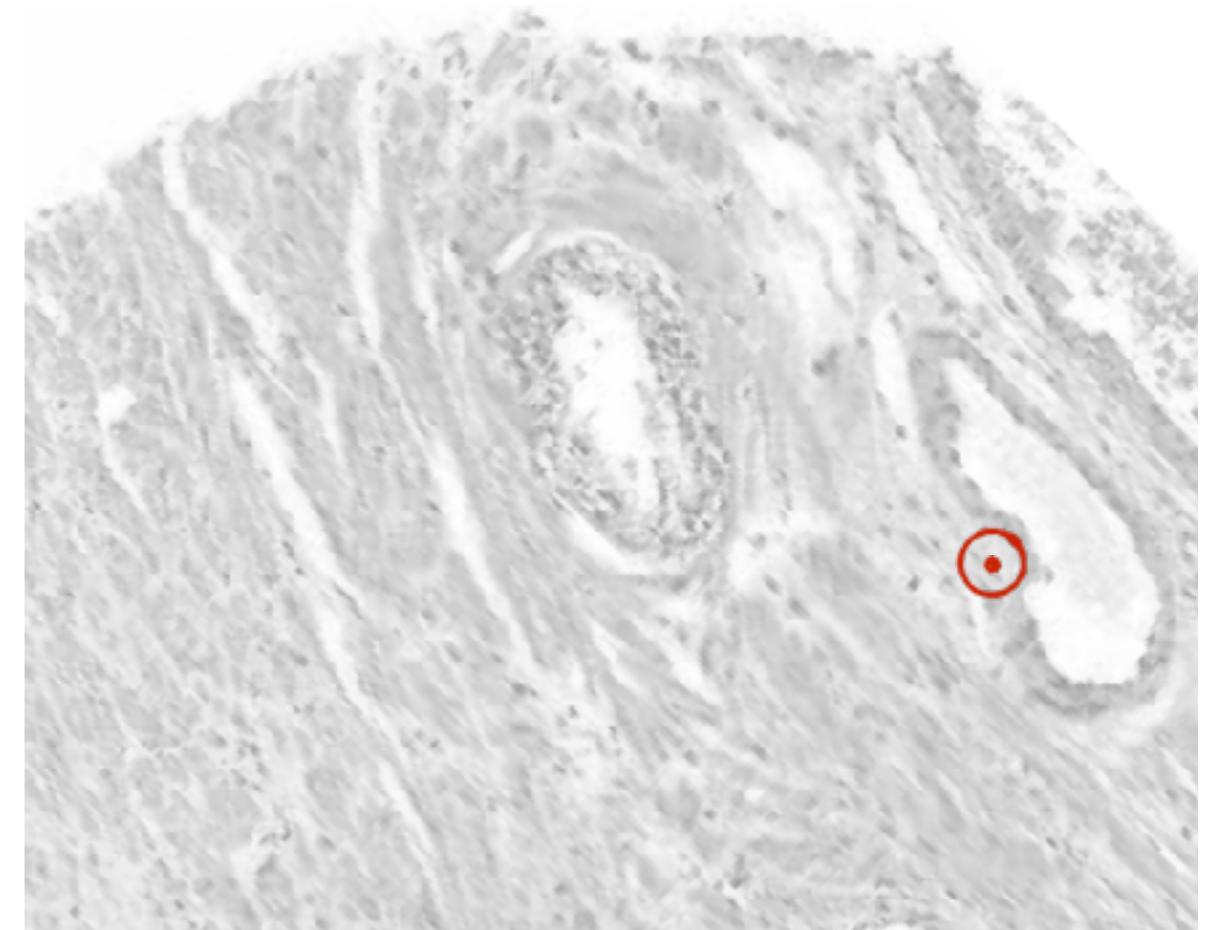
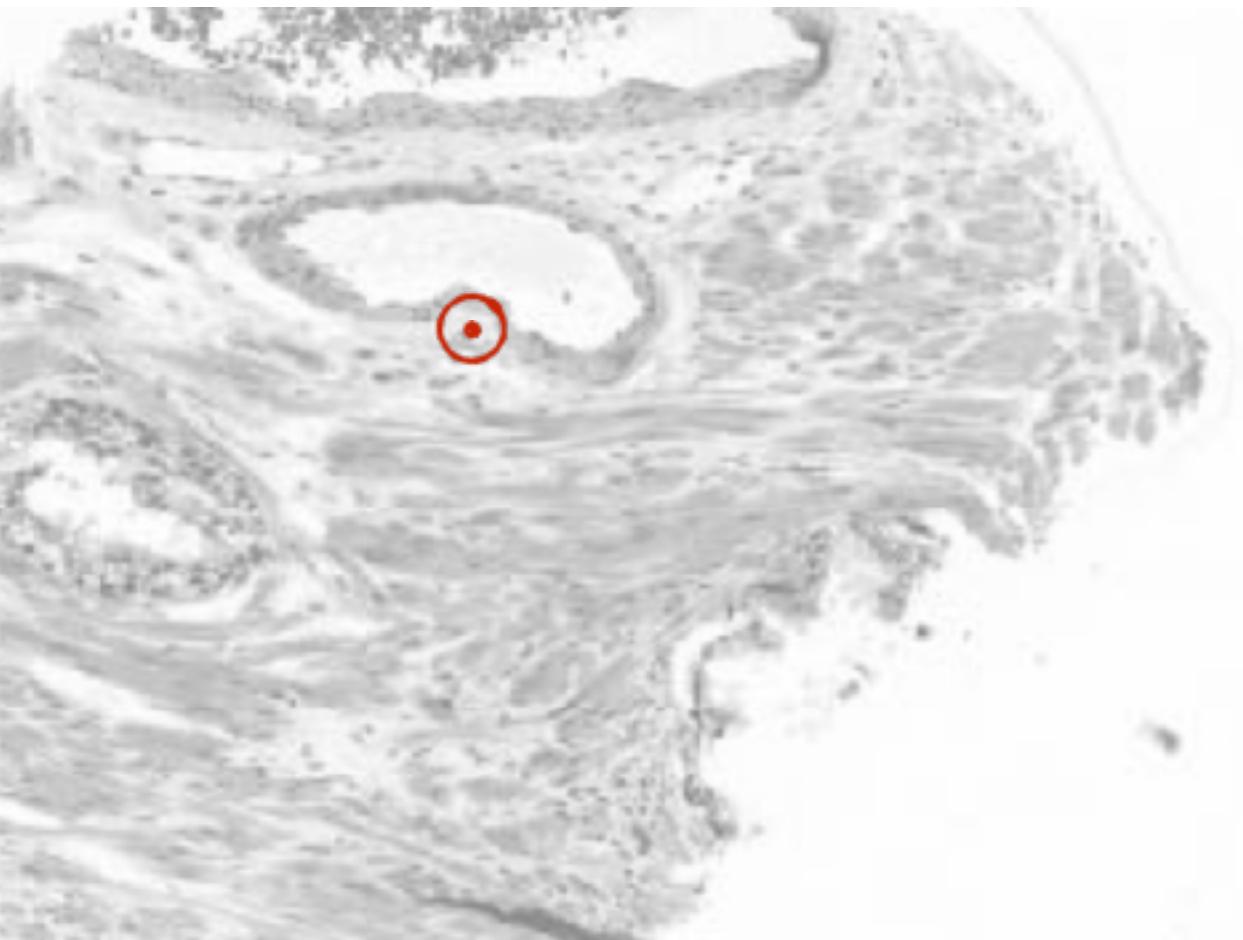
# Canonical Patches



# Challenge: Rotations

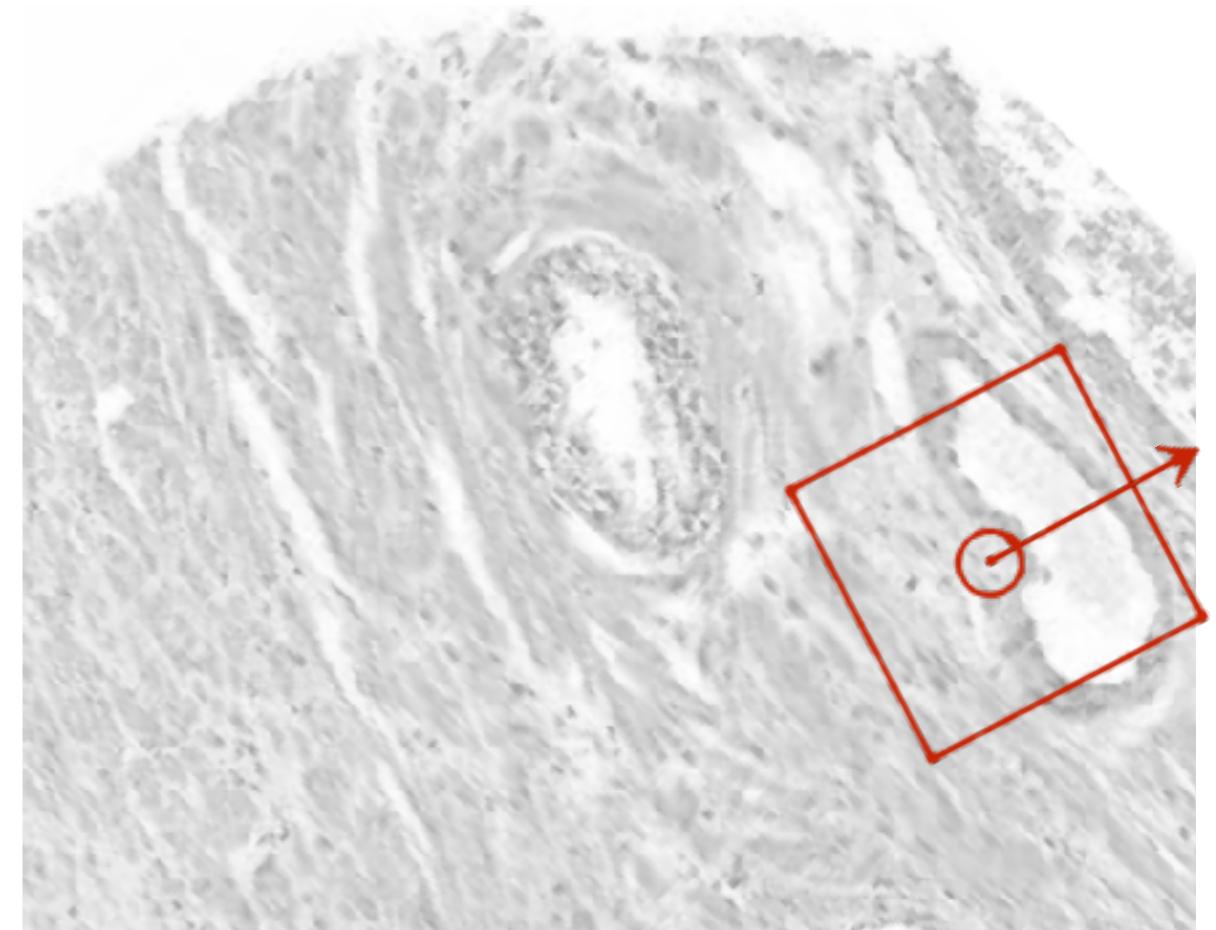
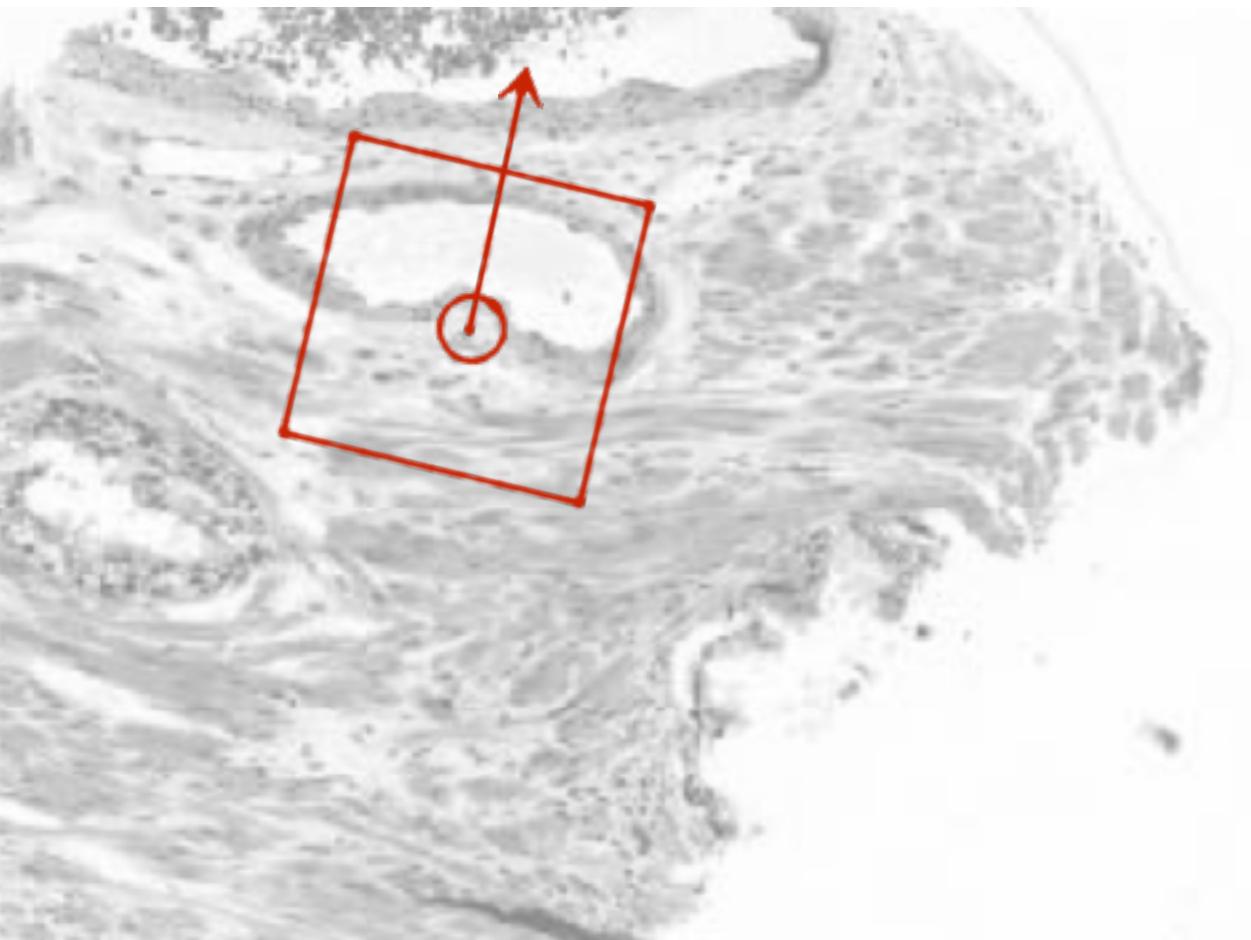


# Challenge: Rotations



Estimate a dominant direction

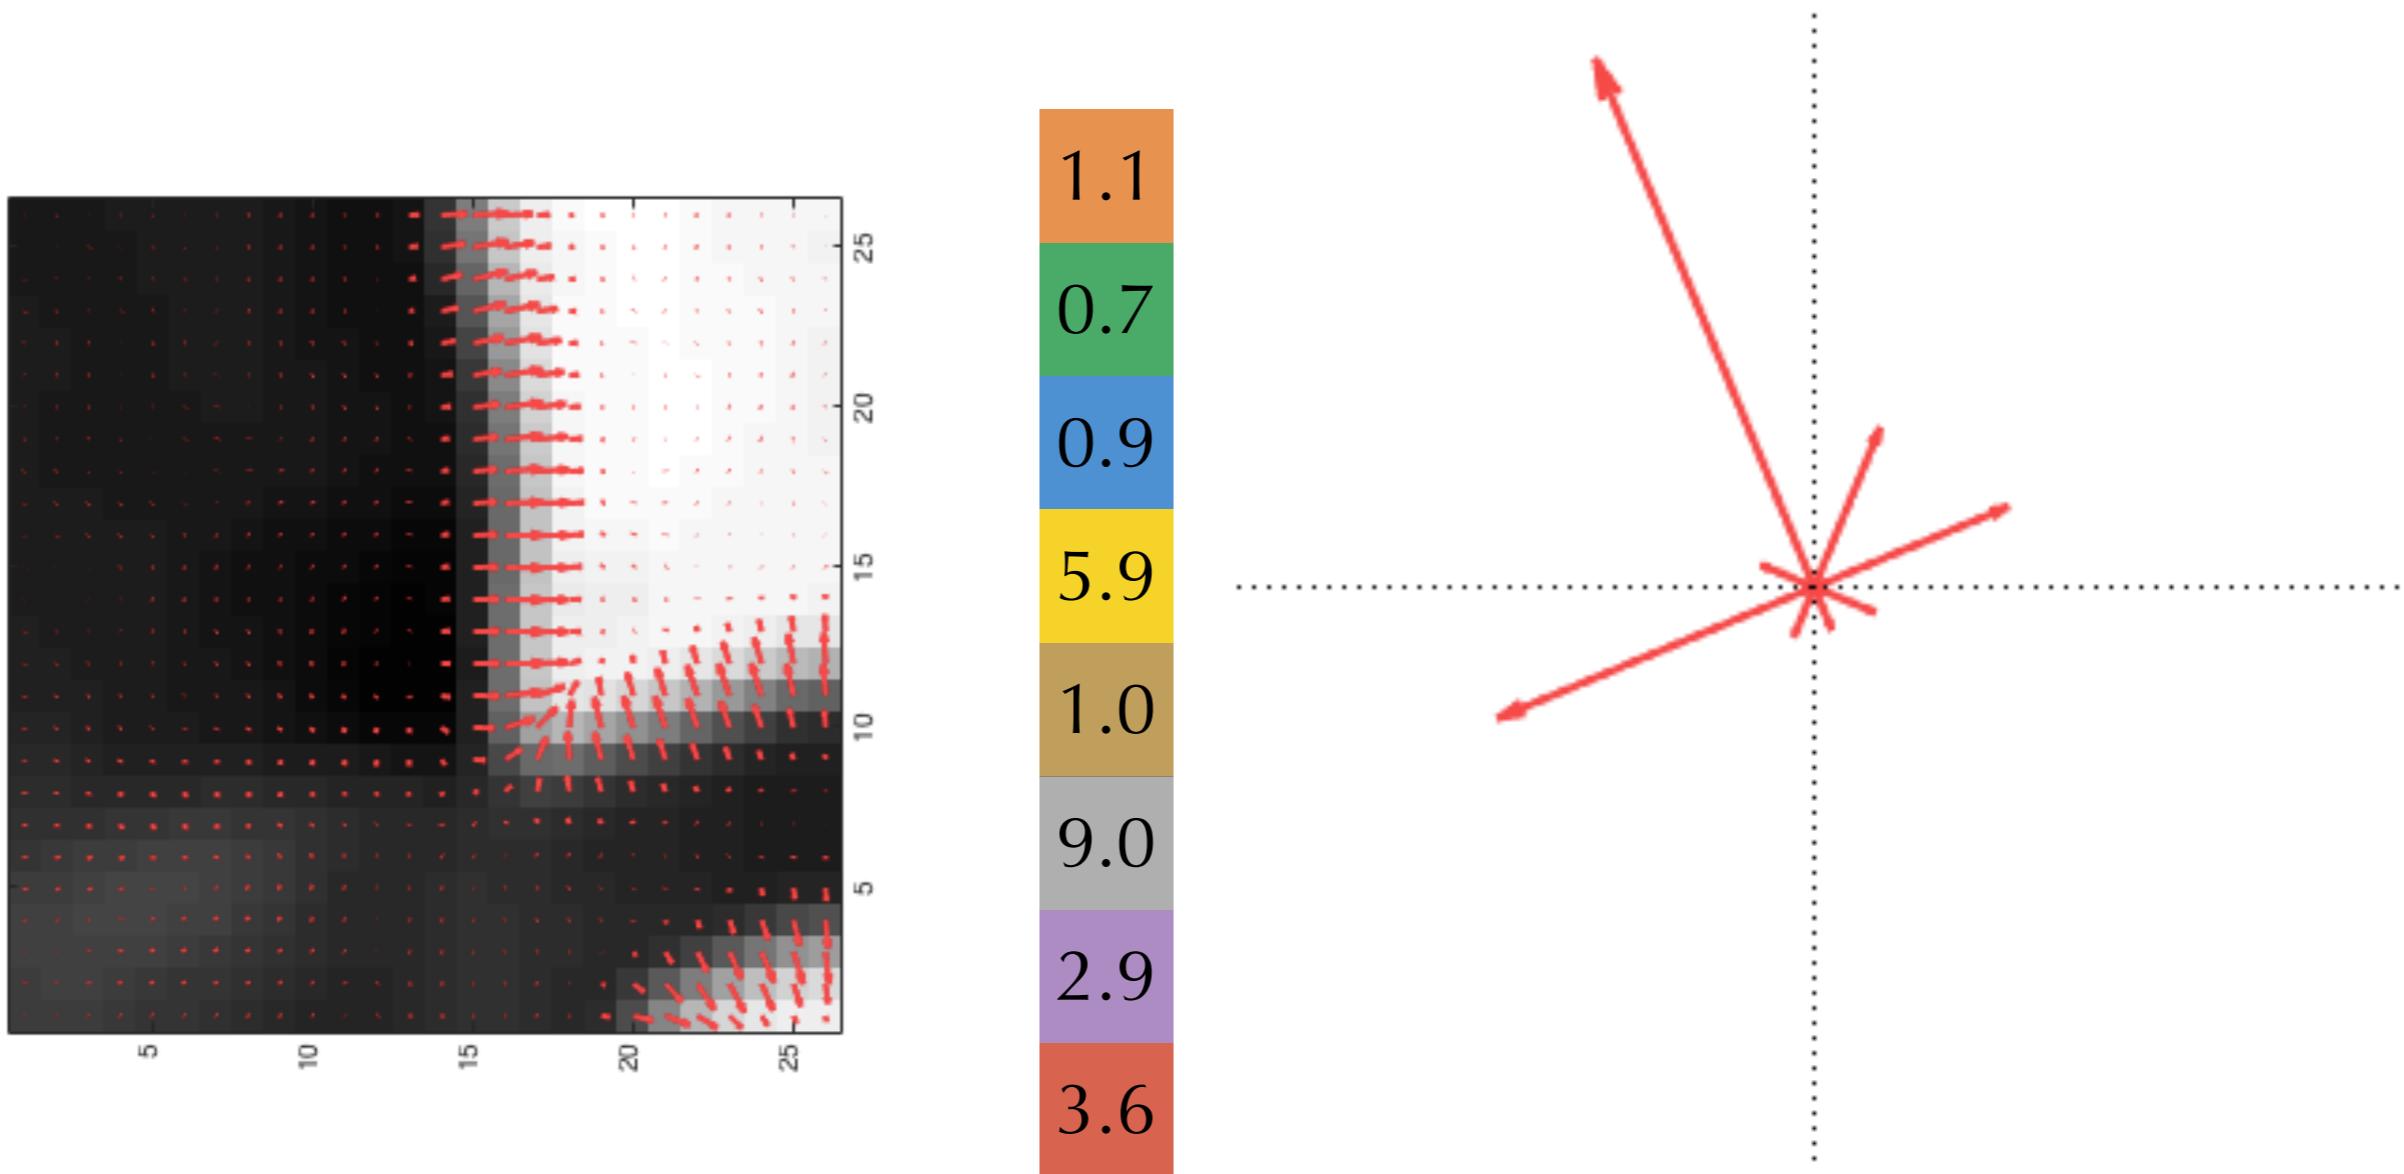
# Challenge: Rotations



Estimate a dominant direction

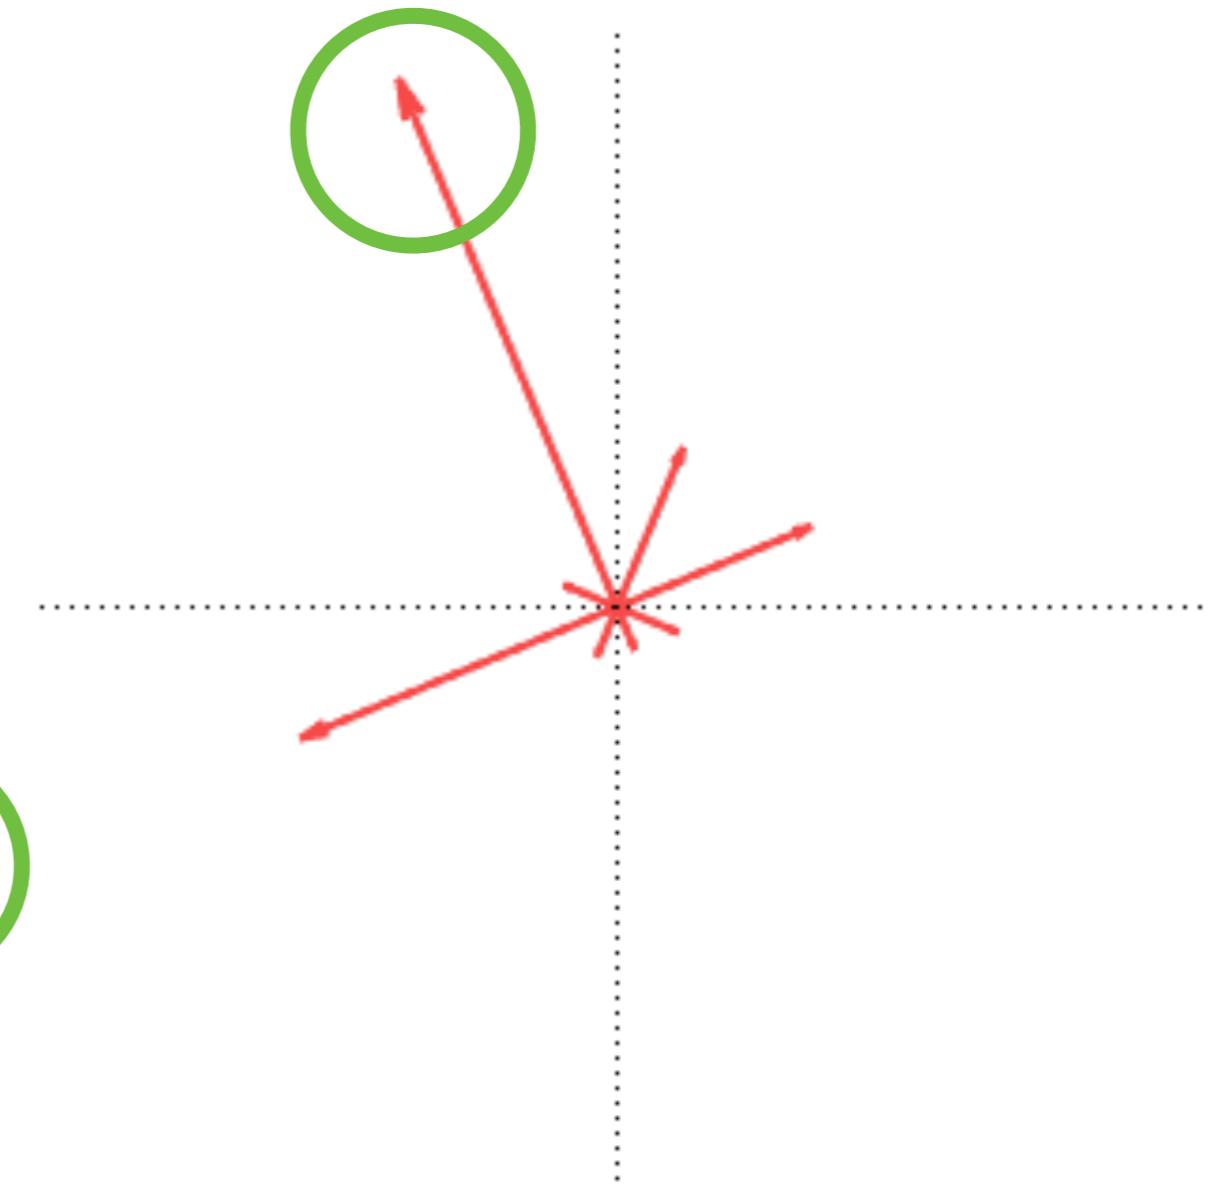
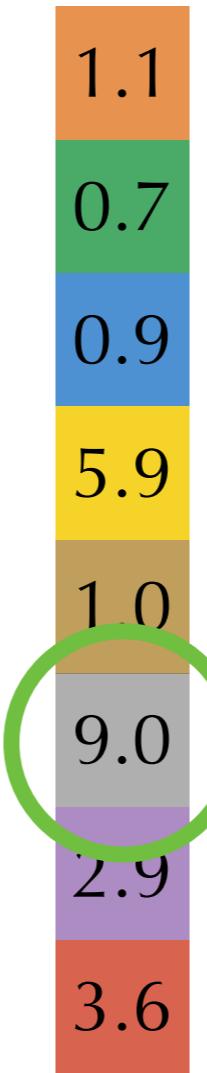
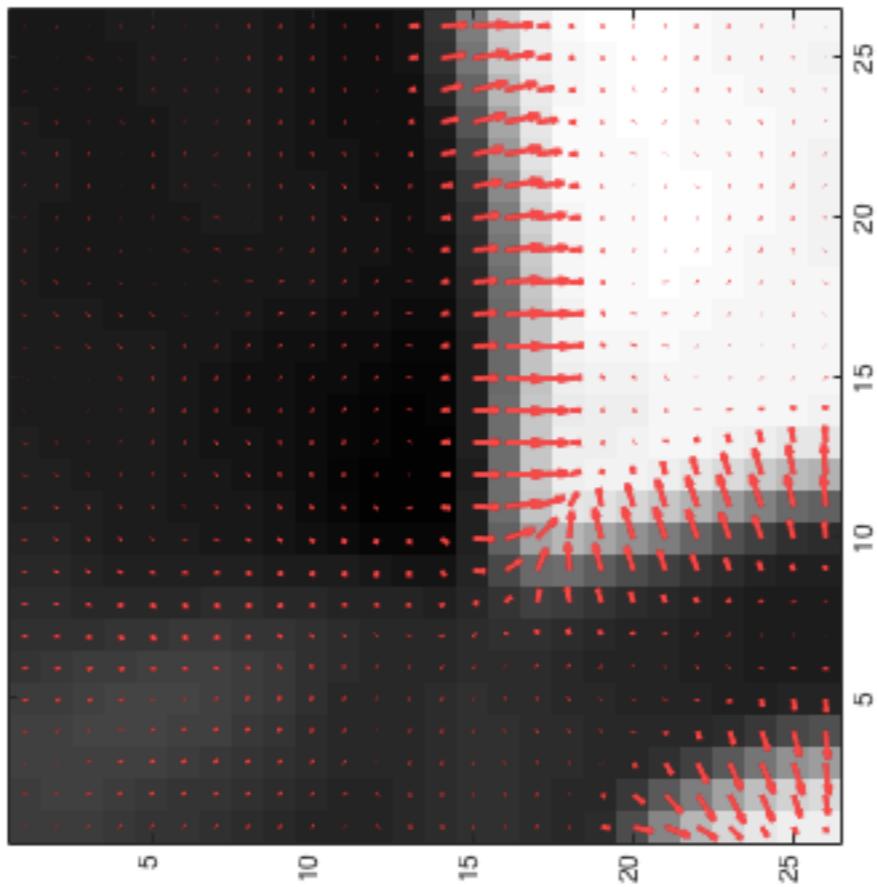
Align patches to this direction

# Orientation Assignment



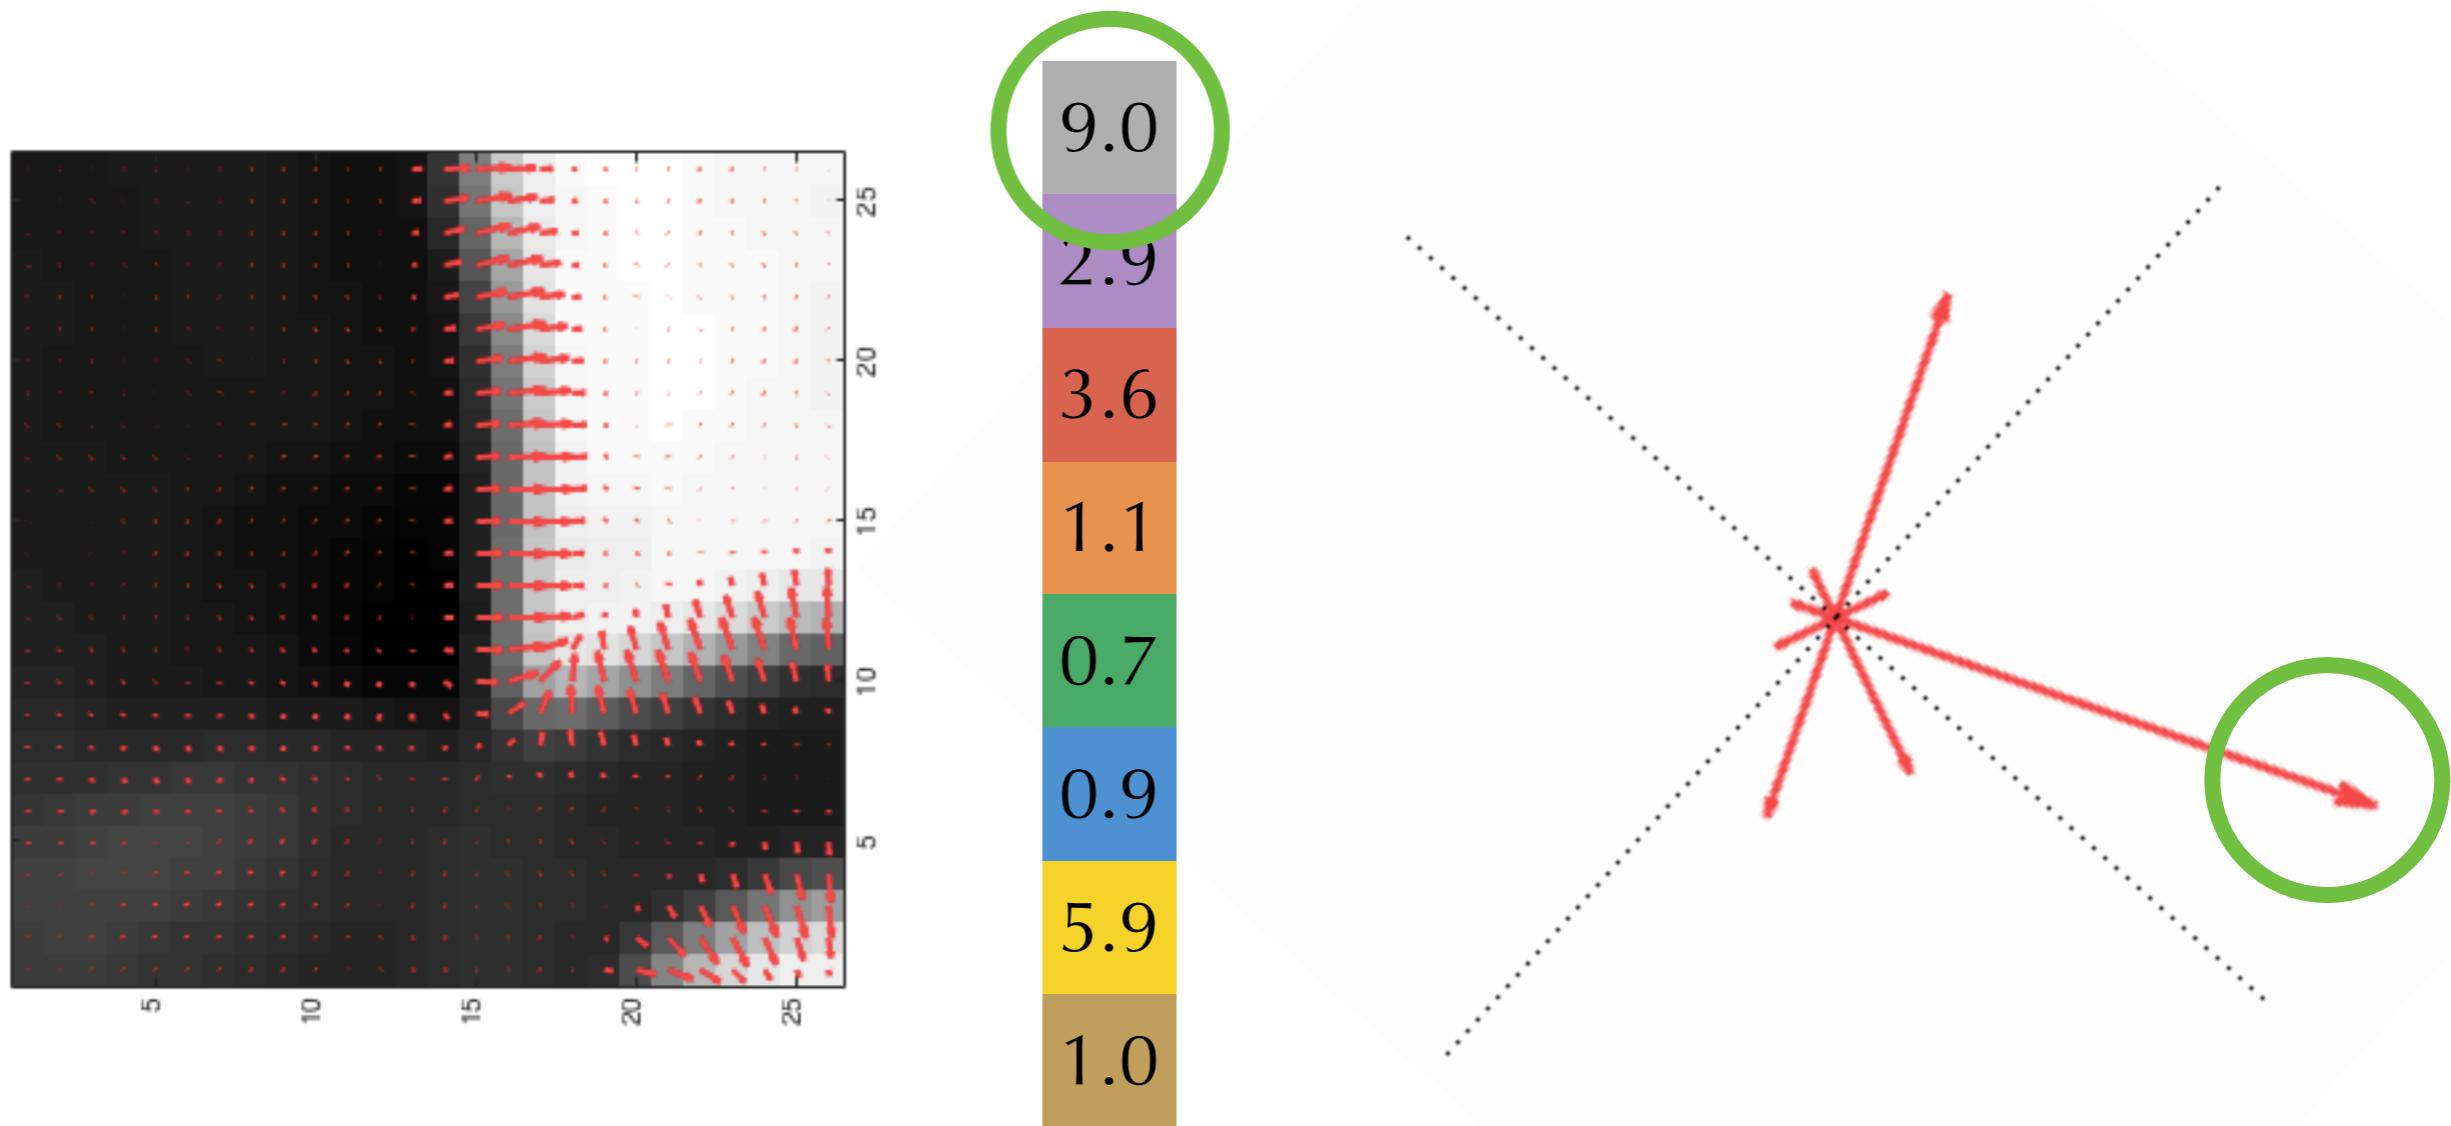
Histogram of gradients around the point.  
Use the dominant orientation.

# Orientation Assignment



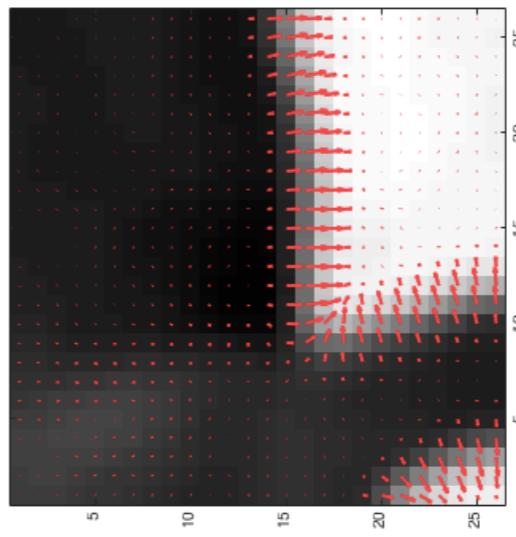
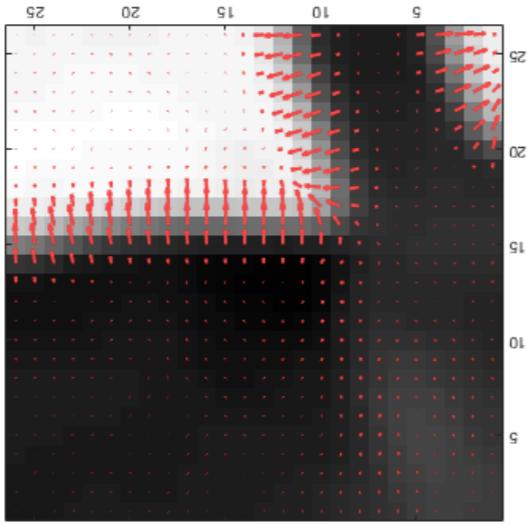
Histogram of gradients around the point.  
Use the dominant orientation.

# Orientation Assignment

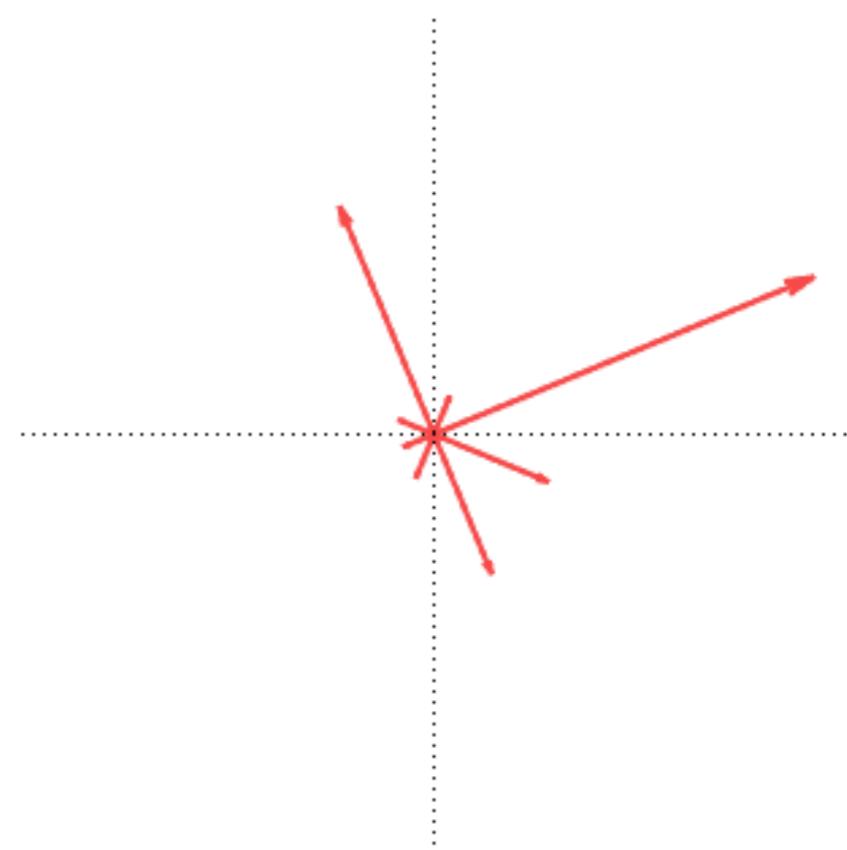
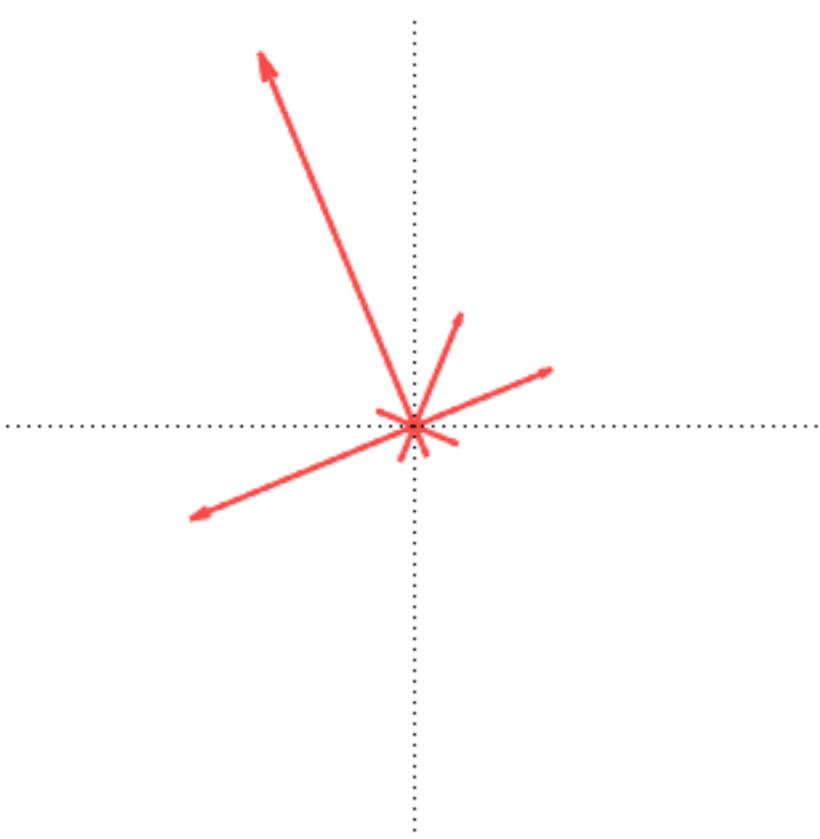
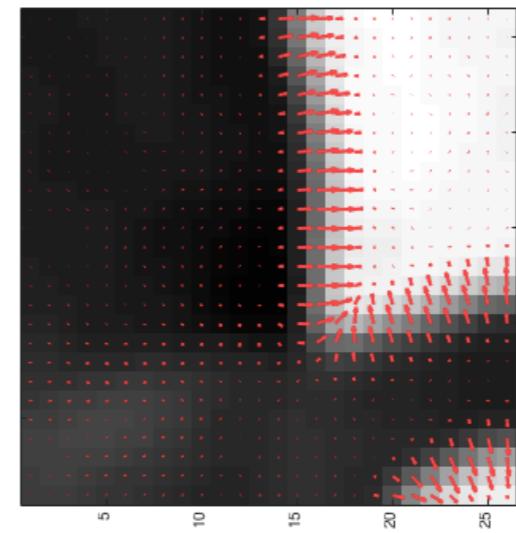
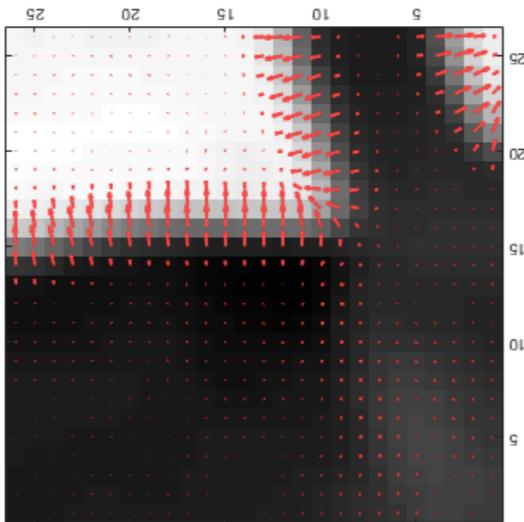


Histogram of gradients around the point.  
Use the dominant orientation.

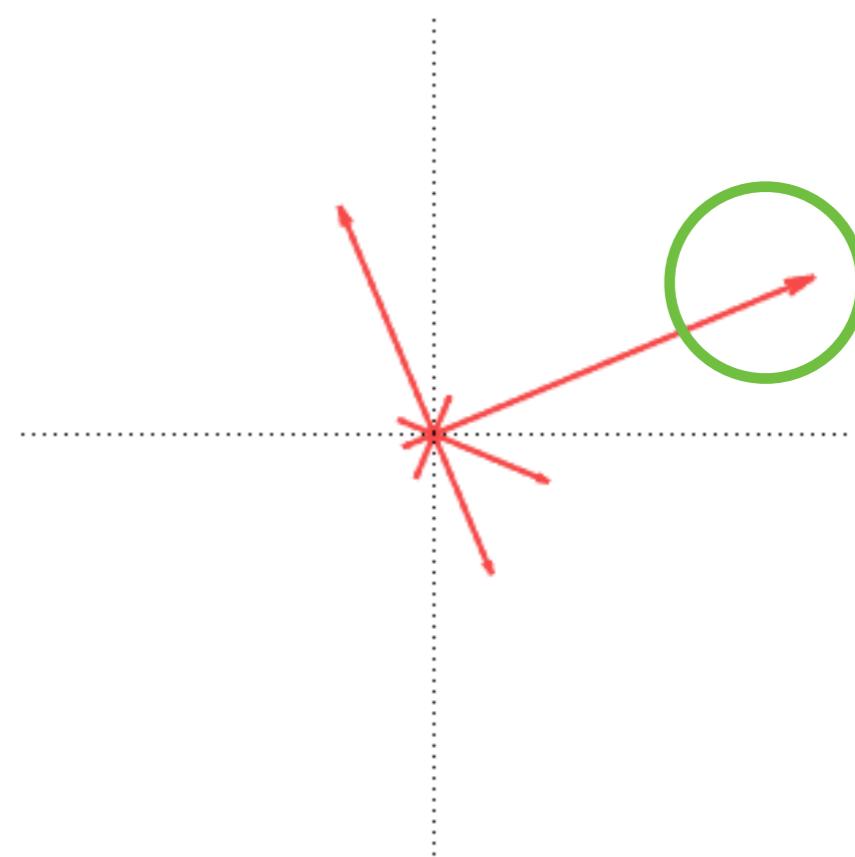
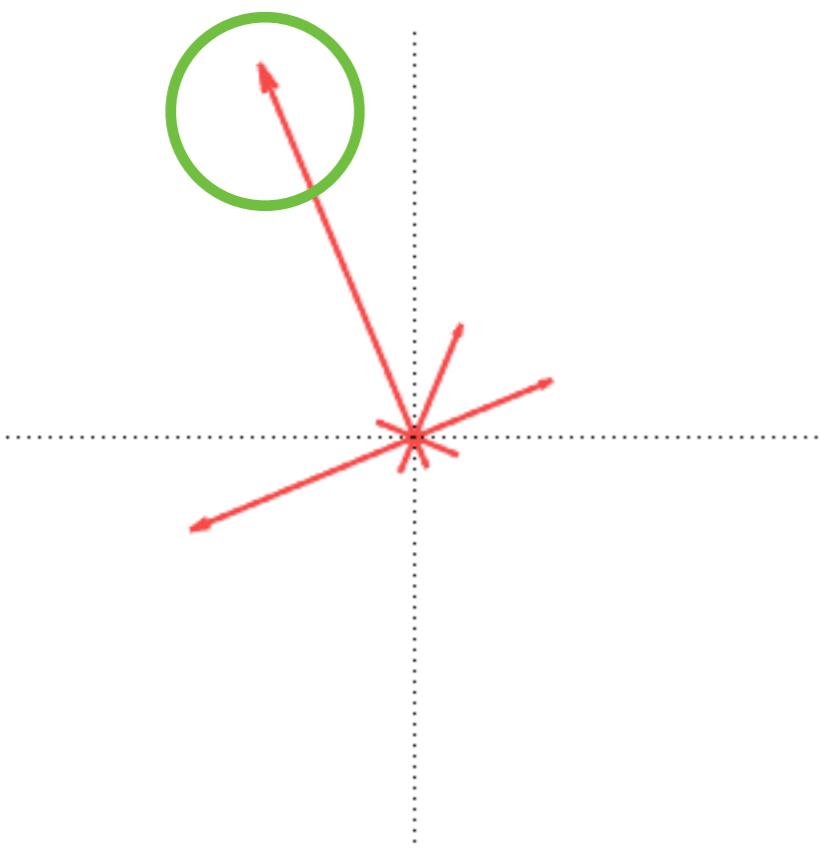
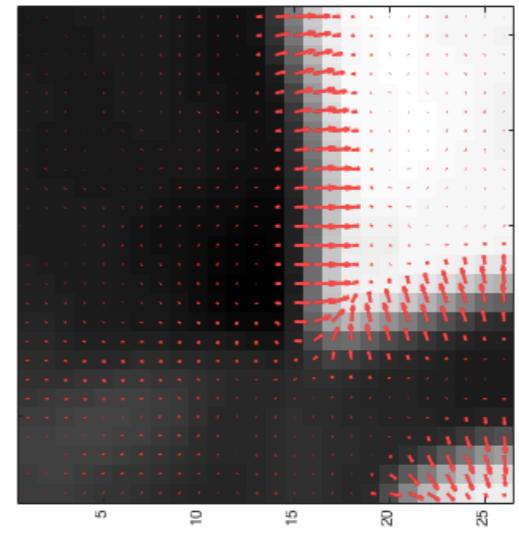
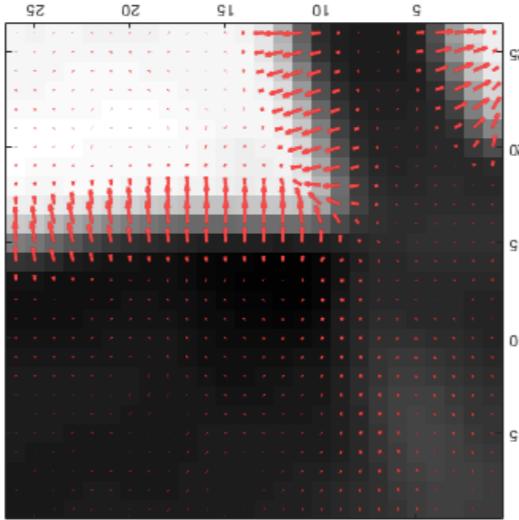
# Orientation Assignment



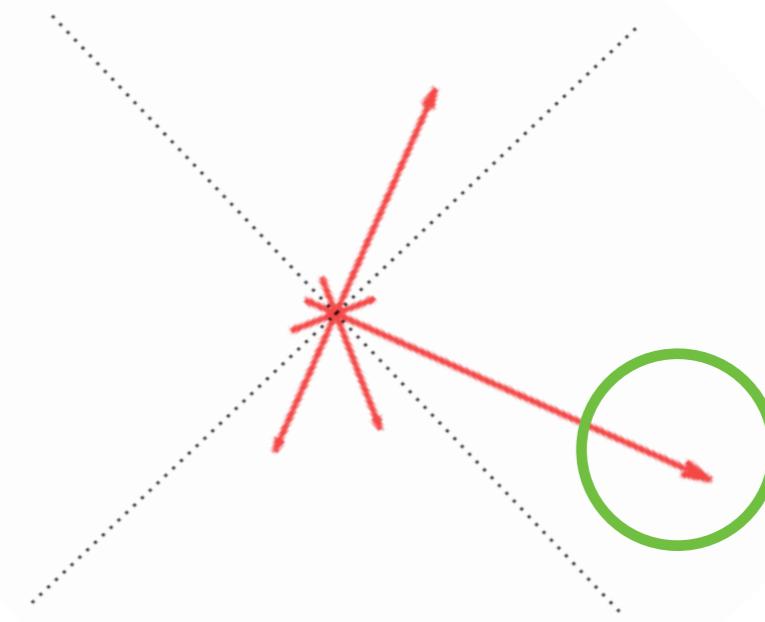
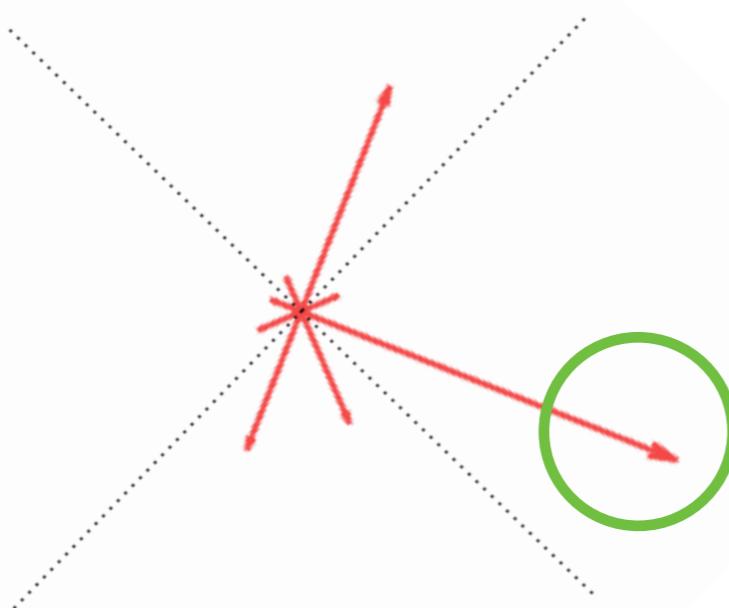
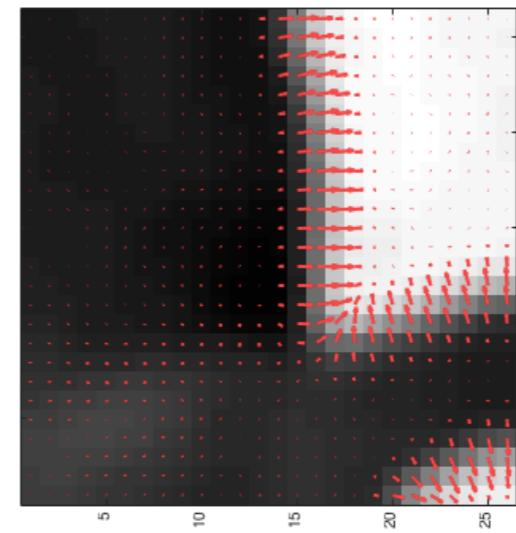
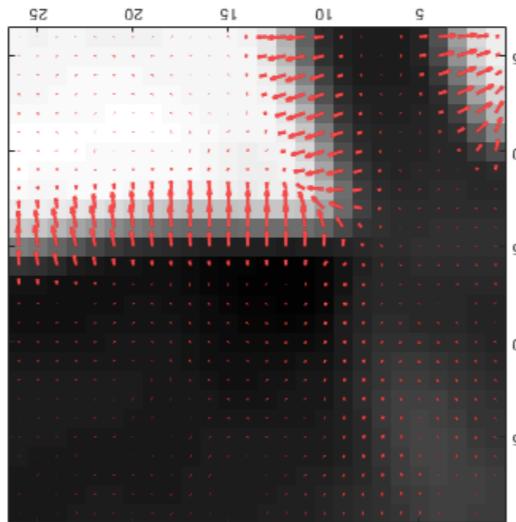
# Orientation Assignment



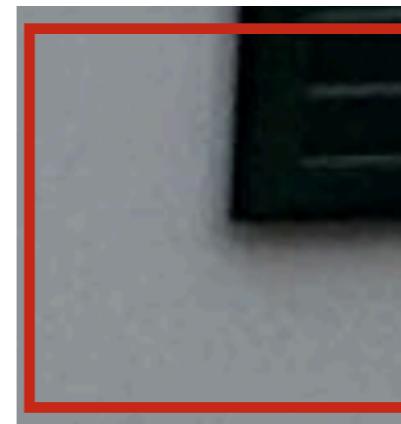
# Orientation Assignment



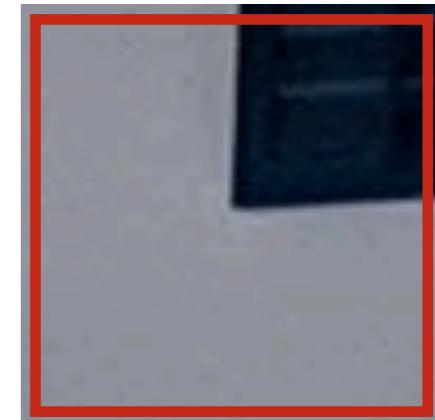
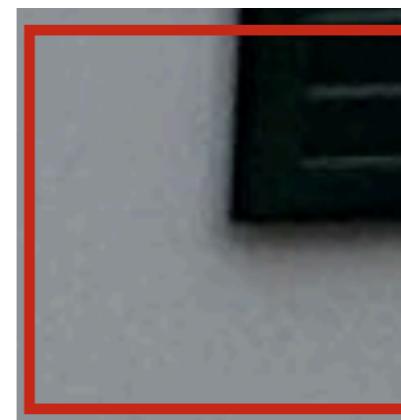
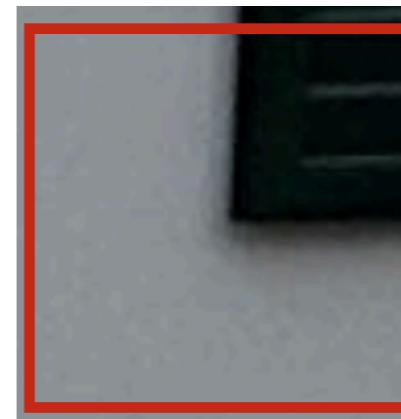
# Orientation Assignment



# The Perils of Rotation Invariance



# The Perils of Rotation Invariance



# Is Your Visual System Rotation Invariant?

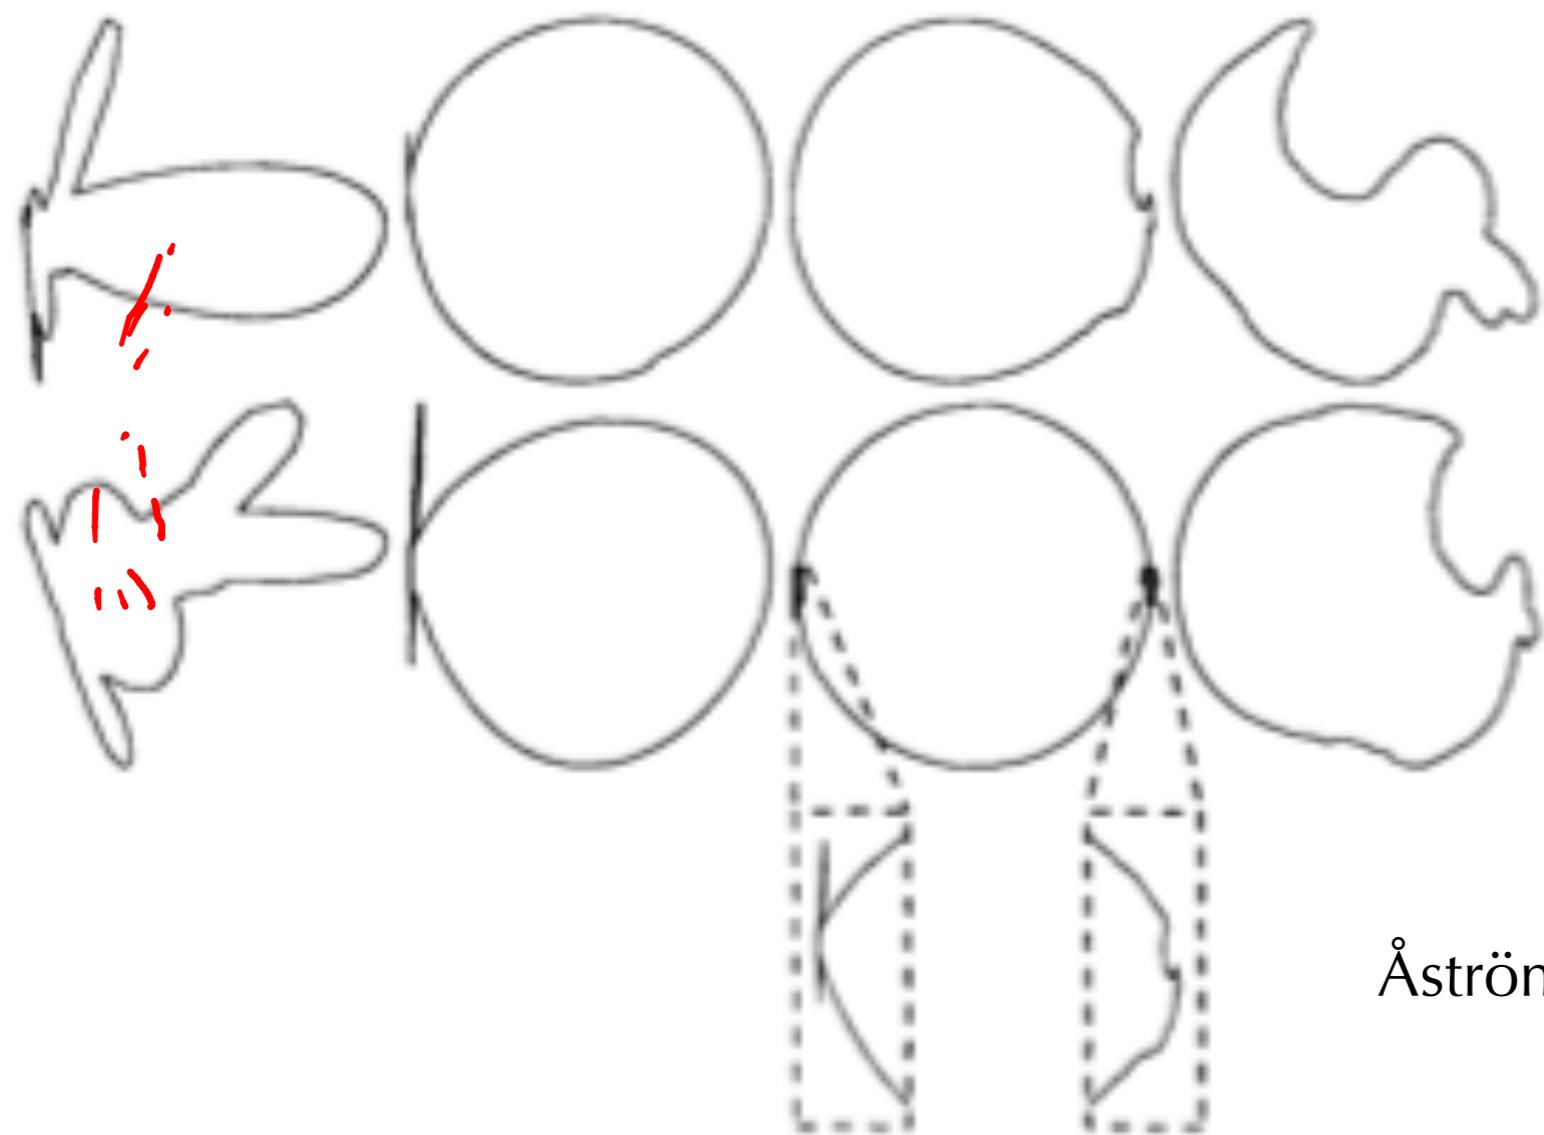


# Is Your Visual System Rotation Invariant?



Thatcher illusion

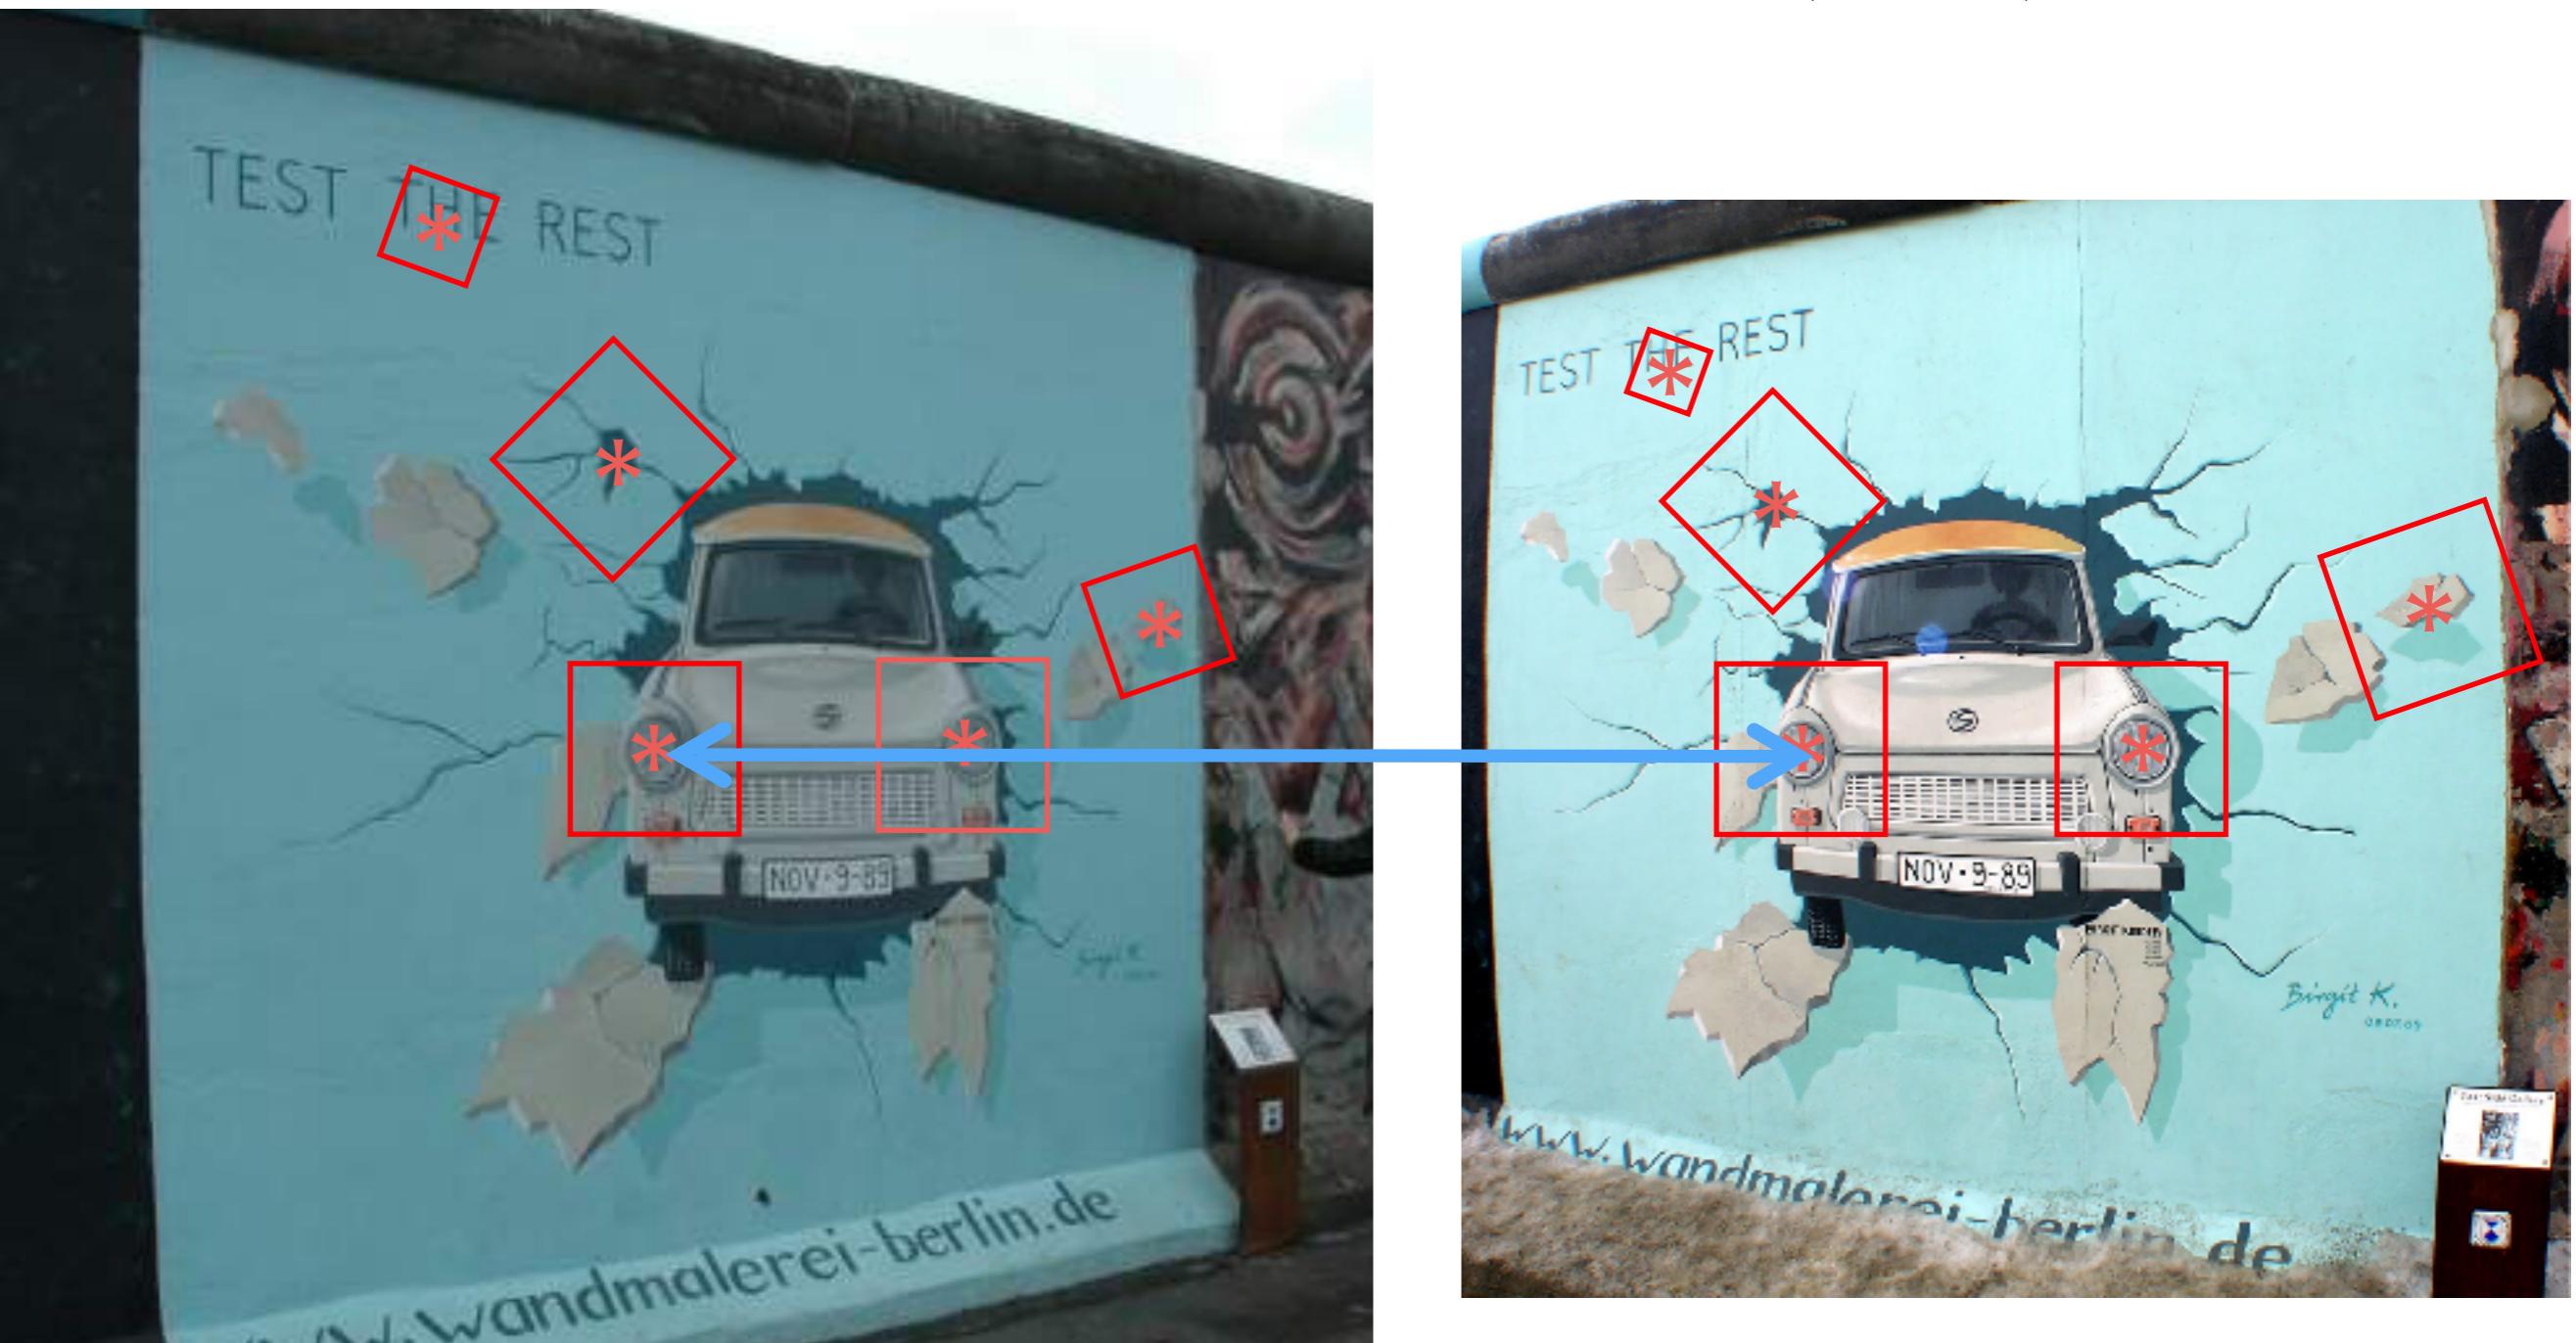
# Invariance vs. Robustness



Invariance = no influence on result

Robustness = we can deal with a moderate amount

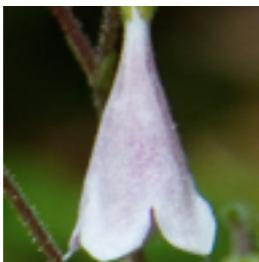
# Scale-Invariant Feature Transform (SIFT) Features



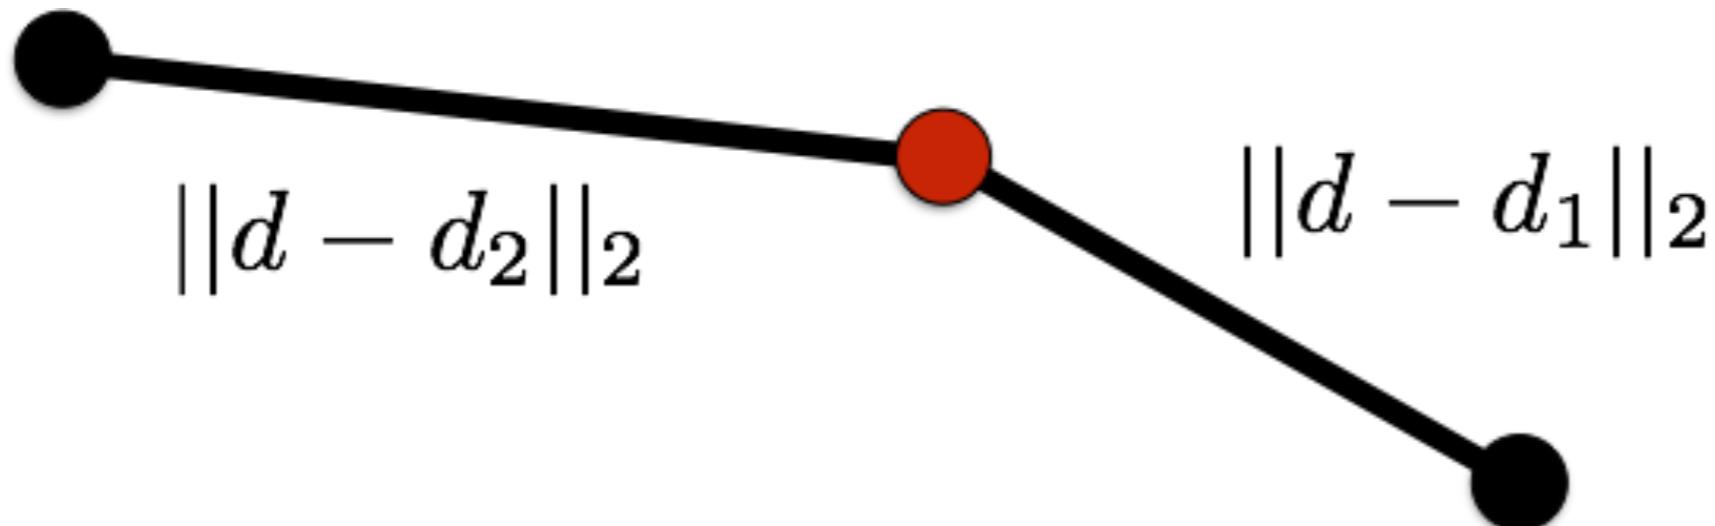
Use descriptor similarity ( $L_2$  distance) to find corresponding points in different images

# $L_2$ Similar = Correct?

query:



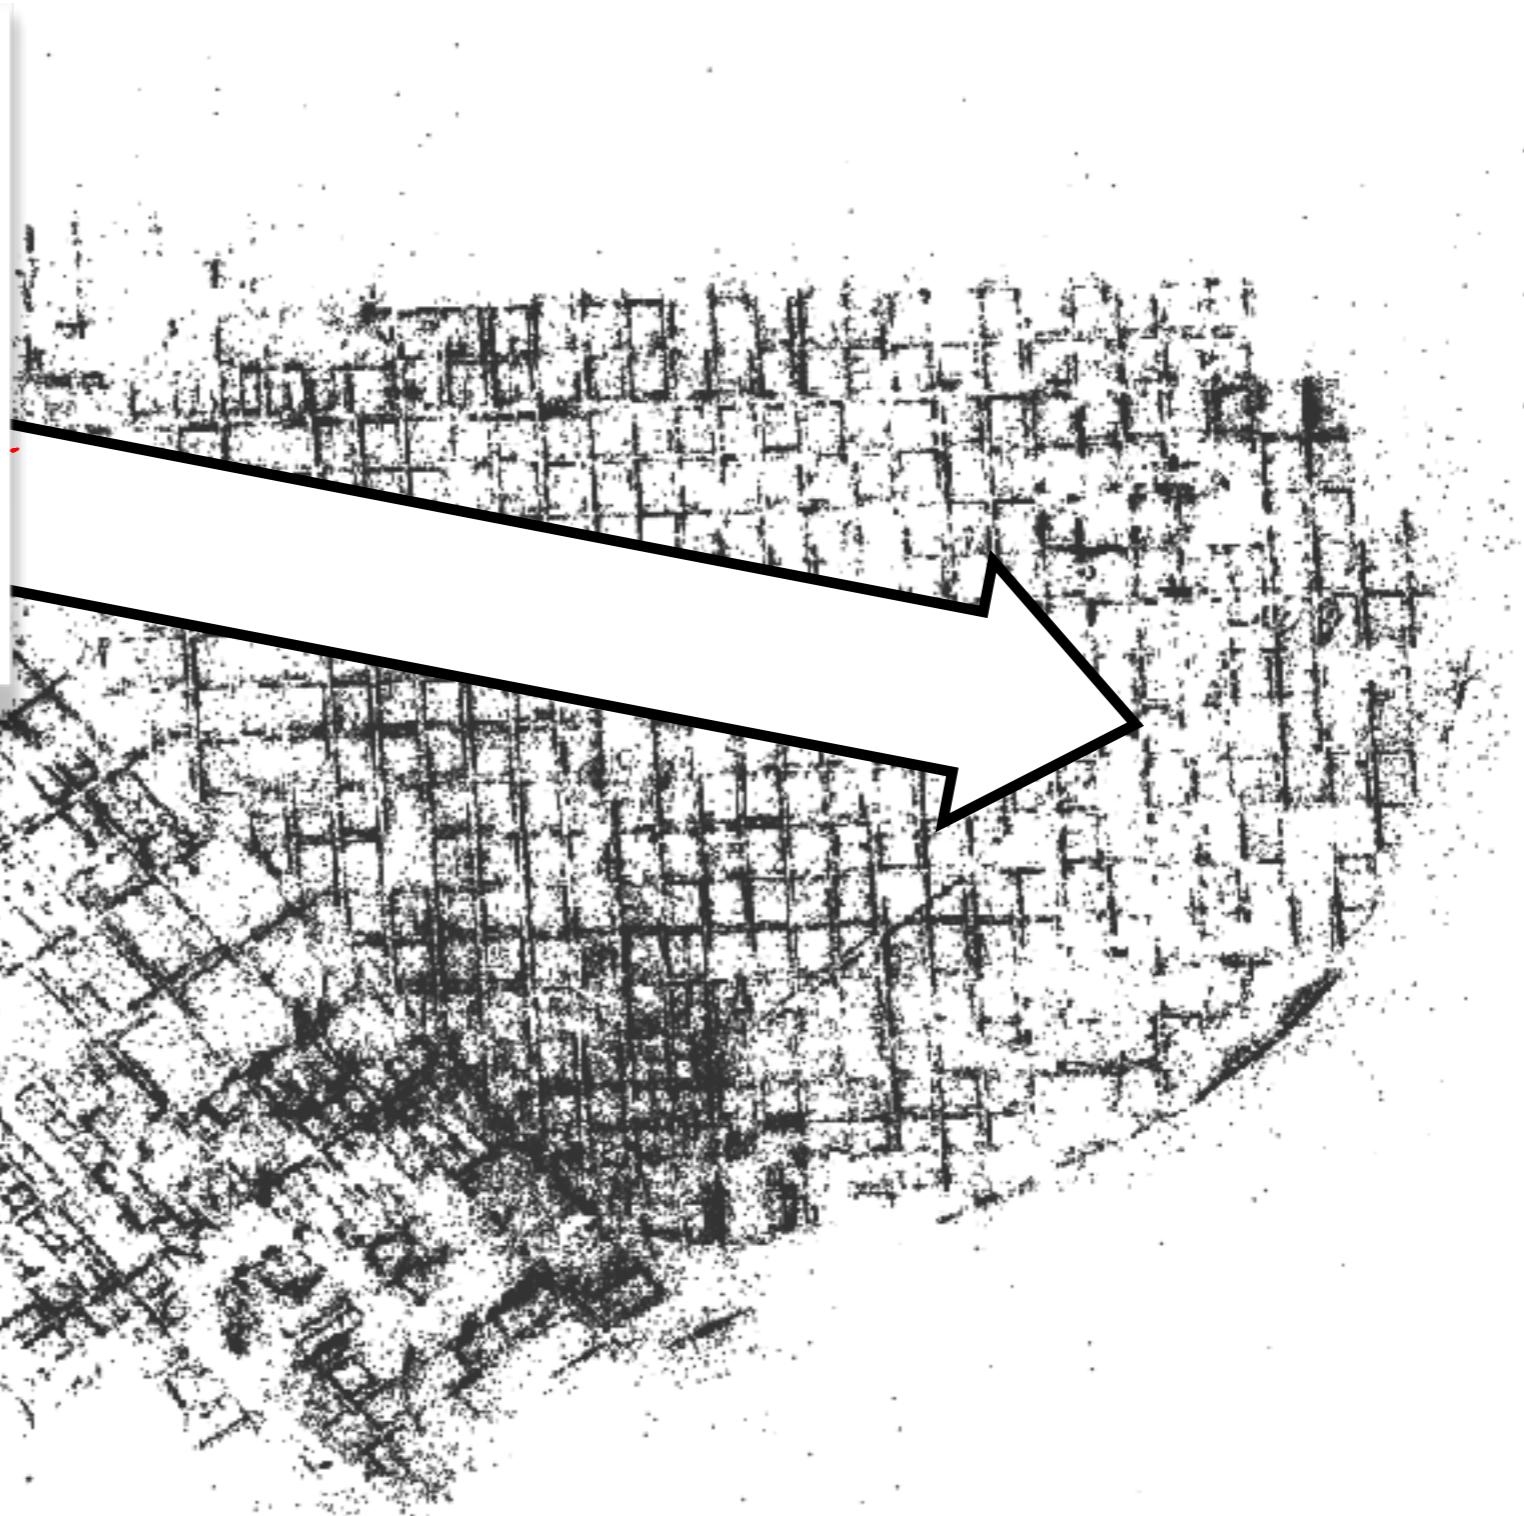
# Lowe's Ratio Test



Accept match with 1<sup>st</sup> neighbor if  $\frac{\|d - d_1\|_2}{\|d - d_2\|_2} < 0.8$

$0.9$   
 $> 0.6$

# City-Scale Localization



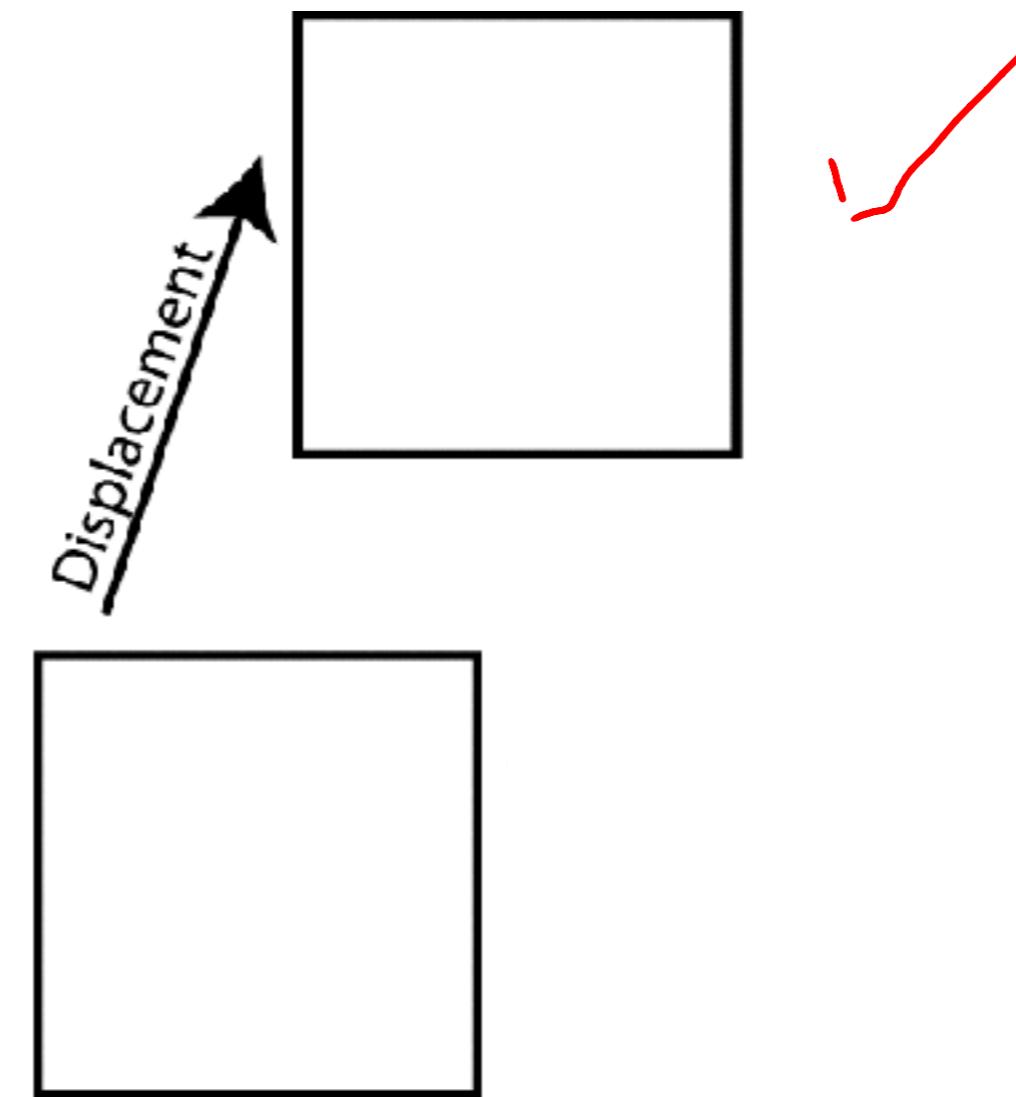
# SIFT Features

- Skipped some details:
  - Down-weighting of pixels far away from center of spatial bin
  - Bilinear interpolation between histogram bins
  - Bilinear interpolation between spatial bins
  - Thresholding on histogram entries & re-normalization
  - For further details, read Lowe's 2004 IJCV paper

# Today

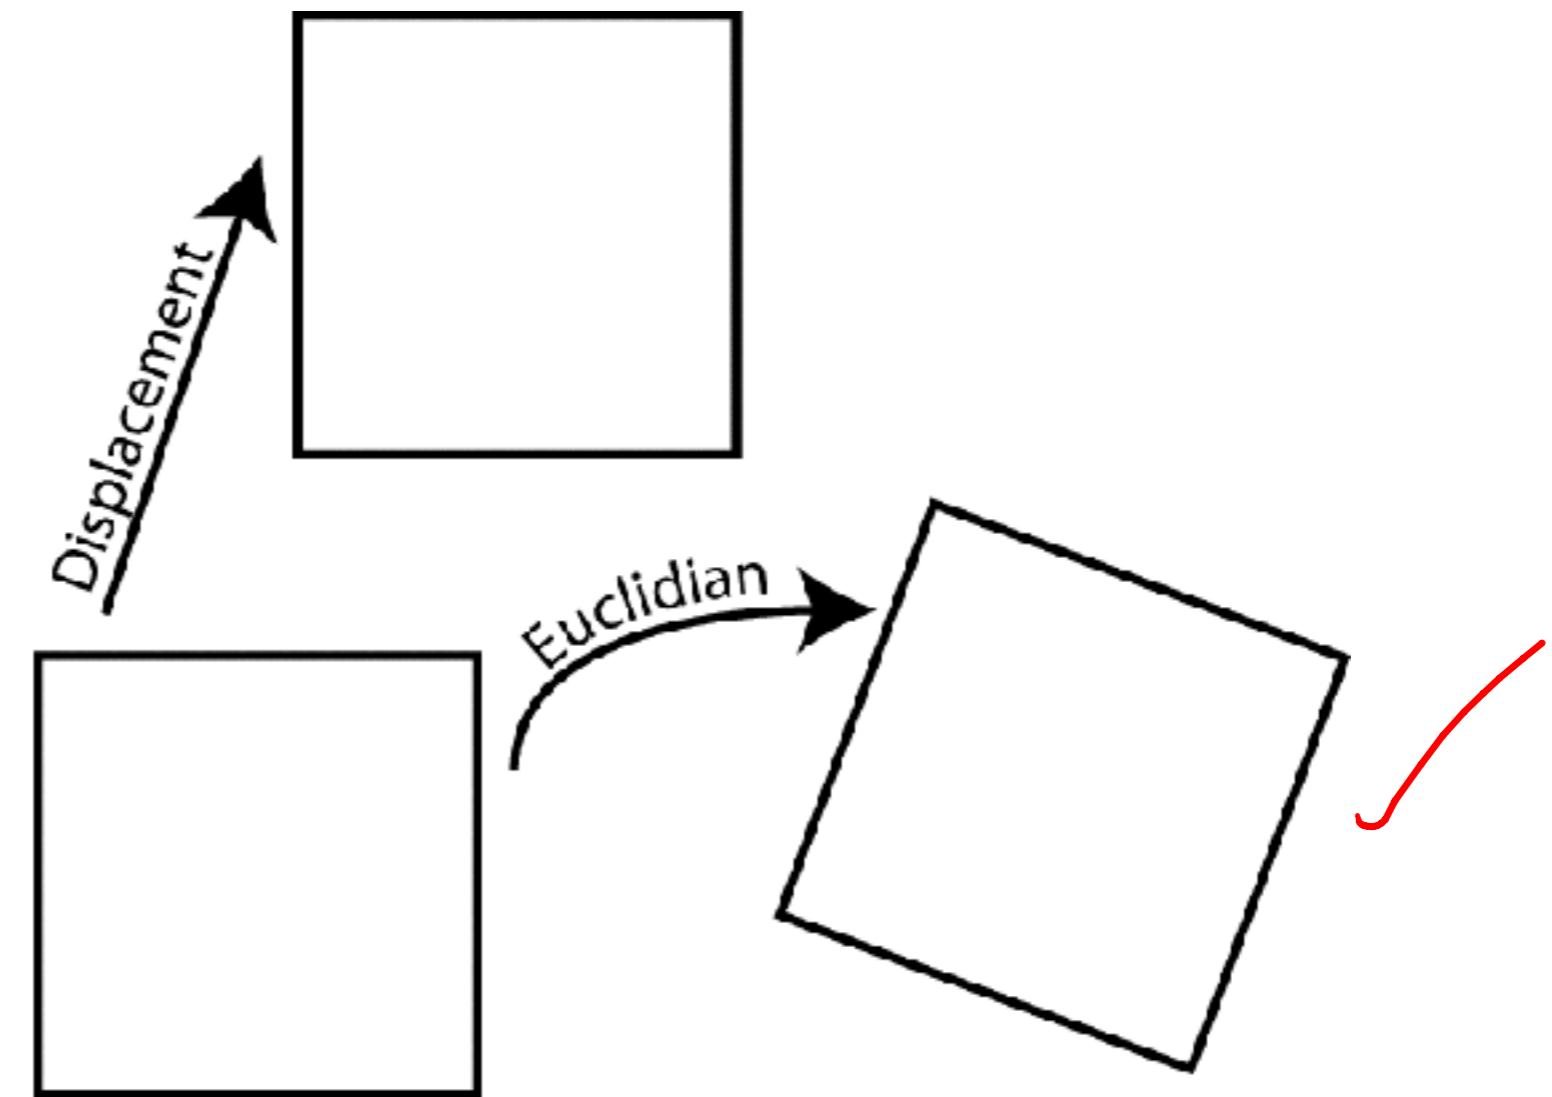
- Scale-Invariant Local Features
- Invariances\*

# Patch Transformations



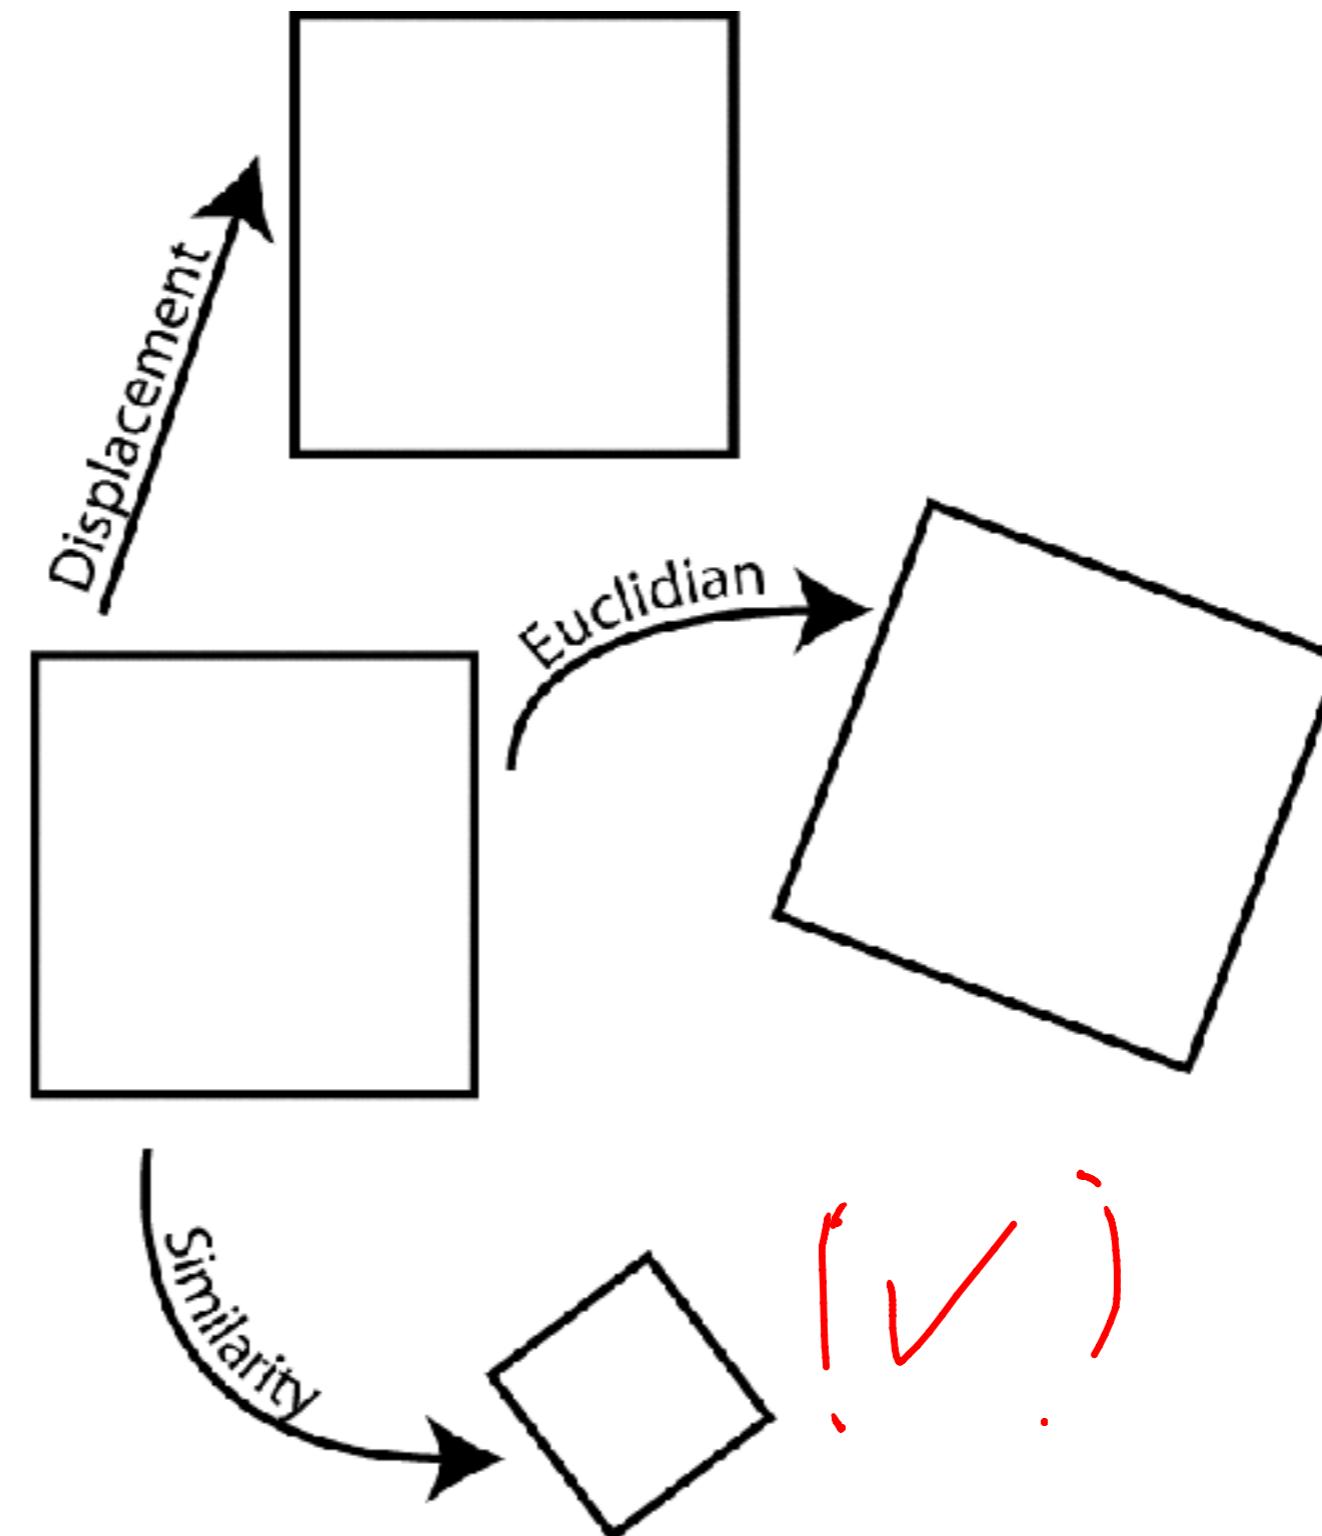
slide credit: Marc Pollefeys

# Patch Transformations



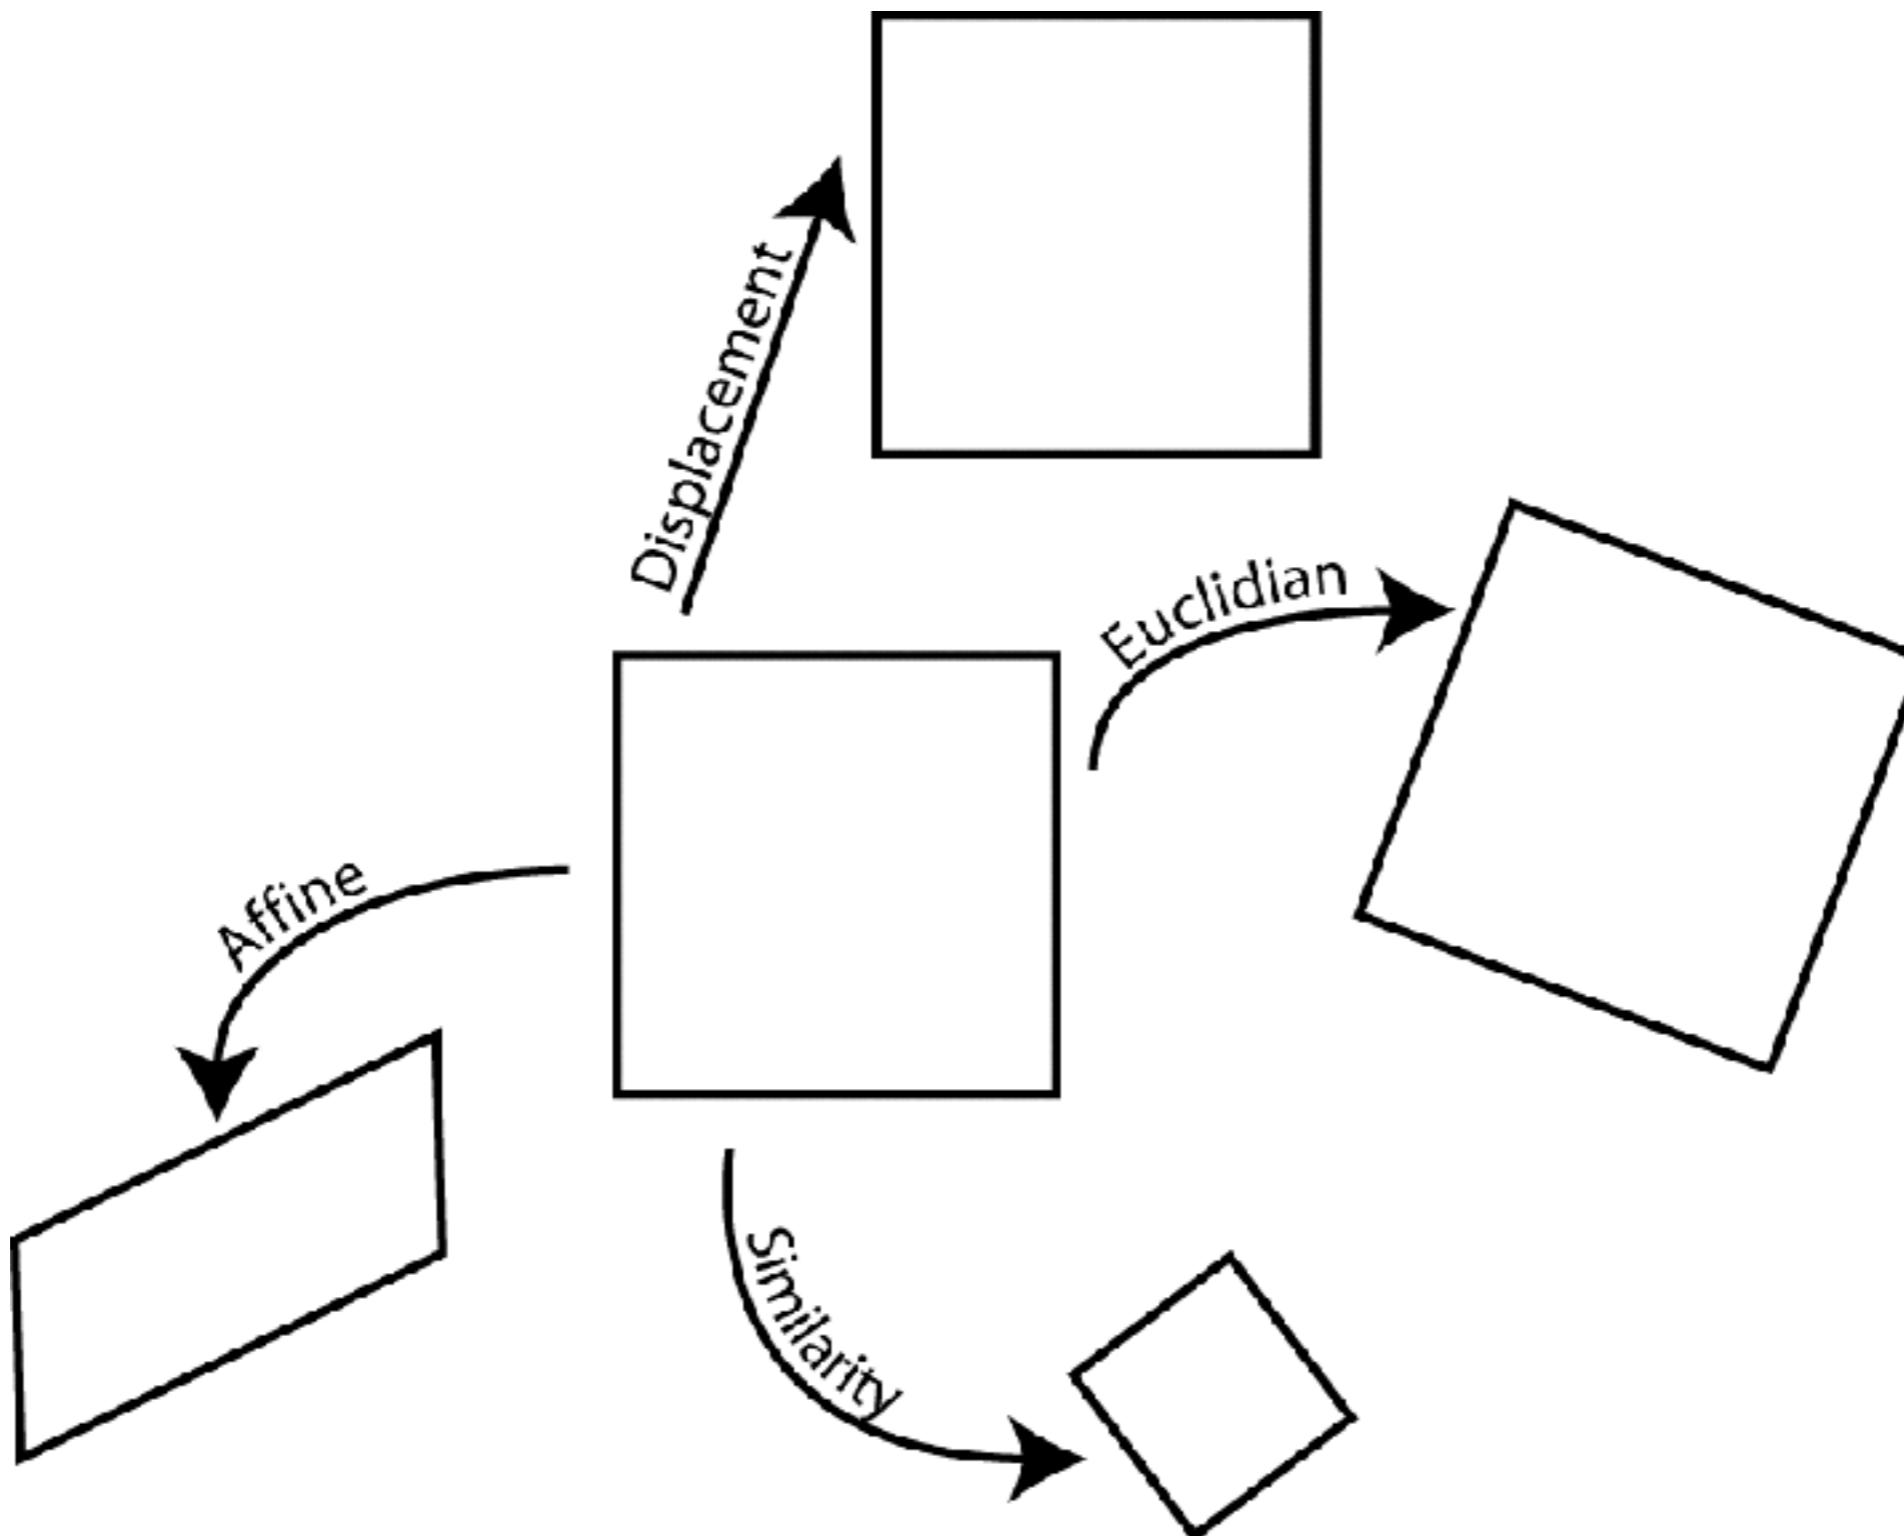
slide credit: Marc Pollefeys

# Patch Transformations



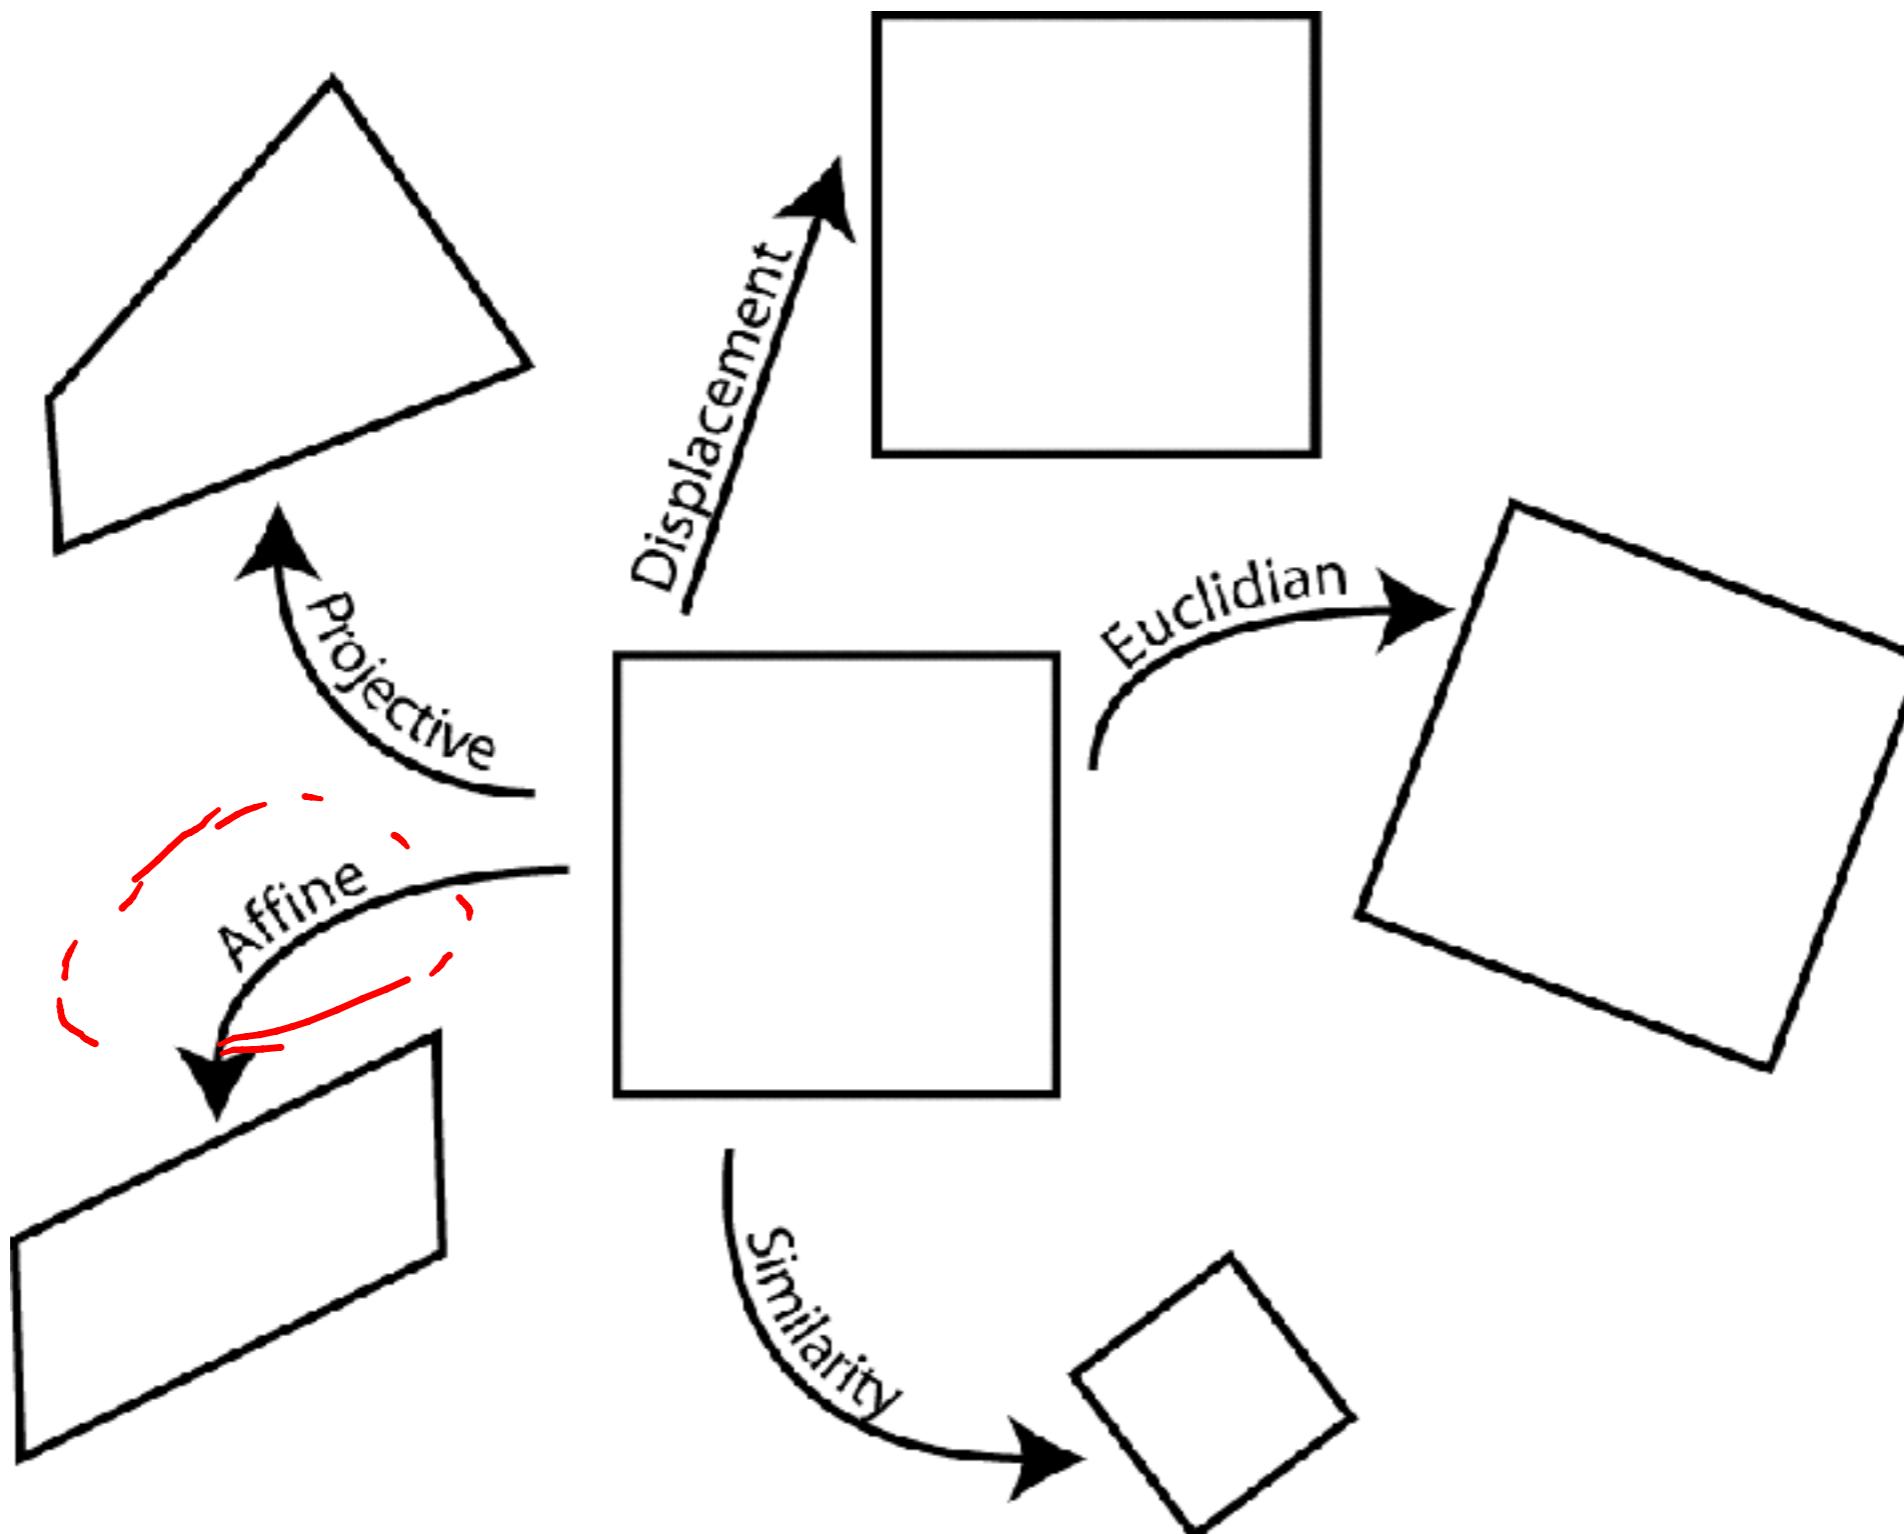
slide credit: Marc Pollefeys

# Patch Transformations



slide credit: Marc Pollefeys

# Patch Transformations



slide credit: Marc Pollefeys

# Projective Transformations



Affine transformations can locally approximate perspective transformations quite well

slide credit: Marc Pollefeyns

# Projective Transformations



Affine transformations can locally approximate perspective transformations quite well

slide credit: Marc Pollefey

# Projective Transformations



Affine transformations can locally approximate perspective transformations quite well

slide credit: Marc Pollefeyns

# Lessons Learned

- Main lessons from this lecture
  - SIFT feature detector: Difference of Gaussians over scale space
  - Extraction of oriented and scale-dependent patches
- SIFT feature descriptor: Concatenation of gradient histograms

# Lessons Learned

- Main lessons from this lecture
  - SIFT feature detector: Difference of Gaussians over scale space
  - Extraction of oriented and scale-dependent patches
- SIFT feature descriptor: Concatenation of gradient histograms
- Next lecture: Learning classifiers