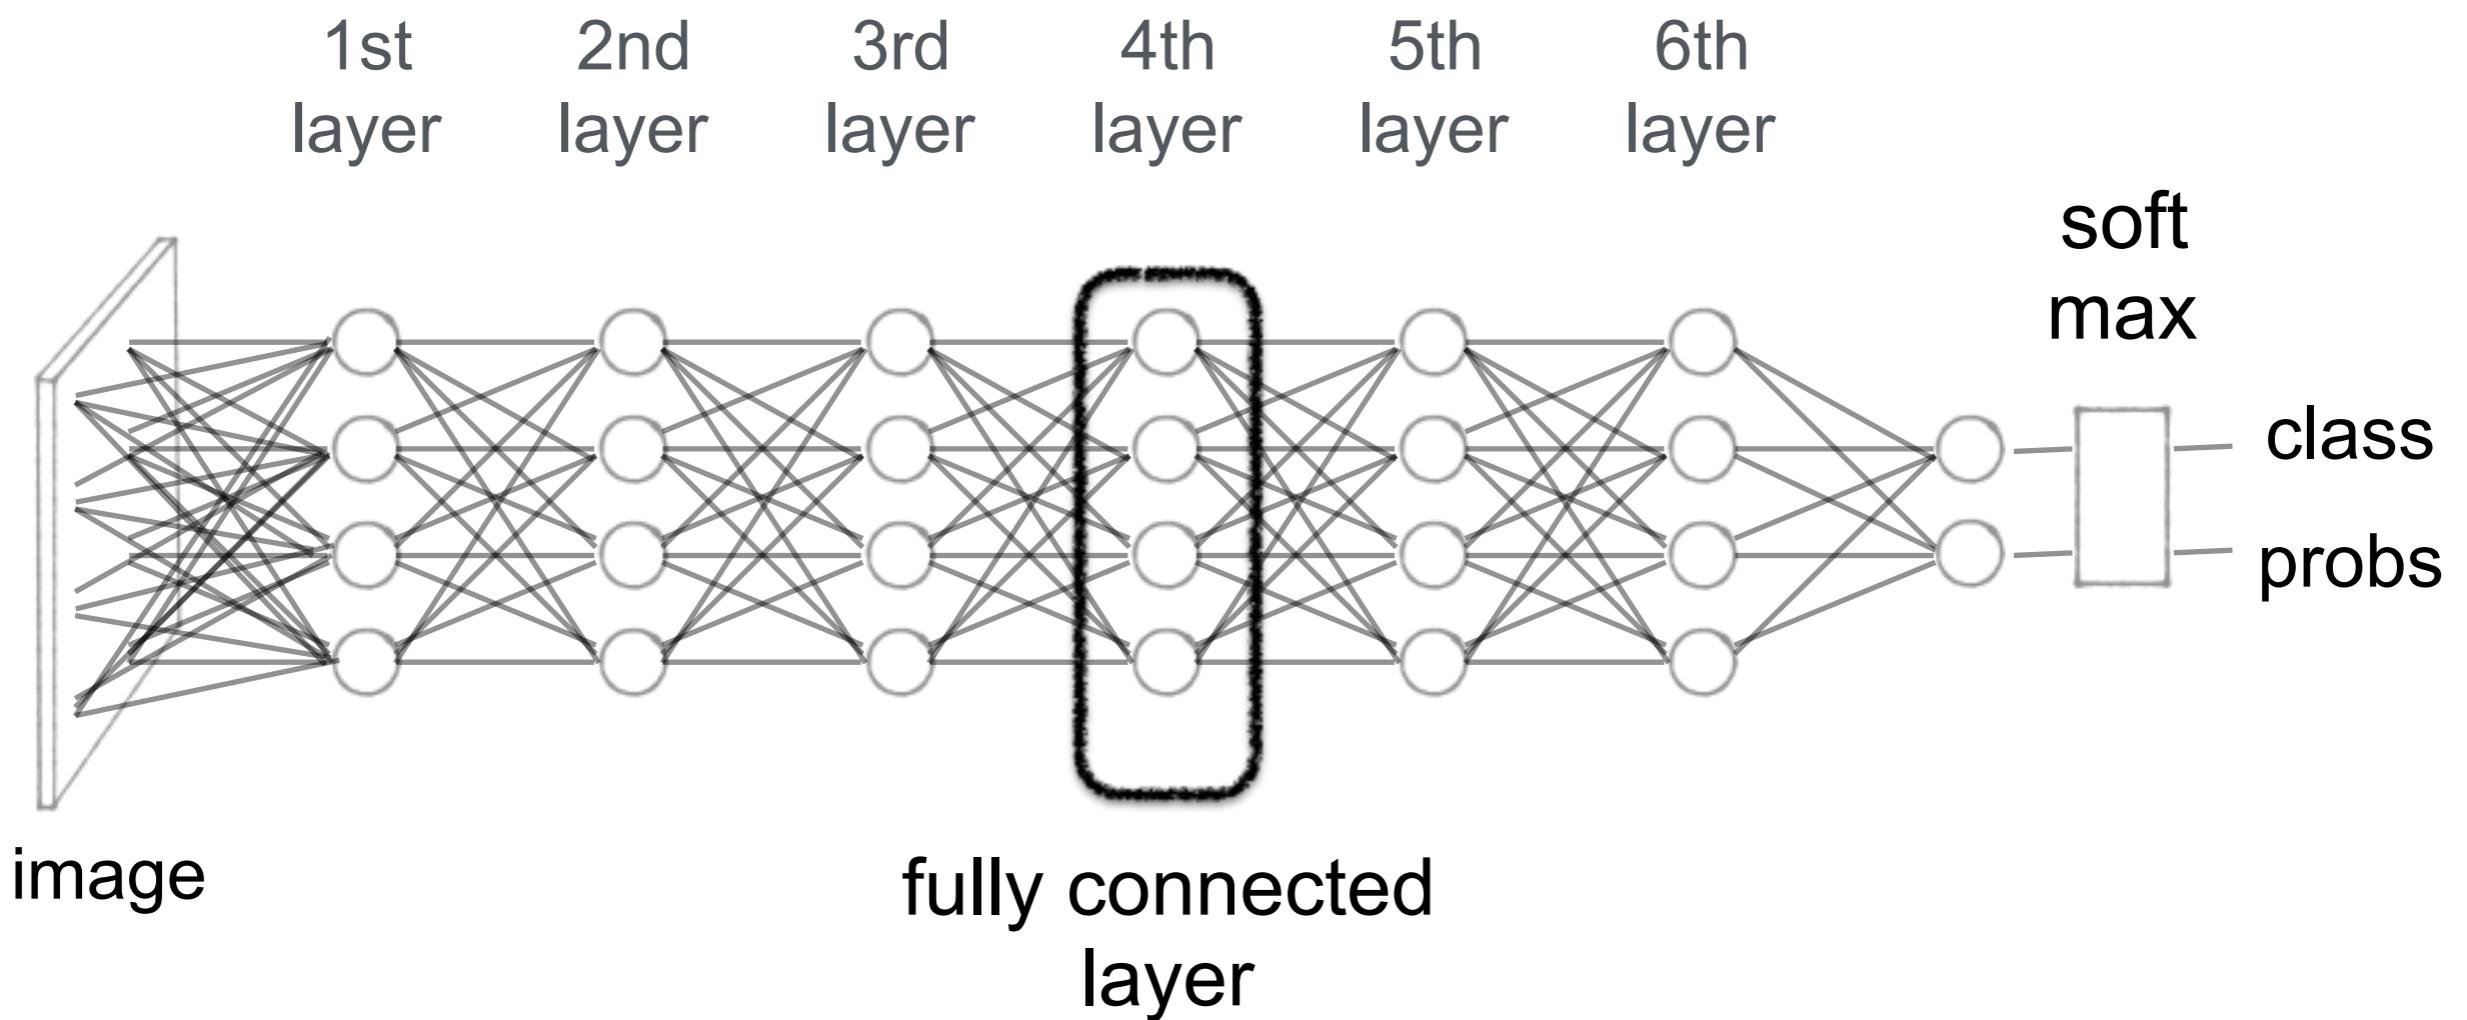


SSY097 - Image Analysis

Lecture 6 - (More) Convolutional Neural Networks

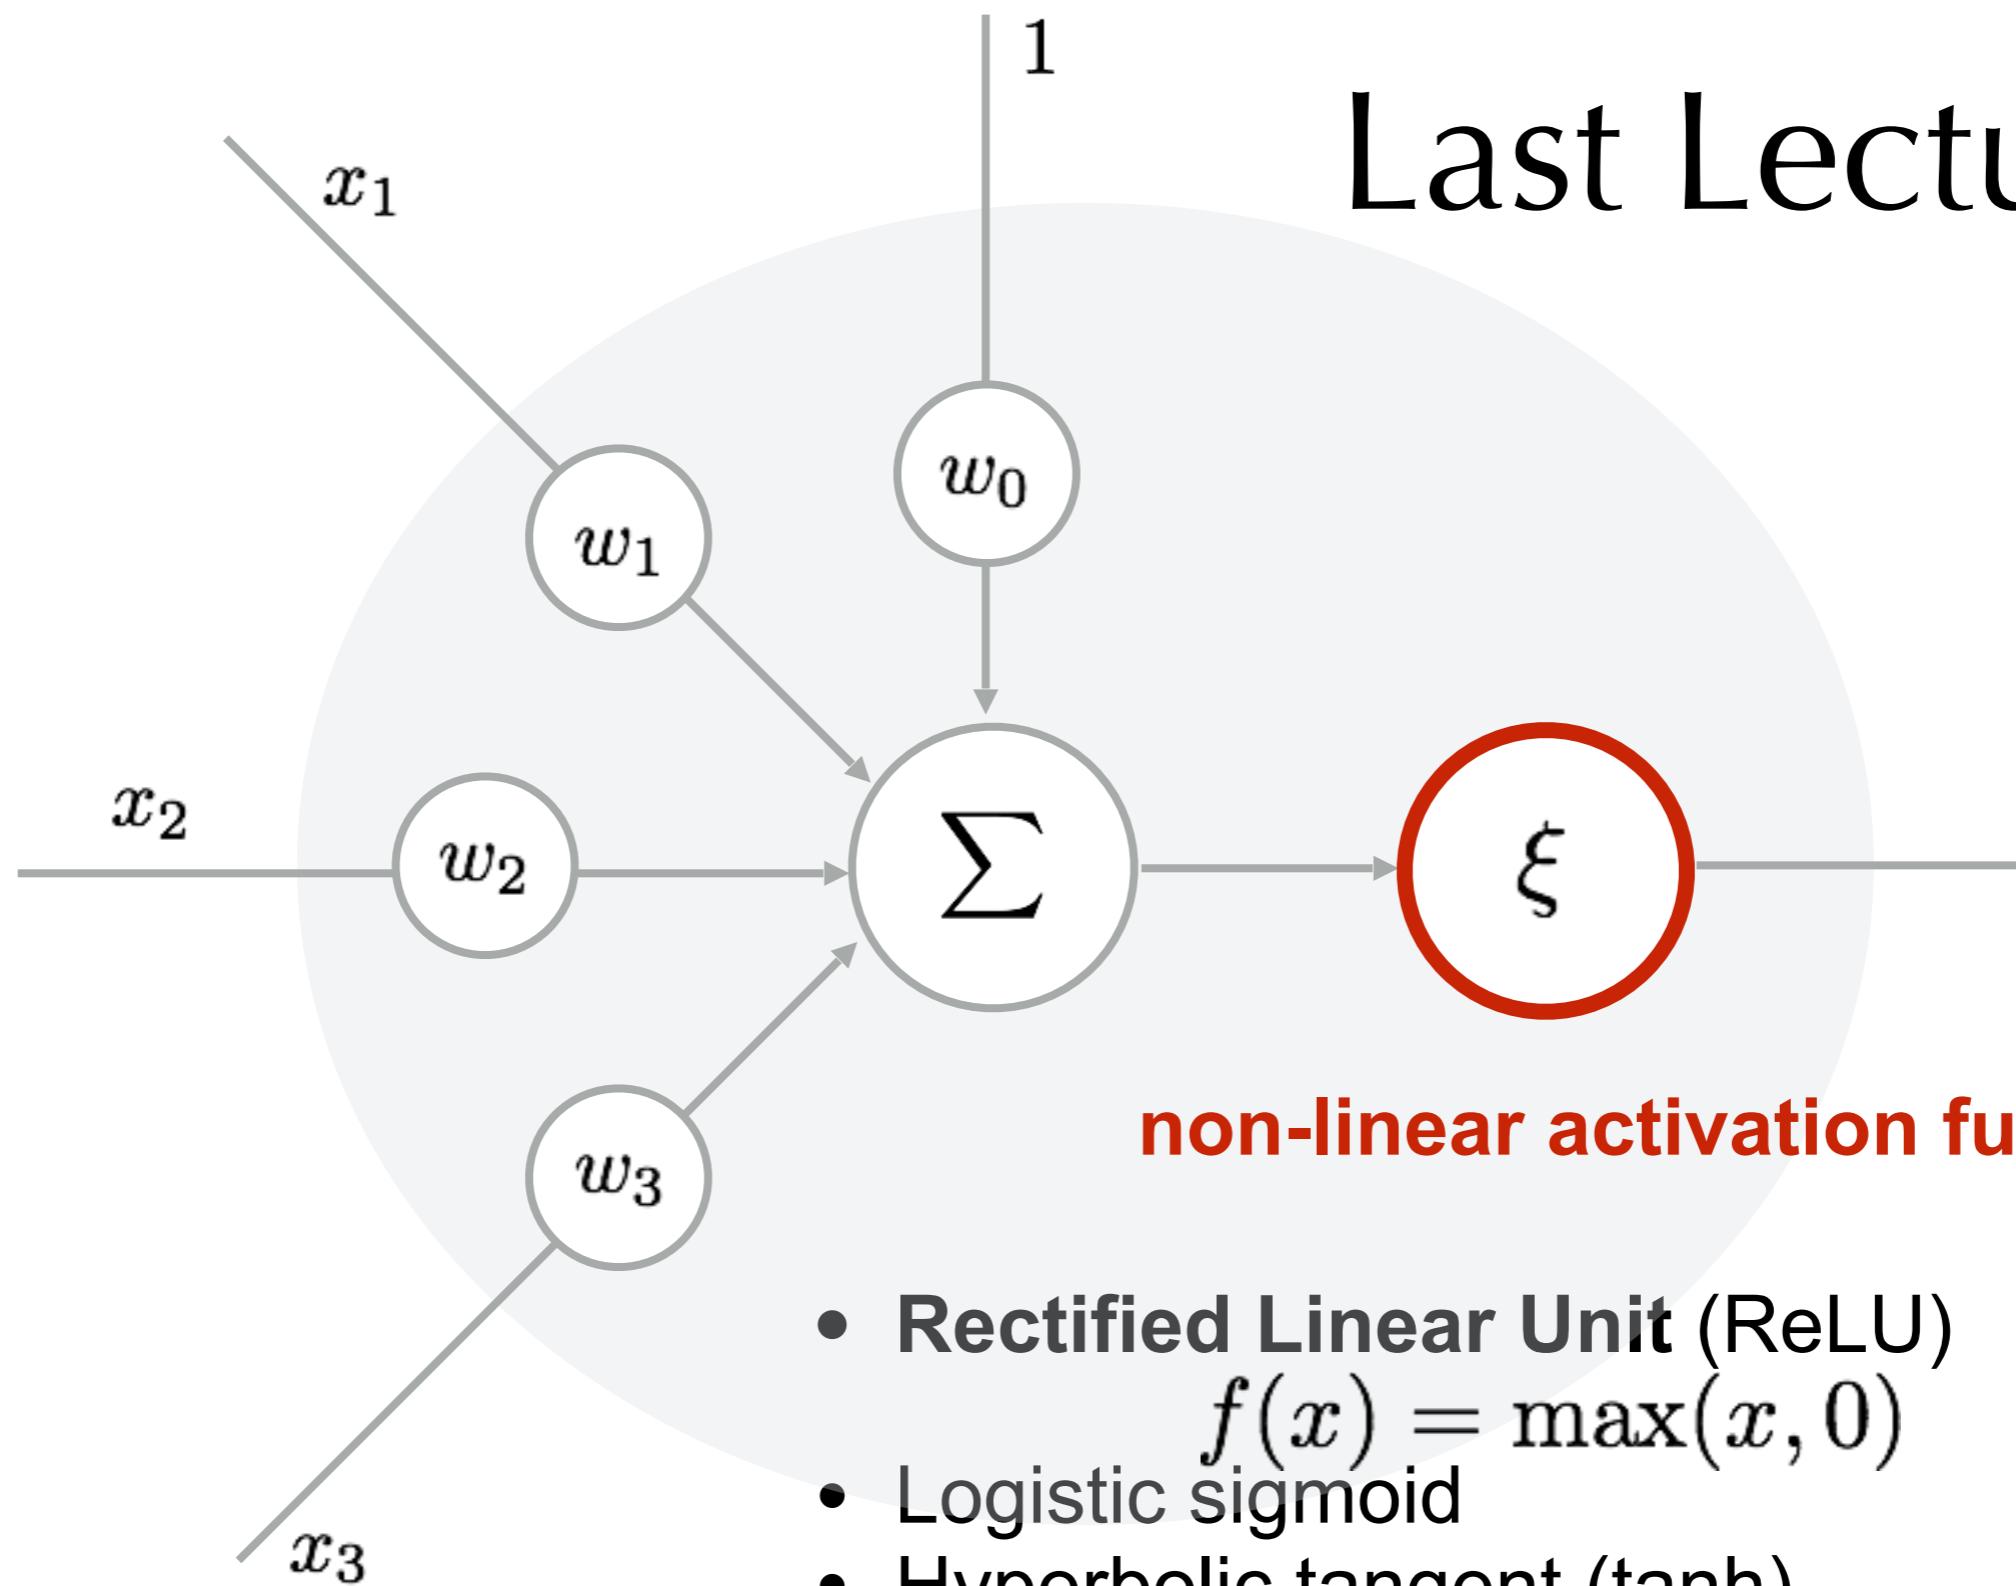
*Torsten Sattler
(slides adapted from Olof Enqvist)*

Last Lecture



Neural Networks

Last Lecture



non-linear activation function

- Rectified Linear Unit (ReLU)
 $f(x) = \max(x, 0)$
- Logistic sigmoid
- Hyperbolic tangent (tanh)
- ...

A neuron

Last Lecture

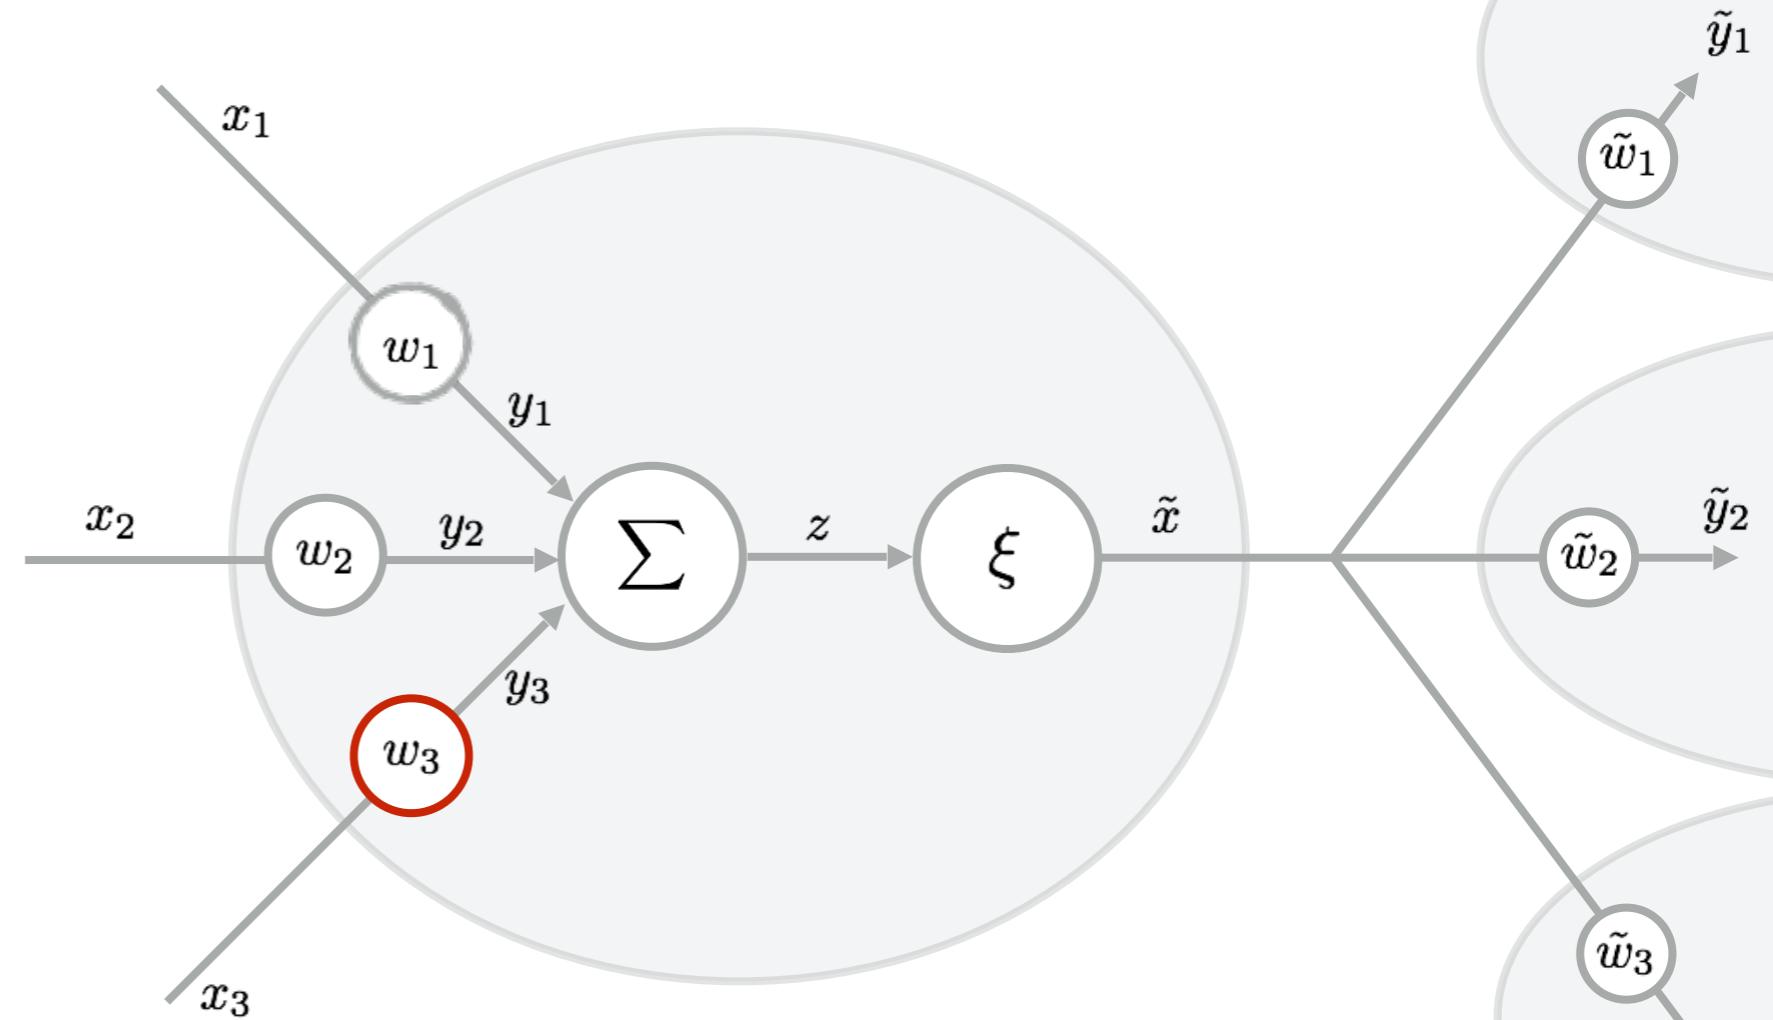
- Optimize all network parameters via gradient descent:

$$\theta^{(k+1)} = \theta^{(k)} - \mu \sum_i \nabla L_i(\theta) \approx \theta^{(k)} - \mu \nabla L_i(\theta)$$

- Compute derivative of partial loss for each parameter w_k

Training a neural network

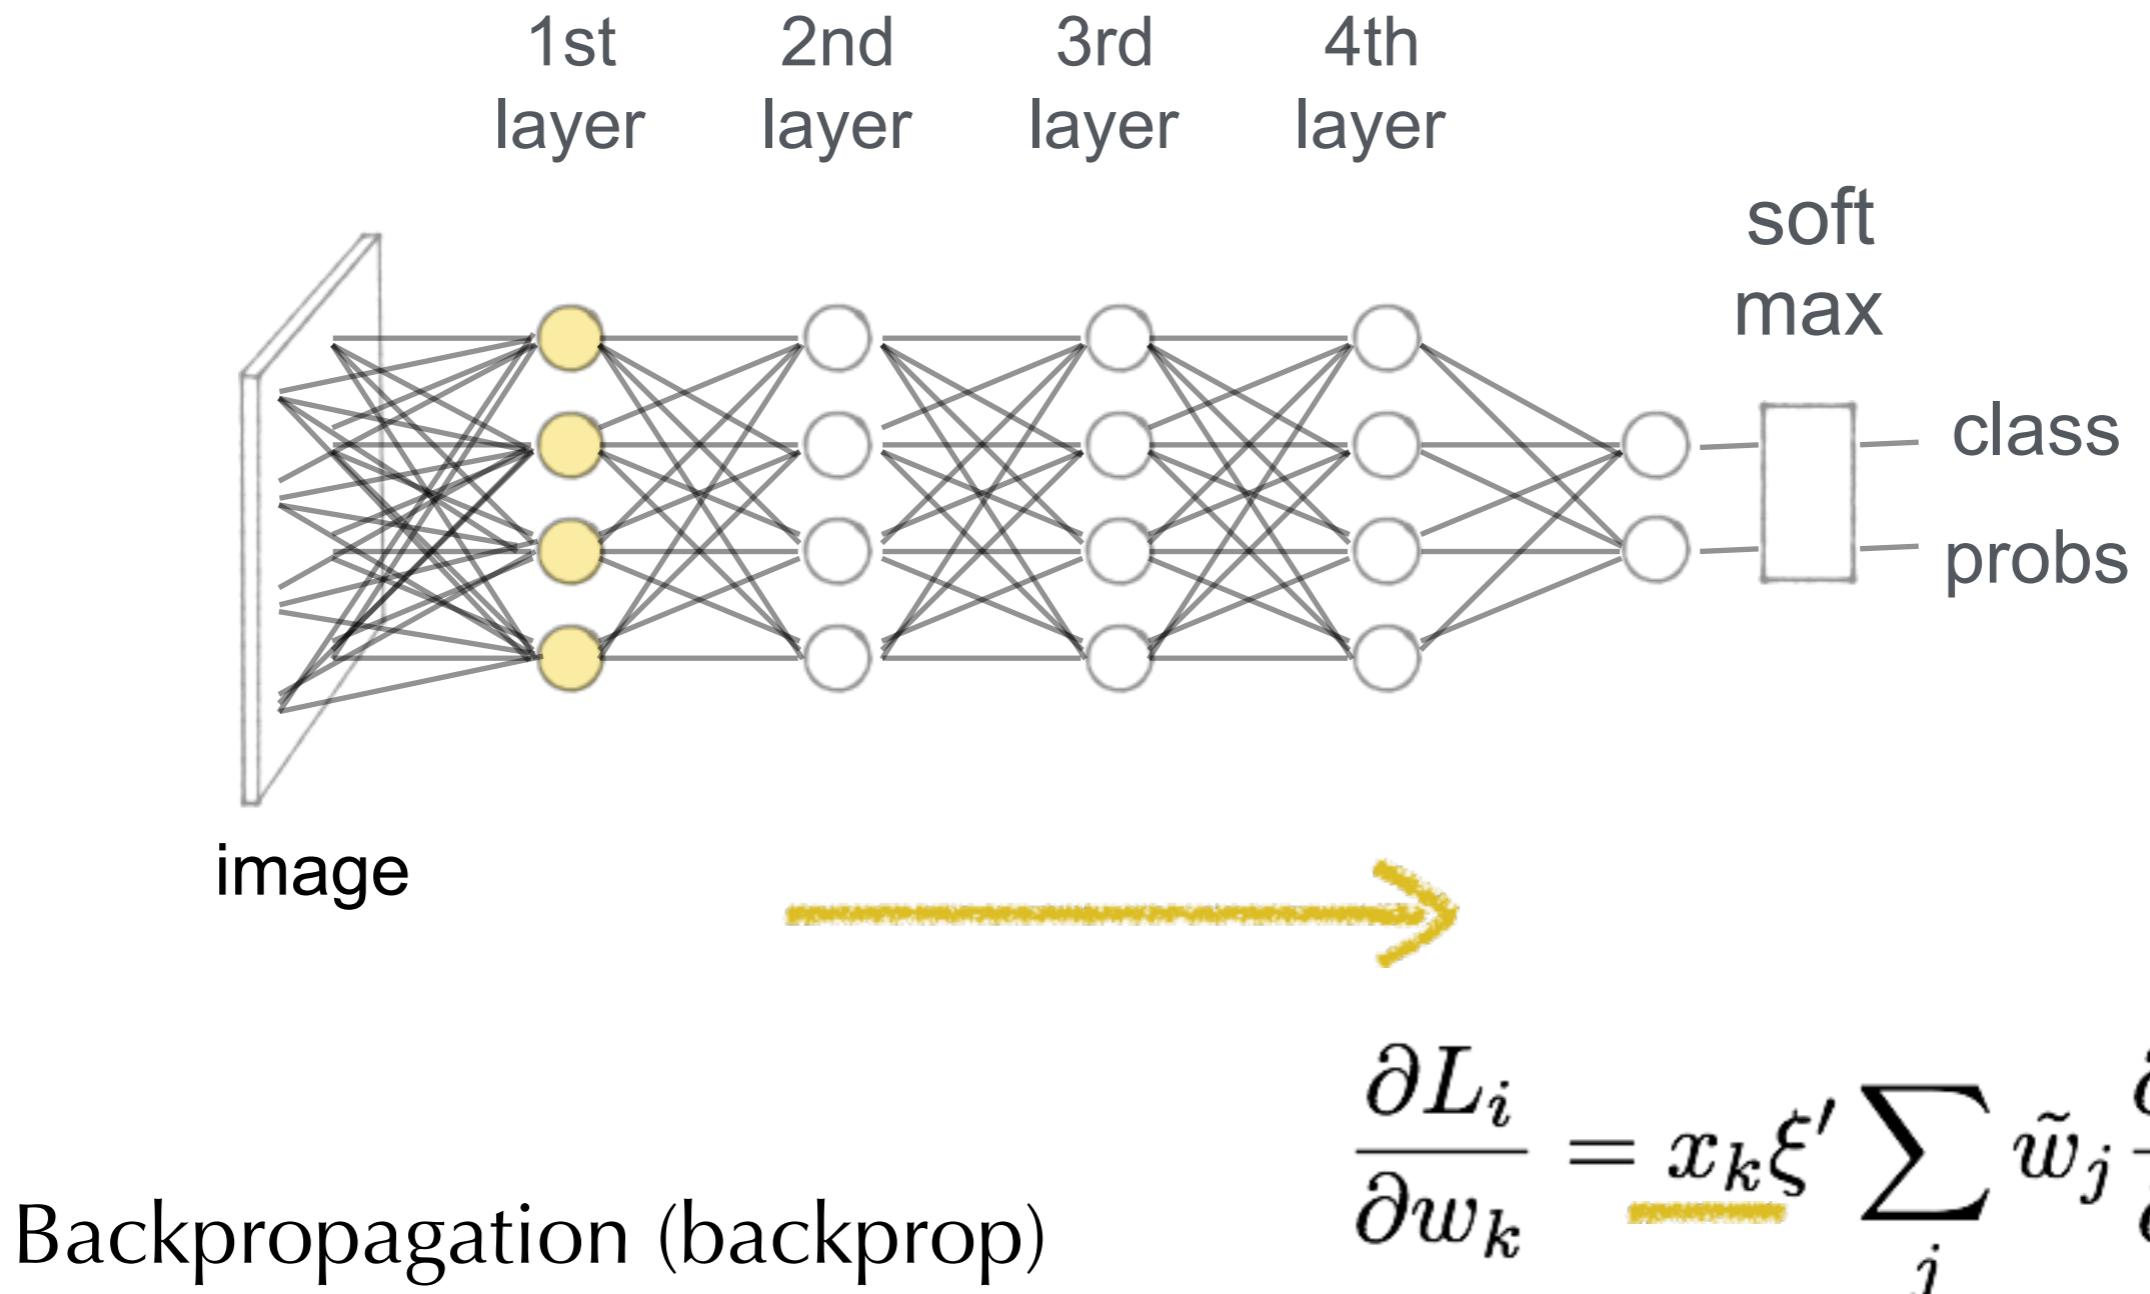
Last Lecture



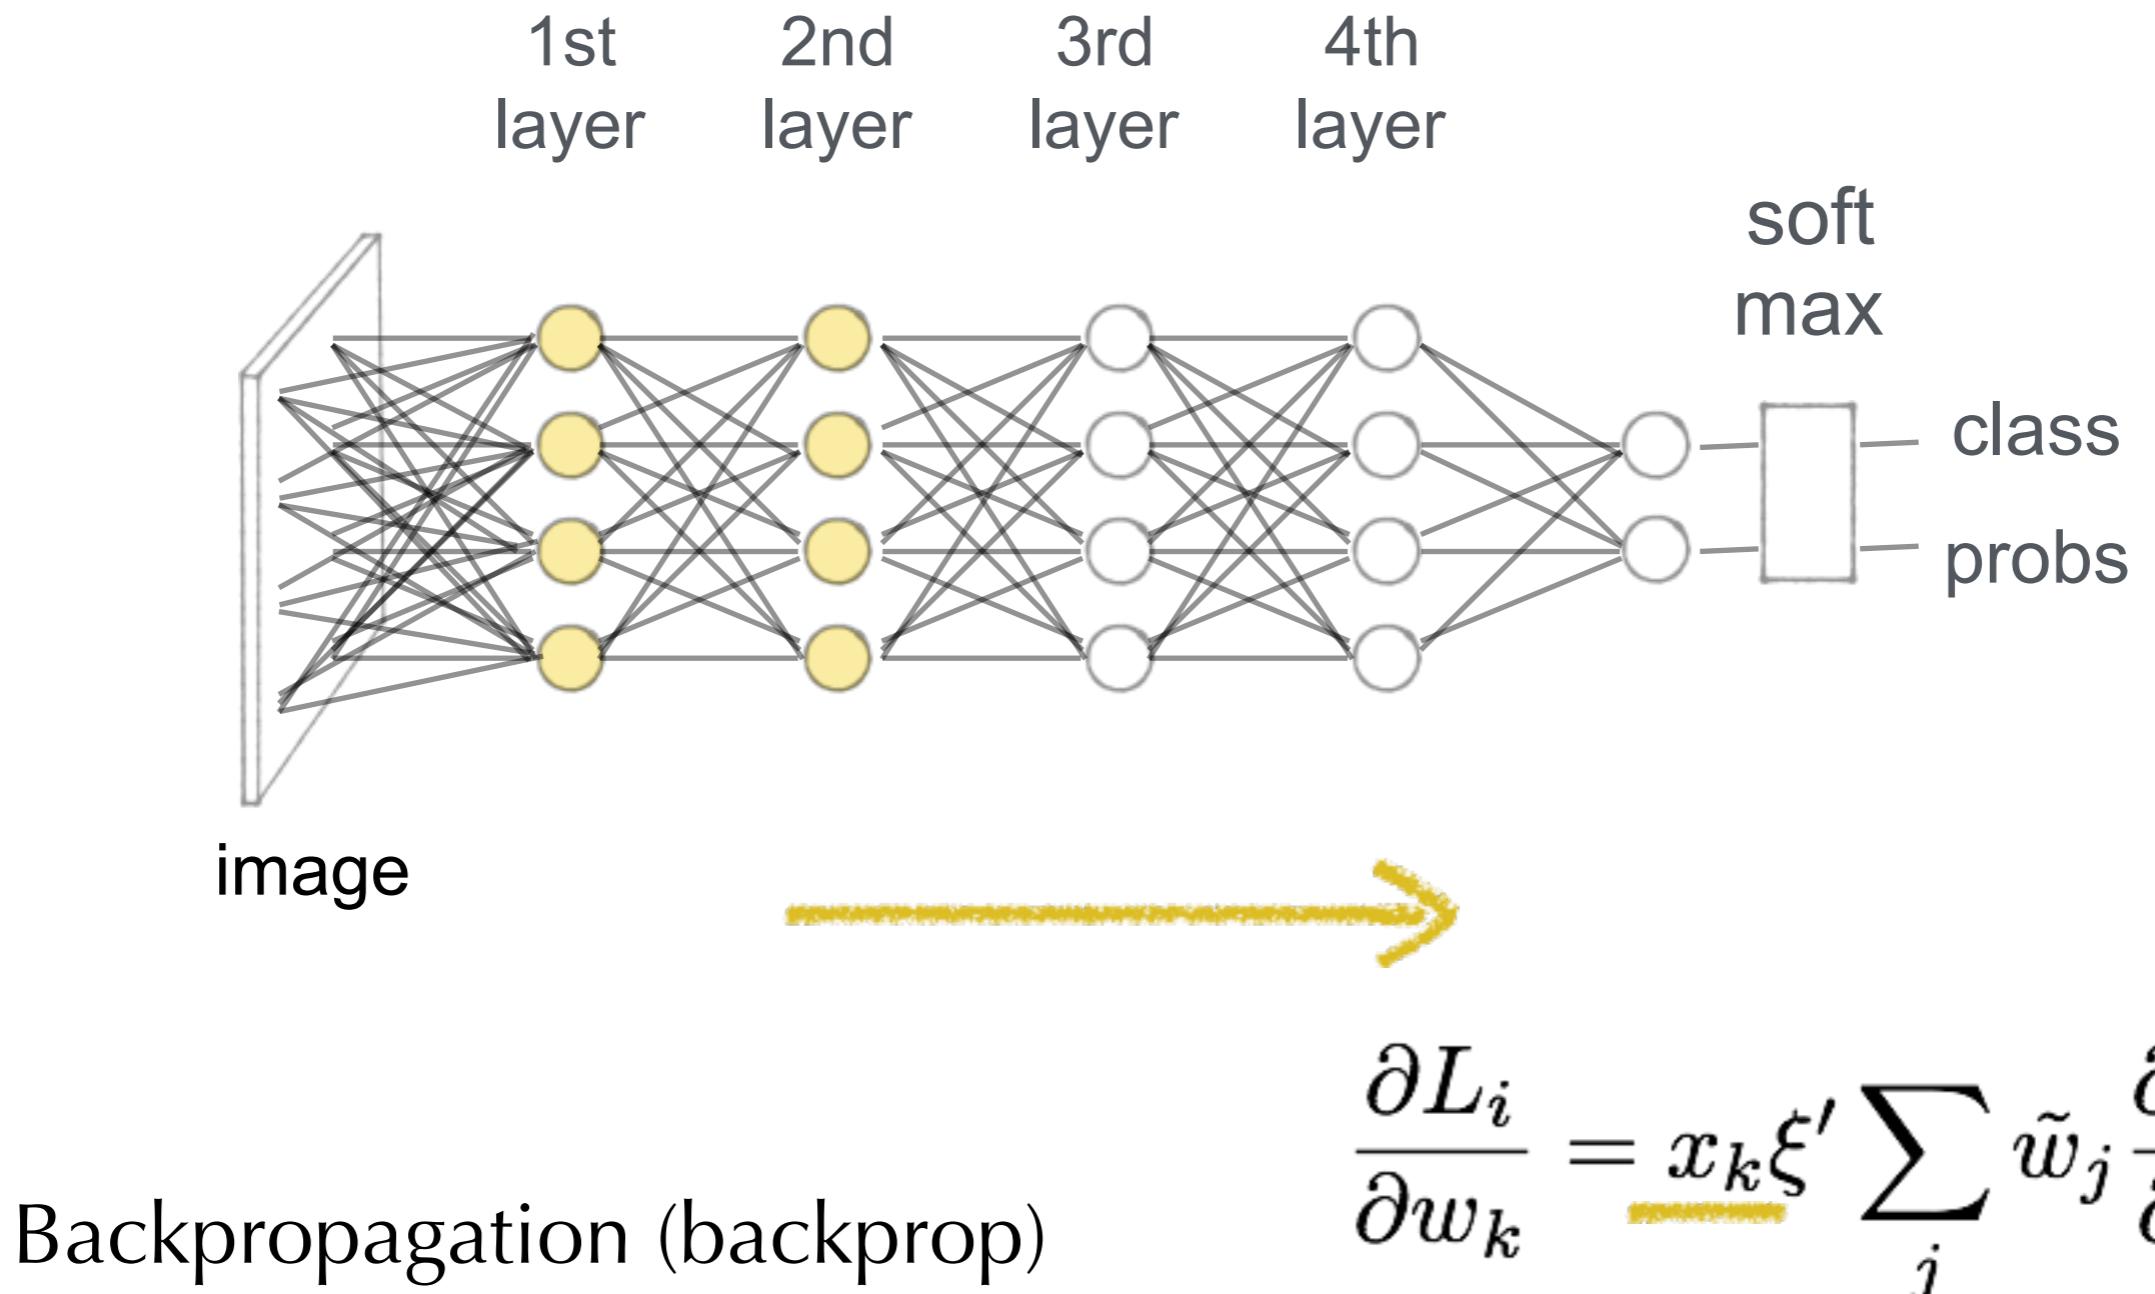
$$\frac{\partial L_i}{\partial w_k} = x_k \xi' \sum_j \tilde{w}_j \frac{\partial L_i}{\partial \tilde{y}_j}$$

Backpropagation (backprop)

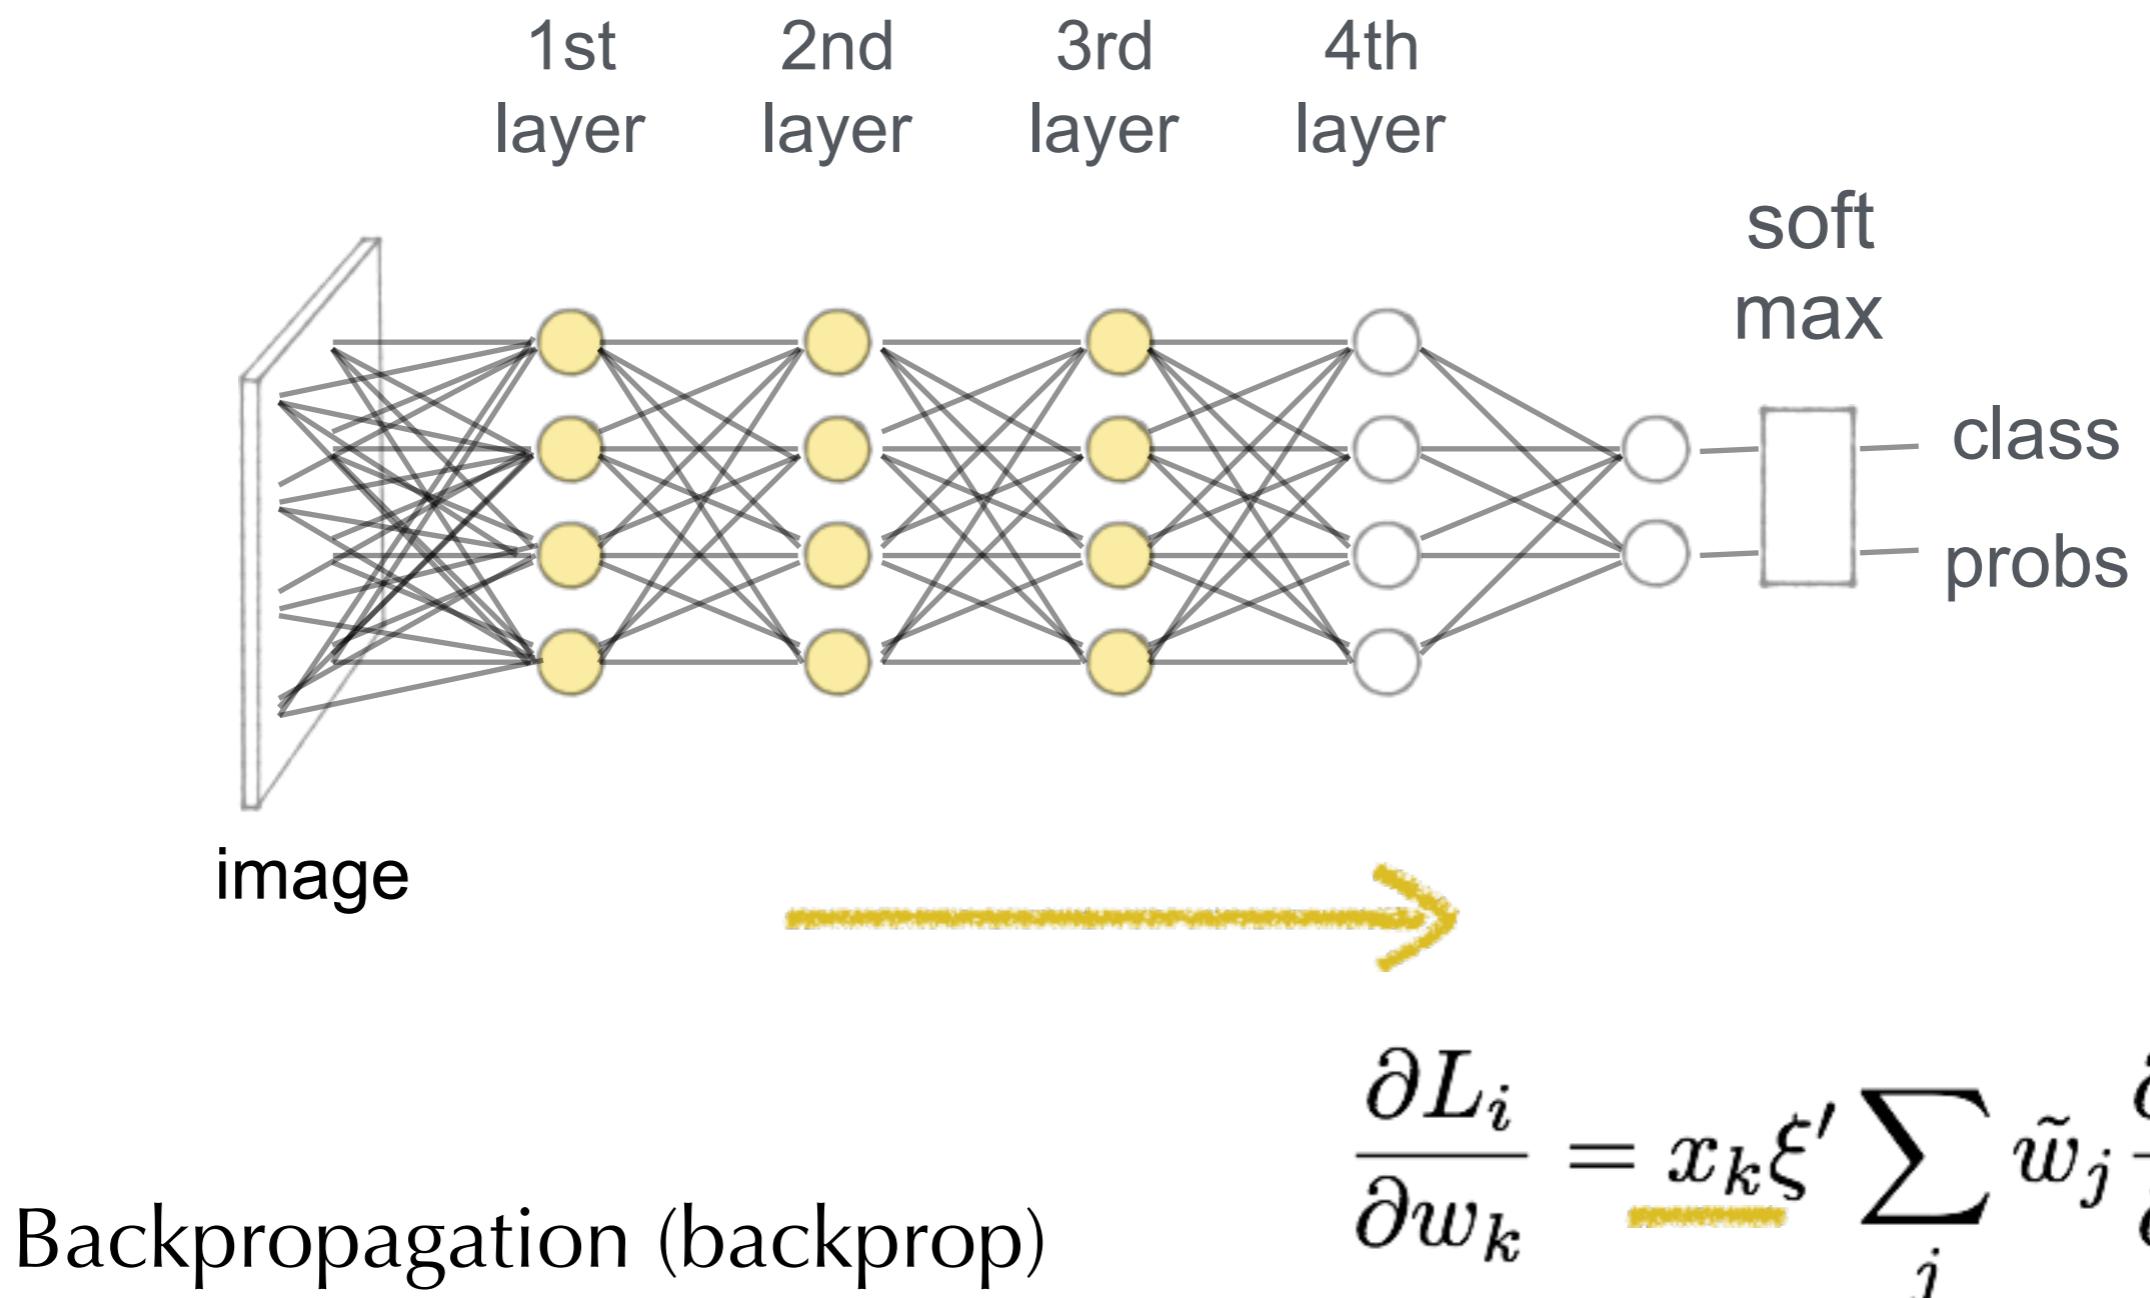
Last Lecture



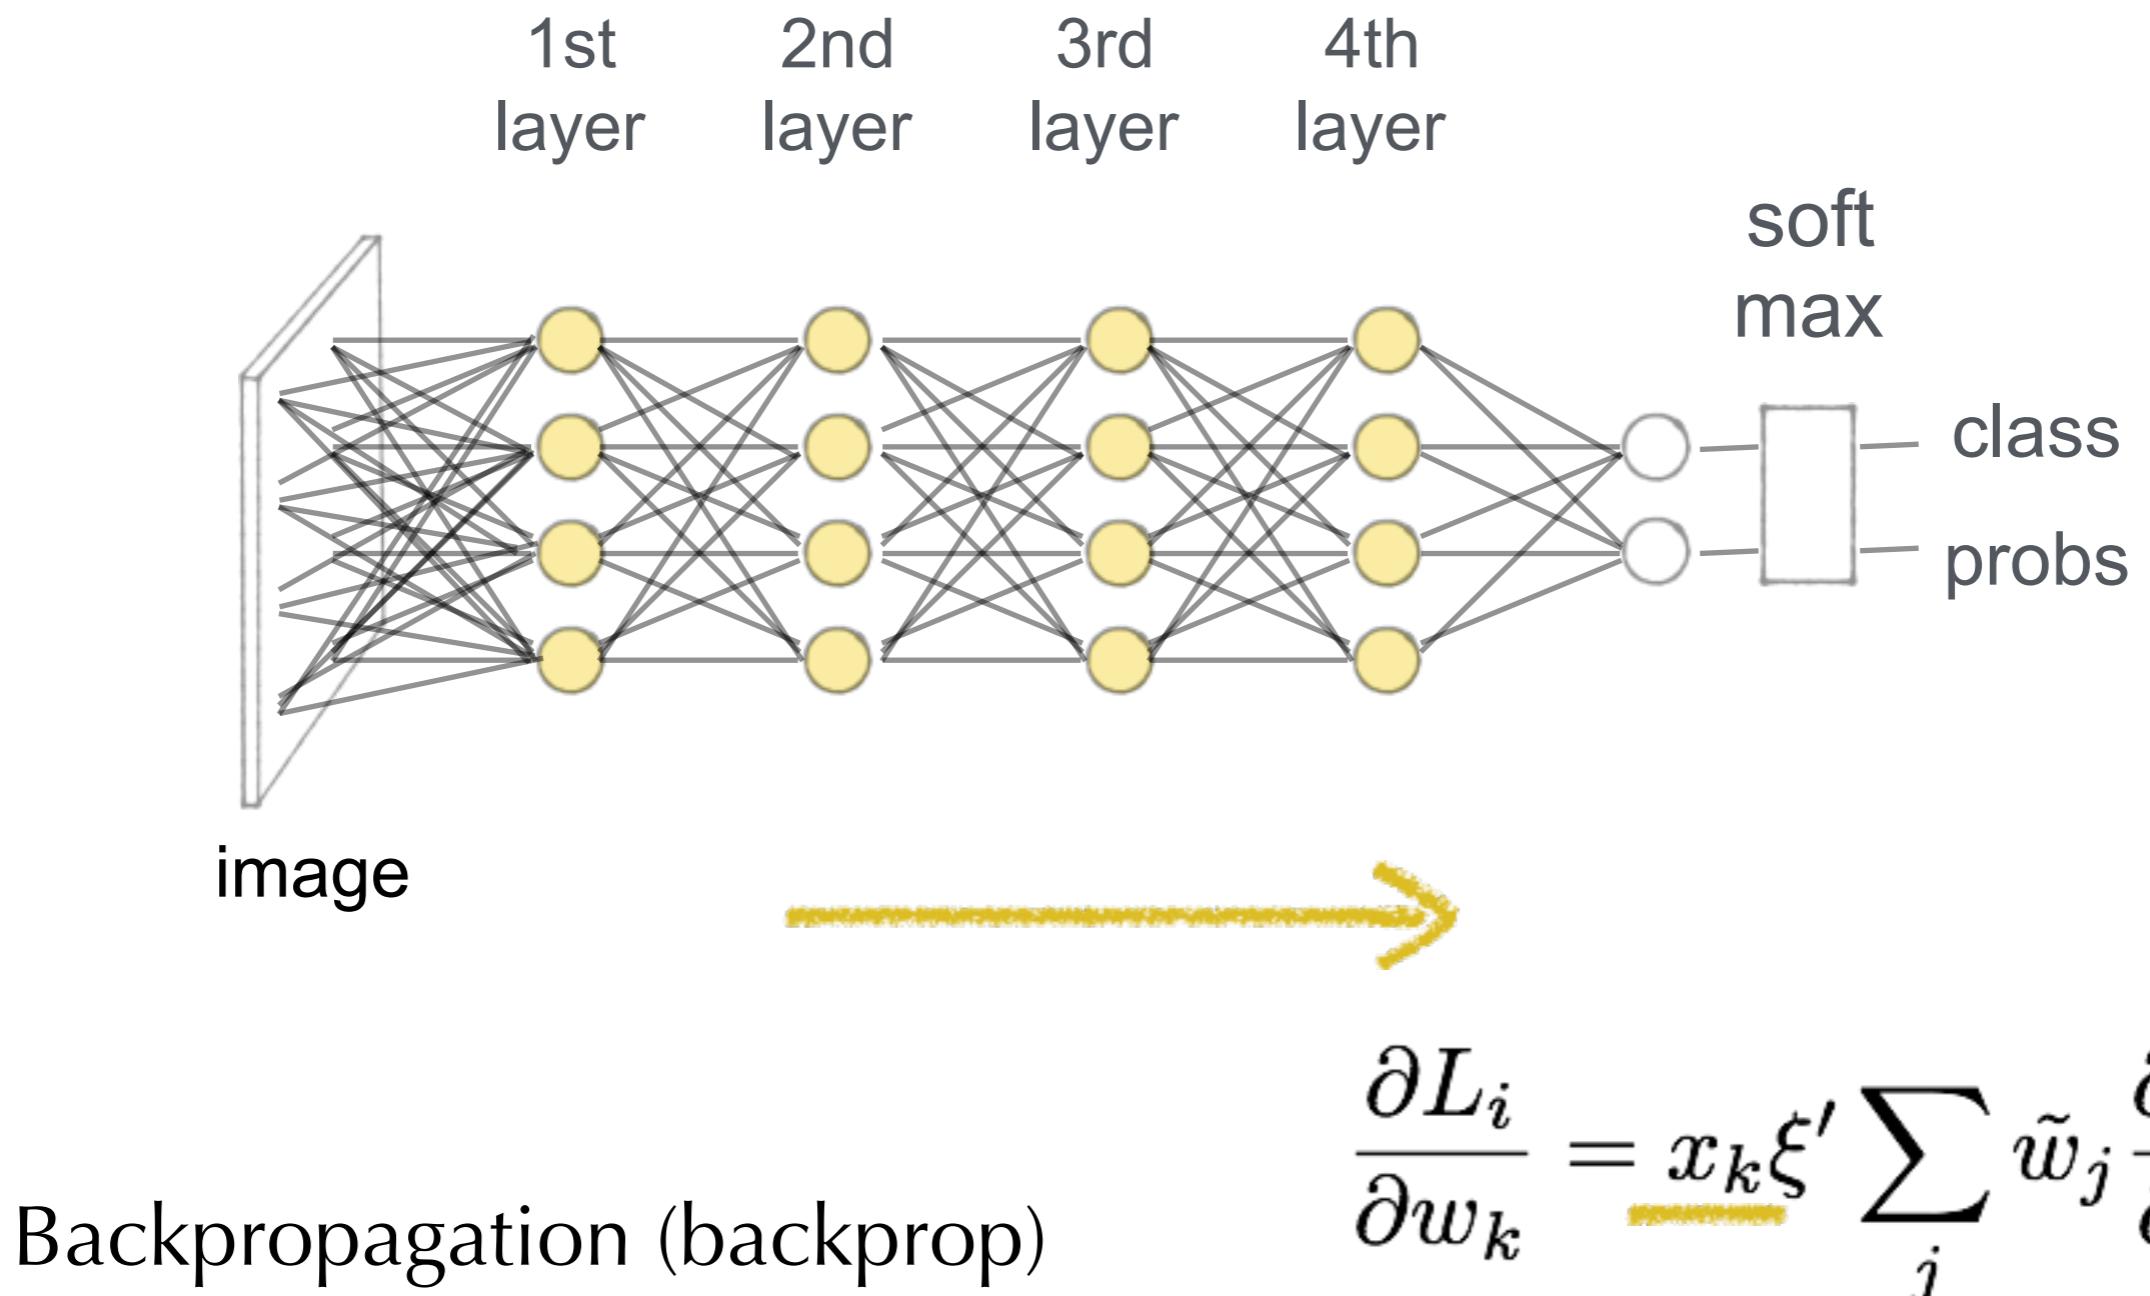
Last Lecture



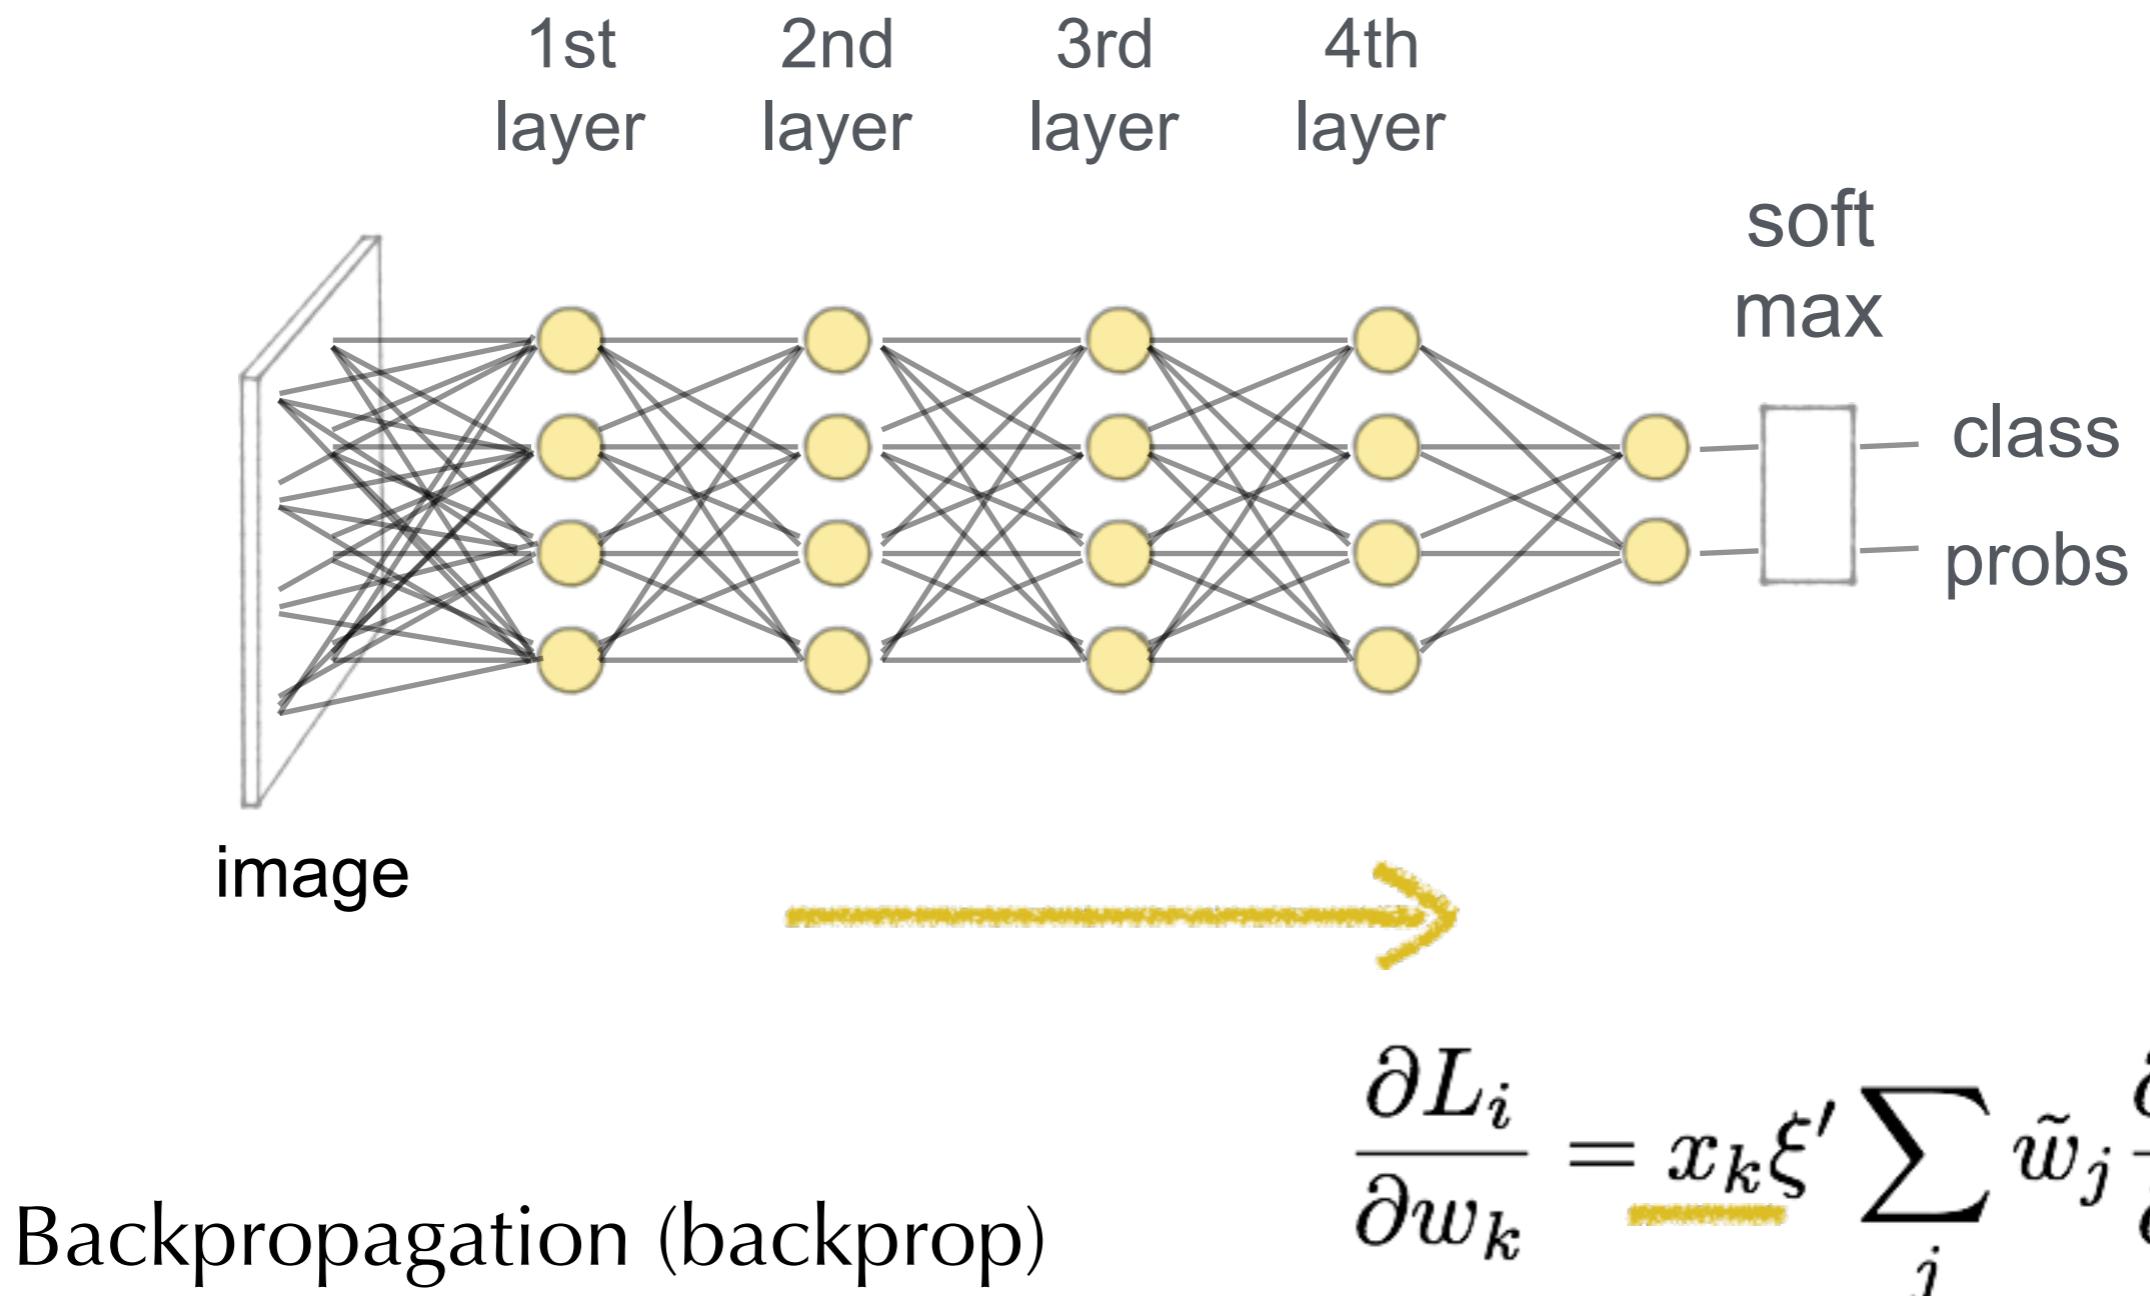
Last Lecture



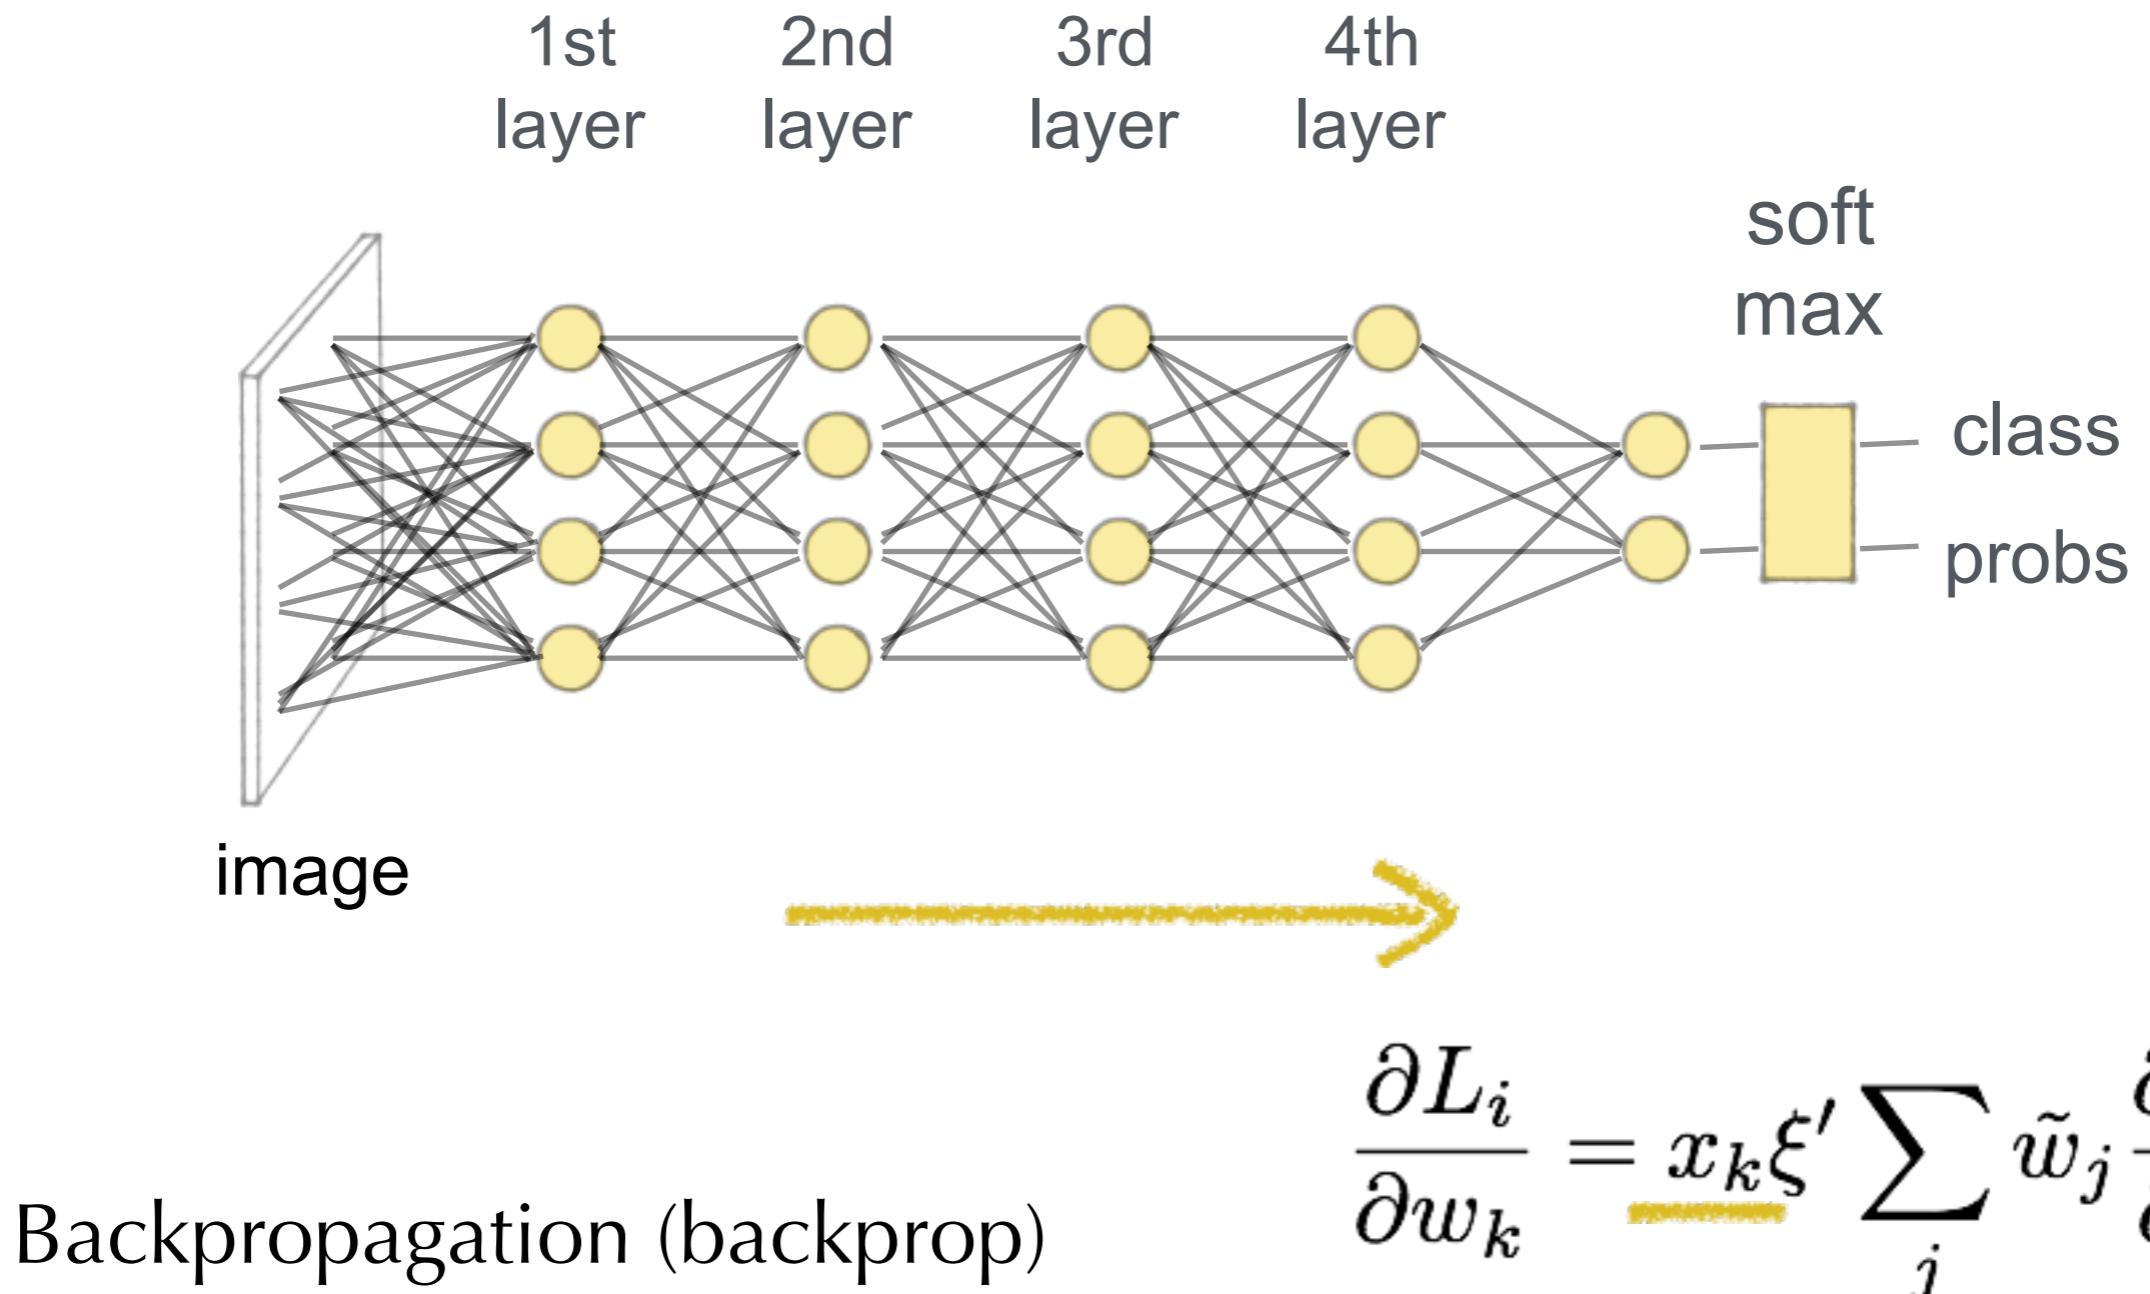
Last Lecture



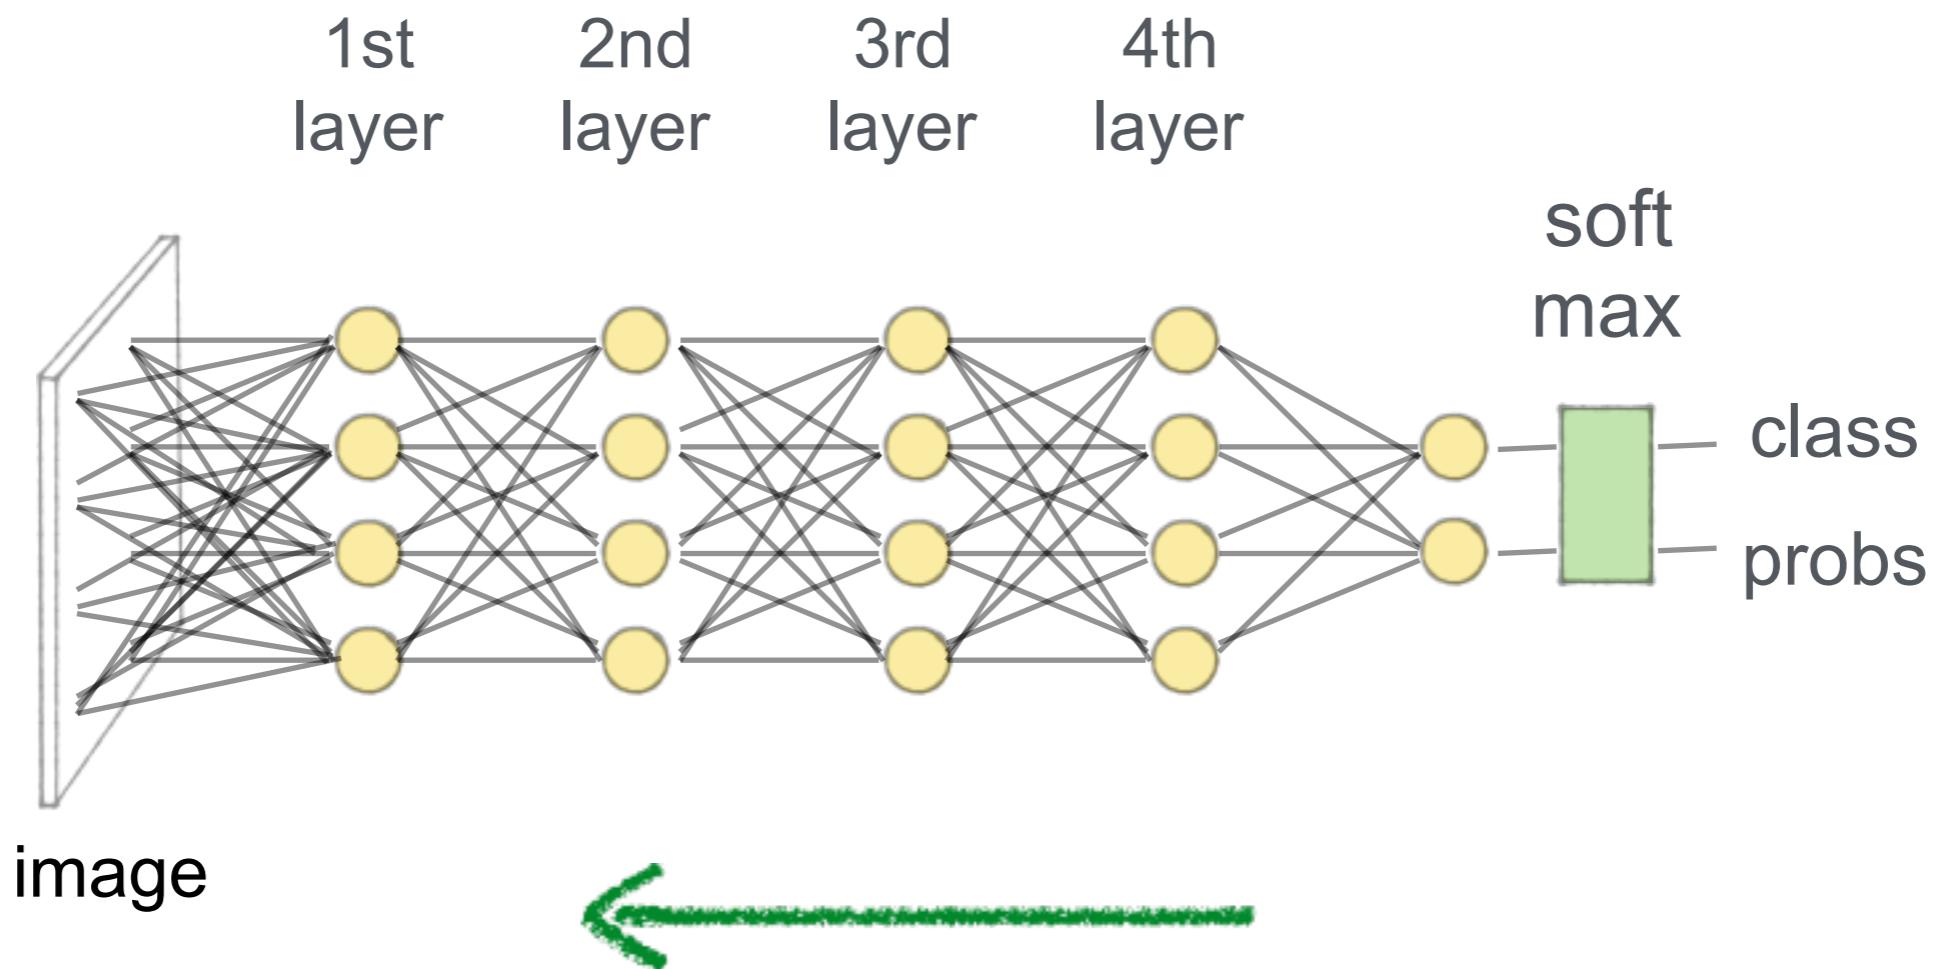
Last Lecture



Last Lecture



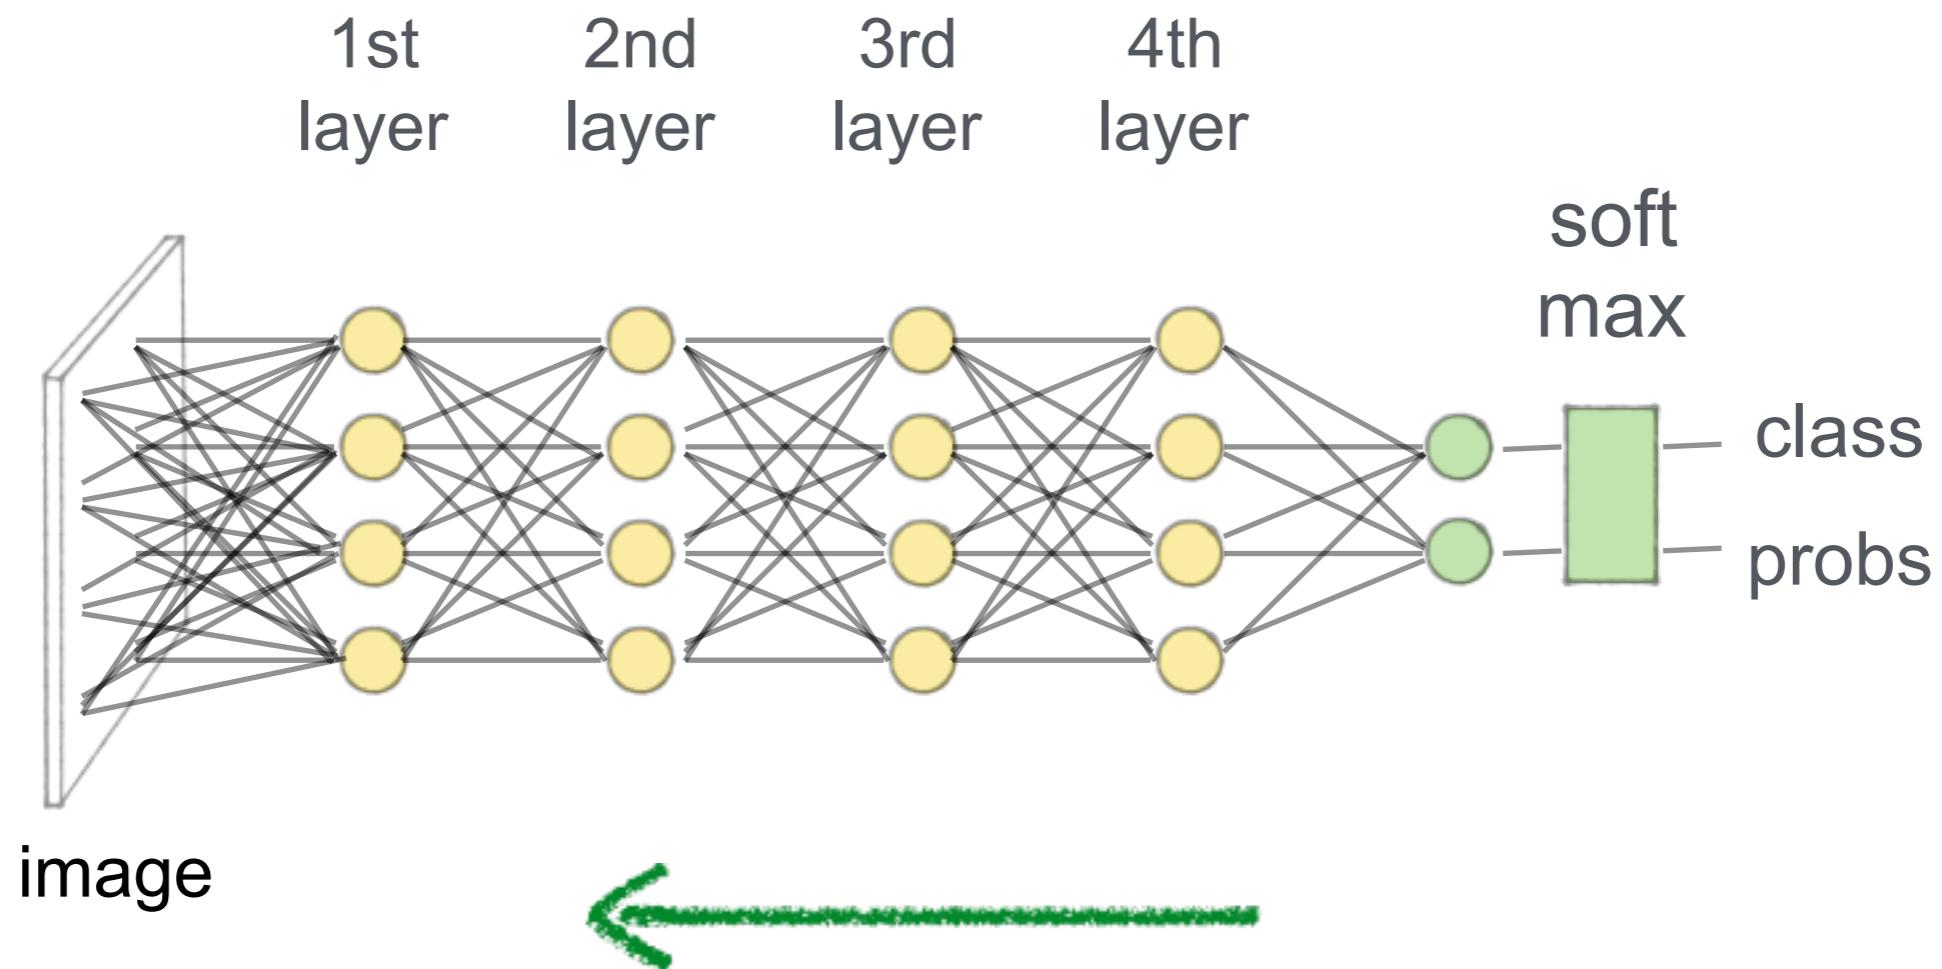
Last Lecture



Backpropagation (backprop)

$$\frac{\partial L_i}{\partial w_k} = \underline{x_k} \xi' \sum_j \tilde{w}_j \frac{\partial L_i}{\partial \tilde{y}_j}$$

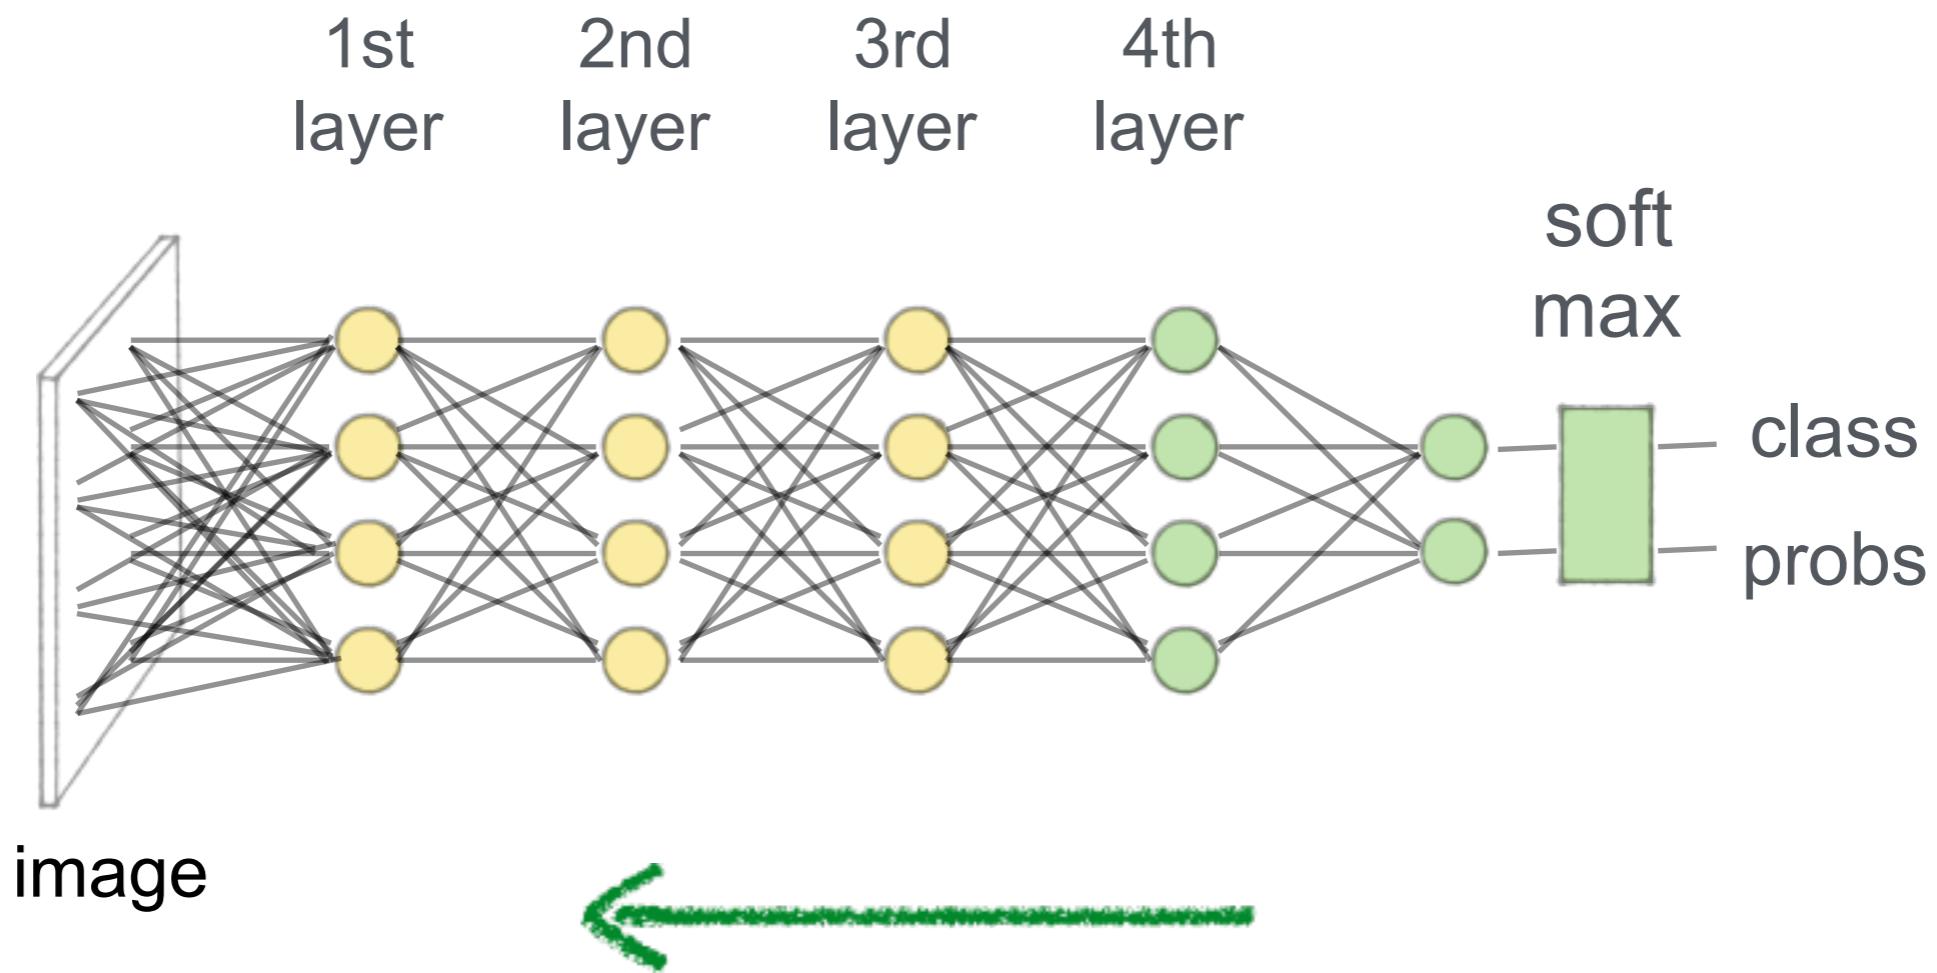
Last Lecture



Backpropagation (backprop)

$$\frac{\partial L_i}{\partial w_k} = \underline{x_k \xi'} \sum_j \tilde{w}_j \frac{\partial L_i}{\partial \tilde{y}_j}$$

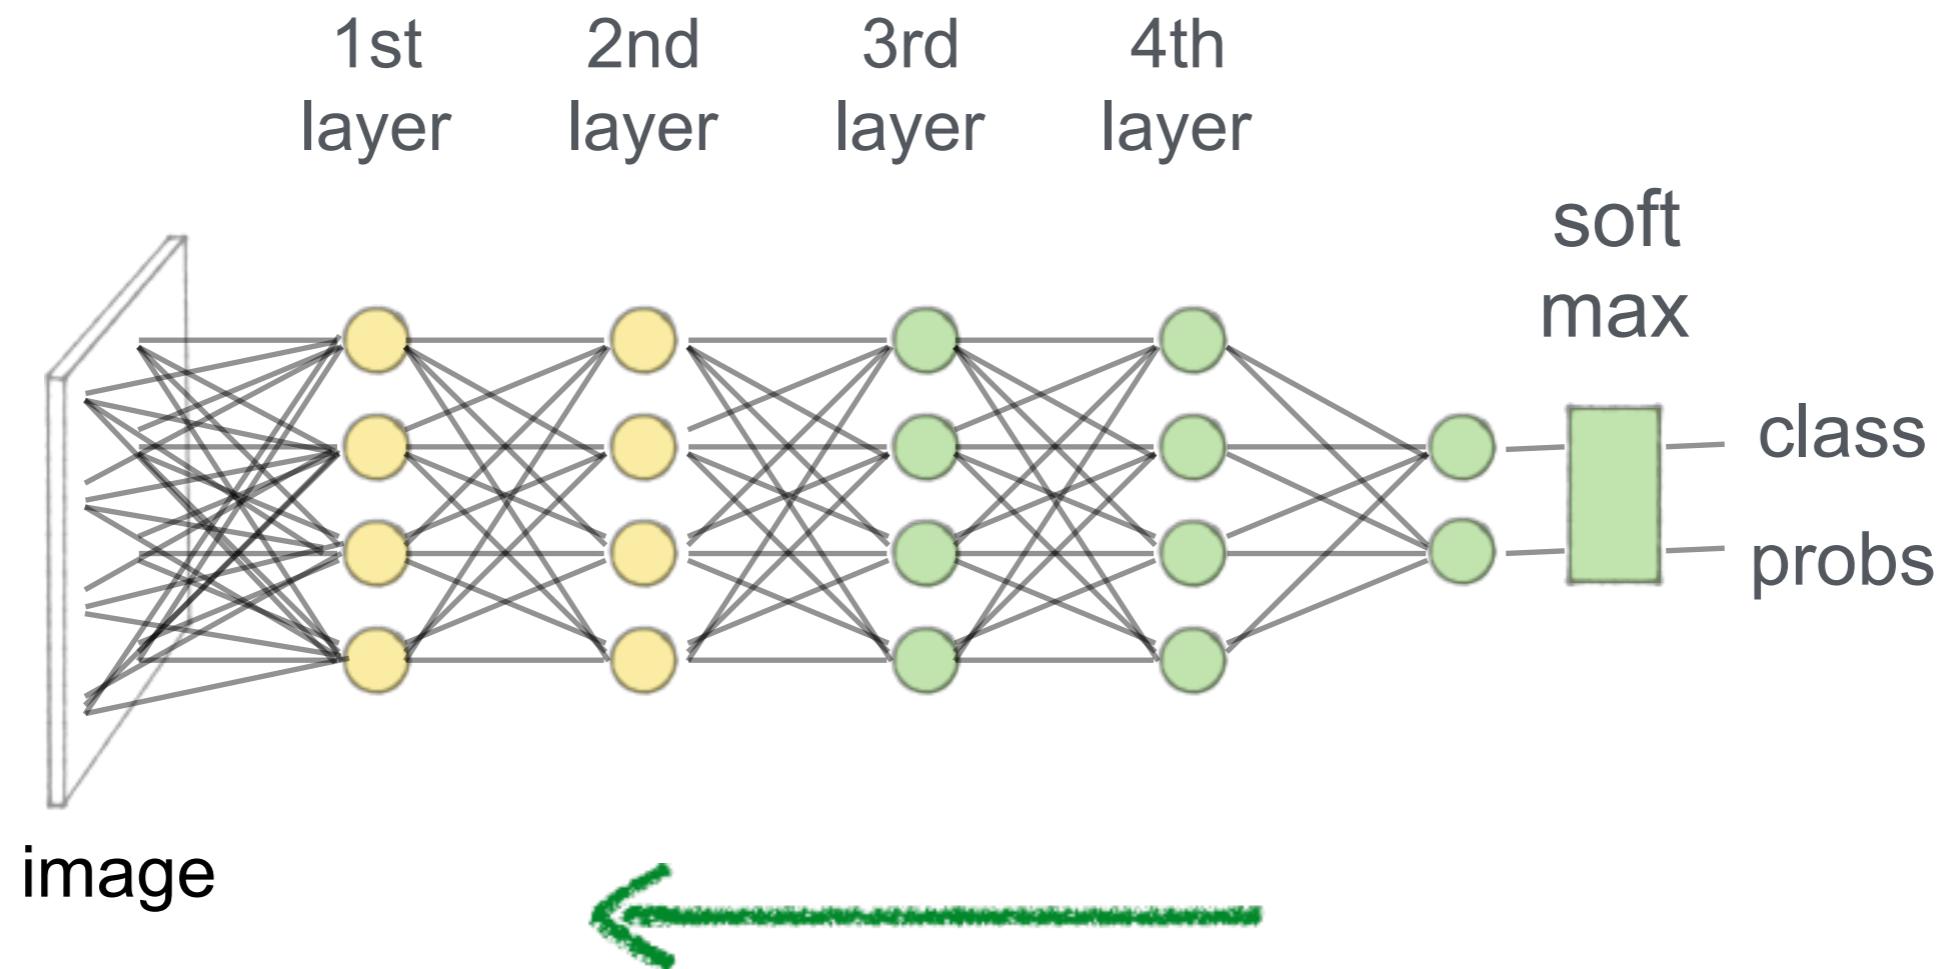
Last Lecture



Backpropagation (backprop)

$$\frac{\partial L_i}{\partial w_k} = \underline{x_k \xi'} \sum_j \tilde{w}_j \frac{\partial L_i}{\partial \tilde{y}_j}$$

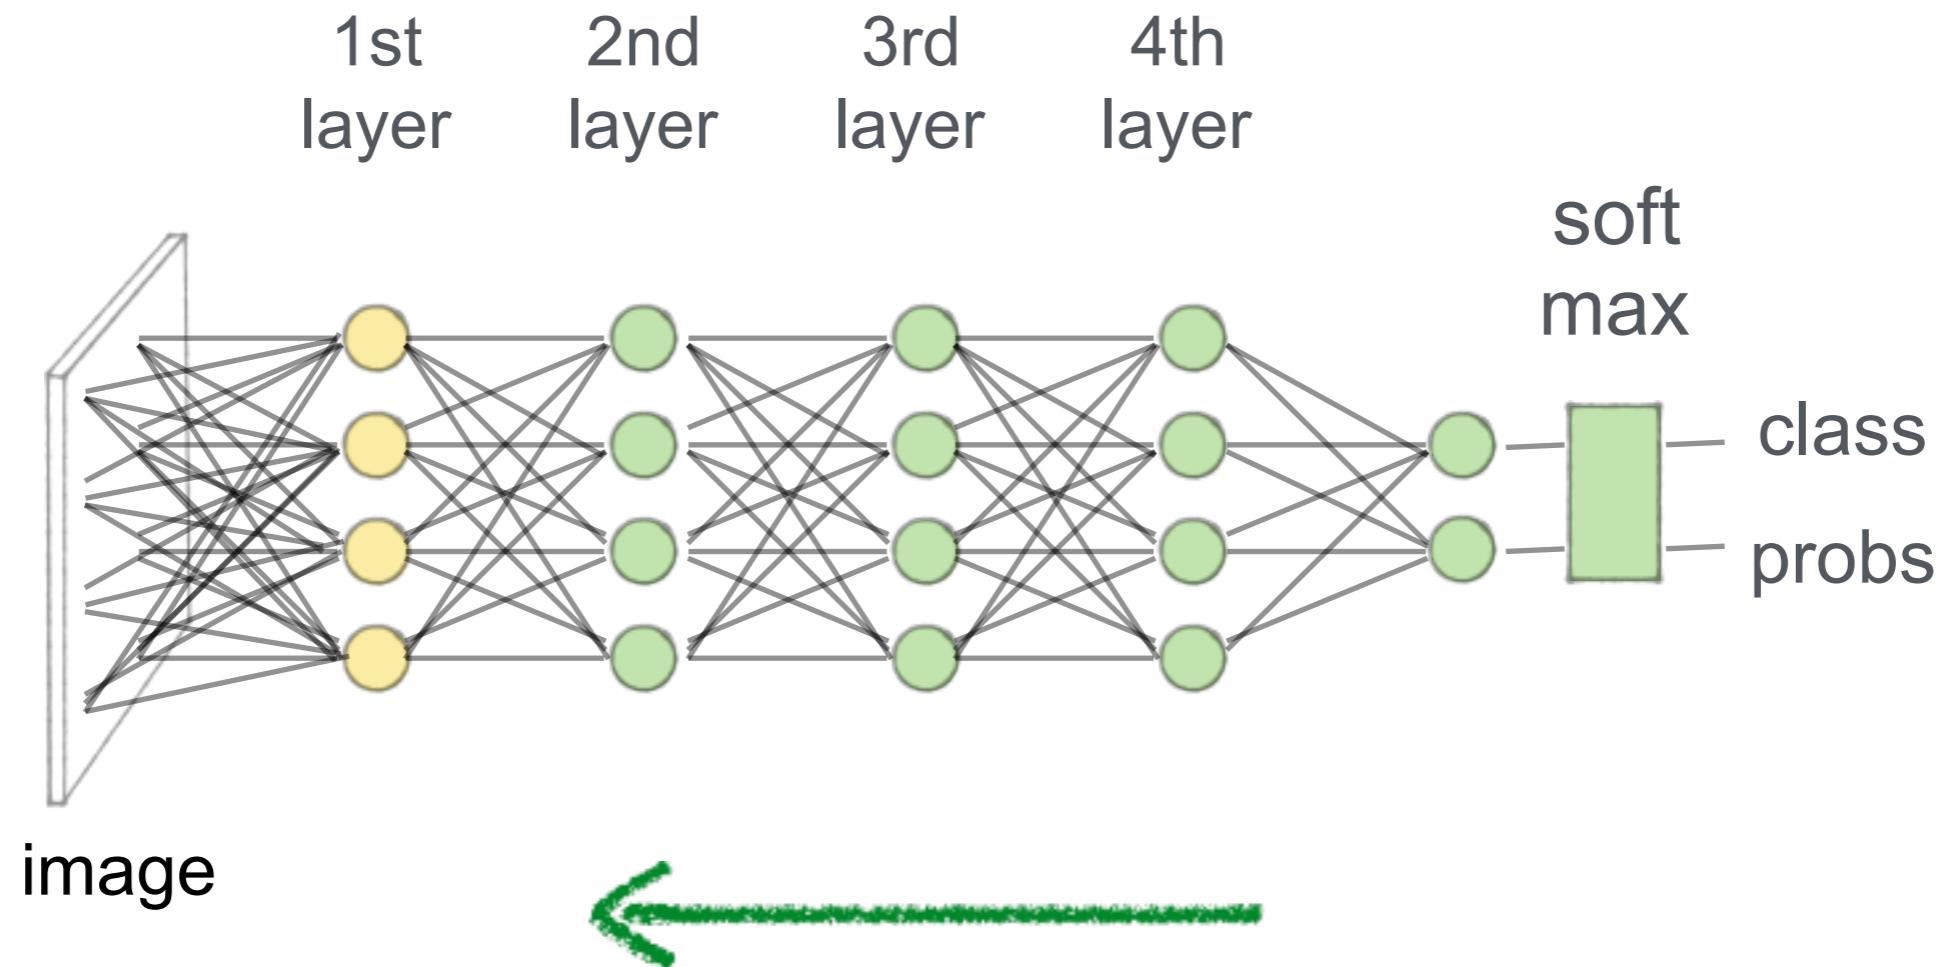
Last Lecture



Backpropagation (backprop)

$$\frac{\partial L_i}{\partial w_k} = \underline{x_k \xi'} \sum_j \tilde{w}_j \frac{\partial L_i}{\partial \tilde{y}_j}$$

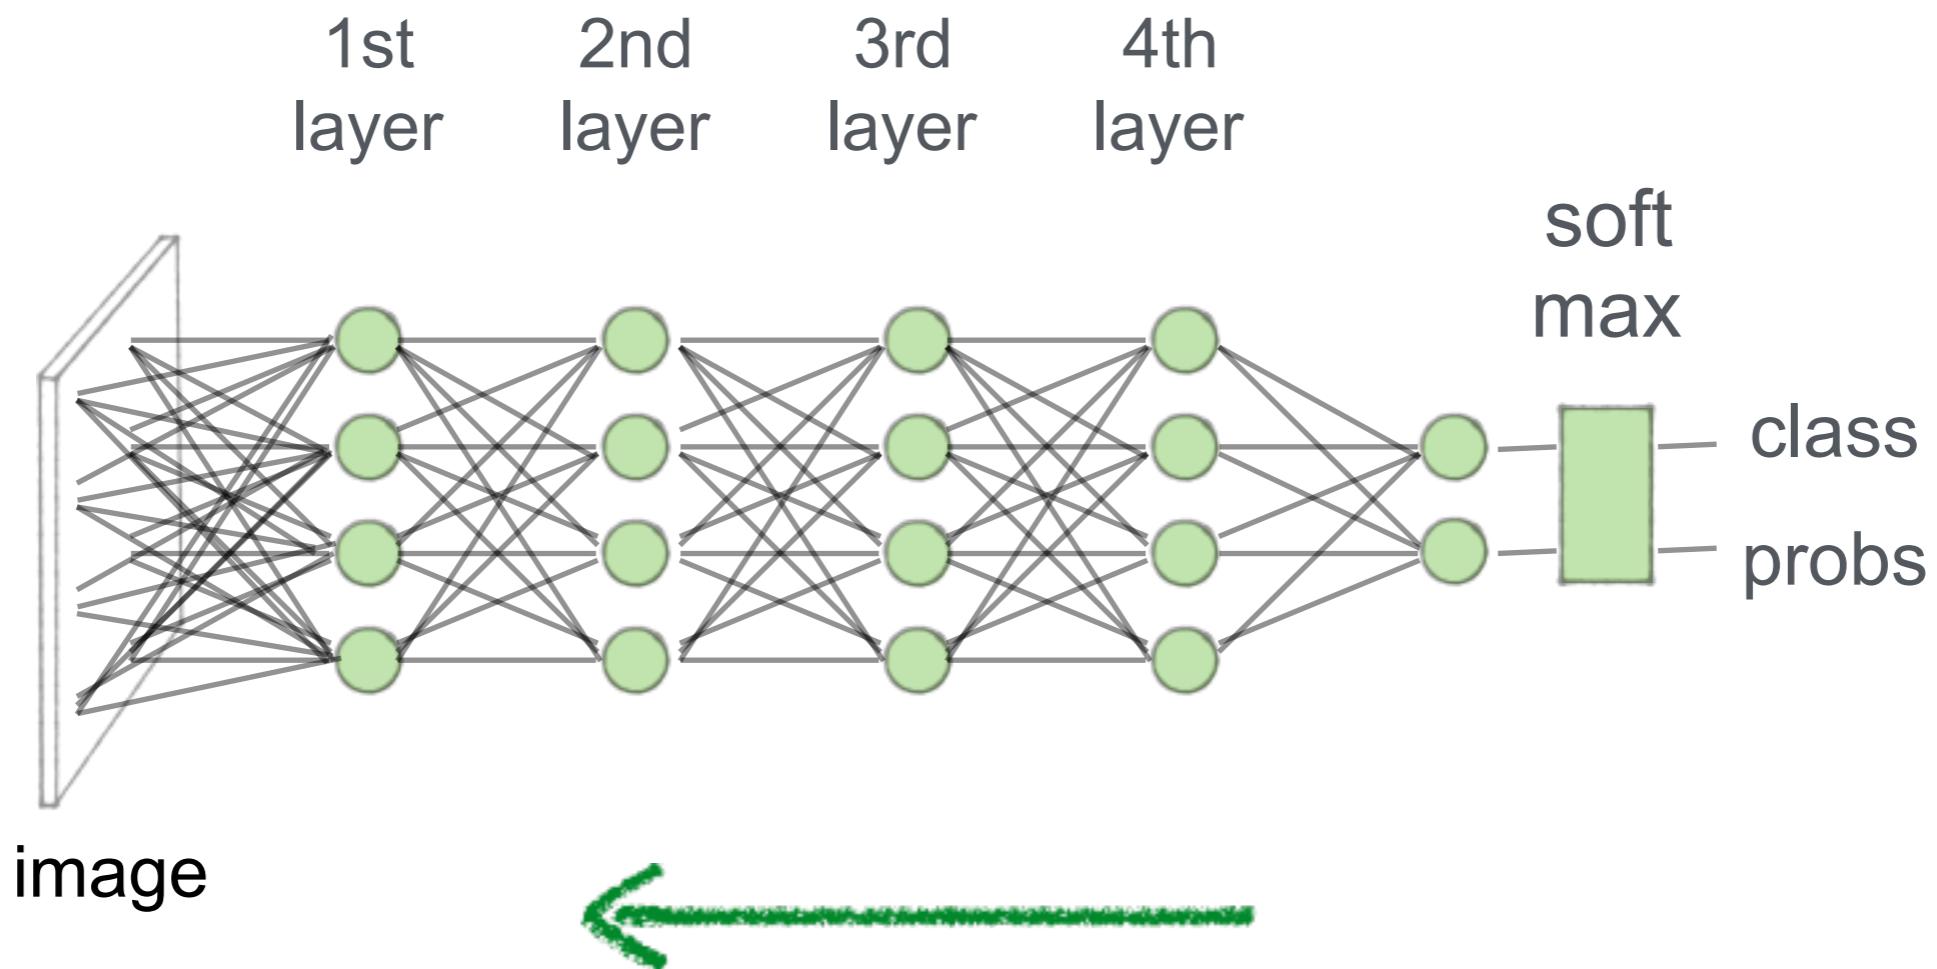
Last Lecture



Backpropagation (backprop)

$$\frac{\partial L_i}{\partial w_k} = \underline{x_k} \xi' \sum_j \tilde{w}_j \frac{\partial L_i}{\partial \tilde{y}_j}$$

Last Lecture



Backpropagation (backprop)

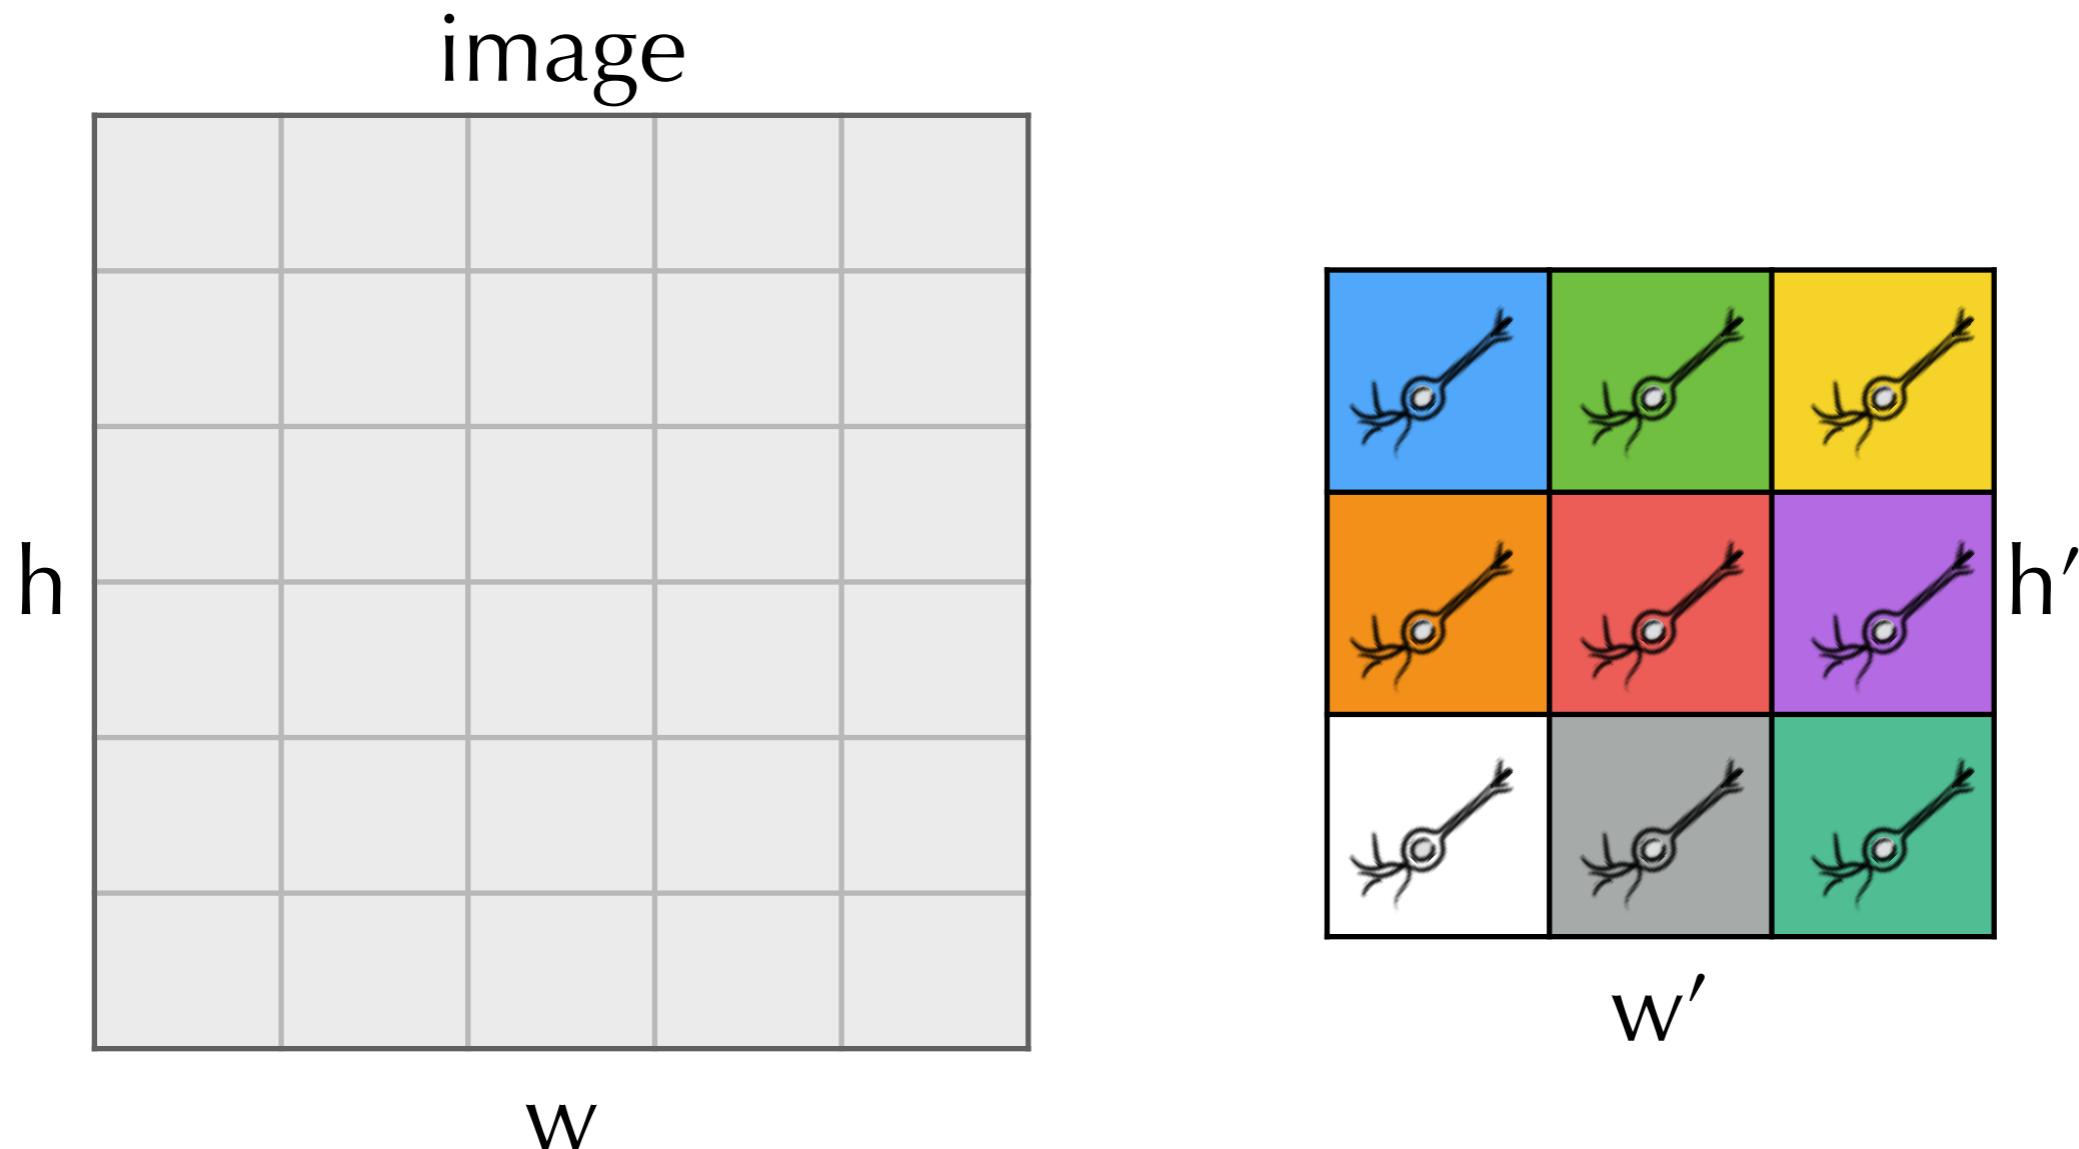
$$\frac{\partial L_i}{\partial w_k} = \underline{x_k \xi'} \sum_j \tilde{w}_j \frac{\partial L_i}{\partial \tilde{y}_j}$$

Today

- Convolutional Neural Networks
- Fully Convolutional Neural Networks
- Convolutional Layers
- Overfitting, Part II

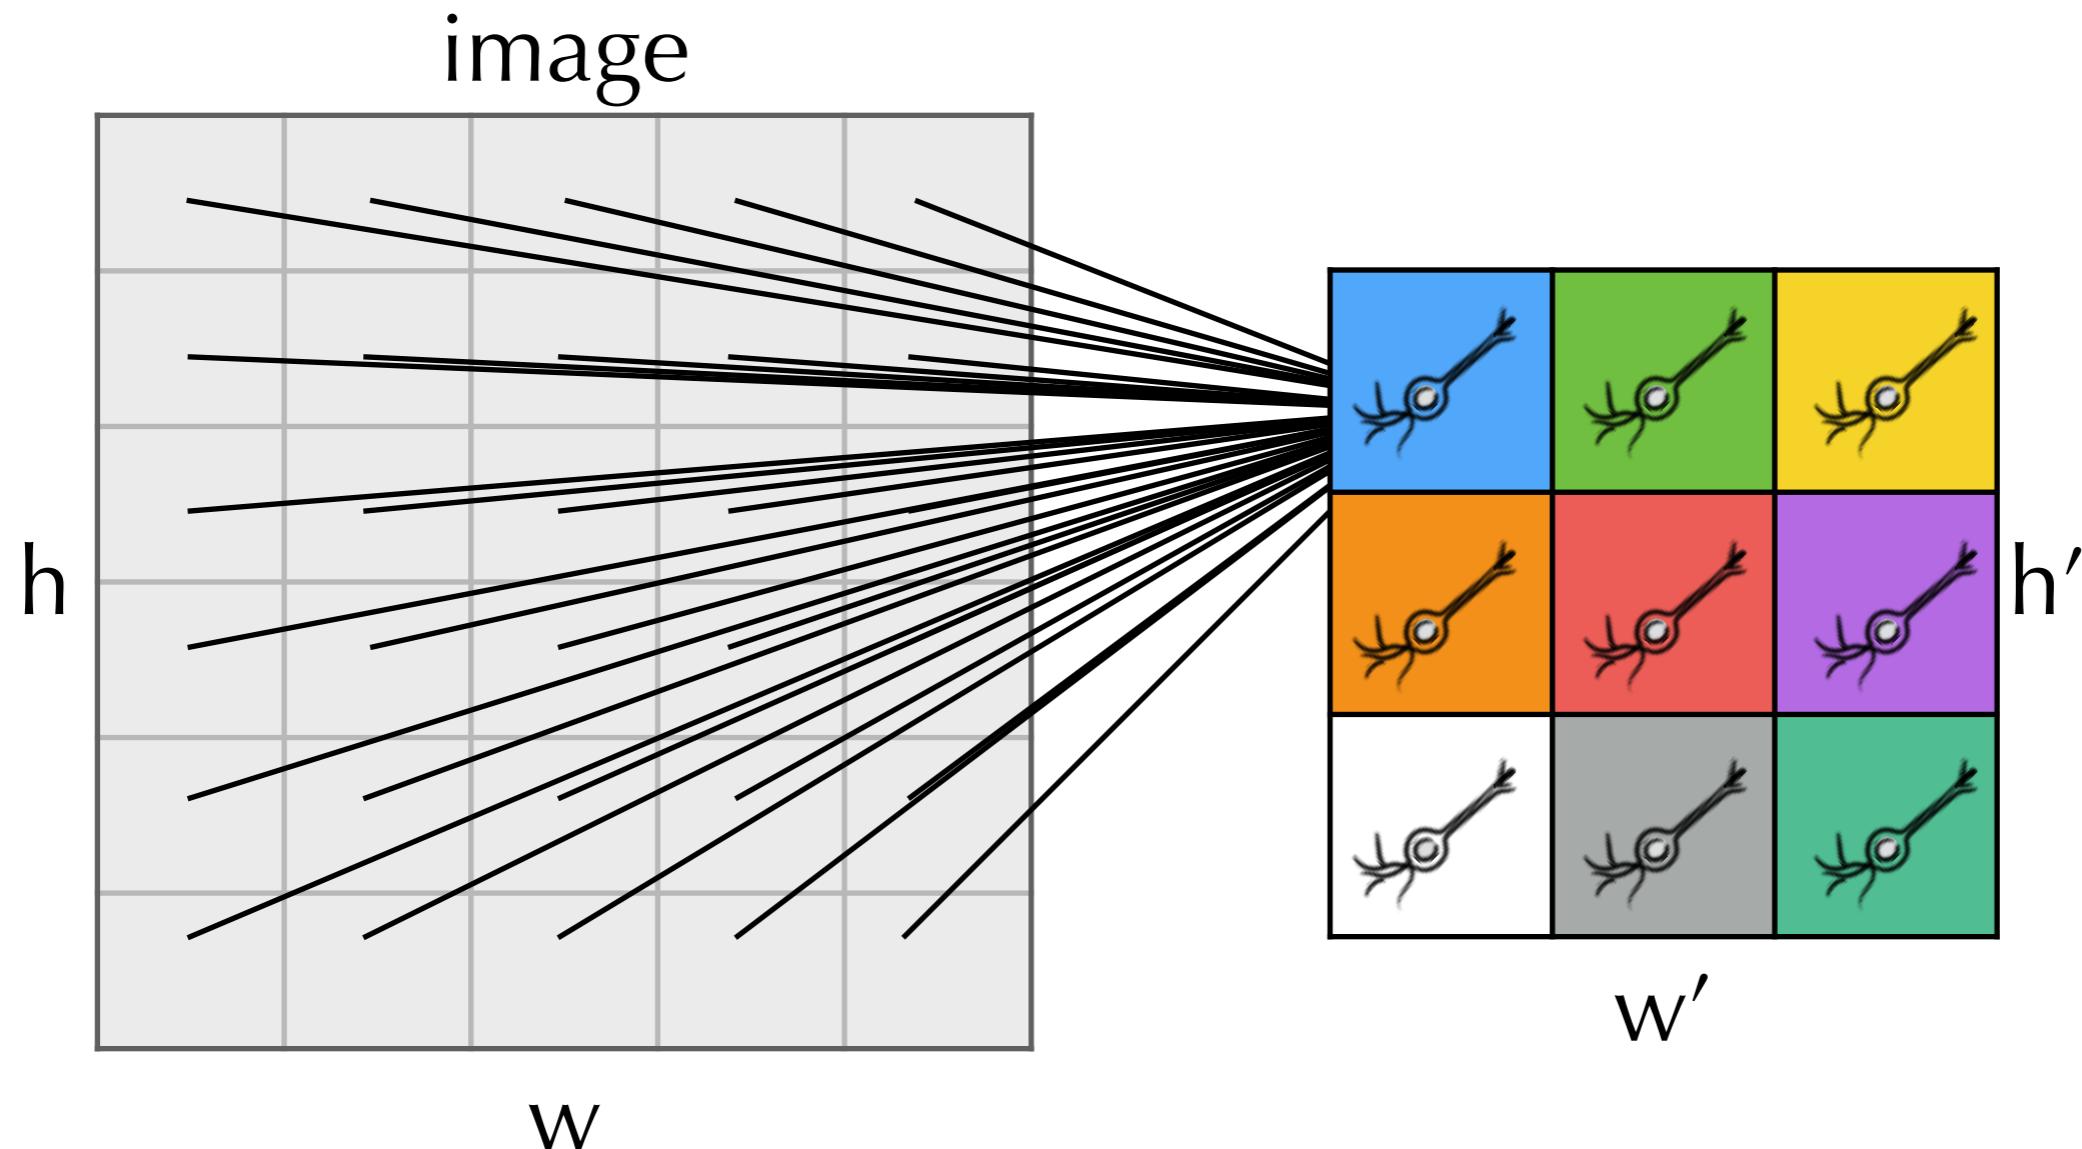
Convolutional Neural Networks

Fully Connected Layers



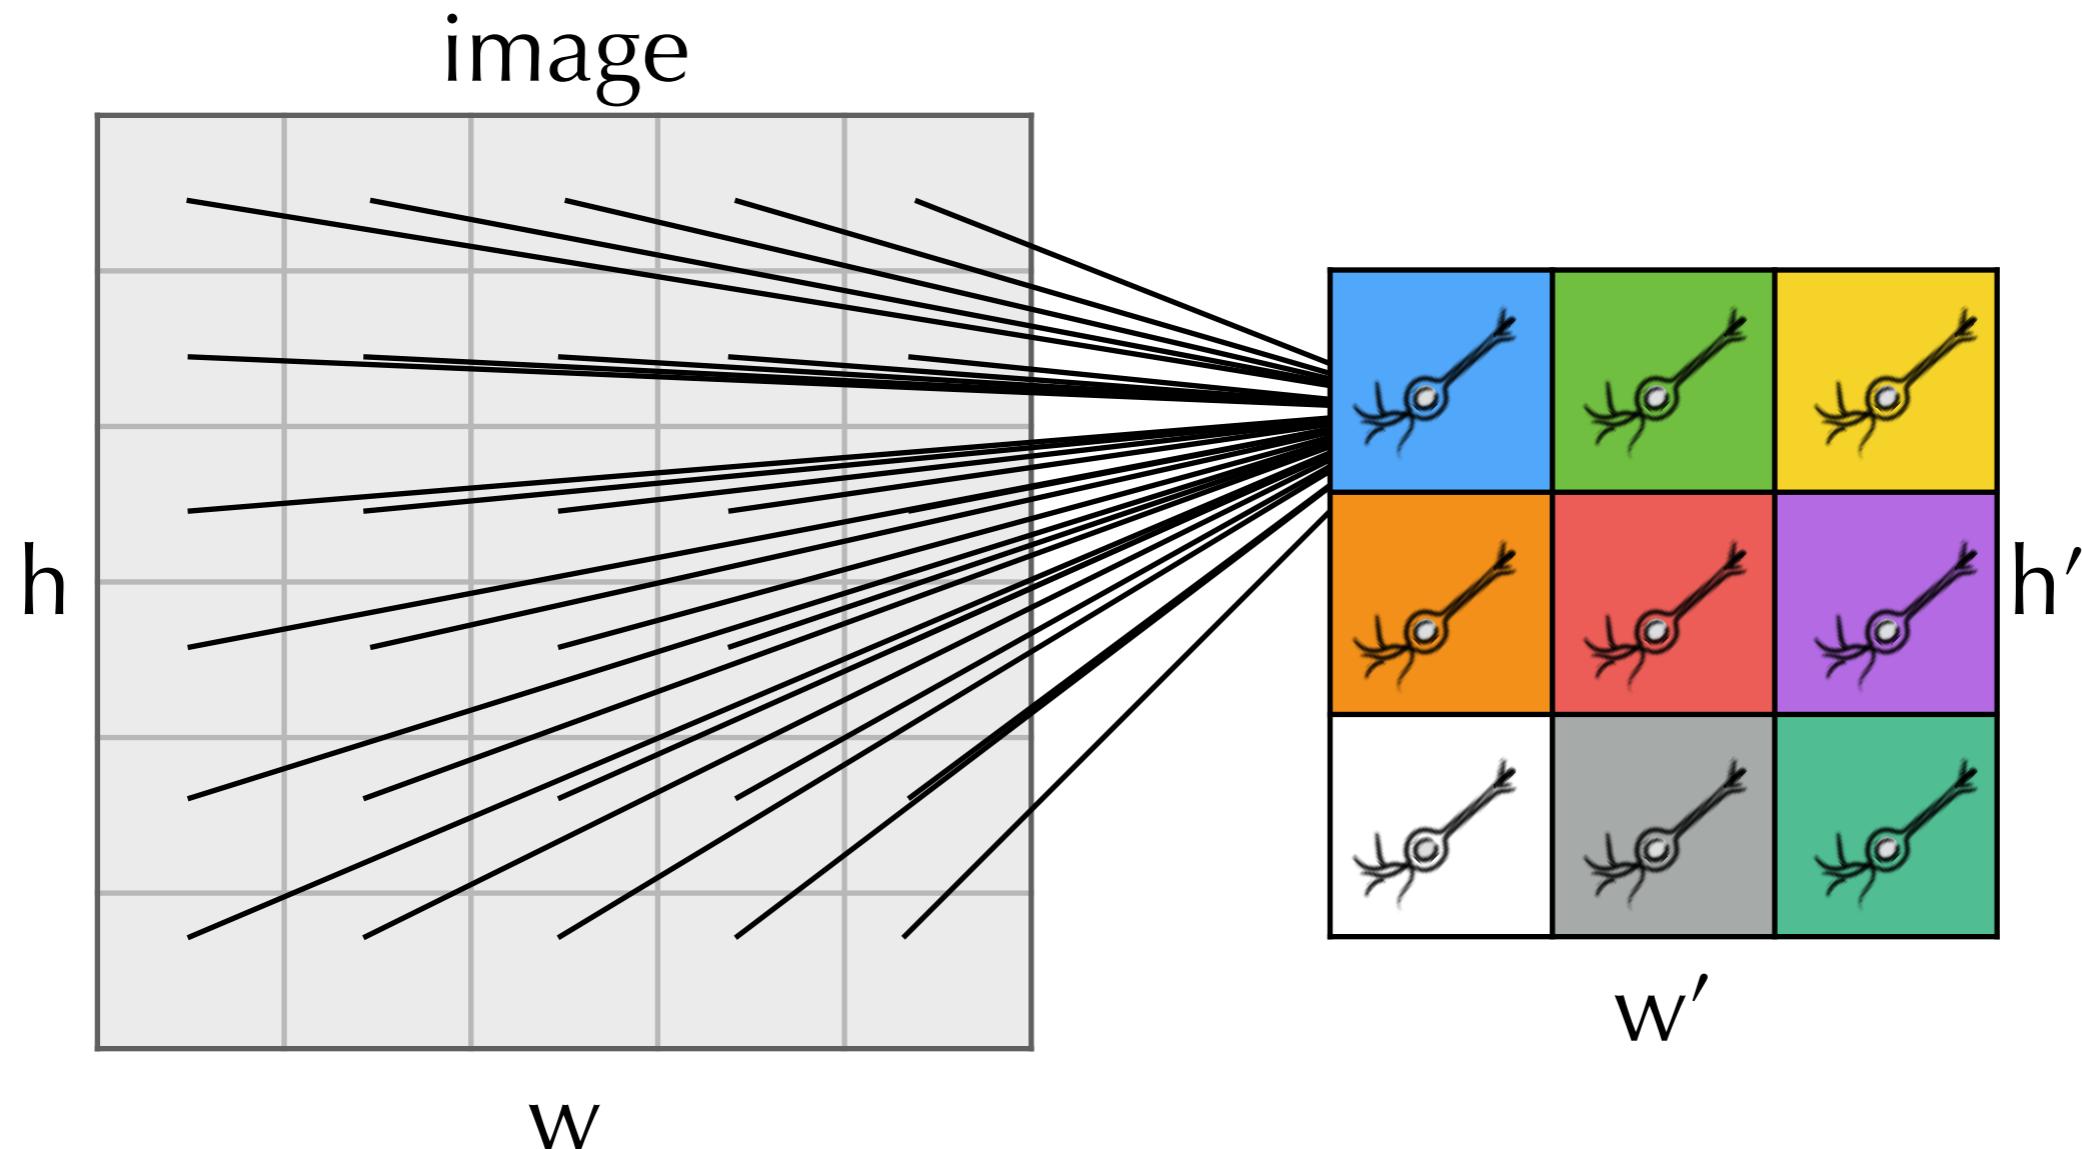
$$h \cdot w \cdot h' \cdot w'$$

Fully Connected Layers



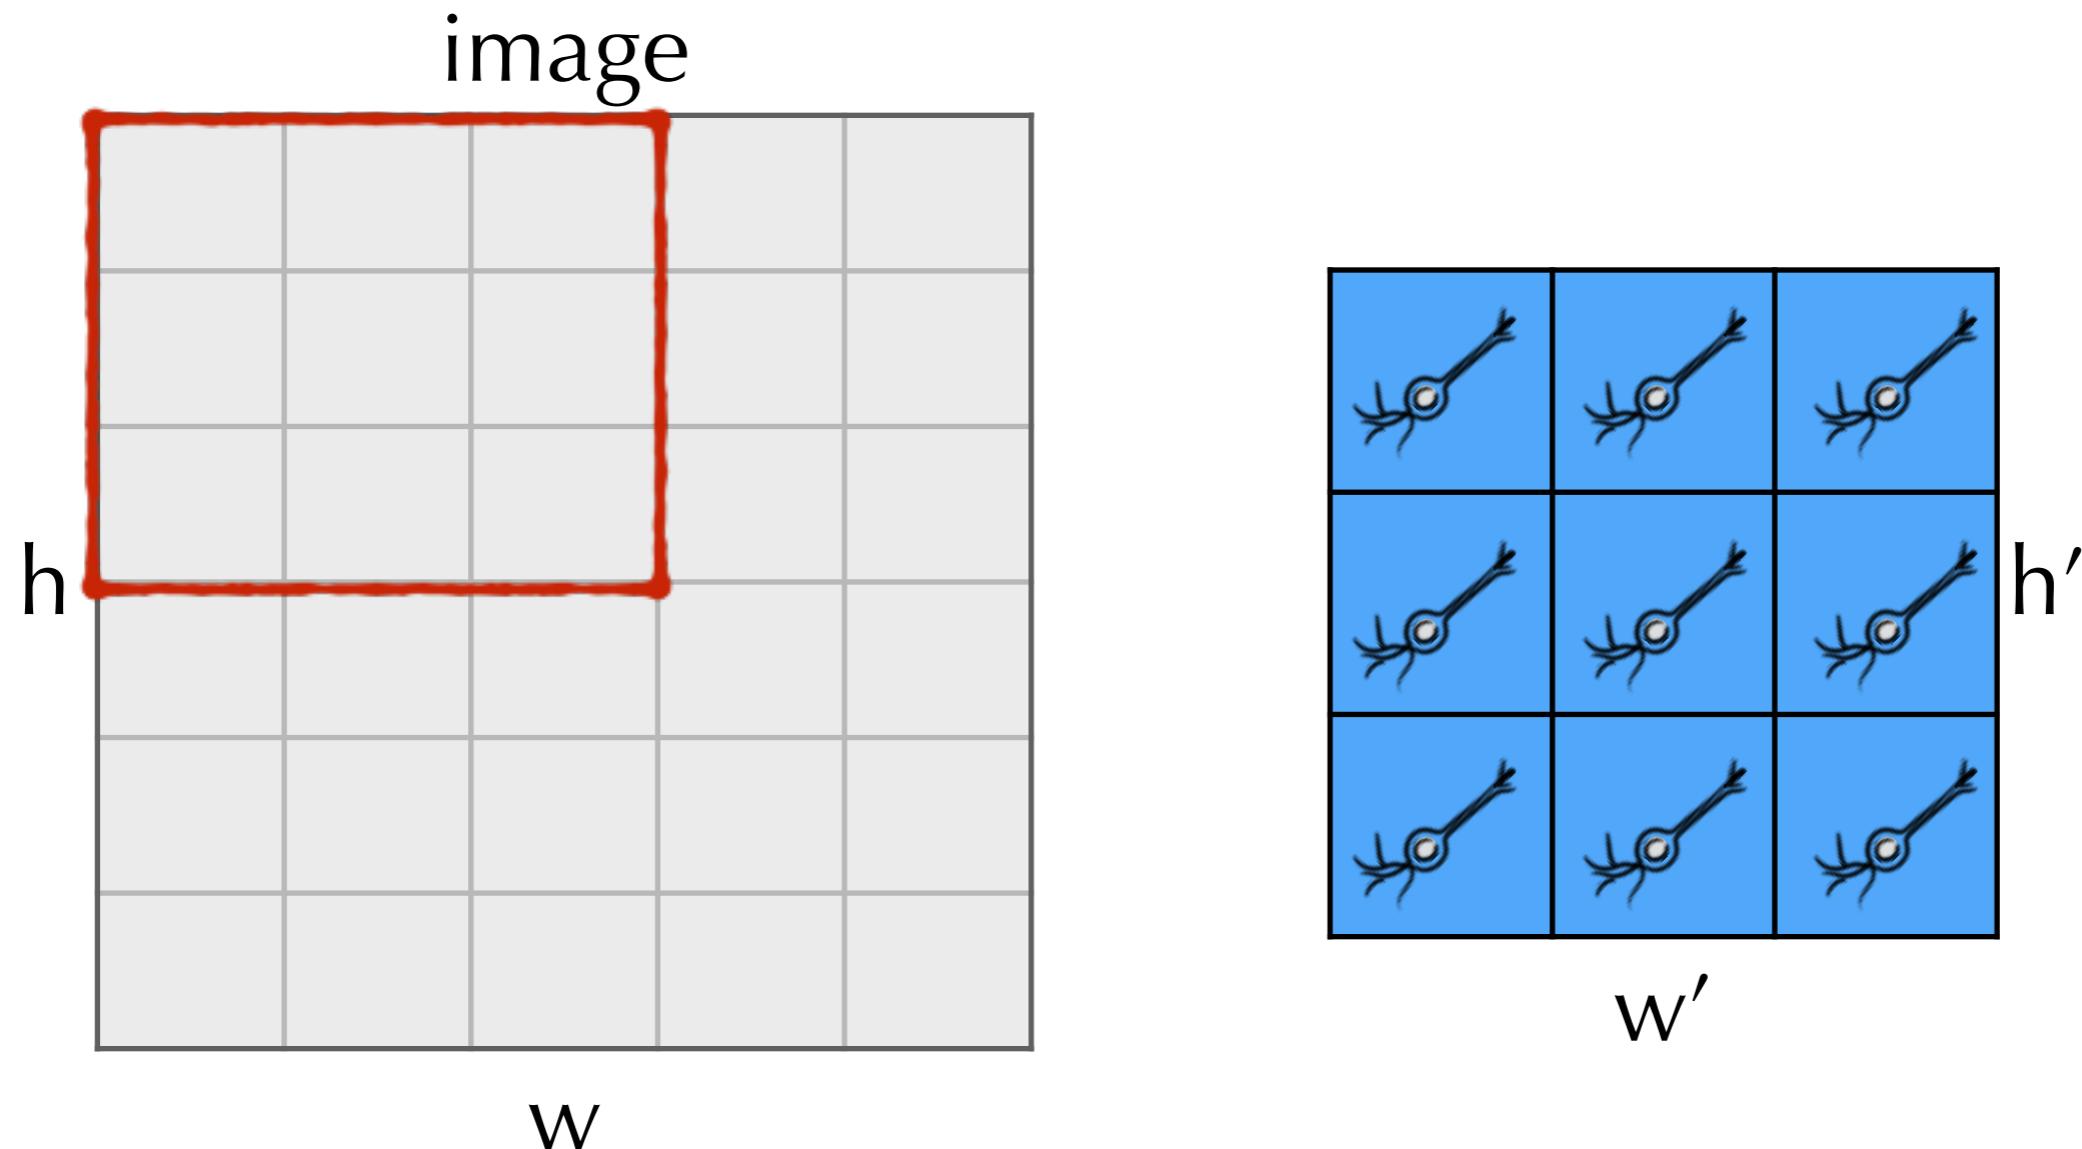
$$h \cdot w \cdot h' \cdot w'$$

Fully Connected Layers

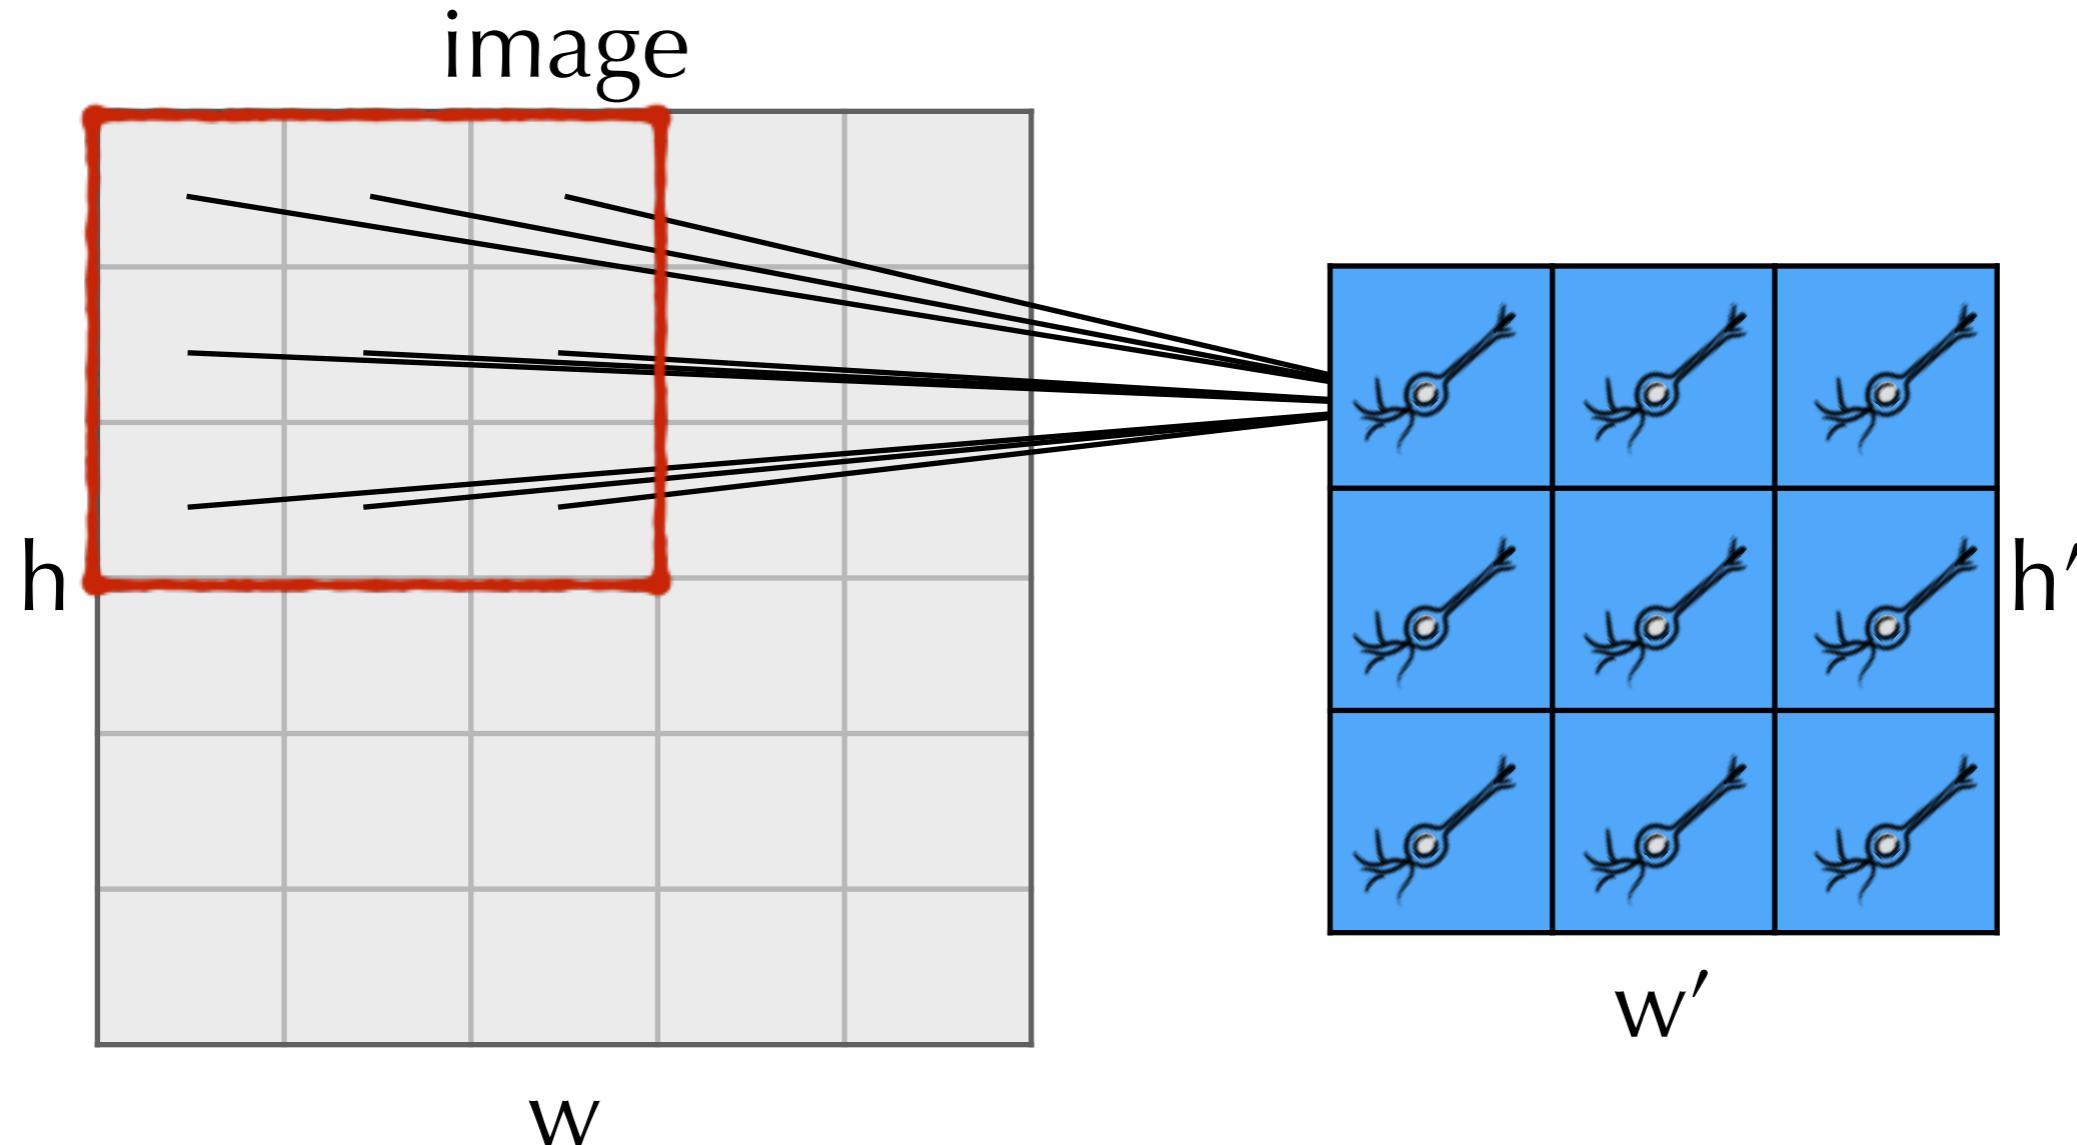


Number parameters: $h \cdot w \cdot h' \cdot w'$

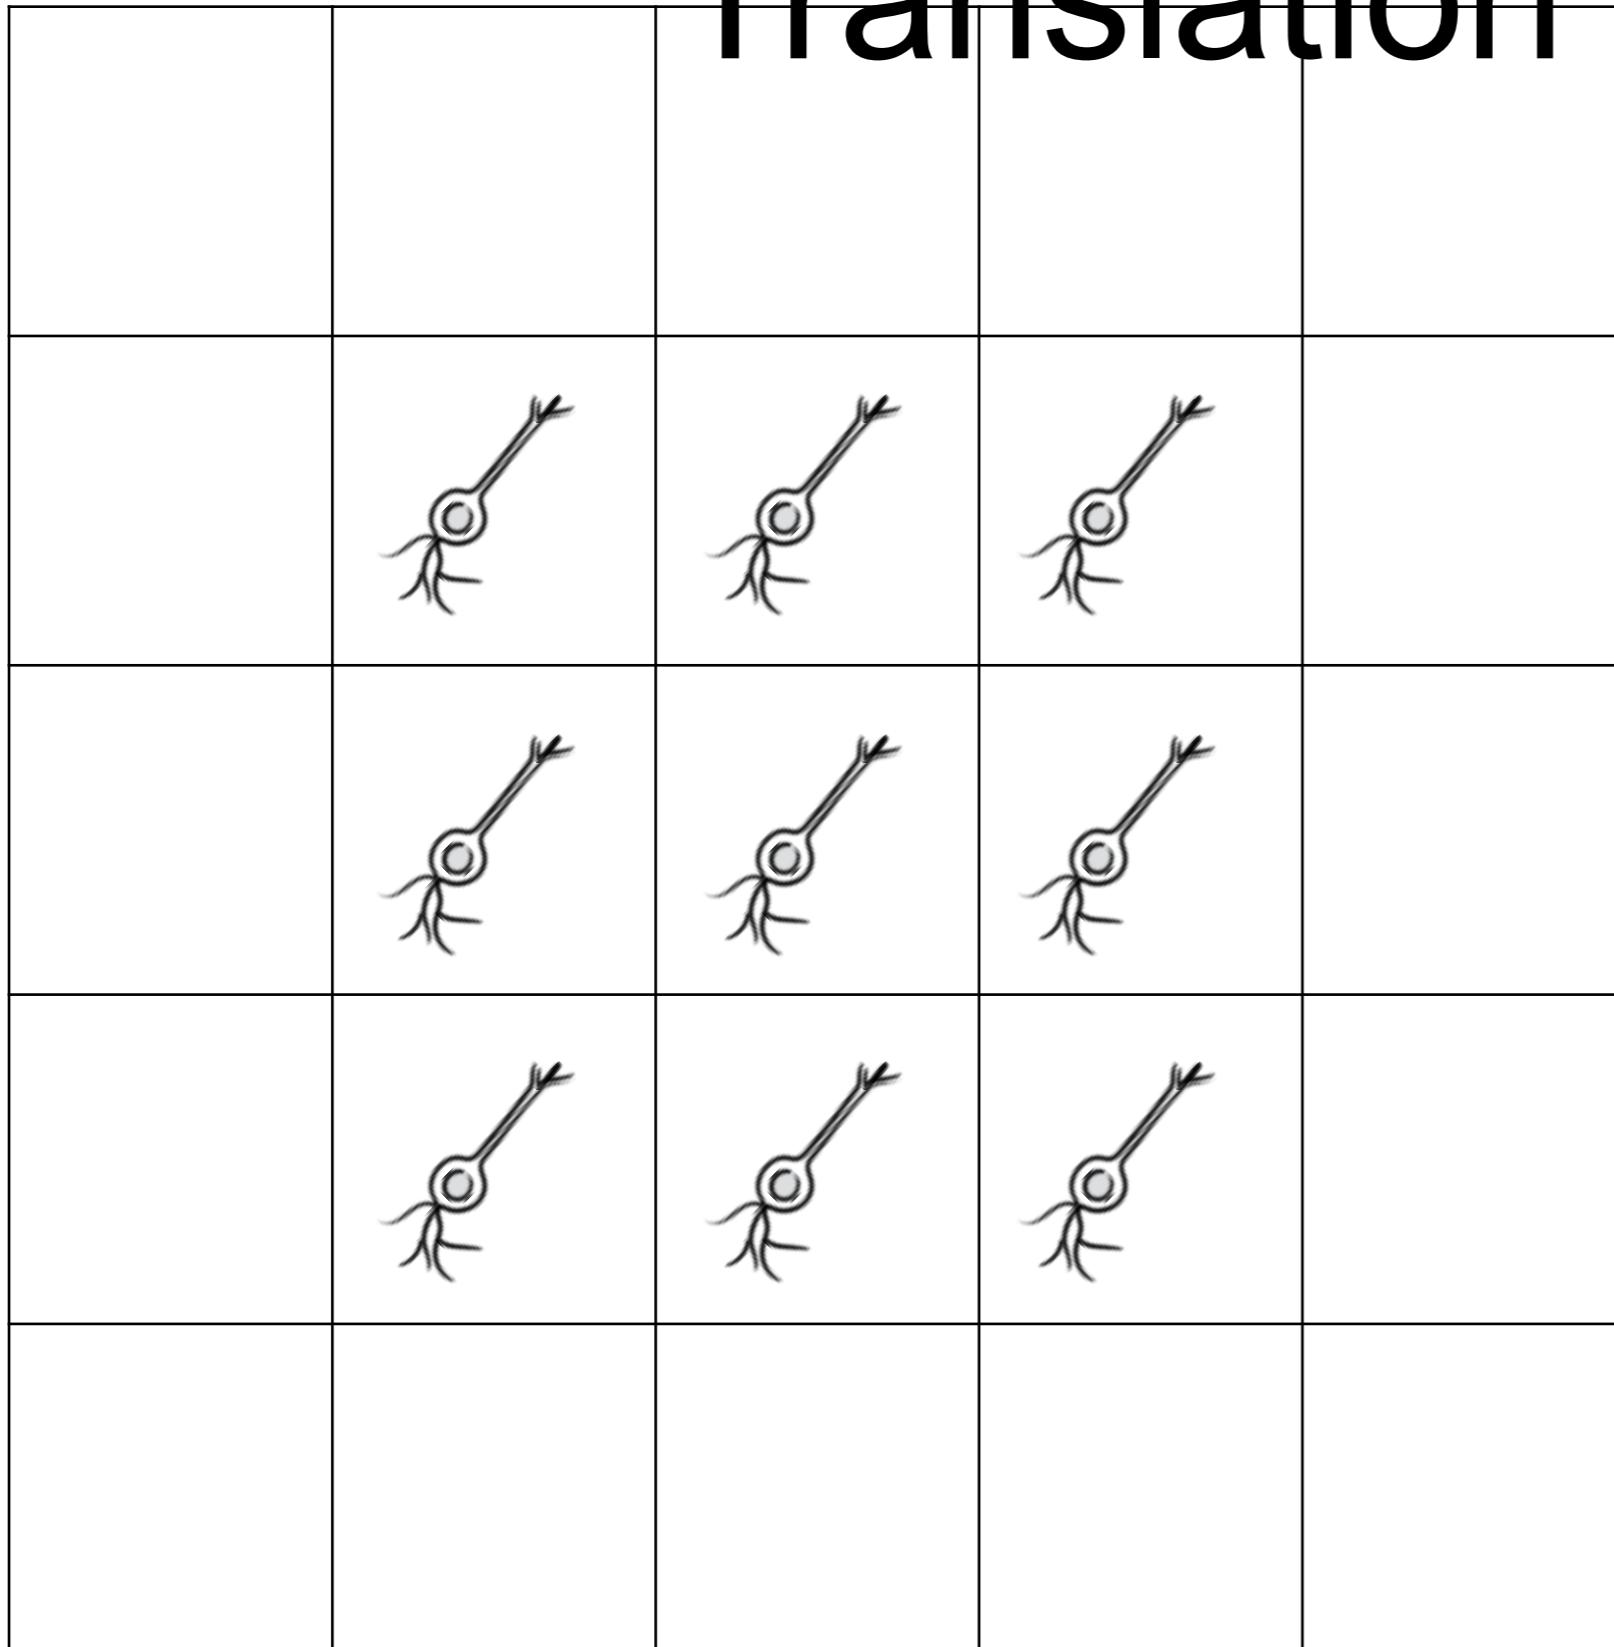
Convolutional Layers



Convolutional Layers

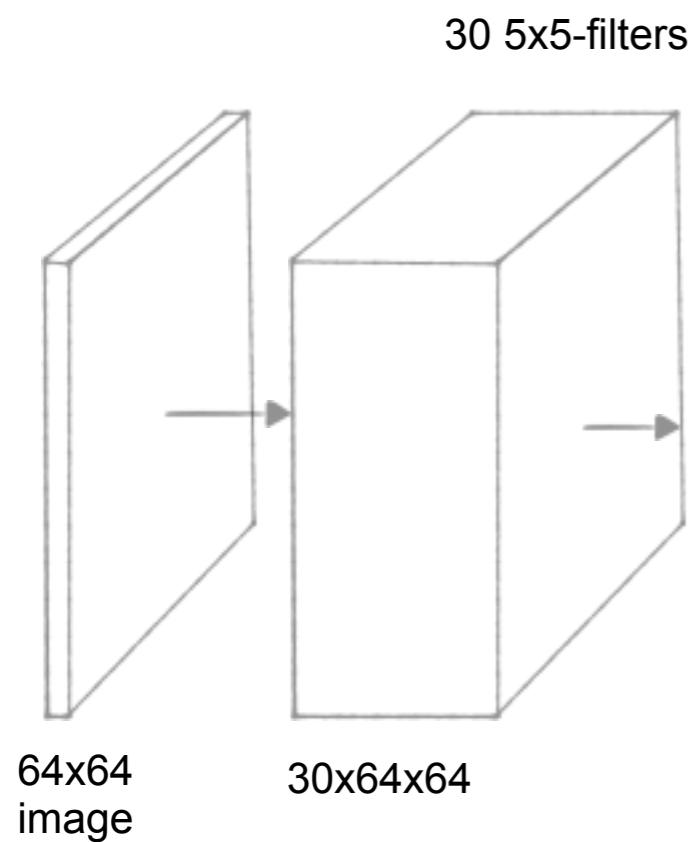


Translation Invariance



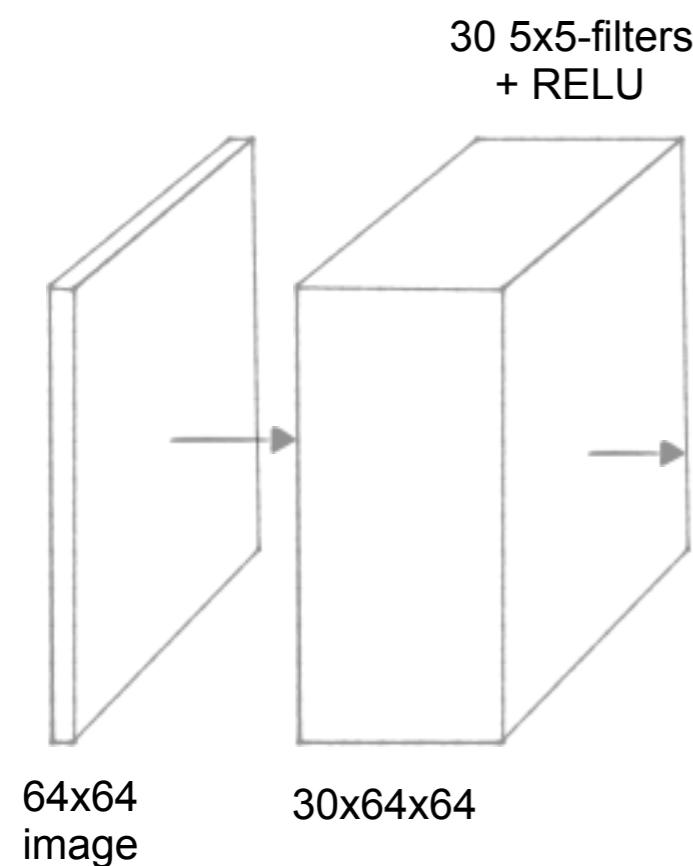
$$I \star \begin{pmatrix} a & b \\ c & d \end{pmatrix}$$

Network Structure



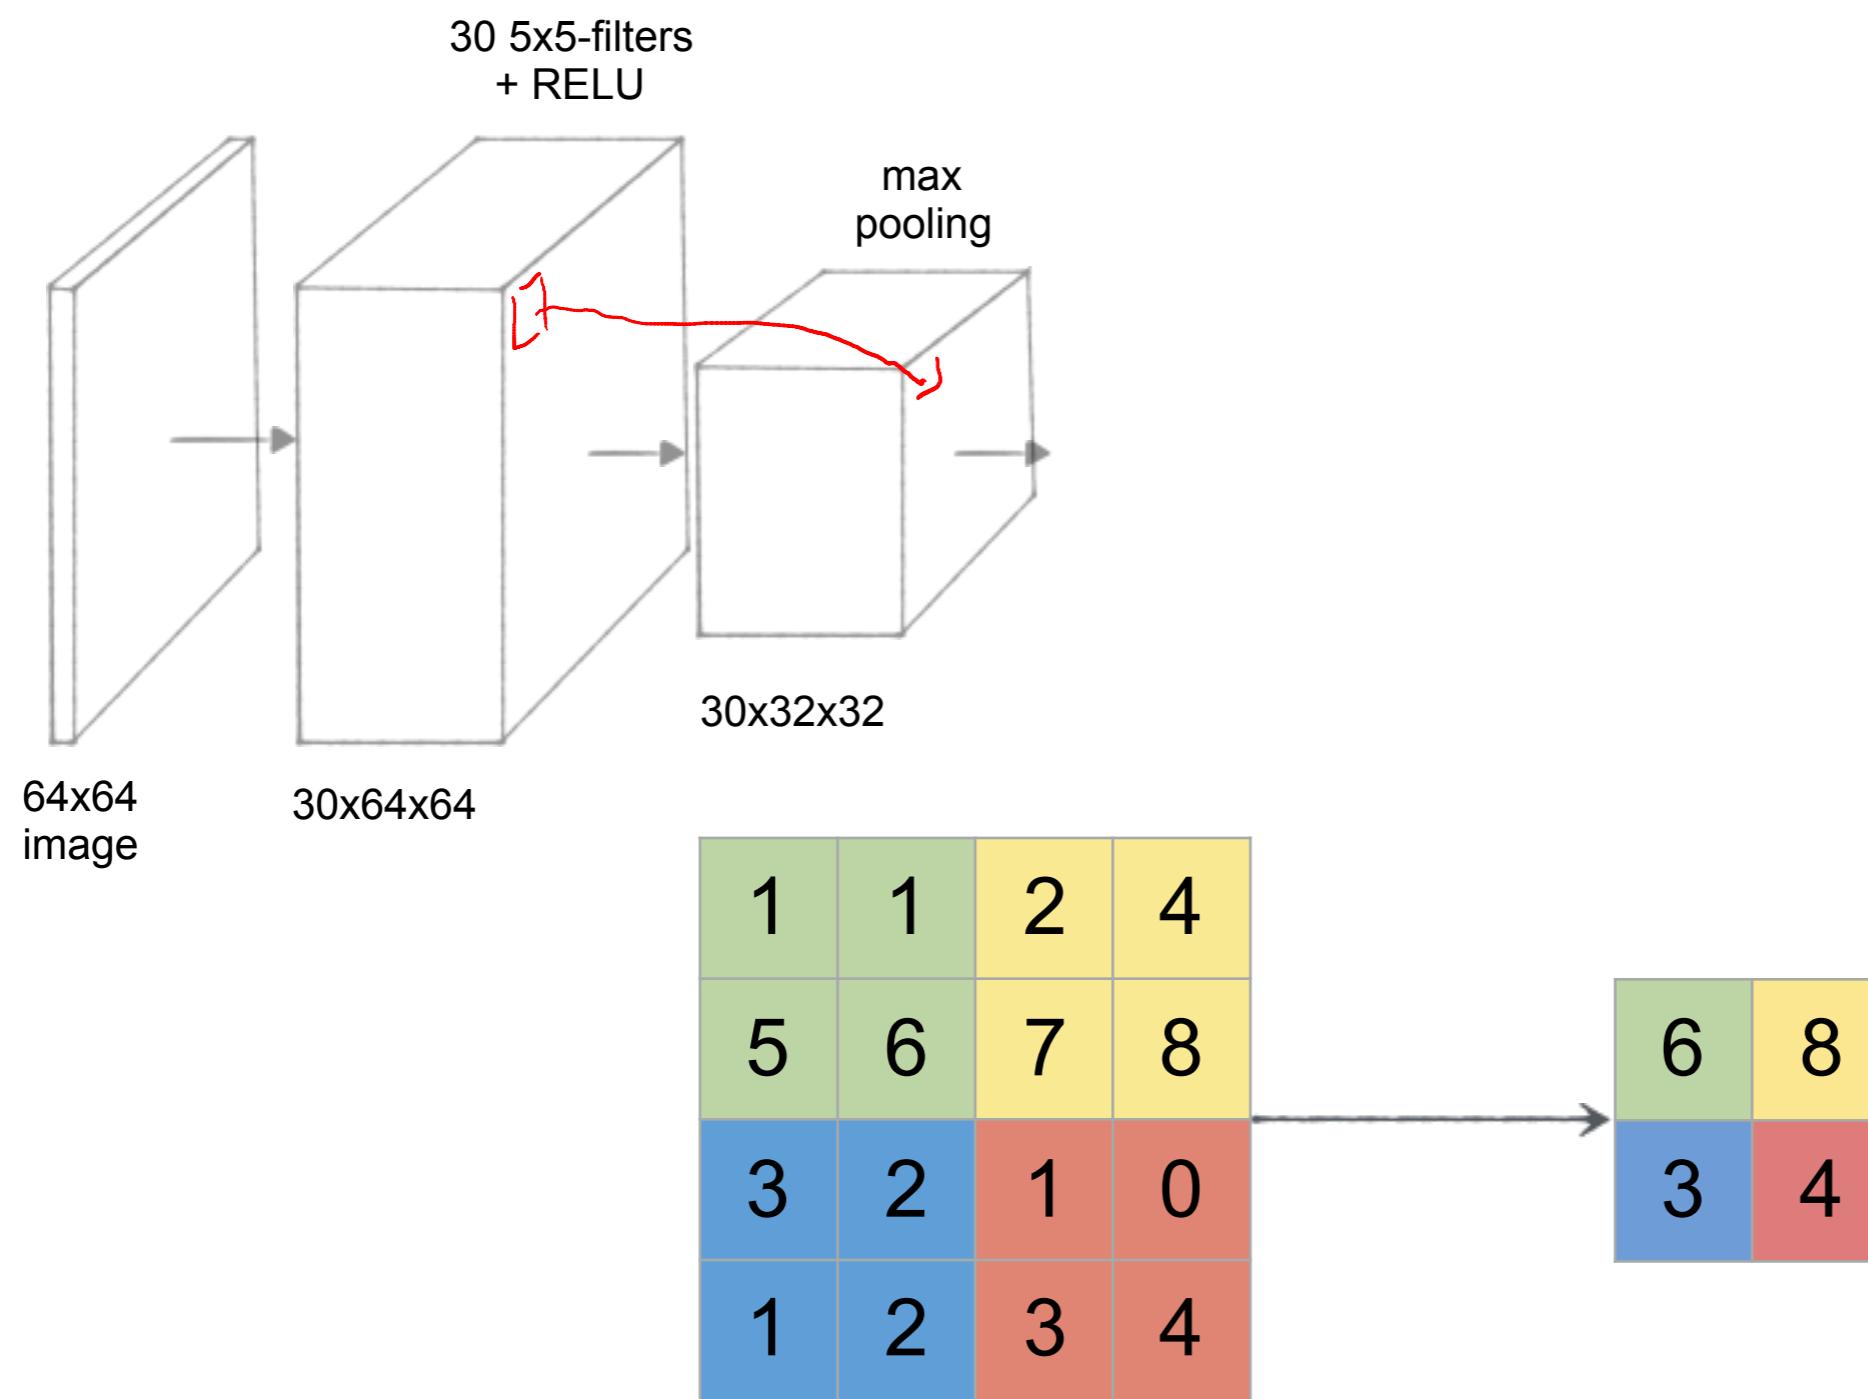
$$I \star w$$

Network Structure

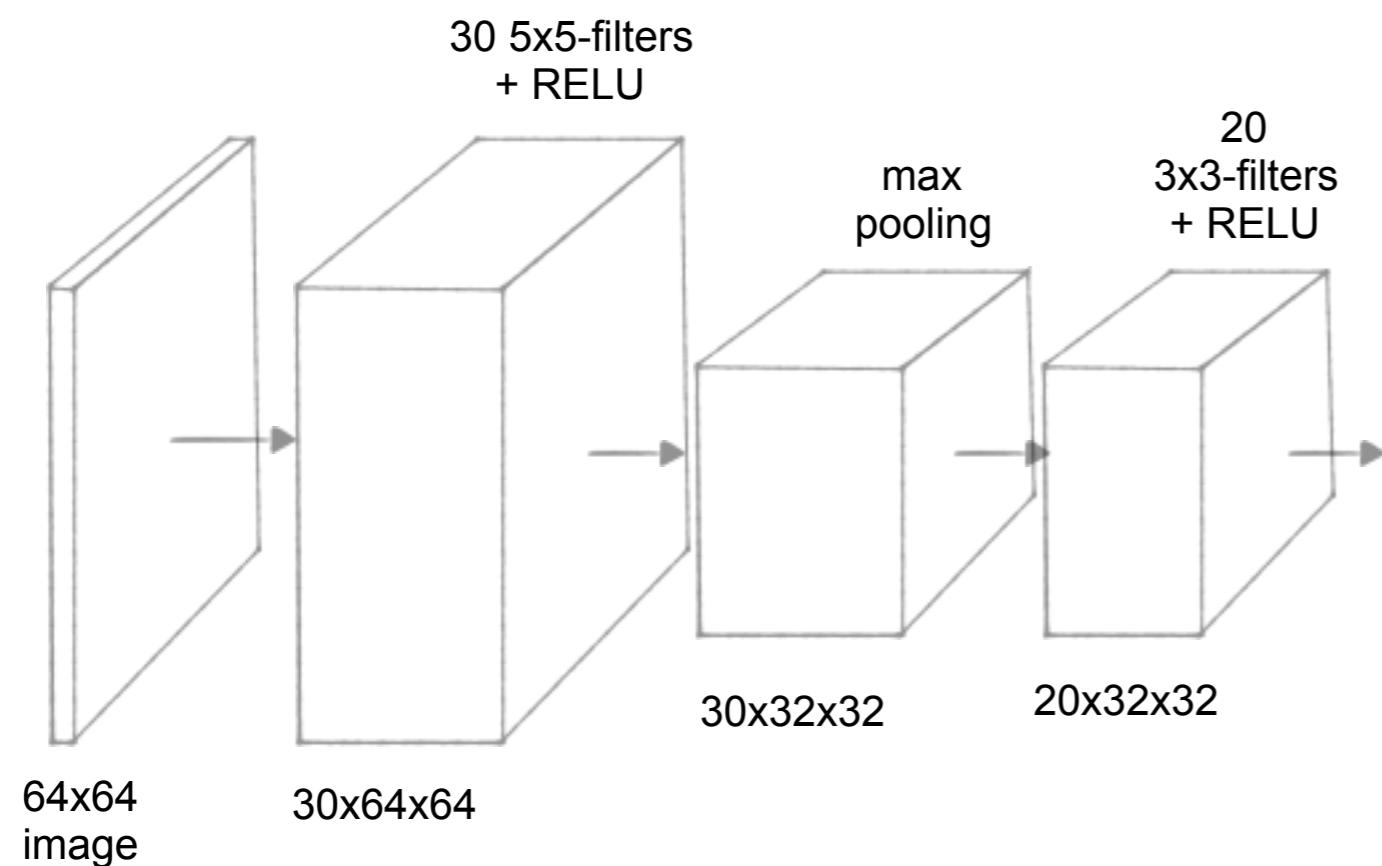


$$\max\{I \star w, 0\}$$

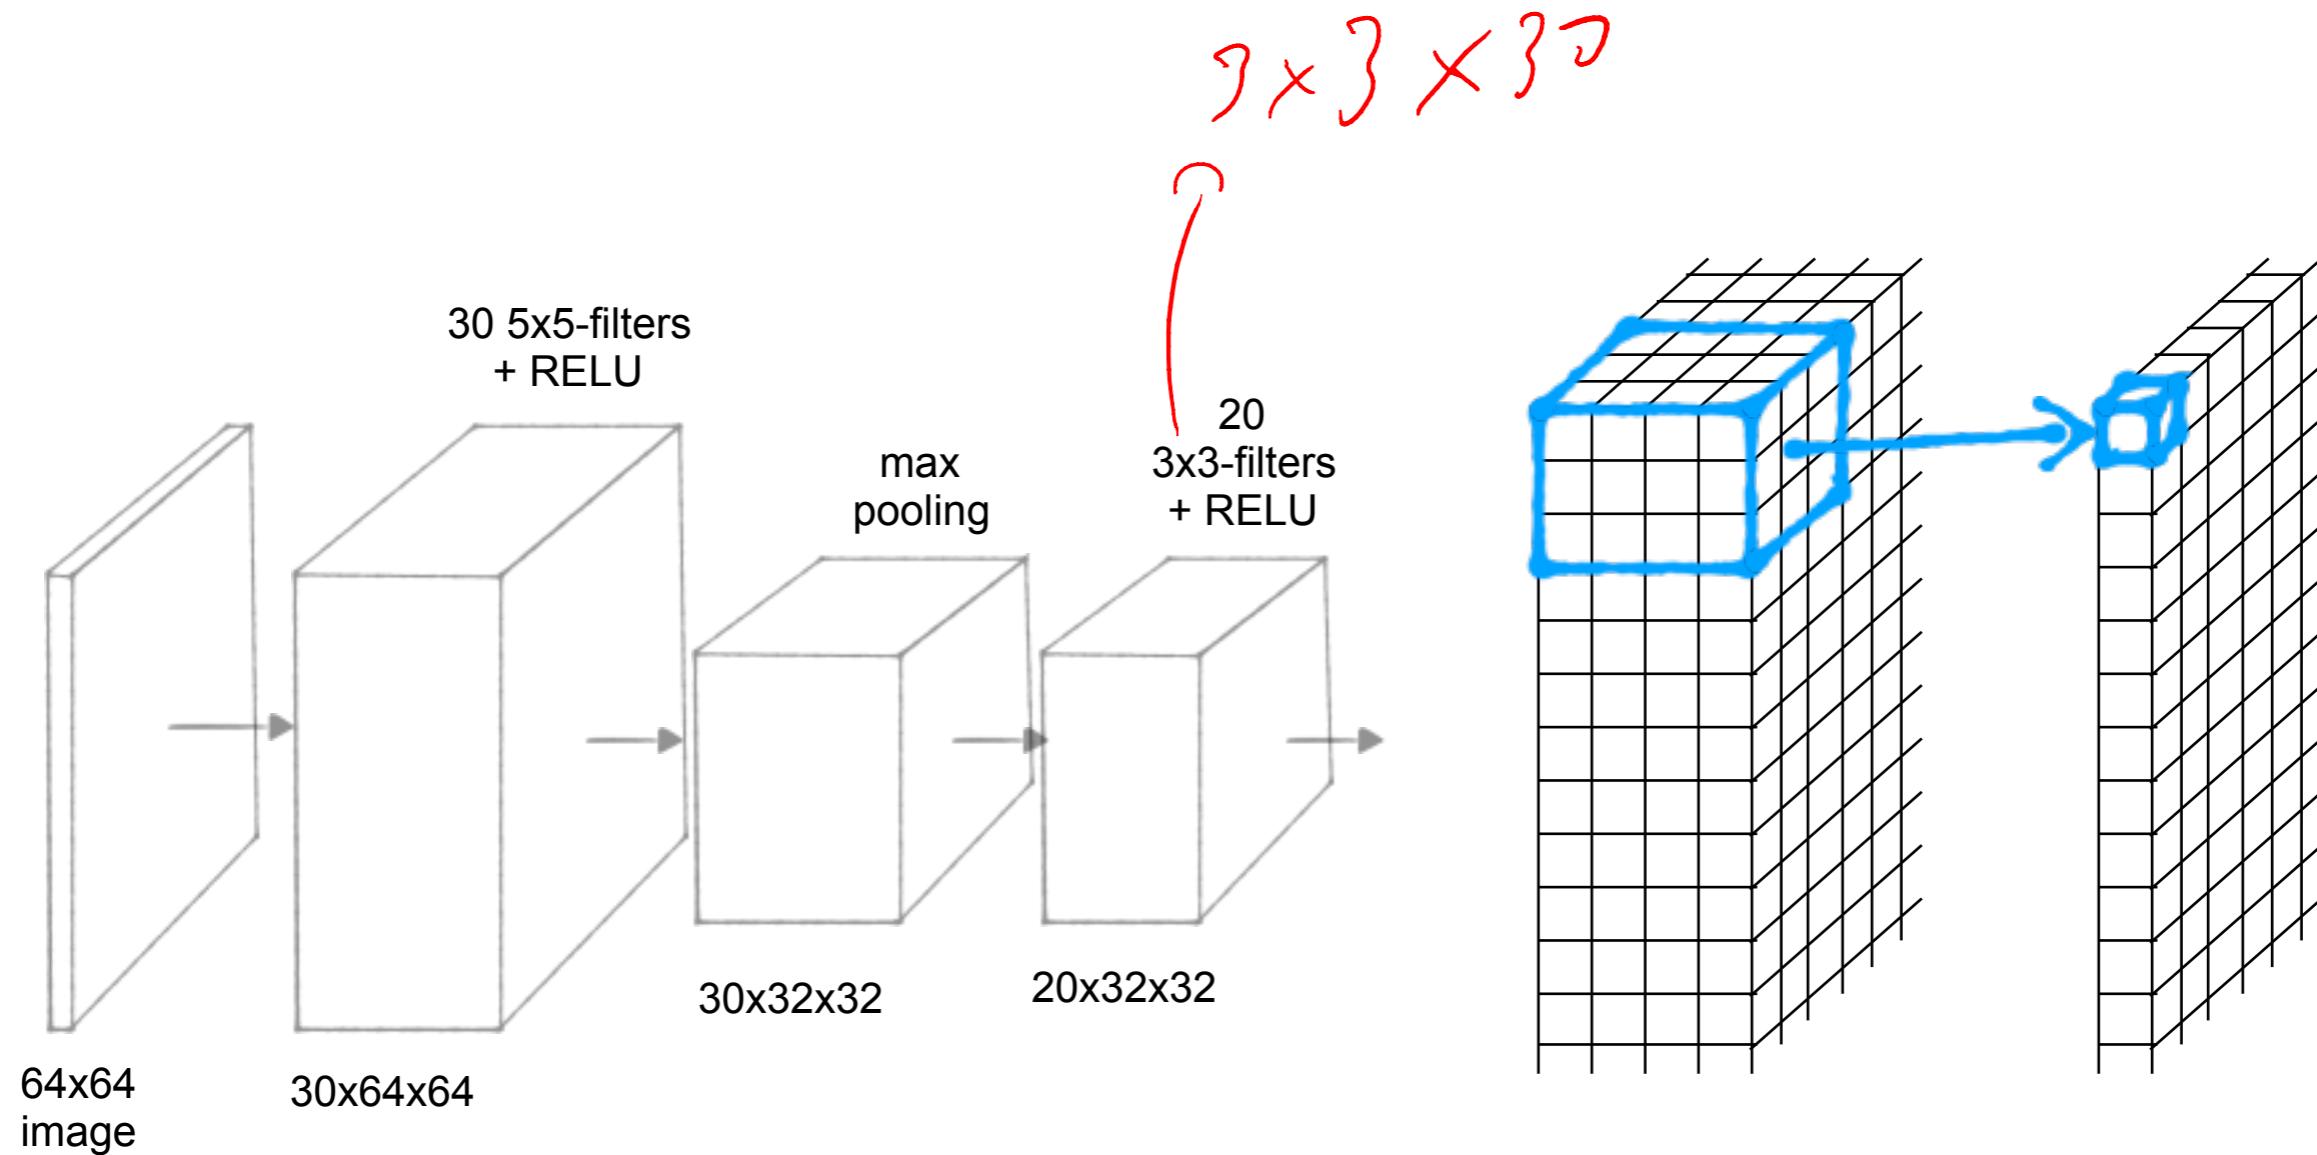
Network Structure



Network Structure



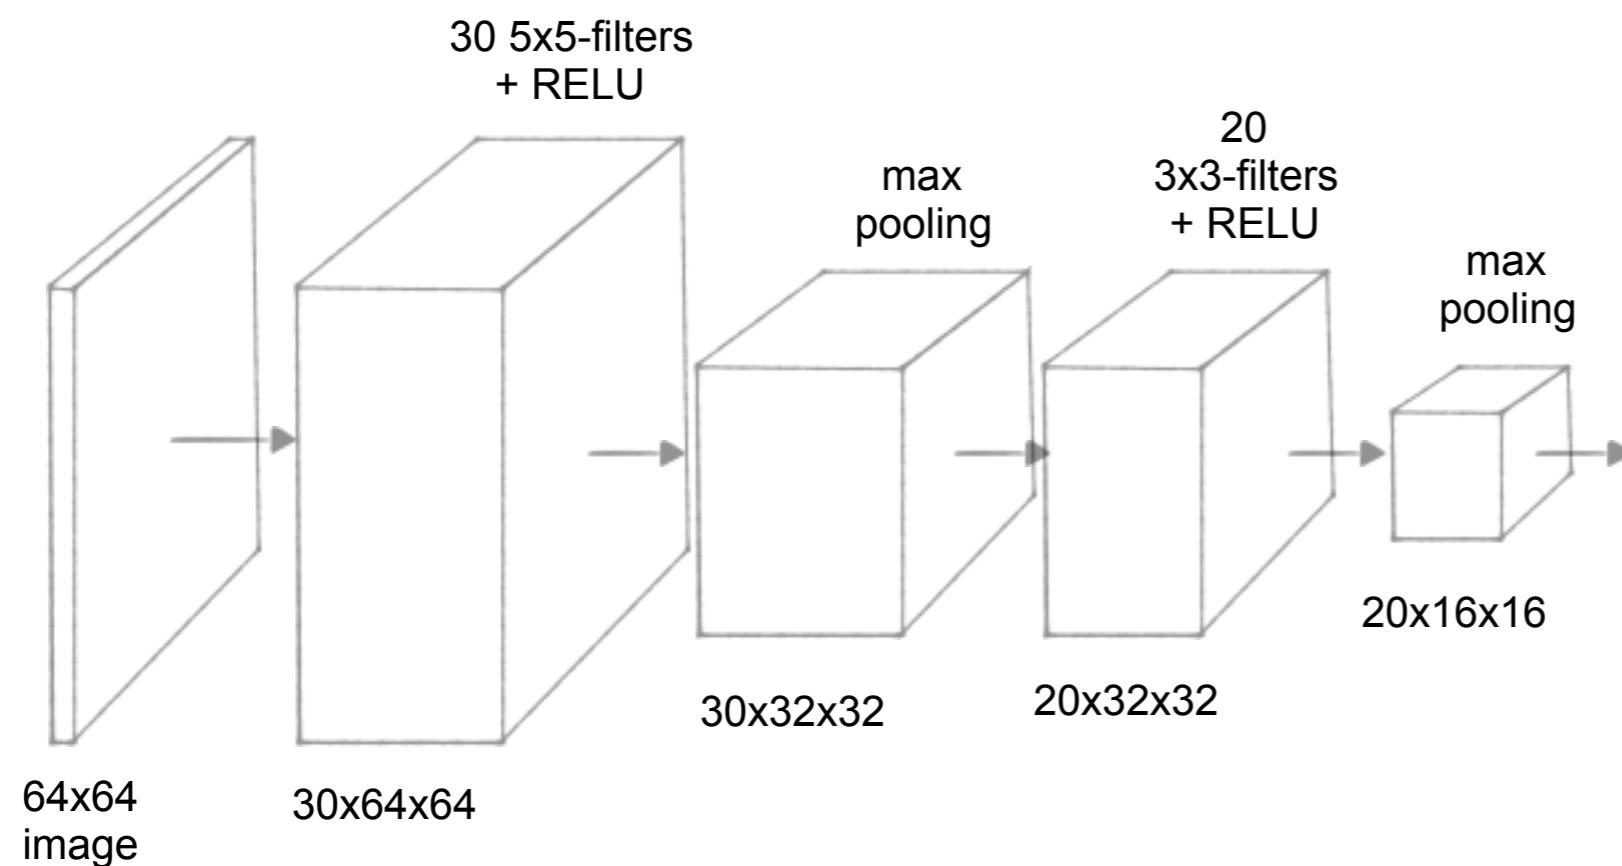
Network Structure



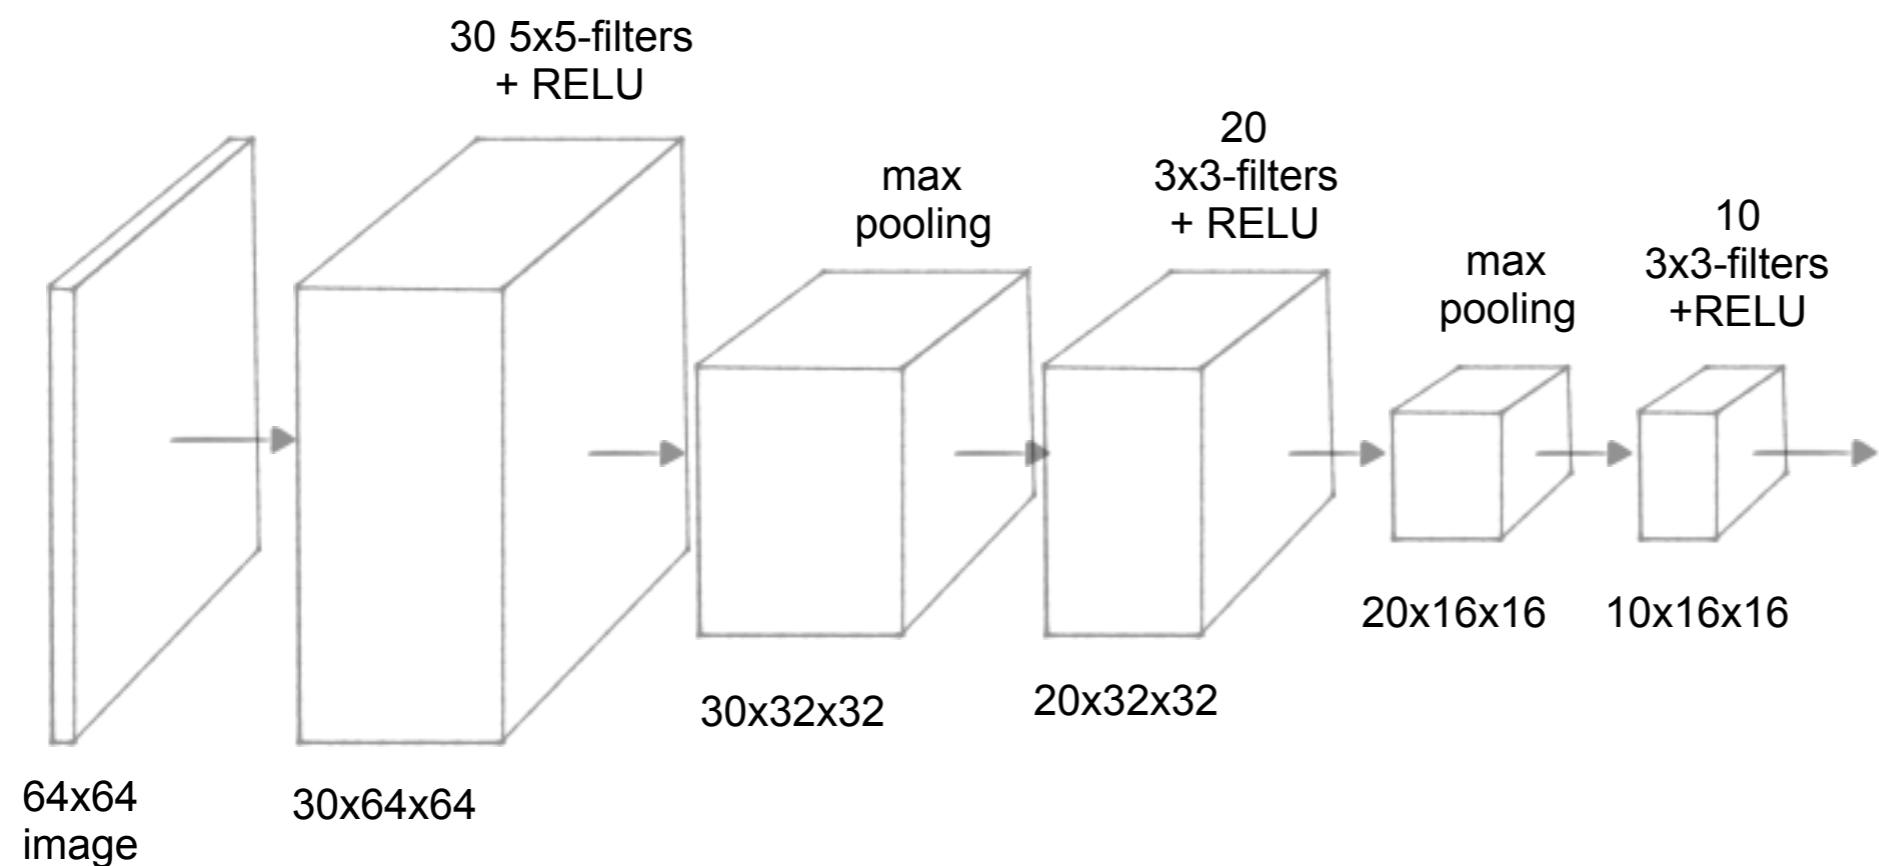
implemented as matrix multiplication:

$$\text{vec}(\mathcal{J}) = F \text{vec}(\mathcal{I})$$

Network Structure

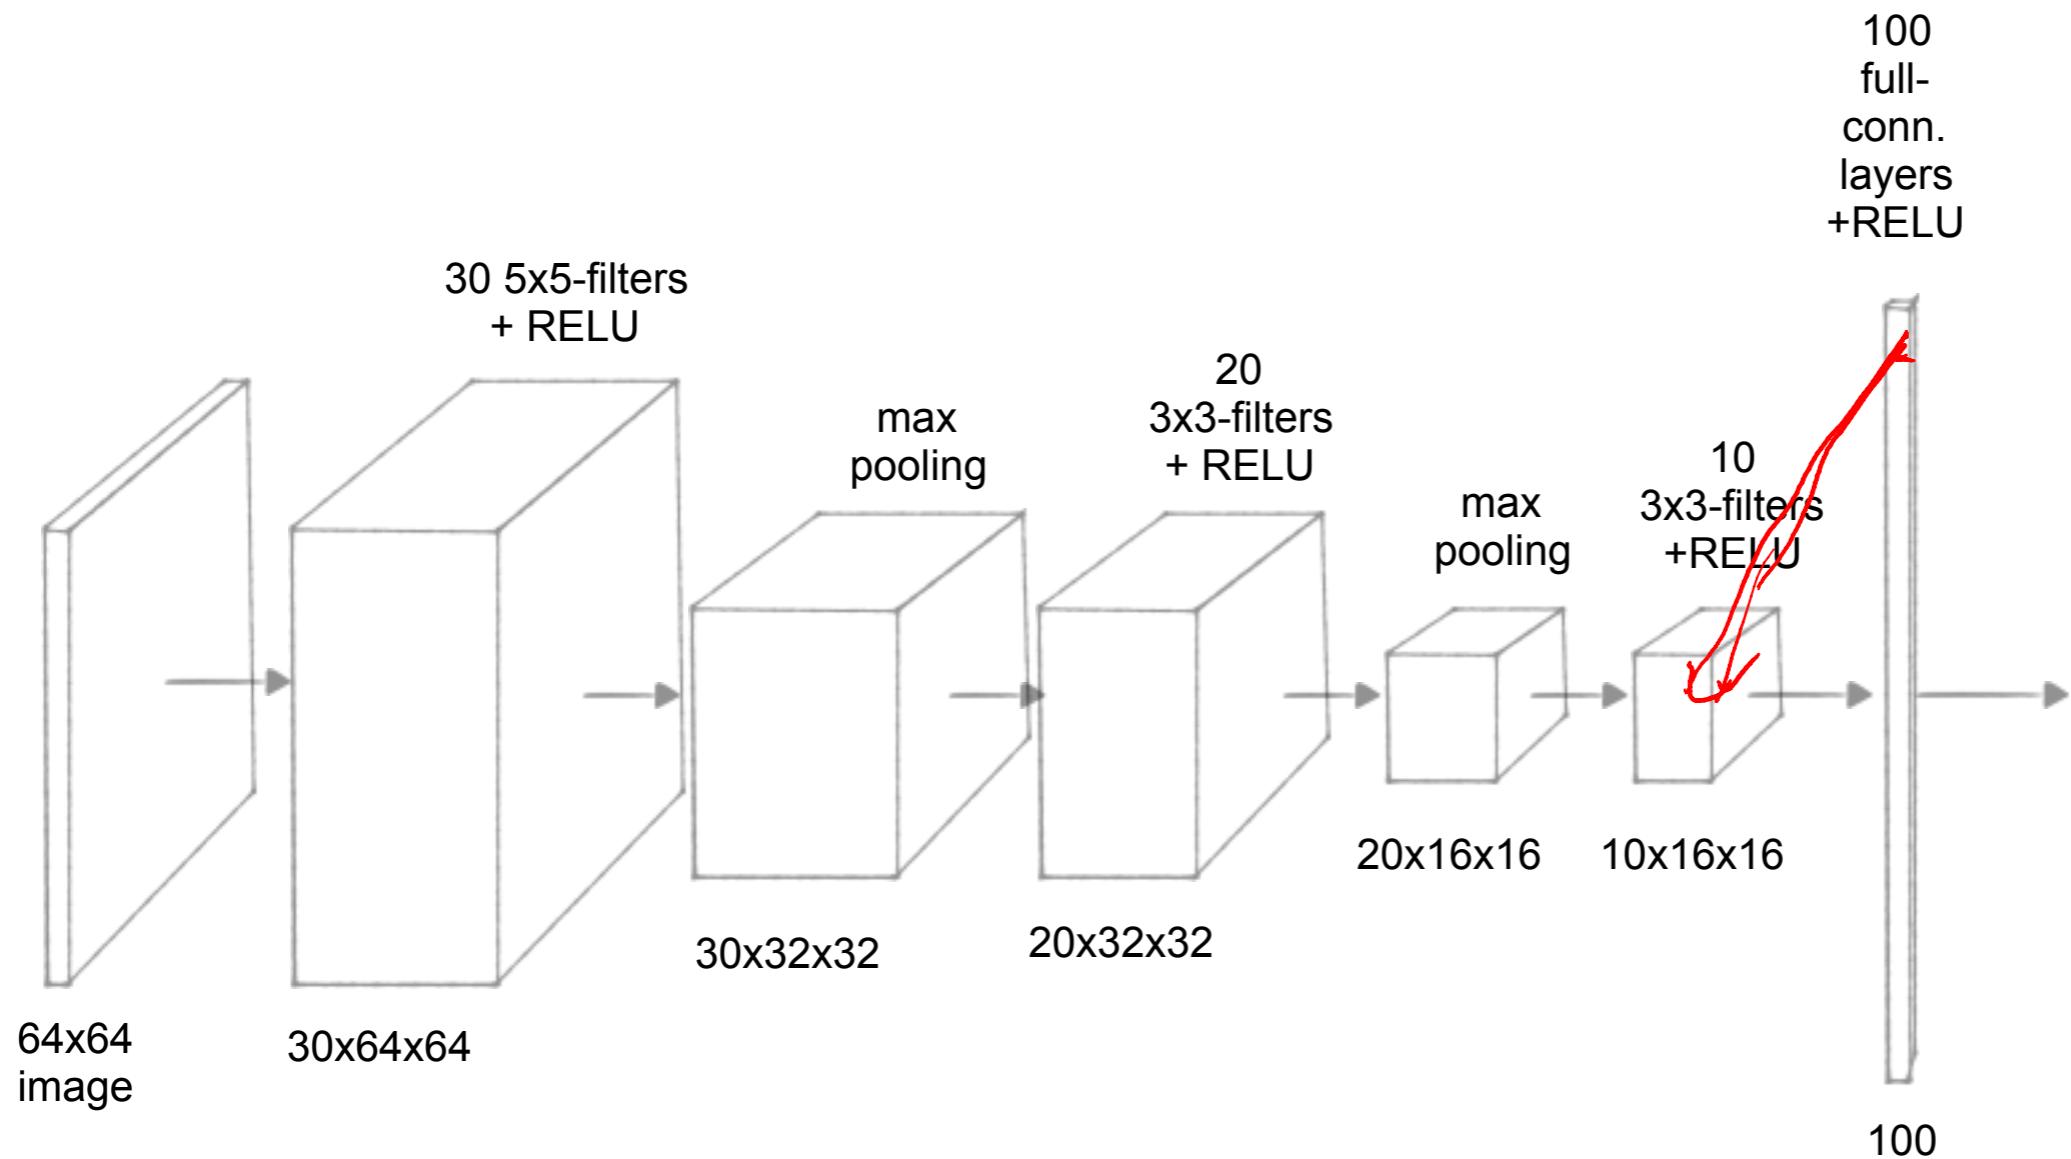


Network Structure



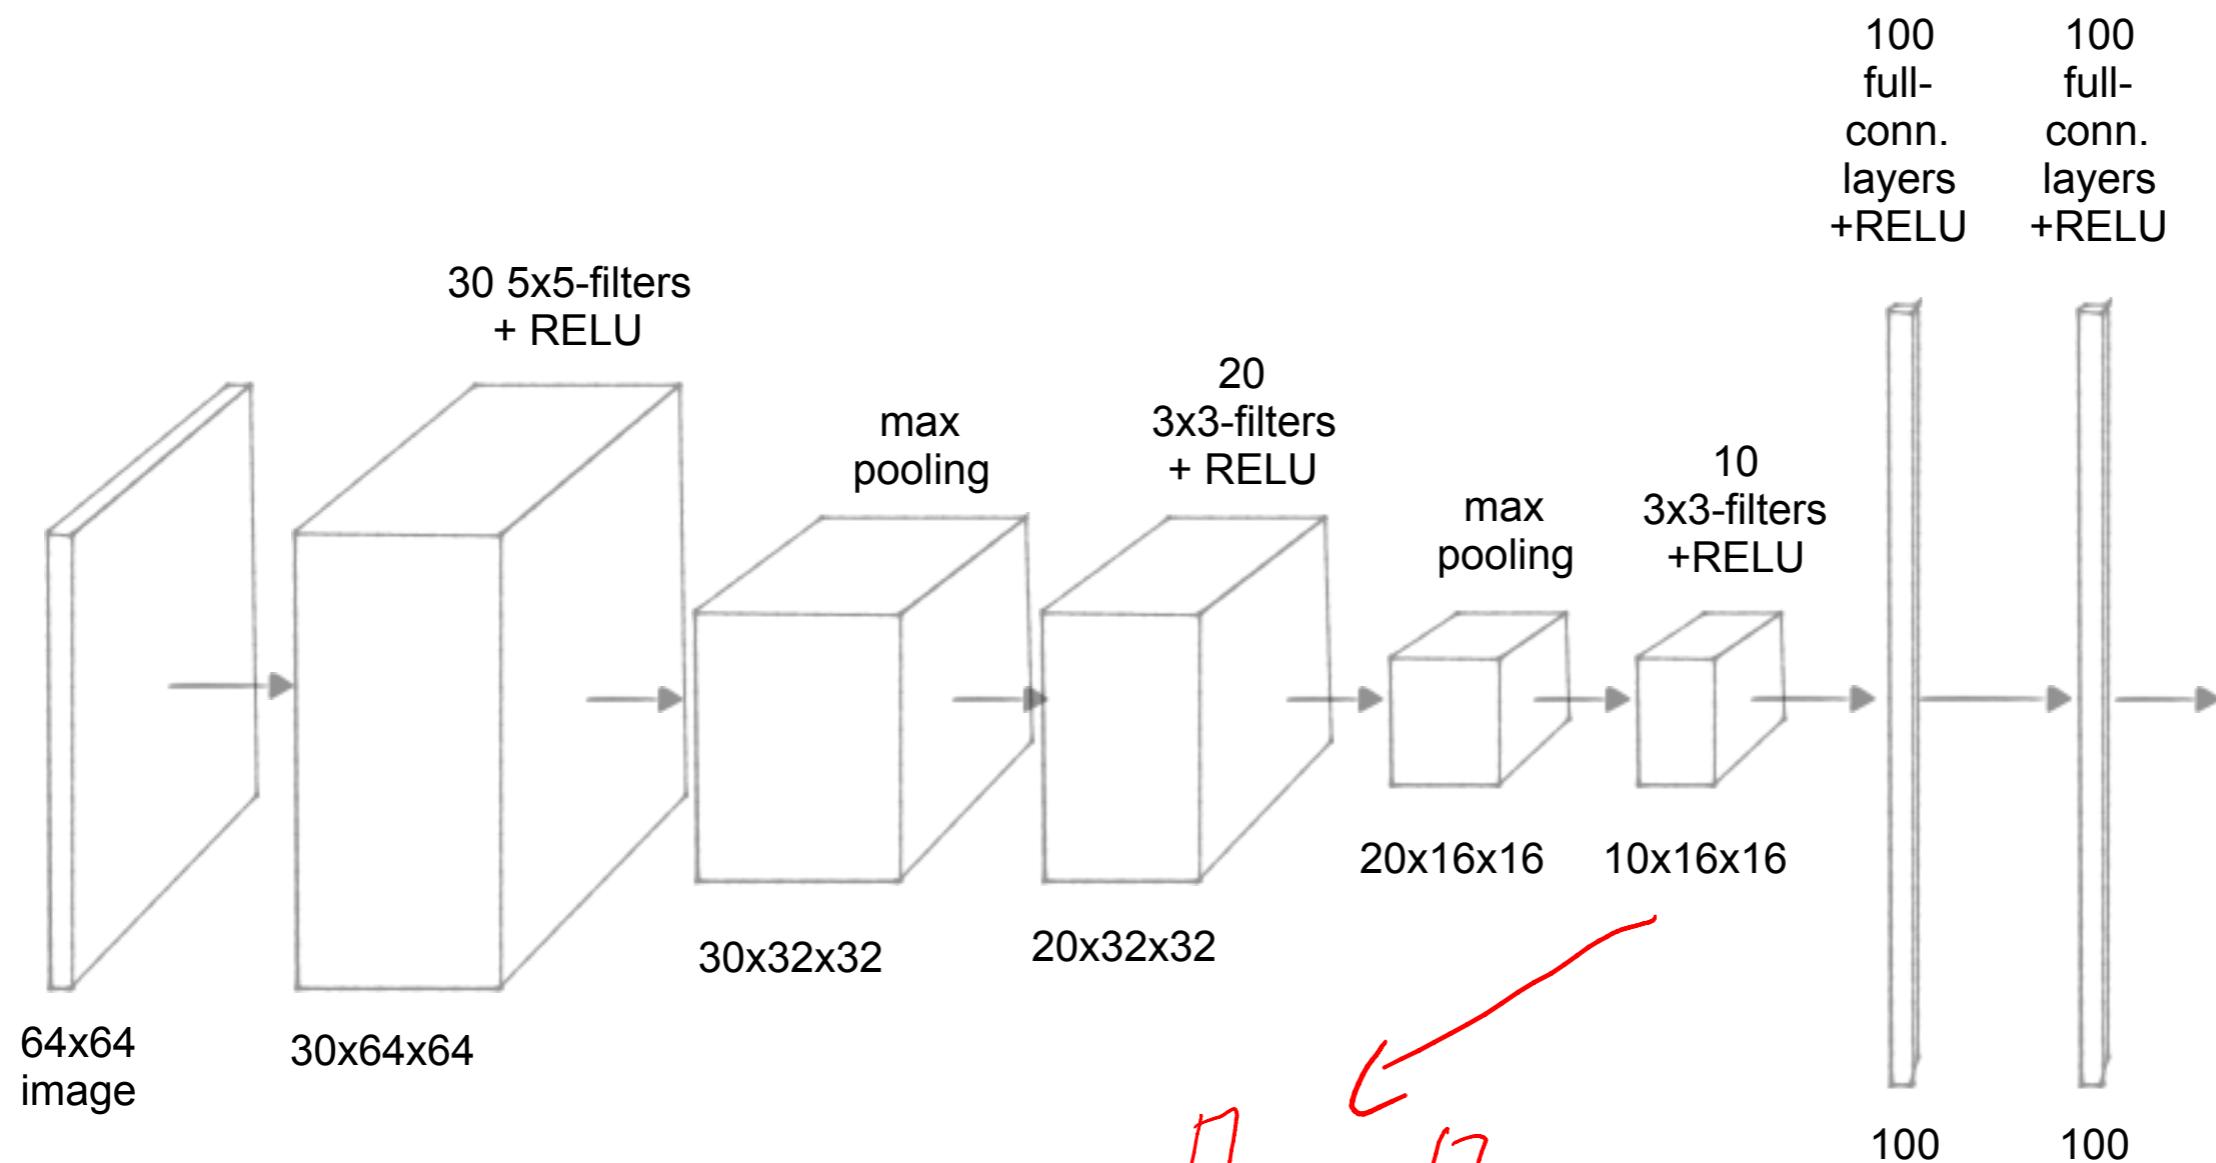
$$\max\{I \otimes f, 0\}$$

Network Structure



$$\sum_{i=1}^{10} \sum_{u=1}^{16} \sum_{v=1}^{16} w_{ijk} x_{ijk}$$

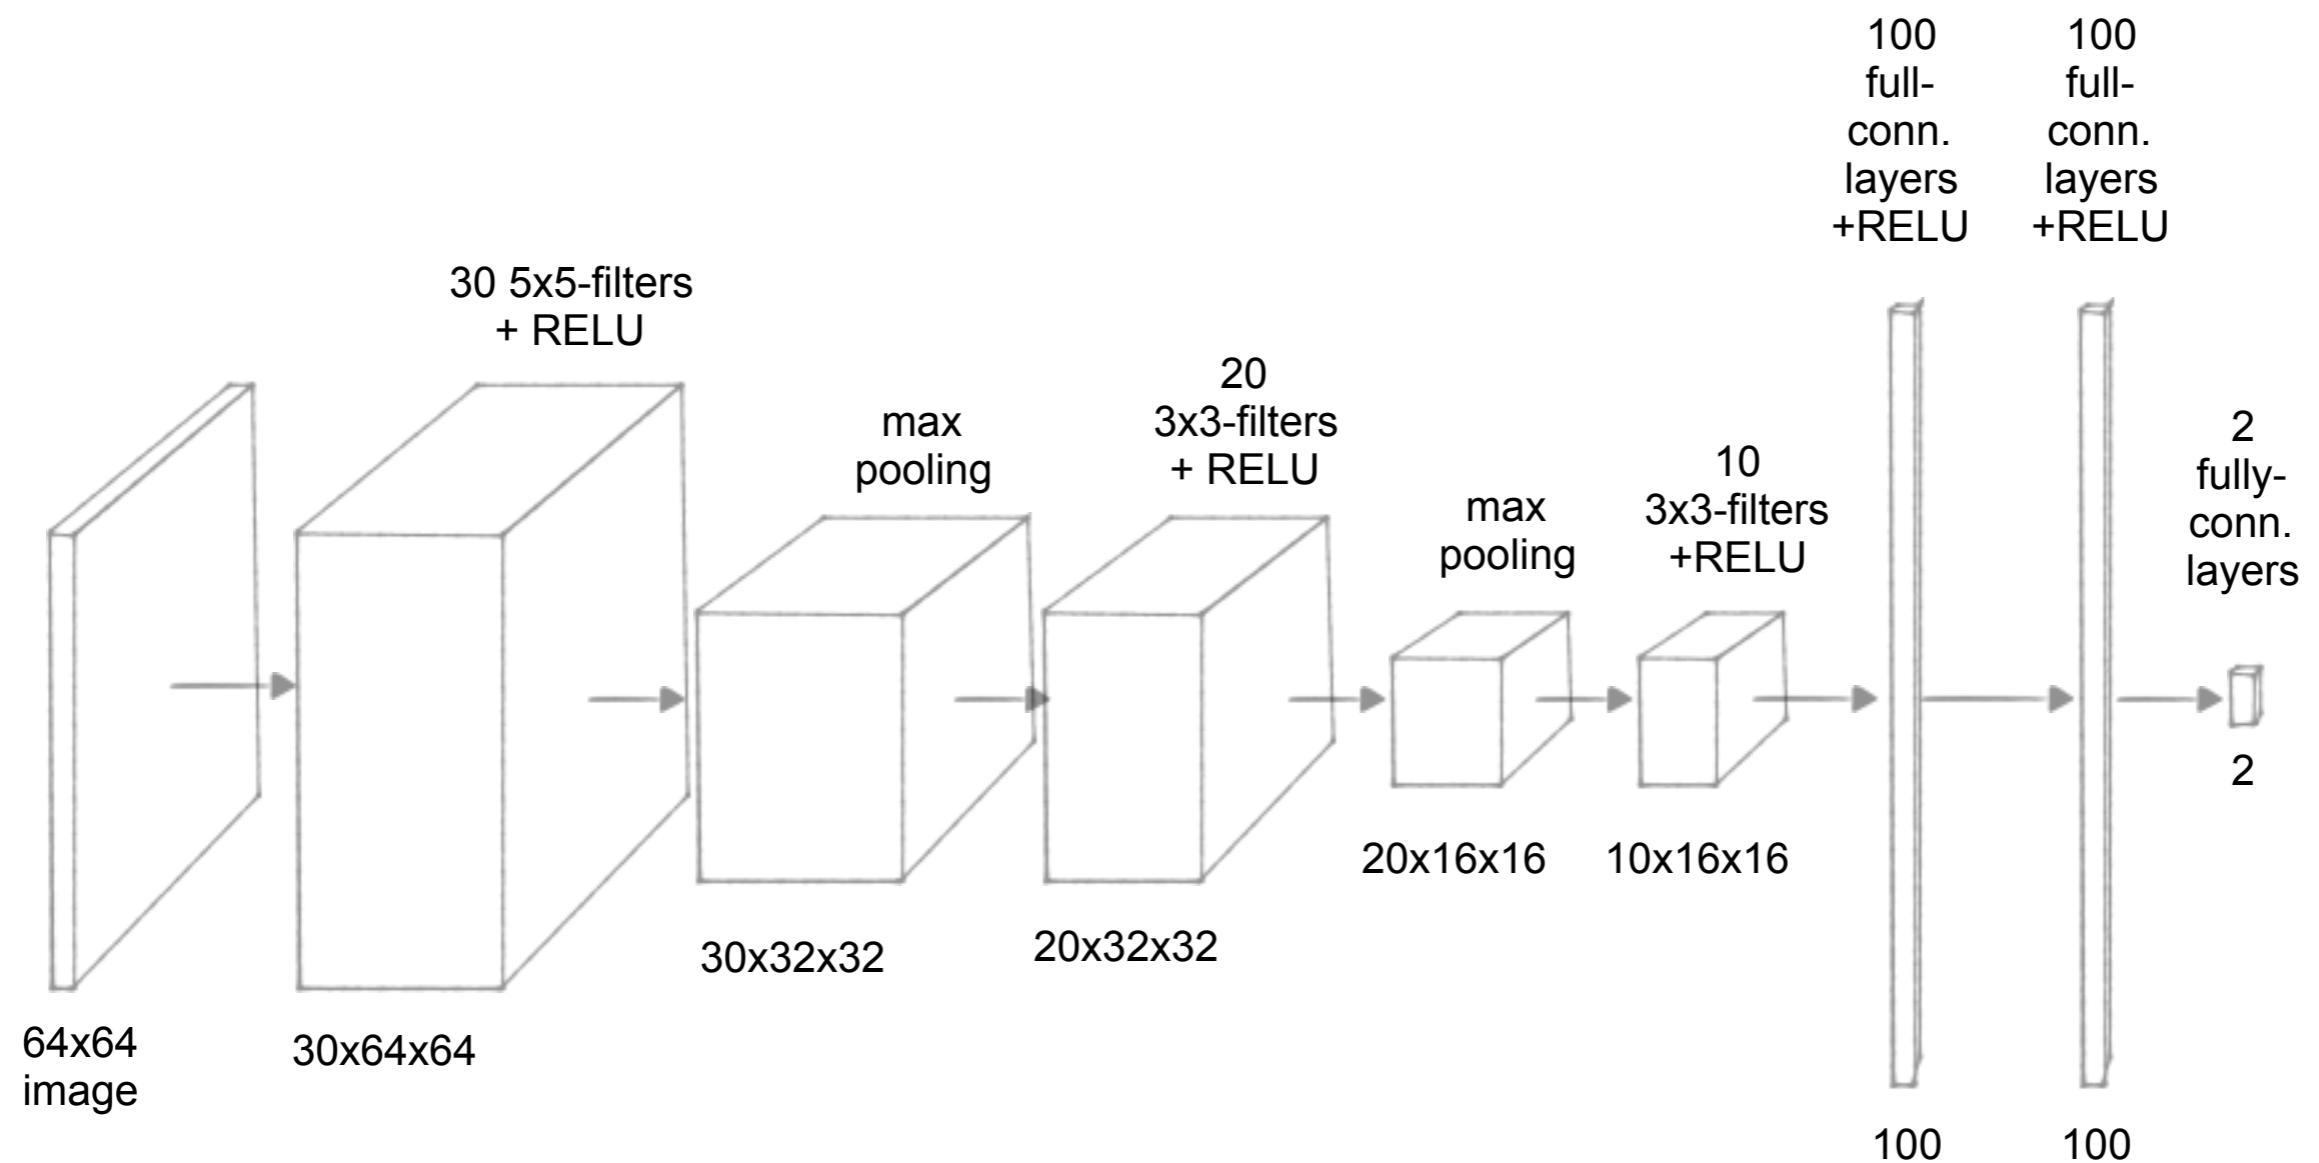
Network Structure



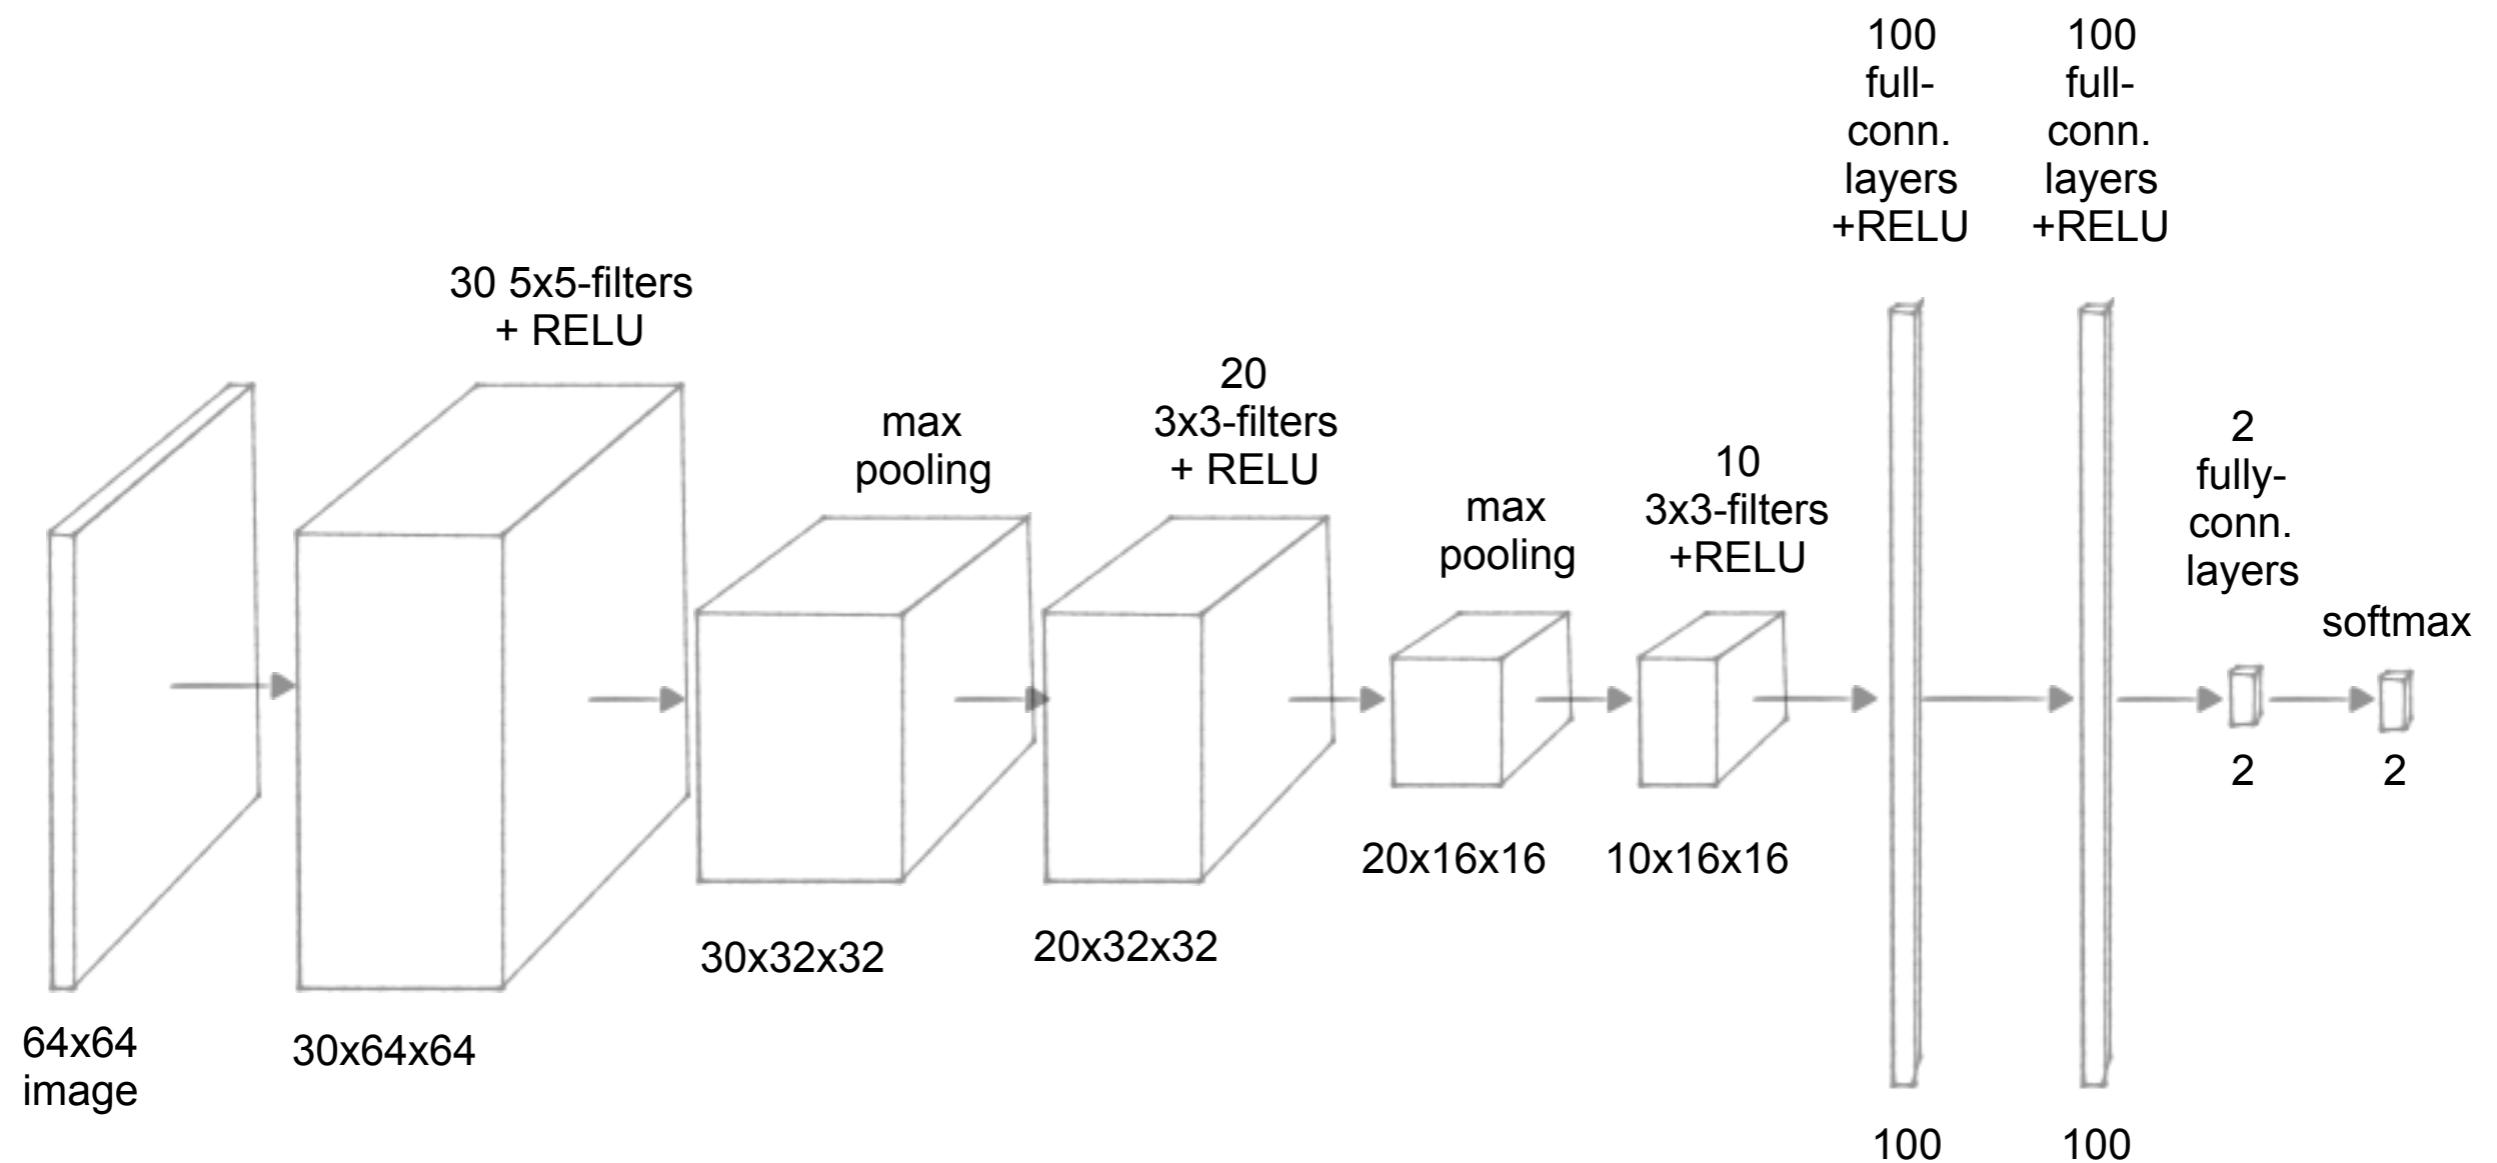
Red annotations explain the dimensionality of the 10x16x16 layer:

- A red bracket indicates a height of 100 units.
- A red bracket indicates a width of $10 \times 16 = 160$.
- A red bracket indicates a depth of $10 \times 16 \times 16$.
- An equals sign ($=$) followed by "160" indicates the total number of units in the layer.

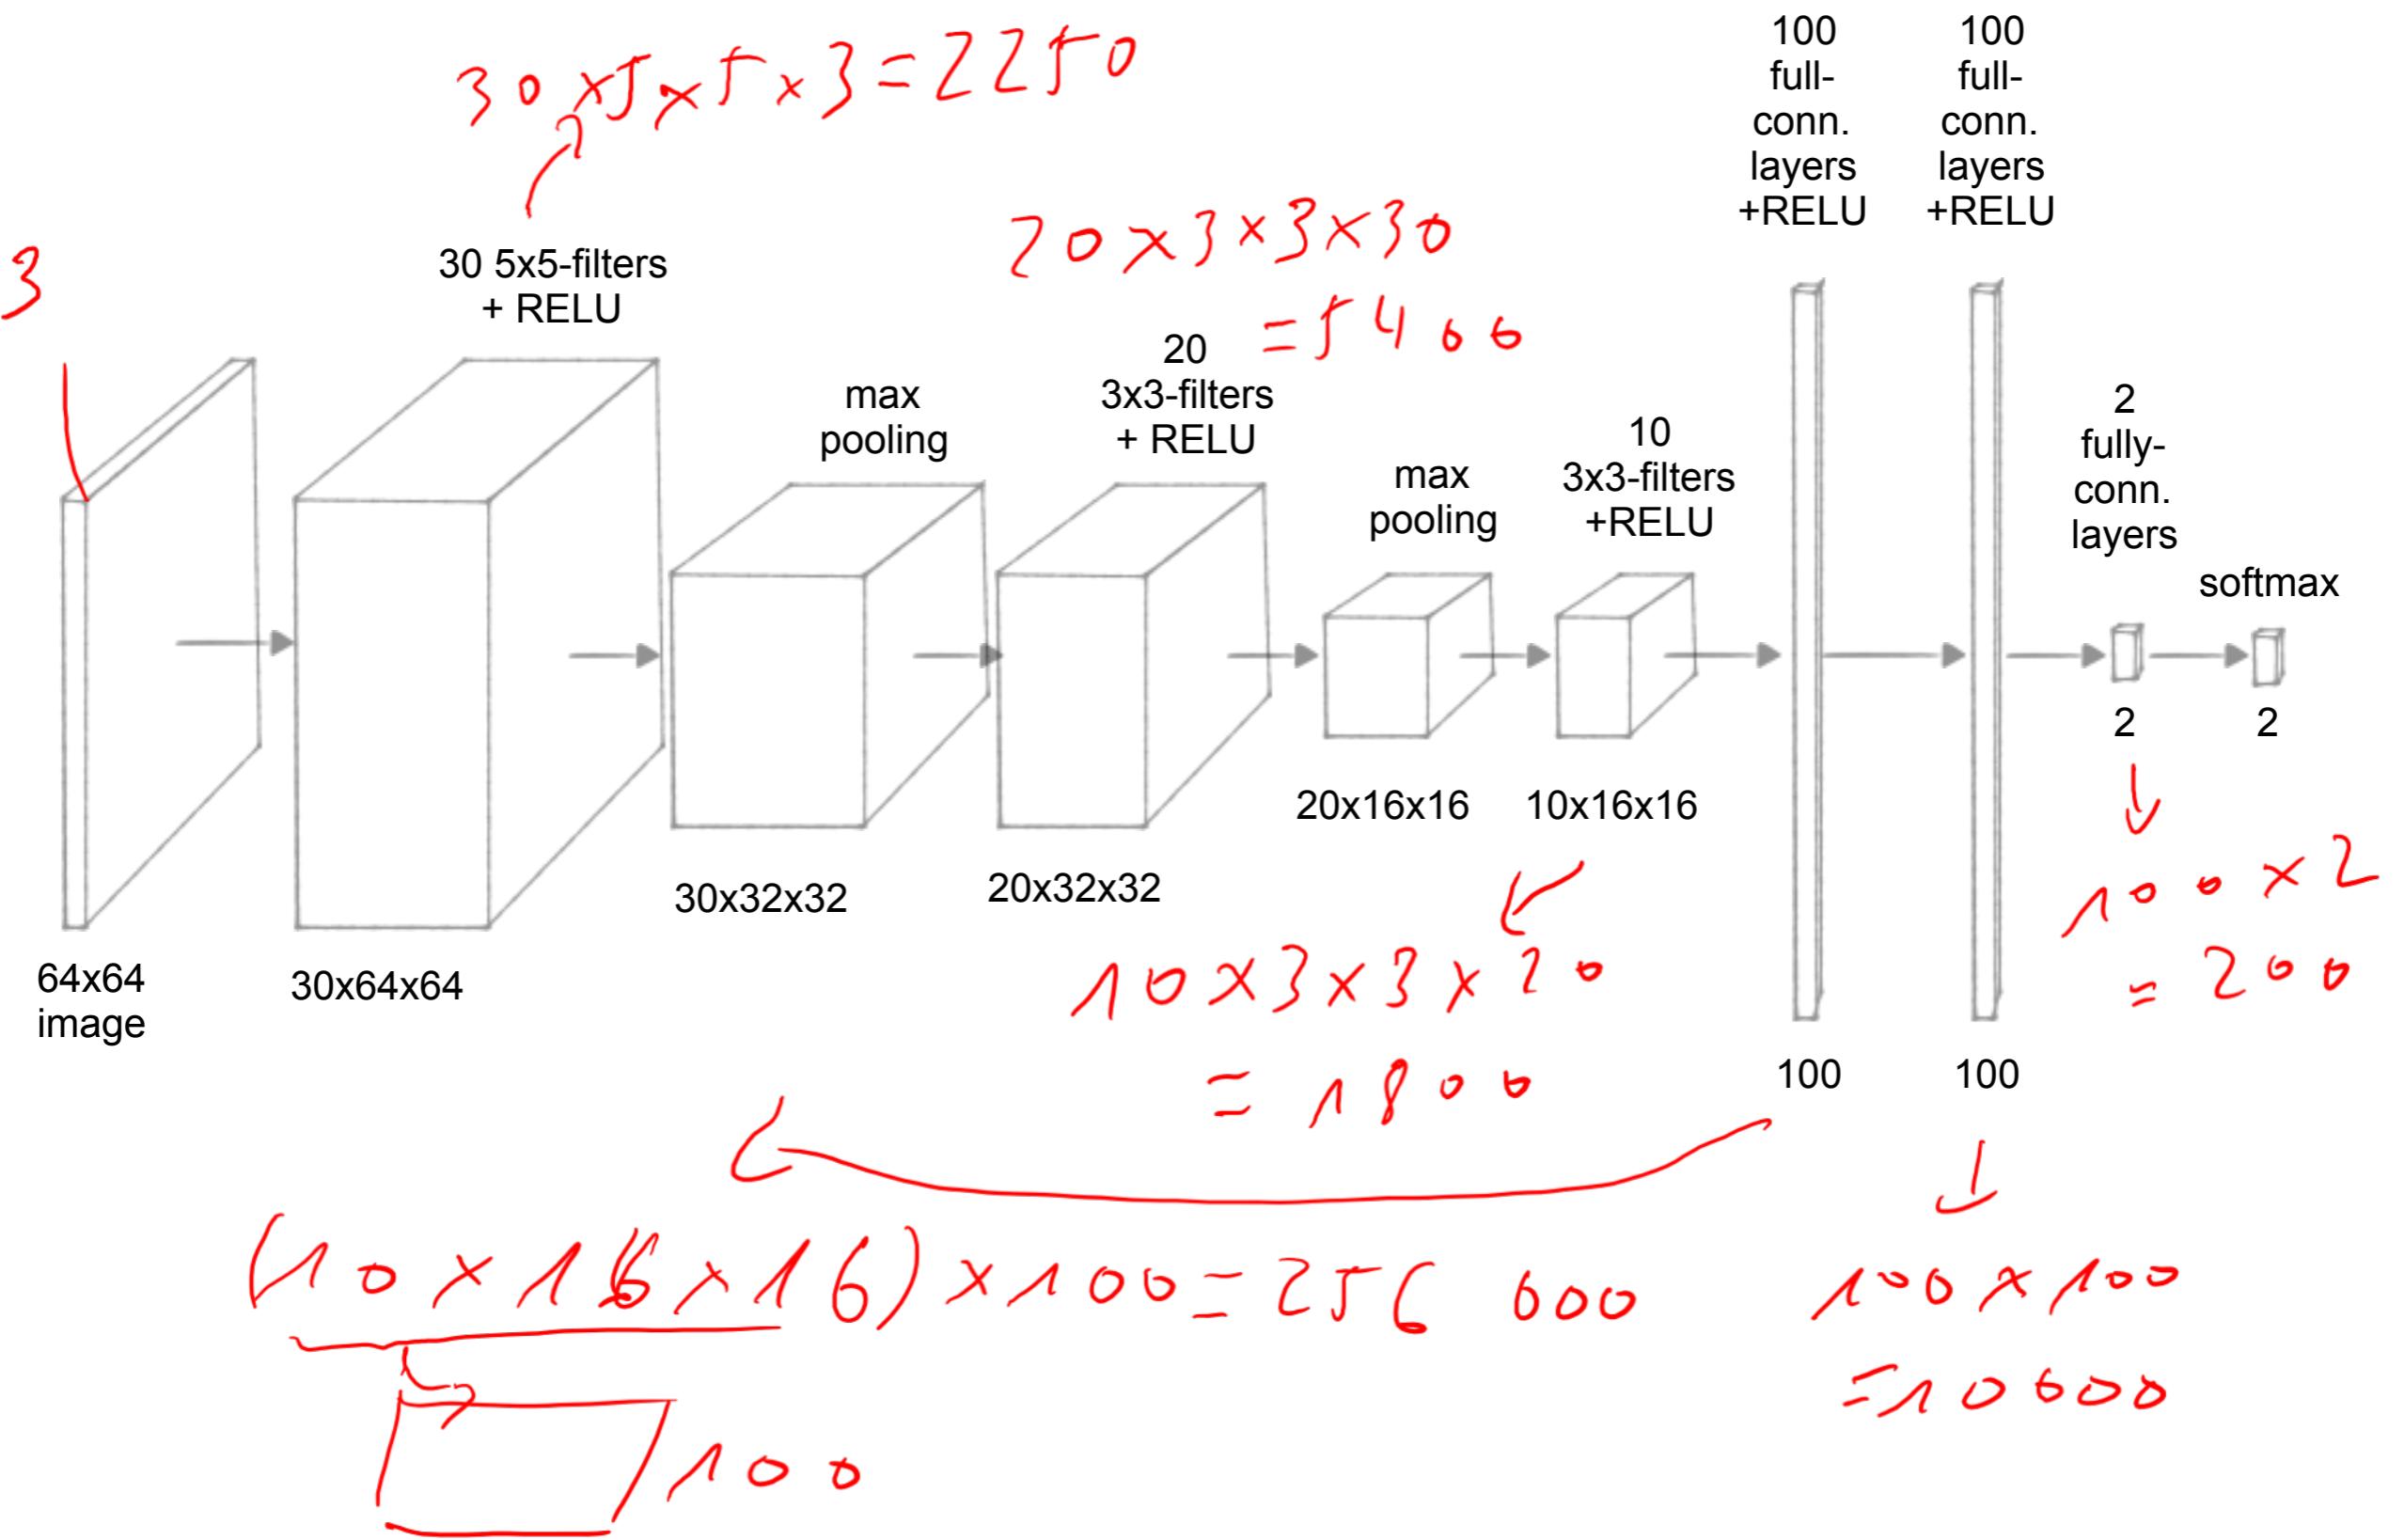
Network Structure



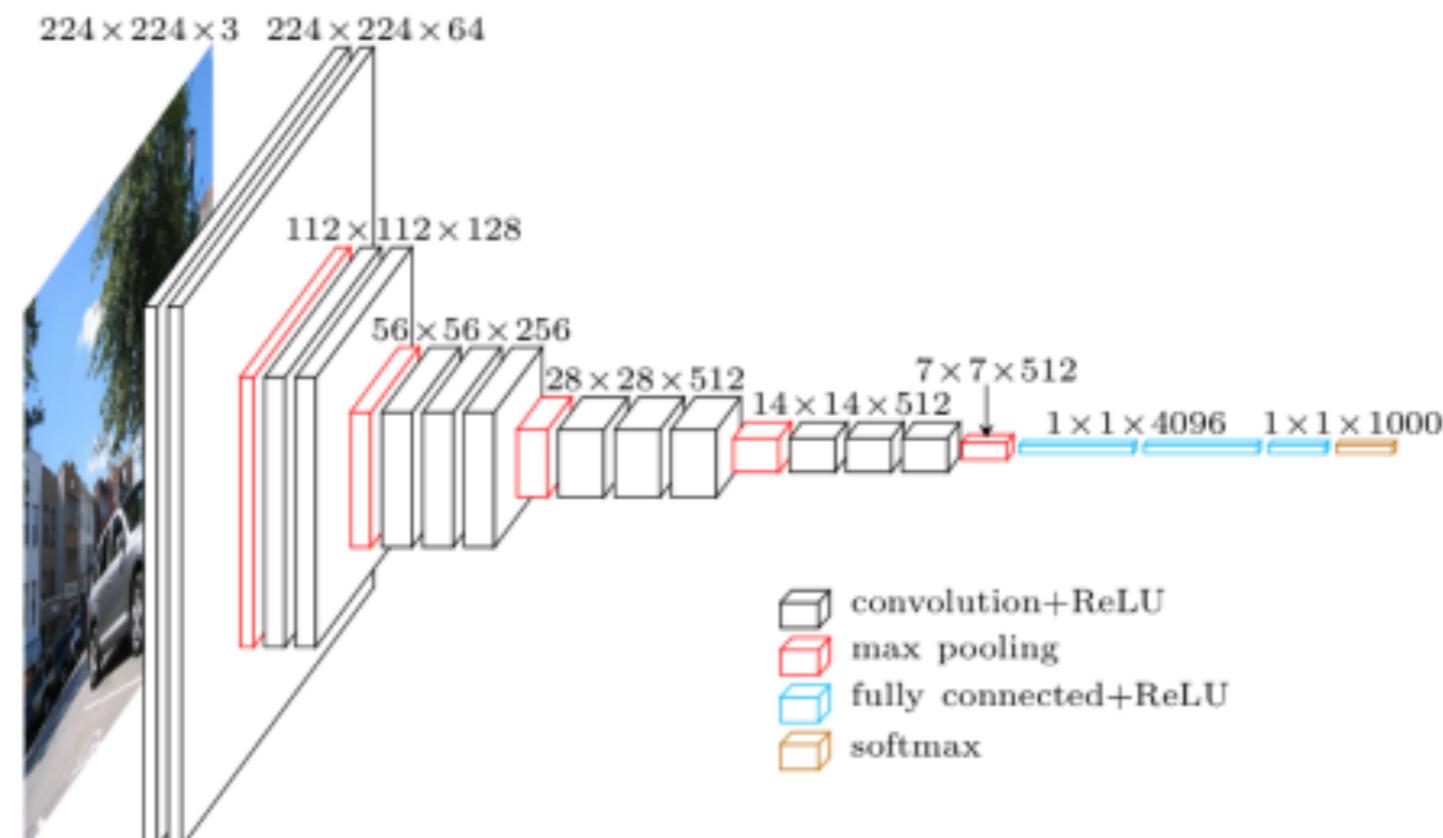
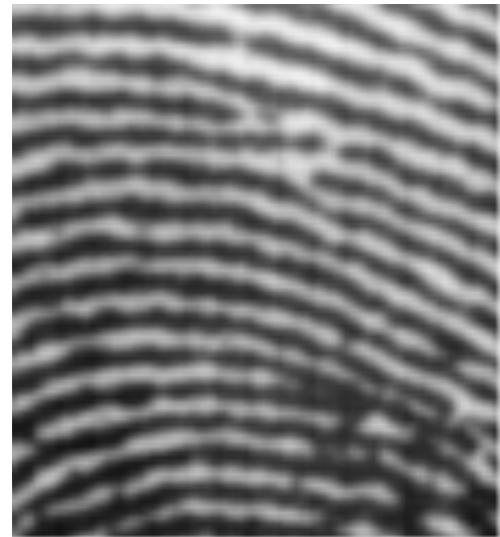
Network Structure



Number Parameters?



Live or Spoof?

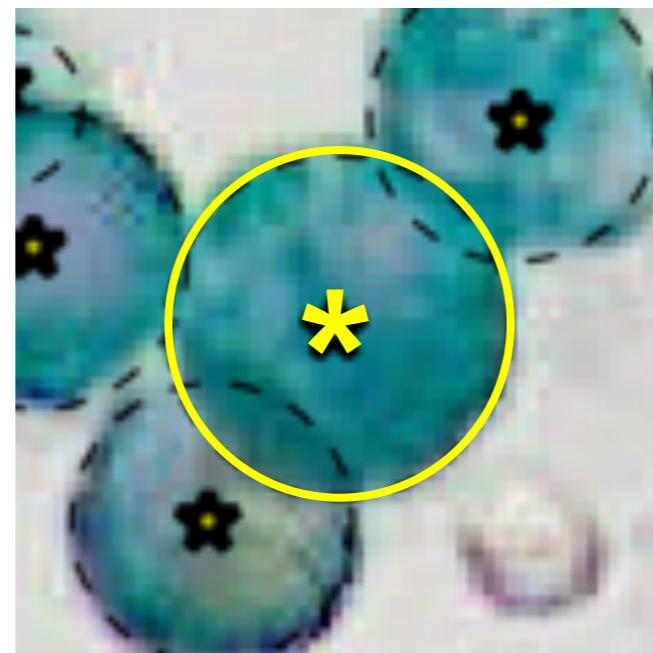


VGG-16

[Simonyan & Zisserman, Very Deep Convolutional Networks for Large-Scale Image Recognition, ICLR 2015]

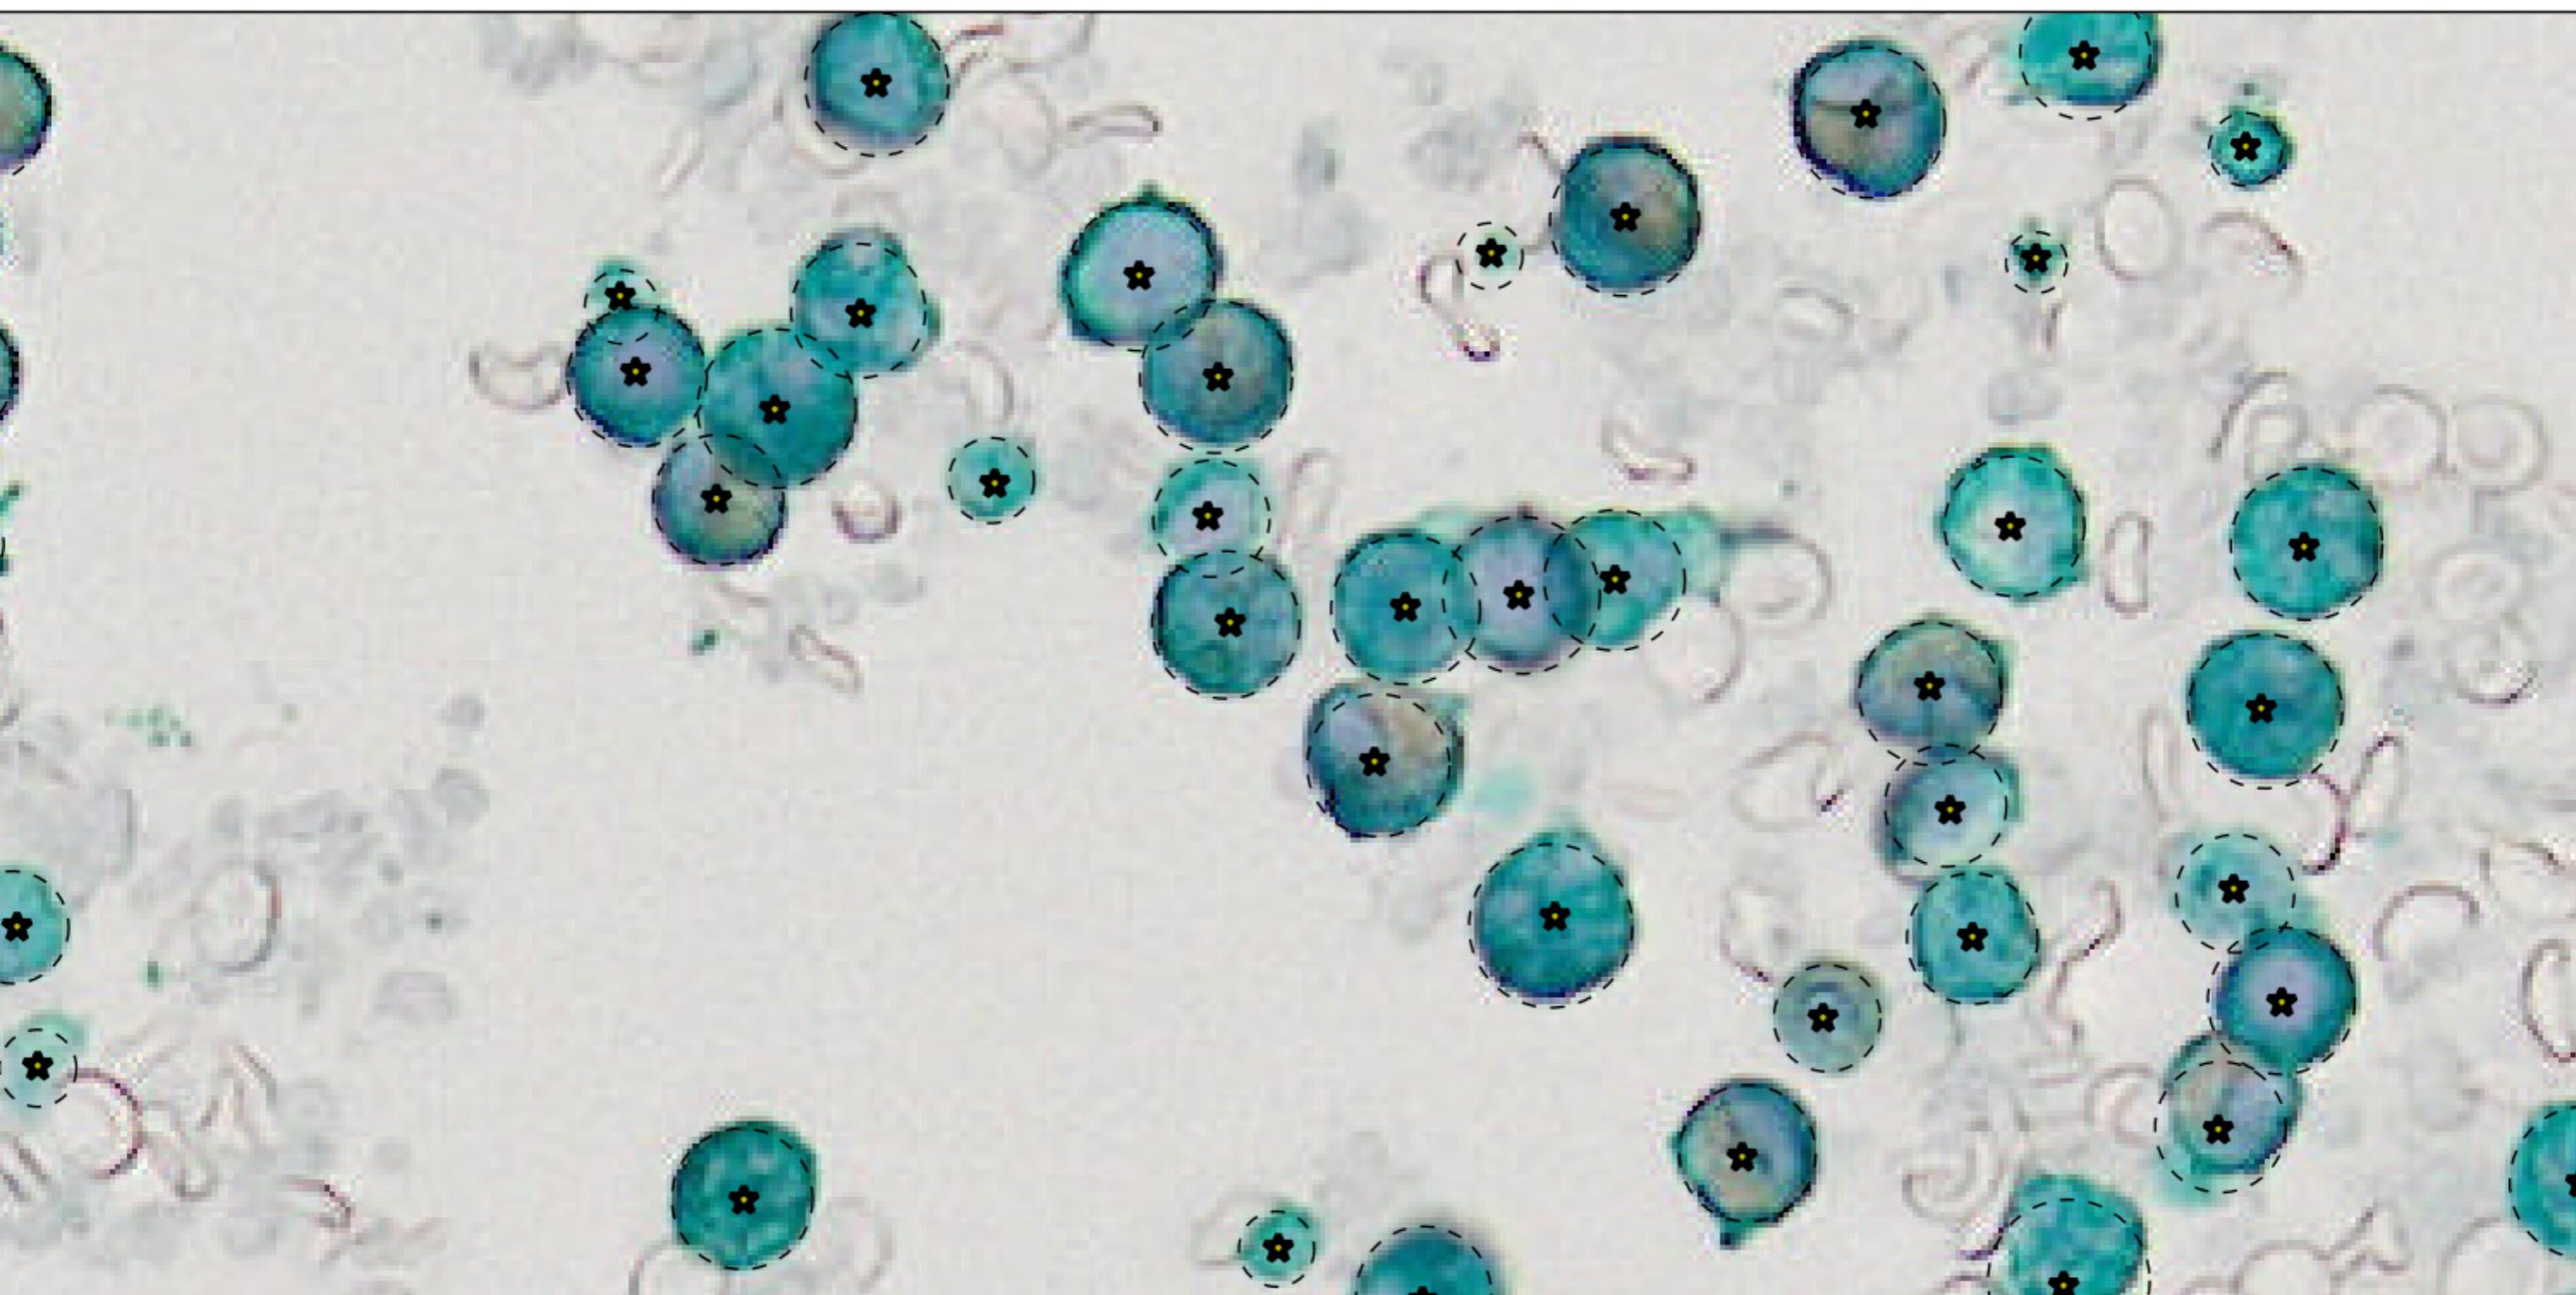
Centre and Radius

- No softmax
- Quadratic loss function

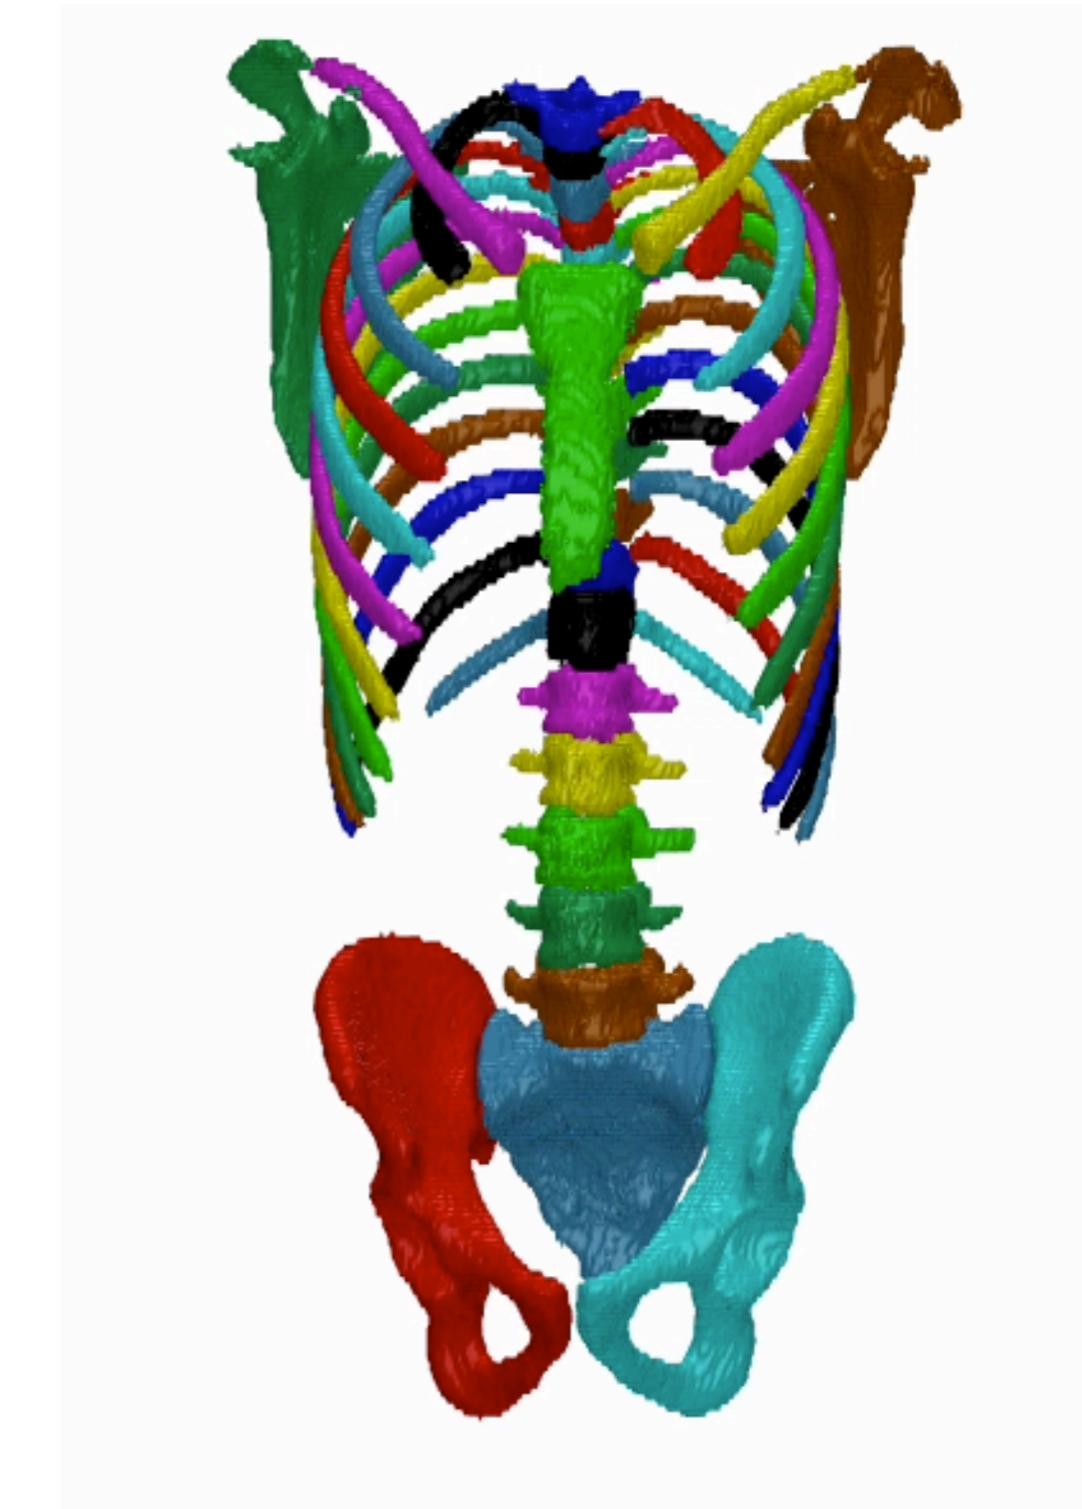
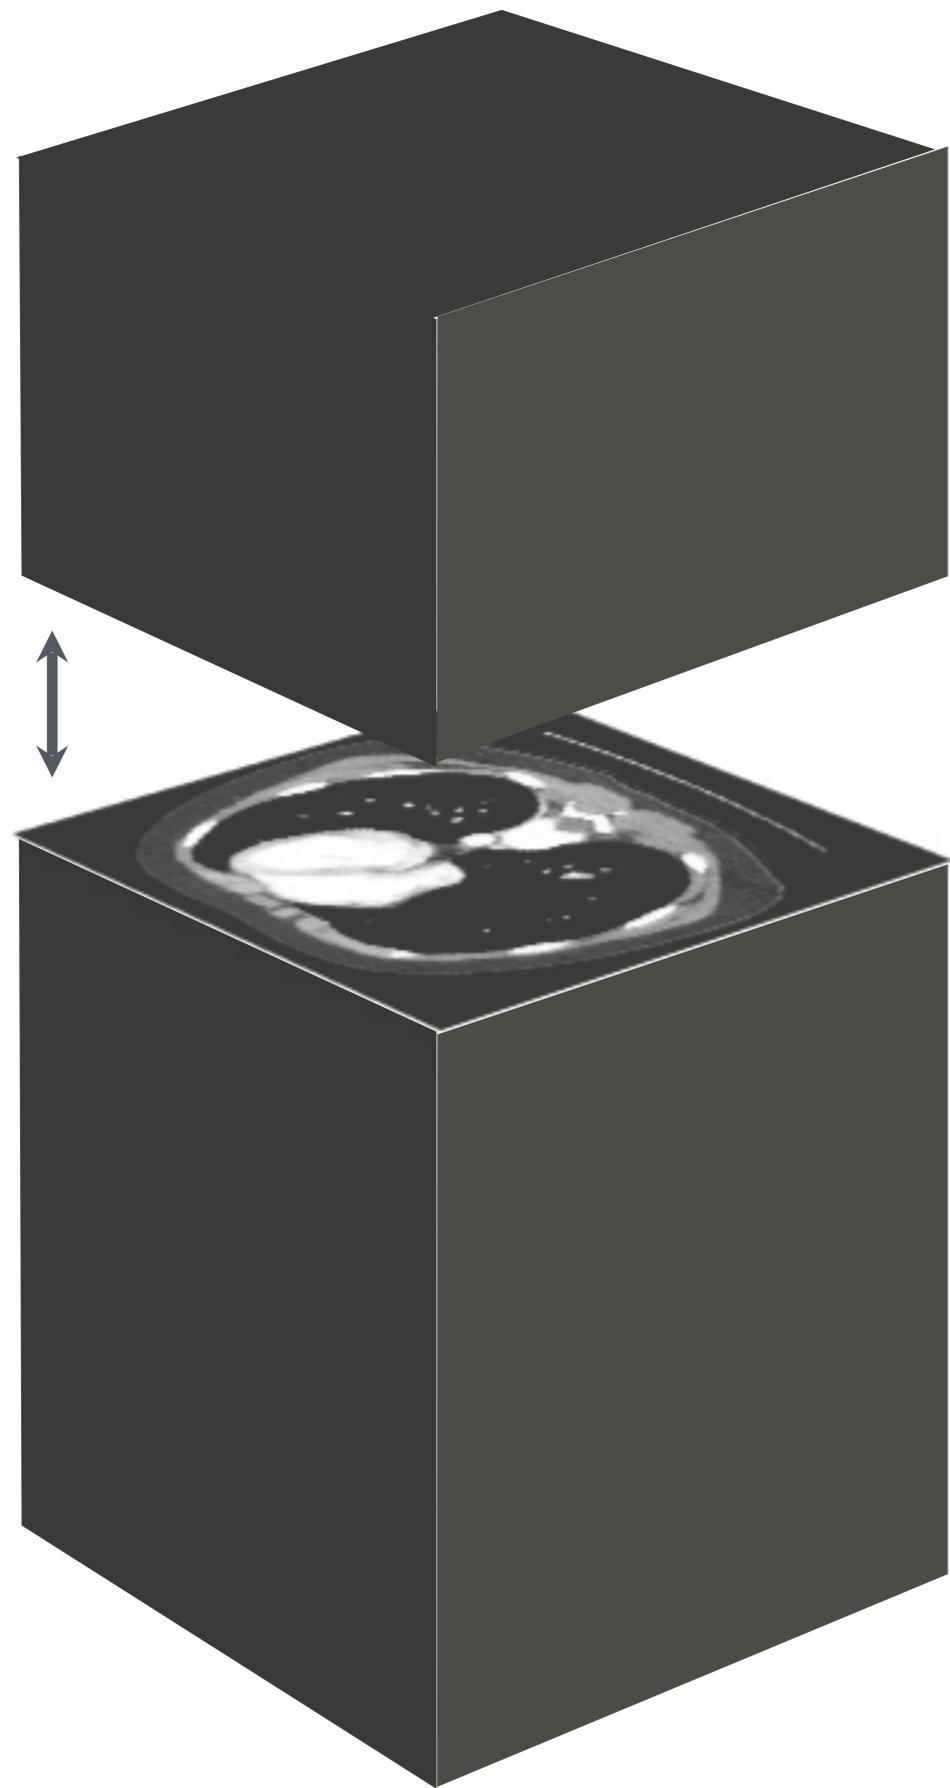


Centre and Radius

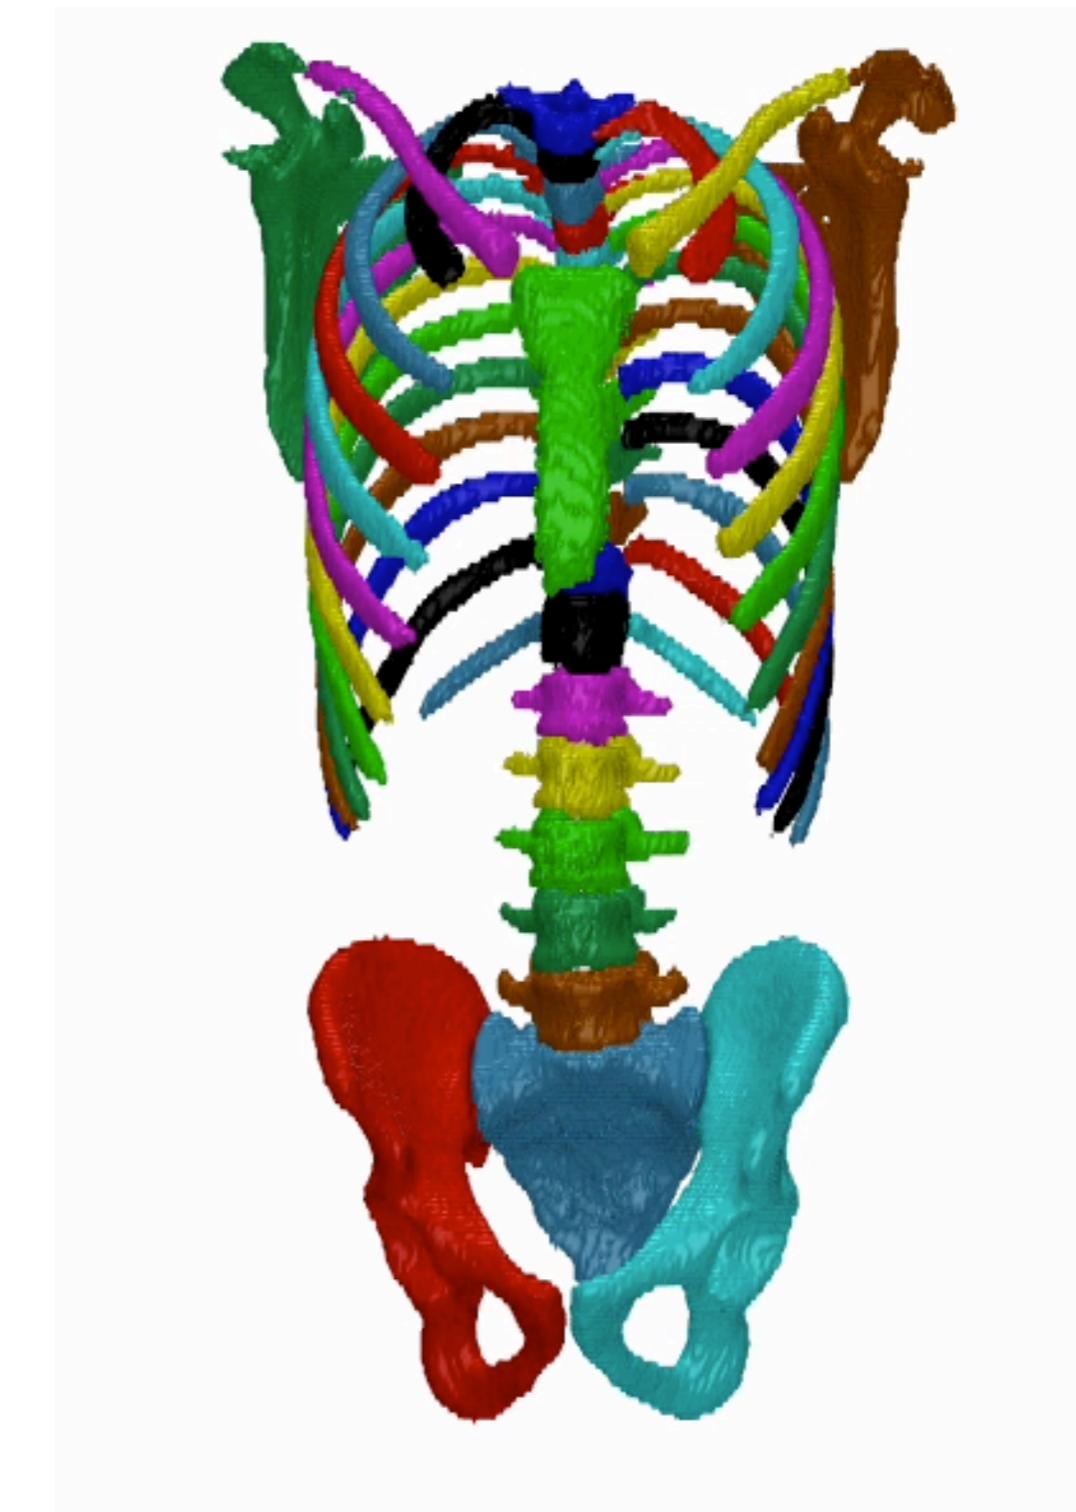
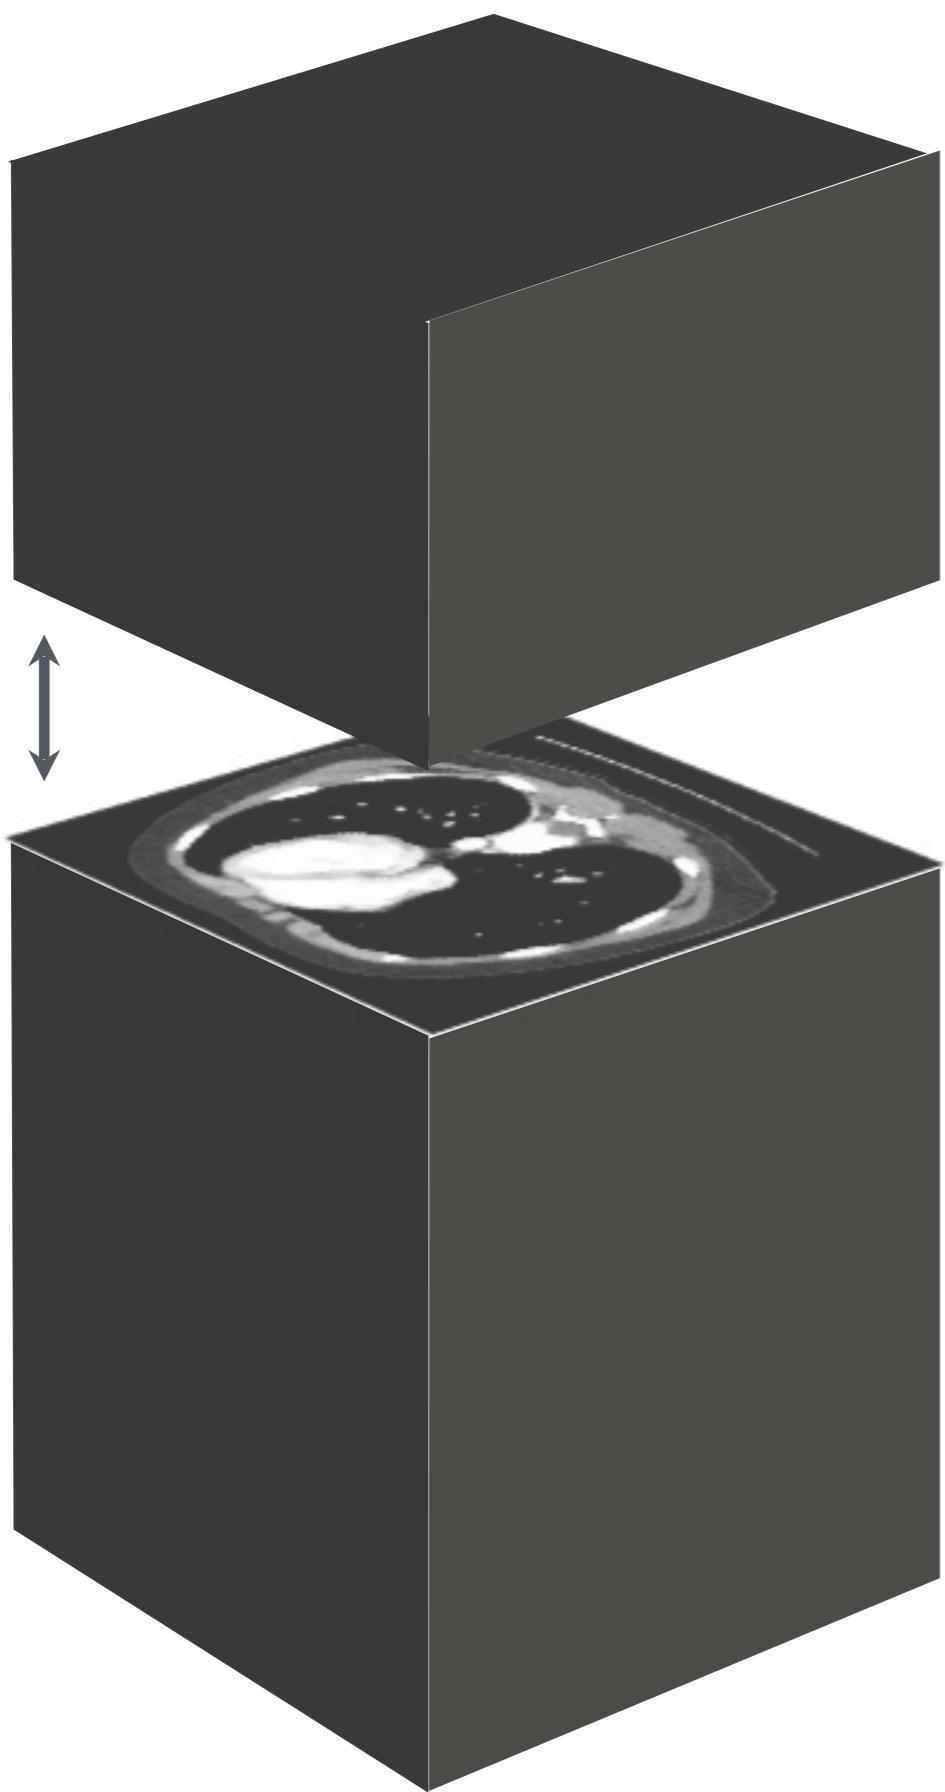
- No softmax
- Quadratic loss function



Segmentation

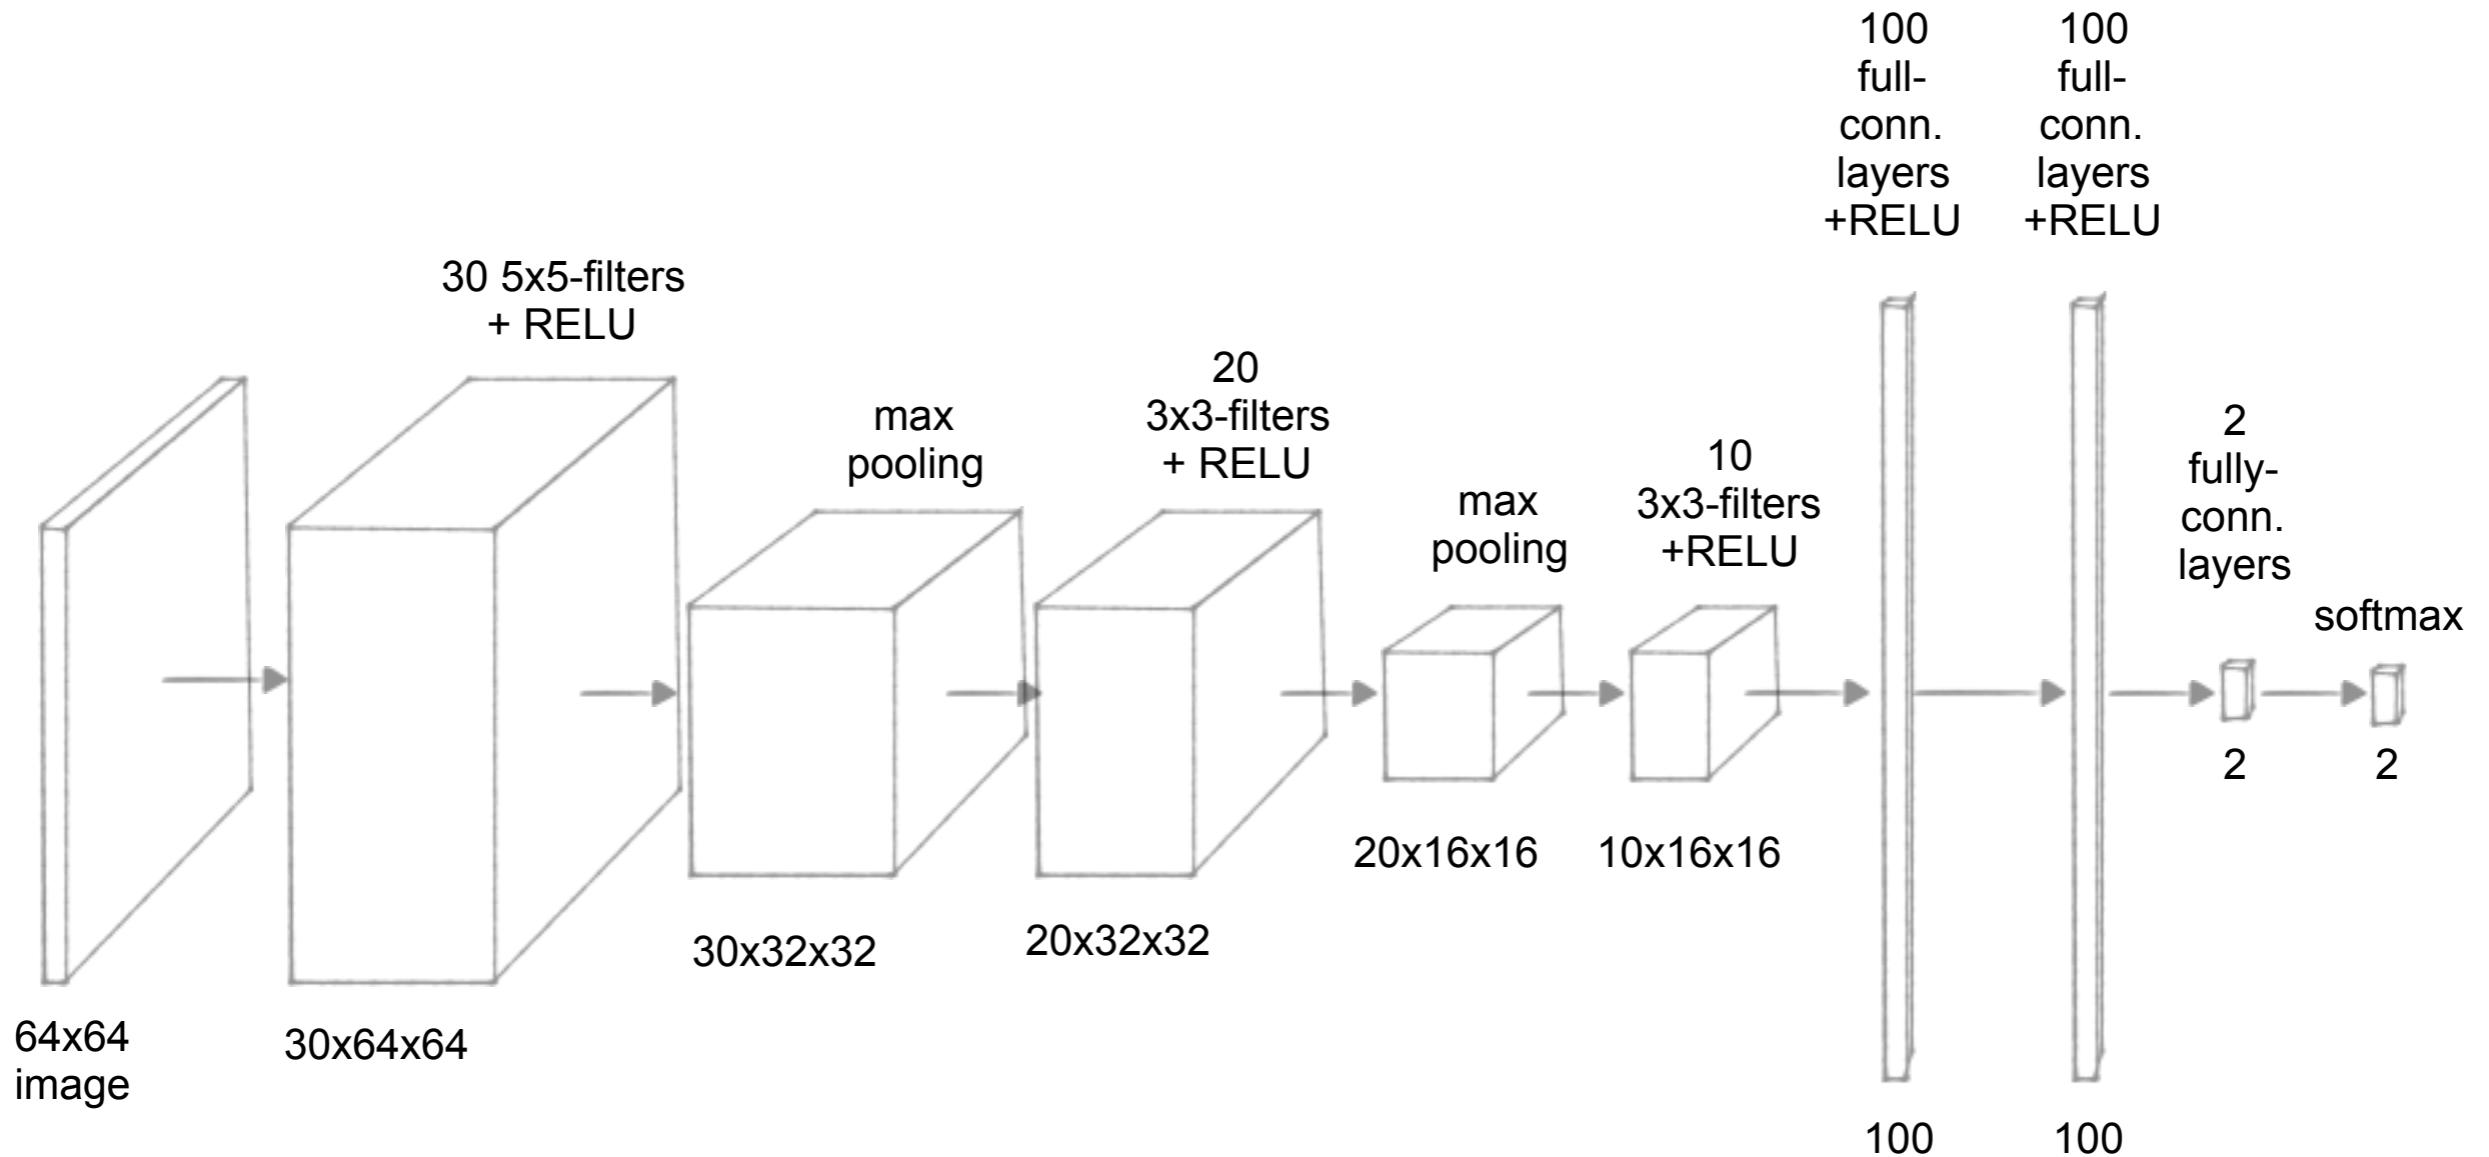


Segmentation



Fully Convolutional Neural Networks

Convolutional Neural Network for Classification

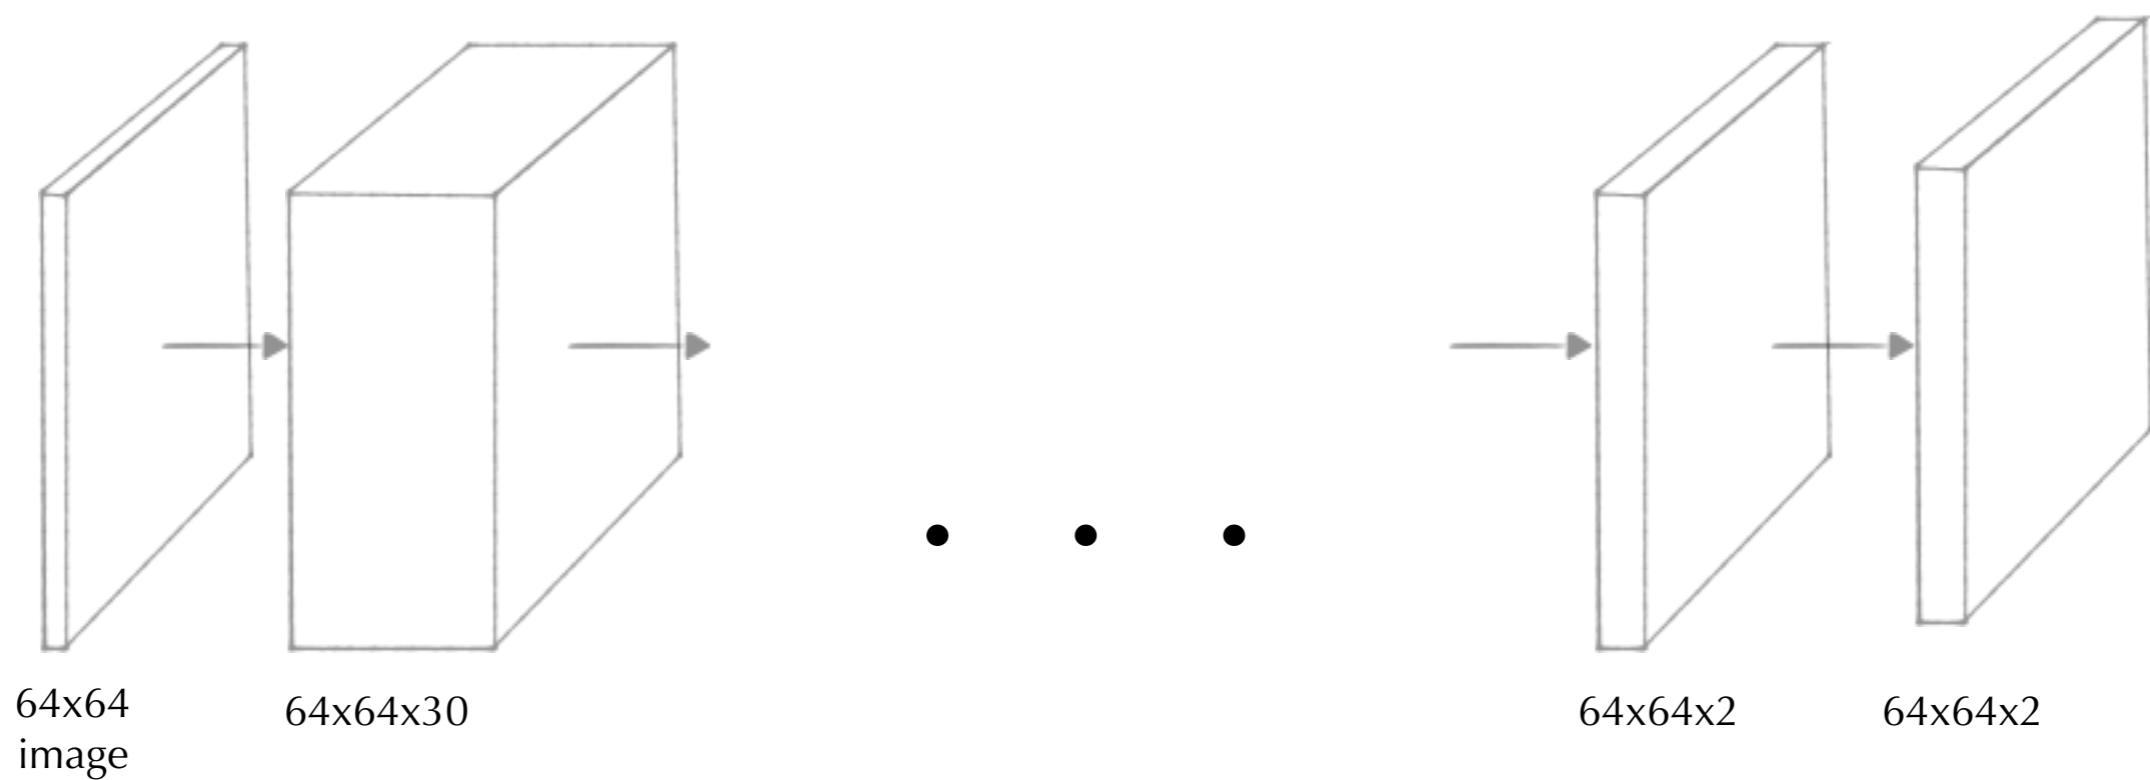


How to get a per-pixel decision?

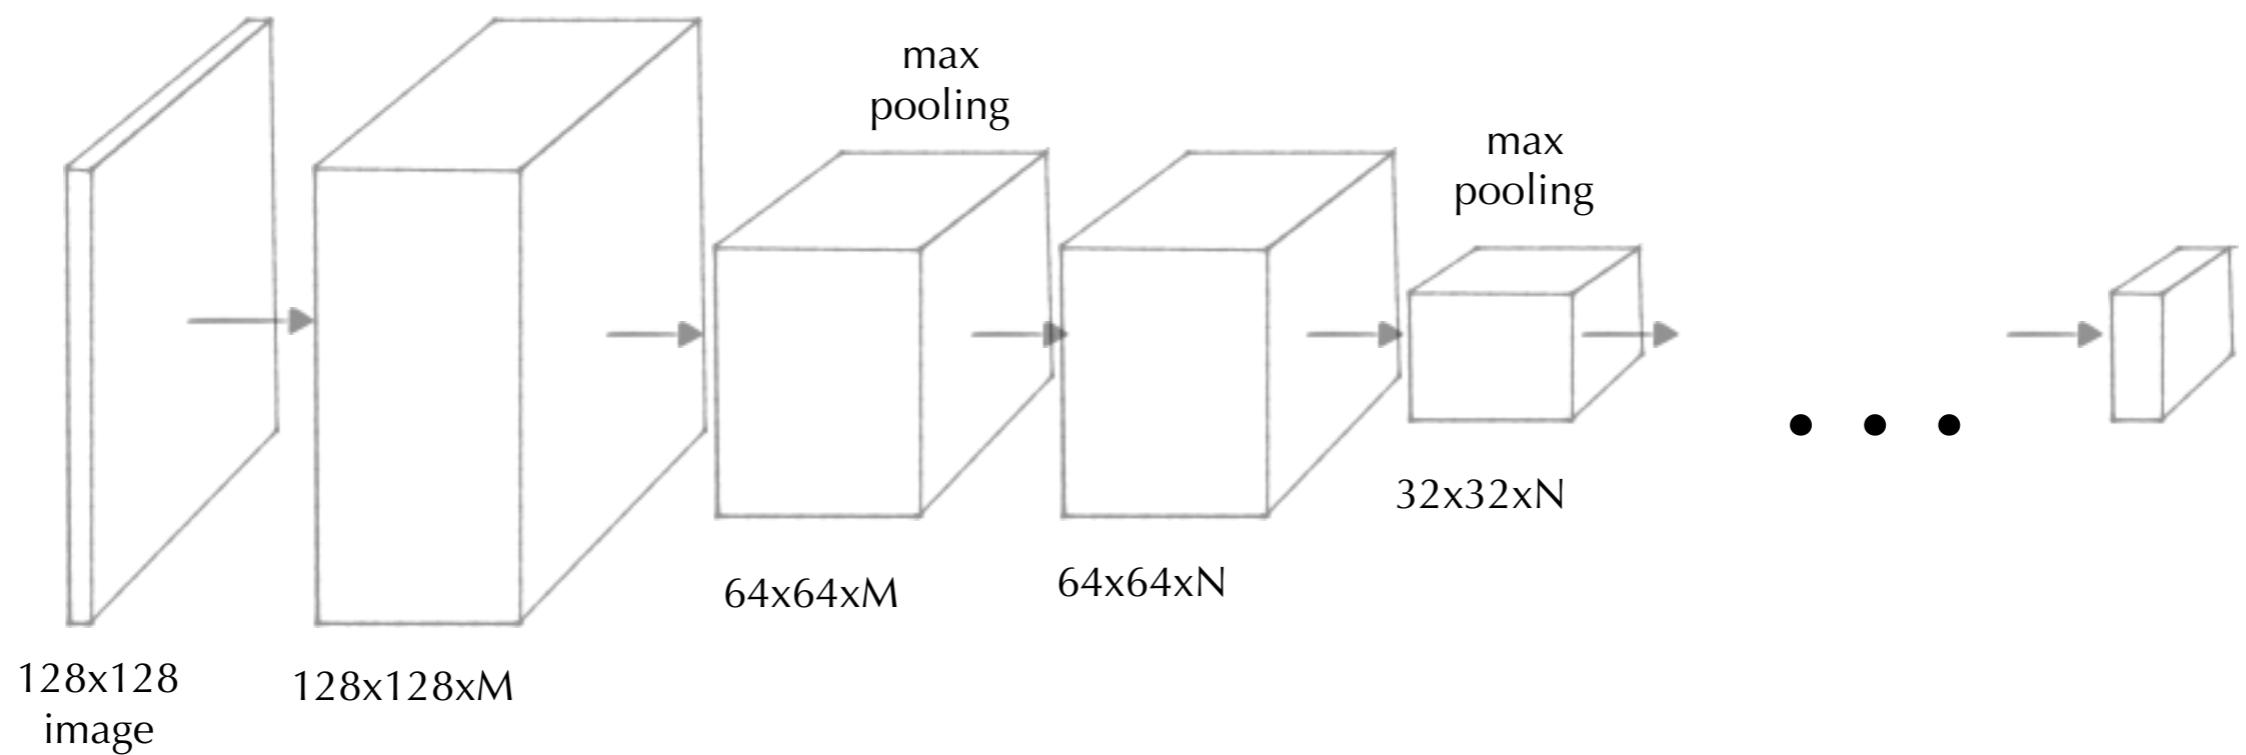
Human Pose Estimation



Fully-Convolutional Networks

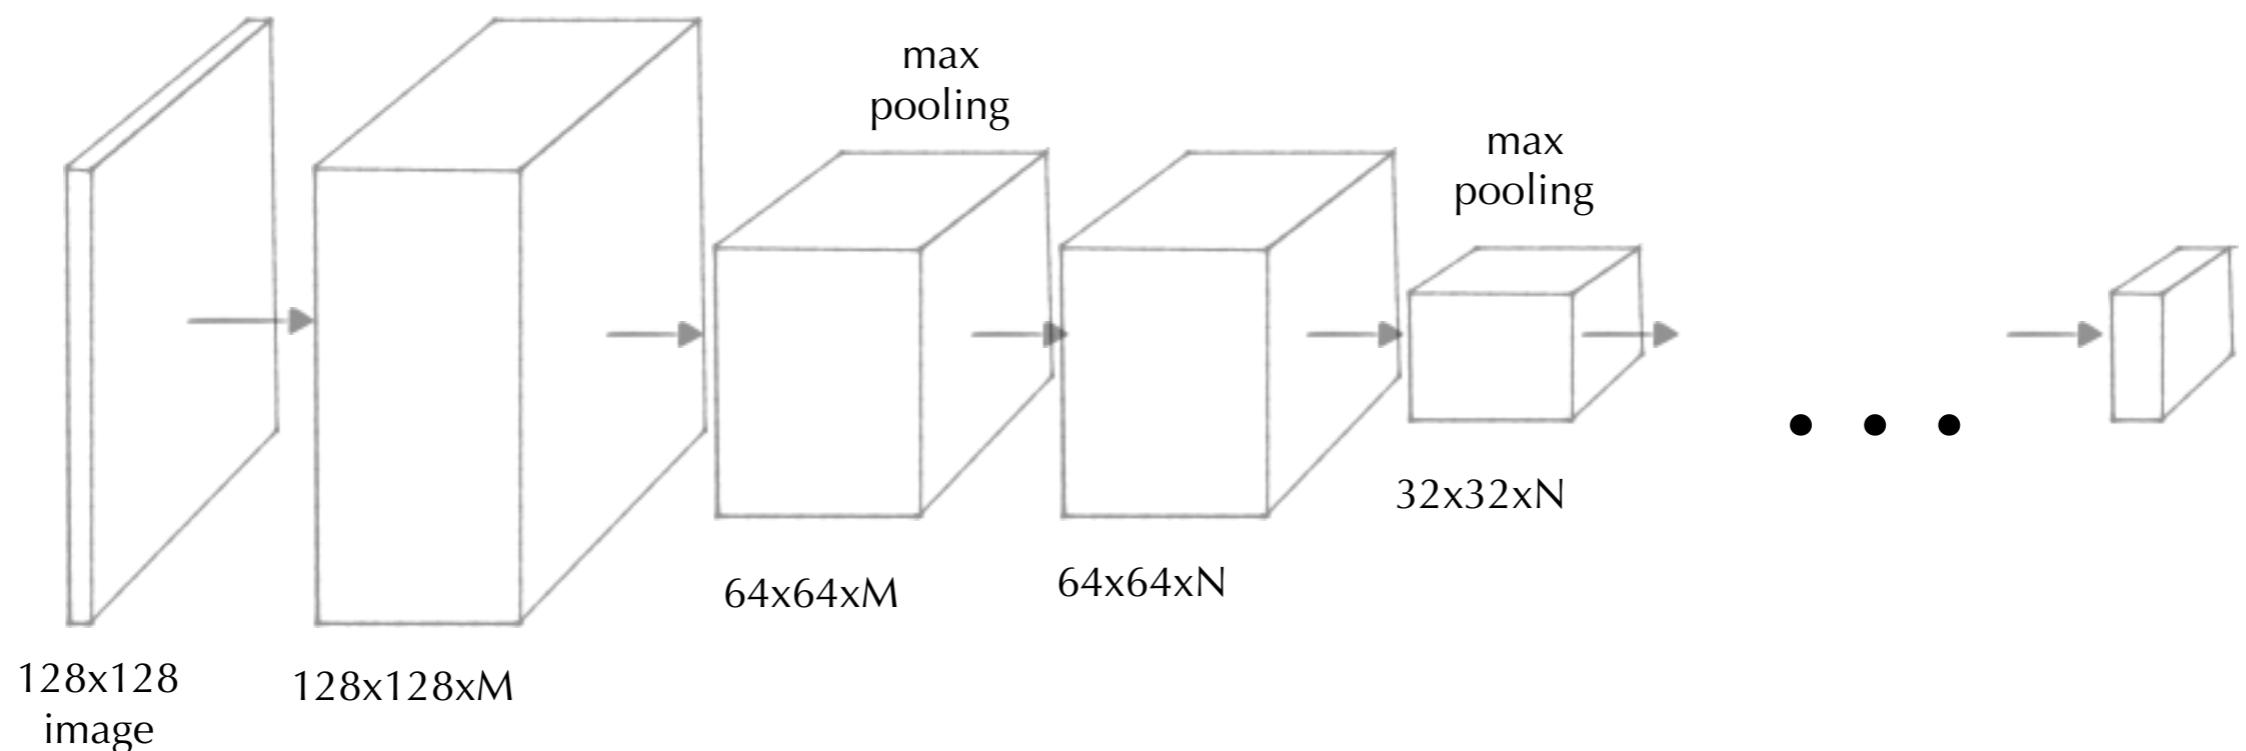


Fully-Convolutional Network for Detection



Fully-Convolutional Network for Detection

filters 32 64 128 256



Fully-Convolutional Network for Detection

filters 32 64 128 256

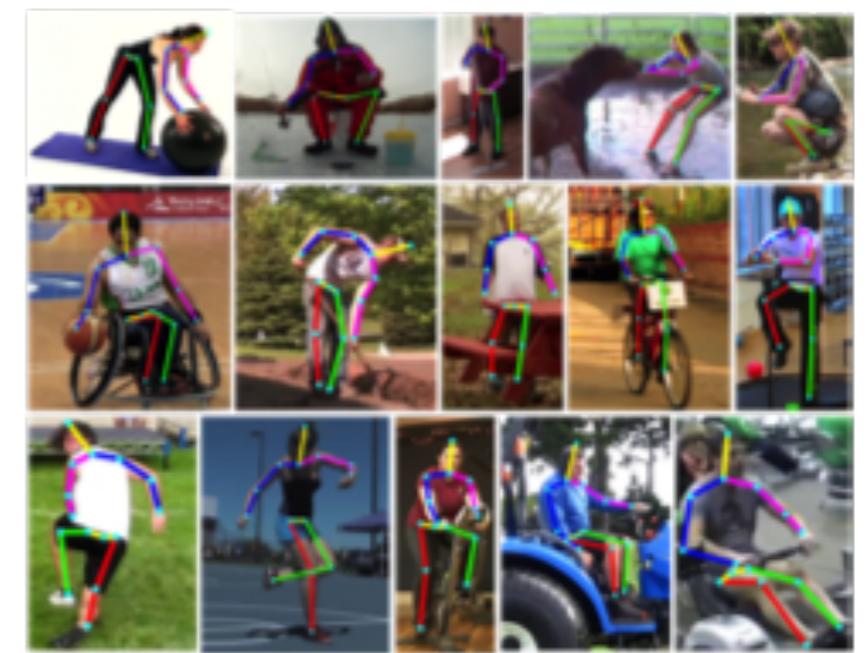
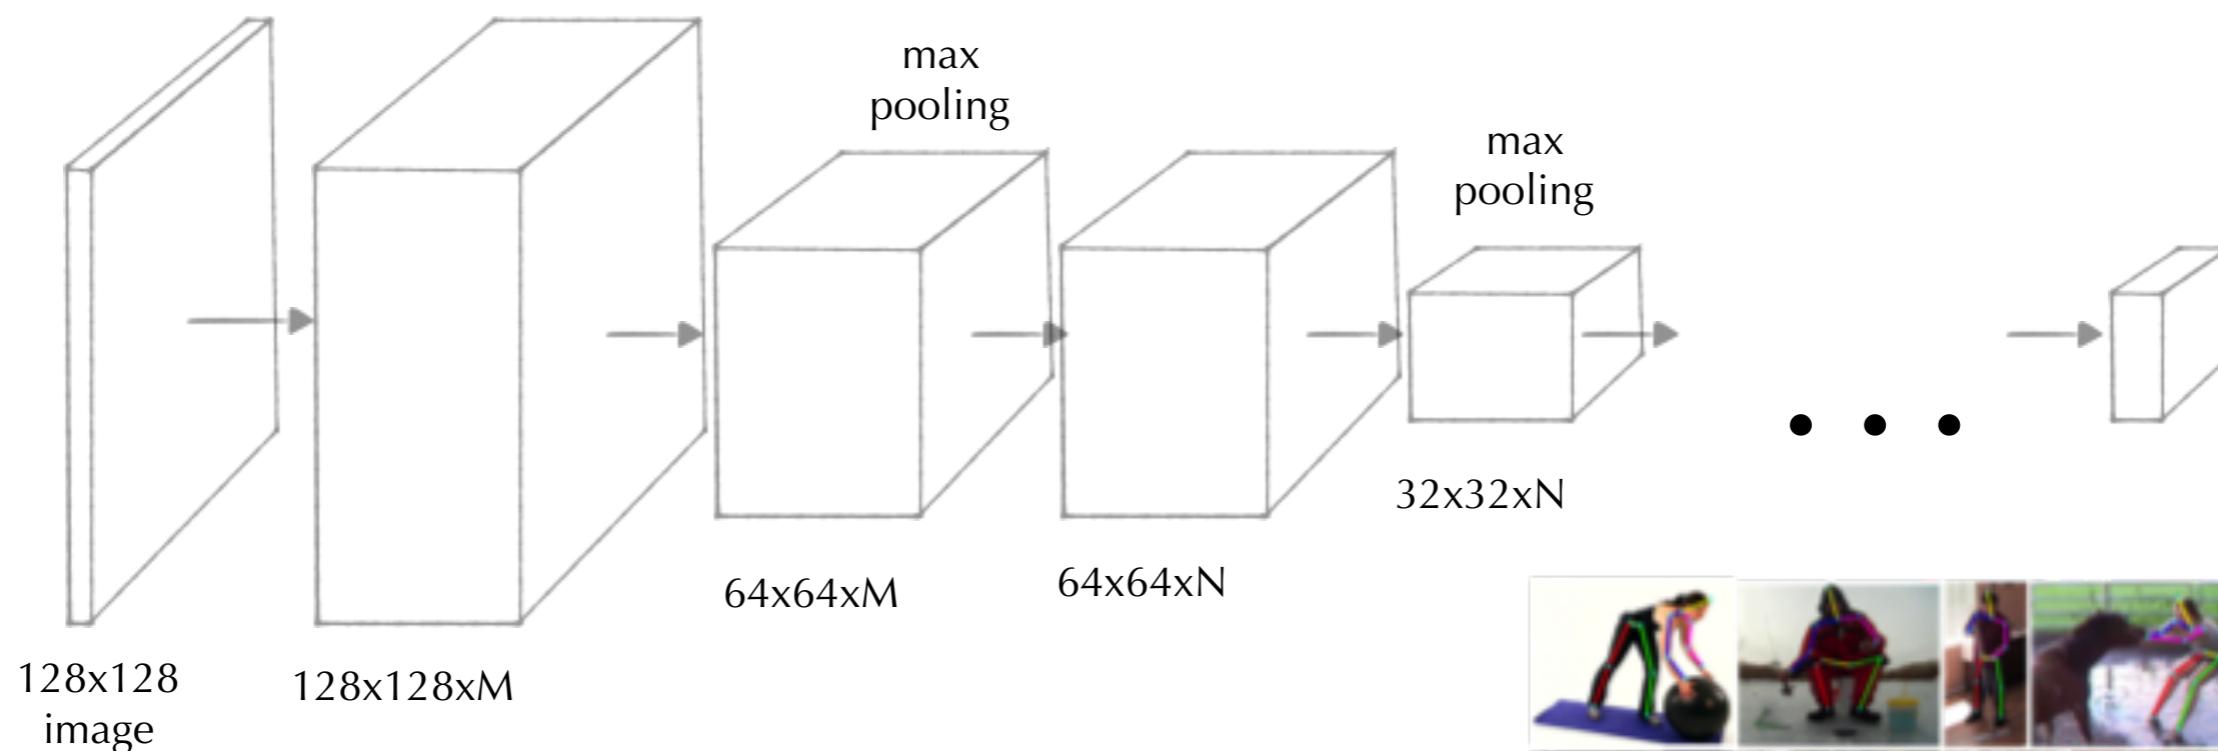
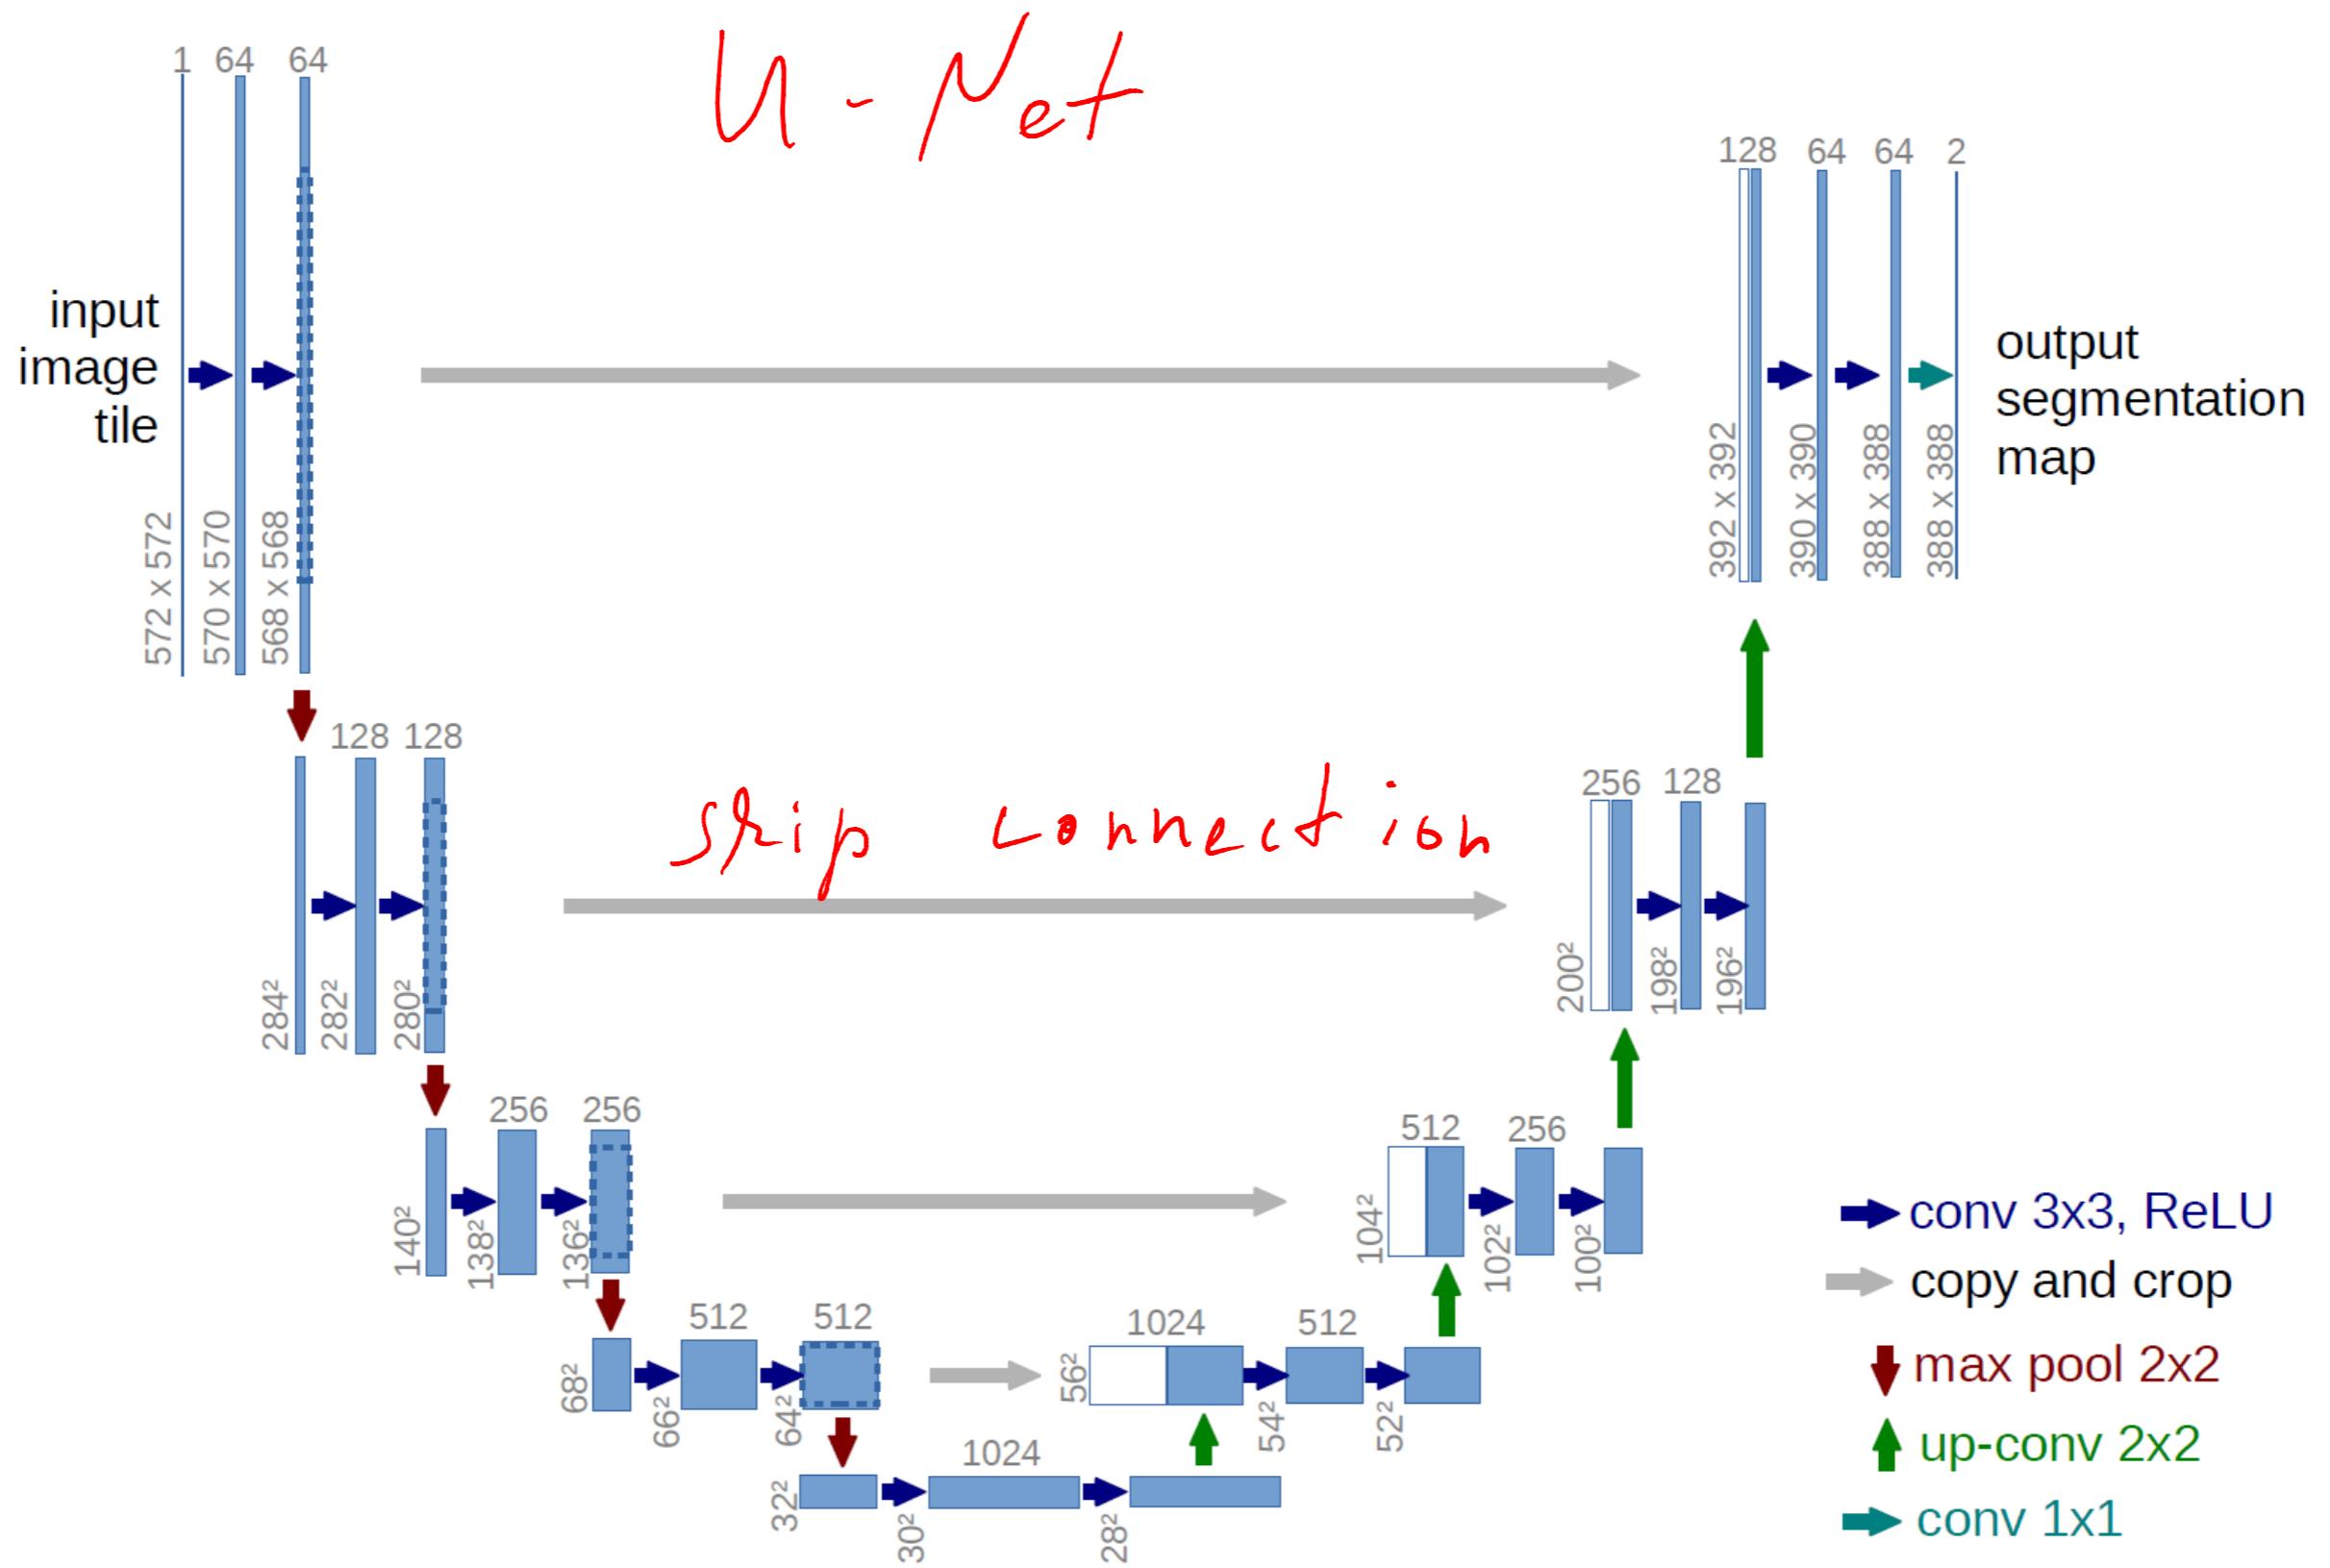


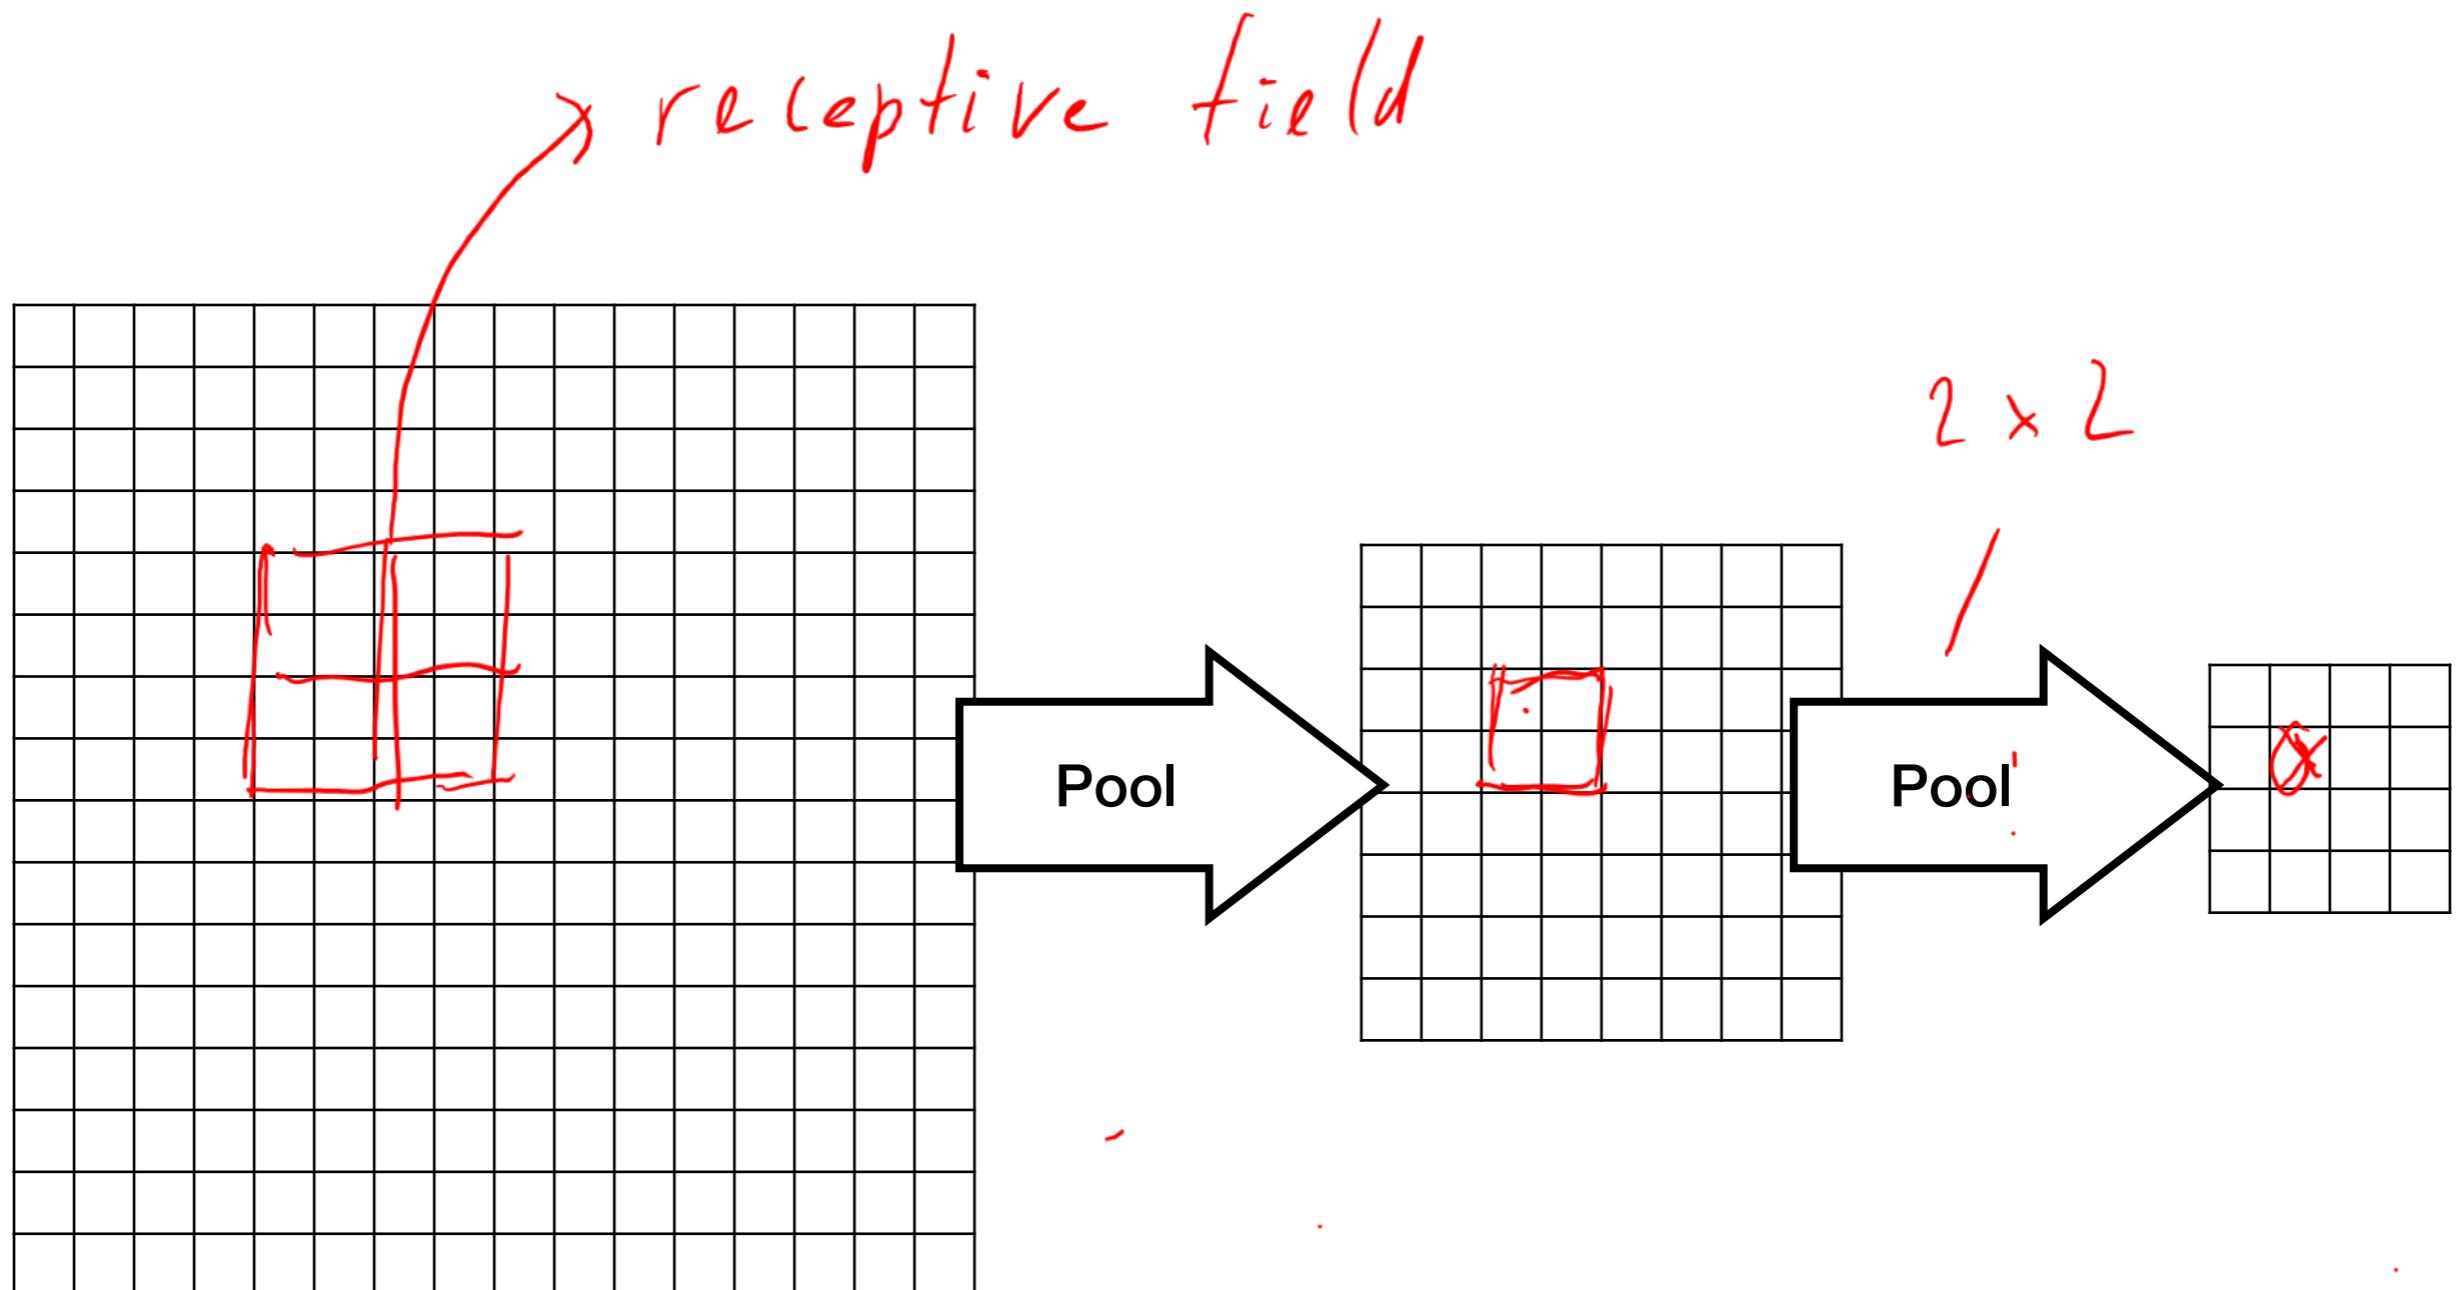
Image Segmentation



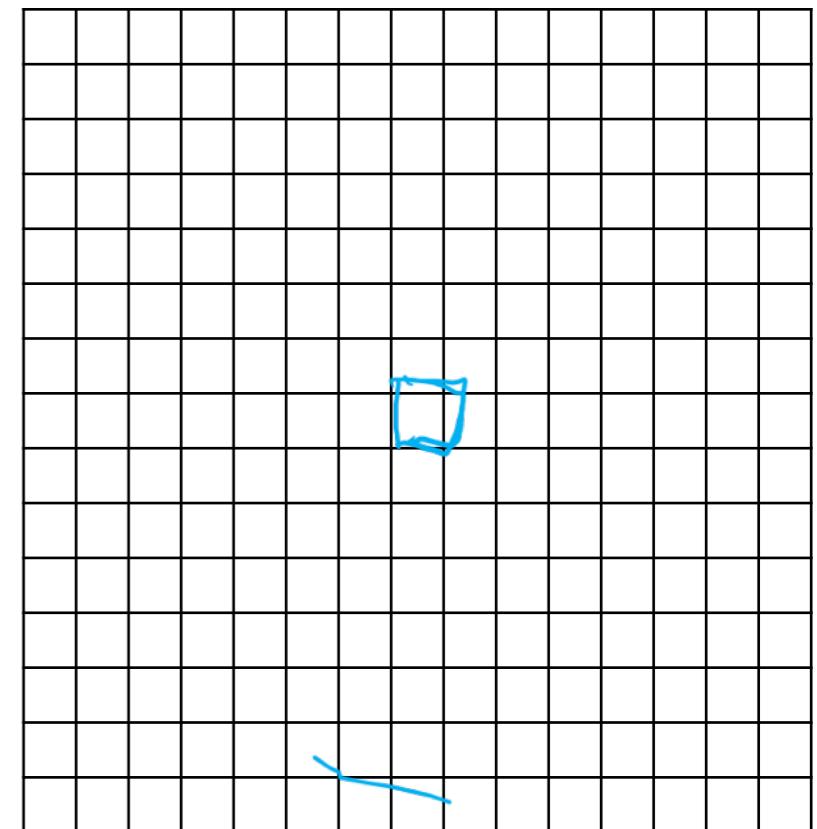
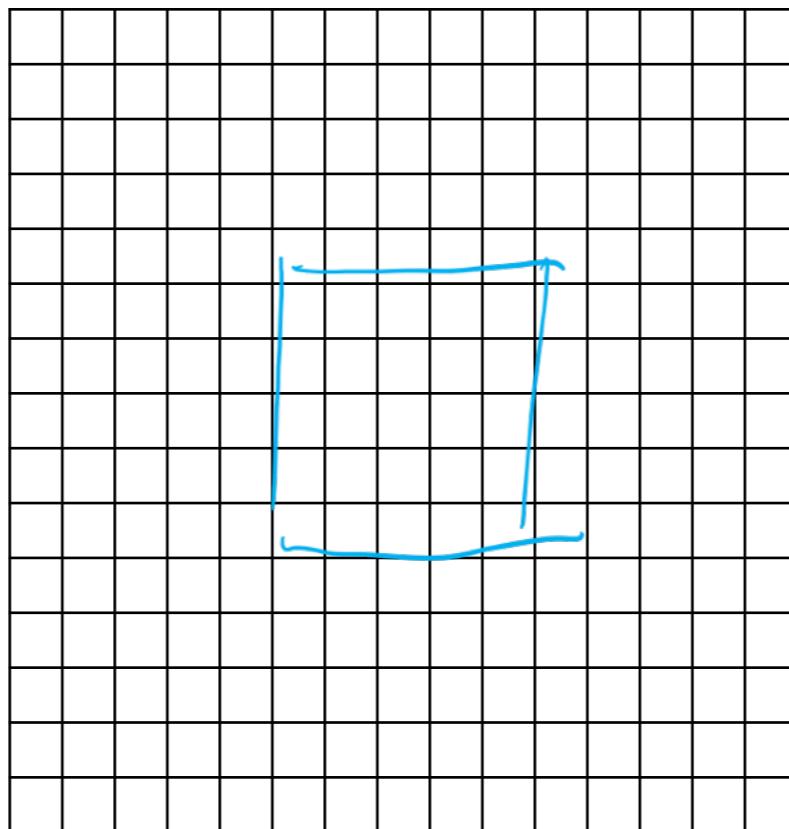
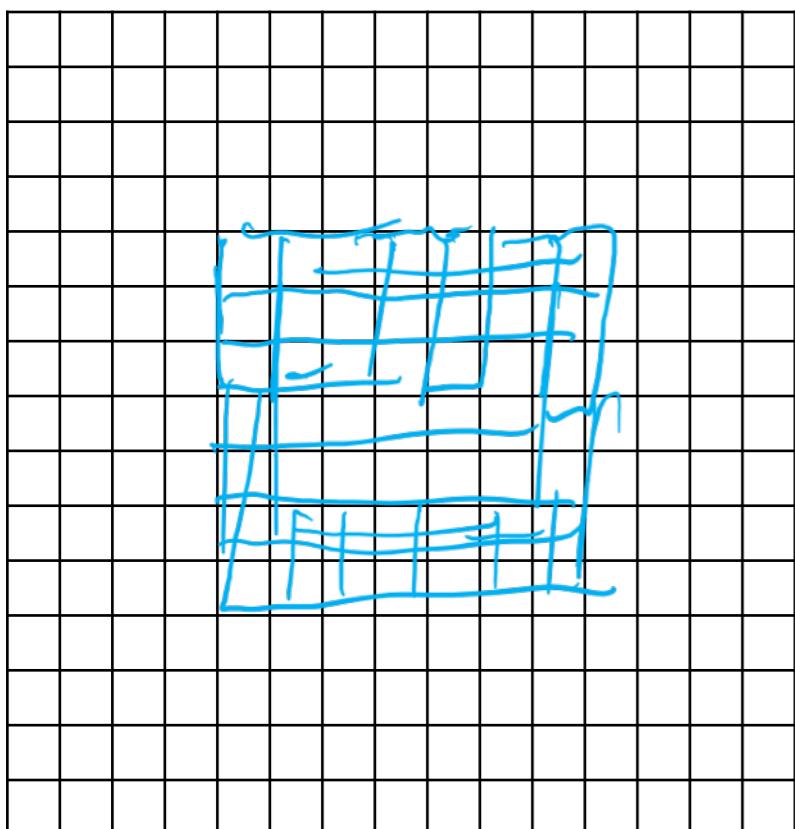
Decrease and Increase Resolution



Effect of (Max-)Pooling

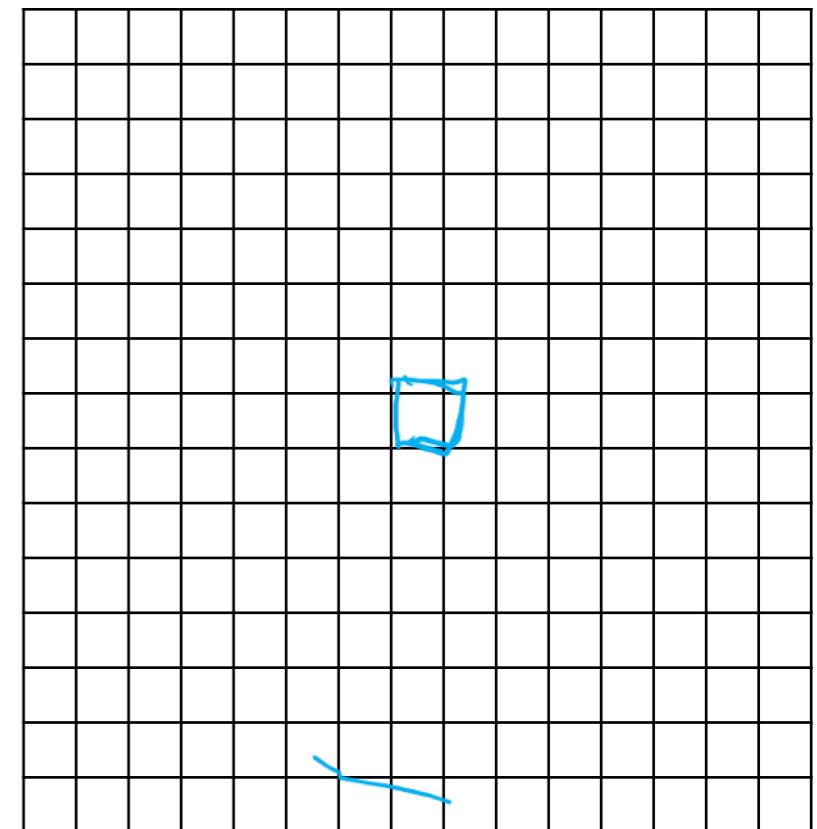
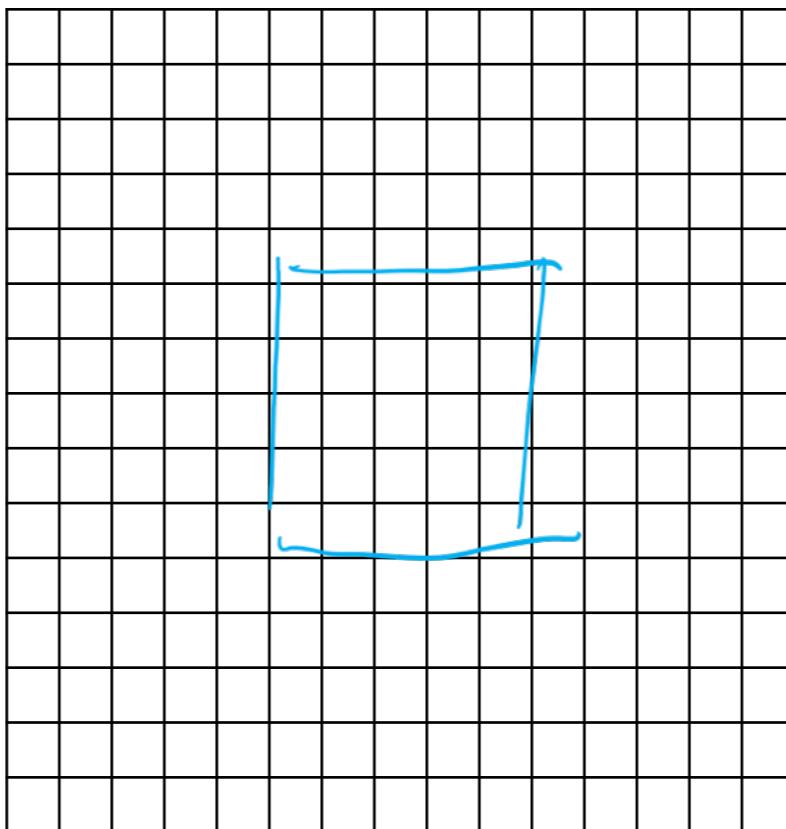
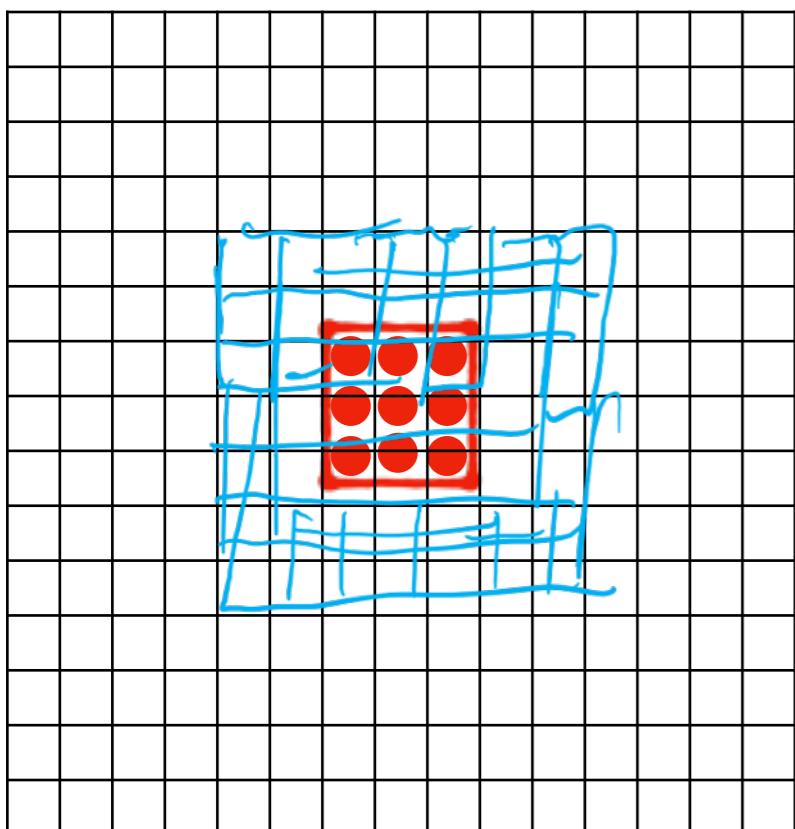


Dilated Convolutions



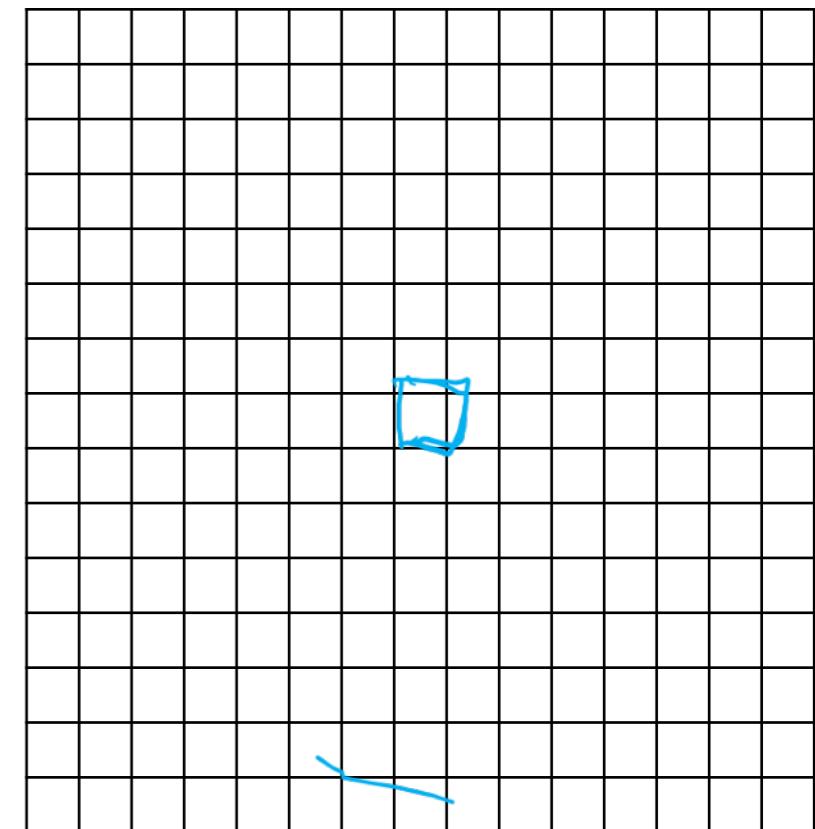
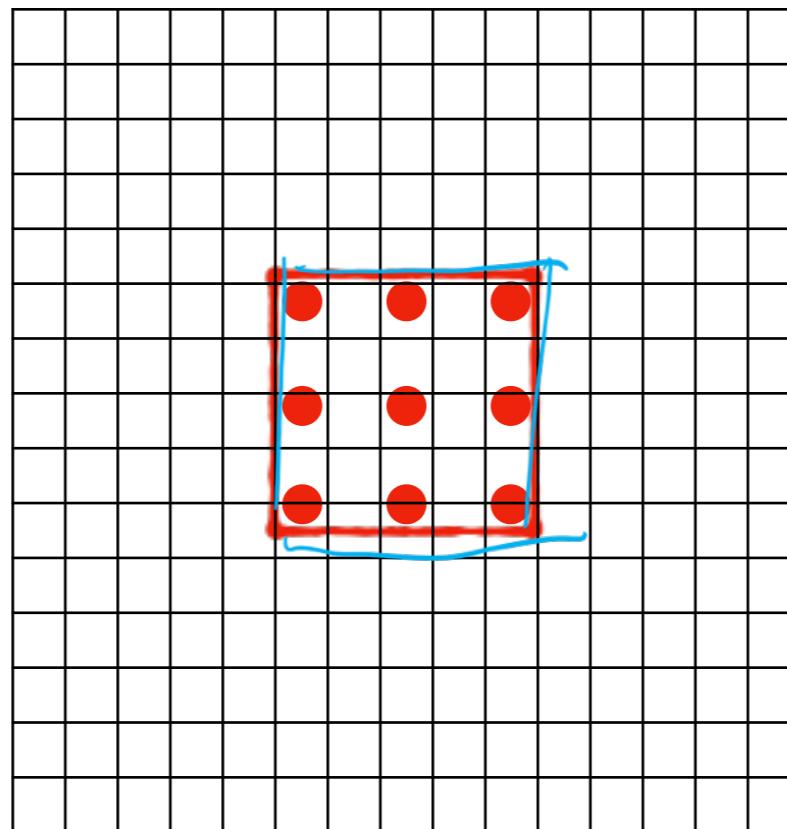
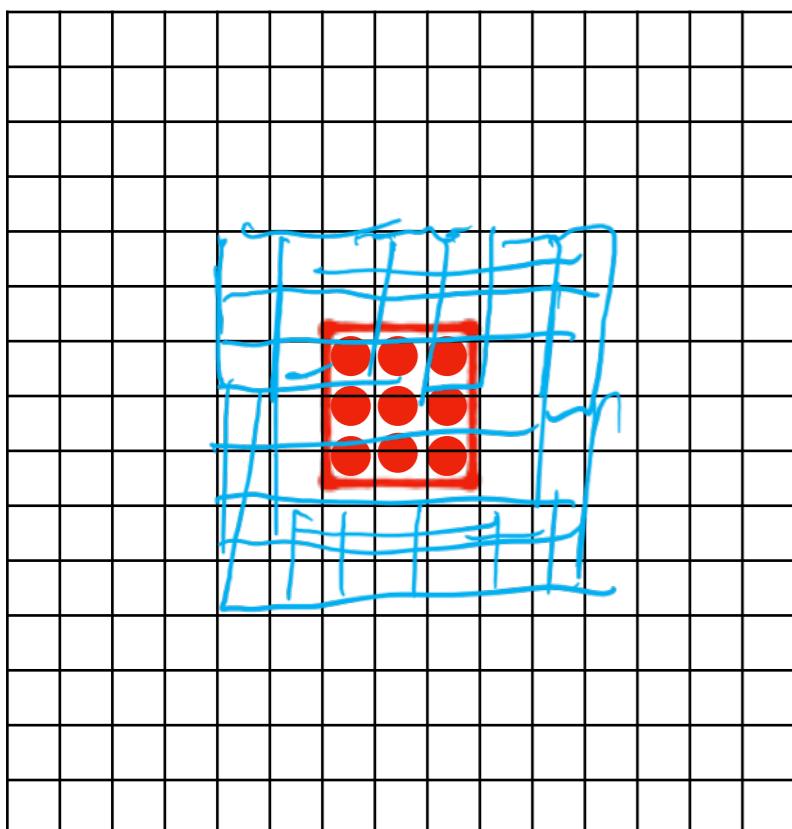
Increase dilation

Dilated Convolutions



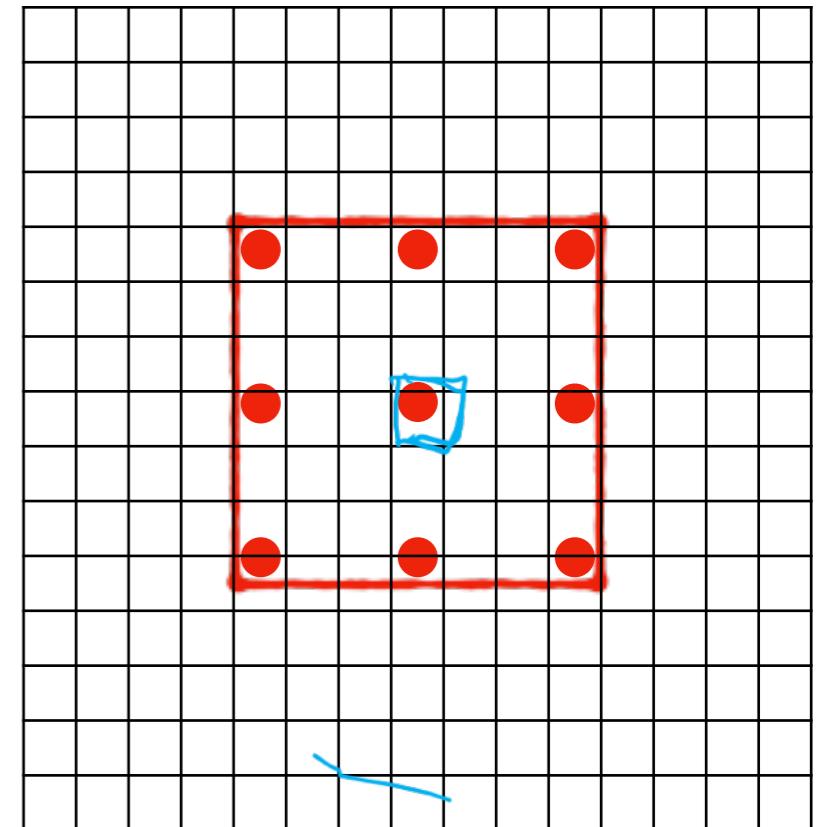
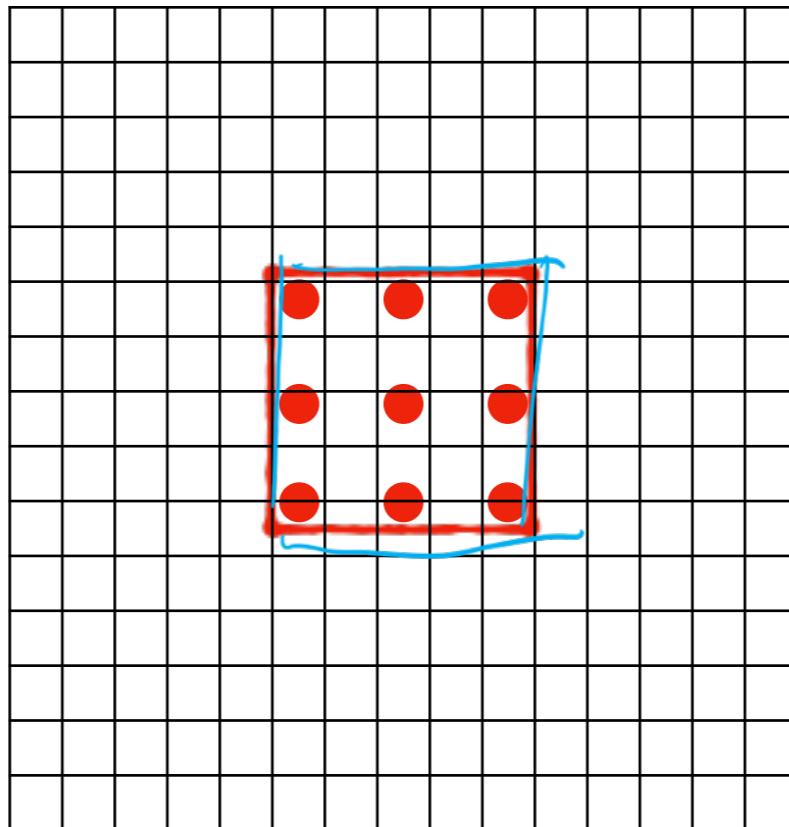
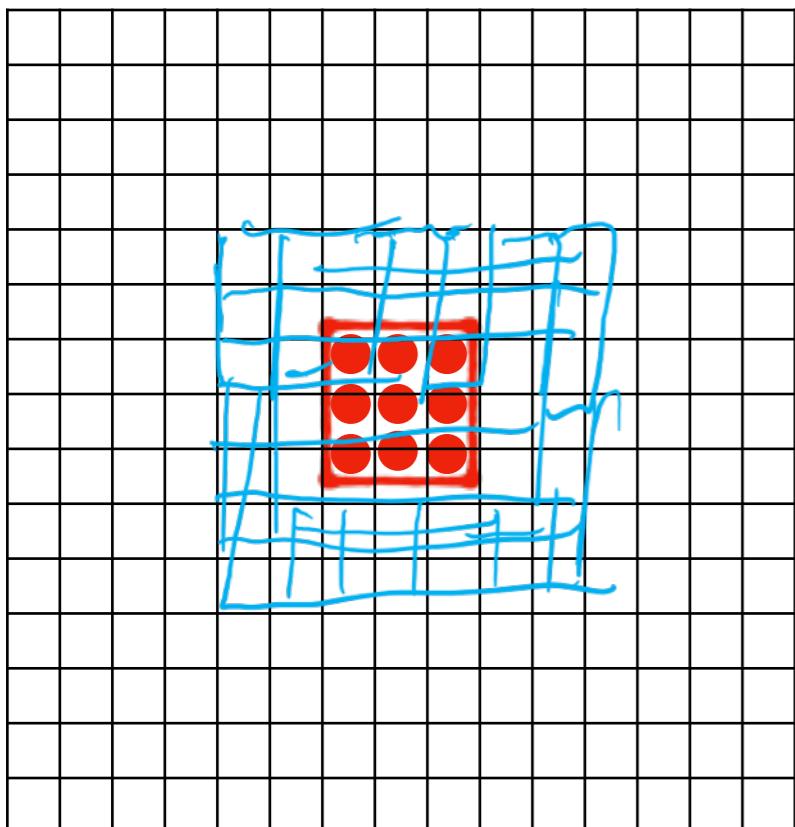
Increase dilation

Dilated Convolutions



Increase dilation

Dilated Convolutions



Increase dilation

Convolutional Layers

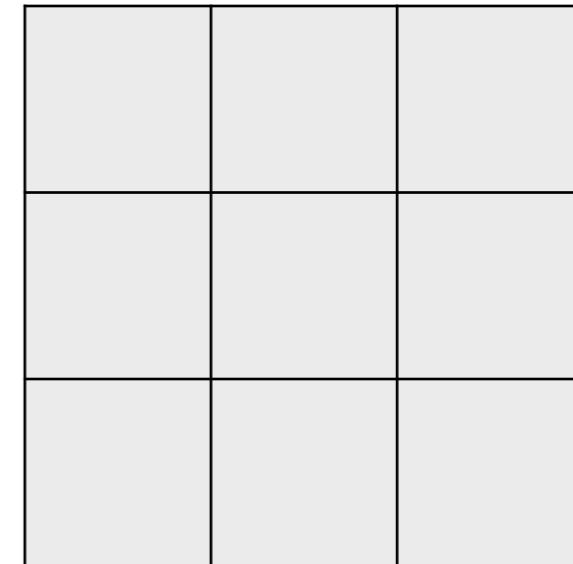
Strided Filtering

$$w = \frac{1}{5}$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0	1	0
1	1	1
0	1	0



$I \star w$

Strided Filtering

$$w = \frac{1}{5}$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

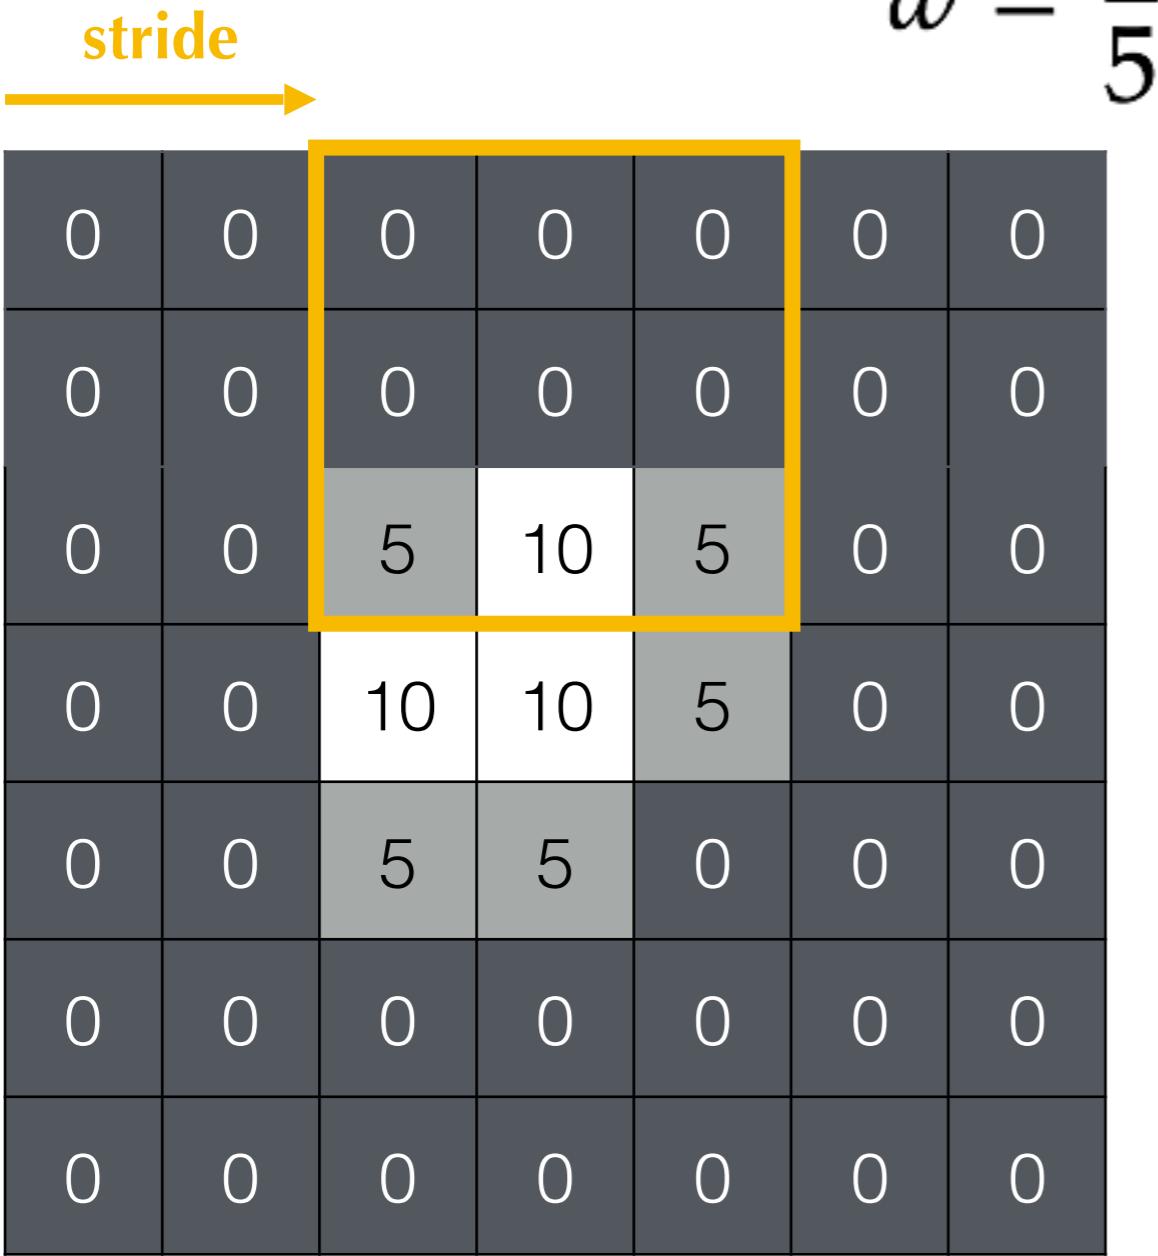
I

0	1	0
1	1	1
0	1	0

0		

$I \star w$

Strided Filtering



I

$$w = \frac{1}{5}$$

0	1	0
1	1	1
0	1	0

0	2	

$I \star w$

Strided Filtering

$$w = \frac{1}{5}$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0	1	0
1	1	1
0	1	0

0	2	0

$I \star w$

Strided Filtering

$$w = \frac{1}{5}$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0	1	0
1	1	1
0	1	0

0	2	0
2		

$I \star w$

Strided Filtering

$$w = \frac{1}{5}$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0	1	0
1	1	1
0	1	0

0	2	0
2	8	

$I \star w$

Strided Filtering

$$w = \frac{1}{5}$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0	1	0
1	1	1
0	1	0

0	2	0
2	8	1

$I \star w$

Strided Filtering

$$w = \frac{1}{5}$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

I

0	1	0
1	1	1
0	1	0

0	2	0
2	8	1
0		

$I \star w$

Strided Filtering

$$w = \frac{1}{5}$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	10	5	0	0
0	0	10	10	5	0	0
0	0	5	5	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

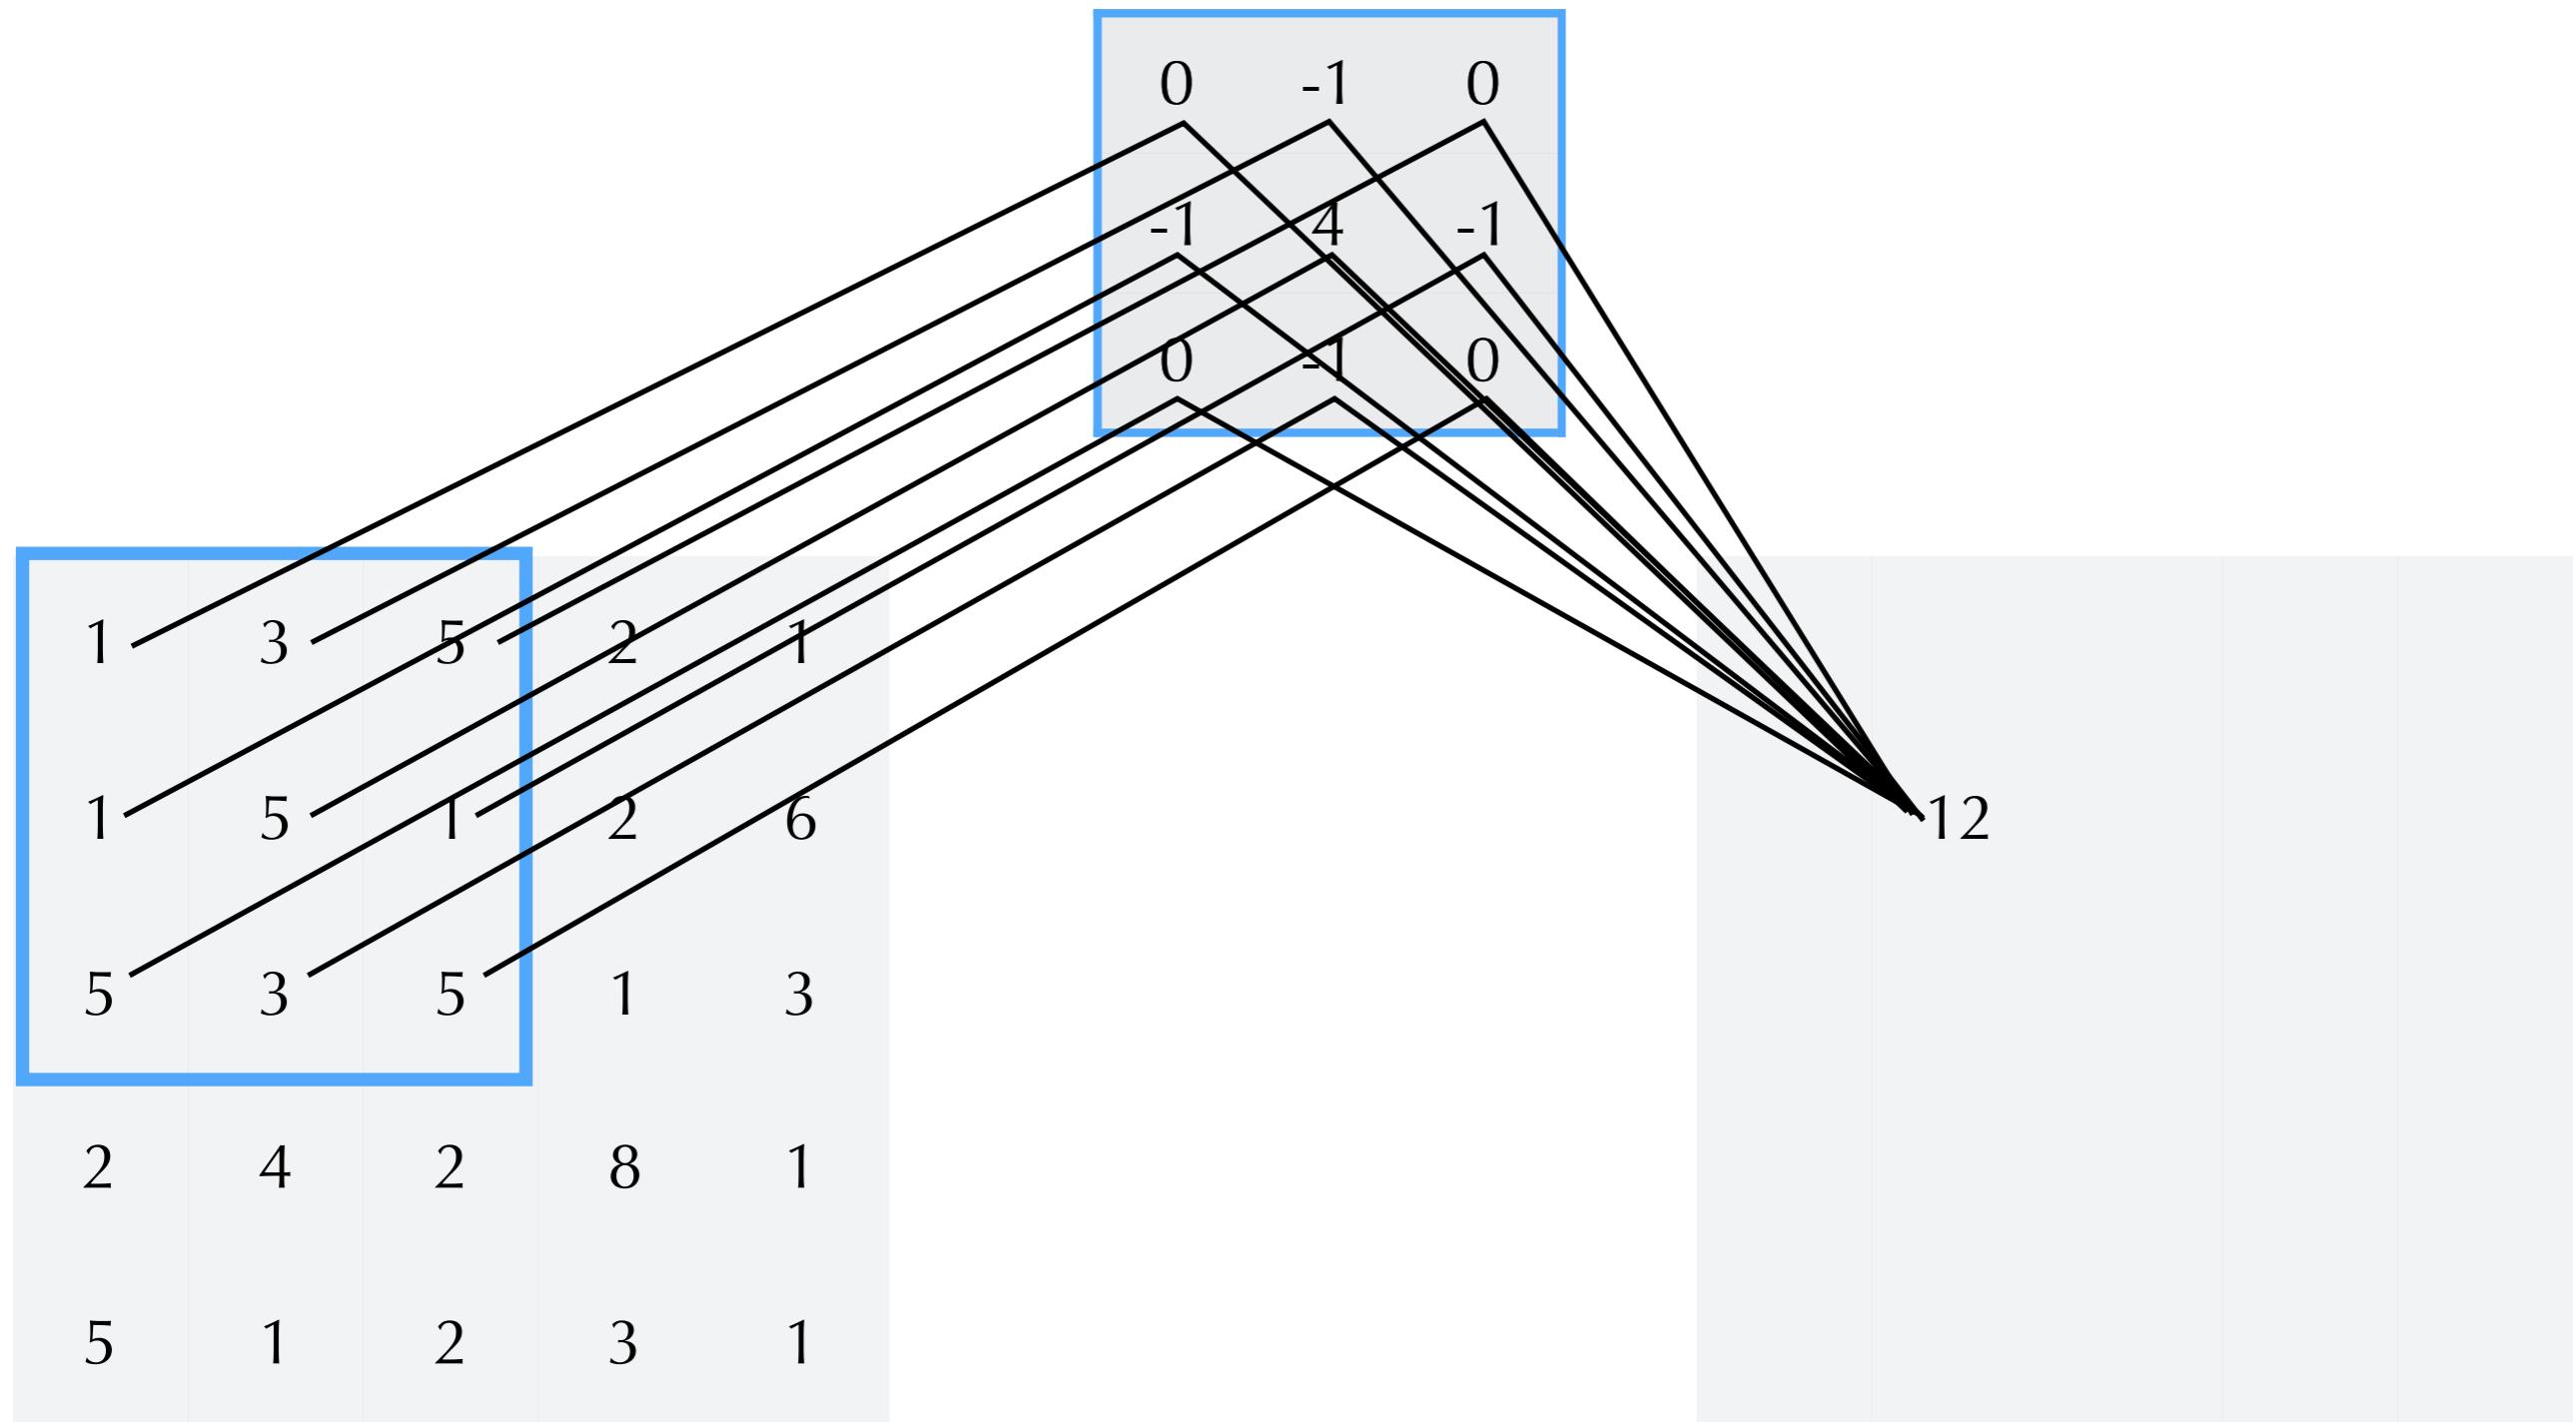
I

0	1	0
1	1	1
0	1	0

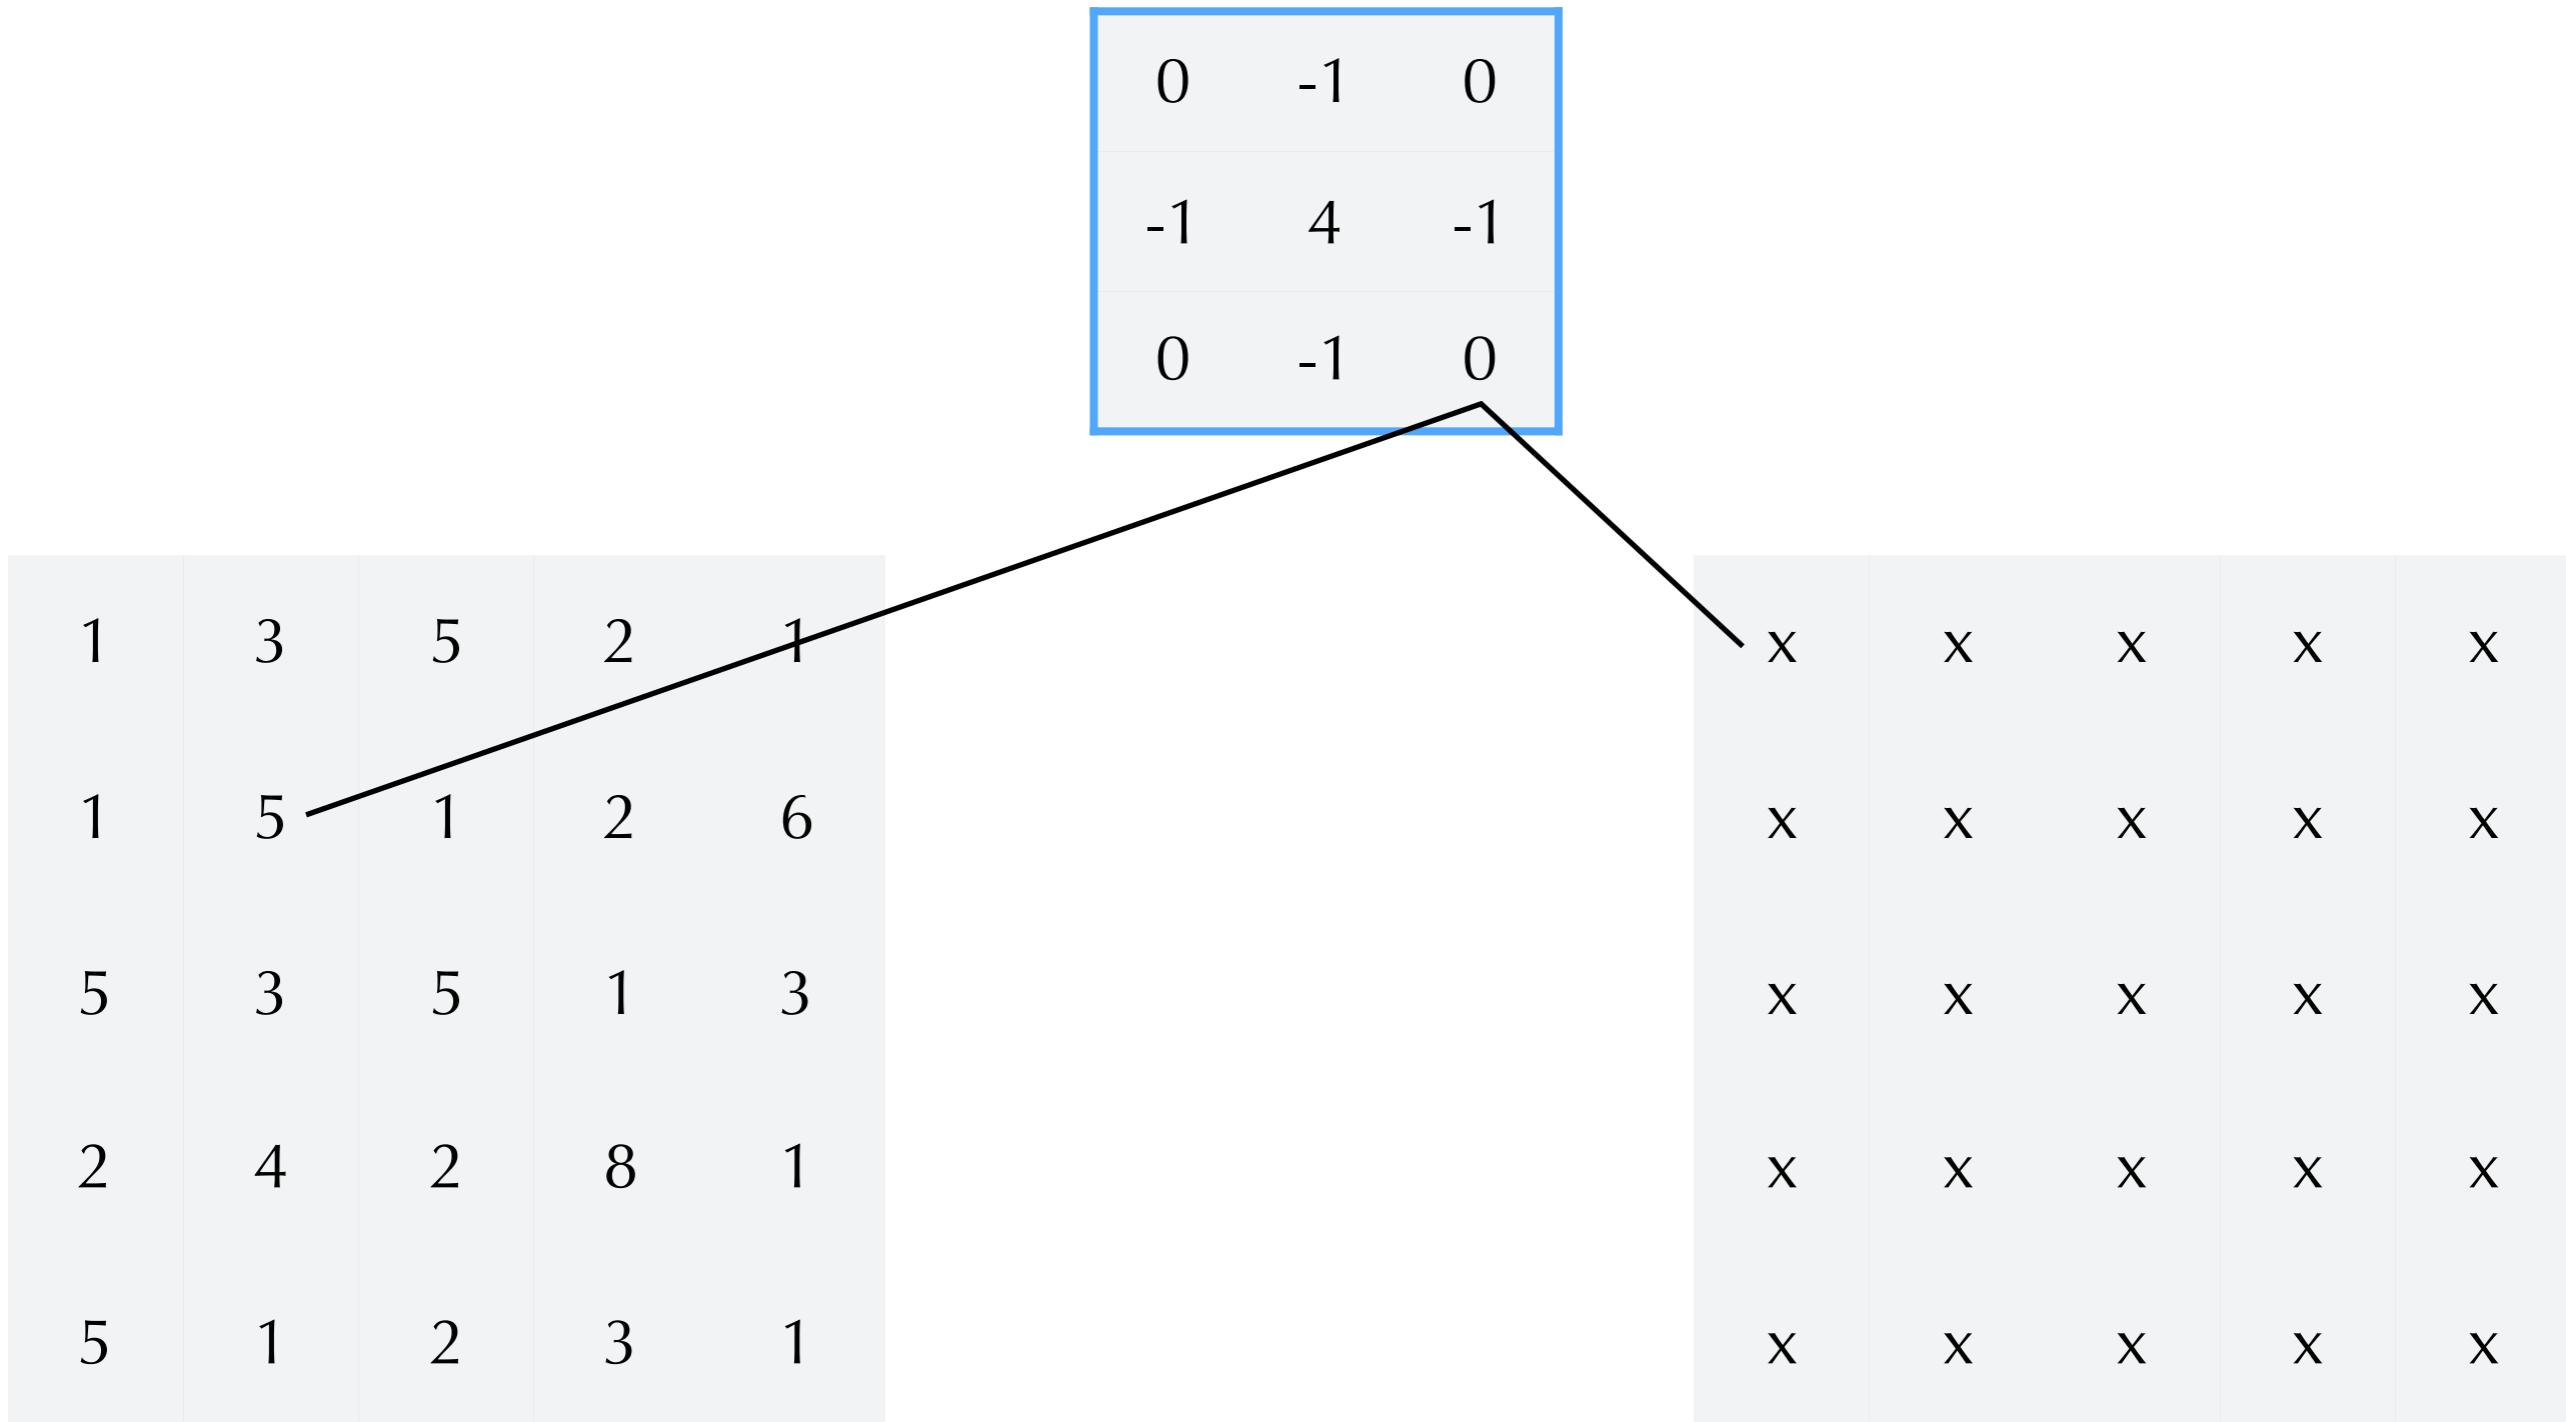
0	2	0
2	8	1
0	1	

$I \star w$

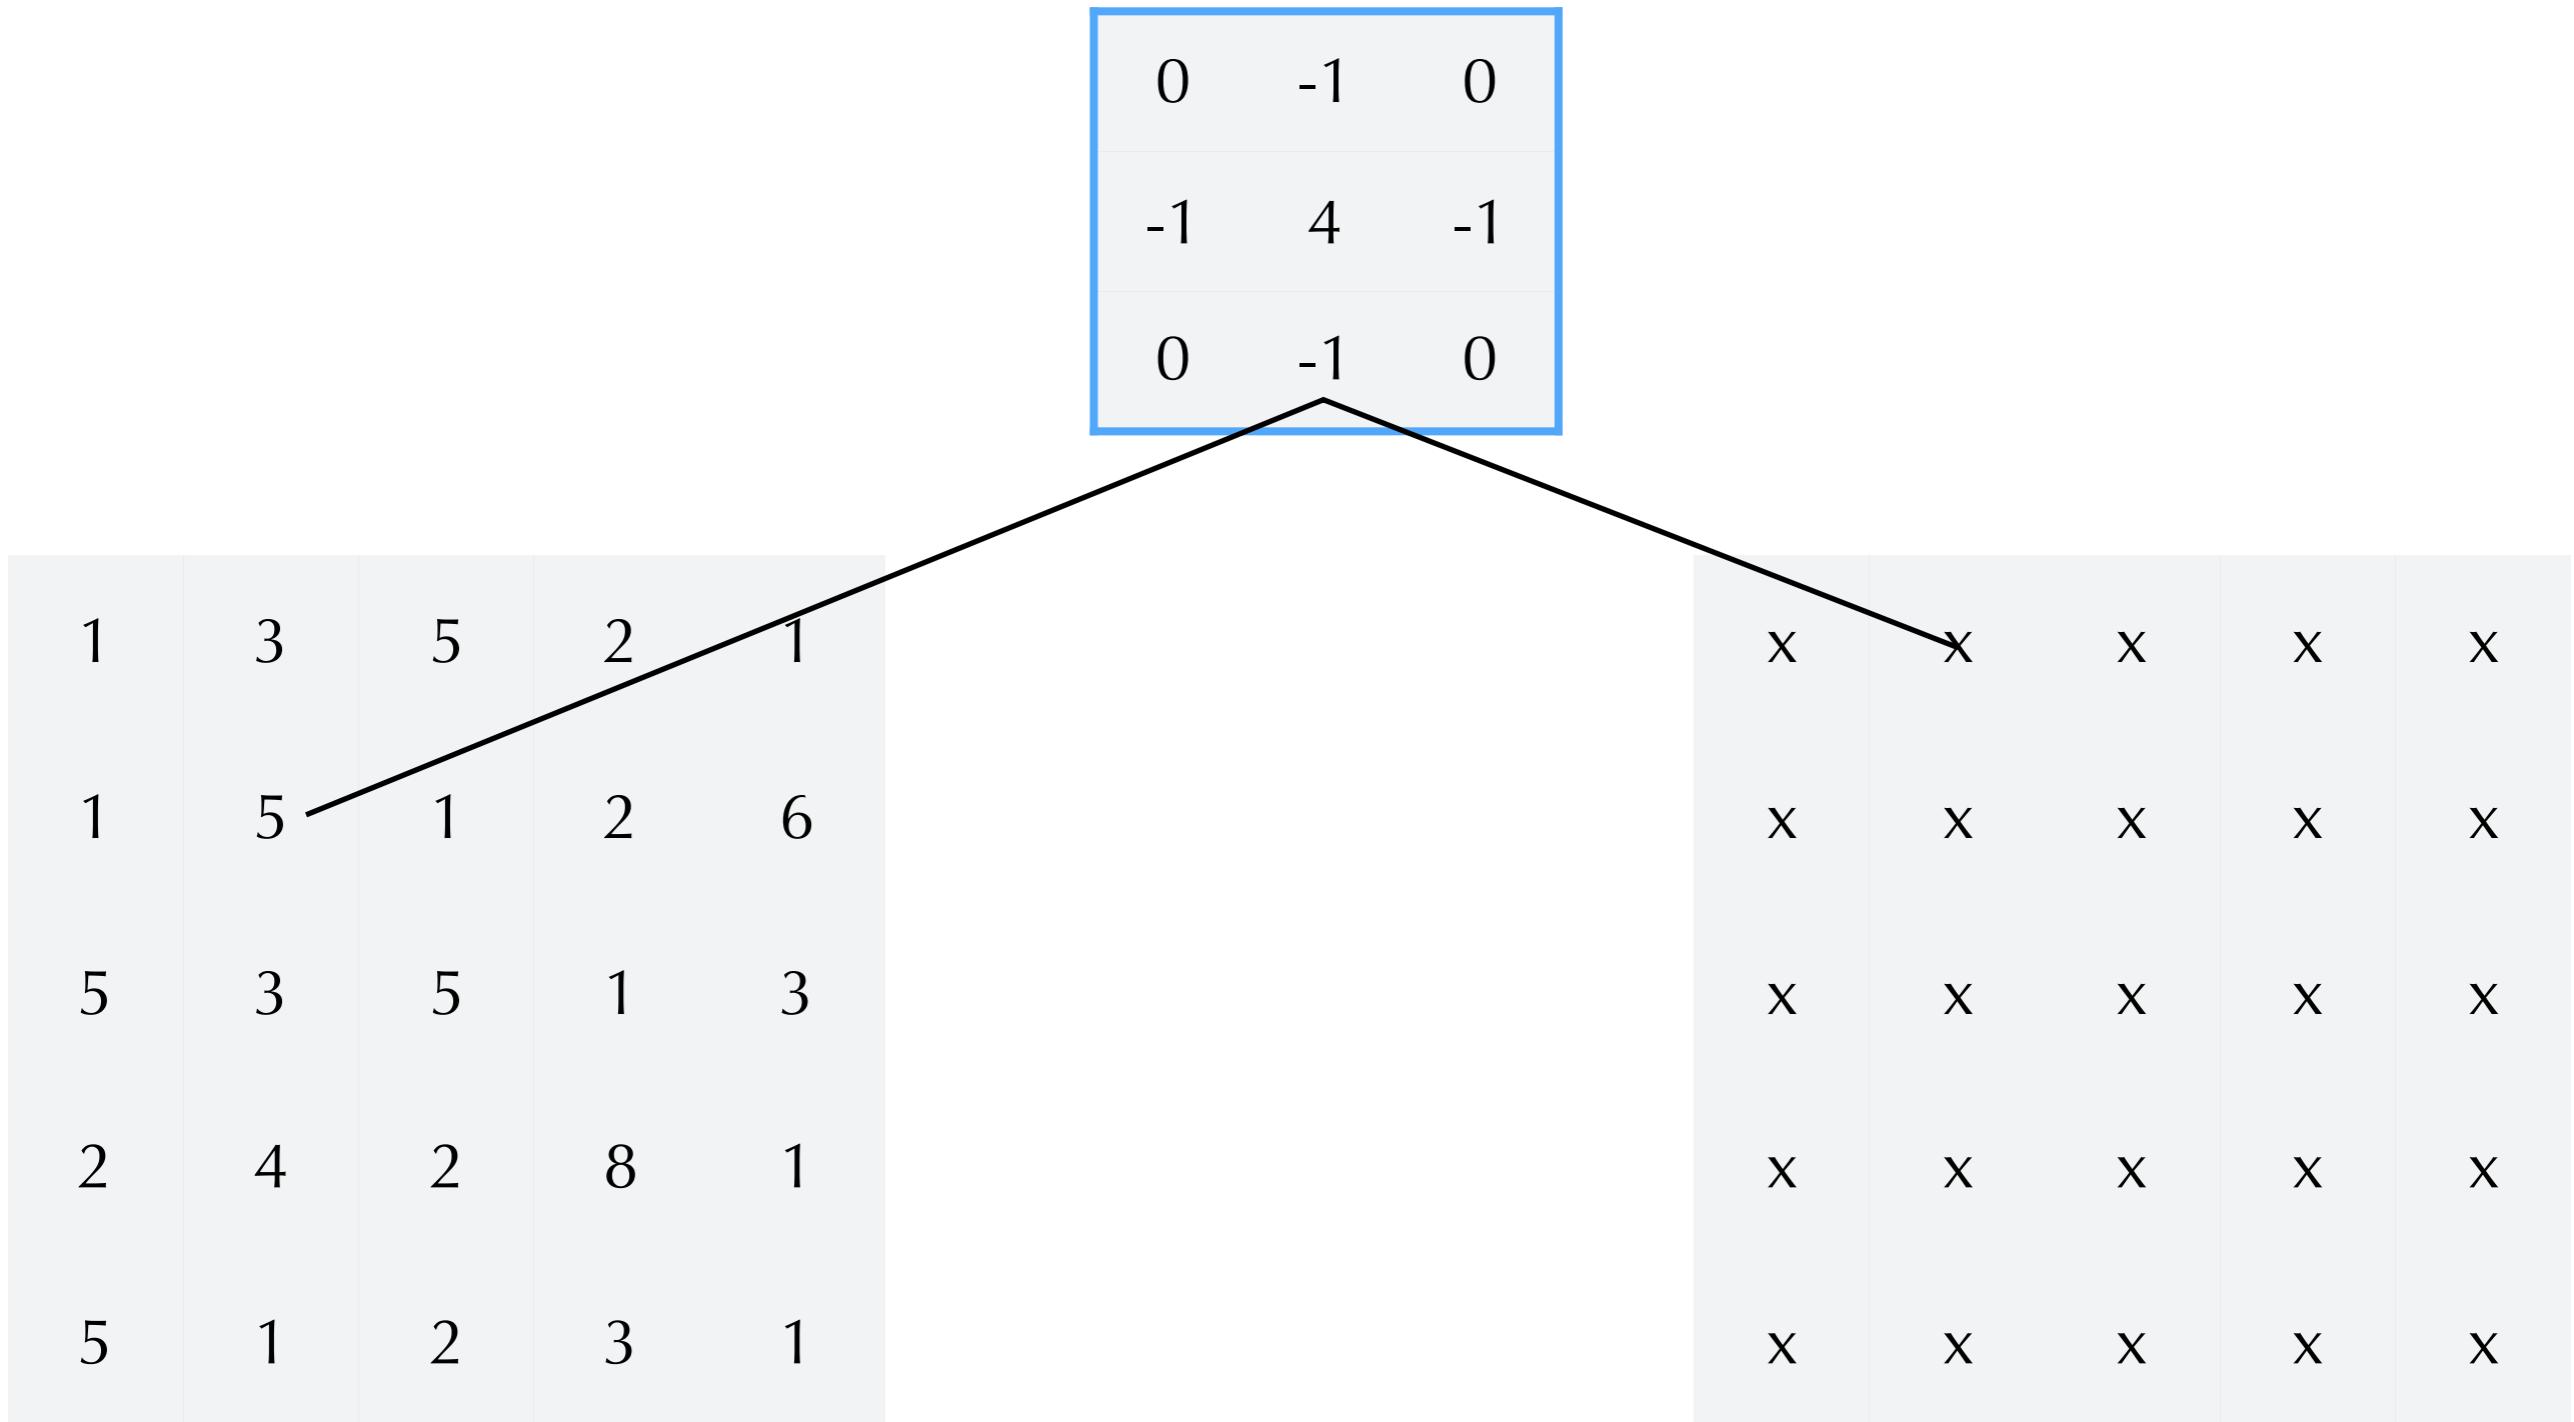
Different View



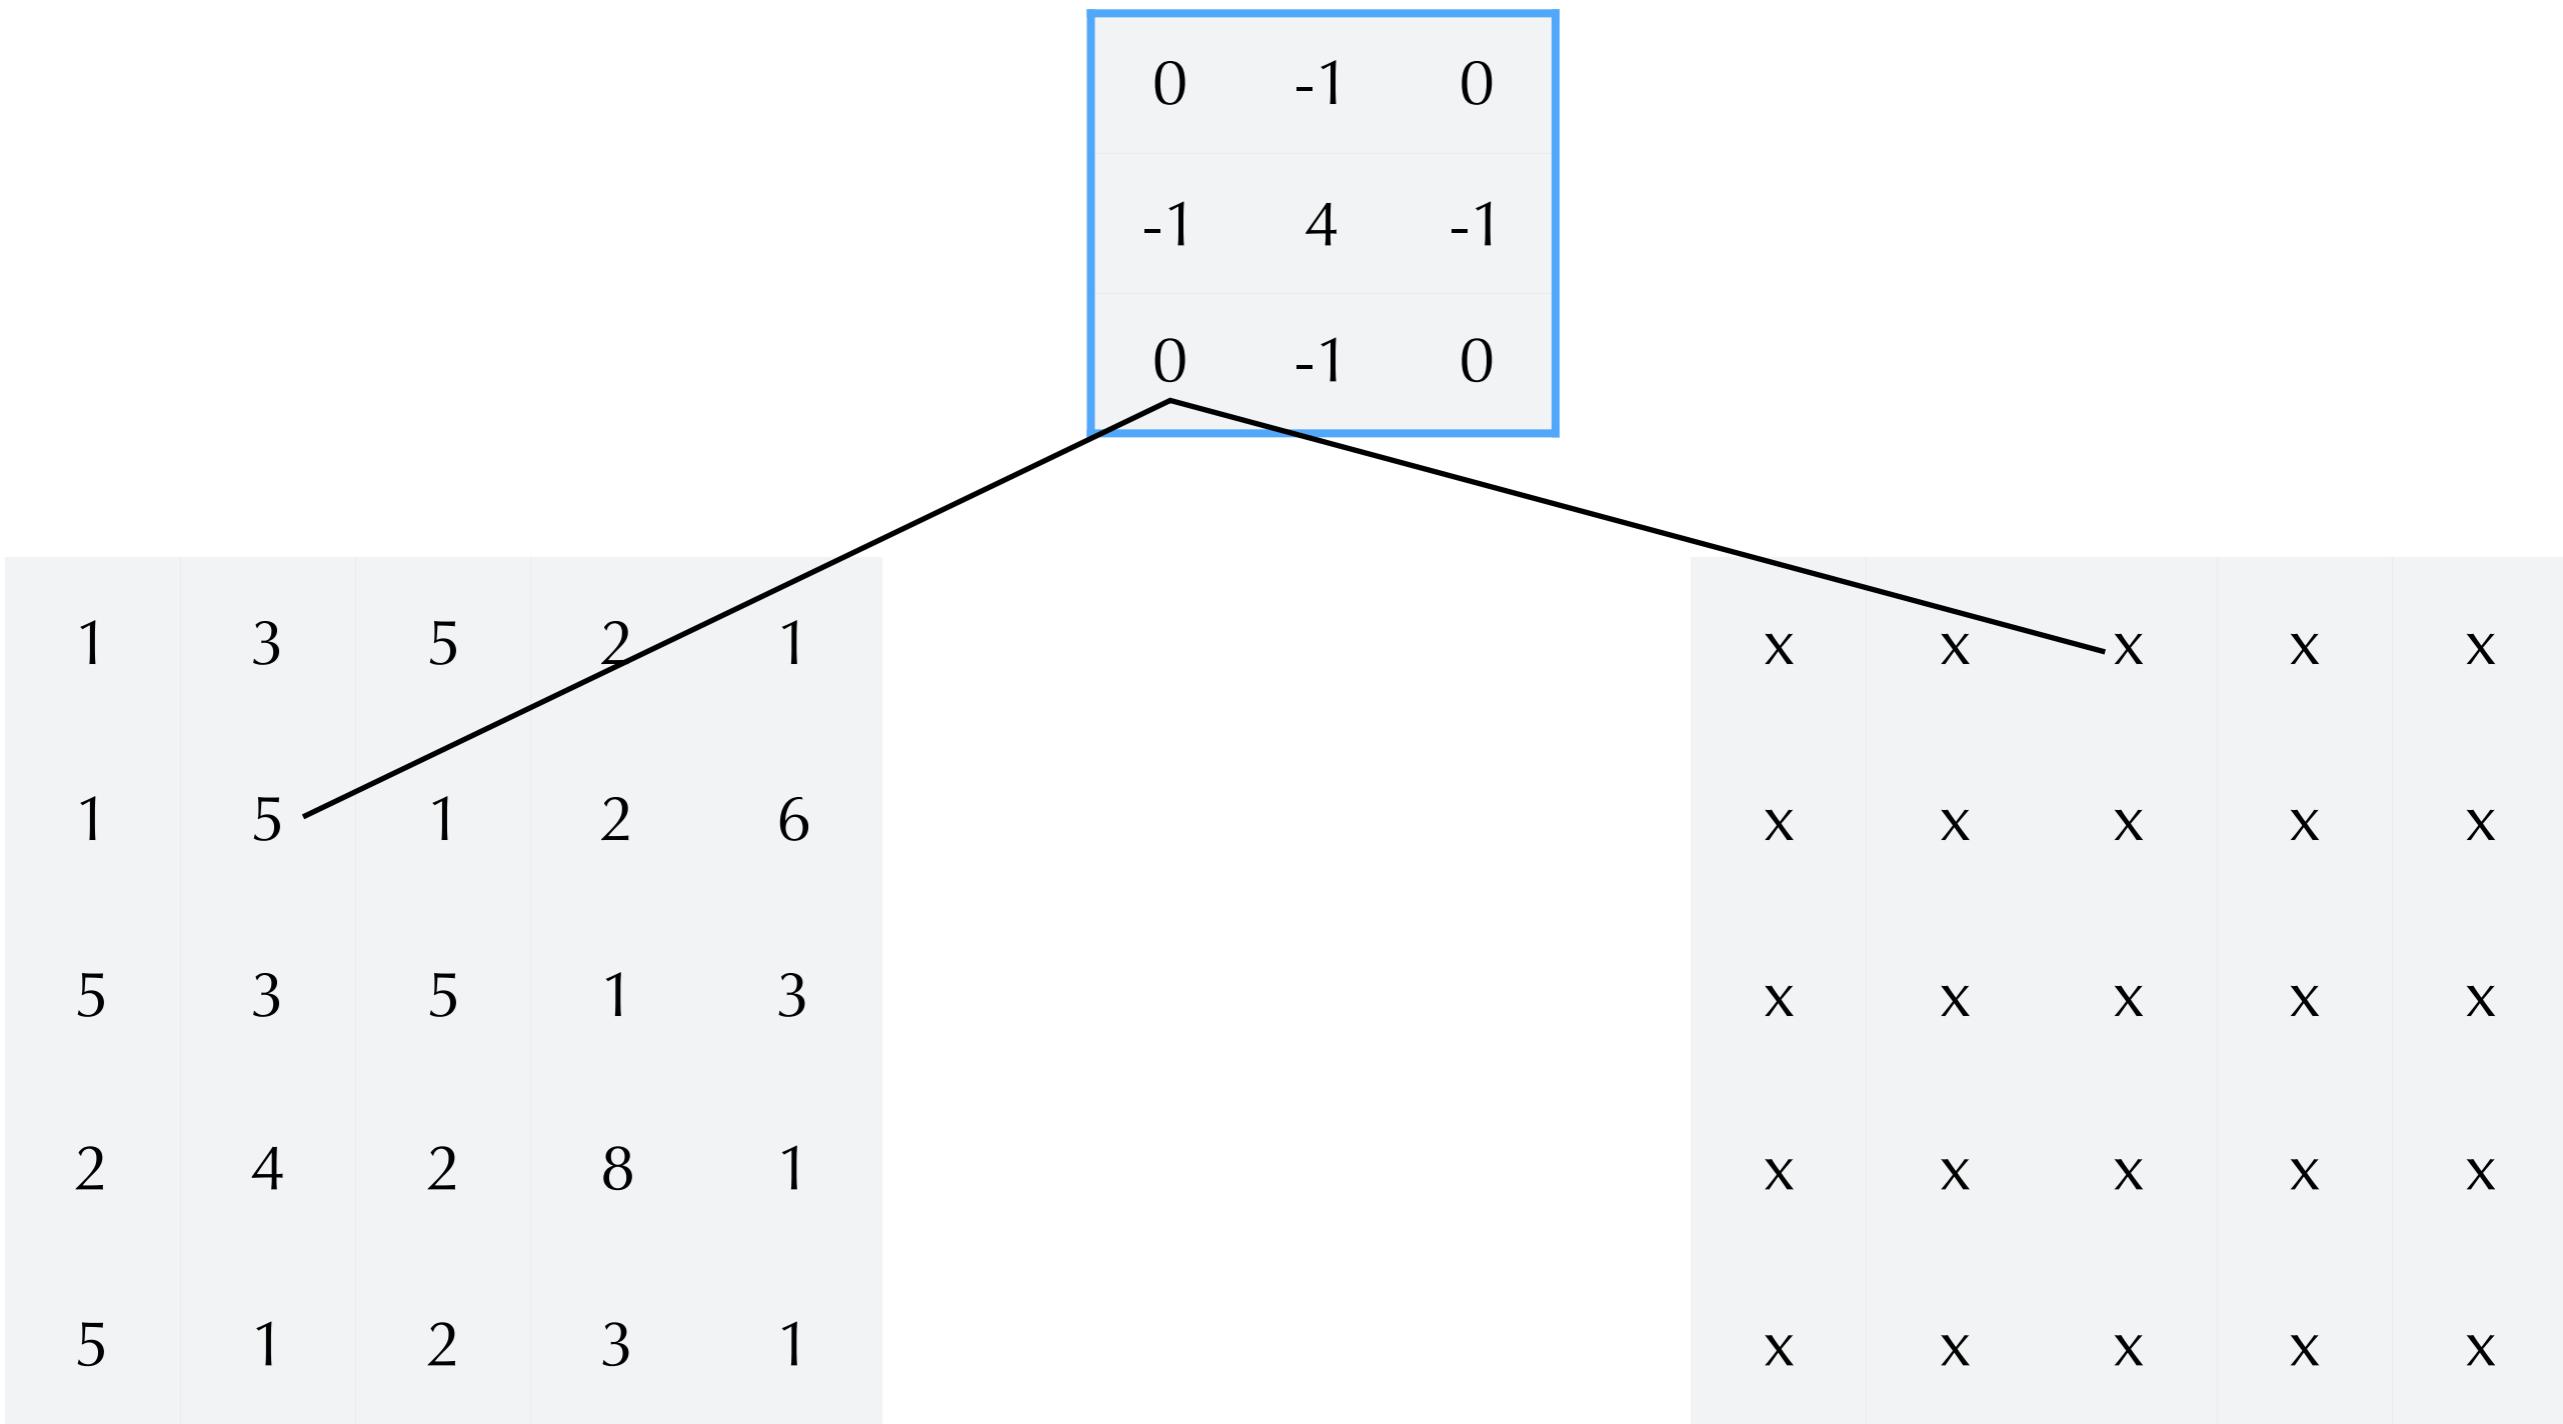
Different View



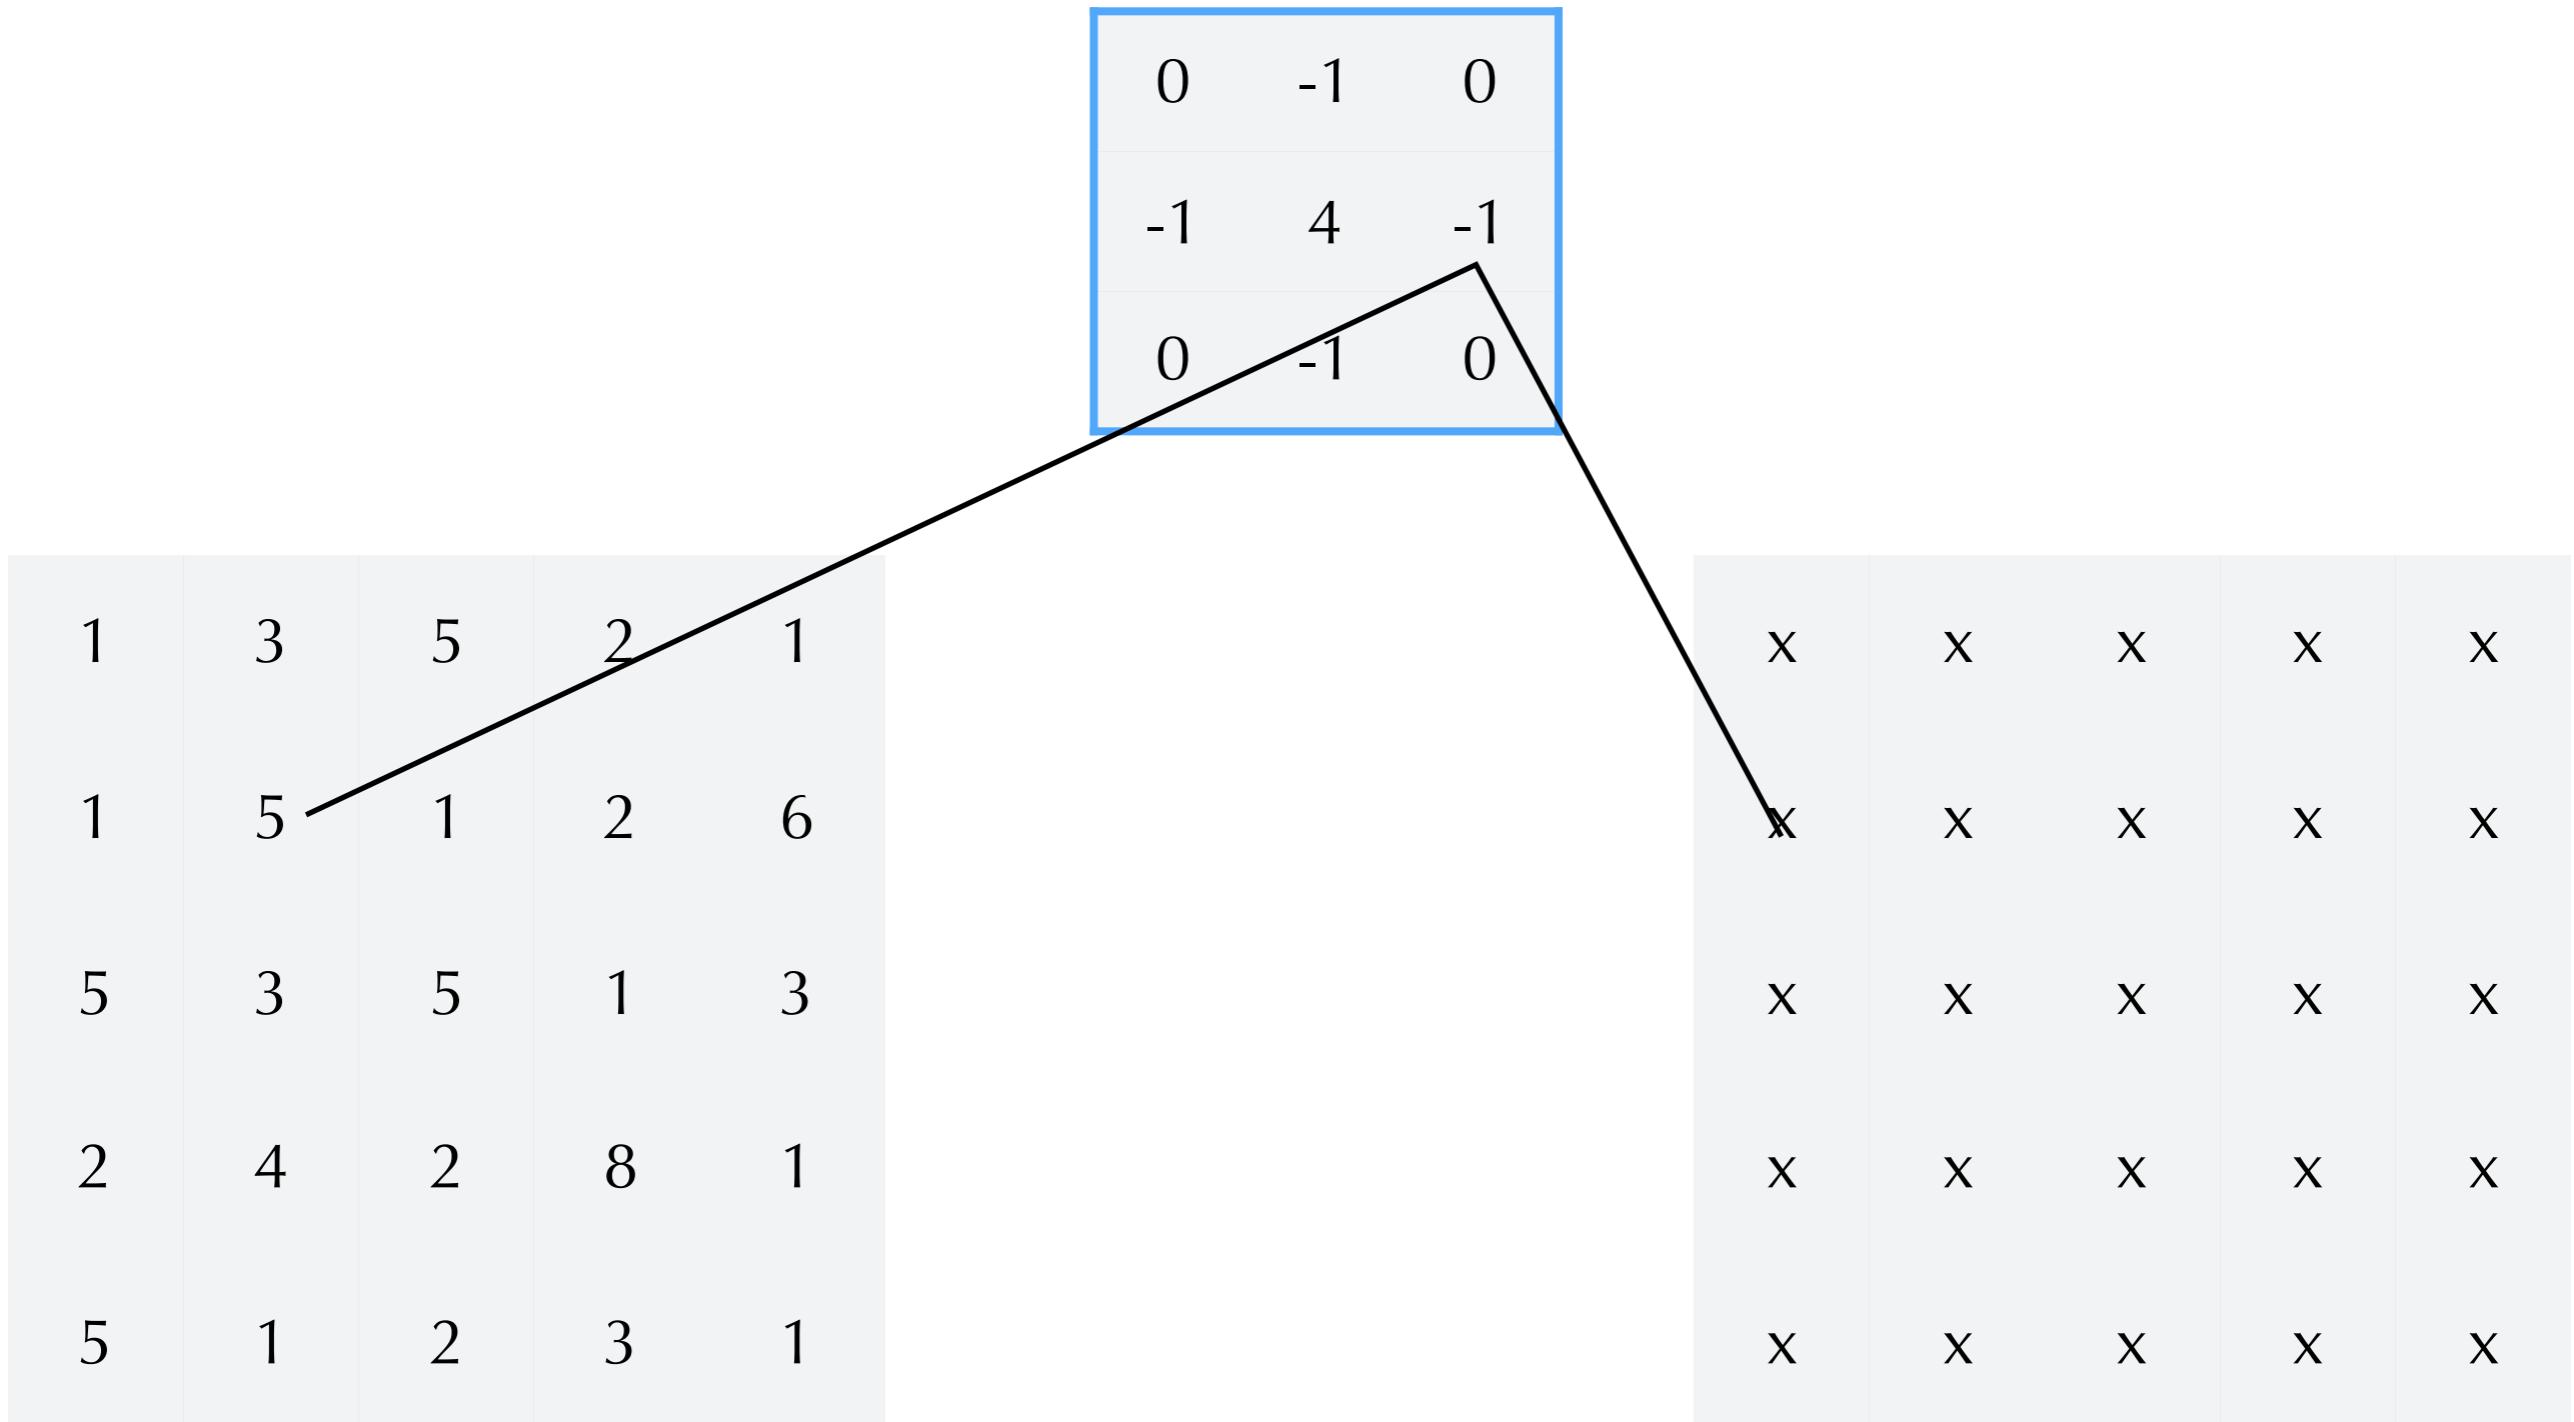
Different View



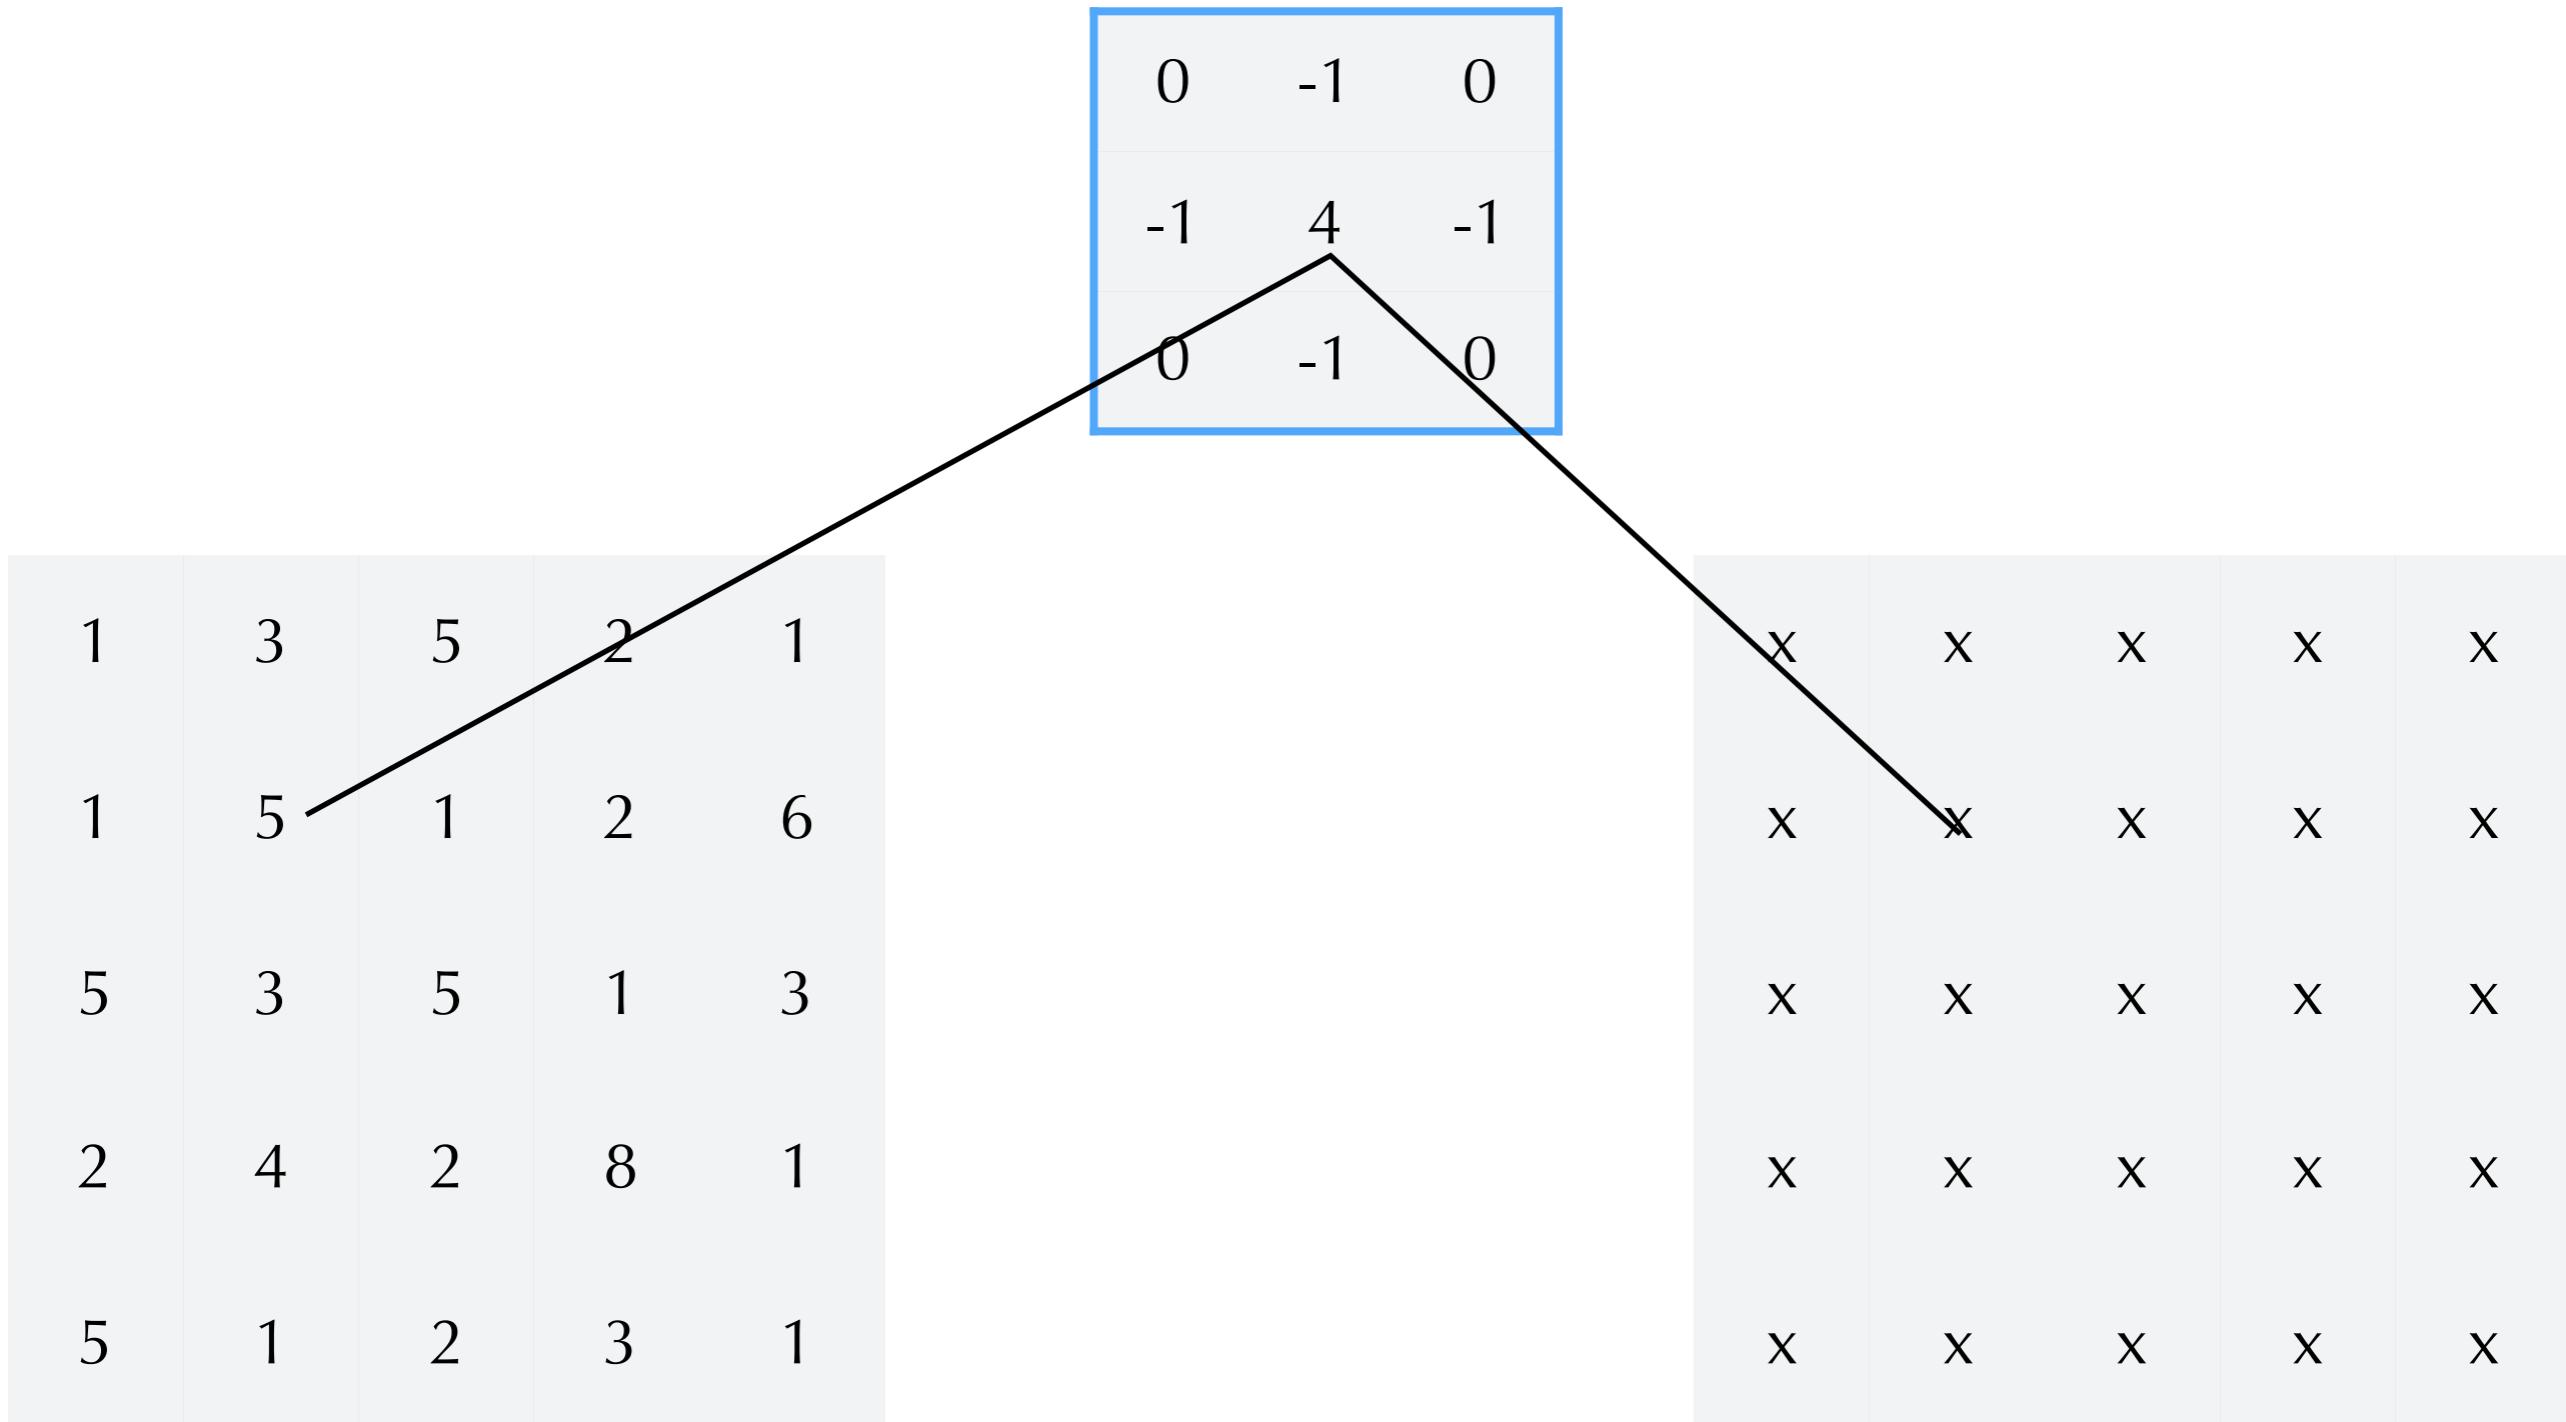
Different View



Different View



Different View



Different View

0	-1	0
-1	4	-1
0	-1	0

x	x	x	x	x
x	x	x	x	x
x	x	x	x	x
x	x	x	x	x
x	x	x	x	x

x	x	x	x	x
x	x	x	x	x
x	x	x	x	x
x	x	x	x	x
x	x	x	x	x

Different View

0	-1	0
-1	4	-1
0	-1	0

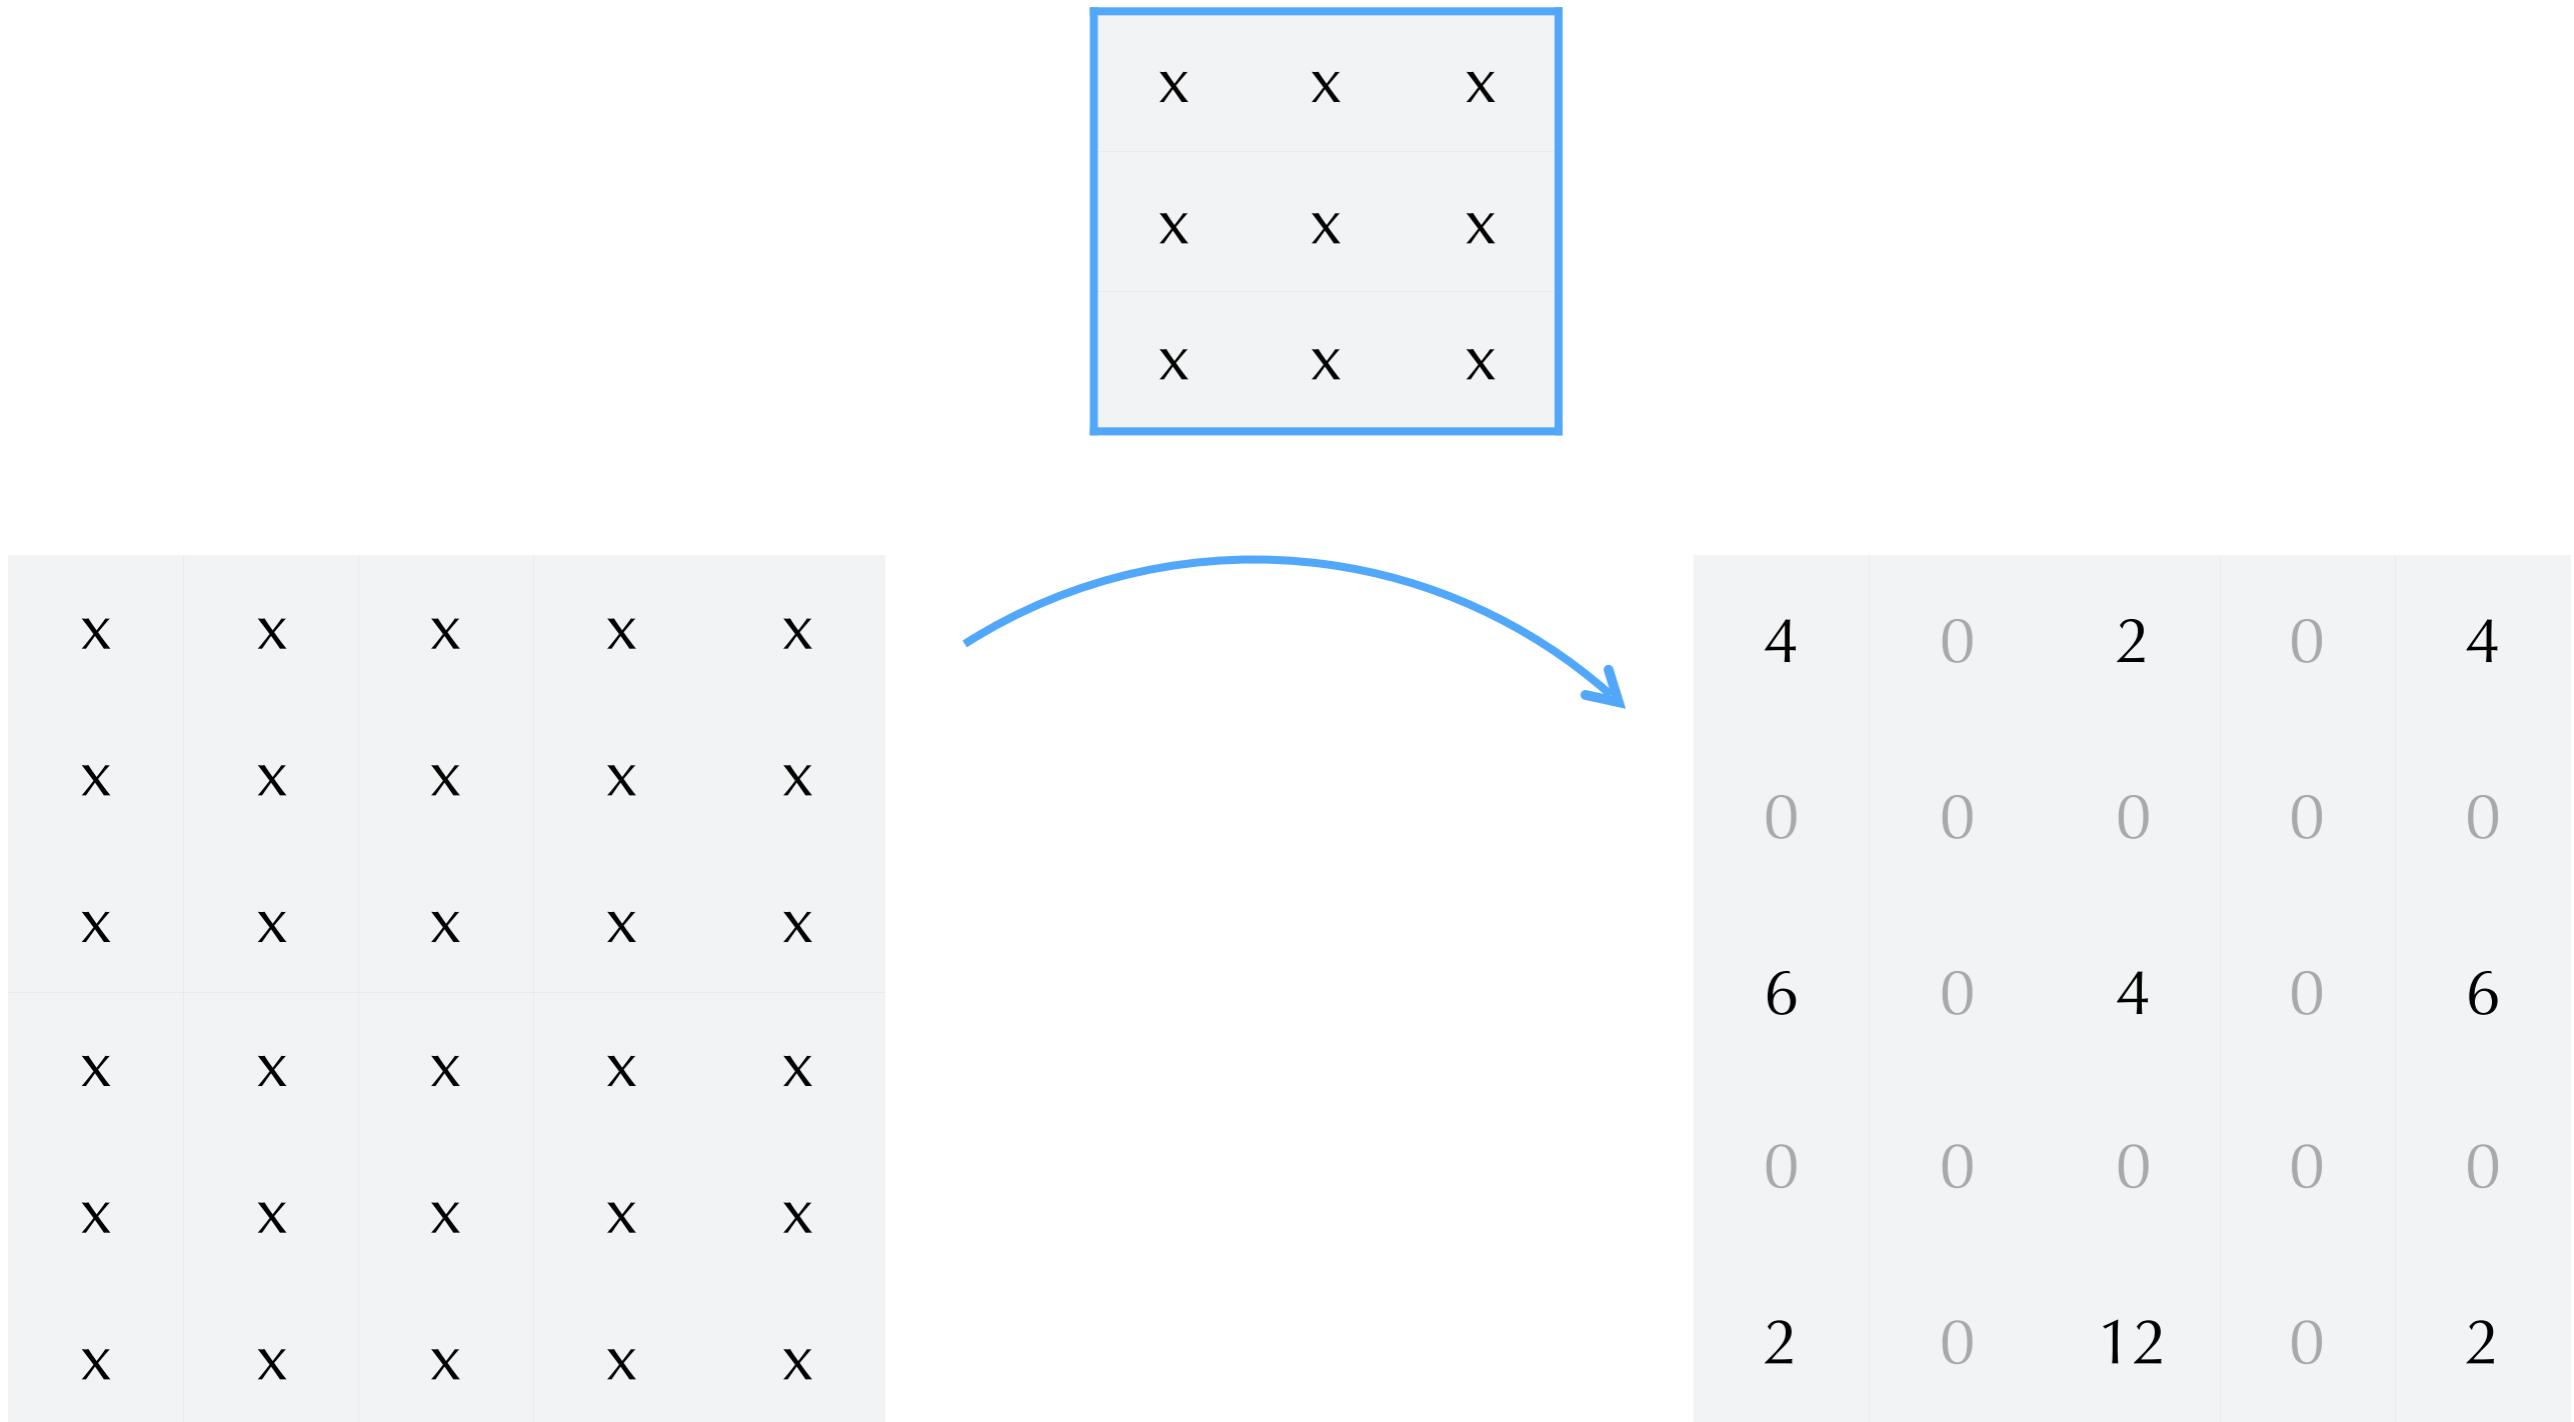
x	x	x	x	x
x	x	x	x	x
x	x	x	x	x
x	x	x	x	x
x	x	x	x	x



x	x	x	x	x
x	x	x	x	x
x	x	x	x	x
x	x	x	x	x
x	x	x	x	x

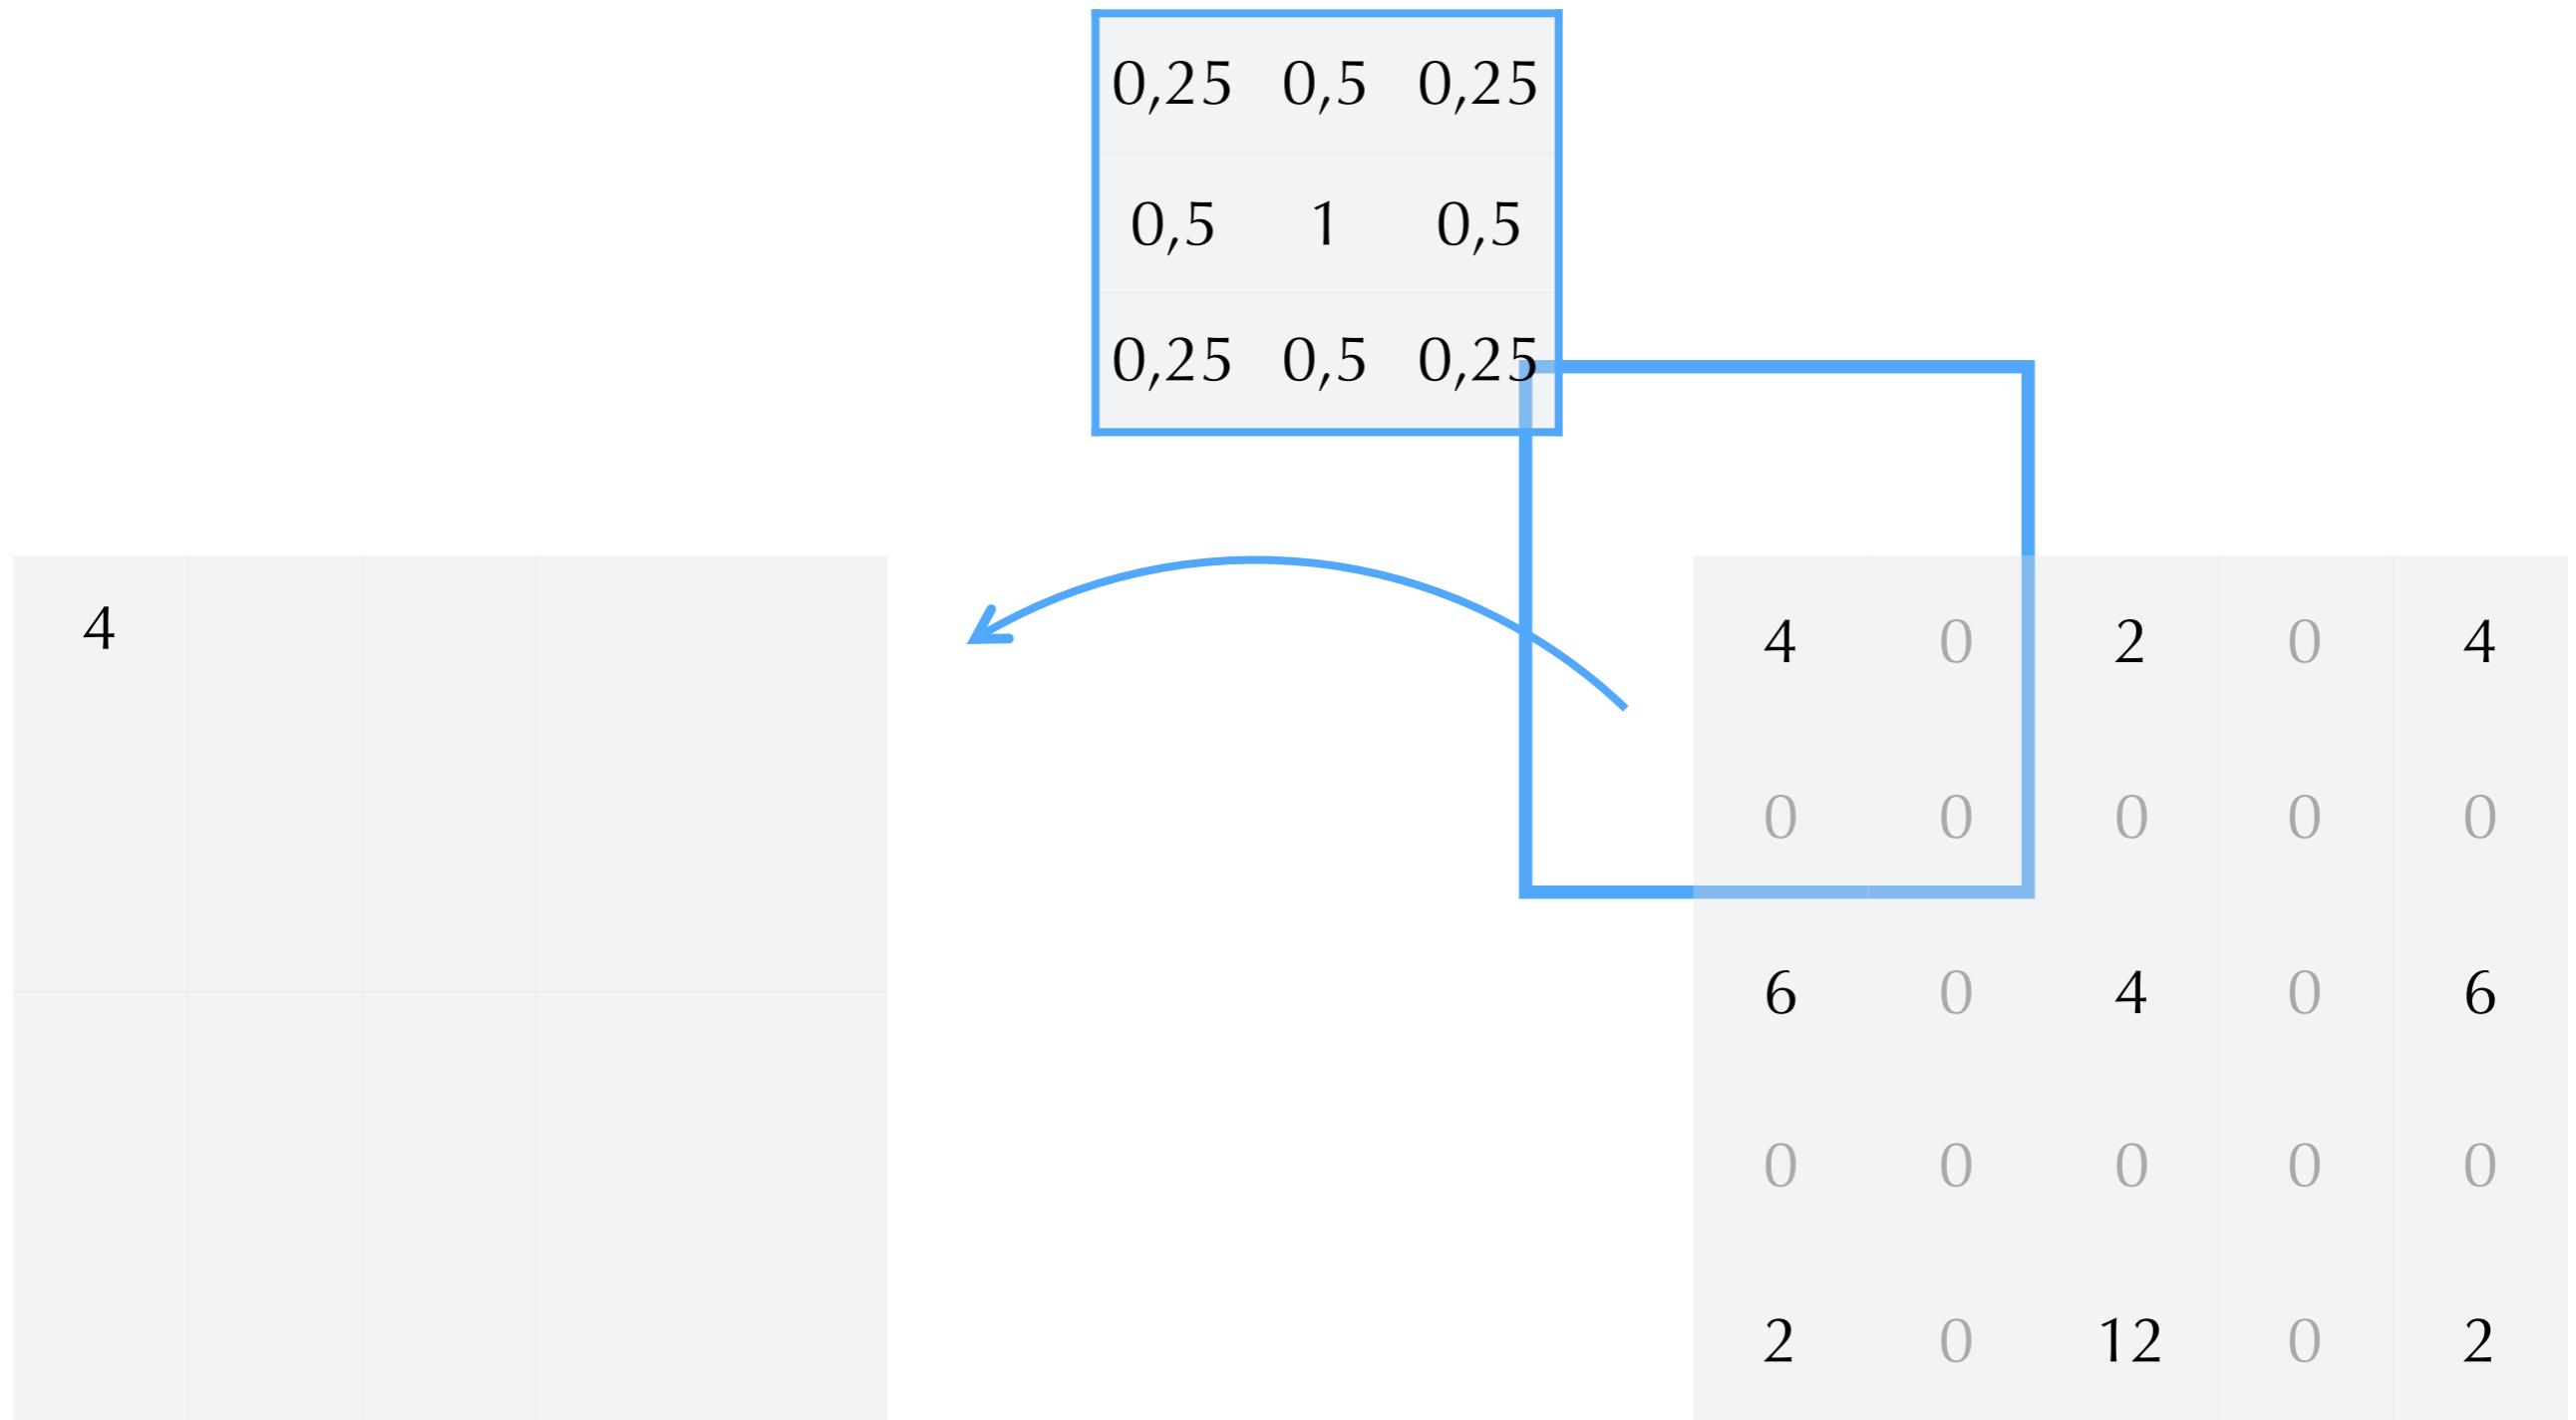
Transposed convolution

What's the point?



Strided filtering

What's the point?

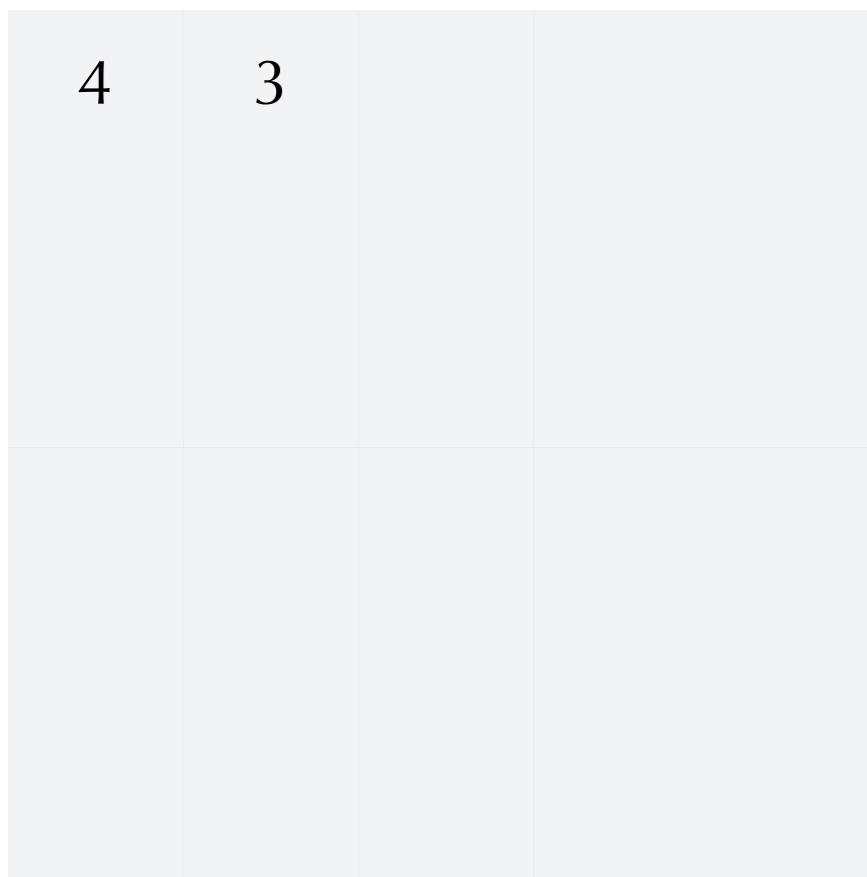


Transposed strided filtering

What's the point?

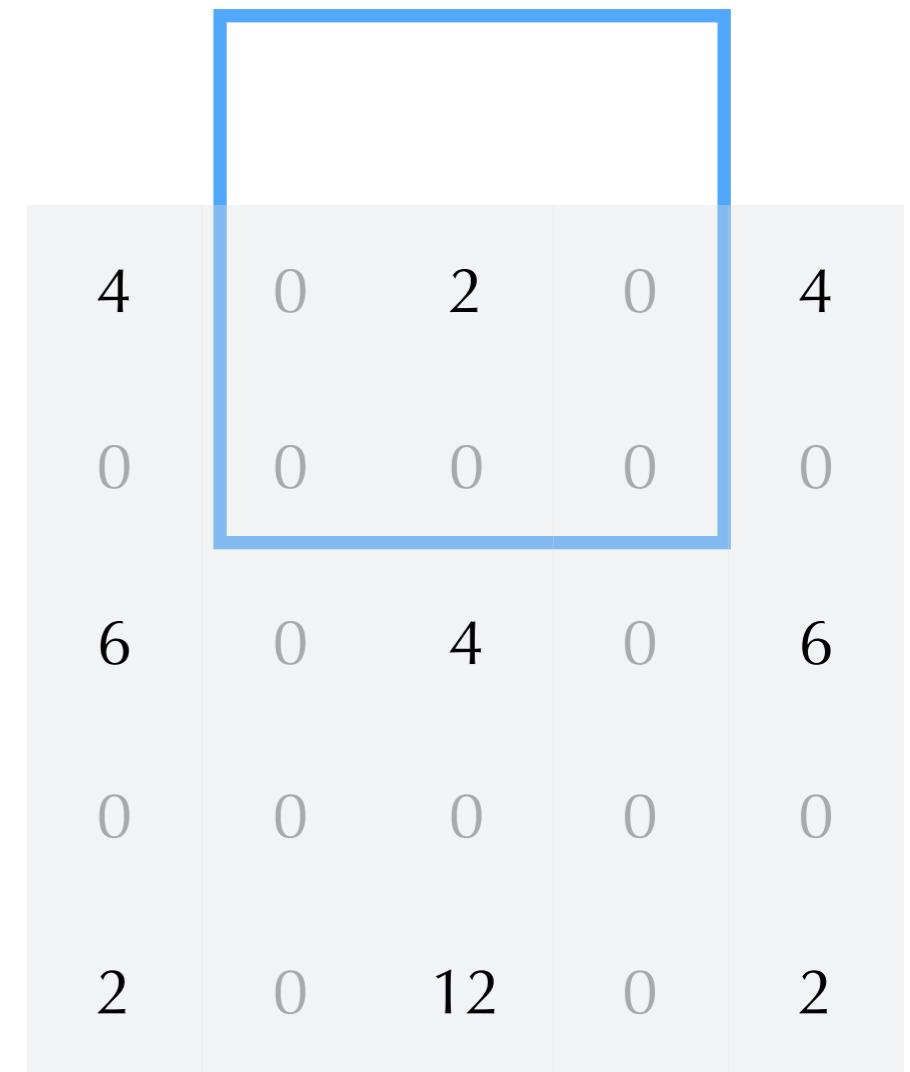
0,25	0,5	0,25
0,5	1	0,5
0,25	0,5	0,25

4	0	2	0	4
0	0	0	0	0
6	0	4	0	6
0	0	0	0	0
2	0	12	0	2



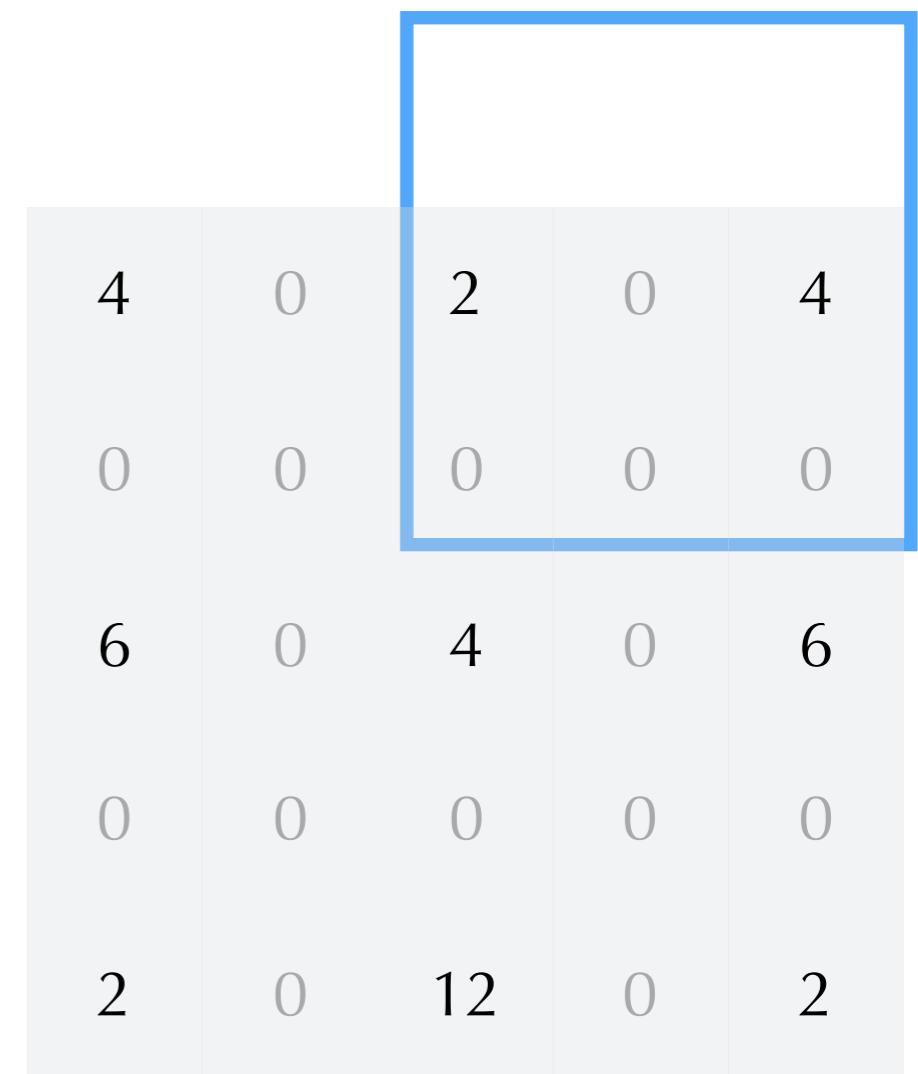
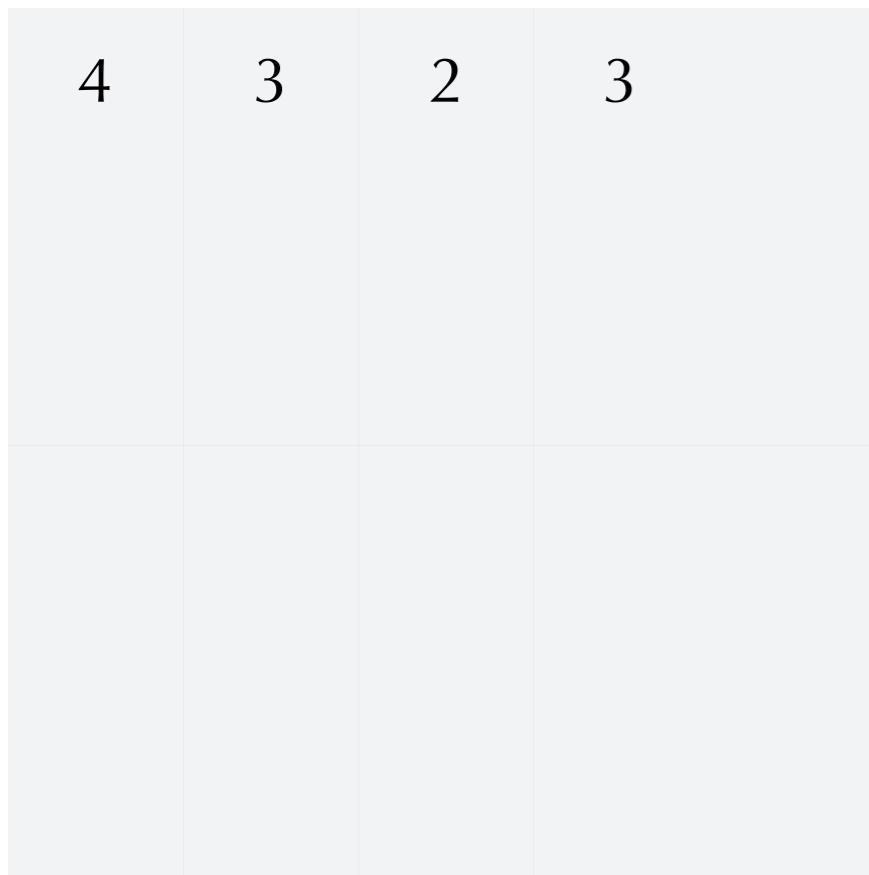
What's the point?

0,25	0,5	0,25
0,5	1	0,5
0,25	0,5	0,25



What's the point?

0,25	0,5	0,25
0,5	1	0,5
0,25	0,5	0,25



What's the point?

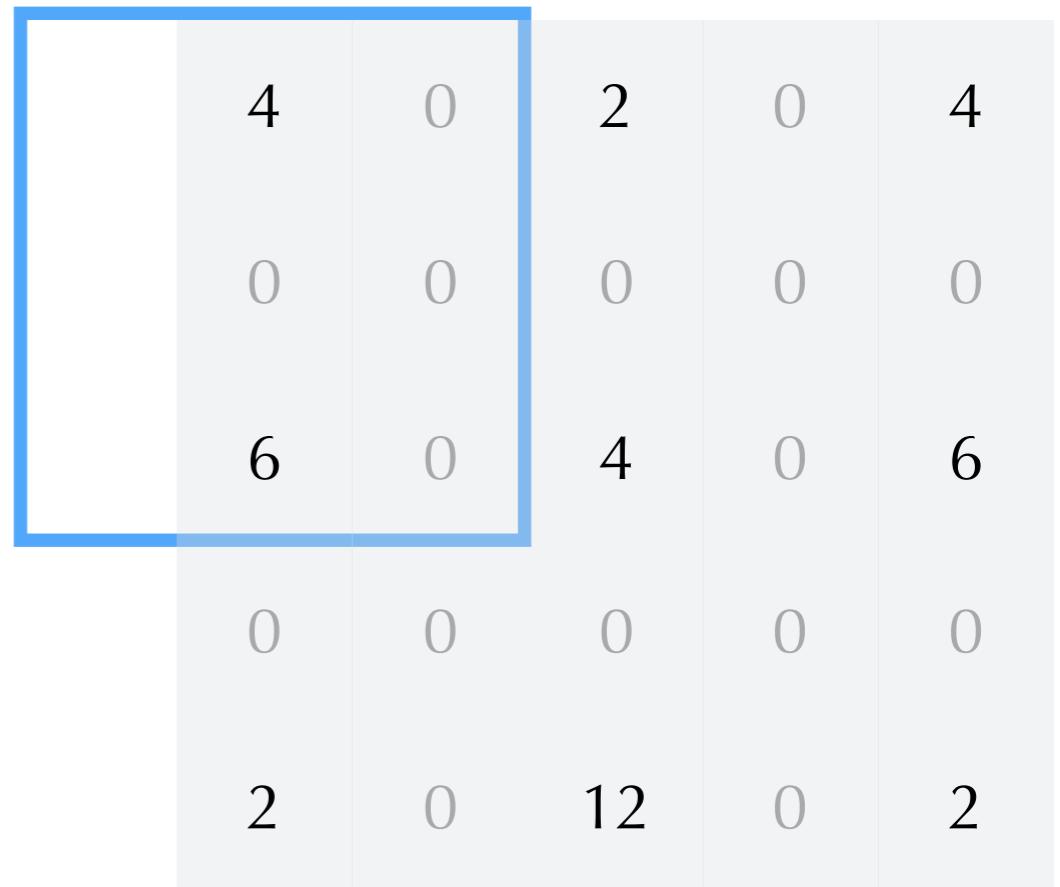
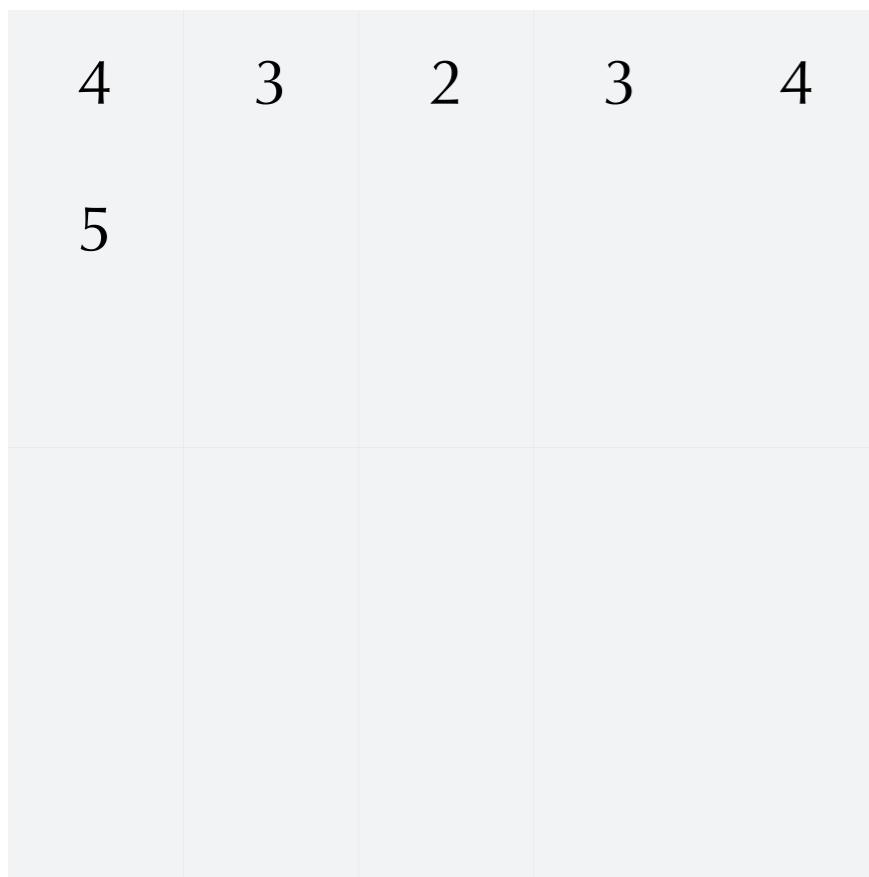
0,25	0,5	0,25
0,5	1	0,5
0,25	0,5	0,25

4 3 2 3 4

4	0	2	0	4
0	0	0	0	0
6	0	4	0	6
0	0	0	0	0
2	0	12	0	2

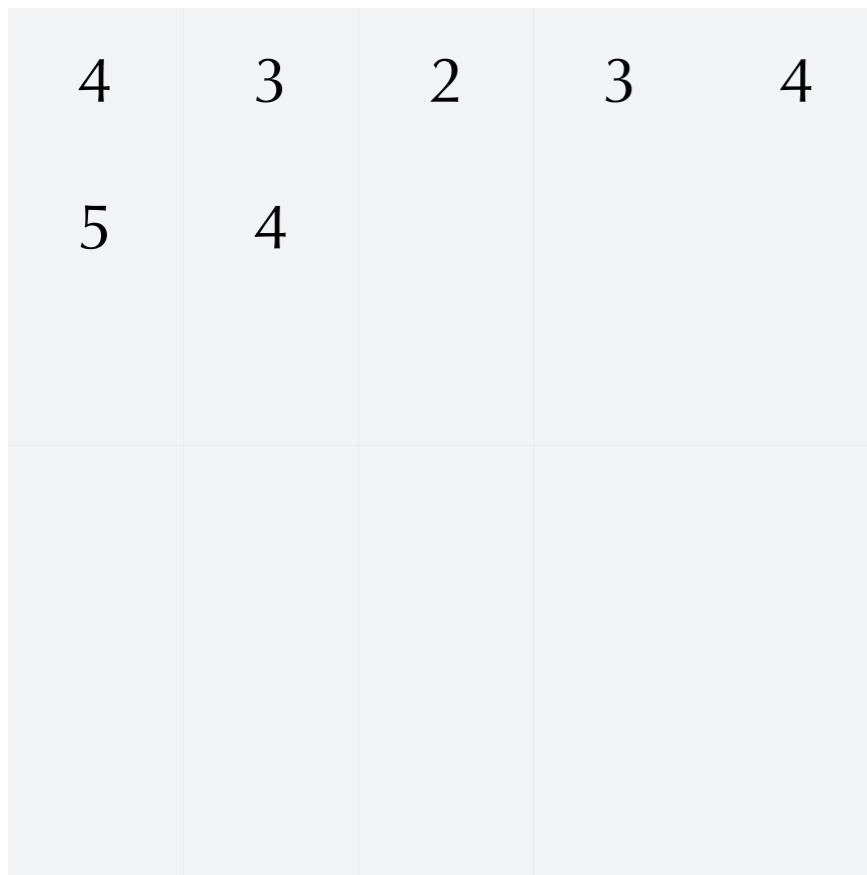
What's the point?

0,25	0,5	0,25
0,5	1	0,5
0,25	0,5	0,25



What's the point?

0,25	0,5	0,25
0,5	1	0,5
0,25	0,5	0,25



4	0	2	0	4
0	0	0	0	0
6	0	4	0	6
0	0	0	0	0
2	0	12	0	2

What's the point?

0,25	0,5	0,25
0,5	1	0,5
0,25	0,5	0,25

4	3	2	3	4
5	4	3	4	5
6	5	4	5	6
4	6	8	6	4
2	7	12	7	2

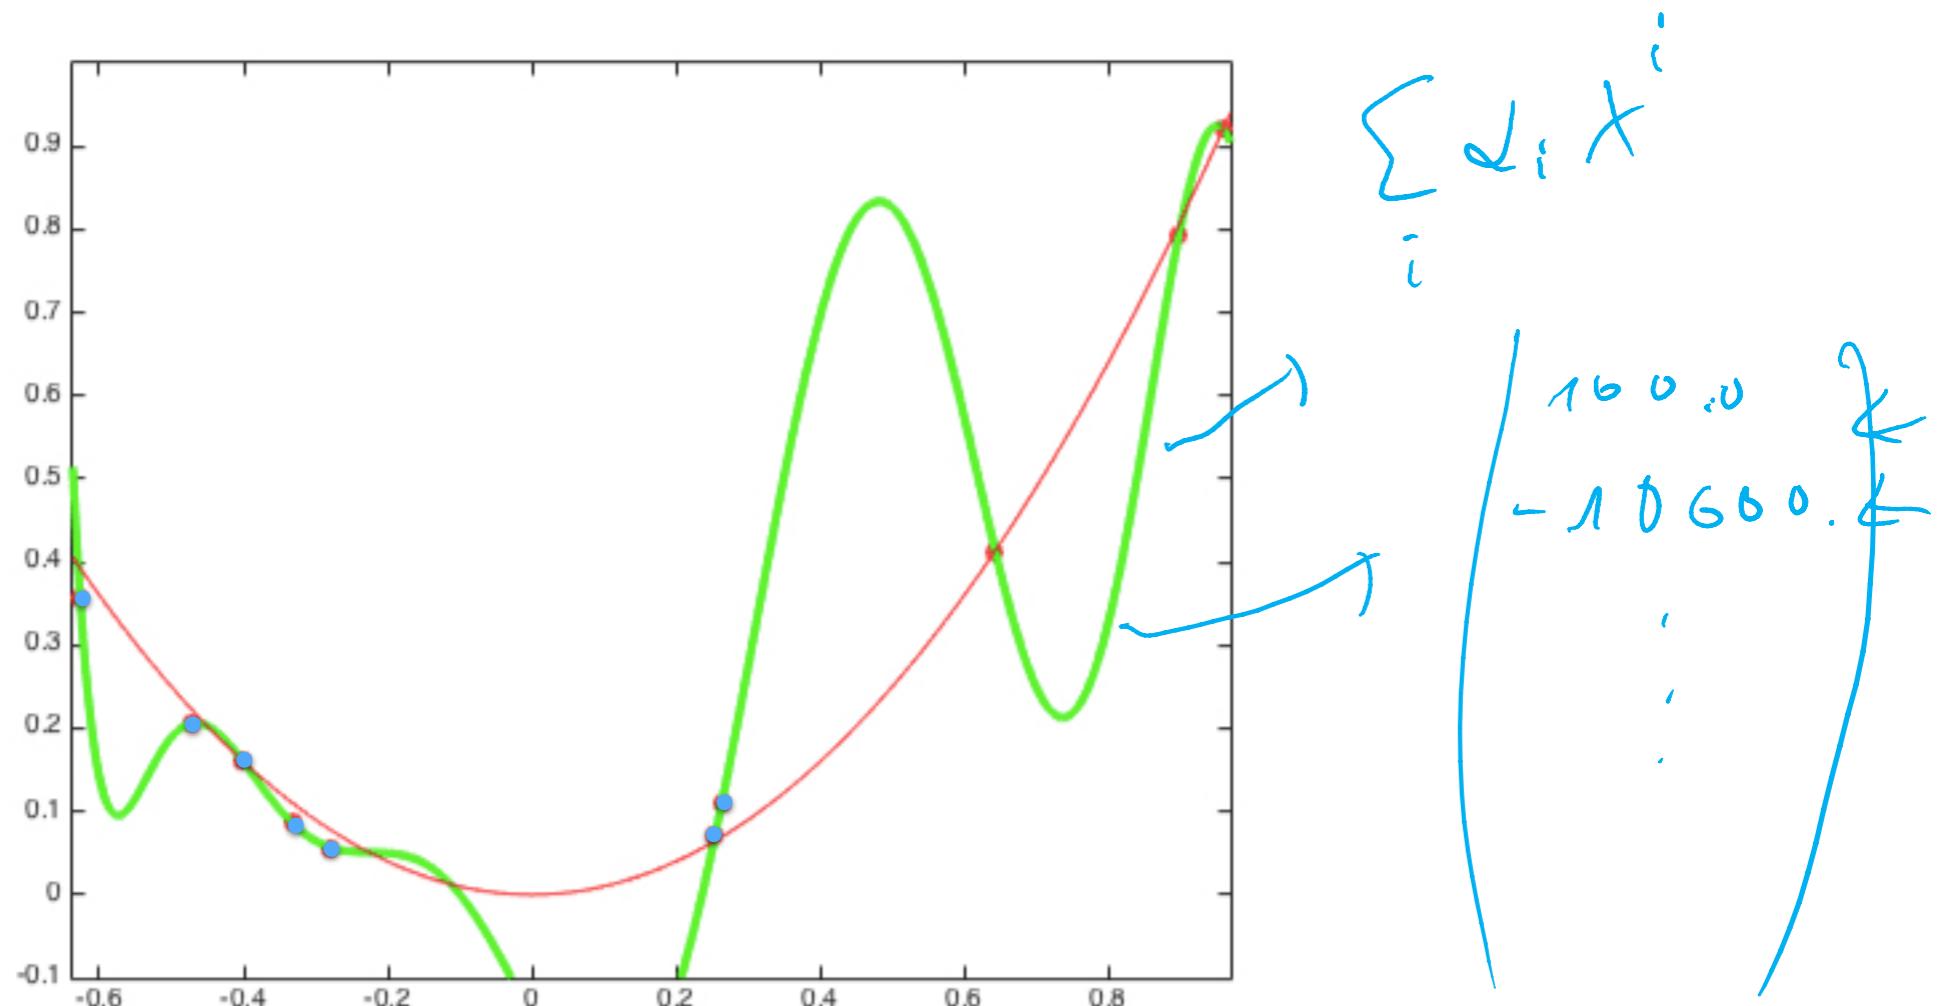
4	0	2	0	4
0	0	0	0	0
6	0	4	0	6
0	0	0	0	0
2	0	12	0	2

Transposed convolution

Overfitting, Part II

Overfitting

10 data points - Fitting 10th degree polynomial



$$y = x^2$$

Overfitting

VGG-16 architecture



224

224

Overfitting

VGG-16 architecture



224



conv 1

224

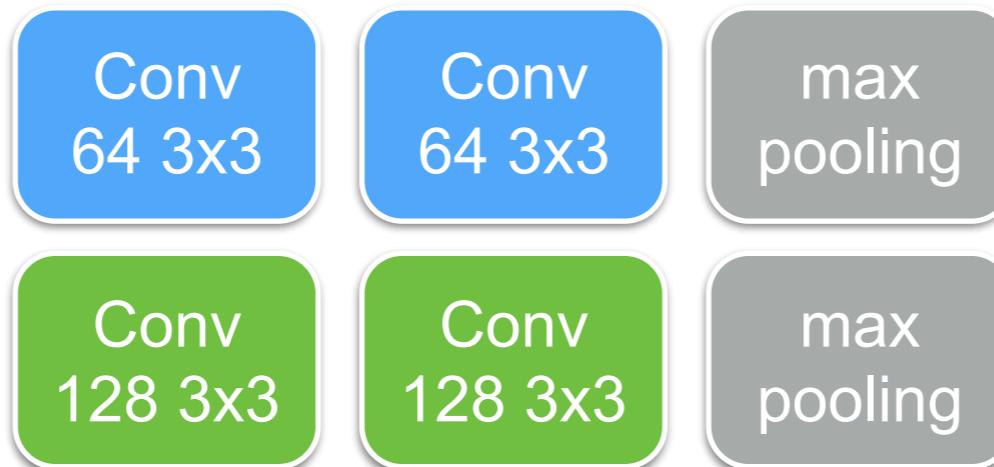
Overfitting

VGG-16 architecture



224

224



conv 1

conv 2

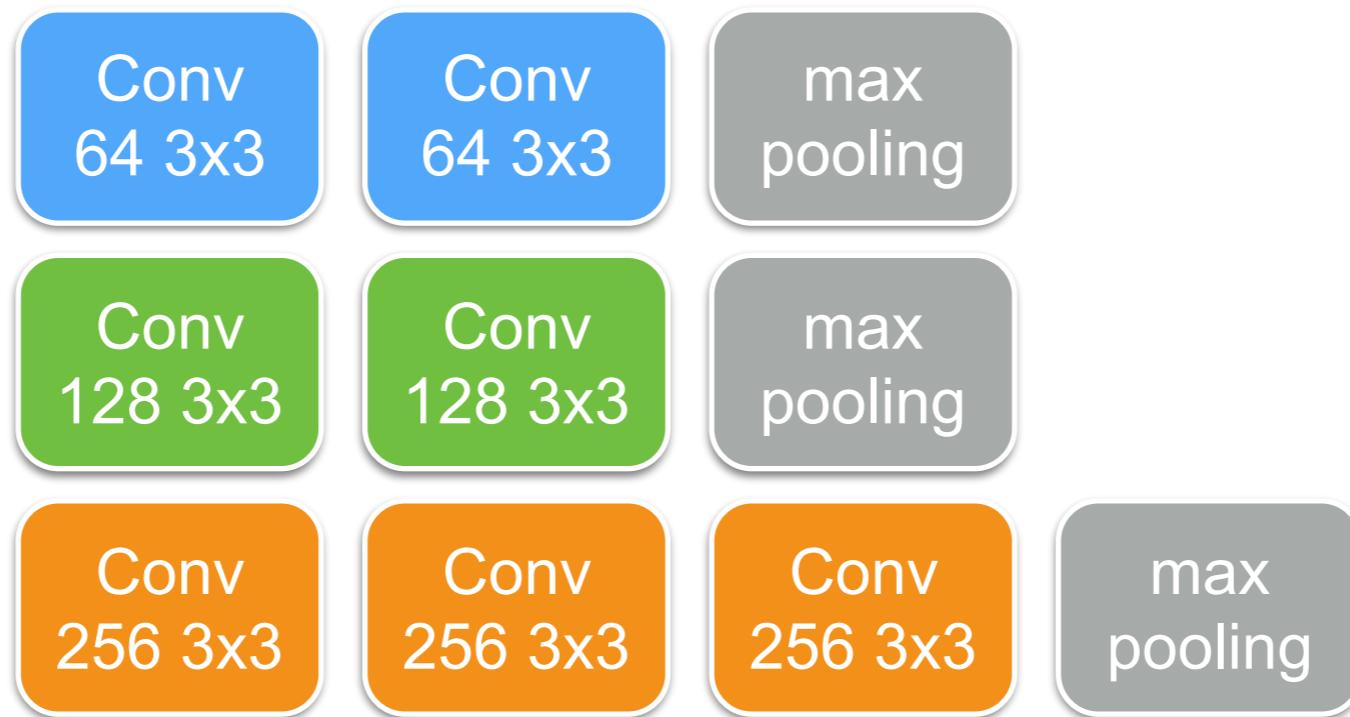
Overfitting

VGG-16 architecture



224

224



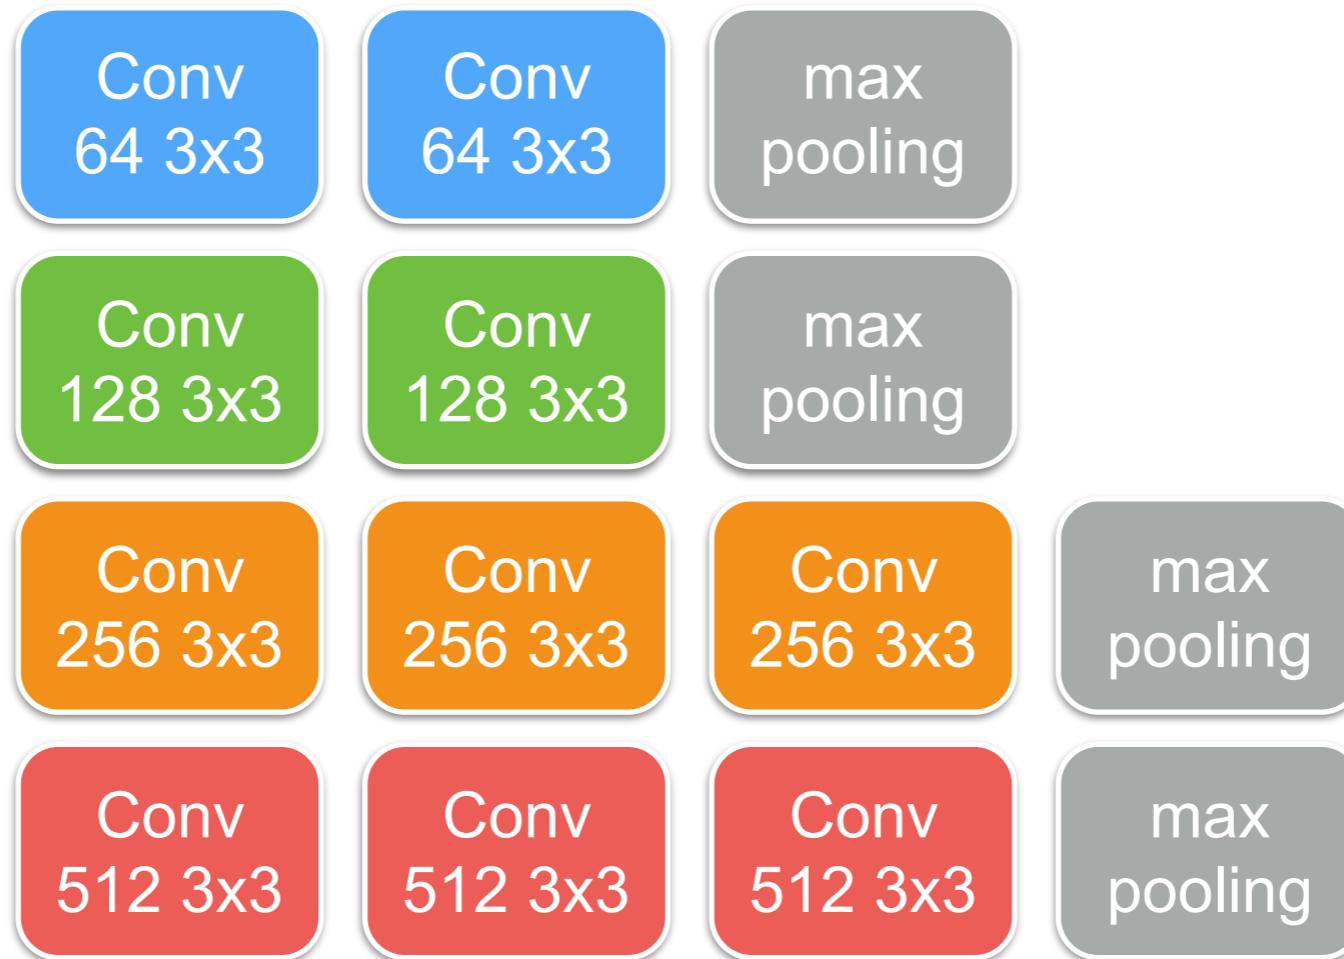
Overfitting

VGG-16 architecture



224

224



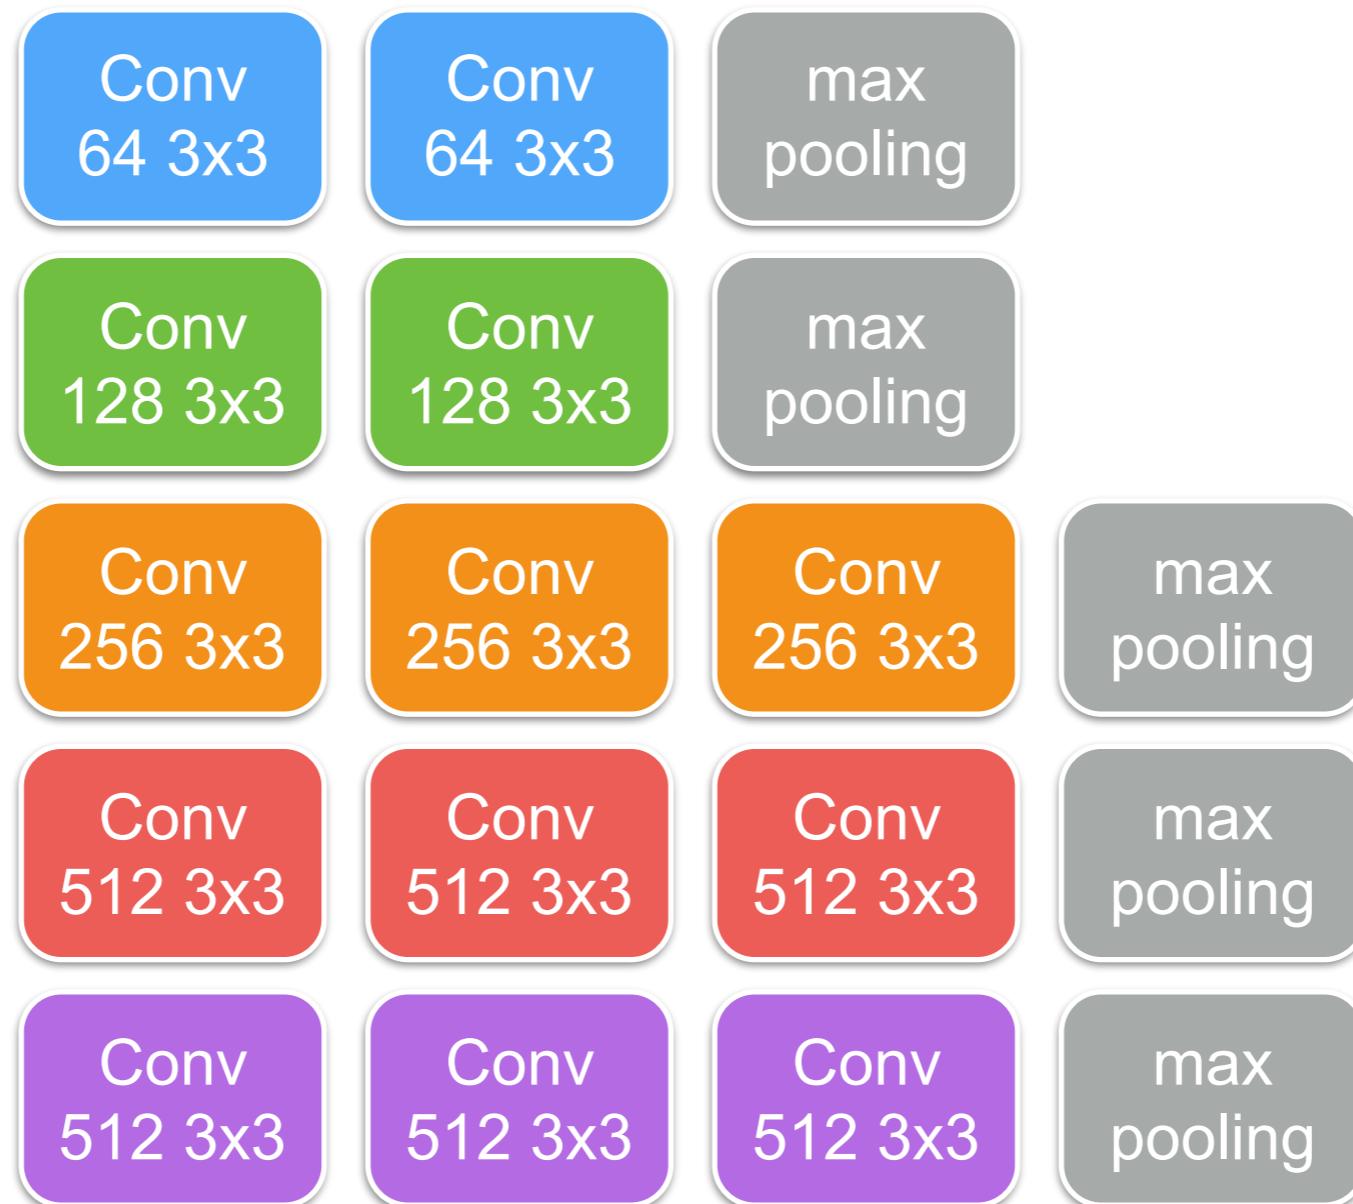
Overfitting

VGG-16 architecture



224

224



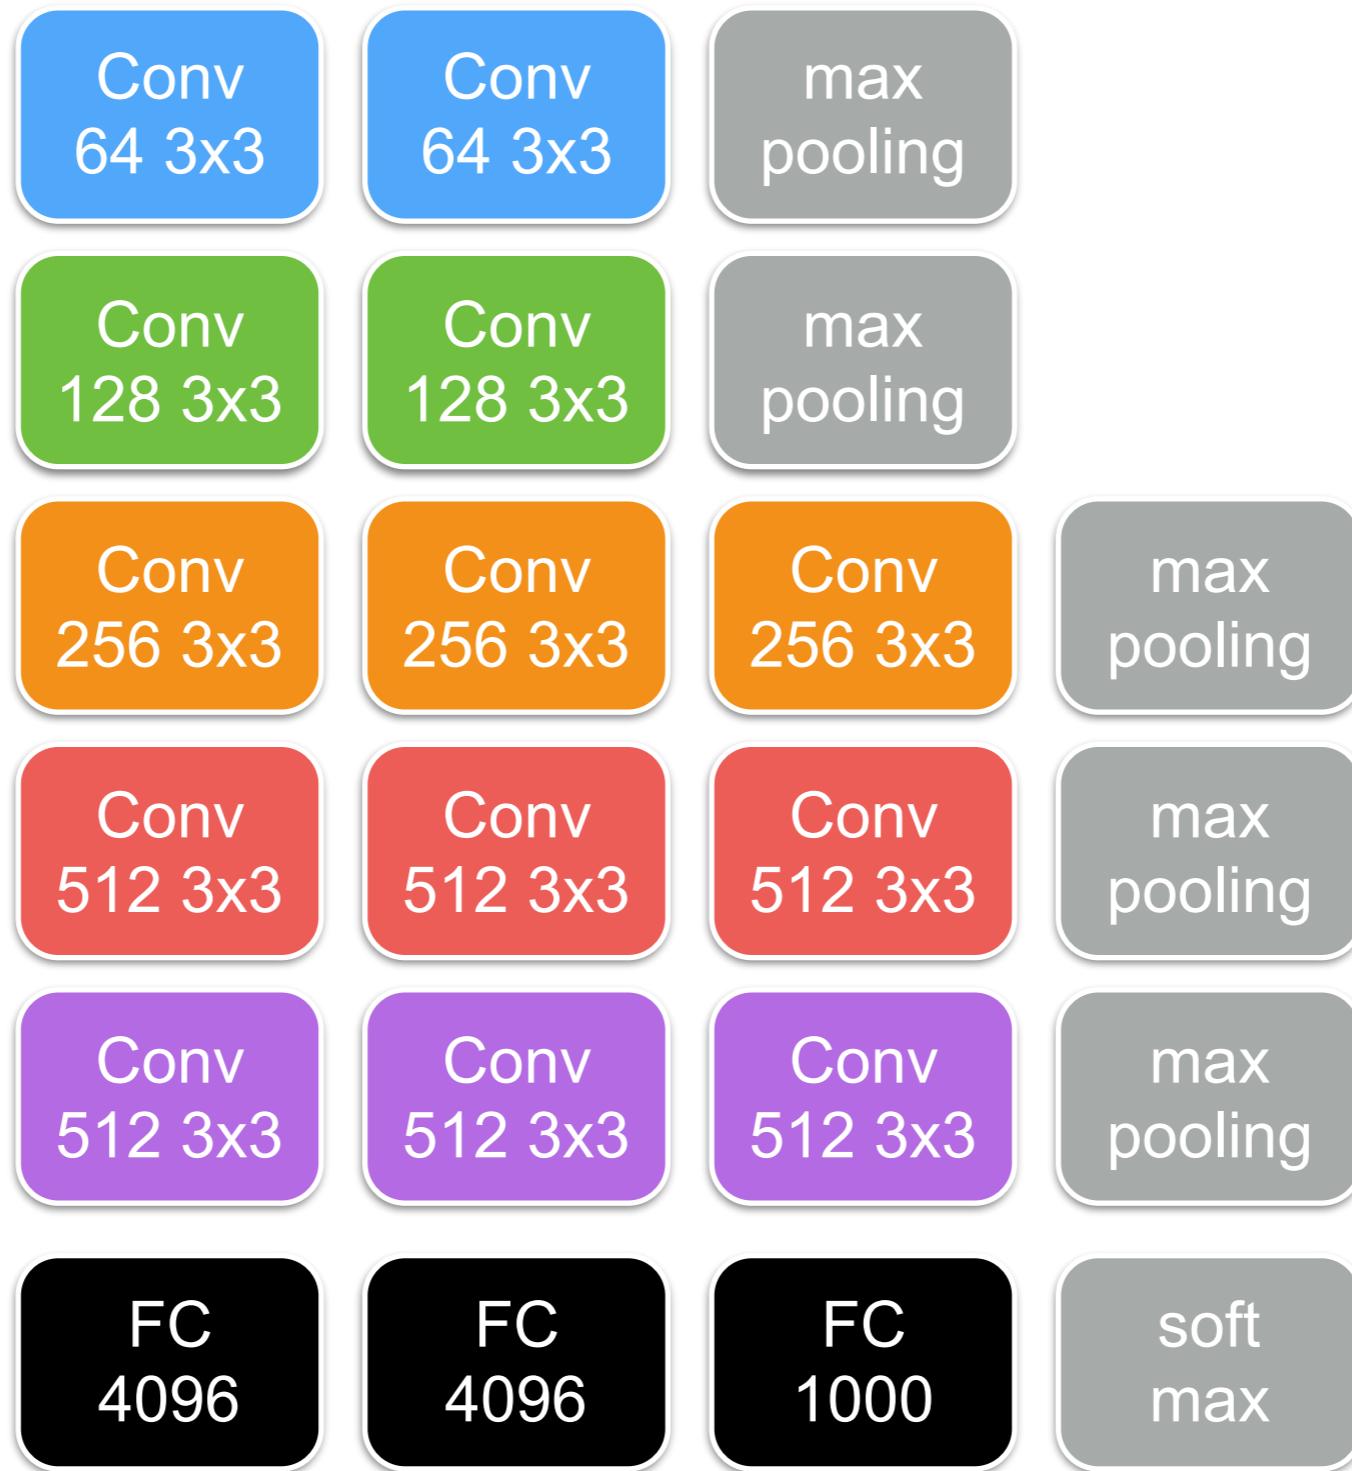
Overfitting

VGG-16 architecture



224

224



Overfitting

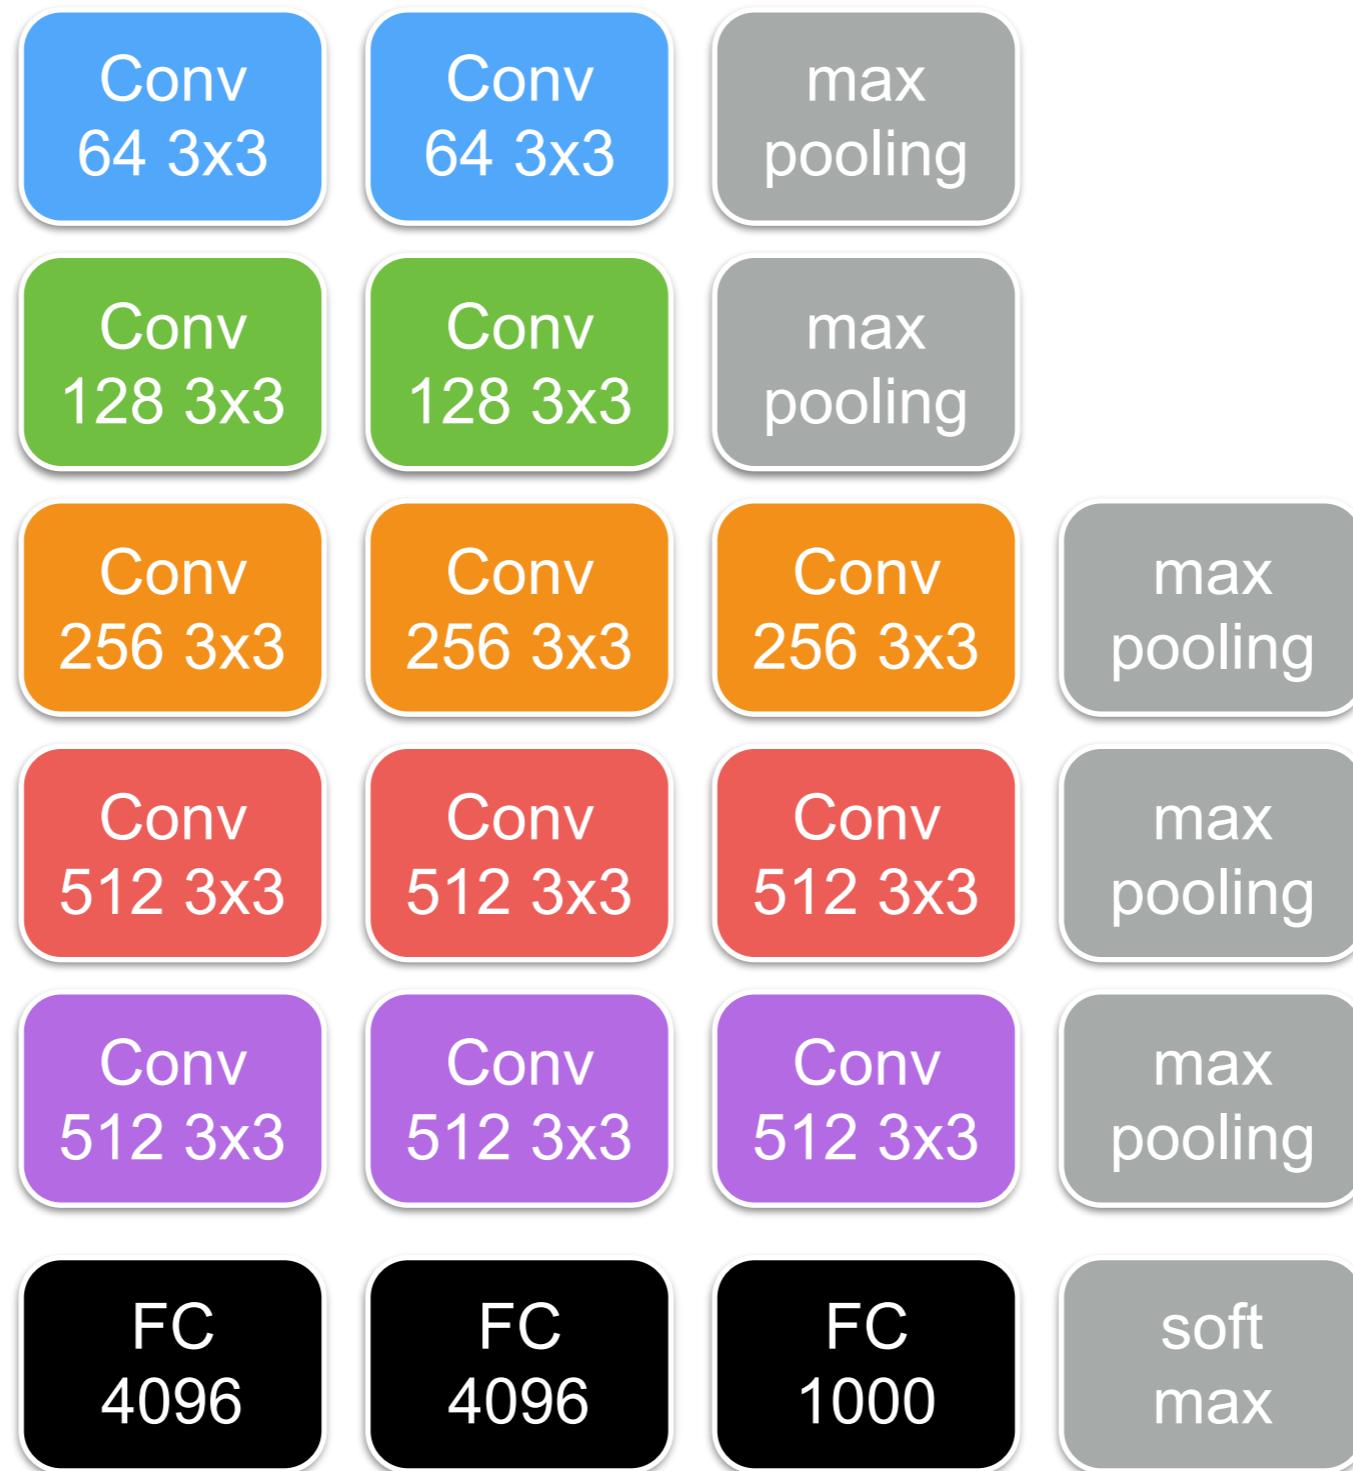
VGG-16 architecture



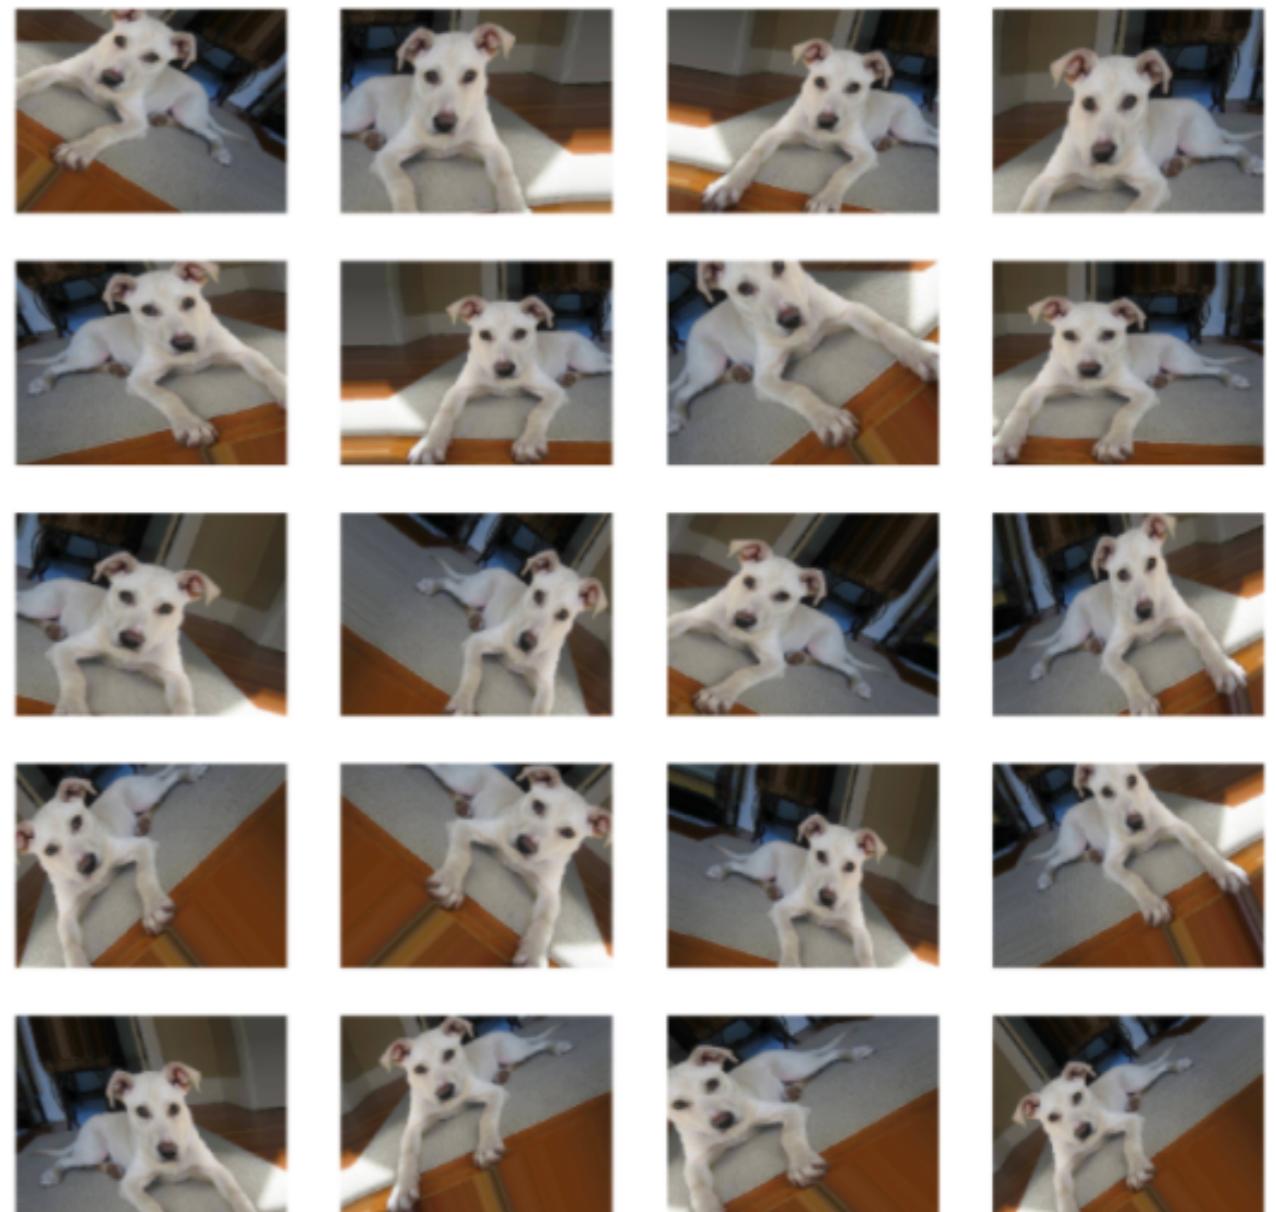
224

224

**138M
parameters!**



Data Augmentation



Regularization

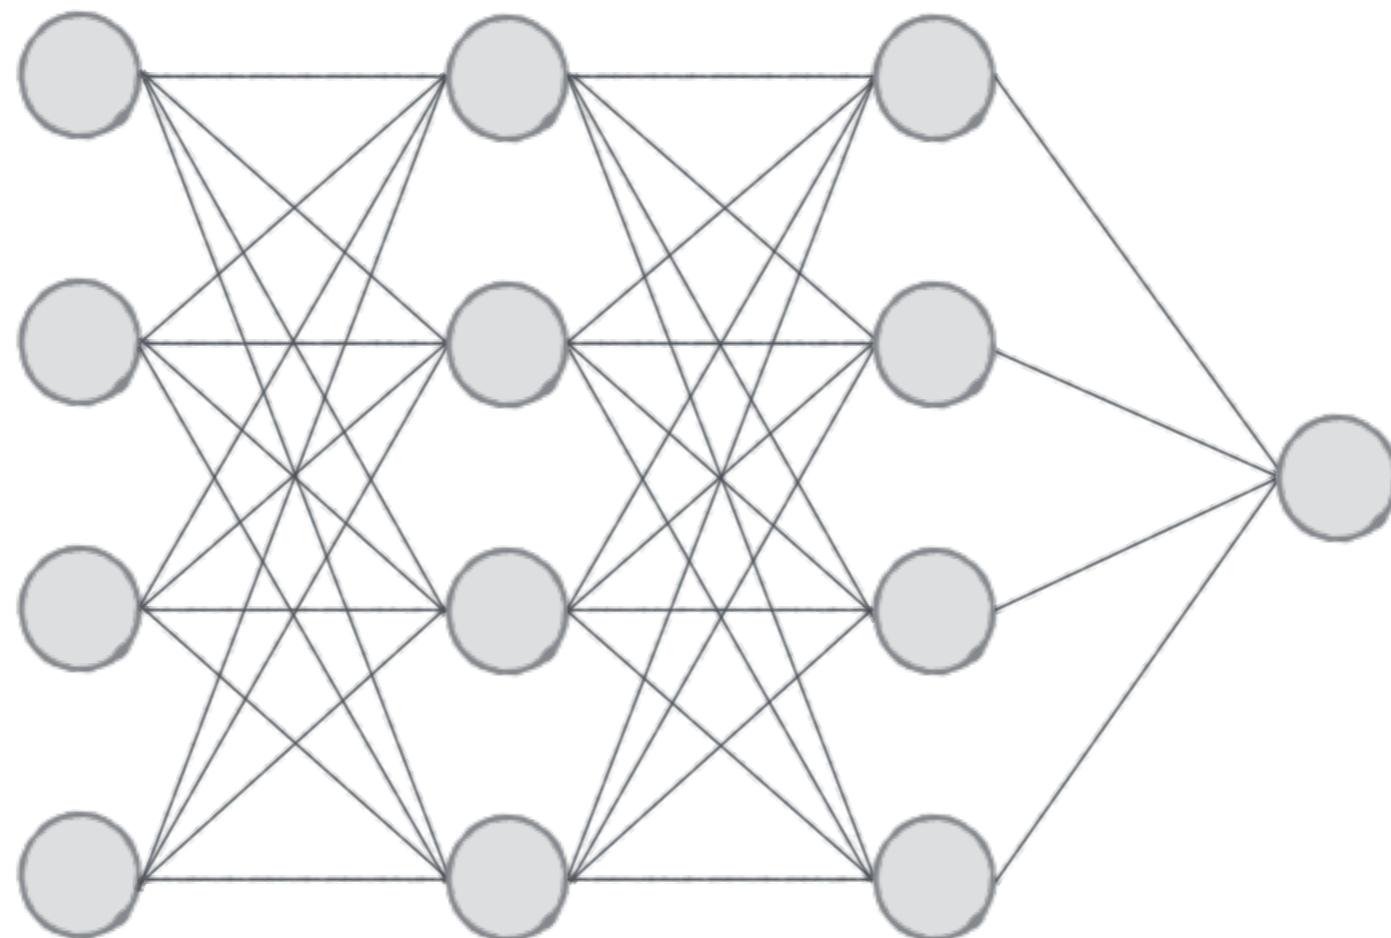
$$\min_{w_1, \dots, w_n} \left(L(w_1, \dots, w_n) + c \sum_i w_i^2 \right)$$

loss

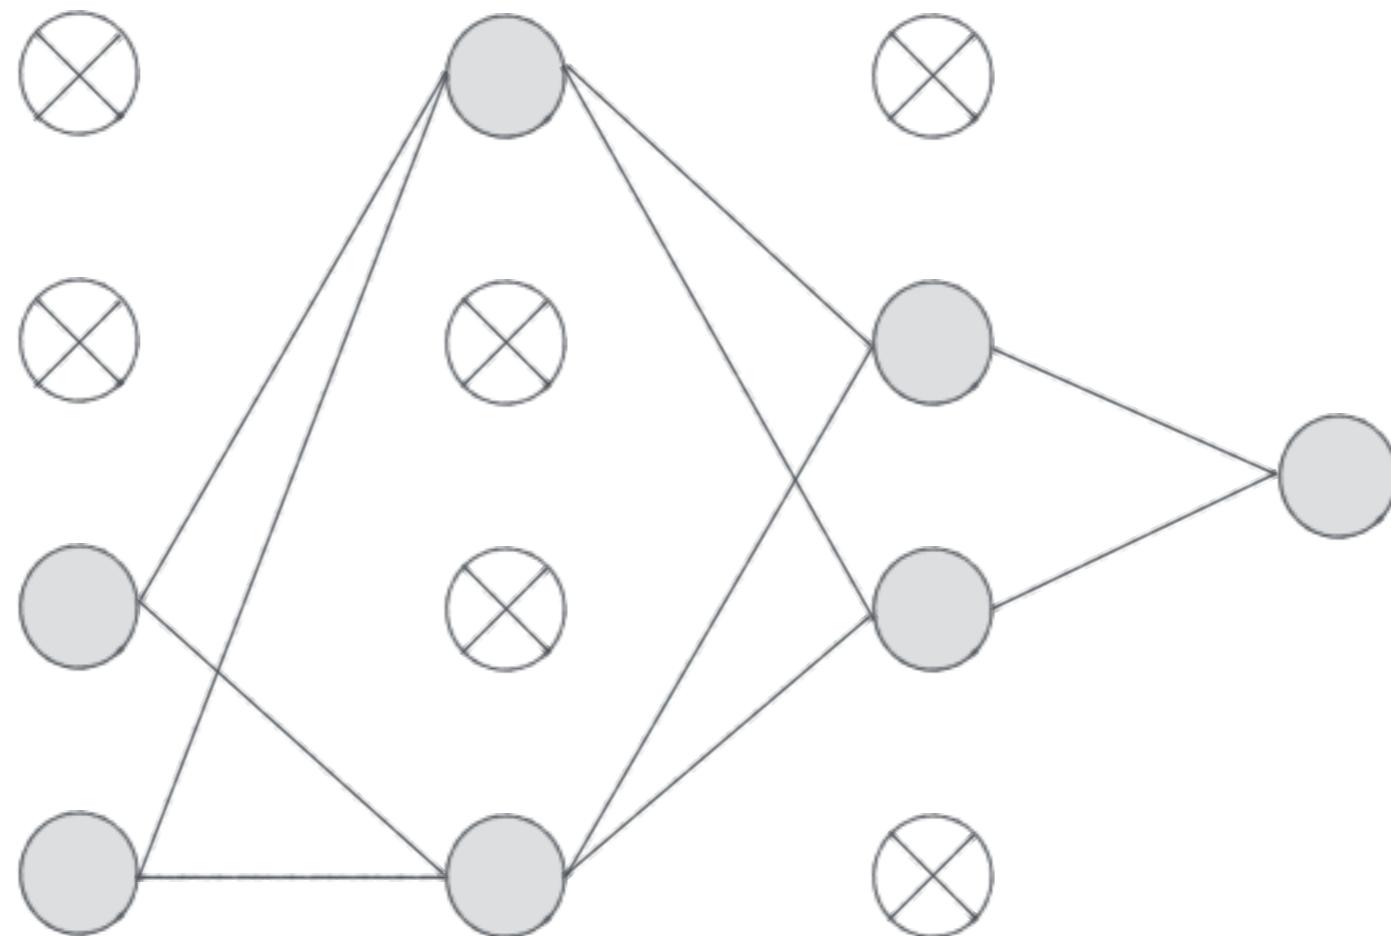
weighting term

regularization term

Dropout



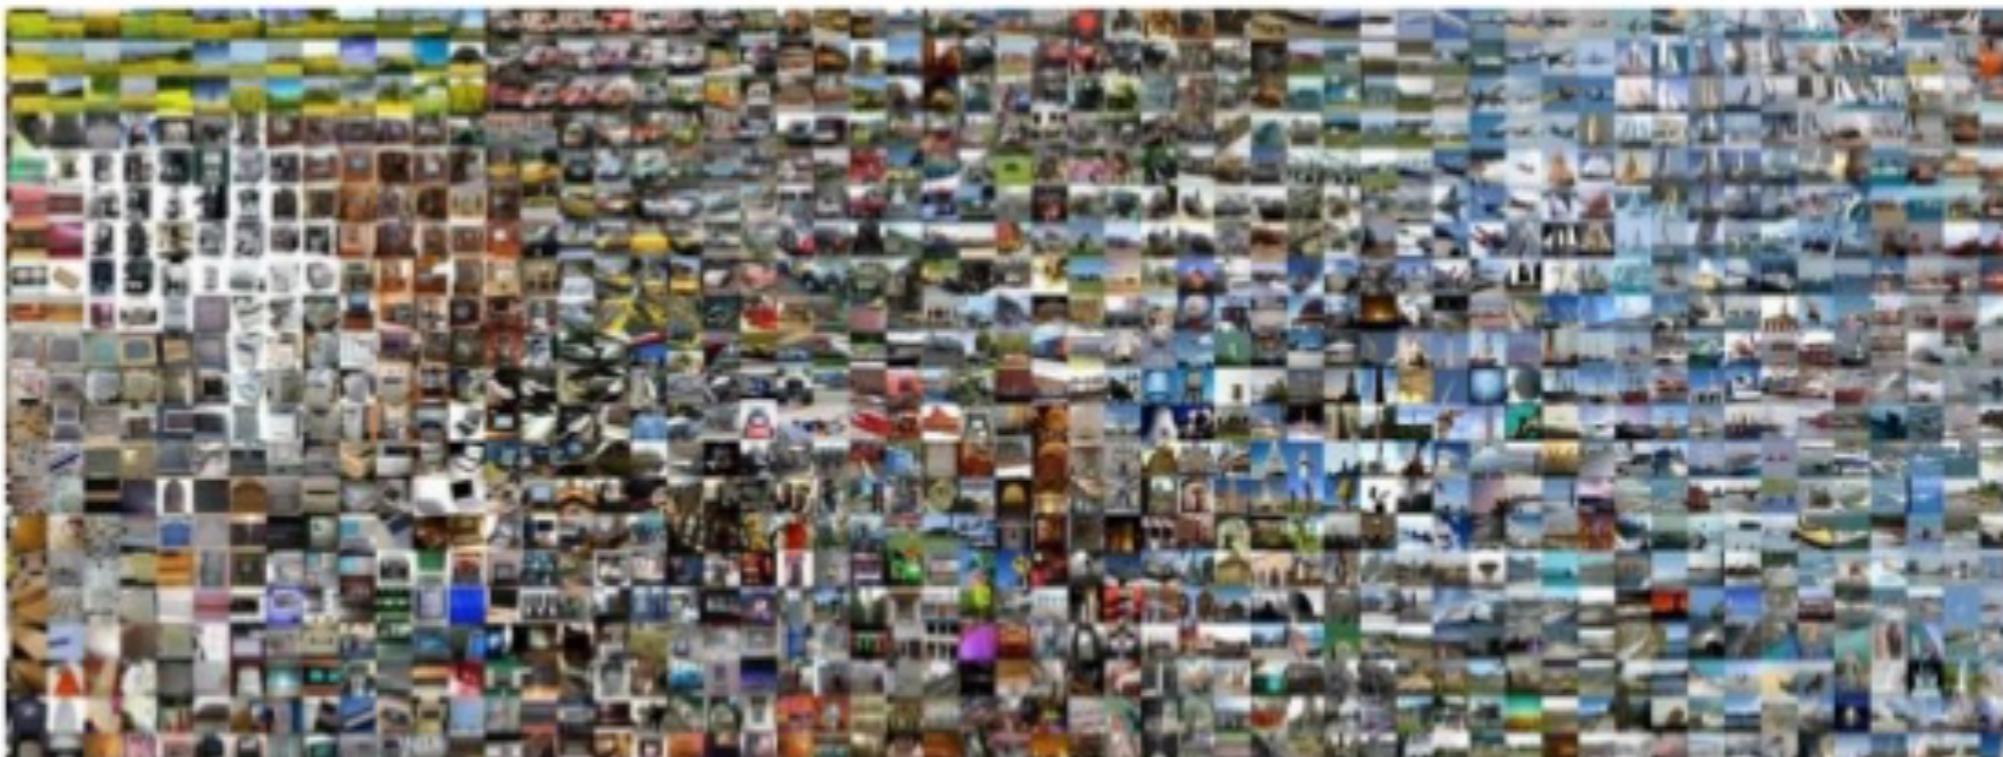
Dropout



Randomly disable neurons for each minibatch

Transfer Learning

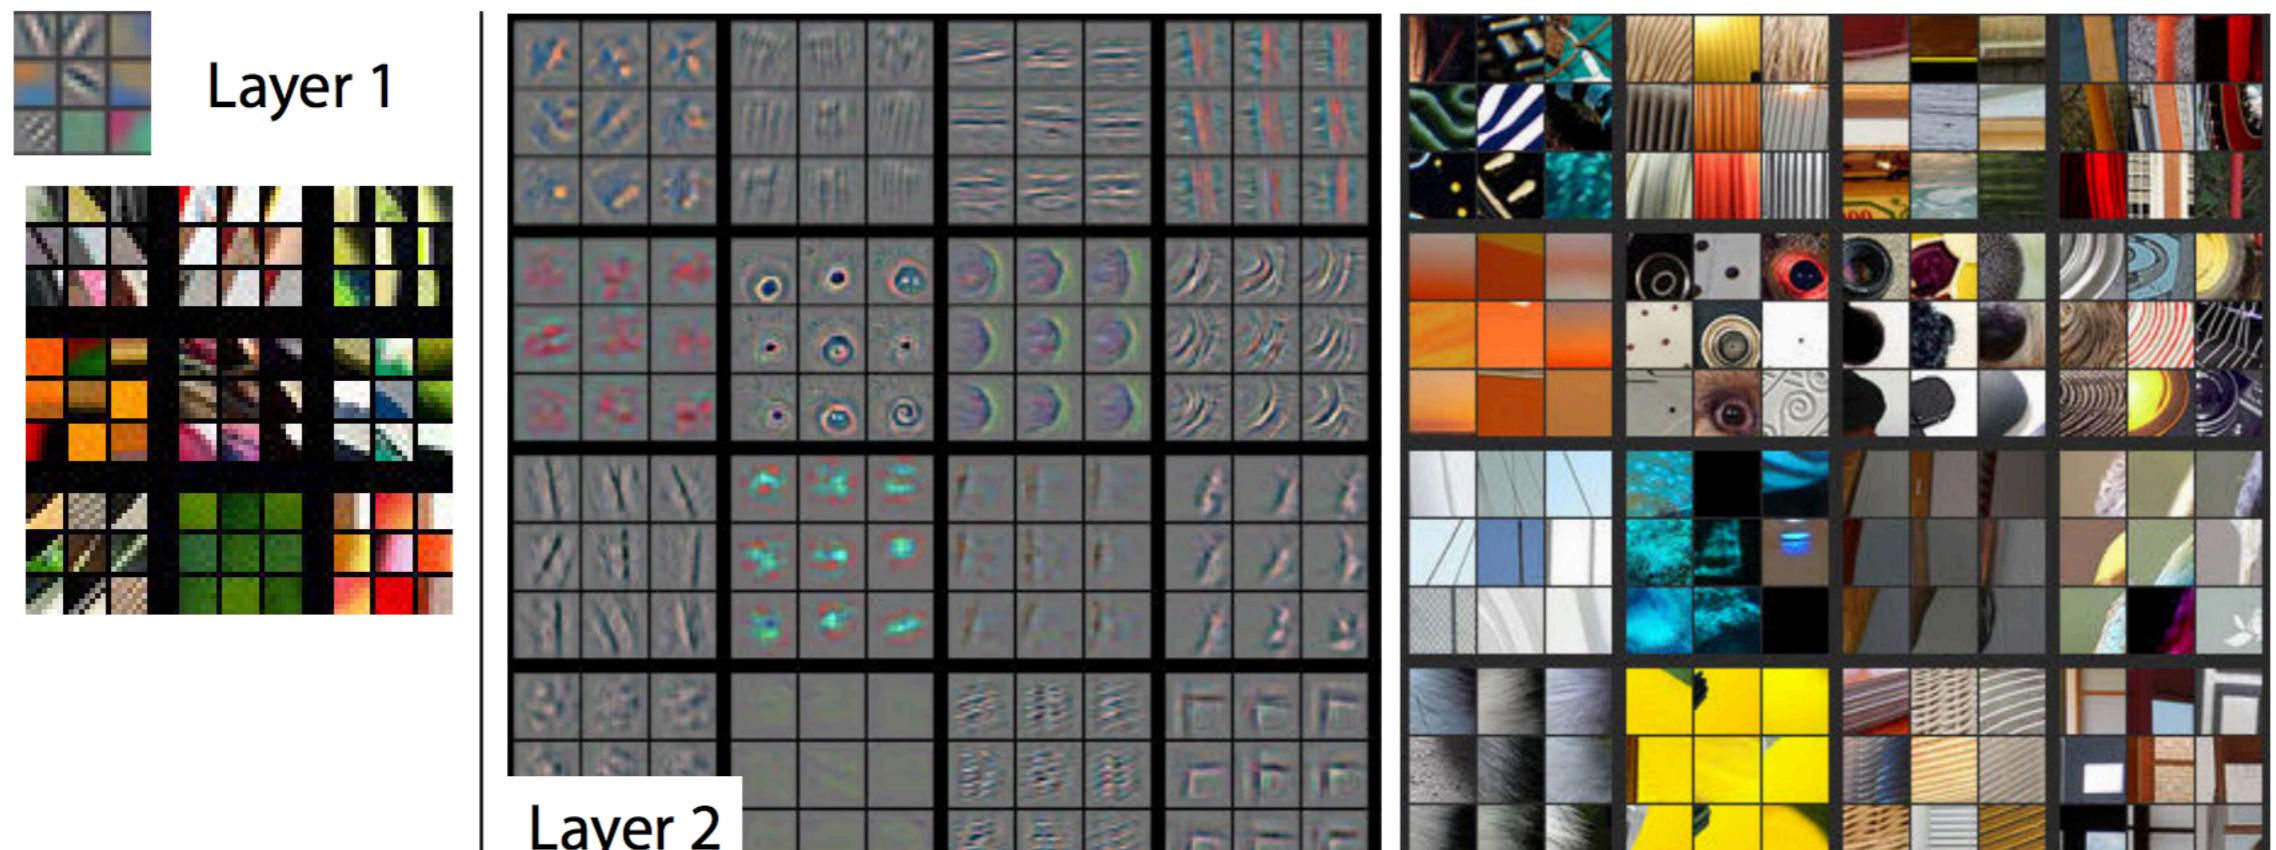
IMAGENET



Russakovsky, O., Deng, J., Su, H., Krause, J., Satheesh, S., Ma, S., ... & Fei-Fei, L. (2015). [Imagenet large scale visual recognition challenge](#). *arXiv preprint arXiv:1409.0575*. [\[web\]](#)

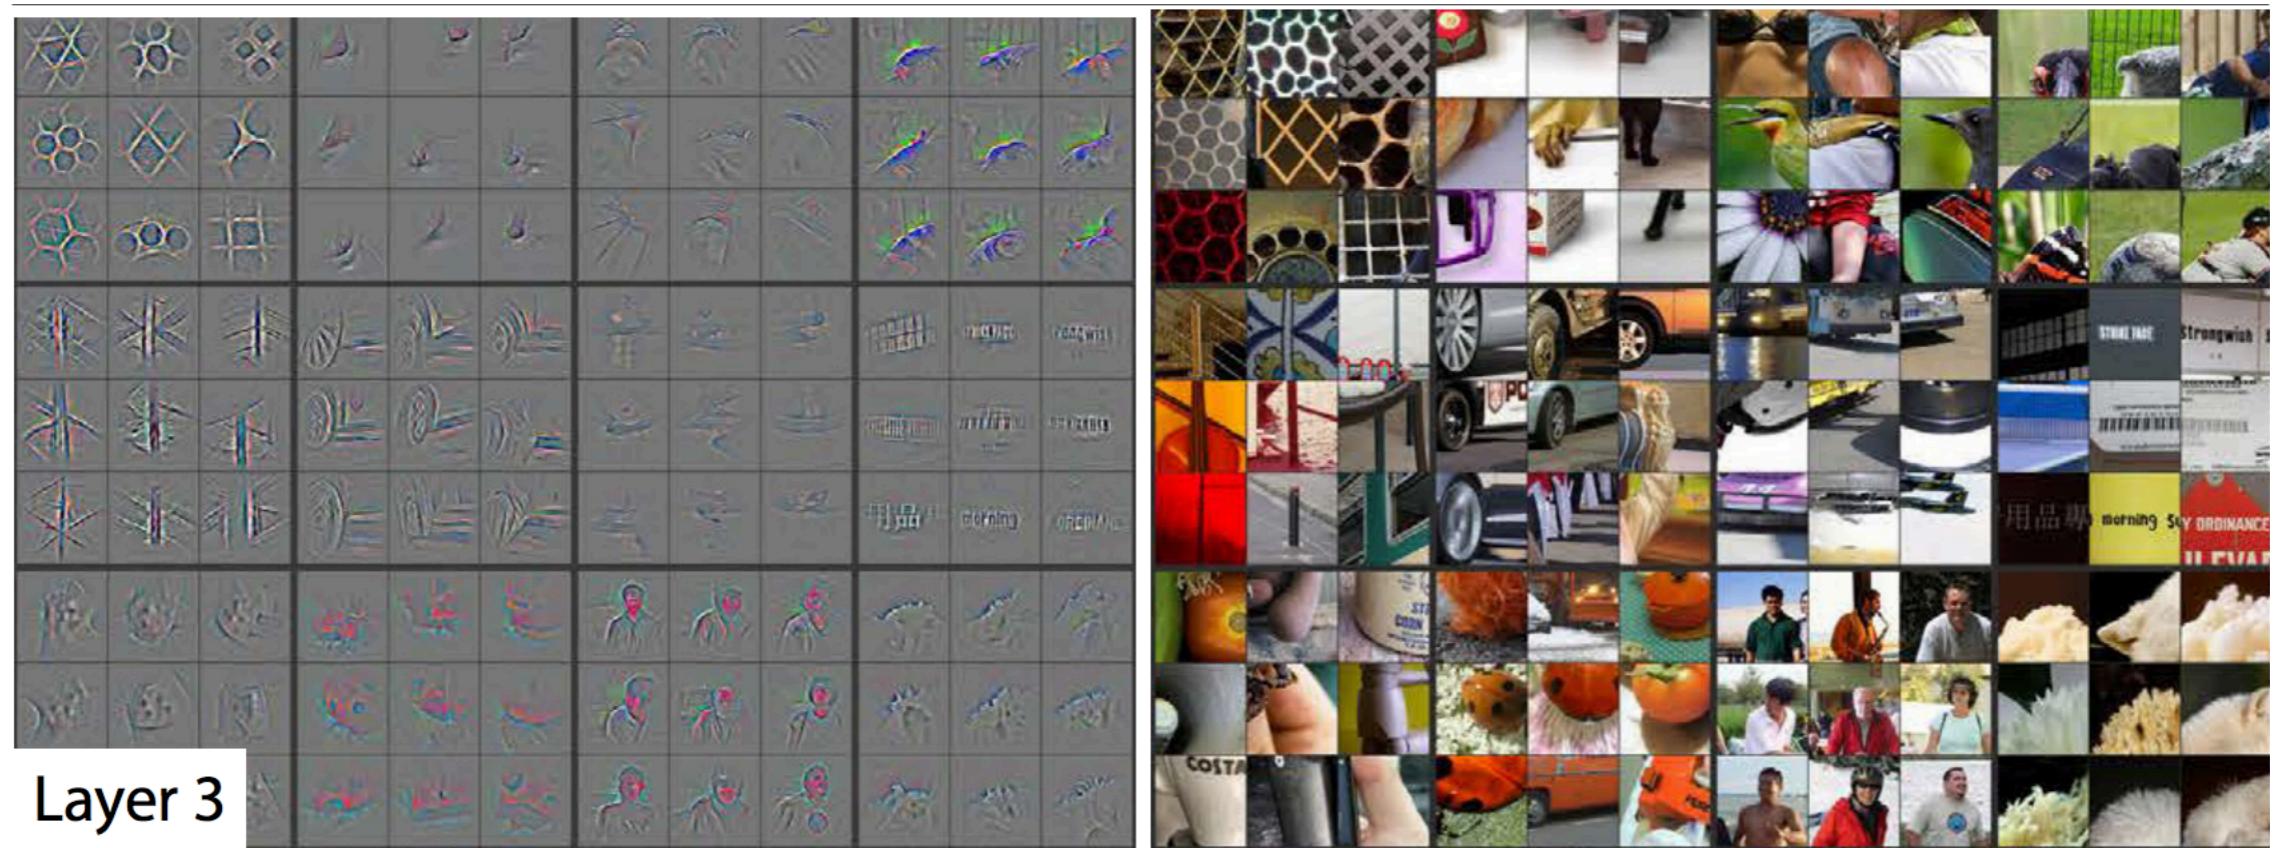
14 million images of 1000 classes

Transfer Learning



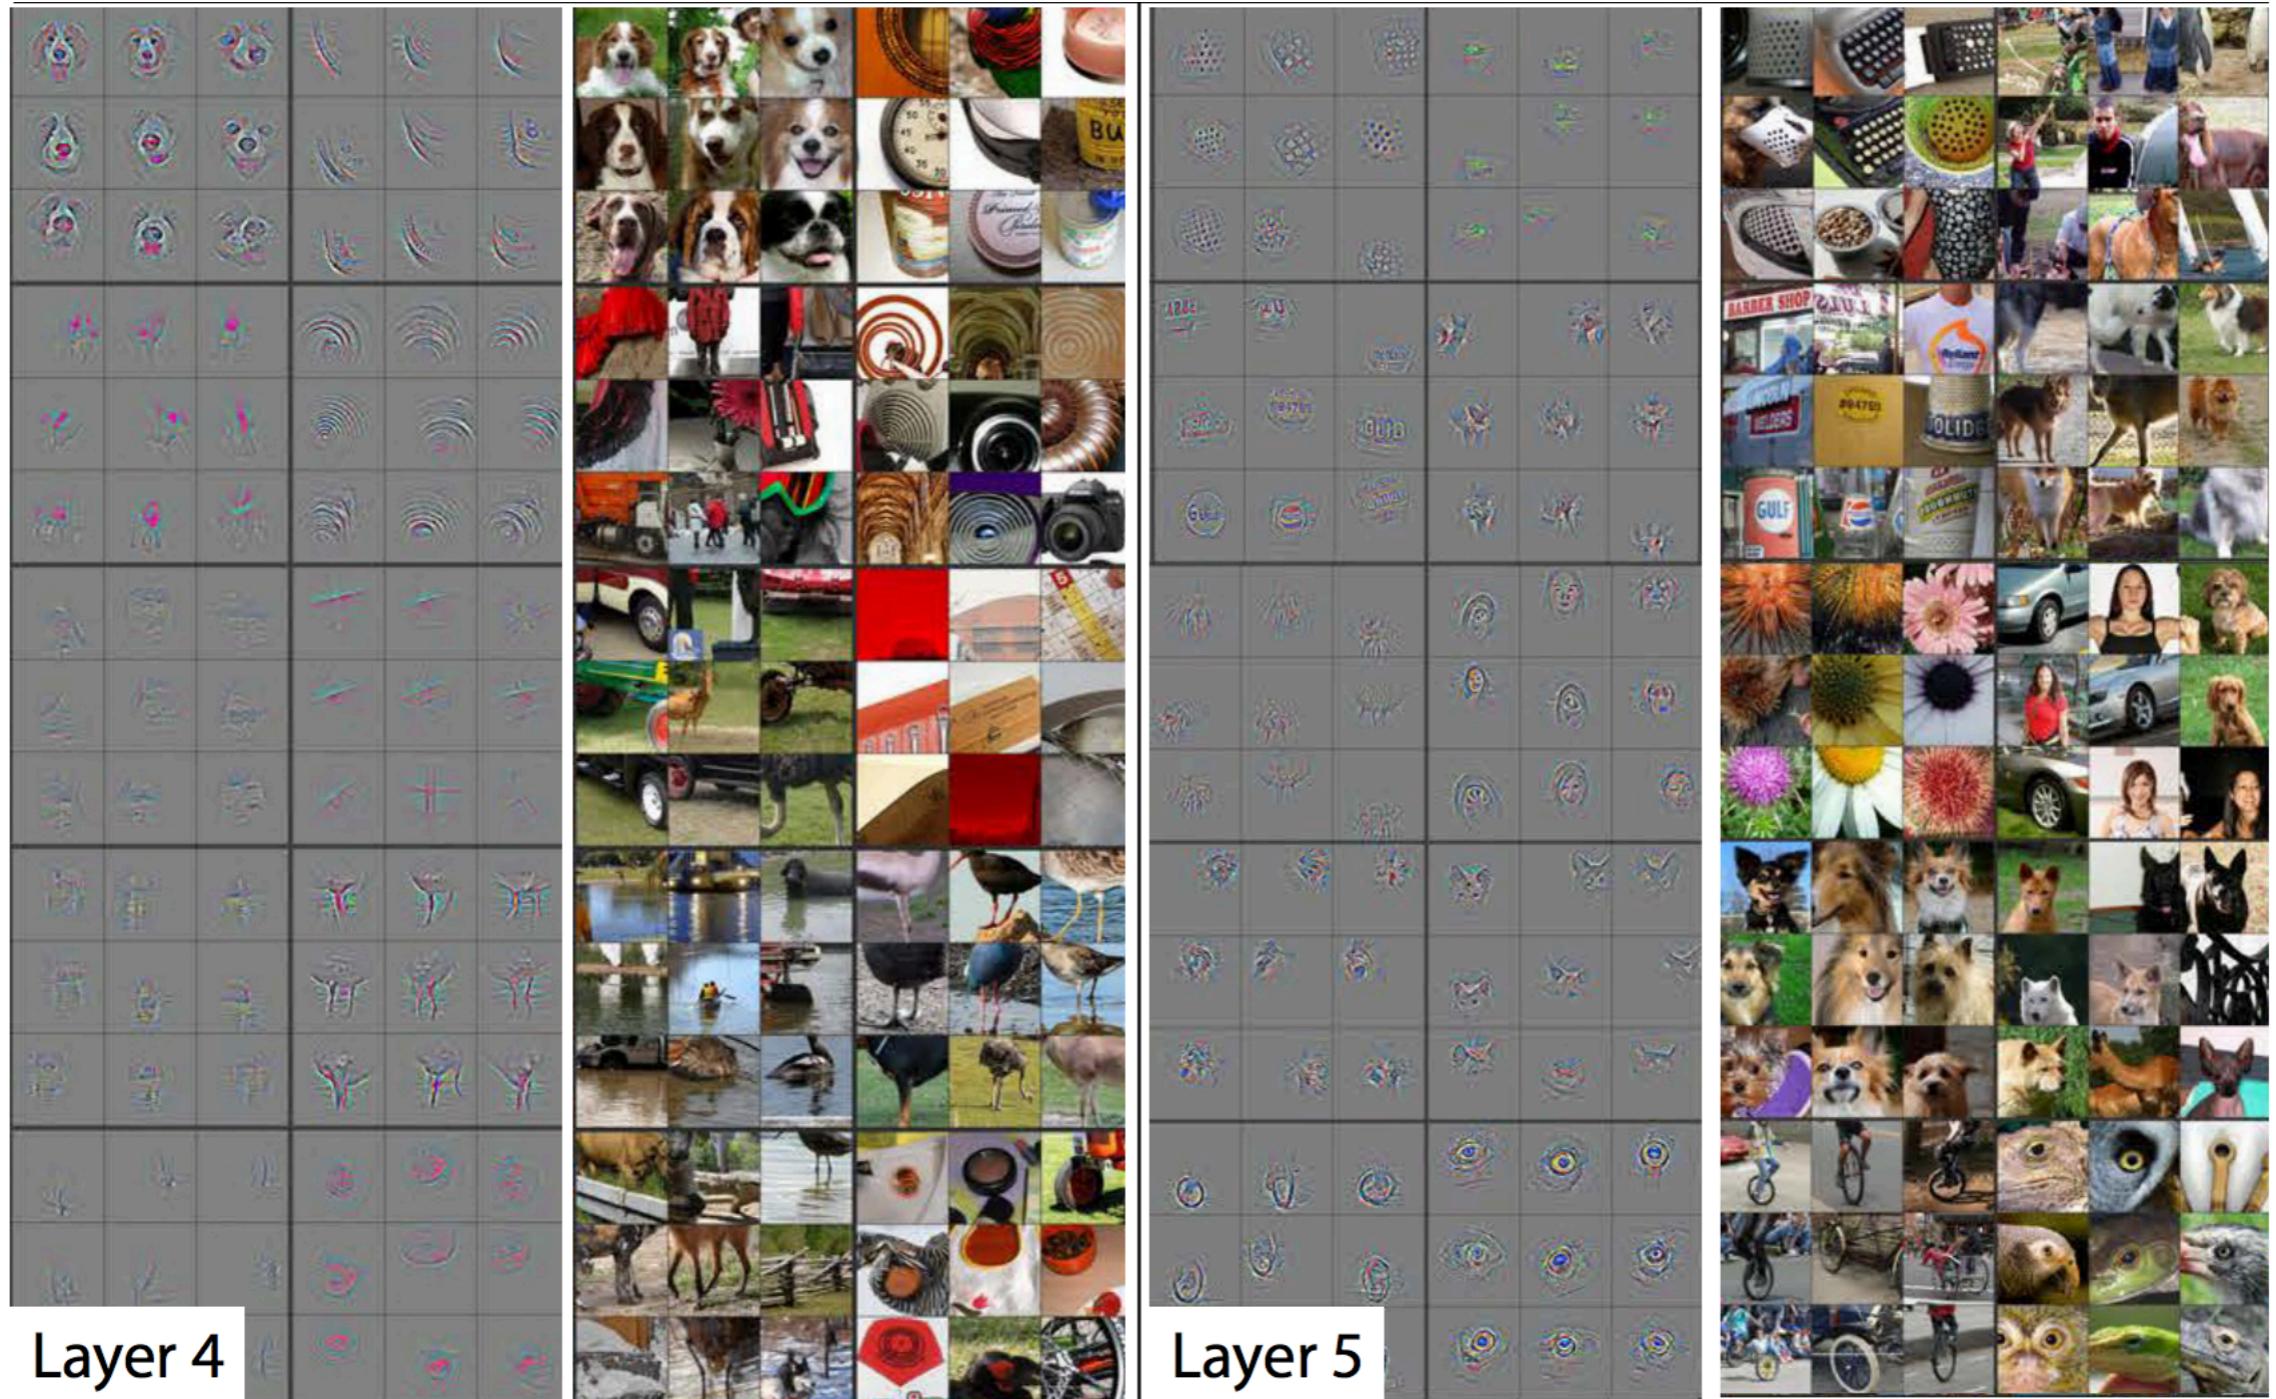
Low-level features: corners, edges, ...

Transfer Learning



Mid-level features

Transfer Learning



Higher-level features: object parts & objects

[Zeiler & Fergus, Visualizing and Understanding Convolutional Networks, ECCV 2014]

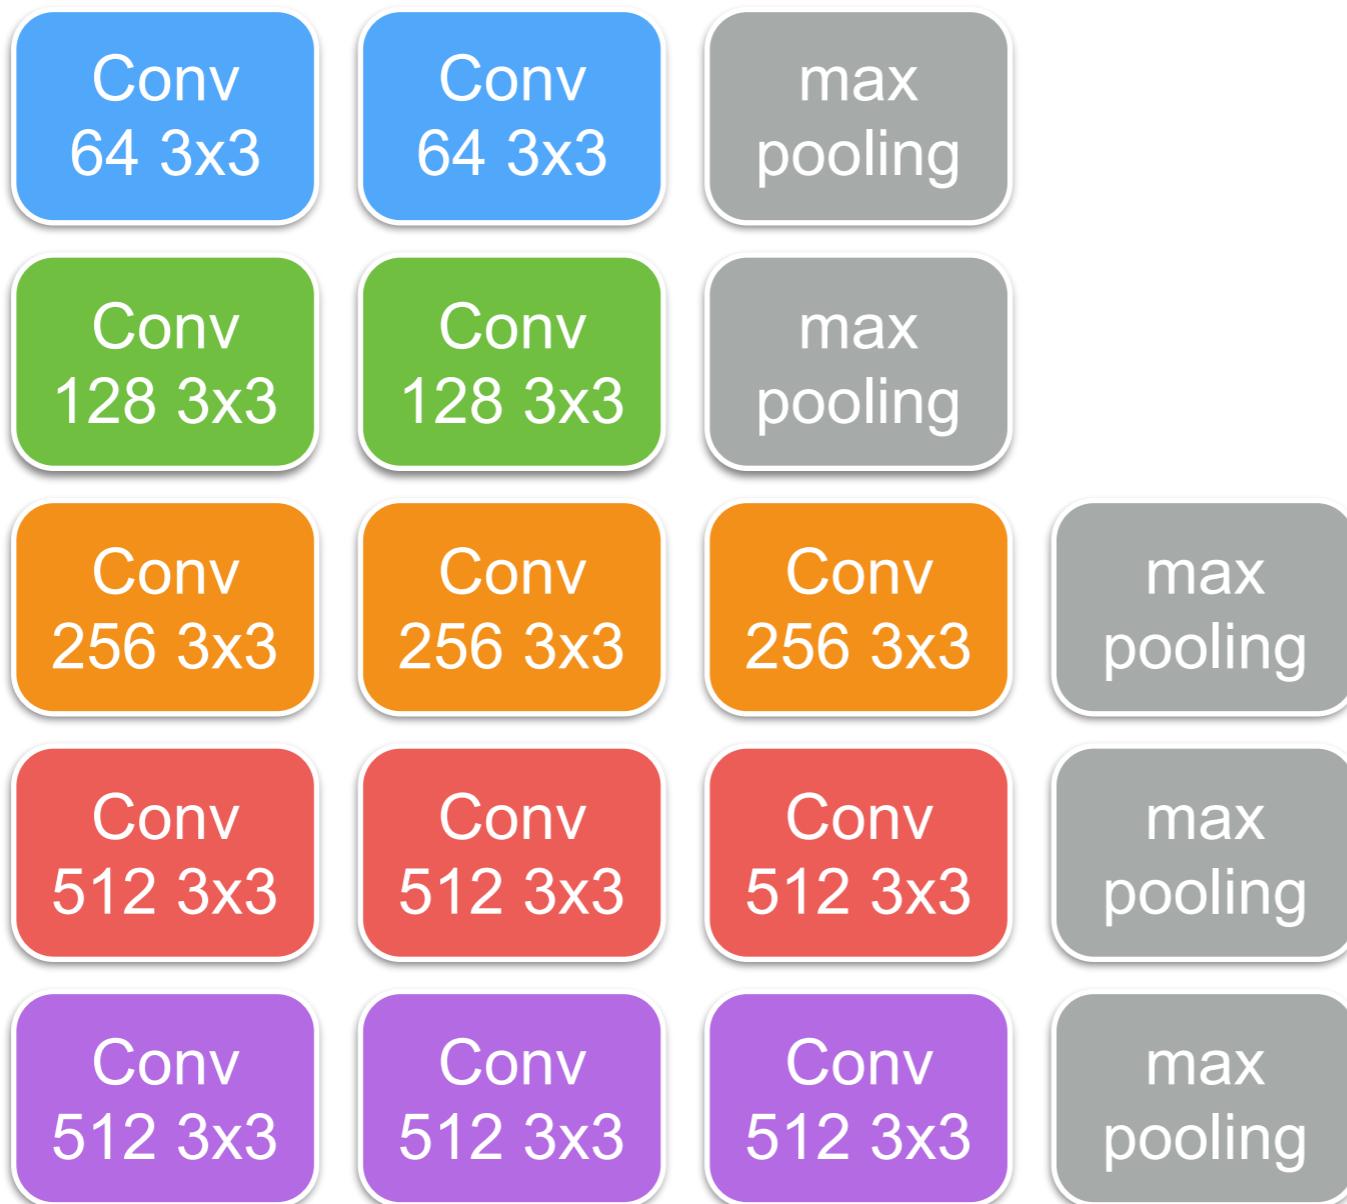
Transfer Learning

VGG-16 architecture



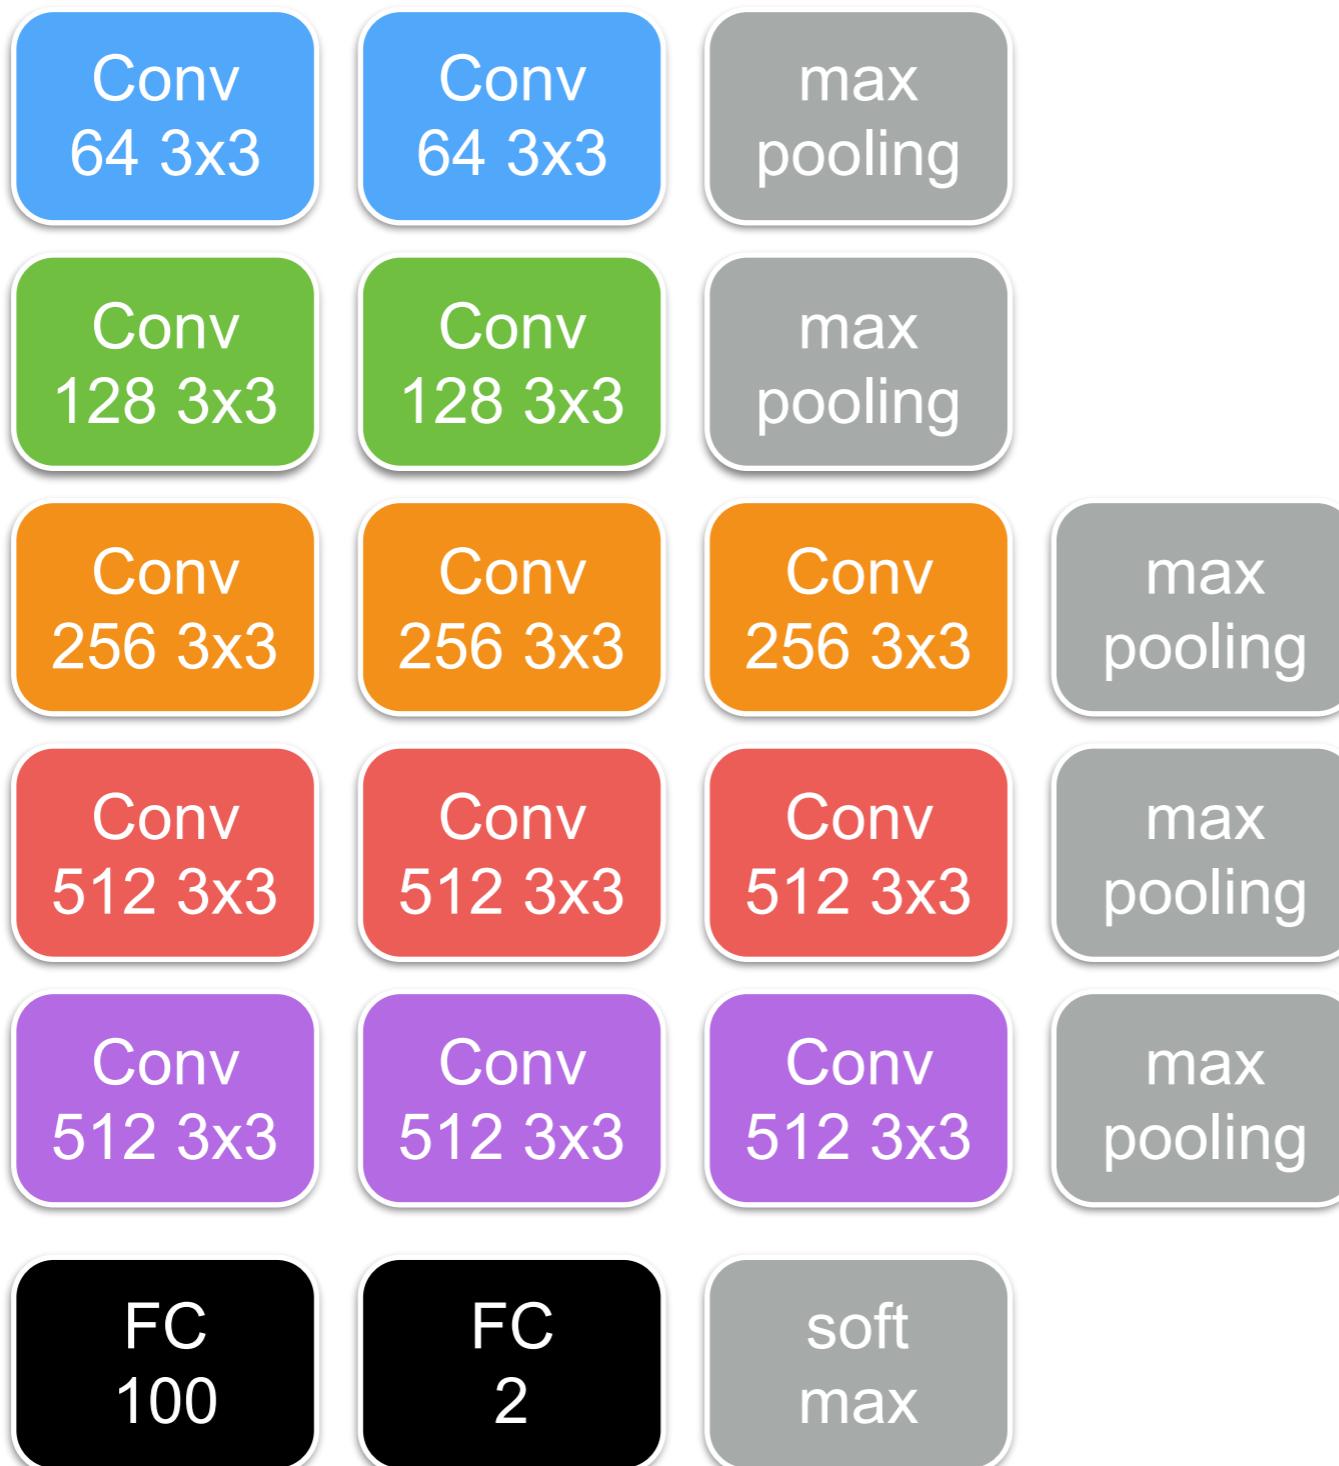
Transfer Learning

VGG-16 architecture



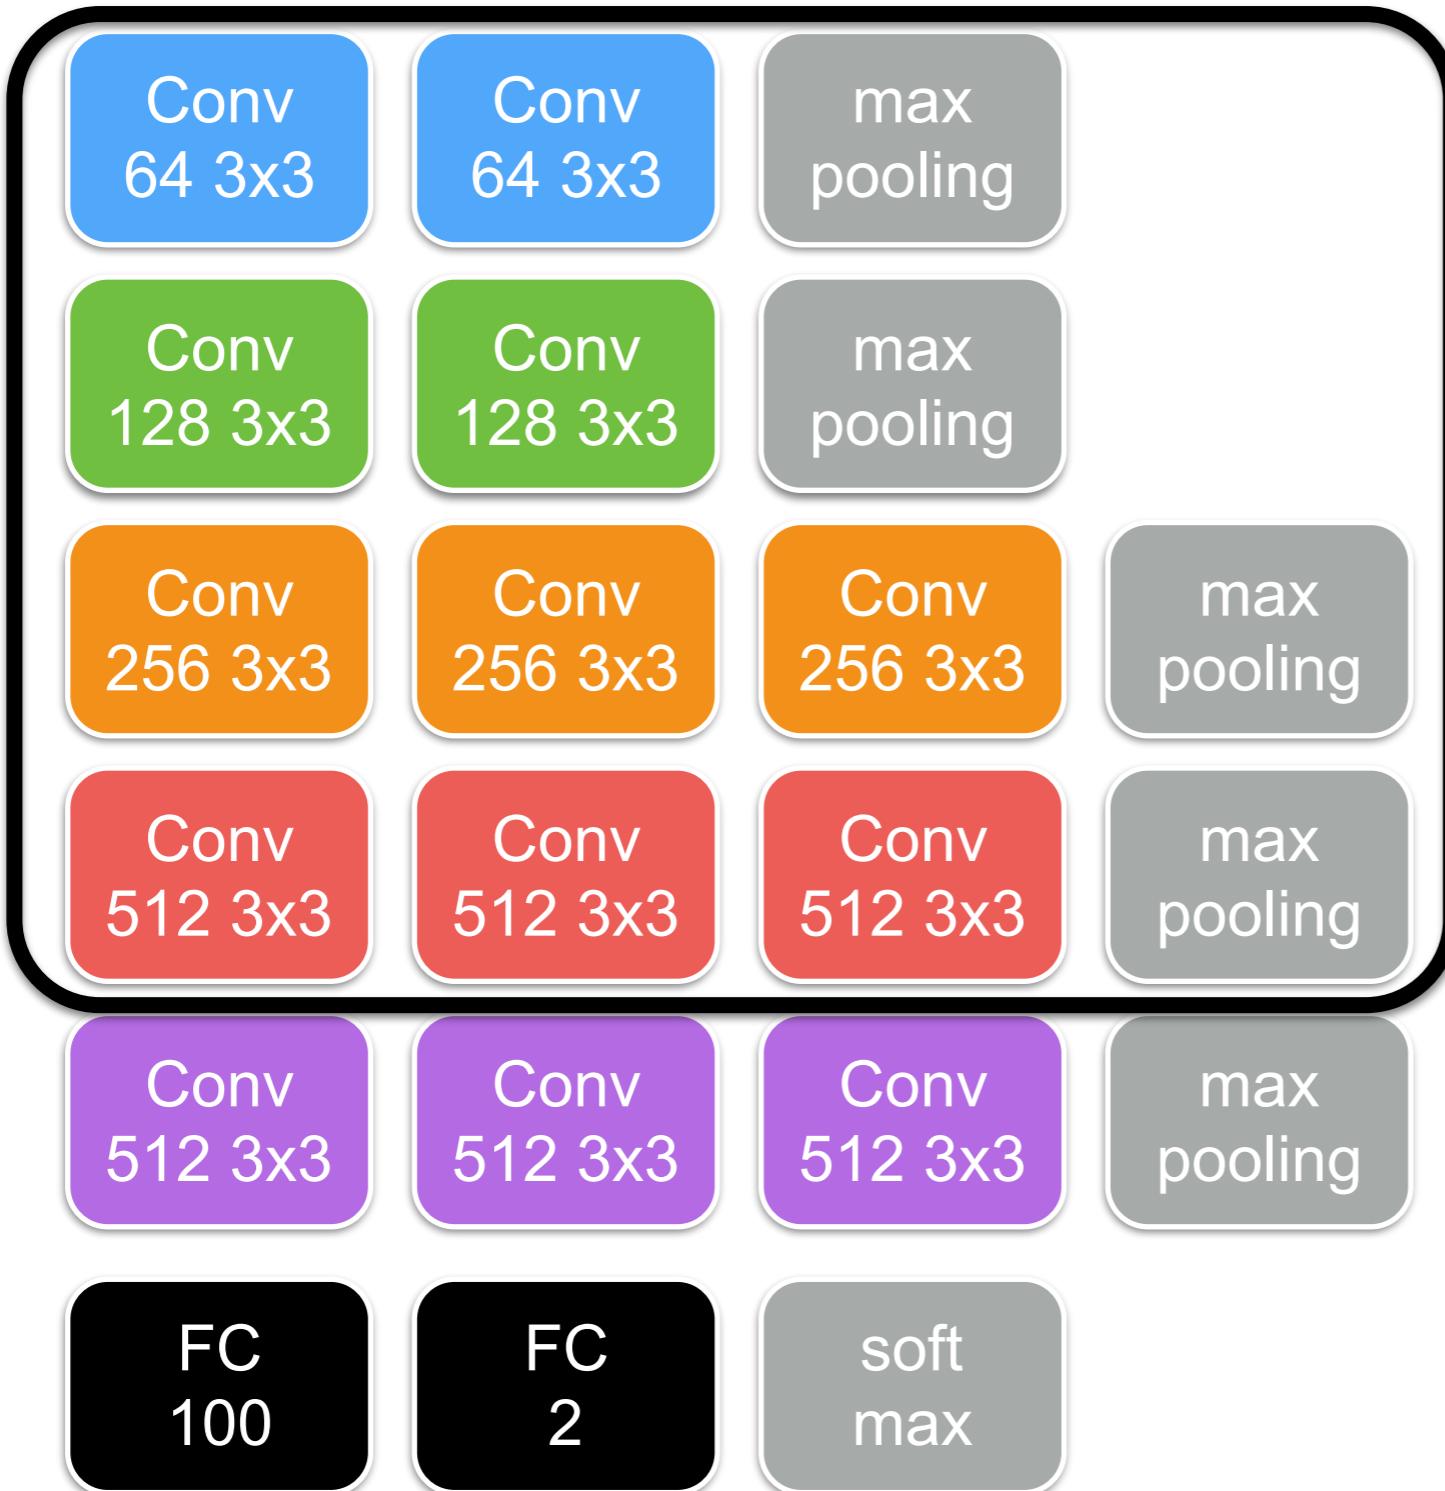
Transfer Learning

VGG-16 architecture



Transfer Learning

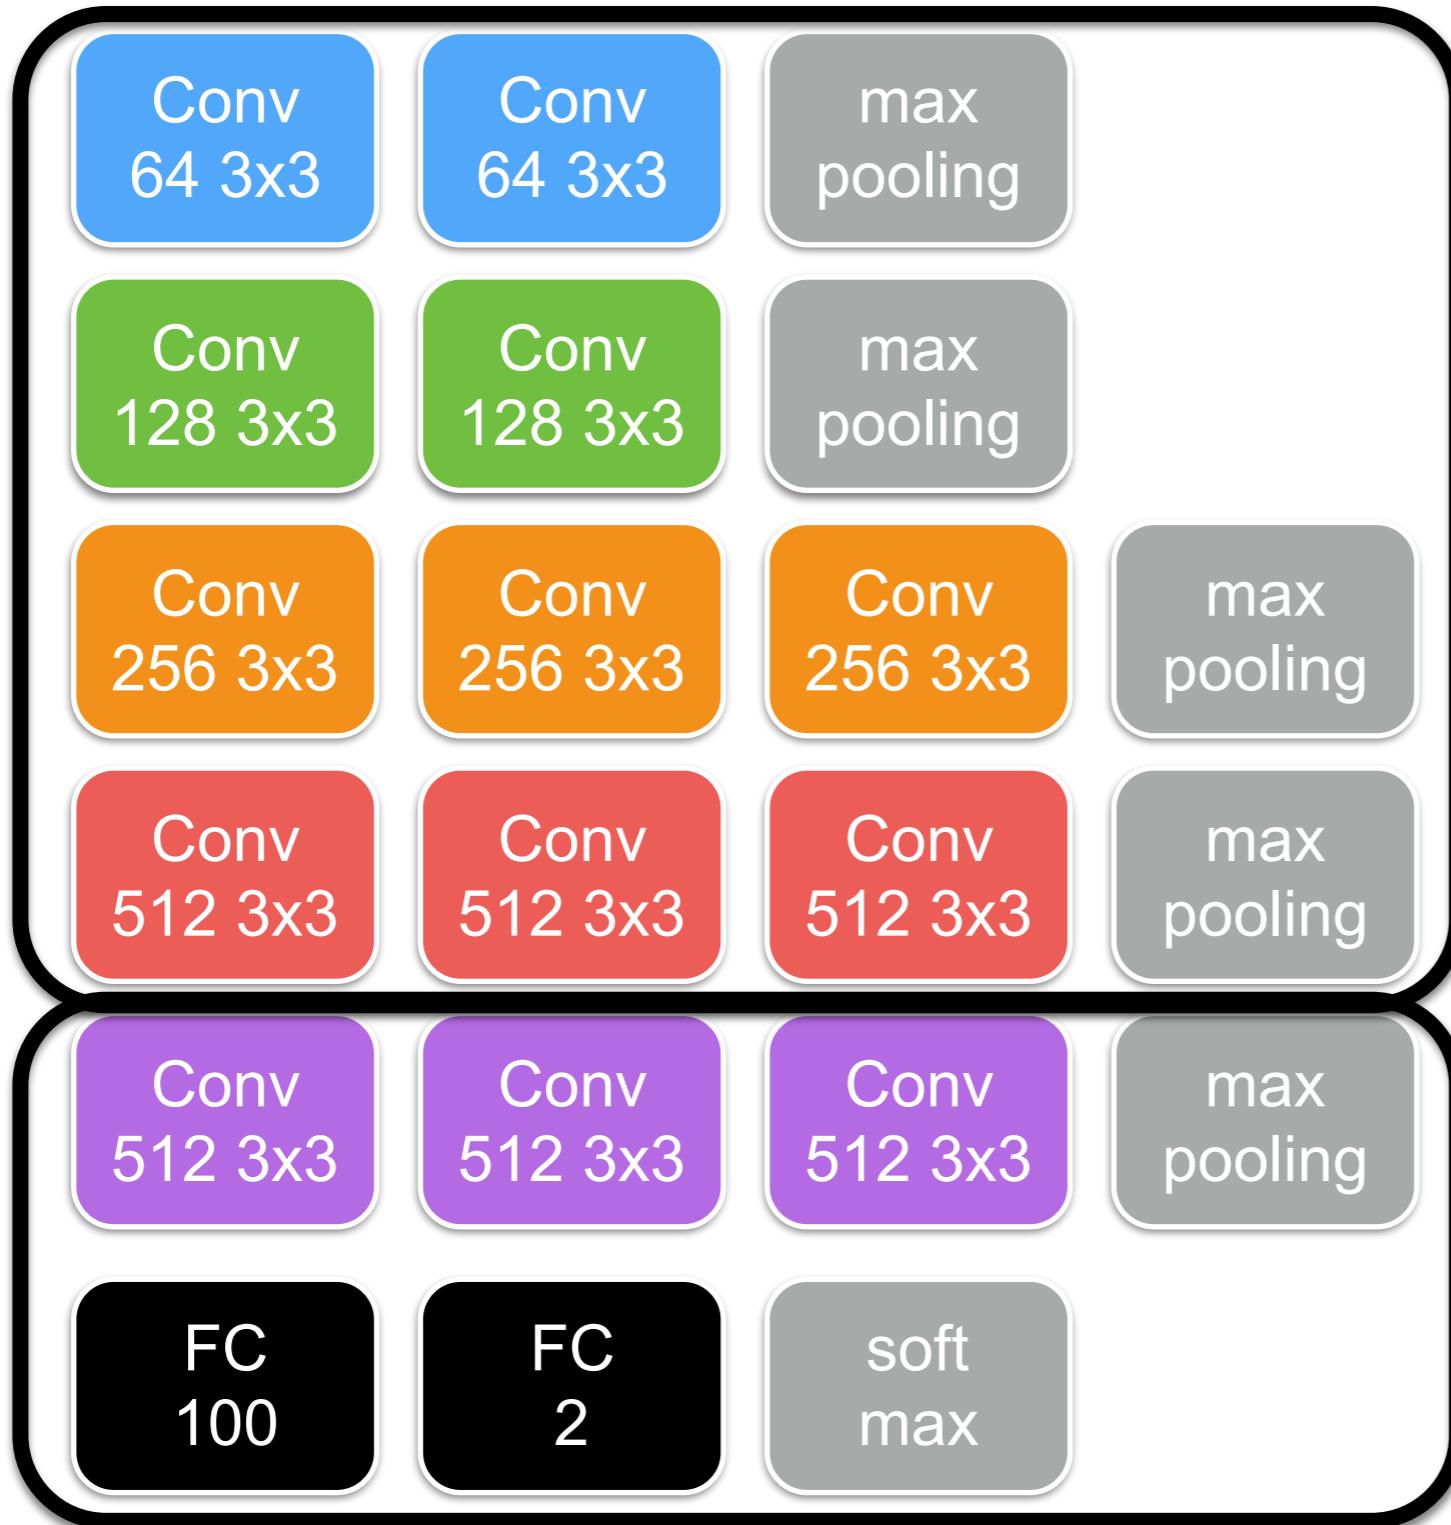
VGG-16 architecture



Fix weights

Transfer Learning

VGG-16 architecture



Fix weights

Train on limited data

Lessons Learned

- Main lessons from this lecture
 - Convolutional neural networks
 - Layers: Fully connected, convolutional, activation functions, max pooling, strides & dilated filters
 - Fully convolutional networks
 - Handling overfitting

Lessons Learned

- Main lessons from this lecture
 - Convolutional neural networks
 - Layers: Fully connected, convolutional, activation functions, max pooling, strides & dilated filters
 - Fully convolutional networks
 - Handling overfitting
- Next lecture: Robust Model Fitting