

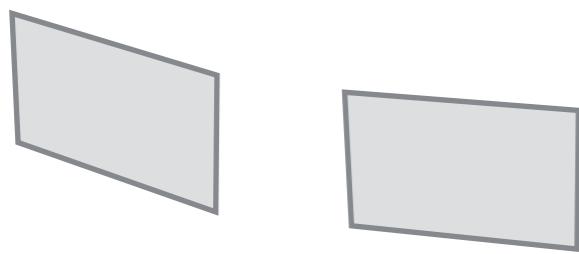
# **SSY097 - Image Analysis**

Lecture 11 - Generative Neural Networks

*Torsten Sattler*

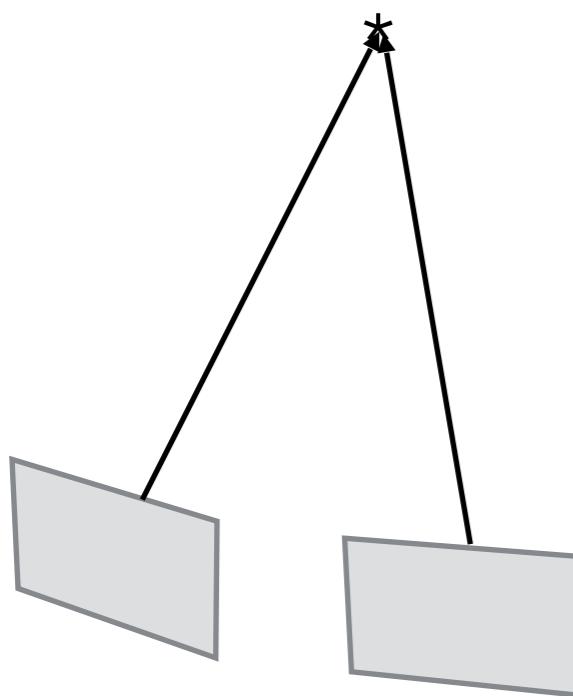
# Last Lecture

Initialize motion from two views



Relative pose for two images

# Last Lecture

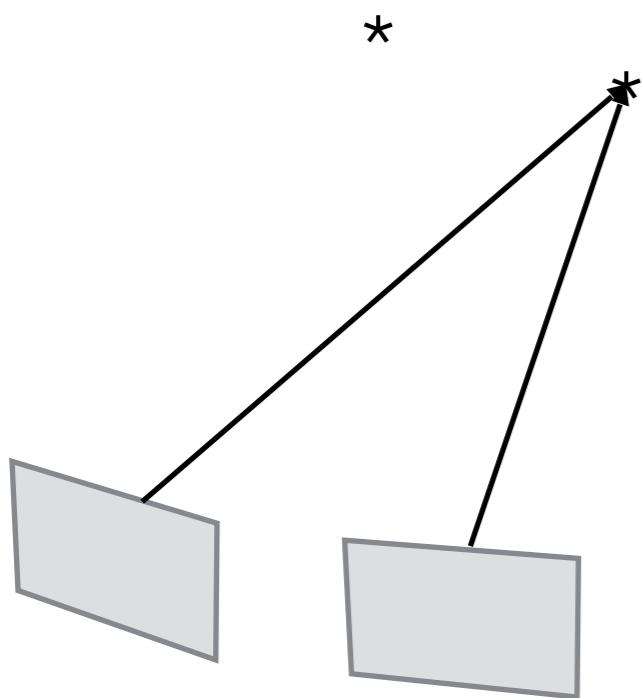


Initialize motion from two views

Initialize structure from two views

Triangulate 3D points

# Last Lecture

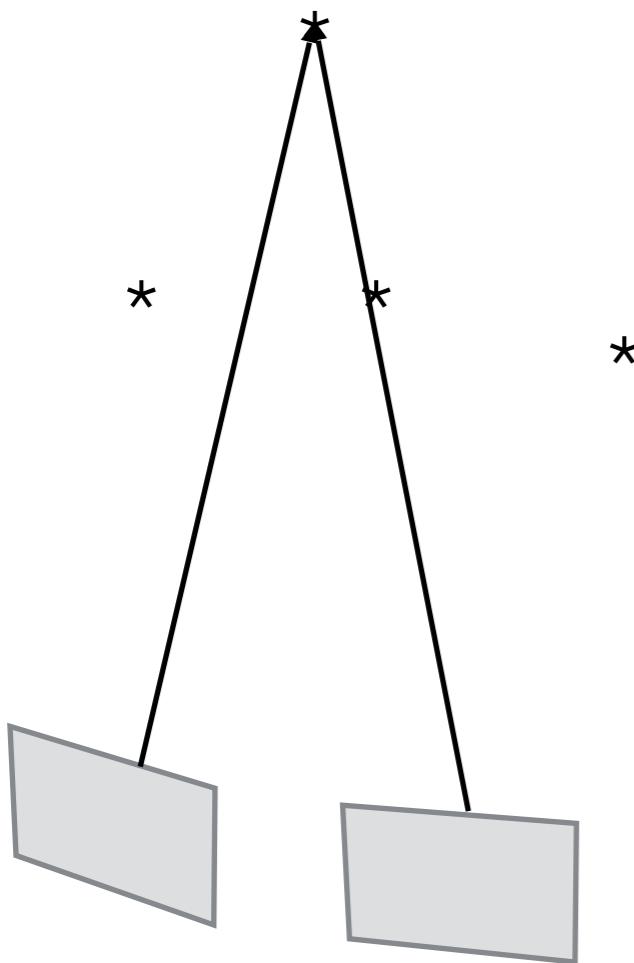


Initialize motion from two views

Initialize structure from two views

Triangulate 3D points

# Last Lecture



Initialize motion from two views

Initialize structure from two views

Triangulate 3D points

# Last Lecture

\*

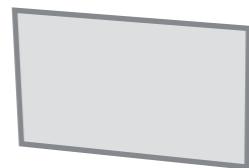
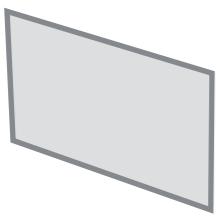
\*

\*

\*

Initialize motion from two views

Initialize structure from two views



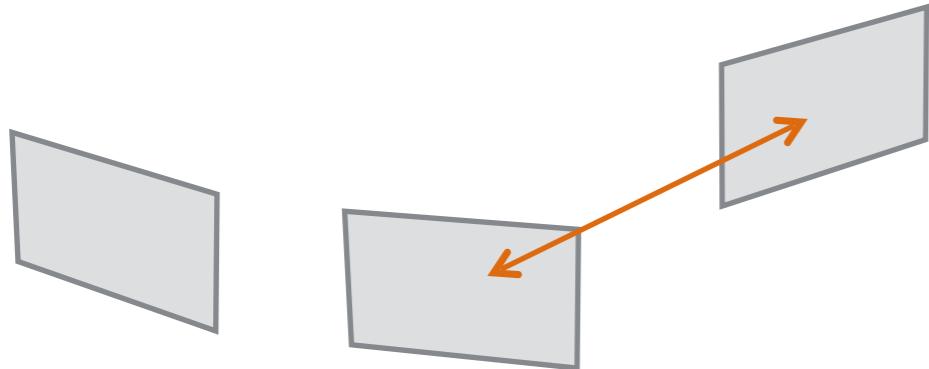
# Last Lecture

\*

\*

\*

\*



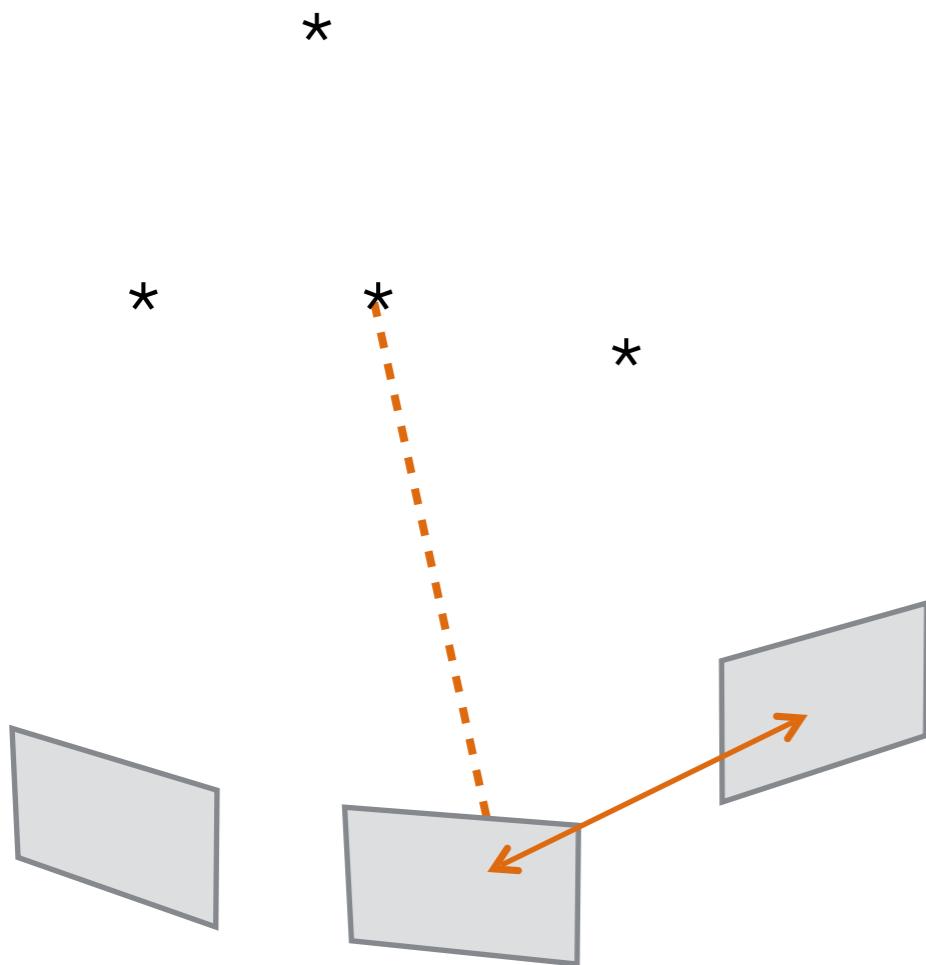
Initialize motion from two views

Initialize structure from two views

Extend motion

Match features

# Last Lecture



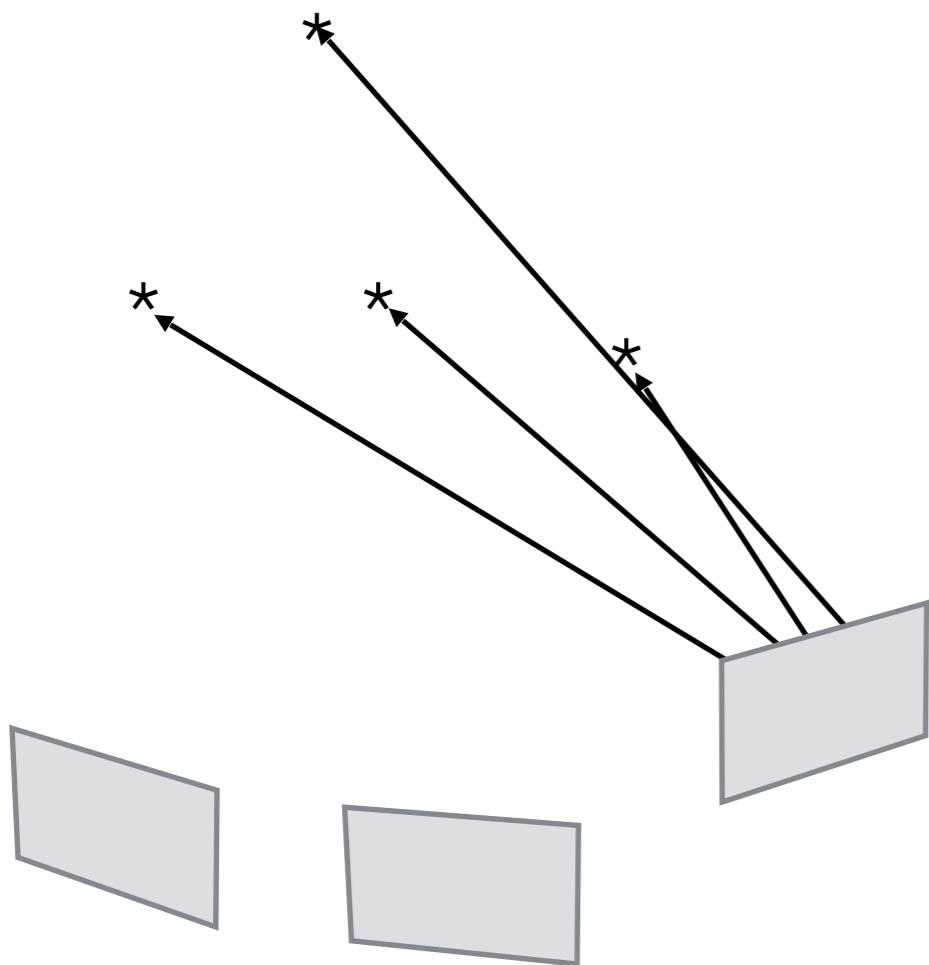
Initialize motion from two views

Initialize structure from two views

Extend motion

Transfer matches to 3D

# Last Lecture



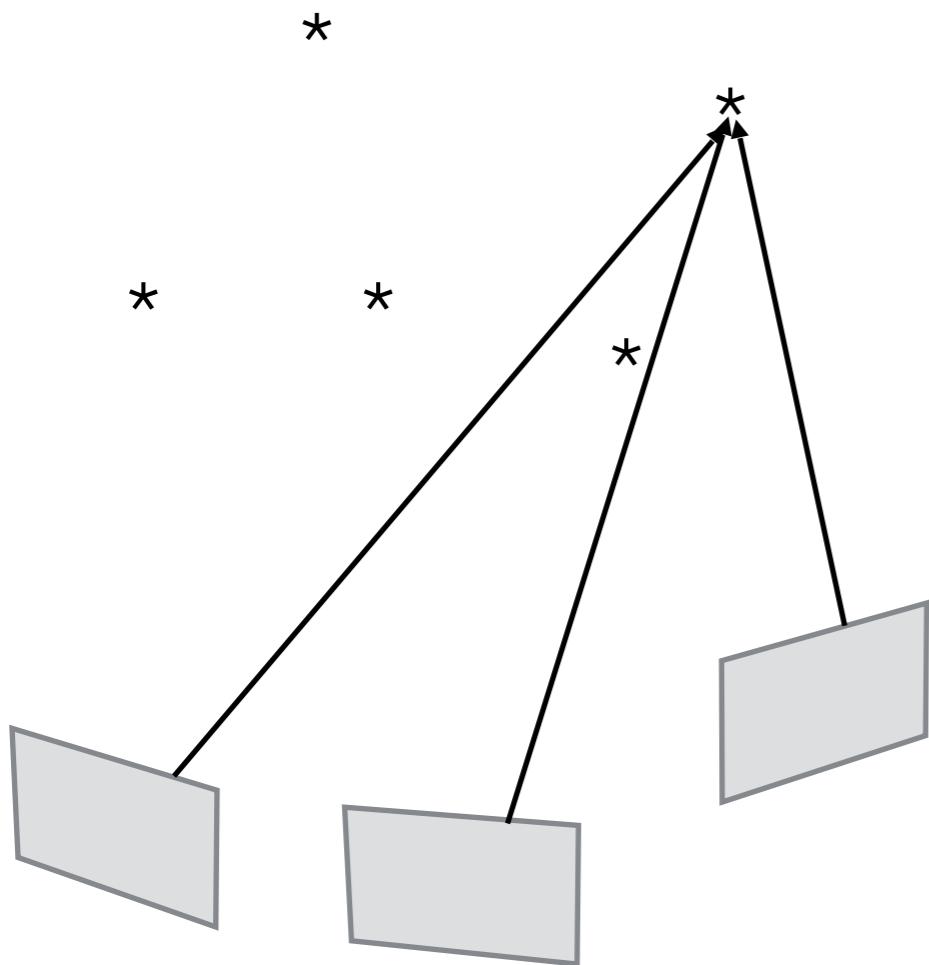
Initialize motion from two views

Initialize structure from two views

Extend motion

Camera pose for third camera

# Last Lecture



Initialize motion from two views

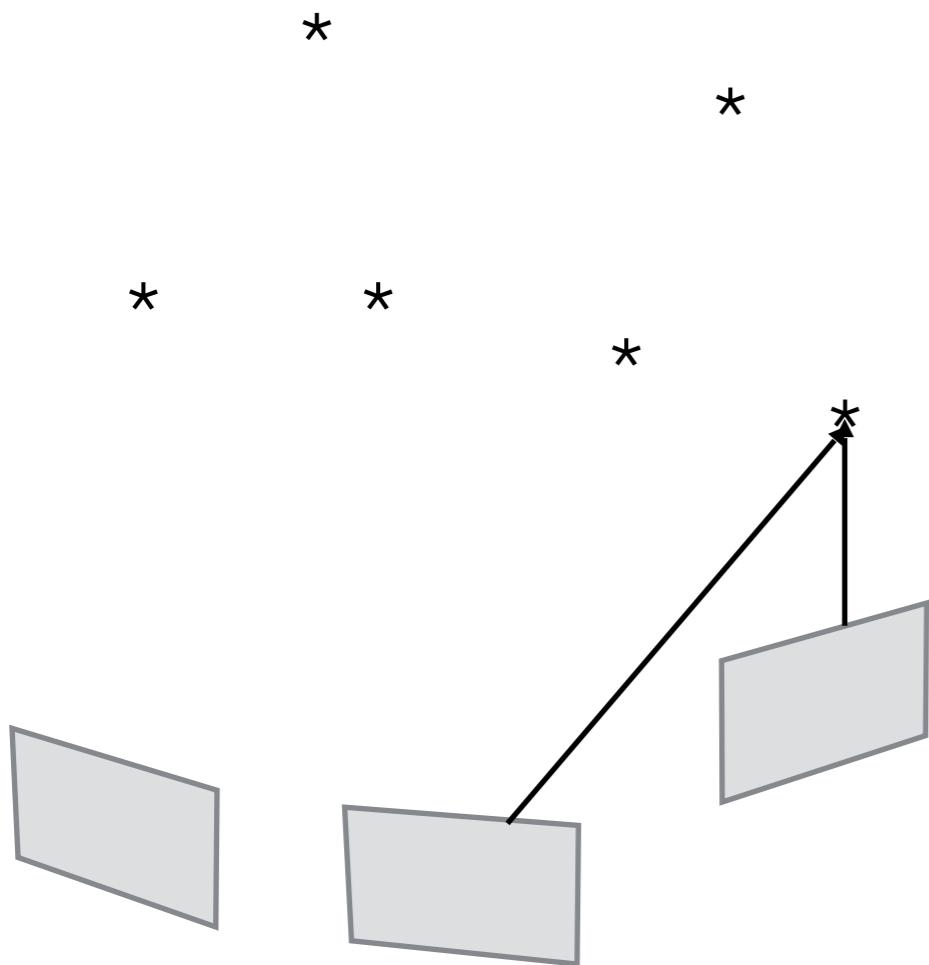
Initialize structure from two views

Extend motion

Extend structure

Triangulate points

# Last Lecture



Initialize motion from two views

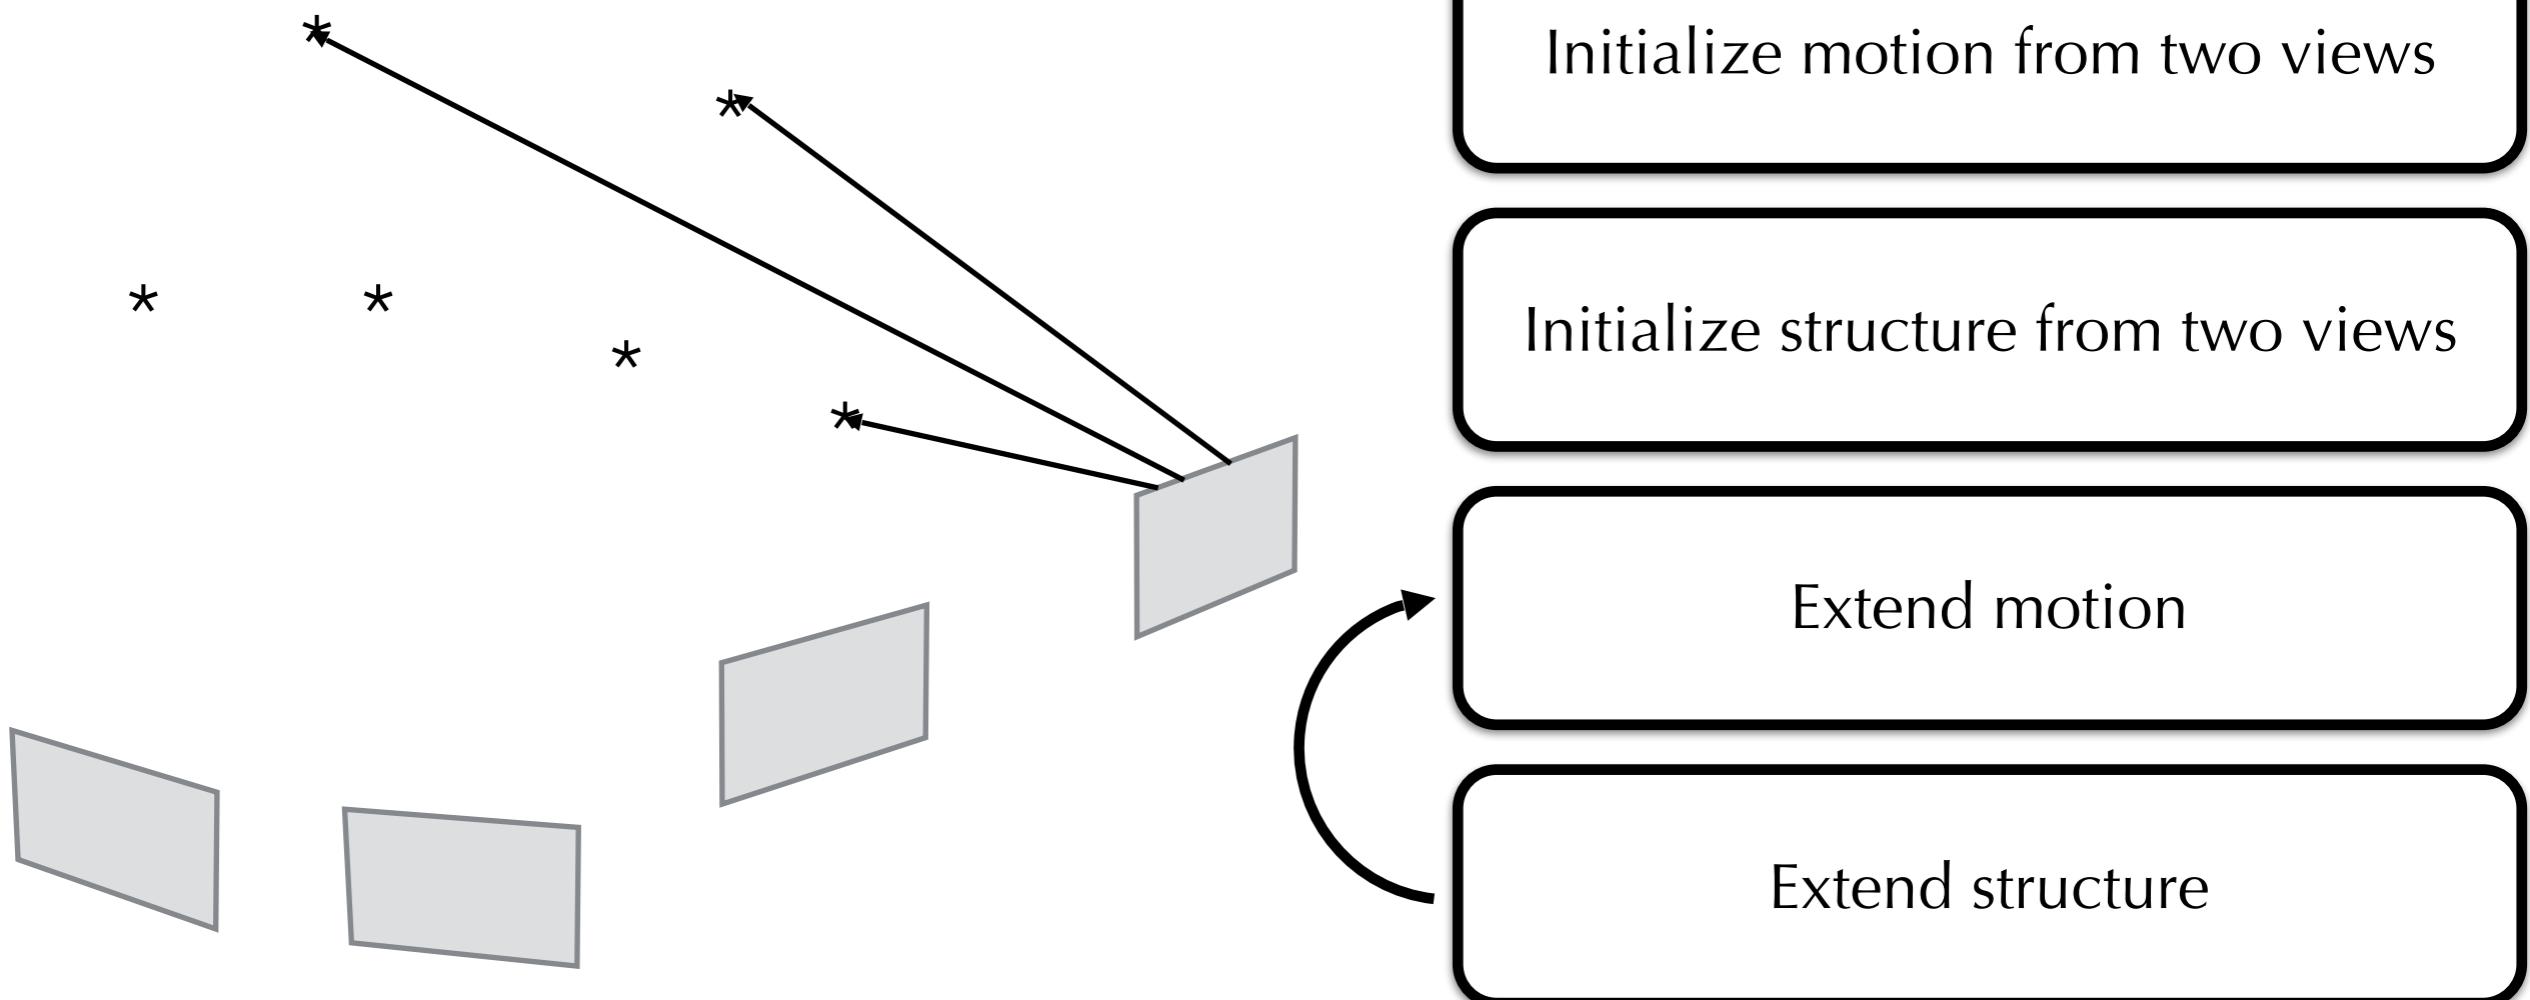
Initialize structure from two views

Extend motion

Extend structure

Triangulate points

# Last Lecture



Camera pose for fourth image

# Last Lecture

3D point  $\mathbf{x}$  seen from camera with pose  $R, \mathbf{t}$ , intrinsics  $K$

$$K(R\mathbf{x} + \mathbf{t}) = \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \underline{x} \\ \underline{z} \\ \underline{y} \\ z \end{pmatrix}$$

3D reconstructions defined up to scale

# Last Lecture

3D point  $\mathbf{x}$  seen from camera with pose  $R, \mathbf{t}$ , intrinsics  $K$

$$K(R\mathbf{X} + \mathbf{t}) = \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \underline{x} \\ \underline{z} \\ \underline{y} \\ z \end{pmatrix}$$

Scale 3D scene by arbitrary factor  $s \neq 0$

$$K(R(s\mathbf{X}) + s\mathbf{t}) = sK(R\mathbf{X} + \mathbf{t})$$

3D reconstructions defined up to scale

# Last Lecture

3D point  $\mathbf{x}$  seen from camera with pose  $R, \mathbf{t}$ , intrinsics  $K$

$$K(R\mathbf{X} + \mathbf{t}) = \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \underline{x} \\ \underline{z} \\ \underline{y} \\ z \end{pmatrix}$$

Scale 3D scene by arbitrary factor  $s \neq 0$

$$\begin{aligned} K(R(s\mathbf{X}) + s\mathbf{t}) &= sK(R\mathbf{X} + \mathbf{t}) \\ &= s \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \underline{sx} \\ \underline{sz} \\ \underline{sy} \\ sz \end{pmatrix} \end{aligned}$$

3D reconstructions defined up to scale

# Last Lecture

3D point  $\mathbf{x}$  seen from camera with pose  $\mathbf{R}$ ,  $\mathbf{t}$ , intrinsics  $\mathbf{K}$

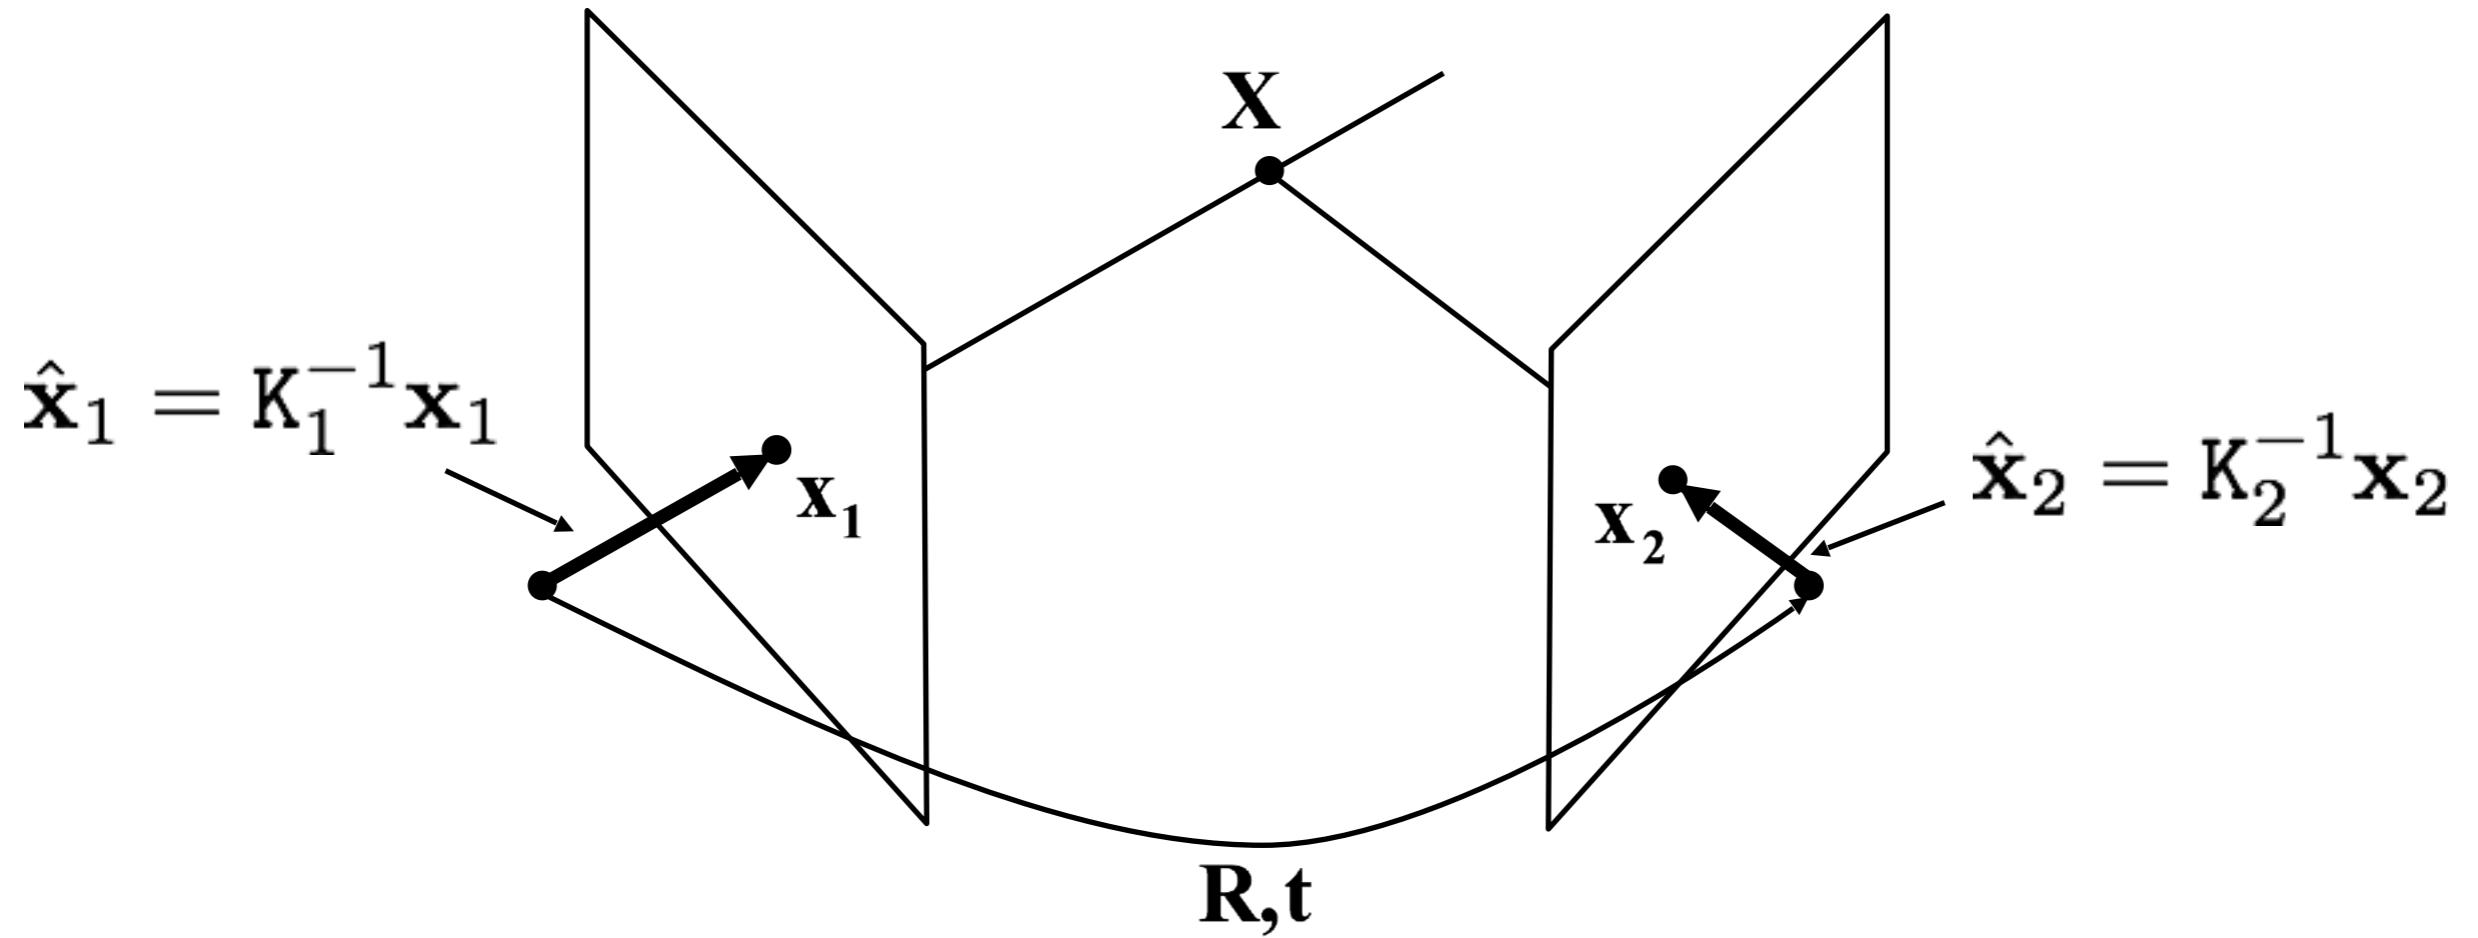
$$\mathbf{K}(\mathbf{R}\mathbf{x} + \mathbf{t}) = \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \underline{x} \\ \underline{z} \\ \underline{y} \\ z \end{pmatrix}$$

Scale 3D scene by arbitrary factor  $s \neq 0$

$$\begin{aligned} \mathbf{K}(\mathbf{R}(s\mathbf{x}) + s\mathbf{t}) &= s\mathbf{K}(\mathbf{R}\mathbf{x} + \mathbf{t}) \\ &= s \begin{pmatrix} x \\ y \\ z \end{pmatrix} \mapsto \begin{pmatrix} \frac{sx}{sz} \\ \frac{sy}{sz} \\ \underline{y} \\ z \end{pmatrix} = \begin{pmatrix} \underline{x} \\ \underline{z} \\ \underline{y} \\ z \end{pmatrix} \end{aligned}$$

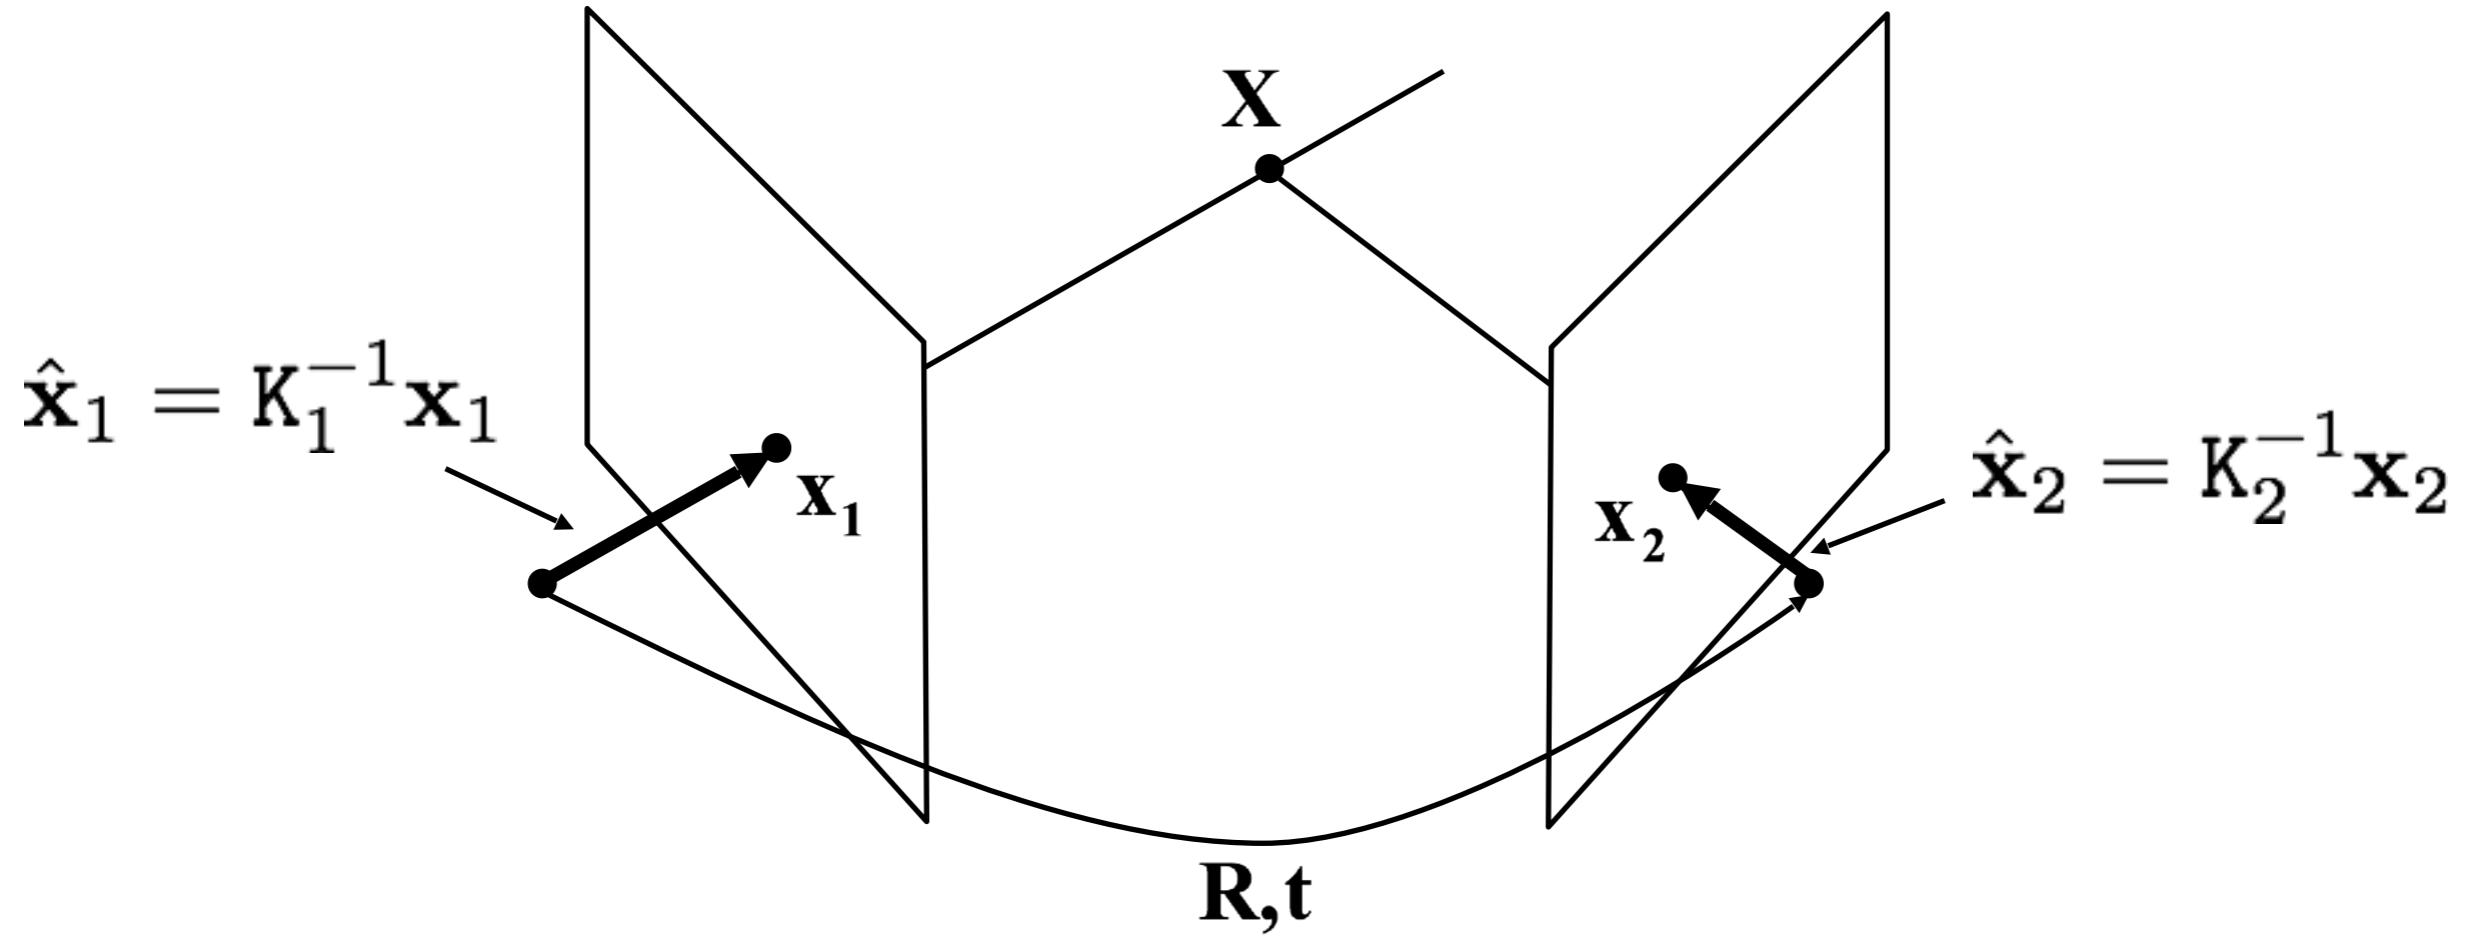
3D reconstructions defined up to scale

# Last Lecture



Epipolar geometry

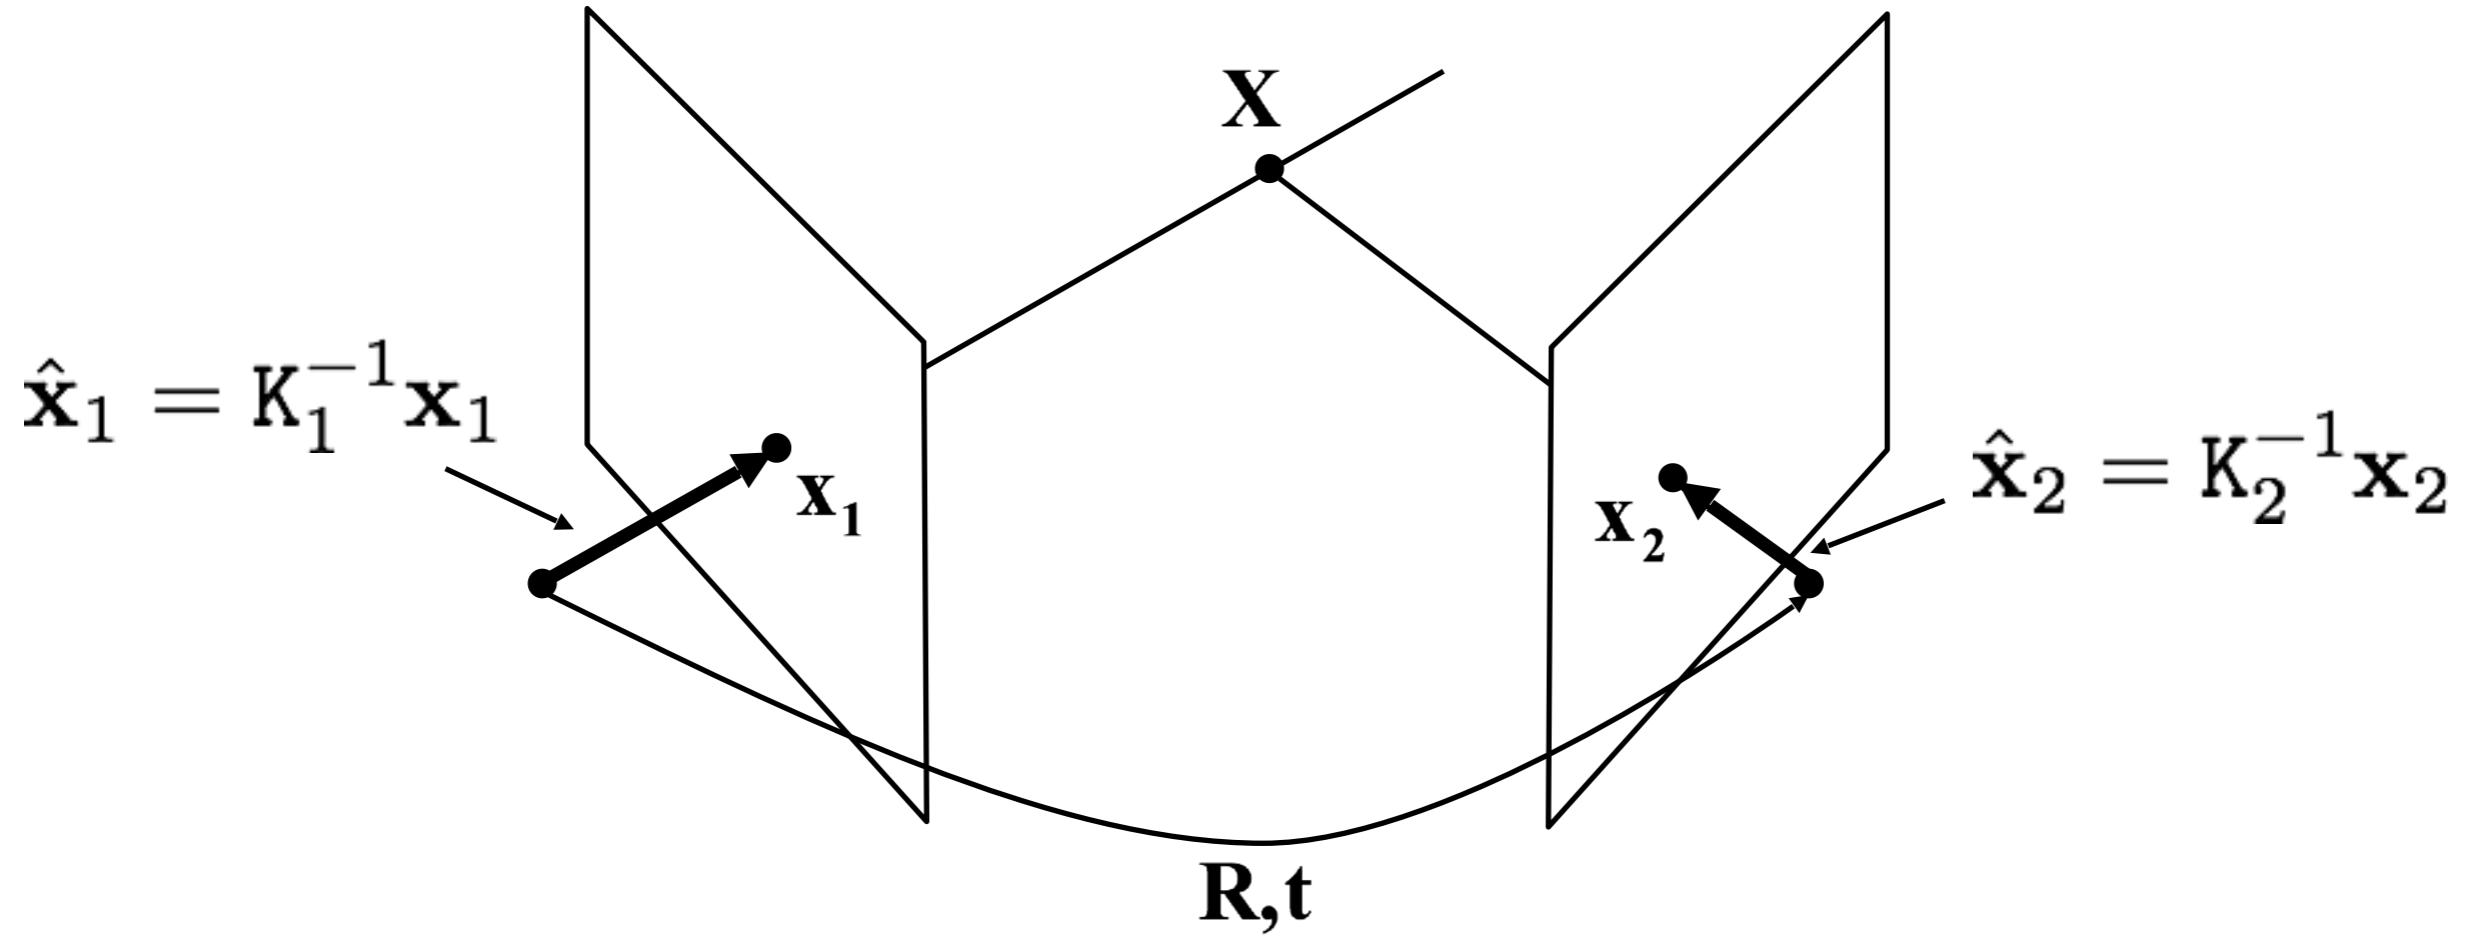
# Last Lecture



$$\mathbf{0} = \hat{\mathbf{x}}_2^T \mathbf{E} \hat{\mathbf{x}}_1$$

Epipolar geometry

# Last Lecture

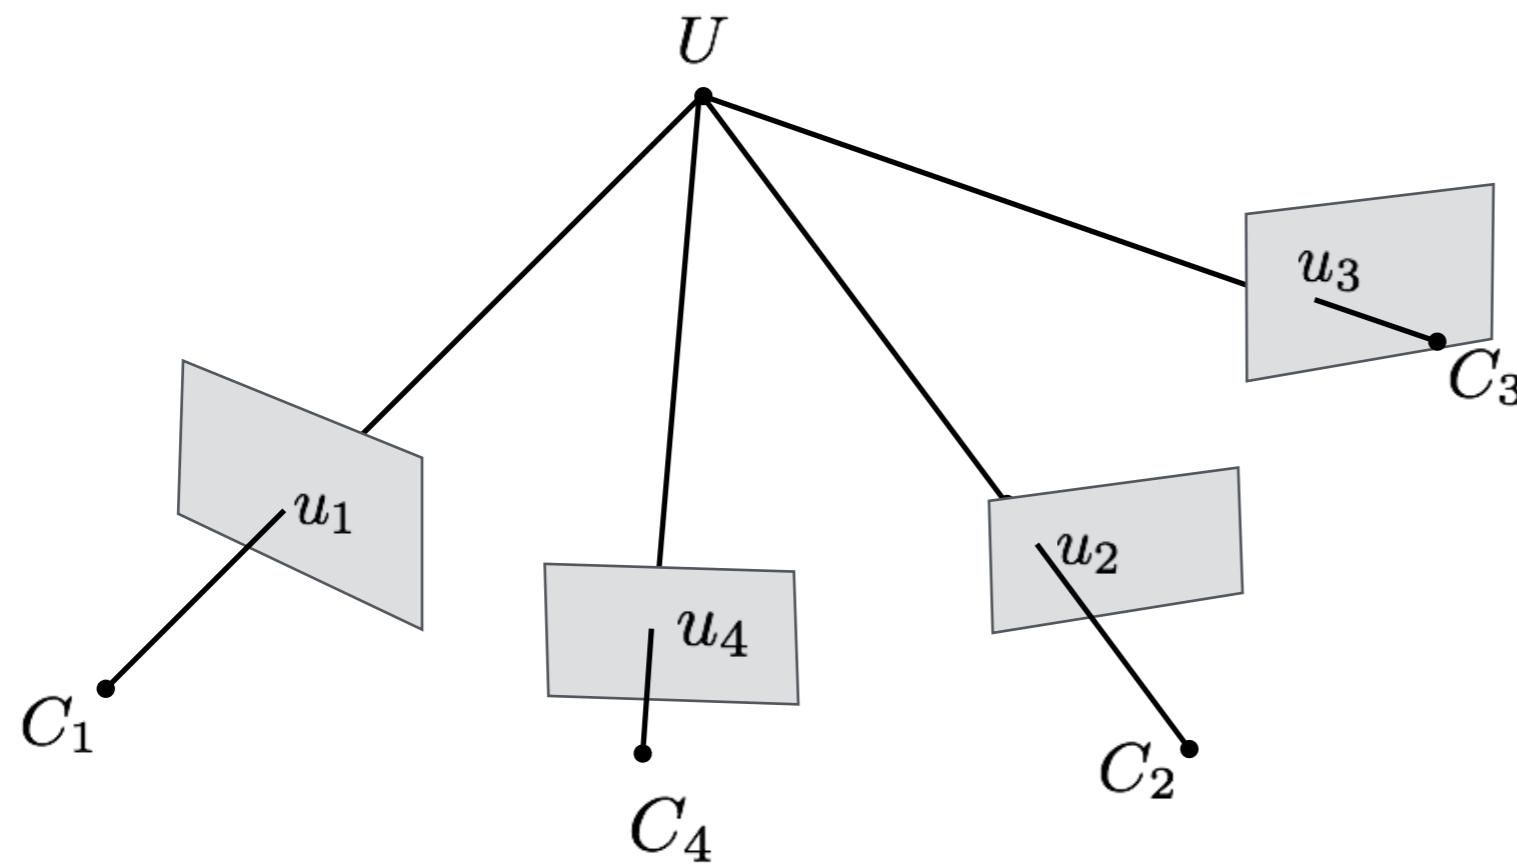


$$0 = \hat{\mathbf{x}}_2^T \mathbf{E} \hat{\mathbf{x}}_1$$

$$0 = \mathbf{x}_2^T \mathbf{F} \mathbf{x}_1$$

Epipolar geometry

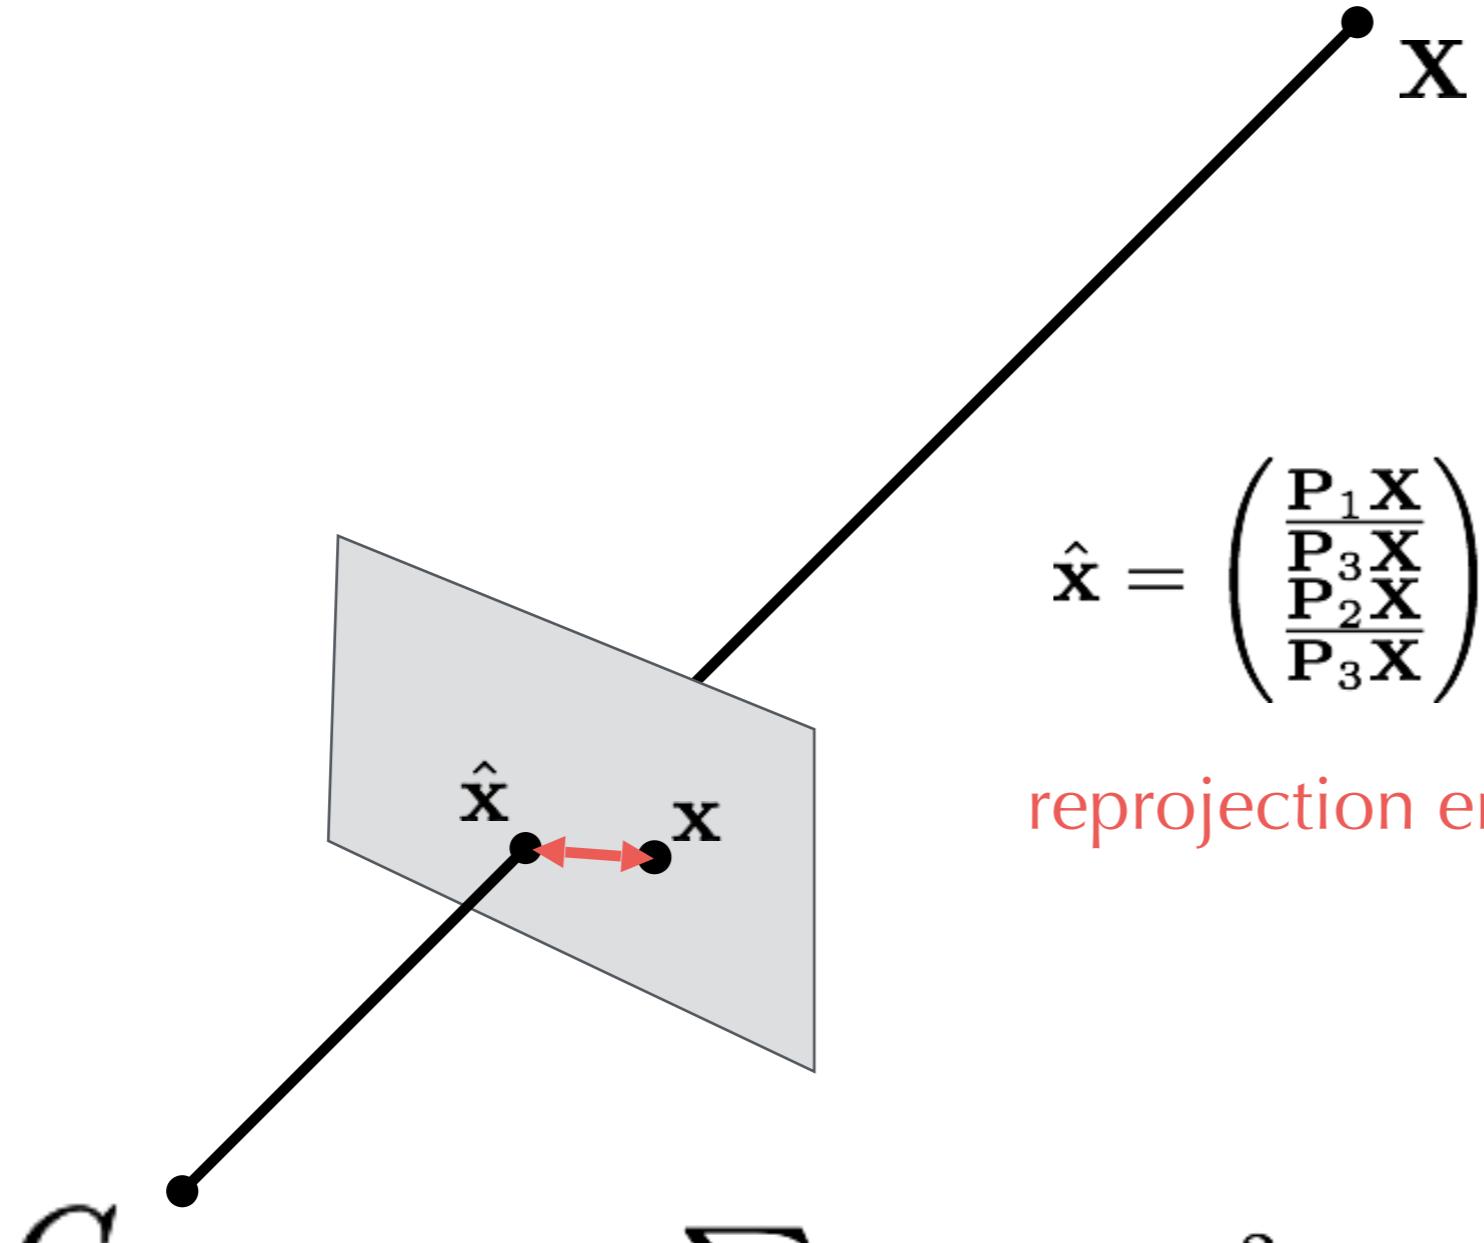
# Last Lecture



$$\begin{pmatrix} \mathbf{P}_3x - \mathbf{P}_1 \\ \mathbf{P}_3y - \mathbf{P}_2 \\ \mathbf{P}'_3x - \mathbf{P}'_1 \\ \mathbf{P}'_3y - \mathbf{P}'_2 \end{pmatrix} \mathbf{X} = \mathbf{0}$$

Triangulation

# Last Lecture



$$\hat{\mathbf{x}} = \begin{pmatrix} \frac{\mathbf{P}_1 \mathbf{X}}{\mathbf{P}_3 \mathbf{X}} \\ \frac{\mathbf{P}_3 \mathbf{X}}{\mathbf{P}_2 \mathbf{X}} \\ \frac{\mathbf{P}_2 \mathbf{X}}{\mathbf{P}_3 \mathbf{X}} \end{pmatrix}$$

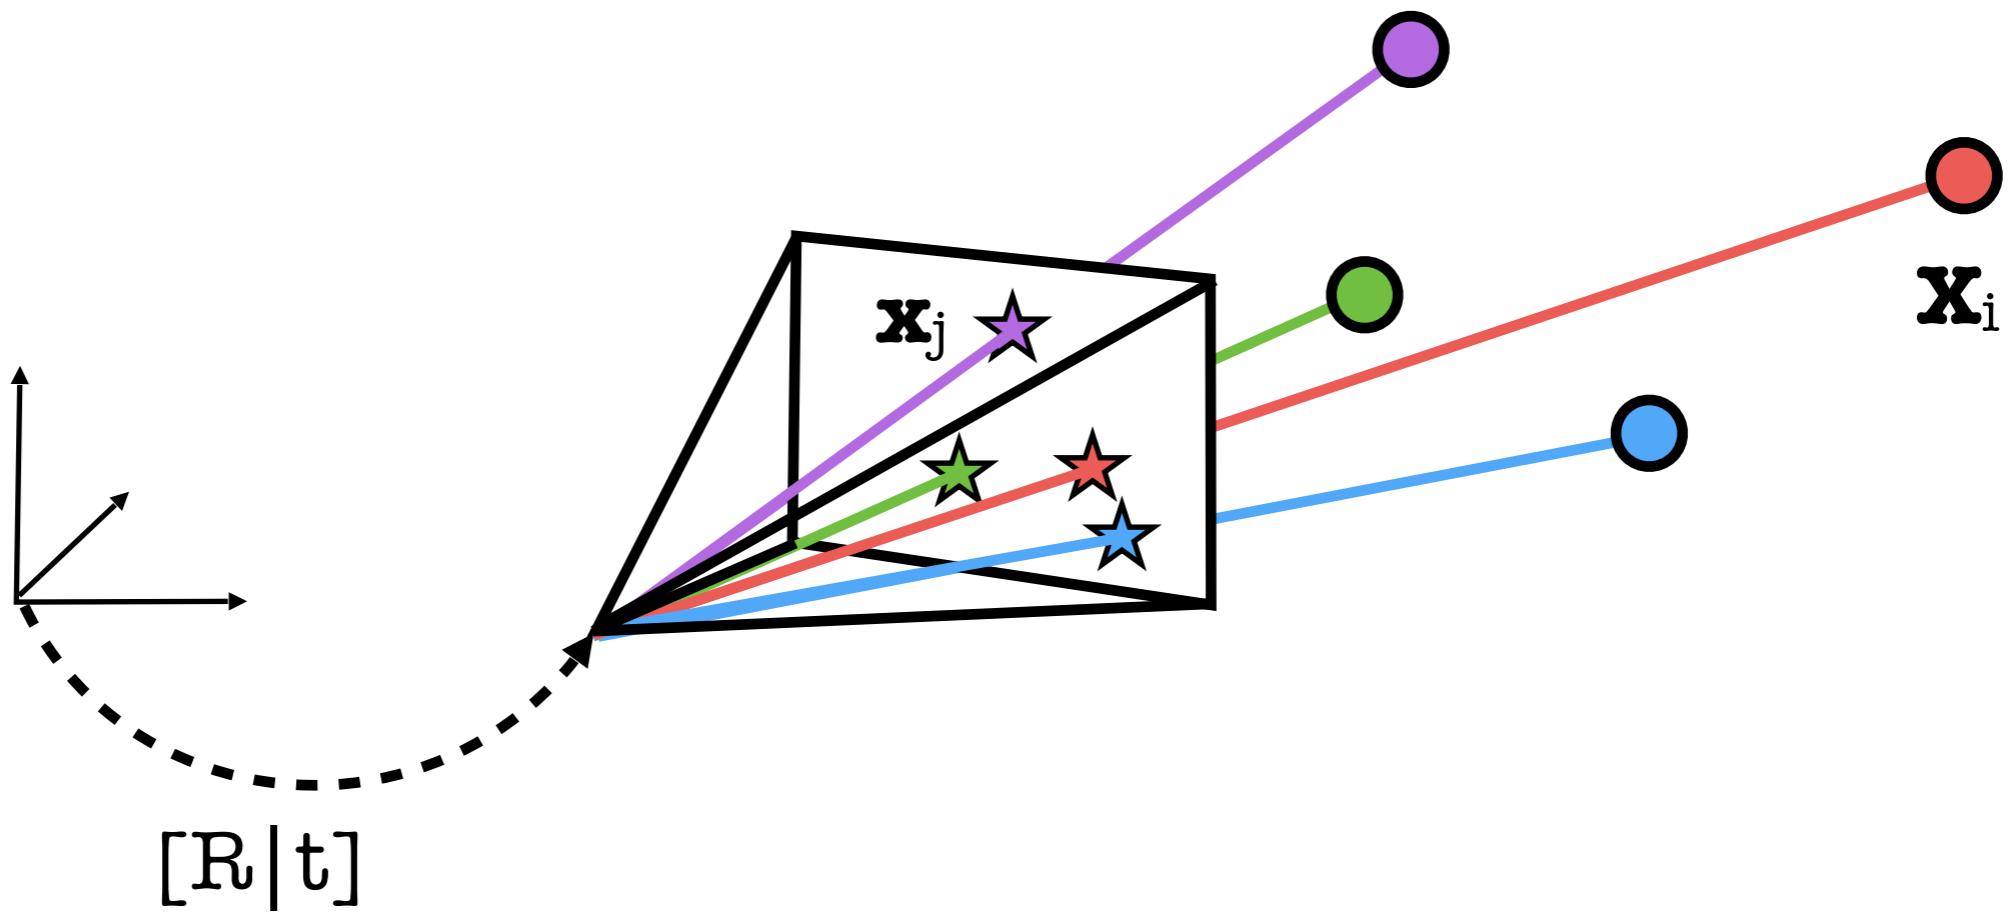
reprojection error  $||\mathbf{x} - \hat{\mathbf{x}}||$

$C_1$

$$\min_{\mathbf{X}} \sum_i ||\mathbf{x}_i - \hat{\mathbf{x}}_i||^2 = \min_{\mathbf{X}} \sum_i ||\mathbf{x}_i - \begin{pmatrix} \frac{\mathbf{P}_1^i \mathbf{X}}{\mathbf{P}_3^i \mathbf{X}} \\ \frac{\mathbf{P}_3^i \mathbf{X}}{\mathbf{P}_2^i \mathbf{X}} \\ \frac{\mathbf{P}_2^i \mathbf{X}}{\mathbf{P}_3^i \mathbf{X}} \end{pmatrix}||^2$$

Minimizing reprojection error

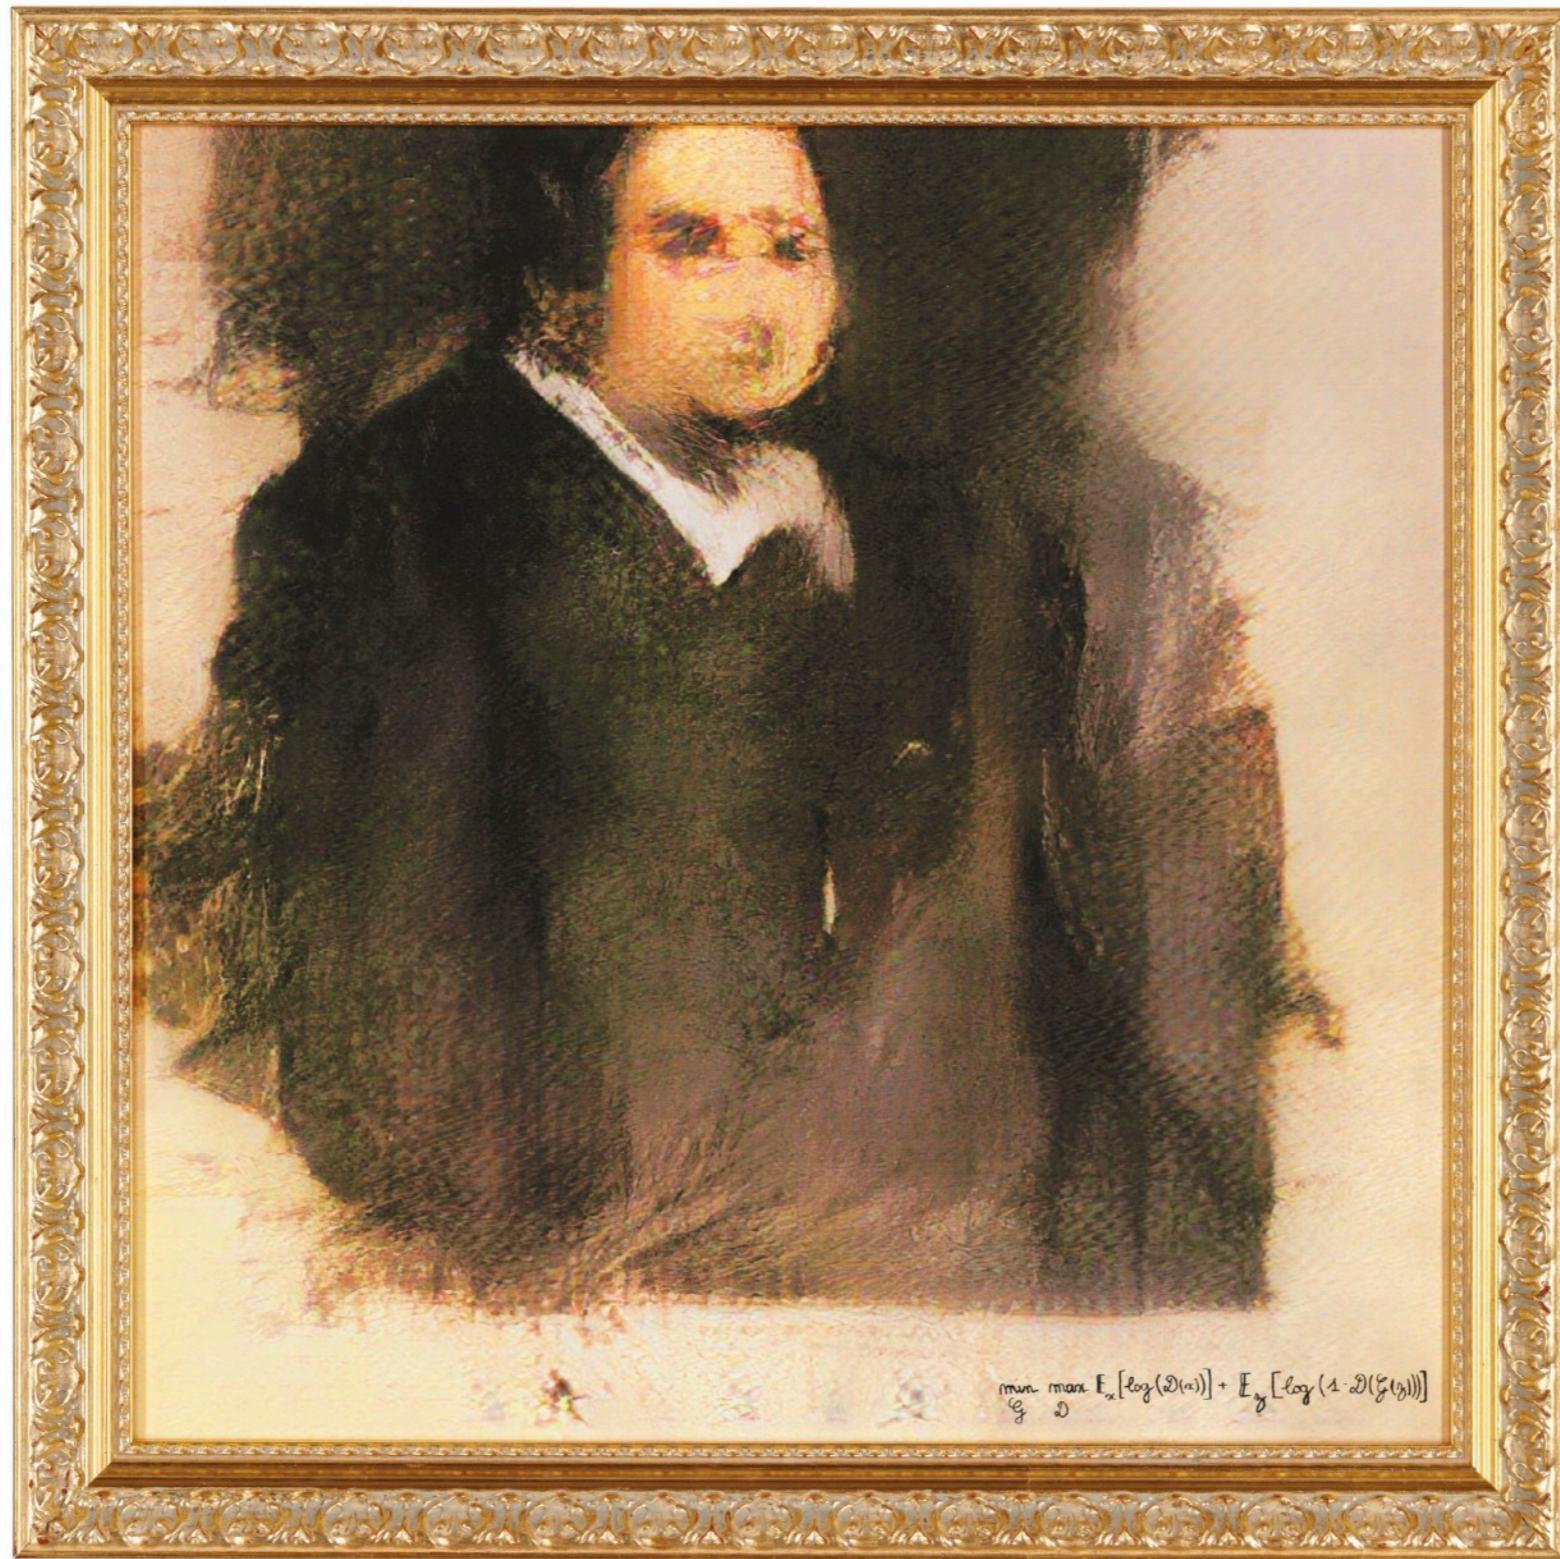
# Last Lecture



Absolute pose estimation

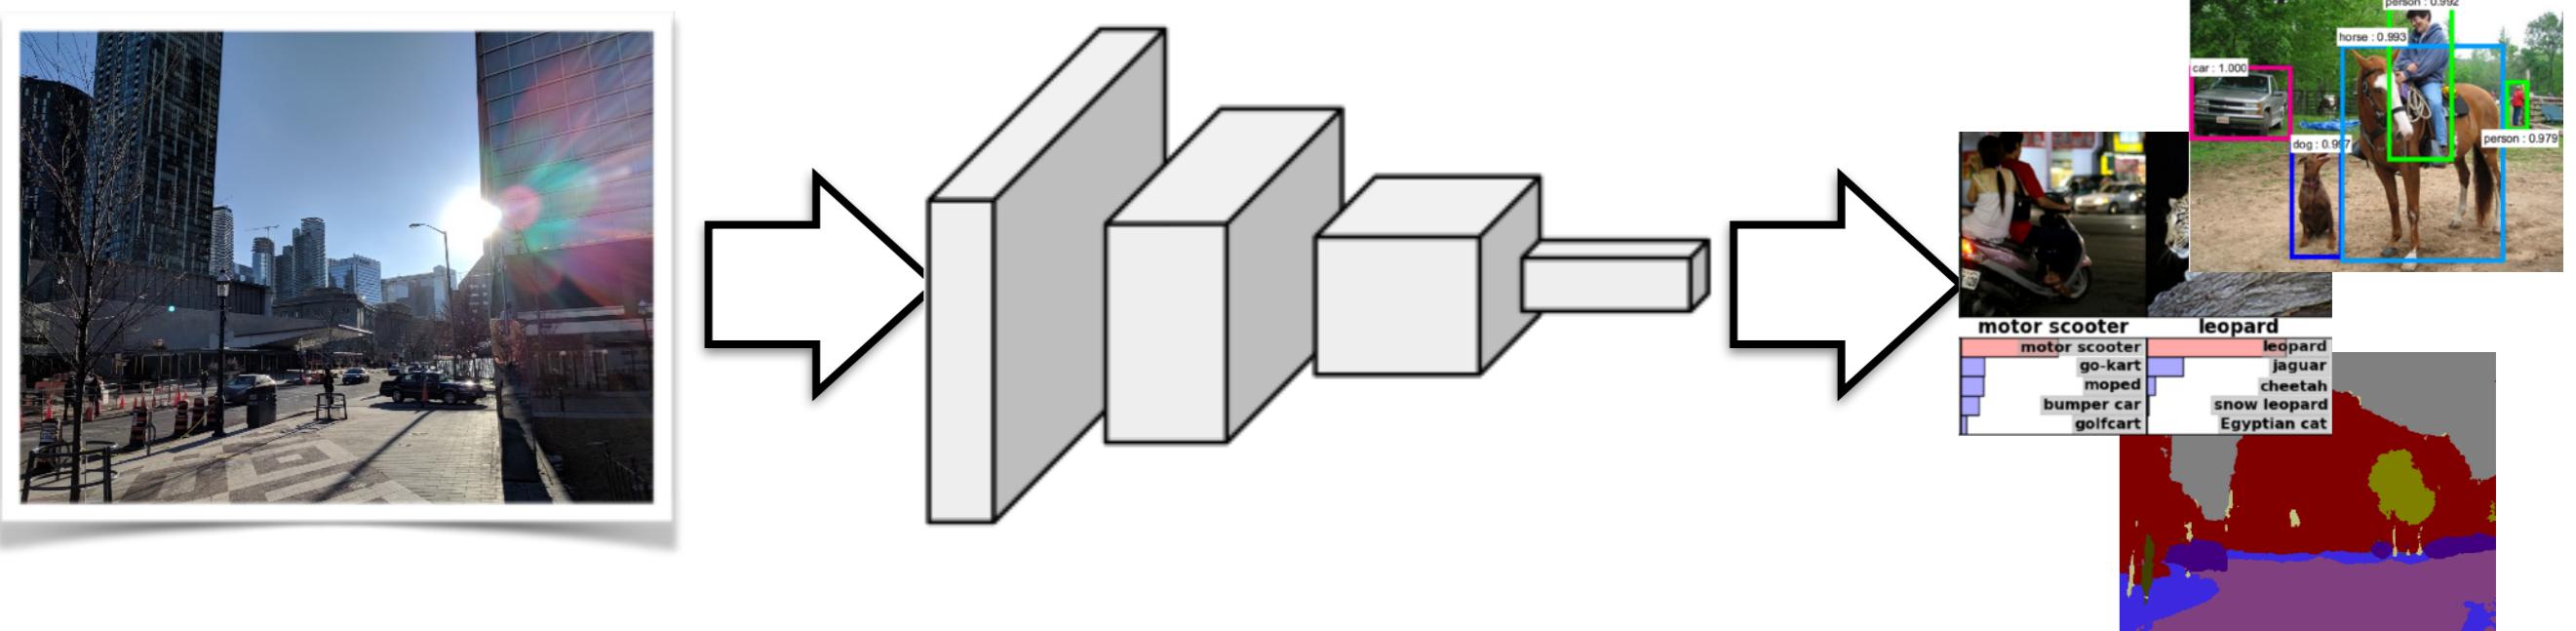
Today

# Generative Neural Networks



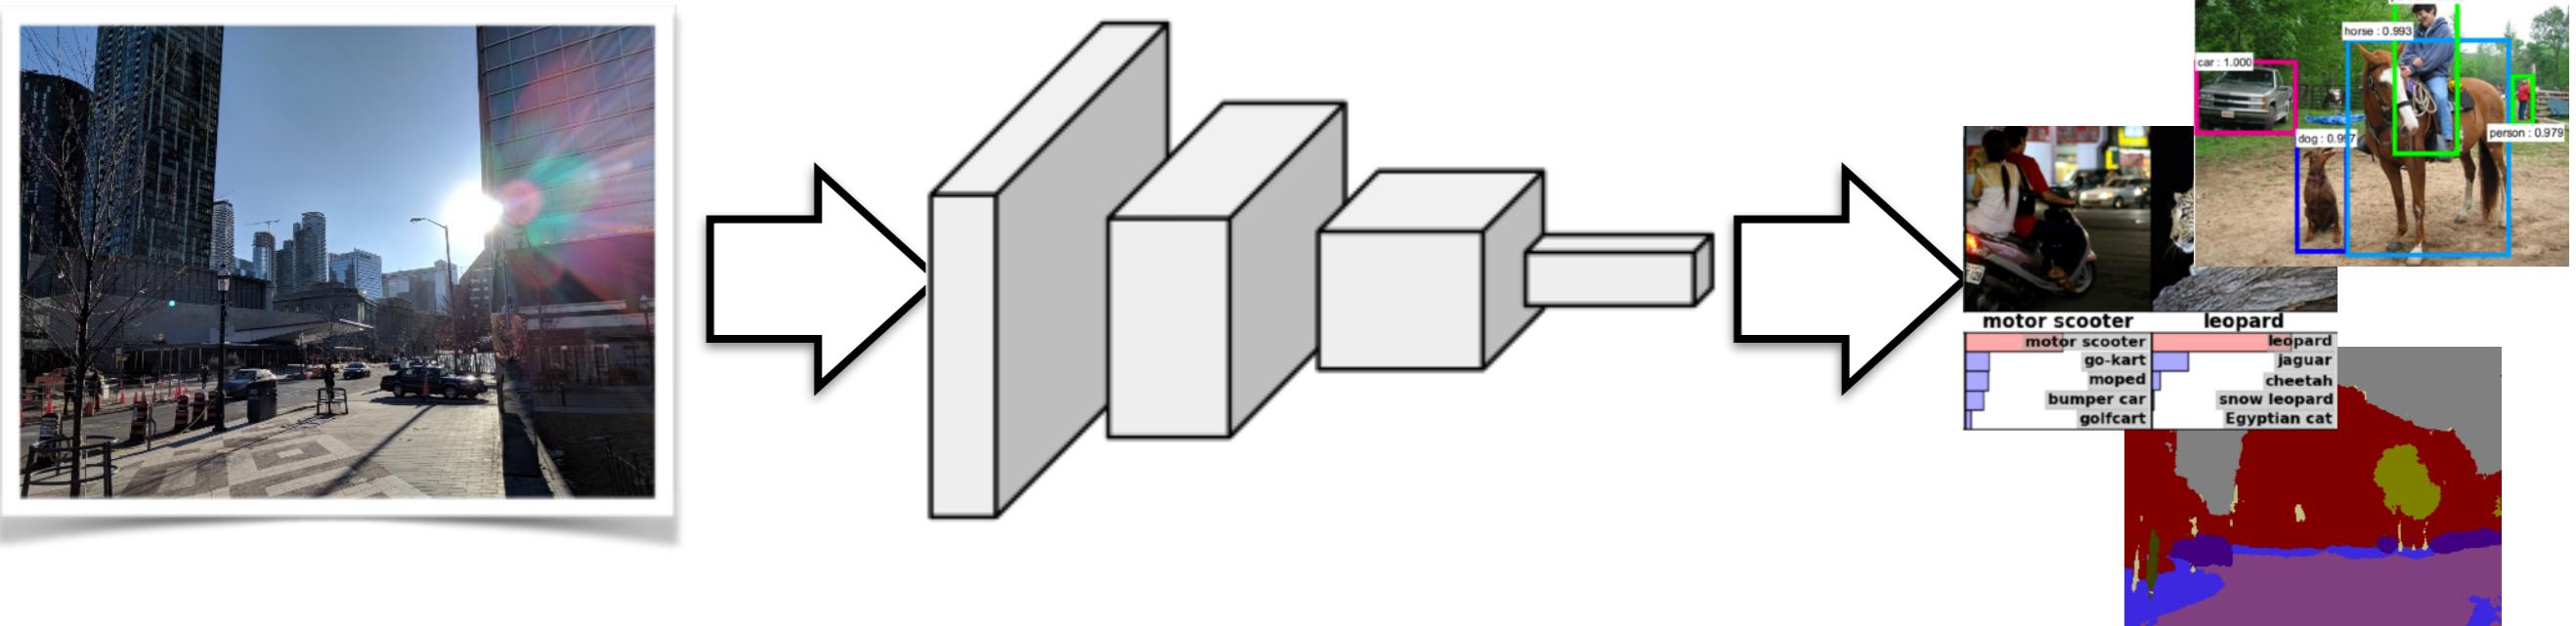
[Portrait of Edmond Belamy, 2018](#). Sold for \$432,500 on 25 October at Christie's in New York. Image © Obvious

# Recap: CNNs



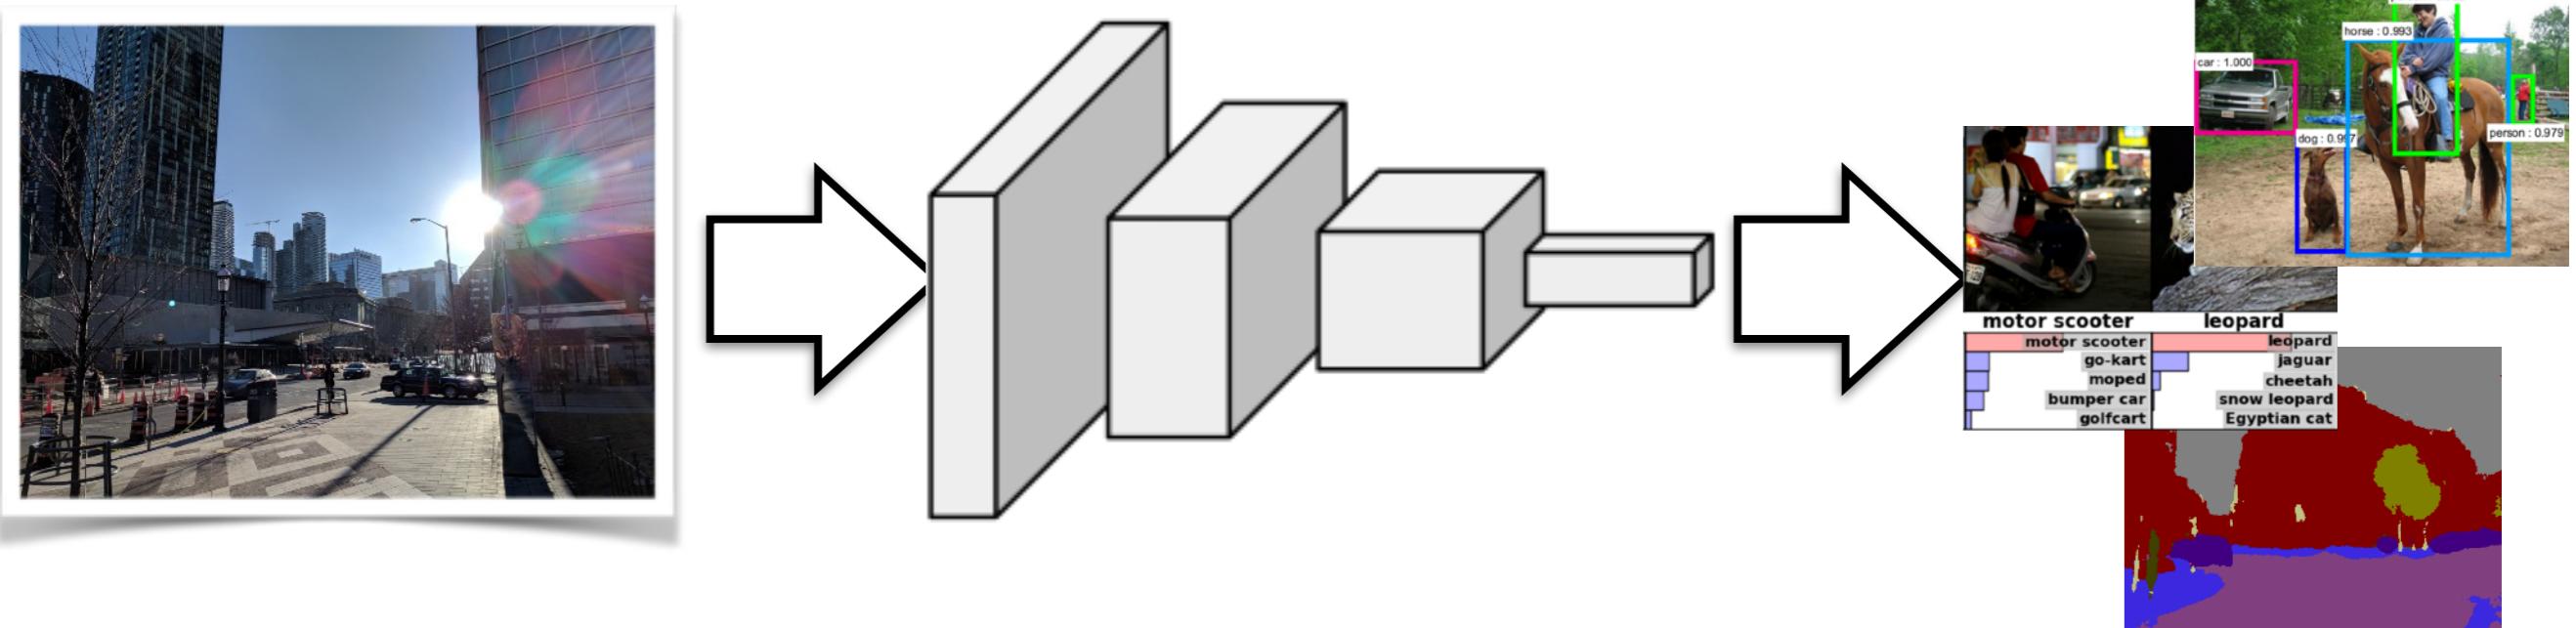
- Convolutional neural networks for classification, segmentation, regression, etc.

# Recap: CNNs



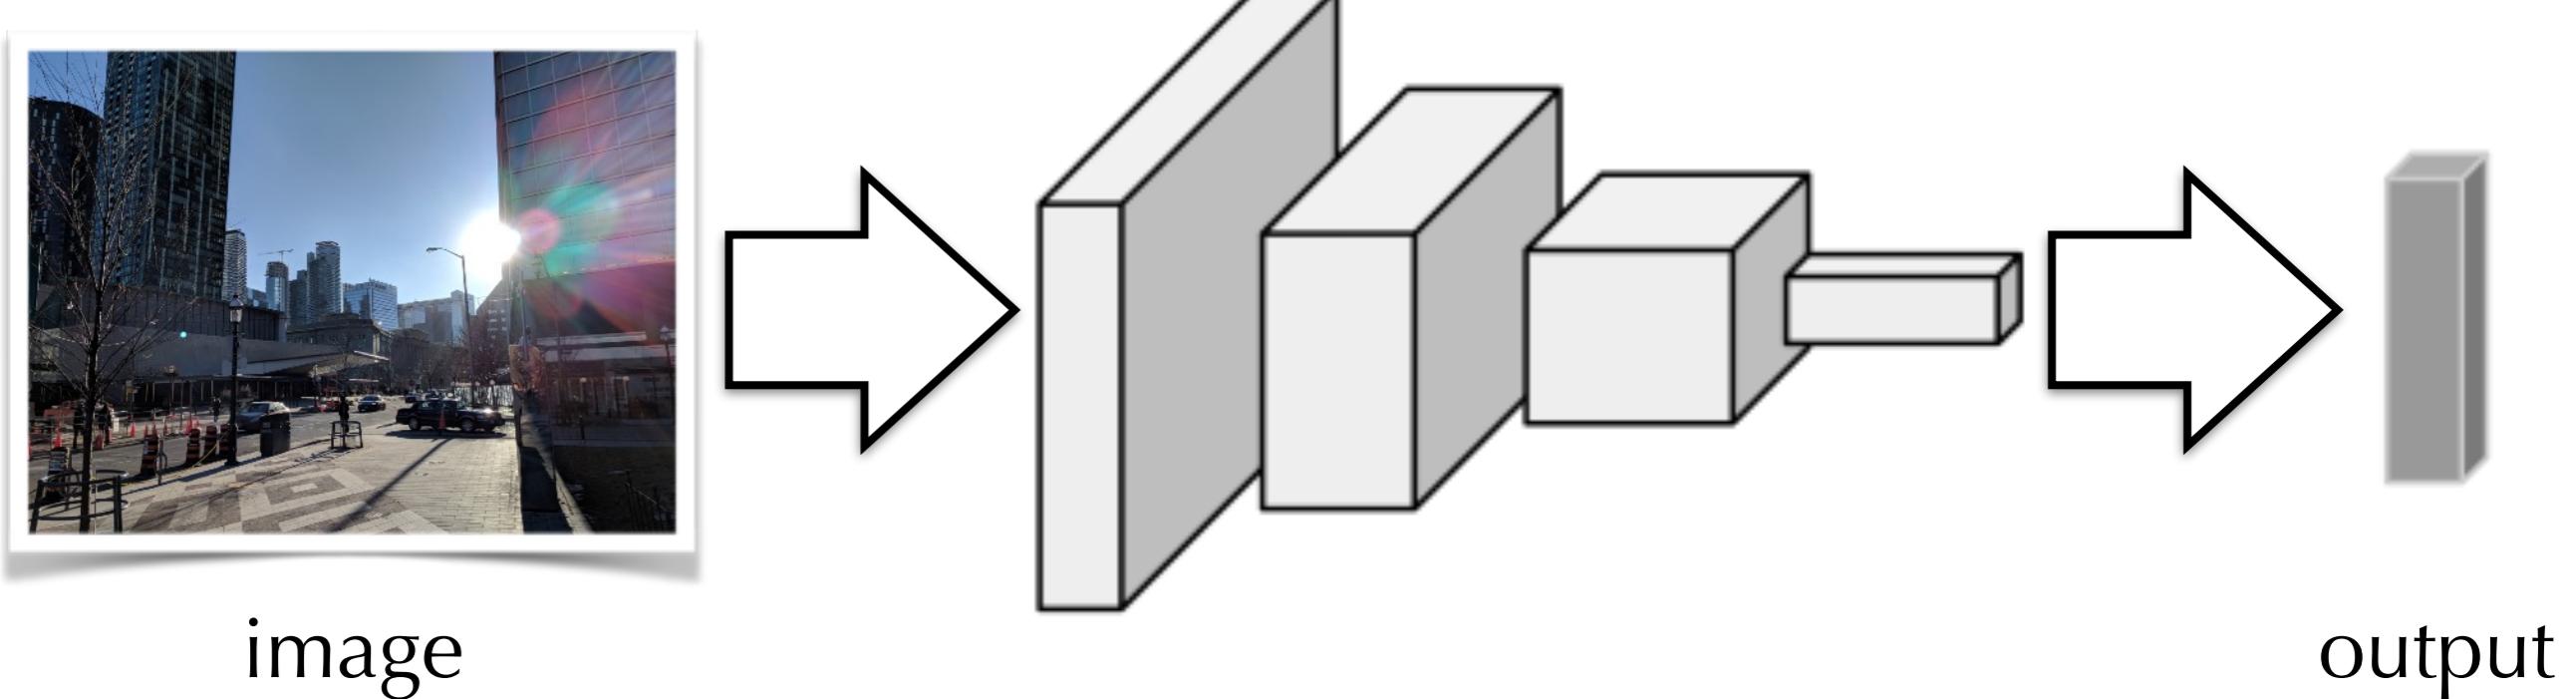
- Convolutional neural networks for classification, segmentation, regression, etc.
- Supervised learning: Training from (image, label) pairs

# Recap: CNNs

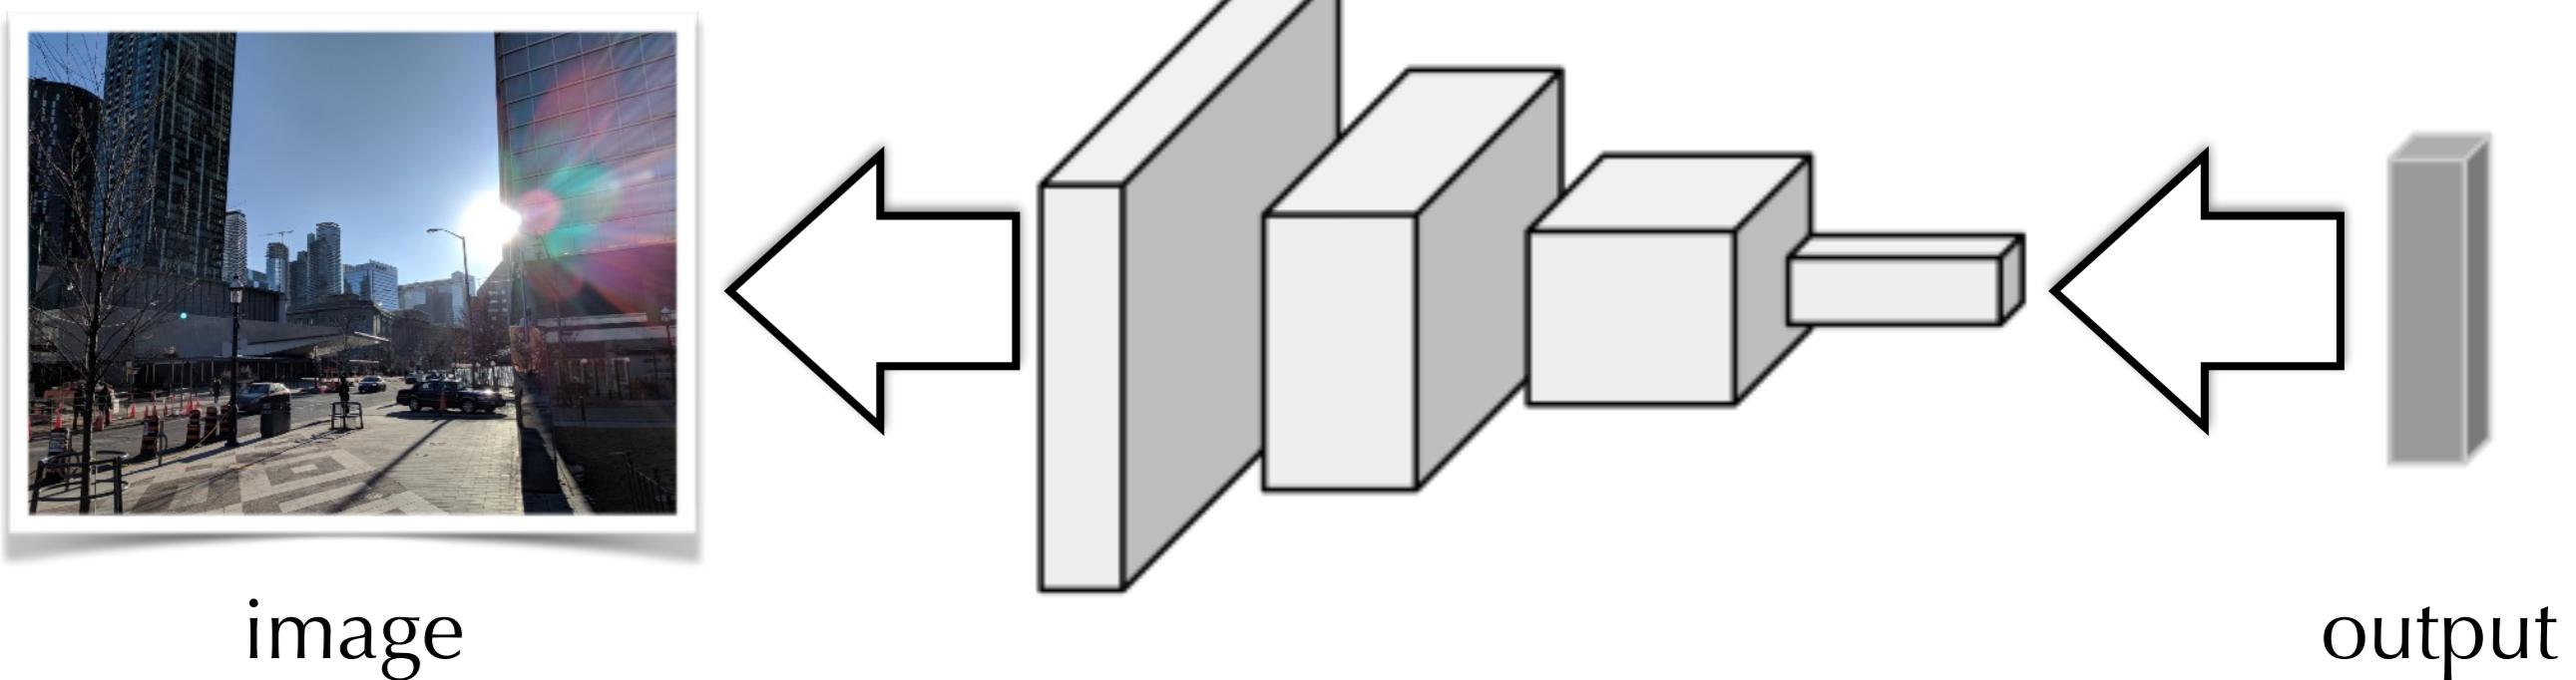


- Convolutional neural networks for classification, segmentation, regression, etc.
- Supervised learning: Training from (image, label) pairs
- Training via stochastic gradient descent (or variants)

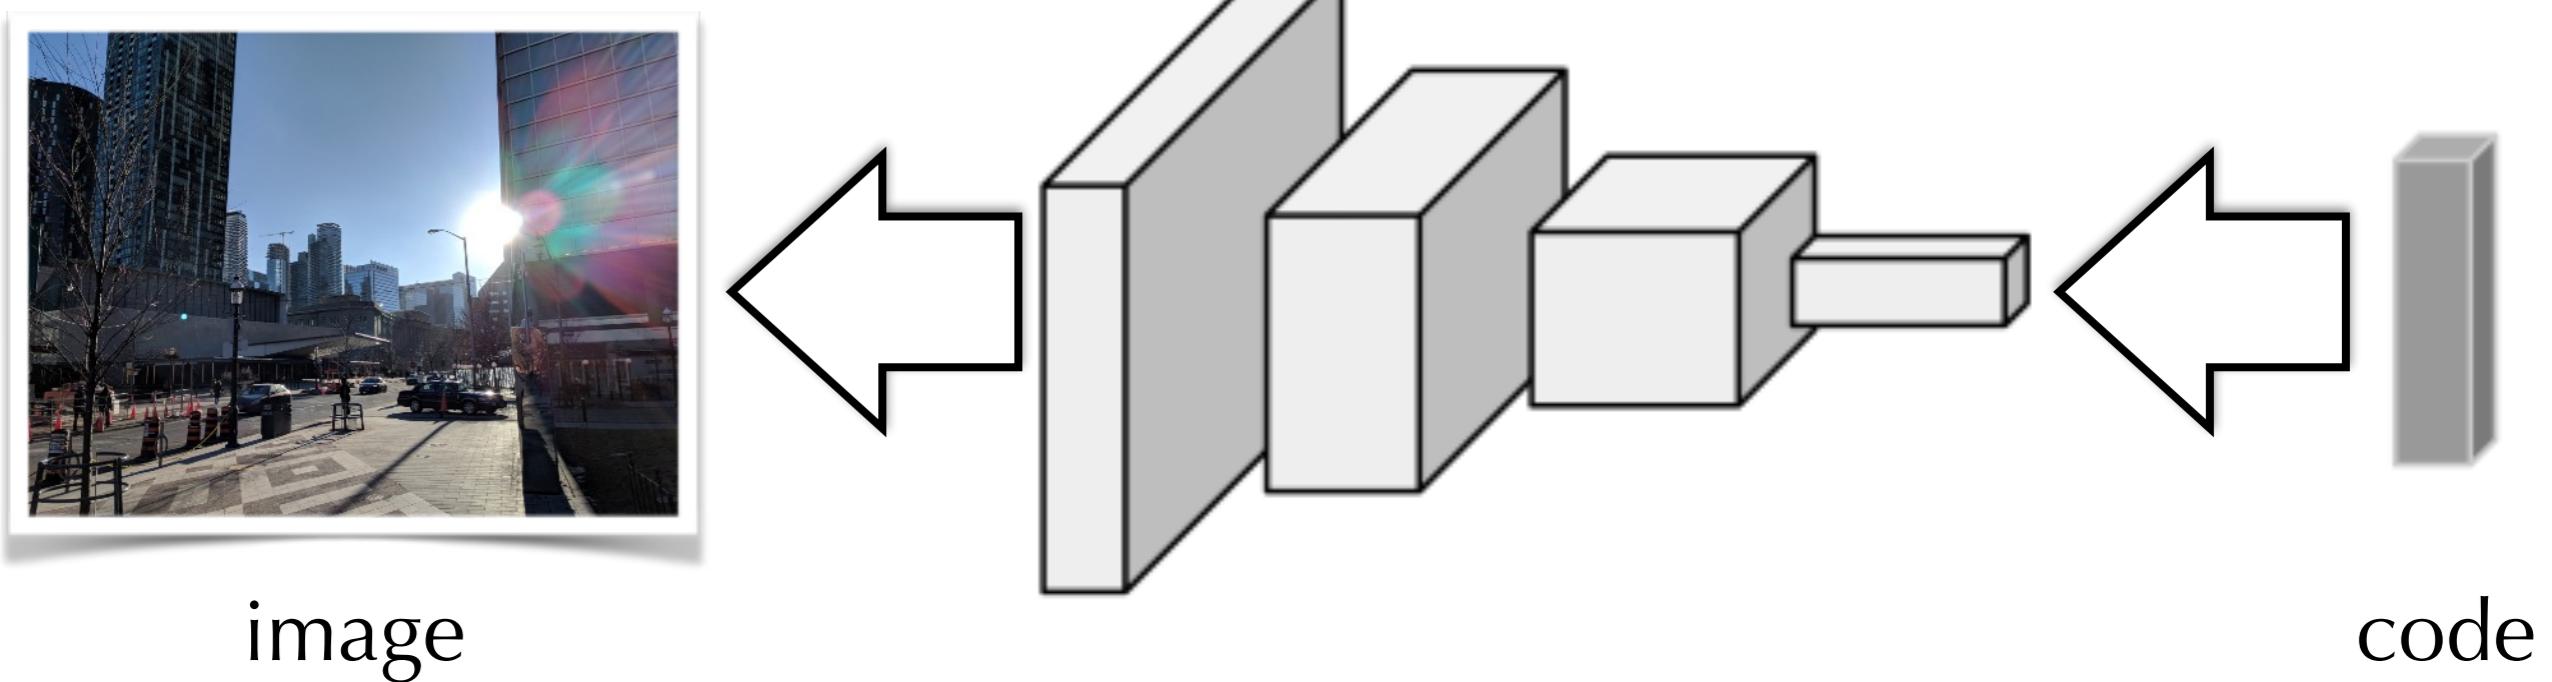
# Generative Neural Networks



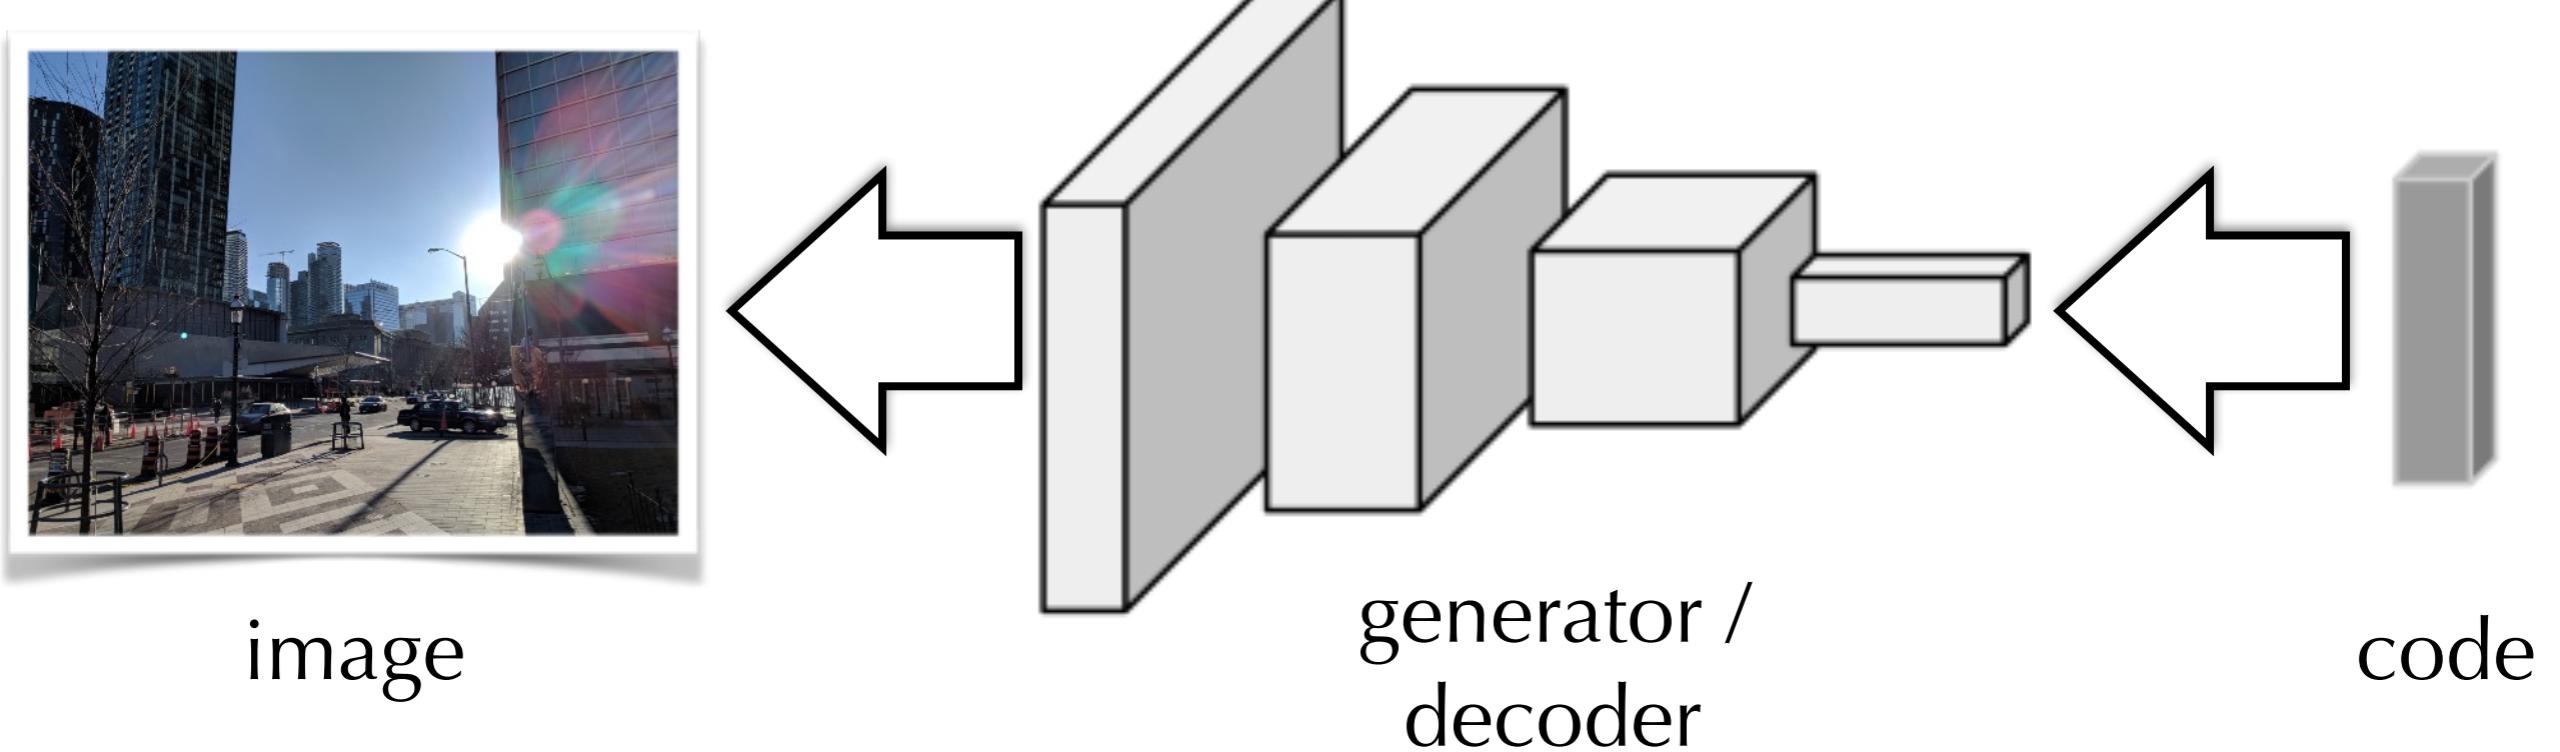
# Generative Neural Networks



# Generative Neural Networks



# Generative Neural Networks



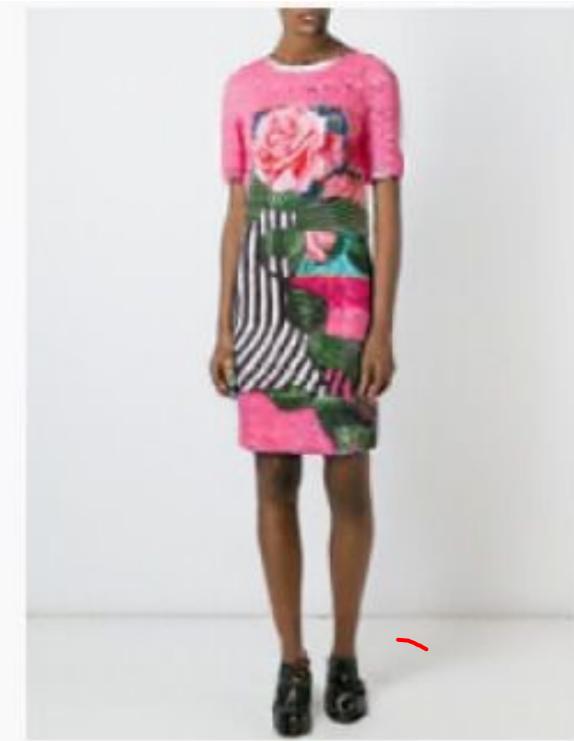
# Automatic Image Generation



[Nguyen et al., Plug & Play Generative Networks: Conditional Iterative Generation of Images in Latent Space, CVPR 2017]

# Commercial Applications

*real*



*real*

# Image-to-Image Translation



[Zhu et al., Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, ICCV 2017]

# Image-to-Image Translation



[Zhu et al., Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, ICCV 2017]

# Image-to-Image Translation



[Zhu et al., Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, ICCV 2017]

# Artistic Style Transfer

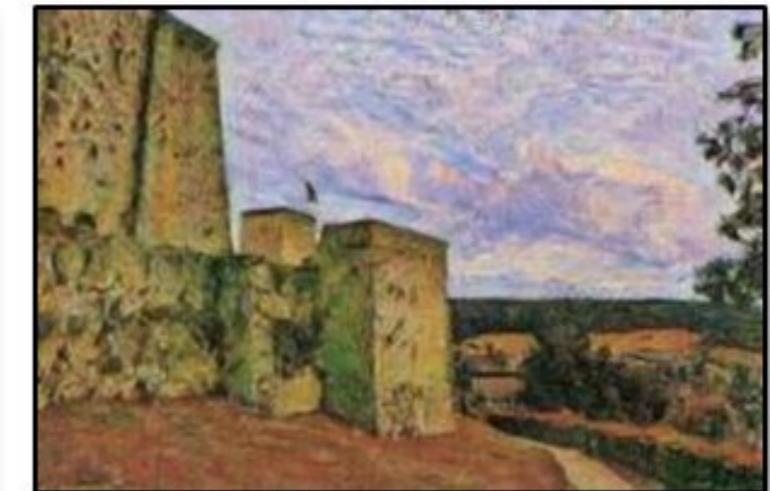
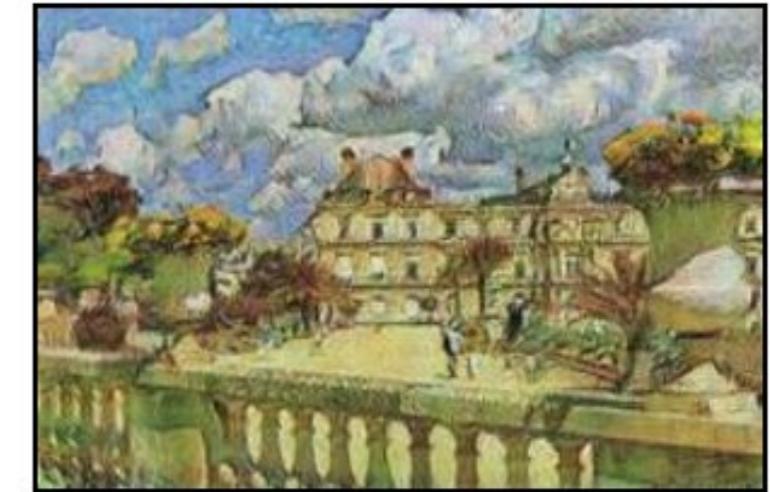
Input



Monet



Van Gogh



[Zhu et al., Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, ICCV 2017]

# Training Data Generation

Input sunny image



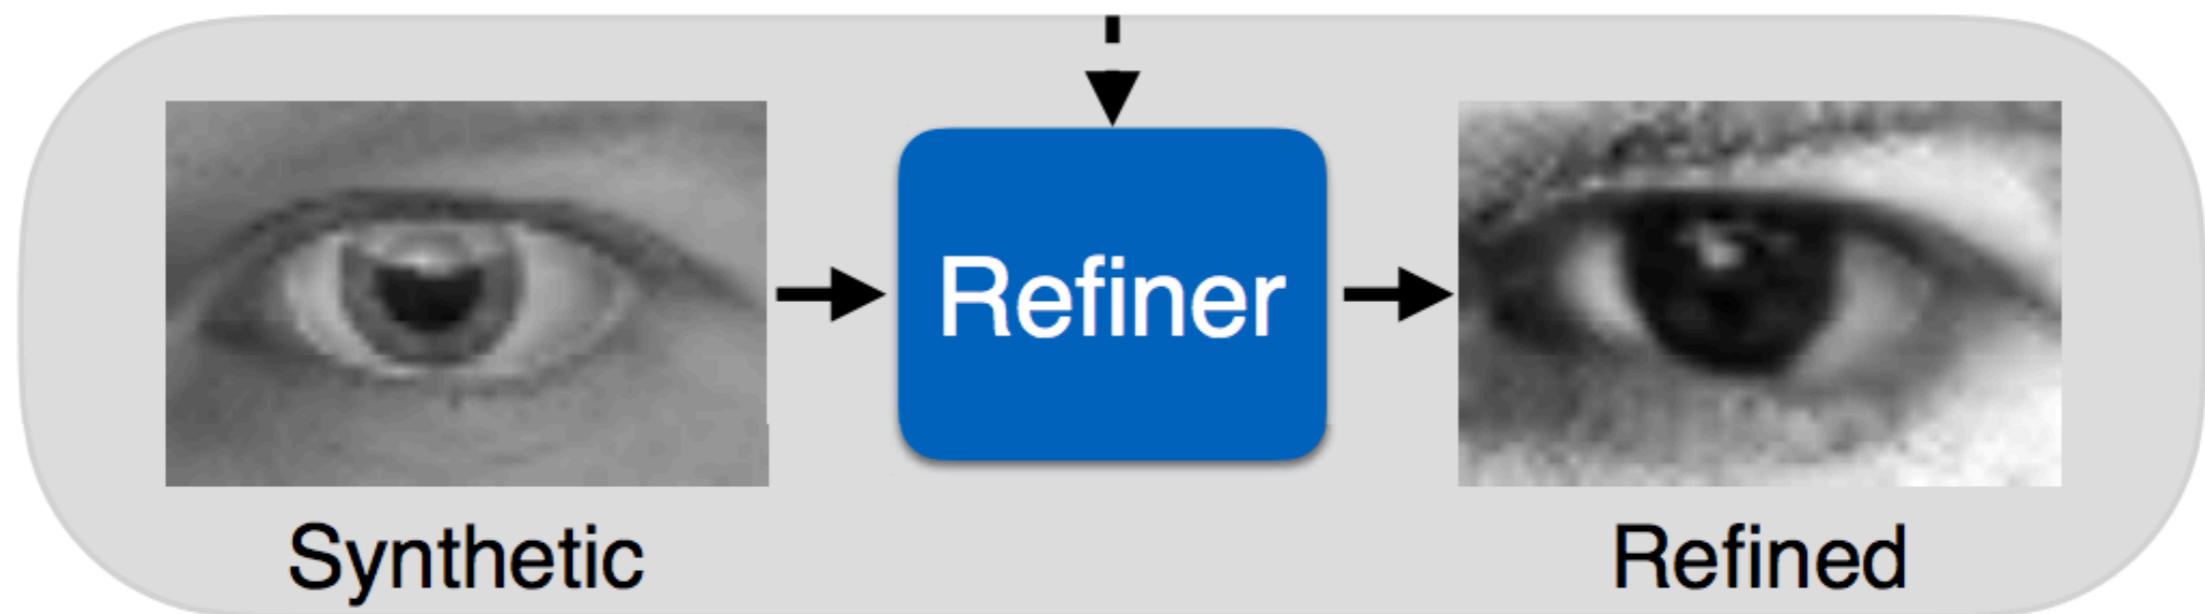
*AI-generated rainy image*



[Liu et al., Unsupervised Image-to-Image Translation Networks, NIPS 2017]

# Training Data Generation

Unlabeled Real Images



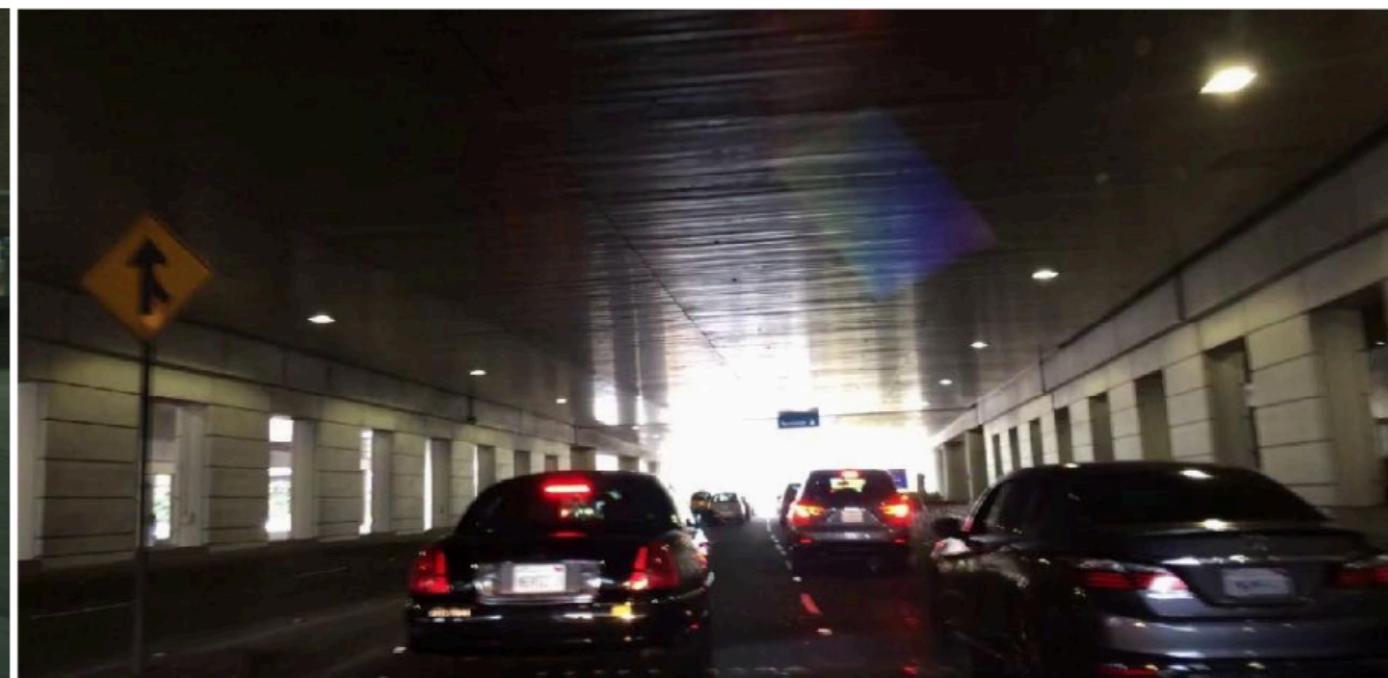
[Shrivastava et al., Learning from Simulated and Unsupervised Images through Adversarial Training, CVPR 2017]

# Domain Adaptation

**Train**



**Test**



**CityScapes (Germany)**

Source domain:  
Plenty of labelled data

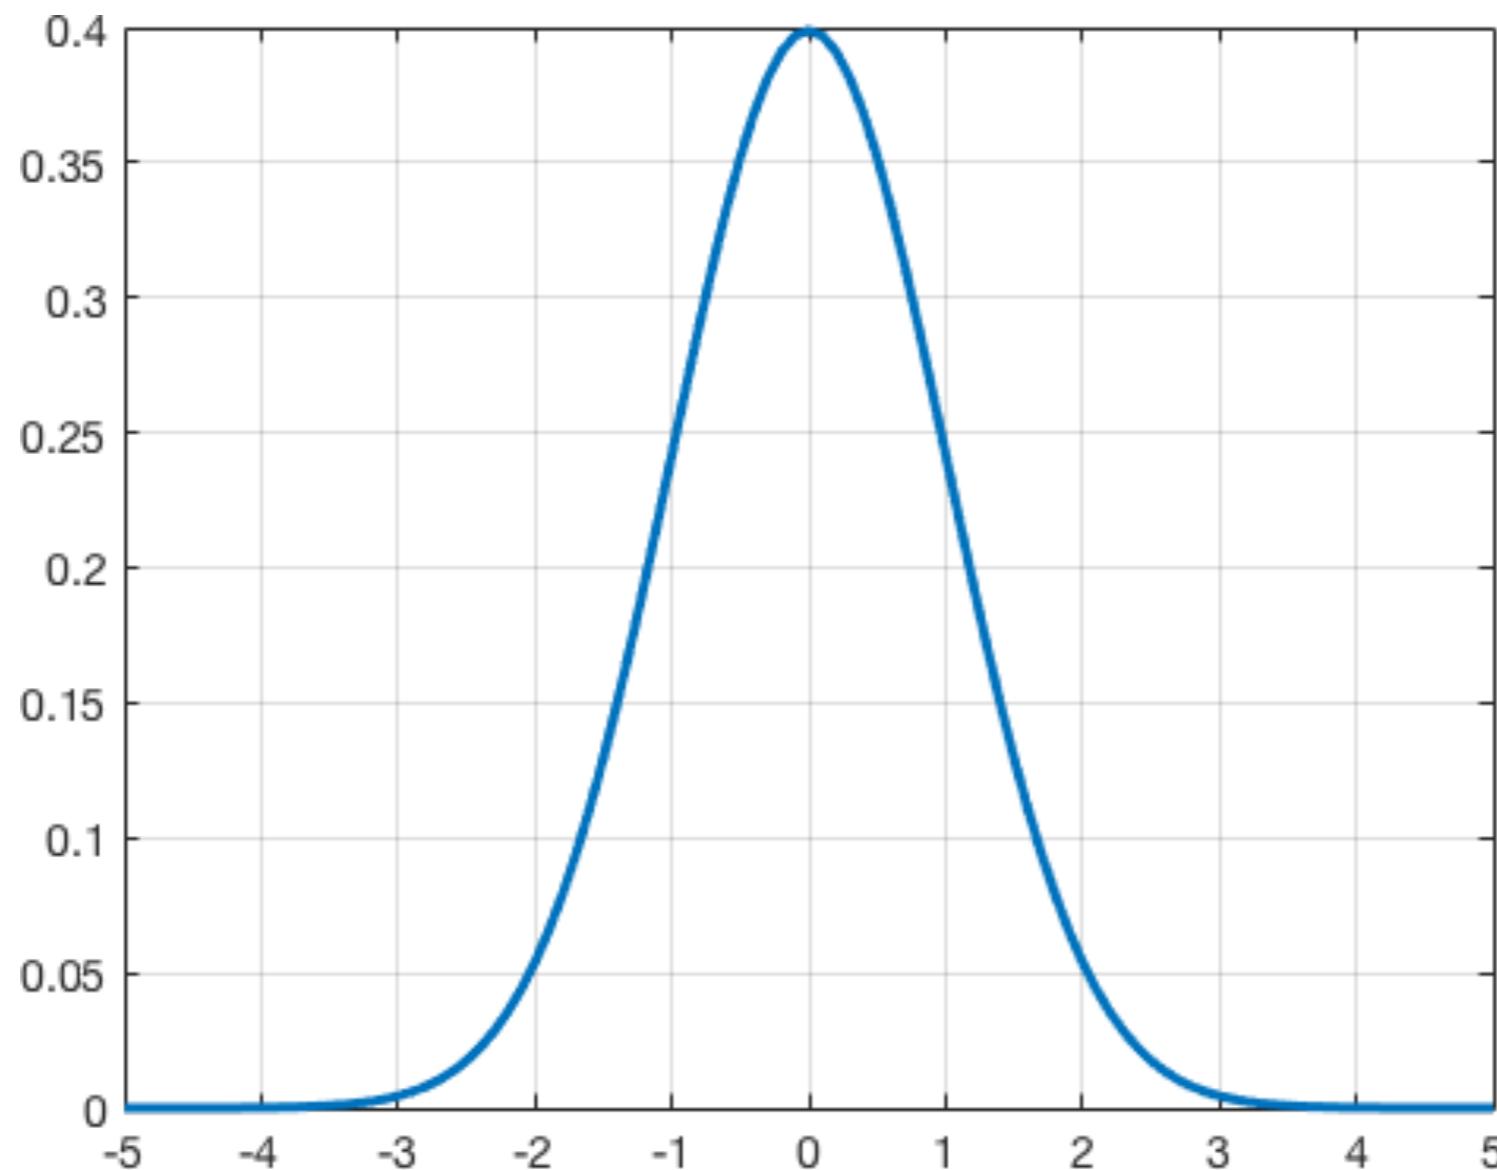
**San Francisco**

Target domain:  
Little to no labelled data

slide credit: Judy Hoffman

# Generation via Sampling

Example: Sampling from normal distribution



# Generation via Sampling

Example: Sampling from normal distribution

Box-Muller Transform:

# Generation via Sampling

Example: Sampling from normal distribution

Box-Muller Transform:

- Input:  $U, V$  drawn uniformly from  $(0,1)$

# Generation via Sampling

Example: Sampling from normal distribution

Box-Muller Transform:

- Input:  $U, V$  drawn uniformly from  $(0,1)$
- Output: Random numbers  $X, Y$  drawn from  $\mathcal{N}(0, 1)$

# Generation via Sampling

Example: Sampling from normal distribution

Box-Muller Transform:

- Input:  $U, V$  drawn uniformly from  $(0,1)$
- Output: Random numbers  $X, Y$  drawn from  $\mathcal{N}(0, 1)$
- Algorithm:

$$X = \sqrt{-2 \ln U} \cos(2\pi V)$$

$$Y = \sqrt{-2 \ln U} \sin(2\pi V)$$

# Sampling from Distribution of Images?

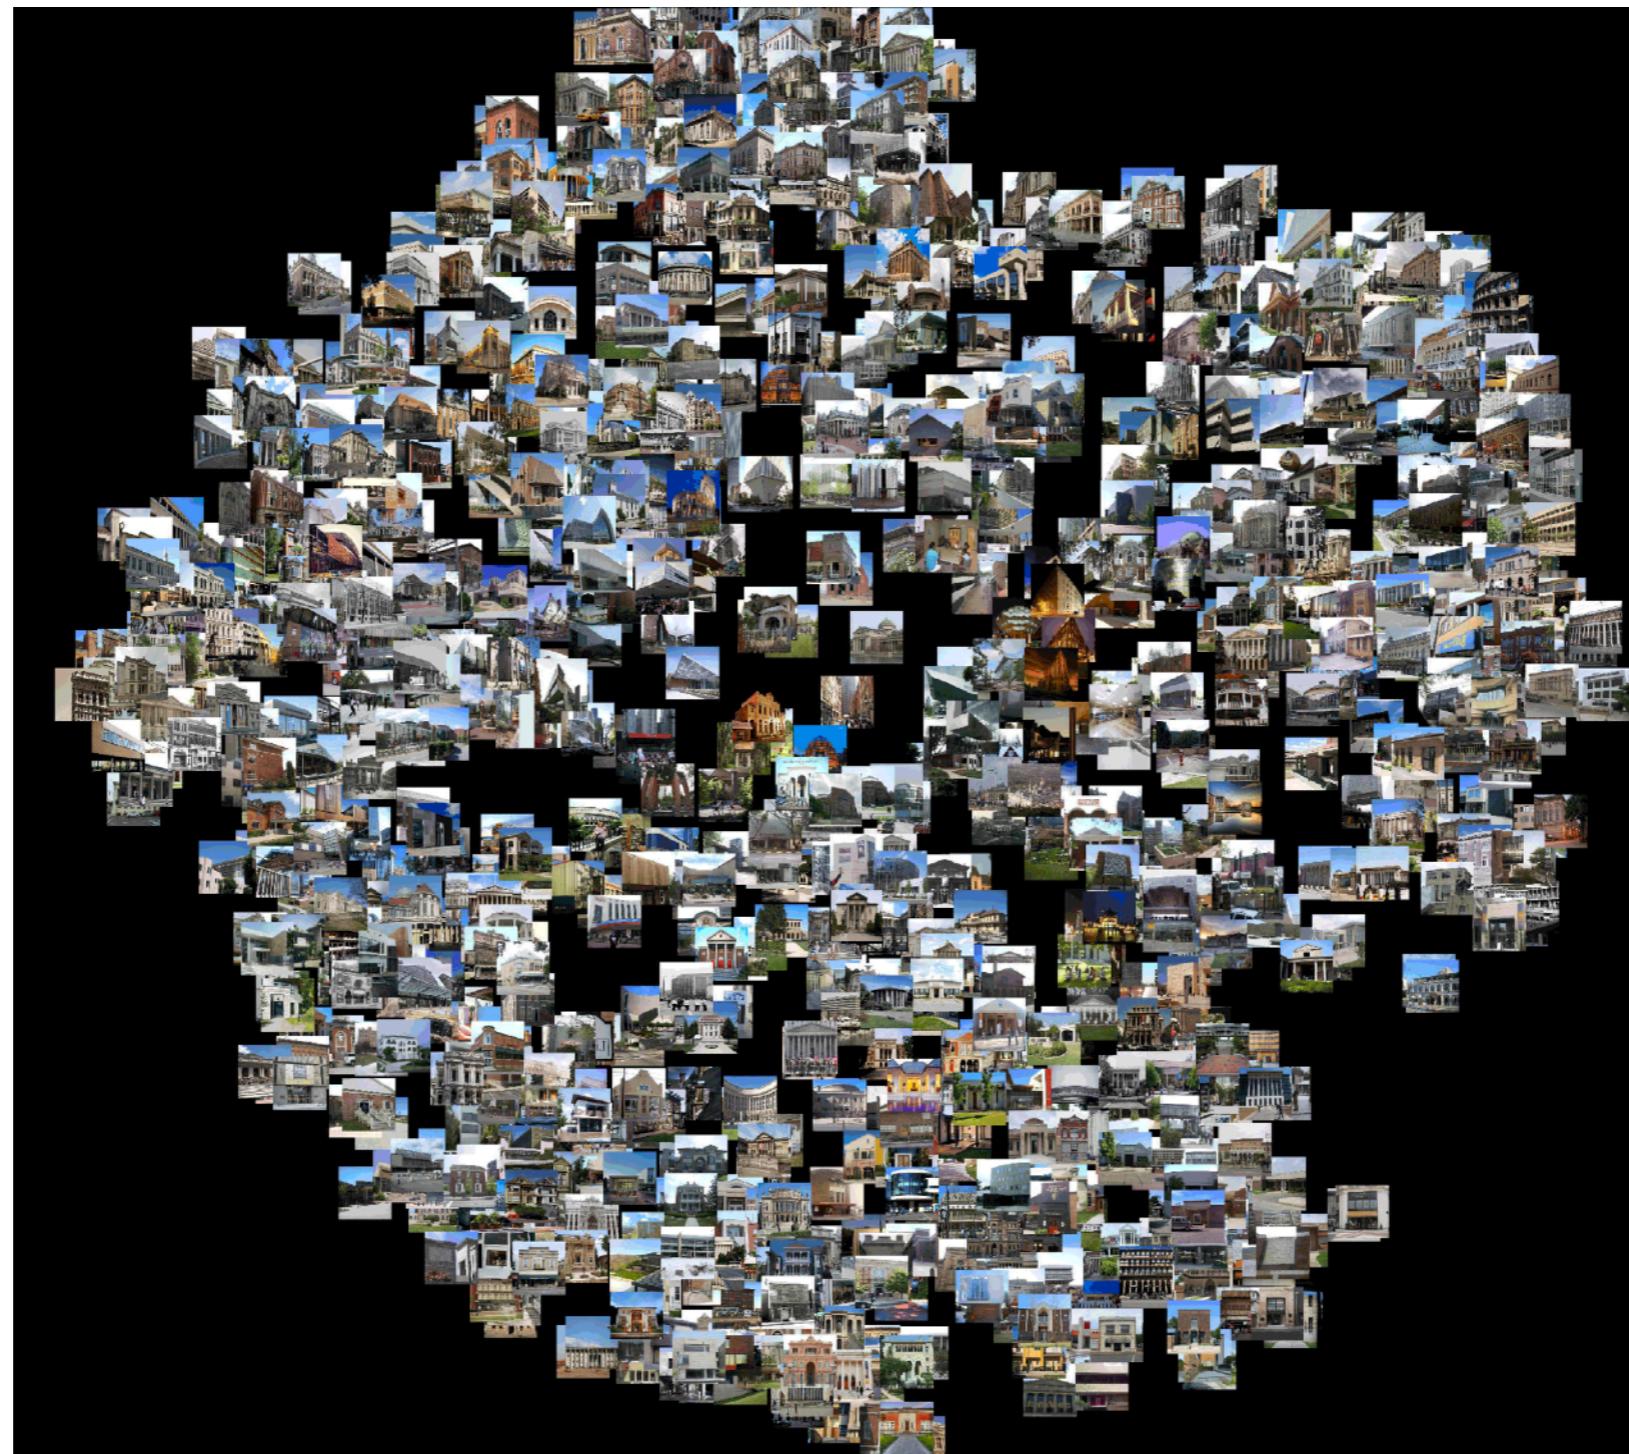


image from <http://3drepresentation.stanford.edu/>

# Generative Neural Networks

code



$$\mathbf{z} \sim p_{\text{model}}$$

# Generative Neural Networks

code

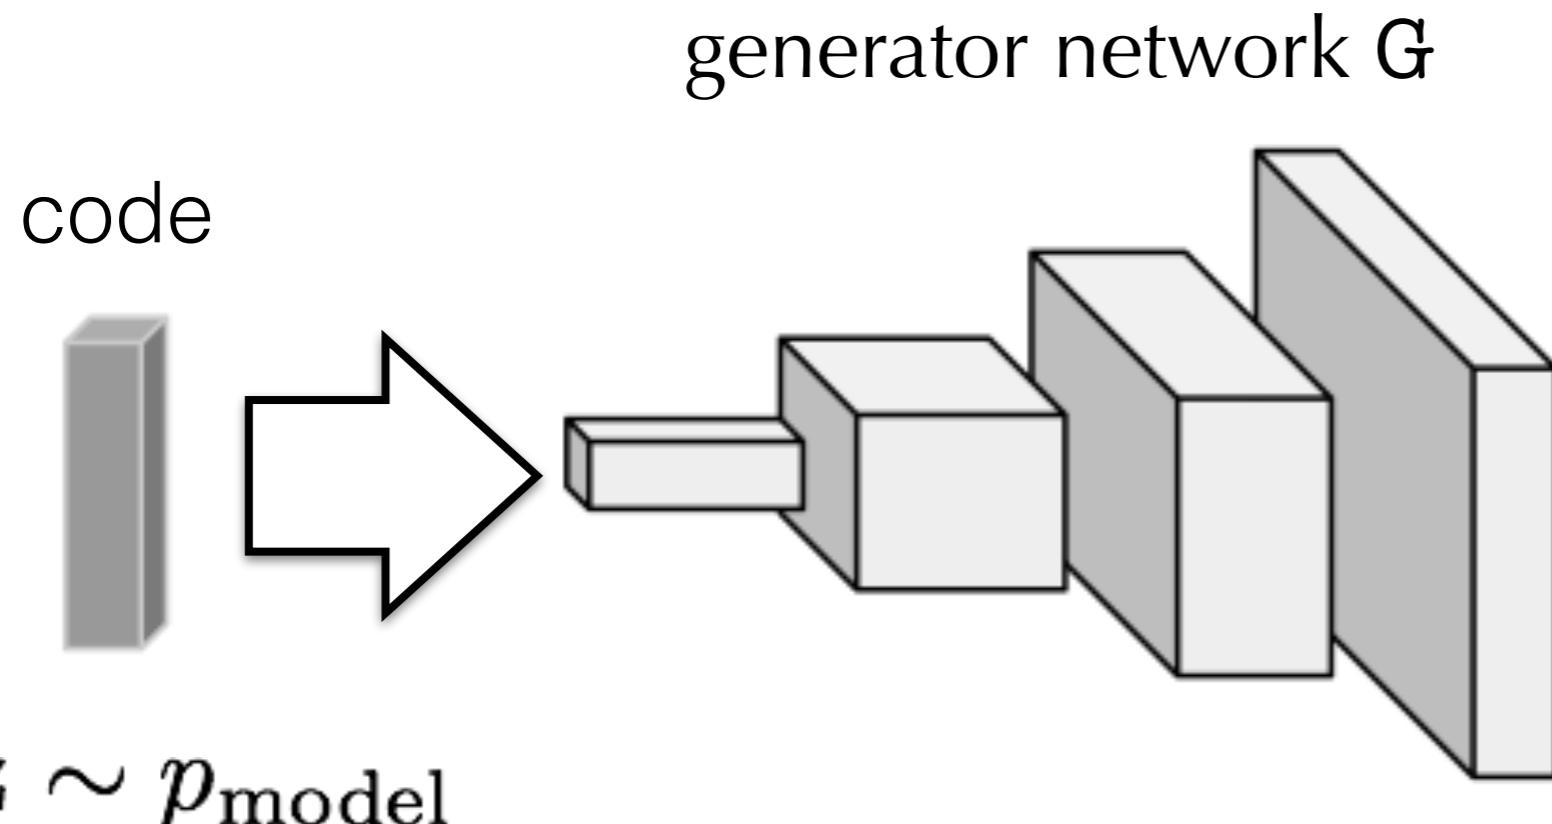


$$\mathbf{z} \sim p_{\text{model}}$$

normal distribution,  
uniform distribution,

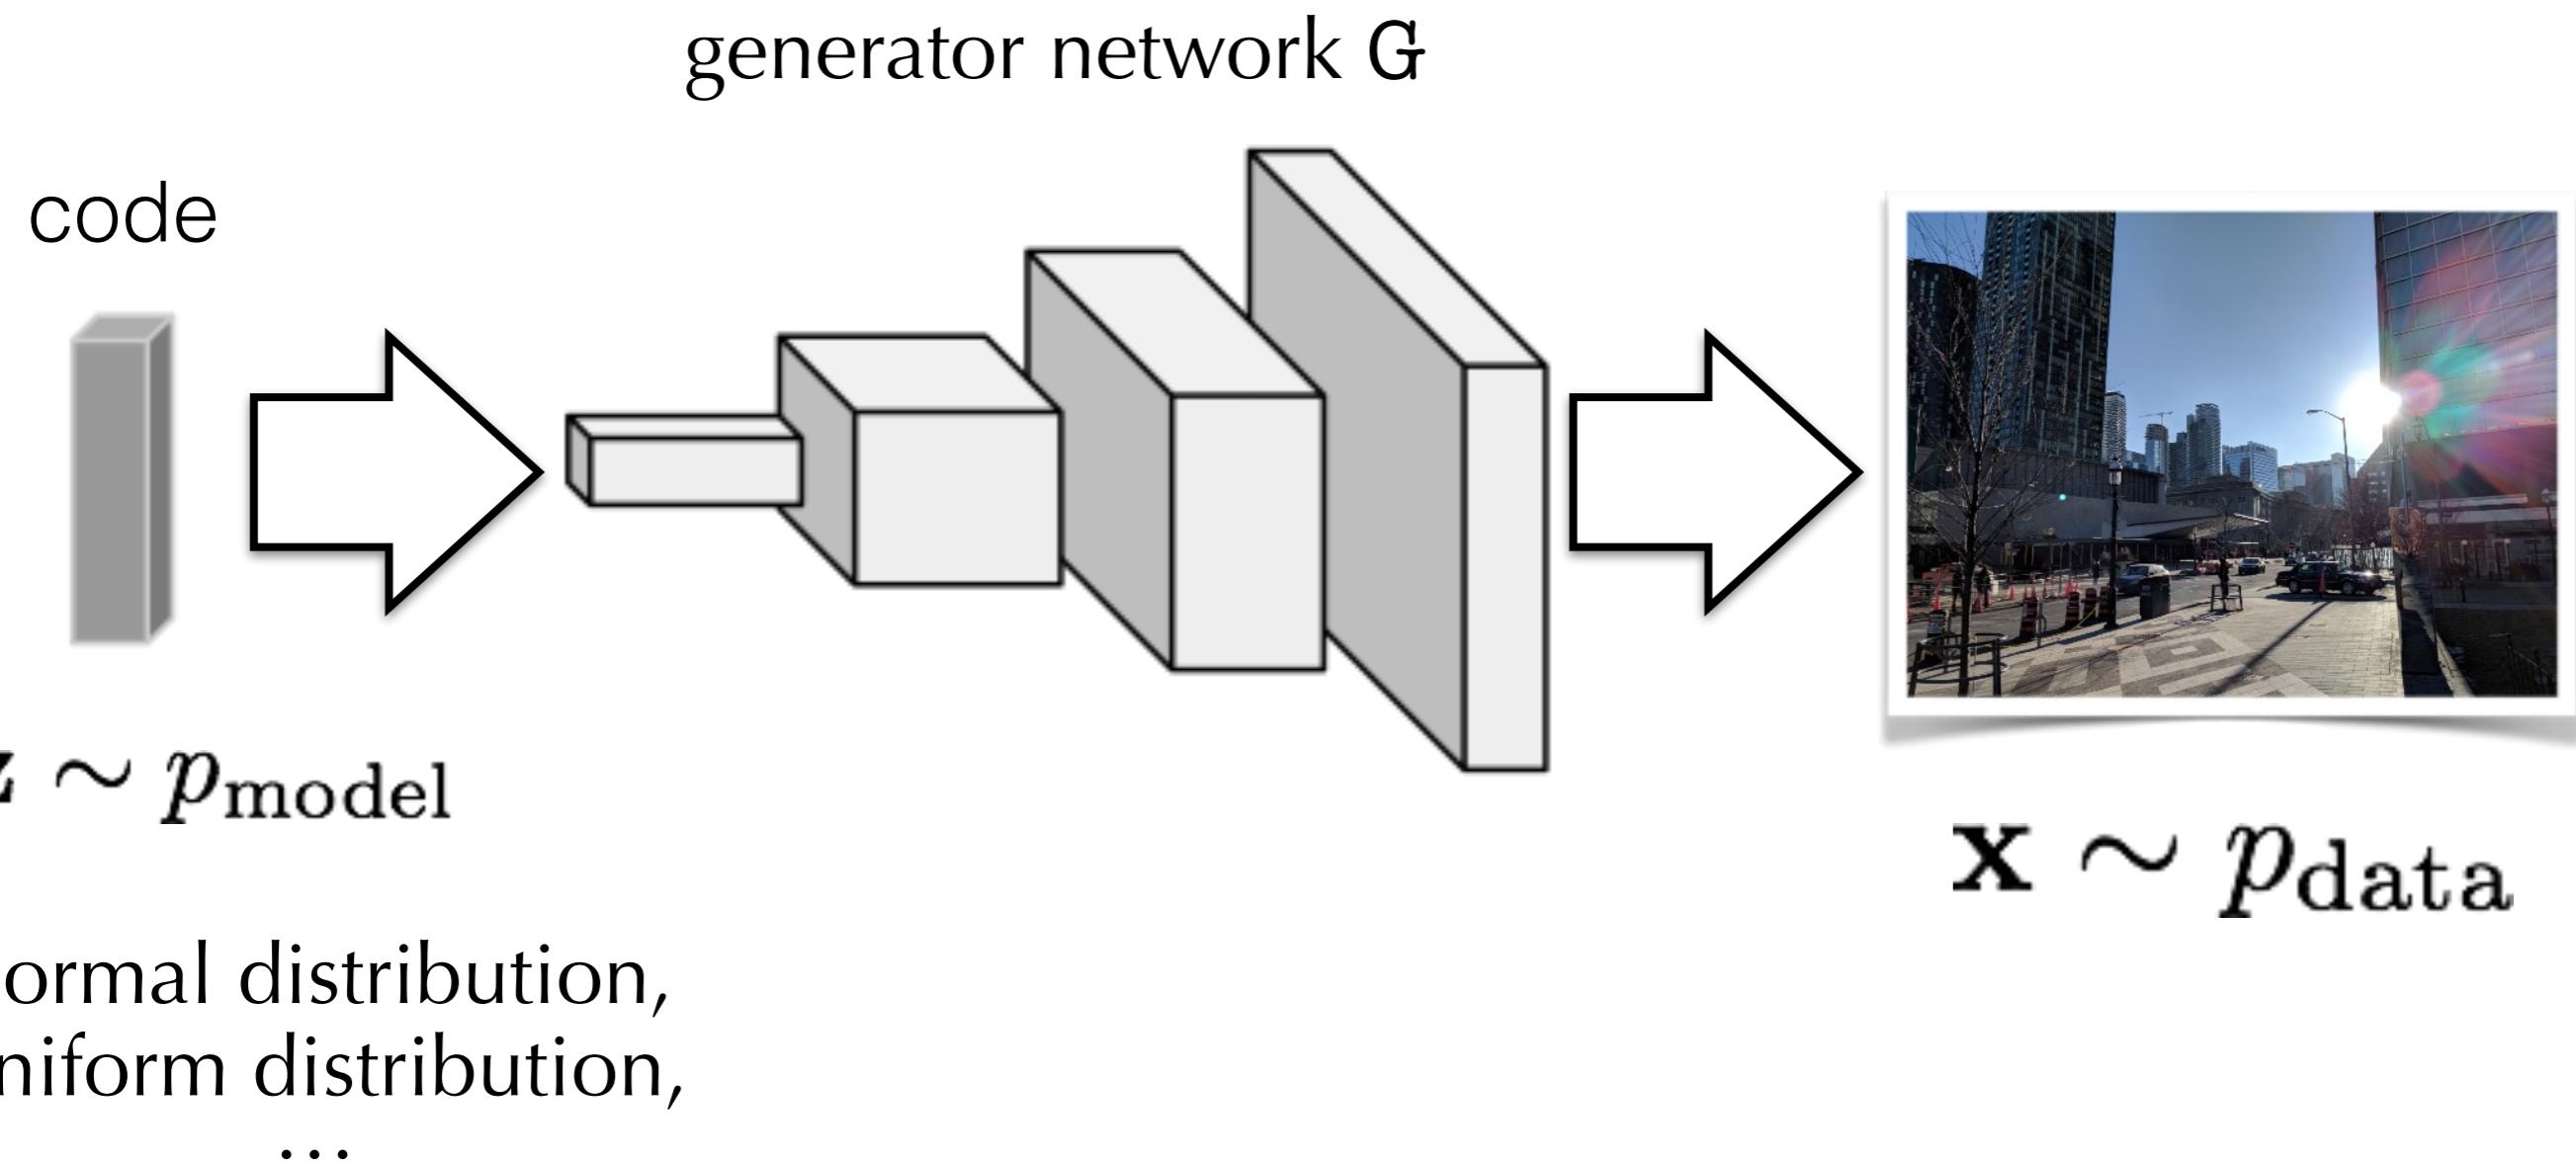
...

# Generative Neural Networks

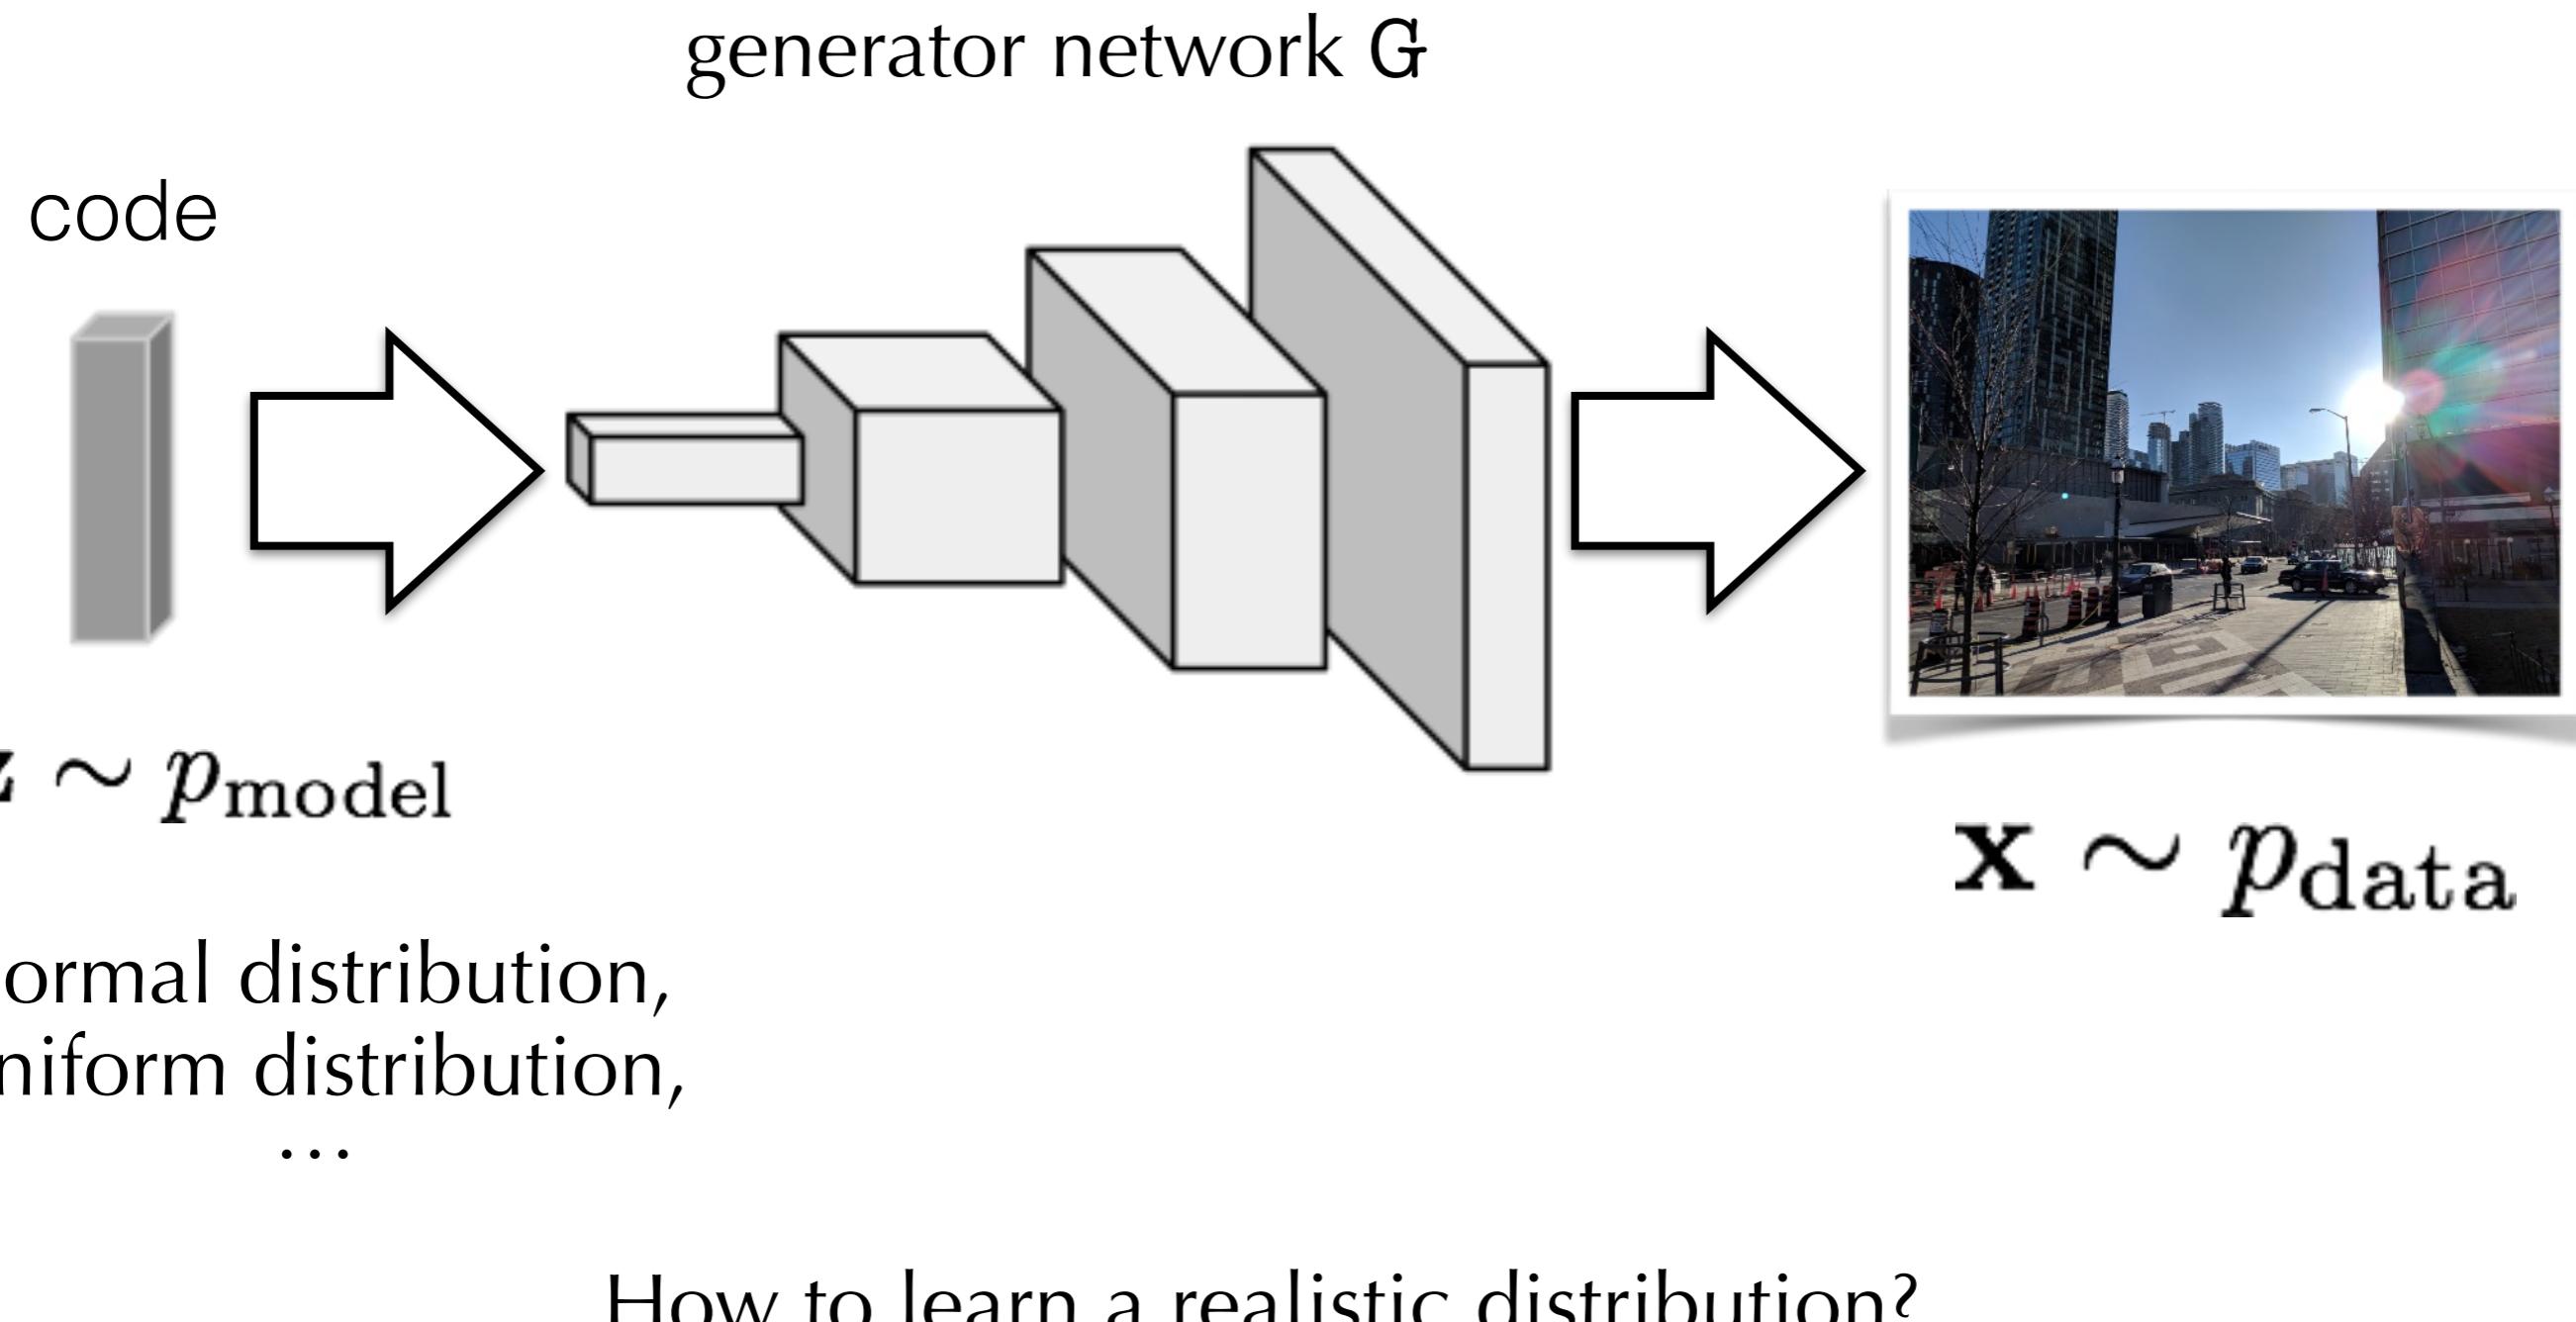


normal distribution,  
uniform distribution,  
...

# Generative Neural Networks



# Generative Neural Networks



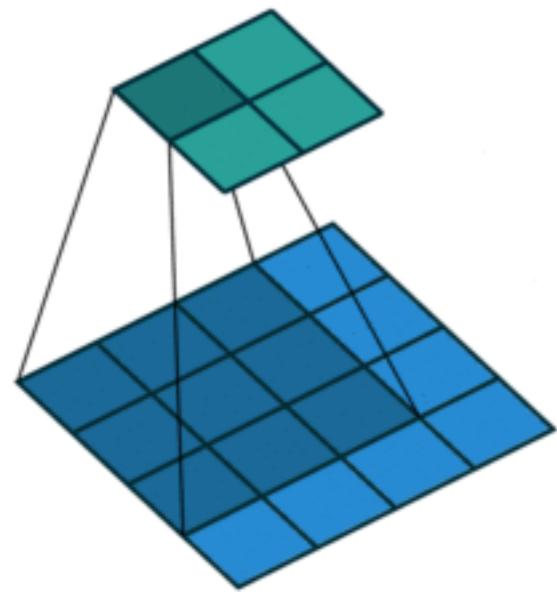
# Today

- Variational Autoencoders (VAE)
- Generative Adversarial Networks (GANs)
- Unsupervised Image Translation

# Today

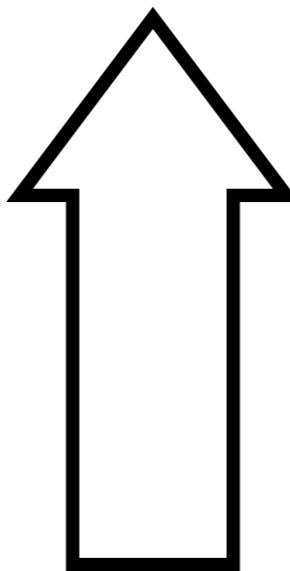
- **Variational Autoencoders (VAE)**
- Generative Adversarial Networks (GANs)
- Unsupervised Image Translation

# Recap: (Transposed) Convolutions



convolution

Output

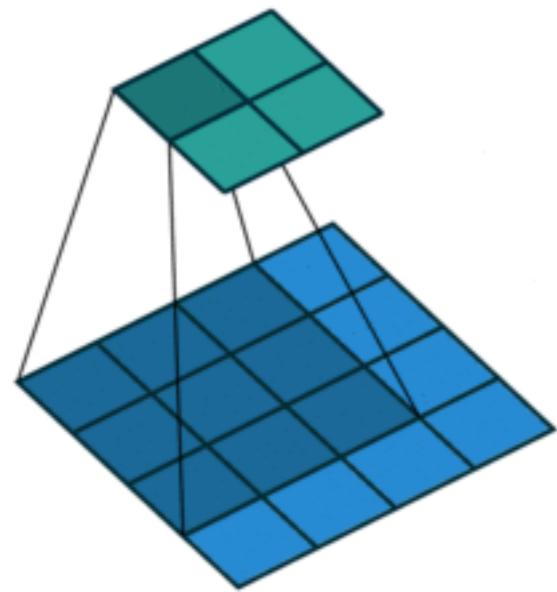


Input

illustrations from [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)

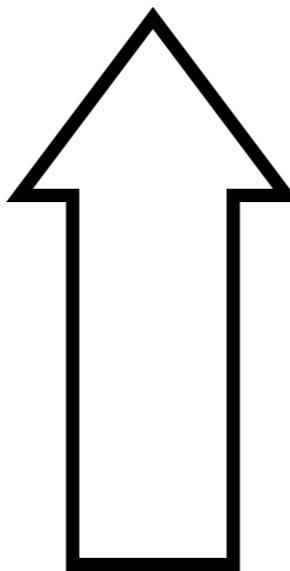
slide credit: Laura Leal-Taixé

# Recap: (Transposed) Convolutions



convolution

Output

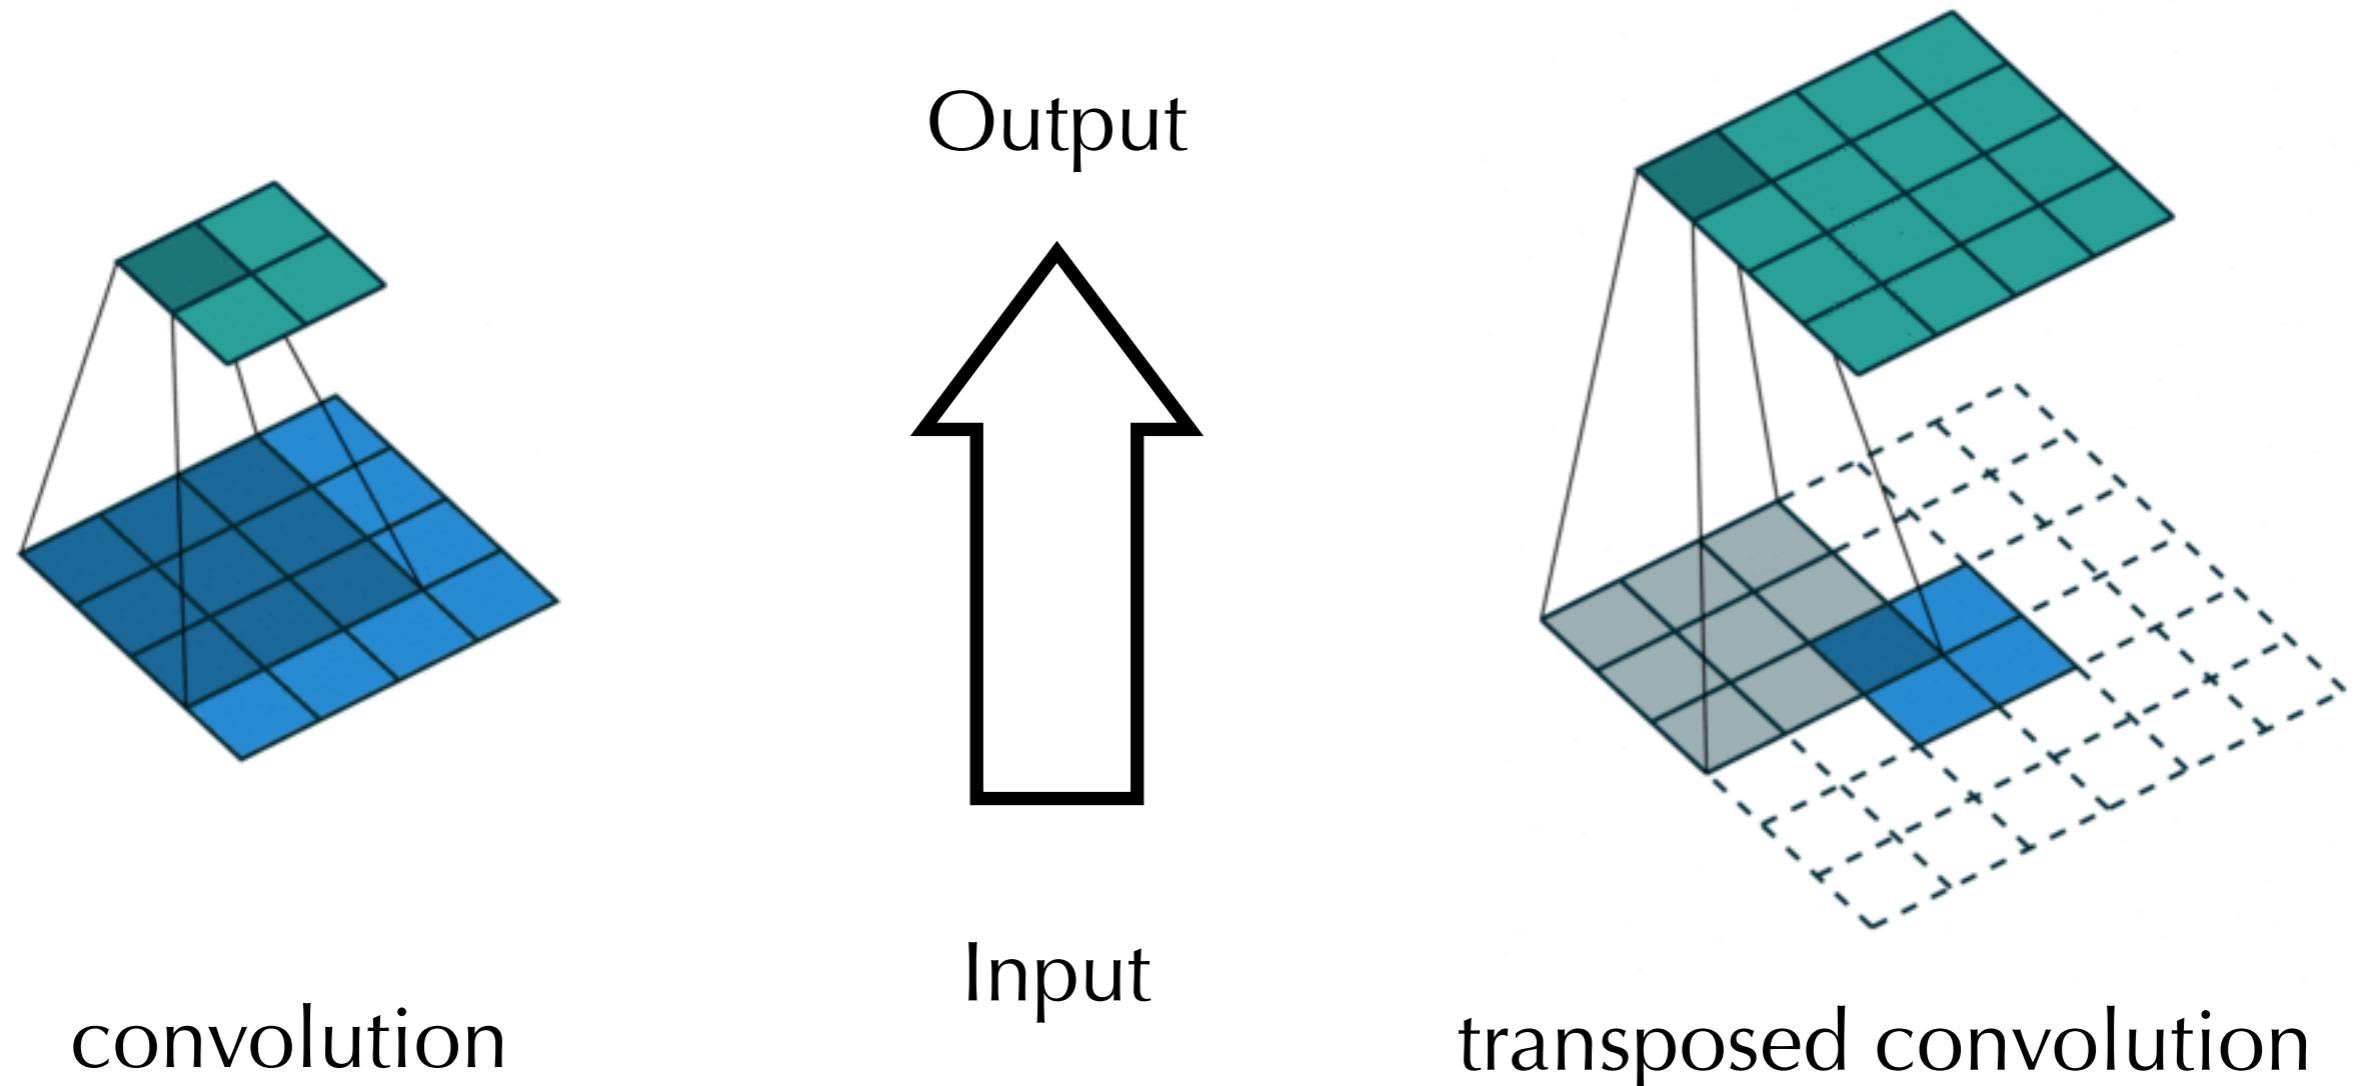


Input

illustrations from [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)

slide credit: Laura Leal-Taixé

# Recap: (Transposed) Convolutions

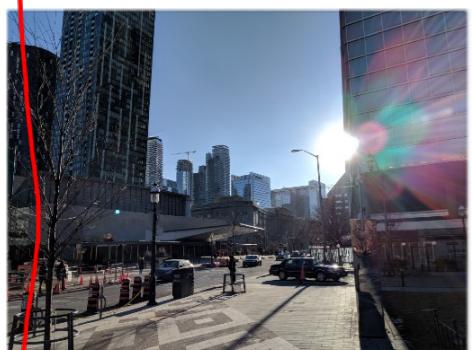


illustrations from [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)

slide credit: Laura Leal-Taixé

# Autoencoder

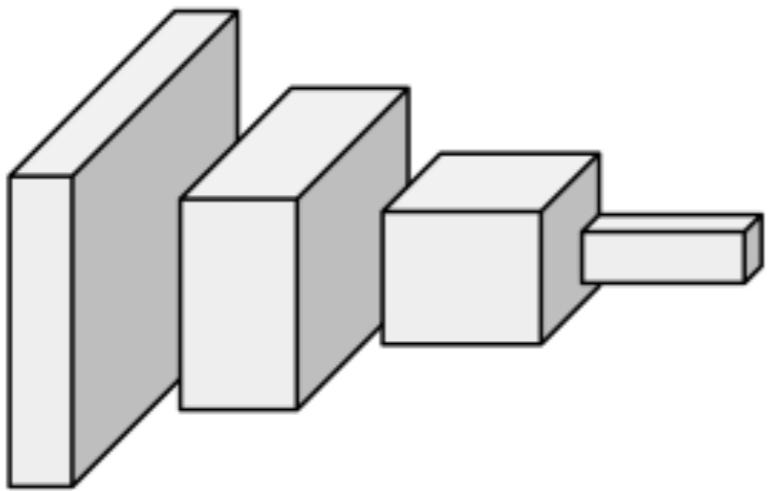
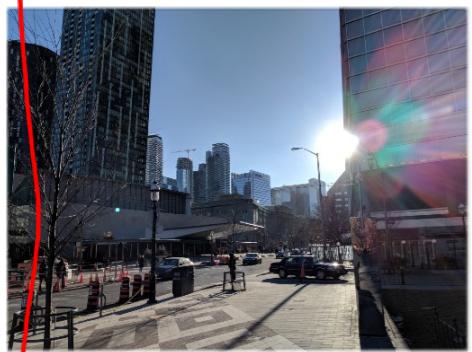
in practice /  
at test time



use rather  
low-dimensional  
space

# Autoencoder

in practice /  
at test time



encoder  
(convolutions)

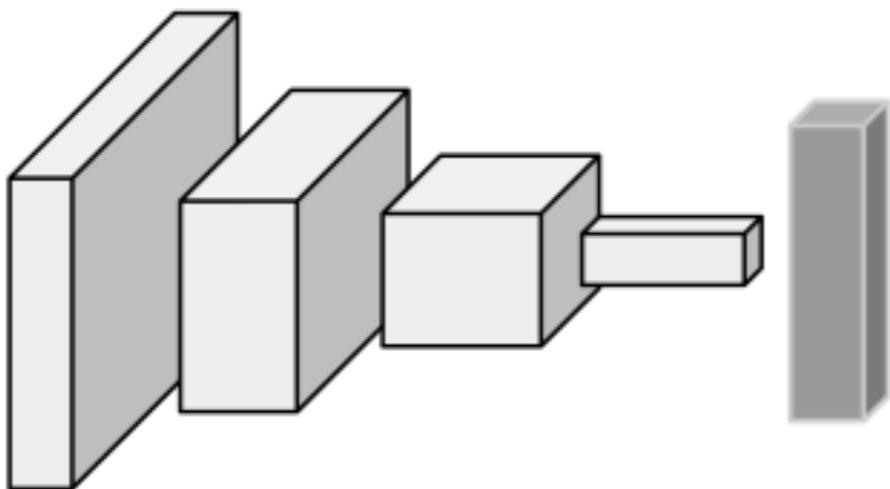
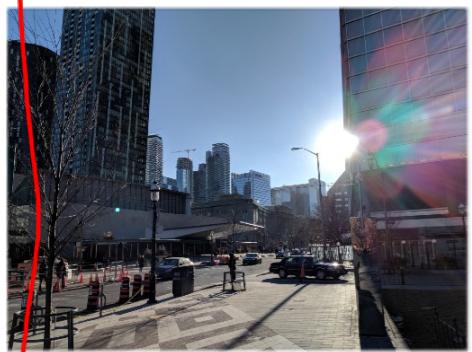
use rather  
low-dimensional  
space

# Autoencoder

in practice /  
at test time

latent  
space

use rather  
low-dimensional  
space



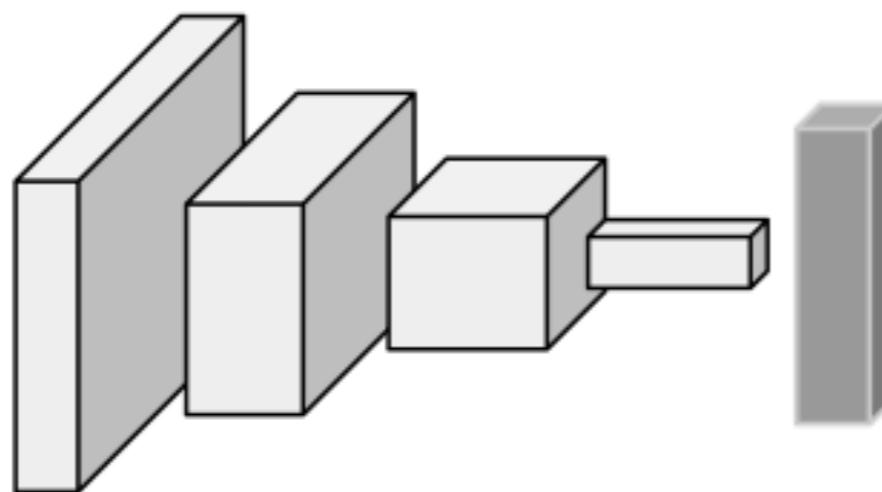
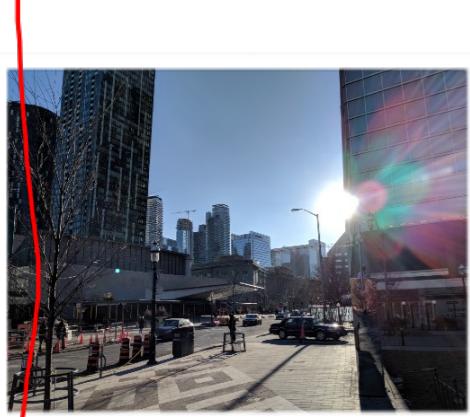
encoder  
(convolutions)

# Autoencoder

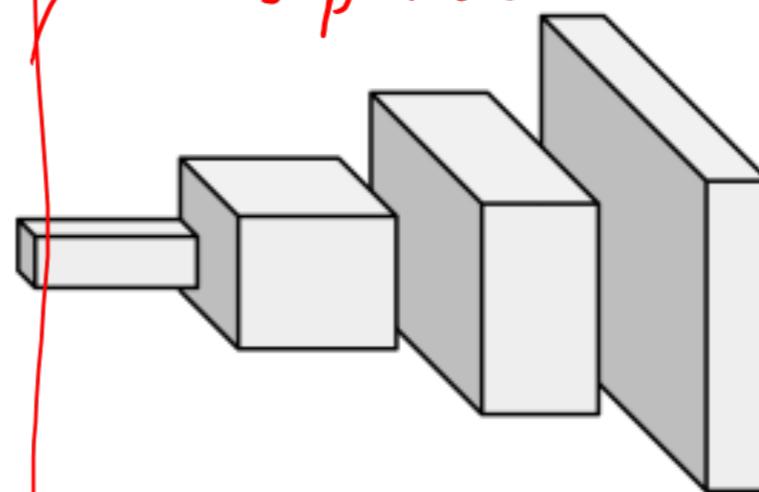
in practice /  
at test time

latent  
space

use rather  
low-dimensional  
space



encoder  
(convolutions)



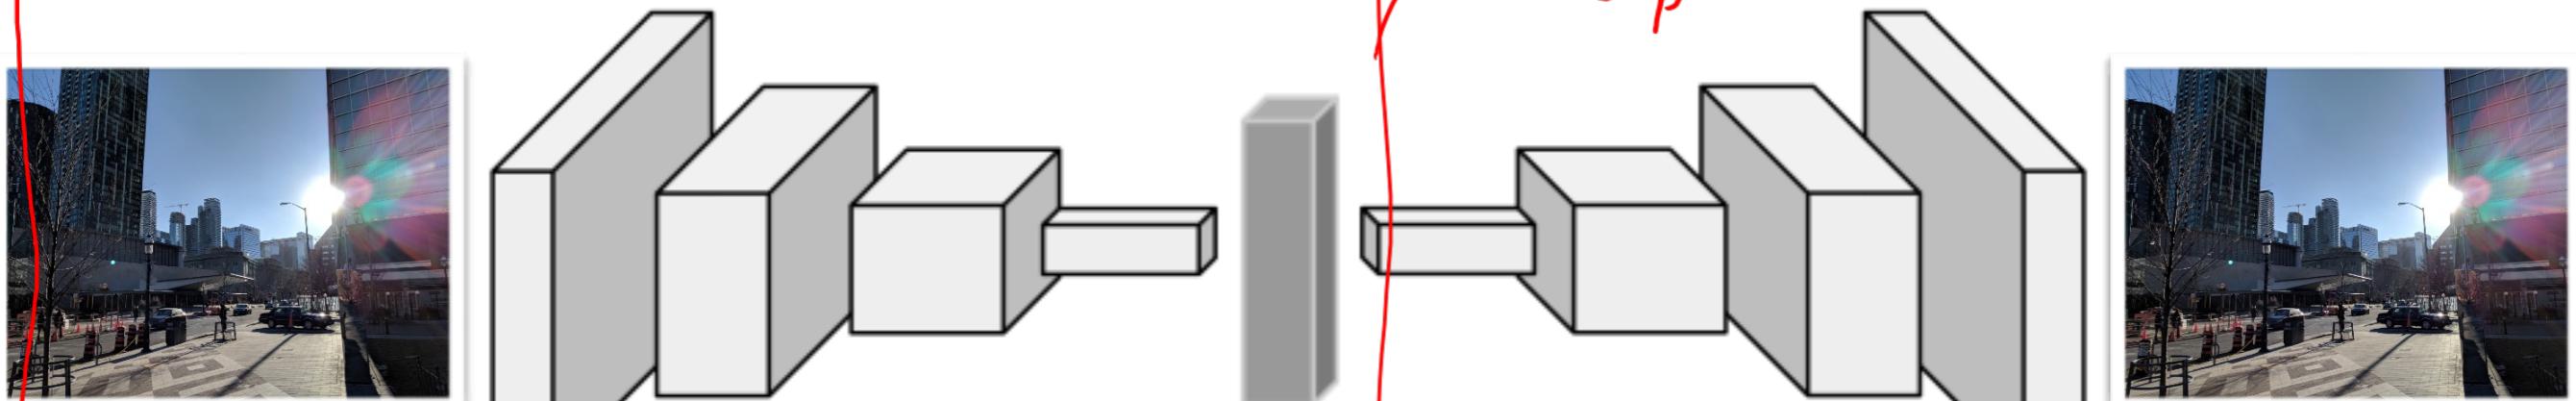
decoder  
(transp. convolutions)

# Autoencoder

in practice /  
at test time

latent  
space

use rather  
low-dimensional  
space



encoder  
(convolutions)

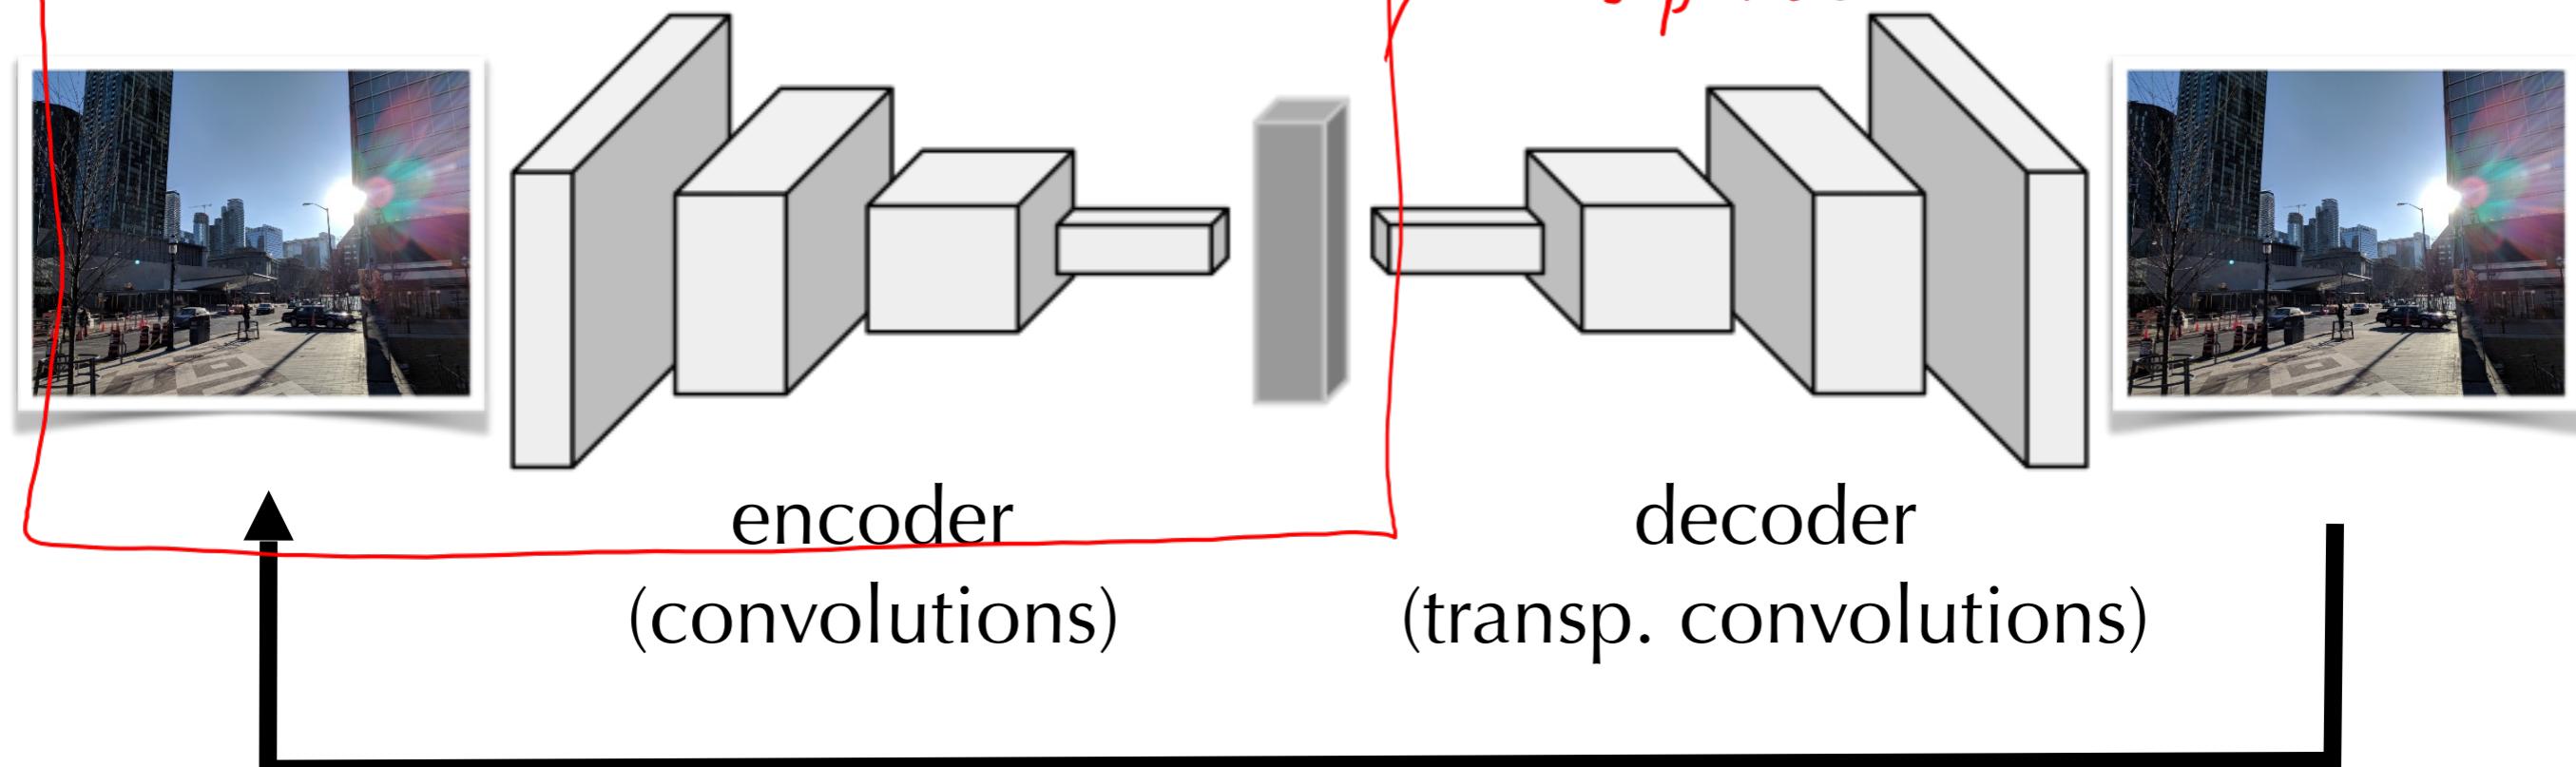
decoder  
(transp. convolutions)

# Autoencoder

in practice /  
at test time

latent  
space

use rather  
low-dimensional  
space

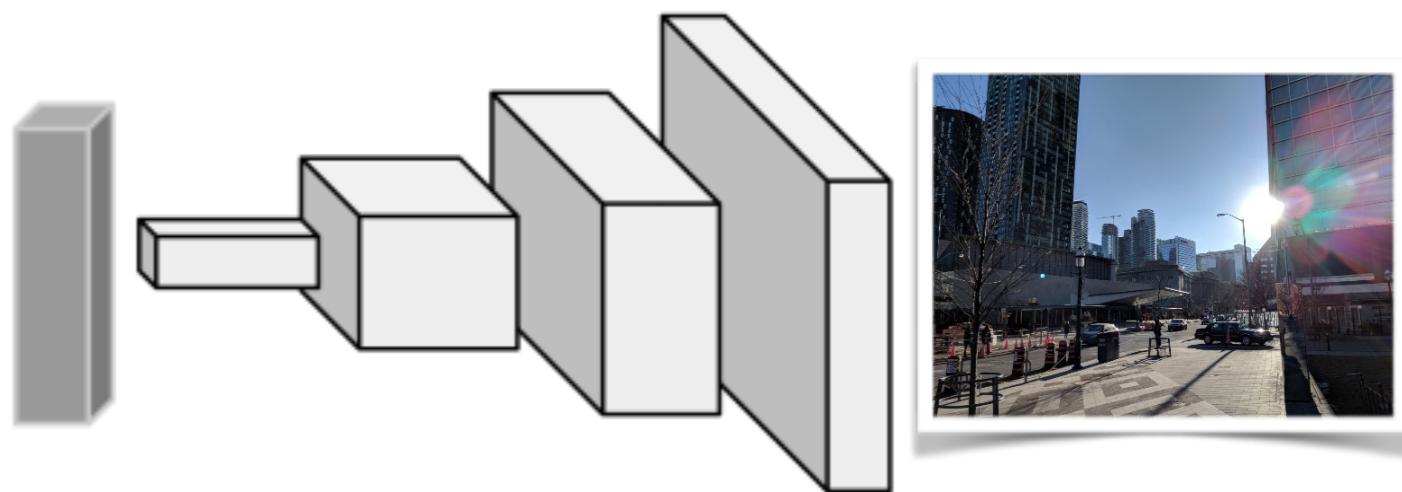


reconstruction loss

# Variational Autoencoder

sample point latent space

$$\mathbf{z} \sim p_{\text{model}}$$



decoder  
(transp. convolutions)

$$p(\mathbf{x}|\mathbf{z})$$

$$\mathbf{x} \sim p_{\text{data}}$$

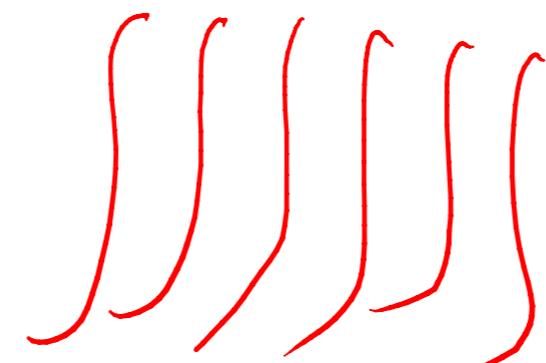
---

learn mapping from simple  
to complex distribution

# Theory

Want to maximize  
*training example*

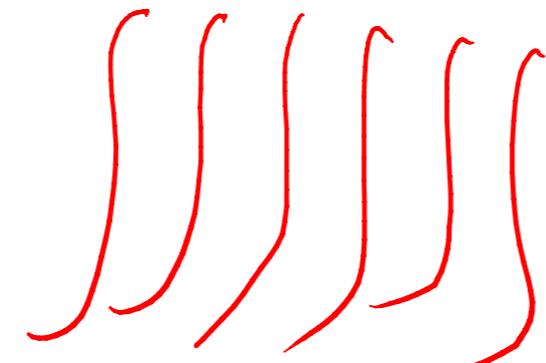
$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$



# Theory

Want to maximize  
*training example*

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

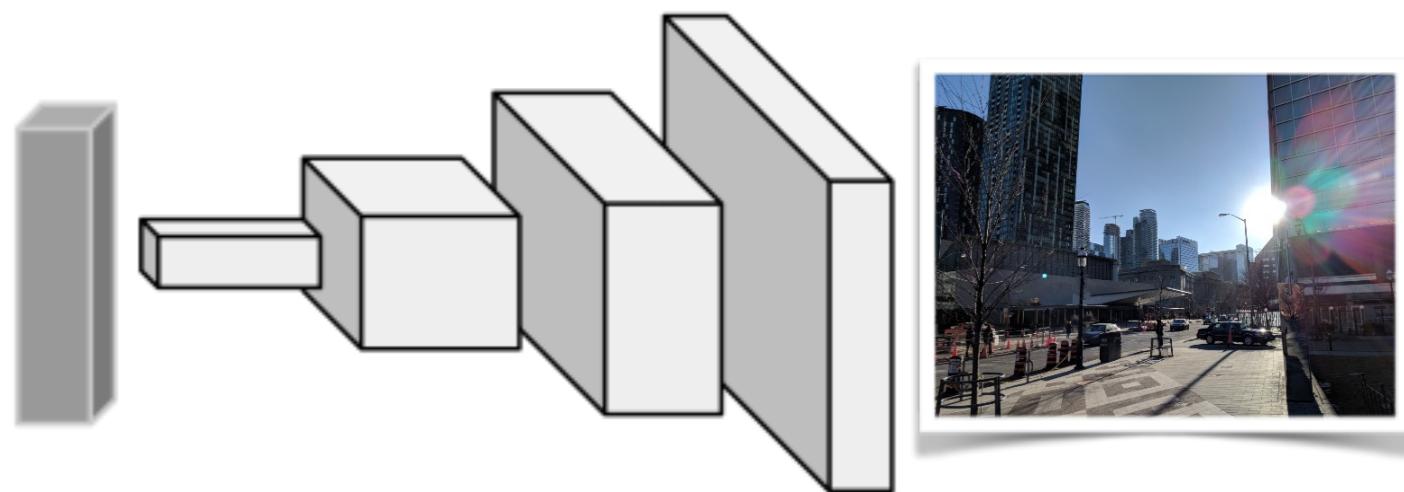


Intractable due to high dimensionality of latent space

# Variational Autoencoder

latent space

$$\mathbf{z} \sim p_{\text{model}}$$



decoder  
(transp. convolutions)

$$p(\mathbf{x}|\mathbf{z})$$

$$\mathbf{x} \sim p_{\text{data}}$$

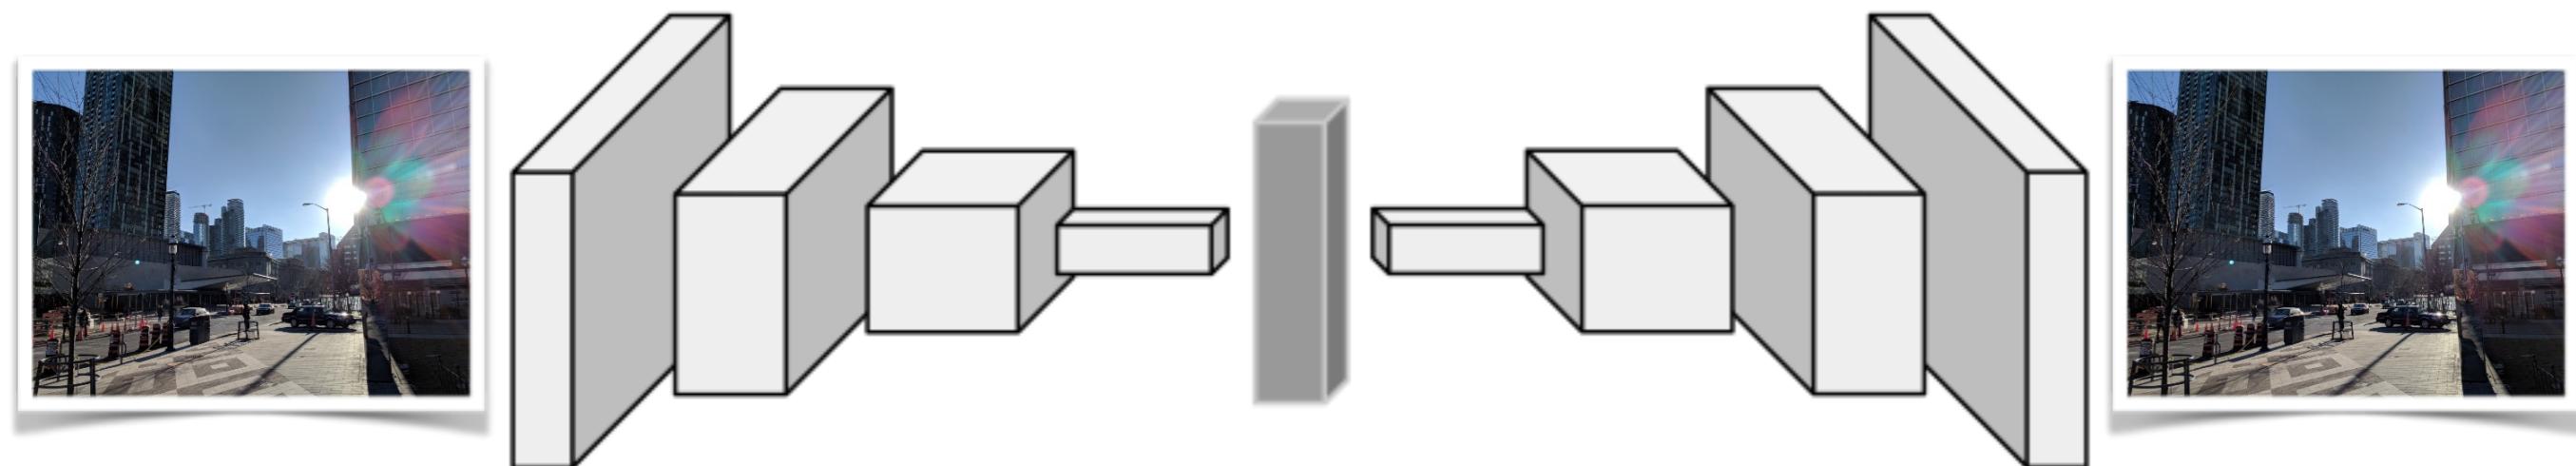
---

learn mapping from simple  
to complex distribution

# Variational Autoencoder

latent space

$$\mathbf{z} \sim p_{\text{model}}$$



encoder  
(convolutions)

decoder  
(transp. convolutions)

$$p(\mathbf{x}|\mathbf{z})$$

$$\mathbf{x} \sim p_{\text{data}}$$

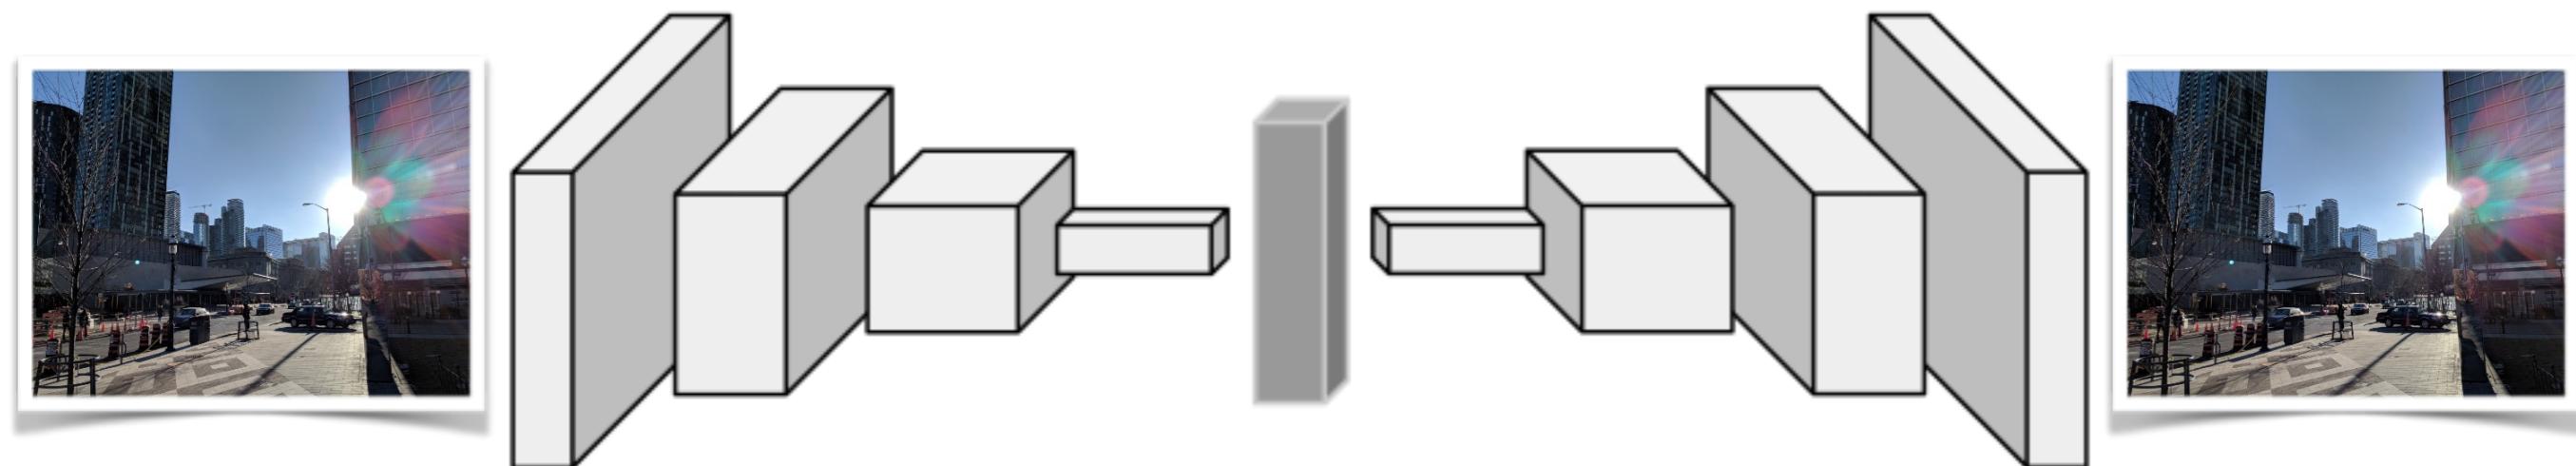
---

learn mapping from simple  
to complex distribution

# Variational Autoencoder

latent space

$$\mathbf{z} \sim p_{\text{model}}$$



encoder  
(convolutions)

$$\mathbf{x} \sim p_{\text{data}}$$

decoder  
(transp. convolutions)

$$p(\mathbf{x}|\mathbf{z})$$

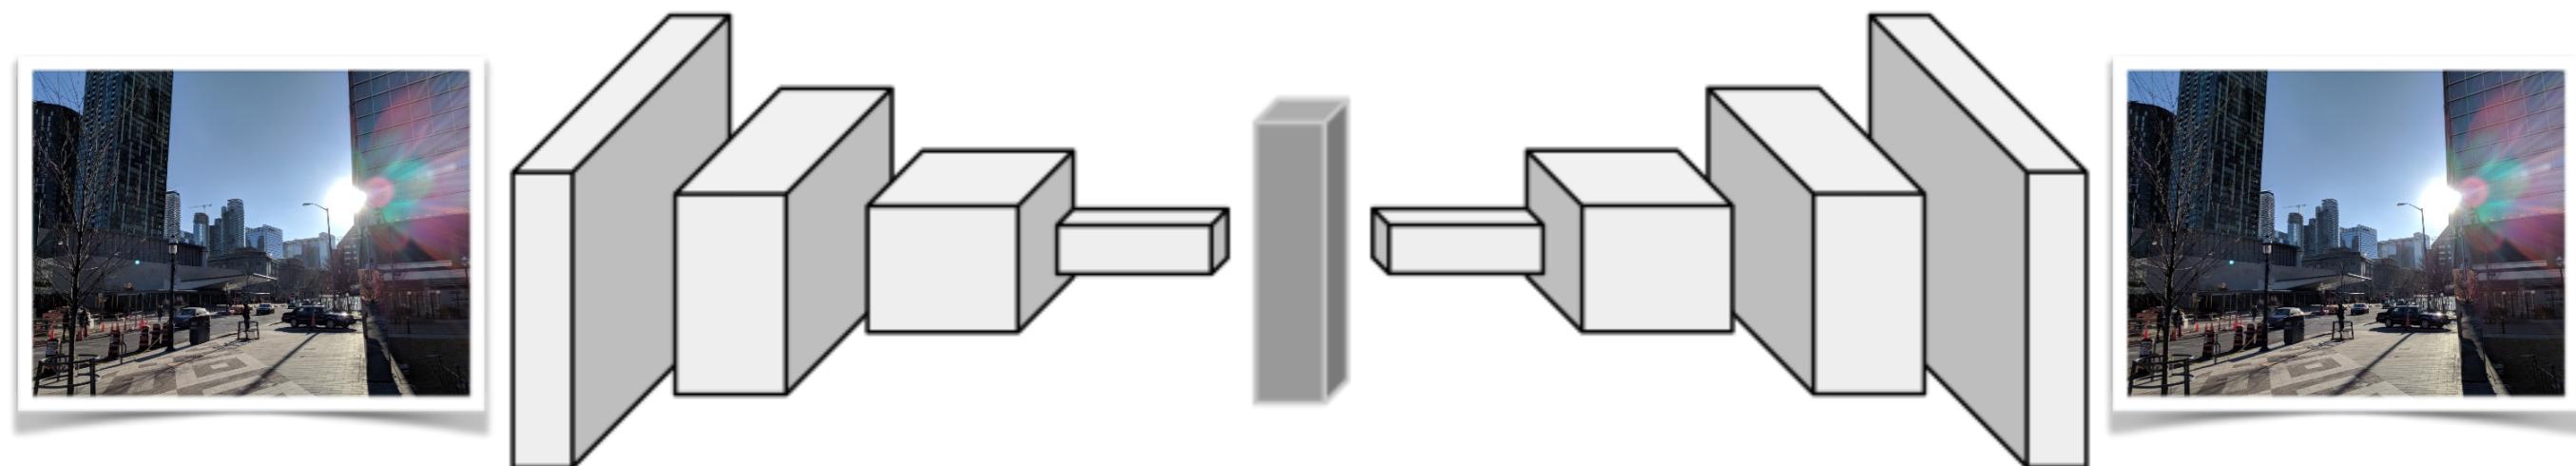
$$\mathbf{x} \sim p_{\text{data}}$$

learn mapping from simple  
to complex distribution

# Variational Autoencoder

latent space

$$\mathbf{z} \sim p_{\text{model}}$$



encoder  
(convolutions)

$$\mathbf{x} \sim p_{\text{data}}$$

$$q(\mathbf{z}|\mathbf{x})$$

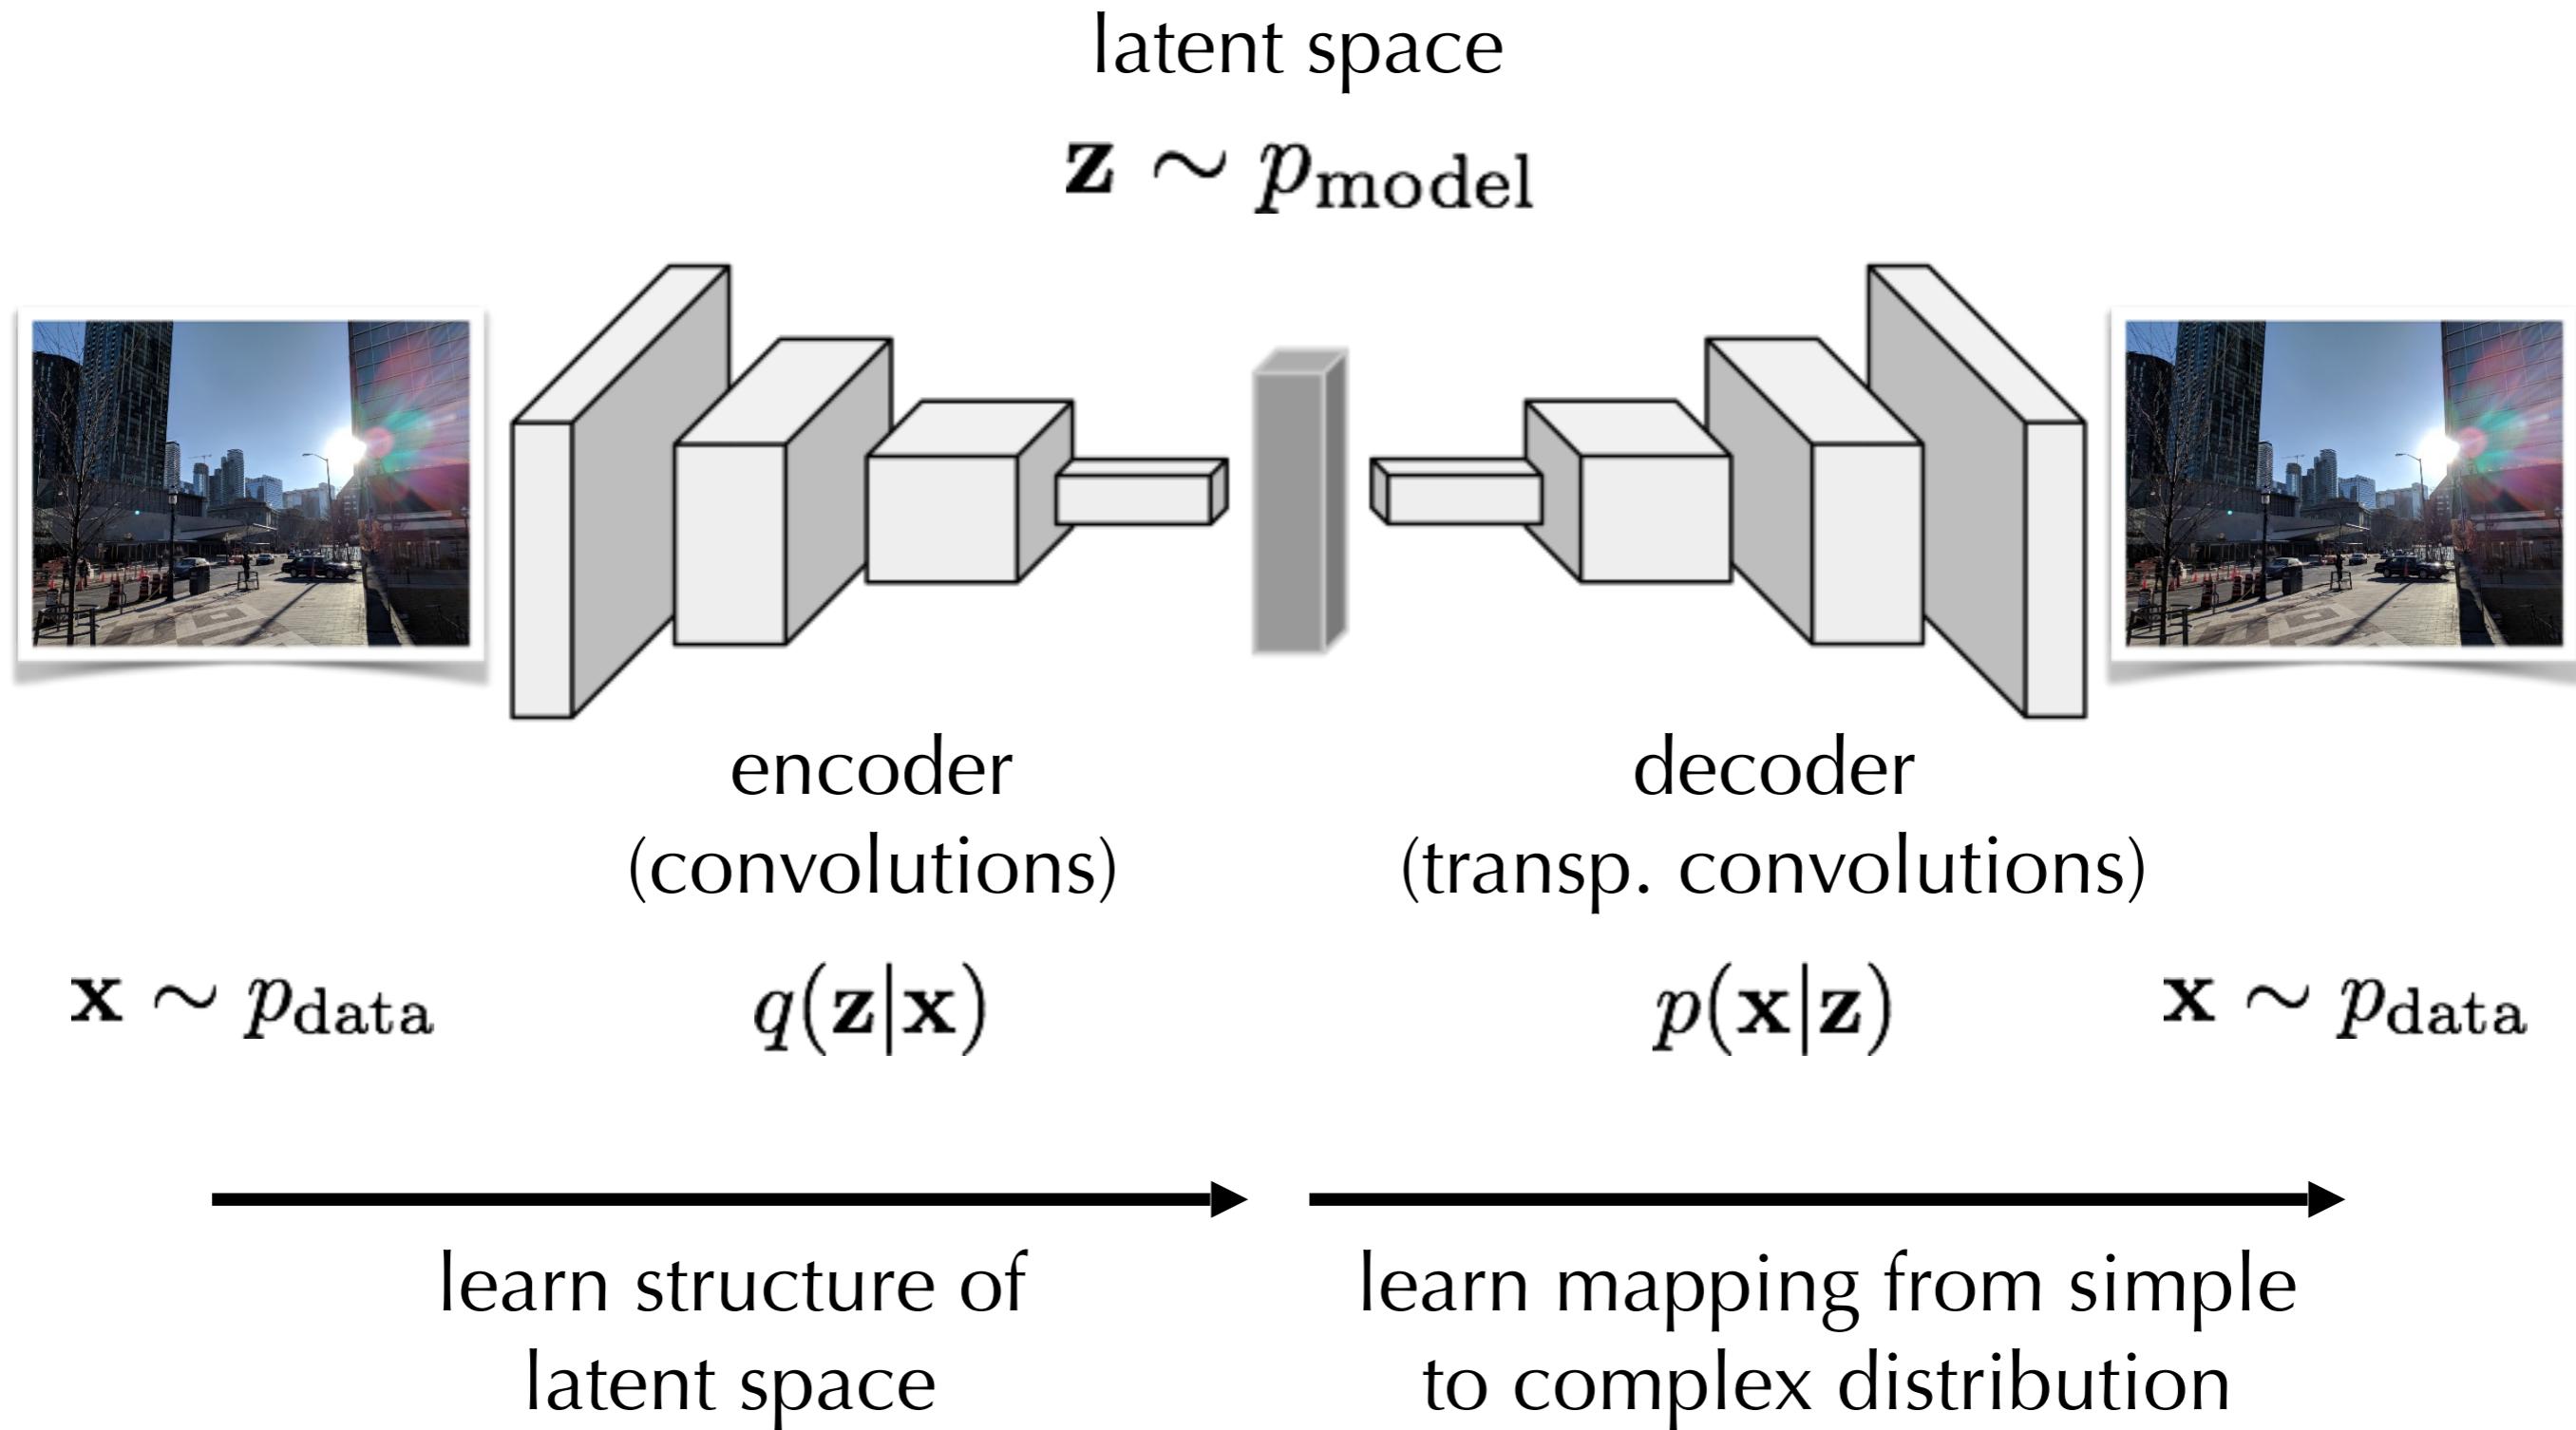
decoder  
(transp. convolutions)

$$p(\mathbf{x}|\mathbf{z})$$

$$\mathbf{x} \sim p_{\text{data}}$$

learn mapping from simple  
to complex distribution

# Variational Autoencoder



# Theory

Use encoder to generate probable samples, simplifying

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

to

$$p(\mathbf{x}) = E_{\mathbf{z} \sim q} p(\mathbf{x}|\mathbf{z}) \cancel{d\mathbf{z}}$$

# Theory

Use encoder to generate probable samples, simplifying

$$p(\mathbf{x}) = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

to

$$p(\mathbf{x}) = E_{\mathbf{z} \sim q} p(\mathbf{x}|\mathbf{z}) \cancel{d\mathbf{z}}$$

Problem: Learned distribution  $q$  might not follow  $p_{\text{model}}$

# KL Divergence

**Kullback-Leibler divergence:** ``distance'' between probability distributions

$$D_{\text{KL}}(P||Q) = \int P(x) \log \left( \frac{P(x)}{Q(x)} \right) dx \geq 0$$

$\geq 0 \Leftrightarrow P(x) \in Q(x)$   
 $\forall x$

$$D_{\text{KL}}(P||Q) \neq D_{\text{KL}}(Q||P)$$

# KL Divergence

**Kullback-Leibler divergence:** ``distance'' between probability distributions

$$D_{\text{KL}}(P||Q) = \int P(x) \log \left( \frac{P(x)}{Q(x)} \right) dx \geq 0$$

$\geq 0 \Leftrightarrow P(x) \in Q(x)$   
 $\forall x$

Equivalent to (when sampling):

$$D_{\text{KL}}(P||Q) = E_{x \sim P} [\log(P(x)) - \log(Q(x))]$$

$$D_{\text{KL}}(P||Q) \neq D_{\text{KL}}(Q||P)$$

# KL Divergence

**Kullback-Leibler divergence:** ``distance'' between probability distributions

$$D_{\text{KL}}(P||Q) = \int P(x) \log \left( \frac{P(x)}{Q(x)} \right) dx \geq 0$$

$\geq 0 \Leftrightarrow P(x) \in Q(x)$   
 $\forall x$

Equivalent to (when sampling):

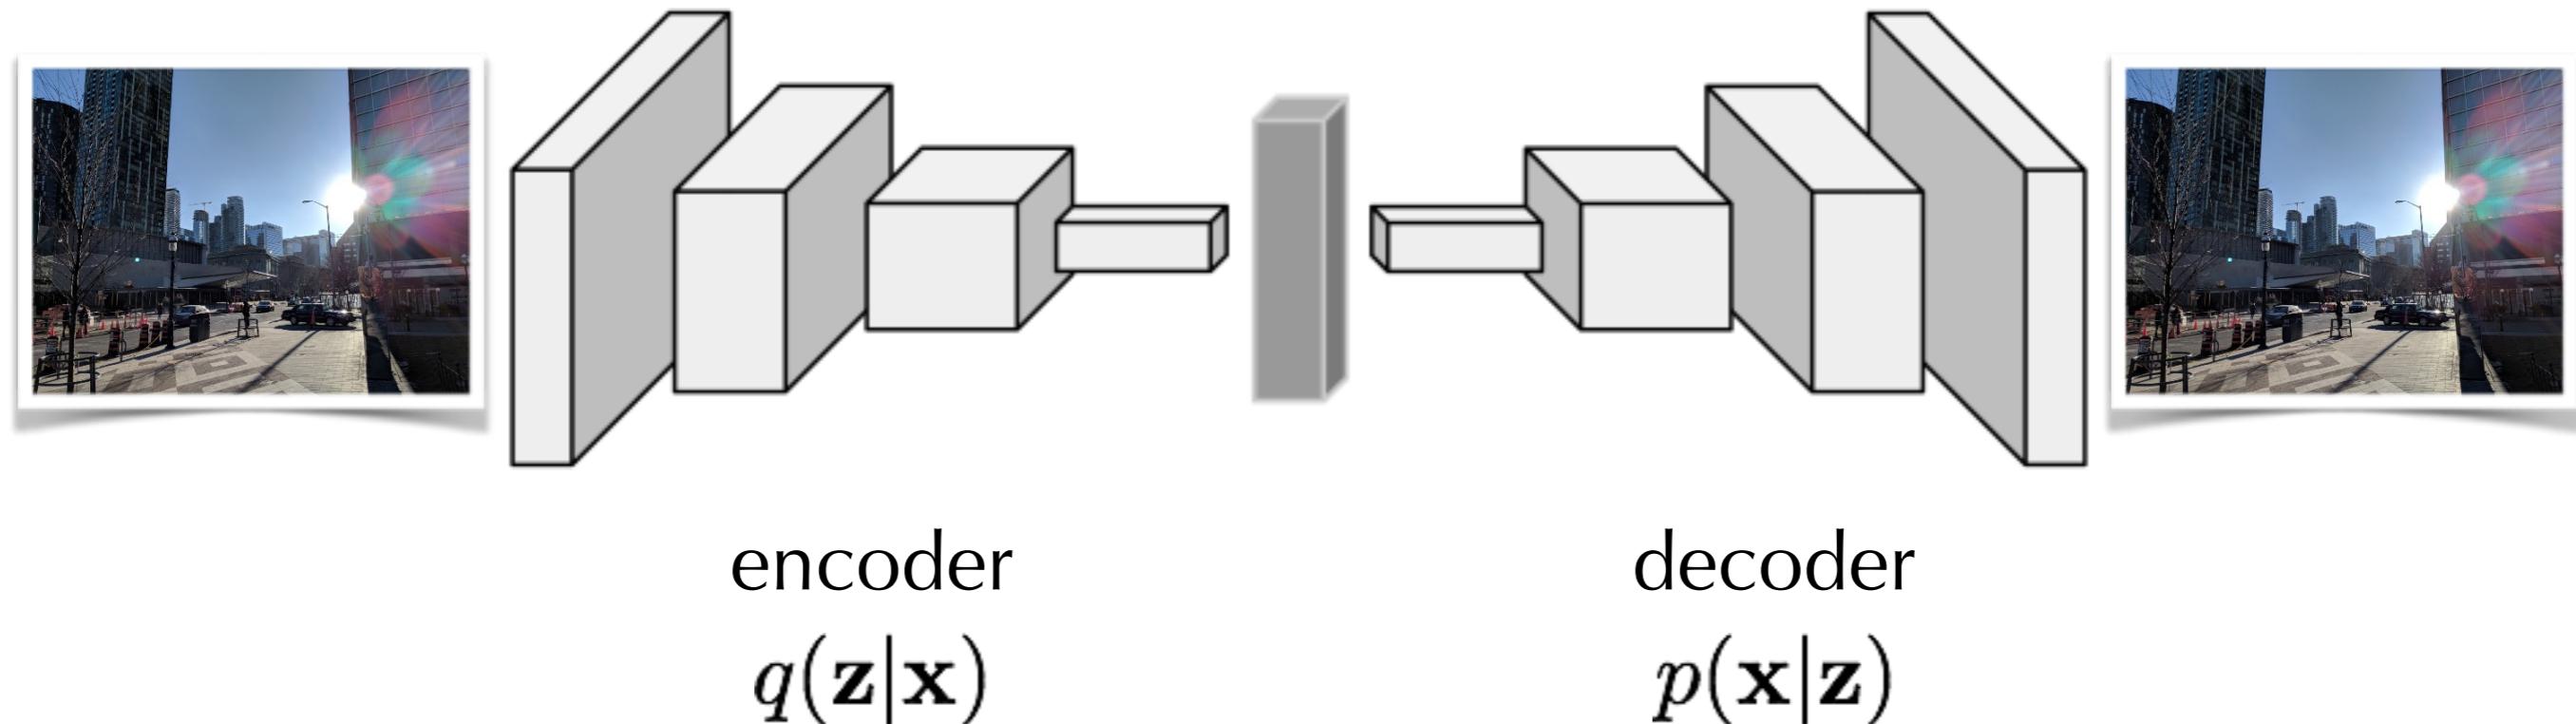
$$D_{\text{KL}}(P||Q) = E_{x \sim P} [\log(P(x)) - \log(Q(x))]$$

Notice that KL divergence is asymmetric

$$D_{\text{KL}}(P||Q) \neq D_{\text{KL}}(Q||P)$$

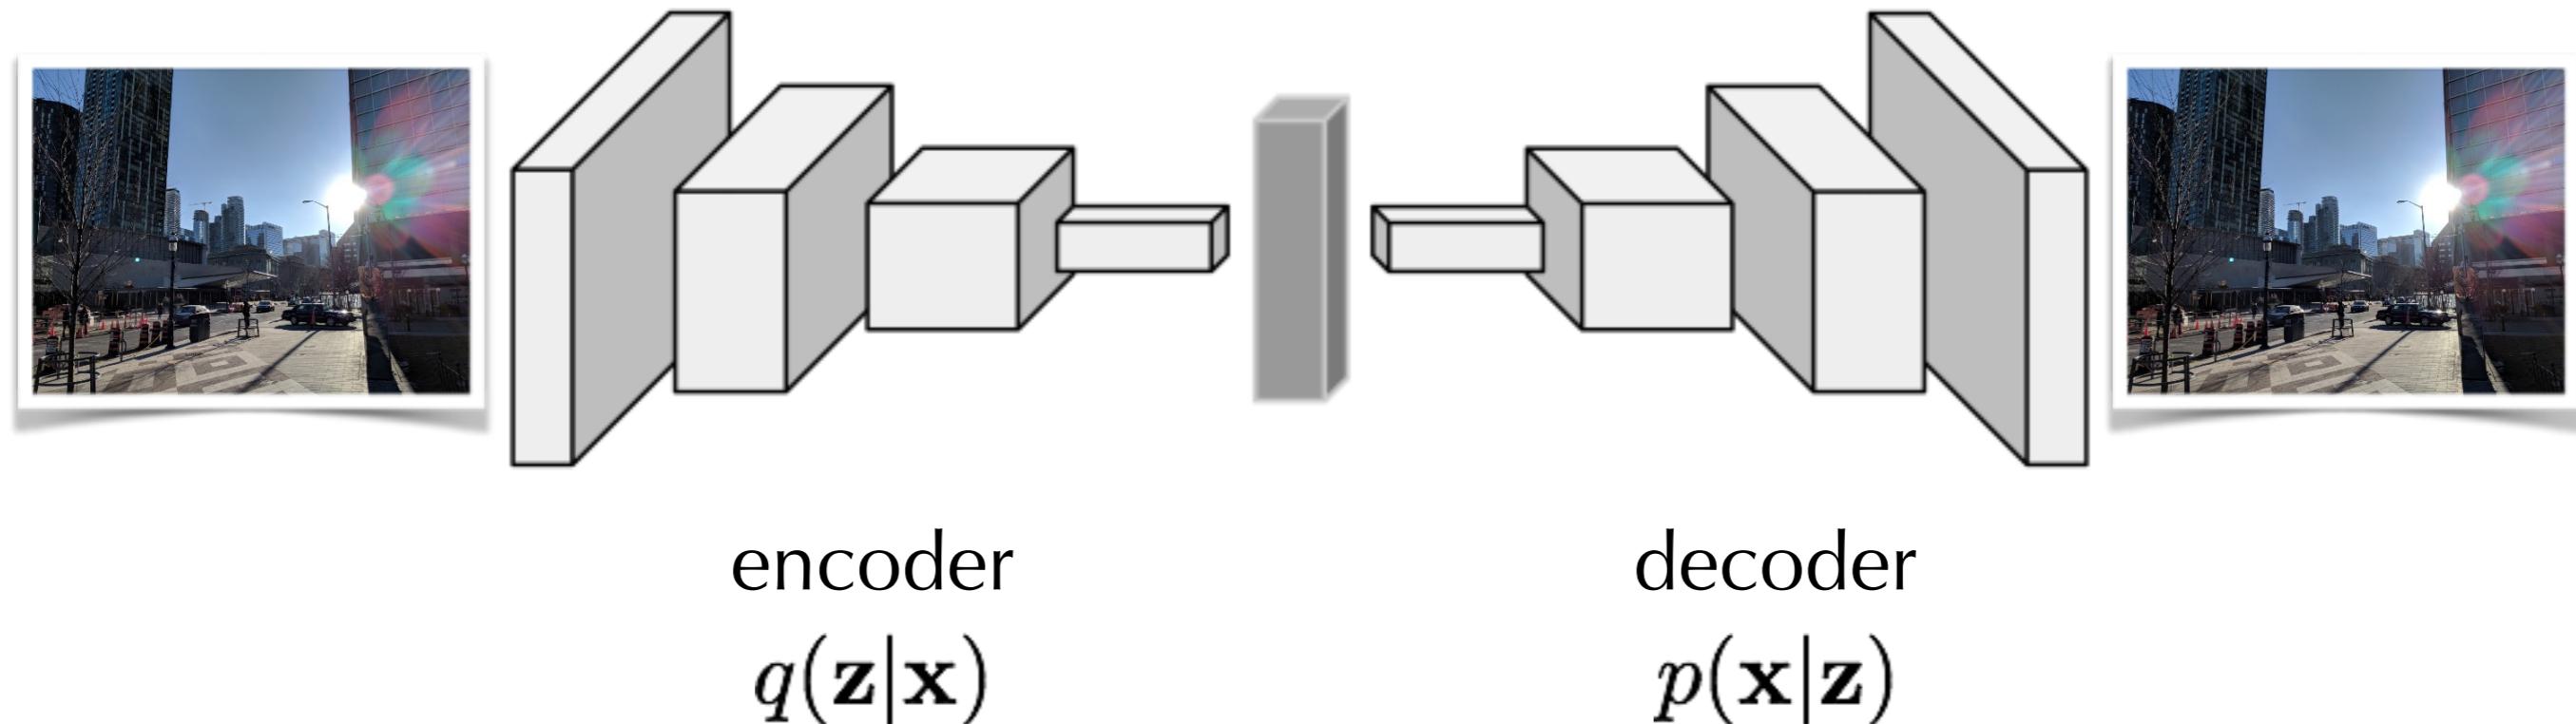
# Theory

$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$



# Theory

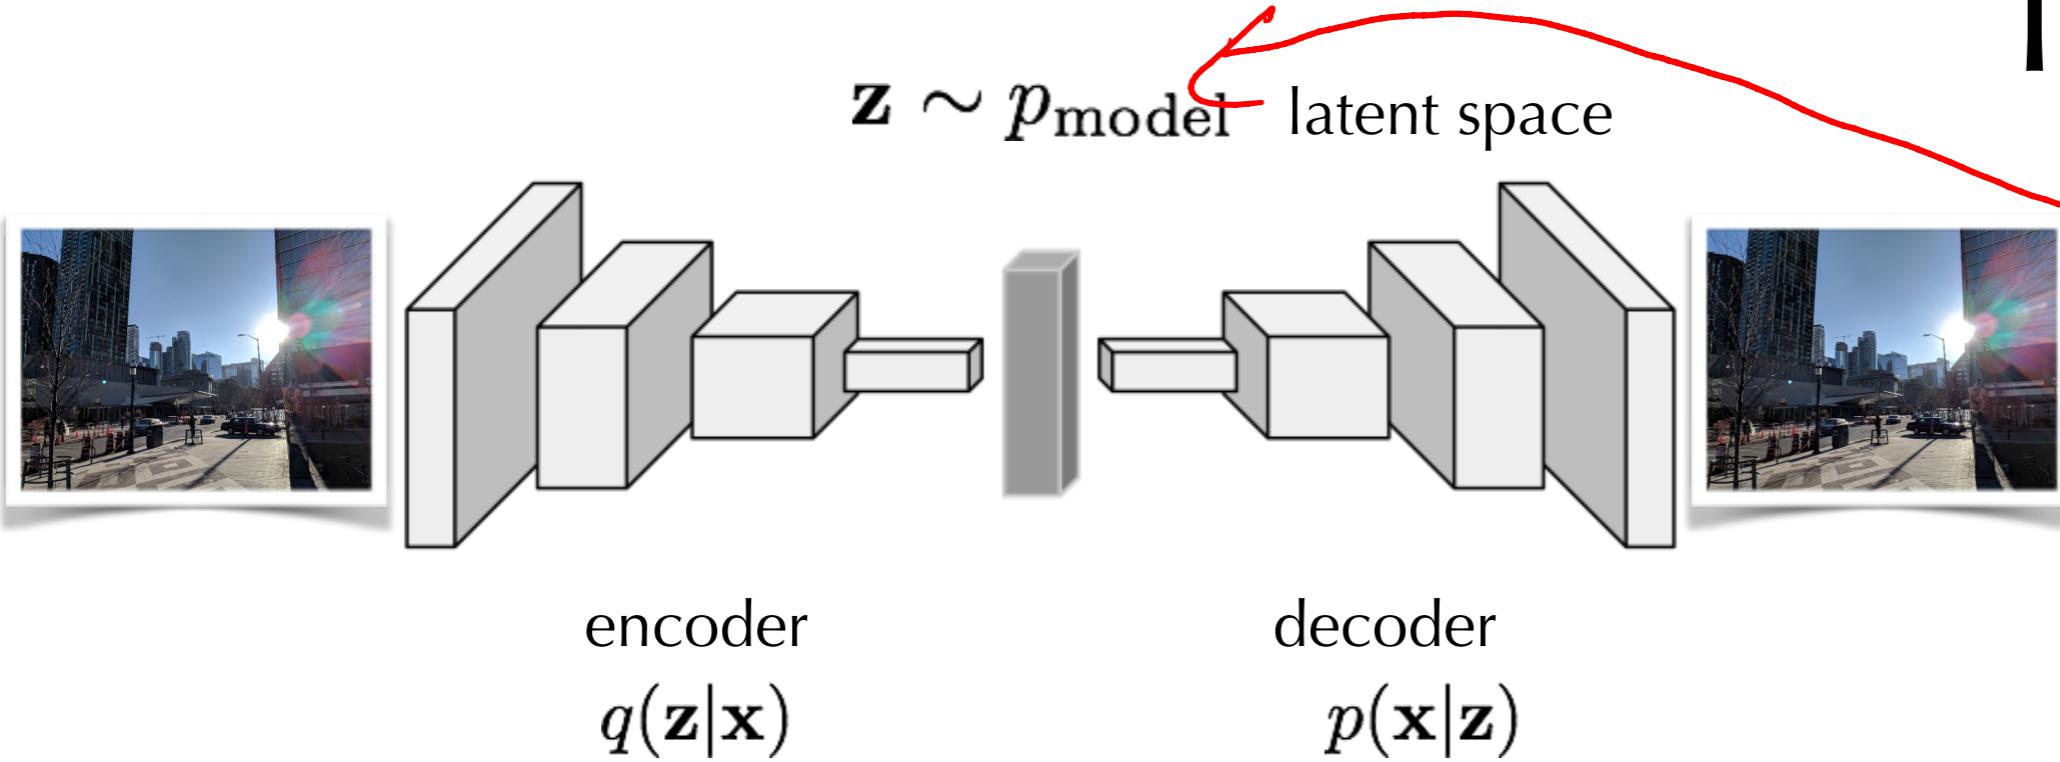
$\mathbf{z} \sim p_{\text{model}}$  latent space



Want to ensure that  $q(\mathbf{z}|\mathbf{x})$  and  $p(\mathbf{z}|\mathbf{x})$  are the same:

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}|\mathbf{x}))]$$

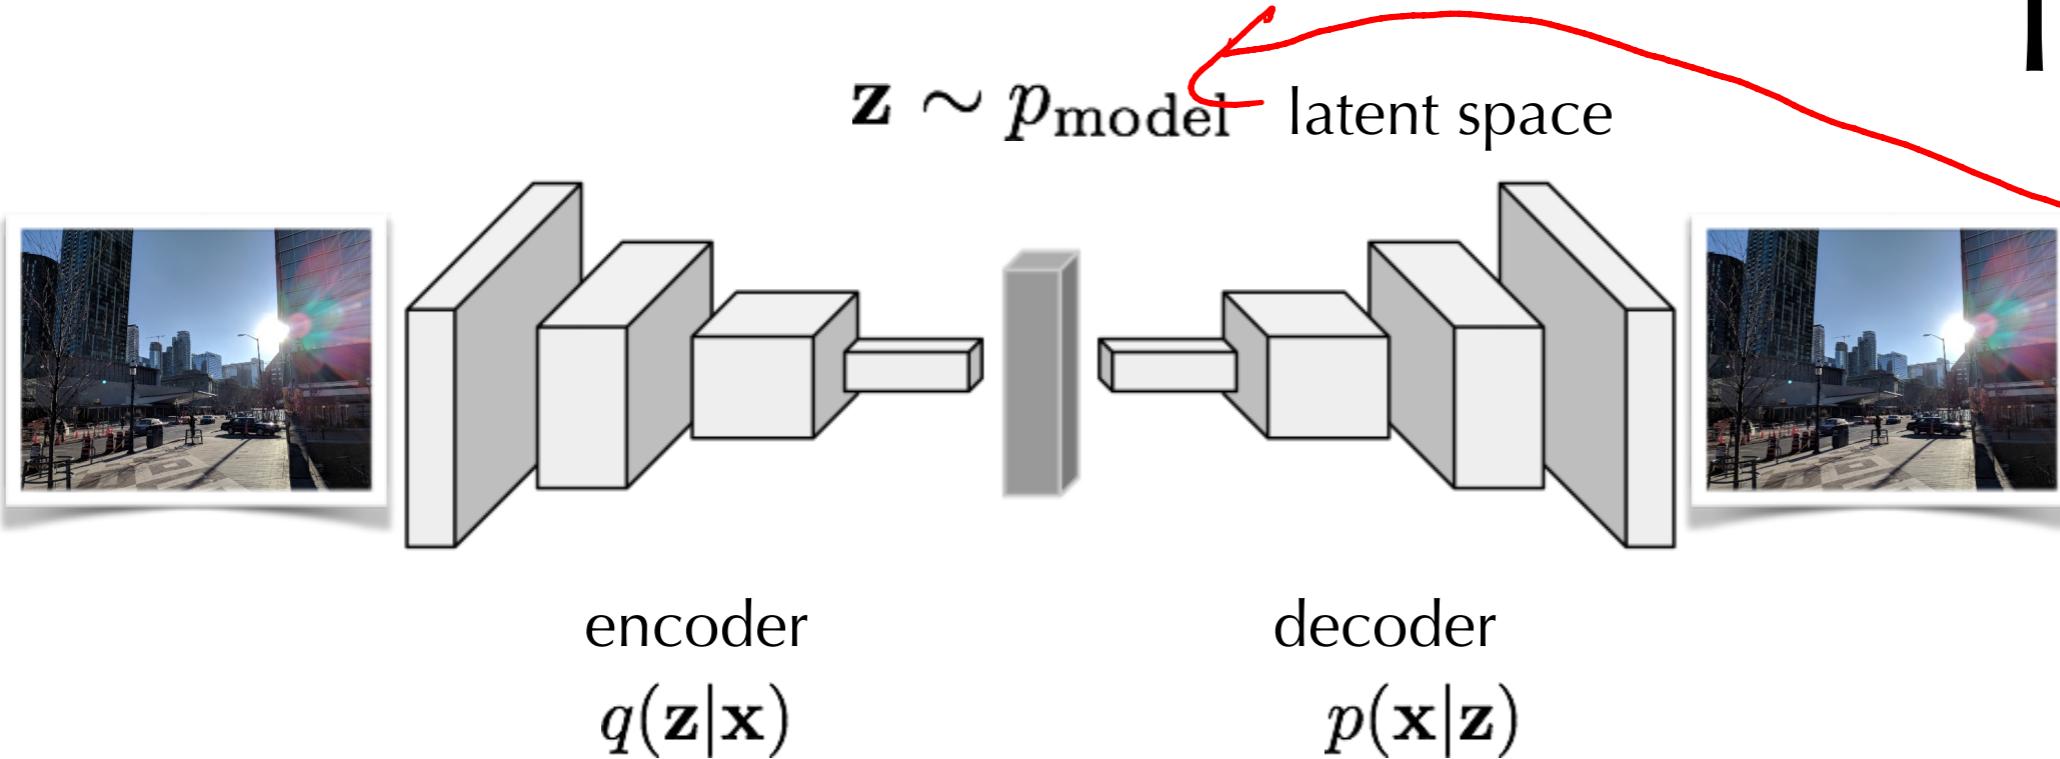
# Theory



Use Bayes rule to include  $p(\mathbf{x})$  and  $p(\mathbf{x}|\mathbf{z})$ :

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}|\mathbf{x}))]$$

# Theory



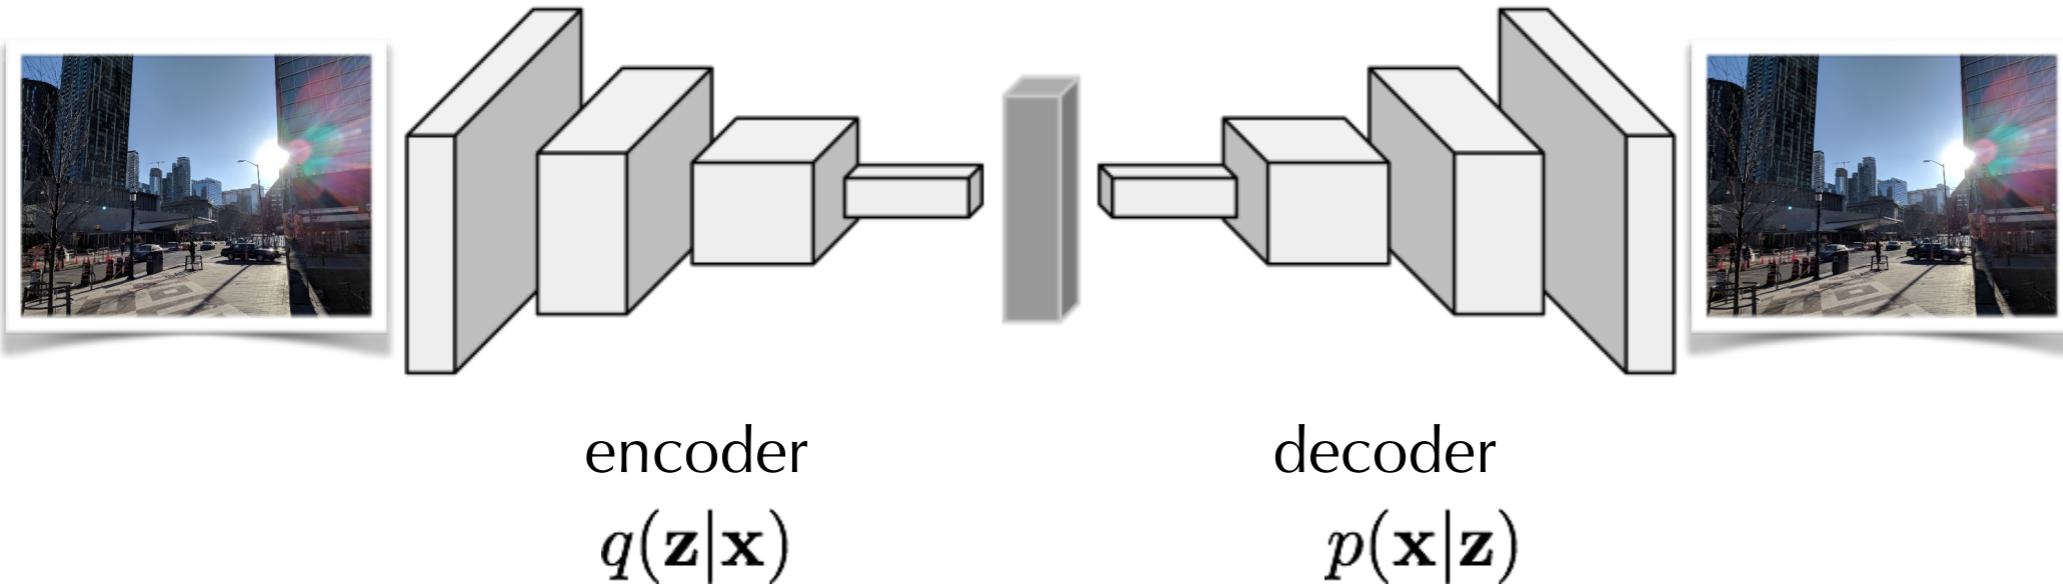
Use Bayes rule to include  $p(\mathbf{x})$  and  $p(\mathbf{x}|\mathbf{z})$ :

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}|\mathbf{x}))]$$

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log\left(\frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}\right)]$$

# Theory

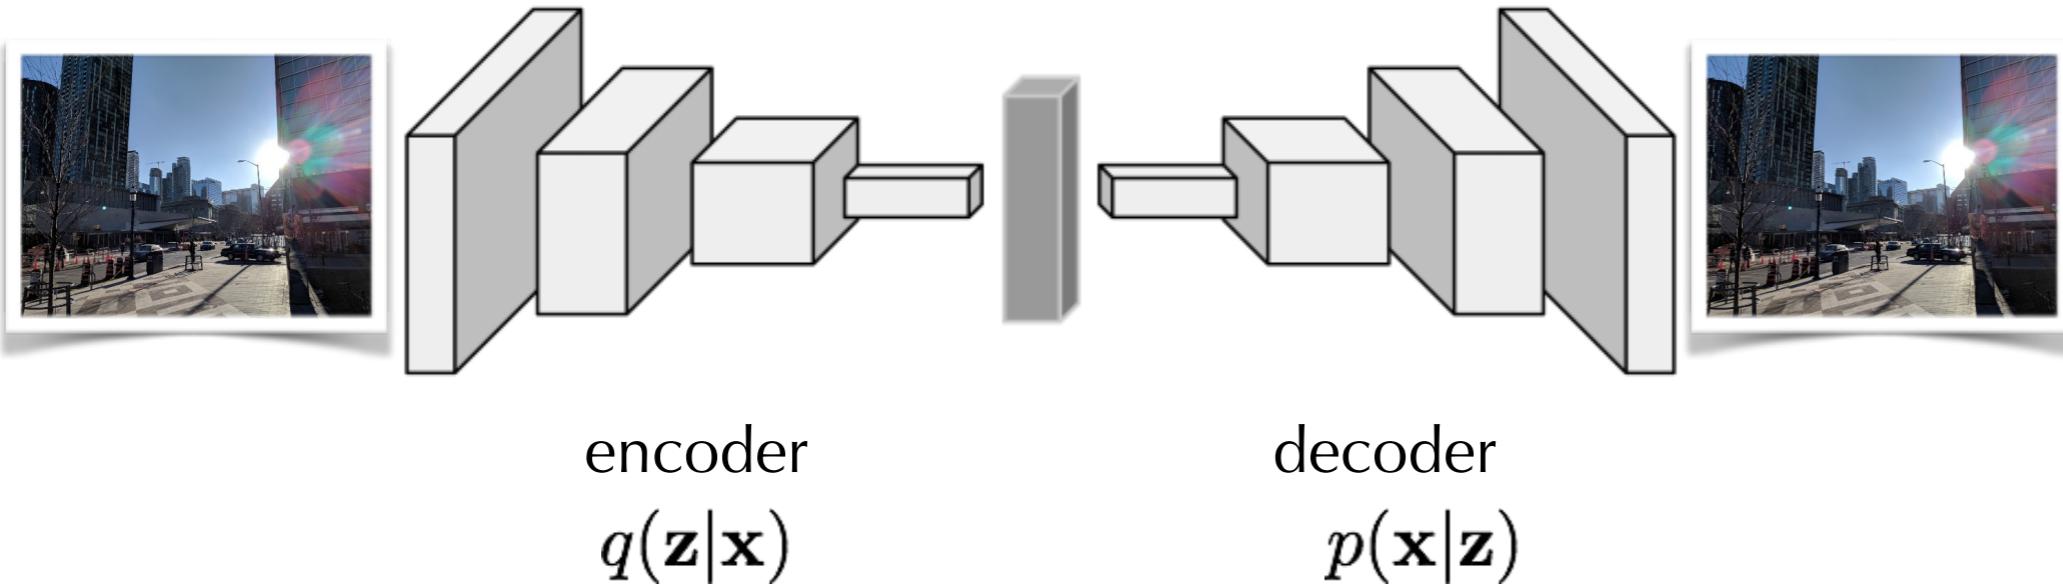
$\mathbf{z} \sim p_{\text{model}}$  latent space



$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log\left(\frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}\right)]$$

# Theory

$\mathbf{z} \sim p_{\text{model}}$  latent space

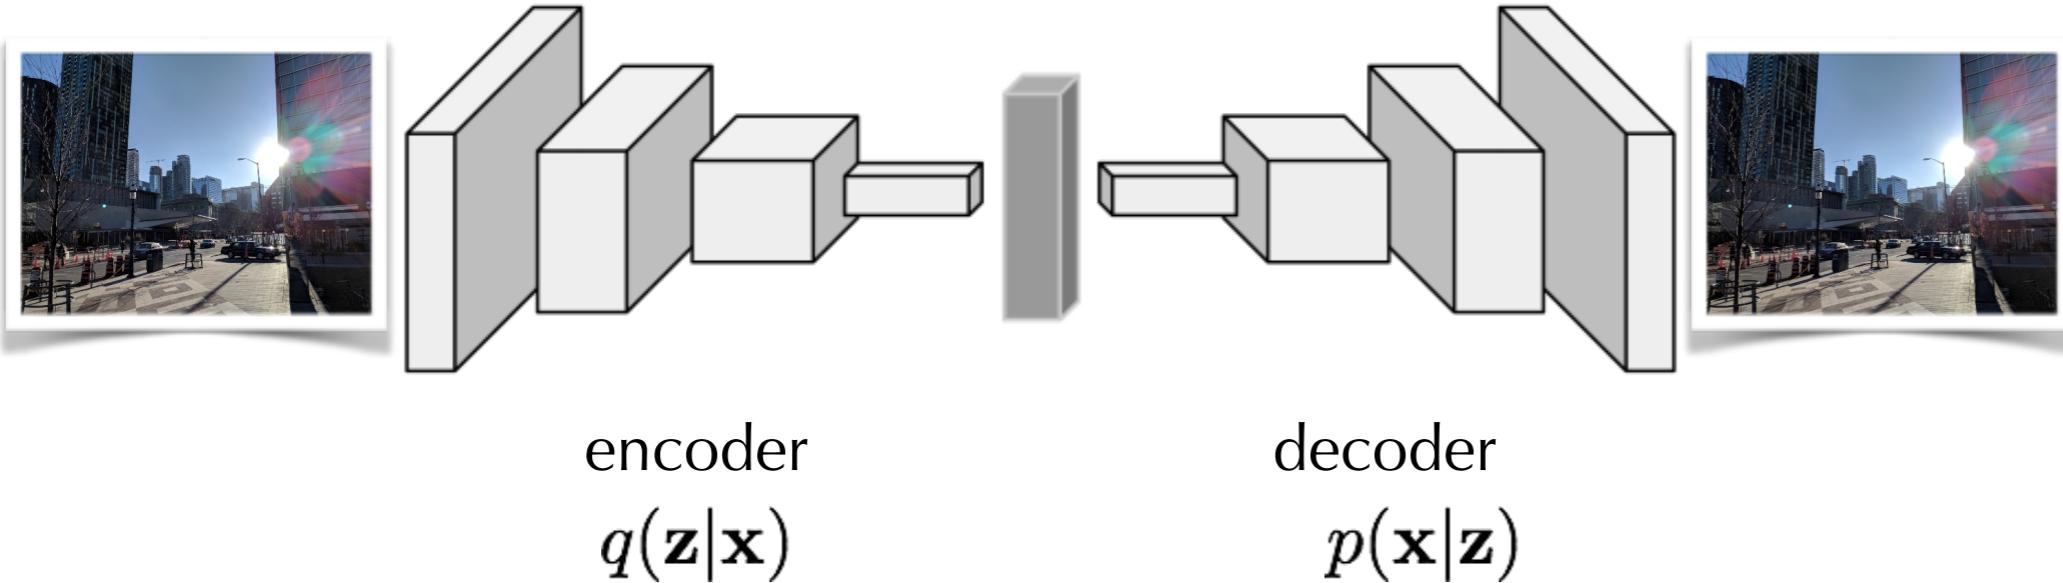


$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log\left(\frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}\right)]$$

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{x}|\mathbf{z})) - \log(p(\mathbf{z})) + \log(p(\mathbf{x}))]$$

# Theory

$\mathbf{z} \sim p_{\text{model}}$  latent space



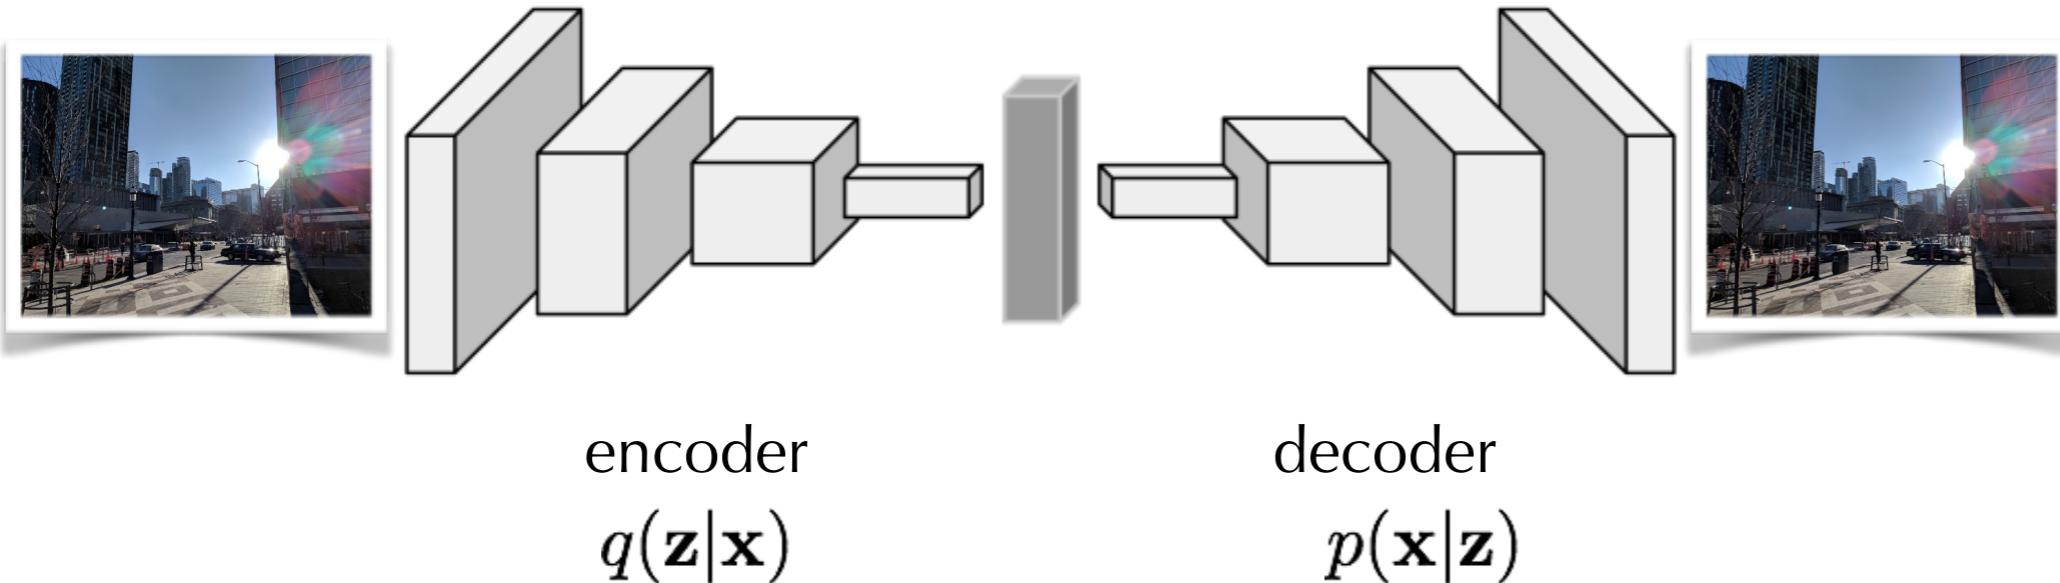
$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log\left(\frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}\right)]$$

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{x}|\mathbf{z})) - \log(p(\mathbf{z})) + \log(p(\mathbf{x}))]$$

$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

# Theory

$\mathbf{z} \sim p_{\text{model}}$  latent space



$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

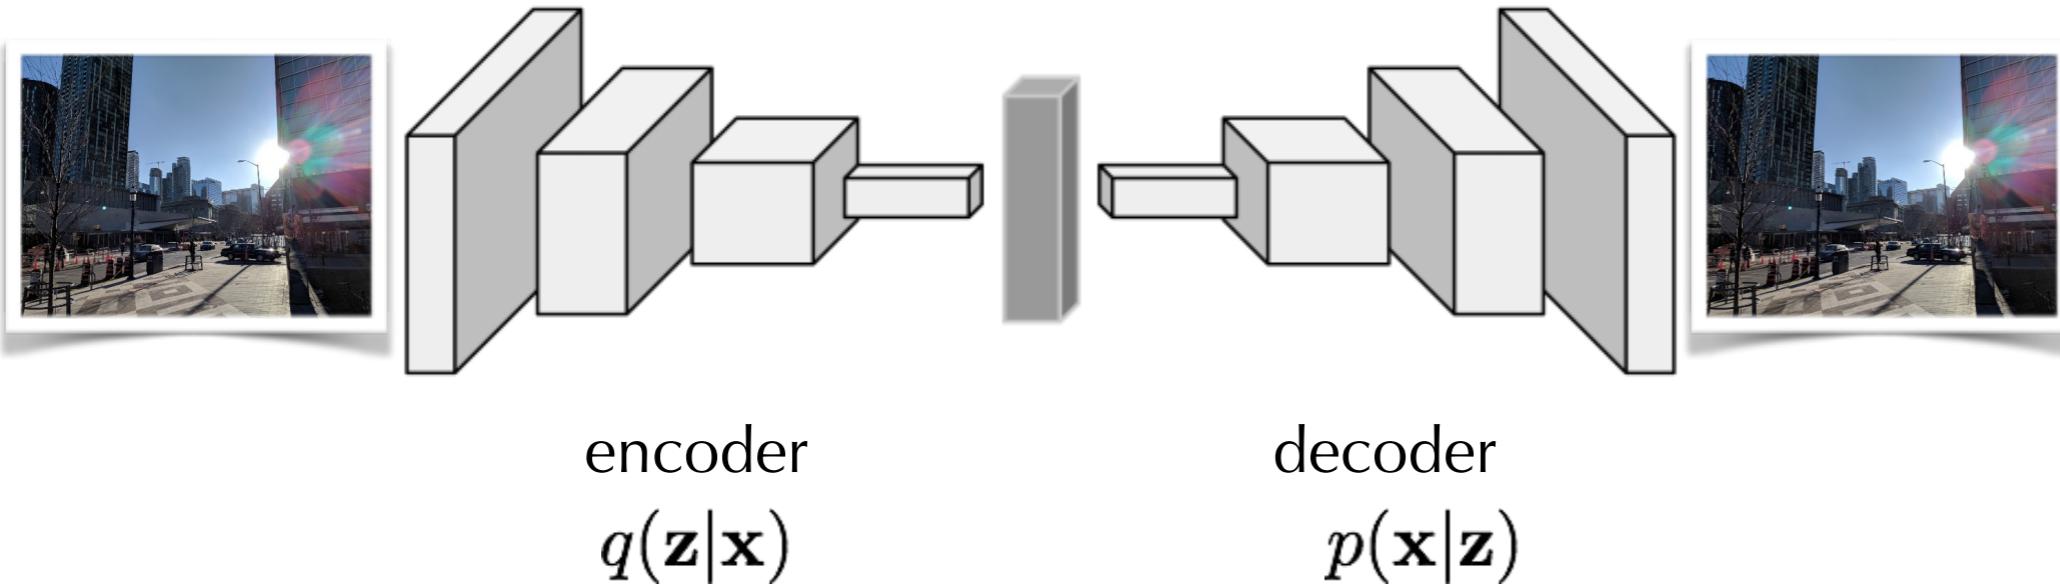
$, \geq 0$

no idea

decoder

# Theory

$\mathbf{z} \sim p_{\text{model}}$  latent space

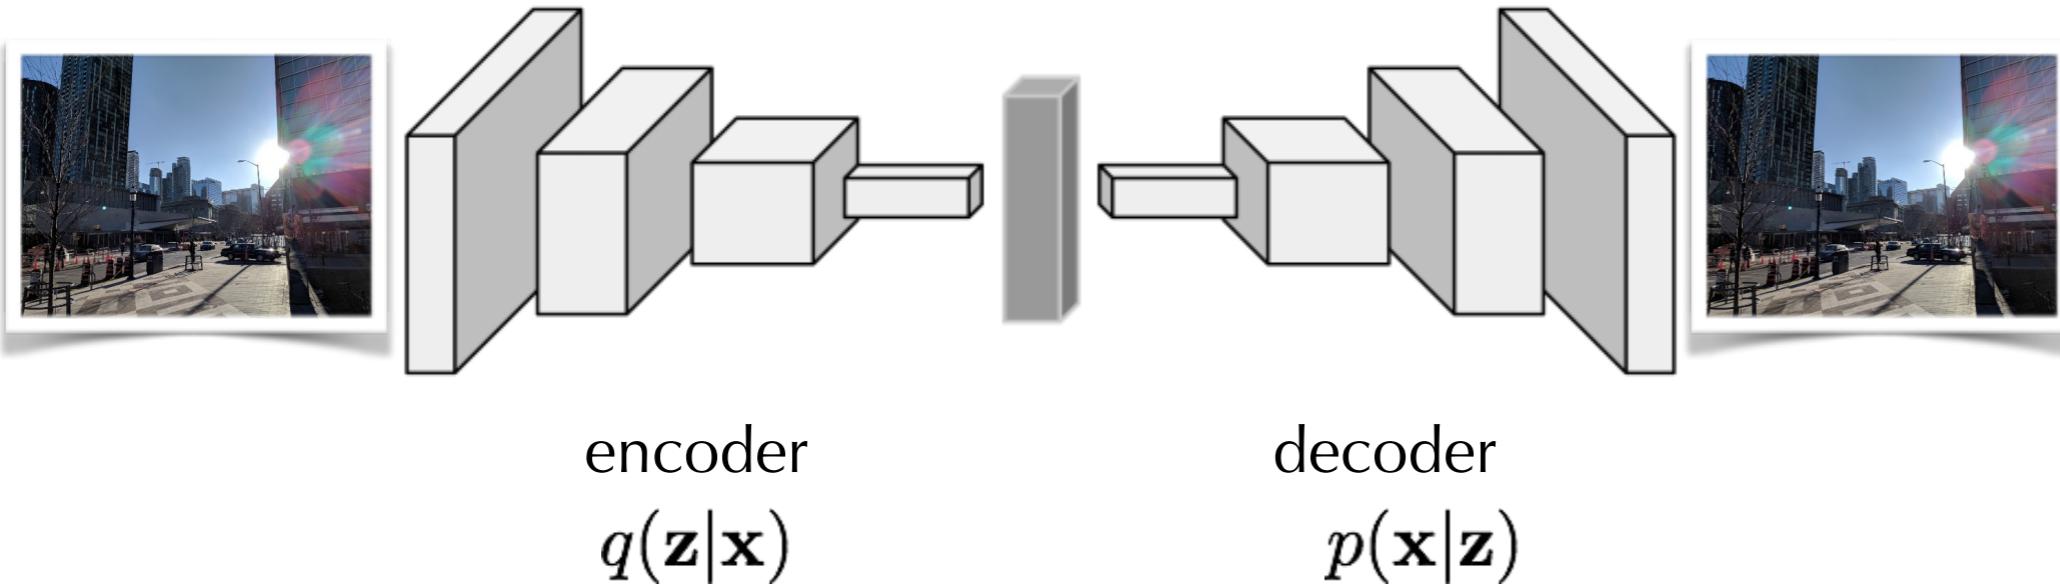


$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

$$\log(p(\mathbf{x})) = \underbrace{D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x}))}_{\text{no idea}} + \underbrace{E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))]}_{\text{decoder!}} - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

# Theory

$\mathbf{z} \sim p_{\text{model}}$  latent space



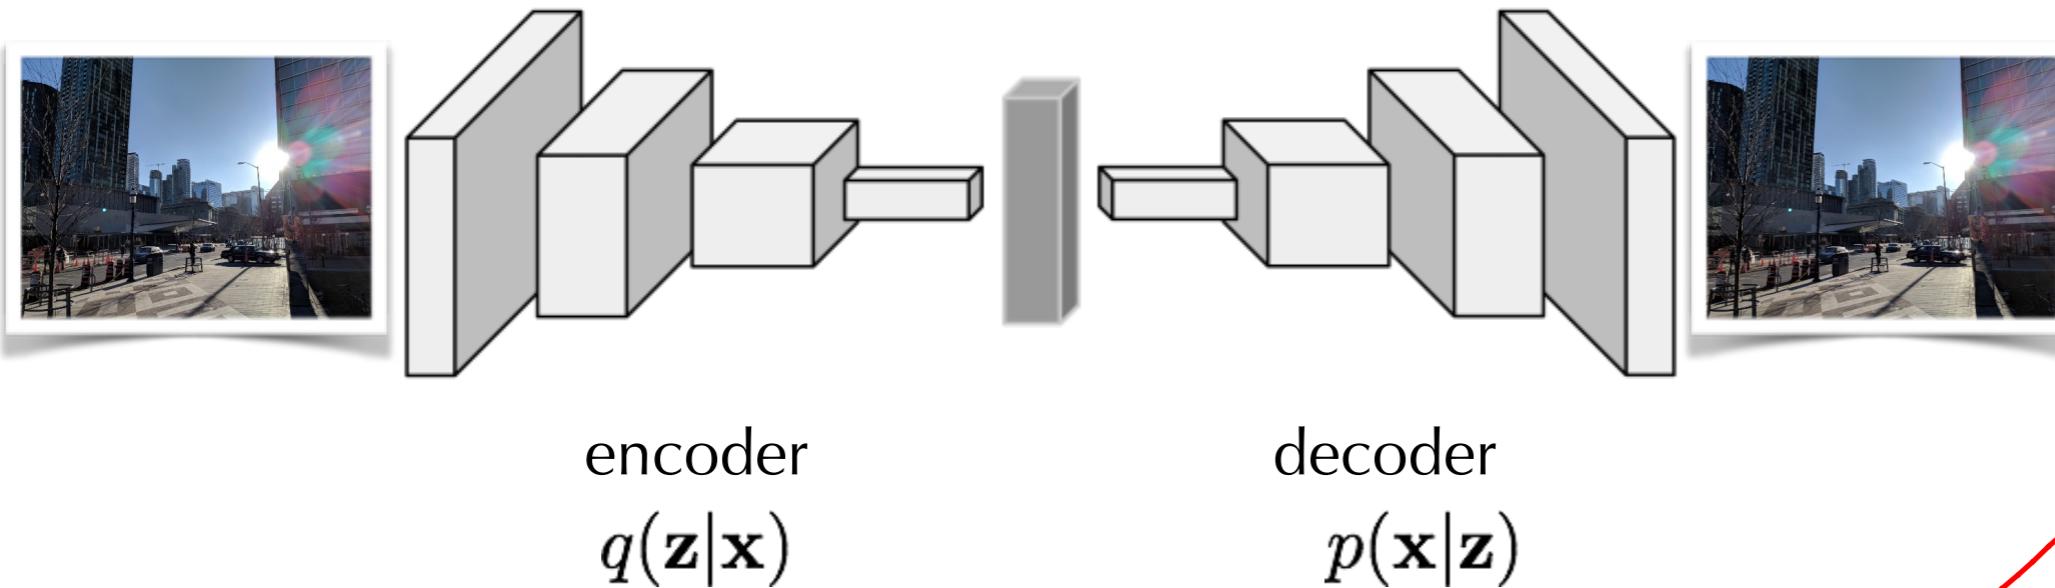
$$D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) = -E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] + E_{\mathbf{z} \sim q}[\log(q(\mathbf{z}|\mathbf{x})) - \log(p(\mathbf{z}))] + \log(p(\mathbf{x}))$$

$$\log(p(\mathbf{x})) = \underbrace{D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x}))}_{\text{no idea}} + \underbrace{E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))]}_{\text{decoder!}} - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

$$\geq E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$$

# Objective Function

$\mathbf{z} \sim p_{\text{model}}$  latent space

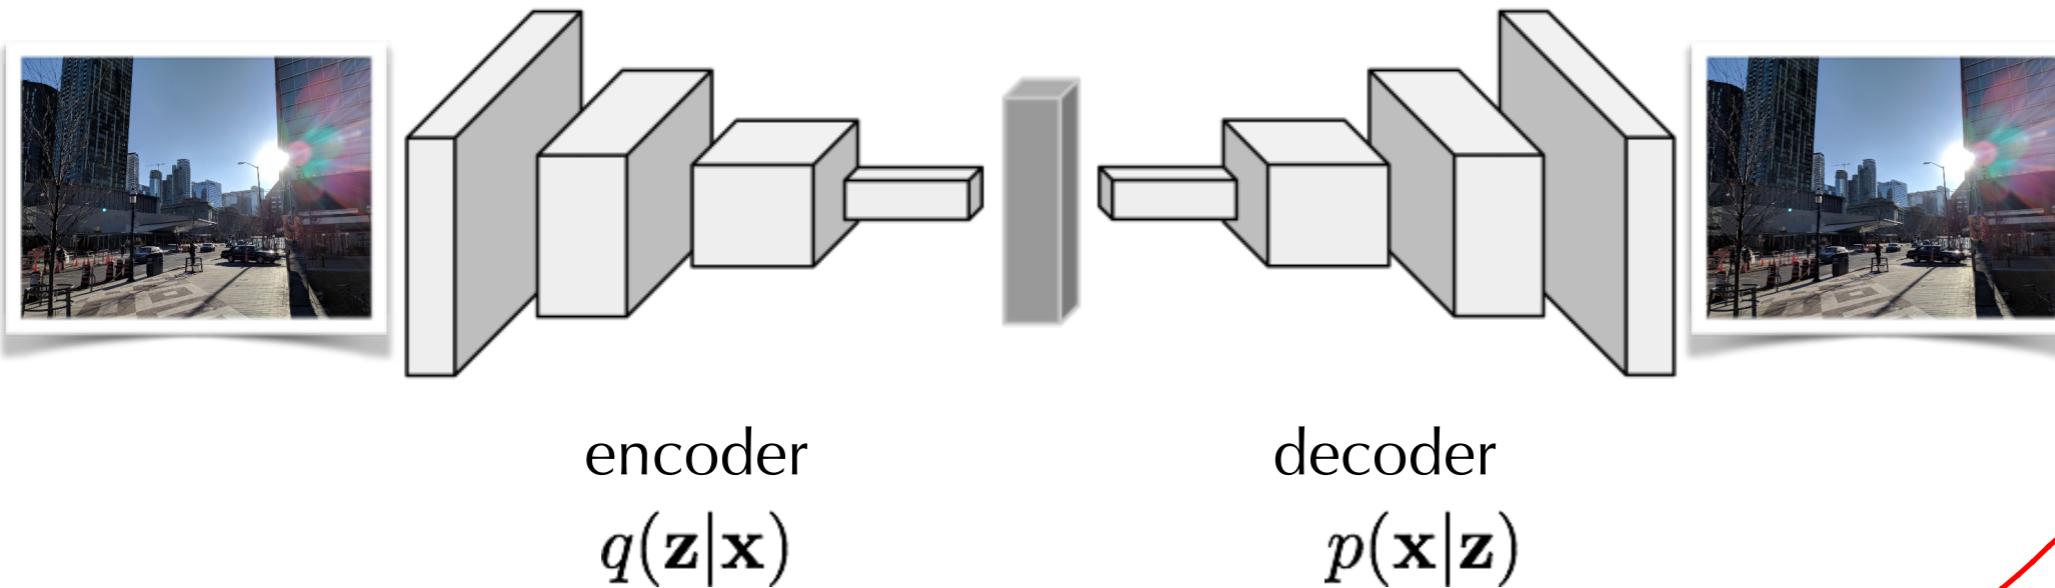


Maximize

$$E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$

# Objective Function

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



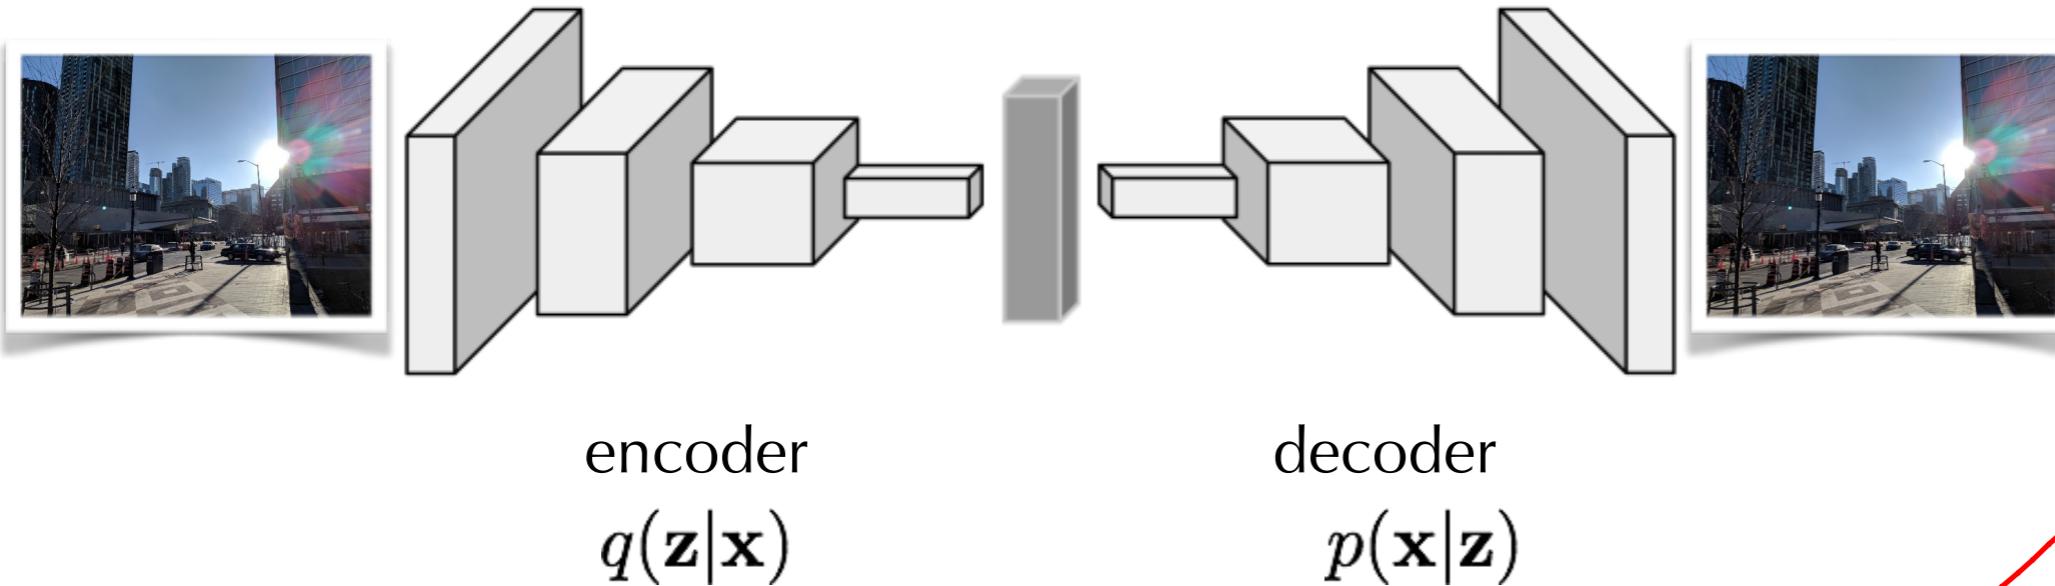
Maximize

$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$

reconstruction ``loss'',  
e.g., L2 norm (normal distribution)

# Objective Function

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



Maximize

$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$

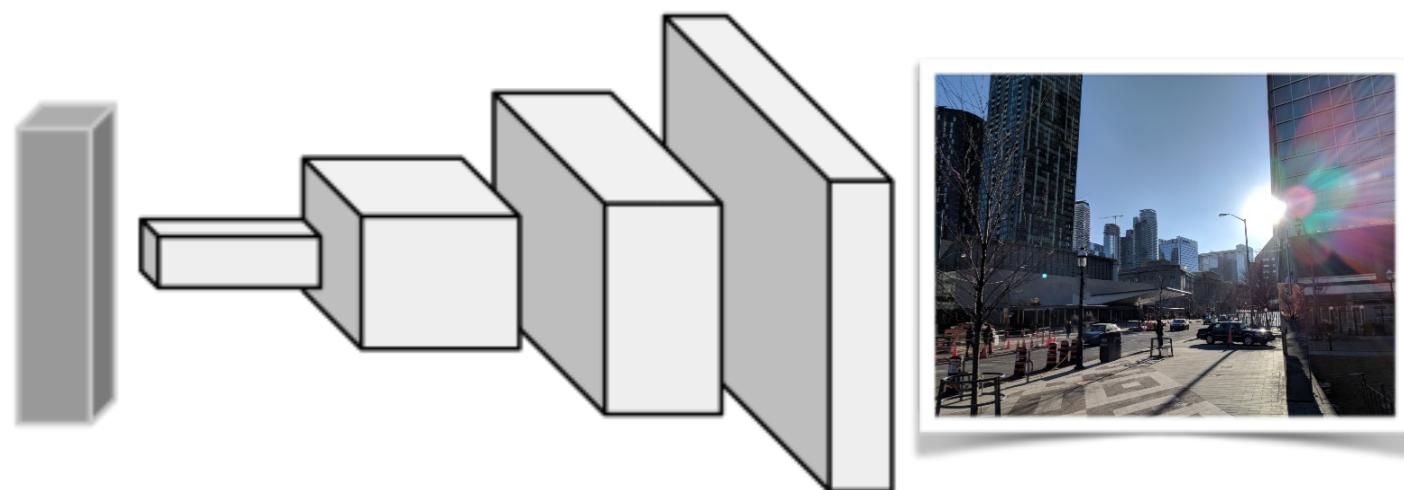
reconstruction ``loss'',  
e.g., L2 norm (normal distribution)

similarity between learned  
and sampling distribution

# At Test Time

sample point latent space

$$\mathbf{z} \sim p_{\text{model}}$$



decoder  
(transp. convolutions)

$$p(\mathbf{x}|\mathbf{z})$$

$$\mathbf{x} \sim p_{\text{data}}$$



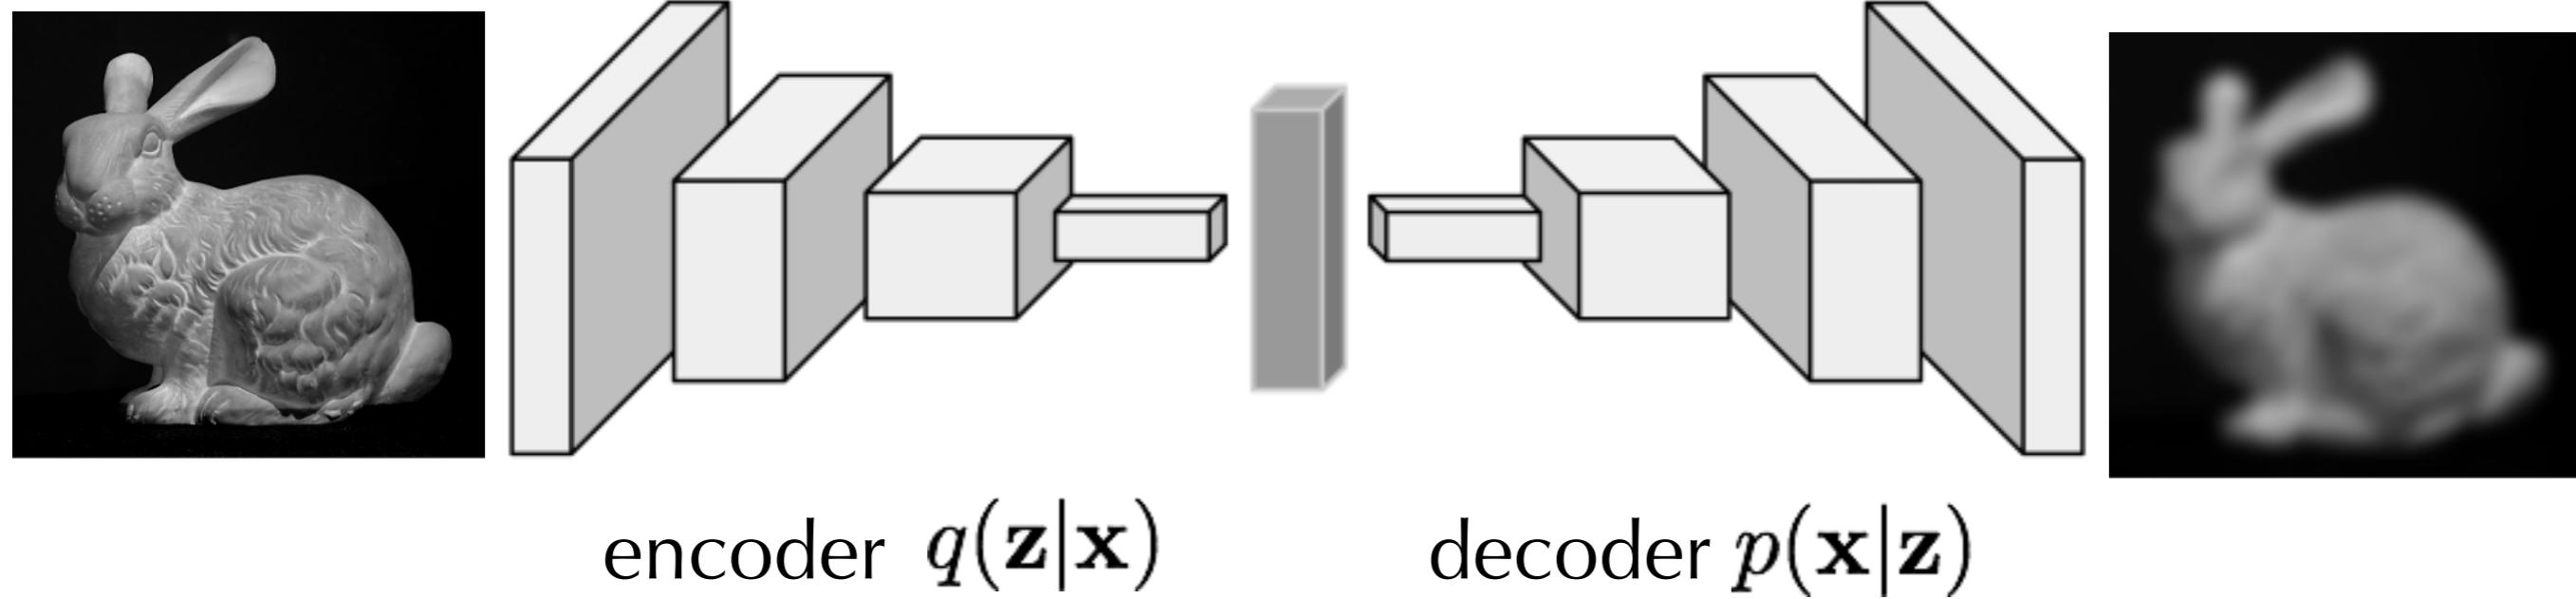
use decoder to ``translate''  
to image

# Today

- Variational Autoencoders (VAE)
- **Generative Adversarial Networks (GANs)**
- Unsupervised Image Translation

# Variational Autoencoder

$\mathbf{z} \sim p_{\text{model}}$  latent space

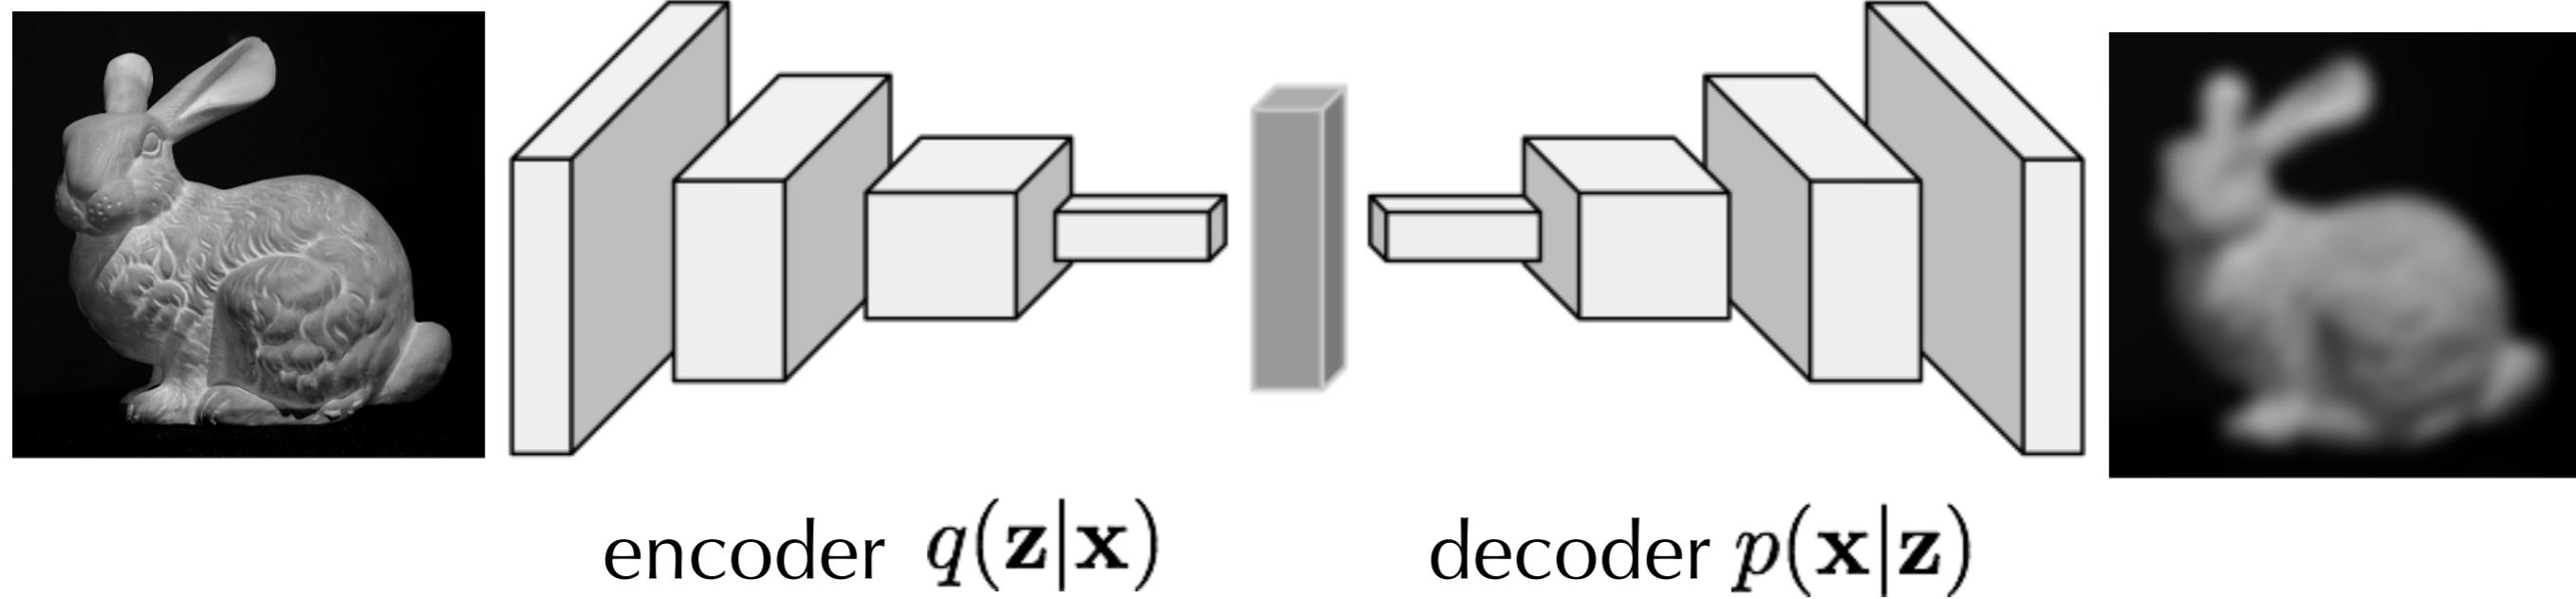


Maximize

$$E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$

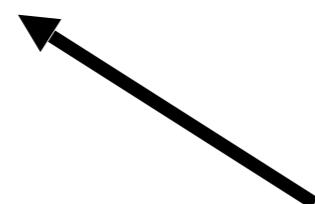
# Variational Autoencoder

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



Maximize

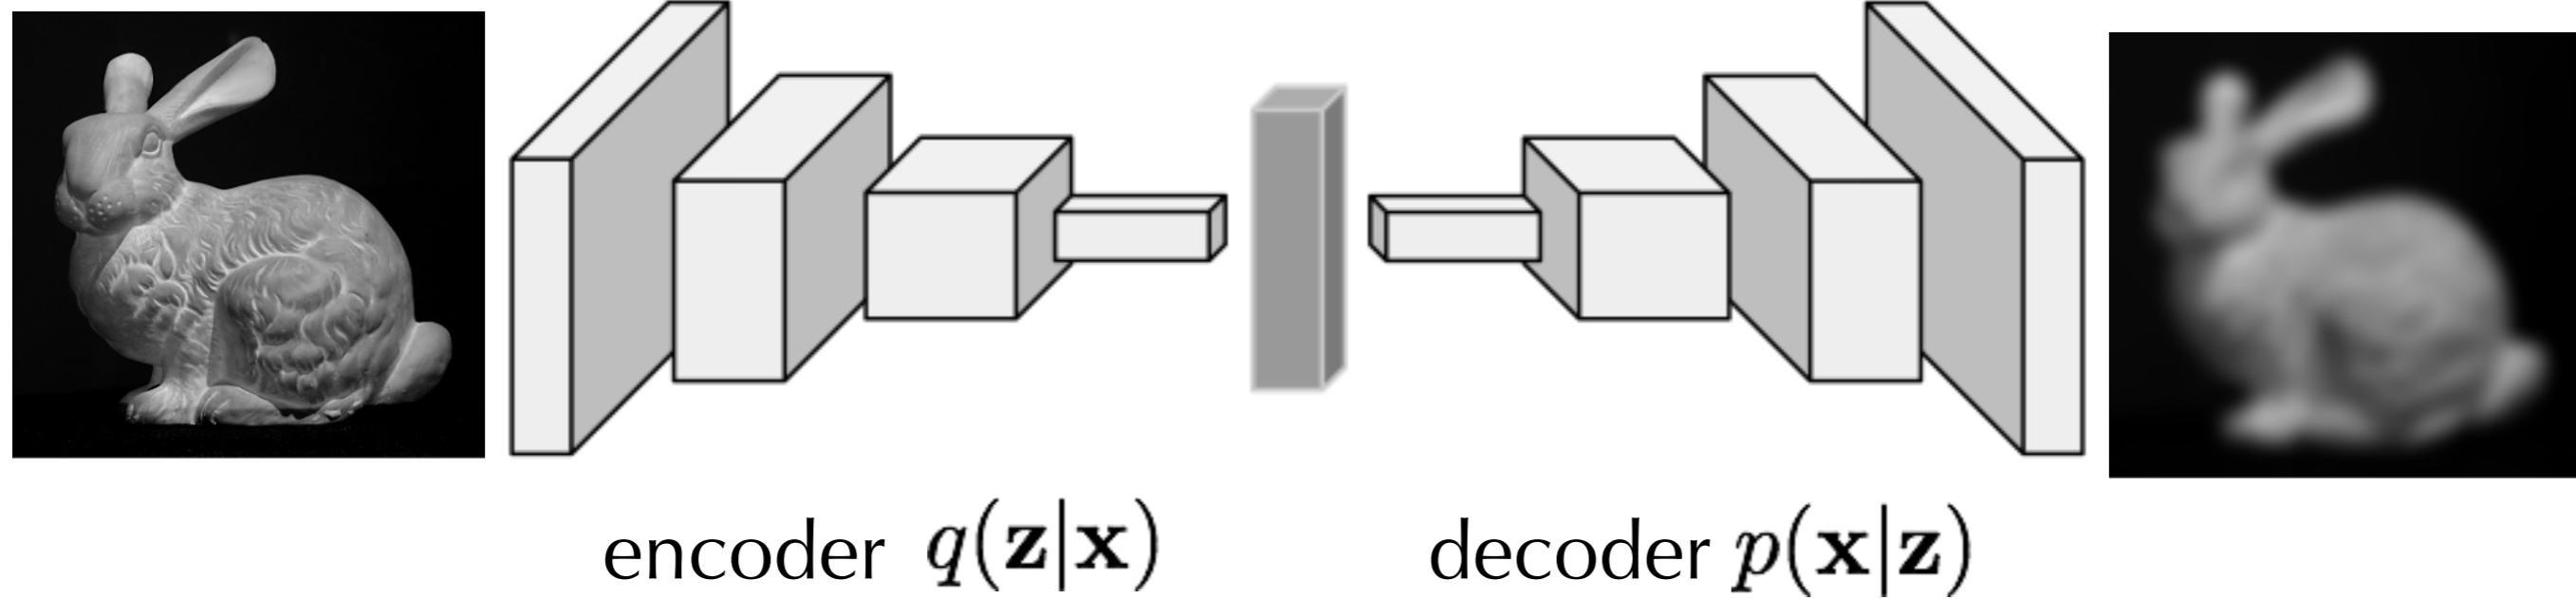
$$E_{\mathbf{z} \sim q} [\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$



Typically a normal distribution  $\rightarrow$  blurry results

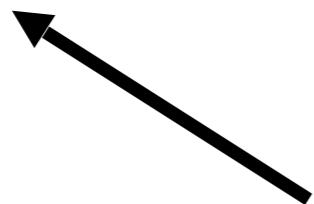
# Variational Autoencoder

$$\mathbf{z} \sim p_{\text{model}} \text{ latent space}$$



Maximize

$$E_{\mathbf{z} \sim q}[\log(p(\mathbf{x}|\mathbf{z}))] - D_{\text{KL}}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z})) \leq \log(p(\mathbf{x}))$$



**Better loss functions?**

Typically a normal distribution → blurry results

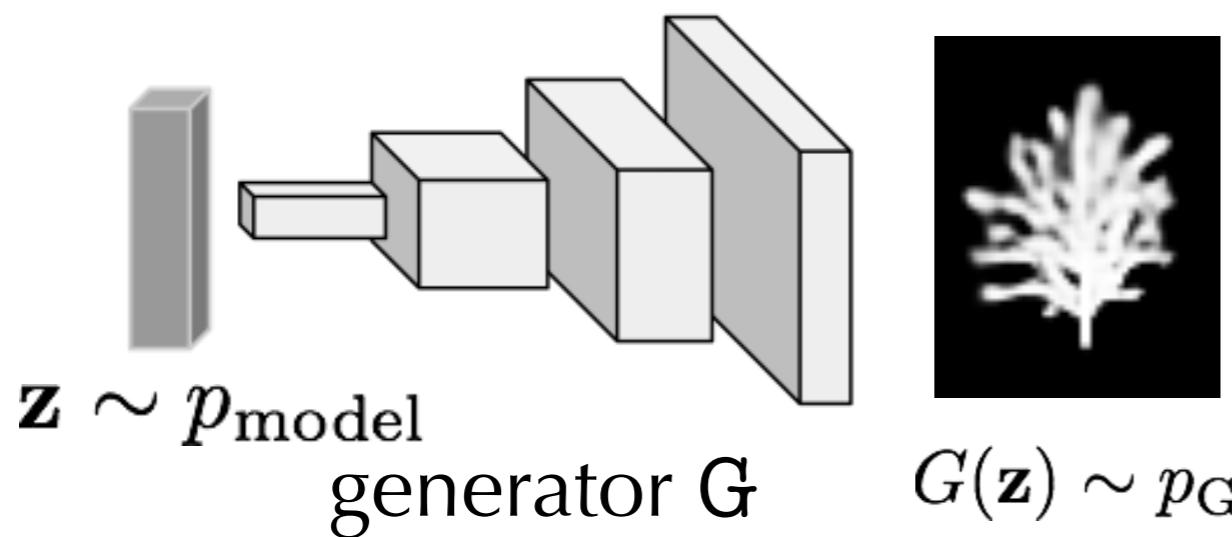
# Generative Adversarial Networks (GANs)

Motivation: Learn loss function

[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

# Generative Adversarial Networks (GANs)

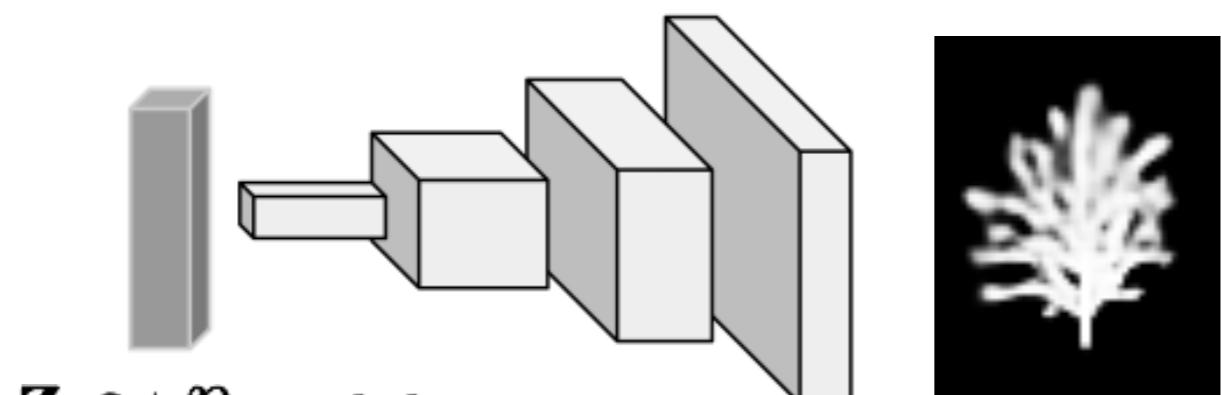
Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

# Generative Adversarial Networks (GANs)

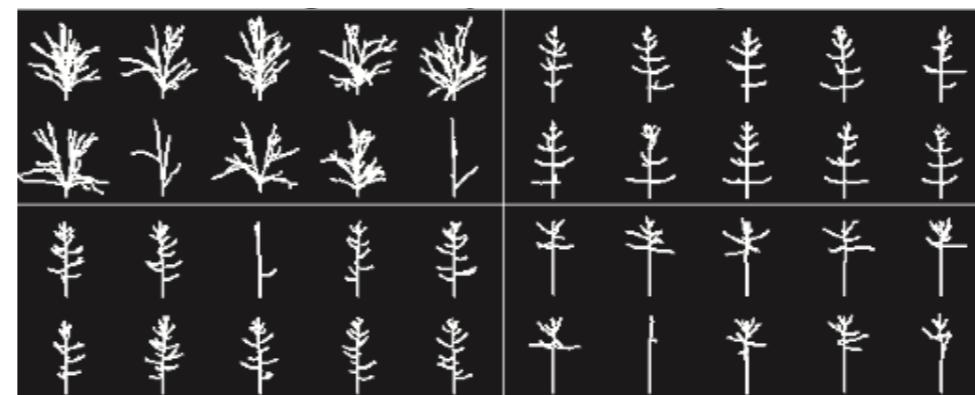
Motivation: Learn loss function



$\mathbf{z} \sim p_{\text{model}}$   
generator  $G$

$G(\mathbf{z}) \sim p_G$

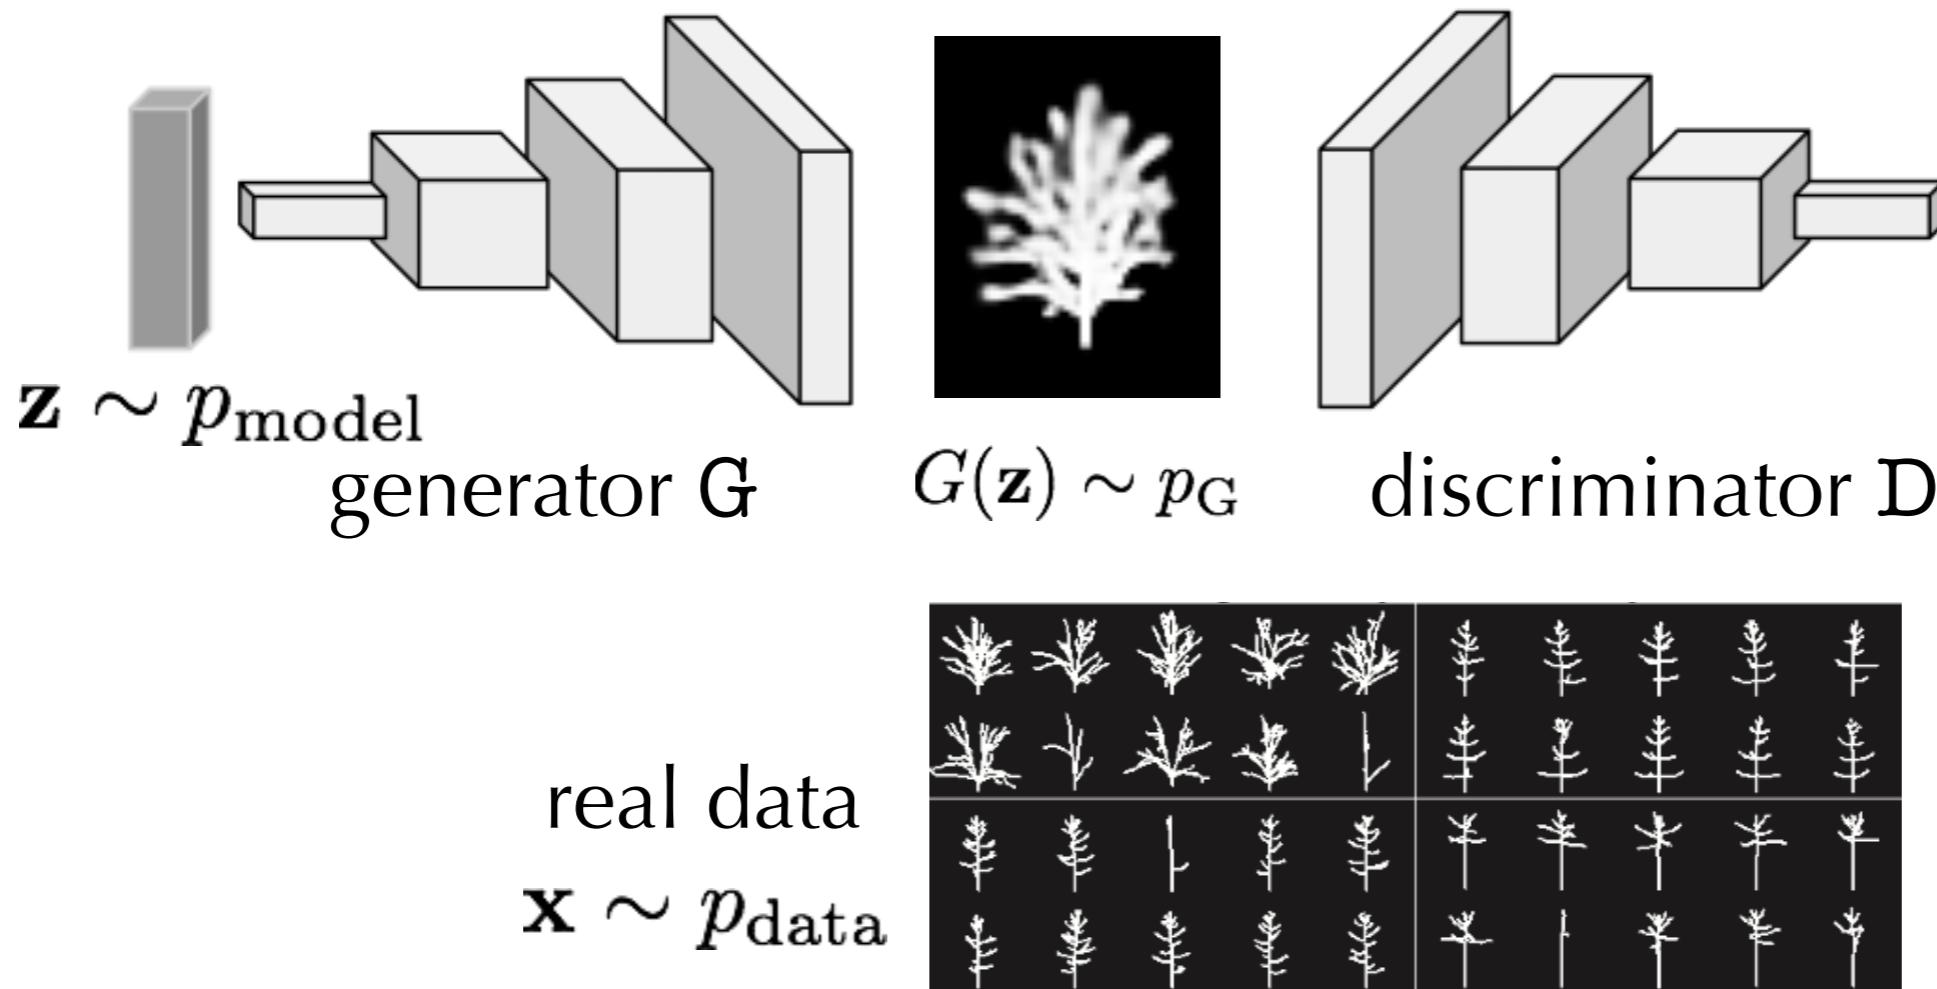
real data  
 $\mathbf{x} \sim p_{\text{data}}$



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

# Generative Adversarial Networks (GANs)

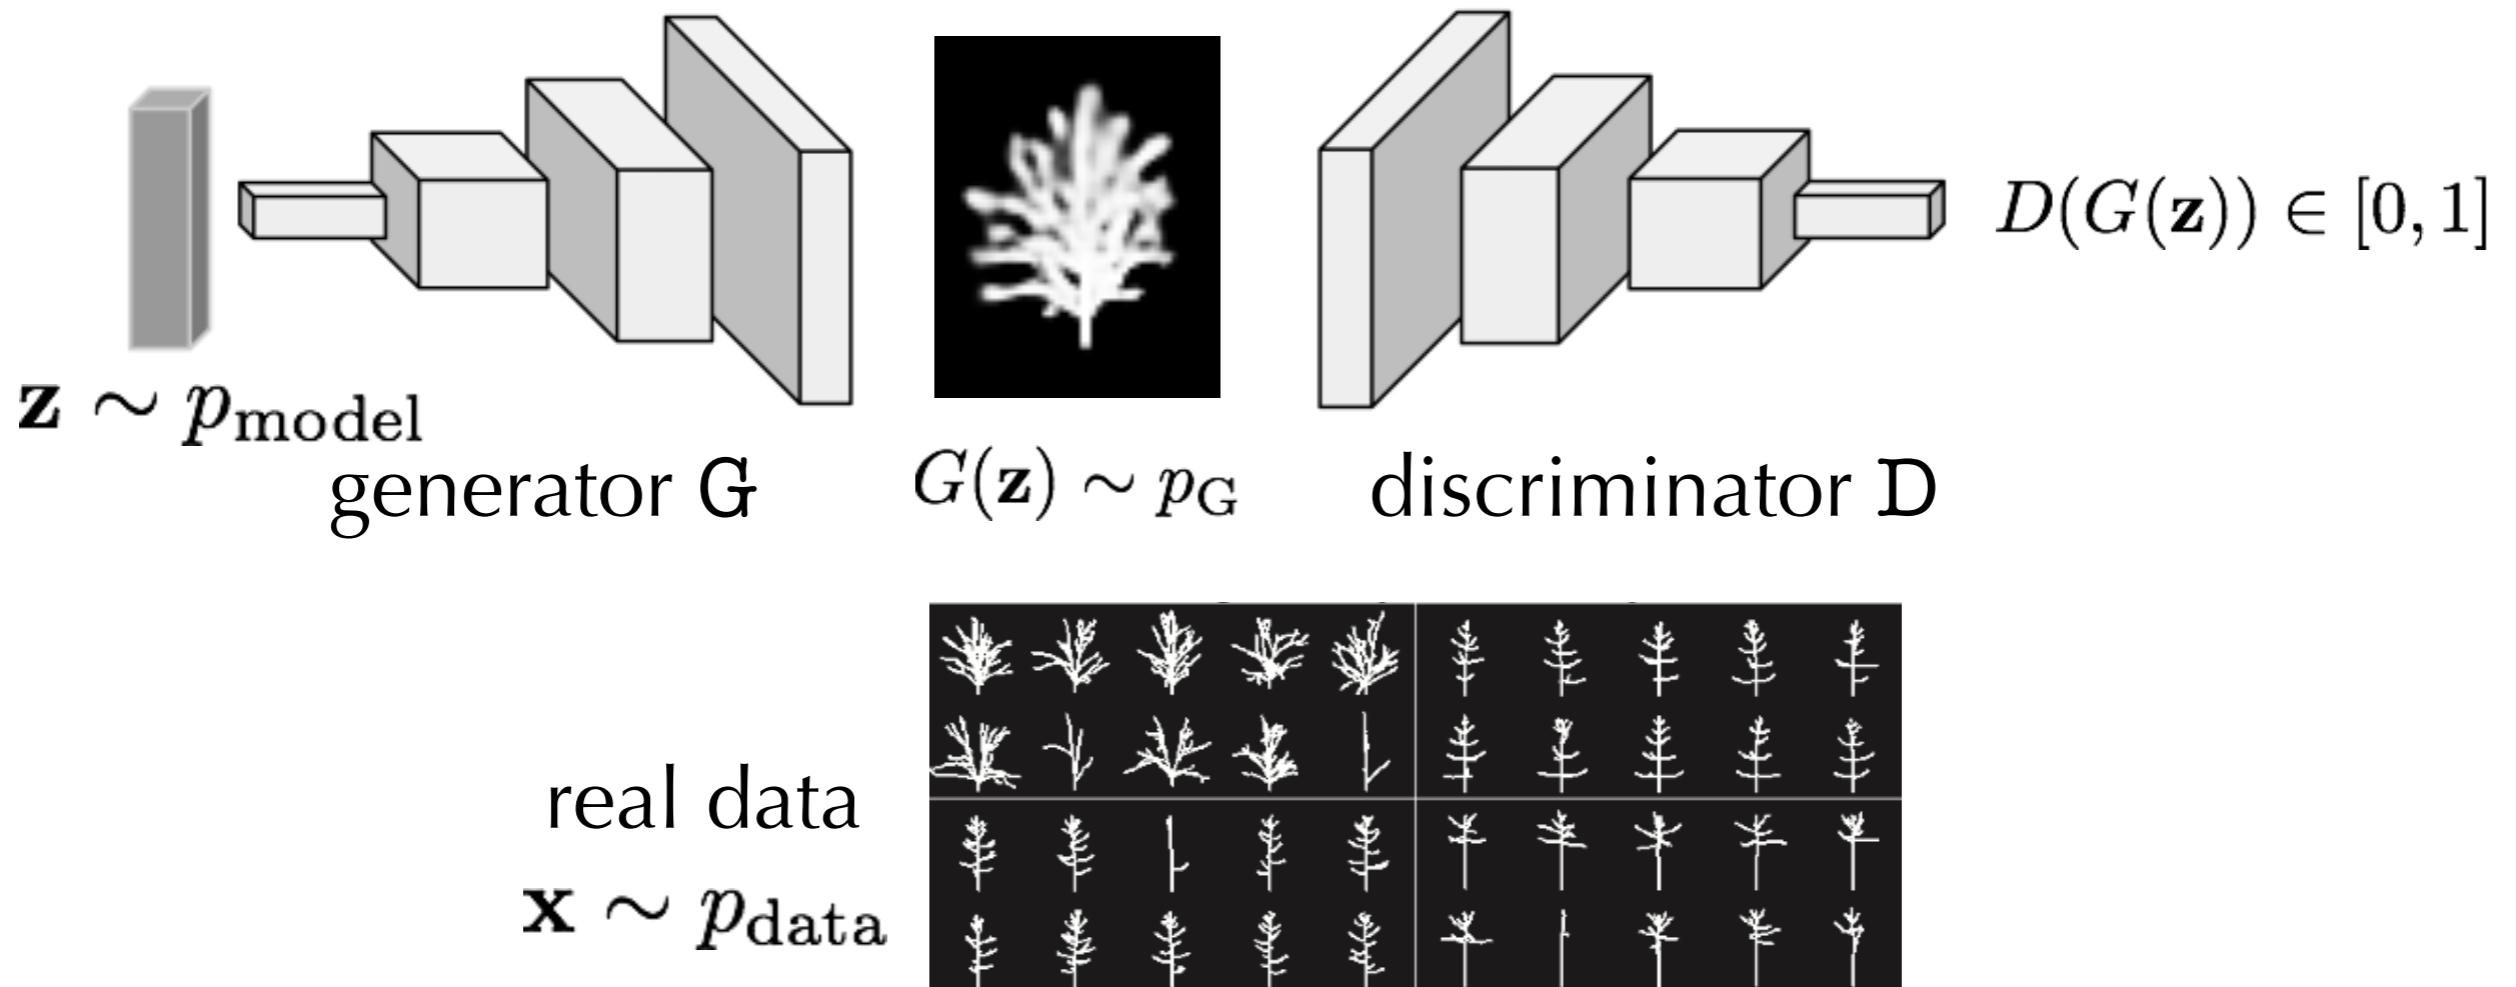
Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

# Generative Adversarial Networks (GANs)

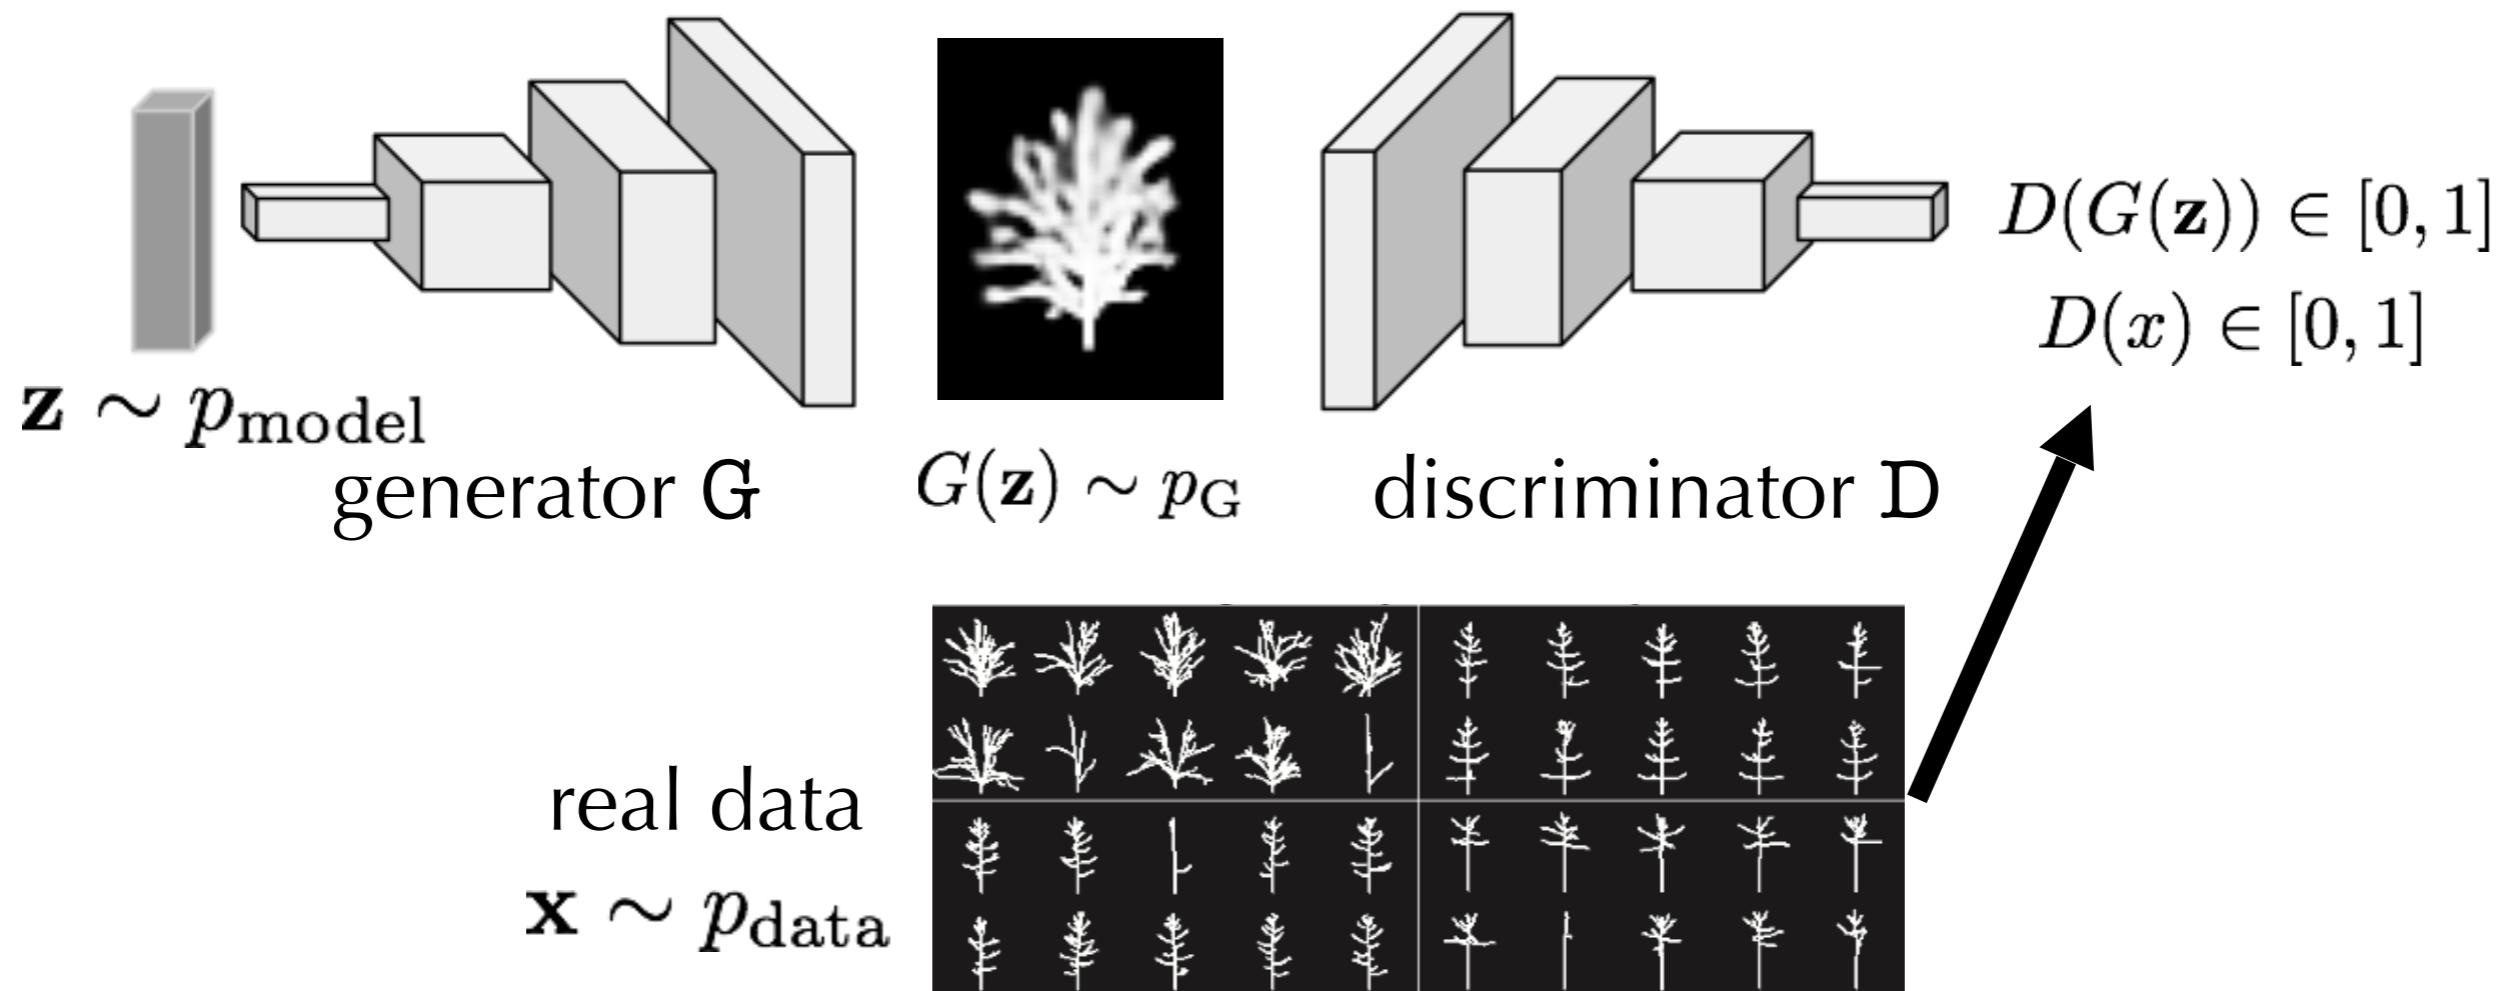
Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

# Generative Adversarial Networks (GANs)

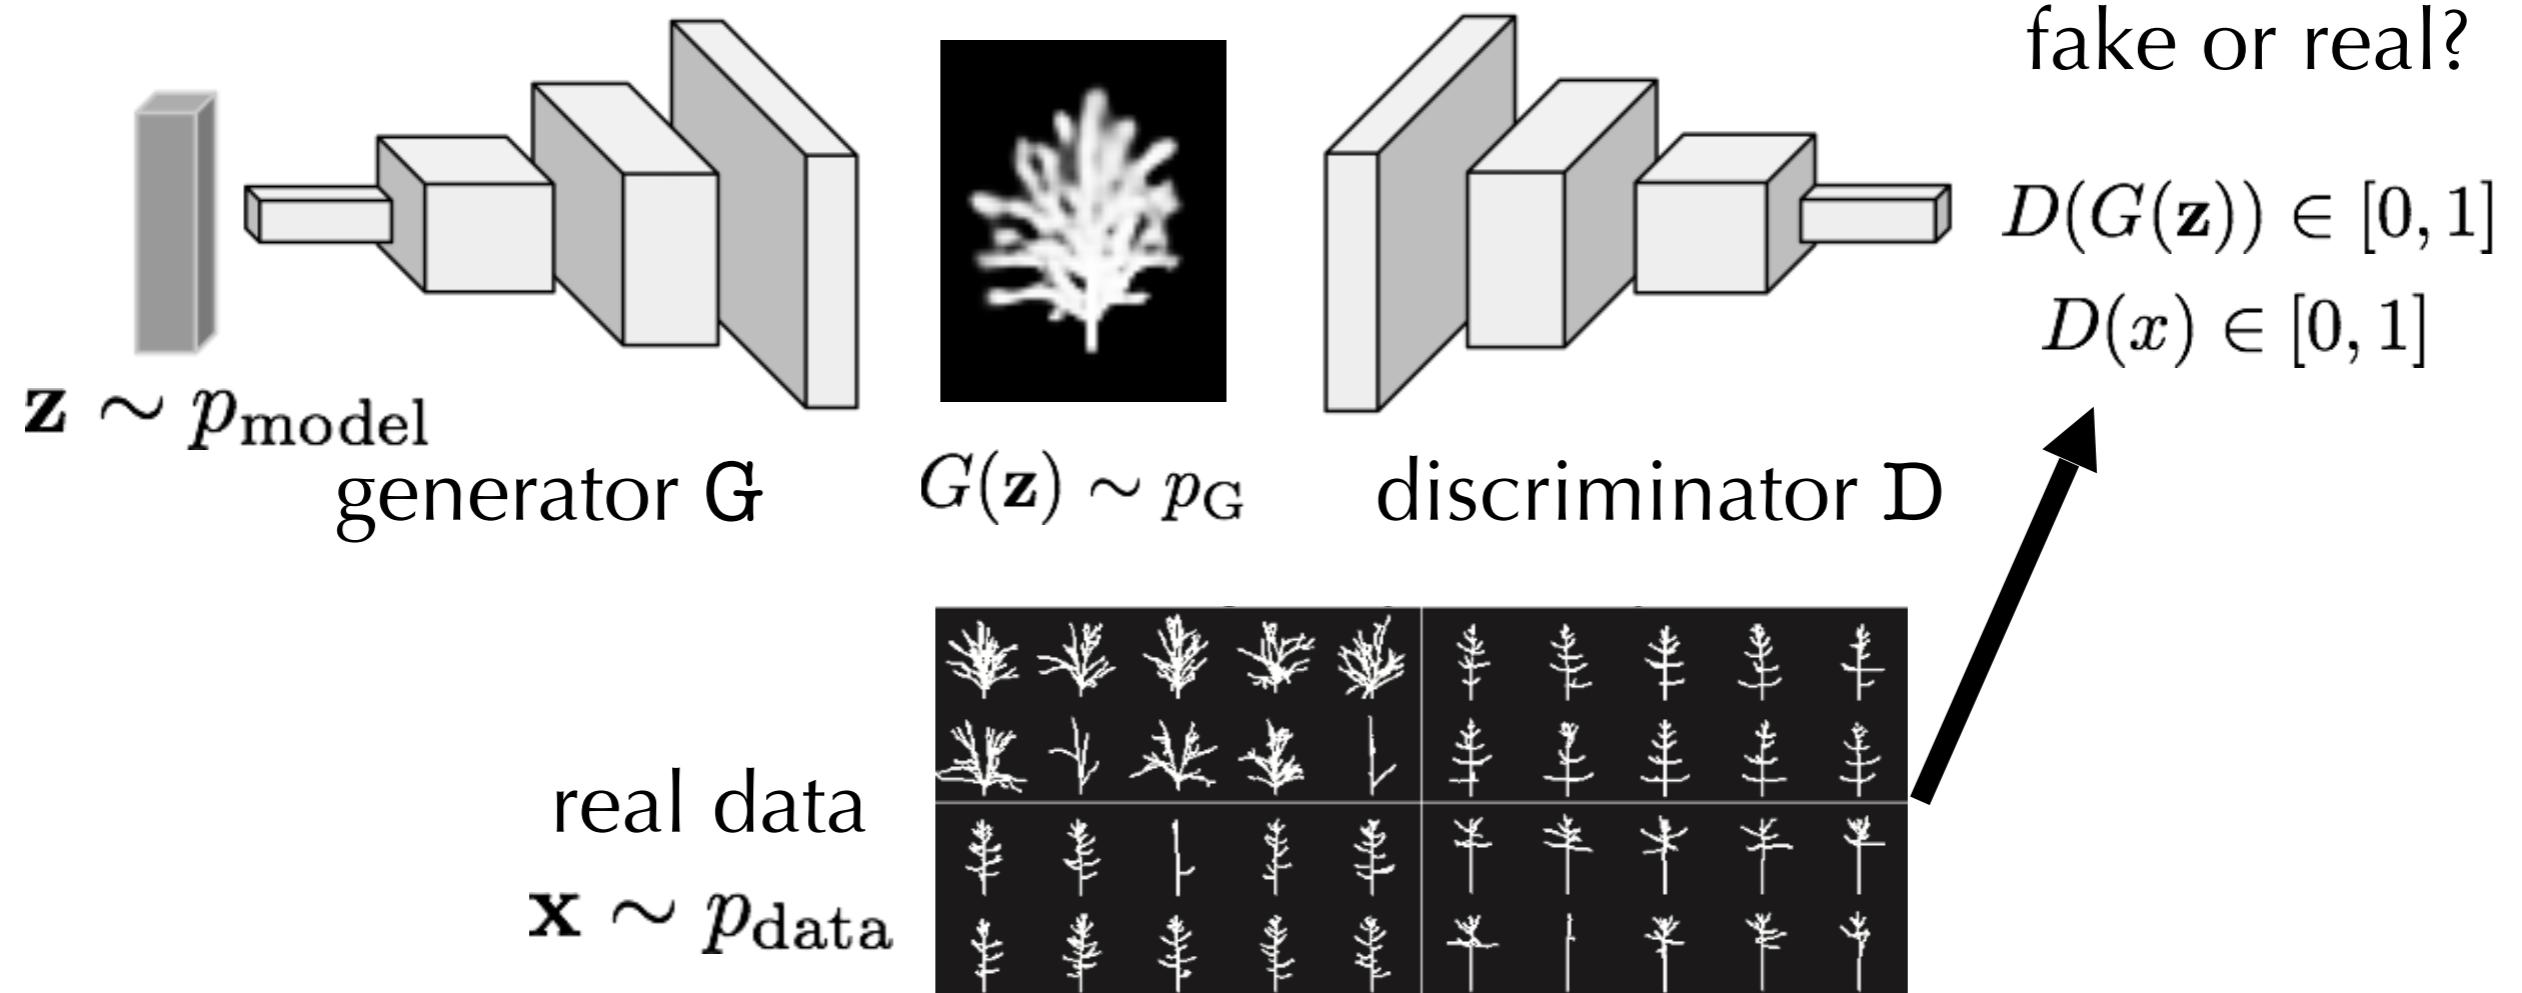
Motivation: Learn loss function



[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

# Generative Adversarial Networks (GANs)

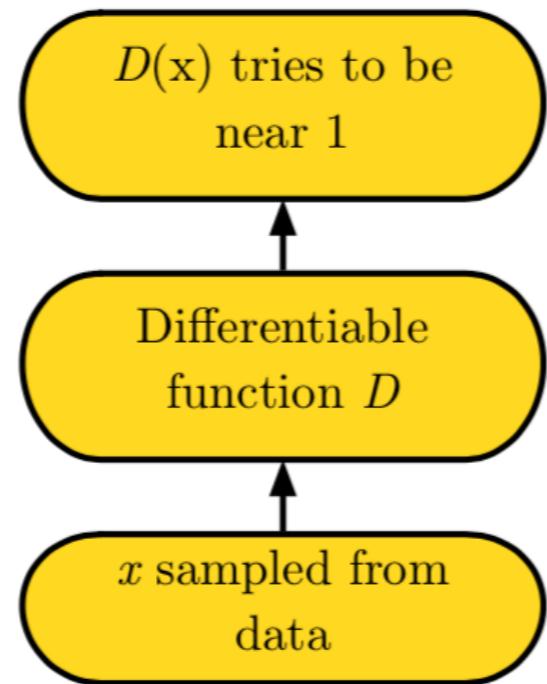
Motivation: Learn loss function



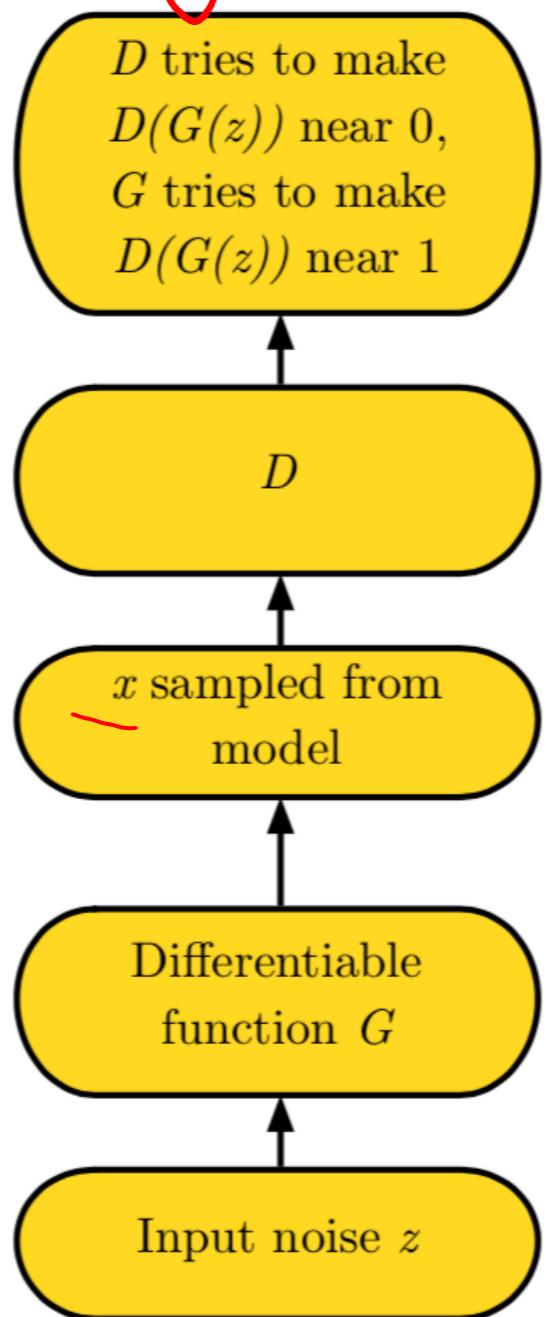
[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

# Generative Adversarial Networks (GANs)

*discriminador*



*generator* ✓



(Goodfellow et al., 2014)

slide credit: Ian Goodfellow

# Game Theoretic View

- Discriminator D tries to distinguish real and fake images

# Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”

# Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

# Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\max_D \quad \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

# Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\max_D \quad \underline{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))}$$

# Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\min_G \max_D \quad \underline{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))}$$

# Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\min_G \max_D \quad \underbrace{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))}_{\text{red line}}$$

# Game Theoretic View

- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\min_G \max_D \quad \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

- Equilibrium / global optimum at  $p_{\text{data}} = p_G$ , i.e., when real and synthetic images are indistinguishable for D:

# Game Theoretic View

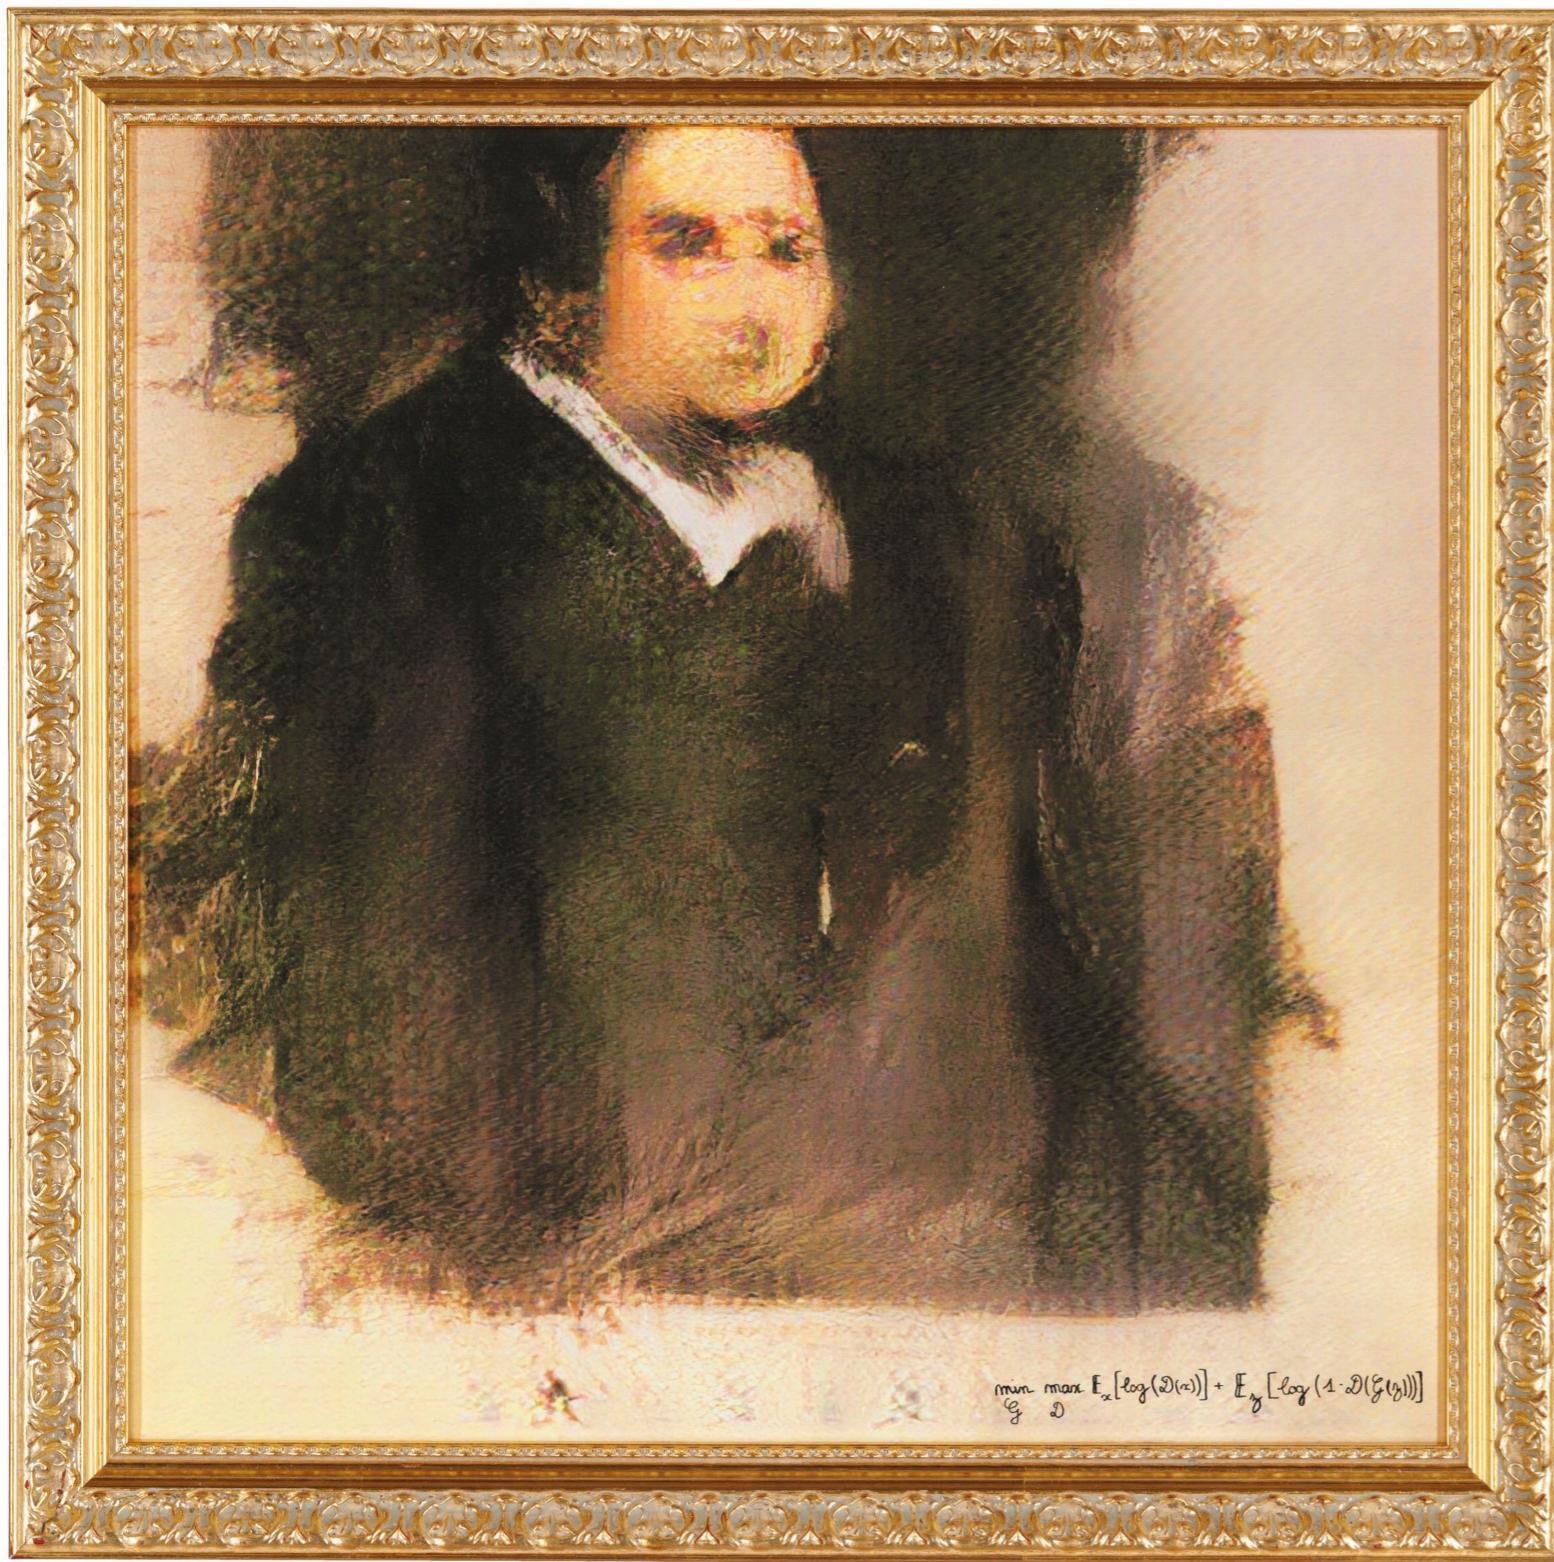
- Discriminator D tries to distinguish real and fake images
- Generator G tries to fool D, tries to make synthetic images look “as real as possible”
- G and D play zero-sum minimax game:

$$\min_G \max_D \quad \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

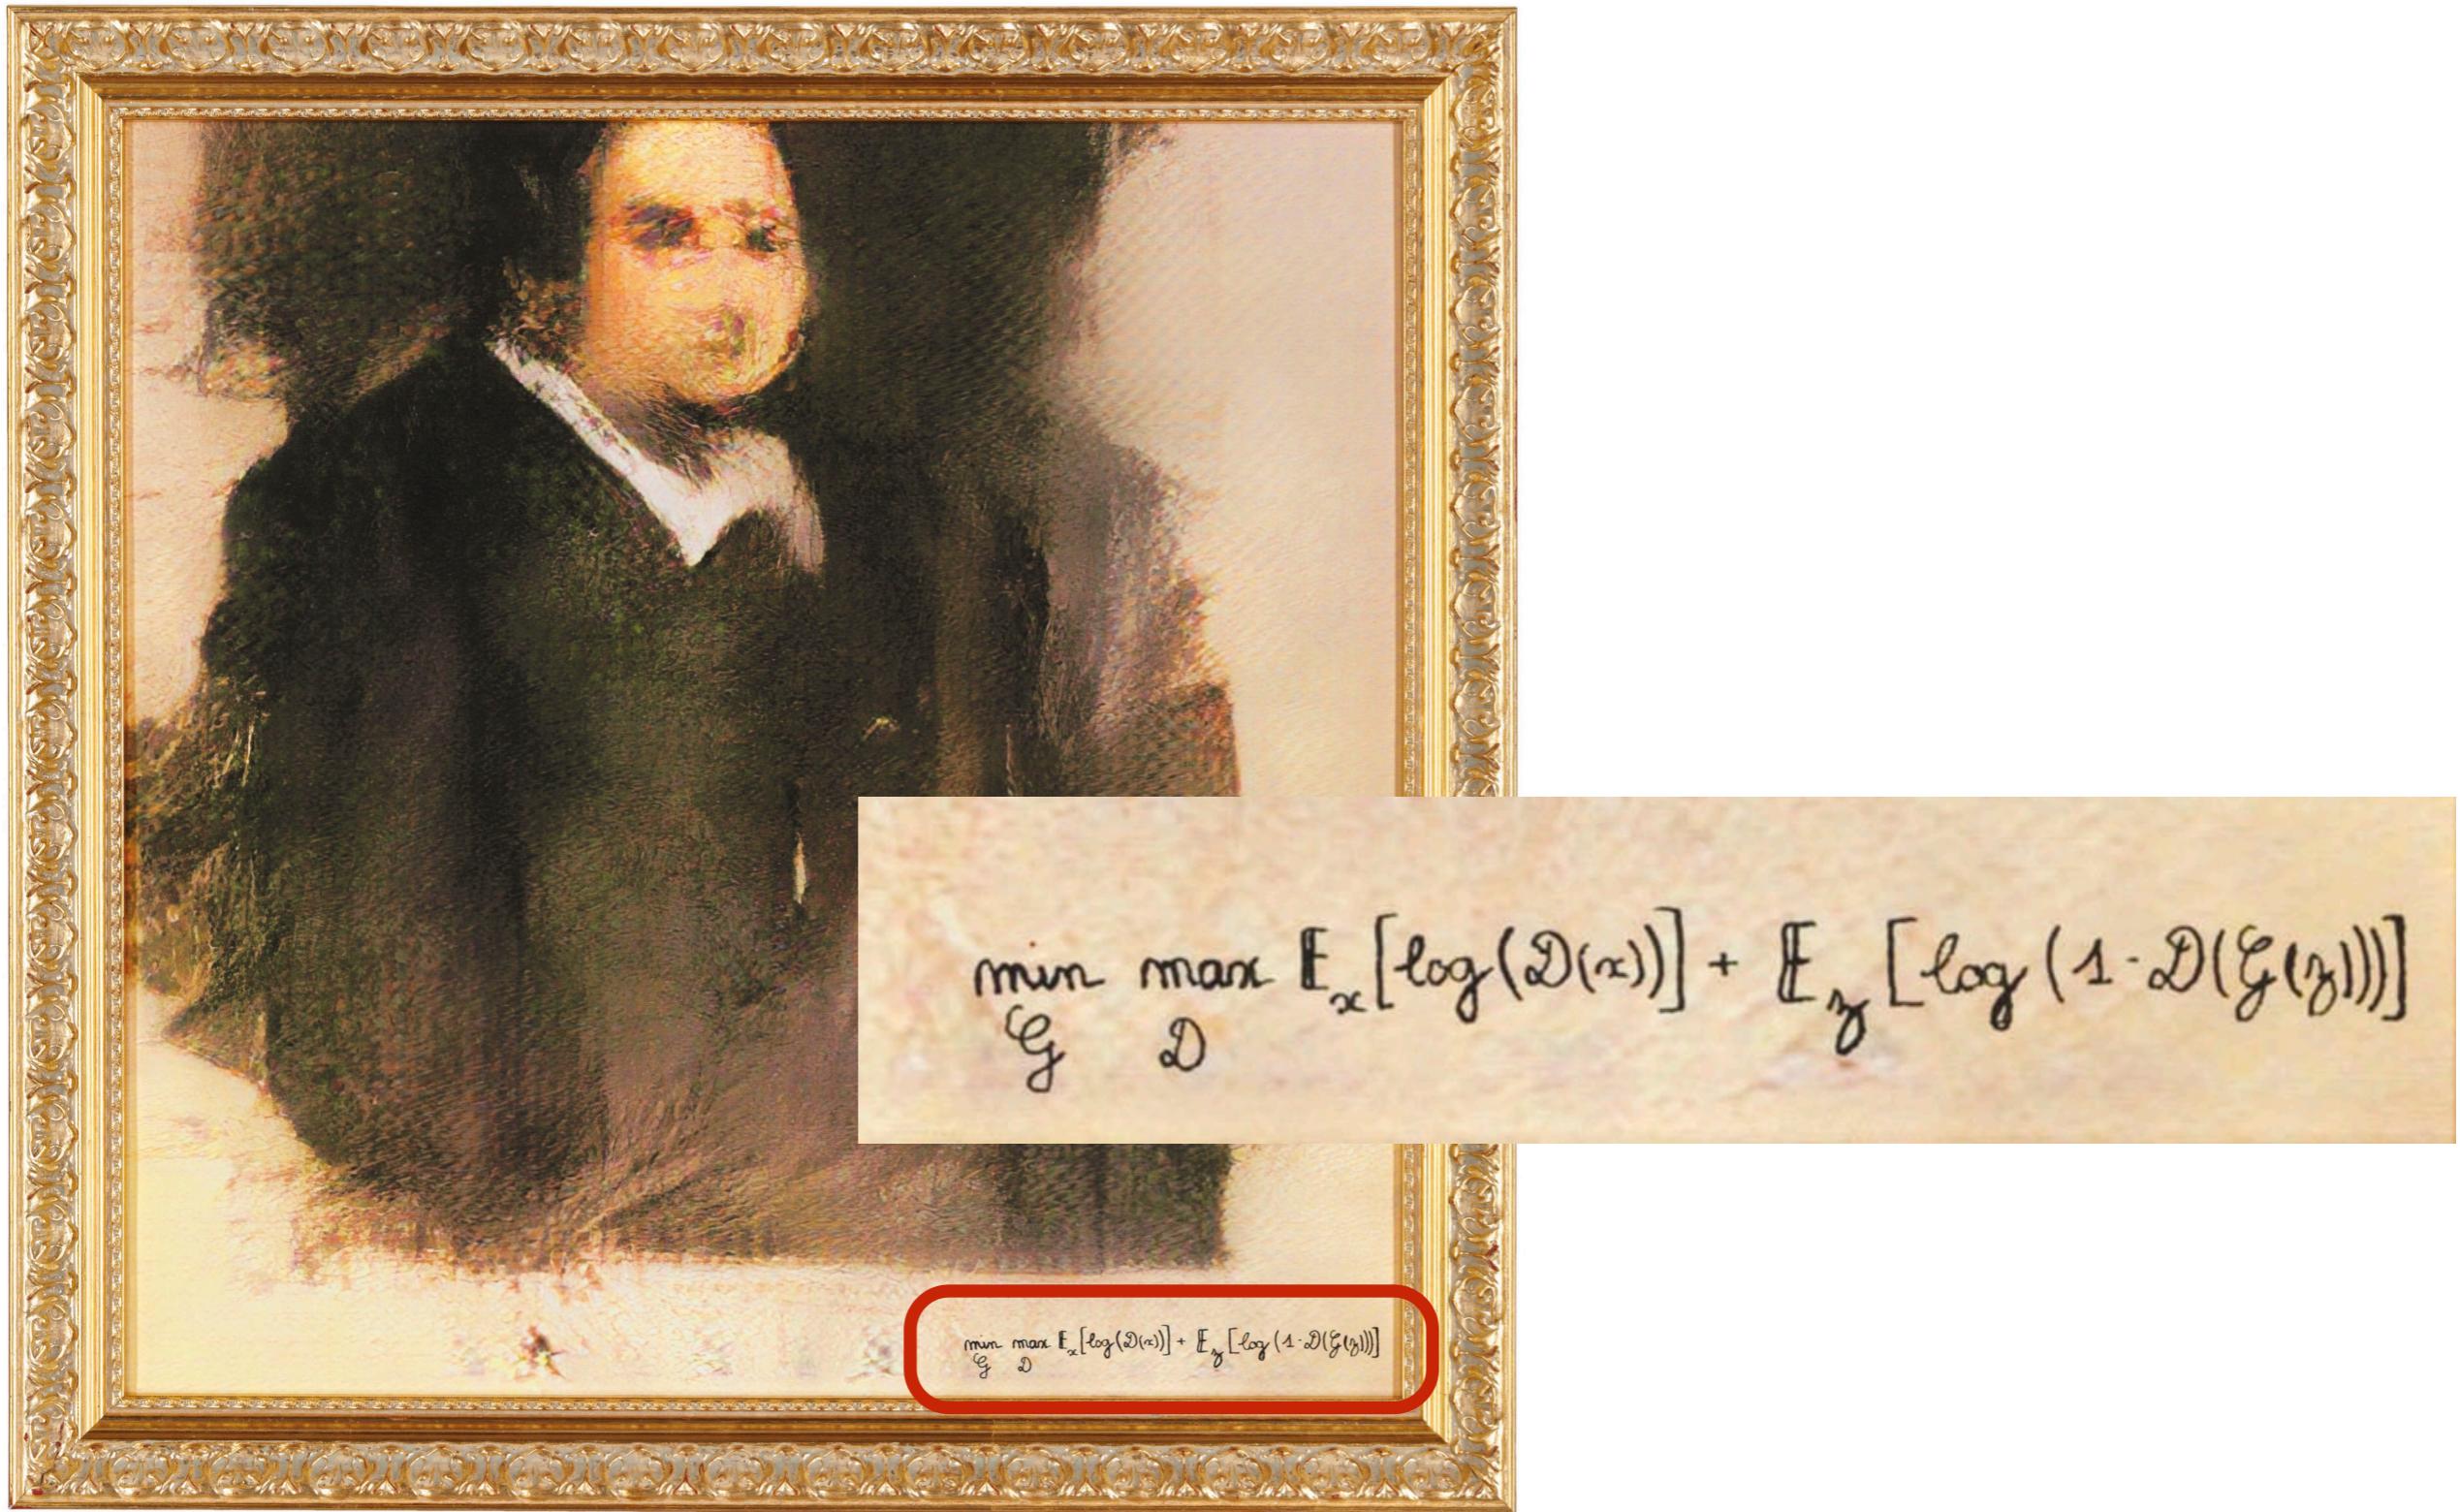
- Equilibrium / global optimum at  $p_{\text{data}} = p_G$ , i.e., when real and synthetic images are indistinguishable for D:

$$d(\mathbf{x}) = 1/2 \quad \forall \mathbf{x} \sim p_{\text{data}}, \mathbf{x} \sim p_G$$

# Generating Art Revisited



# Generating Art Revisited



# Training GANs

---

```
for number of training iterations do  
    for  $k$  steps do
```

```
end for
```

## Train discriminator D

## Train generator G

```
end for
```

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

# Training GANs

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

**end for**

## Train generator G

**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

# Training GANs

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))).$$

**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

# Training GANs is Hard

- Need to carefully **balance the generator and discriminator**

# Training GANs is Hard

- Need to carefully **balance the generator and discriminator**
  - D becomes too powerful too soon → G cannot improve (gradient vanishes)

# Training GANs is Hard

- Need to carefully **balance the generator and discriminator**
  - D becomes too powerful too soon → G cannot improve (gradient vanishes)
  - D not powerful enough → G does not improve

# Training GANs is Hard

- Need to carefully **balance the generator and discriminator**
  - D becomes too powerful too soon → G cannot improve (gradient vanishes)
  - D not powerful enough → G does not improve
  - Black magic: Schedule training of generator and discriminator

# Training GANs is Hard

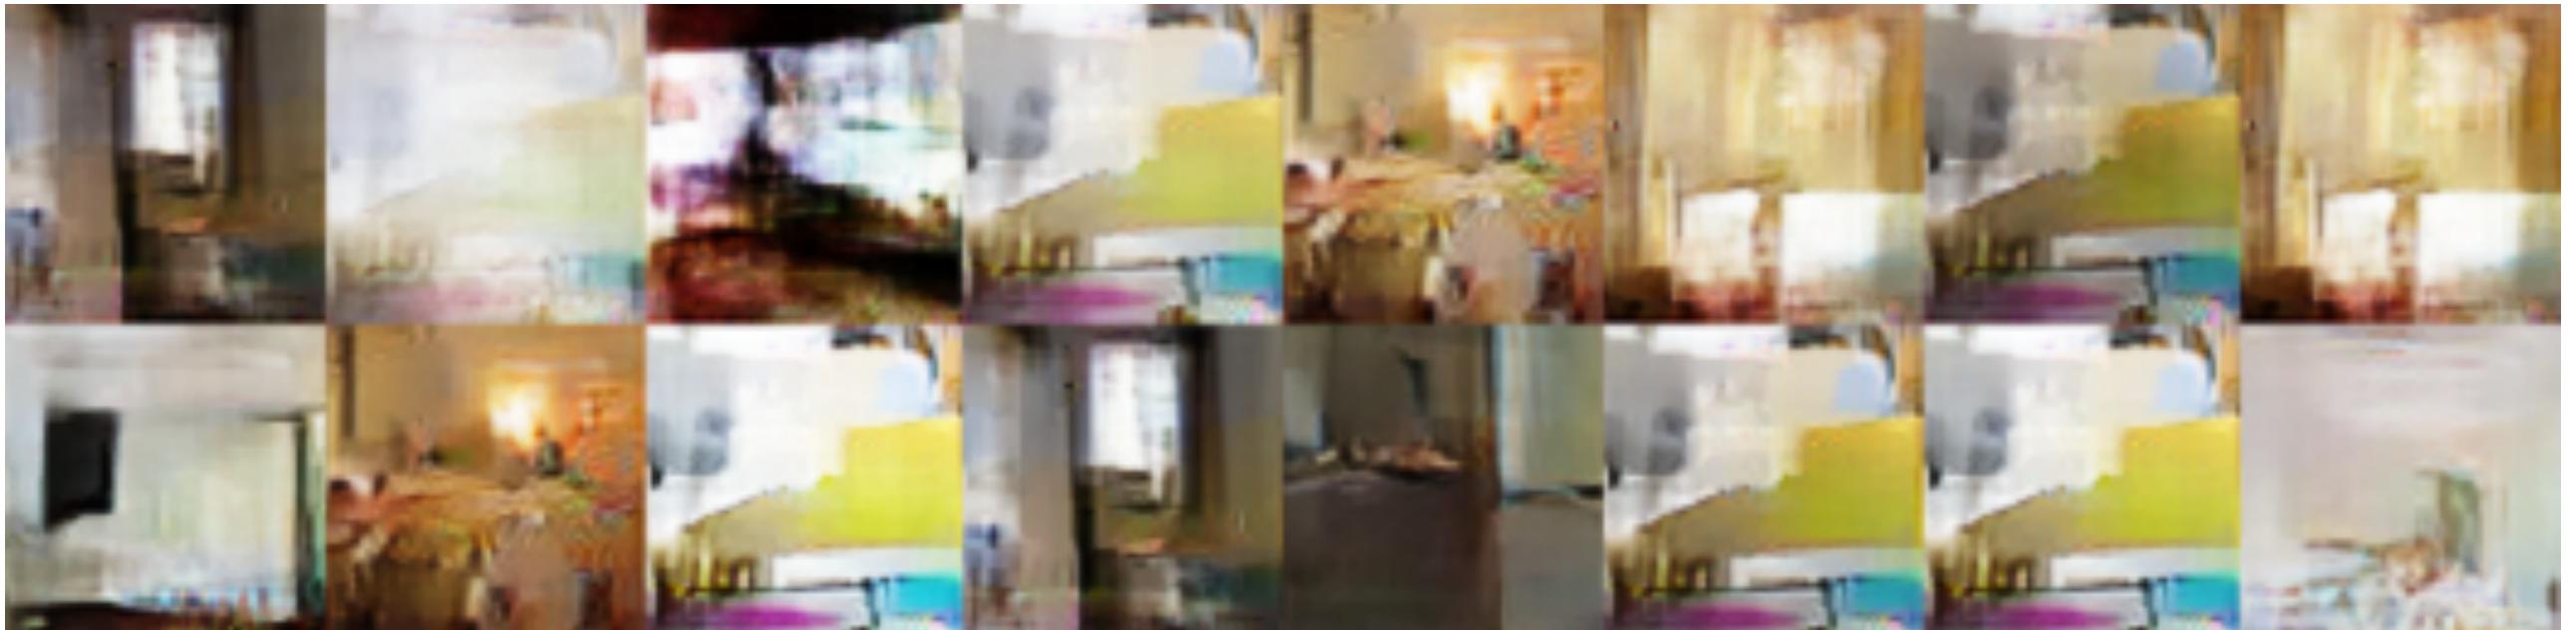
- Need to carefully **balance the generator and discriminator**
  - D becomes too powerful too soon → G cannot improve (gradient vanishes)
  - D not powerful enough → G does not improve
  - Black magic: Schedule training of generator and discriminator
- **Divergence:** GANs might not train at all

# Training GANs is Hard

- Need to carefully **balance the generator and discriminator**
  - D becomes too powerful too soon → G cannot improve (gradient vanishes)
  - D not powerful enough → G does not improve
  - Black magic: Schedule training of generator and discriminator
- **Divergence:** GANs might not train at all
- [List](#) of tips & tricks for training

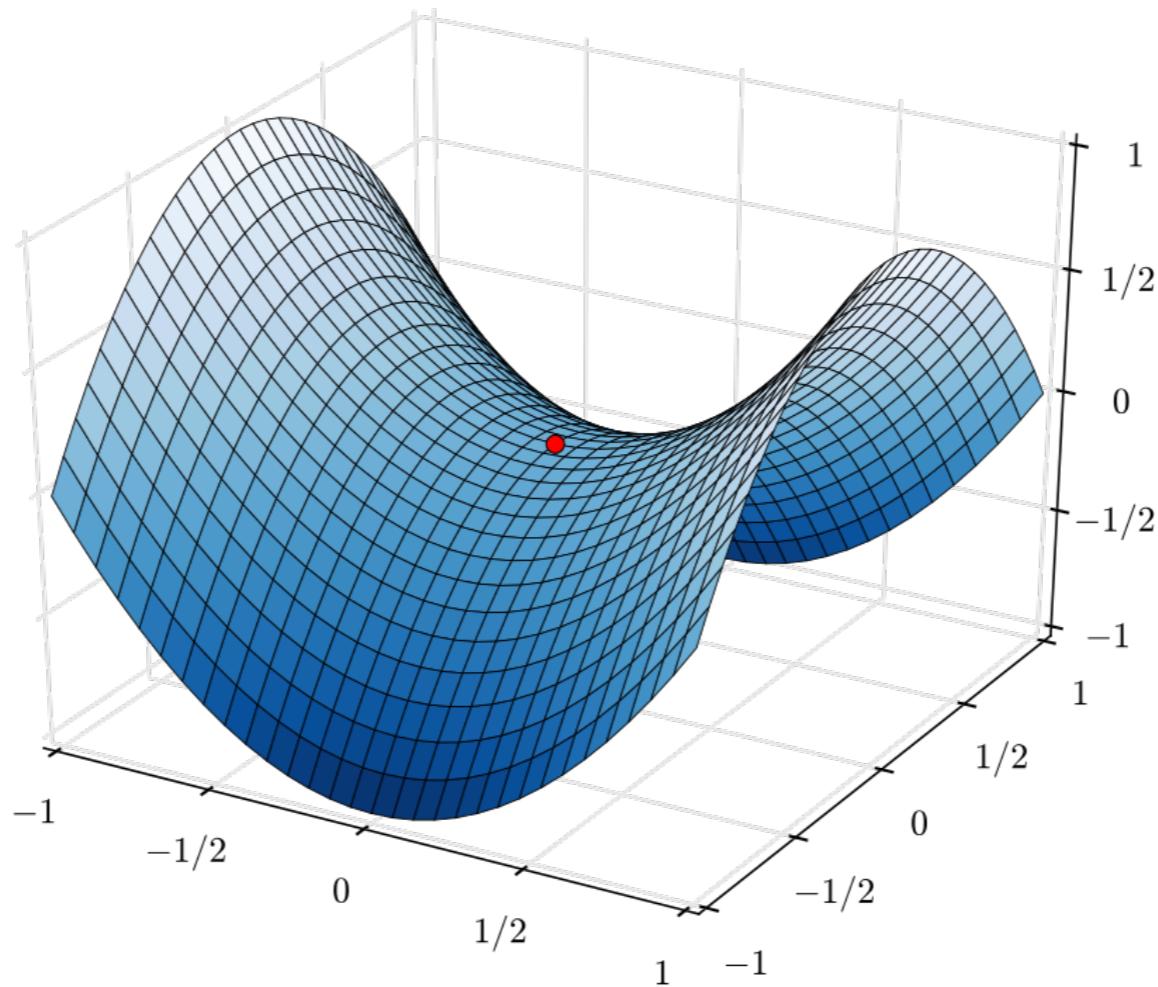
# Mode Collapse

Generated images span only small subspace



[Arjovsky et al., Wasserstein GAN, arXiv:1701.07875]

# Why is Training GANs so Hard?



[image](#): © Nicoguaro (CC Attribution Unported 3.0)

GAN equilibria are saddle points

$$\min_G \max_D \quad \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

[Goodfellow et al., Generative Adversarial Nets, NIPS 2014]

# Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

# Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

- Optimal discriminator  $D^*$  for given  $G$ :

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$$

# Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

- Optimal discriminator  $D^*$  for given  $G$ :

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$$

- GAN objective becomes

$$\min_G \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log \left( \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right) + \mathbb{E}_{\mathbf{x} \sim p_G} \log \left( \frac{p_G(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right)$$

# Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

- Optimal discriminator  $D^*$  for given  $G$ :

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$$

- GAN objective becomes

$$\min_G \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log \left( \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right) + \mathbb{E}_{\mathbf{x} \sim p_G} \log \left( \frac{p_G(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right)$$

# Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

- Optimal discriminator  $D^*$  for given  $G$ :

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$$

- GAN objective becomes

$$\min_G \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log \left( \frac{p_{\text{data}}(\mathbf{x})}{\text{back}(p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x}))} \right) + \mathbb{E}_{\mathbf{x} \sim p_G} \log \left( \frac{p_G(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right)$$

$$\min_G \text{KL} \left( p_{\text{data}} \parallel \frac{p_{\text{data}} + p_G}{2} \right) + \text{KL} \left( p_G \parallel \frac{p_{\text{data}} + p_G}{2} \right) - \log(4)$$

# Why is Training GANs so Hard?

- GAN objective:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log(d(\mathbf{x})) + \mathbb{E}_{\mathbf{z} \sim p_{\text{model}}} \log(1 - d(g(\mathbf{z})))$$

- Optimal discriminator  $D^*$  for given  $G$ :

$$D_G^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})}$$

- GAN objective becomes

$$\min_G \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log \left( \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right) + \mathbb{E}_{\mathbf{x} \sim p_G} \log \left( \frac{p_G(\mathbf{x})}{p_{\text{data}}(\mathbf{x}) + p_G(\mathbf{x})} \right)$$

$$\min_G \text{KL} \left( p_{\text{data}} \parallel \frac{p_{\text{data}} + p_G}{2} \right) + \text{KL} \left( p_G \parallel \frac{p_{\text{data}} + p_G}{2} \right) - \log(4)$$

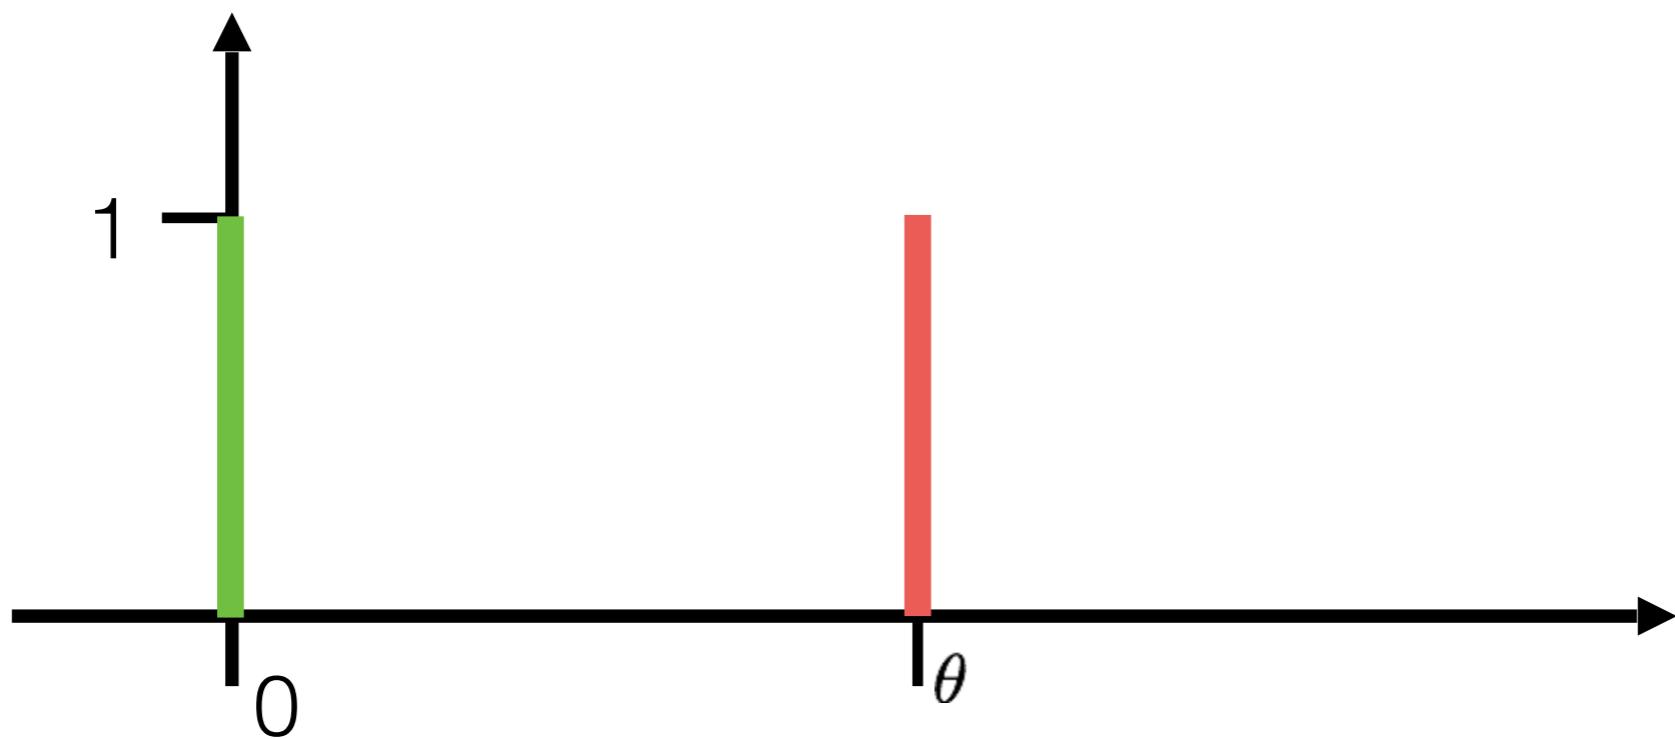
  
 $2 \cdot \text{JSD}(p_{\text{data}} \parallel p_G)$  Jensen-Shannon divergence

# Why is Training GANs so Hard?

**Problem:** Jensen-Shannon divergence gradient not always useful → problem when using gradient descent

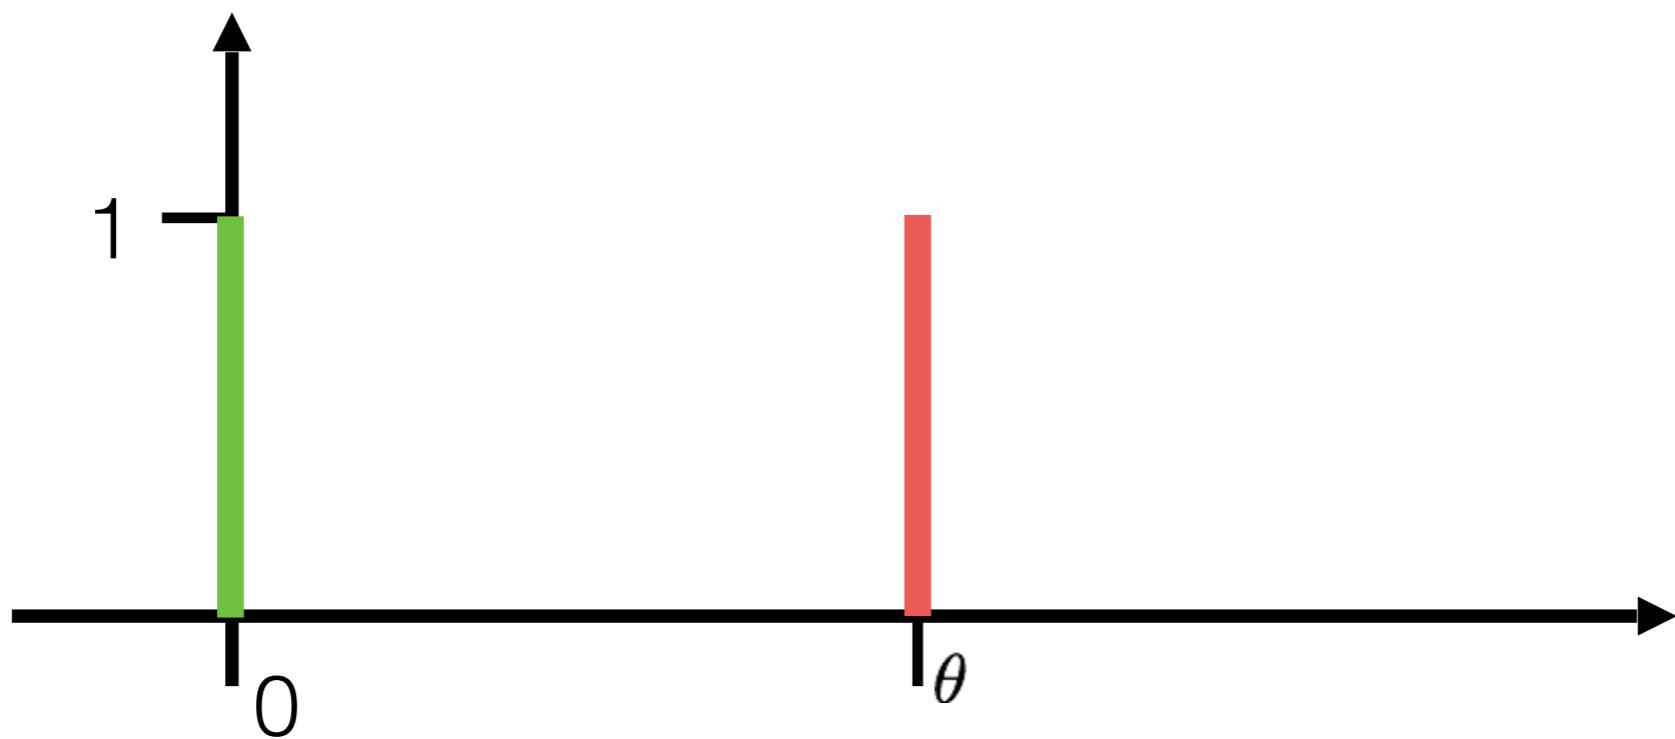
# Why is Training GANs so Hard?

**Problem:** Jensen-Shannon divergence gradient not always useful → problem when using gradient descent



# Why is Training GANs so Hard?

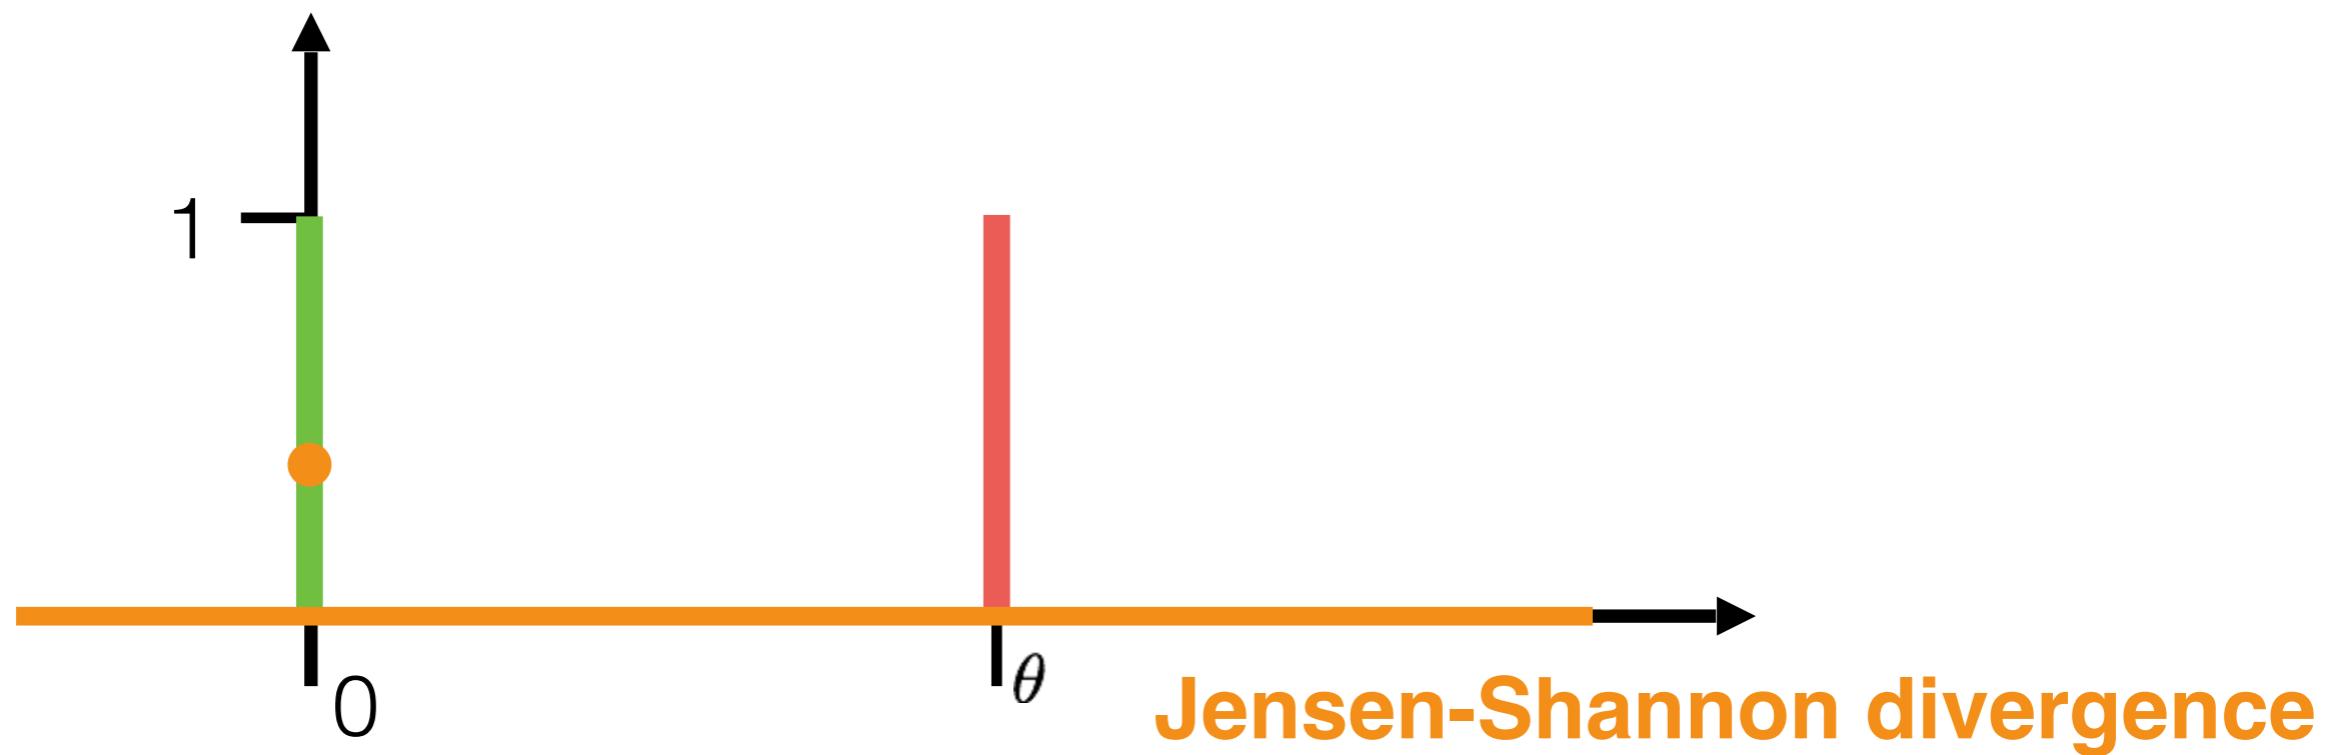
**Problem:** Jensen-Shannon divergence gradient not always useful → problem when using gradient descent



Simple one-parameter distribution  $p_\theta(z) = (\theta, z)$  , with  $\sim U[0, 1]$

# Why is Training GANs so Hard?

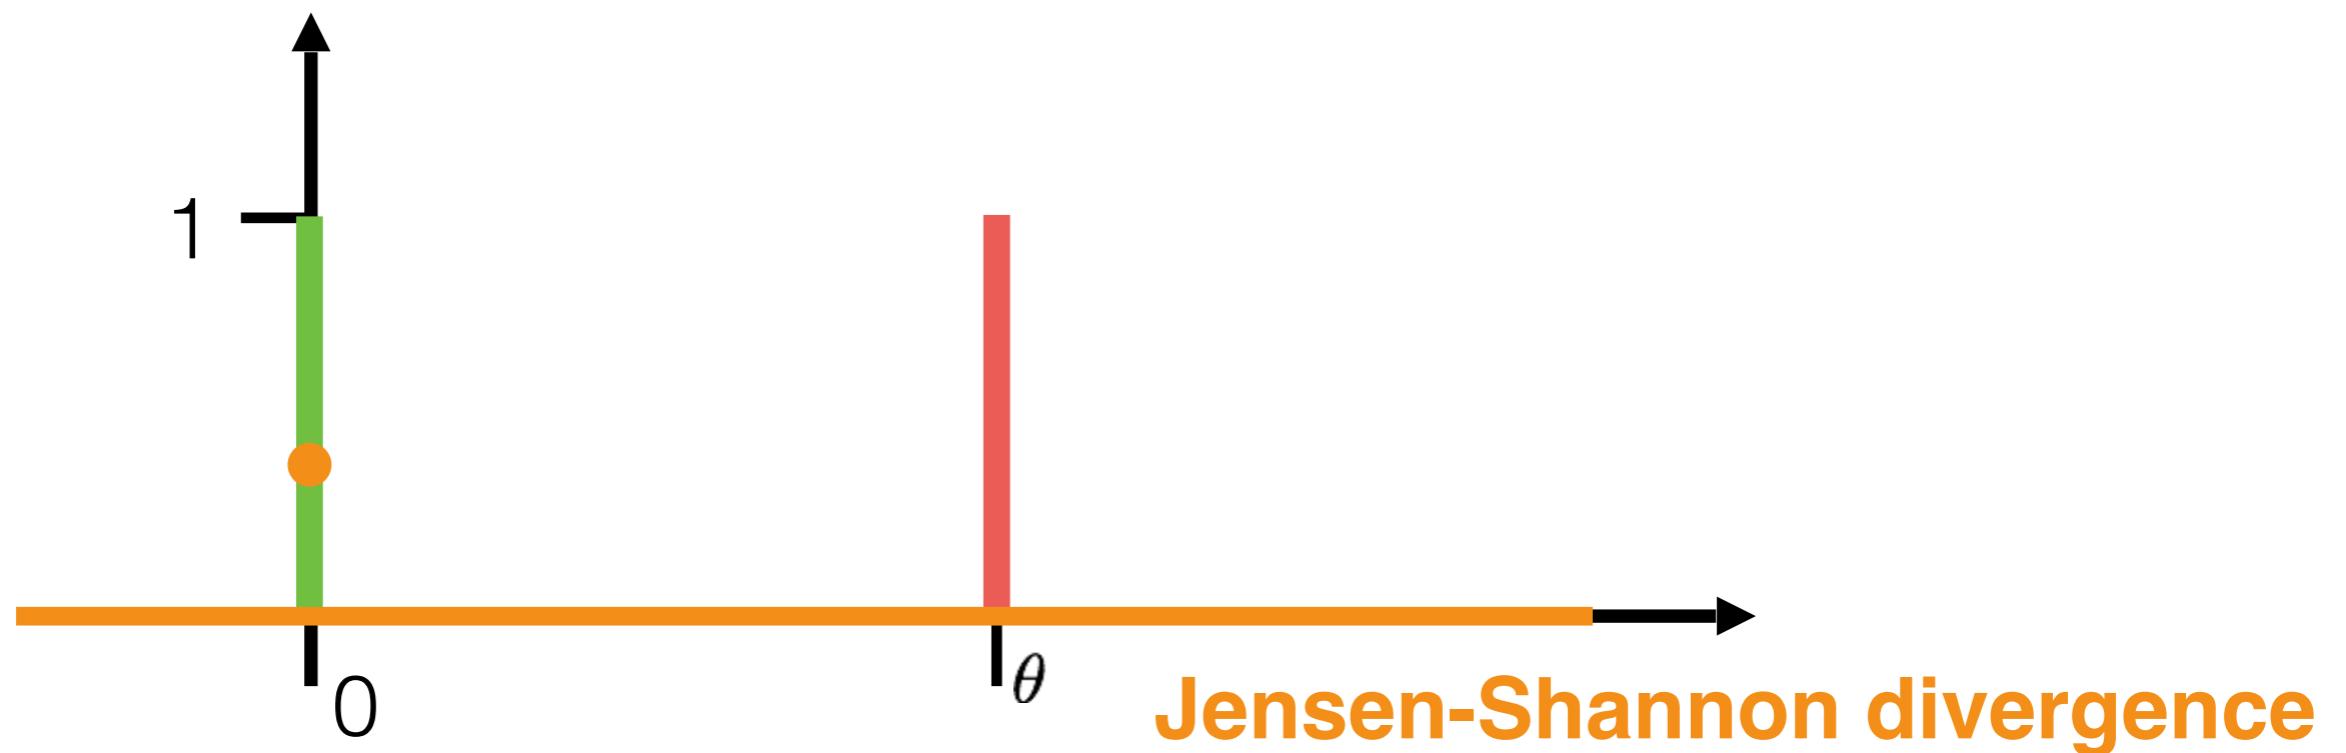
**Problem:** Jensen-Shannon divergence gradient not always useful → problem when using gradient descent



Simple one-parameter distribution  $p_\theta(z) = (\theta, z)$  , with  $\sim U[0, 1]$

# Why is Training GANs so Hard?

**Problem:** Jensen-Shannon divergence gradient not always useful → problem when using gradient descent

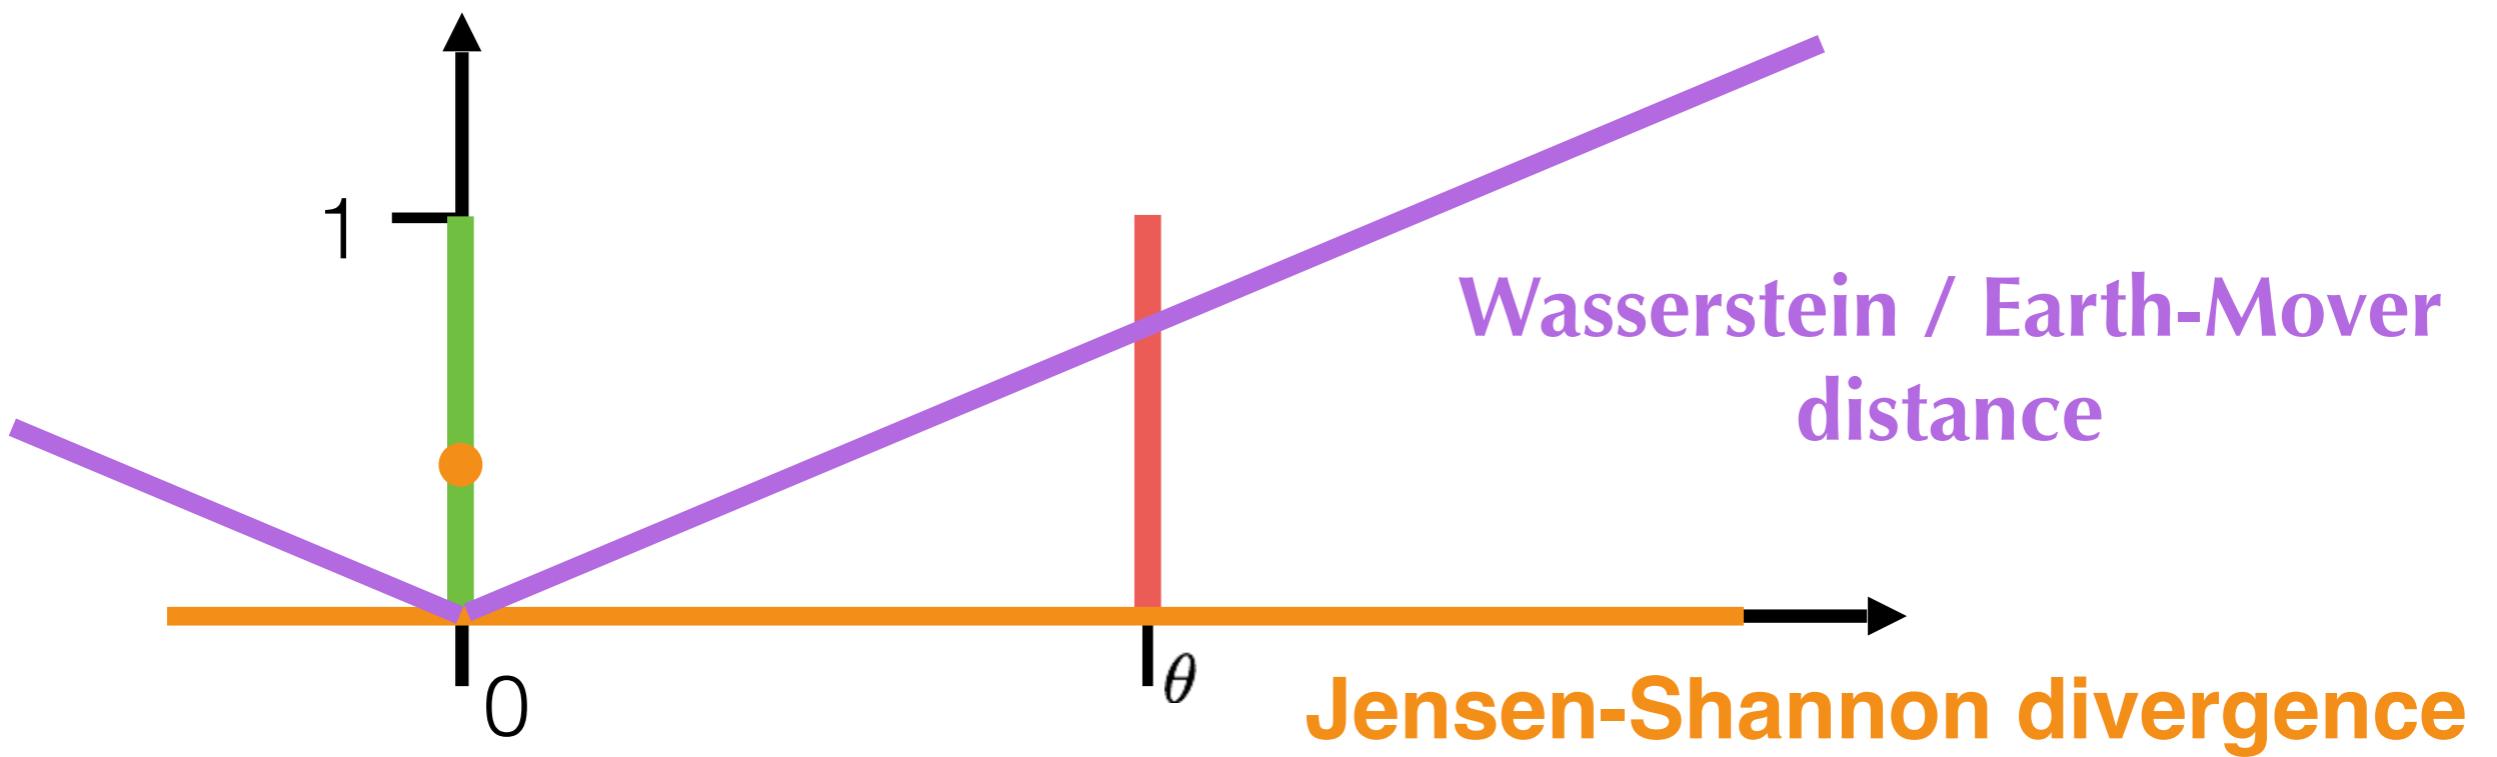


Simple one-parameter distribution  $p_\theta(z) = (\theta, z)$ , with  $\sim U[0, 1]$

Practical relevance: Operating on disjoint manifolds

# Why is Training GANs so Hard?

**Problem:** Jensen-Shannon divergence gradient not always useful → problem when using gradient descent

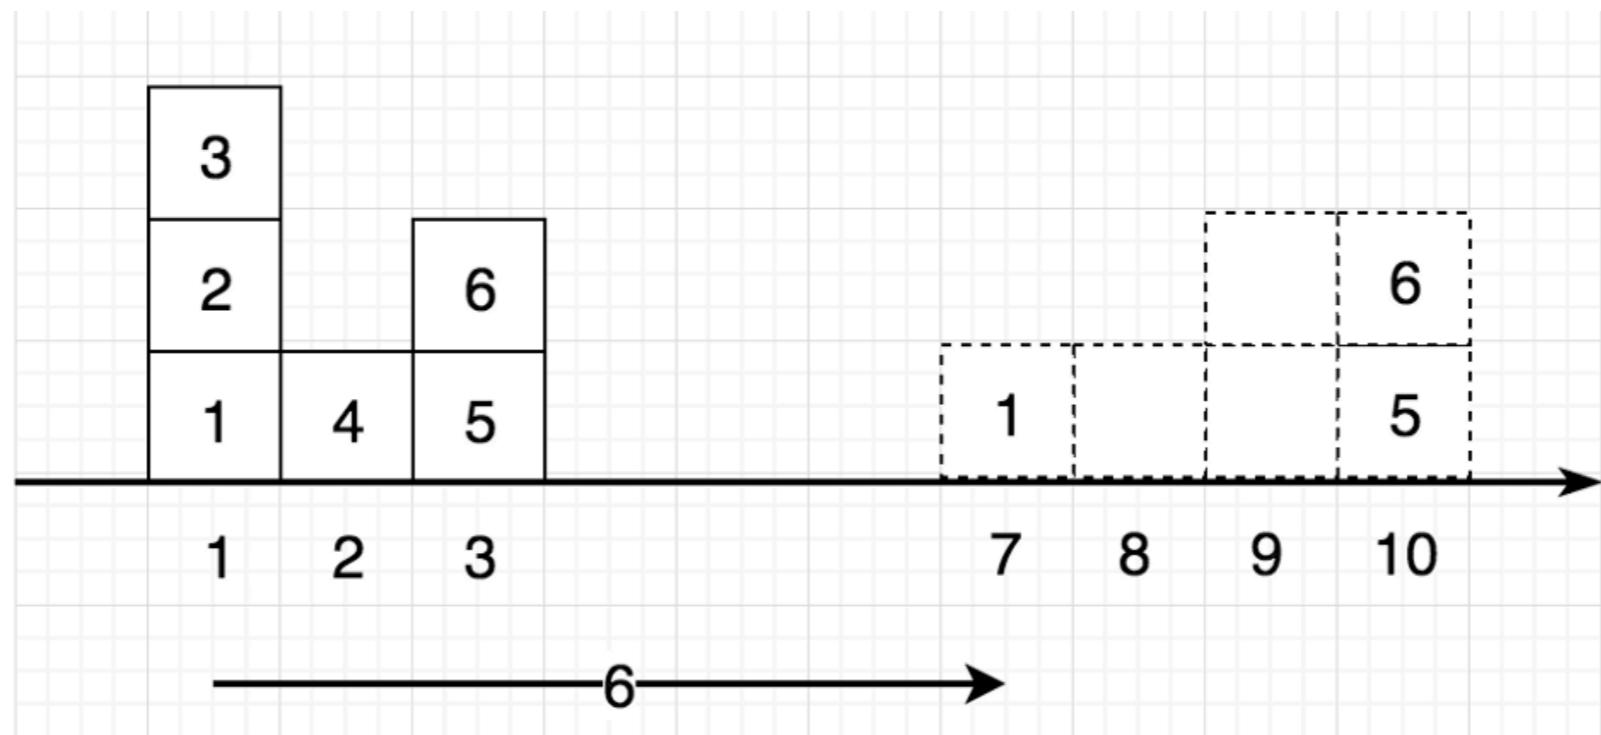


Simple one-parameter distribution  $p_\theta(z) = (\theta, z)$  , with  $\sim U[0, 1]$

Practical relevance: Operating on disjoint manifolds

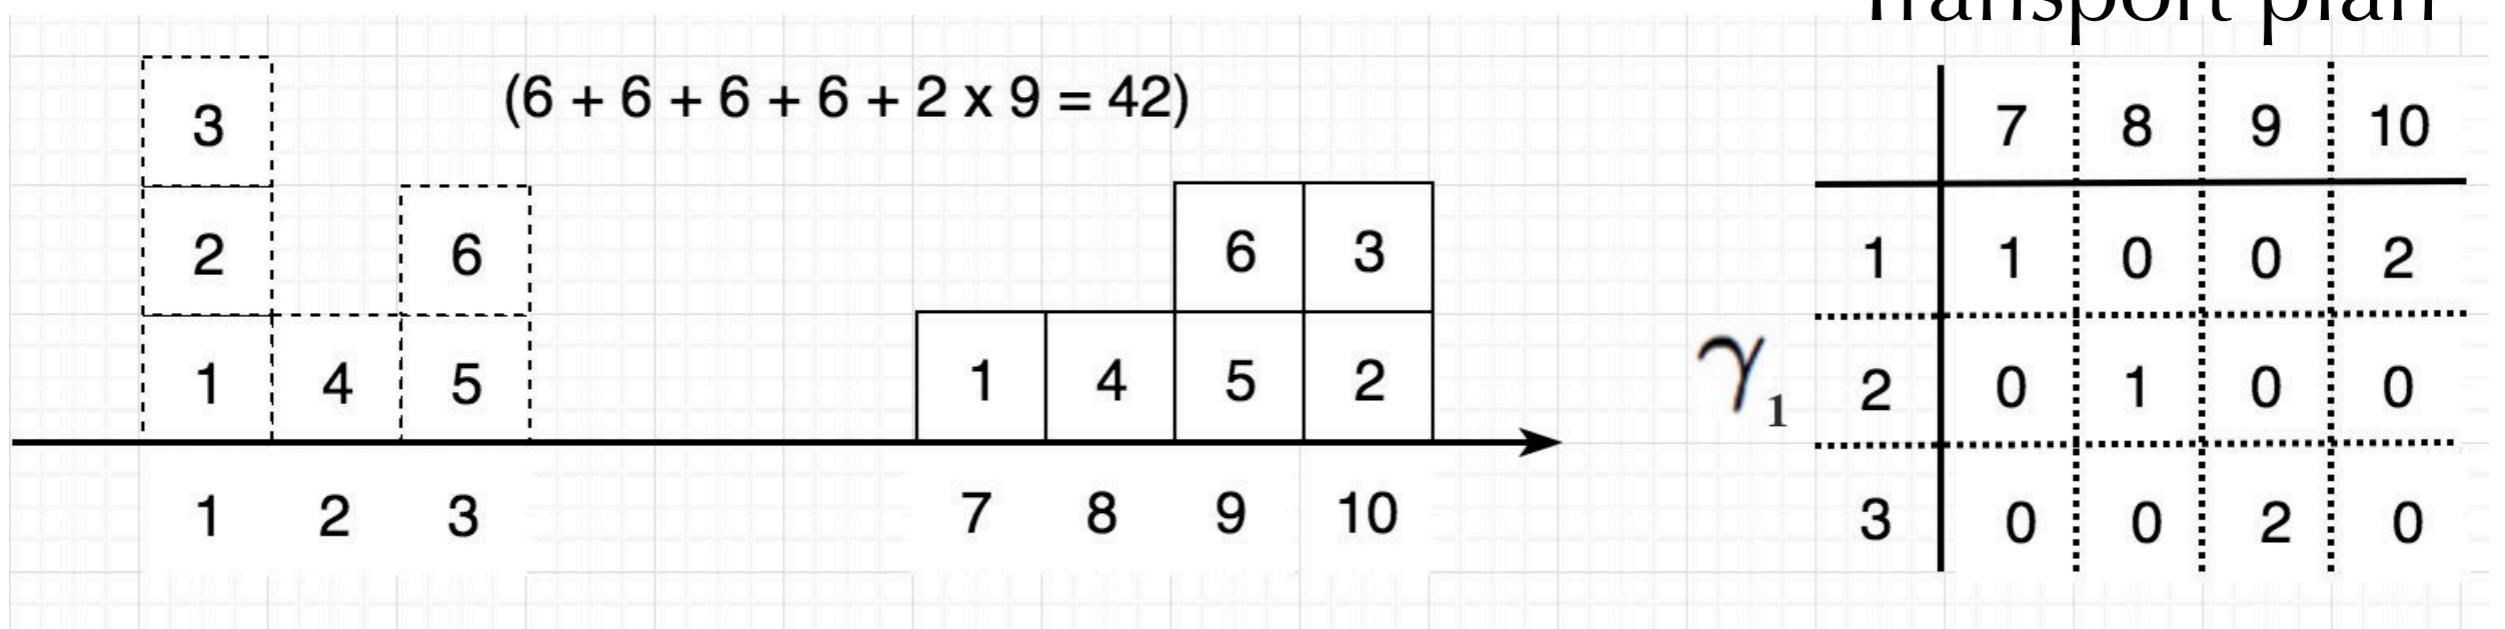
# Wasserstein GANs (WGANS)

**Earth-Mover distance:** Minimum amount of work to move weight from  $p(x)$  to  $q(x)$



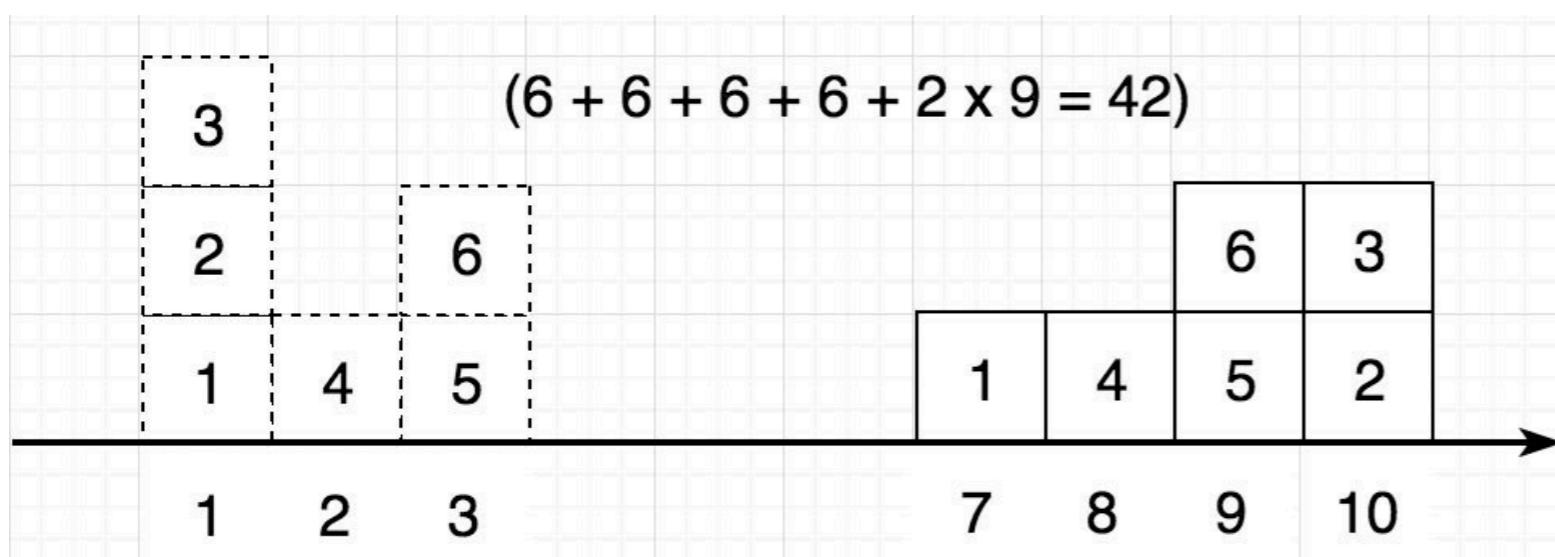
# Wasserstein GANs (WGANS)

**Earth-Mover distance:** Minimum amount of work to move weight from  $p(x)$  to  $q(x)$



# Wasserstein GANs (WGANS)

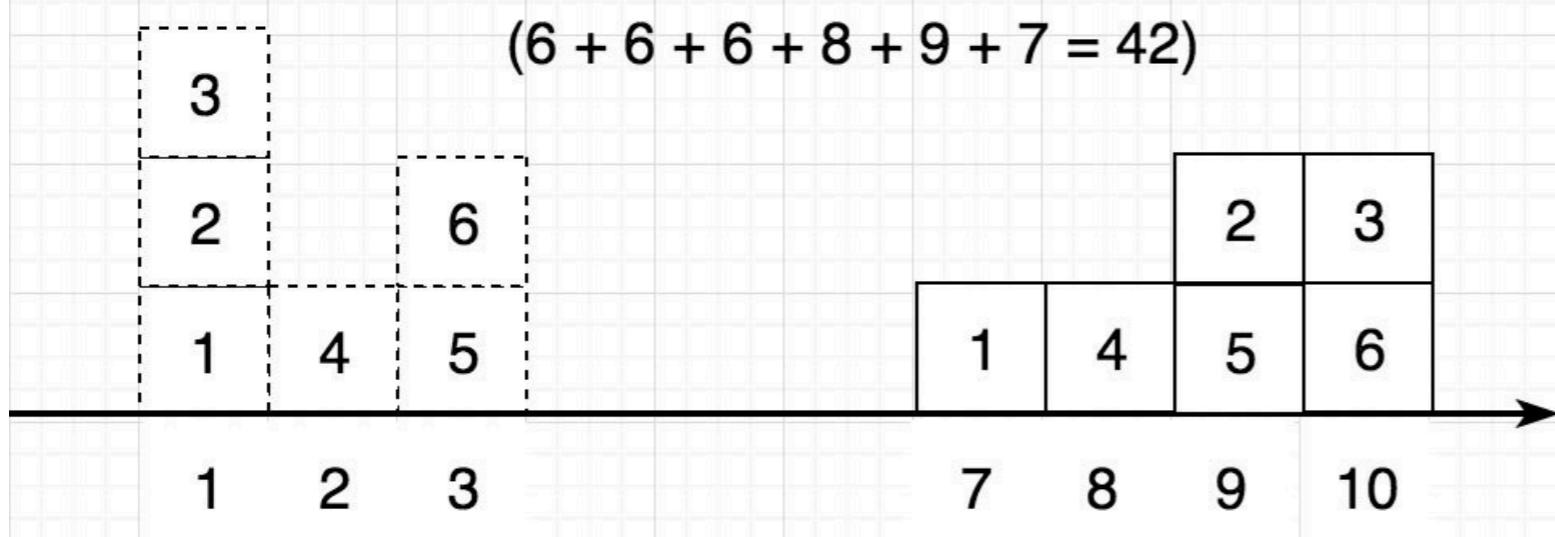
**Earth-Mover distance:** Minimum amount of work to move weight from  $p(x)$  to  $q(x)$



Transport plan

	7	8	9	10
1	1	0	0	2
2	0	1	0	0
3	0	0	2	0

$\gamma_1$



	7	8	9	10
1	1	0	1	1
2	0	1	0	0
3	0	0	1	1

$\gamma_2$

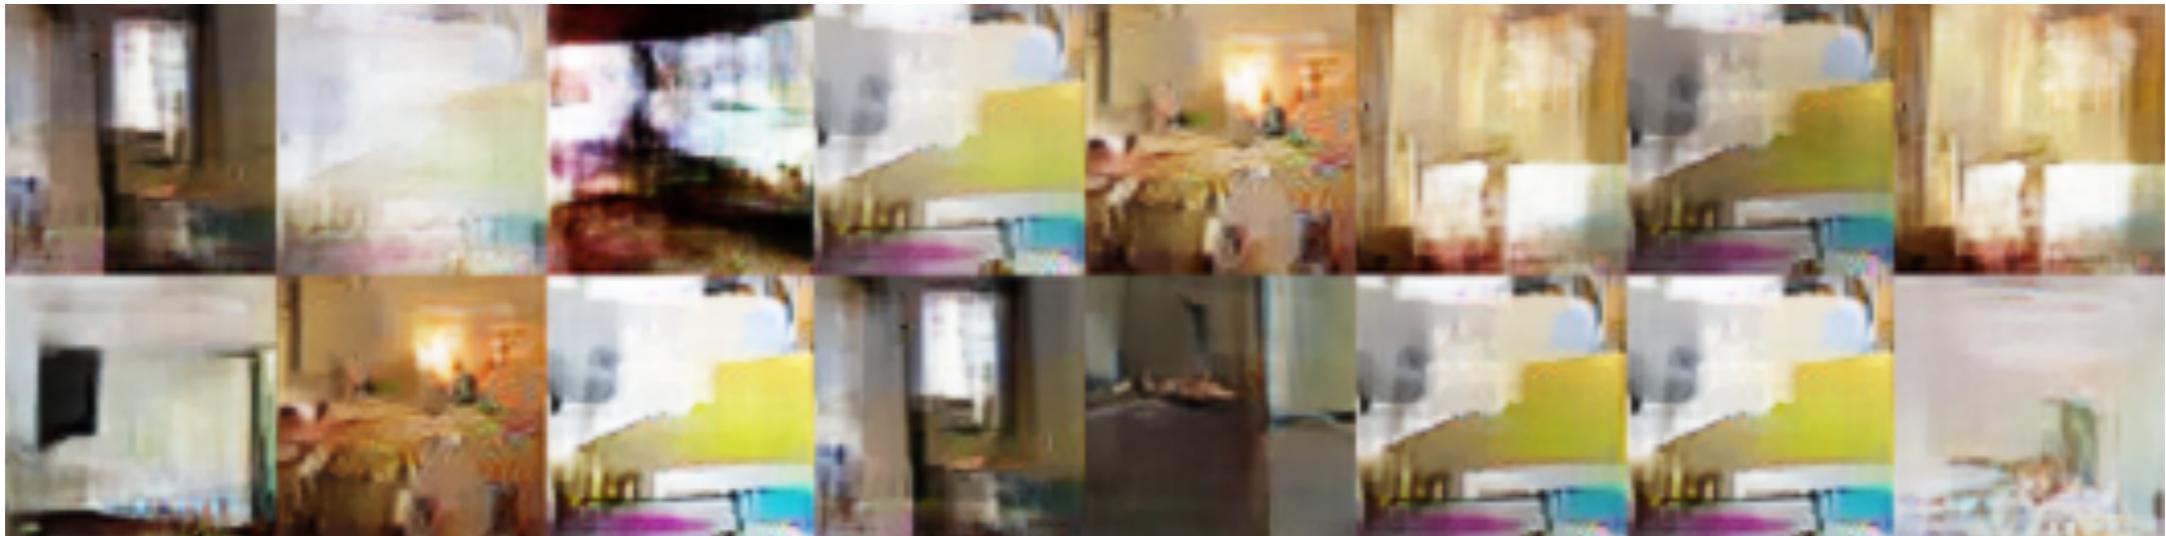
# Wasserstein GANs (WGANS)

**Earth-Mover distance:** Minimum amount of work to move weight from  $p(x)$  to  $q(x)$

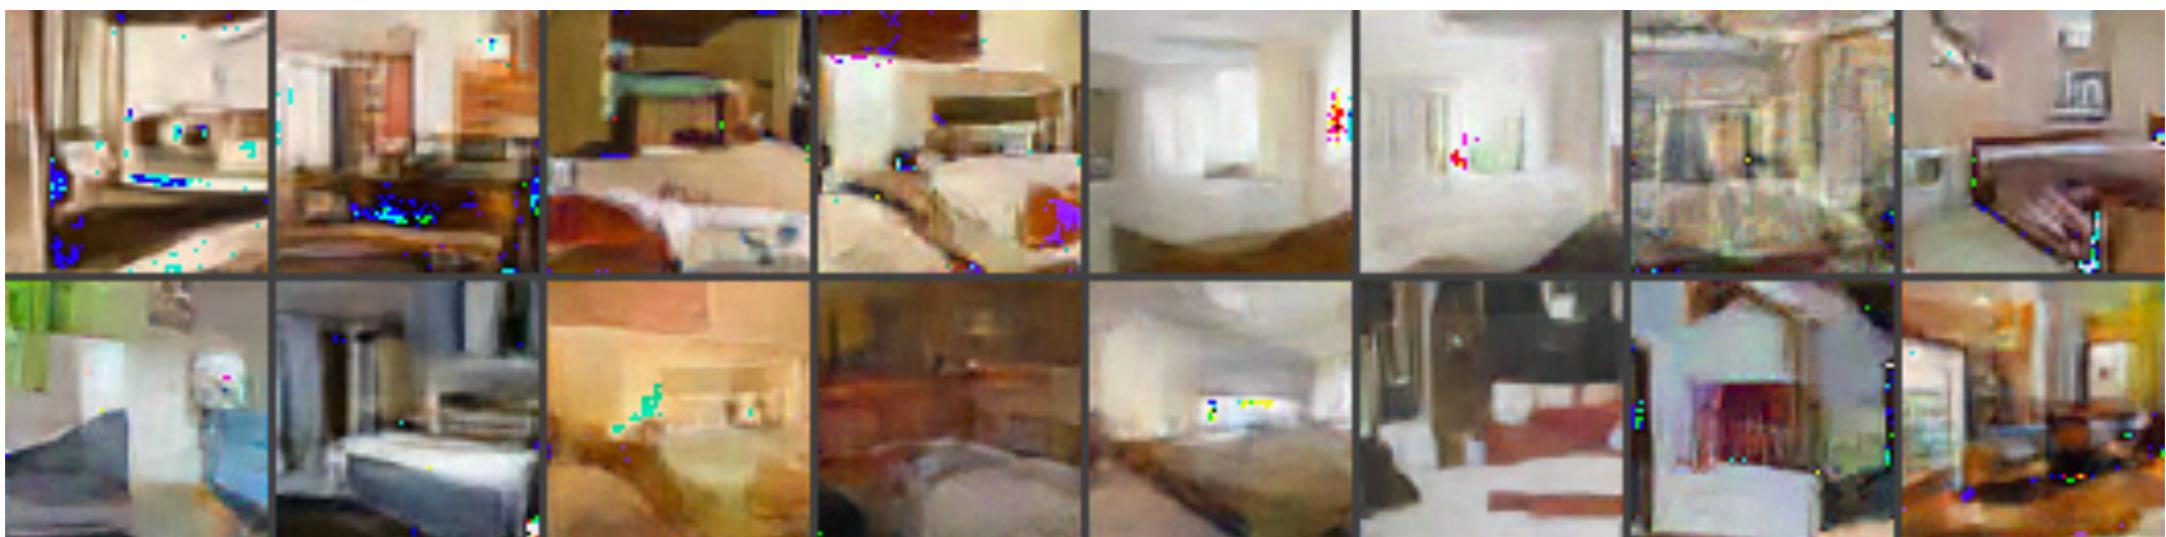
$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

# WGANS vs. GANs

GAN



WGAN



[Arjovsky et al., Wasserstein GAN, arXiv:1701.07875]

# State-of-the-art Results

## PROGRESSIVE GROWING OF GANs FOR IMPROVED QUALITY, STABILITY, AND VARIATION

Tero Karras  
NVIDIA

Timo Aila  
NVIDIA

Samuli Laine  
NVIDIA

Jaakko Lehtinen  
NVIDIA  
Aalto University



# State-of-the-art Results

## PROGRESSIVE GROWING OF GANs FOR IMPROVED QUALITY, STABILITY, AND VARIATION

Tero Karras  
NVIDIA

Timo Aila  
NVIDIA

Samuli Laine  
NVIDIA

Jaakko Lehtinen  
NVIDIA  
Aalto University

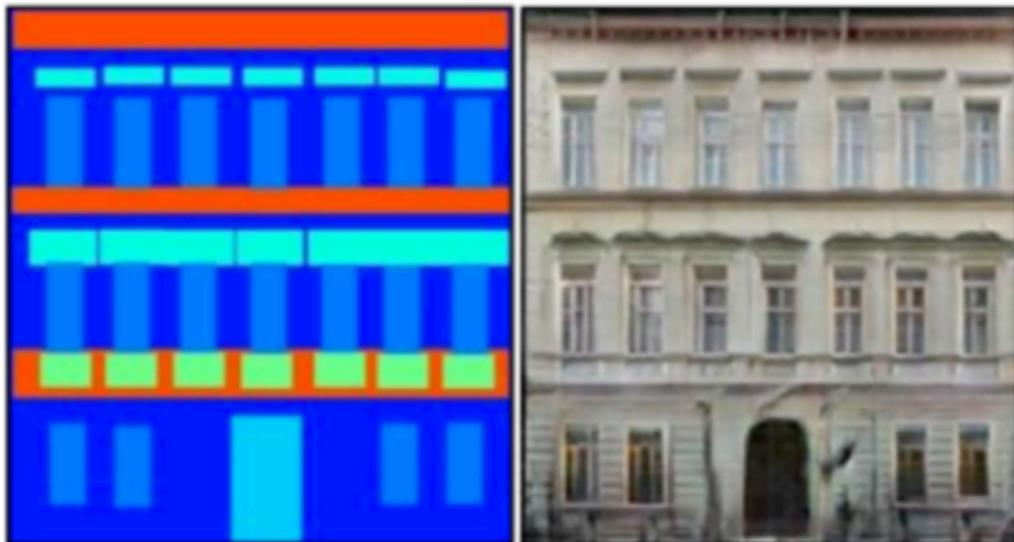


# Today

- Variational Autoencoders (VAE)
- Generative Adversarial Networks (GANs)
- **Unsupervised Image Translation**

# Image-To-Image Translation

## Labels to Facade

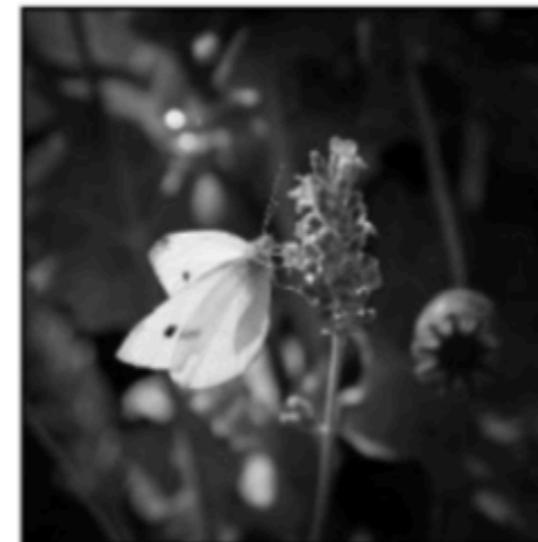


## input

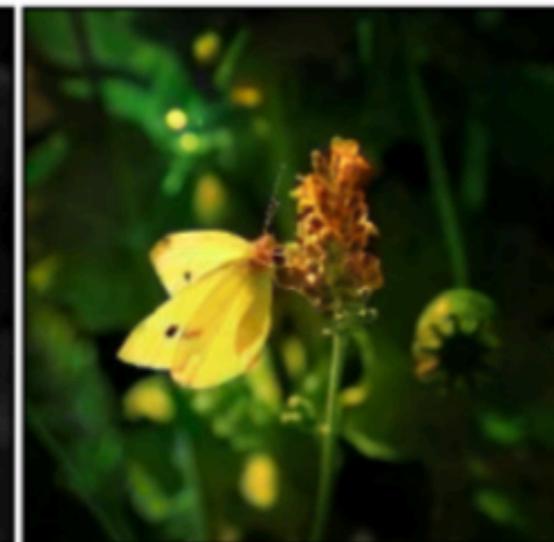
## output

## Day to Night

## BW to Color



## input



## output

A scenic view of a winding road through a forested mountain area. The road curves to the right, with a yellow center line. The ground is covered in patches of snow. In the background, there are tall evergreen trees on the left and a town with houses and buildings nestled among trees on a hillside under a cloudy sky.

## input



## output

## Edges to Photo



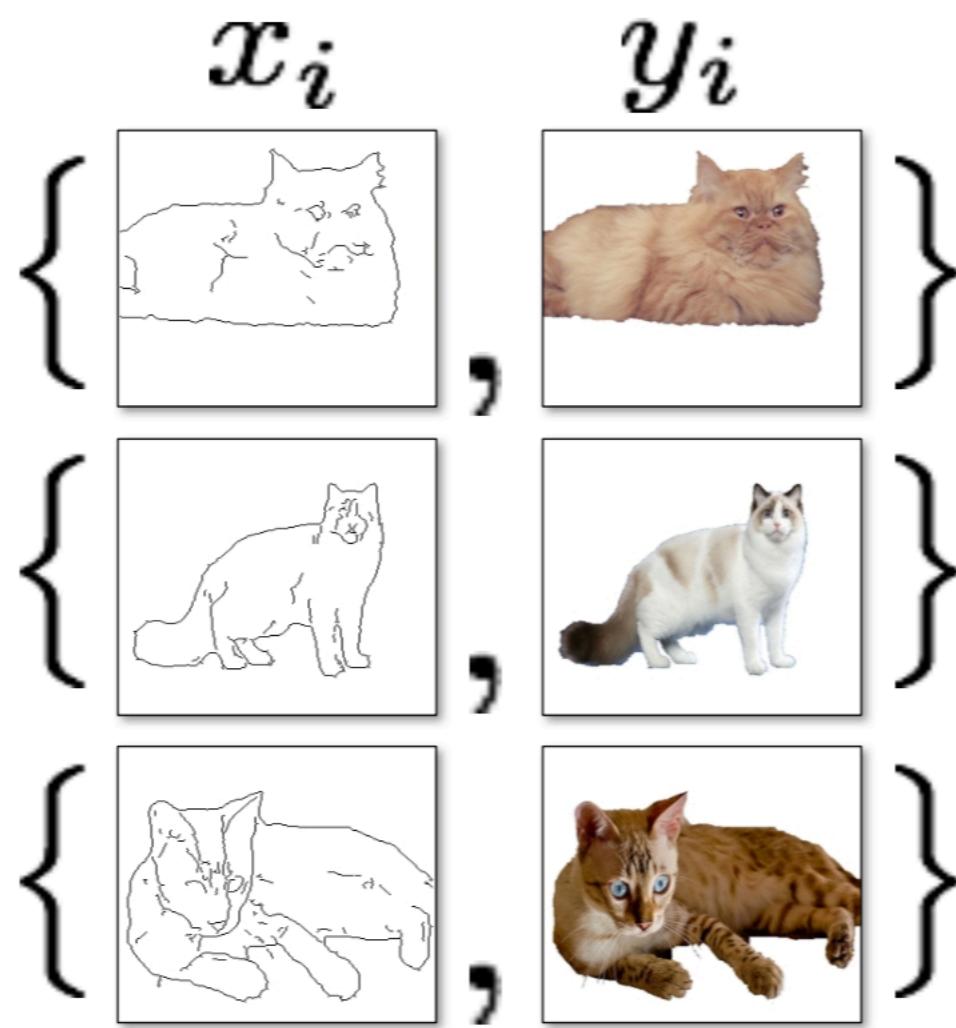
## input



## output

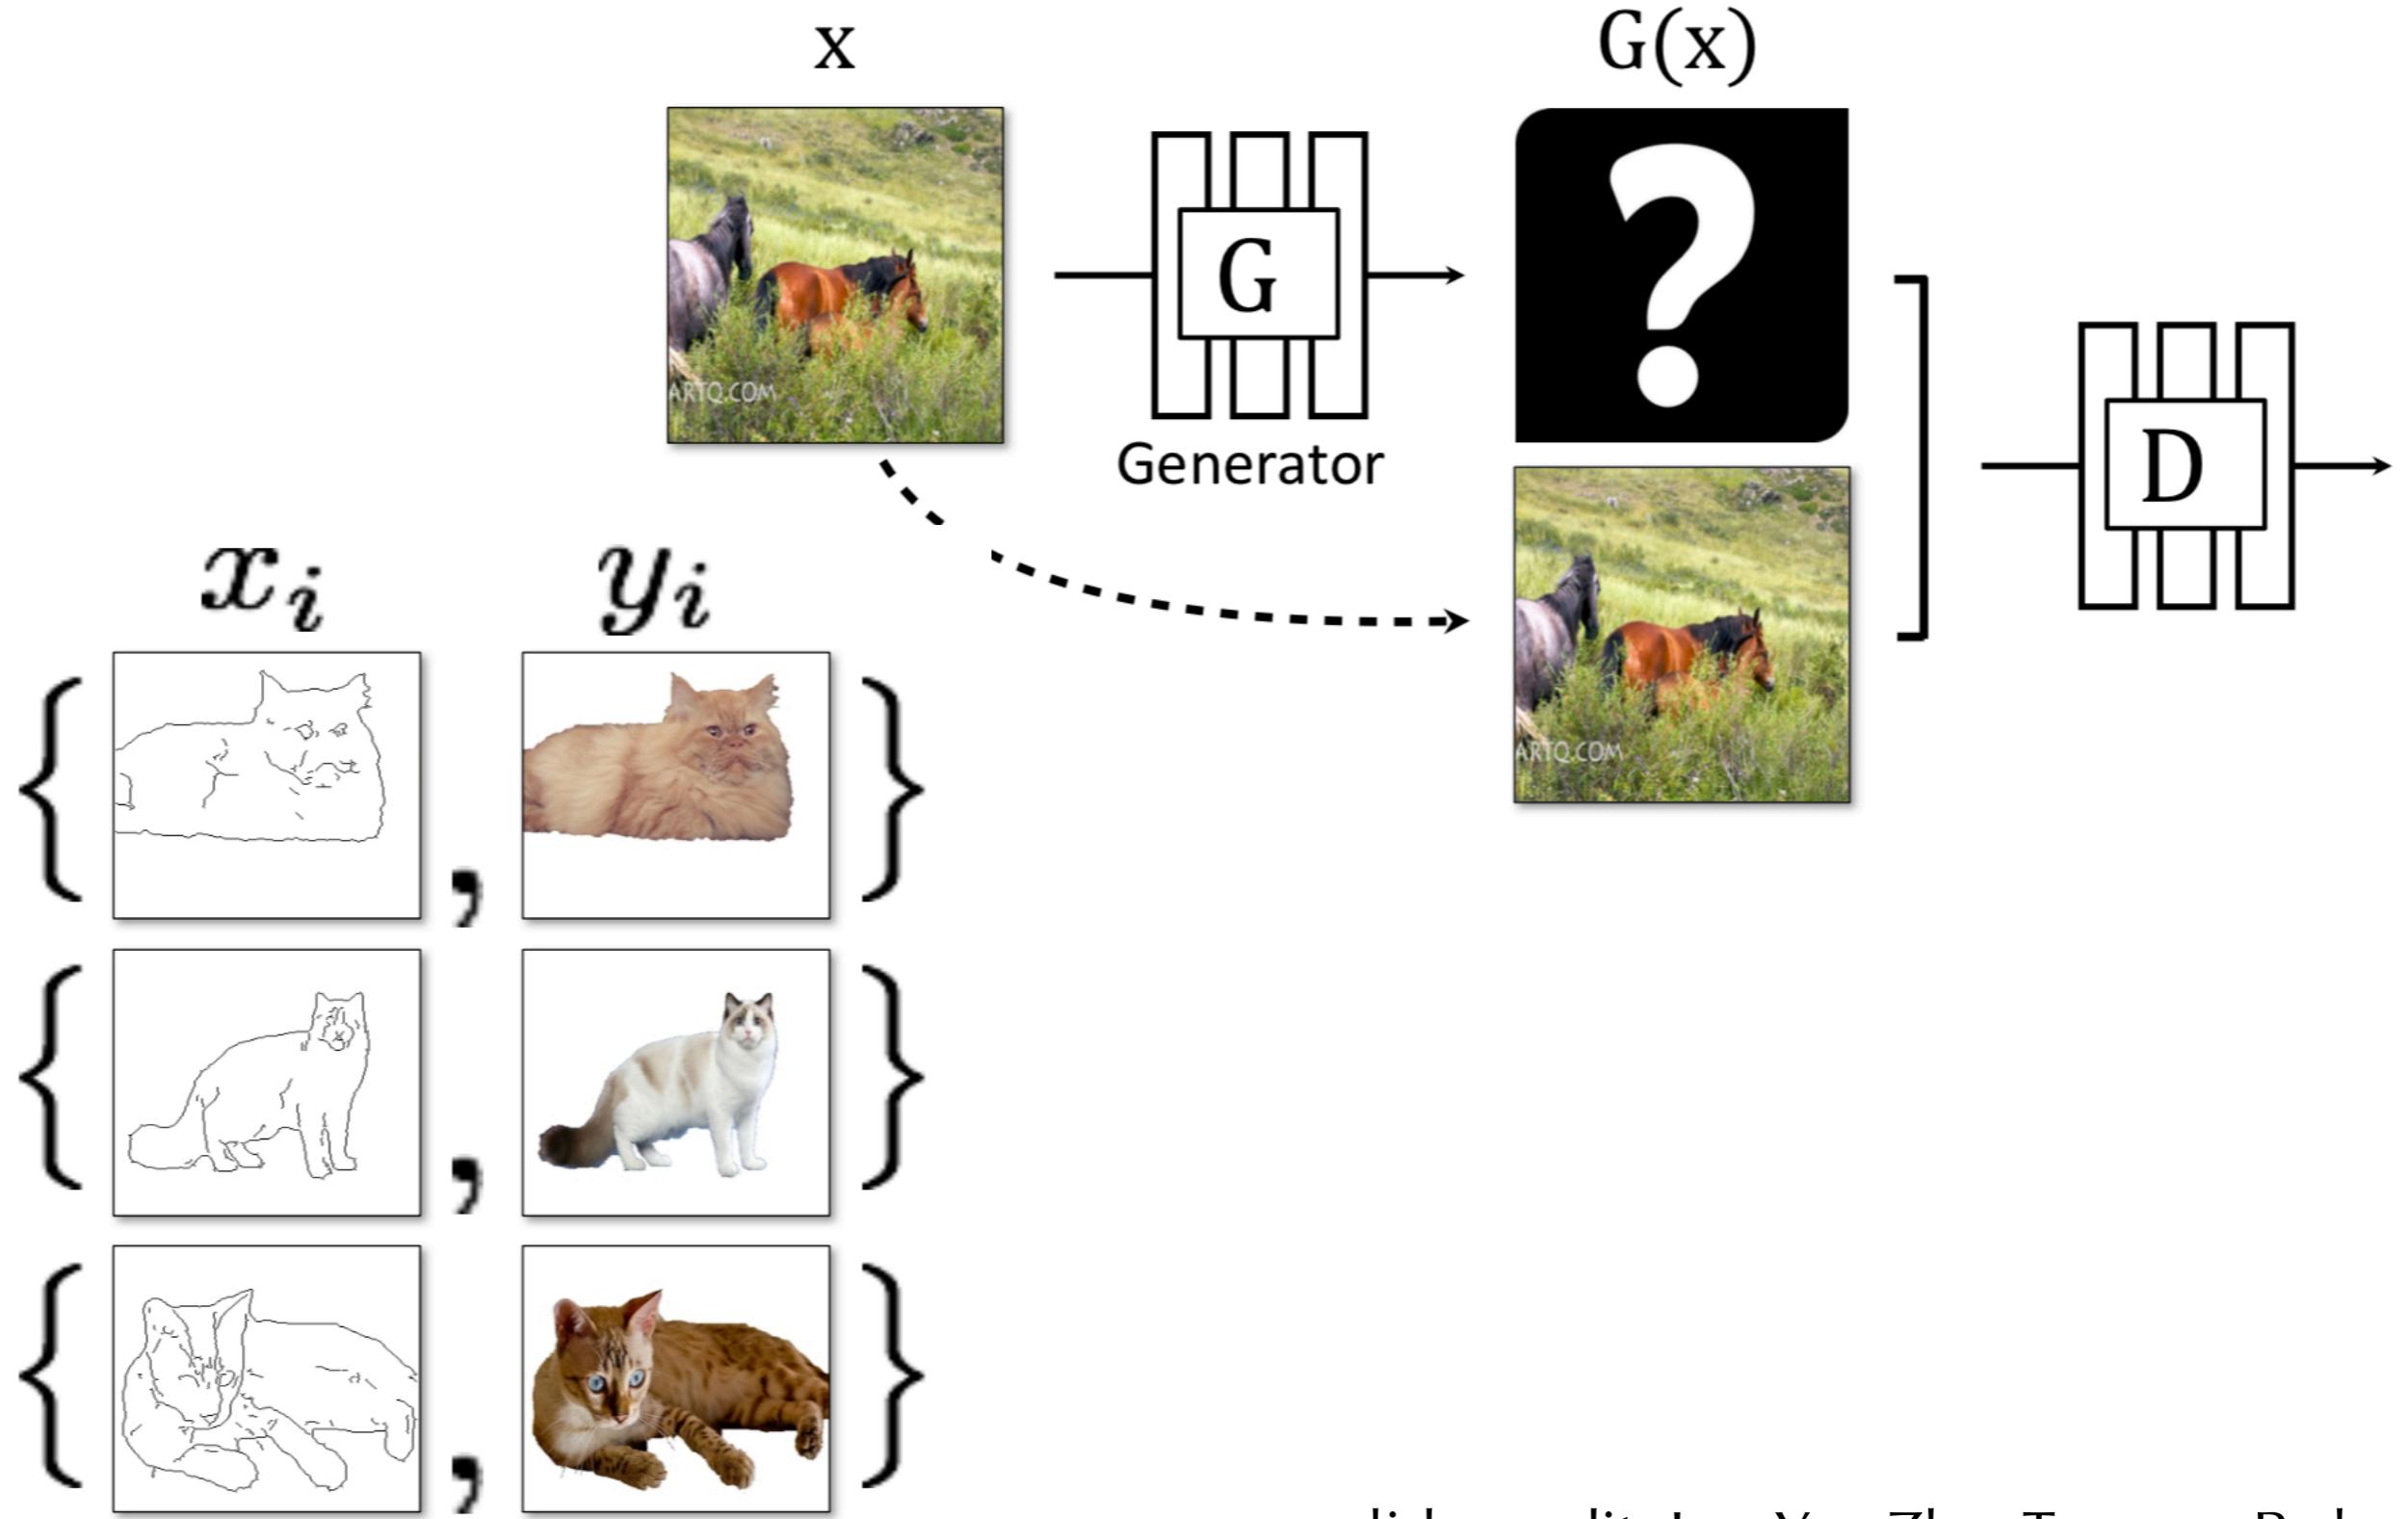
slide credit: Jun-Yan Zhu, Taesung Park

# Supervised Image-To-Image Translation



slide credit: Jun-Yan Zhu, Taesung Park

# Supervised Image-To-Image Translation



slide credit: Jun-Yan Zhu, Taesung Park

# Supervised Image-To-Image Translation



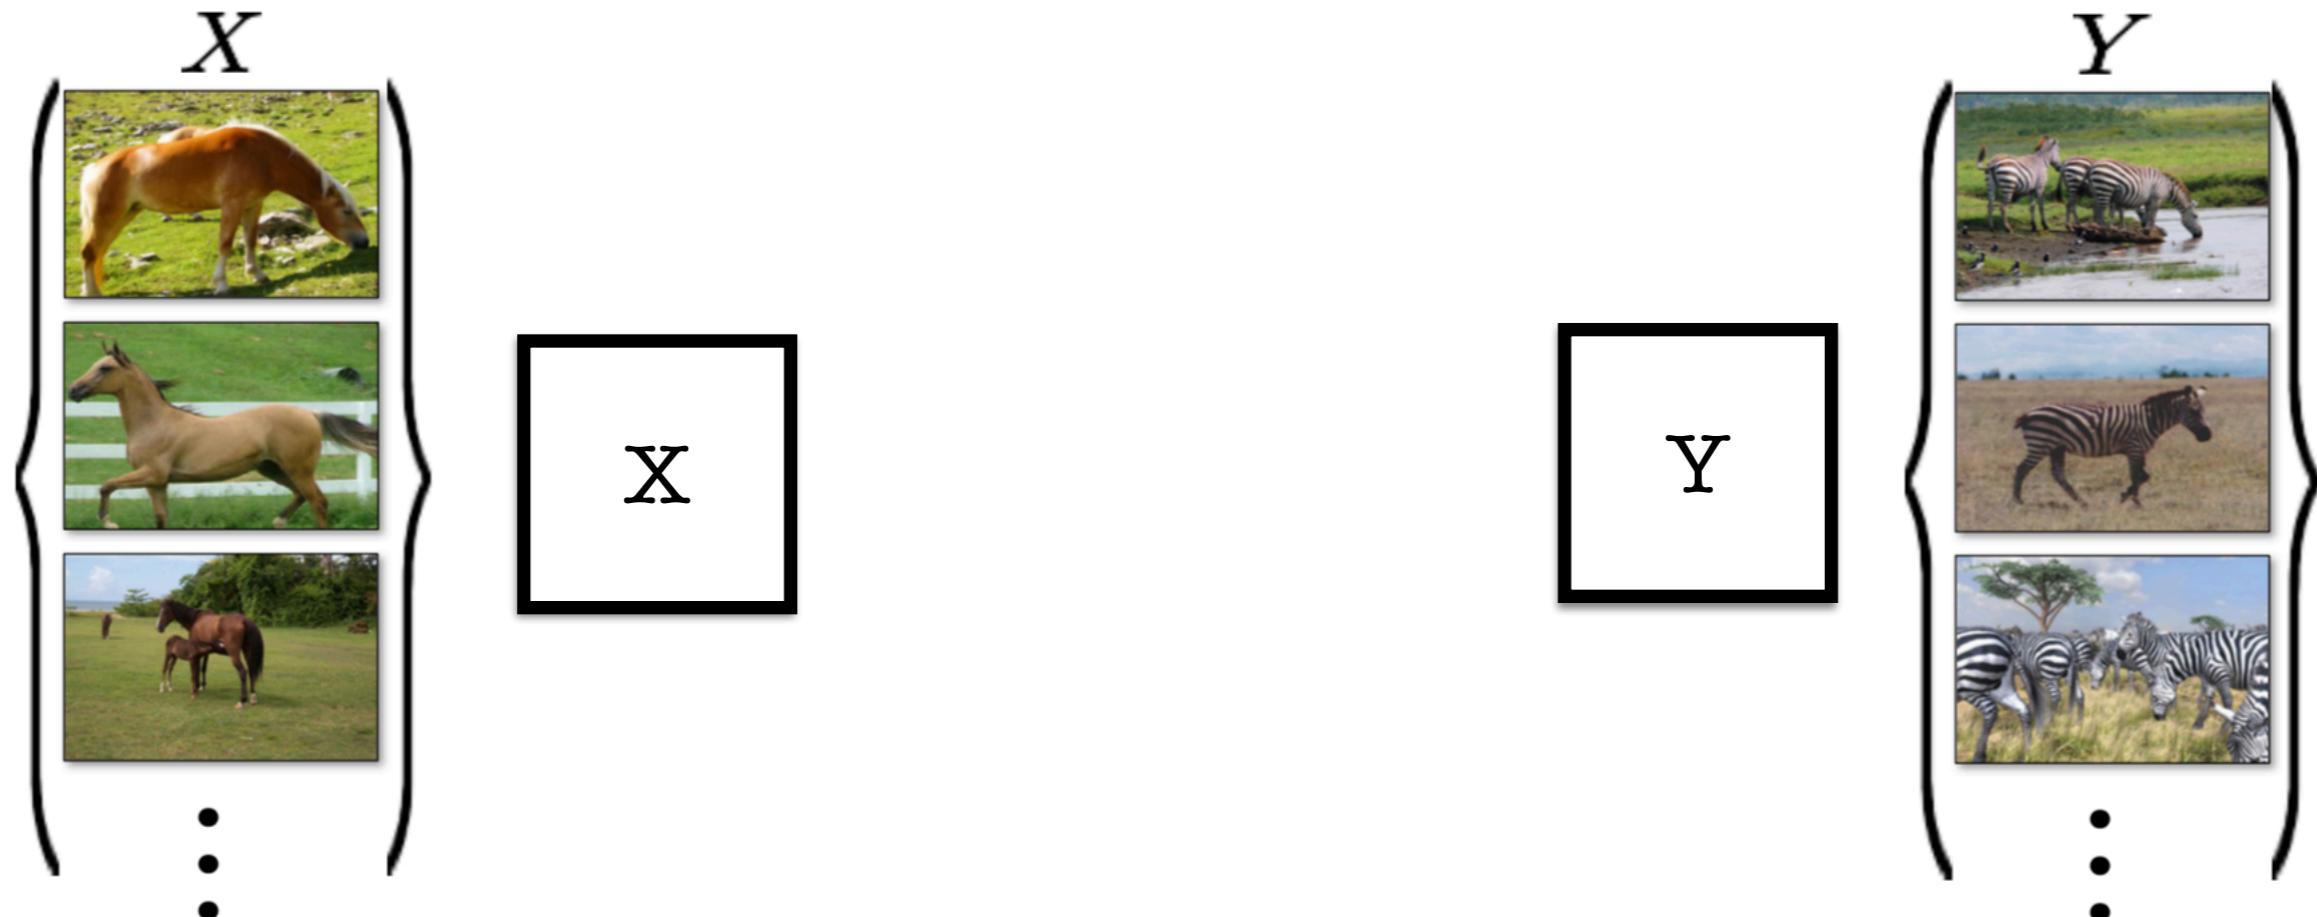
Label  $\leftrightarrow$  photo: per-pixel labeling



Labelled data often hard or impossible to obtain

slide credit: Jun-Yan Zhu, Taesung Park

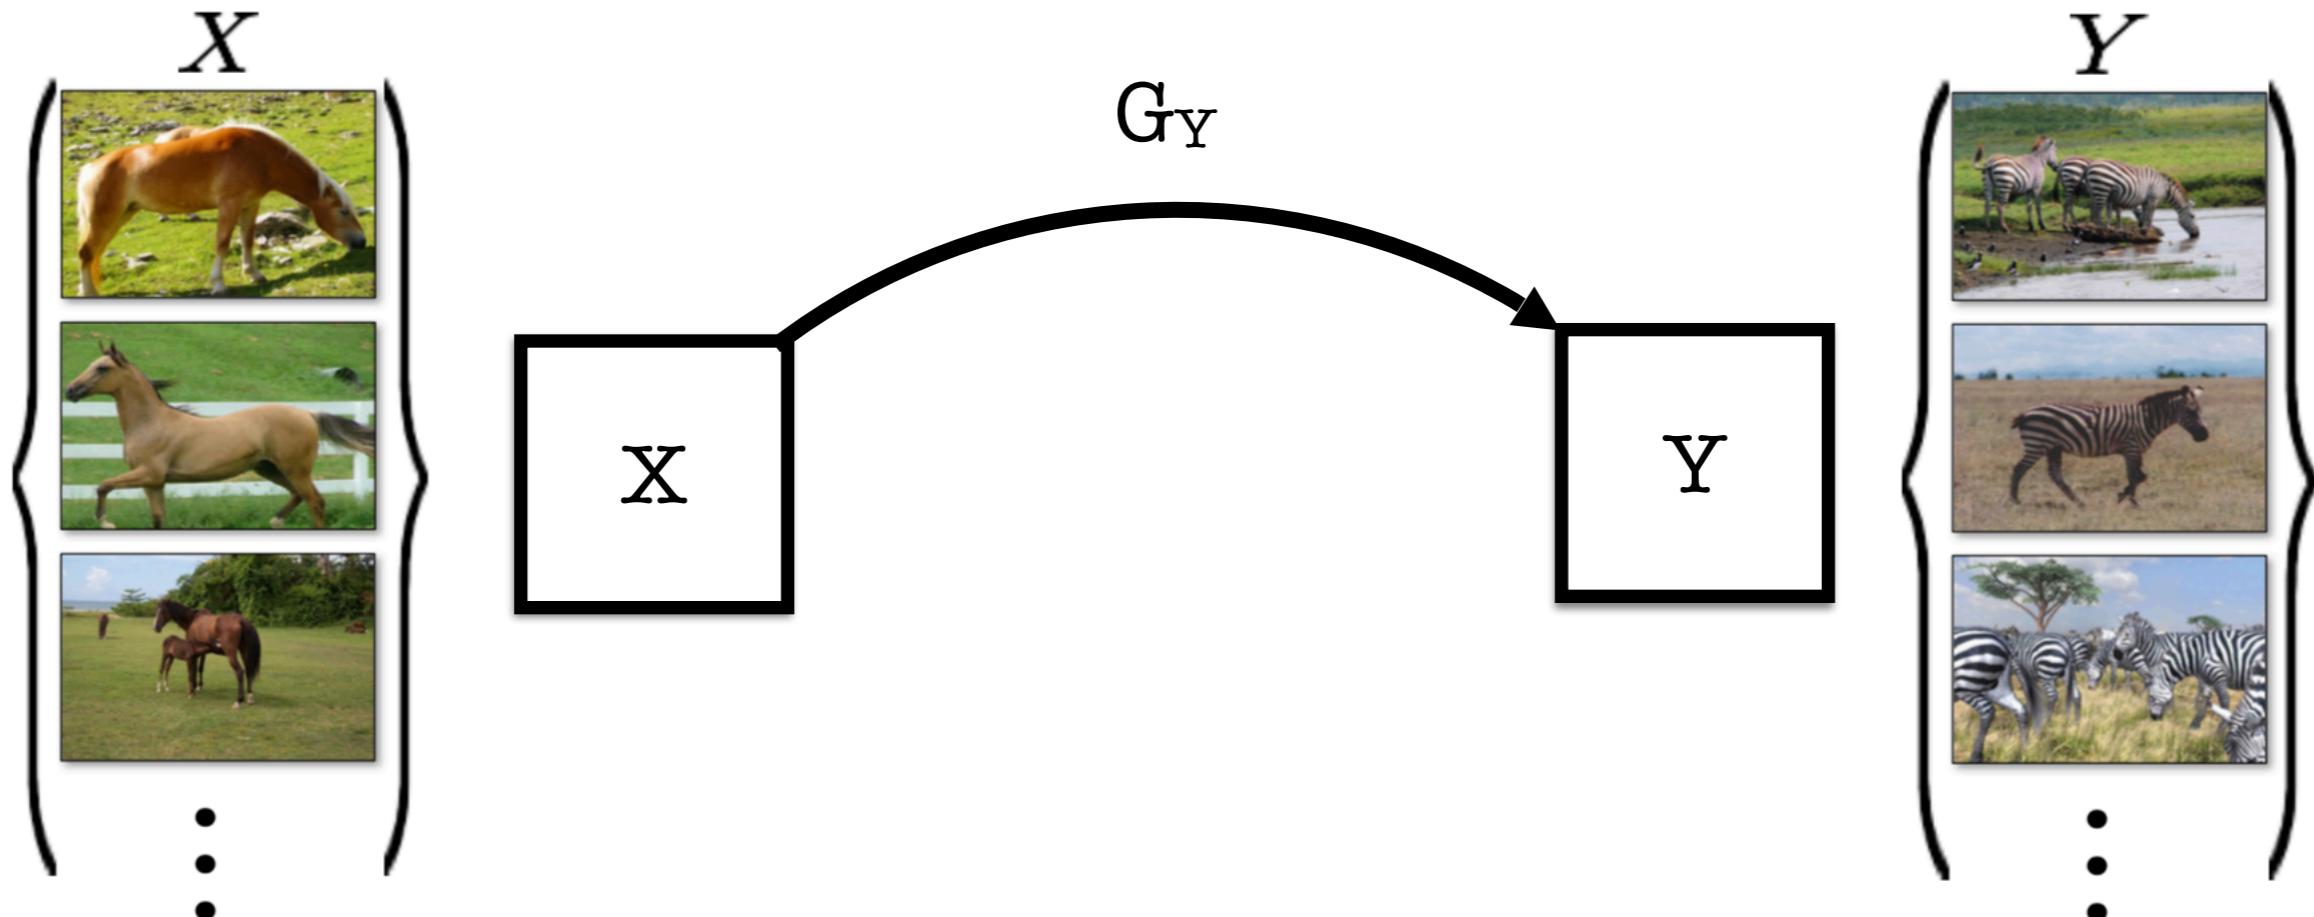
# CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-  
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

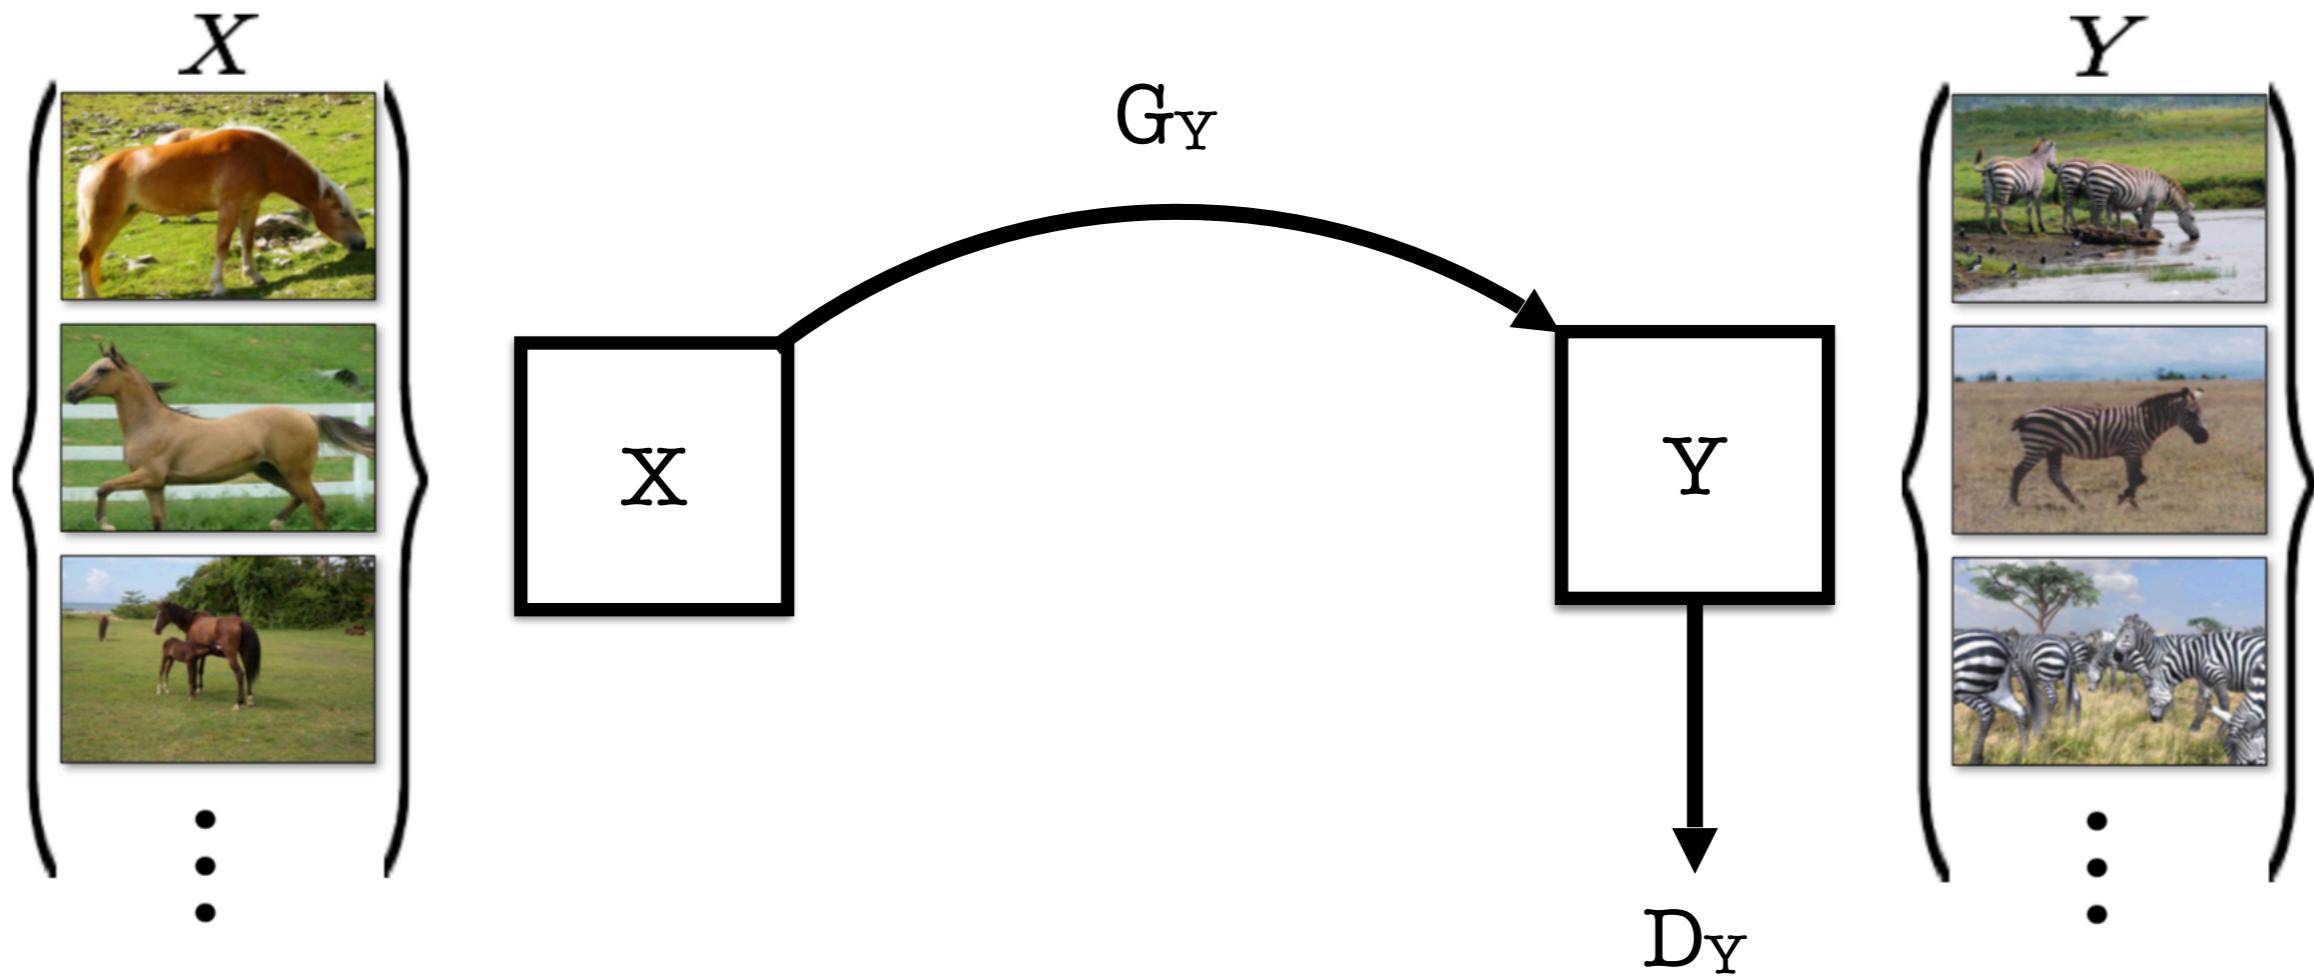
# CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-  
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

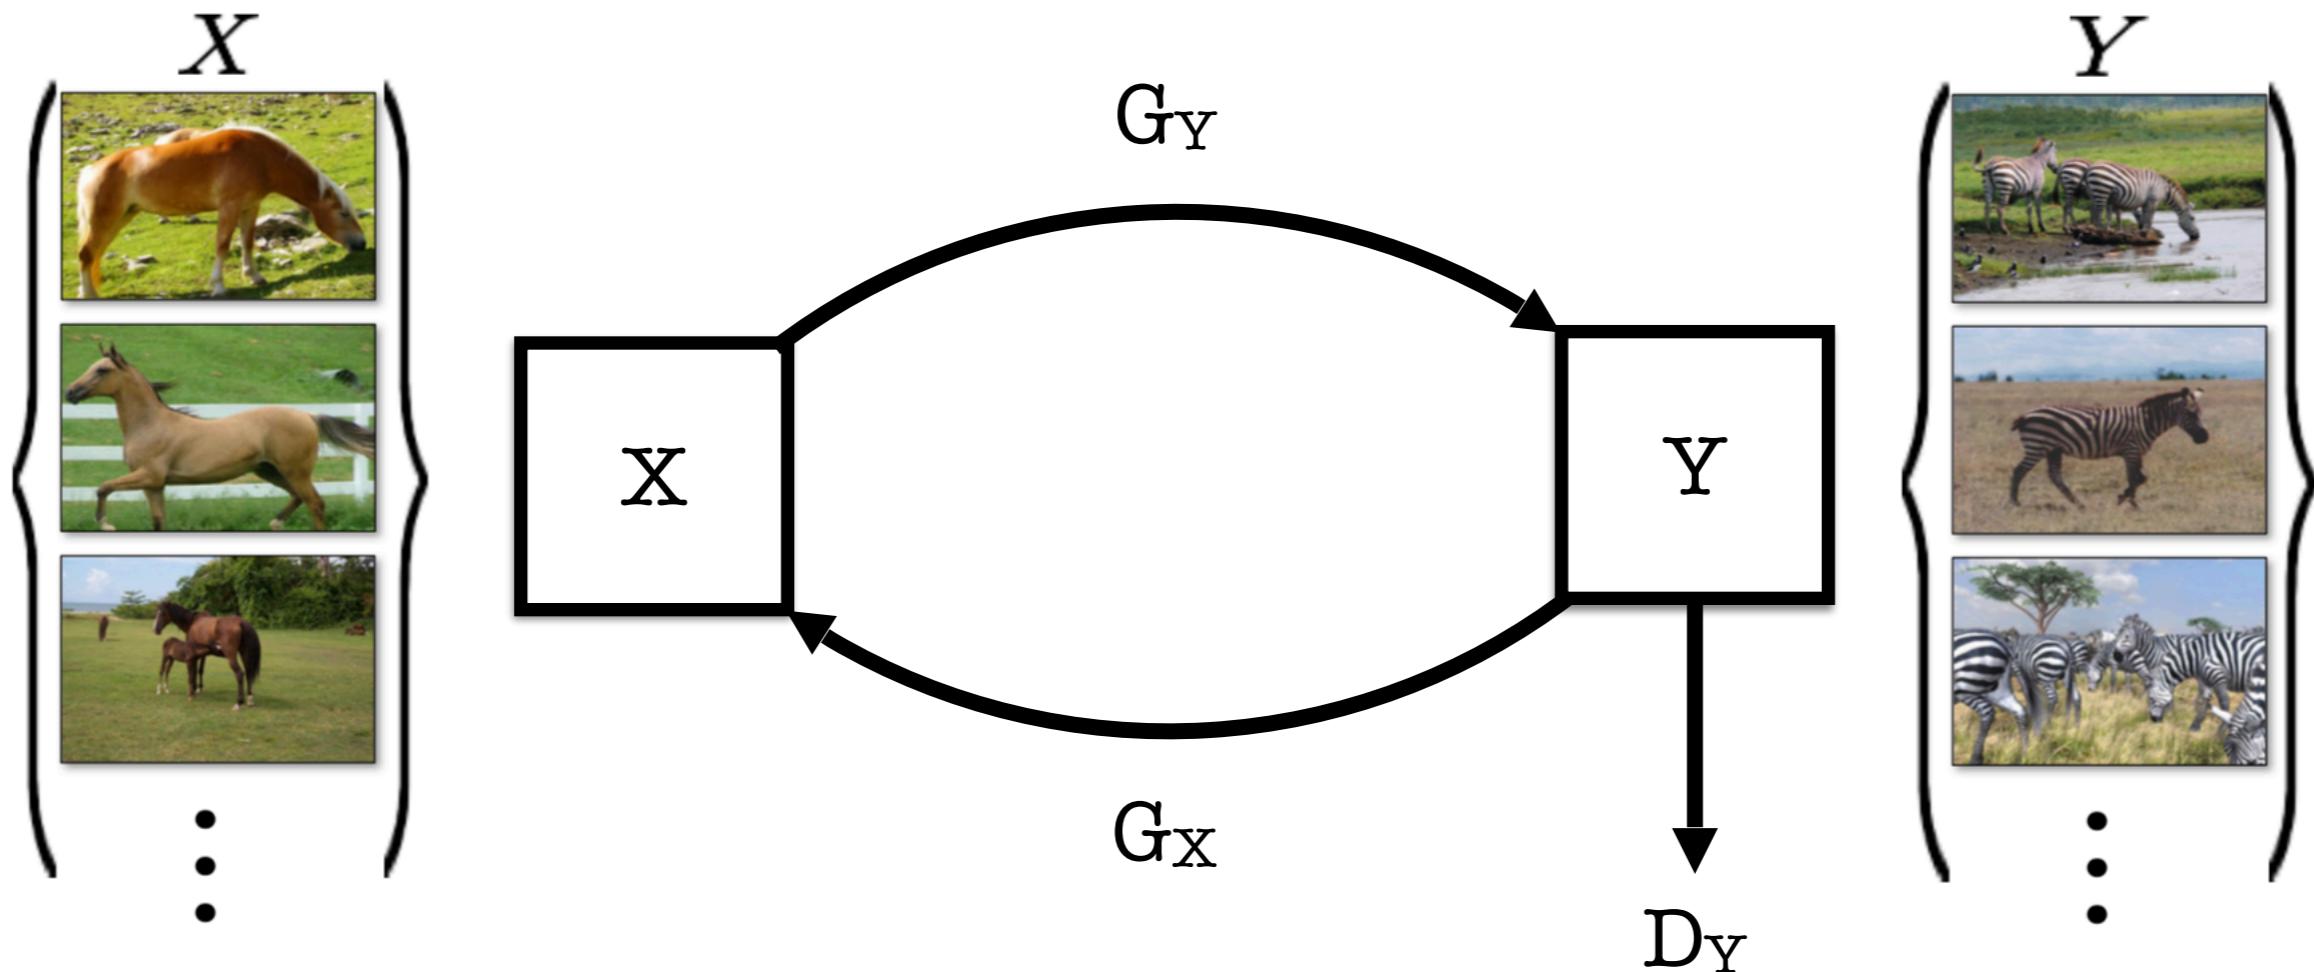
# CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-  
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

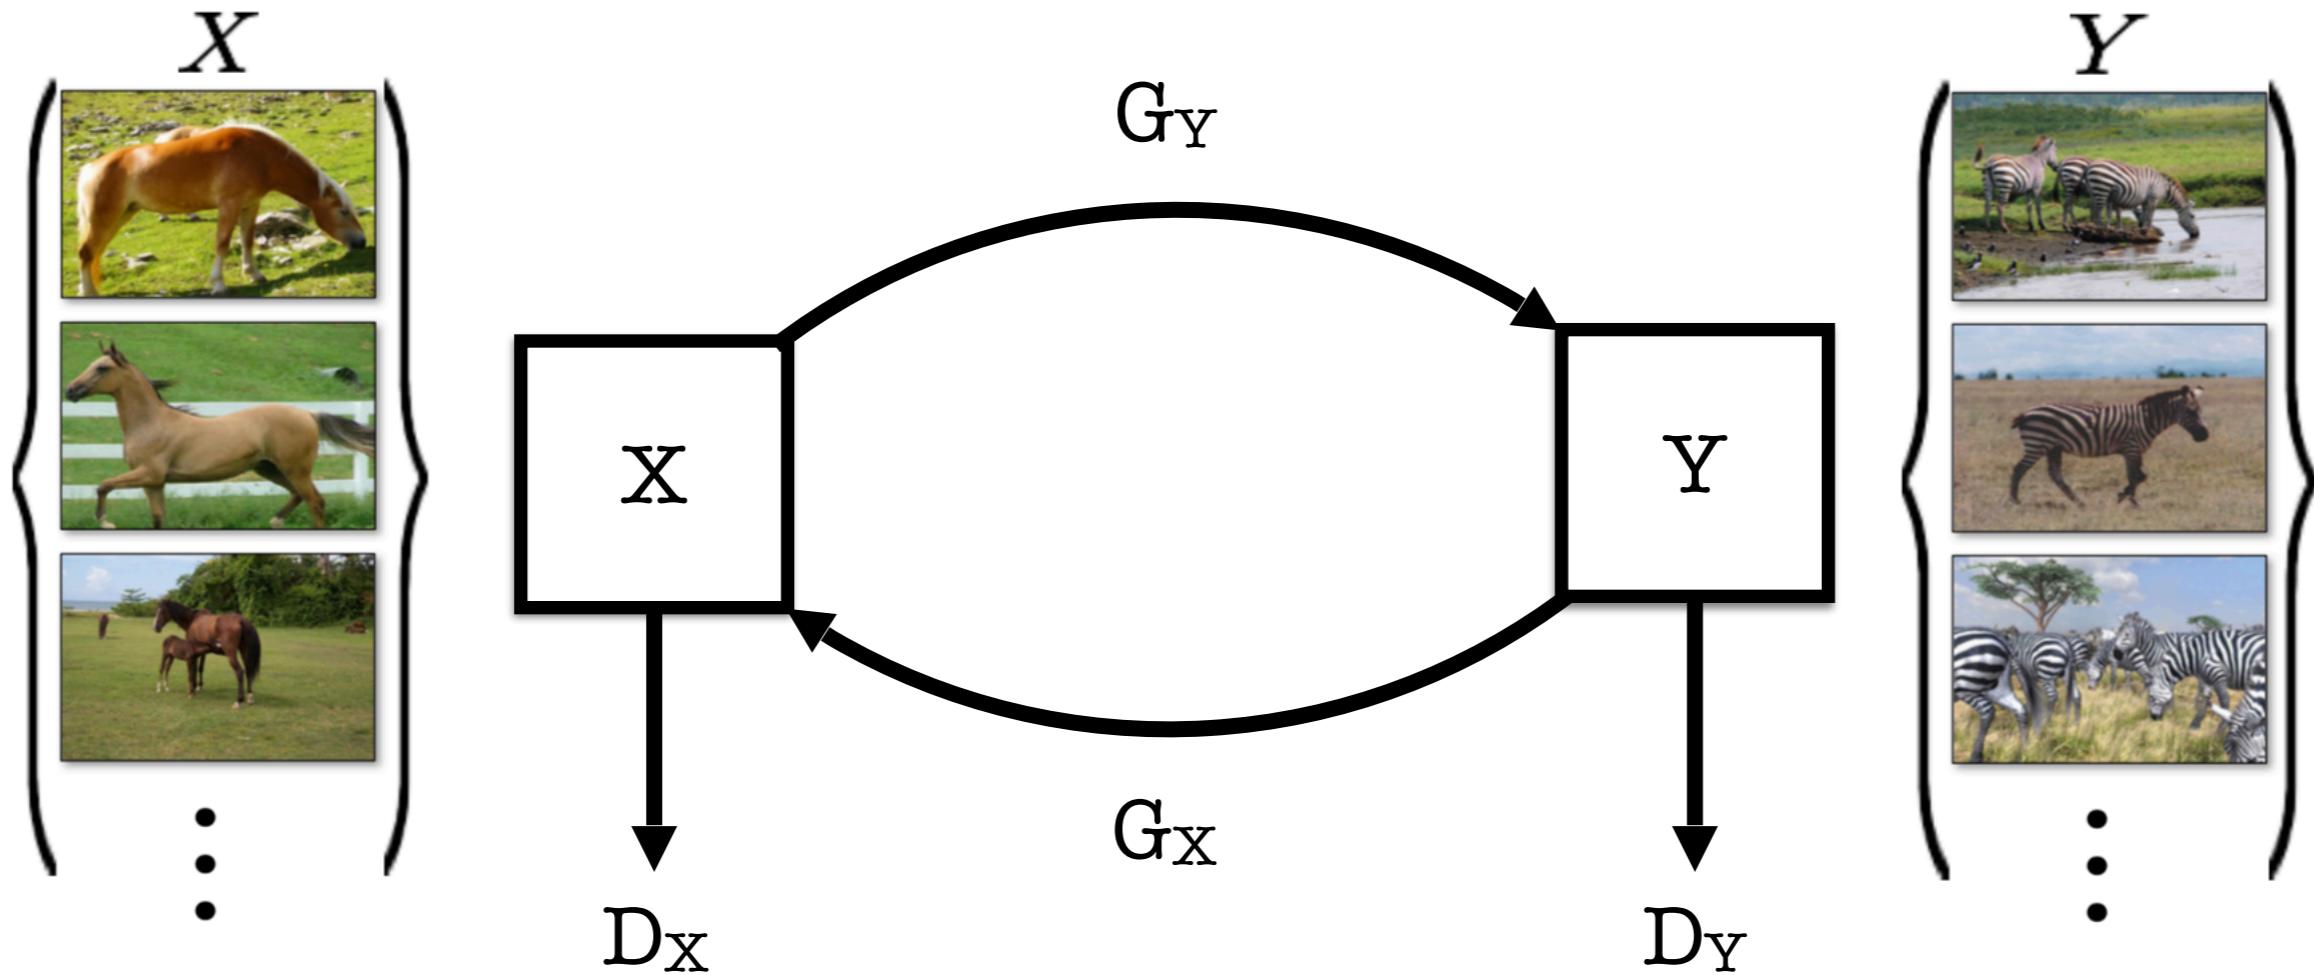
# CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-  
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

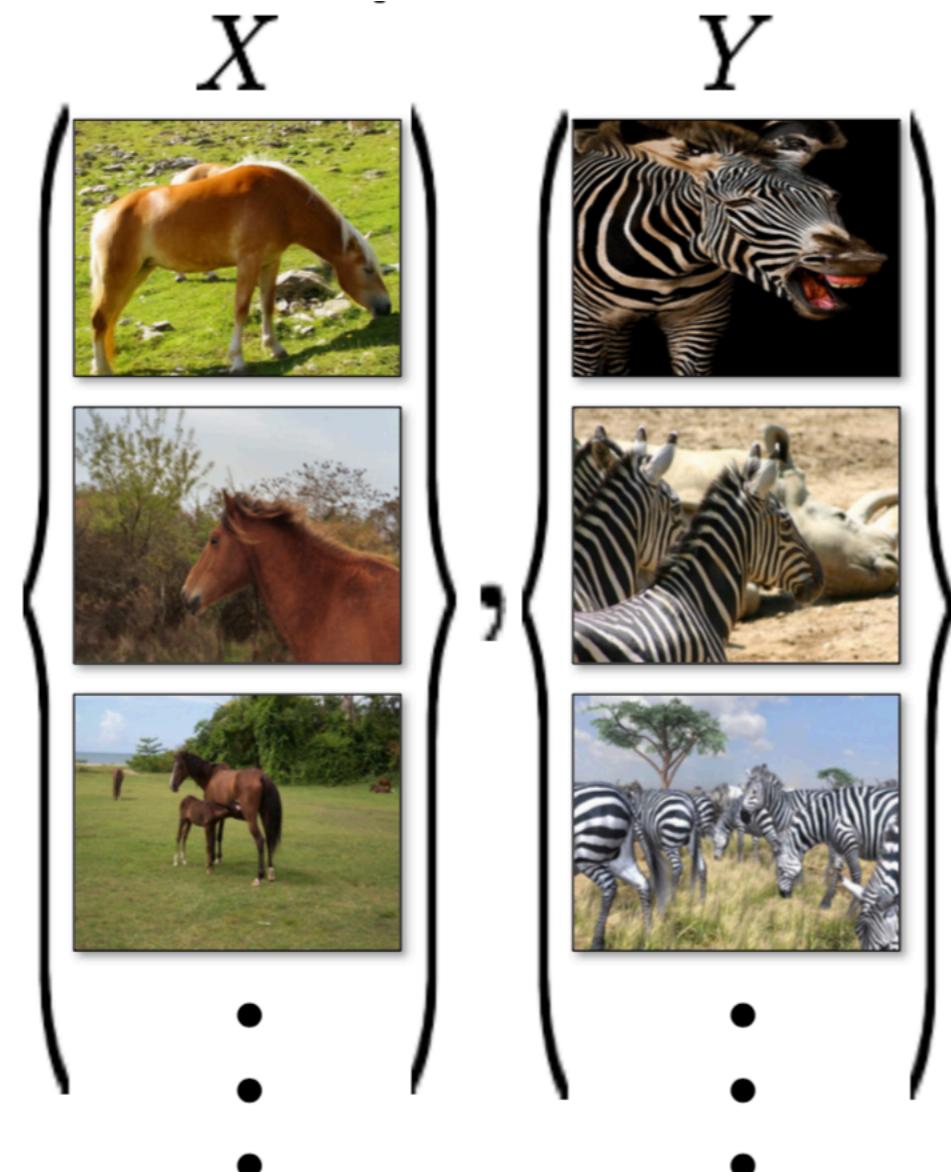
# CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-  
Consistent Adversarial Networks, ICCV 2017]

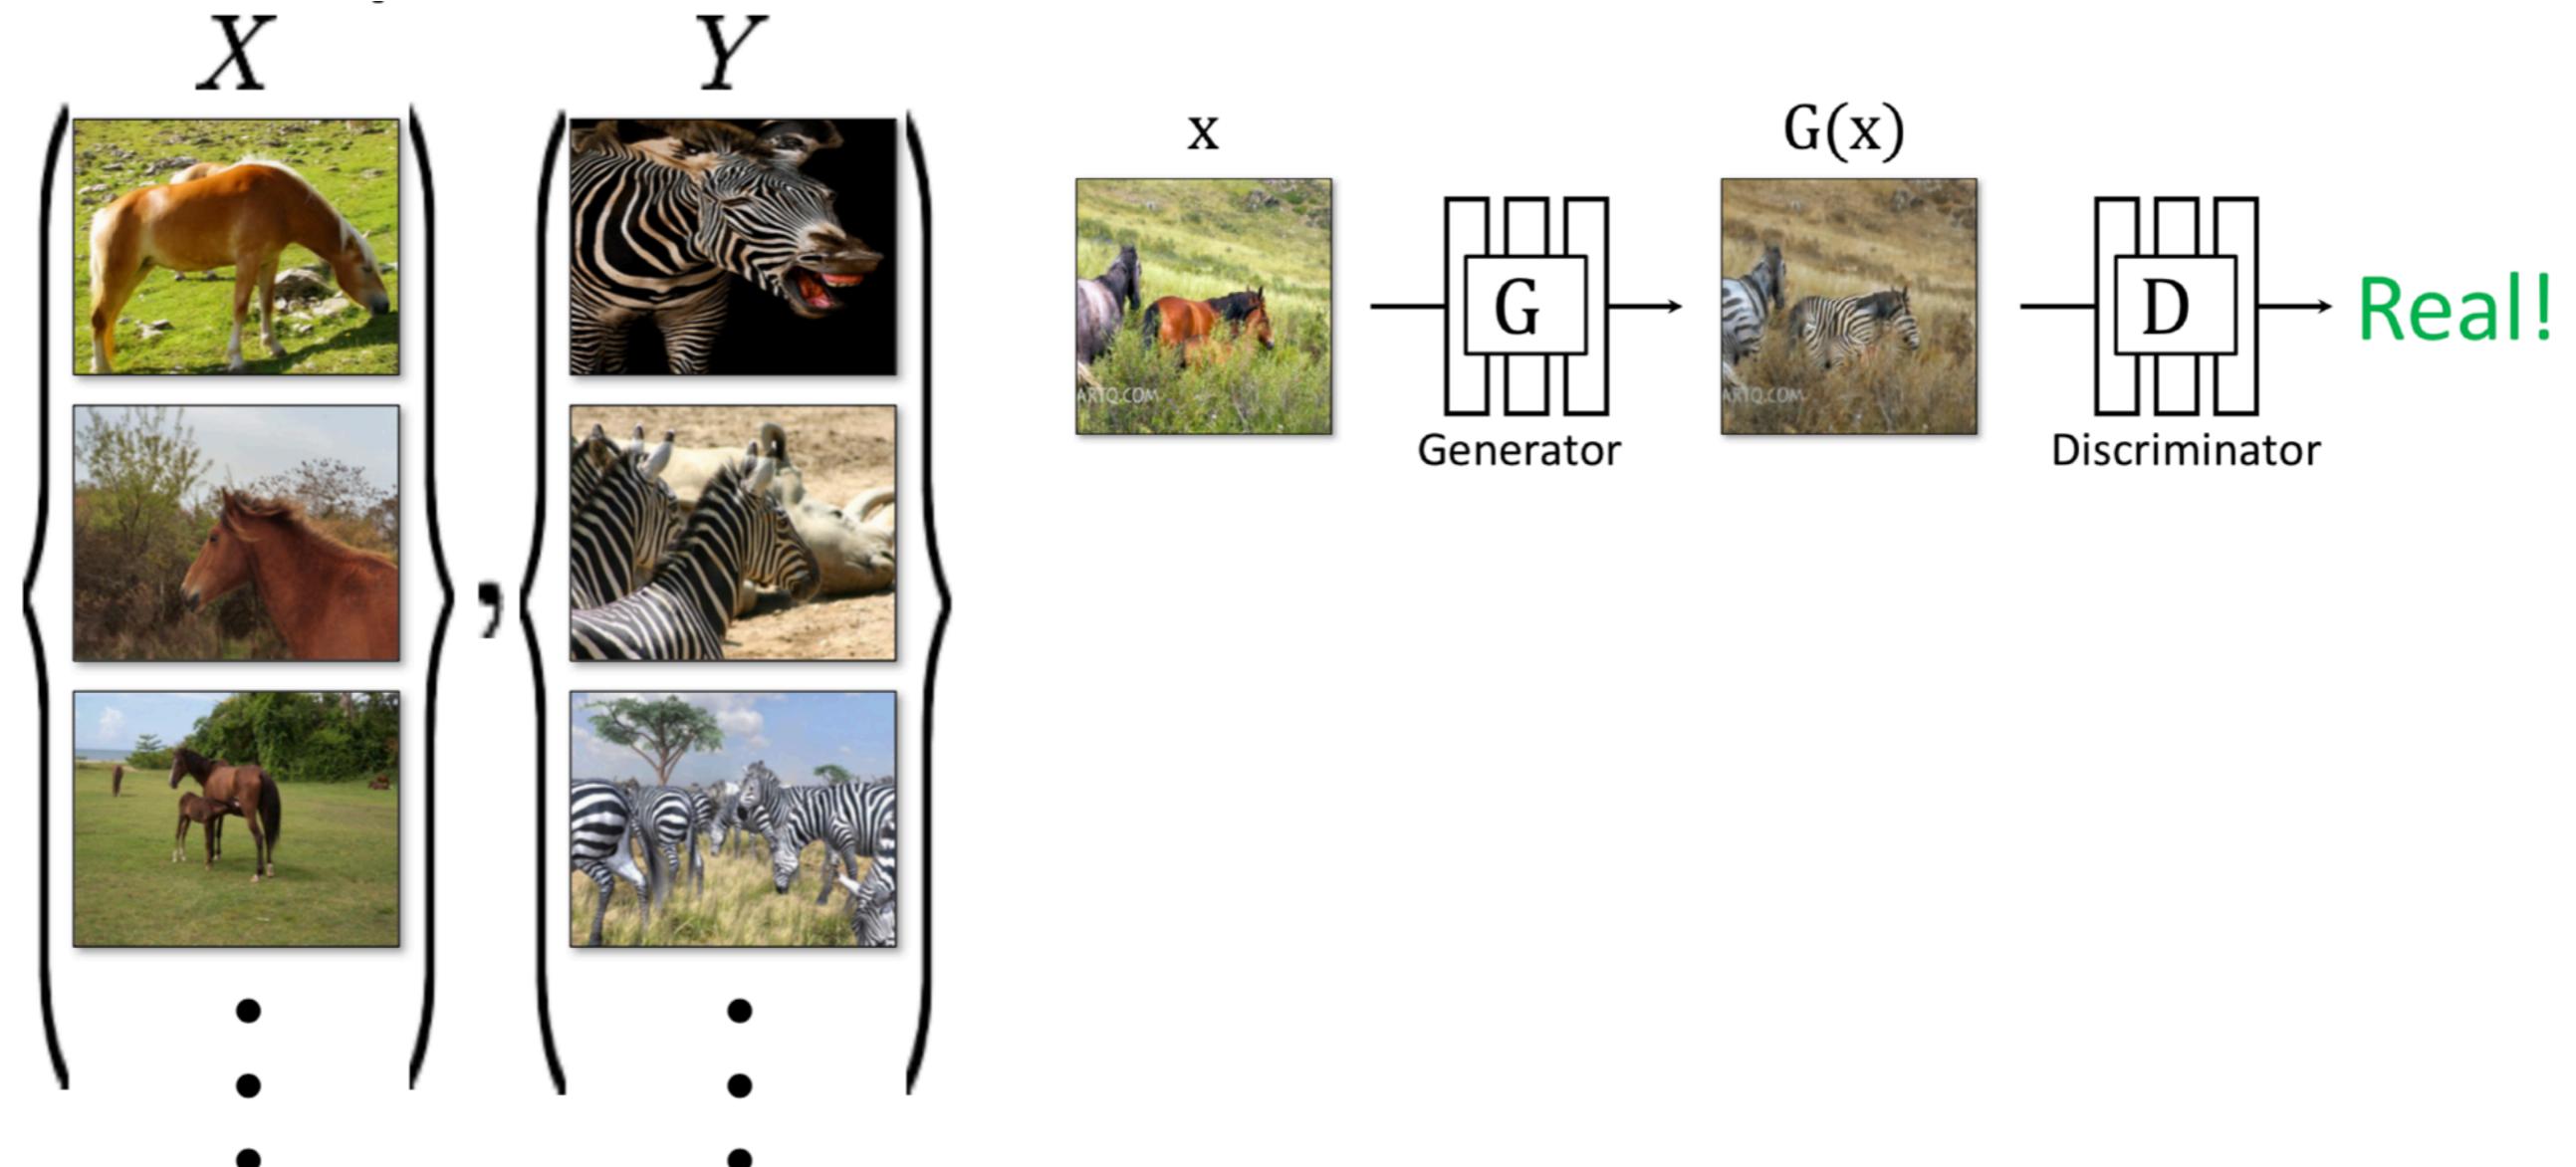
slide credit: Jun-Yan Zhu, Taesung Park

# Unsupervised Image-To-Image Translation



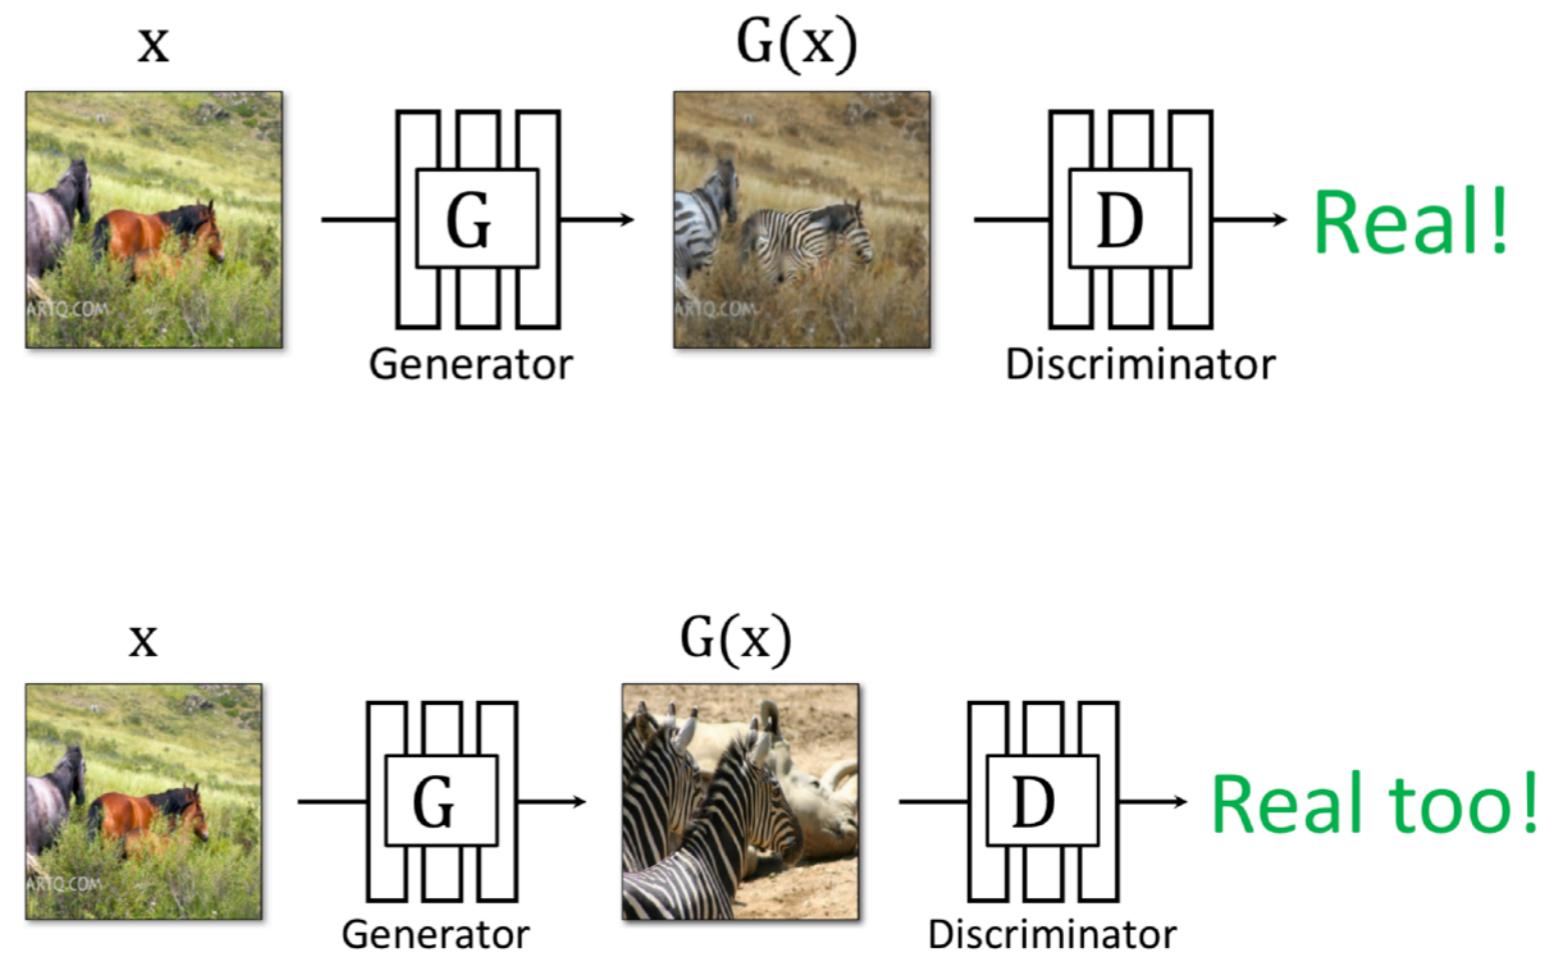
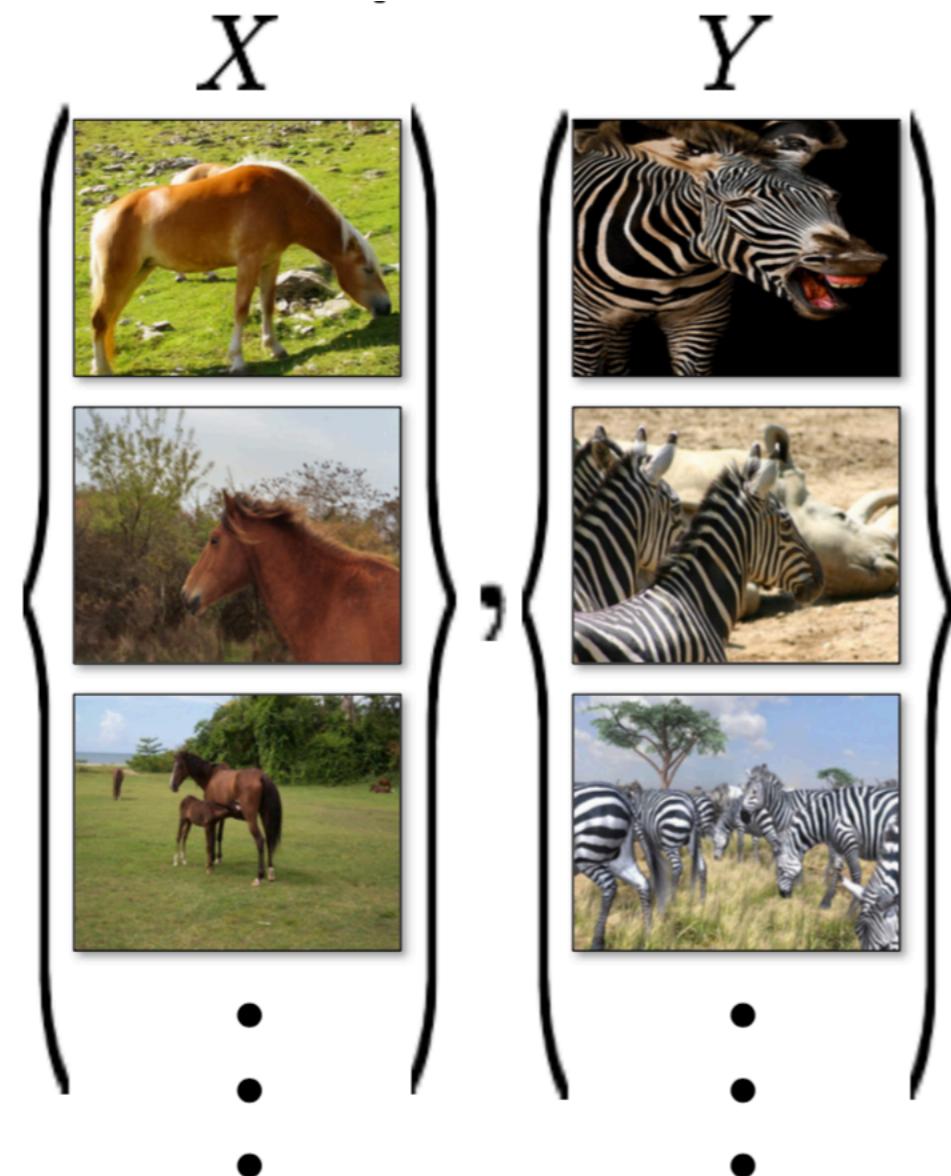
slide credit: Jun-Yan Zhu, Taesung Park

# Unsupervised Image-To-Image Translation



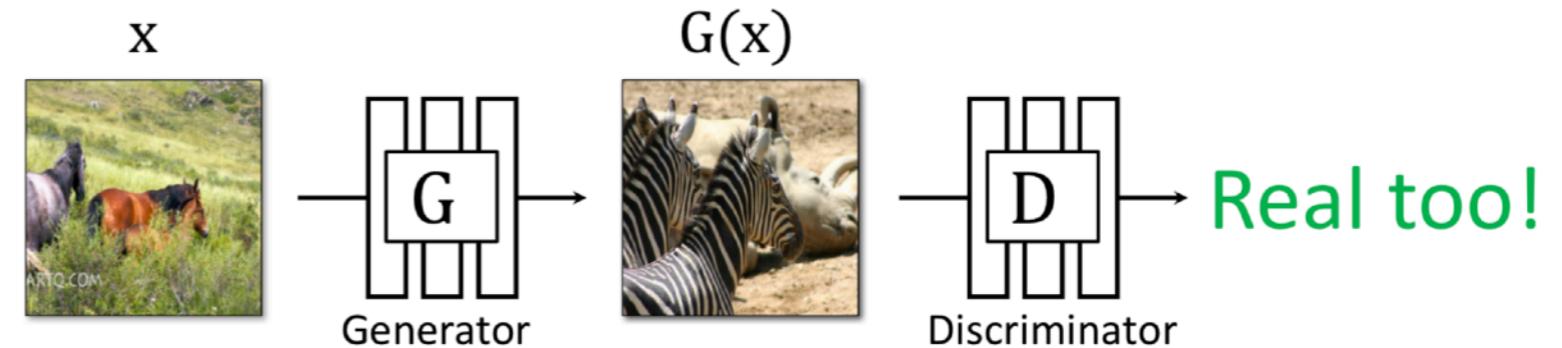
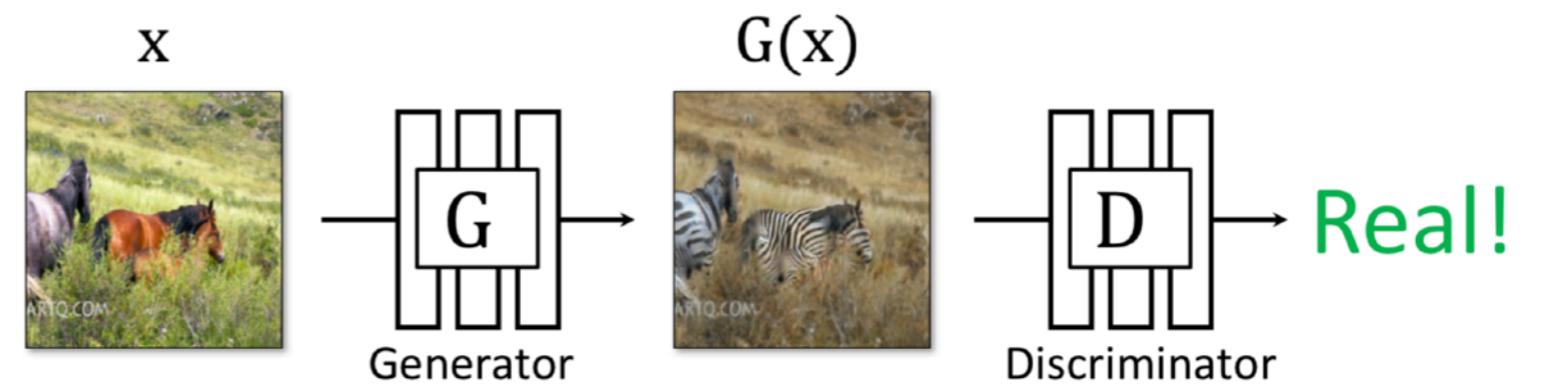
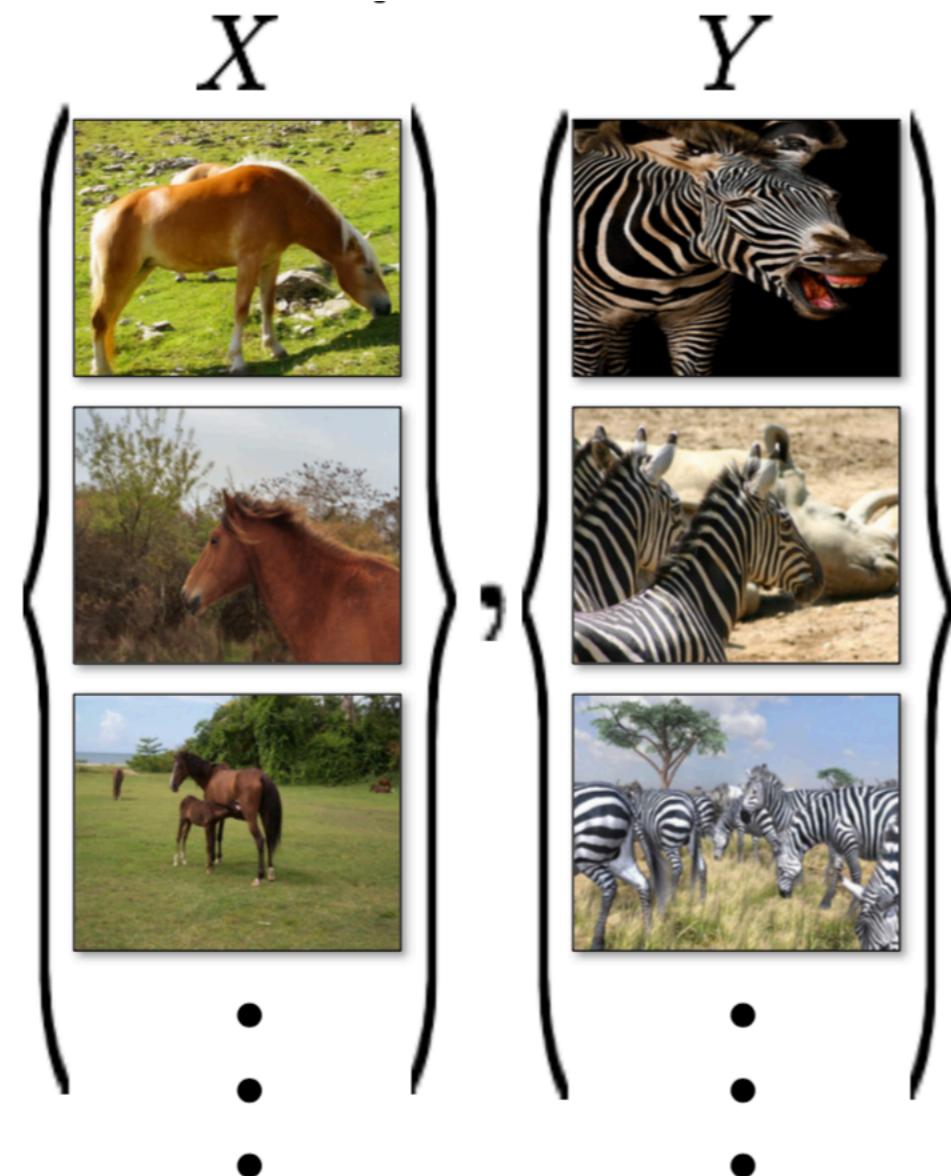
slide credit: Jun-Yan Zhu, Taesung Park

# Unsupervised Image-To-Image Translation



slide credit: Jun-Yan Zhu, Taesung Park

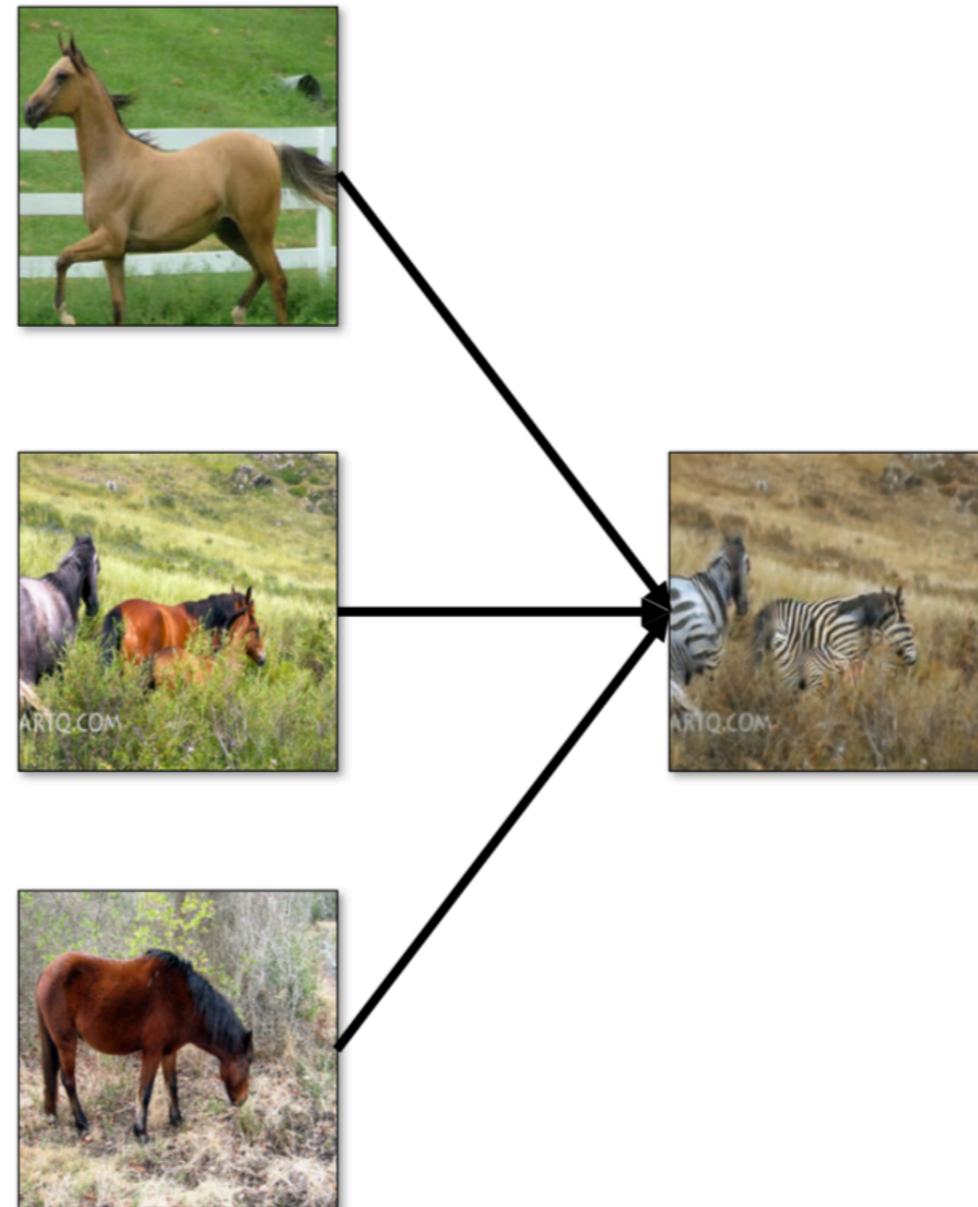
# Unsupervised Image-To-Image Translation



mode collapse

# CycleGan: Cycle -Consistent Adversarial Networks

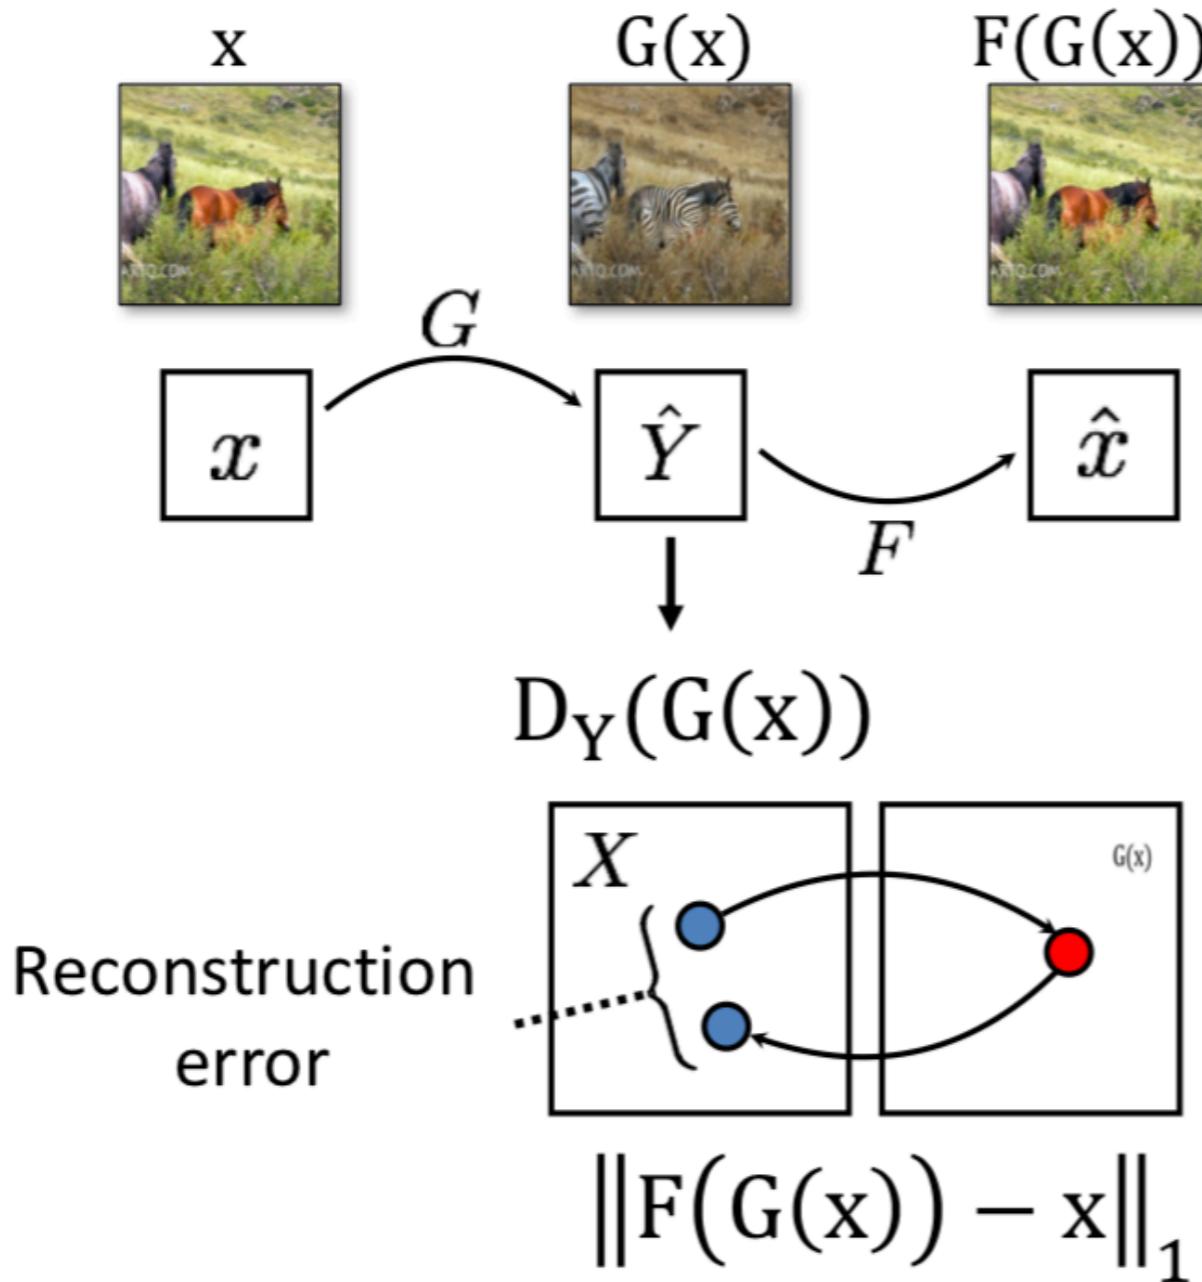
Does not prevent mode collapse



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-  
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

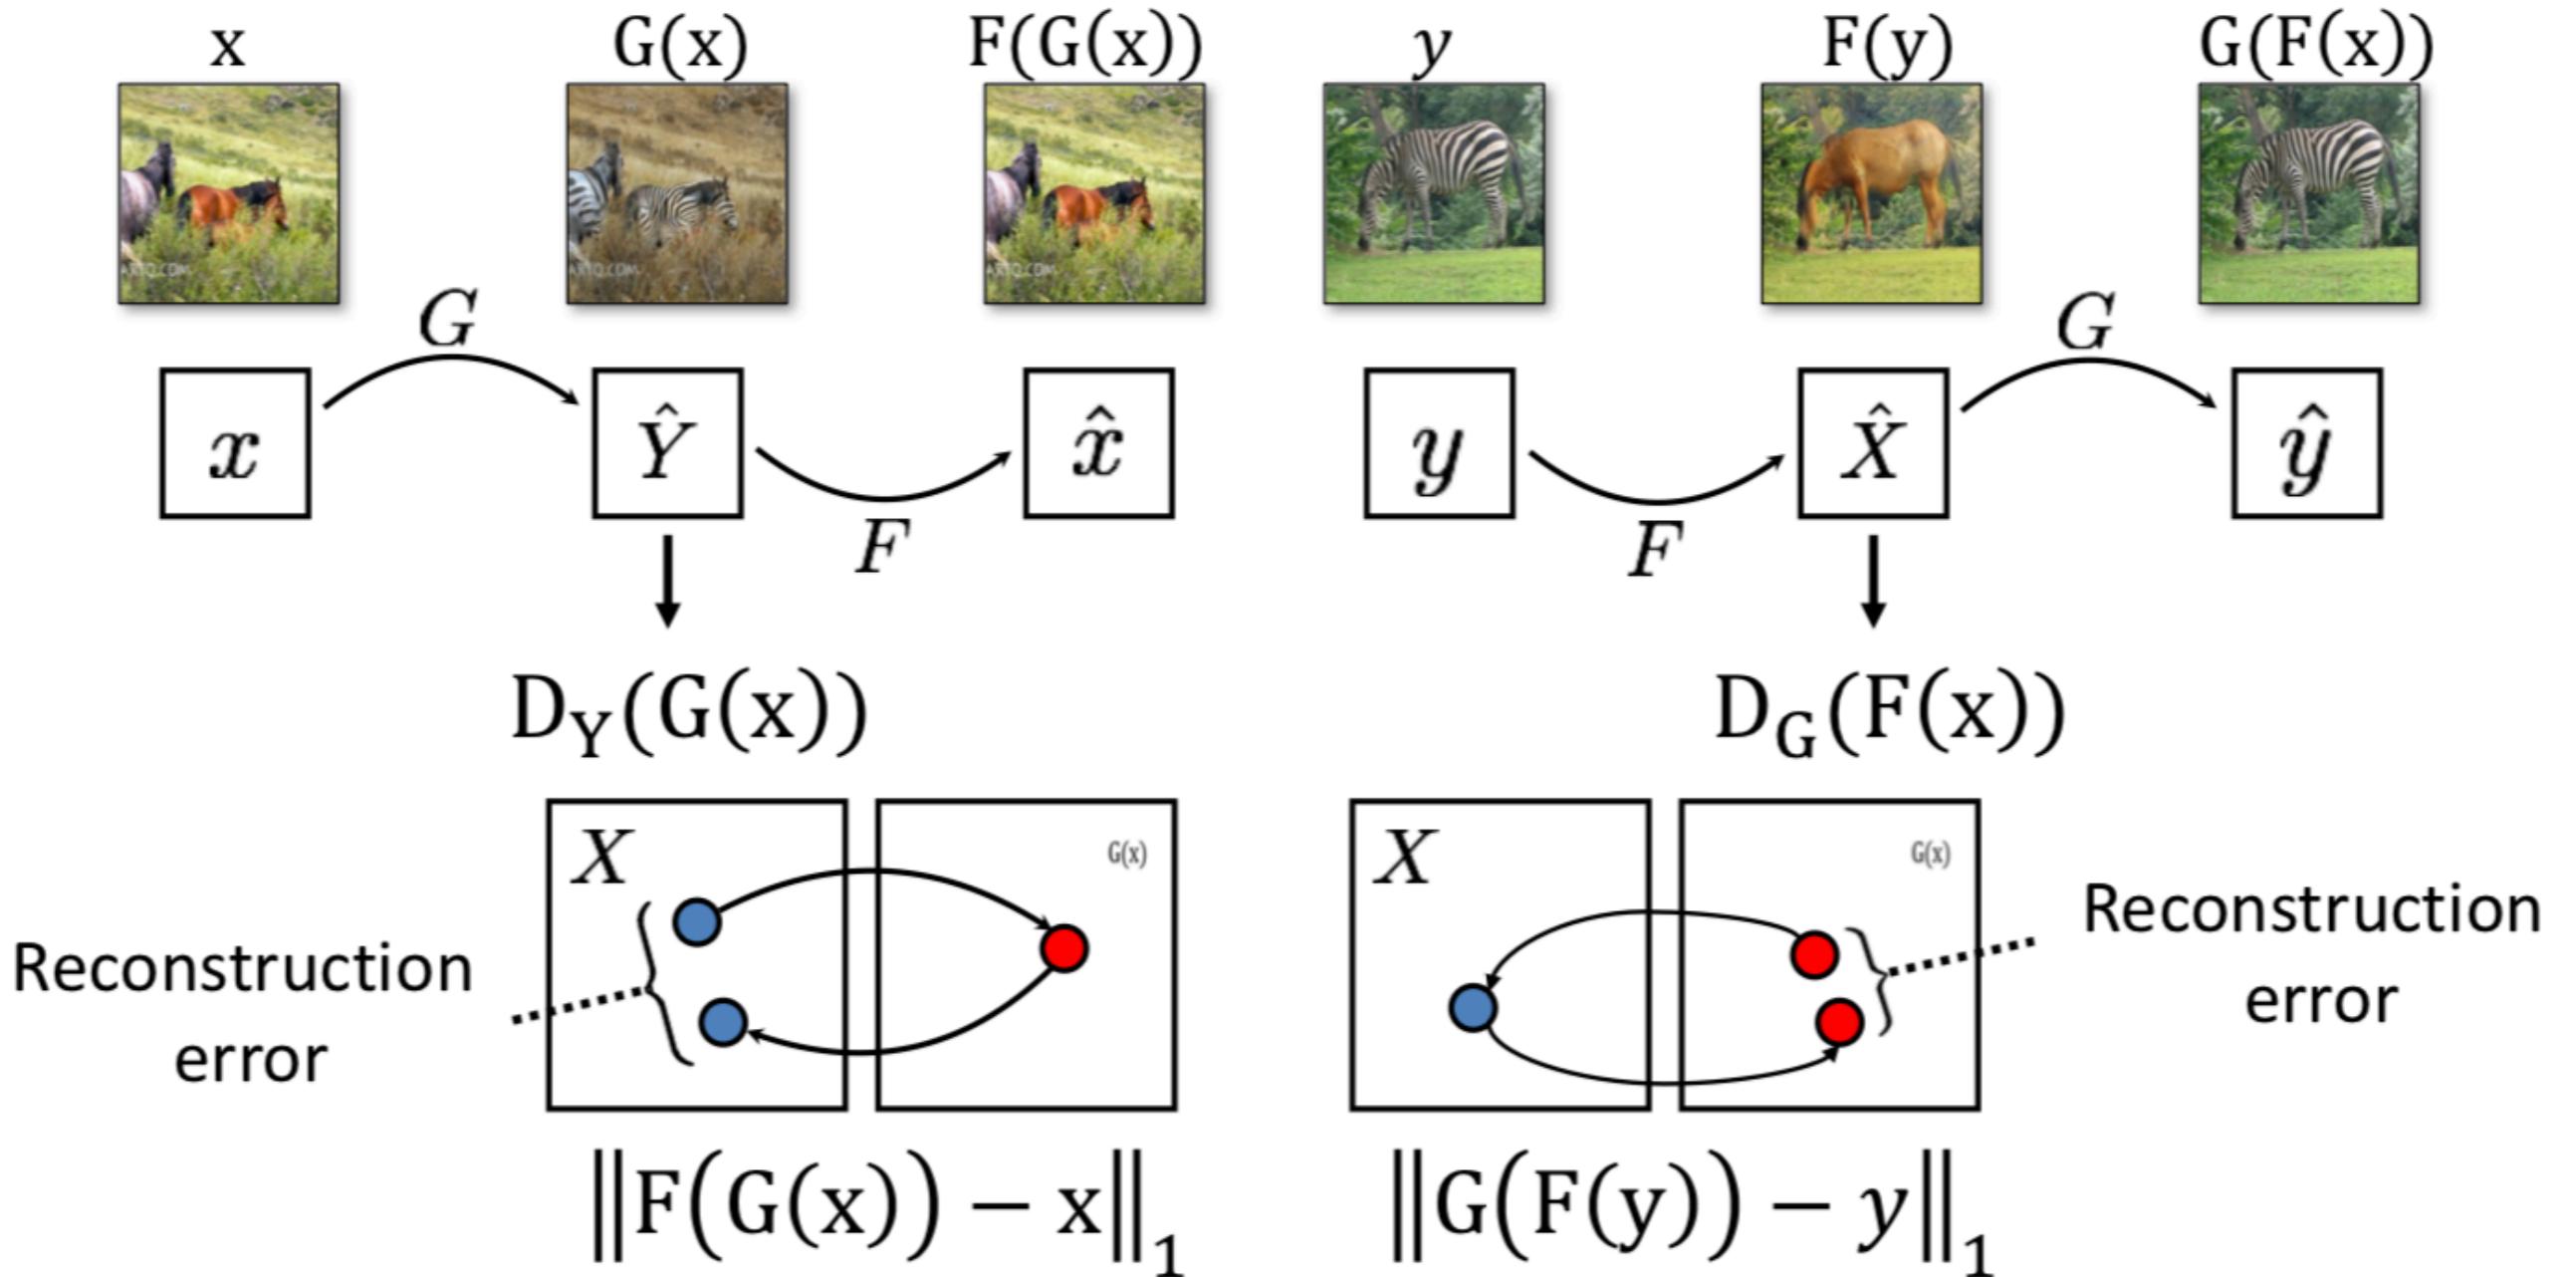
# CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-  
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

# CycleGan: Cycle -Consistent Adversarial Networks



[Zhu & Park et al., Unpaired Image-to-Image Translation using Cycle-  
Consistent Adversarial Networks, ICCV 2017]

slide credit: Jun-Yan Zhu, Taesung Park

# CycleGan Application: Rendering to Image

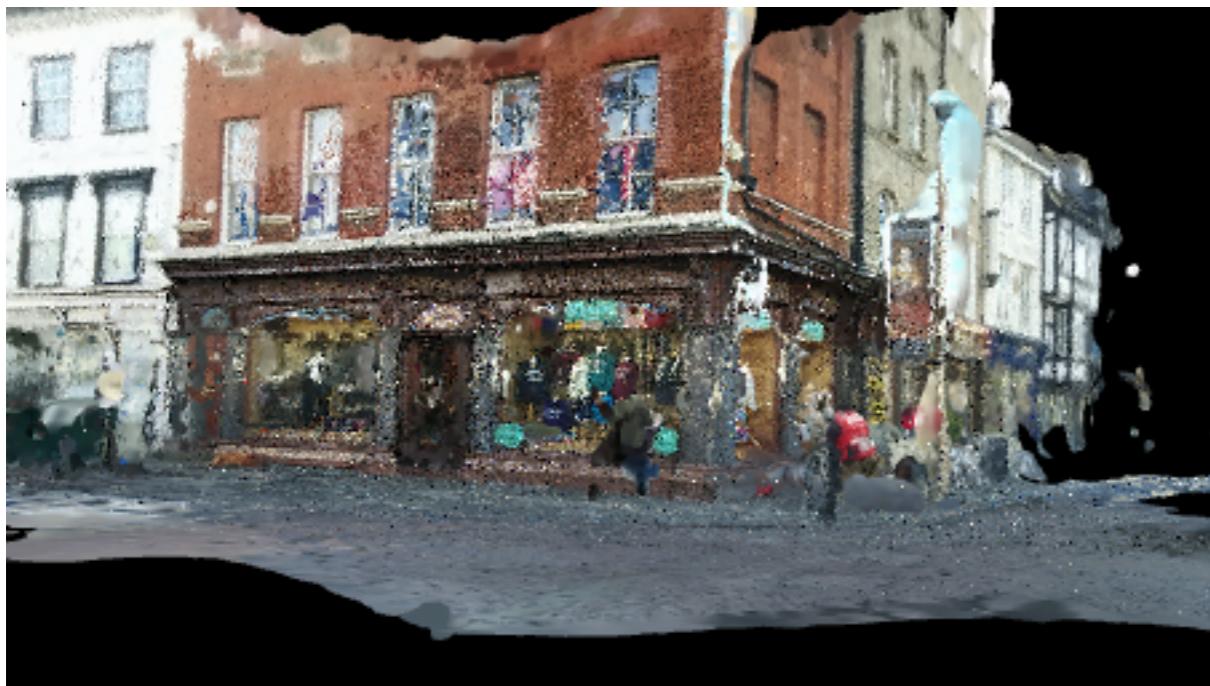


image credit: Markus Müller

# CycleGan Application: Image Retrieval



[Maddern, Pascoe, Linegar, Newman, 1 Year, 1000km: The Oxford RobotCar Dataset, IJRR 2017]

[Sattler et al., Benchmarking 6DOF Outdoor Visual Localization in Changing Conditions, CVPR 2018]

# CycleGan Application: Image Retrieval



[Maddern, Pascoe, Linegar, Newman, 1 Year, 1000km: The Oxford RobotCar Dataset, IJRR 2017]

[Sattler et al., Benchmarking 6DOF Outdoor Visual Localization in Changing Conditions, CVPR 2018]

# CycleGan Application: Image Retrieval

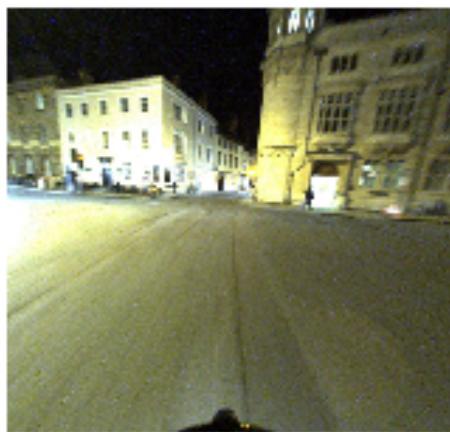


[Maddern, Pascoe, Linegar, Newman, 1 Year, 1000km: The Oxford RobotCar Dataset, IJRR 2017]

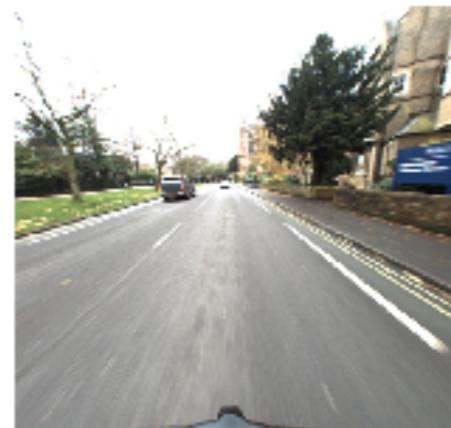
[Sattler et al., Benchmarking 6DOF Outdoor Visual Localization in Changing Conditions, CVPR 2018]

# CycleGan Application: Image Retrieval

**Night-Time  
Query**



**Retrieved Day  
Image**

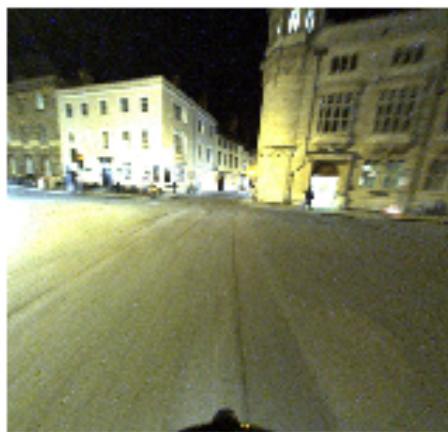


[Anoosheh et al., Night-to-Day Image Translation for Retrieval-based Localization, ICRA 2019]

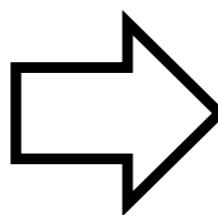
slide credit: Asha Anoosheh

# CycleGan Application: Image Retrieval

**Night-Time  
Query**



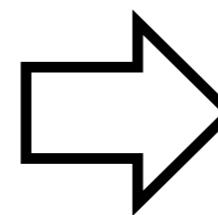
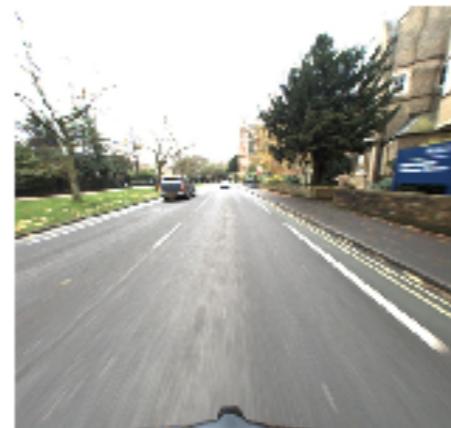
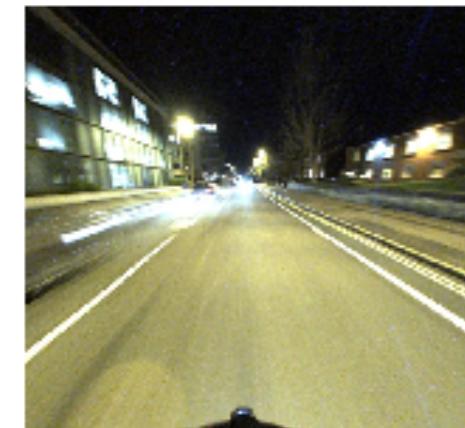
**Retrieved Day  
Image**



**Translated  
Query**



**Retrieved Day  
Image**



[Anoosheh et al., Night-to-Day Image Translation for Retrieval-based Localization, ICRA 2019]

slide credit: Asha Anoosheh

# Lessons Learned

- Main lessons from this lecture
  - Generative modeling via sampling from probability distributions
  - Generative networks learn mapping from simple probability distribution to target probability distribution
  - Adversarial networks: Learn the loss function
  - Unsupervised image-to-image translation via cycle consistency

# Lessons Learned

- Main lessons from this lecture
  - Generative modeling via sampling from probability distributions
  - Generative networks learn mapping from simple probability distribution to target probability distribution
  - Adversarial networks: Learn the loss function
  - Unsupervised image-to-image translation via cycle consistency
- Next lecture: Medical Image Analysis by Jennifer Alvén

# Reading Materials

- Start: [Goodfellow, Bengio, Courville, Deep Learning, MIT press \(Chapter 20.10.4\)](#)
- Continue: Blog posts on [Wasserstein GANs](#), [Tutorial on VAEs](#)
- Advanced material:
  - ICCV 2017 GAN Tutorial ([website](#) with slides and videos)
  - CVPR 2018 GAN Tutorial ([website](#) with slides and videos)
  - Tensorflow implementations of [various GANs](#)
- GANs are an active research topic, be prepared to read a lot